

Information technology — Programming languages — C

Reply To: JeanHeyd Meneide <wg14@soasis.org>

Freek Wiedijk <freek@cs.ru.nl>

Abstract

(This cover sheet to be replaced by ISO.)

This document specifies the form and establishes the interpretation of programs expressed in the programming language C. Its purpose is to promote portability, reliability, maintainability, and efficient execution of C language programs on a variety of computing systems.

Clauses are included that detail the C language itself and the contents of the C language execution library. Annexes summarize aspects of both of them, and enumerate factors that influence the portability of C programs.

Although this document is intended to guide knowledgeable C language programmers as well as implementers of C language translation systems, the document itself is not designed to serve as a tutorial.

Recipients of this draft are invited to submit, with their comments, notification of any relevant patent rights of which they are aware and to provide supporting documentation.

After the end of the C23 cycle, work began on so-called “C2y”, the next projected edition of this document on 2024 January 24.

2024 January

The following changes were made on or after the 2024 January Strasbourg, France meeting:

- N3192: Sequential hexdigits
- Editorial: Adjusted a footnote in Annex K from “reserved” to “potentially reserved”

2024 June

The following changes were made on or after the 2024 June Virtual meeting:

- N3064: Writing to multibyte character files
- N3232: Round-trip rounding
- N3233: Recommendation for printf rounding
- N3239: Some constants are literally literals, v2
- N3242: Problematic use of “correctly rounded”
- N3244: Slay Some Earthly Demons I (6.3.2.1 change, 6.7.2 change option 1, 6.7 change, 6.7.5 change, 6.7.6 change)
- N3247: **fopen** “p” and bring **fopen**’s mode closer to POSIX 202x
- N3254: Accessing byte arrays, v4
- N3259: Support ++ and -- on complex values
- N3260: **_Generic** selection expression with a type operand

- N3273: **alignof** of an incomplete array type
- N3274: Remove imaginary types, v3
- Editorial: “may” -> “can” for alignment with ISO/IEC Directives
- Editorial: Table headings changed to match ISO/IEC Directives

2024 September/October

The following changes were made on or after the 2024 September/October Minneapolis, USA meeting:

- N3272: **strftime** broken-down structure usage (Option 1: Undefined Behavior)
- N3286: Floating-point exception for Macro Replacements
- N3287: Nonsensical Parenthetical in Mathematics Specification
- N3291: Decimal Floating-Point Number Term Misuse
- N3298: Introduce Complex Literals (without changes proposed in 4.2)
- N3303: **HUGE_VAL** Corrections
- N3305: Leftover **WANT_** . . . Macros for `<math.h>` and Decimal Floating-Point
- N3312: Relax Atomic Alignment Requirements
- N3322: Allow Zero Length operations on Null Pointers (Including in the Library)
- N3323: How Do You Add One To Something? (By Using The Proper Type)
- N3324: "pole-error" Wording Fix
- N3326: Standardize **strnlen** and **wcsnlen**
- N3340: Slay Some Earthly Demons II
- N3341: Slay Some Earthly Demons III
- N3342: Slay Some Earthly Demons IV
- N3344: Slay Some Earthly Demons VI
- N3345: Slay Some Earthly Demons VII
- N3346: Slay Some Earthly Demons VIII
- N3349: **abs** Without Undefined Behavior
- N3353: Obsolete Octal and Provide New, Proper Escape Sequences
- N3355: Named/Labeled Loops
- N3356: **if** Declarations
- N3364: **SNAN** Initialization
- N3366: Restartable Functions for Efficient Character Conversions
- N3367: More Modern Bit Utilities
- N3369: The **_Lengthof** Operator
- N3370: Case Ranges in **switch** Statements

Contents

Foreword	xiv
Introduction	xv
1 Scope	1
2 Normative references	2
3 Terms, definitions, and symbols	3
4 Conformance	9
5 Environment	11
5.1 Introduction	11
5.2 Conceptual models	11
5.2.1 Translation environment	11
5.2.2 Execution environments	12
5.3 Environmental considerations	19
5.3.1 Character sets	19
5.3.2 Multibyte characters	20
5.3.3 Character display semantics	21
5.3.4 Signals and interrupts	21
5.3.5 Environmental limits	21
6 Language	36
6.1 Notation	36
6.2 Concepts	36
6.2.1 Scopes of identifiers, type names, and compound literals	36
6.2.2 Linkages of identifiers	37
6.2.3 Name spaces of identifiers	38
6.2.4 Storage durations of objects	38
6.2.5 Types	39
6.2.6 Representations of types	43
6.2.7 Compatible type and composite type	44
6.2.8 Alignment of objects	46
6.2.9 Encodings	46
6.3 Conversions	47
6.3.1 Introduction	47
6.3.2 Arithmetic operands	47

6.3.3	Other operands	51
6.4	Lexical elements	54
6.4.1	General	54
6.4.2	Keywords	55
6.4.3	Identifiers	56
6.4.4	Universal character names	58
6.4.5	Constants	60
6.4.6	String literals	71
6.4.7	Punctuators	72
6.4.8	Header names	73
6.4.9	Preprocessing numbers	74
6.4.10	Comments	74
6.5	Expressions	76
6.5.1	General	76
6.5.2	Primary expressions	77
6.5.3	Postfix operators	79
6.5.4	Unary operators	85
6.5.5	Cast operators	88
6.5.6	Multiplicative operators	89
6.5.7	Additive operators	90
6.5.8	Bitwise shift operators	91
6.5.9	Relational operators	92
6.5.10	Equality operators	92
6.5.11	Bitwise AND operator	93
6.5.12	Bitwise exclusive OR operator	94
6.5.13	Bitwise inclusive OR operator	94
6.5.14	Logical AND operator	94
6.5.15	Logical OR operator	94
6.5.16	Conditional operator	95
6.5.17	Assignment operators	96
6.5.18	Comma operator	99
6.6	Constant expressions	100
6.6.1	General	100
6.6.2	Constant range expressions	101
6.7	Declarations	104
6.7.1	General	104
6.7.2	Storage-class specifiers	106
6.7.3	Type specifiers	110
6.7.4	Type qualifiers	127
6.7.5	Function specifiers	131

6.7.6	Alignment specifier	132
6.7.7	Declarators	133
6.7.8	Type names	139
6.7.9	Type definitions	140
6.7.10	Type inference	142
6.7.11	Initialization	143
6.7.12	Static assertions	149
6.7.13	Attributes	149
6.8	Statements and blocks	160
6.8.1	General	160
6.8.2	Labeled statements	161
6.8.3	Compound statement	161
6.8.4	Expression and null statements	162
6.8.5	Selection statements	162
6.8.6	Iteration statements	165
6.8.7	Jump statements	166
6.9	External definitions	171
6.9.1	General	171
6.9.2	Function definitions	171
6.9.3	External object definitions	173
6.10	Preprocessing directives	175
6.10.1	General	175
6.10.2	Conditional inclusion	178
6.10.3	Source file inclusion	182
6.10.4	Binary resource inclusion	183
6.10.5	Macro replacement	190
6.10.6	Line control	197
6.10.7	Diagnostic directives	198
6.10.8	Pragma directive	198
6.10.9	Null directive	199
6.10.10	Predefined macro names	199
6.10.11	Pragma operator	201
6.11	Future language directions	203
6.11.1	Floating types	203
6.11.2	Linkages of identifiers	203
6.11.3	External names	203
6.11.4	Character escape sequences	203
6.11.5	Octal integer literals	203
6.11.6	Storage-class specifiers	203
6.11.7	Pragma directives	203

6.11.8	Predefined macro names	203
7	Library	204
7.1	Introduction	204
7.1.1	Definitions of terms	204
7.1.2	Standard headers	204
7.1.3	Reserved identifiers	205
7.1.4	Use of library functions	206
7.2	Diagnostics <assert.h>	208
7.2.1	General	208
7.2.2	Program diagnostics	208
7.3	Complex arithmetic <complex.h>	209
7.3.1	Introduction	209
7.3.2	Conventions	209
7.3.3	Branch cuts	209
7.3.4	The CX_LIMITED_RANGE pragma	209
7.3.5	Trigonometric functions	210
7.3.6	Hyperbolic functions	211
7.3.7	Exponential and logarithmic functions	213
7.3.8	Power and absolute-value functions	214
7.3.9	Manipulation functions	214
7.4	Character handling <ctype.h>	217
7.4.1	General	217
7.4.2	Character classification functions	217
7.4.3	Character case mapping functions	219
7.5	Errors <errno.h>	221
7.6	Floating-point environment <fenv.h>	222
7.6.1	General	222
7.6.2	The FENV_ACCESS pragma	224
7.6.3	The FENV_ROUND pragma	225
7.6.4	The FENV_DEC_ROUND pragma	227
7.6.5	Floating-point exceptions	228
7.6.6	Rounding and other control modes	230
7.6.7	Environment	233
7.7	Characteristics of floating types <float.h>	235
7.8	Format conversion of integer types <inttypes.h>	236
7.8.1	General	236
7.8.2	Macros for format specifiers	236
7.8.3	Functions for greatest-width integer types	237
7.9	Alternative spellings <iso646.h>	239

7.10	Characteristics of integer types <code><limits.h></code>	240
7.11	Localization <code><locale.h></code>	241
7.11.1	General	241
7.11.2	The setlocale function	241
7.11.3	Numeric formatting convention inquiry	242
7.12	Mathematics <code><math.h></code>	247
7.12.1	General	247
7.12.2	Treatment of error conditions	250
7.12.3	The FP_CONTRACT pragma	251
7.12.4	Classification macros	251
7.12.5	Trigonometric functions	254
7.12.6	Hyperbolic functions	259
7.12.7	Exponential and logarithmic functions	261
7.12.8	Power and absolute-value functions	269
7.12.9	Error and gamma functions	272
7.12.10	Nearest integer functions	274
7.12.11	Remainder functions	278
7.12.12	Manipulation functions	279
7.12.13	Maximum, minimum, and positive difference functions	282
7.12.14	Fused multiply-add	287
7.12.15	Functions that round result to narrower type	287
7.12.16	Quantum and quantum exponent functions	289
7.12.17	Decimal re-encoding functions	291
7.12.18	Comparison macros	293
7.13	Non-local jumps <code><setjmp.h></code>	296
7.13.1	General	296
7.13.2	Save calling environment	296
7.13.3	Restore calling environment	296
7.14	Signal handling <code><signal.h></code>	298
7.14.1	General	298
7.14.2	Specify signal handling	298
7.14.3	Send signal	300
7.15	Alignment <code><stdalign.h></code>	301
7.16	Varying arguments <code><stdarg.h></code>	302
7.16.1	General	302
7.16.2	Varying argument list access macros	302
7.17	Atomics <code><stdatomic.h></code>	307
7.17.1	Introduction	307
7.17.2	Initialization	308
7.17.3	Order and consistency	309

7.17.4	Fences	311
7.17.5	Lock-free property	312
7.17.6	Atomic integer types	313
7.17.7	Operations on atomic types	314
7.17.8	Atomic flag type and operations	316
7.18	Bit and byte utilities <stdbit.h>	318
7.18.1	General	318
7.18.2	Endian	318
7.18.3	Count Leading Zeros	319
7.18.4	Count Leading Ones	319
7.18.5	Count Trailing Zeros	319
7.18.6	Count Trailing Ones	320
7.18.7	First Leading Zero	320
7.18.8	First Leading One	321
7.18.9	First Trailing Zero	321
7.18.10	First Trailing One	322
7.18.11	Count Zeros	323
7.18.12	Count Ones	323
7.18.13	Single-bit Check	323
7.18.14	Bit Width	324
7.18.15	Bit Floor	324
7.18.16	Bit Ceiling	325
7.18.17	Rotate Left	326
7.18.18	Rotate Right	326
7.18.19	8-bit Memory Reversal	327
7.18.20	Exact-width 8-bit Memory Reversal	327
7.18.21	Endian-Aware 8-bit Load	328
7.18.22	Endian-Aware 8-bit Store	328
7.19	Boolean type and values <stdbool.h>	330
7.20	Checked Integer Arithmetic <stdckdint.h>	331
7.20.1	General	331
7.20.2	Checked Integer Operation Type-generic Macros	331
7.21	Common definitions <stddef.h>	332
7.21.1	General	332
7.21.2	The unreachable macro	333
7.21.3	The nullptr_t type	334
7.22	Integer types <stdint.h>	335
7.22.1	General	335
7.22.2	Integer types	335
7.22.3	Widths of specified-width integer types	337

7.22.4	Width of other integer types	337
7.22.5	Macros for constants of integer type	338
7.22.6	Maximal and minimal values of integer types	338
7.23	Input/output <stdio.h>	339
7.23.1	Introduction	339
7.23.2	Streams	341
7.23.3	Files	342
7.23.4	Operations on files	343
7.23.5	File access functions	345
7.23.6	Formatted input/output functions	348
7.23.7	Character input/output functions	366
7.23.8	Direct input/output functions	369
7.23.9	File positioning functions	370
7.23.10	Error-handling functions	371
7.24	General utilities <stdlib.h>	373
7.24.1	General	373
7.24.2	Numeric conversion functions	373
7.24.3	Pseudo-random sequence generation functions	380
7.24.4	Memory management functions	381
7.24.5	Communication with the environment	384
7.24.6	Searching and sorting utilities	387
7.24.7	Integer arithmetic functions	388
7.24.8	Multibyte/wide character conversion functions	389
7.24.9	Multibyte/wide string conversion functions	390
7.24.10	Alignment of memory	391
7.25	Text transcoding utilities <stdmchar.h>	393
7.25.1	General	393
7.25.2	Single Unit Sized Conversion Functions	396
7.25.3	Multi Unit Sized Conversion Functions	400
7.26	_Noreturn <stdnoreturn.h>	406
7.27	String handling <string.h>	407
7.27.1	String function conventions	407
7.27.2	Copying functions	407
7.27.3	Concatenation functions	409
7.27.4	Comparison functions	409
7.27.5	Search functions	411
7.27.6	Miscellaneous functions	414
7.28	Type-generic math <tgmath.h>	416
7.29	Threads <threads.h>	421
7.29.1	Introduction	421

7.29.2	Initialization functions	422
7.29.3	Condition variable functions	422
7.29.4	Mutex functions	424
7.29.5	Thread functions	426
7.29.6	Thread-specific storage functions	428
7.30	Date and time <time.h>	431
7.30.1	Components of time	431
7.30.2	Time manipulation functions	432
7.30.3	Time conversion functions	435
7.31	Unicode utilities <uchar.h>	441
7.31.1	General	441
7.31.2	Restartable multibyte/wide character conversion functions	441
7.32	Extended multibyte and wide character utilities <wchar.h>	446
7.32.1	Introduction	446
7.32.2	Formatted wide character input/output functions	447
7.32.3	Wide character input/output functions	461
7.32.4	General wide string utilities	464
7.32.4.1	General	464
7.32.4.2	Wide string numeric conversion functions	464
7.32.4.3	Wide string copying functions	469
7.32.4.4	Wide string concatenation functions	470
7.32.4.5	Wide string comparison functions	471
7.32.4.6	Wide string search functions	472
7.32.4.7	Miscellaneous functions	475
7.32.5	Wide character time conversion functions	476
7.32.6	Extended multibyte/wide character conversion utilities	476
7.32.6.1	General	476
7.32.6.2	Single-byte/wide character conversion functions	477
7.32.6.3	Conversion state functions	477
7.32.6.4	Restartable multibyte/wide character conversion functions	478
7.32.6.5	Restartable multibyte/wide string conversion functions	479
7.33	Wide character classification and mapping utilities <wctype.h>	482
7.33.1	Introduction	482
7.33.2	Wide character classification utilities	482
7.33.2.1	General	482
7.33.2.2	Wide character classification functions	482
7.33.2.3	Extensible wide character classification functions	485
7.33.3	Wide character case mapping utilities	486
7.33.3.1	Wide character case mapping functions	486
7.33.3.2	Extensible wide character case mapping functions	486

7.34 Future library directions	488
7.34.1 General	488
7.34.2 Complex arithmetic <complex.h>	488
7.34.3 Character handling <ctype.h>	488
7.34.4 Errors <errno.h>	488
7.34.5 Floating-point environment <fenv.h>	488
7.34.6 Characteristics of floating types <float.h>	488
7.34.7 Format conversion of integer types <inttypes.h>	488
7.34.8 Localization <locale.h>	488
7.34.9 Mathematics <math.h>	488
7.34.10 Signal handling <signal.h>	489
7.34.11 Atomics <stdatomic.h>	489
7.34.12 Boolean type and values <stdbool.h>	489
7.34.13 Bit and byte utilities <stdbit.h>	489
7.34.14 Checked Arithmetic Functions <stdckdint.h>	489
7.34.15 Integer types <stdint.h>	489
7.34.16 Input/output <stdio.h>	489
7.34.17 General utilities <stdlib.h>	489
7.34.18 String handling <string.h>	490
7.34.19 Date and time <time.h>	490
7.34.20 Threads <threads.h>	490
7.34.21 Extended multibyte and wide character utilities <wchar.h>	490
7.34.22 Wide character classification and mapping utilities <wctype.h>	490
Annex A (informative) Language syntax summary	491
Annex B (informative) Library summary	507
Annex C (informative) Sequence points	552
Annex D (informative) Universal character names for identifiers	553
Annex E (informative) Implementation limits	555
Annex F (normative) ISO/IEC 60559 floating-point arithmetic	558
Annex G (normative) ISO/IEC 60559-compatible complex arithmetic	589
Annex H (normative) ISO/IEC 60559 interchange and extended types	598
Annex I (informative) Common warnings	631
Annex J (informative) Portability issues	632

Annex K (normative) Bounds-checking interfaces	671
Annex L (normative) Analyzability	719
Annex M (informative) Change History	720
Bibliography	727
Index	728

Foreword

- 1 ISO (the International Organization for Standardization) and IEC (the International Electrotechnical Commission) form the specialized system for worldwide standardization. National bodies that are members of ISO or IEC participate in the development of International Standards through technical committees established by the respective organization to deal with particular fields of technical activity. ISO and IEC technical committees collaborate in fields of mutual interest. Other international organizations, governmental and non-governmental, in liaison with ISO and IEC, also take part in the work.
- 2 The procedures used to develop this document and those intended for its further maintenance are described in the ISO/IEC Directives, Part 1. In particular, the different approval criteria needed for the different types of document should be noted. This document was drafted in accordance with the editorial rules of the ISO/IEC Directives, Part 2 (see www.iso.org/directives or www.iec.ch/members_experts/refdocs).
- 3 ISO and IEC draw attention to the possibility that the implementation of this document may involve the use of (a) patent(s). ISO and IEC take no position concerning the evidence, validity or applicability of any claimed patent rights in respect thereof. As of the date of publication of this document, ISO and IEC had not received notice of (a) patent(s) which may be required to implement this document. However, implementers are cautioned that this may not represent the latest information, which may be obtained from the patent database available at www.iso.org/patents and patents.iec.ch. ISO and IEC shall not be held responsible for identifying any or all such patent rights.
- 4 Any trade name used in this document is information given for the convenience of users and does not constitute an endorsement.
- 5 For an explanation of the voluntary nature of standards, the meaning of ISO specific terms and expressions related to conformity assessment, as well as information about ISO's adherence to the World Trade Organization (WTO) principles in the Technical Barriers to Trade (TBT) see www.iso.org/iso/foreword.html. In the IEC, see www.iec.ch/understanding-standards.
- 6 This document was prepared by Joint Technical Committee ISO/IEC JTC1, Information technology, Subcommittee SC22, Programming languages, their environments and system software interfaces.
- 7 This fifth edition cancels and replaces the fourth edition (ISO/IEC 9899:2018), which has been technically revised. The main changes are contained in Annex M.
- 8 Any feedback or questions on this document should be directed to the user's national standards body. A complete listing of these bodies can be found at www.iso.org/members.html and www.iec.ch/national-committees.

Introduction

- 1 With the introduction of new devices and extended character sets, new features could be added to future editions of this document. Subclauses in the language and library clauses warn implementers and programmers of usages which, though valid in themselves, could conflict with future additions.
- 2 Certain features are *obsolescent*, which means that they could be considered for withdrawal in future revisions of this document. They are retained because of their widespread use, but their use in new implementations (for implementation features) or new programs (for language [6.11] or library features [7.34]) is discouraged.
- 3 This document is divided into four major subdivisions:
 - preliminary elements (Clauses 1–4);
 - the characteristics of environments that translate and execute C programs (Clause 5);
 - the language syntax, constraints, and semantics (Clause 6);
 - the library facilities (Clause 7).
- 4 In any given subsequent clause or subclause, there are section delineations in **bold** to describe the semantics, restrictions, and behaviors of programs for this language and potentially the use of its library clauses in this document:
 - **Syntax**
which pertains to the spelling and organization of the language and library;
 - **Constraints**
which detail and enumerate various requirements for the correct interpretation of the language and library, typically during translation;
 - **Semantics**
which explain the behavior of language features and similar constructs;
 - **Description**
which explain the behavior of library usage and similar constructs;
 - **Returns**
which describes the effects of constructs provided back to a user of the library;
 - **Runtime-constraints**
which detail and enumerate various requirements that are expected to be checked and which shall not be violated, typically during execution;
 - **Environmental limits**
which list limitations an implementation may impose on a library or language construct which might otherwise be unlimited;
 - **Recommended practice**
which provides guidance and important considerations for implementers of this document.
- 5 Examples are provided to illustrate possible forms of the constructions described. Footnotes are provided to emphasize consequences of the rules described in that subclause or elsewhere in this document. References are used to refer to other related subclauses. Recommendations are provided to give advice or guidance to implementers. Annexes define optional features, provide additional

information and summarize the information contained in this document. A bibliography lists documents that were referred to during the preparation of this document.

- 6 The language clause (Clause 6) is derived from “The C Reference Manual”[15].
- 7 The library clause (Clause 7) is based on the *1984 /usr/group Standard*[16].

Information technology — Programming languages — C

1. Scope

- 1 This document specifies the form and establishes the interpretation of programs written in the C programming language. It is designed to promote the portability of C programs among a variety of data-processing systems. It is intended for use by implementers and programmers. It specifies:
 - the representation of C programs;
 - the syntax and constraints of the C language;
 - the semantic rules for interpreting C programs;
 - the representation of input data to be processed by C programs;
 - the representation of output data produced by C programs;
 - the restrictions and limits imposed by a conforming implementation of C.
- 2 This document does not specify:
 - the mechanism by which C programs are transformed for use by a data-processing system;
 - the mechanism by which C programs are invoked for use by a data-processing system;
 - the mechanism by which input data are transformed for use by a C program;
 - the mechanism by which output data are transformed after being produced by a C program;
 - the size or complexity of a program and its data that will exceed the capacity of any specific data-processing system or the capacity of a particular processor;
 - all minimal requirements of a data-processing system that is capable of supporting a conforming implementation.
- 3 Annex J gives an overview of portability issues that a C program can encounter.

2. Normative references

- 1 The following documents are referred to in the text in such a way that some or all their content constitutes requirements of this document. For dated references, only the edition cited applies. For undated references, the latest edition of the referenced document (including any amendments) applies.
- 2 ISO/IEC 2382:2015, *Information technology — Vocabulary*.
- 3 ISO 4217, *Codes for the representation of currencies*.
- 4 ISO 8601 series, *Data elements and interchange formats — Information interchange — Representation of dates and times*.
- 5 ISO/IEC 10646, *Information technology — Universal Coded Character Set (UCS)*.
- 6 ISO/IEC 60559:2020, *Information technology — Microprocessor Systems — Floating-Point arithmetic*.
- 7 ISO 80000-2, *Quantities and units — Part 2: Mathematics*.
- 8 The Unicode Consortium. *Unicode Standard Annex, UAX #44, Unicode Character Database [online]*. Edited by Ken Whistler. Available at <https://www.unicode.org/reports/tr44>.
- 9 The Unicode Consortium. *The Unicode Standard, Derived Core Properties*. Available at <https://www.unicode.org/Public/UCD/latest/ucd/DerivedCoreProperties.txt>.

3. Terms, definitions, and symbols

- 1 For the purposes of this document, the terms and definitions given in ISO/IEC 2382, ISO 80000–2, and the following apply.
- 2 ISO and IEC maintain terminology databases for use in standardization at the following addresses:
 - ISO Online browsing platform: available at <https://www.iso.org/obp>
 - IEC Electropedia: available at <https://www.electropedia.org/>
- 3 Additional terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this document are not to be presumed to refer implicitly to similar terms defined elsewhere.

3.1

- 1 **access (verb)**
〈execution-time action〉 read or modify the value of an object
- 2 Note 1 to entry: Where only one of these two actions is meant, “read” or “modify” is used.
- 3 Note 2 to entry: “Modify” includes the case where the new value being stored is the same as the previous value.
- 4 Note 3 to entry: Expressions that are not evaluated do not access objects.

3.2

- 1 **alignment**
requirement that objects of a particular type be located on storage boundaries with addresses that are particular multiples of a byte (3.7) address

3.3

- 1 **argument**
actual argument
DEPRECATED: actual parameter
expression in the comma-separated list bounded by the parentheses in a function call expression, or a sequence of preprocessing tokens in the comma-separated list bounded by the parentheses in a function-like macro invocation

3.4

- 1 **arithmetically negate**
produce the negative of a given number
- 2 Note 1 to entry: For a floating-point number (5.3.5.3.3), this changes the sign; for an integer, this is equivalent to subtracting from zero.

3.5

- 1 **behavior**
external appearance or action

3.5.1

- 1 **implementation-defined behavior**
unspecified behavior where each implementation documents how the choice is made

- 2 Note 1 to entry: J.3 gives an overview over properties of C programs that lead to implementation-defined behavior.
- 3 EXAMPLE An example of implementation-defined behavior is the propagation of the high-order bit when a signed integer is shifted right.

3.5.2

1 **locale-specific behavior**

behavior that depends on local conventions of nationality, culture, and language that each implementation documents

- 2 Note 1 to entry: J.4 gives an overview over properties of C programs that lead to locale-specific behavior.
- 3 EXAMPLE An example of locale-specific behavior is whether the **islower** function returns true for characters other than the 26 lowercase Latin letters.

3.5.3

1 **undefined behavior**

behavior, upon use of a nonportable or erroneous program construct or of erroneous data, for which this document imposes no requirements

- 2 Note 1 to entry: Possible undefined behavior ranges from ignoring the situation completely with unpredictable results, to behaving during translation or program execution in a documented manner characteristic of the environment (with or without the issuance of a diagnostic message), to terminating a translation or execution (with the issuance of a diagnostic message).
- 3 Note 2 to entry: J.2 gives an overview over properties of C programs that lead to undefined behavior.
- 4 Note 3 to entry: Any other behavior during execution of a program is only affected as a direct consequence of the concrete behavior that occurs when encountering the erroneous or non-portable program construct or data. In particular, all observable behavior (5.2.2.4) appears as specified in this document when it happens before an operation with undefined behavior in the execution of the program.
- 5 EXAMPLE An example of undefined behavior is the behavior on dereferencing a null pointer.

3.5.4

1 **unspecified behavior**

behavior, that results from the use of an unspecified value, or other behavior upon which this document provides two or more possibilities and imposes no further requirements on which is chosen in any instance

- 2 Note 1 to entry: J.1 gives an overview over properties of C programs that lead to unspecified behavior.
- 3 EXAMPLE An example of unspecified behavior is the order in which the arguments to a function are evaluated.

3.6

1 **bit**

unit of data storage in the execution environment large enough to hold an object that can have one of two values

- 2 Note 1 to entry: It is possible that each individual bit of an object is not addressable.

3.7

1 **byte**

addressable unit of data storage large enough to hold any member of the basic character set of the execution environment

- 2 Note 1 to entry: Each individual byte of an object is uniquely addressable.

3.8

1 **byte array**

object having either no declared type or an array of objects declared with a byte type

3.9

1 **byte type**

non-atomic character type

3.10

1 **low-order bit**

the least significant bit within a byte

3.11

1 **high-order bit**

the most significant bit within a byte

3.12

1 **character**

⟨abstract⟩ member of a set of elements used for the organization, control, or representation of data

3.12.1

1 **character**

single-byte character

⟨C⟩ bit representation that fits in a byte

3.12.2

1 **multibyte character**

sequence of one or more bytes representing a member of the extended character set of either the source or the execution environment

2 Note 1 to entry: The extended character set is a superset of the basic character set.

3.12.3

1 **wide character**

value representable by an object of type `wchar_t`

3.13

1 **constraint**

restriction, either syntactic or semantic, by which the exposition of language elements is interpreted

3.14

1 **correctly rounded result**

representation in the result format that is nearest in value, subject to the current rounding mode, to what the result would be given unlimited range and precision

2 Note 1 to entry: In this document, the words “correctly rounded” may apply to an operation that produces a correctly rounded result.

3 Note 2 to entry: ISO/IEC 60559 or implementation-defined rules apply for extreme magnitude results if the result format contains infinity.

3.15

1 **diagnostic message**

message belonging to an implementation-defined subset of the implementation’s message output

3.16**1 forward reference**

reference to a subsequent subclause in this document that contains additional information relevant to the subclause containing the reference

3.17**1 implementation**

particular set of software, running in a particular translation environment under particular control options, that performs translation of programs for, and supports execution of functions in, a particular execution environment

3.18**1 implementation limit**

restriction imposed upon programs by the implementation

3.19**1 memory location**

either an object of scalar type, or a maximal sequence of adjacent bit-fields all having nonzero width

2 Note 1 to entry: Two threads of execution can update and access separate memory locations without interfering with each other.

3 Note 2 to entry: A bit-field and an adjacent non-bit-field member are in separate memory locations. The same applies to two bit-fields, if one is declared inside a nested structure declaration and the other is not, or if the two are separated by a zero-length bit-field declaration, or if they are separated by a non-bit-field member declaration. It is not safe to concurrently update two non-atomic bit-fields in the same structure if all members declared between them are also (nonzero-length) bit-fields, no matter what the sizes of those intervening bit-fields happen to be.

4 EXAMPLE A structure declared as

```
struct {
    char a;
    int b:5, c:11,:0, d:8;
    struct { int ee:8; } e;
}
```

contains four separate memory locations: The member **a**, and bit-fields **d** and **e.ee** are each separate memory locations, and can be modified concurrently without interfering with each other. The bit-fields **b** and **c** together constitute the fourth memory location. The bit-fields **b** and **c** cannot be concurrently modified, but **b** and **a**, for example, can be.

3.20**1 object**

region of data storage in the execution environment, the contents of which can represent values

2 Note 1 to entry: When referenced, an object can be interpreted as having a particular type; see 6.3.3.1.

3.21**1 parameter**

formal parameter

DEPRECATED: formal argument

object declared as part of a function declaration or definition that acquires a value on entry to the function, or an identifier from the comma-separated list bounded by the parentheses immediately following the macro name in a function-like macro definition

3.22**1 recommended practice**

specification that is strongly recommended as being in keeping with the intent of the standard, but that may be impractical for some implementations

3.23**1 runtime-constraint**

requirement on a program when calling a library function

- 2 Note 1 to entry: Despite the similar terms, a runtime-constraint is not a kind of constraint as defined by 3.13, and it is not necessary for it to be diagnosed at translation time.
- 3 Note 2 to entry: Implementations that support the extensions in Annex K are required to verify that the runtime-constraints for a library function are not violated by the program; see K.3.1.4.
- 4 Note 3 to entry: Implementations that support Annex L are permitted to invoke a runtime-constraint handler when they perform a trap.

3.24**1 value**

precise meaning of the contents of an object when interpreted as having a specific type

3.24.1**1 implementation-defined value**

unspecified value where each implementation documents how the choice is made

3.24.2**1 unspecified value**

valid value of the relevant type where this document imposes no requirements on which value is chosen in any instance

3.25**1 indeterminate representation**

object representation that either represents an unspecified value or is a non-value representation

3.26**1 non-value representation**

an object representation that does not represent a value of the object type

3.27**1 perform a trap**

interrupt execution of the program such that no further operations are performed

- 2 Note 1 to entry: Fetching a non-value representation permits an implementation to perform a trap but is not required to (see 6.2.6.1).
- 3 Note 2 to entry: Implementations that support Annex L are permitted to invoke a runtime-constraint handler when they perform a trap.

3.28**1 $\lceil x \rceil$**

ceiling of x

the least integer greater than or equal to x

- 2 EXAMPLE $\lceil 2.4 \rceil$ is 3, $\lceil -2.4 \rceil$ is -2 .

3.29

- 1 $\lfloor x \rfloor$
floor of x
the greatest integer less than or equal to x

- 2 EXAMPLE $\lfloor 2.4 \rfloor$ is 2, $\lfloor -2.4 \rfloor$ is -3 .

3.30

- 1 **wraparound**
the process by which a value is reduced modulo 2^N , where N is the width of the resulting type

3.31

- 1 **out-of-bounds store**
(attempted) access (3.1) that, at run time, for a given computational state, would modify (or, for an object declared **volatile**, fetch) one or more bytes that lie outside the bounds permitted by this document.

3.32

- 1 **bounded undefined behavior**
undefined behavior (3.5.3) that does not perform an out-of-bounds store.
- 2 Note 1 to entry: The behavior can perform a trap.
- 3 Note 2 to entry: Any values produced can be unspecified values, and the representation of objects that are written to can become indeterminate.

3.33

- 1 **critical undefined behavior**
undefined behavior that is not bounded undefined behavior.
- 2 Note 1 to entry: The behavior can perform an out-of-bounds store or perform a trap.

4. Conformance

- 1 In this document, “shall” is to be interpreted as a requirement on an implementation or on a program; conversely, “shall not” is to be interpreted as a prohibition.
- 2 If a “shall” or “shall not” requirement that appears outside of a constraint or runtime-constraint is violated, the behavior is undefined. Undefined behavior is otherwise indicated in this document by the words “undefined behavior” or by the omission of any explicit definition of behavior. There is no difference in emphasis among these three; they all describe “behavior that is undefined”.
- 3 A program that is correct in all other aspects, operating on correct data, containing unspecified behavior shall be a correct program and act in accordance with 5.2.2.4.
- 4 The implementation shall not successfully translate a preprocessing translation unit containing a **#error** preprocessing directive unless it is part of a group skipped by conditional inclusion.
- 5 A *strictly conforming program* shall use only those features of the language and library specified in this document. It shall not produce output dependent on any unspecified, undefined, or implementation-defined behavior, and shall not exceed any minimum implementation limit.
- 6 EXAMPLE A strictly conforming program can use conditional features (see 6.10.10.4) provided the use is guarded by an appropriate conditional inclusion preprocessing directive using the related macro. For example:

```
#ifdef __STDC_IEC_60559_BFP__ /* FE_UPWARD defined */
    /* ... */
    fesetround(FE_UPWARD);
    /* ... */
#endif
```

- 7 The two forms of *conforming implementation* are hosted and freestanding. A *conforming hosted implementation* shall accept any strictly conforming program. A *conforming freestanding implementation* shall accept any strictly conforming program in which the use of the features specified in the library clause (Clause 7) is confined to the contents of the standard headers `<float.h>`, `<iso646.h>`, `<limits.h>`, `<stdalign.h>`, `<stdarg.h>`, `<stdbit.h>`, `<stdbool.h>`, `<stddef.h>`, `<stdint.h>`, and `<stdnoreturn.h>`. Additionally, a *conforming freestanding implementation* shall accept any strictly conforming program where:
 - the features specified in the header `<string.h>` are used, except the following functions: **strcoll**, **strdup**, **strerror**, **strndup**, **strtok**, **strxfrm**; and/or,
 - the selected function **memalign** from `<stdlib.h>` is used.

A conforming implementation may have extensions (including additional library functions), provided they do not alter the behavior of any strictly conforming program.¹⁾

- 8 The strictly conforming programs that shall be accepted by a conforming freestanding implementation that defines `__STDC_IEC_60559_BFP__` or `__STDC_IEC_60559_DFP__` can also use features in the contents of the standard headers `<fenv.h>`, `<math.h>`, and the **strt*** floating-point numeric conversion functions (7.24.2) of the standard header `<stdlib.h>`, provided the program does not set the state of the **FENV_ACCESS** pragma to “on”.

All identifiers that are reserved when `<stdlib.h>` is included in a hosted implementation are reserved when it is included in a freestanding implementation.

- 9 A *conforming program* is one that is acceptable to a conforming implementation.²⁾

¹⁾This implies that a conforming implementation reserves no identifiers other than those explicitly reserved in this document.

²⁾Strictly conforming programs are intended to be maximally portable among conforming implementations. Conforming programs can depend upon nonportable features of a conforming implementation.

- 10 An implementation shall be accompanied by a document that defines all implementation-defined and locale-specific characteristics and all extensions.

Forward references: conditional inclusion (6.10.2), error directive (6.10.7), characteristics of floating types `<float.h>` (7.7), alternative spellings `<iso646.h>` (7.9), sizes of integer types `<limits.h>` (7.10), alignment `<stdalign.h>` (7.15), varying arguments `<stdarg.h>` (7.16), boolean type and values `<stdbool.h>` (7.19), common definitions `<stddef.h>` (7.21), integer types `<stdint.h>` (7.22), `<stdnoreturn.h>` (7.26).

5. Environment

5.1 Introduction

- 1 An implementation translates C source files and executes C programs in two data-processing-system environments, which will be called the *translation environment* and the *execution environment* in this document. Their characteristics define and constrain the results of executing conforming C programs constructed according to the syntactic and semantic rules for conforming implementations.
- 2 NOTE In this clause, only a few of many possible forward references have been noted.

5.2 Conceptual models

5.2.1 Translation environment

5.2.1.1 Program structure

- 1 A C program is not required to be translated in its entirety at the same time. The text of the program is kept in units called *source files*, (or *preprocessing files*) in this document. A source file together with all the headers and source files included via the preprocessing directive **#include** is known as a *preprocessing translation unit*. After preprocessing, a preprocessing translation unit is called a *translation unit*. Previously translated translation units can be preserved individually or in libraries. The separate translation units of a program communicate by (for example) calls to functions whose identifiers have external linkage, manipulation of objects whose identifiers have external linkage, or manipulation of data files. Translation units may be separately translated and then later linked to produce an executable program.

Forward references: linkages of identifiers (6.2.2), external definitions (6.9), preprocessing directives (6.10).

5.2.1.2 Translation phases

- 1 The precedence among the syntax rules of translation is specified by the following phases.³⁾
 1. Physical source file multibyte characters are mapped, in an implementation-defined manner, to the source character set (introducing new-line characters for end-of-line indicators) if necessary.
 2. Each instance of a backslash character (\) immediately followed by a new-line character is deleted, splicing physical source lines to form logical source lines. Only the last backslash on any physical source line shall be eligible for being part of such a splice. A source file that is not empty shall end in a new-line character, which shall not be immediately preceded by a backslash character before any such splicing takes place.
 3. The source file is decomposed into preprocessing tokens⁴⁾ and sequences of white-space characters (including comments). A source file shall not end in a partial preprocessing token or in a partial comment. Each comment is replaced by one space character. New-line characters are retained. Whether each nonempty sequence of white-space characters other than new-line is retained or replaced by one space character is implementation-defined.
 4. Preprocessing directives are executed, macro invocations are expanded, and **_Pragma** unary operator expressions are executed. If a character sequence that matches the syntax of a universal character name is produced by token concatenation (6.10.5.4), the behavior is undefined. A **#include** preprocessing directive causes the named header or source file to be processed from phase 1 through phase 4, recursively. All preprocessing directives are then deleted.

³⁾This requires implementations to behave as if these separate phases occur, even though many are typically folded together in practice. Source files, translation units, and translated translation units necessarily can be stored as files or through/within any other implementation-defined medium. There is no expectation of any one-to-one correspondence between these entities and any external representation. The description is conceptual only, and does not specify any particular implementation.

⁴⁾As described in 6.4, the process of dividing a source file's characters into preprocessing tokens is context-dependent. For example, see the handling of < within a **#include** preprocessing directive.

5. Each source character set member and escape sequence in character literals and string literals is converted to the corresponding member of the execution character set. Each instance of a source character or escape sequence for which there is no corresponding member is converted in an implementation-defined manner to some member of the execution character set other than the null (wide) character.⁵⁾
6. Adjacent string literal tokens are concatenated.
7. White-space characters separating tokens are no longer significant. Each preprocessing token is converted into a token. The resulting tokens are syntactically and semantically analyzed and translated as a translation unit.
8. All external object and function references are resolved. Library components are linked to satisfy external references to functions and objects not defined in the current translation. All such translator output is collected into a program image which contains information needed for execution in its execution environment.

Forward references: universal character names (6.4.4), lexical elements (6.4), preprocessing directives (6.10), external definitions (6.9).

5.2.1.3 Diagnostics

- 1 A conforming implementation shall produce at least one diagnostic message (identified in an implementation-defined manner) if a preprocessing translation unit or translation unit contains a violation of any syntax rule or constraint, even if the behavior is also explicitly specified as undefined or implementation-defined. Diagnostic messages are not required to be produced in other circumstances.
- 2 **EXAMPLE** An implementation is required to issue a diagnostic for the translation unit:

```
char i;  
int i;
```

because in those cases where wording in this document describes the behavior for a construct as being both a constraint error and resulting in undefined behavior, the constraint error is still required to be diagnosed.

Recommended practice

- 3 An implementation is encouraged to identify the nature of, and where possible localize, each violation. Of course, an implementation is free to produce any number of diagnostic messages, often referred to as warnings, as long as a valid program is still correctly translated. It can also successfully translate an invalid program. Annex I lists a few of the more common warnings.

5.2.2 Execution environments

5.2.2.1 General

- 1 Two execution environments are defined: *freestanding* and *hosted*. In both cases, *program startup* occurs when a designated C function is called by the execution environment. All objects with static storage duration shall be *initialized* (set to their initial values) before program startup. The manner and timing of such initialization are otherwise unspecified. *Program termination* returns control to the execution environment.

Forward references: storage durations of objects (6.2.4), initialization (6.7.11).

5.2.2.2 Freestanding environment

- 1 In a freestanding environment (in which C program execution can take place without any benefit of an operating system), the name and type of the function called at program startup are implementation-defined. Any library facilities available to a freestanding program, other than the minimal set required by Clause 4, are implementation-defined.
- 2 The effect of program termination in a freestanding environment is implementation-defined.

⁵⁾An implementation can convert each instance of the same non-corresponding source character to a different member of the execution character set.

5.2.2.3 Hosted environment

5.2.2.3.1 General

- 1 A hosted environment is not required to be provided, but shall conform to the following specifications if present.

5.2.2.3.2 Program startup

- 1 The function called at program startup is named **main**. The implementation declares no prototype for this function. It shall be defined with a return type of **int** and with no parameters:

```
int main(void) { /* ... */ }
```

or with two parameters (referred to here as **argc** and **argv**, though any names can be used, as they are local to the function in which they are declared):

```
int main(int argc, char *argv[]) { /* ... */ }
```

or equivalent;⁶⁾ or in some other implementation-defined manner.

- 2 If they are declared, the parameters to the **main** function shall obey the following constraints:
 - The value of **argc** shall be nonnegative.
 - **argv[argc]** shall be a null pointer.
 - If the value of **argc** is greater than zero, the array members **argv[0]** through **argv[argc-1]** inclusive shall contain pointers to strings, which are given implementation-defined values by the host environment prior to program startup. The intent is to supply to the program information determined prior to program startup from elsewhere in the hosted environment. If the host environment is not capable of supplying strings with letters in both uppercase and lowercase, the implementation shall ensure that the strings are received in lowercase.
 - If the value of **argc** is greater than zero, the string pointed to by **argv[0]** represents the *program name*; **argv[0][0]** shall be the null character if the program name is not available from the host environment. If the value of **argc** is greater than one, the strings pointed to by **argv[1]** through **argv[argc-1]** represent the *program parameters*.
 - The parameters **argc** and **argv** and the strings pointed to by the **argv** array shall be modifiable by the program, and retain their last-stored values between program startup and program termination.

5.2.2.3.3 Program execution

- 1 In a hosted environment, a program can use all the functions, macros, type definitions, and objects described in the library clause (Clause 7).

5.2.2.3.4 Program termination

- 1 If the return type of the **main** function is a type compatible with **int**, a return from the initial call to the **main** function is equivalent to calling the **exit** function with the value returned by the **main** function as its argument;⁷⁾ reaching the **}** that terminates the **main** function returns a value of 0. If the return type is not compatible with **int**, the termination status returned to the host environment is unspecified.

Forward references: definition of terms (7.1.1), the **exit** function (7.24.5.4).

⁶⁾Thus, **int** can be replaced by a typedef name defined as **int**, or the type of **argv** can be written as **char ** argv**, or the return type can be specified by **typeof(1)**, and so on.

⁷⁾In accordance with 6.2.4, the lifetimes of objects with automatic storage duration declared in **main** will have ended in the former case, even where they would not have in the latter.

5.2.2.4 Program semantics

- 1 The semantic descriptions in this document describe the behavior of an abstract machine in which issues of optimization are irrelevant.
- 2 An access to an object through the use of an lvalue of volatile-qualified type is a *volatile access*. A volatile access to an object, modifying an object, modifying a file, or calling a function that does any of those operations are all *side effects*,⁸⁾ which are changes in the state of the execution environment. *Evaluation* of an expression in general includes both value computations and initiation of side effects. Value computation for an lvalue expression includes determining the identity of the designated object.
- 3 *Sequenced before* is an asymmetric, transitive, pair-wise relation between evaluations executed by a single thread, which induces a partial order among those evaluations. Given any two evaluations *A* and *B*, if *A* is sequenced before *B*, then the execution of *A* shall precede the execution of *B*. (Conversely, if *A* is sequenced before *B*, then *B* is *sequenced after A*.) If *A* is not sequenced before or after *B*, then *A* and *B* are *unsequenced*. Evaluations *A* and *B* are *indeterminately sequenced* when *A* is sequenced either before or after *B*, but it is unspecified which.⁹⁾ The presence of a *sequence point* between the evaluation of expressions *A* and *B* implies that every value computation and side effect associated with *A* is sequenced before every value computation and side effect associated with *B*. (A summary of the sequence points is given in Annex C.)
- 4 In the abstract machine, all expressions are evaluated as specified by the semantics. An actual implementation is not required to evaluate part of an expression if it can deduce that its value is not used and that no needed side effects are produced (including any caused by calling a function or through volatile access to an object).
- 5 When the processing of the abstract machine is interrupted by receipt of a signal, the values of objects that are neither lock-free atomic objects nor of type **volatile sig_atomic_t** are unspecified, as is the state of the dynamic floating-point environment. The representation of any object modified by the handler that is neither a lock-free atomic object nor of type **volatile sig_atomic_t** becomes indeterminate when the handler exits, as does the state of the dynamic floating-point environment if it is modified by the handler and not restored to its original state.
- 6 The least requirements on a conforming implementation are:
 - Volatile accesses to objects are evaluated strictly according to the rules of the abstract machine.
 - At program termination, all data written into files shall be identical to the result that execution of the program according to the abstract semantics would have produced.
 - The input and output dynamics of interactive devices shall take place as specified in 7.23.3. The intent of these requirements is that unbuffered or line-buffered output appear as soon as possible, to ensure that prompting messages appear prior to a program waiting for input.

This is the *observable behavior* of the program.

- 7 What constitutes an interactive device is implementation-defined.
- 8 More stringent correspondences between abstract and actual semantics may be defined by each implementation.
- 9 EXAMPLE 1 An implementation can define a one-to-one correspondence between abstract and actual semantics: at every sequence point, the values of the actual objects would agree with those specified by the abstract semantics. The keyword **volatile** would then be redundant.
- 10 Alternatively, an implementation can perform various optimizations within each translation unit, such that the actual semantics would agree with the abstract semantics only when making function calls across translation

⁸⁾ISO/IEC 60559 requires certain user-accessible status flags and control modes. Floating-point operations implicitly set the status flags; modes affect result values of floating-point operations. Implementations that support such floating-point state are required to regard changes to it as side effects — see Annex F for details. The floating-point environment library <fenv.h> provides a programming facility for indicating when these side effects matter, freeing the implementations in other cases.

⁹⁾The executions of unsequenced evaluations can interleave. Indeterminately sequenced evaluations cannot interleave, but can be executed in any order.

unit boundaries. In such an implementation, at the time of each function entry and function return where the calling function and the called function are in different translation units, the values of all externally linked objects and of all objects accessible via pointers therein would agree with the abstract semantics. Furthermore, at the time of each such function entry the values of the parameters of the called function and of all objects accessible via pointers therein would agree with the abstract semantics. In this type of implementation, objects referred to by interrupt service routines activated by the **signal** function would require explicit specification of **volatile** storage, as well as other implementation-defined restrictions.

- 11 EXAMPLE 2 In executing the fragment

```
char c1, c2;
/* ... */
c1 = c1 + c2;
```

the “integer promotions” require that the abstract machine promote the value of each variable to **int** size and then add the two **ints** and truncate the sum. Provided the addition of two **chars** can be done without integer overflow, or with integer overflow wrapping silently to produce the correct result, the actual execution need only produce the same result, possibly omitting the promotions.

- 12 EXAMPLE 3 Similarly, in the fragment

```
float f1, f2;
double d;
/* ... */
f1 = f2 * d;
```

the multiplication can be executed using **float** arithmetic if the implementation can ascertain that the result would be the same as if it were executed using **double** arithmetic (for example, if **d** were replaced by the literal 2.0, which has type **double**).

- 13 EXAMPLE 4 Implementations employing wide registers are expected to take care to honor appropriate semantics. Values are independent of whether they are represented in a register or in memory. For example, an implicit *spilling* of a register cannot alter the value. Also, an explicit *store and load* rounds to the precision of the storage type. In particular, casts and assignments perform their specified conversion. For the fragment

```
double d1, d2;
float f;
d1 = f = expression;
d2 = (float) expression;
```

the values assigned to **d1** and **d2** are required to have been converted to **float**.

- 14 EXAMPLE 5 Rearrangement for floating-point expressions is often restricted because of limitations in precision as well as range. The implementation cannot generally apply the mathematical associative rules for addition or multiplication, nor the distributive rule, because of roundoff error, even in the absence of overflow and underflow. Likewise, implementations cannot generally replace decimal literals to rearrange expressions. In the following fragment, rearrangements suggested by mathematical rules for real numbers are often not valid (see F.9).

```
double x, y, z;
/* ... */
x = (x * y) * z; // not equivalent to x *= y * z;
z = (x - y) + y; // not equivalent to z = x;
z = x + x * y;   // not equivalent to z = x * (1.0 + y);
y = x / 5.0;     // not equivalent to y = x * 0.2;
```

- 15 EXAMPLE 6 To illustrate the grouping behavior of expressions, in the following fragment

```
int a, b;
/* ... */
a = a + 32760 + b + 5;
```

the expression statement behaves exactly the same as

```
a = (((a + 32760) + b) + 5);
```

due to the associativity and precedence of these operators. Thus, the result of the sum ($a + 32760$) is next added to b , and that result is then added to 5 which results in the value assigned to a . On a machine in which integer overflows produce an explicit trap and in which the range of values representable by an `int` is $[-32768, +32767]$, the implementation cannot rewrite this expression as

```
a = ((a + b) + 32765);
```

since if the values for a and b were, respectively, -32754 and -15 , the sum $a + b$ would produce a trap while the original expression would not; nor can the expression be rewritten either as

```
a = ((a + 32765) + b);
```

or

```
a = (a + (b + 32765));
```

since the values for a and b can be, respectively, 4 and -8 or -17 and 12. However, on an implementation in which integer overflow silently generates some value and where positive and negative integer overflows cancel, the preceding expression statement can be rewritten by the implementation in any of the previously specified ways because the same result will occur.

- 16 EXAMPLE 7 The grouping of an expression does not completely determine its evaluation. In the following fragment:

```
#include <stdio.h>
int sum;
char *p;
/* ... */
sum = sum * 10 - '0' + (*p++ = getchar());
```

the expression statement is grouped as if it were written as

```
sum = (((sum * 10) - '0') + ((*p++) = (getchar())));
```

but the actual increment of p can occur at any time between the previous sequence point and the next sequence point (the `;`), and the call to `getchar` can occur at any point prior to the need of its returned value.

Forward references: expressions (6.5.1), type qualifiers (6.7.4), statements (6.8), floating-point environment `<fenv.h>` (7.6), the `signal` function (7.14), files (7.23.3).

5.2.2.5 Multi-threaded executions and data races

- 1 Under a hosted implementation that does not define `__STDC_NO_THREADS__`, a program can have more than one *thread of execution* (or *thread*) running concurrently. The execution of each thread proceeds as defined by the remainder of this document. The execution of the entire program consists of an execution of all its threads.¹⁰ Under a freestanding implementation, it is implementation-defined whether a program can have more than one thread of execution.
- 2 The value of an object visible to a thread T at a particular point is the initial value of the object, a value stored in the object by T , or a value stored in the object by another thread, according to the rules in the rest of this subclause.
- 3 NOTE 1 In some cases, there can instead be undefined behavior. Much of this subclause is motivated by the desire to support atomic operations with explicit and detailed visibility constraints. However, it also implicitly supports a simpler view for more restricted programs.
- 4 Two expression evaluations *conflict* if one of them modifies a memory location and the other one reads or modifies the same memory location.

¹⁰The execution can usually be viewed as an interleaving of all the threads. However, some kinds of atomic operations, for example, allow executions inconsistent with a simple interleaving as described in this subclause.

- 5 The library defines *atomic operations* (7.17) and operations on mutexes (7.29.4) that are specially identified as synchronization operations. These operations play a special role in making assignments in one thread visible to another. A *synchronization operation* on one or more memory locations is one of an *acquire operation*, a *release operation*, both an acquire and release operation, or a *consume operation*. A synchronization operation without an associated memory location is a *fence* and can be either an acquire fence, a release fence, or both an acquire and release fence. In addition, there are *relaxed atomic operations*, which are not synchronization operations, and atomic *read-modify-write operations*, which have special characteristics.
- 6 NOTE 2 For example, a call that acquires a mutex will perform an acquire operation on the locations composing the mutex. Correspondingly, a call that releases the same mutex will perform a release operation on those same locations. Informally, performing a release operation on *A* forces prior side effects on other memory locations to become visible to other threads that later perform an acquire or consume operation on *A*. Relaxed atomic operations are not included as synchronization operations although, like synchronization operations, they cannot contribute to data races.
- 7 All modifications to a particular atomic object *M* occur in some particular total order, called the *modification order* of *M*. If *A* and *B* are modifications of an atomic object *M*, and *A* happens before *B*, then *A* shall precede *B* in the modification order of *M*, which is defined later in this subclause.
- 8 NOTE 3 This states that the modification orders are expected to respect the “happens before” relation.
- 9 NOTE 4 There is a separate order for each atomic object. There is no requirement that these can be combined into a single total order for all objects. In general this will be impossible since different threads can observe modifications to different variables in inconsistent orders.
- 10 A *release sequence* headed by a release operation *A* on an atomic object *M* is a maximal contiguous sub-sequence of side effects in the modification order of *M*, where the first operation is *A* and every subsequent operation either is performed by the same thread that performed the release or is an atomic read-modify-write operation.
- 11 Certain library calls *synchronize with* other library calls performed by another thread. In particular, an atomic operation *A* that performs a release operation on an object *M* synchronizes with an atomic operation *B* that performs an acquire operation on *M* and reads a value written by any side effect in the release sequence headed by *A*.
- 12 NOTE 5 Except in the specified cases, reading a later value does not necessarily ensure visibility as described later in this subclause. Such a requirement would sometimes interfere with efficient implementation.
- 13 NOTE 6 The specifications of the synchronization operations define when one reads the value written by another. For atomic variables, the definition is clear. All operations on a given mutex occur in a single total order. Each mutex acquisition “reads the value written” by the last mutex release.
- 14 An evaluation *A* carries a dependency¹¹⁾ to an evaluation *B* if:
- the value of *A* is used as an operand of *B*, unless:
 - *B* is an invocation of the **kill_dependency** macro,
 - *A* is the left operand of a && or || operator,
 - *A* is the left operand of a ?: operator, or
 - *A* is the left operand of a , operator;
 - or
 - *A* writes a scalar object or bit-field *M*, *B* reads from *M* the value written by *A*, and *A* is sequenced before *B*, or
 - for some evaluation *X*, *A* carries a dependency to *X* and *X* carries a dependency to *B*.
- 15 An evaluation *A* is *dependency-ordered before*¹²⁾ an evaluation *B* if:

¹¹⁾The “carries a dependency” relation is a subset of the “sequenced before” relation, and is similarly strictly intra-thread.

¹²⁾The “dependency-ordered before” relation is analogous to the “synchronizes with” relation, but uses release/consume in place of release/acquire.

- *A* performs a release operation on an atomic object *M*, and, in another thread, *B* performs a consume operation on *M* and reads a value written by any side effect in the release sequence headed by *A*, or
 - for some evaluation *X*, *A* is dependency-ordered before *X* and *X* carries a dependency to *B*.
- 16 An evaluation *A* *inter-thread happens before* an evaluation *B* if *A* synchronizes with *B*, *A* is dependency-ordered before *B*, or, for some evaluation *X*:
- *A* synchronizes with *X* and *X* is sequenced before *B*,
 - *A* is sequenced before *X* and *X* inter-thread happens before *B*, or
 - *A* inter-thread happens before *X* and *X* inter-thread happens before *B*.
- 17 NOTE 7 The “inter-thread happens before” relation describes arbitrary concatenations of “sequenced before”, “synchronizes with”, and “dependency-ordered before” relationships, with two exceptions. The first exception is that a concatenation is not permitted to end with “dependency-ordered before” followed by “sequenced before”. The reason for this limitation is that a consume operation participating in a “dependency-ordered before” relationship provides ordering only with respect to operations to which this consume operation carries a dependency. The reason that this limitation applies only to the end of such a concatenation is that any subsequent release operation will provide the required ordering for a prior consume operation. The second exception is that a concatenation is not permitted to consist entirely of “sequenced before”. The reasons for this limitation are (1) to permit “inter-thread happens before” to be transitively closed and (2) the “happens before” relation, defined subsequently in this subclause, provides for relationships consisting entirely of “sequenced before”.
- 18 An evaluation *A* *happens before* an evaluation *B* if *A* is sequenced before *B* or *A* inter-thread happens before *B*. The implementation shall ensure that no program execution demonstrates a cycle in the “happens before” relation.
- 19 NOTE 8 This cycle would otherwise be possible only through the use of consume operations.
- 20 A *visible side effect* *A* on an object *M* with respect to a value computation *B* of *M* satisfies the conditions:
- *A* happens before *B*, and
 - there is no other side effect *X* to *M* such that *A* happens before *X* and *X* happens before *B*.
- The value of a non-atomic scalar object *M*, as determined by evaluation *B*, shall be the value stored by the visible side effect *A*.
- 21 NOTE 9 If there is ambiguity about which side effect to a non-atomic object is visible, then there is a data race and the behavior is undefined.
- 22 NOTE 10 This states that operations on ordinary variables are not visibly reordered. This is not detectable without data races, but ensures that data races, as defined here, and with suitable restrictions on the use of atomics, correspond to data races in a simple interleaved (sequentially consistent) execution.
- 23 The value of an atomic object *M*, as determined by evaluation *B*, shall be the value stored by some side effect *A* that modifies *M*, where *B* does not happen before *A*.
- 24 NOTE 11 The set of side effects from which a given evaluation can take its value is also restricted by the rest of the rules described here, and in particular, by the coherence requirements subsequently in this subclause.
- 25 If an operation *A* that modifies an atomic object *M* happens before an operation *B* that modifies *M*, then *A* shall be earlier than *B* in the modification order of *M*.
- 26 NOTE 12 Such a requirement is known as “write-write coherence”.
- 27 If a value computation *A* of an atomic object *M* happens before a value computation *B* of *M*, and *A* takes its value from a side effect *X* on *M*, then the value computed by *B* shall either be the value stored by *X* or the value stored by a side effect *Y* on *M*, where *Y* follows *X* in the modification order of *M*.

- 28 NOTE 13 Such a requirement is known as “read-read coherence”.
- 29 If a value computation *A* of an atomic object *M* happens before an operation *B* on *M*, then *A* shall take its value from a side effect *X* on *M*, where *X* precedes *B* in the modification order of *M*.
- 30 NOTE 14 Such a requirement is known as “read-write coherence”.
- 31 If a side effect *X* on an atomic object *M* happens before a value computation *B* of *M*, then the evaluation *B* shall take its value from *X* or from a side effect *Y* that follows *X* in the modification order of *M*.
- 32 NOTE 15 Such a requirement is known as “write-read coherence”.
- 33 NOTE 16 This effectively disallows compiler reordering of atomic operations to a single object, even if both operations are “relaxed” loads. By doing so, it effectively makes the “cache coherence” guarantee provided by most hardware available to C atomic operations.
- 34 NOTE 17 The value observed by a load of an atomic object depends on the “happens before” relation, which in turn depends on the values observed by loads of atomic objects. The intended reading is that there exists an association of atomic loads with modifications they observe that, together with suitably chosen modification orders and the “happens before” relation derived as described previously, satisfy the resulting constraints as imposed here.
- 35 The execution of a program contains a *data race* if it contains two conflicting actions in different threads, at least one of which is not atomic, and neither happens before the other. Any such data race results in undefined behavior.
- 36 NOTE 18 It can be shown that programs that correctly use simple mutexes and `memory_order_seq_cst` operations to prevent all data races, and use no other synchronization operations, behave as though the operations executed by their constituent threads were simply interleaved, with each value computation of an object being the last value stored in that interleaving. This is normally referred to as “sequential consistency”. However, this applies only to data-race-free programs, and data-race-free programs cannot observe most program transformations that do not change single-threaded program semantics. In fact, most single-threaded program transformations continue to be allowed, since any program that behaves differently as a result of such transformations necessarily has undefined behavior even before such a transformation is applied.
- 37 NOTE 19 Compiler transformations that introduce assignments to a potentially shared memory location that would not be modified by the abstract machine are generally precluded by this document, since such an assignment can overwrite another assignment by a different thread in cases in which an abstract machine execution would not have encountered a data race. This includes implementations of data member assignment that overwrite adjacent members in separate memory locations. Reordering of atomic loads in cases in which the atomics in question can alias is also generally precluded, since this can violate the coherence requirements.
- 38 NOTE 20 Transformations that introduce a speculative read of a potentially shared memory location possibly will not preserve the semantics of the program as defined in this document, since they potentially introduce a data race. However, they are typically valid in the context of an optimizing compiler that targets a specific machine with well-defined semantics for data races. They would be invalid for a hypothetical machine that is not tolerant of races or provides hardware race detection.

5.3 Environmental considerations

5.3.1 Character sets

- 1 Two sets of characters and their associated *collating sequences* shall be defined: the set in which source files are written (the *source character set*), and the set interpreted in the execution environment (the *execution character set*). Each set is further divided into a *basic character set*, whose contents are given by this subclause, and a set of zero or more locale-specific members (which are not members of the basic character set) called *extended characters*. The combined set is also called the *extended character set*. The values of the members of the execution character set are implementation-defined.
- 2 In a character literal or string literal, members of the execution character set shall be represented by corresponding members of the source character set or by *escape sequences* consisting of the backslash `\` followed by one or more characters. A byte with all bits set to 0, called the *null character*, shall exist in the basic execution character set; it is used to terminate a character string.
- 3 Both the basic source and basic execution character sets shall have the following members: the 26 *uppercase letters* of the *Latin alphabet*

A	B	C	D	E	F	G	H	I	J	K	L	M
N	O	P	Q	R	S	T	U	V	W	X	Y	Z

the 26 *lowercase letters* of the Latin alphabet

a	b	c	d	e	f	g	h	i	j	k	l	m
n	o	p	q	r	s	t	u	v	w	x	y	z

the 10 decimal *digit*

0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

the following 32 *graphic characters*

!	"	#	%	&	'	(	)	*	+	,	-	.	/	:
;	<	=	>	?	[	\	]	^	_	{		}	~	
@	\$	`												

the space character, and control characters representing horizontal tab, vertical tab, and form feed. The representation of each member of the source and execution basic character sets shall fit in a byte. In both the source and execution basic character sets, the value of each character after 0 in the preceding list of decimal digits shall be one greater than the value of the previous. In source files, there shall be some way of indicating the end of each line of text; this document treats such an end-of-line indicator as if it were a single new-line character. In the basic execution character set, there shall be control characters representing alert, backspace, carriage return, and new line. If any other characters are encountered in a source file (except in an identifier, a character literal, a string literal, a header name, a comment, or a preprocessing token that is never converted to a token), the behavior is undefined.

- 4 The value of each character after **a**, up to and including **f**, in the prior specified list of lowercase letters, shall be one greater than the value of the previous.
- 5 The value of each character after **A**, up to and including **F**, in the prior specified list of uppercase letters, shall be one greater than the value of the previous.
- 6 A *letter* is an uppercase letter or a lowercase letter as defined previously in this subclause; in this document the term does not include other characters that are letters in other alphabets.
- 7 The universal character name construct provides a way to name other characters.

Forward references: universal character names (6.4.4), character literals (6.4.5.5), preprocessing directives (6.10), string literals (6.4.6), comments (6.4.10), string (7.1.1).

5.3.2 Multibyte characters

- 1 The source character set can contain multibyte characters, used to represent members of the extended character set. The execution character set can also contain multibyte characters, which are not required to have the same encoding as for the source character set. For both character sets, the following shall hold:
 - The basic character set shall be present and each character shall be encoded as a single byte.
 - The presence, meaning, and representation of any additional members is locale-specific.
 - A multibyte character set can have a *state-dependent encoding*, wherein each sequence of multibyte characters begins in an *initial shift state* and enters other locale-specific *shift states* when specific multibyte characters are encountered in the sequence. While in the initial shift state, all single-byte characters retain their usual interpretation and do not alter the shift state. The interpretation for subsequent bytes in the sequence is a function of the current shift state.
 - A byte with all bits zero shall be interpreted as a null character independent of shift state. Such a byte shall not occur as part of any other multibyte character.

- 2 For source files, the following shall hold:
 - An identifier, comment, string literal, character literal, or header name shall begin and end in the initial shift state.
 - An identifier, comment, string literal, character literal, or header name shall consist of a sequence of valid multibyte characters.

5.3.3 Character display semantics

- 1 The *active position* is that location on a display device where the next character output by the **fputc** function would appear. The intent of writing a printing character (as defined by the **isprint** function) to a display device is to display a graphic representation of that character at the active position and then advance the active position to the next position on the current line. The direction of writing is locale-specific. If the active position is at the final position of a line (if there is one), the behavior of the display device is unspecified.
- 2 Alphabetic escape sequences representing non-graphic characters in the execution character set are intended to produce actions on display devices as follows:
 - \a** (*alert*) Produces an audible or visible alert without changing the active position.
 - \b** (*backspace*) Moves the active position to the previous position on the current line. If the active position is at the initial position of a line, the behavior of the display device is unspecified.
 - \f** (*form feed*) Moves the active position to the initial position at the start of the next logical page.
 - \n** (*new line*) Moves the active position to the initial position of the next line.
 - \r** (*carriage return*) Moves the active position to the initial position of the current line.
 - \t** (*horizontal tab*) Moves the active position to the next horizontal tabulation position on the current line. If the active position is at or past the last defined horizontal tabulation position, the behavior of the display device is unspecified.
 - \v** (*vertical tab*) Moves the active position to the initial position of the next vertical tabulation position. If the active position is at or past the last defined vertical tabulation position, the behavior of the display device is unspecified.
- 3 Each of these escape sequences shall produce a unique implementation-defined value which can be stored in a single **char** object. The external representations in a text file are not necessarily identical to the internal representations, and are outside the scope of this document.

Forward references: the **isprint** function (7.4.2.9), the **fputc** function (7.23.7.3).

5.3.4 Signals and interrupts

- 1 Functions shall be implemented such that they can be interrupted at any time by a signal, or can be called by a signal handler, or both, with no alteration to earlier, but still active, invocations' control flow (after the interruption), function return values, or objects with automatic storage duration. All such objects shall be maintained outside the *function image* (the instructions that compose the executable representation of a function) on a per-invocation basis.

5.3.5 Environmental limits

5.3.5.1 General

- 1 Both the translation and execution environments constrain the implementation of language translators and libraries. The following summarizes the language-related environmental limits on a conforming implementation; the library-related limits are discussed in Clause 7.

5.3.5.2 Translation limits

- 1 The implementation shall be able to translate and execute a program that uses but does not exceed the following limitations for these constructs and entities:¹³⁾
 - 127 nesting levels of blocks
 - 63 nesting levels of conditional inclusion
 - 12 pointer, array, and function declarators (in any combinations) modifying an arithmetic, structure, union, or **void** type in a declaration
 - 63 nesting levels of parenthesized declarators within a full declarator
 - 63 nesting levels of parenthesized expressions within a full expression
 - 63 significant initial characters in an internal identifier or a macro name (each universal character name or extended source character is considered a single character)
 - 31 significant initial characters in an external identifier (each universal character name specifying a short identifier of 00FFFF or less is considered 6 characters, each universal character name specifying a short identifier of 010000 or more is considered 10 characters, and each extended source character is considered the same number of characters as the corresponding universal character name, if any)¹⁴⁾
 - 4095 external identifiers in one translation unit
 - 511 identifiers with block scope declared in one block
 - 4095 macro identifiers simultaneously defined in one preprocessing translation unit
 - 127 parameters in one function definition
 - 127 arguments in one function call
 - 127 parameters in one macro definition
 - 127 arguments in one macro invocation
 - 4095 characters in a logical source line
 - 4095 characters in a string literal (after concatenation)
 - 32767 bytes in an object (in a hosted environment only)
 - 15 nesting levels for **#included** files
 - 1023 **case** labels for a **switch** statement (excluding those for any nested **switch** statements)
 - 1023 members in a single structure or union
 - 1023 enumeration constants in a single enumeration
 - 63 levels of nested structure or union definitions in a single member declaration list

5.3.5.3 Numerical limits

5.3.5.3.1 General

- 1 An implementation is required to document all the limits specified in this subclause, which are specified in the headers `<limits.h>` and `<float.h>`. Additional limits are specified in `<stdint.h>`.
Forward references: integer types `<stdint.h>` (7.22).

¹³⁾Implementations are encouraged to avoid imposing fixed translation limits whenever possible.

¹⁴⁾See “future language directions” (6.11.3).

5.3.5.3.2 Characteristics of integer types <limits.h>

- 1 The values given subsequently shall be replaced by constant expressions suitable for use in conditional expression inclusion preprocessing directives. Their implementation-defined values shall be equal or greater to those shown.

- width for an object of type **bool**¹⁵⁾

BOOL_WIDTH	1
-------------------	---

- number of bits for smallest object that is not a bit-field (byte)

CHAR_BIT	8
-----------------	---

The macros **CHAR_WIDTH**, **SCHAR_WIDTH**, and **UCHAR_WIDTH** that represent the width of the types **char**, **signed char** and **unsigned char** shall expand to the same value as **CHAR_BIT**.

- width for an object of type **unsigned short int**

USHRT_WIDTH	16
--------------------	----

The macro **SHRT_WIDTH** represents the width of the type **short int** and shall expand to the same value as **USHRT_WIDTH**.

- width for an object of type **unsigned int**

UINT_WIDTH	16
-------------------	----

The macro **INT_WIDTH** represents the width of the type **int** and shall expand to the same value as **UINT_WIDTH**.

- width for an object of type **unsigned long int**

ULONG_WIDTH	32
--------------------	----

The macro **LONG_WIDTH** represents the width of the type **long int** and shall expand to the same value as **ULONG_WIDTH**.

- width for an object of type **unsigned long long int**

ULLONG_WIDTH	64
---------------------	----

The macro **LLONG_WIDTH** represents the width of the type **long long int** and shall expand to the same value as **ULLONG_WIDTH**.

- maximum width of a bit-precise integer type

BITINT_MAXWIDTH	<i>/* see the following */</i>
------------------------	--------------------------------

The macro **BITINT_MAXWIDTH** represents the maximum width N supported by the declaration of a bit-precise integer (6.2.5) in the type specifier **_BitInt**(N). The value **BITINT_MAXWIDTH** shall expand to a value that is greater than or equal to the value of **ULLONG_WIDTH**.

- maximum number of bytes in a multibyte character, for any supported locale

MB_LEN_MAX	1
-------------------	---

¹⁵⁾This value is exact.

- 2 For all unsigned integer types for which `<limits.h>` or `<stdint.h>` define a macro with suffix `_WIDTH` holding its width N , there is a macro with suffix `_MAX` holding the maximal value $2^N - 1$ that is representable by the type and that has the same type as would an expression that is an object of the corresponding type converted according to the integer promotions. If the value is in the range of the type `uintmax_t` (7.22.2.6) the macro is suitable for use in conditional expression inclusion preprocessing directives.
- 3 For all signed integer types for which `<limits.h>` or `<stdint.h>` define a macro with suffix `_WIDTH` holding its width N , there are macros with suffix `_MIN` and `_MAX` holding the minimal and maximal values -2^{N-1} and $2^{N-1} - 1$ that are representable by the type and that have the same type as would an expression that is an object of the corresponding type converted according to the integer promotions. If the values are in the range of the type `intmax_t` (7.22.2.6) the macros are suitable for use in conditional expression inclusion preprocessing directives.
- 4 If an object of type `char` can hold negative values, the value of `CHAR_MIN` shall be the same as that of `SCHAR_MIN` and the value of `CHAR_MAX` shall be the same as that of `SCHAR_MAX`. Otherwise, the value of `CHAR_MIN` shall be 0 and the value of `CHAR_MAX` shall be the same as that of `UCHAR_MAX` (see 6.2.5.).
Forward references: representations of types (6.2.6), conditional inclusion (6.10.2), integer types `<stdint.h>` (7.22).

5.3.5.3.3 Characteristics of floating types `<float.h>`

- 1 The characteristics of floating types are defined in terms of a model that describes a representation of floating-point numbers and allows other values. The characteristics provide information about an implementation's floating-point arithmetic.¹⁶⁾ An implementation that defines `__STDC_IEC_60559_BFP__` or `__STDC_IEC_559__` shall implement floating types and arithmetic conforming to ISO/IEC 60559 as specified in Annex F of this document. An implementation that defines `__STDC_IEC_60559_COMPLEX__` shall implement complex types and arithmetic conforming to ISO/IEC 60559 as specified in Annex G of this document.
- 2 The following parameters are used to define the model for each floating type:

s	sign (± 1)
b	base or radix of exponent representation (an integer > 1)
e	exponent (an integer between a minimum $e_{\min}$ and a maximum $e_{\max}$)
p	precision (the number of base- b digits in the significand)
f_k	nonnegative integers less than b (the significand digits)

For each floating type, the parameters b , p , $e_{\min}$, and $e_{\max}$ are fixed constants.

- 3 For each floating type, a *floating-point number* (x) is defined by the following model:

$$x = sb^e \sum_{k=1}^p f_k b^{-k}, \quad e_{\min} \leq e \leq e_{\max}$$

- 4 Model floating-point numbers x with $f_1 > 0$ are called *normalized floating-point numbers*.
- 5 Model floating-point numbers $x \neq 0$ with $f_1 = 0$ and $e = e_{\min}$ are called *subnormal floating-point numbers*.
- 6 Model floating-point numbers $x \neq 0$ with $f_1 = 0$ and $e > e_{\min}$ are called *unnormalized floating-point numbers*.
- 7 Model floating-point numbers x with all $f_k = 0$ are zeros.
- 8 Floating types shall be able to represent signed zeros or an unsigned zero and all normalized floating-point numbers. In addition, floating types may be able to contain other kinds of floating-point numbers, such as subnormal floating-point numbers and unnormalized floating-point numbers, and values that are not floating-point numbers, such as NaNs and (signed and unsigned) infinities.
- 9 NOTE 1 Some implementations have types that include finite numbers with range and/or precision that are not covered by the model.

¹⁶⁾The floating-point model is intended to clarify the description of each floating-point characteristic and does not require the floating-point arithmetic of the implementation to be identical.

- 10 A *NaN* is a value signifying Not-a-Number. A *quiet NaN* propagates through almost every arithmetic operation without raising a floating-point exception; a *signaling NaN* generally raises a floating-point exception when occurring as an arithmetic operand.
- 11 NOTE 2 ISO/IEC 60559 specifies quiet and signaling NaNs. For implementations that do not support ISO/IEC 60559, the terms quiet NaN and signaling NaN are intended to apply to values with similar behavior.
- 12 Wherever values are unsigned, any requirement in this document to get the sign shall produce an unspecified sign, and any requirement to set the sign shall be ignored, unless otherwise specified.
- 13 NOTE 3 Bit representations of floating-point values can include a sign bit, even if the values can be regarded as unsigned; ISO/IEC 60559 NaNs are such values.
- 14 Whether and in what cases subnormal numbers are treated as zeros is implementation-defined. Subnormal numbers that in some cases are treated by arithmetic operations as zeros are properly classified as subnormal. However, object representations that could represent subnormal numbers but that are always treated by arithmetic operations as zeros are non-canonical zeros, and the values are properly classified as zero, not subnormal. ISO/IEC 60559 arithmetic (with default exception handling) always treats subnormal numbers as nonzero.
- 15 A value is negative if and only if it compares less than 0. Thus, negative zeros and NaNs are not negative values.
- 16 An implementation can prefer particular representations of values that have multiple representations in a floating type, 6.2.6.1 notwithstanding. The preferred representations of a floating type, including unique representations of values in the type, are called *canonical*. A floating type can also contain *non-canonical* representations, for example, redundant representations of some or all its values, or representations that are extraneous to the floating-point model. Typically, floating-point operations deliver results with canonical representations. ISO/IEC 60559 operations deliver results with canonical representations, unless specified otherwise.
- 17 NOTE 4 The library operations **iscanonical** and **canonicalize** distinguish canonical (preferred) representations, but this distinction alone does not imply that canonical and non-canonical representations are of different values.
- 18 NOTE 5 Some of the values in the ISO/IEC 60559 decimal formats have non-canonical representations (as well as a canonical representation).
- 19 The minimum range of representable values for a floating type is the most negative finite floating-point number representable in that type through the most positive finite floating-point number representable in that type. In addition, if negative infinity is representable in a type, the range of that type is extended to all negative real numbers; likewise, if positive infinity is representable in a type, the range of that type is extended to all positive real numbers.
- 20 The accuracy of the floating-point operations (+, -, *, /) and of most of the library functions in `<math.h>` and `<complex.h>` that return floating-point results is implementation-defined, as is the accuracy of the conversion between floating-point internal representations and string representations performed by the library functions in `<stdio.h>`, `<stdlib.h>`, and `<wchar.h>`. The implementation may state that the accuracy is unknown. Decimal floating-point operations have stricter requirements.
- 21 All integer values in the `<float.h>` header, except **FLT_ROUNDS**, shall be constant expressions suitable for use in conditional expression inclusion preprocessing directives; all floating values shall be arithmetic constant expressions. All except **CR_DECIMAL_DIG** (F.5), **DECIMAL_DIG**, **DEC_EVAL_METHOD**, **FLT_EVAL_METHOD**, **FLT_RADIX**, and **FLT_ROUNDS** have separate names for all floating types. The floating-point model representation is provided for all values except **DEC_EVAL_METHOD**, **FLT_EVAL_METHOD** and **FLT_ROUNDS**.
- 22 The remainder of this subclause specifies characteristics of standard floating types.
- 23 The rounding mode for floating-point addition for standard floating types is characterized by the implementation-defined value of **FLT_ROUNDS**. Evaluation of **FLT_ROUNDS** correctly reflects any execution-time change of rounding mode through the function **fesetround** in `<fenv.h>`.

- 1 indeterminate
- 0 toward zero
- 1 to nearest, ties to even
- 2 toward positive infinity
- 3 toward negative infinity
- 4 to nearest, ties away from zero

All other values for **FLT_ROUNDS** characterize implementation-defined rounding behavior.

- 24 Whether a type has the same base (b), precision (p), and exponent range ($e_{\min} \leq e \leq e_{\max}$) as an ISO/IEC 60559 format is characterized by the implementation-defined values of **FLT_IS_IEC_60559**, **DBL_IS_IEC_60559** and **LDBL_IS_IEC_60559** (this does not imply conformance to Annex F):

- 0 type does not have the precision and exponent range of an ISO/IEC 60559 format
- 1 type has the precision and exponent range of an ISO/IEC 60559 format

- 25 NOTE 6 Outside of the normalized floating-point numbers, the representability of values (e.g. negative zero) of the ISO/IEC 60559 format is not implied.

- 26 The values of floating type yielded by operators subject to the usual arithmetic conversions, including the values yielded by the implicit conversion of operands, and the values of floating literals are evaluated to a format whose range and precision can be greater than required by the type. Such a format is called an *evaluation format*. In all cases, assignment and cast operators yield values in the format of the type. The extent to which evaluation formats are used is characterized by the value of **FLT_EVAL_METHOD**:¹⁷⁾

- 1 indeterminate;
- 0 evaluate all operations and literals just to the range and precision of the type;
- 1 evaluate operations and literals of type **float** and **double** to the range and precision of the **double** type, evaluate **long double** operations and literals to the range and precision of the **long double** type;
- 2 evaluate all operations and literals to the range and precision of the **long double** type.

All other negative values for **FLT_EVAL_METHOD** characterize implementation-defined behavior. The value of **FLT_EVAL_METHOD** does not characterize values returned by function calls (see 6.8.7.5, F.6).

- 27 The presence or absence of subnormal numbers is characterized by the implementation-defined values of **FLT_HAS_SUBNORM**, **DBL_HAS_SUBNORM**, and **LDBL_HAS_SUBNORM**:

- 1 indeterminate
- 0 absent (type does not support subnormal numbers)
- 1 present (type does support subnormal numbers)

The use of **FLT_HAS_SUBNORM**, **DBL_HAS_SUBNORM**, and **LDBL_HAS_SUBNORM** macros is an obsolescent feature.

- 28 Each of the signaling NaN macros

¹⁷⁾The evaluation method determines evaluation formats of expressions involving all floating types, not just real types. For example, if **FLT_EVAL_METHOD** is 1, then the product of two **float _Complex** operands is represented in the **double _Complex** format, and its parts are evaluated to **double**.

FLT_SNAN
DBL_SNAN
LDBL_SNAN

is defined if and only if the respective type contains signaling NaNs. They expand to a constant expression of the respective type representing a signaling NaN. If a signaling NaN macro, optionally preceded by the unary + or - operator, is used as an initializer that is evaluated at translation time to initialize an object of the same type, the object is initialized with a signaling NaN value.

29 The macro

INFINITY

is defined if and only if the implementation supports an infinity for the type **float**. It expands to a constant expression of type **float** representing positive or unsigned infinity.

30 The macro

NAN

is defined if and only if the implementation supports quiet NaNs for the **float** type. It expands to a constant expression of type **float** representing a quiet NaN.

31 The values given in the following list shall be replaced by constant expressions with implementation-defined values that are greater or equal in magnitude (absolute value) to those shown, with the same sign:

- radix of exponent representation, b

FLT_RADIX	2
------------------	---

- number of base-**FLT_RADIX** digits in the floating-point significand, p

FLT_MANT_DIG
DBL_MANT_DIG
LDBL_MANT_DIG

- number of decimal digits, n , such that any floating-point number with p radix b digits can be rounded to a floating-point number with n decimal digits and back again, using to-nearest rounding for both roundings, without change to the value,

$$\begin{cases} p \log_{10} b & \text{if } b \text{ is a power of } 10 \\ \lceil 1 + p \log_{10} b \rceil & \text{otherwise} \end{cases}$$

FLT_DECIMAL_DIG	6
DBL_DECIMAL_DIG	10
LDBL_DECIMAL_DIG	10

- number of decimal digits, n , such that any floating-point number in the widest of the supported floating types and the supported ISO/IEC 60559 encodings with $p_{\max}$ radix b digits can be rounded to a floating-point number with n decimal digits and back again, using to-nearest rounding for both roundings, without change to the value,

$$\begin{cases} p_{\max} \log_{10} b & \text{if } b \text{ is a power of } 10 \\ \lceil 1 + p_{\max} \log_{10} b \rceil & \text{otherwise} \end{cases}$$

DECIMAL_DIG	10
--------------------	----

This is an obsolescent feature, see 7.34.9.

- number of decimal digits, q , such that any floating-point number with q decimal digits can be rounded into a floating-point number with p radix b digits and back again, using to-nearest rounding for both roundings, without change to the q decimal digits,

$$\begin{cases} p \log_{10} b & \text{if } b \text{ is a power of } 10 \\ \lfloor (p-1) \log_{10} b \rfloor & \text{otherwise} \end{cases}$$

FLT_DIG	6
DBL_DIG	10
LDBL_DIG	10

- minimum negative integer such that **FLT_RADIX** raised to one less than that power is a normalized floating-point number, $e_{\min}$

FLT_MIN_EXP
DBL_MIN_EXP
LDBL_MIN_EXP

- minimum negative integer such that 10 raised to that power is in the range of normalized floating-point numbers, $\lceil \log_{10} b^{e_{\min}} - 1 \rceil$

FLT_MIN_10_EXP	-37
DBL_MIN_10_EXP	-37
LDBL_MIN_10_EXP	-37

- maximum integer such that **FLT_RADIX** raised to one less than that power is a representable finite floating-point number; if that representable finite floating-point number is normalized, the value of the macro is $e_{\max}$

FLT_MAX_EXP
DBL_MAX_EXP
LDBL_MAX_EXP

- maximum integer such that 10 raised to that power is in the range of representable finite floating-point numbers, $\lfloor \log_{10}((1 - b^{-p})b^{e_{\max}}) \rfloor$

FLT_MAX_10_EXP	+37
DBL_MAX_10_EXP	+37
LDBL_MAX_10_EXP	+37

- 32 The values given in the following list shall be replaced by constant expressions with implementation-defined values that are greater than or equal to those shown:

- maximum representable finite floating-point number; if that number is normalized, its value is $(1 - b^{-p})b^{e_{\max}}$

FLT_MAX	1E+37
DBL_MAX	1E+37
LDBL_MAX	1E+37

- maximum normalized floating-point number, $(1 - b^{-p})b^{e_{\max}}$

FLT_NORM_MAX	1E+37
DBL_NORM_MAX	1E+37
LDBL_NORM_MAX	1E+37

- 33 The values given in the following list shall be replaced by constant expressions with implementation-defined (positive) values that are less than or equal to those shown:

- the difference between 1 and the least normalized value greater than 1 that is representable in the given floating type, b^{1-p}

FLT_EPSILON	1E-5
DBL_EPSILON	1E-9
LDBL_EPSILON	1E-9

- minimum normalized positive floating-point number, $b^{e_{\min}-1}$

FLT_MIN	1E-37
DBL_MIN	1E-37
LDBL_MIN	1E-37

- minimum positive floating-point number

FLT_TRUE_MIN	1E-37
DBL_TRUE_MIN	1E-37
LDBL_TRUE_MIN	1E-37

Recommended practice

- 34 Conversion between real floating type and decimal character sequence with at most $T_DECIMAL_DIG$ digits should be correctly rounded, where T is the macro prefix for the type. This assures conversion from real floating type to decimal character sequence with $T_DECIMAL_DIG$ digits and back, using to-nearest rounding, is the identity function.
- 35 EXAMPLE 1 The following describes an artificial floating-point representation that meets the minimum requirements of this document, and the appropriate values in a `<float.h>` header for type **float**:

$$x = s16^e \sum_{k=1}^6 f_k 16^{-k}, \quad -31 \leq e \leq +32$$

FLT_RADIX	16
FLT_MANT_DIG	6
FLT_EPSILON	9.53674316E-07F
FLT_DECIMAL_DIG	9
FLT_DIG	6
FLT_MIN_EXP	-31
FLT_MIN	2.93873588E-39F
FLT_MIN_10_EXP	-38
FLT_MAX_EXP	+32
FLT_MAX	3.40282347E+38F
FLT_MAX_10_EXP	+38

- 36 EXAMPLE 2 The following describes floating-point representations that also meet the requirements for binary32 and binary64 numbers in ISO/IEC 60559,¹⁸⁾ and the appropriate values in a `<float.h>` header for types **float** and **double**. The decimal floating literals can possibly not give correct values (and hence are not appropriate values in a `<float.h>` header) if **FLT_EVAL_METHOD** is not 0 or if a translation-time rounding mode other than

¹⁸⁾The floating-point model in ISO/IEC 60559 sums powers of b from zero, so the values of the exponent limits are one less than shown here.

the ISO/IEC 60559 default is supported (either as the default or as a constant rounding mode set by an **FENV_ROUND** pragma). The hexadecimal floating literals are correct in all such cases because their values are exactly representable in the type.

$$x_f = s2^e \sum_{k=1}^{24} f_k 2^{-k}, \quad -125 \leq e \leq +128$$

$$x_d = s2^e \sum_{k=1}^{53} f_k 2^{-k}, \quad -1021 \leq e \leq +1024$$

FLT_IS_IEC_60559	1
FLT_RADIX	2
FLT_MANT_DIG	24
FLT_EPSILON	1.19209290E-07F // decimal floating literal
FLT_EPSILON	0X1P-23F // hex floating literal
FLT_DECIMAL_DIG	9
FLT_DIG	6
FLT_MIN_EXP	-125
FLT_MIN	1.17549435E-38F // decimal floating literal
FLT_MIN	0X1P-126F // hex floating literal
FLT_TRUE_MIN	1.40129846E-45F // decimal floating literal
FLT_TRUE_MIN	0X1P-149F // hex floating literal
FLT_HAS_SUBNORM	1
FLT_MIN_10_EXP	-37
FLT_MAX_EXP	+128
FLT_MAX	3.40282347E+38F // decimal floating literal
FLT_MAX	0X1.fffffeP127F // hex floating literal
FLT_MAX_10_EXP	+38
DBL_MANT_DIG	53
DBL_IS_IEC_60559	1
DBL_EPSILON	2.2204460492503131E-16 // decimal floating literal
DBL_EPSILON	0X1P-52 // hex floating literal
DBL_DECIMAL_DIG	17
DBL_DIG	15
DBL_MIN_EXP	-1021
DBL_MIN	2.2250738585072014E-308 // decimal floating literal
DBL_MIN	0X1P-1022 // hex floating literal
DBL_TRUE_MIN	4.9406564584124654E-324 // decimal floating literal
DBL_TRUE_MIN	0X1P-1074 // hex floating literal
DBL_HAS_SUBNORM	1
DBL_MIN_10_EXP	-307
DBL_MAX_EXP	+1024
DBL_MAX	1.7976931348623157E+308 // decimal floating literal
DBL_MAX	0X1.fffffffffffffP1023 // hex floating literal
DBL_MAX_10_EXP	+308

Forward references: conditional inclusion (6.10.2), predefined macro names (6.10.10), complex arithmetic <complex.h> (7.3), extended multibyte and wide character utilities <wchar.h> (7.32), floating-point environment <fenv.h> (7.6), general utilities <stdlib.h> (7.24), input/output <stdio.h> (7.23), mathematics <math.h> (7.12), ISO/IEC 60559 floating-point arithmetic (Annex F), ISO/IEC 60559-compatible complex arithmetic (Annex G).

5.3.5.3.4 Characteristics of decimal floating types in <float.h>

- 1 This subclause specifies macros in <float.h> that provide characteristics of decimal floating types (an optional feature) in terms of the model presented in 5.3.5.3.3. An implementation shall provide these macros if and only if it defines **__STDC_IEC_60559_DFP__**. The prefixes **DEC32_**, **DEC64_**, and **DEC128_** denote the types **_Decimal32**, **_Decimal64**, and **_Decimal128** respectively.
- 2 **DEC_EVAL_METHOD** is the decimal floating-point analog of **FLT_EVAL_METHOD** (5.3.5.3.3). Its implementation-defined value characterizes the use of evaluation formats for decimal floating

types:

- 1 indeterminate;
 - 0 evaluate all operations and literals just to the range and precision of the type;
 - 1 evaluate operations and literals of type **_Decimal32** and **_Decimal64** to the range and precision of the **_Decimal64** type, evaluate **_Decimal128** operations and literals to the range and precision of the **_Decimal128** type;
 - 2 evaluate all operations and literals to the range and precision of the **_Decimal128** type.
- 3 Each of the decimal signaling NaN macros

DEC32_SNAN
DEC64_SNAN
DEC128_SNAN

expands to a constant expression of the respective decimal floating type representing a signaling NaN. If an optional unary + or - operator followed by a signaling NaN macro is used for initializing an object of the same type that has static or thread storage duration, the object is initialized with a signaling NaN value.

- 4 The macro

DEC_INFINITY

expands to a constant expression of type **_Decimal32** representing positive infinity.

- 5 The macro

DEC_NAN

expands to a constant expression of type **_Decimal32** representing a quiet NaN.

- 6 The integer values given in the following lists shall be replaced by constant expressions suitable for use in conditional expression inclusion preprocessing directives:

- radix of exponent representation, $b(=10)$

For the standard floating types, this value is implementation-defined and is specified by the macro **FLT_RADIX**. For the decimal floating types there is no corresponding macro, since the value 10 is an inherent property of the types. Wherever **FLT_RADIX** appears in a description of a function that has versions that operate on decimal floating types, it is noted that for the decimal floating-point versions the value used is implicitly 10, rather than **FLT_RADIX**.

- number of digits in the coefficient

DEC32_MANT_DIG	7
DEC64_MANT_DIG	16
DEC128_MANT_DIG	34

- minimum exponent

DEC32_MIN_EXP	-94
DEC64_MIN_EXP	-382
DEC128_MIN_EXP	-6142

- maximum exponent

DEC32_MAX_EXP	97
DEC64_MAX_EXP	385
DEC128_MAX_EXP	6145

- maximum representable finite decimal floating-point number (there are 6, 15 and 33 9's after the decimal points respectively)

[illegible]

- the difference between 1 and the least value greater than 1 that is representable in the given floating type

DEC32_EPSILON	1E-6DF
DEC64_EPSILON	1E-15DD
DEC128_EPSILON	1E-33DL

- minimum normalized positive decimal floating-point number

DEC32_MIN	1E-95DF
DEC64_MIN	1E-383DD
DEC128_MIN	1E-6143DL

- minimum positive subnormal decimal floating-point number

DEC32_TRUE_MIN	0.000001E-95DF
DEC64_TRUE_MIN	0.0000000000000001E-383DD
DEC128_TRUE_MIN	0.0000000000000000000000000000000000000001E-6143DL

- 7 For decimal floating-point arithmetic, it is often convenient to consider an alternate equivalent model where the significand is represented with integer rather than fraction digits. With s, b, e, p , and f_k as defined in 5.3.5.3.3, a floating-point number x is defined by the model:

$$x = s \cdot b^{(e-p)} \sum_{k=1}^p f_k \cdot b^{(p-k)}$$

- 8 With b fixed to 10, a decimal floating-point number x is thus:

$$x = s \cdot 10^{(e-p)} \sum_{k=1}^p f_k \cdot 10^{(p-k)}$$

The *quantum exponent* is $q = e - p$ and the *coefficient* is $c = f_1 f_2 \cdots f_p$, which is an integer between 0 and $10^p - 1$, inclusive. Thus, $x = s \cdot c \cdot 10^q$ is represented by the triple of integers (s, c, q) . The *quantum* of x is 10^q , which is the value of a unit in the last place of the coefficient. Table 5.1 shows the range of quantum exponents.

Table 5.1 — Quantum exponent ranges

Type	<code>_Decimal32</code>	<code>_Decimal64</code>	<code>_Decimal128</code>
Maximum Quantum Exponent ($q_{\max}$)	90	369	6111
Minimum Quantum Exponent ($q_{\min}$)	-101	-398	-6176

- 9 For binary floating-point arithmetic following ISO/IEC 60559, representations in the model described in 5.3.5.3.3 that have the same numerical value are indistinguishable in the arithmetic. However, for decimal floating-point arithmetic, representations that have the same numerical value but different quantum exponents, e.g. $(+1, 10, -1)$ representing 1.0 and $(+1, 100, -2)$ representing 1.00, are distinguishable. To facilitate exact fixed-point calculation, operation results that are of decimal floating type have a *preferred quantum exponent*, as specified in ISO/IEC 60559, which is determined by the quantum exponents of the operands if they have decimal floating types (or by specific rules for conversions from other types). Table 5.2 gives rules for determining preferred quantum exponents for results of ISO/IEC 60559 operations, and for other operations specified in this document. When exact, these operations produce a result with their preferred quantum exponent, or as close to it as possible within the limitations of the type. When inexact, these operations produce a result with the least possible quantum exponent. For example, the preferred quantum exponent for addition is the minimum of the quantum exponents of the operands. Hence $(+1, 123, -2) + (+1, 4000, -3) = (+1, 5230, -3)$ or $1.23 + 4.000 = 5.230$.
- 10 Table 5.2 shows, for each operation delivering a result in decimal floating-point format, how the preferred quantum exponents of the operands, $Q(x)$, $Q(y)$, etc., determine the preferred quantum exponent of the operation result, provided the table formula is defined for the arguments. For the cases where the formula is undefined and the function result is $\pm\infty$, the preferred quantum exponent is immaterial because the quantum exponent of $\pm\infty$ is defined to be infinity. For the other cases where the formula is undefined and the function result is finite, the preferred quantum exponent is unspecified.
- 11 NOTE Although unspecified in ISO/IEC 60559, a preferred quantum exponent of 0 for the cases where the formula is undefined and the function result is finite would be a reasonable implementation choice.

Table 5.2 — Preferred quantum exponents

Operation	Preferred quantum exponent of result
roundeven, round, trunc, ceil, floor, rint, nearbyint	$\max(Q(x), 0)$
nextup, nextdown, nextafter, nexttoward	least possible
remainder	$\min(Q(x), Q(y))$
fmin, fmax, fminimum, fmaximum, fminimum_mag, fmaximum_mag, fminimum_num, fmaximum_num, fminimum_mag_num, fmaximum_mag_num	$Q(x)$ if x gives the result, $Q(y)$ if y gives the result
scalbn, scalbln	$Q(x) + n$
ldexp	$Q(x) + p$
logb	0
postfix ++ operator, postfix -- operator, prefix ++ operator, prefix -- operator	$\min(Q(x), 0)$
+, d32add, d64add	$\min(Q(x), Q(y))$
-, d32sub, d64sub	$\min(Q(x), Q(y))$
*, d32mul, d64mul	$Q(x) + Q(y)$
/, d32div, d64div	$Q(x) - Q(y)$
sqrt, d32sqrt, d64sqrt	$\lfloor Q(x)/2 \rfloor$

fma, d32fma, d64fma	$\min(Q(x) + Q(y), Q(z))$
conversion from integer type	0
exact conversion from non-decimal floating type	0
inexact conversion from non-decimal floating type	least possible
conversion between decimal floating types	$Q(x)$
*cx returned by canonicalize	$Q(*x)$
strt0, wcsto, scanf , floating literals of decimal floating type	see 7.24.2.7
$-(x), +(x)$	$Q(x)$
fabs	$Q(x)$
copysign	$Q(x)$
quantize	$Q(y)$
quantum	$Q(x)$
*encptr returned by encodedec, encodebin	$Q(*xptr)$
*xptr returned by decodedec, decodebin	$Q(*encptr)$
fmod	$\min(Q(x), Q(y))$
fdim	$\min((Q(x), Q(y)) \text{ if } x > y, 0 \text{ if } x \leq y)$
cbrt	$\lfloor Q(x)/3 \rfloor$
hypot	$\min(Q(x), Q(y))$
pow	$\lfloor y \times Q(x) \rfloor$
modf	$Q(\text{value})$
*iptr returned by modf	$\max(Q(\text{value}), 0)$
frexp	$Q(\text{value})$ if $\text{value} = 0$, $-(\text{length of coefficient of value})$ otherwise
*res returned by setpayload, setpayloadsig	0 if pl does not represent a valid payload, not applicable otherwise (NaN returned)
getpayload	0 if *x is a NaN, unspecified otherwise
compoundn	$\lfloor n \times \min(0, Q(x)) \rfloor$
pown	$\lfloor n \times Q(x) \rfloor$
powr	$\lfloor y \times Q(x) \rfloor$
rootn	$\lfloor Q(x)/n \rfloor$
rsqrt	$-\lfloor Q(x)/2 \rfloor$
transcendental functions	0

A function family listed in Table 5.2 indicates the functions for all decimal floating types, where the

function family is represented by the name of the functions without a suffix. For example, **ceil** indicates the functions **ceild32**, **ceild64**, and **ceild128**.

Forward references: extended multibyte and wide character utilities <wchar.h> (7.32), floating-point environment <fenv.h> (7.6), general utilities <stdlib.h> (7.24), input/output <stdio.h> (7.23), mathematics <math.h> (7.12), type-generic mathematics <tgmath.h> (7.28), ISO/IEC 60559 floating-point arithmetic (Annex F).

6. Language

6.1 Notation

- 1 In the syntax notation used in this clause, syntactic categories (nonterminals) are indicated by *italic type*, and literal words and character set members (terminals) by **bold type**. A colon (:) following a nonterminal introduces its definition. Alternative definitions are listed on separate lines, except when prefaced by the words “one of”. An optional symbol is indicated by the subscript “opt”, so that

{ *expression*_{opt} }

indicates an optional expression enclosed in braces.

- 2 When syntactic categories are referred to in the main text, they are not italicized and words are separated by spaces instead of hyphens.
- 3 A summary of the language syntax is given in Annex A.

6.2 Concepts

6.2.1 Scopes of identifiers, type names, and compound literals

- 1 An identifier can denote:

- a standard attribute, an attribute prefix, or an attribute name;
- an object;
- a function;
- a tag or a member of a structure, union, or enumeration;
- a typedef name;
- a label name;
- a macro name;
- or, a macro parameter.

The same identifier can denote different entities at different points in the program. A member of an enumeration is called an *enumeration constant*. Macro names and macro parameters are not considered further here, because prior to the semantic phase of program translation any occurrences of macro names in the source file are replaced by the preprocessing token sequences that constitute their macro definitions.

- 2 For each different entity that an identifier designates, the identifier is *visible* (i.e. can be used) only within a region of program text called its *scope*. Different entities designated by the same identifier either have different scopes or are in different name spaces. There are four kinds of scopes: function, file, block, and function prototype. (A *function prototype* is a declaration of a function.)
- 3 A label name is the only kind of identifier that has *function scope*. It can be used (in a **goto** statement) anywhere in the function in which it appears, and is declared implicitly by its syntactic appearance (followed by a : and a statement).
- 4 Every other identifier has scope determined by the placement of its declaration (in a declarator or type specifier). If the declarator or type specifier that declares the identifier appears outside of any block or list of parameters, the identifier has *file scope*, which terminates at the end of the translation unit. If the declarator or type specifier that declares the identifier appears inside a block or within the list of parameter declarations in a function definition, the identifier has *block scope*, which terminates at the end of the associated block. If the declarator or type specifier that declares

the identifier appears within the list of parameter declarations in a function prototype (not part of a function definition), the identifier has *function prototype scope*, which terminates at the end of the function declarator. If an identifier designates two different entities in the same name space, the scopes can overlap. If so, the scope of one entity (the *inner scope*) will end strictly before the scope of the other entity (the *outer scope*). Within the inner scope, the identifier designates the entity declared in the inner scope; the entity declared in the outer scope is *hidden* (and not visible) within the inner scope.

- 5 Unless explicitly stated otherwise, where this document uses the term “identifier” to refer to an entity (as opposed to the syntactic construct), it refers to the entity in the relevant name space whose declaration is visible at the point the identifier occurs.
- 6 Two identifiers have the *same scope* if and only if their scopes terminate at the same point.
- 7 Structure, union, and enumeration tags have scope that begins just after the appearance of the tag in a type specifier that declares the tag. Each enumeration constant has scope that begins just after the appearance of its defining enumerator in an enumerator list. An ordinary identifier that has an underspecified definition has scope that starts when the definition is completed; if the same ordinary identifier declares another entity with a scope that encloses the current block, that declaration is hidden as soon as the inner declarator is completed.¹⁹⁾ Any other identifier has scope that begins just after the completion of its declarator.
- 8 As a special case, a type name (which is not a declaration of an identifier) is considered to have a scope that begins just after the place within the type name where the omitted identifier would appear were it not omitted. A compound literal (which is an expression that provides access to an anonymous object) is associated with the scope of the type name used in its definition; that scope is either file scope, function prototype scope, or block scope.

Forward references: declarations (6.7), function calls (6.5.3.3), function calls (6.5.3.6), function definitions (6.9.2), identifiers (6.4.3), macro replacement (6.10.5), name spaces of identifiers (6.2.3), source file inclusion (6.10.3), statements and blocks (6.8).

6.2.2 Linkages of identifiers

- 1 An identifier declared in different scopes or in the same scope more than once can be made to refer to the same object or function by a process called *linkage*.²⁰⁾ There are three kinds of linkage: external, internal, and none.
- 2 In the set of translation units and libraries that constitutes an entire program, each declaration of a particular identifier with *external linkage* denotes the same object or function. Within one translation unit, each declaration of an identifier with *internal linkage* denotes the same object or function. Each declaration of an identifier with *no linkage* denotes a unique entity.
- 3 If the declaration of a file scope identifier for:
 - an object contains any of the storage-class specifiers **static** or **constexpr**;
 - or, a function contains the storage-class specifier **static**,

then the identifier has internal linkage.²¹⁾

- 4 For an identifier declared with the storage-class specifier **extern** in a scope in which a prior declaration of that identifier is visible,²²⁾ if the prior declaration specifies internal or external linkage, the linkage of the identifier at the later declaration is the same as the linkage specified at the prior declaration. If no prior declaration is visible, or if the prior declaration specifies no linkage, then the identifier has external linkage.
- 5 If the declaration of an identifier for a function has no storage-class specifier, its linkage is determined exactly as if it were declared with the storage-class specifier **extern**. If the declaration of an identifier

¹⁹⁾That means, that the outer declaration is not visible for the initializer.

²⁰⁾There is no linkage between different identifiers.

²¹⁾A function declaration can contain the storage-class specifier **static** only if it is at file scope; see 6.7.2.

²²⁾As specified in 6.2.1, the later declaration can hide the prior declaration.

for an object has file scope and does not contain the storage-class specifier **static** or **constexpr**, its linkage is external.

- 6 The following identifiers have no linkage: an identifier declared to be anything other than an object or a function; an identifier declared to be a function parameter; a block scope identifier for an object declared without the storage-class specifier **extern**.
- 7 If, within a translation unit, the same identifier appears with both internal and external linkage, the behavior is undefined.

Forward references: declarations (6.7), expressions (6.5.1), external definitions (6.9), statements (6.8).

6.2.3 Name spaces of identifiers

- 1 If more than one declaration of a particular identifier is visible at any point in a translation unit, the syntactic context disambiguates uses that refer to different entities. Thus, there are separate *name spaces* for various categories of identifiers, as follows:

- *label names* (disambiguated by the syntax of the label declaration and use);
- the *tags* of structures, unions, and enumerations (disambiguated by following any²³⁾ of the keywords **struct**, **union**, or **enum**);
- the *members* of structures or unions; each structure or union has a separate name space for its members (disambiguated by the type of the expression used to access the member via the `.` or `->` operator);
- standard attributes and attribute prefixes (disambiguated by the syntax of the attribute specifier and name of the attribute token) (6.7.13);
- the trailing identifier in an attribute prefixed token; each attribute prefix has a separate name space for the implementation-defined attributes that it introduces (disambiguated by the attribute prefix and the trailing identifier token);
- all other identifiers, called *ordinary identifiers* (declared in ordinary declarators or as enumeration constants).

Forward references: enumeration specifiers (6.7.3.3), labeled statements (6.8.2), structure and union specifiers (6.7.3.2), structure and union members (6.5.3.4), tags (6.7.3.4), the **goto** statement (6.8.7.2).

6.2.4 Storage durations of objects

- 1 An object has a *storage duration* that determines its lifetime. There are four storage durations: static, thread, automatic, and allocated. Allocated storage is described in 7.24.4.
- 2 The *lifetime* of an object is the portion of program execution during which storage is guaranteed to be reserved for it. An object exists, has a constant address,²⁴⁾ and retains its last-stored value throughout its lifetime.²⁵⁾ If an object is referred to outside of its lifetime, the behavior is undefined. If a pointer value is used in an evaluation after the object the pointer points to (or just past) reaches the end of its lifetime, the behavior is undefined. The representation of a pointer object becomes indeterminate when the object the pointer points to (or just past) reaches the end of its lifetime.
- 3 An object whose identifier is declared without the storage-class specifier **thread_local**, and either with external or internal linkage or with the storage-class specifier **static**, has *static storage duration*. Its lifetime is the entire execution of the program and its stored value is initialized only once, prior to program startup.
- 4 An object whose identifier is declared with the storage-class specifier **thread_local** has *thread storage duration*. Its explicit or implicit initializer is evaluated prior to program execution, its lifetime

²³⁾There is only one name space for tags even though three are possible.

²⁴⁾The term “constant address” means that two pointers to the object constructed at possibly different times will compare equal. The address can be different during two different executions of the same program.

²⁵⁾In the case of a volatile object, the last store is not required to be explicit in the program.

is the entire execution of the thread for which it is created, and its stored value is initialized with the previously determined value when the thread is started. There is a distinct object per thread, and use of the declared name in an expression refers to the object associated with the thread evaluating the expression. The result of attempting to indirectly access an object with thread storage duration from a thread other than the one with which the object is associated is implementation-defined.

- 5 An object whose identifier is declared with no linkage and without the storage-class specifier **static** has *automatic storage duration*, as do some compound literals. The result of attempting to indirectly access an object with automatic storage duration from a thread other than the one with which the object is associated is implementation-defined.
- 6 For such an object that does not have a variable length array type, its lifetime extends from entry into the block with which it is associated until execution of that block ends in any way. (Entering an enclosed block or calling a function suspends, but does not end, execution of the current block.) If the block is entered recursively, a new instance of the object is created each time. The initial representation of the object is indeterminate. If an initialization is specified for the object and it is not specified with **constexpr**, it is performed each time the declaration or compound literal is reached in the execution of the block; if it is specified with **constexpr** the initializer is evaluated once at translation time and the new instance of the object is initialized to that fixed value each time the specification is reached; otherwise, the representation of the object becomes indeterminate each time the declaration is reached.
- 7 For such an object that does have a variable length array type, its lifetime extends from the declaration of the object until execution of the program leaves the scope of the declaration.²⁶⁾ If the scope is entered recursively, a new instance of the object is created each time. The initial representation of the object is indeterminate.
- 8 A non-lvalue expression with structure or union type, where the structure or union contains a member with array type (including, recursively, members of all contained structures and unions) refers to an object with automatic storage duration and *temporary lifetime*.²⁷⁾ Its lifetime begins when the expression is evaluated and its initial value is the value of the expression. Its lifetime ends when the evaluation of the containing full expression ends. Any attempt to modify an object with temporary lifetime results in undefined behavior. An object with temporary lifetime behaves as if it were declared with the type of its value for the purposes of effective type. Such an object may not have a unique address.

Forward references: array declarators (6.7.7.3), compound literals (6.5.3.6), declarators (6.7.7), function calls (6.5.3.3), initialization (6.7.11), statements (6.8), effective type (6.5.1).

6.2.5 Types

- 1 The meaning of a value stored in an object or returned by a function is determined by the *type* of the expression used to access it. (An identifier declared to be an object is the simplest such expression; the type is specified in the declaration of the identifier.) Types are partitioned into *object types* (types that describe objects) and *function types* (types that describe functions). At various points within a translation unit an object type can be *incomplete*²⁸⁾ (lacking sufficient information to determine the size of objects of that type) or *complete* (having sufficient information).²⁹⁾
- 2 An object declared as type **bool** is large enough to store the values **false** and **true**.
- 3 An object declared as type **char** is large enough to store any member of the basic execution character set. If a member of the basic execution character set is stored in a **char** object, its value is guaranteed to be nonnegative. If any other character is stored in a **char** object, the resulting value is implementation-defined but shall be within the range of values that can be represented in that type.

²⁶⁾ Leaving the innermost block containing the declaration, or jumping to a point in that block or an embedded block prior to the declaration, leaves the scope of the declaration.

²⁷⁾ The address of such an object is taken implicitly when an array member is accessed.

²⁸⁾ An incomplete type can only be used when the size of an object of that type is not needed. It is not needed, for example, when a typedef name is declared to be a specifier for a structure or union, or when a pointer to or a function returning a structure or union is being declared. The specification has to be complete before such a function is called or defined.

²⁹⁾ A type can be incomplete or complete throughout an entire translation unit, or it can change states at different points within a translation unit.

- 4 There are five *standard signed integer types*, designated as **signed char**, **short int**, **int**, **long int**, and **long long int**. (These and other types can be designated in several additional ways, as described in 6.7.3.)
- 5 A *bit-precise signed integer type* is designated as **_BitInt**(*N*) where *N* is an integer constant expression that specifies the number of bits that are used to represent the type, including the sign bit. Each value of *N* designates a distinct type.³⁰⁾
- 6 There may also be implementation-defined *extended signed integer types*.³¹⁾ The standard signed integer types, bit-precise signed integer types, and extended signed integer types are collectively called *signed integer types*.³²⁾
- 7 An object declared as type **signed char** occupies the same amount of storage as a “plain” **char** object. A “plain” **int** object has the natural size suggested by the architecture of the execution environment (large enough to contain any value in the range **INT_MIN** to **INT_MAX** as defined in the header `<limits.h>`).
- 8 For each of the signed integer types, there is a *corresponding* (but different) *unsigned integer type* (designated with the keyword **unsigned**) that uses the same amount of storage (including sign information) and has the same alignment requirements. The type **bool** and the unsigned integer types that correspond to the standard signed integer types are the *standard unsigned integer types*. The unsigned integer types that correspond to the extended signed integer types are the *extended unsigned integer types*. In addition to the unsigned integer types that correspond to the bit-precise signed integer types there is the type **unsigned _BitInt**(1), which uses one bit to represent the type. Collectively, **unsigned _BitInt**(1) and the unsigned integer types that correspond to the bit-precise signed integer types are the *bit-precise unsigned integer types*. The standard unsigned integer types, bit-precise unsigned integer types, and extended unsigned integer types are collectively called *unsigned integer types*.³³⁾
- 9 The standard signed integer types and standard unsigned integer types are collectively called the *standard integer types*; the bit-precise signed integer types and bit-precise unsigned integer types are collectively called the *bit-precise integer types*; the extended signed integer types and extended unsigned integer types are collectively called the *extended integer types*.
- 10 For any two integer types with the same signedness and different integer conversion rank (see 6.3.2.1), the range of values of the type with smaller integer conversion rank is a subrange of the values of the other type.
- 11 The range of nonnegative values of a signed integer type is a subrange of the corresponding unsigned integer type, and the representation of the same value in each type is the same.³⁴⁾ The range of representable values for the unsigned type is 0 to $2^N - 1$ (inclusive). A computation involving unsigned operands can never produce an overflow, because arithmetic for the unsigned type is performed modulo 2^N .
- 12 There are three *standard floating types*, designated as **float**, **double**, and **long double**.³⁵⁾ The set of values of the type **float** is a subset of the set of values of the type **double**; the set of values of the type **double** is a subset of the set of values of the type **long double**.
- 13 There are three *decimal floating types*, designated as **_Decimal32**, **_Decimal64**, and **_Decimal128**. Respectively, they have the ISO/IEC 60559 formats: decimal32,³⁶⁾ decimal64, and decimal128. (Decimal floating types are a conditional feature that implementations may not support; see 6.10.10.4.)

³⁰⁾Thus, **_BitInt**(3) is not the same type as **_BitInt**(4).

³¹⁾Implementation-defined keywords have the form of an identifier reserved for any use as described in 7.1.3.

³²⁾Any statement in this document about signed integer types also applies to the bit-precise signed integer types and the extended signed integer types, unless otherwise noted.

³³⁾Any statement in this document about unsigned integer types also applies to the bit-precise unsigned integer types and the extended unsigned integer types, unless otherwise specified.

³⁴⁾The same representation and alignment requirements are meant to imply interchangeability as arguments to functions, return values from functions, and members of unions.

³⁵⁾See “future language directions” (6.11.1).

³⁶⁾ISO/IEC 60559 specifies decimal32 as a data-interchange format that does not require arithmetic support; however, **_Decimal32** is a fully supported arithmetic type.

- 14 The standard floating types and the decimal floating types are collectively called the *real floating types*.
- 15 There are three *complex types*, designated as **float _Complex**, **double _Complex**, and **long double _Complex**. (Complex types are a conditional feature that implementations may not support; see 6.10.10.4.) The real floating and complex types are collectively called the *floating types*.
- 16 For each floating type there is a *corresponding real type*, which is always a real floating type. For real floating types, it is the same type. For complex types, it is the type given by deleting the keyword **_Complex** from the type name.
- 17 Each complex type has the same representation and alignment requirements as an array type containing exactly two elements of the corresponding real type; the first element is equal to the real part, and the second element to the imaginary part, of the complex number.
- 18 The type **char**, the signed and unsigned integer types, and the floating types are collectively called the *basic types*. The basic types are complete object types. Even if the implementation defines two or more basic types to have the same representation, they are nevertheless distinct types.
- 19 NOTE An implementation can define new keywords that provide alternative ways to designate a basic (or any other) type; this does not violate the requirement that all basic types be different. Implementation-defined keywords have the form of an identifier reserved for any use as described in 7.1.3.
- 20 The three types **char**, **signed char**, and **unsigned char** are collectively called the *character types*. The implementation shall define **char** to have the same range, representation, and behavior as either **signed char** or **unsigned char**.³⁷⁾
- 21 An *enumeration* comprises a set of named constants of integer type. Each distinct enumeration constitutes a different *enumerated type*.
- 22 The type **char**, the signed and unsigned integer types, and the enumerated types are collectively called *integer types*. The integer and real floating types are collectively called *real types*.
- 23 Integer and floating types are collectively called *arithmetic types*. Each arithmetic type belongs to one *type domain*: the *real type domain* comprises the real types, the *complex type domain* comprises the complex types.
- 24 The **void** type comprises an empty set of values; it is an incomplete object type that cannot be completed.
- 25 Any number of *derived types* can be constructed from the object and function types, as follows:
- An *array type* describes a contiguously allocated nonempty set of objects with a particular member object type, called the *element type*. The element type shall be complete whenever the array type is specified. Array types are characterized by their element type and by the number of elements in the array. An array type is said to be derived from its element type, and if its element type is *T*, the array type is sometimes called “array of *T*”. The construction of an array type from an element type is called “array type derivation”.
 - A *structure type* describes a sequentially allocated nonempty set of member objects (and, in certain circumstances, an incomplete array), each of which has an optionally specified name and possibly distinct type.
 - A *union type* describes an overlapping nonempty set of member objects, each of which has an optionally specified name and possibly distinct type.
 - A *function type* describes a function with specified return type. A function type is characterized by its return type and the number and types of its parameters. A function type is said to be derived from its return type, and if its return type is *T*, the function type is sometimes called “function returning *T*”. The construction of a function type from a return type is called “function type derivation”.

³⁷⁾ **CHAR_MIN**, defined in <limits.h>, will have one of the values 0 or **SCHAR_MIN**, and this can be used to distinguish the two options. Irrespective of the choice made, **char** is a separate type from the other two and is not compatible with either.

- A *pointer type* can be derived from a function type or an object type, called the *referenced type*. A pointer type describes an object whose value provides a reference to an entity of the referenced type. A pointer type derived from the referenced type *T* is sometimes called “pointer to *T*”. The construction of a pointer type from a referenced type is called “pointer type derivation”. A pointer type is a complete object type.
- An *atomic type* describes the type designated by the construct **_Atomic**(*type-name*). (Atomic types are a conditional feature that implementations may not support; see 6.10.10.4.)

These methods of constructing derived types can be applied recursively.

- 26 Arithmetic types, pointer types, and the **nullptr_t** type are collectively called *scalar types*. Array and structure types are collectively called *aggregate types*.³⁸⁾
- 27 An array type of unknown size is an incomplete type. It is completed, for an identifier of that type, by specifying the size in a later declaration (with internal or external linkage). A structure or union type of unknown content (as described in 6.7.3.4) is an incomplete type. It is completed, for all declarations of that type, by declaring the same structure or union tag with its defining content later in the same scope.
- 28 A complete type shall have a size that is less than or equal to **SIZE_MAX**. A type has *known constant size* if it is complete and is not a variable length array type.
- 29 Array, function, and pointer types are collectively called *derived declarator types*. A *declarator type derivation* from a type *T* is the construction of a derived declarator type from *T* by the application of an array-type, a function-type, or a pointer-type derivation to *T*.
- 30 A type is characterized by its *type category*, which is either the outermost derivation of a derived type (as noted previously in this subclause in the construction of derived types), or the type itself if the type consists of no derived types.
- 31 Any type so far mentioned is an *unqualified type*. Each unqualified type has several *qualified versions* of its type,³⁹⁾ corresponding to the combinations of one, two, or all three of the **const**, **volatile**, and **restrict** qualifiers. The qualified or unqualified versions of a type are distinct types that belong to the same type category and have the same representation and alignment requirements.⁴⁰⁾ An array and its element type are always considered to be identically qualified.⁴¹⁾ Any other derived type is not qualified by the qualifiers (if any) of the type from which it is derived.
- 32 Further, there is the **_Atomic** qualifier. The presence of the **_Atomic** qualifier designates an atomic type. The size, representation, and alignment of an atomic type may not be the same as those of the corresponding unqualified type. Therefore, this document explicitly uses the phrase “atomic, qualified, or unqualified type” whenever the atomic version of a type is permitted along with the other qualified versions of a type. The phrase “qualified or unqualified type”, without specific mention of atomic, does not include the atomic types.
- 33 A pointer to **void** shall have the same representation and alignment requirements as a pointer to a character type.⁴⁰⁾ Similarly, pointers to qualified or unqualified versions of compatible types shall have the same representation and alignment requirements. All pointers to structure types shall have the same representation and alignment requirements as each other. All pointers to union types shall have the same representation and alignment requirements as each other. Pointers to other types may not have the same representation or alignment requirements.
- 34 EXAMPLE 1 The type designated as “**float ***” has type “pointer to **float**”. Its type category is pointer, not a floating type. The const-qualified version of this type is designated as “**float * const**” whereas the type designated as “**const float ***” is not a qualified type — its type is “pointer to const-qualified **float**” and is a pointer to a qualified type.

³⁸⁾Note that aggregate type does not include union type because an object with union type can only contain one member at a time.

³⁹⁾See 6.7.4 regarding qualified array and function types.

⁴⁰⁾The same representation and alignment requirements are meant to imply interchangeability as arguments to functions, return values from functions, and members of unions.

⁴¹⁾This does not apply to the **_Atomic** qualifier. Qualifiers do not have any direct effect on the array type itself, but affect conversion rules for pointer types that reference an array type.

- 35 EXAMPLE 2 The type designated as “**struct tag** (*[5])(**float**)” has type “array of pointer to function returning **struct tag**”. The array has length five and the function has a single parameter of type **float**. Its type category is array.

Forward references: compatible type and composite type (6.2.7), declarations (6.7).

6.2.6 Representations of types

6.2.6.1 General

- 1 The representations of all types are unspecified except as stated in this subclause.
- 2 Except for bit-fields, objects are composed of contiguous sequences of one or more bytes, the number, order, and encoding of which are either explicitly specified or implementation-defined.
- 3 Values stored in unsigned bit-fields and objects of type **unsigned char** shall be represented using a pure binary notation.
- 4 Values stored in non-bit-field objects of any other object type are represented using $n \times \text{CHAR_BIT}$ bits, where n is the size of an object of that type, in bytes. An object that has the value can be copied into an object of type **unsigned char** [n] (e.g. by **memcpy**); the resulting set of bytes is called the *object representation* of the value. Values stored in bit-fields consist of m bits, where m is the size specified for the bit-field. The object representation is the set of m bits the bit-field comprises in the addressable storage unit holding it. Two values (other than NaNs) with the same object representation compare equal, but values that compare equal may have different object representations.
- 5 Certain object representations do not represent a value of the object type. If such a representation is read by an lvalue expression that does not have character type, the behavior is undefined. If such a representation is produced by a side effect that modifies all or any part of the object by an lvalue expression that does not have character type, the behavior is undefined.⁴²⁾ Such a representation is called a non-value representation.
- 6 When a value is stored in an object of structure or union type, including in a member object, the bytes of the object representation that correspond to any padding bytes take unspecified values (e.g. structure and union assignment can fail to copy any padding bits). The object representation of a structure or union object is never a non-value representation, even though the byte range corresponding to a member of the structure or union object can be a non-value representation for that member.
- 7 When a value is stored in a member of an object of union type, the bytes of the object representation that do not correspond to that member but do correspond to other members take unspecified values.
- 8 Where an operator is applied to a value that has more than one object representation, which object representation is used shall not affect the value of the result.⁴³⁾ Where a value is stored in an object using a type that has more than one object representation for that value, it is unspecified which representation is used, but a non-value representation shall not be generated.
- 9 Loads and stores of objects with atomic types are done with **memory_order_seq_cst** semantics.

Forward references: declarations (6.7), expressions (6.5.1), lvalues, arrays, and function designators (6.3.3.1), order and consistency (7.17.3).

6.2.6.2 Integer types

- 1 For unsigned integer types the bits of the object representation shall be divided into two groups: value bits and padding bits. If there are N value bits, each bit shall represent a different power of 2 between 1 and 2^{N-1} , so that objects of that type shall be capable of representing values from 0 to $2^N - 1$ using a pure binary representation; this shall be known as the value representation. The values of any padding bits are unspecified. The number of value bits N is called the *width* of the unsigned integer type. The type **bool** shall have one value bit and $(\text{sizeof}(\text{bool}) * \text{CHAR_BIT}) - 1$

⁴²⁾Thus, an automatic variable can be initialized to a non-value representation without causing undefined behavior, but the value of the variable cannot be used until a proper value is stored in it.

⁴³⁾It is possible for objects **x** and **y** with the same effective type **T** to have the same value when they are accessed as objects of type **T**, but to have different values in other contexts. In particular, if **==** is defined for type **T**, then **x == y** does not imply that **memcmp(&x, &y, sizeof(T)) == 0**. Furthermore, **x == y** does not necessarily imply that **x** and **y** have the same value; other operations on values of type **T** can distinguish between them.

padding bits. Otherwise, there is no requirement to have any padding bits; **unsigned char** shall not have any padding bits.

- 2 For signed integer types, the bits of the object representation shall be divided into three groups: value bits, padding bits, and the sign bit. If the corresponding unsigned type has width N , the signed type uses the same number of N bits, its *width*, as value bits and sign bit. $N - 1$ are value bits and the remaining bit is the sign bit. Each bit that is a value bit shall have the same value as the same bit in the object representation of the corresponding unsigned type. If the sign bit is zero, it shall not affect the resulting value. If the sign bit is one, it has value $-(2^{N-1})$. There can be padding bits; **signed char** shall not have any padding bits.
- 3 The values of any padding bits are unspecified. A valid object representation of a signed integer type where the sign bit is zero is a valid object representation of the corresponding unsigned type, and shall represent the same value. For any integer type, the object representation where all the bits are zero shall be a representation of the value zero in that type.
- 4 The *precision* of an integer type is the number of value bits.
- 5 NOTE 1 Some combinations of padding bits can generate non-value representations, for example, if one padding bit is a parity bit. Regardless, no arithmetic operation on valid values can generate a non-value representation other than as part of an exceptional condition such as an integer overflow. All other combinations of padding bits are alternative object representations of the value specified by the value bits.
- 6 NOTE 2 The sign representation defined in this document is called *two's complement*. Previous editions of this document (specifically ISO/IEC 9899:2018 and prior editions) additionally allowed other sign representations.
- 7 NOTE 3 For unsigned integer types the width and precision are the same, while for signed integer types the width is one greater than the precision.

6.2.7 Compatible type and composite type

- 1 Two types are *compatible types* if they are the same. Additional rules for determining whether two types are compatible are described in 6.7.3 for type specifiers, in 6.7.4 for type qualifiers, and in 6.7.7 for declarators.⁴⁴⁾ Moreover, two complete structure, union, or enumerated types declared with the same tag are compatible if members satisfy the following requirements:
 - there shall be a one-to-one correspondence between their members such that each pair of corresponding members are declared with compatible types;
 - if one member of the pair is declared with an alignment specifier, the other is declared with an equivalent alignment specifier;
 - and, if one member of the pair is declared with a name, the other is declared with the same name.

For two structures, corresponding members shall be declared in the same order. For two unions declared in the same translation unit, corresponding members shall be declared in the same order. For two structures or unions, corresponding bit-fields shall have the same widths. For two enumerations, corresponding members shall have the same values; if one has a fixed underlying type, then the other shall have a compatible fixed underlying type. For determining type compatibility, anonymous structures and unions are considered a regular member of the containing structure or union type, and the type of an anonymous structure or union is considered compatible with the type of another anonymous structure or union, respectively, if their members fulfill the preceding requirements.

Furthermore, two structure, union, or enumerated types declared in separate translation units are compatible in the following cases:

- both are declared without tags and they fulfill the preceding requirements;
- both have the same tag and are completed somewhere in their respective translation units and they fulfill the preceding requirements;

⁴⁴⁾Two types are not expected to be identical to be compatible.

— both have the same tag and at least one of the two types is not completed in its translation unit.

Otherwise, the structure, union, or enumerated types are incompatible.⁴⁵⁾

- 2 All declarations that refer to the same object or function shall have compatible type; otherwise, the behavior is undefined.
- 3 A *composite type* can be constructed from two types that are compatible. If both types are the same type, the composite type is this type. Otherwise, it is a type that is compatible with both and satisfies the following conditions:
 - If both types are structure types or both types are union types, the composite type is determined recursively by forming the composite types of their members.
 - If both types are array types, the following rules are applied:
 - If one type is an array of known constant size, the composite type is an array of that size.
 - Otherwise, if one type is a variable length array whose size is specified by an expression that is not evaluated, the behavior is undefined.
 - Otherwise, if one type is a variable length array whose size is specified, the composite type is a variable length array of that size.
 - Otherwise, if one type is a variable length array of unspecified size, the composite type is a variable length array of unspecified size.
 - Otherwise, both types are arrays of unknown size and the composite type is an array of unknown size.

The element type of the composite type is the composite type of the two element types.

- If both types are function types, the type of each parameter in the composite parameter type list is the composite type of the corresponding parameters.
- If one of the types has a standard attribute, the composite type also has that attribute.
- If both types are enumerated types, the composite type is an enumerated type.
- If one type is an enumerated type and the other is an integer type other than an enumerated type, it is implementation-defined whether or not the composite type is an enumerated type.

These rules apply recursively to the types from which the two types are derived.

- 4 If any of the original types satisfies all requirements of the composite type, it is unspecified whether the composite type is one of these types or a different type that satisfies the requirements.⁴⁶⁾
- 5 For an identifier with internal or external linkage declared in a scope in which a prior declaration of that identifier is visible,⁴⁷⁾ if the prior declaration specifies internal or external linkage, the type of the identifier at the later declaration becomes the composite type.
- 6 EXAMPLE Given the following two file scope declarations:

```
int f(int (*)(char *), double (*)[3]);
int f(int (*)(char *), double (*)[]);
```

The resulting composite type for the function is:

```
int f(int (*)(char *), double (*)[3]);
```

Forward references: array declarators (6.7.7.3).

⁴⁵⁾A structure, union, or enumerated type without a tag or an incomplete structure, union or enumerated type is not compatible with any other structure, union or enum type declared in the same translation unit.

⁴⁶⁾The notion of “same type” affects redeclarations of typedef names and tags in the same scope.

⁴⁷⁾As specified in 6.2.1, the later declaration can hide the prior declaration.

6.2.8 Alignment of objects

- 1 Complete object types have alignment requirements which place restrictions on the addresses at which objects of that type can be allocated. An alignment is an implementation-defined integer value representing the number of bytes between successive addresses at which a given object can be allocated. An object type imposes an alignment requirement on every object of that type: stricter alignment can be requested using the **alignas** keyword.
- 2 A *fundamental alignment* is a valid alignment less than or equal to **alignof(max_align_t)**. Fundamental alignments shall be supported by the implementation for objects of all storage durations. The alignment requirements of the following types shall be fundamental alignments:
 - all qualified or unqualified basic types;
 - all qualified or unqualified enumerated types;
 - all qualified or unqualified pointer types;
 - all array types whose element type has a fundamental alignment requirement;
 - all types specified in Clause 7 as complete object types;
 - all structure or union types whose elements have types with fundamental alignment requirements and none of whose elements have an alignment specifier specifying an alignment that is not a fundamental alignment.

Whether any atomic types have fundamental alignment is implementation-defined.

- 3 An *extended alignment* is represented by an alignment greater than **alignof(max_align_t)**. It is implementation-defined whether any extended alignments are supported and the storage durations for which they are supported. A type having an extended alignment requirement is an *over-aligned type*.⁴⁸⁾
- 4 Alignments are represented as values of the type **size_t**. Valid alignments include only fundamental alignments, plus an additional implementation-defined set of values, which can be empty. Every valid alignment value shall be a nonnegative integral power of two.
- 5 Alignments have an order from *weaker* to *stronger* or *stricter* alignments. Stricter alignments have larger alignment values. An address that satisfies an alignment requirement also satisfies any weaker valid alignment requirement.
- 6 The alignment requirement of a complete type can be queried using an **alignof** expression. The types **char**, **signed char**, and **unsigned char** shall have the weakest alignment requirement.
- 7 Comparing alignments is meaningful and provides the obvious results:
 - Two alignments are equal when their numeric values are equal.
 - Two alignments are different when their numeric values are not equal.
 - When an alignment is larger than another it represents a stricter alignment.

6.2.9 Encodings

- 1 The *literal encoding* is an implementation-defined mapping of the characters of the execution character set to the values in a character literal (6.4.5.5) or string literal (6.4.6). It shall support a mapping from all the basic execution character set values into the implementation-defined encoding. It can contain multibyte character sequences (5.3.2).
- 2 The *wide literal encoding* is an implementation-defined mapping of the characters of the execution character set to the values in a **wchar_t** character literal (6.4.5.5) or a **wchar_t** string literal (6.4.6). It

⁴⁸⁾ Every over-aligned type is, or contains, a structure or union type with a member to which an extended alignment has been applied, or an atomic type that does not have a fundamental alignment.

shall support a mapping from all the basic execution character set values into the implementation-defined encoding. The mapping shall produce values identical to the literal encoding for all the basic execution character set values if an implementation does not define `__STDC_MB_MIGHT_NEQ_WC__`. One or more values may map to one or more values of the extended execution character set.

- 3 A *code unit* is a single compositional unit of encoded information, usually of type `char`, `wchar_t`, `char8_t`, `char16_t`, or `char32_t`.
- 4 A *code point* is a single compositional unit of decoded information. Code points are generally used as the single complete decoded output, or as an intermediary to transcode to other code units. A Unicode code point is a single compositional unit of decoded information as defined in ISO/IEC 10646, typically used to convert to or from UTF-8, UTF-16, and UTF-32.
- 5 The *narrow execution encoding* is the implementation-defined, `LC_CTYPE`-influenced (7.11.1), locale-based execution environment encoding. The *wide execution encoding* is the implementation-defined, `LC_CTYPE`-influenced, locale-based wide execution environment encoding. Both of these encodings are called the *execution encodings*.
- 6 An *unrecorded encoding error* occurs when an encoding, decoding, or transcoding function encounters an input sequence of code units or code points that
 - does not form a valid sequence according to the encoding being associated with the sequence
 - or, is not representable in the output encoding or coded character set.
- 7 An *encoding error* is the same as an unrecorded encoding error, except that the value of the macro `EILSEQ` (7.5) is stored in `errno` when such an error occurs during execution of the functions defined in this document unless otherwise specified.

6.3 Conversions

6.3.1 Introduction

- 1 Several operators convert operand values from one type to another automatically. This subclause specifies the result required from such an *implicit conversion*, as well as those that result from a cast operation (an *explicit conversion*). The list in 6.3.2.8 summarizes the conversions performed by most ordinary operators; it is supplemented as required by the discussion of each operator in 6.5.1.
- 2 Unless explicitly stated otherwise, conversion of an operand value to a compatible type causes no change to the value or the representation.

Forward references: cast operators (6.5.5).

6.3.2 Arithmetic operands

6.3.2.1 Boolean, characters, and integers

- 1 Every integer type has an *integer conversion rank* defined as follows:
 - No two signed integer types shall have the same rank, even if they have the same representation.
 - The rank of a signed integer type shall be greater than the rank of any signed integer type with less precision.
 - The rank of `long long int` shall be greater than the rank of `long int`, which shall be greater than the rank of `int`, which shall be greater than the rank of `short int`, which shall be greater than the rank of `signed char`.
 - The rank of a bit-precise signed integer type shall be greater than the rank of any standard integer type with less width or any bit-precise integer type with less width.
 - The rank of any unsigned integer type shall equal the rank of the corresponding signed integer type, if any.

- The rank of any standard integer type shall be greater than the rank of any extended integer type with the same width or bit-precise integer type with the same width.
- The rank of any bit-precise integer type relative to an extended integer type of the same width is implementation-defined.
- The rank of **char** shall equal the rank of **signed char** and **unsigned char**.
- The rank of **bool** shall be less than the rank of all other standard integer types.
- The rank of any enumerated type shall equal the rank of the compatible integer type (see 6.7.3.3).
- The rank of any extended signed integer type relative to another extended signed integer type with the same precision is implementation-defined, but still subject to the other rules for determining the integer conversion rank.
- For all integer types **T1**, **T2**, and **T3**, if **T1** has greater rank than **T2** and **T2** has greater rank than **T3**, then **T1** has greater rank than **T3**.

2 The following can be used in an expression wherever an **int** or **unsigned int** can be used:

- An object or expression with an integer type (other than **int** or **unsigned int**) whose integer conversion rank is less than or equal to the rank of **int** and **unsigned int**.
- A bit-field of type **bool**, **int**, **signed int**, or **unsigned int**.

The value from a bit-field of a bit-precise integer type is converted to the corresponding bit-precise integer type. If the original type is not a bit-precise integer type (6.2.5): if an **int** can represent all values of the original type (as restricted by the width, for a bit-field), the value is converted to an **int**;⁴⁹⁾ otherwise, it is converted to an **unsigned int**. These are called the *integer promotions*. All other types are unchanged by the integer promotions.

3 NOTE The integer promotions are applied only:

1. as part of the usual arithmetic conversions,
2. to certain argument expressions,
3. to the operands of the unary **+**, **-**, and **~** operators,
4. and to both operands of the shift operators,

as specified by their respective subclauses.

4 The integer promotions preserve value including sign. As discussed earlier, whether a “plain” **char** can hold negative values is implementation-defined.

Forward references: enumeration specifiers (6.7.3.3), structure and union specifiers (6.7.3.2).

6.3.2.2 Boolean type

1 When any scalar value is converted to **bool**, the result is **false** if the value is a zero (for arithmetic types), null (for pointer types), or the scalar has type **nullptr_t**; otherwise, the result is **true**.

6.3.2.3 Signed and unsigned integers

- 1 When a value with integer type is converted to another integer type other than **bool**, if the value can be represented by the new type, it is unchanged.
- 2 Otherwise, if the new type is unsigned, the value is converted by repeatedly adding or subtracting one more than the maximum value that can be represented in the new type until the value is in the range of the new type.⁵⁰⁾
- 3 Otherwise, the new type is signed and the value cannot be represented in it; either the result is implementation-defined or an implementation-defined signal is raised.

⁴⁹⁾E.g. **unsigned _BitInt(7)**: 2 is a bit-field that can hold the values 0, 1, 2, 3, and converts to **unsigned _BitInt(7)**.

⁵⁰⁾The rules describe arithmetic on the mathematical value, not the value of a given type of expression.

6.3.2.4 Real floating and integer

- 1 When a finite value of standard floating type is converted to an integer type other than **bool**, the fractional part is discarded (i.e. the value is truncated toward zero). If the value of the integral part cannot be represented by the integer type, the behavior is undefined.⁵¹⁾
- 2 When a finite value of decimal floating type is converted to an integer type other than **bool**, the fractional part is discarded (i.e. the value is truncated toward zero). If the value of the integral part cannot be represented by the integer type, the “invalid” floating-point exception shall be raised and the result of the conversion is unspecified.
- 3 When a value of integer type is converted to a standard floating type, if the value being converted can be represented exactly in the new type, it is unchanged. If the value being converted is in the range of values that can be represented but cannot be represented exactly, the result is either the nearest higher or nearest lower representable value, chosen in an implementation-defined manner. If the value being converted is outside the range of values that can be represented, the behavior is undefined. Results of some implicit conversions can be represented in greater range and precision than that required by the new type (see 6.3.2.8 and 6.8.7.5).
- 4 When a value of integer type is converted to a decimal floating type, if the value being converted can be represented exactly in the new type, it is unchanged. If the value being converted cannot be represented exactly, the result shall be correctly rounded with exceptions raised as specified in ISO/IEC 60559.

6.3.2.5 Real floating types

- 1 When a value of real floating type is converted to a real floating type, if the value being converted can be represented exactly in the new type, it is unchanged.
- 2 When a value of real floating type is converted to a standard floating type, if the value being converted is in the range of values that can be represented but cannot be represented exactly, the result is either the nearest higher or nearest lower representable value, chosen in an implementation-defined manner. If the value being converted is outside the range of values that can be represented, the behavior is undefined.
- 3 When a value of real floating type is converted to a decimal floating type, if the value being converted cannot be represented exactly, the result is correctly rounded with exceptions raised as specified in ISO/IEC 60559.
- 4 Results of some implicit conversions may be represented in greater range and precision than that required by the new type (see 6.3.2.8 and 6.8.7.5).

6.3.2.6 Complex types

- 1 When a value of complex type is converted to another complex type, both the real and imaginary parts follow the conversion rules for the corresponding real types.

6.3.2.7 Real and complex

- 1 When a value of real type is converted to a complex type, the real part of the complex result value is determined by the rules of conversion to the corresponding real type and the imaginary part of the complex result value is a positive zero or an unsigned zero.
- 2 When a value of complex type is converted to a real type other than **bool**,⁵²⁾ the imaginary part of the complex value is discarded and the value of the real part is converted according to the conversion rules for the corresponding real type.

6.3.2.8 Usual arithmetic conversions

- 1 Many operators that expect operands of arithmetic type cause conversions and yield result types in a similar way. The purpose is to determine a *common real type* for the operands and result. For the specified operands, each operand is converted, without change of type domain, to a type whose

⁵¹⁾The remaindering operation performed when a value of integer type is converted to unsigned type is not necessary to be performed when a value of real floating type is converted to unsigned type. Thus, the range of portable real floating values is $(-1, \text{Utype_MAX} + 1)$.

⁵²⁾See 6.3.2.2.

corresponding real type is the common real type. Unless explicitly stated otherwise, the common real type is also the corresponding real type of the result, whose type domain is the type domain of the operands if they are the same, and complex otherwise. This pattern is called the *usual arithmetic conversions*:

If one operand has decimal floating type, the other operand shall not have standard floating, or complex type.

First, if the type of either operand is **_Decimal128**, the other operand is converted to **_Decimal128**.

Otherwise, if the type of either operand is **_Decimal64**, the other operand is converted to **_Decimal64**.

Otherwise, if the type of either operand is **_Decimal32**, the other operand is converted to **_Decimal32**.

Otherwise, if the corresponding real type of either operand is **long double**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **long double**.

Otherwise, if the corresponding real type of either operand is **double**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **double**.

Otherwise, if the corresponding real type of either operand is **float**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **float**.⁵³⁾

Otherwise, if any of the two types is an enumeration, it is converted to its underlying type. Then, the integer promotions are performed on both operands. Next, the following rules are applied to the promoted operands:

If both operands have the same type, then no further conversion is needed.

Otherwise, if both operands have signed integer types or both have unsigned integer types, the operand with the type of lesser integer conversion rank is converted to the type of the operand with greater rank.

Otherwise, if the operand that has unsigned integer type has rank greater or equal to the rank of the type of the other operand, then the operand with signed integer type is converted to the type of the operand with unsigned integer type.

Otherwise, if the type of the operand with signed integer type can represent all the values of the type of the operand with unsigned integer type, then the operand with unsigned integer type is converted to the type of the operand with signed integer type.

Otherwise, both operands are converted to the unsigned integer type corresponding to the type of the operand with signed integer type.

- 2 The values of floating operands and of the results of floating expressions may be represented in greater range and precision than that required by the type; the types are not changed thereby. See 5.3.5.3.3 regarding evaluation formats.
- 3 EXAMPLE One consequence of **_BitInt** being exempt from the integer promotion rules (6.3.2) is that a **_BitInt** operand of a binary operator is not always promoted to an **int** or **unsigned int** as part of the usual arithmetic conversions. Instead, a lower-ranked operand is converted to the higher-rank operand type and the result of the operation is the higher-ranked type.

```
_BitInt(2) a2 = 1;
_BitInt(3) a3 = 2;
_BitInt(33) a33 = 1;
signed char c = 3;
```

⁵³⁾For example, addition of a **double** **_Complex** and a **float** entails just the conversion of the **float** operand to **double** (and yields a **double** **_Complex** result).

```

a2 * a3; /* As part of the multiplication, a2 is converted to
          _BitInt(3) and the result type is _BitInt(3). */
a2 * c; /* As part of the multiplication, c is promoted to int,
          a2 is converted to int and the result type is int. */
a33 * c; /* As part of the multiplication, c is promoted to int.
          Then, provided int has a width of at most 32,
          it is converted to _BitInt(33) and the result type
          is _BitInt(33). */

void func(_BitInt(8) a8, _BitInt(24) a24) {
    /* Cast one of the operands to 32-bits to guarantee the
       result of the multiplication can contain all possible values. */
    _BitInt(32) a32 = a8 * (_BitInt(32))a24;
}

```

6.3.3 Other operands

6.3.3.1 Lvalues, arrays, and function designators

- 1 An *lvalue* is an expression (with an object type other than **void**) that potentially designates an object;⁵⁴⁾ if an lvalue does not designate an object when it is evaluated, the behavior is undefined. When an object is said to have a particular type, the type is specified by the lvalue used to designate the object. A *modifiable lvalue* is an lvalue that does not have array type, does not have an incomplete type, does not have a const-qualified type, and if it is a structure or union, does not have any member (including, recursively, any member or element of all contained aggregates or unions) with a const-qualified type.
- 2 Except when it is the operand of of:
 - the **sizeof** operator,
 - the typeof operators,
 - the unary & operator,
 - the ++ operator,
 - the -- operator,
 - the left operand of the . operator,
 - or, an assignment operator

an lvalue that does not have array type is converted to the value stored in the designated object (and is no longer an lvalue); this is called *lvalue conversion*. If the lvalue has qualified type, the value has the unqualified version of the type of the lvalue; additionally, if the lvalue has atomic type, the value has the non-atomic version of the type of the lvalue; otherwise, the value has the type of the lvalue. If the lvalue has an incomplete type and does not have array type, the behavior is undefined. If the lvalue (that does not have array type) designates an object of automatic storage duration that never had its address taken, and that object is uninitialized (not declared with an initializer and no assignment to it has been performed prior to use), the behavior is undefined. An address of an object is taken by application of the address operator to an lvalue designating the object, if the object is an element of an array that was designated by an expression that is converted to a pointer as described below, or by being a member of a structure or union object whose address is taken.

- 3 Except when it is the operand of the **sizeof** operator, or the **_Lengthof** operator, or typeof operators, or the unary & operator, or is a string literal used to initialize an array, an expression that has type

⁵⁴⁾The name “lvalue” comes originally from the assignment expression **E1** = **E2**, in which the left operand **E1** is required to be a (modifiable) lvalue. It is perhaps better considered as representing an object “locator value”. What is sometimes called “rvalue” is in this document described as the “value of an expression”.

An obvious example of an lvalue is an identifier of an object. As a further example, if **E** is a unary expression that is a pointer to an object, ***E** is an lvalue that designates the object to which **E** points.

“array of *type*” is converted to an expression with type “pointer to *type*” that points to the initial element of the array object and is not an lvalue. If the array object has register storage class, the behavior is implementation-defined.

- 4 A *function designator* is an expression that has function type. Except when it is the operand of the **sizeof** operator,⁵⁵⁾ a **typeof** operator, or the unary **&** operator, a function designator with type “function returning *type*” is converted to an expression that has type “pointer to function returning *type*”.

Forward references: address and indirection operators (6.5.4.3), assignment operators (6.5.17), common definitions `<stddef.h>` (7.21), initialization (6.7.11), postfix increment and decrement operators (6.5.3.5), prefix increment and decrement operators (6.5.4.2), the **sizeof**, **_Lengthof**, and **alignof** operators (6.5.4.5), structure and union members (6.5.3.4).

6.3.3.2 void

- 1 The (nonexistent) value of a *void expression* (an expression that has type **void**) shall not be used in any way, and implicit or explicit conversions (except to **void**) shall not be applied to such an expression. If an expression of any other type is evaluated as a void expression, its value or designator is discarded. (A void expression is evaluated for its side effects.)

6.3.3.3 Pointers

- 1 A pointer to **void** can be converted to or from a pointer to any object type. A pointer to any object type can be converted to a pointer to **void** and back again; the result shall compare equal to the original pointer.
- 2 For any qualifier *q*, a pointer to a non-*q*-qualified type can be converted to a pointer to the *q*-qualified version of the type; the values stored in the original and converted pointers shall compare equal.
- 3 An integer constant expression with the value 0, such an expression cast to type **void ***, or the predefined constant **nullptr** is called a *null pointer constant*.⁵⁶⁾ If a null pointer constant or a value of the type **nullptr_t** (which is necessarily the value **nullptr**) is converted to a pointer type, the resulting pointer, called a *null pointer*, is guaranteed to compare unequal to a pointer to any object or function.
- 4 Conversion of a null pointer to another pointer type yields a null pointer of that type. Any two null pointers shall compare equal.
- 5 An integer can be converted to any pointer type. Except as previously specified, the result is implementation-defined, possibly not correctly aligned, can possibly not point to an entity of the referenced type, and can produce an indeterminate representation when stored into an object.⁵⁷⁾
- 6 Any pointer type can be converted to an integer type. Except as previously specified, the result is implementation-defined. If the result cannot be represented in the integer type, the behavior is undefined. The result is not required to be in the range of values of any integer type.
- 7 A pointer to an object type can be converted to a pointer to a different object type. If the resulting pointer is not correctly aligned⁵⁸⁾ for the referenced type, the behavior is undefined. Otherwise, when converted back again, the result shall compare equal to the original pointer. When a pointer to an object is converted to a pointer to a character type, the result points to the lowest addressed byte of the object. Successive increments of the result, up to the size of the object, yield pointers to the remaining bytes of the object.
- 8 A pointer to a function of one type can be converted to a pointer to a function of another type and back again; the result shall compare equal to the original pointer. If a converted pointer is used to call a function whose type is not compatible with the referenced type, the behavior is undefined.

⁵⁵⁾Because this conversion does not occur, the operand of the **sizeof** operator remains a function designator and violates the constraints in 6.5.4.5

⁵⁶⁾The macro **NULL** is defined in `<stddef.h>` (and other headers) as a null pointer constant; see 7.21.

⁵⁷⁾The mapping functions for converting a pointer to an integer or an integer to a pointer are intended to be consistent with the addressing structure of the execution environment.

⁵⁸⁾In general, the concept “correctly aligned” is transitive: if a pointer to type *A* is correctly aligned for a pointer to type *B*, which in turn is correctly aligned for a pointer to type *C*, then a pointer to type *A* is correctly aligned for a pointer to type *C*.

6.3.3.4 `nullptr_t`

- 1 The type `nullptr_t` can be converted to `void`, `bool` or to a pointer type; the result is a `void` expression, `false`, or a null pointer value, respectively.
- 2 A null pointer constant or value of type `nullptr_t` can be converted to `nullptr_t`.

Forward references: cast operators (6.5.5), equality operators (6.5.10), integer types capable of holding object pointers (7.22.2.5), simple assignment (6.5.17.2), the `nullptr_t` type (7.21.3).

6.4 Lexical elements

6.4.1 General

Syntax

- 1 *token*:
- keyword*
 - identifier*
 - constant*
 - string-literal*
 - punctuator*
- preprocessing-token*:
- header-name*
 - identifier*
 - pp-number*
 - character-literal*
 - string-literal*
 - punctuator*
 - each universal character name that cannot be one of the above
 - each non-white-space character that cannot be one of the above

Constraints

- 2 Each preprocessing token that is converted to a token shall have the lexical form of a keyword, an identifier, a constant, a string literal, or a punctuator. A single universal character name shall match one of the other preprocessing token categories.

Semantics

- 3 A *token* is the minimal lexical element of the language in translation phases 7 and 8 (5.2.1.2). The categories of tokens are: keywords, identifiers, constants, string literals, and punctuators. A preprocessing token is the minimal lexical element of the language in translation phases 3 through 6. The categories of preprocessing tokens are: header names, identifiers, preprocessing numbers, character literals, string literals, punctuators, and both single universal character names as well as single non-white-space characters that do not lexically match the other preprocessing token categories.⁵⁹⁾ If a ' or a " character matches the last category, the behavior is undefined. Preprocessing tokens can be separated by *white space*; this consists of comments (described later), or *white-space characters* (space, horizontal tab, new-line, vertical tab, and form-feed), or both. As described in 6.10, in certain circumstances during translation phase 4, white space (or the absence thereof) serves as more than preprocessing token separation. White space may appear within a preprocessing token only as part of a header name or between the quotation characters in a character literal or string literal.
- 4 If the input stream has been parsed into preprocessing tokens up to a given character, the next preprocessing token is the longest sequence of characters that could constitute a preprocessing token. There is one exception to this rule: header name preprocessing tokens are recognized only within **#include** and **#embed** preprocessing directives, in **__has_include** and **__has_embed** expressions, as well as in implementation-defined locations within **#pragma** directives. In such contexts, a sequence of characters that could be either a header name or a string literal is recognized as the former.
- 5 EXAMPLE 1 The program fragment **1Ex** is parsed as a preprocessing number token (one that is not a valid floating or integer literal), even though a parse as the pair of preprocessing tokens **1** and **Ex** can produce a valid expression (for example, if **Ex** were a macro defined as +1). Similarly, the program fragment **1E1** is parsed as a preprocessing number (one that is a valid floating literal), whether or not **E** is a macro name.
- 6 EXAMPLE 2 The program fragment **x+++++y** is parsed as **x ++ ++ + y**, which violates a constraint on increment operators, even though the parse **x ++ + ++ y** can yield a correct expression.

⁵⁹⁾An additional category, placemarkers, is used internally in translation phase 4 (see 6.10.5.4); it cannot occur in source files.

Forward references: character literals (6.4.5.5), comments (6.4.10), expressions (6.5.1), floating literals (6.4.5.3), header names (6.4.8), macro replacement (6.10.5), postfix increment and decrement operators (6.5.3.5), prefix increment and decrement operators (6.5.4.2), preprocessing directives (6.10), preprocessing numbers (6.4.9), string literals (6.4.6).

6.4.2 Keywords

Syntax

- 1 *keyword*: one of

alignas	do	int	struct	while
alignof	double	long	switch	_Atomic
auto	else	nullptr	thread_local	_BitInt
bool	enum	register	true	_Complex
break	extern	restrict	typedef	_Decimal128
case	false	return	typeof	_Decimal32
char	float	short	typeof_unqual	_Decimal64
const	for	signed	union	_Generic
constexpr	goto	sizeof	unsigned	_Lengthof
continue	if	static	void	_Noreturn
default	inline	static_assert	volatile	

Semantics

- 2 The previously listed tokens (case sensitive) are reserved (in translation phases 7 and 8) for use as keywords except in an attribute token, and shall not be used otherwise.
- 3 Table 6.1 provides alternate spellings for certain keywords. These can be used wherever the keyword can.⁶⁰⁾

Table 6.1 — Keywords and their spellings

Keyword	Alternative Spelling
alignas	_Alignas
alignof	_Alignof
bool	_Bool
static_assert	_Static_assert
thread_local	_Thread_local

The spelling of these keywords, their alternate forms, and of **false** and **true** inside expressions that are subject to the **#** and **##** preprocessing operators is unspecified.⁶¹⁾

⁶⁰⁾These alternative keywords are obsolescent features and should not be used for new code and development.

⁶¹⁾The intent of this specification is to allow but not force the implementation of the corresponding feature by means of a predefined macro.

6.4.3 Identifiers

6.4.3.1 General

Syntax

1 *identifier*:

identifier-start
identifier identifier-continue

identifier-start:

nondigit
 XID_Start character
 universal character name of class XID_Start

identifier-continue:

digit
nondigit
 XID_Continue character
 universal character name of class XID_Continue

nondigit: one of

_	a	b	c	d	e	f	g	h	i	j	k	l	m
	n	o	p	q	r	s	t	u	v	w	x	y	z
	A	B	C	D	E	F	G	H	I	J	K	L	M
	N	O	P	Q	R	S	T	U	V	W	X	Y	Z

digit: one of

0 1 2 3 4 5 6 7 8 9

Semantics

- 2 An XID_Start character is an implementation-defined character whose corresponding code point in ISO/IEC 10646 has the XID_Start property. An XID_Continue character is an implementation-defined character whose corresponding code point in ISO/IEC 10646 has the XID_Continue property. An identifier is a sequence of one identifier start character followed by 0 or more identifier continue characters, which designates one or more entities as described in 6.2.1. It is implementation-defined if a \$ (U+0024, DOLLAR SIGN) may be used as a nondigit character. Lowercase and uppercase letters are distinct. There is no specific limit on the maximum length of an identifier.
- 3 The character classes XID_Start and XID_Continue are Derived Core Properties as described by UAX #44.⁶²⁾ Each character and universal character name in an identifier shall designate a character whose encoding in ISO/IEC 10646 has the XID_Continue property. The initial character (which can be a universal character name) shall designate a character whose encoding in ISO/IEC 10646 has the XID_Start property. An identifier shall conform to Normalization Form C as specified in ISO/IEC 10646. Annex D provides an overview of the conforming identifiers.
- 4 NOTE 1 Uppercase and lowercase letters are considered different for all identifiers.
- 5 NOTE 2 In translation phase 4, the term identifier also includes those preprocessing tokens (6.4.9) differentiated as keywords (6.4.2) in the later translation phase 7 (5.2.1.2).

⁶²⁾On systems that cannot accept extended characters in external identifiers, an encoding of the universal character name can be used in forming such identifiers. For example, some otherwise unused character or sequence of characters can be used to encode the \u in a universal character name.

- 6 When preprocessing tokens are converted to tokens during translation phase 7, if a preprocessing token could be converted to either a keyword or an identifier, it is converted to a keyword except in an attribute token.
- 7 Some identifiers are reserved.
 - All identifiers that begin with a double underscore (__) or begin with an underscore (_) followed by an uppercase letter are reserved for any use, except those identifiers which are lexically identical to keywords.⁶³⁾
 - All identifiers that begin with an underscore are reserved for use as identifiers with file scope in both the ordinary and tag name spaces.

Other identifiers may be reserved, see 7.1.3.

- 8 If the program declares or defines an identifier in a context in which it is reserved (other than as allowed by 7.1.4), the behavior is undefined.
- 9 If the program defines a reserved identifier or standard attribute token described in 6.7.13.2 as a macro name, or removes (with **#undef**) any macro definition of an identifier in the first group listed previously or standard attribute token described in 6.7.13.2, the behavior is undefined.
- 10 Some identifiers may be potentially reserved. A *potentially reserved identifier* is an identifier which is not reserved unless made so by an implementation providing the identifier (7.1.3) but is anticipated to become reserved by an implementation or a future version of this document. An identifier that this document describes as optional:
 - If it is defined as a macro it is reserved.
 - Otherwise, if the definition is given in clauses 1 to 6 it is reserved.
 - Otherwise, it is potentially reserved.

Recommended practice

- 11 Implementations are encouraged to issue a diagnostic message when a potentially reserved identifier is declared or defined for any use that is not implementation-compatible (see subsequent description in this subclause) in a context where the potentially reserved identifier may be reserved under a conforming implementation. This brings attention to a potential conflict when porting a program to a future edition of this document.
- 12 An implementation-compatible use of a potentially reserved identifier is a declaration of an external name where the name is provided by the implementation as an external name and where the declaration declares an object or function with a type that is compatible with the type of the object or function provided by the implementation under that name.

Implementation limits

- 13 As discussed in 5.3.5.2, an implementation may limit the number of significant initial characters in an identifier; the limit for an *external name* (an identifier that has external linkage) may be more restrictive than that for an *internal name* (a macro name or an identifier that does not have external linkage). The number of significant characters in an identifier is implementation-defined.
- 14 Any identifiers that differ in a significant character are different identifiers. If two identifiers differ only in nonsignificant characters, the behavior is undefined.

Forward references: universal character names (6.4.4), macro replacement (6.10.5), reserved library identifiers (7.1.3), use of library functions (7.1.4), attributes (6.7.13.2).

6.4.3.2 Predefined identifiers

Semantics

- 1 The identifier **__func__** shall be implicitly declared by the translator as if, immediately following the opening brace of each function definition, the declaration

⁶³⁾This allows a reserved identifier that matches the spelling of a keyword to be used as a macro name by the program.

```
static const char __func__[] = "function-name";
```

appeared, where *function-name* is the name of the lexically-enclosing function.⁶⁴⁾

- 2 This name is encoded as if the implicit declaration had been written in the source character set and then translated into the execution character set as indicated in translation phase 5.
- 3 EXAMPLE The following code fragment can be used as an example:

```
#include <stdio.h>
void myfunc(void)
{
    printf("%s\n", __func__);
    /* ... */
}
```

Each time the function is called, it will print to the standard output stream:

```
myfunc
```

Forward references: function definitions (6.9.2).

6.4.4 Universal character names

Syntax

- 1 *universal-character-name*:
 - `\u` *hex-quad*
 - `\U` *hex-quad hex-quad*
 - `\u{ simple-hexadecimal-digit-sequence }`
 - `\U{ simple-hexadecimal-digit-sequence }`

hex-quad:

hexadecimal-digit hexadecimal-digit hexadecimal-digit hexadecimal-digit

simple-hexadecimal-digit-sequence:

hexadecimal-digit
simple-hexadecimal-digit-sequence hexadecimal-digit

Constraints

- 2 A universal character name shall not designate a code point where the hexadecimal value is:
 - in the range D800 through DFFF inclusive; or
 - greater than 10FFFF.⁶⁵⁾

A universal character name outside the c-char sequence of a character literal, or the s-char sequence of a string literal shall not designate a control character or a character in the basic character set.

⁶⁴⁾Since the name `__func__` is reserved for any use by the implementation (7.1.3), if any other identifier is explicitly declared using the name `__func__`, the behavior is undefined.

⁶⁵⁾The disallowed characters are the characters in the basic character set and the code positions reserved by ISO/IEC 10646 for control characters, the character DELETE, the S-zone (reserved for use by UTF-16), and characters too large to be encoded by ISO/IEC 10646. Disallowed universal character escape sequences can still be specified with hexadecimal and octal escape sequences (6.4.5.5).

Description

- 3 Universal character names can be used in identifiers, character literals, and string literals to designate characters that are not in the basic character set.

Semantics

- 4 A universal character name designates the character in ISO/IEC 10646 whose code point is the hexadecimal value represented by the sequence of hexadecimal digits in the universal character name.

6.4.5 Constants

6.4.5.1 General

Syntax

- 1 *constant*:
- integer-literal*
 - floating-literal*
 - enumeration-constant*
 - character-literal*
 - predefined-constant*

Constraints

- 2 Each constant shall have a type and the value of a constant shall be in the range of representable values for its type.

Semantics

- 3 Each constant has a type, determined by its form and value, as detailed later.

6.4.5.2 Integer literals

Syntax

- 1 *integer-literal*:
- decimal-literal integer-suffix_{opt}*
 - octal-literal integer-suffix_{opt}*
 - hexadecimal-literal integer-suffix_{opt}*
 - binary-literal integer-suffix_{opt}*

decimal-literal:

nonzero-digit
decimal-literal ' _{opt} *digit*

hexadecimal-literal:

hexadecimal-prefix hexadecimal-digit-sequence

octal-literal:

prefixed-octal-literal
unprefixed-octal-literal

unprefixed-octal-literal:

0
0 ' _{opt} *octal-digit-sequence*

prefixed-octal-literal:

octal-prefix octal-digit-sequence

binary-literal:

binary-prefix binary-digit
binary-literal ' _{opt} *binary-digit*

hexadecimal-prefix: one of
0x 0X

octal-prefix: one of
0o 0O

binary-prefix: one of
0b 0B

nonzero-digit: one of
1 2 3 4 5 6 7 8 9

hexadecimal-digit-sequence:
hexadecimal-digit
hexadecimal-digit-sequence ' _{opt} *hexadecimal-digit*

octal-digit-sequence:
octal-digit
octal-digit-sequence ' _{opt} *octal-digit*

hexadecimal-digit: one of
0 1 2 3 4 5 6 7 8 9
a b c d e f
A B C D E F

octal-digit: one of
0 1 2 3 4 5 6 7

binary-digit: one of
0 1

integer-suffix:
unsigned-suffix *long-suffix*_{opt}
unsigned-suffix *long-long-suffix*
unsigned-suffix *bit-precise-int-suffix*
long-suffix *unsigned-suffix*_{opt}
long-long-suffix *unsigned-suffix*_{opt}
bit-precise-int-suffix *unsigned-suffix*_{opt}

bit-precise-int-suffix: one of
wb WB

unsigned-suffix: one of
u U

long-suffix: one of
l L

long-long-suffix: one of
ll LL

Description

- 2 An integer literal begins with a digit, but has no period or exponent part. It can have a prefix that specifies its base and a suffix that specifies its type. An optional separating single quote character (') in an integer or floating literal is called a digit separator. Digit separators are ignored when determining the value of the literal.
- 3 EXAMPLE 1 The following integer literals use digit separators; the comment associated with each literal shows the equivalent literal without digit separators.

```
0b11'10'11'01 /* 0b11101101 */
'1'2 /* character literal '1' followed by integer literal 2,
      not the integer literal 12 */
11'22 /* 1122 */
0x'FFFF'FFFF /* invalid hexadecimal literal (' cannot appear after 0x) */
0x1'2'3'4AB'C'D /* 0x1234ABCD */
```

- 4 A decimal literal begins with a nonzero digit and consists of a sequence of decimal digits. An octal literal consists of the prefix **0o** or **0O** followed by a sequence of the digits **0** through **7** only. A hexadecimal literal consists of the prefix **0x** or **0X** followed by a sequence of the decimal digits and the letters **a** (or **A**) through **f** (or **F**) with values 10 through 15 respectively. A binary literal consists of the prefix **0b** or **0B** followed by a sequence of the digits **0** or **1**.
- 5 An *unprefixed octal literal* begins with the digit **0** optionally followed by a sequence of the digits **0** through **7** only. Use of an unprefixed octal literal with digits other than **0** is an obsolescent feature.

Semantics

- 6 The value of a decimal literal is computed base 10; that of an octal literal, base 8; that of a hexadecimal literal, base 16; that of a binary literal, base 2. The lexically first digit is the most significant.
- 7 The type of an integer literal is the first of the corresponding Table 6.2 in which its value can be represented.
- 8 If an integer literal that does not have suffixes **wb**, **WB**, **uwb**, or **UWB** cannot be represented by any type in its list, it may have an extended integer type, if the extended integer type can represent its value. If all the types in the list for the literal are signed, the extended integer type shall be signed. If all the types in the list for the literal are unsigned, the extended integer type shall be unsigned. If the list contains both signed and unsigned types, the extended integer type may be signed or unsigned. If an integer literal cannot be represented by any type in its list and has no extended integer type, then the integer literal has no type.
- 9 EXAMPLE 2 The **wb** suffix results in an **_BitInt** that includes space for the sign bit even if the value of the literal is positive or was specified in binary, octal, or hexadecimal notation.

```
-3wb /* Yields a _BitInt(3) that is then arithmetically negated;
      two value bits, one sign bit */
-0x3wb /* Yields a _BitInt(3) that is then arithmetically negated;
```

Table 6.2 — Relationship between constants, suffixes, and types

Suffix	Decimal Literal	Octal, Hexadecimal or Binary Literal
none	int long int long long int	int unsigned int long int unsigned long int long long int unsigned long long int
u or U	unsigned int unsigned long int unsigned long long int	unsigned int unsigned long int unsigned long long int
l or L	long int long long int	long int unsigned long int long long int unsigned long long int
Both u or U and l or L	unsigned long int unsigned long long int	unsigned long int unsigned long long int
ll or LL	long long int	long long int unsigned long long int
Both u or U and ll or LL	unsigned long long int	unsigned long long int
wb or WB	_BitInt(N) where the width N is the smallest N greater than 1 which can accommodate the value and the sign bit.	_BitInt(N) where the width N is the smallest N greater than 1 which can accommodate the value and the sign bit.
Both u or U and wb or WB	unsigned _BitInt(N) where the width N is the smallest N greater than 0 which can accommodate the value.	unsigned _BitInt(N) where the width N is the smallest N greater than 0 which can accommodate the value.

```

        two value bits, one sign bit */
3wb    /* Yields a _BitInt(3); two value bits, one sign bit */
3uwb   /* Yields an unsigned _BitInt(2) */
-3uwb  /* Yields an unsigned _BitInt(2) that is then arithmetically
        negated, resulting in wraparound */

```

Forward references: preprocessing numbers (6.4.9), numeric conversion functions (7.24.2).

6.4.5.3 Floating literals

Syntax

1 *floating-literal*:

decimal-floating-literal
hexadecimal-floating-literal

decimal-floating-literal:

fractional-literal *exponent-part*_{opt} *floating-suffix*_{opt}
digit-sequence *exponent-part* *floating-suffix*_{opt}

hexadecimal-floating-literal:

hexadecimal-prefix *hexadecimal-fractional-literal*

binary-exponent-part floating-suffix_{opt}
hexadecimal-prefix hexadecimal-digit-sequence
binary-exponent-part floating-suffix_{opt}

fractional-literal:

digit-sequence_{opt} . digit-sequence
digit-sequence .

exponent-part:

e *sign_{opt} digit-sequence*
E *sign_{opt} digit-sequence*

sign: one of

+ -

digit-sequence:

digit
digit-sequence ' _{opt} digit

hexadecimal-fractional-literal:

hexadecimal-digit-sequence_{opt} . hexadecimal-digit-sequence
hexadecimal-digit-sequence .

binary-exponent-part:

p *sign_{opt} digit-sequence*
P *sign_{opt} digit-sequence*

real-floating-suffix: one of

f l F L df dd dL DF DD DL

floating-suffix:

real-floating-suffix complex-suffix_{opt}
complex-suffix real-floating-suffix_{opt}

complex-suffix: one of

i I j J

Constraints

- 2 A floating suffix **df**, **dd**, **dL**, **DF**, **DD**, or **DL** shall not be used in a hexadecimal floating literal. A floating suffix shall not designate a type that the implementation does not provide.

Description

- 3 A floating literal has a *significant part* that can be followed by an *exponent part* and a suffix that specifies its type. The components of the significant part can include a digit sequence representing

the whole-number part, followed by a period (.), followed by a digit sequence representing the fraction part. Digit separators (6.4.5.2) are ignored when determining the value of the literal. The components of the exponent part are an **e**, **E**, **p**, or **P** followed by an exponent consisting of an optionally signed digit sequence. Either the whole-number part or the fraction part has to be present; for decimal floating literals, either the period or the exponent part has to be present.

Semantics

- 4 The significand part is interpreted as a (decimal or hexadecimal) rational number; the digit sequence in the exponent part is interpreted as a decimal integer. For decimal floating literals, the exponent indicates the power of 10 by which the significand part is to be scaled. For hexadecimal floating literals, the exponent indicates the power of 2 by which the significand part is to be scaled. For decimal floating literals, and also for hexadecimal floating literals when **FLT_RADIX** is not a power of 2, the result is either the nearest representable value, or the larger or smaller representable value immediately adjacent to the nearest representable value, chosen in an implementation-defined manner. For hexadecimal floating literals when **FLT_RADIX** is a power of 2, the result is correctly rounded.
- 5 The presence of a complex suffix indicates that the type of the literal is the complex type that corresponds to the type of the literal where the complex suffix is removed. The described floating value of such a complex literal designates the imaginary part of an arithmetic constant expression of the indicated complex type, the real part is (positive) zero.⁶⁶⁾ An unsuffixed floating literal has type **double**. If suffixed by a real floating suffix it has a type according to Table 6.3.

Table 6.3 — Suffixes for floating literals

Suffix	Type
f , F	float
l , L	long double
df , DF	_Decimal32
dd , DD	_Decimal64
dll , DL	_Decimal128

- 6 The values of floating literals may be represented in greater range and precision than that required by the type (determined by the suffix); the types are not changed thereby. See 5.3.5.3.3 regarding evaluation formats.⁶⁷⁾
- 7 Floating literals of decimal floating type that have the same numerical value but different quantum exponents have distinguishable internal representations. The value shall be correctly rounded as specified in ISO/IEC 60559. The coefficient c and the quantum exponent q of a finite converted decimal floating-point number (see 5.3.5.3.4) are determined as follows:

- q is set to the value of $sign_{opt} digit-sequence$ in the exponent part, if any, or to 0, otherwise.
- If there is a fractional literal, q is decreased by the number of digits to the right of the period and the period is removed to form a digit sequence.
- c is set to the value of the digit sequence (after any period has been removed).
- Rounding required because of insufficient precision or range in the type of the result will round c to the full precision available in the type, and will adjust q accordingly within the limits of the type, provided the rounding does not yield an infinity (in which case the result is an appropriately signed internal representation of infinity). If the full precision of the type

⁶⁶⁾Implementations that define the macro **__STDC_NO_COMPLEX__** may not support complex types. For these implementations, a literal with a complex suffix presents a constraint violation.

⁶⁷⁾Hexadecimal floating literals can be used to obtain exact values in the semantic type that are independent of the evaluation format. Casts produce values in the semantic type, though depend on the rounding mode and can raise the inexact floating-point exception.

would require q to be smaller than the minimum for the type, then q is pinned at the minimum and c is adjusted through the subnormal range accordingly, perhaps to zero.

- 8 Floating literals are converted to internal format as if at translation-time. The conversion of a floating literal shall not raise an exceptional condition or a floating-point exception at execution time. All floating literals of the same source form⁶⁸⁾ shall convert to the same internal format and, provided they are subject to the same translation-time rounding direction (either the default or a constant rounding mode set by an **FENV_ROUND** or **FENV_DEC_ROUND** pragma), to the same value.
- 9 EXAMPLE Following are floating literals of type **_Decimal64** and their values as triples (s, c, q) . For **_Decimal64**, the precision (maximum coefficient length) is 16 and the quantum exponent range is $-398 \leq q \leq 369$.

0.dd	(+1, 0, 0)
0.00dd	(+1, 0, -2)
123.dd	(+1, 123, 0)
1.23E3dd	(+1, 123, 1)
1.23E+3dd	(+1, 123, 1)
12.3E+7dd	(+1, 123, 6)
12.0dd	(+1, 120, -1)
12.3dd	(+1, 123, -1)
0.00123dd	(+1, 123, -5)
1.23E-12dd	(+1, 123, -14)
1234.5E-4dd	(+1, 12345, -5)
0E+7dd	(+1, 0, 7)
12345678901234567890.dd	(+1, 1234567890123457, 4) assuming default rounding and DEC_EVAL_METHOD is 0 or 1 ⁶⁹⁾
1234E-400dd	(+1, 12, -398) assuming default rounding and DEC_EVAL_METHOD is 0 or 1
1234E-402dd	(+1, 0, -398) assuming default rounding and DEC_EVAL_METHOD is 0 or 1
1000.dd	(+1, 1000, 0)
.0001dd	(+1, 1, -4)
1000.e0dd	(+1, 1000, 0)
.0001e0dd	(+1, 1, -4)
1000.0dd	(+1, 10000, -1)
0.0001dd	(+1, 1, -4)
1000.00dd	(+1, 100 000, -2)
00.0001dd	(+1, 1, -4)
001000.dd	(+1, 1000, 0)
001000.0dd	(+1, 10000, -1)
001000.00dd	(+1, 100 000, -2)
00.00dd	(+1, 0, -2)
00.dd	(+1, 0, 0)
.00dd	(+1, 0, -2)
00.00e-5dd	(+1, 0, -7)
00.e-5dd	(+1, 0, -5)
.00e-5dd	(+1, 0, -7)

Recommended practice

- 10 The implementation should produce a diagnostic message if a hexadecimal floating literal cannot be represented exactly in its evaluation format; the implementation should then proceed with the translation of the program.
- 11 The translation-time conversion of floating literals should match the execution-time conversion of character strings by library functions, such as **strtod**, given matching inputs suitable for both

⁶⁸⁾1.23, 1.230, 123e-2, 123e-02, and 1.23L are all different source forms and thus can convert to a different internal format and value (though they can use the same internal format and value).

⁶⁹⁾That is, assuming the default translation rounding-direction mode is not changed by an **FENV_DEC_ROUND** pragma (7.6.4).

conversions, the same result format, and default execution-time rounding.⁷⁰⁾

- 12 NOTE Floating literals do not include a sign and are arithmetically negated by the unary `-` operator (6.5.4.4) which arithmetically negates the rounded value of the literal. In contrast, the numeric conversion functions in the `strto` family (7.24.2.6, 7.24.2.7) can include the sign as part of the input value and convert and round the arithmetically negated input; implementations conforming to Annex F have this behavior. Negating before rounding and negating after rounding can yield different results, depending on the rounding direction and whether the results are correctly rounded. For example, the results are the same when both are correctly rounded using rounding to nearest or rounding toward zero, but the results are different when they are inexact and correctly rounded using rounding toward positive infinity or rounding toward negative infinity.

Conversions yielding exact results are not affected by the order of negating and rounding. For types with radix 10, decimal floating literals expressed within the precision and range of the evaluation format convert exactly. For types whose radix is a power of 2, hexadecimal floating literals expressed within the precision and range of the evaluation format convert exactly.

Forward references: preprocessing numbers (6.4.9), numeric conversion functions (7.24.2), the `strto` function family (7.24.2.6, 7.24.2.7).

6.4.5.4 Enumeration constants

Syntax

- 1 *enumeration-constant:*
 identifier

Semantics

- 2 An identifier declared as an enumeration constant for an enumeration without a fixed underlying type has either type **int** or the enumerated type, as defined in 6.7.3.3. An identifier declared as an enumeration constant for an enumeration with a fixed underlying type has the associated enumerated type.

An enumeration constant can be used in an expression (or constant expression) wherever a value of an integer type may be used.

Forward references: enumeration specifiers (6.7.3.3).

6.4.5.5 Character literals

Syntax

- 1 *character-literal:*
 *encoding-prefix*_{opt} ' *c-char-sequence* '

encoding-prefix: one of

u8 **u** **U** **L**

c-char-sequence:

c-char
c-char-sequence *c-char*

c-char:

any member of the source character set except
 the single-quote ' , backslash \, or new-line character
escape-sequence

escape-sequence:

simple-escape-sequence

⁷⁰⁾The specification for the library functions recommends more accurate conversion than required for floating literals (see 7.24.2.6).

octal-escape-sequence
hexadecimal-escape-sequence
universal-character-name

simple-escape-sequence: one of

`\' \" \? \\`
`\a \b \f \n \r \t \v`

octal-escape-sequence:

`\ octal-digit`
`\ octal-digit octal-digit`
`\ octal-digit octal-digit octal-digit`
`\o{ simple-octal-digit-sequence }`

simple-octal-digit-sequence:

`octal-digit`
`simple-octal-digit-sequence octal-digit`

hexadecimal-escape-sequence:

`\x simple-hexadecimal-digit-sequence`
`\x{ simple-hexadecimal-digit-sequence }`

Description

- 2 An *ordinary character literal* is a sequence of one or more multibyte characters enclosed in single-quotes, as in `'x'`. A *UTF-8 character literal* is the same, except prefixed by **u8**. Together ordinary and UTF-8 character literal are *narrow character literals*. A **wchar_t** *character literal* is prefixed by the letter **L**. A *UTF-16 character literal* is prefixed by the letter **u**. A *UTF-32 character literal* is prefixed by the letter **U**. Collectively, **wchar_t**, UTF-16, and UTF-32 character literals are called *wide character literals*. With a few exceptions detailed later, the elements of the sequence are any members of the source character set; they are mapped in an implementation-defined manner to members of the execution character set.
- 3 The single-quote `'`, the double-quote `"`, the question-mark `?`, the backslash `\`, and arbitrary integer values are representable according to Table 6.6:

Table 6.6 — Escape sequences

single quote <code>'</code>	<code>\'</code>
double quote <code>"</code>	<code>\"</code>
question mark <code>?</code>	<code>\?</code>
backslash <code>\</code>	<code>\\</code>
octal character	<code>\octal digits</code>
hexadecimal character	<code>\x hexadecimal digits</code>

- 4 The double-quote `"` and question-mark `?` are representable either by themselves or by the escape sequences `\"` and `\?`, respectively, but the single-quote `'` and the backslash `\` shall be represented, respectively, by the escape sequences `\'` and `\\`.
- 5 The octal digits that follow the backslash in an octal escape sequence are taken to be part of the construction of a single character for a narrow character literal or of a single wide character for a wide character literal. The numerical value of the octal integer so formed specifies the value of the desired character or wide character.
- 6 The hexadecimal digits that follow the backslash and the letter `x` in a hexadecimal escape sequence are taken to be part of the construction of a single character for a narrow character literal or of a single wide character for a wide character literal. The numerical value of the hexadecimal integer so formed specifies the value of the desired character or wide character.
- 7 Each octal or hexadecimal escape sequence is the longest sequence of characters that can constitute the escape sequence.
- 8 In addition, characters not in the basic character set are representable by universal character names and certain non-graphic characters are representable by escape sequences consisting of the backslash `\` followed by a lowercase letter: `\a`, `\b`, `\f`, `\n`, `\r`, `\t`, and `\v`.⁷¹⁾

Constraints

- 9 The value of an octal or hexadecimal escape sequence shall be in the range of representable values for the corresponding type, as dictated by Table 6.7:

Table 6.7 — Types provided by prefixes

Prefix	Corresponding Type
none	unsigned char
u8	char8_t
L	the unsigned type corresponding to wchar_t
u	char16_t
U	char32_t

- 10 A UTF-8, UTF-16, or UTF-32 character literal shall not contain more than one character.⁷²⁾ The value shall be representable with a single UTF-8, UTF-16, or UTF-32 code unit, respectively.

Semantics

- 11 An ordinary character literal has type **int**. The value of an ordinary character literal containing a single character that maps to a single value in the literal encoding (6.2.9) is the numerical value of the representation of the mapped character in the literal encoding interpreted as an integer. The value of an ordinary character literal containing more than one character (e.g. `'ab'`), or containing a character or escape sequence that does not map to a single value in the literal encoding, is implementation-defined. If an ordinary character literal contains a single character or escape sequence, its value is the one that results when an object with type **char** whose value is that of the single character or escape sequence is converted to type **int**.

⁷¹⁾The semantics of these characters were discussed in 5.3.3. If any other character follows a backslash, the result is not a token and a diagnostic is required. See "future language directions" (6.11.4).

⁷²⁾For example `u8'ab'` violates this constraint.

- 12 A UTF-8 character literal has type **char8_t**. If the UTF-8 character literal is not produced through a hexadecimal or octal escape sequence, the value of a UTF-8 character literal is equal to its ISO/IEC 10646 code point value, provided that the code point value can be encoded as a single UTF-8 code unit. Otherwise, the value of the UTF-8 character literal is the numeric value specified in the hexadecimal or octal escape sequence.
- 13 A UTF-16 character literal has type **char16_t** which is an unsigned integer type defined in the `<uchar.h>` header. If the UTF-16 character literal is not produced through a hexadecimal or octal escape sequence, the value of a UTF-16 character literal is equal to its ISO/IEC 10646 code point value, provided that the code point value can be encoded as a single UTF-16 code unit. Otherwise, the value of the UTF-16 character literal is the numeric value specified in the hexadecimal or octal escape sequence.
- 14 A UTF-32 character literal has type **char32_t** which is an unsigned integer type defined in the `<uchar.h>` header. If the UTF-32 character literal is not produced through a hexadecimal or octal escape sequence, the value of a UTF-32 character literal is equal to its ISO/IEC 10646 code point value, provided that the code point value can be encoded as a single UTF-32 code unit. Otherwise, the value of the UTF-32 character literal is the numeric value specified in the hexadecimal or octal escape sequence.
- 15 A **wchar_t** character literal has type **wchar_t**, an integer type defined in the `<stddef.h>` header. The value of a **wchar_t** character literal containing a single multibyte character that maps to a single member of the extended execution character set is the wide character corresponding to that multibyte character in the implementation-defined wide literal encoding (6.2.9). The value of a **wchar_t** character literal containing more than one multibyte character or a single multibyte character that maps to multiple members of the extended execution character set, or containing a multibyte character or escape sequence not represented in the extended execution character set, is implementation-defined.
- 16 EXAMPLE 1 The construction `'\0'` is commonly used to represent the null character.
- 17 EXAMPLE 2 Implementations that use eight bits for objects that have type **char** can furnish certain values in an variety of ways. In an implementation in which type **char** has the same range of values as **signed char**, the ordinary character literal `'\xFF'` has the value -1 ; if type **char** has the same range of values as **unsigned char**, the ordinary character literal `'\xFF'` has the value $+255$.
- 18 EXAMPLE 3 Even if eight bits are used for objects that have type **char**, the construction `'\x123'` specifies an ordinary character literal containing only one character, since a hexadecimal escape sequence is terminated only by a non-hexadecimal character. To specify an ordinary character literal containing the two characters whose values are `'\x12'` and `'3'`, the construction `'\0223'` can be used, since an octal escape sequence is terminated after three octal digits. (The value of this two-character ordinary character literal is implementation-defined.)
- 19 EXAMPLE 4 Even if 12 or more bits are used for objects that have type **wchar_t**, the construction `L'\1234'` specifies the implementation-defined value that results from the combination of the values `0123` and `'4'`.

Forward references: common definitions `<stddef.h>` (7.21), the **mbtowc** function (7.24.8.3), Unicode utilities `<uchar.h>` (7.31).

6.4.5.6 Predefined constants

Syntax

- 1 *predefined-constant*:

false
true
nullptr

Description

- 2 Some keywords represent constants of a specific value and type.
- 3 The keywords **false** and **true** are constants of type **bool** with a value of 0 for **false** and 1 for **true**.⁷³⁾

⁷³⁾The constants **false** and **true** promote to type **int**, see 6.3.2.1. When used for arithmetic, in translation phase 4 (5.2.1.2),

- 4 The keyword **nullptr** represents a null pointer constant. Details of its type are described in 7.21.3.

6.4.6 String literals

Syntax

- 1 *string-literal*:
- $$\text{encoding-prefix}_{\text{opt}} \text{ " } s\text{-char-sequence}_{\text{opt}} \text{ "}$$
- s-char-sequence*:
- $$s\text{-char}$$
- $$s\text{-char-sequence } s\text{-char}$$
- s-char*:
- any member of the source character set except
the double-quote " , backslash \ , or new-line character
- escape-sequence*

Constraints

- 2 If a sequence of adjacent string literal tokens includes prefixed string literal tokens, the prefixed tokens shall all have the same prefix.

Description

- 3 An *ordinary string literal* is a sequence of zero or more multibyte characters enclosed in double-quotes, as in "xyz". A *UTF-8 string literal* is the same, except prefixed by **u8**. A **wchar_t** *string literal* is the same, except prefixed by **L**. A *UTF-16 string literal* is the same, except prefixed by **U**. A *UTF-32 string literal* is the same, except prefixed by **U**. Collectively, **wchar_t**, UTF-16, and UTF-32 string literals are called *wide string literals*.
- 4 The same considerations apply to each element of the sequence in a string literal as if it were in an narrow character literal (for a character or UTF-8 string literal) or a wide character literal (for a wide string literal), except that the single-quote ' is representable either by itself or by the escape sequence \', but the double-quote " shall be represented by the escape sequence \".

Semantics

- 5 In translation phase 6 (5.2.1.2), the multibyte character sequences specified by any sequence of adjacent character and identically-prefixed string literal tokens are concatenated into a single multibyte character sequence. If any of the tokens has an encoding prefix, the resulting multibyte character sequence is treated as having the same prefix; otherwise, it is treated as an ordinary string literal.
- 6 In translation phase 7 (5.2.1.2), a byte or code of value zero is appended to each multibyte character sequence that results from a string literal or literals.⁷⁴⁾ The multibyte character sequence is then used to initialize an array of static storage duration and length just sufficient to contain the sequence. For ordinary string literals, the array elements have type **char**, and are initialized with the individual bytes of the multibyte character sequence corresponding to the literal encoding (6.2.9). For UTF-8 string literals, the array elements have type **char8_t**, and are initialized with the characters of the multibyte character sequence, as encoded in UTF-8. For wide string literals prefixed by the letter **L**, the array elements have type **wchar_t** and are initialized with the sequence of wide characters corresponding to the wide literal encoding. For wide string literals prefixed by the letter **u** or **U**, the array elements have type **char16_t** or **char32_t**, respectively, and are initialized sequence of wide characters corresponding to UTF-16 and UTF-32 encoded text, respectively. The value of a string literal containing a multibyte character or escape sequence not represented in the execution

they are signed values and the result of such arithmetic is consistent with the results of later translation phases.

⁷⁴⁾A string literal may not be a string (see 7.1.1), because a null character can be embedded in it by a \0 escape sequence.

character set is implementation-defined. Any hexadecimal escape sequence or octal escape sequence specified in a **u8**, **u**, or **U** string specifies a single **char8_t**, **char16_t**, or **char32_t** value and can result in the full character sequence not being valid UTF-8, UTF-16, or UTF-32.

- 7 It is unspecified whether these arrays are distinct provided their elements have the appropriate values. If the program attempts to modify such an array, the behavior is undefined.
- 8 **EXAMPLE 1** This pair of adjacent ordinary string literals

```
"\x12" "3"
```

produces a single ordinary string literal containing the two characters whose values are `'\x12'` and `'3'`, because escape sequences are converted into single members of the execution character set just prior to adjacent string literal concatenation.

- 9 **EXAMPLE 2** Each of the sequences of adjacent string literal tokens

```
"a" "b" L"c"
"a" L"b" "c"
L"a" "b" L"c"
L"a" L"b" L"c"
```

is equivalent to the string literal

```
L"abc"
```

Likewise, each of the sequences

```
"a" "b" u"c"
"a" u"b" "c"
u"a" "b" u"c"
u"a" u"b" u"c"
```

is equivalent to

```
u"abc"
```

Forward references: common definitions `<stddef.h>` (7.21), the **mbstowcs** function (7.24.9.2), Unicode utilities `<uchar.h>` (7.31).

6.4.7 Punctuators

Syntax

- 1 *punctuator*: one of
- ```
[] () { } . ->
++ -- & * + - ~ !
/ % << >> < > <= >= == != ^ | && ||
? : :: ; ...
= *= /= %= += -= <<= >>= &= ^= |=
, # ##
<: :> <% %> %: %:::
```

### Semantics

- 2 A punctuator is a symbol that has independent syntactic and semantic significance. Depending on context, it may specify an operation to be performed (which in turn may yield a value or a function designator, produce a side effect, or some combination thereof) in which case it is known as an *operator* (other forms of operator also exist in some contexts). An *operand* is an entity on which an operator acts.



- 3 In all aspects of the language, the six tokens<sup>75)</sup>

`<:    :>   <%   %>   %:   %: %:`

behave, respectively, the same as the six tokens

`[   ]   {   }   #   ##`

except for their spelling.<sup>76)</sup>

**Forward references:** expressions (6.5.1), declarations (6.7), preprocessing directives (6.10), statements (6.8).

### 6.4.8 Header names

#### Syntax

- 1 *header-name*:

`< h-char-sequence >`  
`" q-char-sequence "`

*h-char-sequence*:

`h-char`  
`h-char-sequence h-char`

*h-char*:

any member of the source character set except  
the new-line character and `>`

*q-char-sequence*:

`q-char`  
`q-char-sequence q-char`

*q-char*:

any member of the source character set except  
the new-line character and `"`

#### Semantics

- 2 The sequences in both forms of header names are mapped in an implementation-defined manner to headers or external source file names as specified in 6.10.3.
- 3 If the characters `'`, `\`, `"`, `/**`, or `/*` occur in the sequence between the `<` and `>` delimiters, the behavior is undefined. Similarly, if the characters `'`, `\`, `/**`, or `/*` occur in the sequence between the `"` delimiters, the behavior is undefined.<sup>77)</sup>

Header name preprocessing tokens are recognized only within **#include** and **#embed** preprocessing directives, in **\_\_has\_include** and **\_\_has\_embed** expressions, as well as in implementation-defined locations within **#pragma** directives.<sup>78)</sup>

- 4 EXAMPLE The following sequence of characters:

<sup>75)</sup>These tokens are sometimes called “digraphs”.

<sup>76)</sup>Thus `[` and `<` behave differently when “stringized” (see 6.10.5.3), but can otherwise be freely interchanged.

<sup>77)</sup>Thus, sequences of characters that resemble escape sequences cause undefined behavior.

<sup>78)</sup>For an example of a header name preprocessing token used in a **#pragma** directive, see 6.10.11.

```
0x3<1/a.h>1e2
#include <1/a.h>
#define const.member@$
```

forms the following sequence of preprocessing tokens (with each individual preprocessing token delimited by a { on the left and a } on the right).

```
{0x3}{<}{1}{/}{a}{.}{h}{>}{1e2}
{#}{include} {<1/a.h>}
{#}{define} {const}{.}{member}{@}{$}
```

**Forward references:** source file inclusion (6.10.3).

## 6.4.9 Preprocessing numbers

### Syntax

- 1 *pp-number*:
  - digit*
  - .* *digit*
  - pp-number* *identifier-continue*
  - pp-number* ' *digit*
  - pp-number* ' *nondigit*
  - pp-number* **e** *sign*
  - pp-number* **E** *sign*
  - pp-number* **p** *sign*
  - pp-number* **P** *sign*
  - pp-number* **.**

### Description

- 2 A preprocessing number begins with a digit optionally preceded by a period (.) and may be followed by valid identifier characters and the character sequences **e+**, **e-**, **E+**, **E-**, **p+**, **p-**, **P+**, or **P-**.
- 3 Preprocessing number tokens lexically include all floating and integer literals.

### Semantics

- 4 A preprocessing number does not have type or a value; it acquires both after a successful conversion (as part of translation phase 7 (5.2.1.2)) to a floating or integer literal. This notwithstanding for the evaluation of expressions within conditional source inclusion (6.10.2) and to determine the **limit** parameter for binary resource inclusion (6.10.4), preprocessing numbers which have the form of an integer literal are interpreted as such. For the determination of a line number in a **#line** directive (6.10.6), digit sequences that match the requirements for a preprocessing number are interpreted as numbers as well and the interpretation is of a decimal integer, even if the leading digit is 0.

## 6.4.10 Comments

- 1 Except within a character literal, a string literal, or a comment, the characters /\* introduce a comment. The contents of such a comment are examined only to identify multibyte characters and to find the characters \*/ that terminate it.<sup>79)</sup>
- 2 Except within a character literal, a string literal, or a comment, the characters // introduce a comment that includes all multibyte characters up to, but not including, the next new-line character. The contents of such a comment are examined only to identify multibyte characters and to find the terminating new-line character.
- 3 EXAMPLE

```
"a//b" // four-character ordinary string literal
#include "//e" // undefined behavior
```

<sup>79)</sup>Thus, /\* ... \*/ comments do not nest.

```
// */ // comment, not syntax error
f = g/**/h; // equivalent to f = g / h;
//\
i(); // part of a two-line comment
/\
/ j(); // part of a two-line comment
#define glue(x,y) x##y
glue(/,/) k(); // syntax error, not comment
/***/ l(); // equivalent to l();
m = n/**/o
 + p; // equivalent to m = n + p;
```

## 6.5 Expressions

### 6.5.1 General

- 1 An *expression* is a sequence of operators and operands that specifies computation of a value, or that designates an object or a function, or that generates side effects, or that performs a combination thereof. The value computations of the operands of an operator are sequenced before the value computation of the result of the operator.
- 2 If a side effect on a scalar object is unsequenced relative to either a different side effect on the same scalar object or a value computation using the value of the same scalar object, the behavior is undefined. If there are multiple allowable orderings of the subexpressions of an expression, the behavior is undefined if such an unsequenced side effect occurs in any of the orderings.<sup>80)</sup>
- 3 The grouping of operators and operands is indicated by the syntax.<sup>81)</sup> Except as specified later, side effects and value computations of subexpressions are unsequenced.<sup>82)</sup>
- 4 Some operators (the unary operator `~`, and the binary operators `<<`, `>>`, `&`, `^`, and `|`, collectively described as *bitwise operators*) are required to have operands that have integer type. These operators yield values that depend on the internal representations of integers, and have implementation-defined and undefined aspects for signed types.
- 5 If an *exceptional condition* occurs during the evaluation of an expression (that is, if the result is not mathematically defined or not in the range of representable values for its type), the behavior is undefined.
- 6 The *effective type* of an object that is not a byte array, for an access to its stored value, is the declared type of the object.<sup>83)</sup> If a value is stored into a byte array through an lvalue having a type that is not a byte type, then the type of the lvalue becomes the effective type of the object for that access and for subsequent accesses that do not modify the stored value. If a value is copied into a byte array using `memcpy` or `memmove`, or is copied as an array of byte type, then the effective type of the modified object for that access and for subsequent accesses that do not modify the value is the effective type of the object from which the value is copied, if it has one. For all other accesses to a byte array, the effective type of the object is simply the type of the lvalue used for the access.<sup>84)</sup>
- 7 Following a byte-level copy operation, an object of atomic type (6.7.3.5) has an indeterminate representation and a copy of a **FILE** (7.23.3) object is not required to serve in place of the original.
- 8 An object shall have its stored value accessed only by an lvalue expression that has one of the following types:<sup>85)</sup>
  - a type compatible with the effective type of the object,

<sup>80)</sup>This paragraph renders undefined statement expressions such as

```
i = ++i + 1;
a[i++] = i;
```

while allowing

```
i = i + 1;
a[i] = i;
```

<sup>81)</sup>The syntax specifies the precedence of operators in the evaluation of an expression, which is the same as the order of the major subclauses of this subclause, highest precedence first. Thus, for example, the expressions allowed as the operands of the binary `+` operator (6.5.7) are those expressions defined in 6.5.2 through 6.5.7. The exceptions are cast expressions (6.5.5) as operands of unary operators (6.5.4), and an operand contained between any of the following pairs of operators: grouping parentheses `()` (6.5.2), generic selection parentheses `()` (6.5.2.1), subscripting brackets `code[]` (6.5.3.2), function-call parentheses `()` (6.5.3.3), and the conditional operator `?:` (6.5.16).

Within each major subclause, the operators have the same precedence. Left- or right-associativity is indicated in each subclause by the syntax for the expressions discussed therein.

<sup>82)</sup>In an expression that is evaluated more than once during the execution of a program, unsequenced and indeterminately sequenced evaluations of its subexpressions can be performed inconsistently in different evaluations.

<sup>83)</sup>Allocated objects have no declared type.

<sup>84)</sup>The object needs to have valid alignment and size for the effective type to be accessed and the type of the lvalue needs to have at least the qualifications of the declared type of the object, if any.

<sup>85)</sup>The intent of this list is to specify those circumstances in which an object can or cannot be aliased.

- a qualified version of a type compatible with the effective type of the object,
  - the signed or unsigned type corresponding to a type compatible with the underlying type of the effective type of the object,
  - the signed or unsigned type corresponding to a type compatible with a qualified version of the underlying type of the effective type of the object,
  - a byte type, or
  - an aggregate or union type that includes one of the aforementioned types among its members (including, recursively, a member of a subaggregate or contained union).
- 9 A floating expression may be *contracted*, that is, evaluated as though it were a single operation, thereby omitting rounding errors implied by the source code and the expression evaluation method.<sup>86)</sup>

The **FP\_CONTRACT** pragma in `<math.h>` provides a way to disallow contracted expressions. Otherwise, whether and how expressions are contracted is implementation-defined.<sup>87)</sup>

- 10 Operators involving decimal floating types are evaluated according to the semantics of ISO/IEC 60559, including production of results with the preferred quantum exponent as specified in ISO/IEC 60559.

**Forward references:** the **FP\_CONTRACT** pragma (7.12.3), copying functions (7.27.2).

## 6.5.2 Primary expressions

### Syntax

- 1 *primary-expression:*
- identifier*
  - constant*
  - string-literal*
  - ( expression )*
  - generic-selection*

### Constraints

- 2 The identifier in an identifier primary expression shall have a visible declaration as an ordinary identifier that declares an object or a function.<sup>88)</sup>

### Semantics

- 3 An identifier primary expression designating an object is an lvalue. An identifier primary expression designating a function is a function designator.
- 4 A constant is a primary expression. Its type depends on its form and value, as detailed in 6.4.5.
- 5 A string literal is a primary expression. It is an lvalue with type as detailed in 6.4.6.
- 6 A *parenthesized expression* is a primary expression. Its type, value, and semantics are identical to those of the unparenthesized expression.
- 7 A generic selection is a primary expression. Its type, value, and semantics depend on the selected generic association, as detailed in the following subclause.

**Forward references:** declarations (6.7).

<sup>86)</sup>The intermediate operations in the contracted expression are evaluated as if to infinite range and precision, while the final operation is rounded to the format determined by the expression evaluation method. A contracted expression can also omit the raising of floating-point exceptions.

<sup>87)</sup>This license is specifically intended to allow implementations to exploit fast machine instructions that combine multiple C operators. As contractions potentially undermine predictability, and can even decrease accuracy for containing expressions, their use needs to be well-defined and clearly documented.

<sup>88)</sup>An identifier designating an enumeration constant is a primary expression through the constant production, not the identifier production.

### 6.5.2.1 Generic selection

#### Syntax

- 1 *generic-selection*:
- ```

    _Generic ( generic-controlling-operand , generic-assoc-list )
generic-controlling-operand:
    assignment-expression
    type-name
generic-assoc-list:
    generic-association
    generic-assoc-list , generic-association
generic-association:
    type-name : assignment-expression
    default : assignment-expression

```

Constraints

- 2 A generic selection shall have no more than one **default** generic association. The type name in a generic association shall specify a type other than a variably modified type. No two generic associations in the same generic selection shall specify compatible types. If the generic controlling operand is an assignment expression, the *controlling type* of the generic selection expression is the type of the assignment expression as if it had undergone an lvalue conversion,⁸⁹⁾ array to pointer conversion, or function to pointer conversion. Otherwise, the controlling type of the generic selection expression is the type designated by the type name. The controlling type shall be compatible with at most one of the types named in the generic association list. If a generic selection has no **default** generic association, its controlling type shall be compatible with exactly one of the types named in its generic association list.

Semantics

- 3 The generic controlling operand is not evaluated. If a generic selection has a generic association with a type name that is compatible with the controlling type, then the result expression of the generic selection is the expression in that generic association. Otherwise, the result expression of the generic selection is the expression in the **default** generic association. None of the expressions from any other generic association of the generic selection is evaluated.
- 4 The type and value of a generic selection are identical to those of its result expression. It is an lvalue, a function designator, or a void expression if its result expression is, respectively, an lvalue, a function designator, or a void expression.
- 5 EXAMPLE 1 A **cbrt** type-generic macro can be implemented as follows:

```

#define cbrt(X) _Generic((X),
                        long double: cbrtL,
                        default: cbrt,
                        float: cbrtf
                        )(X)

```

7.28 shows how such a macro can be implemented with the required rounding properties.

- 6 EXAMPLE 2 The following two generic selection expressions select different associations because the assignment expression operand undergoes lvalue conversion while the type name operand is unchanged:

```

int func(const int i) {
    return _Generic(i,
        int : 0, // 'int' is selected
        const int : 1,
        default : 2) +
    _Generic(sizeof(i),
        int : 0,

```

⁸⁹⁾An lvalue conversion drops type qualifiers.

```

    const int : 1, // 'const int' is selected
    default : 2);
}

```

Thus, this function returns 1.

6.5.3 Postfix operators

6.5.3.1 General

Syntax

- 1 *postfix-expression*:
 - primary-expression*
 - postfix-expression* [*expression*]
 - postfix-expression* (*argument-expression-list*_{opt})
 - postfix-expression* . *identifier*
 - postfix-expression* -> *identifier*
 - postfix-expression* ++
 - postfix-expression* --
 - compound-literal*
- argument-expression-list*:
- assignment-expression*
 - argument-expression-list* , *assignment-expression*

6.5.3.2 Array subscripting

Constraints

- 1 One of the expressions shall have type “pointer to complete object *type*”, the other expression shall have integer type, and the result has type “*type*”.

Semantics

- 2 A postfix expression followed by an expression in square brackets [] is a subscripted designation of an element of an array object. The definition of the subscript operator [] is that **E1[E2]** is identical to **(*((E1)+(E2)))**. Because of the conversion rules that apply to the binary + operator, if **E1** is an array object (equivalently, a pointer to the initial element of an array object) and **E2** is an integer, **E1[E2]** designates the **E2**-th element of **E1** (counting from zero).
- 3 Successive subscript operators designate an element of a multidimensional array object. If **E** is an *n*-dimensional array ($n \geq 2$) with dimensions $i \times j \times \dots \times k$, then **E** (used as other than an lvalue) is converted to a pointer to an $(n - 1)$ -dimensional array with dimensions $j \times \dots \times k$. If the unary * operator is applied to this pointer explicitly, or implicitly as a result of subscripting, the result is the referenced $(n - 1)$ -dimensional array, which itself is converted into a pointer if used as other than an lvalue. It follows from this that arrays are stored in row-major order (last subscript varies fastest).
- 4 EXAMPLE The following snippet has an array object defined by the declaration:

```
int x[3][5];
```

Here **x** is a 3×5 array of objects of type **int**; more precisely, **x** is an array of three element objects, each of which is an array of five objects of type **int**. In the expression **x[i]**, which is equivalent to **(*((x)+(i)))**, **x** is first converted to a pointer to the initial array of five objects of type **int**. Then **i** is adjusted according to the type of **x**, which conceptually entails multiplying **i** by the size of the object to which the pointer points, namely an array of five **int** objects. The results are added and indirection is applied to yield an array of five objects of type **int**. When used in the expression **x[i][j]**, that array is in turn converted to a pointer to the first of the objects of type **int**, so **x[i][j]** yields an **int**.

Forward references: additive operators (6.5.7), address and indirection operators (6.5.4.3), array declarators (6.7.7.3).

6.5.3.3 Function calls

Constraints

- 1 The expression that denotes the called function⁹⁰⁾ shall have type pointer to function returning **void** or returning a complete object type other than an array type.
- 2 The number of arguments shall agree with the number of parameters. Each argument shall have a type such that its value may be assigned to an object with the unqualified version of the type of its corresponding parameter.

Semantics

- 3 A postfix expression followed by parentheses () containing a possibly empty, comma-separated list of expressions is a function call. The postfix expression denotes the called function. The list of expressions specifies the arguments to the function.
- 4 An argument can be an expression of any complete object type. In preparing for the call to a function, the arguments are evaluated, and each parameter is assigned the value of the corresponding argument.⁹¹⁾
- 5 If the expression that denotes the called function has type pointer to function returning an object type, the function call expression has the same type as that object type, and has the value determined as specified in 6.8.7.5. Otherwise, the function call has type **void**.
- 6 The arguments are implicitly converted, as if by assignment, to the types of the corresponding parameters, taking the type of each parameter to be the unqualified version of its declared type. The ellipsis notation in a variadic function declarator (6.7.7.4) causes argument type conversion to stop after the last declared parameter, if present. The integer promotions are performed on each trailing argument, and trailing arguments that have type **float** are promoted to **double**. These are called the *default argument promotions*. No other conversions are performed implicitly.
- 7 If the function is defined with a type that is not compatible with the type (of the expression) pointed to by the expression that denotes the called function, the behavior is undefined.
- 8 There is a sequence point after the evaluations of the function designator and the actual arguments but before the actual call. Every evaluation in the calling function (including other function calls) that is not otherwise specifically sequenced before or after the execution of the body of the called function is indeterminately sequenced with respect to the execution of the called function.⁹²⁾
- 9 Recursive function calls shall be permitted, both directly and indirectly through any chain of other functions.
- 10 **EXAMPLE** In the function call

```
(*pf[f1()]) (f2(), f3() + f4())
```

the functions **f1**, **f2**, **f3**, and **f4** can be called in any order. All side effects are completed before the function pointed to by **pf[f1()]** is called.

Forward references: function declarators (6.7.7.4), function definitions (6.9.2), the **return** statement (6.8.7.5), simple assignment (6.5.17.2).

6.5.3.4 Structure and union members

Constraints

- 1 The first operand of the **.** operator shall have an atomic, qualified, or unqualified structure or union type, and the second operand shall name a member of that type.
- 2 The first operand of the **->** operator shall have type “pointer to atomic, qualified, or unqualified structure” or “pointer to atomic, qualified, or unqualified union”, and the second operand shall

⁹⁰⁾Most often, this is the result of converting an identifier that is a function designator.

⁹¹⁾A function can change the values of its parameters, but these changes cannot affect the values of the arguments. On the other hand, it is possible to pass a pointer to an object, and the function can then change the value of the object pointed to. A parameter declared to have array or function type is adjusted to have a pointer type as described in 6.7.7.4.

⁹²⁾In other words, function executions do not interleave with each other.

name a member of the type pointed to.

Semantics

- 3 A postfix expression followed by the `.` operator and an identifier designates a member of a structure or union object. The value is that of the named member,⁹³⁾ and is an lvalue if the first expression is an lvalue. If the first expression has qualified type, the result has the so-qualified version of the type of the designated member.
- 4 A postfix expression followed by the `->` operator and an identifier designates a member of a structure or union object. The value is that of the named member of the object to which the first expression points, and is an lvalue.⁹⁴⁾ If the first expression is a pointer to a qualified type, the result has the so-qualified version of the type of the designated member.
- 5 Accessing a member of an atomic structure or union object results in undefined behavior.⁹⁵⁾
- 6 One special guarantee is made to simplify the use of unions: if a union contains several structures that share a common initial sequence (see following sentence), and if the union object currently contains one of these structures, it is permitted to inspect the common initial part of any of them anywhere that a declaration of the completed type of the union is visible. Two structures share a *common initial sequence* if corresponding members have compatible types (and, for bit-fields, the same widths) for a sequence of one or more initial members.
- 7 EXAMPLE 1 If `f` is a function returning a structure or union, and `x` is a member of that structure or union, `f().x` is a valid postfix expression but is not an lvalue.
- 8 EXAMPLE 2 In:

```
struct s { int i; const int ci; };
struct s s;
const struct s cs;
volatile struct s vs;
```

the various members have the types:

```
s.i      int
s.ci     const int
cs.i     const int
cs.ci    const int
vs.i     volatile int
vs.ci    volatile const int
```

- 9 EXAMPLE 3 The following is a valid fragment:

```
union {
    struct {
        int    alltypes;
    } n;
    struct {
        int    type;
        int    intnode;
    } ni;
    struct {
        int    type;
        double doublenode;
    } nf;
} u;
```

⁹³⁾If the member used to read the contents of a union object is not the same as the member last used to store a value in the object the appropriate part of the object representation of the value is reinterpreted as an object representation in the new type as described in 6.2.6 (a process sometimes called type punning). This can possibly be a non-value representation.

⁹⁴⁾If `&E` is a valid pointer expression (where `&` is the address of operator, which generates a pointer to its operand), the expression `(&E) ->MOS` is the same as `E.MOS`.

⁹⁵⁾For example, a data race would occur if access to the entire structure or union in one thread conflicts with access to a member from another thread, where at least one access is a modification. Members can be safely accessed using a non-atomic object which is assigned to or from the atomic object.

```

u.nf.type = 1;
u.nf.doublenode = 3.14;
/* ... */
if (u.n.alltypes == 1)
    if (sin(u.nf.doublenode) == 0.0)
        /* ... */

```

The following is not a valid fragment (because the union type is not visible within function `f`):

```

struct t1 { int m; };
struct t2 { int m; };
int f(struct t1 *p1, struct t2 *p2)
{
    if (p1->m < 0)
        p2->m = -p2->m;
    return p1->m;
}
int g()
{
    union {
        struct t1 s1;
        struct t2 s2;
    } u;
    /* ... */
    return f(&u.s1, &u.s2);
}

```

Forward references: address and indirection operators (6.5.4.3), structure and union specifiers (6.7.3.2).

6.5.3.5 Postfix increment and decrement operators

Constraints

- 1 The operand of the postfix increment or decrement operator shall have atomic, qualified, or unqualified arithmetic or pointer type, and shall be a modifiable lvalue.

Semantics

- 2 The result of the postfix ++ operator is the value of the operand. As a side effect, the value of the operand object is incremented by the adjustment. See the discussions of additive operators and compound assignment for information on constraints, types, and conversions and the effects of operations on pointers. The value computation of the result is sequenced before the side effect of updating the stored value of the operand. With respect to an indeterminately sequenced function call, the operation of postfix ++ is a single evaluation. Postfix ++ on an object with atomic type is a read-modify-write operation with `memory_order_seq_cst` memory order semantics.⁹⁶⁾
- 3 The *adjustment* used to increment or decrement the operand is the value 1 with a type and value representation as follows:

— if the operand has a pointer type, the adjustment has type `ptrdiff_t`;

⁹⁶⁾Where a pointer to an atomic object can be formed and `E` has integer type, `E++` is equivalent to the following code sequence where `T` is the type of `E`:

```

T *addr = &E;
T old = *addr;
T new;
do {
    new = old + 1;
} while (!atomic_compare_exchange_strong(addr, &old, new));

```

with `old` being the result of the operation.

Special care is necessary if `E` has floating type; see 6.5.17.3.

- if the operand has complex type, the adjustment has the corresponding real type of the operand;
 - if the operand has decimal floating type, the adjustment has the same type as the operand with a quantum exponent of 0;
 - otherwise, the adjustment has the same type as the operand.
- 4 The postfix `--` operator is analogous to the postfix `++` operator, except that the value of the operand is decremented by the adjustment.

Forward references: additive operators (6.5.7), compound assignment (6.5.17.3).

6.5.3.6 Compound literals

Syntax

- 1 *compound-literal:*
 (*storage-class-specifiers*_{opt} *type-name*) *braced-initializer*
- storage-class-specifiers:*
 storage-class-specifier
 storage-class-specifiers storage-class-specifier

Constraints

- 2 The type name shall specify a complete object type or an array of unknown size, but not a variable length array type.
- 3 All the constraints for initializer lists in 6.7.11 also apply to compound literals.
- 4 If the compound literal is associated with file scope or block scope (see 6.2.1) the storage-class specifiers **SC** (possibly empty),⁹⁷⁾ type name **T**, and initializer list, if any, shall be such that they are valid specifiers for an object definition in file scope or block scope, respectively, of the following form,

SC **typeof**(**T**) **ID** = { **IL** };

where **ID** is an identifier that is unique for the whole program and where **IL** is a (possibly empty) initializer list with nested structure, designators, values and types as the initializer list of the compound literal. All the constraints for storage-class specifiers in 6.7.2 also apply correspondingly to compound literals. If the compound literal is associated with function prototype scope, constraints as if in block scope apply.

Semantics

- 5 For a *compound literal* associated with function prototype scope:
- the type is determined as if in block scope and no object is created;
 - if it is a compound literal constant it is evaluated at translation time;
 - if it is not a compound literal constant, neither the compound literal as a whole nor any of the initializers are evaluated.

Otherwise, a compound literal provides access to an unnamed object whose value, type, storage duration, initializer, and other properties are as if given by the definition syntax in the constraints.

- 6 If the storage duration is automatic, the lifetime of the instance of the unnamed object is the current execution of the enclosing block.⁹⁸⁾

⁹⁷⁾If the storage-class specifiers contain the same storage-class specifier more than once, the following constraint is violated.

⁹⁸⁾This differs from a cast expression. For example, a cast specifies a conversion to scalar types or **void** only, and the result of a cast expression is not an lvalue.

- 7 If the storage-class specifiers are absent or contain **constexpr**, **static**, **register**, or **thread_local** the behavior is as if the object were declared and initialized in the corresponding scope with these storage-class specifiers; if another storage-class specifier is present, the behavior is undefined. If the storage-class specifier **constexpr** is present, the initializer is evaluated at translation time. Otherwise, if the storage duration is automatic, the initializer is evaluated at each evaluation of the compound literal; if the storage duration is static or thread the initializer is (as if) evaluated once prior to program startup.
- 8 The value of the compound literal is that of an lvalue corresponding to the unnamed object.
- 9 All the semantic rules for initializer lists in 6.7.11 also apply to compound literals.⁹⁹⁾
- 10 String literals, and compound literals with **const**-qualified types, including those specified with **constexpr**, are not required to designate distinct objects.¹⁰⁰⁾
- 11 EXAMPLE 1 The following two functions can be used as an example:

```
int f(int*);
int g(char * para[f((int[27]){ 0, })]) {
    /* ... */
    return 0;
}
```

Here, each call to **g** creates an unnamed object of type **int[27]** to determine the variably modified type of **para** for the duration of the call. During that determination, a pointer to the object is passed into a call to the function **f**. If a pointer to the object is kept by **f**, access to that object is possible during the whole execution of the call to **g**. The lifetime of the object ends with the end of the call to **g**; for any access after that, the behavior is undefined.

- 12 EXAMPLE 2 The file scope definition

```
int *p = (int []){2, 4};
```

initializes **p** to point to the first element of an array of two **int** s, the first having the value two and the second having the value four. The expressions in this compound literal are expected to be constant. The unnamed object has static storage duration.

- 13 EXAMPLE 3 In contrast, in

```
void f(void)
{
    int *p;
    /* ... */
    p = (int [2]){*p};
    /* ... */
}
```

p is assigned the address of the first element of an array of two **int** s, the first having the value previously pointed to by **p** and the second, zero. The initializer expression in this compound literal is not required to be constant. The unnamed object has automatic storage duration.

- 14 EXAMPLE 4 Initializers with designations can be combined with compound literals. Structure objects created using compound literals can be passed to functions without depending on member order:

```
drawline((struct point){.x=1, .y=1},
        (struct point){.x=3, .y=4});
```

Or, if **drawline** instead expected pointers to **struct point**:

```
drawline(&(struct point){.x=1, .y=1},
        &(struct point){.x=3, .y=4});
```

⁹⁹⁾For example, subobjects without explicit initializers are initialized to zero.

¹⁰⁰⁾This allows implementations to share storage for string literals and constant compound literals with the same or overlapping representations.

- 15 EXAMPLE 5 A read-only compound literal can be specified through constructions like:

```
(const float []){1e0, 1e1, 1e2, 1e3, 1e4, 1e5, 1e6}
```

- 16 EXAMPLE 6 The following three expressions have different meanings:

```
"/tmp/fileXXXXXX"
(char []){"/tmp/fileXXXXXX"}
(const char []){"/tmp/fileXXXXXX"}
```

The first always has static storage duration and has type array of **char**, but can be modifiable; the last two have automatic storage duration when they occur within the body of a function, and the first of these two is modifiable.

- 17 EXAMPLE 7 Like string literals, const-qualified compound literals can be placed into read-only memory and can even be shared. For example,

```
(const char []){"abc"} == "abc"
```

can yield 1 if the literals' storage is shared.

- 18 EXAMPLE 8 Since compound literals are unnamed, a single compound literal cannot specify a circularly linked object. For example, there is no way to write a self-referential compound literal that can be used as the function argument in place of the named object `endless_zeros` in the following snippet:

```
struct int_list { int car; struct int_list *cdr; };
struct int_list endless_zeros = {0, &endless_zeros};
eval(endless_zeros);
```

- 19 EXAMPLE 9 Each compound literal creates only a single object in a given scope:

```
struct s { int i; };

int f (void)
{
    struct s *p = 0, *q;
    int j = 0;

    again:
        q = p, p = &((struct s){ j++ });
        if (j < 2) goto again;

    return p == q && q->i == 1;
}
```

The function `f()` always returns the value 1.

- 20 If an iteration statement were used instead of an explicit **goto** and a label, the lifetime of the unnamed object would be the body of the loop only, and on entry next time around `p` would have indeterminate representation, which would result in undefined behavior.

Forward references: type names (6.7.8), initialization (6.7.11).

6.5.4 Unary operators

6.5.4.1 General

Syntax

- 1 *unary-expression:*
- postfix-expression*
 - ++** *unary-expression*
 - *unary-expression*
 - unary-operator* *cast-expression*
 - _Lengthof** *unary-expression*

```

    _Lengthof ( type-name )
    sizeof unary-expression
    sizeof ( type-name )
    alignof ( type-name )

```

unary-operator: one of

& * + - ~ !

6.5.4.2 Prefix increment and decrement operators

Constraints

- 1 The operand of the prefix increment or decrement operator shall have atomic, qualified, or unqualified arithmetic or pointer type, and shall be a modifiable lvalue.

Semantics

- 2 The value of the operand of the prefix ++ operator is incremented. The result is the new value of the operand after incrementation. The expression ++*E* is similar to (*E*+=1), except where the value 1 is of the adjustment (6.5.3.5). See the discussions of additive operators and compound assignment for information on constraints, types, side effects, and conversions and the effects of operations on pointers.
- 3 The prefix -- operator is analogous to the prefix ++ operator, except that the value of the operand is decremented.

Forward references: additive operators (6.5.7), compound assignment (6.5.17.3).

6.5.4.3 Address and indirection operators

Constraints

- 1 The operand of the unary & operator shall be either a function designator, the result of a [] or unary * operator, or an lvalue that designates an object that is not a bit-field and is not declared with the **register** storage-class specifier.
- 2 The operand of the unary * operator shall have pointer type.

Semantics

- 3 The unary & operator yields the address of its operand. If the operand has type "*type*", the result has type "pointer to *type*". If the operand is the result of a unary * operator, neither that operator nor the & operator is evaluated and the result is as if both were omitted, except that the constraints on the operators still apply and the result is not an lvalue. Similarly, if the operand is the result of a [] operator, neither the & operator nor the unary * that is implied by the [] is evaluated and the result is as if the & operator were removed and the [] operator were changed to a + operator. Otherwise, the result is a pointer to the object or function designated by its operand.
- 4 The unary * operator denotes indirection. If the operand points to a function, the result is a function designator; if it points to an object, the result is an lvalue designating the object. If the operand has type "pointer to *type*", the result has type "*type*". If an invalid value has been assigned to the pointer, the behavior of the unary * operator is undefined.¹⁰¹⁾

Forward references: storage-class specifiers (6.7.2), structure and union specifiers (6.7.3.2).

¹⁰¹⁾Thus, &**E* is equivalent to *E* (even if *E* is a null pointer), and &(*E*1[*E*2]) to ((*E*1)+(*E*2)). It is always true that if *E* is a function designator or an lvalue that is a valid operand of the unary & operator, *&*E* is a function designator or an lvalue equal to *E*. If **P* is an lvalue and *T* is the name of an object pointer type, *(*T*)*P* is an lvalue that has a type compatible with that to which *T* points.

Among the invalid values for dereferencing a pointer by the unary * operator are a null pointer, an address inappropriately aligned for the type of object pointed to, and the address of an object after the end of its lifetime.

6.5.4.4 Unary arithmetic operators

Constraints

- 1 The operand of the unary + or - operator shall have arithmetic type; of the ~ operator, integer type; of the ! operator, scalar type.

Semantics

- 2 The result of the unary + operator is the value of its (promoted) operand. The integer promotions are performed on the operand, and the result has the promoted type.
- 3 The result of the unary - operator is the negative of its (promoted) operand. The integer promotions are performed on the operand, and the result has the promoted type.
- 4 The result of the ~ operator is the bitwise complement of its (promoted) operand (that is, each bit in the result is set if and only if the corresponding bit in the converted operand is not set). The integer promotions are performed on the operand, and the result has the promoted type. If the promoted type is an unsigned type, the expression ~E is equivalent to the maximum value representable in that type minus E.
- 5 The result of the logical negation operator ! is 0 if the value of its operand compares unequal to 0, 1 if the value of its operand compares equal to 0. The result has type **int**. The expression !E is equivalent to (0==E).

6.5.4.5 The **sizeof**, **_Lengthof**, and **alignof** operators

Constraints

- 1 The **sizeof** operator shall not be applied to an expression that has function type or an incomplete type, to the parenthesized name of such a type, or to an expression that designates a bit-field member. The **_Lengthof** operator shall only be applied to an expression that has a complete array type, or to the parenthesized name of such a type. The **alignof** operator shall be applied to a complete object type or an array thereof.

Semantics

- 2 The **sizeof** operator yields the size (in bytes) of its operand, which can be an expression or the parenthesized name of a type. The size is determined from the type of the operand. The result is an integer. If the type of the operand is a variable length array type, the operand is evaluated; otherwise, the operand is not evaluated and the result is an integer constant expression.
- 3 The **alignof** operator yields the alignment requirement of its operand type. The operand is not evaluated and the result is an integer constant expression. When applied to an array type, the result is the alignment requirement of the element type.
- 4 When **sizeof** is applied to an operand that has type **char**, **unsigned char**, or **signed char**, (or a qualified version thereof) the result is 1. When applied to an operand that has array type, the result is the total number of bytes in the array.¹⁰²⁾ When applied to an operand that has structure or union type, the result is the total number of bytes in such an object, including internal and trailing padding.
- 5 The **_Lengthof** operator yields the number of elements of its operand. The number of elements is determined from the type of the operand. The result is an integer. If the number of elements of the array type is variable, the operand is evaluated; otherwise, the operand is not evaluated and the result is an integer constant expression.
- 6 The value of the result of all operators is implementation-defined, and its type (an unsigned integer type) is **size_t**, defined in `<stddef.h>` (and other headers).
- 7 EXAMPLE 1 A principal use of the **sizeof** operator is in communication with routines such as storage allocators and I/O systems. A storage-allocation function can accept a size (in bytes) of an object to allocate and return a pointer to **void**. For example:

```
extern void *alloc(size_t);
```

¹⁰²⁾When applied to a parameter declared to have array or function type, the **sizeof** operator yields the size of the adjusted (pointer) type (see 6.9.2).

```
double *dp = alloc(sizeof *dp);
```

The implementation of the `alloc` function presumably ensures that its return value is aligned suitably for conversion to a pointer to `double`.

- 8 EXAMPLE 2 The use of `_Lengthof` operator is to compute the number of elements in an array. Similar results can be computed from the `sizeof` operator, but the `_Lengthof` operator is more concise and is encouraged:

```
double array[42];
static_assert(_Lengthof array == 42); // number of elements
static_assert(sizeof array == _Lengthof array * sizeof(double)); /* number of
    bytes in a one-dimensional array */
static_assert(_Lengthof array == sizeof array / sizeof array[0]); // equivalent
```

The `_Lengthof` operator only computes the number of elements in the array of the outermost dimension, which is not the same as `sizeof`:

```
double array[4][42];
static_assert(_Lengthof array == 4); // number of elements in
    // the last array's dimension
static_assert(sizeof array != _Lengthof array * sizeof(double)); /* NOT
    equivalent; sizeof is the byte size of the whole
    multidimensional array object */
static_assert(_Lengthof array == sizeof array / sizeof array[0]); // equivalent
```

- 9 EXAMPLE 3 In this example, the size of a variable length array is computed and returned from a function:

```
#include <stddef.h>

size_t fsize3(int n)
{
    char b[n+3]; // variable length array
    return sizeof b; // execution time sizeof
}

int main(void)
{
    size_t size;
    size = fsize3(10); // fsize3 returns 13
    return 0;
}
```

- 10 EXAMPLE 4 The `alignof` operator can be applied to an incomplete array type, as shown in the following snippet:

```
int a1 = alignof(int[]); // 0k
int a2 = alignof(int[][1]); // 0k
```

Forward references: common definitions `<stddef.h>` (7.21), declarations (6.7), structure and union specifiers (6.7.3.2), type names (6.7.8), array declarators (6.7.7.3).

6.5.5 Cast operators

Syntax

- 1 *cast-expression*:
- ```
 unary-expression
 (type-name) cast-expression
```



**Constraints**

- 2 Unless the type name specifies a void type, the type name shall specify atomic, qualified, or unqualified scalar type, and the operand shall have scalar type.
- 3 Conversions that involve pointers, other than where permitted by the constraints of 6.5.17.2, shall be specified by means of an explicit cast.
- 4 A pointer type shall be converted only to **void**, an integer type, or a pointer type. Only a pointer, integer, or **nullptr\_t** type shall be converted to a pointer type. The type **nullptr\_t** shall not be converted to any type other than **void**, **bool** or a pointer type. If the target type is **nullptr\_t**, the cast expression shall be a null pointer constant or have type **nullptr\_t**.

**Semantics**

- 5 Size expressions and typeof operators contained in a type name used with a cast operator are evaluated whenever the cast expression is evaluated.
- 6 Preceding an expression by a parenthesized type name converts the value of the expression to the unqualified, non-atomic version of the named type. This construction is called a *cast*.<sup>103)</sup> A cast that specifies no conversion has no effect on the type or value of an expression.
- 7 If the value of the expression is represented with greater range or precision than required by the type named by the cast (6.3.2.8), then the cast specifies a conversion even if the type of the expression is the same as the named type and removes any extra range and precision.

**Forward references:** equality operators (6.5.10), function declarators (6.7.7.4), simple assignment (6.5.17.2), type names (6.7.8).

**6.5.6 Multiplicative operators****Syntax**

- 1 *multiplicative-expression:*  
     *cast-expression*  
     *multiplicative-expression* \* *cast-expression*  
     *multiplicative-expression* / *cast-expression*  
     *multiplicative-expression* % *cast-expression*

**Constraints**

- 2 Each of the operands shall have arithmetic type. The operands of the % operator shall have integer type.
- 3 If either operand has decimal floating type, the other operand shall not have standard floating type, or complex type type.

**Semantics**

- 4 The usual arithmetic conversions are performed on the operands.
- 5 The result of the binary \* operator is the product of the operands.
- 6 The result of the / operator is the quotient from the division of the first operand by the second; the result of the % operator is the remainder. In both operations, if the value of the second operand is zero, the behavior is undefined.
- 7 When integers are divided, the result of the / operator is the algebraic quotient with any fractional part discarded.<sup>104)</sup> If the quotient **a/b** is representable, the expression **(a/b)\*b + a%b** shall equal **a**; otherwise, the behavior of both **a/b** and **a%b** is undefined.

<sup>103)</sup>A cast does not yield an lvalue.

<sup>104)</sup>This is often called “truncation toward zero”.

## 6.5.7 Additive operators

### Syntax

- 1 *additive-expression*:
- multiplicative-expression*
  - additive-expression* + *multiplicative-expression*
  - additive-expression* - *multiplicative-expression*

### Constraints

- 2 For addition, either both operands shall have arithmetic type, or one operand shall be a pointer to a complete object type and the other shall have integer type. (Incrementing is equivalent to adding 1.)
- 3 For subtraction, one of the following shall hold:
- both operands have arithmetic type;
  - both operands are pointers to qualified or unqualified versions of compatible complete object types; or
  - the left operand is a pointer to a complete object type and the right operand has integer type.

(Decrementing is equivalent to subtracting 1.)

- 4 If either operand has decimal floating type, the other operand shall not have standard floating type, or complex type.

### Semantics

- 5 If both operands have arithmetic type, the usual arithmetic conversions are performed on them.
- 6 The result of the binary + operator is the sum of the operands.
- 7 The result of the binary - operator is the difference resulting from the subtraction of the second operand from the first.
- 8 For the purposes of these operators, a pointer to an object that is not an element of an array behaves the same as a pointer to the first element of an array of length one with the type of the object as its element type.
- 9 When an expression that has integer type is added to or subtracted from a pointer, the result has the type of the pointer operand. If the pointer operand points to an element of an array object, and the array is large enough, the result points to an element offset from the original element such that the difference of the subscripts of the resulting and original array elements equals the integer expression. In other words, if the expression **P** points to the *i*-th element of an array object, the expressions (**P**) + **N** (equivalently, **N** + (**P**)) and (**P**) - **N** (where **N** has the value *n*) point to, respectively, the *i* + *n*-th and *i* - *n*-th elements of the array object, provided they exist. Moreover, if the expression **P** points to the last element of an array object, the expression (**P**) + 1 points one past the last element of the array object, and if the expression **Q** points one past the last element of an array object, the expression (**Q**) - 1 points to the last element of the array object. If the pointer operand is not null, and the pointer operand and result do not point to elements of the same array object or one past the last element of the array object, the behavior is undefined. If the pointer operand is a null pointer value and the integer operand is nonzero, the behavior is undefined.<sup>105)</sup> If the addition or subtraction produces an overflow, the behavior is undefined. If the pointer operand is null or the result points one past the last element of the array object, it shall not be used as the operand of a unary \* operator that is evaluated.
- 10 When two pointers are subtracted, both shall be null pointer values, point to elements of the same array object, or one past the last element of the array object. If both pointers are null pointer values, the result is zero. Otherwise, the result is the difference of the subscripts of the two array elements.

<sup>105)</sup> Thus, the expression **ptr** + **N** (where **ptr** is a null pointer value) is defined to result in a null pointer value when **N** is zero and has undefined behavior otherwise.

The size of the result is implementation-defined, and its type (a signed integer type) is `ptrdiff_t` defined in the `<stddef.h>` header. If the result is not representable in an object of that type, the behavior is undefined. In other words, if the expressions `P` and `Q` point to, respectively, the  $i$ -th and  $j$ -th elements of an array object, the expression `(P) - (Q)` has the value  $i - j$  provided the value fits in an object of type `ptrdiff_t`. Moreover, if the expression `P` points either to an element of an array object or one past the last element of an array object, and the expression `Q` points to the last element of the same array object, the expression `((Q)+1) - (P)` has the same value as `((Q) - (P))+1` and as `-((P) - ((Q)+1))`, and has the value zero if the expression `P` points one past the last element of the array object, even though the expression `(Q)+1` does not point to an element of the array object.

- 11 NOTE Another way to approach pointer arithmetic is first to convert the pointer(s) to character pointer(s): In this scheme the integer expression added to or subtracted from the converted pointer is first multiplied by the size of the object originally pointed to, and the resulting pointer is converted back to the original type. For pointer subtraction, the result of the difference between the character pointers is similarly divided by the size of the object originally pointed to. When viewed in this way, an implementation need only provide one extra byte (which can overlap another object in the program) just after the end of the object to satisfy the “one past the last element” requirements.
- 12 EXAMPLE Pointer arithmetic is well defined with pointers to variable length array types.

```

{
 int n = 4, m = 3;
 int a[n][m];
 int (*p)[m] = a; // p == &a[0]
 p += 1; // p == &a[1]
 (*p)[2] = 99; // a[1][2] == 99
 n = p - a; // n == 1
}

```

- 13 If array `a` in the preceding example were declared to be an array of known constant size, and pointer `p` were declared to be a pointer to an array of the same known constant size (pointing to `a`), the results would be the same.

**Forward references:** array declarators (6.7.7.3), common definitions `<stddef.h>` (7.21).

## 6.5.8 Bitwise shift operators

### Syntax

- 1 *shift-expression:*

```

additive-expression
shift-expression << additive-expression
shift-expression >> additive-expression

```

### Constraints

- 2 Each of the operands shall have integer type.

### Semantics

- 3 The integer promotions are performed on each of the operands. The type of the result is that of the promoted left operand. If the value of the right operand is negative or is greater than or equal to the width of the promoted left operand, the behavior is undefined.
- 4 The result of `E1 << E2` is `E1` left-shifted `E2` bit positions; vacated bits are filled with zeros. If `E1` has an unsigned type, the value of the result is  $E1 \times 2^{E2}$ , wrapped around. If `E1` has a signed type and nonnegative value, and  $E1 \times 2^{E2}$  is representable in the result type, then that is the resulting value; otherwise, the behavior is undefined.
- 5 The result of `E1 >> E2` is `E1` right-shifted `E2` bit positions. If `E1` has an unsigned type or if `E1` has a signed type and a nonnegative value, the value of the result is the integral part of the quotient of  $E1/2^{E2}$ . If `E1` has a signed type and a negative value, the resulting value is implementation-defined.

### 6.5.9 Relational operators

#### Syntax

- 1 *relational-expression*:
- ```

    shift-expression
    relational-expression < shift-expression
    relational-expression > shift-expression
    relational-expression <= shift-expression
    relational-expression >= shift-expression
  
```

Constraints

- 2 One of the following shall hold:
- both operands have real type; or
 - both operands are pointers to qualified or unqualified versions of compatible object types.
- 3 If either operand has decimal floating type, the other operand shall not have standard floating type.

Semantics

- 4 If both of the operands have arithmetic type, the usual arithmetic conversions are performed. Positive zeros compare equal to negative zeros.
- 5 For the purposes of these operators, a pointer to an object that is not an element of an array behaves the same as a pointer to the first element of an array of length one with the type of the object as its element type.
- 6 When two pointers are compared, the result depends on the relative locations in the address space of the objects pointed to, if any. If two pointers to object types are both null pointer values, both point to the same object, or both point one past the last element of the same array object, they compare equal. If the objects pointed to are members of the same aggregate object, pointers to structure members declared later compare greater than pointers to members declared earlier in the structure, and pointers to array elements with larger subscript values compare greater than pointers to elements of the same array with lower subscript values. All pointers to members of the same union object compare equal. If the expression **P** points to an element of an array object and the expression **Q** points to the last element of the same array object, the pointer expression **Q+1** compares greater than **P**. In all other cases, the behavior is undefined.
- 7 Each of the operators < (less than), > (greater than), <= (less than or equal to), and >= (greater than or equal to) shall yield 1 if the specified relation is true and 0 if it is false.¹⁰⁶⁾ The result has type **int**.

6.5.10 Equality operators

Syntax

- 1 *equality-expression*:
- ```

 relational-expression
 equality-expression == relational-expression
 equality-expression != relational-expression

```

#### Constraints

- 2 One of the following shall hold:
- both operands have arithmetic type;
  - both operands are pointers to qualified or unqualified versions of compatible types;

<sup>106)</sup>The expression **a<b<c** is not interpreted as in ordinary mathematics. As the syntax indicates, it means (**a<b**)<**c**; in other words, “if **a** is less than **b**, compare 1 to **c**; otherwise, compare 0 to **c**”.

- one operand is a pointer to an object type and the other is a pointer to a qualified or unqualified version of **void**;
  - both operands have type **nullptr\_t**;
  - one operand has type **nullptr\_t** and the other is a null pointer constant; or,
  - one operand is a pointer and the other is a null pointer constant or has type **nullptr\_t**.
- 3 If either operand has decimal floating type, the other operand shall not have standard floating type, or complex type.

#### Semantics

- 4 The == (equal to) and != (not equal to) operators are analogous to the relational operators except for their lower precedence.<sup>107)</sup> Each of the operators yields 1 if the specified relation is true and 0 if it is false. The result has type **int**. For any pair of operands, exactly one of the relations is true.
- 5 If both of the operands have arithmetic type, the usual arithmetic conversions are performed. Positive zeros compare equal to negative zeros. Values of complex types are equal if and only if both their real parts are equal and also their imaginary parts are equal. Any two values of arithmetic types from different type domains are equal if and only if the results of their conversions to the (complex) result type determined by the usual arithmetic conversions are equal. If both operands have type **nullptr\_t** or one operand has type **nullptr\_t** and the other is a null pointer constant, they compare equal.
- 6 Otherwise, at least one operand is a pointer. If one operand is a pointer and the other is a null pointer constant or has type **nullptr\_t**, they compare equal if the former is a null pointer. If one operand is a pointer to an object type and the other is a pointer to a qualified or unqualified version of **void**, the former is converted to the type of the latter.
- 7 Two pointers compare equal if and only if both are null pointers, both are pointers to the same object (including a pointer to an object and a subobject at its beginning) or function, both are pointers to one past the last element of the same array object, or one is a pointer to one past the end of one array object and the other is a pointer to the start of a different array object that happens to immediately follow the first array object in the address space.<sup>108)</sup>
- 8 For the purposes of these operators, a pointer to an object that is not an element of an array behaves the same as a pointer to the first element of an array of length one with the type of the object as its element type.

### 6.5.11 Bitwise AND operator

#### Syntax

- 1 *AND-expression:*
- equality-expression*  
*AND-expression & equality-expression*

#### Constraints

- 2 Each of the operands shall have integer type.

#### Semantics

- 3 The usual arithmetic conversions are performed on the operands.
- 4 The result of the binary & operator is the bitwise AND of the operands (that is, each bit in the result is set if and only if each of the corresponding bits in the converted operands is set).

<sup>107)</sup> Because of the precedences, **a < b == c < d** is 1 whenever **a < b** and **c < d** have the same truth-value.

<sup>108)</sup> Two objects can be adjacent in memory because they are adjacent elements of a larger array or adjacent members of a structure with no padding between them, or because the implementation chose to place them so, even though they are unrelated. If prior invalid pointer operations (such as accesses outside array bounds) produced undefined behavior, subsequent comparisons also produce undefined behavior.

### 6.5.12 Bitwise exclusive OR operator

#### Syntax

- 1 *exclusive-OR-expression*:  
     AND-expression  
     *exclusive-OR-expression* ^ AND-expression

#### Constraints

- 2 Each of the operands shall have integer type.

#### Semantics

- 3 The usual arithmetic conversions are performed on the operands.  
 4 The result of the ^ operator is the bitwise exclusive OR of the operands (that is, each bit in the result is set if and only if exactly one of the corresponding bits in the converted operands is set).

### 6.5.13 Bitwise inclusive OR operator

#### Syntax

- 1 *inclusive-OR-expression*:  
     *exclusive-OR-expression*  
     *inclusive-OR-expression* | *exclusive-OR-expression*

#### Constraints

- 2 Each of the operands shall have integer type.

#### Semantics

- 3 The usual arithmetic conversions are performed on the operands.  
 4 The result of the | operator is the bitwise inclusive OR of the operands (that is, each bit in the result is set if and only if at least one of the corresponding bits in the converted operands is set).

### 6.5.14 Logical AND operator

#### Syntax

- 1 *logical-AND-expression*:  
     *inclusive-OR-expression*  
     *logical-AND-expression* && *inclusive-OR-expression*

#### Constraints

- 2 Each of the operands shall have scalar type.

#### Semantics

- 3 The && operator shall yield 1 if both of its operands compare unequal to 0; otherwise, it yields 0. The result has type **int**.  
 4 Unlike the bitwise binary & operator, the && operator guarantees left-to-right evaluation; if the second operand is evaluated, there is a sequence point between the evaluations of the first and second operands. If the first operand compares equal to 0, the second operand is not evaluated.

### 6.5.15 Logical OR operator

#### Syntax

- 1 *logical-OR-expression*:  
     *logical-AND-expression*  
     *logical-OR-expression* || *logical-AND-expression*

**Constraints**

- 2 Each of the operands shall have scalar type.

**Semantics**

- 3 The `||` operator shall yield 1 if either of its operands compare unequal to 0; otherwise, it yields 0. The result has type **int**.
- 4 Unlike the bitwise `|` operator, the `||` operator guarantees left-to-right evaluation; if the second operand is evaluated, there is a sequence point between the evaluations of the first and second operands. If the first operand compares unequal to 0, the second operand is not evaluated.

**6.5.16 Conditional operator****Syntax**

- 1 *conditional-expression:*  
     *logical-OR-expression*  
     *logical-OR-expression* ? *expression* : *conditional-expression*

**Constraints**

- 2 The first operand shall have scalar type.
- 3 One of the following shall hold for the second and third operands:<sup>109)</sup>
- both operands have arithmetic type;
  - both operands have compatible structure or union type;
  - both operands have void type;
  - both operands are pointers to qualified or unqualified versions of compatible types;
  - both operands have **nullptr\_t** type;
  - one operand is a pointer and the other is a null pointer constant or has type **nullptr\_t**; or
  - one operand is a pointer to an object type and the other is a pointer to a qualified or unqualified version of **void**.
- 4 If either of the second or third operands has decimal floating type, the other operand shall not have standard floating type, or complex type.

**Semantics**

- 5 The first operand is evaluated; there is a sequence point between its evaluation and the evaluation of the second or third operand (whichever is evaluated). The second operand is evaluated only if the first compares unequal to 0; the third operand is evaluated only if the first compares equal to 0; the result is the value of the second or third operand (whichever is evaluated), converted to the type described subsequently in this subclause.<sup>110)</sup>
- 6 If the second and third operands have arithmetic type, the result type is the same as if the usual arithmetic conversions were applied to both operands. If both the operands have structure or union type, the result is the composite type. If both operands have void type, the result has void type.
- 7 If both the second and third operands are pointers, the result type is a pointer to a type qualified with all the type qualifiers of the types referenced by both operands; if one is a null pointer constant (other than a pointer) or has type **nullptr\_t** and the other is a pointer, the result type is the pointer type; if both the second and third operands have **nullptr\_t** type, the result also has that type. Furthermore, if both operands are pointers to compatible types or to differently qualified versions of

<sup>109)</sup>If a second or third operand of type **nullptr\_t** is used and the other operand is not a pointer and does not have type **nullptr\_t** itself, a constraint is violated even if that other operand is a null pointer constant such as 0.

<sup>110)</sup>A conditional expression does not yield an lvalue.

compatible types, the result type is a pointer to an appropriately qualified version of the composite type; if one operand is a null pointer constant, the result has the type of the other operand; otherwise, one operand is a pointer to **void** or a qualified version of **void**, in which case the result type is a pointer to an appropriately qualified version of **void**.

- 8 If one operand is a pointer to a variably modified type and the other operand is a null pointer constant or has type **nullptr\_t**, the behavior is undefined if the type depends on an array size expression that is not evaluated.
- 9 **EXAMPLE** The common type that results when the second and third operands are pointers is determined in two independent stages. The appropriate qualifiers, for example, do not depend on whether the two pointers have compatible types.
- 10 Given the declarations

```
const void *c_vp;
void *vp;
const int *c_ip;
volatile int *v_ip;
int *ip;
const char *c_cp;
```

the third column in the Table 6.8 is the common type that is the result of a conditional expression in which the first two columns are the second and third operands (in either order):

**Table 6.8 — Common type for conditional expression evaluations**

|             |             |                             |
|-------------|-------------|-----------------------------|
| <b>c_vp</b> | <b>c_ip</b> | <b>const void *</b>         |
| <b>v_ip</b> | <b>0</b>    | <b>volatile int *</b>       |
| <b>c_ip</b> | <b>v_ip</b> | <b>const volatile int *</b> |
| <b>vp</b>   | <b>c_cp</b> | <b>const void *</b>         |
| <b>ip</b>   | <b>c_ip</b> | <b>const int *</b>          |
| <b>vp</b>   | <b>ip</b>   | <b>void *</b>               |

## 6.5.17 Assignment operators

### 6.5.17.1 General

#### Syntax

- 1 *assignment-expression*:  
*conditional-expression*  
*unary-expression assignment-operator assignment-expression*

*assignment-operator*: one of

**=   \*=   /=   %=   +=   -=   <<=   >>=   &=   ^=   |=**

#### Constraints

- 2 An assignment operator shall have a modifiable lvalue as its left operand.

#### Semantics

- 3 An assignment operator stores a value in the object designated by the left operand. An assignment expression has the value of the left operand after the assignment,<sup>111)</sup> but is not an lvalue. The type of an assignment expression is the type the left operand would have after lvalue conversion. The side effect of updating the stored value of the left operand is sequenced after the value computations of the left and right operands. The evaluations of the operands are unsequenced.

<sup>111)</sup>The implementation is permitted to read the object to determine the value but is not required to, even when the object has volatile-qualified type.



### 6.5.17.2 Simple assignment

#### Constraints

- 1 One of the following shall hold:<sup>112)</sup>
  - the left operand has atomic, qualified, or unqualified arithmetic type, and the right operand has arithmetic type;
  - the left operand has an atomic, qualified, or unqualified version of a structure or union type compatible with the type of the right operand;
  - the left operand has atomic, qualified, or unqualified pointer type, and (considering the type the left operand would have after lvalue conversion) both operands are pointers to qualified or unqualified versions of compatible types, and the type pointed to by the left operand has all the qualifiers of the type pointed to by the right operand;
  - the left operand has atomic, qualified, or unqualified pointer type, and (considering the type the left operand would have after lvalue conversion) one operand is a pointer to an object type, and the other is a pointer to a qualified or unqualified version of **void**, and the type pointed to by the left operand has all the qualifiers of the type pointed to by the right operand;
  - the left operand has an atomic, qualified, or unqualified version of the **nullptr\_t** type and the right operand is a null pointer constant or its type is **nullptr\_t**;
  - the left operand is an atomic, qualified, or unqualified pointer, and the right operand is a null pointer constant or its type is **nullptr\_t**; or
  - the left operand has type atomic, qualified, or unqualified **bool**, and the right operand is a pointer or its type is **nullptr\_t**.

#### Semantics

- 2 In *simple assignment* (**=**), the value of the right operand is converted to the type of the assignment expression and replaces the value stored in the object designated by the left operand.<sup>113)</sup>
- 3 If the value being stored in an object is read from another object that overlaps in any way the storage of the first object, then the two objects shall occupy exactly the same storage and shall have qualified or unqualified versions of a compatible type; otherwise, the behavior is undefined.
- 4 EXAMPLE 1 In the program fragment

```
int f(void);
char c;
/* ... */
if ((c = f()) == -1)
 /* ... */
```

the **int** value returned by the function can be truncated when stored in the **char**, and then converted back to **int** width prior to the comparison. In an implementation in which “plain” **char** has the same range of values as **unsigned char** (and **char** is narrower than **int**), the result of the conversion cannot be negative, so the operands of the comparison can never compare equal. Therefore, for full portability, the variable **c** would be declared as **int**.

- 5 EXAMPLE 2 In the fragment:

```
char c;
int i;
long l;

l = (c = i);
```

<sup>112)</sup>The asymmetric appearance of these constraints with respect to type qualifiers is due to the conversion (specified in 6.3.3.1) that changes lvalues to “the value of the expression” and thus removes any type qualifiers that were applied to the type category of the expression (for example, it removes **const** but not **volatile** from the type **int volatile \* const**).

<sup>113)</sup>As described in 6.2.6.1, a store to an object with atomic type is done with **memory\_order\_seq\_cst** semantics.

the value of `i` is converted to the type of the assignment expression `c = i`, that is, `char` type. The value of the expression enclosed in parentheses is then converted to the type of the outer assignment expression, that is, `long int` type.

- 6 EXAMPLE 3 The following fragment can be used as an example:

```
const char **cpp;
char *p;
const char c = 'A';

cpp = &p; // constraint violation
*cpp = &c; // valid
*p = 0; // valid
```

The first assignment is unsafe because it would allow the following valid code to attempt to change the value of the const object `c`.

### 6.5.17.3 Compound assignment

#### Constraints

- For the operators `+=` and `-=` only, either the left operand shall be an atomic, qualified, or unqualified pointer to a complete object type, and the right shall have integer type; or the left operand shall have atomic, qualified, or unqualified arithmetic type, and the right shall have arithmetic type.
- For the other operators, the left operand shall have atomic, qualified, or unqualified arithmetic type, and (considering the type the left operand would have after lvalue conversion) each operand shall have arithmetic type consistent with those allowed by the corresponding binary operator.
- If either operand has decimal floating type, the other operand shall not have standard floating type, or complex type.

#### Semantics

- A *compound assignment* of the form `E1 op= E2` is equivalent to the simple assignment expression `E1 = E1 op (E2)`, except that the lvalue `E1` is evaluated only once, and with respect to an indeterminately sequenced function call, the operation of a compound assignment is a single evaluation. If `E1` has an atomic type, compound assignment is a read-modify-write operation with `memory_order_seq_cst` memory order semantics.
- NOTE Where a pointer to an atomic object can be formed and `E1` and `E2` have integer type, this is equivalent to the following code sequence where `T1` is the type of `E1` and `T2` is the type of `E2`:

```
T1 *addr = &E1;
T2 val = (E2);
T1 old = *addr;
T1 new;
do {
 new = old op val;
} while (!atomic_compare_exchange_strong(addr, &old, new));
```

with `new` being the result of the operation.

If `E1` or `E2` has floating type, then exceptional conditions or floating-point exceptions encountered during discarded evaluations of `new` would also be discarded to satisfy the equivalence of `E1 op= E2` and `E1 = E1 op (E2)`. For example, if Annex F is in effect, the floating types involved have ISO/IEC 60559 binary formats, and `FLT_EVAL_METHOD` is 0, the equivalent code would be:

```
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
/* ... */
fenv_t fenv;
T1 *addr = &E1;
T2 val = E2;
T1 old = *addr;
T1 new;
```

```

 feholdexcept(&fenv);
 for (;;) {
 new = old op val;
 if (atomic_compare_exchange_strong(addr, &old, new))
 break;
 feclearexcept(FE_ALL_EXCEPT);
 }
 feupdateenv(&fenv);

```

If **FLT\_EVAL\_METHOD** is not 0, then *T2* is expected to be a type with the range and precision to which **E2** is evaluated to satisfy the equivalence.

### 6.5.18 Comma operator

#### Syntax

- 1 *expression*:
 

*assignment-expression*  
*expression* , *assignment-expression*

#### Semantics

- 2 The left operand of a comma operator is evaluated as a void expression; there is a sequence point between its evaluation and that of the right operand. Then the right operand is evaluated; the result has its type and value.<sup>114)</sup>
- 3 **EXAMPLE** As indicated by the syntax, the comma operator (as described in this subclause) cannot appear in contexts where a comma is used to separate items in a list (such as arguments to functions or lists of initializers). On the other hand, it can be used within a parenthesized expression or within the second expression of a conditional operator in such contexts. In the function call

```
f(a, (t=3, t+2), c)
```

the function has three arguments, the second of which has the value 5.

**Forward references:** initialization (6.7.11).

<sup>114)</sup>A comma operator does not yield an lvalue.

## 6.6 Constant expressions

### 6.6.1 General

#### Syntax

- 1 *constant-expression:*  
                                   *conditional-expression*

#### Description

- 2 A constant expression can be evaluated during translation rather than runtime, and accordingly can be used in any place that a constant can be.

#### Constraints

- 3 Constant expressions shall not contain assignment, increment, decrement, function-call, or comma operators, except when they are contained within a subexpression that is not evaluated.<sup>115)</sup>  
 4 Each constant expression shall evaluate to a constant that is in the range of representable values for its type.

#### Semantics

- 5 An expression that evaluates to a constant is required in several contexts. If a floating expression is evaluated in the translation environment, the arithmetic range and precision shall be at least as great as if the expression were being evaluated in the execution environment.<sup>116)</sup>  
 6 A compound literal with storage-class specifier **constexpr** is a *compound literal constant*, as is a postfix expression that applies the **.** member access operator to a compound literal constant of structure or union type, even recursively. A compound literal constant is a constant expression with the type and value of the unnamed object.  
 7 An identifier that is:  
     — an enumeration constant,  
     — a predefined constant, or  
     — declared with storage-class specifier **constexpr** and has an object type,

is a *named constant*, as is a postfix expression that applies the **.** member access operator to a named constant of structure or union type, even recursively. For enumeration and predefined constants, their value and type are defined in the respective clauses; for **constexpr** objects, such a named constant is a constant expression with the type and value of the declared object.

- 8 An *integer constant expression*<sup>117)</sup> shall have integer type and shall only have operands that are integer literals, named and compound literal constants of integer type, character literals, **sizeof** or **\_Lengthof** expressions whose results are integer constant expressions, **alignof** expressions, and floating, named, or compound literal constants of arithmetic type that are the immediate operands of casts. Cast operators in an integer constant expression shall only convert arithmetic types to integer types, except as part of an operand to the **typeof** operators, **sizeof** operator, **\_Lengthof** operator, or **alignof** operator.  
 9 More latitude is permitted for constant expressions in initializers. Such a constant expression shall be, or evaluate to, one of the following:  
     — a named constant,

<sup>115)</sup>The operand of a **typeof** (6.7.3.6), **sizeof**, **\_Lengthof**, or **alignof** operator is usually not evaluated (6.5.4.5).

<sup>116)</sup>The use of evaluation formats as characterized by **FLT\_EVAL\_METHOD** and **DEC\_EVAL\_METHOD** also applies to evaluation in the translation environment.

<sup>117)</sup>An integer constant expression is required in contexts such as the size of a bit-field member of a structure, the value of an enumeration constant, and the size of a non-variable length array. Further constraints that apply to the integer constant expressions used in conditional-inclusion preprocessing directives are discussed in 6.10.2.

- a compound literal constant,
  - an arithmetic constant expression,
  - a null pointer constant,
  - an address constant, or
  - an address constant for a complete object type plus or minus an integer constant expression.
- 10 An *arithmetic constant expression* shall have arithmetic type and shall only have operands that are integer literals, floating literals, named or compound literal constants of arithmetic type, character literals, **sizeof** and **\_Lengthof** expressions whose results are integer constant expressions, and **alignof** expressions. Cast operators in an arithmetic constant expression shall only convert arithmetic types to arithmetic types, except as part of an operand to the **typeof** operators, **sizeof** operator, **\_Lengthof** operator, or **alignof** operator.
- 11 An *address constant* is a null pointer,<sup>118)</sup> a pointer to an lvalue designating an object of static storage duration, or a pointer to a function designator; it shall be created explicitly using the unary & operator or an integer constant cast to pointer type, or implicitly using an expression of array or function type.
- 12 The array-subscript [ ] and member-access -> operator, the address & and indirection \* unary operators, and pointer casts can be used in the creation of an address constant, but the value of an object shall not be accessed by use of these operators.<sup>119)</sup>
- 13 A *structure or union constant* is a named constant or compound literal constant with structure or union type, respectively.
- 14 An implementation may accept other forms of constant expressions, called *extended constant expressions*. It is implementation-defined whether extended constant expressions are usable in the same manner as the constant expressions defined in this document, including whether or not extended integer constant expressions are considered to be integer constant expressions.<sup>120)</sup>
- 15 Starting from a structure or union constant, the member-access . operator can be used to form a named constant or compound literal constant as described previously in this subclause.
- 16 If the member-access operator . accesses a member of a union constant, the accessed member shall be the same as the member that is initialized by the union constant's initializer.
- 17 The semantic rules for the evaluation of a constant expression are the same as for nonconstant expressions.<sup>121)</sup>

**Forward references:** array declarators (6.7.7.3), initialization (6.7.11).

## 6.6.2 Constant range expressions

### Syntax

- 1 *constant-range-expression:*  
                                   *constant-expression* ... *constant-expression*

<sup>118)</sup> A named constant or compound literal constant of integer type and value zero is a null pointer constant. A named constant or compound literal constant with a pointer type and a value null is a null pointer but not a null pointer constant; it can only be used to initialize a pointer object if its type implicitly converts to the target type.

<sup>119)</sup> Named constants or compound literal constants with arithmetic type, including names of **constexpr** objects, are valid in offset computations such as array subscripts or in pointer casts, as long as the expressions in which they occur form integer constant expressions. In contrast, names of other objects, even if **const**-qualified and with static storage duration, are not valid.

<sup>120)</sup> For example, in the declaration **int arr\_or\_vla[ (int)+1.0 ] ;**, while possible to be computed by some implementations as an array with a size of one, it is implementation-defined whether this results in a variable length array declaration or a declaration of an array of known constant size of automatic storage duration. The choice depends on whether **(int)+1.0** is an extended integer constant expression.

<sup>121)</sup> Thus, in the following initialization,

```
static int i = 2 || 1 / 0;
```

the expression is a valid integer constant expression with value one.

**Description**

- 2 A *constant range expression* is a special syntactic form that describes a sequence of contiguously incrementing integer values ranging from the arithmetic value of the first constant expression to the second, without listing intermediate values explicitly. The operator is not a value expression and is only permitted in specific contexts, such as the operand to a **case** label to indicate that a single label matches multiple values.

**Constraints**

- 3 The constant expressions shall be integer constant expressions.

**Semantics**

- 4 The values described by the `...` operator form a *closed range*, which contains all integer values in sequence starting from and including the first, low value, up to and including the second, high value.<sup>122)</sup>
- 5 If the arithmetic value of the first constant expression is greater than the one of the second, the range described by the constant range expression is empty.

**Recommended practice**

- 6 Implementations are encouraged to emit a diagnostic message when a range is empty.
- 7 EXAMPLE 1 Range expressions may be used as the operand to a case label to concisely describe a sequence of matching values in one label:

```
extern int n;
extern void f(void);

switch (n) {
 case 2 ... 5: // equivalent to
 // case 2:
 // case 3:
 // case 4:
 // case 5:
 f ();
 break;
}
```

- 8 EXAMPLE 2 Because a range expression describes a closed range, it is possible to match past-the-end values such as the size of an array:

```
constexpr const int N = 42;
int arr[N];

switch (i) {
 case 0 ... N: // matches the past-the-end range of arr
 f (arr[i]); // not OK, will dereference arr[N]
 g (&arr[i]); // may be OK depending on purpose
 break;
}

switch (i) {
 case 0 ... N - 1: // only matches the valid element range of arr
 f (arr[i]); // OK
 break;
}
```

<sup>122)</sup>A range is not itself usable as a value and therefore does not have any specific type or representation, or perform any type conversion.

- 9 EXAMPLE 3 This range expression describes a range with elements that have different types. The implementation creates a range with the values from  $-10$  to  $+10$ , inclusive, in whatever representation it needs in order to preserve the semantics of the expressed range:

```
extern int n;

switch (n) {
 case -10 ... 10ul: // case covers all values from -10 to 10
 break;
}
```

**Forward references:** the **switch** statement (6.8.5.3).

## 6.7 Declarations

### 6.7.1 General

#### Syntax

- 1 *declaration:*
- ```

        declaration-specifiers init-declarator-listopt ;
        attribute-specifier-sequence declaration-specifiers init-declarator-list ;
        static_assert-declaration
        attribute-declaration

declaration-specifiers:
        declaration-specifier attribute-specifier-sequenceopt
        declaration-specifier declaration-specifiers

declaration-specifier:
        storage-class-specifier
        type-specifier-qualifier
        function-specifier

init-declarator-list:
        init-declarator
        init-declarator-list , init-declarator

init-declarator:
        declarator
        declarator = initializer

attribute-declaration:
        attribute-specifier-sequence ;

simple-declaration:
        attribute-specifier-sequenceopt declaration-specifiers declarator = initializer

```

Constraints

- 2 If a declaration other than a `static_assert` or attribute declaration does not include an init declarator list, its declaration specifiers shall include one of the following:
- a struct or union specifier or enum specifier that includes a tag, with the declaration being of a form specified in 6.7.3.4 to declare that tag;
 - an enum specifier that includes an enumerator list.
- 3 EXAMPLE 1 The following are invalid, because the declared tag or enumeration constants are in a nested construct, rather than a declaration specifier of the declaration being of one of the given forms:
- ```

struct { struct s2 { int x2a; } x2b; };
typedef (struct s3 { int x3; });
alignas (struct s4 { int x4; }) int;
typedef (struct s5 *);
typedef (enum { E6 });
struct { void (*p)(struct s7 *); };

```
- 4 If an identifier has no linkage, there shall be no more than one declaration of the identifier (in a declarator or type specifier) with the same scope and in the same name space, except that:
- a typedef name can be redefined to denote the same type as it currently does, provided that type is not a variably modified type;
  - enumeration constants and tags can be redeclared as specified in 6.7.3.3 and 6.7.3.4, respectively.



- 5 All declarations in the same scope that refer to the same object or function shall specify compatible types.
- 6 If an identifier for an object that does not have function prototype scope is declared with no linkage, the type for the object shall be complete by the end of its declarator, or by the end of its init-declarator if it has an initializer.<sup>123)</sup>

### Semantics

- 7 A declaration specifies the interpretation and properties of a set of identifiers. A *definition* of an identifier is a declaration for that identifier that for:
  - an object, causes storage to be reserved for that object,
  - a function, includes the function body,<sup>124)</sup>
  - an enumeration constant, is the first (or only) declaration of the identifier, or
  - a typedef name, is the first (or only) declaration of the identifier.
- 8 The declaration specifiers consist of a sequence of specifiers, followed by an optional attribute specifier sequence. The declaration specifiers indicate the linkage, storage duration, and part of the type of the entities that the declarators denote. The init declarator list is a comma-separated sequence of declarators, each of which can have additional type information, or an initializer, or both. The declarators contain the identifiers (if any) being declared. The optional attribute specifier sequence in a declaration appertains to each of the entities declared by the declarators of the init declarator list.
- 9 The optional attribute specifier sequence terminating a sequence of declaration specifiers appertains to the type determined by the preceding sequence of declaration specifiers. The attribute specifier sequence affects the type only for the declaration it appears in, not other declarations involving the same type.
- 10 Except where specified otherwise, the meaning of an attribute declaration is implementation-defined.
- 11 EXAMPLE 2 In the declaration for an entity, attributes appertaining to that entity can appear at the start of the declaration and after the identifier for that declaration.

```
[[deprecated]] void f [[deprecated]] (void); // valid
```

- 12 A declaration such that the declaration specifiers contain no type specifier or that is declared with **constexpr** is said to be *underspecified*. If such a declaration is not a definition, if it declares no or more than one ordinary identifier, if the declared identifier already has a declaration in the same scope, if the declared entity is not an object, or if anywhere within the sequence of tokens making up the declaration identifiers that are not ordinary are declared, the behavior is implementation-defined.<sup>125)</sup>
- 13 A *simple declaration* is a declaration<sup>126)</sup> that can appear in place of the controlling expression of a selection statement.
- 14 EXAMPLE 3 As declarations using **constexpr** are underspecified, the following has implementation-defined behavior because tokens within the declaration declare **s** which is not an ordinary identifier:

```
constexpr typeof(struct s *) x = 0;
```

**Forward references:** declarators (6.7.7), enumeration specifiers (6.7.3.3), initialization (6.7.11), storage-class specifiers (6.7.2), type inference (6.7.10), type names (6.7.8), type qualifiers (6.7.4), selection statements (6.8.5).

<sup>123)</sup>In function definitions (but not in function declarations that are not part of a function definition) the adjusted type of a parameter needs to be complete (see 6.7.7.4).

<sup>124)</sup>Function definitions have a different syntax, described in 6.9.2.

<sup>125)</sup>It is recommended that implementations that accept such declarations follow the semantics of the corresponding feature in ISO/IEC 14882.

<sup>126)</sup>subject to the same constraints as all other declarations with initializers

## 6.7.2 Storage-class specifiers

### Syntax

1 *storage-class-specifier*:

```

 auto
 constexpr
 extern
 register
 static
 thread_local
 typedef

```

### Constraints

- 2 At most, one storage-class specifier can be given in the declaration specifiers in a declaration, except that:
- **thread\_local** can appear with **static** or **extern**,
  - **auto** can appear with all the others except **typedef**,<sup>127)</sup> and
  - **constexpr** can appear with **auto**, **register**, or **static**.
- 3 In the declaration of an object with block scope, if the declaration specifiers include **thread\_local**, they shall also include either **static** or **extern**. If **thread\_local** appears in any declaration of an object, it shall be present in every declaration of that object.
- 4 **thread\_local** shall not appear in the declaration specifiers of a function declaration. **auto** shall only appear in the declaration specifiers of an identifier with file scope or along with other storage-class specifiers if the type is to be inferred from an initializer.
- 5 An object declared with storage-class specifier **constexpr** or any of its members, even recursively, shall not have an atomic type, or a variably modified type, or a type that is **volatile** or **restrict** qualified. An initializer of floating type shall be evaluated with the translation-time floating-point environment. The declaration shall be a definition and shall have an initializer.<sup>128)</sup> The value of any constant expressions or of any character in a string literal of the initializer shall be exactly representable in the corresponding target type; no change of value shall be applied.<sup>129)</sup>
- 6 If an object or subobject declared with storage-class specifier **constexpr** has pointer, integer, or arithmetic type, any explicit initializer value for it shall be null,<sup>130)</sup> an integer constant expression, or an arithmetic constant expression, respectively. If the object declared has real floating type, the initializer shall have integer or real floating type. If the initializer has decimal floating type, the object declared shall have decimal floating type and the conversion shall preserve the quantum of the initializer. If the initializer has real type and a signaling NaN value, the unqualified versions of the type of the initializer and the corresponding real type of the object declared shall be compatible.
- 7 The declaration of an identifier for a function that has block scope shall have no explicit storage-class specifier other than **extern**.
- 8 EXAMPLE 1 Although in the following the expression **A.p** is not a null pointer constant, only a constant expression with pointer type and a null pointer value, the member-wise initialization of **B** with **A** is valid.

```
struct s { void *p; };
```

<sup>127)</sup>See “future language directions” (6.11.6).

<sup>128)</sup>All assignment expressions of such an initializer, if any, are constant expressions or string literals, see 6.7.11.

<sup>129)</sup>In the context of arithmetic conversions, 6.3.2 describes the details of changes of value that occur if values of arithmetic expressions are stored in the objects that for example have a different signedness, excess precision or quantum exponent. Whenever such a change of value is necessary, the constraint is violated.

<sup>130)</sup>The named constant or compound literal constant corresponding to an object declared with storage-class specifier **constexpr** and pointer type is a constant expression with a value null, and thus a null pointer and an address constant. Thus, such a named constant is a valid initializer for other **constexpr** declarations, provided the pointer types match accordingly. However, even if it has type **void\*** it is not a null pointer constant.

```
constexpr struct s A = { nullptr };
constexpr struct s B = A;
```

- 9 EXAMPLE 2 Pointers can be initialized to eligible constant expressions, such as a null pointer constant:

```
constexpr int *p = {}; // Default initialization with a null pointer
```

- 10 EXAMPLE 3

```
void f (void) {
 constexpr float f = 1.0f;
 constexpr float g = 3.0f;
 fesetround(FE_TOWARDSZERO); // does not affect
 // the following initialization
 // of "h"
 constexpr float h = f / g;
 // ...
}
```

## Semantics

- 11 Storage-class specifiers specify various properties of identifiers and declared features:
- storage duration (**static** in block scope, **thread\_local**, **auto**, **register**),
  - linkage (**extern**, **static** and **constexpr** in file scope, **typedef**),
  - value (**constexpr**), and
  - type (**typedef**).
- 12 The meanings of the various linkages and storage durations were discussed in 6.2.2 and 6.2.4, **typedef** is discussed in 6.7.9, and type inference using **auto** is discussed in 6.7.10.
- 13 A declaration of an identifier for an object with storage-class specifier **register** suggests that access to the object be as fast as possible. The extent to which such suggestions are effective is implementation-defined.<sup>131)</sup>
- 14 If an aggregate or union object is declared with a storage-class specifier other than **typedef**, the properties resulting from the storage-class specifier, except with respect to linkage, also apply to the members of the object, including recursively for any aggregate or union member objects.
- 15 If **auto** appears with another storage-class specifier, or if it appears in a declaration at file scope, it is ignored for the purposes of determining a storage duration or linkage. In this case, it indicates only that the declared type can be inferred.
- 16 An object declared with a storage-class specifier **constexpr** has its value permanently fixed at translation-time; if not yet present, a **const**-qualification is implicitly added to the object's type. The declared identifier is considered a constant expression of the respective kind, see 6.6.1.
- 17 NOTE 1 An object declared in block scope with a storage-class specifier **constexpr** and without **static** has automatic storage duration, the identifier has no linkage, and each instance of the object has a unique address obtainable with & (if it is not declared with the **register** specifier), if any. Such an object in file scope has static storage duration, the corresponding identifier has internal linkage, and each translation unit that sees the same textual definition implements a separate object with a distinct address.
- 18 NOTE 2 The constraints for **constexpr** objects are intended to enforce checks for portability at translation time.

<sup>131)</sup>The implementation can treat any **register** declaration simply as an **auto** declaration. However, whether or not addressable storage is used, the address of any part of an object declared with storage-class specifier **register** cannot be computed, either explicitly (by use of the unary & operator as discussed in 6.5.4.3) or implicitly (by converting an array name to a pointer as discussed in 6.3.3.1). Thus, the only operators that can be applied to an array declared with storage-class specifier **register** are **sizeof**, **\_Lengthof**, and the typeof operators.

```
constexpr unsigned int minusOne = -1; // constraint violation
constexpr unsigned int uint_max = -1U; // ok
constexpr double onethird = 1.0/3.0; // possible constraint violation
constexpr double onethirdtrunc = (double)(1.0/3.0); // ok
constexpr _Decimal32 small = DEC64_TRUE_MIN * 0; // constraint violation
```

If a truncation of excess precision changes the value in the initializer of `onethird`, a constraint is violated and a diagnostic is required. In contrast to that, the explicit conversion in the initializer for `onethirdtrunc` ensures that the definition is valid. Similarly, the initializer of `small` has a quantum exponent that is larger than the largest possible quantum exponent for `_Decimal32`.

- 19 NOTE 3 Similarly, implementation-defined behavior related to the `char` type of the elements of the string literal `"\xFF"` can cause constraint violations at translation time:

```
constexpr char string[] = { "\xFF", }; // ok
constexpr char8_t u8string[] = { u8"\xFF", }; // ok
constexpr unsigned char ucstring[] = { "\xFF", }; // possible constraint
 // violation
```

In both the string and `ucstring` initializers, the initializer is a (brace-enclosed) string literal of type `char`. If the type `char` is capable of representing negative values and its width is 8, then the preceding code is equivalent to:

```
constexpr char string[] = { -1, 0, }; // ok
constexpr char8_t u8string[] = { 255, 0, }; // ok
constexpr unsigned char ucstring[] = { -1, 0, }; // constraint violation
```

The hexadecimal escape sequence results in a value of 255. For an initializer of type `char`, it is converted to a signed 8-bit integer, making a value of -1. A negative value does not fit within the range of values for `unsigned char`, and therefore the initialization of `ucstring` is a constraint violation under the previously stated implementation conditions. In the case where `char` is not capable of representing negative values, the original snippet is equivalent to the following and there is no constraint violation.

```
constexpr char string[] = { 255, 0, }; // ok
constexpr char8_t u8string[] = { 255, 0, }; // ok
constexpr unsigned char ucstring[] = { 255, 0, }; // ok
```

- 20 EXAMPLE 4 An identifier declared with the `constexpr` specifier can have its value used in constant expressions:

```
constexpr int K = 47;
enum {
 A = K, // valid, constant initialization
};
constexpr int L = K; // valid, constexpr initialization
static int b = K + 1; // valid, static initialization
int array[K]; // not a VLA
```

- 21 EXAMPLE 5 This example illustrates `constexpr` initializations involving different type domains, decimal and non-decimal floating types, NaNs and infinities, and quanta in decimal floating types.

```
#include <float.h>
#include <complex.h>

constexpr float _Complex fc1 = 1.0; // ok
constexpr float _Complex fc2 = 0.1; // constraint violation, unless double
 // has the same precision as float
 // and is evaluated with the same
 // precision
constexpr float _Complex fc3 = 3*I; // ok

constexpr double d1 = (double _Complex)1.0; // constraint violation
```

```

constexpr float f1 = (long double)INFINITY; // ok
constexpr float f2 = (long double)NAN; // ok, quiet NaNs in real floating
 // types are considered the same
 // value, regardless of payloads

constexpr double d2 = DBL_SNaN; // ok
constexpr double d3 = FLT_SNaN; // constraint violation, even if float
 // and double have the same format

constexpr double _Complex dc1 = DBL_SNaN; // ok
constexpr double _Complex dc2 = CMPLX(DBL_SNaN, 0.); // ok
constexpr double _Complex dc3 = CMPLX(0., DBL_SNaN); // ok

constexpr _Decimal32 d321 = 1.0; // ok
constexpr _Decimal32 d322 = 1; // ok
constexpr _Decimal32 d323 = INFINITY; // ok
constexpr _Decimal32 d324 = NAN; // ok
constexpr _Decimal64 d641 = DEC64_SNaN; // ok
constexpr _Decimal64 d642 = DEC32_SNaN; // constraint violation
constexpr float f3 = 1.DF; // constraint violation
constexpr float f4 = DEC_INFINITY; // constraint violation
constexpr double d4 = DEC_NaN; // constraint violation
constexpr _Decimal32 d325 = DEC64_TRUE_MIN * 0; // constraint violation,
 // quantum not preserved

```

- 22 EXAMPLE 6 An object declared with the **constexpr** specifier stores the exact value of its initializer, no implicit value change is applied:

```

#include <float.h>

constexpr int A = 42LL; // valid, 42 always fits in an int
constexpr signed short B = ULLONG_MAX; // constraint violation, value never
 // fits
constexpr float C = 47u; // valid, exactly representable
 // in float

#if FLT_MANT_DIG > 24
constexpr float D = 5369000000; // constraint violation if float is
 // ISO/IEC 60559 binary32
#endif

#if (FLT_MANT_DIG == DBL_MANT_DIG) \
 && (0 <= FLT_EVAL_METHOD) \
 && (FLT_EVAL_METHOD <= 1)
constexpr float E = 1.0 / 3.0; // only valid if double expressions
 // and float objects have the same
 // precision
#endif

#if FLT_EVAL_METHOD == 0
constexpr float F = 1.0f / 3.0f; // valid, same type and precision
#else
constexpr float F = (float)(1.0f / 3.0f); // needs cast to truncate the
 // excess precision
#endif

```

- 23 EXAMPLE 7 This recursively applies to initializers for all elements of an aggregate object declared with the **constexpr** specifier:

```

constexpr static unsigned short array[] = {
 3000, // valid, fits in unsigned short range
 300000, // constraint violation if short is 16-bit

```

```

-1 // constraint violation, target type is unsigned
};

struct S {
 int x, y;
};
constexpr struct S s = {
 .x = INT_MAX, // valid
 .y = UINT_MAX, // constraint violation
};

```

Forward references: type definitions (6.7.9), type inference (6.7.10).

## 6.7.3 Type specifiers

### 6.7.3.1 General

#### Syntax

- 1 *type-specifier*:
  - void**
  - char**
  - short**
  - int**
  - long**
  - float**
  - double**
  - signed**
  - unsigned**
  - \_BitInt** ( *constant-expression* )
  - bool**
  - \_Complex**
  - \_Decimal32**
  - \_Decimal64**
  - \_Decimal128**
  - atomic-type-specifier*
  - struct-or-union-specifier*
  - enum-specifier*
  - typedef-name*
  - typeof-specifier*

#### Constraints

- 2 Except where the type is inferred (6.7.10), at least one type specifier shall be given in the declaration specifiers in each declaration, and in the specifier-qualifier list in each member declaration and type name. Each list of type specifiers shall be one of the following multisets (delimited by commas, when there is more than one multiset per item); the type specifiers can occur in any order, possibly intermixed with the other declaration specifiers.

- **void**
- **char**
- **signed char**
- **unsigned char**
- **short, signed short, short int, or signed short int**
- **unsigned short, or unsigned short int**
- **int, signed, or signed int**

- **unsigned**, or **unsigned int**
- **long**, **signed long**, **long int**, or **signed long int**
- **unsigned long**, or **unsigned long int**
- **long long**, **signed long long**, **long long int**, or **signed long long int**
- **unsigned long long**, or **unsigned long long int**
- **\_BitInt**( *constant-expression* ), or **signed \_BitInt**( *constant-expression* )
- **unsigned \_BitInt**( *constant-expression* )
- **float**
- **double**
- **long double**
- **\_Decimal32**
- **\_Decimal64**
- **\_Decimal128**
- **bool**
- **float \_Complex**
- **double \_Complex**
- **long double \_Complex**
- atomic type specifier
- struct or union specifier
- enum specifier
- typedef name
- typeof specifier

- 3 The type specifier **\_Complex** shall not be used if the implementation does not support complex types, and the type specifiers **\_Decimal32**, **\_Decimal64**, and **\_Decimal128** shall not be used if the implementation does not support decimal floating types (see 6.10.10.4).
- 4 The parenthesized constant expression that follows the **\_BitInt** keyword shall be an integer constant expression *N* that specifies the width (6.2.6.2) of the type. The value of *N* for **unsigned \_BitInt** shall be greater than or equal to 1. The value of *N* for **\_BitInt** shall be greater than or equal to 2. The value of *N* shall be less than or equal to the value of **BITINT\_MAXWIDTH** (see 5.3.5.3.2).

#### Semantics

- 5 Specifiers for structures, unions, enumerations, atomic types, and typeof specifiers are discussed in 6.7.3.2 through 6.7.3.6. Declarations of typedef names are discussed in 6.7.9. The characteristics of the other types are discussed in 6.2.5.
- 6 For a declaration such that the declaration specifiers contain no type specifier a mechanism to infer the type from an initializer is discussed in 6.7.10. In such a declaration, optional elements, if any, of a sequence of declaration specifiers appertain to the inferred type (for qualifiers and attribute specifiers) or to the declared objects (for alignment specifiers).
- 7 Each of the comma-separated multisets designates the same type, except that for bit-fields, it is implementation-defined whether the specifier **int** designates the same type as **signed int** or the same type as **unsigned int**.

**Forward references:** atomic type specifiers (6.7.3.5), enumeration specifiers (6.7.3.3), structure and union specifiers (6.7.3.2), tags (6.7.3.4), type definitions (6.7.9).

### 6.7.3.2 Structure and union specifiers

#### Syntax

- 1 *struct-or-union-specifier*:
- ```

struct-or-union attribute-specifier-sequenceopt identifieropt { member-declaration-list }
struct-or-union attribute-specifier-sequenceopt identifier

```

struct-or-union:

struct
union

member-declaration-list:

member-declaration
member-declaration-list member-declaration

member-declaration:

*attribute-specifier-sequence*_{opt} *specifier-qualifier-list* *member-declarator-list*_{opt} ;
static_assert-declaration

specifier-qualifier-list:

type-specifier-qualifier *attribute-specifier-sequence*_{opt}
type-specifier-qualifier *specifier-qualifier-list*

type-specifier-qualifier:

type-specifier
type-qualifier
alignment-specifier

member-declarator-list:

member-declarator
member-declarator-list , *member-declarator*

member-declarator:

declarator
*declarator*_{opt} : *constant-expression*

Constraints

- 2 A member declaration that does not declare an anonymous structure or anonymous union shall contain a member declarator list.
- 3 A structure or union shall not contain a member with incomplete or function type (hence, a structure shall not contain an instance of itself, but can contain a pointer to an instance of itself), except that the last member of a structure with more than one named member can have incomplete array type; such a structure (and any union containing, possibly recursively, a member that is such a structure) shall not be a member of a structure or an element of an array.
- 4 The expression that specifies the width of a bit-field shall be an integer constant expression with a nonnegative value that does not exceed the width of an object of the type that would be specified were the colon and expression omitted.¹³²⁾ If the value is zero, the declaration shall have no declarator.

¹³²⁾While the number of bits in a **bool** object is at least **CHAR_BIT**, the width of a **bool** is just 1 bit.

- 5 A bit-field shall have a type that is a qualified or unqualified **bool**, **signed int**, **unsigned int**, a bit-precise integer type, or other implementation-defined type. It is implementation-defined whether atomic types are permitted.
- 6 An attribute specifier sequence shall not appear in a struct-or-union specifier without a member declaration list, except in a declaration of the form:

struct-or-union attribute-specifier-sequence identifier ;

The attributes in the attribute specifier sequence, if any, are thereafter considered attributes of the **struct** or **union** whenever it is named.

Semantics

- 7 As discussed in 6.2.5, a structure is a type consisting of a sequence of members, whose storage is allocated in an ordered sequence, and a union is a type consisting of a sequence of members whose storage overlap.
- 8 Structure and union specifiers have the same form. The keywords **struct** and **union** indicate that the type being specified is, respectively, a structure type or a union type.
- 9 The optional attribute specifier sequence in a struct-or-union specifier appertains to the structure or union type being declared. The optional attribute specifier sequence in a member declaration appertains to each of the members declared by the member declarator list; it shall not appear if the optional member declarator list is omitted. The optional attribute specifier sequence in a specifier qualifier list appertains to the type denoted by the preceding type specifier qualifiers. The attribute specifier sequence affects the type only for the member declaration or type name it appears in, not other types or declarations involving the same type.
- 10 The member declaration list is a sequence of declarations for the members of the structure or union. If the member declaration list does not contain any named members, either directly or via an anonymous structure or anonymous union, the behavior is implementation-defined.¹³³⁾
- 11 A member of a structure or union can have any complete object type other than a variably modified type.¹³⁴⁾ In addition, a member can be declared to consist of a specified number of bits (including a sign bit, if any). Such a member is called a *bit-field*;¹³⁵⁾ its width is preceded by a colon.
- 12 A bit-field is interpreted as having a signed or unsigned integer type consisting of the specified number of bits.¹³⁶⁾ If the value 0 or 1 is stored into a nonzero-width bit-field of type **bool**, the value of the bit-field shall compare equal to the value stored; a **bool** bit-field has the semantics of a **bool**.
- 13 An implementation can allocate any addressable storage unit large enough to hold a bit-field. If enough space remains, a bit-field that immediately follows another bit-field in a structure shall be packed into adjacent bits of the same unit. If insufficient space remains, whether a bit-field that does not fit is put into the next unit or overlaps adjacent units is implementation-defined. The order of allocation of bit-fields within a unit (high-order to low-order or low-order to high-order) is implementation-defined. The alignment of the addressable storage unit is unspecified.
- 14 A bit-field declaration with no declarator, but only a colon and a width, indicates an unnamed bit-field.¹³⁷⁾ As a special case, a bit-field structure member with a width of zero indicates that no further bit-field is to be packed into the unit in which the previous bit-field, if any, was placed.
- 15 An unnamed member whose type specifier is a structure specifier with no tag is called an *anonymous structure*; an unnamed member whose type specifier is a union specifier with no tag is called an *anonymous union*. The members of an anonymous structure or union are members of the containing structure or union, keeping their structure or union layout. This applies recursively if the containing

¹³³⁾For further rules affecting compatibility and completeness of structure or union types, see 6.2.7 and 6.7.3.4.

¹³⁴⁾A structure or union cannot contain a member with a variably modified type because member names are not ordinary identifiers as defined in 6.2.3.

¹³⁵⁾The unary & (address-of) operator cannot be applied to a bit-field object; thus, there are no pointers to or arrays of bit-field objects.

¹³⁶⁾As specified in 6.7.3, if the actual type specifier used is **int** or a typedef-name defined as **int**, then it is implementation-defined whether the bit-field is signed or unsigned. This includes an **int** type specifier produced using the typedef specifiers (6.7.3.6).

¹³⁷⁾An unnamed bit-field structure member is useful for padding to conform to externally imposed layouts.

structure or union is also anonymous.

- 16 Each non-bit-field member of a structure or union object is aligned in an implementation-defined manner appropriate to its type.
- 17 Within a structure object, the non-bit-field members and the units in which bit-fields reside have addresses that increase in the order in which they are declared. A pointer to a structure object, suitably converted, points to its initial member (or if that member is a bit-field, then to the unit in which it resides), and vice versa. There can be unnamed padding within a structure object, but not at its beginning.
- 18 The size of a union is sufficient to contain the largest of its members. The value of at most one of the members can be stored in a union object at any time. A pointer to a union object, suitably converted, points to each of its members (or if a member is a bit-field, then to the unit in which it resides), and vice versa. The members of a union object overlap in such a way that pointers to them when converted to pointers to character type point to the same byte. There can be unnamed padding at the end of a union object, but not at its beginning.
- 19 There can be unnamed padding at the end of a structure or union.
- 20 As a special case, the last member of a structure with more than one named member can have an incomplete array type; this is called a *flexible array member*. In most situations, the flexible array member is ignored. In particular, the size of the structure is as if the flexible array member were omitted except that it may have more trailing padding than the omission would imply. However, when a `.` (or `->`) operator has a left operand that is (a pointer to) a structure with a flexible array member and the right operand names that member, it behaves as if that member were replaced with the longest array (with the same element type) that would not make the structure larger than the object being accessed; the offset of the array shall remain that of the flexible array member, even if this would differ from that of the replacement array. If this array would have no elements, it behaves as if it had one element but the behavior is undefined if any attempt is made to access that element or to generate a pointer one past it.
- 21 EXAMPLE 1 The following declarations illustrate the behavior when an attribute is written on a tag declaration:

```

struct [[deprecated]] S; // valid, [[deprecated]] appertains to struct S
void f(struct S *s);      // valid, the struct S type has the [[deprecated]]
                          // attribute
struct S {                // valid, struct S inherits the [[deprecated]] attribute
    int a;                  // from the previous declaration
};
void g(struct [[deprecated]] S s); // invalid

```

- 22 EXAMPLE 2 The following illustrates anonymous structures and unions:

```

struct v {
    union {                // anonymous union
        struct { int i, j; }; // anonymous structure
        struct { long k, l; } w;
    };
    int m;
} v1;

v1.i = 2; // valid
v1.k = 3; // invalid: inner structure is not anonymous
v1.w.k = 5; // valid

```

- 23 EXAMPLE 3 After the declaration:

```
struct s { int n; double d[]; };
```

the structure **struct** s has a flexible array member d. A typical way to use this is:

```
int m = /* some value */;
struct s *p = malloc(sizeof(struct s) + sizeof(double [m]));
```

and assuming that the call to `malloc` succeeds, the object pointed to by `p` behaves, for most purposes, as if `p` had been declared as:

```
struct { int n; double d[m]; } *p;
```

(there are circumstances in which this equivalence is broken; in particular, the offsets of member `d` can sometimes not be the same).

- 24 Following the declaration of the previous example:

```
struct s t1 = { 0 };           // valid
struct s t2 = { 1, { 4.2 } }; // invalid
t1.n = 4;                     // valid
t1.d[0] = 4.2;                // can be undefined behavior
```

The initialization of `t2` is invalid (and violates a constraint) because `struct s` is treated as if it did not contain member `d`. The assignment to `t1.d[0]` has probably undefined behavior, but it is possible that

```
sizeof(struct s) >= offsetof(struct s, d) + sizeof(double)
```

in which case the assignment would be legitimate. Nevertheless, it cannot appear in strictly conforming code.

- 25 After the further declaration:

```
struct ss { int n; };
```

the expressions:

```
sizeof(struct s) >= sizeof(struct ss)
sizeof(struct s) >= offsetof(struct s, d)
```

are always equal to 1.

- 26 If `sizeof(double)` is 8, then after the following code is executed:

```
struct s *s1;
struct s *s2;
s1 = malloc(sizeof(struct s) + 64);
s2 = malloc(sizeof(struct s) + 46);
```

and assuming that the calls to `malloc` succeed, the objects pointed to by `s1` and `s2` behave, for most purposes, as if the identifiers had been declared as:

```
struct { int n; double d[8]; } *s1;
struct { int n; double d[5]; } *s2;
```

- 27 Following the further successful assignments:

```
s1 = malloc(sizeof(struct s) + 10);
s2 = malloc(sizeof(struct s) + 6);
```

they then behave as if the declarations were:

```
struct { int n; double d[1]; } *s1, *s2;
```

and:

```
double *dp;
dp = &(s1->d[0]); // valid
```

```

*dp = 42;           // valid
dp = &(s2->d[0]);   // valid
*dp = 42;           // undefined behavior

```

- 28 The assignment:

```
*s1 = *s2;
```

only copies the member **n**; if any of the array elements are within the first **sizeof(struct s)** bytes of the structure, they are set to an indeterminate representation, that can sometimes not coincide with a copy of the representation of the elements of the source array.

- 29 EXAMPLE 4 Because members of anonymous structures and unions are considered to be members of the containing structure or union, **struct s** in the following example has more than one named member and thus the use of a flexible array member is valid:

```

struct s {
    struct { int i; };
    int a[];
};

```

Forward references: declarators (6.7.7), tags (6.7.3.4).

6.7.3.3 Enumeration specifiers

Syntax

- 1 *enum-specifier*:
- ```

enum attribute-specifier-sequenceopt identifieropt enum-type-specifieropt
 { enumerator-list }
enum attribute-specifier-sequenceopt identifieropt enum-type-specifieropt
 { enumerator-list , }
enum identifier enum-type-specifieropt

```
- enumerator-list*:
- ```

enumerator
enumerator-list , enumerator

```
- enumerator*:
- ```

enumeration-constant attribute-specifier-sequenceopt
enumeration-constant attribute-specifier-sequenceopt = constant-expression

```
- enum-type-specifier*:
- ```

: specifier-qualifier-list

```

- 2 All enumerations have an *underlying type*. The underlying type can be explicitly specified using an enum type specifier and is its *fixed underlying type*. If it is not explicitly specified, the underlying type is the enumeration's compatible type, which is either **char** or a standard or extended signed or unsigned integer type.

Constraints

- 3 For an enumeration with a fixed underlying type, the integer constant expression defining the value of the enumeration constant shall be representable in that fixed underlying type. If the value of an enumeration constant without a defining constant expression for an enumeration with fixed underlying type is obtained by adding 1 to the previous enumeration constant, the value of that previous enumeration constant shall not be the maximum value of the underlying type.
- 4 For an enumeration without a fixed underlying type, the expression that defines the value of an enumeration constant shall be an integer constant expression. For all the integer constant expressions which make up the values of the enumeration constants, there shall be a type capable of representing all the values that is a standard or extended signed or unsigned integer type, or **char**.
- 5 If an enum type specifier is present, then the longest possible sequence of tokens that can be interpreted as a specifier qualifier list is interpreted as part of the enum type specifier. It shall name

an integer type that is neither an enumeration nor a bit-precise integer type. The underlying type of the enumeration is the unqualified, non-atomic version of the type specified by the type specifiers in the specifier qualifier list.¹³⁸⁾

- 6 An enum specifier of the form

enum *identifier* *enum-type-specifier*

shall not appear except in a declaration of the form

enum *identifier* *enum-type-specifier* ;

unless it is immediately followed by an opening brace, an enumerator list (with an optional ending comma), and a closing brace.

- 7 If two enum specifiers that include an enum type specifier declare the same type, the underlying types shall be compatible.
- 8 An enumeration with a fixed underlying type shall be defined with an enum type specifier. No enum specifier for an enumeration without a fixed underlying type shall include an enum type specifier.

Semantics

- 9 The optional attribute specifier sequence in the **enum** specifier appertains to the enumeration; the attributes in that attribute specifier sequence are thereafter considered attributes of the enumeration whenever it is named. The optional attribute specifier sequence in the enumerator appertains to that enumerator.
- 10 The identifiers in an enumerator list are declared as constants of the types specified later in this subclause and can appear wherever such are permitted.¹³⁹⁾ An enumerator with = defines its enumeration constant as the value of the constant expression. If the first enumerator has no =, the value of its enumeration constant is zero. Each subsequent enumerator with no = defines its enumeration constant as the value of the constant expression obtained by adding 1 to the value of the previous enumeration constant. (The use of enumerators with = may produce enumeration constants with values that duplicate other values in the same enumeration.) The enumerators of an enumeration are also known as its members.
- 11 The type for the members of an enumeration is called the *enumeration member type*.
- 12 During the processing of each enumeration constant in the enumerator list, the type of the enumeration constant shall be:

- the previously declared type, if it is a redeclaration of the same enumeration constant; or,
- the enumerated type, for an enumeration with fixed underlying type; or,
- **int**, if there are no previous enumeration constants in the enumerator list and no explicit = with a defining integer constant expression; or,
- **int**, if given explicitly with = and the value of the integer constant expression is representable by an **int**; or,
- the type of the integer constant expression, if given explicitly with = and if the value of the integer constant expression is not representable by **int**; or,
- the type of the value from the previous enumeration constant with one added to it. If such an integer constant expression would overflow or wraparound the value of the previous enumeration constant from the addition of one, the type takes on either:

¹³⁸⁾The specifier qualifier list is not a context listed in 6.7.6 as permitted for alignment specifiers, so the presence of an alignment specifier in the list violates a constraint.

¹³⁹⁾Thus, the identifiers of enumeration constants declared in the same scope are all required to be distinct from each other and from other identifiers declared in ordinary declarators.

- a suitably sized signed integer type, excluding the bit-precise signed integer types, capable of representing the value of the previous enumeration constant plus one; or,
- a suitably sized unsigned integer type, excluding the bit-precise unsigned integer types, capable of representing the value of the previous enumeration constant plus one.

A signed integer type is chosen if the previous enumeration constant being added is of signed integer type. An unsigned integer type is chosen if the previous enumeration constant is of unsigned integer type. If there is no suitably sized integer type described previously which can represent the new value, then the enumeration has no type which can represent all its values.¹⁴⁰⁾

- 13 For all enumerations without a fixed underlying type, each enumerated type shall be compatible with **char** or a signed or an unsigned integer type that is not **bool** or a bit-precise integer type. The choice of type is implementation-defined,¹⁴¹⁾ but shall be capable of representing the values of all the members of the enumeration.¹⁴²⁾
- 14 Enumeration constants can be redefined in the same scope with the same value as part of a redeclaration of the same enumerated type.
- 15 The enumeration member type for an enumerated type without fixed underlying type upon completion is:
 - **int** if all the values of the enumeration are representable as an **int**; or,
 - the enumerated type.¹⁴³⁾
- 16 The enumeration member type for an enumerated type with fixed underlying type is the enumerated type. The enumerated type is compatible with the underlying type of the enumeration. After possible lvalue conversion a value of the enumerated type behaves the same as the value with the underlying type, in particular with all aspects of promotion, conversion, and arithmetic. Conversion to the enumerated type has the same semantics as conversion to the underlying type.¹⁴⁴⁾
- 17 EXAMPLE 1 The following fragment:

```
enum hue { chartreuse, burgundy, claret=20, winedark };
enum hue col, *cp;
col = claret;
cp = &col;
if (*cp != burgundy)
    /* ... */
```

makes **hue** the tag of an enumeration, and then declares **col** as an object that has that type and **cp** as a pointer to an object that has that type. The enumerated values are in the set {0, 1, 20, 21}.

- 18 EXAMPLE 2 Even if the value of an enumeration constant is generated by the implicit addition of one, an enumeration with a fixed underlying type does not exhibit typical overflow behavior:

```
#include <limits.h>

enum us : unsigned short {
    us_max = USHRT_MAX,
    us_violation, /* Constraint violation:
```

¹⁴⁰⁾Therefore, a constraint has been violated.

¹⁴¹⁾An implementation can delay the choice of which integer type until all enumeration constants have been seen.

¹⁴²⁾For further rules affecting compatibility and completeness of enumerated types see 6.2.7 and 6.7.3.4.

¹⁴³⁾The integer type selected during processing of the enumerator list (before completion) of the enumeration can sometimes not be the same as the compatible implementation-defined integer type selected for the completed enumeration.

¹⁴⁴⁾This means in particular that if the compatible type is **bool**, values of the enumerated type behave in all aspects the same as **bool**, conversion to the enumerated type behaves the same as **bool** (6.3.2.2), and the members only have values **false** and **true**. If it is a signed integer type and the constant expression of an enumeration constant overflows, a constraint for constant expressions (6.6.1) is violated.

```

        USHRT_MAX + 1 would wraparound. */
    us_violation_2 = us_max + 1, /* Maybe constraint violation:
        USHRT_MAX + 1 can be promoted to "int", and
        result is too wide for the
        underlying type. */
    us_wraparound_to_zero = (unsigned short)(USHRT_MAX + 1) /* Okay:
        conversion done in constant expression@/before
        conversion to underlying type:unsigned
        semantics okay.

};

enum ui : unsigned int {
    ui_max = UINT_MAX,
    ui_violation, /* Constraint violation:
        UINT_MAX + 1 would wraparound. */
    ui_no_violation = ui_max + 1, /* Okay: Arithmetic performed as typical
        unsigned integer arithmetic: conversion
        from a value that is already 0 to 0. */
    ui_wraparound_to_zero = (unsigned int)(UINT_MAX + 1) /* Okay: conversion
        done in constant expression before conversion to
        underlying type: unsigned semantics okay. */

};

int main () {
    // Same as return 0;
    return ui_wraparound_to_zero + us_wraparound_to_zero;
}

```

19 EXAMPLE 3 The following fragment:

```

#include <limits.h>

enum E1: short;
enum E2: short;
enum E3; /* Constraint violation: E3 forward declaration. */
enum E4 : unsigned long long;

enum E1 : short { m11, m12 };
enum E1 x = m11;

enum E2 : long { m21, m22 }; /* Constraint violation: different underlying types
*/

enum E3 {
    m31,
    m32,
    m33 = sizeof(enum E3) /* Constraint violation: E3 is not complete here. */
};
enum E3 : int; /* Constraint violation: E3 previously had no underlying type */

enum E4 : unsigned long long {
    m40 = sizeof(enum E4),
    m41 = ULLONG_MAX,
    m42 /* Constraint violation: unrepresentable value (wraparound) */
};

enum E5 y; /* Constraint violation: incomplete type */
enum E6 : long int z; /* Constraint violation: enum-type-specifier
        with identifier in declarator */
enum E7 : long int = 0; /* Syntax violation:
        enum-type-specifier with initializer */

```

demonstrates many of the properties of multiple declarations of enumerations with underlying types. Particularly, **enum E3** is declared and defined without an underlying type first, therefore a redeclaration with an underlying type second is a violation. Because it not complete at that time within its enumerator list, **sizeof(enum E3)** is a constraint violation within the **enum E3** definition. **enum E4** is complete as it is being defined, therefore **sizeof(enum E4)** is not a constraint violation.

20 EXAMPLE 4 The following fragment:

```
enum no_underlying {
    a0
};

int main (void) {
    int a = _Generic(a0,
        int: 2,
        unsigned char: 1,
        default: 0
    );
    int b = _Generic((enum no_underlying)a0,
        int: 2,
        unsigned char: 1,
        default: 0
    );
    return a + b;
}
```

demonstrates the implementation-defined nature of the underlying type of enumerations using generic selection (6.5.2.1). The value of **a** after its initialization is 2. The value of **b** after its initialization is implementation-defined: the enumeration is compatible with a type large enough to fit the values of its enumeration constants due to the constraints. Because the only value is 0 for **a0**, **b** can hold any of 2, 1, or 0.

The following fragment is similar, but uses a fixed underlying type:

```
enum underlying : unsigned char {
    b0
};

int main (void) {
    int a = _Generic(b0,
        int: 2,
        unsigned char: 1,
        default: 0
    );
    int b = _Generic((enum underlying)b0,
        int: 2,
        unsigned char: 1,
        default: 0
    );
    return 0;
}
```

Here, we are guaranteed that **a** and **b** are both initialized to 1. This makes enumerations with a fixed underlying type more portable.

21 EXAMPLE 5 Enumerations with a fixed underlying type have braces and the enumerator list specified as part of their declaration if they are not a standalone declaration:

```
void f1 (enum a : long b); /* Constraint violation */
void f2 (enum c : long { x } d);
enum e : int f3(); /* Constraint violation */

typedef enum t u; /* Constraint violation: forward declaration of t. */
typedef enum v : short W; /* Constraint violation */
```



```

typedef enum q : short { s } R;

struct s1 {
    int x;
    enum e : int : 1; /* Constraint violation */
    int y;
};

enum forward; /* Constraint violation */
extern enum forward fwd_val0; /* Constraint violation: incomplete type */
extern enum forward* fwd_ptr0; /* Constraint violation: enums cannot be
                                used like other incomplete types */
extern int* fwd_ptr0; /* Constraint violation: incompatible
                      with incomplete type. */

enum forward1 : int;
extern enum forward1 fwd_val1;
extern int fwd_val1;
extern enum forward1* fwd_ptr1;
extern int* fwd_ptr1;

int main () {
    enum e : short;
    enum e : short f = 0; /* Constraint violation */
    enum g : short { y } h = y;
    return 0;
}

```

- 22 EXAMPLE 6 Enumerations with a fixed underlying type are complete when the enum type specifier for that specific enumeration is complete. The enumeration `e` in this snippet:

```
enum e : typeof ((enum e : short { A })0, (short)0);
```

`enum e` is considered complete by the first opening brace within the `typeof` in this snippet.

Forward references: generic selection (6.5.2.1), tags (6.7.3.4), declarations (6.7), declarators (6.7.7), function declarators (6.7.7.4), type names (6.7.8).

6.7.3.4 Tags

Constraints

- 1 Where two declarations that use the same tag declare the same type, they shall both use the same choice of **struct**, **union**, or **enum**. If two declarations of the same type have a member-declaration or enumerator-list, one shall not be nested within the other and both declarations shall fulfill all requirements of compatible types (6.2.7) with the additional requirement that corresponding members of structure or union types shall have the same (and not merely compatible) types.
- 2 A type specifier of the form

enum *identifier*

without an enumerator list shall only appear after the type it specifies is complete.

- 3 A type specifier of the form

struct-or-union *attribute-specifier-sequence*_{opt} *identifier*

shall not contain an attribute specifier sequence.¹⁴⁵⁾

Semantics

- 4 All declarations of structure, union, or enumerated types that have the same scope and use the same tag declare the same type.

¹⁴⁵⁾As specified in 6.7.3.2, the type specifier can be followed by a ; or a member declaration list.

- 5 Irrespective of whether there is a tag or what other declarations of the type are in the same translation unit, the type (except enumerated types with a fixed underlying type) is incomplete until immediately after the closing brace of the list defining the content for the first time and complete thereafter.
- 6 Enumerated types with fixed underlying type (6.7.3.3) are complete immediately after their first associated enum type specifier ends.
- 7 EXAMPLE 1 The following example shows allowed redeclarations of the same structure, union, or enumerated type in the same scope:

```

struct foo { struct { int x; }; };
struct foo { struct { int x; }; };
union bar { int x; float y; };
union bar { int x; float y; };
typedef struct q { int x; } q_t;
typedef struct q { int x; } q_t;
void foo(void)
{
    struct S { int x; };
    struct T { struct S s; };
    struct S { int x; };
    struct T { struct S s; };
}
enum X { A = 1, B = 1 + 1 };
enum X { B = 2, A = 1 };

enum Q { C = 1 };
enum Q { C = C };           // ok!

```

- 8 EXAMPLE 2 The following example shows invalid redeclarations of the same structure, union, or enumerated type in the same scope:

```

struct foo { int (*p)[3]; };
struct foo { int (*p)[]; };    // member has different type

union bar { int x; float y; };
union bar { int z; float y; }; // member has different name

union purr { int x; float y; };
union purr { float y; int x; }; // members have different order
// "purr" only valid if each union "purr" is in
// two different translation units

typedef struct { int x; } q_t;
typedef struct { int x; } q_t; // not the same type

struct S { int x; };

void foo(void)
{
    struct T { struct S s; };
    struct S { int x; };
    struct T { struct S s; }; // struct S not the same type
}

enum X { A = 1, B = 2 };
enum X { A = 1, B = 3 };    // different enumeration constant

enum R { C = 1 };
enum Q { C = 1 };          // conflicting enumeration constant

```

- 9 Two declarations of structure, union, or enumerated types which are in different scopes or use different tags declare distinct types. Each declaration of a structure, union, or enumerated type which does not include a tag declares a distinct type.

- 10 A type specifier of the form

struct-or-union *attribute-specifier-sequence*_{opt} *identifier*_{opt} { *member-declaration-list* }

or

enum *attribute-specifier-sequence*_{opt} *identifier*_{opt} *enum-type-specifier*_{opt} { *enumerator-list* }

or

enum *attribute-specifier-sequence*_{opt} *identifier*_{opt} *enum-type-specifier*_{opt} { *enumerator-list* , }

declares a structure, union, or enumerated type. The list defines the *structure content*, *union content*, or *enumeration content*. If an identifier is provided,¹⁴⁶⁾ the type specifier also declares the identifier to be the tag of that type. The optional attribute specifier sequence appertains to the structure, union, or enumerated type being declared; the attributes in that attribute specifier sequence are thereafter considered attributes of the structure, union, or enumerated type whenever it is named.

- 11 A declaration of the form

struct-or-union *attribute-specifier-sequence*_{opt} *identifier* ;

or

enum *identifier* *enum-type-specifier* ;

specifies a structure, union, or enumerated type and declares the identifier as a tag of that type.¹⁴⁷⁾ The optional attribute specifier sequence appertains to the structure or union type being declared; the attributes in that attribute specifier sequence are thereafter considered attributes of the structure or union type whenever it is named.

- 12 If a type specifier of the form

struct-or-union *attribute-specifier-sequence*_{opt} *identifier*

occurs other than as part of one of the preceding forms, and no other declaration of the identifier as a tag is visible, then it declares an incomplete structure or union type, and declares the identifier as the tag of that type.¹⁴⁷⁾

- 13 If a type specifier of the form

struct-or-union *attribute-specifier-sequence*_{opt} *identifier*

or

enum *identifier*

occurs other than as part of one of the preceding forms, and a declaration of the identifier as a tag is visible, then it specifies the same type as that other declaration, and does not redeclare the tag.

- 14 EXAMPLE 3 This mechanism allows declaration of a self-referential structure.

```
struct tnode {
    int count;
    struct tnode *left, *right;
};
```

specifies a structure that contains an integer and two pointers to objects of the same type. Once this declaration has been given, the declaration

```
struct tnode s, *sp;
```

¹⁴⁶⁾If there is no identifier, the type can, within the translation unit, only be referred to by the declaration of which it is a part. Of course, when the declaration is of a typedef name, subsequent declarations can make use of that typedef name to declare objects having the specified structure, union, or enumerated type.

¹⁴⁷⁾A similar construction for an **enum** that does not contain a fixed underlying type does not exist. Enumerations with a fixed underlying type are always complete after the enum type specifier.

declares **s** to be an object of the given type and **sp** to be a pointer to an object of the given type. With these declarations, the expression **sp->left** refers to the left **struct tnode** pointer of the object to which **sp** points; the expression **s.right->count** designates the **count** member of the right **struct tnode** pointed to from **s**.

- 15 The following alternative formulation uses the **typedef** mechanism:

```
typedef struct tnode TNODE;
struct tnode {
    int count;
    TNODE *left, *right;
};
TNODE s, *sp;
```

- 16 EXAMPLE 4 To illustrate the use of prior declaration of a tag to specify a pair of mutually referential structures, the declarations

```
struct s1 { struct s2 *s2p; /* ... */ }; // D1
struct s2 { struct s1 *s1p; /* ... */ }; // D2
```

specify a pair of structures that contain pointers to each other. If **s2** were already declared as a tag in an enclosing scope, the declaration **D1** would refer to *it*, not to the tag **s2** declared in **D2**. To eliminate this context sensitivity, the declaration

```
struct s2;
```

can be inserted ahead of **D1**. This declares a new tag **s2** in the inner scope; the declaration **D2** then completes the specification of the new type.

Forward references: declarators (6.7.7), type definitions (6.7.9).

6.7.3.5 Atomic type specifiers

Syntax

- 1 *atomic-type-specifier:*
 _Atomic (*type-name*)

Constraints

- 2 Atomic type specifiers shall not be used if the implementation does not support atomic types (see 6.10.10.4).
- 3 The type name in an atomic type specifier shall not refer to an array type, a function type, an atomic type, or a qualified type.

Semantics

- 4 The properties associated with atomic types are meaningful only for expressions that are lvalues. If the **_Atomic** keyword is immediately followed by a left parenthesis, it is interpreted as a type specifier (with a type name), not as a type qualifier.

6.7.3.6 Typeof specifiers

Syntax

- 1 *typeof-specifier:*
 typeof (*typeof-specifier-argument*)
 typeof_unqual (*typeof-specifier-argument*)
typeof-specifier-argument:
 expression
 type-name

- 2 The **typeof** and **typeof_unqual** tokens are collectively called the *typeof operators*.

Constraints

- 3 The `typeof` operators shall not be applied to an expression that designates a bit-field member.

Semantics

- 4 The `typeof` specifier applies the `typeof` operators to an *expression* (6.5.1) or a type name. If the `typeof` operators are applied to an expression, they yield the type of their operand.¹⁴⁸⁾ Otherwise, they designate the same type as the type name with any nested `typeof` specifier evaluated.¹⁴⁹⁾ If the type of the operand is a variably modified type, the operand is evaluated; otherwise, the operand is not evaluated.
- 5 The result of the **`typeof_unqual`** operator is the non-atomic unqualified version of the type that would result from the **`typeof`** operator.¹⁵⁰⁾ The **`typeof`** operator preserves all qualifiers.
- 6 EXAMPLE 1 Type of an expression.

```
typeof(1+1) main () {
    return 0;
}
```

is equivalent to this program:

```
int main () {
    return 0;
}
```

- 7 EXAMPLE 2 The following program:

```
const _Atomic int purr = 0;
const int meow = 1;
const char* const animals[] = {
    "aardvark",
    "bluejay",
    "catte",
};

typeof_unqual(meow) main (int argc, char* argv[]) {
    typeof_unqual(purr)      plain_purr;
    typeof(_Atomic typeof(meow)) atomic_meow;
    typeof(animals)          animals_array;
    typeof_unqual(animals)   animals2_array;
    return 0;
}
```

is equivalent to this program:

```
const _Atomic int purr = 0;
const int meow = 1;
const char* const animals[] = {
    "aardvark",
    "bluejay",
    "catte",
};

int main (int argc, char* argv[]) {
    int      plain_purr;
    const _Atomic int atomic_meow;
```

¹⁴⁸⁾When applied to a parameter declared to have array or function type, the `typeof` operators yield the adjusted (pointer) type (see 6.9.2).

¹⁴⁹⁾If the `typeof` specifier argument is itself a `typeof` specifier, the operand will be evaluated before evaluating the current `typeof` operator. This happens recursively until a `typeof` specifier is no longer the operand.

¹⁵⁰⁾**`_Atomic`** (*type-name*), with parentheses, is considered an **`_Atomic`**-qualified type.

```

const char* const animals_array[3];
const char*     animals2_array[3];
return 0;
}

```

- 8 EXAMPLE 3 The equivalence between **sizeof** and **typeof**'s deduction of the type means this program has no constraint violations:

```

int main (int argc, char* argv[]) {
    static_assert(sizeof(typeof('p')) == sizeof(int));
    static_assert(sizeof(typeof('p')) == sizeof('p'));
    static_assert(sizeof(typeof((char)'p')) == sizeof(char));
    static_assert(sizeof(typeof((char)'p')) == sizeof((char)'p'));
    static_assert(sizeof(typeof("meow")) == sizeof(char[5]));
    static_assert(sizeof(typeof("meow")) == sizeof("meow"));
    static_assert(sizeof(typeof(argc)) == sizeof(int));
    static_assert(sizeof(typeof(argc)) == sizeof(argc));
    static_assert(sizeof(typeof(argv)) == sizeof(char**));
    static_assert(sizeof(typeof(argv)) == sizeof(argv));

    static_assert(sizeof(typeof_unqual('p')) == sizeof(int));
    static_assert(sizeof(typeof_unqual('p')) == sizeof('p'));
    static_assert(sizeof(typeof_unqual((char)'p')) == sizeof(char));
    static_assert(sizeof(typeof_unqual((char)'p')) == sizeof((char)'p'));
    static_assert(sizeof(typeof_unqual("meow")) == sizeof(char[5]));
    static_assert(sizeof(typeof_unqual("meow")) == sizeof("meow"));
    static_assert(sizeof(typeof_unqual(argc)) == sizeof(int));
    static_assert(sizeof(typeof_unqual(argc)) == sizeof(argc));
    static_assert(sizeof(typeof_unqual(argv)) == sizeof(char**));
    static_assert(sizeof(typeof_unqual(argv)) == sizeof(argv));
    return 0;
}

```

- 9 EXAMPLE 4 The following program with nested **typeof(...)**:

```

int main (int argc, char*[]) {
    float val = 6.0f;
    return (typeof(typeof_unqual(typeof(argc))))val;
}

```

is equivalent to this program:

```

int main (int argc, char*[]) {
    float val = 6.0f;
    return (int)val;
}

```

- 10 EXAMPLE 5 Variable length arrays with **typeof** operators perform the operation at execution time rather than translation time.

```

#include <stddef.h>

size_t vla_size (int n) {
    typedef char vla_type[n + 3];
    vla_type b; // variable length array
    return sizeof(
        typeof_unqual(b)
    ); // execution-time sizeof, translation-time typeof operation
}

int main () {

```

```

    return (int)vla_size(10); // vla_size returns 13
}

```

- 11 EXAMPLE 6 Nested typeof operators, arrays, and pointers do not perform array to pointer decay.

```

int main () {
    typeof(typeof(const char*)[4]) y = {
        "a",
        "b",
        "c",
        "d"
    }; // 4-element array of "pointer to const char"
    return 0;
}

```

- 12 EXAMPLE 7 Function, pointer, and array types can be substituted with typeof operations.

```

void f(int);

typeof(f(5)) g(double x) {           // g has type "void(double)"
    printf("value %g\n", x);
}

typeof(g)* h;                        // h has type "void(*) (double)"
typeof(true ? g : nullptr) k;        // k has type "void(*) (double)"

void j(double A[5], typeof(A)* B);    // j has type "void(double*, double**)"

extern typeof(double[]) D;            // D has an incomplete type
typeof(D) C = { 0.7, 99 };           // C has type "double[2]"

typeof(D) D = { 5, 8.9, 0.1, 99 };    // D is now completed to "double[4]"
typeof(D) E;                         // E has type "double[4]" from D's completed type

```

6.7.4 Type qualifiers

6.7.4.1 General

Syntax

- 1 *type-qualifier*:

```

const
restrict
volatile
_Atomic

```

Constraints

- 2 Types other than pointer types whose referenced type is an object type and (possibly multi-dimensional) array types with such pointer types as element type shall not be restrict-qualified.
- 3 The **_Atomic** qualifier shall not be used if the implementation does not support atomic types (see 6.10.10.4).
- 4 The type modified by the **_Atomic** qualifier shall not be an array type or a function type.

Semantics

- 5 The properties associated with qualified types are meaningful only for expressions that are lvalues.¹⁵¹⁾

¹⁵¹⁾The implementation can place a **const** object that is not **volatile** in a read-only region of storage. Moreover, the implementation is not expected to allocate storage for such an object if its address is never used.

- 6 If the same qualifier appears more than once in the same specifier-qualifier list or as declaration specifiers, either directly, via one or more `typeof` specifiers, or via one or more **typedefs**, the behavior is the same as if it appeared only once. If other qualifiers appear along with the **_Atomic** qualifier the resulting type is the so-qualified atomic type.
- 7 If an attempt is made to modify an object defined with a const-qualified type through use of an lvalue with non-const-qualified type, the behavior is undefined. If an attempt is made to refer to an object defined with a volatile-qualified type through use of an lvalue with non-volatile-qualified type, the behavior is undefined.¹⁵²⁾
- 8 An object that has volatile-qualified type can be modified in ways unknown to the implementation or have other unknown side effects. Therefore, any expression referring to such an object shall be evaluated strictly according to the rules of the abstract machine, as described in 5.2.2.4. Furthermore, at every sequence point the value last stored in the object shall agree with that prescribed by the abstract machine, except as modified by the unknown factors mentioned previously.¹⁵³⁾ What constitutes an access to an object that has volatile-qualified type is implementation-defined.
- 9 An object that is accessed through a restrict-qualified pointer has a special association with that pointer. This association, defined in 6.7.4.2, requires that all accesses to that object use, directly or indirectly, the value of that pointer.¹⁵⁴⁾ The intended use of the **restrict** qualifier (like the **register** storage class) is to promote optimization, and deleting all instances of the qualifier from all preprocessing translation units composing a conforming program does not change its meaning (i.e. observable behavior), unless **_Generic** is used to distinguish whether or not a type has that qualifier.
- 10 If the specification of an array type includes any type qualifiers, both the array and the element type are so-qualified. If the specification of a function type includes any type qualifiers, the behavior is implementation-defined.¹⁵⁵⁾
- 11 For two qualified types to be compatible, both shall have the identically qualified version of a compatible type; the order of type qualifiers within a list of specifiers or qualifiers does not affect the specified type.
- 12 EXAMPLE 1 An object declared

```
extern const volatile int real_time_clock;
```

can be modifiable by hardware, but cannot be assigned to, incremented, or decremented.

- 13 EXAMPLE 2 The following declarations and expressions illustrate the behavior when type qualifiers modify an aggregate type:

```
const struct s { int mem; } cs = { 1 };
struct s ncs; // the object ncs is modifiable
typedef int A[2][3];
const A a = {{4, 5, 6}, {7, 8, 9}}; // array of array of const int
int *pi;
const int *pci;

ncs = cs; // valid
cs = ncs; // violates modifiable lvalue constraint for =
pi = &ncs.mem; // valid
pi = &cs.mem; // violates type constraints for =
pci = &cs.mem; // valid
```

¹⁵²⁾This applies to those objects that behave as if they were defined with qualified types, even if they are never actually defined as objects in the program (such as an object at a memory-mapped input/output address).

¹⁵³⁾A **volatile** declaration can be used to describe an object corresponding to a memory-mapped input/output port or an object accessed by an asynchronously interrupting function. Actions on objects so declared are not allowed to be “optimized out” by an implementation or reordered except as permitted by the rules for evaluating expressions.

¹⁵⁴⁾For example, a statement that assigns a value returned by **malloc** to a single pointer establishes this association between the allocated object and the pointer.

¹⁵⁵⁾This can occur with **typedef** s. This rule does not apply to the **_Atomic** qualifier. Qualifiers do not have any direct effect on the array type itself, but affect conversion rules for pointer types that reference an array type.


```
pi = a[0];    // invalid: a[0] has type "const int *"
```

14 EXAMPLE 3 The declaration

```
_Atomic volatile int *p;
```

specifies that `p` has the type “pointer to volatile atomic int”, a pointer to a volatile-qualified atomic type.

6.7.4.2 Formal definition of **restrict**

- 1 Let **D** be a declaration of an ordinary identifier that provides a means of designating an object **P** as a restrict-qualified pointer to type **T**.
- 2 If **D** appears inside a block and does not have storage class **extern**, let **B** denote the block. If **D** appears in the list of parameter declarations of a function definition, let **B** denote the associated block. Otherwise, let **B** denote the block of **main** (or the block of whatever function is called at program startup in a freestanding environment).
- 3 In what follows, a pointer expression **E** is said to be *based* on object **P** if (at some sequence point in the execution of **B** prior to the evaluation of **E**) modifying **P** to point to a copy of the array object into which it formerly pointed would change the value of **E**.¹⁵⁶ “based” is defined only for expressions with pointer types.
- 4 During each execution of **B**, let **L** be any lvalue that has **&L** based on **P**. If **L** is used to access the value of the object **X** that it designates, and **X** is also modified (by any means), then the following requirements apply: **T** shall not be const-qualified. Every other lvalue used to access the value of **X** shall also have its address based on **P**. Every access that modifies **X** shall be considered also to modify **P**, for the purposes of this subclause. If **P** is assigned the value of a pointer expression **E** that is based on another restricted pointer object **P2**, associated with block **B2**, then either the execution of **B2** shall begin before the execution of **B**, or the execution of **B2** shall end prior to the assignment. If these requirements are not met, then the behavior is undefined.
- 5 Here an execution of **B** means that portion of the execution of the program that would correspond to the lifetime of an object with scalar type and automatic storage duration associated with **B**.
- 6 A translator is free to ignore any or all aliasing implications of uses of **restrict**.
- 7 EXAMPLE 1 The file scope declarations

```
int * restrict a;
int * restrict b;
extern int c[];
```

assert that if an object is accessed using one of `a`, `b`, or `c`, and that object is modified anywhere in the program, then it is never accessed using either of the other two.

8 EXAMPLE 2 The function parameter declarations in the following example

```
void f(int n, int * restrict p, int * restrict q)
{
    while (n-- > 0)
        *p++ = *q++;
}
```

assert that, during each execution of the function, if an object is accessed through one of the pointer parameters, then it is not also accessed through the other. The translator can make this no-aliasing inference based on the parameter declarations alone, without analyzing the function body.

- 9 The benefit of the **restrict** qualifiers is that they enable a translator to make an effective dependence analysis of function **f** without examining any of the calls of **f** in the program. The cost is that the programmer has to

¹⁵⁶In other words, **E** depends on the value of **P** itself rather than on the value of an object referenced indirectly through **P**. For example, if identifier **p** has type **(int **restrict)**, then the pointer expressions **p** and **p+1** are based on the restricted pointer object designated by **p**, but the pointer expressions ***p** and **p[1]** are not.

examine all those calls to ensure that none give undefined behavior. For example, the second call of `f` in `g` has undefined behavior because each of `d[1]` through `d[49]` is accessed through both `p` and `q`.

```
void g(void)
{
    extern int d[100];
    f(50, d + 50, d); // valid
    f(50, d + 1, d); // undefined behavior
}
```

10 EXAMPLE 3 The function parameter declarations

```
void h(int n, int * restrict p, int * restrict q, int * restrict r)
{
    int i;
    for (i = 0; i < n; i++)
        p[i] = q[i] + r[i];
}
```

illustrate how an unmodified object can be aliased through two restricted pointers. If `a` and `b` are disjoint arrays, a call of the form `h(100, a, b, b)` has defined behavior, because array `b` is not modified within function `h`.

11 EXAMPLE 4 The rule limiting assignments between restricted pointers does not distinguish between a function call and an equivalent nested block. With one exception, only “outer-to-inner” assignments between restricted pointers declared in nested blocks have defined behavior.

```
{
    int * restrict p1;
    int * restrict q1;
    p1 = q1; // undefined behavior
    {
        int * restrict p2 = p1; // valid
        int * restrict q2 = q1; // valid
        p1 = q2;                // undefined behavior
        p2 = q2;                // undefined behavior
    }
}
```

12 The one exception allows the value of a restricted pointer to be carried out of the block in which it (or, more precisely, the ordinary identifier used to designate it) is declared when that block finishes execution. For example, this permits `new_vector` to return a `vector`.

```
typedef struct { int n; float * restrict v; } vector;
vector new_vector(int n)
{
    vector t;
    t.n = n;
    t.v = malloc(n * sizeof(float));
    return t;
}
```

13 EXAMPLE 5 Suppose that a programmer knows that references of the form `p[i]` and `q[j]` are never aliases in the body of a function:

```
void f(int n, int *p, int *q) { /* ... */ }
```

There are several ways that this information can be conveyed to a translator using the `restrict` qualifier. Example 2 shows the most effective way, qualifying all pointer parameters, and can be used provided that neither `p` nor `q` becomes based on the other in the function body. A potentially effective alternative is:

```
void f(int n, int * restrict p, int * const q) { /* ... */ }
```

Again, it is possible for a translator to make the no-aliasing inference based on the parameter declarations alone, though now subtler reasoning is used: that the `const`-qualification of `q` precludes it becoming based on `p`. There is also a requirement that `q` is not modified, so this alternative cannot be used for the function in Example 2, as written.

- 14 EXAMPLE 6 Another potentially effective alternative is:

```
void f(int n, int *p, int const * restrict q) { /* ... */ }
```

Again, it is possible for a translator to make the no-aliasing inference based on the parameter declarations alone, though now even subtler reasoning is used: that this combination of `restrict` and `const` means that objects referenced using `q` cannot be modified, and so no modified object can be referenced using both `p` and `q`.

- 15 EXAMPLE 7 The least effective alternative is:

```
void f(int n, int * restrict p, int *q) { /* ... */ }
```

Here the translator can make the no-aliasing inference only by analyzing the body of the function and proving that `q` cannot become based on `p`. Some translator designs can choose to exclude this analysis, given availability of the more effective alternatives described previously. Such a translator is required to assume that aliases are present because assuming that aliases are not present can result in an incorrect translation. Also, a translator that attempts the analysis can potentially not succeed in all cases and consequently need to conservatively assume that aliases are present.

6.7.5 Function specifiers

Syntax

- 1 *function-specifier*:

```
inline
_Noreturn
```

Constraints

- 2 Function specifiers shall be used only in the declaration of an identifier for a function.
- 3 An inline definition of a function with external linkage shall not contain, anywhere in the tokens making up the function definition, a definition of a modifiable object with static or thread storage duration, and shall not contain, anywhere in the tokens making up the function definition, a reference to an identifier with internal linkage.
- 4 In a hosted environment, no function specifier(s) shall appear in a declaration of `main`.
- 5 If a function declaration declaring an identifier with external linkage has an `inline` function specifier, then the function shall also be defined in the same translation unit.

Semantics

- 6 A function specifier can appear more than once; the behavior is the same as if it appeared only once.
- 7 A function declared with an `inline` function specifier is an *inline function*. Making a function an inline function suggests that calls to the function be as fast as possible.¹⁵⁷⁾ The extent to which such suggestions are effective is implementation-defined.¹⁵⁸⁾
- 8 Any function with internal linkage can be an inline function. For a function with external linkage, the following restrictions apply: If all the file scope declarations for a function in a translation unit include the `inline` function specifier without `extern`, then the definition in that translation unit

¹⁵⁷⁾By using, for example, an alternative to the usual function call mechanism, such as “inline substitution”. Inline substitution is not textual substitution, nor does it create a new function. Therefore, for example, the expansion of a macro used within the body of the function uses the definition it had at the point the function body appears, and not where the function is called; and identifiers refer to the declarations in scope where the body occurs. Likewise, the function has a single address, regardless of the number of inline definitions that occur in addition to the external definition.

¹⁵⁸⁾For example, an implementation can possibly never perform inline substitution, or can only perform inline substitutions to calls in the scope of an `inline` declaration.

is an *inline definition*. An inline definition does not provide an external definition for the function and does not forbid an external definition in another translation unit. Inline definitions provide an alternative to external definitions, which a translator can use to implement any call to the function in the same translation unit. It is unspecified whether a call to the function uses the inline definition or the external definition.¹⁵⁹⁾

- 9 A function declared with a **_Noreturn** function specifier shall not return to its caller. The attribute **[[noreturn]]** provides similar semantics. The **_Noreturn** function specifier is an obsolescent feature (6.7.13.7).

Recommended practice

- 10 The implementation should produce a diagnostic message for a function declared with a **_Noreturn** function specifier that appears to be capable of returning to its caller.
- 11 EXAMPLE 1 The declaration of an inline function with external linkage can result in either an external definition, or a definition available for use only within the translation unit. A file scope declaration with **extern** creates an external definition. The following example shows an entire translation unit.

```
inline double fahr(double t)
{
    return (9.0 * t) / 5.0 + 32.0;
}

inline double cels(double t)
{
    return (5.0 * (t - 32.0)) / 9.0;
}

extern double fahr(double); // creates an external definition

double convert(int is_fahr, double temp)
{
    /* A translator can perform inline substitutions */
    return is_fahr ? cels(temp) : fahr(temp);
}
```

- 12 The definition of **fahr** is an external definition because **fahr** is also declared with **extern**, but the definition of **cels** is an inline definition. Because **cels** has external linkage and is referenced, an external definition has to appear in another translation unit (see 6.9); the inline definition and the external definition are distinct and either can be used for the call.
- 13 EXAMPLE 2 The following inline definitions are invalid:

```
static int a;
typedef (a) inline f() { return 0; }
typedef ((int) { 0 }) inline g() { return 0; }
```

Forward references: function definitions (6.9.2).

6.7.6 Alignment specifier

Syntax

- 1 *alignment-specifier*:
- ```
alignas (type-name)
alignas (constant-expression)
```

<sup>159)</sup>Since an inline definition is distinct from the corresponding external definition and from any other corresponding inline definitions in other translation units, all corresponding objects with static storage duration are also distinct in each of the definitions.

**Constraints**

- 2 An alignment specifier shall appear only in the declaration specifiers of a declaration, or in the *specifier-qualifier* list of a member declaration, or in the type name of a compound literal. An alignment specifier shall not be used in conjunction with either of the storage-class specifiers **typedef** or **register**, nor in a declaration of a function or bit-field.
- 3 The constant expression shall be an integer constant expression. It shall evaluate to a valid fundamental alignment, or to a valid extended alignment supported by the implementation for an object of the storage duration (if any) being declared, or to zero.
- 4 An object shall not be declared with an over-aligned type with an extended alignment requirement not supported by the implementation for an object of that storage duration.
- 5 The combined effect of all alignment specifiers in a declaration shall not specify an alignment that is less strict than the alignment that would otherwise be required for the type of the object or member being declared.
- 6 If the definition of an object has an alignment specifier, any other declaration of that object in the same translation unit shall either specify equivalent alignment or have no alignment specifier. If the definition of an object does not have an alignment specifier, any other declaration of that object in the same translation unit shall also have no alignment specifier.

**Semantics**

- 7 The first form is equivalent to **alignas**(**alignof**( *type-name* ) ).
- 8 The alignment requirement of the declared object or member is taken to be the specified alignment. An alignment specification of zero has no effect.<sup>160)</sup> When multiple alignment specifiers occur in a declaration, the effective alignment requirement is the strictest specified alignment.
- 9 If the definition of an object has an alignment specifier, no declaration of the object in a different translation unit shall specify a stricter alignment.

**6.7.7 Declarators****6.7.7.1 General****Syntax**

- 1 *declarator*:
 

*pointer*<sub>opt</sub> *direct-declarator*

*direct-declarator*:
 

*identifier* *attribute-specifier-sequence*<sub>opt</sub>  
 ( *declarator* )  
*array-declarator* *attribute-specifier-sequence*<sub>opt</sub>  
*function-declarator* *attribute-specifier-sequence*<sub>opt</sub>

*array-declarator*:
 

*direct-declarator* [ *type-qualifier-list*<sub>opt</sub> *assignment-expression*<sub>opt</sub> ]  
*direct-declarator* [ **static** *type-qualifier-list*<sub>opt</sub> *assignment-expression* ]  
*direct-declarator* [ *type-qualifier-list* **static** *assignment-expression* ]  
*direct-declarator* [ *type-qualifier-list*<sub>opt</sub> \* ]

*function-declarator*:
 

*direct-declarator* ( *parameter-type-list*<sub>opt</sub> )

*pointer*:
 

\* *attribute-specifier-sequence*<sub>opt</sub> *type-qualifier-list*<sub>opt</sub>  
\* *attribute-specifier-sequence*<sub>opt</sub> *type-qualifier-list*<sub>opt</sub> *pointer*

*type-qualifier-list*:
 

*type-qualifier*

<sup>160)</sup>An alignment specification of zero also does not affect other alignment specifications in the same declaration.

$\text{type-qualifier-list } \text{type-qualifier}$   
 $\text{parameter-type-list:}$   
 $\text{parameter-list}$   
 $\text{parameter-list } , \dots$   
 $\dots$   
 $\text{parameter-list:}$   
 $\text{parameter-declaration}$   
 $\text{parameter-list } , \text{parameter-declaration}$   
 $\text{parameter-declaration:}$   
 $\text{attribute-specifier-sequence}_{\text{opt}} \text{ declaration-specifiers declarator}$   
 $\text{attribute-specifier-sequence}_{\text{opt}} \text{ declaration-specifiers abstract-declarator}_{\text{opt}}$

### Semantics

- 2 Each declarator declares an identifier for a single object, function, or type, within a declaration. The preceding specifiers indicate the type, storage class, or other properties of the identifier or identifiers being declared. Each declarator specifies one declaration and names it and/or modifies the type of the specifiers with operators such as \* (pointer to) and ( ) (function returning).
- 3 A *full declarator* is a declarator that is not part of another declarator. If, in the nested sequence of declarators in a full declarator, there is a declarator specifying a variable length array type, the type specified by the full declarator is said to be *variably modified*. Furthermore, any type derived by declarator type derivation from a variably modified type is itself variably modified.
- 4 In the following subclauses, consider a declaration

#### T D1

where **T** contains the declaration specifiers that specify a type *T* (such as **int**) and **D1** is a declarator that contains an identifier *ident*. The type specified for the identifier *ident* in the various forms of declarator is described inductively using this notation.

- 5 If, in the declaration “**T D1**”, **D1** has the form

$\text{identifier attribute-specifier-sequence}_{\text{opt}}$

then the type specified for *ident* is *T* and the optional attribute specifier sequence appertains to the entity that is declared.

- 6 If, in the declaration “**T D1**”, **D1** has the form

**( D )**

then *ident* has the type specified by the declaration “**T D**”. Thus, a declarator in parentheses is identical to the unparenthesized declarator, but the binding of complicated declarators can be altered by parentheses.

### Implementation limits

- 7 As discussed in 5.3.5.2, an implementation may limit the number of pointer, array, and function declarators that modify an arithmetic, structure, union, or **void** type, either directly or via one or more **typedefs**.

**Forward references:** array declarators (6.7.7.3), type definitions (6.7.9).

#### 6.7.7.2 Pointer declarators

##### Semantics

- 1 If, in the declaration “**T D1**”, **D1** has the form

$* \text{attribute-specifier-sequence}_{\text{opt}} \text{type-qualifier-list}_{\text{opt}} \text{D}$

and the type specified for *ident* in the declaration “**T D**” is “*derived-declarator-type-list T*”, then the type specified for *ident* is “*derived-declarator-type-list type-qualifier-list pointer to T*”. For each type qualifier in the list, *ident* is a so-qualified pointer. The optional attribute specifier sequence appertains to the pointer and not the object pointed to.

- 2 For two pointer types to be compatible, both shall be identically qualified and both shall be pointers to compatible types.
- 3 EXAMPLE The following pair of declarations demonstrates the difference between a “variable pointer to a constant value” and a “constant pointer to a variable value”.

```
const int *ptr_to_constant;
int *const constant_ptr;
```

The contents of any object pointed to by `ptr_to_constant` cannot be modified through that pointer, but `ptr_to_constant` itself can be changed to point to another object. Similarly, the contents of the `int` pointed to by `constant_ptr` can be modified, but `constant_ptr` itself always points to the same location.

- 4 The declaration of the constant pointer `constant_ptr` can be clarified by including a definition for the type “pointer to `int`”.

```
typedef int *int_ptr;
const int_ptr constant_ptr;
```

declares `constant_ptr` as an object that has type “const-qualified pointer to `int`”.

### 6.7.7.3 Array declarators

#### Constraints

- 1 In addition to optional type qualifiers and the keyword **static**, the [ and ] can delimit an expression or \*. If they delimit an expression (which specifies the size of an array), the expression shall have an integer type. If the expression is a constant expression, it shall have a value greater than zero. The element type shall not be an incomplete or function type. The optional type qualifiers and the keyword **static** shall appear only in a declaration of a function parameter with an array type, and then only in the outermost array type derivation.
- 2 If an identifier is declared as having a variably modified type, it shall be an ordinary identifier (as defined in 6.2.3), have no linkage, and have either block scope or function prototype scope. If an identifier is declared to be an object with static or thread storage duration, it shall not have a variable length array type.

#### Semantics

- 3 If, in the declaration “**T D1**”, **D1** has one of the forms:

```
D [type-qualifier-listopt assignment-expressionopt] attribute-specifier-sequenceopt
D [static type-qualifier-listopt assignment-expression] attribute-specifier-sequenceopt
D [type-qualifier-list static assignment-expression] attribute-specifier-sequenceopt
D [type-qualifier-listopt *] attribute-specifier-sequenceopt
```

and the type specified for *ident* in the declaration “**T D**” is “*derived-declarator-type-list T*”, then the type specified for *ident* is “*derived-declarator-type-list array of T*”.<sup>161)162)</sup> The optional attribute specifier sequence appertains to the array. (See 6.7.7.4 for the meaning of the optional type qualifiers and the keyword **static**.)

- 4 If the size is not present, the array type is an incomplete type. If the size is \* instead of being an expression, the array type is a *variable length array* type of unspecified size, which can only be used as part of the nested sequence of declarators or abstract declarators for a parameter declaration, not including anything inside an array size expression in one of those declarators;<sup>163)</sup> such arrays are nonetheless complete types. If the size is an integer constant expression and the element type has a known constant size, the array type is not a variable length array type; otherwise, the array type is a *variable length array* type. (Variable length arrays with automatic storage duration are a conditional feature that implementations may support; see 6.10.10.4.)
- 5 If the size is an expression that is not an integer constant expression: if it occurs in a declaration at function prototype scope, it is treated as if it were replaced by \*; otherwise, each time it is evaluated

<sup>161)</sup>When several “array of” specifications are adjacent, a multidimensional array is declared.

<sup>162)</sup>The array is considered identically qualified to **T** according to 6.2.5.

<sup>163)</sup>They can be used only in function declarations that are not definitions (see 6.7.7.4 and 6.9.2).

it shall have a value greater than zero. The size of each instance of a variable length array type does not change during its lifetime. Where a size expression is part of the operand of the `typeof` or `sizeof` operators and changing the value of the size expression would not affect the result of the operator, it is unspecified whether or not the size expression is evaluated. Where a size expression is part of the operand with a `_Lengthof` operator and changing the value of the size expression would not affect the result of the operator, the size expression is not evaluated. Where a size expression is part of the operand of an `alignof` operator, that expression is not evaluated.

- 6 For two array types to be compatible, both shall have compatible element types, and if both size specifiers are present, and are integer constant expressions, then both size specifiers shall have the same constant value. If the two array types are used in a context which requires them to be compatible, the behavior is undefined if the two size specifiers evaluate to unequal values.

7 EXAMPLE 1

```
float fa[11], *afp[17];
```

declares an array of **float** numbers and an array of pointers to **float** numbers.

8 EXAMPLE 2 There is a distinction between the declarations

```
extern int *x;
extern int y[];
```

The first declares `x` to be a pointer to **int**; the second declares `y` to be an array of **int** of unknown size (an incomplete type), the storage for which is defined elsewhere.

9 EXAMPLE 3 The following declarations demonstrate the compatibility rules for variably modified types.

```
extern int n;
extern int m;

void fcompat(void)
{
 int a[n][6][m];
 int (*p)[4][n+1];
 int c[n][n][6][m];
 int (*r)[n][n][n+1];
 p = a; // invalid: not compatible because 4 != 6
 r = c; // compatible, but defined behavior only if
 // n == 6 and m == n+1
}
```

- 10 EXAMPLE 4 All valid declarations of variably modified (VM) types are either at block scope or function prototype scope. Array objects declared with the **thread\_local**, **static**, or **extern** storage-class specifier cannot have a variable length array (VLA) type. However, an object declared with the **static** storage-class specifier can have a VM type (that is, a pointer to a VLA type). Finally, only ordinary identifiers can be declared with a VM type and identifiers with VM type cannot, therefore, be members of structures or unions.

```
extern int n;
int A[n]; // invalid: file scope VLA
extern int (*p2)[n]; // invalid: file scope VM
int B[100]; // valid: file scope but not VM

void fvla(int m, int C[m][m]); // valid: VLA with prototype scope

void fvla(int m, int C[m][m]) // valid: adjusted to auto pointer to VLA
{
 typedef int VLA[m][m]; // valid: block scope typedef VLA

 struct tag {
 int (*y)[n]; // invalid: y not ordinary identifier
 int z[n]; // invalid: z not ordinary identifier
 };
}
```



```

};
int D[m]; // valid: auto VLA
static int E[m]; // invalid: static block scope VLA
extern int F[m]; // invalid: F has linkage and is VLA
int (*s)[m]; // valid: auto pointer to VLA
extern int (*r)[m]; // invalid: r has linkage and points to VLA
static int (*q)[m] = &B; // valid: q is a static block pointer to VLA
}

```

- 11 EXAMPLE 5 The following is invalid, because the use of `[*]` is inside an array size expression rather than directly part of the nested sequence of abstract declarators for a parameter declaration:

```
void f(int (*)[sizeof(int (*)[*])]);
```

**Forward references:** function declarators (6.7.7.4), function definitions (6.9.2), initialization (6.7.11).

#### 6.7.7.4 Function declarators

##### Constraints

- 1 A function declarator shall not specify a return type that is a function type or an array type.
- 2 The only storage-class specifier that shall occur in a parameter declaration is **register**.
- 3 After adjustment, the parameters in a parameter type list in a function declarator that is part of a definition of that function shall not have incomplete type.
- 4 A parameter declaration shall not specify a void type, except for the special case of a single unnamed parameter of type **void** with no storage class specifier, no type qualifier, and no following ellipsis terminator.

##### Semantics

- 5 If, in the declaration “**T D1**”, **D1** has the form

**D** ( *parameter-type-list*<sub>opt</sub> ) *attribute-specifier-sequence*<sub>opt</sub>

and the type specified for *ident* in the declaration “**T D**” is “*derived-declarator-type-list T*”, then the type specified for *ident* is “*derived-declarator-type-list* function returning the unqualified, non-atomic version of *T*”. The optional attribute specifier sequence appertains to the function type.

- 6 A parameter type list specifies the types of, and can declare identifiers for, the parameters of the function.
- 7 A declaration of a parameter as “array of *type*” shall be adjusted to “qualified pointer to *type*”, where the type qualifiers (if any) are those specified within the [ and ] of the array type derivation. If the keyword **static** also appears within the [ and ] of the array type derivation, then for each call to the function, the value of the corresponding actual argument shall provide access to the first element of an array with at least as many elements as specified by the size expression.
- 8 A declaration of a parameter as “function returning *type*” shall be adjusted to “pointer to function returning *type*”, as in 6.3.3.1.
- 9 If the list terminates with an ellipsis (...), no information about the number or types of the parameters after the comma is supplied.<sup>164</sup> Functions whose declarators use the ellipses in a parameter list are known as *variadic functions*. Arguments supplied to a function whose positions match or come after the ellipsis in the parameter list are its *varying arguments*.
- 10 The special case of an unnamed parameter of type **void** as the only item in the list specifies that the function has no parameters.
- 11 If, in a parameter declaration, an identifier can be treated either as a typedef name or as a parameter name, it shall be taken as a typedef name.
- 12 If the function declarator is not part of a definition of that function, parameters can have incomplete

<sup>164</sup>The macros defined in the `<stdarg.h>` header (7.16) can be used to access arguments that correspond to the ellipsis in variadic functions.

type and can use the `[*]` notation in their sequences of declarator specifiers to specify variable length array types.

- 13 The storage-class specifier in the declaration specifiers for a parameter declaration, if present, is ignored unless the declared parameter is one of the members of the parameter type list for a function definition. The optional attribute specifier sequence in a parameter declaration appertains to the parameter.
- 14 For a function declarator without a parameter type list: the effect is as if it were declared with a parameter type list consisting of the keyword **void**. A function declarator provides a prototype for the function.
- 15 For two function types to be compatible, both shall specify compatible return types. Moreover, the parameter type lists shall agree in the number of parameters and in whether the function is variadic or not; corresponding parameters shall have compatible types. In the determination of type compatibility and of a composite type, each parameter declared with function or array type is taken as having the adjusted type and each parameter declared with qualified type is taken as having the unqualified version of its declared type.
- 16 EXAMPLE 1 The declaration

```
int f(void), *fip(), (*pfi)();
```

declares a function **f** with no parameters returning an **int**, a function **fip** with no parameters returning a pointer to an **int**, and a pointer **pfi** to a function with no parameters returning an **int**. It is especially useful to compare the last two. The binding of `*fip()` is `*(fip())`, so that the declaration suggests, and the same construction in an expression requires, the calling of a function **fip**, and then using indirection through the pointer result to yield an **int**. In the declarator `(*pfi)()`, the extra parentheses are necessary to indicate that indirection through a pointer to a function yields a function designator, which is then used to call the function; it returns an **int**.

- 17 If the declaration occurs outside of any function, the identifiers have file scope and external linkage. If the declaration occurs inside a function, the identifiers of the functions **f** and **fip** have block scope and either internal or external linkage (depending on what file scope declarations for these identifiers are visible), and the identifier of the pointer **pfi** has block scope and no linkage.
- 18 EXAMPLE 2 The declaration

```
int (*apfi[3])(int *x, int *y);
```

declares an array **apfi** of three pointers to functions returning **int**. Each of these functions has two parameters that are pointers to **int**. The identifiers **x** and **y** are declared for descriptive purposes only and go out of scope at the end of the declaration of **apfi**.

- 19 EXAMPLE 3 The declaration

```
int (*fpfi(int (*)(long), int))(int, ...);
```

declares a function **fpfi** that returns a pointer to a function returning an **int**. The function **fpfi** has two parameters: a pointer to a function returning an **int** (with one parameter of type **long int**), and an **int**. The pointer returned by **fpfi** points to a function that has one **int** parameter and accepts zero or more additional arguments of any type.

- 20 EXAMPLE 4 The following prototype has a variably modified parameter.

```
void addscalar(int n, int m,
 double a[n][n*m+300], double x);

int main(void)
{
 double b[4][308];
 addscalar(4, 2, b, 2.17);
 return 0;
}
```

```

void addscalar(int n, int m,
 double a[n][n*m+300], double x)
{
 for (int i = 0; i < n; i++)
 for (int j = 0, k = n*m+300; j < k; j++)
 // a is a pointer to a VLA with n*m+300 elements
 a[i][j] += x;
}

```

21 EXAMPLE 5 The following are all compatible function prototype declarators.

```

double maximum(int n, int m, double a[n][m]);
double maximum(int n, int m, double a[*][*]);
double maximum(int n, int m, double a[][*]);
double maximum(int n, int m, double a[][m]);

```

as are:

```

void f(double (* restrict a)[5]);
void f(double a[restrict][5]);
void f(double a[restrict 3][5]);
void f(double a[restrict static 3][5]);

```

The last declaration also specifies that the argument corresponding to *a* in any call to *f* can be expected to be a non-null pointer to the first of at least three arrays of 5 doubles, which the others do not.

**Forward references:** function definitions (6.9.2), type names (6.7.8).

## 6.7.8 Type names

### Syntax

- 1 *type-name*:
 

*specifier-qualifier-list abstract-declarator*<sub>opt</sub>  
*abstract-declarator*:
 

*pointer*  
*pointer*<sub>opt</sub> *direct-abstract-declarator*  
*direct-abstract-declarator*:
 

**( abstract-declarator )**  
*array-abstract-declarator attribute-specifier-sequence*<sub>opt</sub>  
*function-abstract-declarator attribute-specifier-sequence*<sub>opt</sub>  
*array-abstract-declarator*:
 

*direct-abstract-declarator*<sub>opt</sub> **[ type-qualifier-list**<sub>opt</sub> *assignment-expression*<sub>opt</sub> **]**  
*direct-abstract-declarator*<sub>opt</sub> **[ static type-qualifier-list**<sub>opt</sub> *assignment-expression* **]**  
*direct-abstract-declarator*<sub>opt</sub> **[ type-qualifier-list static assignment-expression ]**  
*direct-abstract-declarator*<sub>opt</sub> **[ \* ]**

*function-abstract-declarator*:
 

*direct-abstract-declarator*<sub>opt</sub> **( parameter-type-list**<sub>opt</sub> **)**

### Semantics

- 2 In several contexts, it is necessary to specify a type. This is accomplished using a *type name*, which is syntactically a declaration for a function or an object of that type that omits the identifier.<sup>165)</sup> The optional attribute specifier sequence in a direct abstract declarator appertains to the preceding array or function type. The attribute specifier sequence affects the type only for the declaration it appears in, not other declarations involving the same type.

<sup>165)</sup>As indicated by the syntax, empty parentheses in a type name are interpreted as “function with no parameters”, rather than redundant parentheses around the omitted identifier.

## 3 EXAMPLE The constructions

```

(a) int
(b) int *
(c) int *[3]
(d) int (*)[3]
(e) int (*)[*]
(f) int *()
(g) int (*)(void)
(h) int (*const [])(unsigned int, ...)

```

name respectively the types

- (a) **int**,
- (b) pointer to **int**,
- (c) array of three pointers to **int**,
- (d) pointer to an array of three **int** s,
- (e) pointer to a variable length array of an unspecified number of **int** s,
- (f) function with no parameters returning a pointer to **int**,
- (g) pointer to function with no parameters returning an **int**, and
- (h) array of an unspecified number of constant pointers to functions, each with one parameter that has type **unsigned int** and an unspecified number of other parameters, returning an **int**.

## 6.7.9 Type definitions

### Syntax

- 1 *typedef-name:*  
*identifier*

### Constraints

- 2 If a typedef name specifies a variably modified type then it shall have block scope.

### Semantics

- 3 In a declaration whose storage-class specifier is **typedef**, each declarator defines an identifier to be a typedef name that denotes the type specified for the identifier in the way described in 6.7.7. Any array size expressions associated with variable length array declarators and typeof operators are evaluated each time the declaration of the typedef name is reached in the order of execution. A **typedef** declaration does not introduce a new type, only a synonym for the type so specified. That is, in the following declarations:

```

typedef T type_ident;
type_ident D;

```

**type\_ident** is defined as a typedef name with the type specified by the declaration specifiers in **T** (known as *T*), and the identifier in **D** has the type “*derived-declarator-type-list T*” where the *derived-declarator-type-list* is specified by the declarators of **D**. A typedef name shares the same name space as other identifiers declared in ordinary declarators. If the identifier is redeclared in an enclosed block, the type of the inner declaration shall not be inferred (6.7.10).

- 4 EXAMPLE 1 After

```

typedef int MILES, KLICKSP();
typedef struct { double hi, lo; } range;

```

the constructions

```

MILES distance;
extern KLICKSP *metricp;
range x;
range z, *zp;

```

are all valid declarations. The type of `distance` is `int`, that of `metricp` is “pointer to function with no parameters returning `int`”, and that of `x` and `z` is the specified structure; `zp` is a pointer to such a structure. The object `distance` has a type compatible with any other `int` object.

- 5 EXAMPLE 2 After the declarations

```

typedef struct s1 { int x; } t1, *tp1;
typedef struct s2 { int x; } t2, *tp2;

```

type `t1` and the type pointed to by `tp1` are compatible. Type `t1` is also compatible with type `struct s1`, but not compatible with the types `struct s2`, `t2`, the type pointed to by `tp2`, or `int`.

- 6 EXAMPLE 3 The following obscure constructions

```

typedef signed int t;
typedef int plain;
struct tag {
 unsigned t:4;
 const t:5;
 plain r:5;
};

```

declare a typedef name `t` with type `signed int`, a typedef name `plain` with type `int`, and a structure with three bit-field members, one named `t` that contains values in the range  $[0, 15]$ , an unnamed const-qualified bit-field which (if it can be accessed) would contain values in the range  $[-16, +15]$ , and one named `r` that contains values in one of the ranges  $[0, 31]$  or  $[-16, +15]$ . (The choice of range is implementation-defined.) The first two bit-field declarations differ in that `unsigned` is a type specifier (which forces `t` to be the name of a structure member), while `const` is a type qualifier (which modifies `t` which is still visible as a typedef name). If these declarations are followed in an inner scope by

```

t f(t (t));
long t;

```

then a function `f` is declared with type “function returning `signed int` with one unnamed parameter with type pointer to function returning `signed int` with one unnamed parameter with type `signed int`”, and an identifier `t` with type `long int`.

- 7 EXAMPLE 4 On the other hand, typedef names can be used to improve code readability. All three of the following declarations of the `signal` function specify exactly the same type, the first without making use of any typedef names.

```

typedef void fv(int), (*pfv)(int);

void (*signal(int, void (*)(int)))(int);
fv *signal(int, fv *);
pfv signal(int, pfv);

```

- 8 EXAMPLE 5 If a typedef name denotes a variable length array type, the length of the array is fixed at the time the typedef name is defined, not each time it is used:

```

void copyt(int n)
{
 typedef int B[n]; // B is n ints, n evaluated now
 n += 1;
 B a; // a is n ints, n without += 1
 int b[n]; // a and b are different sizes
 for (int i = 1; i < n; i++)

```

```

 a[i-1] = b[i];
 }

```

### 6.7.10 Type inference

#### Constraints

- 1 A declaration for which the type is inferred shall contain the storage-class specifier **auto**.

#### Semantics

- 2 For such a declaration that is the definition of an object the init-declarator shall have the form

*direct-declarator* = *assignment-expression*

The inferred type of the declared object is the type of the assignment expression after lvalue, array to pointer or function to pointer conversion, additionally qualified by qualifiers and amended by attributes as they appear in the declaration specifiers, if any.<sup>166)</sup> Implementations can accept a direct declarator that is not of the form

*identifier* *attribute-specifier-sequence*<sub>opt</sub>

optionally enclosed in balanced pairs of parentheses; if a direct declarator of a different form is accepted, the behavior is implementation-defined.<sup>167)</sup>

- 3 NOTE A declaration that also defines a structure or union type has implementation-defined behavior. Here, the identifier **x** which is not ordinary but in the name space of the structure type is declared.

```
auto p = (struct { int x; } *)0;
```

Even a forward declaration of a structure tag

```
struct s;
auto p = (struct s { int x; } *)0;
```

would not change that situation. A direct use of the structure definition as the type specifier ensures portability of the declaration.

```
struct s { int x; } * p = 0;
```

The following also has implementation-defined behavior:

```
auto alignas (struct s *) x = 0;
```

- 4 EXAMPLE 1 The following file scope definitions:

```
static auto a = 3.5;
auto p = &a;
```

are interpreted as if they had been written as:

```
static double a = 3.5;
double * p = &a;
```

So effectively **a** is a **double** and **p** is a **double\***. The restrictions on the syntax of such declarations does not allow the declarator to be **\*p**, but that the final type here nevertheless is a pointer type.

- 5 EXAMPLE 2 The scope of the identifier for which the type is inferred only starts after the end of the initializer (6.2.1), so the assignment expression cannot use the identifier to refer to the object or function that is declared, for example to take its address. Any use of the identifier in the initializer is invalid, even if an entity with the same name exists in an outer scope.

<sup>166)</sup>The scope rules as described in 6.2.1 also prohibit the use of the identifier of the declarator within the assignment expression.

<sup>167)</sup>It is recommended that implementations that accept different forms of direct declarators follow the syntax and semantics of the corresponding feature in ISO/IEC 14882.

```

{
 double a = 7;
 double b = 9;
 {
 double b = b * b; // undefined, uses uninitialized
 // variable without address
 printf("%g\n", a); // valid, uses "a" from outer scope, prints 7
 auto a = a * a; // invalid, "a" from outer scope is not
 // visible during initialization
 }
 {
 auto b = a * a; // valid, uses "a" from outer scope
 auto a = b; // valid, "a" from outer scope not visible now
 // ...
 printf("%g\n", a); // valid, uses "a" from inner scope, prints 49
 }
 // ...
}

```

- 6 EXAMPLE 3 In the following, declarations of `pA` and `qA` are valid. The type of `A` after array-to-pointer conversion is a pointer type, and `qA` is a pointer to array.

```

double A[3] = { 0 };
auto pA = A;
auto qA = &A;

```

- 7 EXAMPLE 4 Type inference can be used to capture the type of a call to a type-generic function. It ensures that the same type as the argument `x` is used.

```

#include <tgmath.h>
auto y = cos(x);

```

If instead the type of `y` is explicitly specified to a different type than `x`, a diagnosis of the mismatch is not enforced.

- 8 EXAMPLE 5 A type-generic macro that generalizes the `div` functions (7.24.7.2) is defined and used as follows.

```

#define div(X, Y) _Generic((X)+(Y), \
 int: div, \
 long: ldiv, \
 long long: lldiv)((X), (Y))
auto z = div(x, y);
auto q = z.quot;
auto r = z.rem;

```

- 9 EXAMPLE 6 Definitions of objects with inferred type are valid in all contexts that allow the initializer syntax as described. In particular they can be used to ensure type safety of `for`-loop controlling expressions.

```

for (auto i = j; i < 2*j; ++i) {
 // ...
}

```

Here, regardless of the integer rank or signedness of the type of `j`, `i` will have the non-atomic unqualified version of `j`'s type. So, after lvalue conversion and possible promotion, the two operands of the `<` operator in the controlling expression are guaranteed to have the same type, and, in particular, the same signedness.

## 6.7.11 Initialization

### Syntax

- 1 *braced-initializer*:
- ```
{ }
```

```
{ initializer-list }
{ initializer-list , }
```

initializer:

```
assignment-expression
braced-initializer
```

initializer-list:

```
designationopt initializer
initializer-list , designationopt initializer
```

designation:

```
designator-list =
```

designator-list:

```
designator
designator-list designator
```

designator:

```
[ constant-expression ]
. identifier
```

- 2 An empty brace pair ({}) is called an *empty initializer* and is referred to as *empty initialization*.

Constraints

- 3 No initializer shall attempt to provide a value for an object not contained within the entity being initialized.
- 4 The type of the entity to be initialized shall be an array of unknown size or a complete object type. An entity of variable length array type shall not be initialized except by an empty initializer. An array of unknown size shall not be initialized by an empty initializer.
- 5 The initializer for a scalar shall be a single expression, optionally enclosed in braces, or it shall be an empty initializer.
- 6 The initializer for an object that has structure or union type shall be either a single expression that has compatible type or a brace-enclosed list of initializers for the elements or named members.
- 7 The initializer for an array shall be either a string literal, optionally enclosed in braces, or a brace-enclosed list of initializers for the elements. An array initialized by a character string literal or UTF-8 string literal shall have a character type as element type. An array initialized with a wide string literal shall have element type compatible with a qualified or unqualified **wchar_t**, **char16_t**, or **char32_t**, and the string literal shall have the corresponding encoding prefix (L, u, or U, respectively).
- 8 All the expressions in an initializer for an object that has static or thread storage duration or is declared with the **constexpr** storage-class specifier shall be constant expressions or string literals.
- 9 If the declaration of an identifier has block scope, and the identifier has external or internal linkage, the declaration shall have no initializer for the identifier.
- 10 If a designator has the form

```
[ constant-expression ]
```

then the current object (defined subsequently in this subclause) shall have array type and the expression shall be an integer constant expression. If the array is of unknown size, any nonnegative value is valid.

- 11 If a designator has the form

```
. identifier
```

then the current object (defined subsequently in this subclause) shall have structure or union type and the identifier shall be the name of a member of that type.

Semantics

- 12 An initializer specifies the initial value stored in an object. For objects with atomic type additional restrictions apply, see 7.17.2 and 7.17.8.
- 13 Except where explicitly stated otherwise, for the purposes of this subclause unnamed members of objects of structure and union type do not participate in initialization. Unnamed members of structure objects have indeterminate representation even after initialization.
- 14 If an object that has automatic storage duration is not initialized explicitly, its representation is indeterminate. If an object that has static or thread storage duration is not initialized explicitly, or any object is initialized with an empty initializer, then it is subject to *default initialization*, which initializes an object as follows:
 - if it has pointer type, it is initialized to a null pointer;
 - if it has decimal floating type, it is initialized to positive zero, and the quantum exponent is implementation-defined;¹⁶⁸⁾
 - if it has arithmetic type, and it does not have decimal floating type, it is initialized to (positive or unsigned) zero;
 - if it is an aggregate, every member is initialized (recursively) according to these rules, and any padding is initialized to zero bits;
 - if it is a union, the first named member is initialized (recursively) according to these rules, and any padding is initialized to zero bits.
- 15 If the initializer for a scalar is not the empty initializer, the initial value of the object is that of the expression (after conversion); the same type constraints and conversions as for simple assignment apply, taking the type of the scalar to be the unqualified version of its declared type.
- 16 The rest of this subclause deals with initializers for objects that have aggregate or union type.
- 17 If the initializer for a struct or a union is a single expression, the initial value of the object, including unnamed members, is that of the expression.¹⁶⁹⁾
- 18 For an array of character type initialized by an ordinary string literal or UTF-8 string literal, successive bytes of the string literal (including the terminating null character if there is room or if the array is of unknown size) initialize the elements of the array.
- 19 For an array initialized by a wide string literal, successive wide characters of the wide string literal (including the terminating null wide character if there is room or if the array is of unknown size) initialize the elements of the array.
- 20 Each brace-enclosed initializer list has an associated *current object*. When no designations are present, subobjects of the current object are initialized in order according to the type of the current object: array elements in increasing subscript order, structure members in declaration order, and the first named member of a union.¹⁷⁰⁾ In contrast, a designation causes the following initializer to begin initialization of the subobject described by the designator. Initialization then continues forward in order, beginning with the next subobject after that described by the designator.¹⁷¹⁾
- 21 Each designator list begins its description with the current object associated with the closest surrounding brace pair. Each item in the designator list (in order) specifies a particular member of its current object and changes the current object for the next designator (if any) to be that member.¹⁷²⁾

¹⁶⁸⁾A representation with all bits zero results in a decimal floating-point zero with the most negative exponent.

¹⁶⁹⁾If the object being initialized does not have automatic storage duration, this case violates a constraint unless the expression is a named constant or compound literal constant (6.6.1).

¹⁷⁰⁾If the initializer list for a subaggregate or contained union does not begin with a left brace, its subobjects are initialized as usual, but the subaggregate or contained union does not become the current object: current objects are associated only with brace-enclosed initializer lists.

¹⁷¹⁾After a union member is initialized, the next object is not the next member of the union; instead, it is the next subobject of an object containing the union.

¹⁷²⁾Thus, a designator can only specify a strict subobject of the aggregate or union that is associated with the surrounding brace pair. Note, too, that each separate designator list is independent.

The current object that results at the end of the designator list is the subobject to be initialized by the following initializer.

- 22 The initialization shall occur in initializer list order, each initializer provided for a particular subobject overriding any previously listed initializer for the same subobject;¹⁷³⁾ all subobjects that are not initialized explicitly are subject to default initialization.
- 23 If the aggregate or union contains elements or members that are aggregates or unions, these rules apply recursively to the subaggregates or contained unions. If the initializer of a subaggregate or contained union begins with a left brace, the initializers enclosed by that brace and its matching right brace initialize the elements or members of the subaggregate or the contained union. Otherwise, only enough initializers from the list are taken to account for the elements or members of the subaggregate or the first member of the contained union; any remaining initializers are left to initialize the next element or member of the aggregate of which the current subaggregate or contained union is a part.
- 24 If there are fewer initializers in a brace-enclosed list than there are elements or members of an aggregate, or fewer characters in a string literal used to initialize an array of known size than there are elements in the array, the remainder of the aggregate is subject to default initialization.
- 25 If an array of unknown size is initialized, its size is determined by the largest indexed element with an explicit initializer. The array type is completed at the end of its initializer list.
- 26 The evaluations of the initialization list expressions are indeterminately sequenced with respect to one another and thus the order in which any side effects occur is unspecified.¹⁷⁴⁾
- 27 EXAMPLE 1 Provided that <complex.h> has been included, the declarations

```
int i = 3.5;
double complex c = 5 + 3 * I;
```

define and initialize `i` with the value 3 and `c` with the value $5.0 + i3.0$.

- 28 EXAMPLE 2 The declaration

```
int x[] = { 1, 3, 5 };
```

defines and initializes `x` as a one-dimensional array object that has three elements, as no size was specified and there are three initializers.

- 29 EXAMPLE 3 The declaration

```
int y[4][3] = {
    { 1, 3, 5 },
    { 2, 4, 6 },
    { 3, 5, 7 },
};
```

is a definition with a fully bracketed initialization: 1, 3, and 5 initialize the first row of `y` (the array object `y[0]`), namely `y[0][0]`, `y[0][1]`, and `y[0][2]`. Likewise the next two lines initialize `y[1]` and `y[2]`. The initializer ends early, so `y[3]` is initialized with zeros. Precisely the same effect can be achieved by

```
int y[4][3] = {
    1, 3, 5, 2, 4, 6, 3, 5, 7
};
```

The initializer for `y[0]` does not begin with a left brace, so three items from the list are used. Likewise the next three are taken successively for `y[1]` and `y[2]`.

- 30 EXAMPLE 4 The declaration

```
int z[4][3] = {
```

¹⁷³⁾Any initializer for the subobject which is overridden and so not used to initialize that subobject can potentially not be evaluated at all.

¹⁷⁴⁾In particular, the evaluation order can be the same or different as the order of subobject initialization.

```
};    { 1 }, { 2 }, { 3 }, { 4 }
```

initializes the first column of **z** as specified and initializes the rest with zeros.

31 EXAMPLE 5 The declaration

```
struct { int a[3], b; } w[] = { { 1 }, 2 };
```

is a definition with an inconsistently bracketed initialization. It defines an array with two element structures: **w[0].a[0]** is 1 and **w[1].a[0]** is 2; all the other elements are zero.

32 EXAMPLE 6 The declaration

```
short q[4][3][2] = {
    { 1 },
    { 2, 3 },
    { 4, 5, 6 }
};
```

contains an incompletely but consistently bracketed initialization. It defines a three-dimensional array object: **q[0][0][0]** is 1, **q[1][0][0]** is 2, **q[1][0][1]** is 3, and 4, 5, and 6 initialize **q[2][0][0]**, **q[2][0][1]**, and **q[2][1][0]**, respectively; all the rest are zero. The initializer for **q[0][0]** does not begin with a left brace, so up to six items from the current list can be used. There is only one, so the values for the remaining five elements are initialized with zero. Likewise, the initializers for **q[1][0]** and **q[2][0]** do not begin with a left brace, so each uses up to six items, initializing their respective two-dimensional subaggregates. If there had been more than six items in any of the lists, a diagnostic message would have been issued. The same initialization result can be achieved by:

```
short q[4][3][2] = {
    1, 0, 0, 0, 0, 0,
    2, 3, 0, 0, 0, 0,
    4, 5, 6
};
```

or by:

```
short q[4][3][2] = {
    {
        { 1 },
    },
    {
        { 2, 3 },
    },
    {
        { 4, 5 },
        { 6 },
    }
};
```

in a fully bracketed form.

33 The fully bracketed and minimally bracketed forms of initialization are, in general, less likely to cause confusion.

34 EXAMPLE 7 One form of initialization that completes array types involves typedef names. Given the declaration

```
typedef int A[]; // OK - declared with block scope
```

the declaration

```
A a = { 1, 2 }, b = { 3, 4, 5 };
```

is identical to

```
int a[] = { 1, 2 }, b[] = { 3, 4, 5 };
```

due to the rules for incomplete types.

35 EXAMPLE 8 The declaration

```
char s[] = "abc", t[3] = "abc";
```

defines “plain” **char** array objects **s** and **t** whose elements are initialized with ordinary string literals. This declaration is identical to

```
char s[] = { 'a', 'b', 'c', '\0' },
t[] = { 'a', 'b', 'c' };
```

The contents of the arrays are modifiable. On the other hand, the declaration

```
char *p = "abc";
```

defines **p** with type “pointer to **char**” and initializes it to point to an object with type “array of **char**” with length 4 whose elements are initialized with an ordinary string literal. If an attempt is made to use **p** to modify the contents of the array, the behavior is undefined.

36 EXAMPLE 9 Arrays can be initialized to correspond to the elements of an enumeration by using designators:

```
enum { member_one, member_two };
const char *nm[] = {
    [member_two] = "member two",
    [member_one] = "member one",
};
```

37 EXAMPLE 10 Structure members can be initialized to nonzero values without depending on their order:

```
div_t answer = { .quot = 2, .rem = -1 };
```

38 EXAMPLE 11 Designators can be used to provide explicit initialization when unadorned initializer lists are difficult to understand:

```
struct { int a[3], b; } w[] =
    { [0].a = {1}, [1].a[0] = 2 };
```

39 EXAMPLE 12

```
struct T {
    int k;
    int l;
};

struct S {
    int i;
    struct T t;
};

struct T x = { .l = 43, .k = 42, };

void f(void)
{
    struct S l = { 1, .t = x, .t.l = 41, };
}
```

The value of **l.t.k** is 42, because implicit initialization does not override explicit initialization.

- 40 EXAMPLE 13 Space can be “allocated” from both ends of an array by using a single designator:

```
int a[A_MAX] = {
    1, 3, 5, 7, 9, [A_MAX-5] = 8, 6, 4, 2, 0
};
```

In the preceding snippet, if **A_MAX** is greater than ten, there will be some zero-valued elements in the middle; if it is less than ten, some of the values provided by the first five initializers will be overridden by the second five.

- 41 EXAMPLE 14 Any member of a union can be initialized:

```
union { /* ... */ } u = {.any_member = 42 };
```

Forward references: common definitions <stddef.h> (7.21).

6.7.12 Static assertions

Syntax

- 1 *static_assert-declaration:*
- ```
static_assert (constant-expression , string-literal) ;
static_assert (constant-expression) ;
```

### Constraints

- 2 The constant expression shall compare unequal to 0.

### Semantics

- 3 The constant expression shall be an integer constant expression. If the value of the constant expression compares unequal to 0, the declaration has no effect. Otherwise, the constraint is violated and the implementation shall produce a diagnostic message which should include the text of the string literal, if present.

**Forward references:** diagnostics (7.2).

## 6.7.13 Attributes

### 6.7.13.1 Introduction

- Attributes specify additional information for various source constructs such as types, objects, identifiers, or blocks. They are identified by an *attribute token*, which can either be a *attribute prefixed token* (for implementation-specific attributes) or a *standard attribute* specified by an identifier (for attributes specified in this document).
- Support for any of the standard attributes specified in this document is implementation-defined and optional. For an attribute token (including an attribute prefixed token) not specified in this document, the behavior is implementation-defined. Any attribute token that is not supported by the implementation is ignored.
- Attributes are said to appertain to some source construct, identified by the syntactic context where they appear, and for each individual attribute, the corresponding clause constrains the syntactic context in which this appurtenance is valid. The attribute specifier sequence appertaining to some source construct shall contain only attributes that are allowed to apply to that source construct.
- In all aspects of the language, a standard attribute specified by this document as an identifier **attr** and an identifier of the form **\_\_attr\_\_** shall behave the same when used as an attribute token, except for the spelling.<sup>175)</sup>
- For all standard attributes specified by this document, the current value when its token sequence is given to the **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) is written in the associated subclause for that attribute. A history of those values can be found in Table M.1.

<sup>175)</sup>Thus, the attributes **[\_\_nodiscard\_\_]** and **[\_\_nodiscard\_\_]** can be freely interchanged. Implementations are encouraged to behave similarly for attribute tokens (including attribute prefixed tokens) they provide.

**Recommended practice**

- 6 It is recommended that implementations support all standard attributes as defined in this document.

**6.7.13.2 General****Syntax**

- 1 *attribute-specifier-sequence*:  
                   *attribute-specifier-sequence*<sub>opt</sub> *attribute-specifier*  
*attribute-specifier*:  
                   [ [ *attribute-list* ] ]  
*attribute-list*:  
                   *attribute*<sub>opt</sub>  
                   *attribute-list* , *attribute*<sub>opt</sub>  
*attribute*:  
                   *attribute-token* *attribute-argument-clause*<sub>opt</sub>  
*attribute-token*:  
                   *standard-attribute*  
                   *attribute-prefixed-token*  
*standard-attribute*:  
                   *identifier*  
  
*attribute-prefixed-token*:  
                   *attribute-prefix* :: *identifier*  
*attribute-prefix*:  
                   *identifier*  
*attribute-argument-clause*:  
                   ( *balanced-token-sequence*<sub>opt</sub> )  
*balanced-token-sequence*:  
                   *balanced-token*  
                   *balanced-token-sequence* *balanced-token*  
*balanced-token*:  
                   ( *balanced-token-sequence*<sub>opt</sub> )  
                   [ *balanced-token-sequence*<sub>opt</sub> ]  
                   { *balanced-token-sequence*<sub>opt</sub> }  
                   any token other than a parenthesis, a bracket, or a brace

**Constraints**

- 2 The identifier in a standard attribute shall be one of:

**deprecated**  
**fallthrough**

**maybe\_unused**  
**nodiscard**

**noreturn**  
**\_Noreturn**

**unsequenced**  
**reproducible**

**Semantics**

- 3 An attribute specifier that contains no attributes has no effect. The order in which attribute tokens appear in an attribute list is not significant. If a keyword (6.4.2) that satisfies the syntactic requirements of an identifier (6.4.3) is contained in an attribute token, it is considered an identifier. A strictly conforming program using a standard attribute remains strictly conforming in the absence of that attribute.<sup>176)</sup>
- 4 NOTE For each standard attribute, the form of the balanced token sequence, if any, will be specified.

**Recommended practice**

- 5 Each implementation should choose a distinctive name for the attribute prefix in an attribute prefixed token. Implementations should not define attributes without an attribute prefix unless it is

<sup>176)</sup>Standard attributes specified by this document can be parsed but ignored by an implementation without changing the semantics of a correct program; the same is not true for attributes not specified by this document.

a standard attribute as specified in this document.

- 6 EXAMPLE 1 Suppose that an implementation chooses the attribute prefix `hal` and provides specific attributes named `daisy` and `rosie`.

```
[[deprecated, hal::daisy]] double nine1000(double);
[[deprecated]] [[hal::daisy]] double nine1000(double);
[[deprecated]] double nine1000 [[hal::daisy]] (double);
```

Then all the following declarations should be equivalent aside from the spelling:

```
[[__deprecated__, __hal__::__daisy__]] double nine1000(double);
[[__deprecated__]] [[__hal__::__daisy__]] double nine1000(double);
[[__deprecated__]] double nine1000 [[__hal__::__daisy__]] (double);
```

These use the alternate spelling that is required for all standard attributes and recommended for prefixed attributes. These can be better-suited for use in header files, where the use of the alternate spelling avoids naming conflicts with user-provided macros.

- 7 EXAMPLE 2 For the same implementation, the following two declarations are equivalent, because the ordering inside attribute lists is not important.

```
[[hal::daisy, hal::rosie]] double nine999(double);
[[hal::rosie, hal::daisy]] double nine999(double);
```

On the other hand the following two declarations are not equivalent, because the ordering of different attribute specifiers can affect the semantics.

```
[[hal::daisy]] [[hal::rosie]] double nine999(double);
[[hal::rosie]] [[hal::daisy]] double nine999(double); // can have different semantics
```

### 6.7.13.3 The **nodiscard** attribute

#### Constraints

- 1 The **nodiscard** attribute shall be applied to a function or to the definition of a structure, union, or enumerated type. If an attribute argument clause is present, it shall have the form:

( *string-literal* )

#### Semantics

- 2 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **nodiscard** as the pp-tokens operand if the implementation supports the attribute.
- 3 A name or entity declared without the **nodiscard** attribute can later be redeclared with the attribute and vice versa. An entity is considered marked after the first declaration that marks it.

#### Recommended practice

- 4 A **nodiscard** call is a function call expression that calls a function previously declared with attribute **nodiscard**, or whose return type is a structure, union, or enumerated type marked with attribute **nodiscard**. Evaluation of a **nodiscard** call as a void expression (6.8.4) is discouraged unless explicitly cast to **void**. Implementations are encouraged to issue a diagnostic in such cases. This is typically because immediately discarding the return value of a **nodiscard** call has surprising consequences.
- 5 The diagnostic message should include text provided by the string literal within the attribute argument clause of any **nodiscard** attribute applied to the name or entity.
- 6 EXAMPLE 1

```
struct [[nodiscard]] error_info { /*...*/ };
struct error_info enable_missile_safety_mode(void);
void launch_missiles(void);
void test_missiles(void) {
 enable_missile_safety_mode();
}
```

```
launch_missiles();
}
```

A diagnostic for the call to `enable_missile_safety_mode` is encouraged.

#### 7 EXAMPLE 2

```
[[nodiscard]] int important_func(void);
void call(void) {
 int i = important_func();
}
```

No diagnostic for the call to `important_func` is encouraged despite the value of `i` not being used.

#### 8 EXAMPLE 3

```
[[nodiscard("armer needs to check armed state")]]
bool arm_detonator(int within);

void call(void) {
 arm_detonator(3);
 detonate();
}
```

A diagnostic for the call to `arm_detonator` using the string literal "armer needs to check armed state" from the attribute argument clause is encouraged.

### 6.7.13.4 The `maybe_unused` attribute

#### Constraints

- 1 The `maybe_unused` attribute shall be applied to the declaration of a structure, a union, a `typedef` name, an object, a structure or union member, a function, an enumeration, an enumerator, or a label. No attribute argument clause shall be present.

#### Semantics

- 2 The `maybe_unused` attribute indicates that a name or entity is possibly intentionally unused.
- 3 The `__has_c_attribute` conditional inclusion expression (6.10.2) shall return the value 202311L when given `maybe_unused` as the pp-tokens operand if the implementation supports the attribute.

A name or entity declared without the `maybe_unused` attribute can later be redeclared with the attribute and vice versa. An entity is considered marked with the attribute after the first declaration that marks it.

#### Recommended practice

- 4 For an entity marked `maybe_unused`, implementations are encouraged not to emit a diagnostic that the entity is unused, or that the entity is used despite the presence of the attribute.

#### 5 EXAMPLE

```
[[maybe_unused]] void f([[maybe_unused]] int i) {
 [[maybe_unused]] int j = i + 100;
 assert(j);
}
```

Implementations are encouraged not to diagnose that `j` is unused, even if `NDEBUG` is defined.

### 6.7.13.5 The `deprecated` attribute

#### Constraints

- 1 The `deprecated` attribute shall be applied to the declaration of a structure, a union, a `typedef` name, an object, a structure or union member, a function, an enumeration, or an enumerator.
- 2 If an attribute argument clause is present, it shall have the form:



( *string-literal* )

### Semantics

- 3 The **deprecated** attribute can be used to mark names and entities whose use is still allowed, but is discouraged for some reason.<sup>177)</sup>
- 4 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **deprecated** as the pp-tokens operand if the implementation supports the attribute.
- 5 A name or entity declared without the **deprecated** attribute can later be redeclared with the attribute and vice versa. An entity is considered marked with the attribute after the first declaration that marks it.

### Recommended practice

- 6 Implementations should use the **deprecated** attribute to produce a diagnostic message in case the program refers to a name or entity other than to declare it, after a declaration that specifies the attribute, when the reference to the name or entity is not within the context of a related deprecated entity. The diagnostic message should include text provided by the string literal within the attribute argument clause of any **deprecated** attribute applied to the name or entity.
- 7 EXAMPLE

```

struct [[deprecated]] S {
 int a;
};

enum [[deprecated]] E1 {
 one
};

enum E2 {
 two [[deprecated("use 'three' instead")]],
 three
};

[[deprecated]] typedef int Foo;

void f1(struct S s) { // Diagnose use of S
 int i = one; // Diagnose use of E1
 int j = two; // Diagnose use of two: "use 'three' instead"
 int k = three;
 Foo f; // Diagnose use of Foo
}

[[deprecated]] void f2(struct S s) {
 int i = one;
 int j = two;
 int k = three;
 Foo f;
}

struct [[deprecated]] T {
 Foo f;
 struct S s;
};

```

Implementations are encouraged to diagnose the use of deprecated entities within a context which is not itself deprecated, as indicated for function **f1**, but not to diagnose within function **f2** and **struct T**, as they are themselves deprecated.

<sup>177)</sup>In particular, **deprecated** is appropriate for names and entities that are obsolescent, insecure, unsafe, or otherwise unfit for purpose.

### 6.7.13.6 The **fallthrough** attribute

#### Constraints

- 1 The attribute token **fallthrough** shall only appear in an attribute declaration (6.7); such a declaration is a *fallthrough declaration*. No attribute argument clause shall be present. A fallthrough declaration can only appear within an enclosing **switch** statement (6.8.5.3). The next block item (6.8.3) that would be encountered after a fallthrough declaration shall be a **case** label or **default** label associated with the innermost enclosing **switch** statement and, if the fallthrough declaration is contained in an iteration statement, the next statement shall be part of the same execution of the secondary block of the innermost enclosing iteration statement.

#### Semantics

- 2 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **fallthrough** as the pp-tokens operand if the implementation supports the attribute.

#### Recommended practice

- 3 The use of a fallthrough declaration is intended to suppress a diagnostic that an implementation can otherwise issue for a **case** or **default** label that is reachable from another **case** or **default** label along some path of execution. Implementations are encouraged to issue a diagnostic if a fallthrough declaration is not dynamically reachable.
- 4 EXAMPLE

```
void f(int n) {
 void g(void), h(void), i(void);
 switch (n) {
 case 1: /* diagnostic on fallthrough discouraged */
 case 2:
 g();
 [[fallthrough]];
 case 3: /* diagnostic on fallthrough discouraged */
 do {
 [[fallthrough]]; /* constraint violation: next statement is not
 part of the same secondary block execution */
 } while(false);
 case 6:
 do {
 [[fallthrough]]; /* constraint violation: next statement is not
 part of the same secondary block execution */
 } while (n--);
 case 7:
 while (false) {
 [[fallthrough]]; /* constraint violation: next statement is not
 part of the same secondary block execution */
 }
 case 5:
 h();
 case 4: /* fallthrough diagnostic encouraged */
 i();
 [[fallthrough]]; /* constraint violation */
 }
}
```

### 6.7.13.7 The **noreturn** and **\_Noreturn** attributes

#### Description

- 1 When **\_Noreturn** is used as an attribute token (instead of a function specifier), the constraints and semantics are identical to that of the **noreturn** attribute token. Use of **\_Noreturn** as an attribute

token is an obsolescent feature.<sup>178)</sup>

### Constraints

- 2 The **noreturn** attribute shall be applied to a function. No attribute argument clause shall be present.

### Semantics

- 3 The first declaration of a function shall specify the **noreturn** attribute if any declaration of that function specifies the **noreturn** attribute. If a function is declared with the **noreturn** attribute in one translation unit and the same function is declared without the **noreturn** attribute in another translation unit, the behavior is undefined.
- 4 If a function **f** is called where **f** was previously declared with the **noreturn** attribute and **f** eventually returns, the behavior is undefined.
- 5 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **noreturn** as the pp-tokens operand if the implementation supports the attribute.

### Recommended practice

- 6 The implementation should produce a diagnostic message for a function declared with a **noreturn** attribute that appears to be capable of returning to its caller.
- 7 EXAMPLE

```
[[noreturn]] void f(void) {
 abort(); // ok
}

[[noreturn]] void g(int i) { // causes undefined behavior if i <= 0
 if (i > 0) abort();
}

[[noreturn]] int h(void);
```

Implementations are encouraged to diagnose the definition of **g()** because it is capable of returning to its caller. Implementations are similarly encouraged to diagnose the declaration of **h()** because it appears capable of returning to its caller due to the non-**void** return type.

## 6.7.13.8 Standard attributes for function types

### 6.7.13.8.1 General

#### Constraints

- 1 The identifier in a standard function type attribute shall be one of:

**unsequenced**                      **reproducible**

- 2 An attribute for a function type shall be applied to a function declarator<sup>179)</sup> or to a type specifier that has a function type. The corresponding attribute is a property of the function type.<sup>180)</sup> No attribute argument clause shall be present.

#### Description

- 3 The main purpose of the function type properties and attributes defined in this clause is to provide the translator with information about the access of objects by a function such that certain properties of function calls can be deduced; the properties distinguish read operations (stateless and independent) and write operations (effectless, idempotent and reproducible) or a combination of

<sup>178)</sup> `[_Nonnull]` and `[[noreturn]]` are equivalent attributes to support code that includes `<stdnoreturn.h>`, because that header defines **noreturn** as a macro that expands to `_Nonnull`.

<sup>179)</sup> That is, they appear in the attributes right after the closing parenthesis of the parameter list, independently of whether the function type is, for example, used directly to declare a function or whether it is used in a pointer to function type.

<sup>180)</sup> If several declarations of the same function or function pointer are visible, regardless whether an attribute is present at several or just one of the declarators, it is attached to the type of the corresponding function definition, function pointer object, or function pointer value.

both (unsequenced). Although semantically attached to a function type, the attributes described are not part of the prototype of such a function, and redeclarations and conversions that drop such an attribute are valid and constitute compatible types. Conversely, if a definition that does not have the asserted property is accessed by a function declaration or a function pointer with a type that has the attribute, the behavior is undefined.<sup>181)</sup>

- 4 To allow reordering of calls to functions as they are described here, possible access to objects with a lifetime that starts before or ends after a call has to be restricted; effects on all objects that are accessed during a function call are restricted to the same thread as the call and the based-on relation between pointer parameters and lvalues (6.7.4.2) models the fact that objects do not change inadvertently during the call. In the following, an operation is said to be sequenced *during* a function call if it is sequenced after the start of the function call<sup>182)</sup> and before the call terminates. An object definition of an object *X* in a function *f* *escapes* if an access to *X* happens while no call to *f* is active. An object is *local* to a call to a function *f* if its lifetime starts and ends during the call or if it is defined by *f* but does not escape. A function call and an object *X* *synchronize* if all accesses to *X* that are not sequenced during the call happen before or after the call. Execution state that is described in the library clause, such as the floating-point environment, conversion state, locale, input/output streams, external files or **errno** are considered as objects for the purposes of these attributes; operations that access this state, even indirectly, are considered as lvalue conversions for the purposes of these attributes, and operations that allow to change this state are considered as store operations, for the purposes of these attributes.
- 5 A function definition *f* is *stateless* if any definition of an object of static or thread storage duration in *f* or in a function that is called by *f* is **const** but not **volatile** qualified.
- 6 An object *X* is *observed* by a function call if both synchronize, if *X* is not local to the call, if *X* has a lifetime that starts before the function call and if an access of *X* is sequenced during the call; the last value of *X*, if any, that is stored before the call is said to be the value of *X* that is observed by the call. A function pointer value *f* is *independent* if for any object *X* that is observed by some call to *f* through an lvalue that is not based on a parameter of the call, then all accesses to *X* in all calls to *f* during the same program execution observe the same value; otherwise if the access is based on a pointer parameter, there shall be a unique such pointer parameter *P* such that any access to *X* shall be to an lvalue that is based on *P*. A function definition is independent if the derived function pointer value is independent.
- 7 A store operation to an object *X* that is sequenced during a function call such that both synchronize is said to be *observable* if *X* is not local to the call, if the lifetime of *X* ends after the call, if the stored value is different from the value observed by the call, if any, and if it is the last value written before the termination of the call. An evaluation of a function call<sup>183)</sup> is *effectless* if any store operation that is sequenced during the call is the modification of an object that synchronizes with the call; if additionally the operation is observable, there shall be a unique pointer parameter *P* of the function such that any access to *X* shall be to an lvalue that is based on *P*. A function pointer value *f* is effectless if any evaluation of a function call that calls *f* is effectless. A function definition is effectless if the derived function pointer value is effectless.
- 8 An evaluation *E* is *idempotent* if a second evaluation of *E* can be sequenced immediately after the original one without changing the resulting value, if any, or the observable state of the execution. A function pointer value *f* is idempotent if any evaluation of a function call<sup>184)</sup> that calls *f* is idempotent. A function definition is idempotent if the derived function pointer value is idempotent.
- 9 A function is *reproducible* if it is effectless and idempotent; it is *unsequenced* if it is stateless, effectless, idempotent and independent.<sup>185)</sup>

<sup>181)</sup>That is, the fact that a function has one of these properties is in general not determined by the specification of the translation unit in which it is found; other translation units and specific run time conditions also condition the possible assertion of the properties.

<sup>182)</sup>The initializations of the parameters is sequenced during the function call.

<sup>183)</sup>This considers the evaluation of the function call itself, not the evaluation of a full function call expression. Such an evaluation is sequenced after all evaluations that determine *f* and the call arguments, if any, have been performed.

<sup>184)</sup>This considers the evaluation of the function call itself, not the evaluation of a full function call expression. Such an evaluation is sequenced after all evaluations that determine *f* and the call arguments, if any, have been performed.

<sup>185)</sup>A function call of an unsequenced function can be executed as early as the function pointer value, the values of the

- 10 NOTE 1 The synchronization requirements with respect to any accessed object *X* for the independence of functions provide boundaries up to which a function call can safely be reordered without changing the semantics of the program. If *X* is **const** but not **volatile** qualified the reordering is unconstrained. If it is an object that is conditioned in an initialization phase, for a single threaded program a synchronization is provided by the sequenced before relation and the reordering can, in principle, move the call just after the initialization. For a multi-threaded program, synchronization guarantees can be given by calls to synchronizing functions of the `<threads.h>` header or by an appropriate call to **atomic\_thread\_fence** at the end of the initialization phase. If a function is known to be independent or effectless, adding **restrict** qualifications to the declarations of all pointer parameters does not change the semantics of any call. Similarly, changing the memory order to **memory\_order\_relaxed** for all atomic operations during a call to such a function preserves semantics.
- 11 NOTE 2 In general the functions provided by the `<math.h>` header do not have the properties that are defined previously in this subclause; many of them change the floating-point state or **errno** when they encounter an error (so they have observable side effects) and the results of most of them depend on execution-wide state such as the rounding direction mode (so they are not independent). Whether a particular C library function is reproducible or unsequenced additionally often depends on properties of the implementation, such as implementation-defined behavior for certain error conditions.

### Recommended practice

- 12 If possible, it is recommended that implementations diagnose if an attribute of this clause is applied to a function definition that does not have the corresponding property. It is recommended that applications that assert the independent or effectless properties for functions qualify pointer parameters with **restrict**.

**Forward references:** errors `<errno.h>` (7.5), floating-point environment `<fenv.h>` (7.6), localization `<locale.h>` (7.11), mathematics `<math.h>` (7.12), fences (7.17.4), input/output `<stdio.h>` (7.23), threads `<threads.h>` (7.29), extended multibyte and wide character utilities `<wchar.h>` (7.32).

#### 6.7.13.8.2 The **reproducible** type attribute

##### Description

- 1 The **reproducible** type attribute asserts that a function or pointed-to function with that type is reproducible.
- 2 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **reproducible** as the pp-tokens operand if the implementation supports the attribute.
- 3 EXAMPLE The attribute in the following function declaration asserts that two consecutive calls to the function will result in the same return value. Changes to the abstract state during the call are possible as long as they are not observable, but no other side effects will occur. Thus the function definition can for example use local objects of static or thread storage duration to keep track of the arguments for which the function has been called and cache their computed return values.

```
size_t hash(char const[static 32]) [[reproducible]];
```

#### 6.7.13.8.3 The **unsequenced** type attribute

##### Description

- 1 The **unsequenced** type attribute asserts that a function or pointed-to function with that type is unsequenced.
- 2 The **\_\_has\_c\_attribute** conditional inclusion expression (6.10.2) shall return the value 202311L when given **unsequenced** as the pp-tokens operand if the implementation supports the attribute.
- 3 NOTE 1 The unsequenced type attribute asserts strong properties for such a function, in particular that certain sequencing requirements for function calls can be relaxed without affecting the state of the abstract machine. Thereby, calls to such functions are natural candidates for optimization techniques such as common subexpression elimination, local memoization or lazy evaluation.
- 4 NOTE 2 A proof of validity of the annotation of a function type with the **unsequenced** attribute can depend

arguments and all objects that are accessible through them, and all values of globally accessible state have been determined, and it can be executed as late as the arguments and the objects they possibly target are unchanged and as any of its return value or modified pointed-to arguments are accessed.

on the property of whether a derived function pointer escapes the translation unit or not. For a function with internal linkage where no function pointer escapes the translation unit, all calling contexts are known and it is possible, in principle, to prove that no control flow exists such that a library function is called with arguments that trigger an exceptional condition. For a function with external linkage such a proof is potentially not possible and such a function therefore only is used when it can be ensured that no exceptional condition results from the provided arguments.

- 5 NOTE 3 The unsequenced property does not necessarily imply that the function is reentrant or that calls can be executed concurrently. This is because an unsequenced function can read from and write to objects of static storage duration, as long as no change is observable after a call terminates.
- 6 EXAMPLE 1 The attribute in the following function declaration asserts that it doesn't depend on any modifiable state of the abstract machine. Calls to the function can be executed out of sequence before the return value is needed and two calls to the function with the same argument value will result in the same return value.

```
bool tendency(signed char) [[unsequenced]];
```

Therefore such a call for a given argument value needs only to be executed once and the returned value can be reused when appropriate. For example, calls for all possible argument values can be executed during program startup and tabulated.

- 7 EXAMPLE 2 The attribute in the following function declaration asserts that it doesn't depend on any modifiable state of the abstract machine. Within the same thread, calls to the function can be executed out of sequence before the return value is needed and two calls to the function will result in the same pointer return value. Therefore such a call needs only to be executed once in a given thread and the returned pointer value can be reused when appropriate. For example, a single call can be executed during thread startup and the return value `p` and the value of the object `*p` of type `toto` `const` can be cached.

```
typedef struct toto toto;
toto const* toto_zero(void) [[unsequenced]];
```

- 8 EXAMPLE 3 The unsequenced property of a function *f* can be locally asserted within a function *g* that uses it. For example the library function `sqrt` is in general not unsequenced because a negative argument will raise a domain error and because the result can depend on the rounding mode. Nevertheless in contexts similar to the following function a user can prove that it will not be called with invalid arguments, and, that the floating-point environment has the same value for all calls.

```
#include <math.h>
#include <fenv.h>

inline double distance (double const x[static 2]) [[reproducible]] {
 #pragma STDC FP_CONTRACT OFF
 #pragma STDC FENV_ROUND FE_TONEAREST
 // We assert that sqrt will not be called with invalid arguments
 // and the result only depends on the argument value.
 extern typeof(sqrt) [[unsequenced]] sqrt;
 return sqrt(x[0]*x[0] + x[1]*x[1]);
}
```

The function `distance` potentially has the side effect of changing the floating-point environment. Nevertheless the floating environment is thread local, thus a change to that state outside the function is sequenced with the change within and additionally the observed value is restored when the function returns. Thus this side effect is not observable for a caller. Overall the function `distance` is stateless, effectless and idempotent and in particular it is reproducible as the attribute indicates. Because the function can be called in a context where the floating-point environment has different state, `distance` is not independent and thus it is also not unsequenced. Nevertheless, adding an unsequenced attribute where this is justified can introduce optimization opportunities.

```
double g (double y[static 1], double const x[static 2]) {
 // We assert that distance will not see different states of the floating
 // point environment.
 extern double distance (double const x[static 2]) [[unsequenced]];
 y[0] = distance(x);
 ...
}
```

```
} return distance(x); // replacement by y[0] is valid
```

## 6.8 Statements and blocks

### 6.8.1 General

#### Syntax

- 1 *statement*:
- labeled-statement*  
*unlabeled-statement*
- unlabeled-statement*:
- expression-statement*  
*attribute-specifier-sequence*<sub>opt</sub> *primary-block*  
*attribute-specifier-sequence*<sub>opt</sub> *jump-statement*
- primary-block*:
- compound-statement*  
*selection-statement*  
*iteration-statement*
- secondary-block*:
- statement*

#### Semantics

- 2 A *statement* specifies an action to be performed. Except as indicated, statements are executed in sequence. The optional attribute specifier sequence appertains to the respective statement.
- 3 A *block* is either a primary block, a secondary block, or the block associated with a function definition; it allows a set of declarations and statements to be grouped into one syntactic unit. Whenever a block *B* appears in the syntax production as part of the definition of an enclosing block *A*, scopes of identifiers and lifetimes of objects that are associated with *B* do not extend to the parts of *A* that are outside of *B*. The initializers of objects that have automatic storage duration, and any size expressions and typeof operators in declarations of ordinary identifiers with block scope, are evaluated and the values are stored in the objects (the representation of objects without an initializer becomes indeterminate) each time the declaration is reached in the order of execution, as if it were a statement, and within each declaration in the order that declarators appear.
- 4 A *full expression* is an expression that is not part of another expression, nor part of a declarator or abstract declarator. There is also an implicit full expression in which the non-constant size expressions for a variably modified type are evaluated; within that full expression, the evaluation of different size expressions are unsequenced with respect to one another. There is a sequence point between the evaluation of a full expression and the evaluation of the next full expression to be evaluated.
- 5 NOTE Each of the following is a full expression:
- a full declarator for a variably modified type,
  - an initializer that is not part of a compound literal,
  - the expression in an expression statement,
  - the controlling expression of a selection statement (**if** or **switch**),
  - the controlling expression of a **while** or **do** statement,
  - each of the (optional) expressions of a **for** statement,
  - the (optional) expression in a **return** statement.

While a constant expression satisfies the definition of a full expression, evaluating it does not depend on nor produce any side effects, so the sequencing implications of being a full expression are not relevant to a constant expression.

**Forward references:** expression and null statements (6.8.4), selection statements (6.8.5), iteration statements (6.8.6), the **return** statement (6.8.7.5).



## 6.8.2 Labeled statements

### Syntax

- 1 *label*:
- attribute-specifier-sequence*<sub>opt</sub> *identifier* :
- attribute-specifier-sequence*<sub>opt</sub> **case** *constant-range-expression* :
- attribute-specifier-sequence*<sub>opt</sub> **case** *constant-expression* :
- attribute-specifier-sequence*<sub>opt</sub> **default** :
- labeled-statement*:
- label* *statement*

### Constraints

- 2 A **case** or **default** label shall appear only in a **switch** statement. Further constraints on such labels are discussed under the **switch** statement.
- 3 Label names shall be unique within a function.

### Semantics

- 4 Any statement or declaration in a compound statement can be preceded by a prefix that declares an identifier as a label name. The optional attribute specifier sequence appertains to the label. Labels in themselves do not alter the flow of control, which continues unimpeded across them.
- 5 A **case** label specifies the value of the following constant expression, if followed by a single constant expression, or each value in the sequence described by the following range, if followed by a constant range expression.
- 6 If a statement following a label is an iteration or **switch** statement, that statement is *named* by the label *label*. If a statement following a label is a labeled statement, the statement named by this label is the same statement named within the nested labeled statement, if any.<sup>186)</sup>
- 7 EXAMPLE Only an iteration or **switch** statement may be named by a label, by appearing as its immediate syntactic operand:

```
loop:
for (int i = 0; i < IK; ++ i) { // this for is named by "loop:"
 /* ... */
}

braces:
{
 // no statement is named by "braces:"
 for (int i = 0; i < IK; ++ i) {
 /* ... */
 }
}
```

**Forward references:** the **goto** statement (6.8.7.2), the **break** statement (6.8.7.4), the **switch** statement (6.8.5.3).

## 6.8.3 Compound statement

### Syntax

- 1 *compound-statement*:
- { *block-item-list*<sub>opt</sub> }
- block-item-list*:
- block-item*
- block-item-list* *block-item*
- block-item*:
- declaration*

<sup>186)</sup>A statement may therefore be named by more than one label.

*unlabeled-statement*  
*label*

### Semantics

- 2 A *compound statement* that is a function body together with the parameter type list and the optional attribute specifier sequence between them forms the block associated with the function definition in which it appears. Otherwise, it is a block that is different from any other block. A label that is not followed by another label or an unlabeled statement shall be translated as if it were followed by a null statement.
- 3 A statement that follows one or more labels without other intervening block items is named by those labels, if it would be eligible to be named by the labels of a labeled statement (6.8.2).

## 6.8.4 Expression and null statements

### Syntax

- 1 *expression-statement*:  
                   *expression*<sub>opt</sub> ;  
                   *attribute-specifier-sequence expression* ;

### Semantics

- 2 The attribute specifier sequence appertains to the expression. The expression in an expression statement is evaluated as a void expression for its side effects.<sup>187)</sup>
- 3 A *null statement* (consisting of just a semicolon) performs no operations.
- 4 EXAMPLE 1 If a function call is evaluated as an expression statement for its side effects only, the discarding of its value can be made explicit by converting the expression to a void expression by means of a cast:

```
int p(int);
/* ... */
(void)p(0);
```

- 5 EXAMPLE 2 In the program fragment

```
char *s;
/* ... */
while (*s++ != '\0')
 ;
```

a null statement is used to supply an empty loop body to the iteration statement.

**Forward references:** iteration statements (6.8.6).

## 6.8.5 Selection statements

### 6.8.5.1 General

#### Syntax

- 1 *selection-statement*:  
                   **if** ( *selection-header* ) *secondary-block*  
                   **if** ( *selection-header* ) *secondary-block* **else** *secondary-block*  
                   **switch** ( *selection-header* ) *secondary-block*

<sup>187)</sup>Such as assignments, and function calls which have side effects.

*selection-header:*

*expression*  
*declaration expression*  
*simple-declaration*

### Semantics

- 2 A selection statement selects among a set of secondary blocks depending on the value of a controlling expression.
- 3 If the selection header is the first or second form, the controlling expression is the expression of the selection header. Otherwise, there is no explicit expression and the controlling expression is the value of the single declared object after initialization.
- 4 If the selection header includes a declaration (the second or third forms), the scope of any identifiers it declares is the remainder of the declaration and the entire selection statement, including the subsequent expression in the second form as well as the secondary block or blocks. In the second form, the declaration is reached in the order of execution before evaluation of the expression.

#### 6.8.5.2 The **if** statement

##### Constraints

- 1 The controlling expression of an **if** statement shall have scalar type.

##### Semantics

- 2 In both forms, the first secondary block is executed if the controlling expression<sup>188)</sup> compares unequal to 0. In the **else** form, the second secondary block is executed if the expression compares equal to 0. If the first secondary block is reached via a label, the second secondary block is not executed.
- 3 An **else** is associated with the lexically nearest preceding **if** that is allowed by the syntax.
- 4 EXAMPLE 1 The controlling expression of an **if** statement that uses a declaration as its selection header is implicitly the value of the declaration:

```

if (int x = get ()) {
 // x is non-0 here
}
else {
 // x is 0 here
}

```

is equivalent to

```

if (int x = get (); x) {
 // x is non-0 here
}
else {
 // x is 0 here
}

```

and therefore to

```

{
 int x = get ();
 if (x) {
 // x is non-0 here
 } else {
 // x is 0 here
 }
}

```

<sup>188)</sup>The controlling expression is implicitly the initialized value of the declared object, if the selection header only has a declaration. The implementation is permitted to re-read the object to determine the value, but is not required to do so, even when the object has **volatile**-qualified type.

```
 }
}
```

- 5 EXAMPLE 2 The controlling expression of any **if** statement is always implicitly compared to 0 by the statement itself:

```
double x = DBL_SNAN;
if (x) {
 // fetestexcept (FE_INVALID) is nonzero because of the comparison
}
```

### 6.8.5.3 The switch statement

#### Constraints

- 1 The controlling expression of a **switch** statement shall have integer type.
- 2 If a **switch** statement has an associated **case** or **default** label within the scope of an identifier with a variably modified type, the entire secondary block of the **switch** statement shall be within the scope of that identifier.<sup>189)</sup>
- 3 The expression of each **case** label shall be an integer constant expression or a constant range expression and no two of the values of **case** constant expressions or constant range expressions associated to the same **switch** statement shall specify the same value after conversion. There can be at most one **default** label associated to a **switch** statement. (Any enclosed **switch** statement can have a **default** label or **case**-specified values that duplicate **case**-specified values in the enclosing **switch** statement.)
- 4 The arithmetic values specified by a constant range expression shall not change as a result of conversion to the promoted type of the controlling expression.

#### Semantics

- 5 A **switch** statement causes control to jump to, into, or past the secondary block that is the *switch body*, depending on the value of a controlling expression, and on the presence of a **default** label and the specified values of any **case** labels on or in the switch body. A **case** or **default** label is accessible only within the closest enclosing **switch** statement.
- 6 The integer promotions are performed on the controlling expression. The values specified by each **case** label are converted to the promoted type of the controlling expression. If a converted value matches that of the promoted controlling expression, control jumps to the statement or declaration following the matched **case** label. Otherwise, if there is a **default** label, control jumps to the statement or declaration following the **default** label. If no converted **case** value matches and there is no **default** label, no part of the switch body is executed.
- 7 A **case** label with a range expression that describes an empty range (which occurs when the first value is greater than the second value) never matches any value of the controlling expression.<sup>190)</sup>

#### Implementation limits

- 8 As discussed in 5.3.5.2, the implementation may limit the number of **case** labels in a **switch** statement.
- 9 EXAMPLE 1 In the artificial program fragment

```
switch (expr)
{
 int i = 4;
 f(i);
case 0:
 i = 17;
 /* falls through into default code */
}
```

<sup>189)</sup>That is, the declaration: precedes the **switch** statement; appears in the selection header; or, it follows the last **case** or **default** label associated with the **switch** that is in the block containing the declaration.

<sup>190)</sup>Even if the controlling expression has the same value as one of the operands to the empty range expression.

```

default:
 printf("%d\n", i);
}

```

the object whose identifier is `i` exists with automatic storage duration (within the block) but is never initialized, and thus if the controlling expression has a nonzero value, the call to the `printf` function will access an object with an indeterminate representation. Similarly, the call to the function `f` cannot be reached.

- 10 EXAMPLE 2 A value specified as part of the sequence in a range expression argument to **case** shall not be repeated by a different label:

```

extern int n;
switch (n) {
 case 1 ... 5:
 /* ... */
 break;
 case 6: // OK, 6 has not been specified
 /* ... */
 case 4: // constraint violation:
 /* ... */ // 4 was specified by 1 ... 5
 break;
 case 3 ... 8: // constraint violation:
 /* ... */ // 3, 4, 5, 6 were specified above
 break;
}

```

- 11 EXAMPLE 3 A constant range expression may express one or no values.

```

extern int n;
switch (n) {
 case 1 ... 1: // OK, same as case 1:
 /* ... */
 break;
 case 8 ... 7: // OK, matches nothing,
 // but diagnostic encouraged
 /* ... */
 break;
}

```

## 6.8.6 Iteration statements

### 6.8.6.1 General

#### Syntax

- 1 *iteration-statement*:
- ```

while ( expression ) secondary-block
do secondary-block while ( expression ) ;
for ( expressionopt ; expressionopt ; expressionopt ) secondary-block
for ( declaration expressionopt ; expressionopt ) secondary-block

```

Constraints

- 2 The controlling expression of an iteration statement shall have scalar type.

Semantics

- 3 An iteration statement causes a secondary block called the *loop body* to be executed repeatedly until the controlling expression compares equal to 0. The repetition occurs regardless of whether the loop body is entered from the iteration statement or by a jump.¹⁹¹⁾

¹⁹¹⁾Code jumped over is not executed. In particular, the controlling expression of a **for** or **while** statement is not evaluated before entering the loop body, nor is *clause-1* (6.8.6.4) of a **for** statement.

- 4 An iteration statement may be assumed by the implementation to terminate if its controlling expression is not a constant expression,¹⁹²⁾ and none of the following operations are performed in its body, controlling expression or (in the case of a **for** statement) its *expression-3*:¹⁹³⁾

- input/output operations
- accessing a volatile object
- synchronization or atomic operations.

6.8.6.2 The **while** statement

- 1 The evaluation of the controlling expression takes place before each execution of the loop body.

6.8.6.3 The **do** statement

- 1 The evaluation of the controlling expression takes place after each execution of the loop body.

6.8.6.4 The **for** statement

- 1 The statement

```
for (clause-1; expression-2; expression-3) secondary-block
```

behaves as follows: The expression *expression-2* is the controlling expression that is evaluated before each execution of the loop body. The expression *expression-3* is evaluated as a void expression after each execution of the loop body. If *clause-1* is a declaration, the scope of any identifiers it declares is the remainder of the declaration and the entire loop, including the other two expressions; it is reached in the order of execution before the first evaluation of the controlling expression. If *clause-1* is an expression, it is evaluated as a void expression before the first evaluation of the controlling expression.¹⁹⁴⁾

- 2 Both *clause-1* and *expression-3* can be omitted. An omitted *expression-2* is replaced by a nonzero constant.

6.8.7 Jump statements

6.8.7.1 General

Syntax

- 1 *jump-statement*:
- ```
goto identifier ;
continue identifieropt ;
break identifieropt ;
return expressionopt ;
```

##### Semantics

- 2 A jump statement causes an unconditional jump to another place.

#### 6.8.7.2 The **goto** statement

##### Constraints

- 1 The identifier in a **goto** statement shall name a label located somewhere in the enclosing function. A **goto** statement shall not jump from outside the scope of an identifier having a variably modified type to inside the scope of that identifier.

<sup>192)</sup>An omitted controlling expression is replaced by a nonzero constant, which is a constant expression.

<sup>193)</sup>This is intended to allow compiler transformations such as removal of empty loops even when termination cannot be proven.

<sup>194)</sup>Thus, *clause-1* specifies initialization for the loop, possibly declaring one or more variables for use in the loop; the controlling expression, *expression-2*, specifies an evaluation made before each iteration, such that execution of the loop continues until the expression compares equal to 0; and *expression-3* specifies an operation (such as incrementing) that is performed after each iteration.

**Semantics**

- 2 A **goto** statement causes an unconditional jump to the statement prefixed by the named label in the enclosing function.
- 3 EXAMPLE 1 It is sometimes convenient to jump into the middle of a complicated set of statements. The following outline presents one possible approach to a problem based on these three assumptions:
  1. The general initialization code accesses objects only visible to the current function.
  2. The general initialization code is too large to warrant duplication.
  3. The code to determine the next operation is at the head of the loop. (To allow it to be reached by **continue** statements, for example.)

```

/* ... */
goto first_time;
for (;;) {
 // determine next operation
 /* ... */
 if (need to reinitialize) {
 // reinitialize-only code
 /* ... */
 first_time:
 // general initialization code
 /* ... */
 continue;
 }
 // handle other operations
 /* ... */
}

```

- 4 EXAMPLE 2 A **goto** statement which jumps past any declarations of objects with variably modified types is not conforming. A jump within the scope, however, is valid.

```

goto lab3; // invalid: going INTO scope of VLA.
{
 double a[n];
 a[j] = 4.4;
lab3:
 a[j] = 3.3;
 goto lab4; // valid: going WITHIN scope of VLA.
 a[j] = 5.5;
lab4:
 a[j] = 6.6;
}
goto lab4; // invalid: going INTO scope of VLA.

```

**6.8.7.3 The continue statement****Constraints**

- 1 A **continue** statement shall appear only in or as a loop body.
- 2 A **continue** statement with an identifier operand shall appear within a statement named by the label with the corresponding identifier.

**Semantics**

- 3 A **continue** statement causes a jump to the loop-continuation portion of an enclosing iteration statement; that is, to the end of the loop body. More precisely, in each of the statements:

```
while (/* ... */) {
 /* ... */
 continue;
 /* ... */
contin:
}
```

```
do {
 /* ... */
 continue;
 /* ... */
contin:;
} while (/* ... */);
```

```
for (/* ... */) {
 /* ... */
 continue;
 /* ... */
contin:
}
```

unless the **continue** statement shown is in an enclosed iteration statement (in which case it is interpreted within that statement), it is equivalent to **goto contin;**<sup>195)</sup>

- 4 If the **continue** statement has an identifier operand, the jump is to the loop-continuation of the iteration statement named by the label with the corresponding identifier. Otherwise, the jump is to the loop-continuation of the innermost enclosing iteration statement.
- 5 EXAMPLE 1 In the following code, **continue** only jumps to the loop-continuation of the inner nested loop, whereas **continue outer** jumps to the loop-continuation of the outermost loop:

```
outer:
for (int i = 0; i < 10; ++ i) {
 for (int j = 0; j < 20; ++ j) {
 continue; // jumps to CONT1
 continue outer; // jumps to CONT2
 } // CONT1
} // CONT2
}
```

```
int i = 0;
outer:
do {
 int j = 0;
 do {
 continue; // jumps to CONT1
 continue outer; // jumps to CONT2
 } while (j < 20); // CONT1
} while (i < 10); // CONT2
```

#### 6.8.7.4 The break statement

##### Constraints

- 1 A **break** statement shall appear only in or as a switch body or loop body.
- 2 A **break** statement with an identifier operand shall appear within a **switch** or iteration statement named by the label with the corresponding identifier.

##### Semantics

- 3 A **break** statement terminates execution of an enclosing **switch** or iteration statement, as if by jumping to a label immediately following it in the surrounding scope using **goto**.
- 4 If the **break** statement has an identifier operand, the jump exits the **switch** or iteration statement named by the label with the corresponding identifier. Otherwise, the jump exits the innermost enclosing **switch** or iteration statement.
- 5 EXAMPLE 1 Therefore, this:

```
while (1) {
```

<sup>195)</sup> Following the **contin:** label in the 2nd example is a null statement. The null statement in the first and third example is implied by the label (6.8.3).



```

 break;
 }
 // break jumps here

```

is identical to this:

```

while (1) {
 goto there;
}
there:

```

- 6 EXAMPLE 2 In the following code, **break** only exits the switch, whereas **break outer** exits the enclosing loop as well:

```

outer:
for (int i = 0; i < IK; ++ i) {
 switch (i) {
 case 1:
 break; // jumps to CONT1
 case 2:
 break outer; // jumps to CONT2
 }
 // CONT1
}
// CONT2

```

#### 6.8.7.5 The return statement

##### Constraints

- 1 A **return** statement with an expression shall not appear in a function whose return type is **void**. A **return** statement without an expression shall only appear in a function whose return type is **void**.

##### Semantics

- 2 A **return** statement terminates execution of the current function and returns control to its caller. A function can have any number of **return** statements.
- 3 If a **return** statement with an expression is executed, the value of the expression is returned to the caller as the value of the function call expression. If the expression has a type different from the return type of the function in which it appears, the value is converted as if by assignment to an object having the return type of the function.<sup>196)</sup>
- 4 EXAMPLE In:

```

struct s { double i; } f(void);
union {
 struct {
 int f1;
 struct s f2;
 } u1;
 struct {
 struct s f3;
 int f4;
 } u2;
} g;

struct s f(void)
{
 return g.u1.f2;
}

```

<sup>196)</sup>The **return** statement is not an assignment. The overlap restriction of 6.5.17.2 does not apply to the case of function return. The representation of floating-point values can have wider range or precision than implied by the type; a cast can be used to remove this extra range and precision.

```
}

/* ... */
g.u2.f3 = f();
```

there is no undefined behavior, although there would be if the assignment were done directly (without using a function call to fetch the value).

## 6.9 External definitions

### 6.9.1 General

#### Syntax

- 1 *translation-unit:*  
                     *external-declaration*  
                     *translation-unit external-declaration*
- external-declaration:*  
                     *function-definition*  
                     *declaration*

#### Constraints

- 2 The storage-class specifier **register** shall not appear in the declaration specifiers in an external declaration. The storage-class specifier **auto** shall only appear in the declaration specifiers in an external declaration if the type is inferred.
- 3 There shall be no more than one external definition for each identifier declared with internal linkage in a translation unit. Moreover, if an identifier declared with internal linkage is used in an expression there shall be exactly one external definition for the identifier in the translation unit, unless it is:
- part of the operand of a **sizeof** operator whose result is an integer constant expression;
  - part of the operand of a **\_Lengthof** operator whose result is an integer constant expression;
  - part of the operand of an **alignof** operator whose result is an integer constant expression;
  - part of the controlling expression of a generic selection;
  - part of the expression in a generic association that is not the result expression of its generic selection;
  - or, part of the operand of any typeof operator whose result is not a variably modified type.

#### Semantics

- 4 As discussed in 5.2.1.1, the unit of program text after preprocessing is a translation unit, which consists of a sequence of external declarations. These are described as “external” because they appear outside any function (and hence have file scope). As discussed in 6.7, a declaration that also causes storage to be reserved for an object or a function named by the identifier is a definition.
- 5 An *external definition* is an external declaration that is also a definition of a function (other than an inline definition) or an object. If an identifier declared with external linkage is used in an expression (other than as part of the operand of a typeof operator whose result is not a variably modified type, part of the controlling expression of a generic selection, part of the expression in a generic association that is not the result expression of its generic selection, or part of a **sizeof**, **\_Lengthof**, or **alignof** operator whose result is an integer constant expression), somewhere in the entire program there shall be exactly one external definition for the identifier; otherwise, there shall be no more than one.<sup>197)</sup>

### 6.9.2 Function definitions

#### Syntax

- 1 *function-definition:*  
                     *attribute-specifier-sequence*<sub>opt</sub> *declaration-specifiers declarator function-body*

<sup>197)</sup> Thus, if an identifier declared with external linkage is not used in an expression, there need be no external definition for it.

*function-body:*  
*compound-statement*

### Constraints

- 2 The identifier declared in a function definition (which is the name of the function) shall have a function type, as specified by the declarator portion of the function definition.
- 3 The return type of a function shall be **void** or a complete object type other than array type.
- 4 The storage-class specifier, if any, in the declaration specifiers shall be either **extern** or **static**.
- 5 If the parameter list consists of a single parameter of type **void**, the parameter declarator shall not include an identifier.
- 6 Variable length array types of unspecified size shall not be used as part of a parameter declaration in a function definition.

### Semantics

- 7 The optional attribute specifier sequence in a function definition appertains to the function.
- 8 The declarator in a function definition specifies the name of the function being defined and the types (and optionally the names) of all the parameters; the declarator also serves as a function prototype for later calls to the same function in the same translation unit. The type of each parameter is adjusted as described in 6.7.7.4.
- 9 If a function that accepts a variable number of arguments is defined without a parameter type list that ends with the ellipsis notation, the behavior is undefined.
- 10 The parameter type list, the attribute specifier sequence of the declarator that follows the parameter type list, and the compound statement of the function body form a single block.<sup>198)</sup> Each parameter has automatic storage duration; its identifier, if any,<sup>199)</sup> is an lvalue.<sup>200)</sup> The layout of the storage for parameters is unspecified.
- 11 On entry to the function, the size expressions of each variably modified parameter and type of operators used in declarations of parameters are evaluated and the value of each argument expression is converted to the type of the corresponding parameter as if by assignment. (Array expressions and function designators as arguments were converted to pointers before the call.)
- 12 After all parameters have been assigned, the compound statement of the function body is executed.
- 13 Unless otherwise specified, if the **}** that terminates the function body is reached, and the value of the function call is used by the caller, the behavior is undefined.
- 14 NOTE In a function definition, the return type of the function and its prototype cannot be inherited from a typedef:

```
typedef int F(void); // type F is "function with no parameters
 // returning int"
F f, g; // f and g both have type compatible with F
F f { /* ... */ } // WRONG: syntax/constraint error
F g() { /* ... */ } // WRONG: declares that g returns a function
int f(void) { /* ... */ } // RIGHT: f has type compatible with F
int g() { /* ... */ } // RIGHT: g has type compatible with F
F *e(void) { /* ... */ } // e returns a pointer to a function
F *((e))(void) { /* ... */ } // same: parentheses irrelevant
int (*fp)(void); // fp points to a function that has type F
F *Fp; // Fp points to a function that has type F
```

<sup>198)</sup>The visibility scope of a parameter in a function definition starts when its declaration is completed, extends to following parameter declarations, to possible attributes that follow the parameter type list, and then to the entire function body. The lifetime of each instance of a parameter starts when the declaration is evaluated starting a call and ends when that call terminates.

<sup>199)</sup>A parameter that has no declared name is inaccessible within the function body.

<sup>200)</sup>A parameter identifier cannot be redeclared in the function body except in an enclosed block.

- 15 EXAMPLE 1 In the following:

```
extern int max(int a, int b)
{
 return a > b ? a: b;
}
```

**extern** is the storage-class specifier and **int** is the type specifier; **max(int a, int b)** is the function declarator; and

```
{ return a > b ? a: b; }
```

is the function body.

- 16 EXAMPLE 2 To pass one function to another, one can say

```
int f(void);
/* ... */
g(f);
```

Then the definition of **g** can read

```
void g(int (*funcp)(void))
{
 /* ... */
 (*funcp)(); /* or funcp(); ...*/
}
```

or, equivalently,

```
void g(int func(void))
{
 /* ... */
 func(); /* or (*func)(); ...*/
}
```

### 6.9.3 External object definitions

#### Semantics

- 1 If the declaration of an identifier for an object has file scope and an initializer, or has file scope and storage-class specifier **thread\_local**, the declaration is an external definition for the identifier.
- 2 A declaration of an identifier for an object that has file scope without an initializer, and without the storage-class specifier **extern** or **thread\_local**, constitutes a *tentative definition*. If a translation unit contains one or more tentative definitions for an identifier, and the translation unit contains no external definition for that identifier, then the behavior is exactly as if the translation unit contains a file scope declaration of that identifier with an empty initializer and a type determined as follows:
  - if the composite type as of the end of the translation unit is an array of unknown size, then an array of size one with the composite element type;
  - otherwise, the composite type at the end of the translation unit.
- 3 If the declaration of an identifier for an object is a tentative definition and has internal linkage, the declared type shall not be an incomplete type.

## 4 EXAMPLE 1

```

int i1 = 1; // definition, external linkage
static int i2 = 2; // definition, internal linkage
extern int i3 = 3; // definition, external linkage
int i4; // tentative definition, external linkage
static int i5; // tentative definition, internal linkage

int i1; // valid tentative definition, refers to previous
int i2; // 6.2.2 renders undefined, linkage disagreement
int i3; // valid tentative definition, refers to previous
int i4; // valid tentative definition, refers to previous
int i5; // 6.2.2 renders undefined, linkage disagreement

extern int i1; // refers to previous, whose linkage is external
extern int i2; // refers to previous, whose linkage is internal
extern int i3; // refers to previous, whose linkage is external
extern int i4; // refers to previous, whose linkage is external
extern int i5; // refers to previous, whose linkage is internal

```

## 5 EXAMPLE 2 If at the end of the translation unit containing

```
int i[];
```

the array `i` still has incomplete type, the implicit initializer causes it to have one element, which is set to zero on program startup.

## 6.10 Preprocessing directives

### 6.10.1 General

#### Syntax

|   |                            |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|---|----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | <i>preprocessing-file:</i> | <i>group</i> <sub>opt</sub>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|   | <i>group:</i>              | <i>group-part</i><br><i>group group-part</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|   | <i>group-part:</i>         | <i>if-section</i><br><i>control-line</i><br><i>text-line</i><br><b># non-directive</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|   | <i>if-section:</i>         | <i>if-group elif-groups</i> <sub>opt</sub> <i>else-group</i> <sub>opt</sub> <i>endif-line</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|   | <i>if-group:</i>           | <b># if</b> <i>constant-expression new-line group</i> <sub>opt</sub><br><b># ifdef</b> <i>identifier new-line group</i> <sub>opt</sub><br><b># ifndef</b> <i>identifier new-line group</i> <sub>opt</sub>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|   | <i>elif-groups:</i>        | <i>elif-group</i><br><i>elif-groups elif-group</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
|   | <i>elif-group:</i>         | <b># elif</b> <i>constant-expression new-line group</i> <sub>opt</sub><br><b># elifdef</b> <i>identifier new-line group</i> <sub>opt</sub><br><b># elifndef</b> <i>identifier new-line group</i> <sub>opt</sub>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|   | <i>else-group:</i>         | <b># else</b> <i>new-line group</i> <sub>opt</sub>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
|   | <i>endif-line:</i>         | <b># endif</b> <i>new-line</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|   | <i>control-line:</i>       | <b># include</b> <i>pp-tokens new-line</i><br><b># embed</b> <i>pp-tokens new-line</i><br><b># define</b> <i>identifier replacement-list new-line</i><br><b># define</b> <i>identifier lparen identifier-list</i> <sub>opt</sub> <b>)</b> <i>replacement-list new-line</i><br><b># define</b> <i>identifier lparen ... ) replacement-list new-line</i><br><b># define</b> <i>identifier lparen identifier-list , ... ) replacement-list new-line</i><br><b># undef</b> <i>identifier new-line</i><br><b># line</b> <i>pp-tokens new-line</i><br><b># error</b> <i>pp-tokens</i> <sub>opt</sub> <i>new-line</i><br><b># warning</b> <i>pp-tokens</i> <sub>opt</sub> <i>new-line</i><br><b># pragma</b> <i>pp-tokens</i> <sub>opt</sub> <i>new-line</i><br><b># new-line</b> |
|   | <i>text-line:</i>          | <i>pp-tokens</i> <sub>opt</sub> <i>new-line</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|   | <i>non-directive:</i>      | <i>pp-tokens new-line</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|   | <i>lparen:</i>             | a ( character not immediately preceded by white space                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |

*replacement-list:*

*pp-tokens*<sub>opt</sub>

*pp-tokens:*

*preprocessing-token*

*pp-tokens preprocessing-token*

*new-line:*

the new-line character

*identifier-list:*

*identifier*

*identifier-list* , *identifier*

*pp-parameter:*

*pp-parameter-name pp-parameter-clause*<sub>opt</sub>

*pp-parameter-name:*

*pp-standard-parameter*

*pp-prefixed-parameter*

*pp-standard-parameter:*

*identifier*

*pp-prefixed-parameter:*

*identifier :: identifier*

*pp-parameter-clause:*

( *pp-balanced-token-sequence*<sub>opt</sub> )

*pp-balanced-token-sequence:*

*pp-balanced-token*

*pp-balanced-token-sequence pp-balanced-token*

*pp-balanced-token:*

( *pp-balanced-token-sequence*<sub>opt</sub> )

[ *pp-balanced-token-sequence*<sub>opt</sub> ]

{ *pp-balanced-token-sequence*<sub>opt</sub> }

any pp-token other than a parenthesis, a bracket, or a brace

*embed-parameter-sequence:*

*pp-parameter*

*embed-parameter-sequence pp-parameter*

## Description

- 2 A *preprocessing directive* consists of a sequence of preprocessing tokens that satisfies the following constraints: The first token in the sequence is a # preprocessing token that (at the start of translation phase 4) is either the first character in the source file (optionally after white space containing no new-line characters) or that follows white space containing at least one new-line character. The last



token in the sequence is the first new-line character that follows the first token in the sequence.<sup>201)</sup> A new-line character ends the preprocessing directive even if it occurs within what would otherwise be an invocation of a function-like macro.

- 3 A text line shall not begin with a **#** preprocessing token. A non-directive shall not begin with any of the directive names appearing in the syntax.
- 4 Some preprocessing directives take additional information using preprocessor parameters. A *preprocessing parameter* (pp-parameter) shall be either a *preprocessor prefixed parameter* (identified by a pp-prefixed-parameter, for implementation-defined preprocessor parameters) or a *preprocessor standard parameter* (identified with a pp-standard-parameter, for pp-parameters specified by this document).
- 5 In all aspects, a preprocessor standard parameter specified by this document as an identifier **pp\_param** and an identifier of the form **\_\_pp\_param\_\_** shall behave the same when used as a preprocessor parameter, except for the spelling.
- 6 EXAMPLE 1 Thus, the preprocessor parameters on the two binary resource inclusion directives (6.10.4):

```
#embed "boop.h" limit(5)
#embed "boop.h" __limit__(5)
```

behave the same, and can be freely interchanged. Implementations are encouraged to behave similarly for preprocessor parameters (including preprocessor prefixed parameters) they provide.

- 7 When in a group that is skipped (6.10.2), the directive syntax is relaxed to allow any sequence of preprocessing tokens to occur between the directive name and the following new-line character.

### Constraints

- 8 The only white-space characters that shall appear between preprocessing tokens within a preprocessing directive (from just after the introducing **#** preprocessing token through just before the terminating new-line character) are space and horizontal-tab (including spaces that have replaced comments or possibly other white-space characters in translation phase 3).
- 9 A preprocessor parameter shall be either a preprocessor standard parameter, or an implementation-defined preprocessor prefixed parameter.<sup>202)</sup>

### Semantics

- 10 The implementation can process and skip sections of source files conditionally, include other source files, and replace macros. These capabilities are called *preprocessing*, because conceptually they occur before translation of the resulting translation unit.
- 11 The preprocessing tokens within a preprocessing directive are not subject to macro expansion unless otherwise stated.
- 12 EXAMPLE 2 In:

```
#define EMPTY
EMPTY # include <file.h>
```

the sequence of preprocessing tokens on the second line is *not* a preprocessing directive, because it does not begin with a **#** at the start of translation phase 4 (5.2.1.2), even though it will do so after the macro **EMPTY** has been replaced.

- 13 The execution of a non-directive preprocessing directive results in undefined behavior.

<sup>201)</sup>Thus, preprocessing directives are commonly called “lines”. These “lines” have no other syntactic significance, as all white space is equivalent except in certain situations during preprocessing (see the **#** ordinary string literal creation operator in 6.10.5.3, for example).

<sup>202)</sup>An unrecognized preprocessor prefixed parameter is a constraint violation, except within `has_embed` expressions (6.10.2).

## 6.10.2 Conditional inclusion

### Syntax

- 1 *defined-macro-expression*:
 

**defined** *identifier*  
**defined** ( *identifier* )
- h-preprocessing-token*:
 

any *preprocessing-token* other than >
- h-pp-tokens*:
 

*h-preprocessing-token*  
*h-pp-tokens* *h-preprocessing-token*
- header-name-tokens*:
 

*string-literal*  
 < *h-pp-tokens* >
- has-include-expression*:
 

**\_\_has\_include** ( *header-name* )  
**\_\_has\_include** ( *header-name-tokens* )
- has-embed-expression*:
 

**\_\_has\_embed** ( *header-name* *embed-parameter-sequence*<sub>opt</sub> )  
**\_\_has\_embed** ( *header-name-tokens* *pp-balanced-token-sequence*<sub>opt</sub> )
- has-c-attribute-express*:
 

**\_\_has\_c\_attribute** ( *pp-tokens* )

- 2 The **#if** and **#elif** directives are collectively known as the *conditional expression inclusion preprocessing directives*. The conditional expression inclusion preprocessing directives, **#ifdef**, **#ifndef**, **#elifdef**, and **#elifndef** directives are collectively known as the *conditional inclusion preprocessing directives*.

### Constraints

- 3 The expression that controls conditional inclusion shall be an integer constant expression except that: identifiers (including those lexically identical to keywords) are interpreted as described subsequently in this subclause<sup>203)</sup> and it can contain zero or more defined macro expressions, **has\_include** expressions, **has\_embed** expressions, and/or **has\_c\_attribute** expressions as unary operator expressions.
- 4 A defined macro expression evaluates to 1 if the identifier is currently defined as a macro name (that is, if it is predefined or if it has been the subject of a **#define** preprocessing directive without an intervening **#undef** directive with the same subject identifier), 0 if it is not.
- 5 The second form of the **has\_include** expression and **has\_embed** expression is considered only if the first form does not match, in which case the preprocessing tokens are processed just as in normal text.
- 6 The header or source file identified by the parenthesized preprocessing token sequence in each contained **has\_include** expression is searched for as if that preprocessing token were the pp-tokens in a **#include** directive, except that no further macro expansion is performed. Such a directive shall satisfy the syntactic requirements of a **#include** directive. The **has\_include** expression evaluates to 1 if the search for the source file succeeds, and to 0 if the search fails.
- 7 The resource (6.10.4) identified by the header-name preprocessing token sequence in each contained **has\_embed** expression is searched for as if those preprocessing token were the pp-tokens in a **#embed** directive, except that no further macro expansion is performed. Such a directive shall satisfy the syntactic requirements of a **#embed** directive. The **has\_embed** expression evaluates to the same value as the following predefined macros (6.10.10.2):

— **\_\_STDC\_EMBED\_NOT\_FOUND\_\_**, if the search fails or if any of the embed parameters in the embed parameter sequence specified are not supported by the implementation for the **#embed**

<sup>203)</sup> Because the controlling constant expression is evaluated during translation phase 4, all identifiers either are or are not macro names — there simply are no keywords, enumeration constants, etc.

directive; or,

- **\_\_STDC\_EMBED\_FOUND\_\_**, if the search for the resource succeeds and all embed parameters in the embed parameter sequence specified are supported by the implementation for the **#embed** directive and the resource is not empty; or,
- **\_\_STDC\_EMBED\_EMPTY\_\_**, if the search for the resource succeeds and all embed parameters in the embed parameter sequence specified are supported by the implementation for the **#embed** directive and the resource is empty.

- 8 NOTE 1 Unrecognized preprocessor prefixed parameters in `has_embed` expressions are not a constraint violation and instead cause the expression to be evaluated to 0, as specified previously.
- 9 Each `has_c_attribute` expression is replaced by a nonzero pp-number matching the form of an integer literal if the implementation supports an attribute with the name specified by interpreting the pp-tokens as an attribute token, and by 0 otherwise. The pp-tokens shall match the form of an attribute token.
- 10 Each preprocessing token that remains (in the list of preprocessing tokens that will become the controlling expression) after all macro replacements have occurred shall be in the lexical form of a token (6.4).

### Semantics

- 11 The **#ifdef**, **#ifndef**, **#elifdef**, and **#elifndef** directives, and the **defined** conditional inclusion operator, shall treat **\_\_has\_include**, **\_\_has\_embed** and **\_\_has\_c\_attribute** as if they were the name of defined macros. The identifiers **\_\_has\_include**, **\_\_has\_embed**, and **\_\_has\_c\_attribute** shall not appear in any context not mentioned in this subclause.
- 12 Preprocessing directives of the forms

```
if constant-expression new-line groupopt
elif constant-expression new-line groupopt
```

check whether the controlling constant expression evaluates to nonzero.

- 13 Prior to evaluation, macro invocations in the list of preprocessing tokens that will become the controlling constant expression are replaced (except for those macro names modified by the **defined** unary operator), just as in normal text. If the token **defined** is generated as a result of this replacement process or use of the **defined** unary operator does not match one of the two specified forms prior to macro replacement, the behavior is undefined. After all replacements due to macro expansion and evaluations of defined macro expressions, `has_include` expressions, `has_embed` expressions, and `has_c_attribute` expressions have been performed, all remaining identifiers other than **true** (including those lexically identical to keywords such as **false**) are replaced with the pp-number 0, **true** is replaced with pp-number 1, and then each preprocessing token is converted into a token. The resulting tokens compose the controlling constant expression which is evaluated according to the rules of 6.6.1. For the purposes of this token conversion and evaluation, all signed integer types and all unsigned integer types act as if they have the same representation as, respectively, the types `intmax_t` and `uintmax_t` defined in the header `<stdint.h>`. This includes interpreting character literals, which can involve converting escape sequences into execution character set members. Whether the numeric value for these character literals matches the value obtained when an identical character literal occurs in an expression (other than within a **#if** or **#elif** directive) is implementation-defined. Whether a single-character character literal may have a negative value is implementation-defined.
- 14 NOTE 2 On an implementation where **INT\_MAX** is 0x7FFF and **UINT\_MAX** is 0xFFFF, the literal 0x8000 is signed and positive within a **#if** expression even though it would be unsigned in translation phase 7 (5.2.1.2).
- 15 NOTE 3 The constant expression in the following **#if** directive and **if** statement is not guaranteed to evaluate to the same value in these two contexts.

```
#if 'z' - 'a' == 25
if ('z' - 'a' == 25)
```

## 16 Preprocessing directives of the forms

```
ifdef identifier new-line groupopt
ifndef identifier new-line groupopt
elifdef identifier new-line groupopt
elifndef identifier new-line groupopt
```

check whether the identifier is or is not currently defined as a macro name. Their conditions are equivalent to **#if defined** *identifier*, **#if !defined** *identifier*, **#elif defined** *identifier*, and **#elif !defined** *identifier* respectively.

- 17 Each directive's condition is checked in order. If it evaluates to false (zero), the group that it controls is skipped: directives are processed only through the name that determines the directive to keep track of the level of nested conditionals; the rest of the directives' preprocessing tokens are ignored, as are the other preprocessing tokens in the group. Only the first group whose control condition evaluates to true (nonzero) is processed; any following groups are skipped and their controlling directives are processed as if they were in a group that is skipped. If none of the conditions evaluates to true, and there is a **#else** directive, the group controlled by the **#else** is processed; lacking a **#else** directive, all the groups until the **#endif** are skipped.<sup>204)</sup>

- 18 EXAMPLE 1 This demonstrates a way to include a header file only if it is available.

```
#if __has_include(<optional.h>)
include <optional.h>
define have_optional 1
#elif __has_include(<experimental/optional.h>)
include <experimental/optional.h>
define have_optional 1
define have_experimental_optional 1
#endif
#ifndef have_optional
define have_optional 0
#endif
```

## 19 EXAMPLE 2

```
/* Fallback for compilers not yet implementing this feature. */
#ifndef __has_c_attribute
#define __has_c_attribute(x) 0
#endif /* __has_c_attribute */

#if __has_c_attribute(fallthrough)
/* Standard attribute is available, use it. */
#define FALLTHROUGH [[fallthrough]]
#elif __has_c_attribute(vendor::fallthrough)
/* Vendor attribute is available, use it. */
#define FALLTHROUGH [[vendor::fallthrough]]
#else
/* Fallback implementation. */
#define FALLTHROUGH
#endif
```

## 20 EXAMPLE 3

```
#ifdef __STDC__
#define TITLE "ISO C Compilation"
#elifndef __cplusplus
#define TITLE "Non-ISO C Compilation"
#else
```

<sup>204)</sup>As indicated by the syntax, no preprocessing tokens are allowed to follow a **#else** or **#endif** directive before the terminating new-line character. However, comments can appear anywhere in a source file, including within a preprocessing directive.

```

/* C++ */
#define TITLE "C++ Compilation"
#endif

```

- 21 EXAMPLE 4 A combination of `__FILE__` (6.10.10.2) and `__has_embed` can be used to check for support of specific implementation extensions for the `#embed` (6.10.4) directive's parameters.

```

#if __has_embed(__FILE__ ext::token(0xB055))
#define DESCRIPTION "Supports extended token embed parameter"
#else
#define DESCRIPTION "Does not support extended token embed parameter"
#endif

```

- 22 EXAMPLE 5 The following snippet uses `__has_embed` to check for support of a specific implementation-defined embed parameter, and otherwise uses standard behavior to produce the same effect.

```

void parse_into_s(short* ptr, unsigned char* ptr_bytes, unsigned long long size);

int main () {
 #if __has_embed ("bits.bin" ds9000::element_type(short))
 /* Implementation extension: create short integers from the */
 /* translation environment resource into */
 /* a sequence of integer literals */
 short meow[] = {
 #embed "bits.bin" ds9000::element_type(short)
 };
 #elif __has_embed ("bits.bin")
 /* no support for implementation-specific */
 /* ds9000::element_type(short) parameter */
 const unsigned char meow_bytes[] = {
 #embed "bits.bin"
 };
 short meow[sizeof(meow_bytes) / sizeof(short)] = {};
 /* parse meow_bytes into short values by-hand! */
 parse_into_s(meow, meow_bytes, sizeof(meow_bytes));
 #else
 #error "cannot find bits.bin resource"
 #endif
 return (int)(meow[0] + meow[_Lengthof(meow) - 1]);
}

```

- 23 EXAMPLE 6 If the search for the resource is successful, this resource is always considered empty due to the `limit(0)` embed parameter, including in `__has_embed` expressions.

```

int main () {
 #if __has_embed(<infinite-resource> limit(0)) == 2
 // if <infinite-resource> exists, this
 // token sequence is always taken.
 return 0;
 #else
 // the 'infinite-resource' resource does not exist
 #error "The resource does not exist"
 #endif
}

```

**Forward references:** macro replacement (6.10.5), source file inclusion (6.10.3), mandatory macros (6.10.10.2), largest integer types (7.22.2.6).

### 6.10.3 Source file inclusion

#### Constraints

- 1 A **#include** directive shall identify a header or source file that can be processed by the implementation.

#### Semantics

- 2 A preprocessing directive of the form

**# include** < *h-char-sequence* > *new-line*

searches a sequence of implementation-defined places for a header identified uniquely by the specified sequence between the < and > delimiters, and causes the replacement of that directive by the entire contents of the header. How the places are specified or the header identified is implementation-defined.

- 3 A preprocessing directive of the form

**# include** " *q-char-sequence* " *new-line*

causes the replacement of that directive by the entire contents of the source file identified by the specified sequence between the " delimiters. The named source file is searched for in an implementation-defined manner. If this search is not supported, or if the search fails, the directive is reprocessed as if it read

**# include** < *h-char-sequence* > *new-line*

with the identical contained sequence (including > characters, if any) from the original directive.

- 4 A preprocessing directive of the form

**# include** *pp-tokens* *new-line*

(that does not match one of the two previous forms) is permitted. The preprocessing tokens after **#include** in the directive are processed just as in normal text. (Each identifier currently defined as a macro name is replaced by its replacement list of preprocessing tokens.) The directive resulting after all replacements shall match one of the two previous forms.<sup>205)</sup> The method by which a sequence of preprocessing tokens between a < and a > preprocessing token pair or a pair of " characters is combined into a single header name preprocessing token is implementation-defined.

- 5 The implementation shall provide unique mappings for sequences consisting of one or more nondigits or digits (6.4.3.1) followed by a period (.) and a single nondigit. The first character shall not be a digit. The implementation can ignore distinctions of alphabetical case and restrict the mapping to eight significant characters before the period.
- 6 A **#include** preprocessing directive may appear in a source file that has been read because of a **#include** directive in another file, up to an implementation-defined nesting limit (see 5.3.5.2).
- 7 EXAMPLE 1 The most common uses of **#include** preprocessing directives are as in the following:

```
#include <stdio.h>
#include "myprog.h"
```

- 8 EXAMPLE 2 This illustrates macro-replaced **#include** directives:

```
#if VERSION == 1
 #define INCFILE "vers1.h"
#elif VERSION == 2
 #define INCFILE "vers2.h" // and so on
#else
 #define INCFILE "versN.h"
```

<sup>205)</sup> Adjacent string literals are not concatenated into a single string literal (see the translation phases in 5.2.1.2); thus, an expansion that results in two string literals is an invalid directive.

```
#endif
#include INCFILE
```

**Forward references:** macro replacement (6.10.5).

## 6.10.4 Binary resource inclusion

### 6.10.4.1 #embed preprocessing directive

#### Description

- 1 A *resource* is a source of data accessible from the translation environment. An *embed parameter* is a single preprocessor parameter in the embed parameter sequence. It has an *implementation resource width*, which is the implementation-defined size in bits of the located resource. It also has a *resource width*, which is either:
  - the number of bits as computed from the optionally-provided **limit** embed parameter (6.10.4.2), if present; or,
  - the implementation resource width.

- 2 An *embed parameter sequence* is a whitespace-delimited list of preprocessor parameters which can modify the result of the replacement for the **#embed** preprocessing directive.

#### Constraints

- 3 An **#embed** directive shall identify a resource that can be processed by the implementation as a binary data sequence given the provided embed parameters.
- 4 Embed parameters not specified in this document shall be implementation-defined. Implementation-defined embed parameters may change the subsequently-defined semantics of the directive; otherwise, **#embed** directives which do not contain implementation-defined embed parameters shall behave as described in this document.
- 5 A resource is considered empty when its resource width is zero.
- 6 Let *embed element width* be either:
  - an integer constant expression greater than zero determined by an implementation-defined embed parameter; or,
  - **CHAR\_BIT** (5.3.5.3.2).

The result of  $(resource\ width) \% (embed\ element\ width)$  shall be zero.<sup>206)</sup>

#### Semantics

- 7 The expansion of a **#embed** directive is a token sequence formed from the list of integer constant expressions described later in this subclause. The group of tokens for each integer constant expression in the list is separated in the token sequence from the group of tokens for the previous integer constant expression in the list by a comma. The sequence neither begins nor ends in a comma. If the list of integer constant expressions is empty, the token sequence is empty. The directive is replaced by its expansion and, with the presence of certain embed parameters, additional or replacement token sequences.
- 8 A preprocessing directive of the form

**# embed** < *h-char-sequence* > *embed-parameter-sequence*<sub>opt</sub> *new-line*

searches a sequence of implementation-defined places for a resource identified uniquely by the specified sequence between the < and >. The search for the named resource is done in an implementation-defined manner.

<sup>206)</sup>This constraint helps ensure data is neither filled with padding values nor truncated in a given environment, and helps ensure the data is portable with respect to usages of **memcpy** (7.27.2.1) with character type arrays initialized from the data.



- 9 A preprocessing directive of the form

**# embed** " *q-char-sequence* " *embed-parameter-sequence*<sub>opt</sub> *new-line*

searches a sequence of implementation-defined places for a resource identified uniquely by the specified sequence between the " delimiters. The search for the named resource is done in an implementation-defined manner. If this search is not supported, or if the search fails, the directive is reprocessed as if it read

**# embed** < *h-char-sequence* > *embed-parameter-sequence*<sub>opt</sub> *new-line*

with the identical contained *q-char-sequence* (including > characters, if any) from the original directive.

- 10 Either form of the **#embed** directive specified previously behaves as specified later in this subclause. The values of the integer constant expressions in the expanded sequence are determined by an implementation-defined mapping of the resource's data. Each integer constant expression's value is in the range from 0 to  $(2^{\text{embed element width}}) - 1$ , inclusive.<sup>207)</sup> If:

- the list of integer constant expressions is used to initialize an array of a type compatible with **unsigned char**, or compatible with **char** if **char** cannot hold negative values; and,
- the embed element width is equal to **CHAR\_BIT** (5.3.5.3.2),

then the contents of the initialized elements of the array are as-if the resource's binary data is **fread** (7.23.8.1) into the array at translation time.

- 11 A preprocessing directive of the form

**# embed** *pp-tokens new-line*

(that does not match one of the two previous forms) is permitted. The preprocessing tokens after **embed** in the directive are processed just as in normal text. (Each identifier currently defined as a macro name is replaced by its replacement list of preprocessing tokens.) The directive resulting after all replacements shall match one of the two previous forms.<sup>208)</sup> The method by which a sequence of preprocessing tokens between a < and a > preprocessing token pair or a pair of " characters is combined into a single resource name preprocessing token is implementation-defined.

- 12 An embed parameter with a preprocessor parameter token that is one of the following is a *standard embed parameter*:

**limit**

**prefix**

**suffix**

**if\_empty**

The significance of these standard embed parameters is specified later in this subclause.

#### Recommended practice

- 13 The **#embed** directive is meant to translate binary data in a resource to a sequence of integer constant expressions in a way that preserves the value of the resource's bit stream where possible.
- 14 A mechanism similar to, but distinct from, the implementation-defined search paths used for source file inclusion (6.10.3) is encouraged.
- 15 Implementations should take into account translation-time bit and byte orders as well as execution-time bit and byte orders to more appropriately represent the resource's binary data from the directive. This maximizes the chance that, if the resource referenced at translation time through the **#embed** directive is the same one accessed through execution-time means, the data that is e.g. **fread** or similar into contiguous storage will compare bit-for-bit equal to an array of character type initialized from an **#embed** directive's expanded contents.
- 16 EXAMPLE 1 Placing a small image resource.

<sup>207)</sup>For example, an embed element width of 8 will yield a range of values from 0 to 255, inclusive.

<sup>208)</sup>Adjacent string literals are not concatenated into a single string literal (see the translation phases in 5.2.1.2); thus, an expansion that results in two string literals is an invalid directive.



```
#include <stddef.h>

void have_you_any_wool(const unsigned char*, size_t);

int main (int, char*[]) {
 static const unsigned char baa_baa[] = {
#embed "black_sheep.ico"
 };

 have_you_any_wool(baa_baa, _Lengthof(baa_baa));

 return 0;
}
```

17 EXAMPLE 2 This snippet:

```
int main (int, char*[]) {
 static const unsigned char coefficients[] = {
#embed "only_8_bits.bin" // potential constraint violation
 };

 return 0;
}
```

can violate the constraint that  $(resource\ width) \% (embed\ element\ width)$  is 0. There is a chance that the implementation-defined 8 bit width of the resource is not evenly divisible by the embed element width (e.g. on a system where **CHAR\_BIT** is 16). Issuing a diagnostic in this case can aid in portability by calling attention to potentially incompatible expectations between implementations and their resources.

18 EXAMPLE 3 Initialization of non-arrays.

```
int main () {
 /* Braces can be kept or elided as per normal initialization rules */
 int i = {
#embed "i.dat"
 }; /* valid if i.dat produces 1 value,
 i value is [0, 2(embed element width)) */
 int i2 =
#embed "i.dat"
 ; /* valid if i.dat produces 1 value,
 i2 value is [0, 2(embed element width)) */
 struct s {
 double a, b, c;
 struct { double e, f, g; };
 double h, i, j;
 };
 struct s x = {
 /* initializes each element in order according to initialization
 rules with comma-separated list of integer constant expressions
 inside of braces */
#embed "s.dat"
 };
 return 0;
}
```

Non-array types can still be initialized since the directive produces a comma-delimited list of integer constant expressions, a single integer constant expression, or nothing.

19 EXAMPLE 4 Equivalency of bit sequence and bit order between a translation-time read and an execution-time read of the same resource/file.

```
#include <string.h>
```

```

#include <stddef.h>
#include <stdio.h>

int main(void) {
 static const unsigned char embed_data[] = {
#embed <data.dat>
 };

 const size_t f_size = sizeof(embed_data);
 unsigned char f_data[f_size];
 FILE* f_source = fopen("data.dat", "rb");
 if (f_source == nullptr)
 return 1;
 char* f_ptr = (char*)&f_data[0];
 if (fread(f_ptr, 1, f_size, f_source) != f_size) {
 fclose(f_source);
 return 1;
 }
 fclose(f_source);

 int is_same = memcmp(&embed_data[0], f_ptr, f_size);
 // if both operations refers to the same resource/file at
 // execution time and translation time, "is_same" should be 0
 return is_same == 0 ? 0 : 1;
}

```

#### 6.10.4.2 **limit** parameter

##### Constraints

- 1 The **limit** standard embed parameter can appear zero times or one time in the embed parameter sequence. Its preprocessor argument clause shall be present and have the form:

( *constant-expression* )

and shall be an integer constant expression. The integer constant expression shall not evaluate to a value less than 0.

- 2 The token **defined** shall not appear within the constant expression.

##### Semantics

- 3 The embed parameter with a preprocessor parameter token limit denotes a balanced preprocessing token sequence that will be used to compute the resource width. Independently of any macro replacement done previously (e.g. when matching the form of **#embed**), the constant expression is evaluated after the balanced preprocessing token sequence is processed as in normal text, using the rules specified for conditional inclusion (6.10.2), with the exception that any defined macro expressions are not permitted.
- 4 The resource width is:
  - 0, if the integer constant expression evaluates to 0; or,
  - the implementation resource width if it is less than the embed element width multiplied by the integer constant expression; or,
  - the embed element width multiplied by the integer constant expression, if it is less than or equal to the implementation resource width.

- 5 EXAMPLE 1 Checking the first 4 elements of a sound resource.

```

#include <assert.h>

int main (int, char*[]) {

```

```

static const char sound_signature[] = {
#embed <sdk/jump.wav> limit(2+2)
};
static_assert(_Lengthof sound_signature == 4,
 "There should only be 4 elements in this array.");

// verify PCM WAV resource
assert(sound_signature[0] == 'R');
assert(sound_signature[1] == 'I');
assert(sound_signature[2] == 'F');
assert(sound_signature[3] == 'F');
assert(sizeof(sound_signature) == 4);

return 0;
}

```

- 6 EXAMPLE 2 Similar to a previous example, except it illustrates macro expansion specifically done for the `limit(...)` parameter.

```

#include <assert.h>

#define TWO_PLUS_TWO 2+2

int main (int, char*[]) {
 const char sound_signature[] = {
 /* the token sequence within the parentheses
 for the "limit" parameter undergoes macro
 expansion, at least once, resulting in
#embed <sdk/jump.wav> limit(2+2)
 */
#embed <sdk/jump.wav> limit(TWO_PLUS_TWO)
 };
 static_assert(_Lengthof(sound_signature) == 4,
 "There should only be 4 elements in this array.");

 // verify PCM WAV resource
 assert(sound_signature[0] == 'R');
 assert(sound_signature[1] == 'I');
 assert(sound_signature[2] == 'F');
 assert(sound_signature[3] == 'F');
 assert(sizeof(sound_signature) == 4);

 return 0;
}

```

- 7 EXAMPLE 3 A potential constraint violation since an environment that has a `CHAR_BIT` whose value is greater than 24 can possibly not get enough data from the resource.

```

int main (int, char*[]) {
 const unsigned char arr[] = {
#embed "24_bits.bin" limit(1) // can be a constraint violation
 };

 return 0;
}

```

- 8 EXAMPLE 4 Resources interfacing with certain implementations can receive an infinite stream of data, such as the (hypothetical) `</owo/uwurandom>` resource used in the following snippet:

```

int main (int, char*[]) {
 const unsigned char arr[] = {

```

```
#embed </owo/uwurandom> limit(513)
};

return 0;
}
```

The **limit** parameter can help process only a portion of that information and prevent exhaustion of an implementation's internal resources when processing such data.

#### 6.10.4.3 suffix parameter

##### Constraints

- 1 The **suffix** standard embed parameter can appear zero times or one time in the embed parameter sequence. Its preprocessor argument clause shall be present and have the form:

( *pp-balanced-token-sequence*<sub>opt</sub> )

##### Semantics

- 2 The embed parameter with a preprocessing parameter token **suffix** denotes a balanced preprocessing token sequence within its preprocessor argument clause that will be placed immediately after the result of the associated **#embed** directive's expansion.
- 3 If the resource is empty, then **suffix** has no effect and is ignored.
- 4 EXAMPLE Extra elements added to array initializer.

```
#include <string.h>

#ifndef SHADER_TARGET
#define SHADER_TARGET "edith-impl.glsl"
#endif

extern char* null_term_shader_data;

void fill_in_data () {
 const char internal_data[] = {
#embed SHADER_TARGET \
 suffix(,)
 0
 };

 strcpy(null_term_shader_data, internal_data);
}
```

#### 6.10.4.4 prefix parameter

##### Constraints

- 1 The **prefix** standard embed parameter can appear zero times or one time in the embed parameter sequence. Its preprocessor parameter clause shall be present and have the form:

( *pp-balanced-token-sequence*<sub>opt</sub> )

##### Semantics

- 2 The embed parameter with a preprocessor parameter token **prefix** denotes a balanced preprocessing token sequence within its preprocessor argument clause that will be placed immediately before the result of the associated **#embed** directive's expansion, if any.
- 3 If the resource is empty, then **prefix** has no effect and is ignored.
- 4 EXAMPLE A null-terminated character array with prefixed and suffixed additional tokens when the resource is not empty, providing null termination and a byte order mark.

```
#include <assert.h>
#include <string.h>
```

```

#ifndef SHADER_TARGET
#define SHADER_TARGET "ches.glsl"
#endif

extern char* merp;

void init_data () {
 const char whl[] = {
#embed SHADER_TARGET
 prefix(0xEF, 0xBB, 0xBF,) /* UTF-8 BOM */ \
 suffix(,)
 0
 };
 // always null terminated,
 // contains BOM if not-empty
 int is_good = (_Lengthof(whl) == 1 && whl[0] == '\0')
 || (whl[0] == '\xEF' && whl[1] == '\xBB'
 && whl[2] == '\xBF' && whl[_Lengthof(whl) - 1] == '\0');
 assert(is_good);
 strcpy(merp, whl);
}

```

#### 6.10.4.5 **if\_empty** parameter

##### Constraints

- 1 The **if\_empty** standard embed parameter can appear zero times or one time in the embed parameter sequence. Its preprocessor argument clause shall be present and have the form:

( *pp-balanced-token-sequence<sub>opt</sub>* )

##### Semantics

- 2 The embed parameter with a preprocessing parameter token **if\_empty** denotes a balanced preprocessing token sequence within its preprocessor argument clause that will replace the **#embed** directive entirely.

If the resource is not empty, then **if\_empty** has no effect and is ignored.

- 3 EXAMPLE 1 If the search for the resource is successful, this resource is always considered empty due to the **limit(0)** embed parameter. This program always returns 0, even if the resource is searched for and found successfully by the implementation and has an implementation resource width greater than 0.

```

int main () {
 return
#embed <some_resource> limit(0) prefix(1) if_empty(0)
 ;
 // becomes:
 // return 0;
}

```

- 4 EXAMPLE 2 An example similar to using the suffix **embed** parameter, but changed slightly.

```

#include <string.h>

#ifndef SHADER_TARGET
#define SHADER_TARGET "edith-impl.glsl"
#endif

extern char* null_term_shader_data;

void fill_in_data () {
 const char internal_data[] = {

```

```
#embed SHADER_TARGET \
 suffix(, 0) \
 if_empty(0)
};

strcpy(null_term_shader_data, internal_data);
}
```

- 5 EXAMPLE 3 This resource is considered empty due to the **limit(0)** embed parameter, meaning an **if\_empty** expression replaces the directive as specified previously. A constraint is still violated if the search for the resource is unsuccessful.

```
int main () {
 return
 #embed <infinite-resource> limit(0) if_empty(45540)
 ;
}
```

becomes:

```
int main () {
 return 45540;
}
```

## 6.10.5 Macro replacement

### 6.10.5.1 General

#### Constraints

- 1 Two replacement lists are identical if and only if the preprocessing tokens in both have the same number, ordering, spelling, and white-space separation, where all white-space separations are considered identical.
- 2 An identifier currently defined as an object-like macro shall not be redefined by another **#define** preprocessing directive unless the second definition is an object-like macro definition and the two replacement lists are identical. Likewise, an identifier currently defined as a function-like macro shall not be redefined by another **#define** preprocessing directive unless the second definition is a function-like macro definition that has the same number and spelling of parameters, and the two replacement lists are identical.
- 3 There shall be white space between the identifier and the replacement list in the definition of an object-like macro.
- 4 If the identifier-list in the macro definition does not end with an ellipsis, the number of arguments (including those arguments consisting of no preprocessing tokens) in an invocation of a function-like macro shall equal the number of parameters in the macro definition. Otherwise, there shall be at least as many arguments in the invocation as there are parameters in the macro definition (excluding the ...). There shall exist a ) preprocessing token that terminates the invocation.
- 5 The identifiers **\_\_VA\_ARGS\_\_** and **\_\_VA\_OPT\_\_** shall occur only in the replacement-list of a function-like macro that uses the ellipsis notation in the parameters.
- 6 A parameter identifier in a function-like macro shall be uniquely declared within its scope.

#### Semantics

- 7 The identifier immediately following the **define** is called the *macro name*. There is one name space for macro names. Any white-space characters preceding or following the replacement list of preprocessing tokens are not considered part of the replacement list for either form of macro.
- 8 If a # preprocessing token, followed by an identifier, occurs lexically at the point at which a preprocessing directive can begin, the identifier is not subject to macro replacement.
- 9 A preprocessing directive of the form

**# define** *identifier replacement-list new-line*

defines an *object-like macro* that causes each subsequent instance of the macro name<sup>209)</sup> to be replaced by the replacement list of preprocessing tokens that constitute the remainder of the directive. The replacement list is then rescanned for more macro names as specified later in this subclause.

- 10 A preprocessing directive of the form

```
define identifier lparen identifier-listopt) replacement-list new-line
define identifier lparen ...) replacement-list new-line
define identifier lparen identifier-list , ...) replacement-list new-line
```

defines a *function-like macro* with parameters, whose use is similar syntactically to a function call. The parameters are specified by the optional list of identifiers, whose scope extends from their declaration in the identifier list until the new-line character that terminates the **#define** preprocessing directive. Each subsequent instance of the function-like macro name followed by a ( as the next preprocessing token introduces the sequence of preprocessing tokens that is replaced by the replacement list in the definition (an invocation of the macro). The replaced sequence of preprocessing tokens is terminated by the matching ) preprocessing token, skipping intervening matched pairs of left and right parenthesis preprocessing tokens. Within the sequence of preprocessing tokens making up an invocation of a function-like macro, new-line is considered a normal white-space character.

- 11 The sequence of preprocessing tokens bounded by the outside-most matching parentheses forms the list of arguments for the function-like macro. The individual arguments within the list are separated by comma preprocessing tokens, but comma preprocessing tokens between matching inner parentheses do not separate arguments. If there are sequences of preprocessing tokens within the list of arguments that would otherwise act as preprocessing directives,<sup>210)</sup> the behavior is undefined.
- 12 If there is a ... in the identifier-list in the macro definition, then the trailing arguments (if any), including any separating comma preprocessing tokens, are merged to form a single item: the *varying arguments*. The number of arguments so combined is such that, following merger, the number of arguments is one more than the number of parameters in the macro definition (excluding the ...), except that if there are as many arguments as named parameters, the macro invocation behaves as if a comma token has been appended to the argument list such that varying arguments are formed that contain no pp-tokens.

**6.10.5.2 Argument substitution****Syntax**

- 1 *va-opt-replacement*:

```
__VA_OPT__ (pp-tokensopt)
```

**Description**

- 2 Argument substitution is a process during macro expansion in which identifiers corresponding to the parameters of the macro definition and the special constructs **\_\_VA\_ARGS\_\_** and **\_\_VA\_OPT\_\_** are replaced with token sequences from the arguments of the macro invocation and possibly of the argument of the feature **\_\_VA\_OPT\_\_**. The latter process allows to control a substitute token sequence that is only expanded if the argument list that corresponds to a trailing ... of the parameter list is present and has a non-empty substitution.

**Constraints**

- 3 The identifier **\_\_VA\_OPT\_\_** shall always occur as part of the preprocessing token sequence *va-opt-replacement*; its closing ) is determined by skipping intervening pairs of matching left and right parentheses in its pp-tokens. The pp-tokens of a *va-opt-replacement* shall not contain **\_\_VA\_OPT\_\_**. The pp-tokens shall form a valid replacement list for the current function-like macro.

<sup>209)</sup>Since, by macro-replacement time, all character literals and string literals are preprocessing tokens, not sequences possibly containing identifier-like subsequences (see 5.2.1.2, translation phases), they are never scanned for macro names or parameters.

<sup>210)</sup>Despite the name, a non-directive is a preprocessing directive.

## Semantics

- 4 After the arguments for the invocation of a function-like macro have been identified, argument substitution takes place. A va-opt-replacement is treated as if it were a parameter. For each parameter in the replacement list that is neither preceded by a **#** or **##** preprocessing token nor followed by a **##** preprocessing token, the preprocessing tokens naming the parameter are replaced by a token sequence determined as follows:

- If the parameter is of the form va-opt-replacement, the replacement preprocessing tokens are the preprocessing token sequence for the corresponding argument, as specified later in this subclause.
- Otherwise, the replacement preprocessing tokens are the preprocessing tokens of the corresponding argument after all macros contained therein have been expanded. The argument's preprocessing tokens are completely macro replaced before being substituted as if they formed the rest of the preprocessing file with no other preprocessing tokens being available.

### 5 EXAMPLE 1

```
#define LPAREN() (
#define G(Q) 42
#define F(R, X, ...) __VA_OPT__(G R X))
int x = F(LPAREN(), 0, <:-); // replaced by int x = 42;
```

- 6 An identifier **\_\_VA\_ARGS\_\_** that occurs in the replacement list is treated as if it were a parameter, and the varying arguments form the preprocessing tokens used to replace it.
- 7 The preprocessing token sequence for the corresponding argument of a va-opt-replacement is defined as follows. If a (hypothetical) substitution of **\_\_VA\_ARGS\_\_** as neither an operand of **#** nor **##** consists of no preprocessing tokens, the argument consists of a single placemaker preprocessing token (6.10.5.4, 6.10.5.5). Otherwise, the argument consists of the results of the expansion of the contained pp-tokens as the replacement list of the current function-like macro before removal of placemaker tokens, rescanning, and further replacement.
- 8 NOTE The placemaker tokens are removed before stringization (6.10.5.3), and can be removed by rescanning and further replacement (6.10.5.5).

### 9 EXAMPLE 2

```
#define F(...) f(0 __VA_OPT__(,) __VA_ARGS__)
#define G(X, ...) f(0, X __VA_OPT__(,) __VA_ARGS__)
#define SDEF(sname, ...) S sname __VA_OPT__(= { __VA_ARGS__ })
#define EMP

F(a, b, c) // replaced by f(0, a, b, c)
F() // replaced by f(0)
F(EMP) // replaced by f(0)

G(a, b, c) // replaced by f(0, a, b, c)
G(a,) // replaced by f(0, a)
G(a) // replaced by f(0, a)

SDEF(foo); // replaced by S foo;
SDEF(bar, 1, 2); // replaced by S bar = { 1, 2 };

#define H1(X, ...) X __VA_OPT__(##) __VA_ARGS__
// error: ## on line above
// cannot appear at the beginning of a replacement
// list (6.10.5.4)

#define H2(X, Y, ...) __VA_OPT__(X ## Y,) __VA_ARGS__
```



```

H2(a, b, c, d) // replaced by ab, c, d

#define H3(X, ...) #__VA_OPT__(X##X X##X)
H3(, 0) // replaced by ""

#define H4(X, ...) __VA_OPT__(a X ## X) ## b
H4(, 1) // replaced by a b

#define H5A(...) __VA_OPT__()/**/__VA_OPT__()
#define H5B(X) a ## X ## b
#define H5C(X) H5B(X)
H5C(H5A()) // replaced by ab

```

### 6.10.5.3 The # operator

#### Constraints

- Each # preprocessing token in the replacement list for a function-like macro shall be followed by a parameter as the next preprocessing token in the replacement list.

#### Semantics

- If, in the replacement list, a parameter is immediately preceded by a # preprocessing token, both are replaced by a single ordinary string literal preprocessing token that contains the spelling of the preprocessing token sequence for the corresponding argument (excluding placemarkers).
- Let the *stringizing argument* be the preprocessing token sequence for the corresponding argument with placemarkers removed. Each occurrence of white space between the stringizing argument's preprocessing tokens becomes a single space character in the ordinary string literal. White space before the first preprocessing token and after the last preprocessing token composing the stringizing argument is deleted. Otherwise, the original spelling of each preprocessing token in the stringizing argument is retained in the ordinary string literal, except for special handling for producing the spelling of string literals and character literals: a \ character is inserted before each " and \ character of a character literal or string literal (including the delimiting " characters), except that it is implementation-defined whether a \ character is inserted before the \ character beginning a universal character name.
- If the replacement that results is not a valid ordinary string literal, the behavior is undefined. The ordinary string literal corresponding to an empty stringizing argument is "". The order of evaluation of # and ## operators is unspecified.

### 6.10.5.4 The ## operator

#### Constraints

- A ## preprocessing token shall not occur at the beginning or at the end of a replacement list for either form of macro definition.

#### Semantics

- If, in the replacement list of a function-like macro, a parameter is immediately preceded or followed by a ## preprocessing token, the parameter is replaced by the corresponding argument's preprocessing token sequence; however, if an argument consists of no preprocessing tokens, the parameter is replaced by a *placemarker* preprocessing token instead.<sup>211)</sup>
- For both object-like and function-like macro invocations, before the replacement list is reexamined for more macro names to replace, each instance of a ## preprocessing token in the replacement list (not from an argument) is deleted and the preceding preprocessing token is concatenated with the following preprocessing token. Placemarker preprocessing tokens are handled specially: concatenation of two placemarkers results in a single placemarker preprocessing token, and concatenation of a placemarker with a non-placemarker preprocessing token results in the non-placemarker preprocessing token. If the result is not a valid preprocessing token, the behavior is undefined. The

<sup>211)</sup>Placemarker preprocessing tokens do not appear in the syntax because they are temporary entities that exist only within translation phase 4 (5.2.1.2).

resulting token is available for further macro replacement. The order of evaluation of **##** operators is unspecified.

- 4 EXAMPLE In the following fragment:

```
#define hash_hash # ## #
#define mkstr(a) # a
#define in_between(a) mkstr(a)
#define join(c, d) in_between(c hash_hash d)

char p[] = join(x, y); // equivalent to
 // char p[] = "x ## y";
```

The expansion produces, at various stages:

```
join(x, y)

in_between(x hash_hash y)

in_between(x ## y)

mkstr(x ## y)

"x ## y"
```

In other words, expanding **hash\_hash** produces a new token, consisting of two adjacent sharp signs, but this new token is not the **##** operator.

#### 6.10.5.5 Rescanning and further replacement

- 1 After all parameters in the replacement list have been substituted and **#** and **##** processing has taken place, all placemarkers preprocessing tokens are removed. The resulting preprocessing token sequence is then rescanned, along with all subsequent preprocessing tokens of the source file, for more macro names to replace.
- 2 If the name of the macro being replaced is found during this scan of the replacement list (not including the rest of the source file's preprocessing tokens), it is not replaced. Furthermore, if any nested replacements encounter the name of the macro being replaced, it is not replaced. These nonreplaced macro name preprocessing tokens are no longer available for further replacement even if they are later (re)examined in contexts in which that macro name preprocessing token would otherwise have been replaced.
- 3 The resulting completely macro-replaced preprocessing token sequence is not processed as a preprocessing directive even if it resembles one, but all pragma unary operator expressions within it are then processed as specified in 6.10.11.
- 4 EXAMPLE There are cases where it is not clear whether a replacement is nested or not. For example, given the following macro definitions:

```
#define f(a) a*g
#define g(a) f(a)
```

the invocation

```
f(2)(9)
```

can expand to either

```
2*f(9)
```

or

```
2*9*g
```

Strictly conforming programs are not permitted to depend on such unspecified behavior.

#### 6.10.5.6 Scope of macro definitions

- 1 A macro definition lasts (independent of block structure) until a corresponding **#undef** directive is encountered or (if none is encountered) until the end of the preprocessing translation unit. Macro definitions have no significance after translation phase 4.

- 2 A preprocessing directive of the form

**# undef** *identifier new-line*

causes the specified identifier no longer to be defined as a macro name. It is ignored if the specified identifier is not currently defined as a macro name.

- 3 EXAMPLE 1 The simplest use of this facility is to define a “manifest constant”, as in

```
#define TABSIZE 100

int table[TABSIZE];
```

- 4 EXAMPLE 2 The following defines a function-like macro whose value is the maximum of its arguments. It has the advantages of working for any compatible types of the arguments and of generating in-line code without the overhead of function calling. It has the disadvantages of evaluating one or the other of its arguments a second time (including side effects) and generating more code than a function if invoked several times. It also cannot have its address taken, as it has none.

```
#define max(a, b) ((a) > (b) ? (a) : (b))
```

The parentheses ensure that the arguments and the resulting expression are bound properly.

- 5 EXAMPLE 3 To illustrate the rules for redefinition and reexamination, the sequence

```
#define x 3
#define f(a) f(x * (a))
#undef x
#define x 2
#define g f
#define z z[0]
#define h g(~
#define m(a) a(w)
#define w 0,1
#define t(a) a
#define p() int
#define q(x) x
#define r(x,y) x ## y
#define str(x) # x

f(y+1) + f(f(z)) % t(t(g)(0) + t)(1);
g(x+(3,4)-w) | h 5) & m
 (f)^m(m);
p() i[q()] = { q(1), r(2,3), r(4,), r(,5), r(,) };
char c[2][6] = { str(hello), str() };
```

results in

```
f(2 * (y+1)) + f(2 * (f(2 * (z[0])))) % f(2 * (0)) + t(1);
f(2 * (2+(3,4)-0,1)) | f(2 * (~ 5)) & f(2 * (0,1))^m(0,1);
int i[] = { 1, 23, 4, 5, };
char c[2][6] = { "hello", "" };
```

- 6 EXAMPLE 4 To illustrate the rules for creating ordinary string literals and concatenating tokens, the sequence

```
#define str(s) # s
```

```

#define xstr(s) str(s)
#define debug(s, t) printf("x" # s " = %d, x" # t " = %s", \
 x ## s, x ## t)
#define INCFILE(n) vers ## n
#define glue(a, b) a ## b
#define xglue(a, b) glue(a, b)
#define HIGHLOW "hello"
#define LOW LOW " , world"

debug(1, 2);
fputs(str(strncmp("abc\0d", "abc", '\4') // this goes away
 == 0) str(: @\n), s);
#include xstr(INCFILE(2).h)
glue(HIGH, LOW);
xglue(HIGH, LOW)

```

results in

```

printf("x" "1" " = %d, x" "2" " = %s", x1, x2);
fputs(
 "strcmp(\"abc\\0d\", \"abc\", '\\4') == 0" ": @\n",
 s);
#include "vers2.h" (after macro replacement, before file access)
"hello";
"hello" " , world"

```

or, after concatenation of the ordinary string literals,

```

printf("x1= %d, x2= %s", x1, x2);
fputs(
 "strcmp(\"abc\\0d\", \"abc\", '\\4') == 0: @\n",
 s);
#include "vers2.h" (after macro replacement, before file access)
"hello";
"hello, world"

```

Space around the # and ## tokens in the macro definition is optional.

- 7 EXAMPLE 5 To illustrate the rules for placemaker preprocessing tokens, the sequence

```

#define t(x,y,z) x ## y ## z
int j[] = { t(1,2,3), t(,4,5), t(6,,7), t(8,9,),
 t(10,,), t(,11,), t(,12), t(,,) };

```

results in

```

int j[] = { 123, 45, 67, 89,
 10, 11, 12, };

```

- 8 EXAMPLE 6 To demonstrate the redefinition rules, the following sequence is valid.

```

#define OBJ_LIKE (1-1)
#define OBJ_LIKE /* white space */ (1-1) /* other */
#define FUNC_LIKE(a) (a)
#define FUNC_LIKE(a) (/* note the white space */ \
 a /* other stuff on this line
 */)

```

But the following redefinitions of the preceding definitions are invalid:

```

#define OBJ_LIKE (0) // different token sequence

```

```
#define OBJ_LIKE (1 - 1) // different white space
#define FUNC_LIKE(b) (a) // different parameter usage
#define FUNC_LIKE(b) (b) // different parameter spelling
```

- 9 EXAMPLE 7 Finally, to show the varying argument list macro facilities:

```
#define debug(...) fprintf(stderr, __VA_ARGS__)
#define showlist(...) puts(__VA_ARGS__)
#define report(test, ...) ((test)?puts(#test):\
 printf(__VA_ARGS__))
debug("Flag");
debug("X = %d\n", x);
showlist(The first, second, and third items.);
report(x>y, "x is %d but y is %d", x, y);
```

results in

```
fprintf(stderr, "Flag");
fprintf(stderr, "X = %d\n", x);
puts("The first, second, and third items.");
((x>y)?puts("x>y"):
 printf("x is %d but y is %d", x, y));
```

## 6.10.6 Line control

### Constraints

- 1 The string literal of a **#line** directive, if present, shall be an ordinary string literal.

### Semantics

- 2 The *line number* of the current source line is one greater than the number of new-line characters read or introduced in translation phase 1 (5.2.1.2) while processing the source file to the current token.
- 3 If a preprocessing token (in particular **\_\_LINE\_\_**) spans two or more physical lines, it is unspecified which of those line numbers is associated with that token. If a preprocessing directive spans two or more physical lines, it is unspecified which of those line numbers is associated with the preprocessing directive. If a macro invocation spans multiple physical lines, it is unspecified which of those line numbers is associated with that invocation. The line number of a preprocessing token is independent of the context (in particular, as a macro argument or in a preprocessing directive). The line number of a **\_\_LINE\_\_** in a macro body is the line number of the macro invocation.
- 4 A preprocessing directive of the form

**# line** *digit-sequence* *new-line*

causes the implementation to behave as if the following sequence of source lines begins with a source line that has a line number as specified by the digit sequence (interpreted as a decimal integer, ignoring any optional digit separators (6.4.5.2) in the digit sequence). The digit sequence shall not specify zero, nor a number greater than 2147483647.

- 5 A preprocessing directive of the form

**# line** *digit-sequence* " *s-char-sequence*<sub>opt</sub> " *new-line*

sets the presumed line number similarly and changes the presumed name of the source file to be the contents of the ordinary string literal.

- 6 A preprocessing directive of the form

**# line** *pp-tokens* *new-line*

(that does not match one of the two previous forms) is permitted. The preprocessing tokens after **line** on the directive are processed just as in normal text (each identifier currently defined as a macro name is replaced by its replacement list of preprocessing tokens). The directive resulting after

all replacements shall match one of the two previous forms and is then processed as appropriate.<sup>212)</sup>

### Recommended practice

- 7 The line number associated with a *pp-token* should be the line number of the first character of the *pp-token*. The line number associated with a preprocessing directive should be the line number of the line with the first **#** token. The line number associated with a macro invocation should be the line number of the first character of the macro name in the invocation.

## 6.10.7 Diagnostic directives

### Semantics

- 1 A preprocessing directive of either form

```
error pp-tokensopt new-line
warning pp-tokensopt new-line
```

causes the implementation to produce a diagnostic message that includes the specified sequence of preprocessing tokens.

## 6.10.8 Pragma directive

### Semantics

- 1 A preprocessing directive of the form

```
pragma pp-tokensopt new-line
```

where the preprocessing token **STDC** does not immediately follow **pragma** in the directive (prior to any macro replacement)<sup>213)</sup> causes the implementation to behave in an implementation-defined manner. The behavior can cause translation to fail or cause the translator or the resulting program to behave in a non-conforming manner. Any such **pragma** that is not recognized by the implementation is ignored.

- 2 If the preprocessing token **STDC** does immediately follow **pragma** in the directive (prior to any macro replacement), then no macro replacement is performed on the directive, and the directive shall have one of the following forms<sup>214)</sup> whose meanings are described elsewhere:

*standard-pragma:*

```
pragma STDC FP_CONTRACT on-off-switch
pragma STDC FENV_ACCESS on-off-switch
pragma STDC FENV_DEC_ROUND dec-direction
pragma STDC FENV_ROUND direction
pragma STDC CX_LIMITED_RANGE on-off-switch
```

*on-off-switch:* one of

```
ON OFF DEFAULT
```

*direction:* one of

```
FE_DOWNWARD FE_TONEAREST FE_TONEARESTFROMZERO
FE_TOWARDZERO FE_UPWARD FE_DYNAMIC
```

<sup>212)</sup>Because a new-line is explicitly included as part of the **#line** directive, the number of new-line characters read while processing to the first *pp-token* can be different depending on whether the implementation uses a one-pass preprocessor. Therefore, there are two possible values for the line number following a directive of the form **#line** LINE *new-line*.

<sup>213)</sup>An implementation is not required to perform macro replacement in pragmas, but it is permitted except for in standard pragmas (where **STDC** immediately follows **pragma**). If the result of macro replacement in a non-standard pragma has the same form as a standard pragma, the behavior is still implementation-defined; an implementation is permitted to behave as if it were the standard pragma, but is not required to.

<sup>214)</sup>See “future language directions” (6.11.7).

*dec-direction*: one of

**FE\_DEC\_DOWNWARD**    **FE\_DEC\_TONEAREST**    **FE\_DEC\_TONEARESTFROMZERO**  
**FE\_DEC\_TOWARDZERO**    **FE\_DEC\_UPWARD**    **FE\_DEC\_DYNAMIC**

### Recommended practice

- Implementations are encouraged to diagnose unrecognized pragmas.

**Forward references:** the **FP\_CONTRACT** pragma (7.12.3), the **FENV\_ACCESS** pragma (7.6.2), the **FENV\_DEC\_ROUND** pragma (7.6.4), the **FENV\_ROUND** pragma (7.6.3), the **CX\_LIMITED\_RANGE** pragma (7.3.4).

## 6.10.9 Null directive

### Semantics

- A preprocessing directive of the form

*# new-line*

has no effect.

## 6.10.10 Predefined macro names

### 6.10.10.1 General

- The values of the predefined macros listed in the following subclauses<sup>215)</sup> (except for **\_\_FILE\_\_** and **\_\_LINE\_\_**) remain constant throughout the translation unit.
- None of the following macro names in this subclause nor the identifiers **defined**, **\_\_has\_c\_attribute**, **\_\_has\_include**, or **\_\_has\_embed** shall be the subject of a **#define** or a **#undef** preprocessing directive. Any other predefined macro names: shall begin with a leading underscore followed by an uppercase letter; or, a second underscore; or, shall be any of the identifiers **alignas**, **alignof**, **bool**, **false**, **static\_assert**, **thread\_local**, or **true**.
- The implementation shall not predefine the macro **\_\_cplusplus**, nor shall it define it in any standard header.

**Forward references:** standard headers (7.1.2).

### 6.10.10.2 Mandatory macros

- The following macro names shall be defined by the implementation:

**\_\_DATE\_\_** The date of translation of the preprocessing translation unit: an ordinary string literal of the form "Mmm dd yyyy", where the names of the months are the same as those generated by the **asctime** function, and the first character of **dd** is a space character if the value is less than 10. If the date of translation is not available, an implementation-defined valid date shall be supplied.

**\_\_FILE\_\_** The presumed name of the current source file (an ordinary string literal).<sup>216)</sup>

**\_\_LINE\_\_** The presumed line number (within the current source file) of the current source line (an integer literal).<sup>216)</sup>

**\_\_STDC\_\_** The integer literal 1, intended to indicate a conforming implementation.

**\_\_STDC\_EMBED\_NOT\_FOUND\_\_**, **\_\_STDC\_EMBED\_FOUND\_\_**, **\_\_STDC\_EMBED\_EMPTY\_\_** The integer literals 0, 1, and 2, respectively.

**\_\_STDC\_HOSTED\_\_** The integer literal 1 if the implementation is a hosted implementation or the integer literal 0 if it is not.

<sup>215)</sup>See "future language directions" (6.11.8).

<sup>216)</sup>The presumed source file name and line number can be changed by the **#line** directive.



**\_\_STDC\_UTF\_16\_\_** The integer literal 1, intended to indicate that values of type **char16\_t** are UTF-16 encoded.

**\_\_STDC\_UTF\_32\_\_** The integer literal 1, intended to indicate that values of type **char32\_t** are UTF-32 encoded.

**\_\_STDC\_VERSION\_\_** The integer literal 202ymmL.<sup>217)</sup>

**\_\_TIME\_\_** The time of translation of the preprocessing translation unit: an ordinary string literal of the form "hh:mm:ss" as in the time generated by the **asctime** function. If the time of translation is not available, an implementation-defined valid time shall be supplied.

**Forward references:** the **asctime** function (7.30.3.2).

### 6.10.10.3 Environment macros

- 1 The following macro names are conditionally defined by the implementation:

**\_\_STDC\_ISO\_10646\_\_** An integer literal of the form **yyyymmL** (for example, 202012L). If this symbol is defined, then every character in the Unicode required set, when stored in an object of type **wchar\_t**, has the same value as the short identifier of that character. The *Unicode required set* consists of all the characters that are defined by ISO/IEC 10646, along with all amendments and technical corrigenda, as of the specified year and month. If some other encoding is used, the macro shall not be defined and the actual encoding used is implementation-defined.

**\_\_STDC\_LITERAL\_UTF8\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **char** (with each the value of **char** treated as an **unsigned char**) from a string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-8 sequence.

**\_\_STDC\_LITERAL\_UTF16\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **char** (with each the value of **char** treated as an **unsigned char**) from a string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-16 sequence.

**\_\_STDC\_LITERAL\_UTF32\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **char** (with each the value of **char** treated as an **unsigned char**) from a string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-32 sequence.

**\_\_STDC\_MB\_MIGHT\_NEQ\_WC\_\_** The integer literal 1, intended to indicate that, in the encoding for **wchar\_t**, a member of the basic character set is not required to have a code value equal to its value when used as the lone character in an ordinary character literal.

**\_\_STDC\_WIDE\_LITERAL\_UTF8\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the wide literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **wchar\_t** from a **wchar\_t** string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-8 sequence.

**\_\_STDC\_WIDE\_LITERAL\_UTF16\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the wide literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **wchar\_t** from a **wchar\_t** string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-16 sequence.

<sup>217)</sup>See Annex M for the values in previous editions of this document. The intention is that this will remain an integer literal of type **long int** that is increased with each edition of this document.



**\_\_STDC\_WIDE\_LITERAL\_UTF32\_\_** A strictly positive integer constant expression if this symbol is defined. If it is defined, then the wide literal encoding (6.2.9) is capable of storing every character in the Unicode required set and every sequence of **wchar\_t** from a **wchar\_t** string literal (excluding non-Universal character escape sequences) forms a valid encoding of a UTF-32 sequence.

**Forward references:** common definitions (7.21), Unicode utilities (7.31).

#### 6.10.10.4 Conditional feature macros

- The following macro names are conditionally defined by the implementation:

**\_\_STDC\_ANALYZABLE\_\_** The integer literal 1, if the implementation conforms to the specifications in Annex L (Analyzability).

**\_\_STDC\_IEC\_60559\_BFP\_\_** The integer literal 202yml, intended to indicate conformance to Annex F (ISO/IEC 60559 floating-point arithmetic) for binary floating-point arithmetic.

**\_\_STDC\_IEC\_559\_\_** The integer literal 1, intended to indicate conformance to the specifications in Annex F (ISO/IEC 60559 floating-point arithmetic) for binary floating-point arithmetic. Use of this macro is an obsolescent feature.

**\_\_STDC\_IEC\_60559\_DFP\_\_** The integer literal 202yml, intended to indicate support of decimal floating types and conformance to Annex F (ISO/IEC 60559 floating-point arithmetic) for decimal floating-point arithmetic.

**\_\_STDC\_IEC\_60559\_COMPLEX\_\_** The integer literal 202yml, intended to indicate conformance to the specifications in Annex G (ISO/IEC 60559 compatible complex arithmetic).

**\_\_STDC\_IEC\_60559\_TYPES\_\_** The integer literal 202yml, intended to indicate conformance to the specification in Annex H (ISO/IEC 60559 interchange and extended types).

**\_\_STDC\_LIB\_EXT1\_\_** The integer literal 202yml, intended to indicate support for the extensions defined in Annex K (Bounds-checking interfaces).

**\_\_STDC\_NO\_ATOMICS\_\_** The integer literal 1, intended to indicate that the implementation does not support atomic types (including the **\_Atomic** type qualifier) and the `<stdatomic.h>` header.

**\_\_STDC\_NO\_COMPLEX\_\_** The integer literal 1, intended to indicate that the implementation does not support complex types or the `<complex.h>` header. An implementation that defines this macro shall not define **\_\_STDC\_IEC\_60559\_COMPLEX\_\_**.

**\_\_STDC\_NO\_THREADS\_\_** The integer literal 1, intended to indicate that the implementation does not support the `<threads.h>` header.

**\_\_STDC\_NO\_VLA\_\_** The integer literal 1, intended to indicate that the implementation does not support variable length arrays with automatic storage duration. Parameters declared with variable length array types are adjusted and then define objects of automatic storage duration with pointer types. Thus, support for such declarations is mandatory.

- NOTE The intention for the macros **\_\_STDC\_LIB\_EXT1\_\_**, **\_\_STDC\_IEC\_60559\_BFP\_\_**, **\_\_STDC\_IEC\_60559\_DFP\_\_**, **\_\_STDC\_IEC\_60559\_COMPLEX\_\_**, and **\_\_STDC\_IEC\_60559\_TYPES\_\_**, with the value 202yml, is that this will remain an integer literal of type **long int** that is increased with each edition of this document.

#### 6.10.11 Pragma operator

##### Semantics

- A unary operator expression of the form:

**\_Pragma ( *string-literal* )**

is processed as follows: The string literal is *destringized* by deleting any encoding prefix, deleting the leading and trailing double-quotes, replacing each escape sequence `\"` by a double-quote, and replacing each escape sequence `\\` by a single backslash. The resulting sequence of characters is processed through translation phase 3 to produce preprocessing tokens that are executed as if they were the *pp-tokens* in a pragma directive. The original four preprocessing tokens in the unary operator expression are removed.

- 2 EXAMPLE A directive of the form:

```
#pragma listing on "..\listing.dir"
```

can also be expressed as:

```
_Pragma ("listing on \"..\listing.dir\"")
```

The latter form is processed in the same way whether it appears literally as shown, or results from macro replacement, as in:

```
#define LISTING(x) PRAGMA(listing on #x)
#define PRAGMA(x) _Pragma(#x)

LISTING (..\listing.dir)
```

## 6.11 Future language directions

### 6.11.1 Floating types

- 1 Future standardization can include additional floating types, including those with greater range, precision, or both than **long double**.

### 6.11.2 Linkages of identifiers

- 1 Declaring an identifier with internal linkage at file scope without the **static** or **constexpr** storage-class specifier is an obsolescent feature.

### 6.11.3 External names

- 1 Restriction of the significance of an external name to fewer than 255 characters (considering each universal character name or extended source character as a single character) is an obsolescent feature that is a concession to existing implementations.

### 6.11.4 Character escape sequences

- 1 Lowercase letters as escape sequences are reserved for future standardization. Other characters may be used in extensions.

### 6.11.5 Octal integer literals

- 1 The use of nonzero octal integer literals without the prefix **0o** or **0O** is an obsolescent feature.

### 6.11.6 Storage-class specifiers

- 1 The placement of a storage-class specifier other than at the beginning of the declaration specifiers in a declaration is an obsolescent feature.
- 2 Future standardization can change the **auto** storage-class specifier to a type specifier.

### 6.11.7 Pragma directives

- 1 Pragmas whose first preprocessing token is **STDC** are reserved for future standardization.

### 6.11.8 Predefined macro names

- 1 Macro names beginning with **\_\_STDC\_\_** are reserved for future standardization.
- 2 Use of the **\_\_STDC\_IEC\_559\_\_** macro is an obsolescent feature.

## 7. Library

### 7.1 Introduction

#### 7.1.1 Definitions of terms

- 1 A *string* is a contiguous sequence of characters terminated by and including the first null character. The term *multibyte string* is sometimes used instead to emphasize special processing given to multibyte characters contained in the string or to avoid confusion with a wide string. A *pointer to a string* is a pointer to its initial (lowest addressed) character. The *length of a string* is the number of bytes preceding the null character and the *value of a string* is the sequence of the values of the contained characters, in order.
- 2 The *decimal-point character* is the character used by functions that convert floating-point numbers to or from character sequences to denote the beginning of the fractional part of such character sequences.<sup>218)</sup> It is represented in the text and examples by a period, but can be changed by the **setlocale** function.
- 3 A *null wide character* is a wide character with code value zero.
- 4 A *wide string* is a contiguous sequence of wide characters terminated by and including the first null wide character. A *pointer to a wide string* is a pointer to its initial (lowest addressed) wide character. The *length of a wide string* is the number of wide characters preceding the null wide character and the *value of a wide string* is the sequence of code values of the contained wide characters, in order.
- 5 A *shift sequence* is a contiguous sequence of bytes within a multibyte string that (potentially) causes a change in shift state (see 5.3.2). A shift sequence shall not have a corresponding wide character; it is instead taken to be an adjunct to an adjacent multibyte character.<sup>219)</sup> In this clause, “*white-space character*” refers to (execution) white-space character as defined by **isspace**. “*White-space wide character*” refers to (execution) white-space wide character as defined by **iswspace**.

**Forward references:** character handling (7.4), the **setlocale** function (7.11.2).

#### 7.1.2 Standard headers

- 1 Each library function is declared in a *header*,<sup>220)</sup> whose contents are made available by the **#include** preprocessing directive. The header declares a set of related functions, plus any types and additional macros needed to facilitate the use of such related functions. In addition to the provisions given in this clause, an implementation that defines **\_\_STDC\_IEC\_60559\_BFP\_\_**, **\_\_STDC\_IEC\_559\_\_**, or **\_\_STDC\_IEC\_60559\_DFP\_\_** shall conform to the specifications in Annex F, one that defines **\_\_STDC\_IEC\_60559\_COMPLEX\_\_** shall conform to the specifications in Annex G, one that defines **\_\_STDC\_IEC\_60559\_TYPES\_\_** shall conform to the specifications in Annex H and one that defines **\_\_STDC\_LIB\_EXT1\_\_** shall conform to the specifications in Annex K, and those Annexes should be read as if they were merged into the parallel structure of named subclauses of this clause. Declarations of types described here, in Annex H, or in Annex K, shall not include type qualifiers, unless explicitly stated otherwise.
- 2 An implementation that does not support decimal floating types (6.10.10.4) may not support interfaces or aspects of interfaces that are specific to these types.
- 3 The standard headers are<sup>221)</sup>

<sup>218)</sup>The functions that make use of the decimal-point character are the numeric conversion functions (7.24.2, 7.32.4.2) and the formatted input/output functions (7.23.6, 7.32.2).

<sup>219)</sup>For state-dependent encodings, the values for **MB\_CUR\_MAX** and **MB\_LEN\_MAX** are thus required to be large enough to count all the bytes in any complete multibyte character plus at least one adjacent shift sequence of maximum length. Whether these counts provide for more than one shift sequence is the implementation's choice.

<sup>220)</sup>A header is not necessarily a source file, nor are the < and > delimited sequences in header names necessarily valid source file names.

<sup>221)</sup>The headers <complex.h>, <stdatomic.h>, and <threads.h> are conditional features that implementations can support but are not required to; see 6.10.10.4.

|              |               |                 |
|--------------|---------------|-----------------|
| <assert.h>   | <setjmp.h>    | <stdlib.h>      |
| <complex.h>  | <signal.h>    | <stdnoreturn.h> |
| <ctype.h>    | <stdalign.h>  | <string.h>      |
| <errno.h>    | <stdarg.h>    | <tgmath.h>      |
| <fenv.h>     | <stdatomic.h> | <threads.h>     |
| <float.h>    | <stdbit.h>    | <time.h>        |
| <inttypes.h> | <stdbool.h>   | <uchar.h>       |
| <iso646.h>   | <stdckdint.h> | <wchar.h>       |
| <limits.h>   | <stddef.h>    | <wctype.h>      |
| <locale.h>   | <stdint.h>    |                 |
| <math.h>     | <stdio.h>     |                 |

- 4 If a file with the same name as one of the preceding entries in < and > delimited sequences, not provided as part of the implementation, is placed in any of the standard places that are searched for included source files, the behavior is undefined.
- 5 Standard headers can be included in any order; each can be included more than once in a given scope, with no effect different from being included only once, except that the effect of including <assert.h> depends on the definition of **NDEBUG** (see 7.2). If used, a header shall be included outside of any external declaration or definition, and it shall first be included before the first reference to any of the functions or objects it declares, or to any of the types or macros it defines. However, if an identifier is declared or defined in more than one header, the second and subsequent associated headers can be included after the initial reference to the identifier. The program shall not have any macros with names lexically identical to keywords currently defined prior to the inclusion of the header or when any macro defined in the header is expanded.
- 6 Some standard headers define or declare identifiers that had not been present in previous versions of this document. To allow implementations and users to adapt to that situation, they also define a version macro for feature test of the form **\_\_STDC\_VERSION\_\_ XXXX\_H\_\_**, where XXXX is the all-caps spelling of the corresponding header <xxxx.h>. These version macros expand to one of many values detailed in the subclauses of Annex M such as 202311L and 202ymmL.
- 7 Any definition of an object-like macro described in this clause or Annex F, Annex G, Annex H, or Annex K shall expand to code that is fully protected by parentheses where necessary, so that it groups in an arbitrary expression as if it were a single identifier.
- 8 The evaluation of the replacement for an object-like macro described in this clause or Annex F, Annex G, Annex H, or Annex K shall not raise a floating-point exception.<sup>222)</sup>
- 9 Any declaration of a library function shall have external linkage.
- 10 A summary of the contents of the standard headers is given in Annex B.

**Forward references:** diagnostics (7.2).

### 7.1.3 Reserved identifiers

- 1 Each header declares or defines all identifiers listed in its associated subclause, and optionally declares or defines identifiers listed in its associated future library directions subclause and identifiers which are always reserved either for any use or for use as file scope identifiers.

— All potentially reserved identifiers (including ones listed in the future library directions) that are provided by an implementation with an external definition are reserved for any use. An implementation shall not provide an external definition of a potentially reserved identifier unless that identifier is reserved for a use where it would have external linkage.<sup>223)</sup> All other potentially reserved identifiers that are provided by an implementation (including in the form of a macro) are reserved for any use when the associated header is included. No other

<sup>222)</sup>As with other evaluations, a floating-point exception could still be raised in the context of an expression that uses that macro replacement.

<sup>223)</sup>All library functions have external linkage.

potentially reserved identifiers are reserved.<sup>224)</sup>

- Each macro name in any of the following subclauses (including the future library directions) is reserved for use as specified if any of its associated headers is included; unless explicitly stated otherwise (see 7.1.4).
- All identifiers with external linkage in any of the following subclauses (including the future library directions) and **errno** are always reserved for use as identifiers with external linkage.<sup>225)</sup>
- Each identifier with file scope listed in any of the following subclauses (including the future library directions) is reserved for use as a macro name and as an identifier with file scope in the same name space if any of its associated headers is included.

#### 7.1.4 Use of library functions

- 1 Each of the following statements applies unless explicitly stated otherwise in the detailed descriptions that follow:

- If an argument to a function has an invalid value (such as a value outside the domain of the function, or a pointer outside the address space of the program, or a null pointer, or a pointer to non-modifiable storage when the corresponding parameter is not const-qualified) or a type (after default argument promotion) not expected by a variadic function, the behavior is undefined.
- If a function argument is described as being an array, the pointer passed to the function shall have a value such that all address computations and accesses to objects (that would be valid if the pointer did point to the first element of such an array) are valid.<sup>226)</sup>
- Any function declared in a header may be additionally implemented as a function-like macro defined in the header, so if a library function is declared explicitly when its header is included, one of the techniques shown later in the next subclause can be used to ensure the declaration is not affected by such a macro. Any macro definition of a function can be suppressed locally by enclosing the name of the function in parentheses, because the name is then not followed by the left parenthesis that indicates expansion of a macro function name. For the same syntactic reason, it is permitted to take the address of a library function even if it is also defined as a macro.<sup>227)</sup> The use of **#undef** to remove any macro definition will also ensure that an actual function is referred to.
- Any invocation of a library function that is implemented as a macro shall expand to code that evaluates each of its arguments exactly once, fully protected by parentheses where necessary, so it is generally safe to use arbitrary expressions as arguments.<sup>228)</sup>
- Likewise, those function-like macros described in the following subclauses may be invoked in an expression anywhere a function with a compatible return type could be called.<sup>229)</sup>

<sup>224)</sup>A potentially reserved identifier becomes a reserved identifier when an implementation begins using it or a future standard reserves it, but is otherwise available for use by the programmer.

<sup>225)</sup>The list of reserved identifiers with external linkage includes **math\_errhandling**, **setjmp**, **va\_copy**, and **va\_end**.

<sup>226)</sup>This includes, for example, passing a valid pointer that points one-past-the-end of an array along with a size of 0, or using any valid pointer with a size of 0. It also applies to functions which take a pointer to the start of array and a size describing that array, where the pointer is the null pointer and the size is exactly 0.

<sup>227)</sup>This means that an implementation is required to provide an actual function for each library function, even if it also provides a macro for that function.

<sup>228)</sup>However, such macros can sometimes not contain the sequence points that the corresponding function calls do.

<sup>229)</sup>Because external identifiers and some macro names beginning with an underscore are reserved, implementations can provide special semantics for such names. For example, the identifier **\_BUILTIN\_abs** can be used to indicate generation of in-line code for the **abs** function. Thus, the appropriate header can specify

```
#define abs(x) _BUILTIN_abs(x)
```

for a compiler whose code generator will accept it.

In this manner, a user desiring to guarantee that a given library function such as **abs** will be a genuine function can write

- All object-like macros listed as expanding to integer constant expressions shall additionally be suitable for use in conditional expression inclusion preprocessing directives.
- 2 Provided that a library function can be declared without reference to any type defined in a header, it is also permissible to declare the function and use it without including its associated header.
- 3 There is a sequence point immediately before a library function returns.
- 4 The functions in the standard library are not guaranteed to be reentrant and may modify objects with static or thread storage duration.<sup>230)</sup>
- 5 Unless explicitly stated otherwise in the detailed descriptions that follow, library functions shall prevent data races as follows: A library function shall not directly or indirectly access objects accessible by threads other than the current thread unless the objects are accessed directly or indirectly via the function's arguments. A library function shall not directly or indirectly modify objects accessible by threads other than the current thread unless the objects are accessed directly or indirectly via the function's non-const arguments.<sup>231)</sup> Implementations may share their own internal objects between threads if the objects are not visible to users and are protected against data races.
- 6 Unless otherwise specified, library functions shall perform all operations solely within the current thread if those operations have effects that are visible to users.<sup>232)</sup>
- 7 EXAMPLE The function `atoi` can be used in any of several ways:

- by use of its associated header (possibly generating a macro expansion)

```
#include <stdlib.h>
const char *str;
/* ... */
i = atoi(str);
```

- by use of its associated header (assuredly generating a true function reference)

```
#include <stdlib.h>
#undef atoi
const char *str;
/* ... */
i = atoi(str);
```

or

```
#include <stdlib.h>
const char *str;
/* ... */
i = (atoi)(str);
```

- by explicit declaration

```
extern int atoi(const char *);
const char *str;
/* ... */
i = atoi(str);
```

---

```
#undef abs
```

whether the implementation's header provides a macro implementation of `abs` or a built-in implementation. The prototype for the function, which precedes and is hidden by any macro definition, is thereby revealed also.

<sup>230)</sup> Thus, a signal handler cannot, in general, call standard library functions.

<sup>231)</sup> This means, for example, that an implementation is not permitted to use a `static` object for internal purposes without synchronization because it can cause a data race even in programs that do not explicitly share objects between threads. Similarly, an implementation of `memcpy` is not permitted to copy bytes beyond the specified length of the destination object and then restore the original values because it can cause a data race if the program shared those bytes between threads.

<sup>232)</sup> This allows implementations to parallelize operations if there are no visible side effects.



## 7.2 Diagnostics <assert.h>

### 7.2.1 General

- 1 The header <assert.h> defines the **assert** and **\_\_STDC\_VERSION\_ASSERT\_H\_\_** macros and refers to another macro,

**NDEBUG**

which is *not* defined by <assert.h>. If **NDEBUG** is defined as a macro name at the point in the source file where <assert.h> is included, the **assert** macro is defined simply as

```
#define assert(...) ((void)0)
```

The **assert** macro is redefined according to the current state of **NDEBUG** each time that <assert.h> is included.

- 2 The **assert** macro shall be implemented as a macro with an ellipsis parameter, not as an actual function. If the macro definition is suppressed to access an actual function, the behavior is undefined.
- 3 NOTE Nevertheless, when **NDEBUG** is not defined, the macro acts as a function taking one parameter as indicated by the prototype as given later in this subclause (7.2.2.1). For both **assert()** and **assert(1, 1)**, the number of arguments does not agree with the number of parameters. A diagnostic is required by 6.5.3.3.
- 4 The macro

**\_\_STDC\_VERSION\_ASSERT\_H\_\_**

is an integer constant expression with a value equivalent to 202311L.

### 7.2.2 Program diagnostics

#### 7.2.2.1 The **assert** macro

##### Synopsis

- 1 

```
#include <assert.h>
void assert(scalar expression);
```

##### Description

- 2 The **assert** macro puts diagnostic tests into programs; it expands to a void expression. When it is executed, if *expression* (which shall have a scalar type) is false (that is, compares equal to 0), the **assert** macro writes information about the particular invocation that failed (including the text of the argument, the name of the source file, the source line number, and the name of the enclosing function — the latter are respectively the values of the preprocessing macros **\_\_FILE\_\_** and **\_\_LINE\_\_** and of the identifier **\_\_func\_\_**) on the standard error stream in an implementation-defined format.<sup>233)</sup>

It then calls the **abort** function.

##### Returns

- 3 The **assert** macro returns no value.

**Forward references:** the **abort** function (7.24.5.1).

<sup>233)</sup>The message written can be of the form:

```
Assertion failed: expression, function abc, file xyz, line nnn.
```



## 7.3 Complex arithmetic <complex.h>

### 7.3.1 Introduction

- 1 The header <complex.h> defines macros and declares functions that support complex arithmetic.<sup>234)</sup>
- 2 Implementations that define the macro `__STDC_NO_COMPLEX__` may not provide this header nor support any of its facilities. The macro `__STDC_VERSION_COMPLEX_H__` is an integer constant expression with a value equivalent to 202ymmL.
- 3 The macro `complex` expands to `_Complex`; the macro `_Complex_I` expands to an arithmetic constant expression of type `float _Complex` with the value of the imaginary unit;<sup>235)</sup> the macro `I` expands to `_Complex_I` or to an arithmetic constant expression of another implementation-defined type with the same value and precision. Notwithstanding the provisions of 7.1.3

- a program may undefine and perhaps then redefine the macros `complex`, and `I`;
- the implementation may define a macro `imaginary` but a program may undefine and perhaps then redefine that identifier.

- 4 Each synopsis, other than for the `CMPLX` macros, specifies a family of functions consisting of a principal function with one or more `double complex` parameters and a `double complex` or `double` return value; and other functions with the same name but with `f` and `l` suffixes which are corresponding functions with `float` and `long double` parameters and return values.

**Forward references:** the `CMPLX` macros (7.3.9.3), ISO/IEC 60559-compatible complex arithmetic (Annex G).

### 7.3.2 Conventions

- 1 Values are interpreted as radians, not degrees. An implementation may set `errno` but is not required to do so.

### 7.3.3 Branch cuts

- 1 Some of the following functions have branch cuts, across which the function is discontinuous. For implementations with a signed zero (including all ISO/IEC 60559 implementations) that follow the specifications of Annex G, the sign of zero distinguishes one side of a cut from another so that the function is continuous (except for format limitations) as the cut is approached from either side. For example, for the square root function, which has a branch cut along the negative real axis, the top of the cut, with imaginary part  $+0$ , maps to the positive imaginary axis, and the bottom of the cut, with imaginary part  $-0$ , maps to the negative imaginary axis.
- 2 Implementations that do not support a signed zero (see Annex F) cannot distinguish the sides of branch cuts. These implementations shall map a cut so that the function is continuous as the cut is approached coming around the finite endpoint of the cut in a counter clockwise direction. (Branch cuts for the functions specified here have just one finite endpoint.) For example, in the square root function, coming counter clockwise around the finite endpoint of the cut along the negative real axis approaches the cut from above, so that the cut maps to the positive imaginary axis.

### 7.3.4 The `CX_LIMITED_RANGE` pragma

#### Synopsis

- 1 

```
#include <complex.h>
#pragma STDC CX_LIMITED_RANGE on-off-switch
```

#### Description

- 2 The usual mathematical formulas for complex multiply, divide, and absolute value are problematic because of their treatment of infinities and because of undue overflow and underflow. The `CX_LIMITED_RANGE` pragma can be used to inform the implementation that (where the state is “on”)

<sup>234)</sup>See “future library directions” (7.34.2).

<sup>235)</sup>The imaginary unit is a number  $i$  such that  $i^2 = -1$ .

the usual mathematical formulas are acceptable.<sup>236)</sup> The pragma can occur either outside external declarations or preceding all explicit declarations and statements inside a compound statement. When outside external declarations, the pragma takes effect from its occurrence until another **CX\_LIMITED\_RANGE** pragma is encountered, or until the end of the translation unit. When inside a compound statement, the pragma takes effect from its occurrence until another **CX\_LIMITED\_RANGE** pragma is encountered (including within a nested compound statement), or until the end of the compound statement; at the end of a compound statement the state for the pragma is restored to its condition just before the compound statement. If this pragma is used in any other context, the behavior is undefined. The default state for the pragma is “off”.

### 7.3.5 Trigonometric functions

#### 7.3.5.1 The **cacos** functions

##### Synopsis

```
1 #include <complex.h>
 double complex cacos(double complex z);
 float complex cacosf(float complex z);
 long double complex cacosl(long double complex z);
```

##### Description

- 2 The **cacos** functions compute the complex arc cosine of **z**, with branch cuts outside the interval  $[-1, +1]$  along the real axis.

##### Returns

- 3 The **cacos** functions return the complex arc cosine value, in the range of a strip mathematically unbounded along the imaginary axis and in the interval  $[0, \pi]$  along the real axis.

#### 7.3.5.2 The **casin** functions

##### Synopsis

```
1 #include <complex.h>
 double complex casin(double complex z);
 float complex casinf(float complex z);
 long double complex casinl(long double complex z);
```

##### Description

- 2 The **casin** functions compute the complex arc sine of **z**, with branch cuts outside the interval  $[-1, +1]$  along the real axis.

##### Returns

- 3 The **casin** functions return the complex arc sine value, in the range of a strip mathematically unbounded along the imaginary axis and in the interval  $[-\frac{\pi}{2}, +\frac{\pi}{2}]$  along the real axis.

#### 7.3.5.3 The **catan** functions

##### Synopsis

```
1 #include <complex.h>
 double complex catan(double complex z);
 float complex catanf(float complex z);
 long double complex catanl(long double complex z);
```

<sup>236)</sup>The purpose of the pragma is to allow the implementation to use the formulas:

$$\begin{aligned}(x + iy) \times (u + iv) &= (xu - yv) + i(yu + xv) \\ (x + iy) / (u + iv) &= [(xu + yv) + i(yu - xv)] / (u^2 + v^2) \\ |x + iy| &= \sqrt{x^2 + y^2}\end{aligned}$$

where the programmer can determine they are safe.

**Description**

- 2 The **catan** functions compute the complex arc tangent of  $z$ , with branch cuts outside the interval  $[-i, +i]$  along the imaginary axis.

**Returns**

- 3 The **catan** functions return the complex arc tangent value, in the range of a strip mathematically unbounded along the imaginary axis and in the interval  $[-\frac{\pi}{2}, +\frac{\pi}{2}]$  along the real axis.

**7.3.5.4 The ccos functions****Synopsis**

```
1 #include <complex.h>
 double complex ccos(double complex z);
 float complex ccosf(float complex z);
 long double complex ccosl(long double complex z);
```

**Description**

- 2 The **ccos** functions compute the complex cosine of  $z$ .

**Returns**

- 3 The **ccos** functions return the complex cosine value.

**7.3.5.5 The csin functions****Synopsis**

```
1 #include <complex.h>
 double complex csin(double complex z);
 float complex csinf(float complex z);
 long double complex csinl(long double complex z);
```

**Description**

- 2 The **csin** functions compute the complex sine of  $z$ .

**Returns**

- 3 The **csin** functions return the complex sine value.

**7.3.5.6 The ctan functions****Synopsis**

```
1 #include <complex.h>
 double complex ctan(double complex z);
 float complex ctanf(float complex z);
 long double complex ctanl(long double complex z);
```

**Description**

- 2 The **ctan** functions compute the complex tangent of  $z$ .

**Returns**

- 3 The **ctan** functions return the complex tangent value.

**7.3.6 Hyperbolic functions****7.3.6.1 The cacosh functions****Synopsis**

```
1 #include <complex.h>
 double complex cacosh(double complex z);
 float complex cacoshf(float complex z);
 long double complex cacoshl(long double complex z);
```

**Description**

- 2 The **cacosh** functions compute the complex arc hyperbolic cosine of **z**, with a branch cut at values less than 1 along the real axis.

**Returns**

- 3 The **cacosh** functions return the complex arc hyperbolic cosine value, in the range of a half-strip of nonnegative values along the real axis and in the interval  $[-i\pi, +i\pi]$  along the imaginary axis.

**7.3.6.2 The **casinh** functions****Synopsis**

```
1 #include <complex.h>
 double complex casinh(double complex z);
 float complex casinhf(float complex z);
 long double complex casinhl(long double complex z);
```

**Description**

- 2 The **casinh** functions compute the complex arc hyperbolic sine of **z**, with branch cuts outside the interval  $[-i, +i]$  along the imaginary axis.

**Returns**

- 3 The **casinh** functions return the complex arc hyperbolic sine value, in the range of a strip mathematically unbounded along the real axis and in the interval  $[-\frac{i\pi}{2}, +\frac{i\pi}{2}]$  along the imaginary axis.

**7.3.6.3 The **catanh** functions****Synopsis**

```
1 #include <complex.h>
 double complex catanh(double complex z);
 float complex catanhf(float complex z);
 long double complex catanhl(long double complex z);
```

**Description**

- 2 The **catanh** functions compute the complex arc hyperbolic tangent of **z**, with branch cuts outside the interval  $[-1, +1]$  along the real axis.

**Returns**

- 3 The **catanh** functions return the complex arc hyperbolic tangent value, in the range of a strip mathematically unbounded along the real axis and in the interval  $[-\frac{i\pi}{2}, +\frac{i\pi}{2}]$  along the imaginary axis.

**7.3.6.4 The **ccosh** functions****Synopsis**

```
1 #include <complex.h>
 double complex ccosh(double complex z);
 float complex ccoshf(float complex z);
 long double complex ccoshl(long double complex z);
```

**Description**

- 2 The **ccosh** functions compute the complex hyperbolic cosine of **z**.

**Returns**

- 3 The **ccosh** functions return the complex hyperbolic cosine value.

**7.3.6.5 The **csinh** functions**

## Synopsis

```

1 #include <complex.h>
 double complex csinh(double complex z);
 float complex csinhf(float complex z);
 long double complex csinhl(long double complex z);

```

## Description

2 The **csinh** functions compute the complex hyperbolic sine of *z*.

## Returns

3 The **csinh** functions return the complex hyperbolic sine value.

### 7.3.6.6 The **ctanh** functions

## Synopsis

```

1 #include <complex.h>
 double complex ctanh(double complex z);
 float complex ctanhf(float complex z);
 long double complex ctanhl(long double complex z);

```

## Description

2 The **ctanh** functions compute the complex hyperbolic tangent of *z*.

## Returns

3 The **ctanh** functions return the complex hyperbolic tangent value.

## 7.3.7 Exponential and logarithmic functions

### 7.3.7.1 The **cexp** functions

## Synopsis

```

1 #include <complex.h>
 double complex cexp(double complex z);
 float complex cexpf(float complex z);
 long double complex cexpl(long double complex z);

```

## Description

2 The **cexp** functions compute the complex base-*e* exponential of *z*.

## Returns

3 The **cexp** functions return the complex base-*e* exponential value.

### 7.3.7.2 The **clog** functions

## Synopsis

```

1 #include <complex.h>
 double complex clog(double complex z);
 float complex clogf(float complex z);
 long double complex clogl(long double complex z);

```

## Description

2 The **clog** functions compute the complex natural (base-*e*) logarithm of *z*, with a branch cut along the negative real axis.

## Returns

3 The **clog** functions return the complex natural logarithm value, in the range of a strip mathematically unbounded along the real axis and in the interval  $[-i\pi, +i\pi]$  along the imaginary axis.

### 7.3.8 Power and absolute-value functions

#### 7.3.8.1 The **cabs** functions

##### Synopsis

```
1 #include <complex.h>
 double cabs(double complex z);
 float cabsf(float complex z);
 long double cabsl(long double complex z);
```

##### Description

- 2 The **cabs** functions compute the complex absolute value (also called norm, modulus, or magnitude) of  $z$ .

##### Returns

- 3 The **cabs** functions return the complex absolute value.

#### 7.3.8.2 The **cpow** functions

##### Synopsis

```
1 #include <complex.h>
 double complex cpow(double complex x, double complex y);
 float complex cpowf(float complex x, float complex y);
 long double complex cpowl(long double complex x, long double complex y);
```

##### Description

- 2 The **cpow** functions compute the complex power function  $x^y$ , with a branch cut for the first parameter along the negative real axis.

##### Returns

- 3 The **cpow** functions return the complex power function value.

#### 7.3.8.3 The **csqrt** functions

##### Synopsis

```
1 #include <complex.h>
 double complex csqrt(double complex z);
 float complex csqrtf(float complex z);
 long double complex csqrtl(long double complex z);
```

##### Description

- 2 The **csqrt** functions compute the complex square root of  $z$ , with a branch cut along the negative real axis.

##### Returns

- 3 The **csqrt** functions return the complex square root value, in the range of the right half-plane (including the imaginary axis).

### 7.3.9 Manipulation functions

#### 7.3.9.1 The **carg** functions

##### Synopsis

```
1 #include <complex.h>
 double carg(double complex z);
 float cargf(float complex z);
 long double cargl(long double complex z);
```

**Description**

- 2 The **carg** functions compute the argument (also called phase (which is an angle)) of  $z$ , with a branch cut along the negative real axis.

**Returns**

- 3 The **carg** functions return the value of the argument in the interval  $[-\pi, +\pi]$ .

**7.3.9.2 The **cimag** functions****Synopsis**

```
1 #include <complex.h>
 double cimag(double complex z);
 float cimagf(float complex z);
 long double cimagl(long double complex z);
```

**Description**

- 2 The **cimag** functions compute the imaginary part of  $z$ .<sup>237)</sup>

**Returns**

- 3 The **cimag** functions return the imaginary part value (as a real).

**7.3.9.3 The **CMPLX** macros****Synopsis**

```
1 #include <complex.h>
 double complex CMPLX(double x, double y);
 float complex CMPLXF(float x, float y);
 long double complex CMPLXL(long double x, long double y);
```

**Description**

- 2 The **CMPLX** macros expand to an expression of the specified complex type, with the real part having the (converted) value of  $x$  and the imaginary part having the (converted) value of  $y$ . The resulting expression shall be suitable for use as an initializer for an object with static or thread storage duration, provided both arguments are likewise suitable. The resulting expression shall be an arithmetic constant expression, provided both arguments are arithmetic constant expressions.

**Returns**

- 3 The **CMPLX** macros return the complex value  $x + iy$ .

**7.3.9.4 The **conj** functions****Synopsis**

```
1 #include <complex.h>
 double complex conj(double complex z);
 float complex conjf(float complex z);
 long double complex conjl(long double complex z);
```

**Description**

- 2 The **conj** functions compute the complex conjugate of  $z$ , by arithmetically negating its imaginary part.

**Returns**

- 3 The **conj** functions return the complex conjugate value.

**7.3.9.5 The **cproj** functions**

<sup>237)</sup>For a complex variable  $z$ ,  $z$  and **CMPLX**(**creal**( $z$ ), **cimag**( $z$ )) are equivalent expressions.

**Synopsis**

```

1 #include <complex.h>
 double complex cproj(double complex z);
 float complex cprojf(float complex z);
 long double complex cprojl(long double complex z);

```

**Description**

- 2 The **cproj** functions compute a projection of **z** onto the Riemann sphere where **z** projects to **z** except that all complex infinities (even those with one infinite part and one NaN part) project to positive infinity on the real axis. If **z** has an infinite part, then **cproj(z)** is equivalent to

```
INFINITY + I * copysign(0.0, cimag(z))
```

**Returns**

- 3 The **cproj** functions return the value of the projection onto the Riemann sphere.

**7.3.9.6 The creal functions****Synopsis**

```

1 #include <complex.h>
 double creal(double complex z);
 float crealf(float complex z);
 long double creall(long double complex z);

```

**Description**

- 2 The **creal** functions compute the real part of **z**.<sup>238)</sup>

**Returns**

- 3 The **creal** functions return the real part value.

<sup>238)</sup>For a complex variable **z**, **z** and **CMPLX(creal(z), cimag(z))** are equivalent expressions.



## 7.4 Character handling <ctype.h>

### 7.4.1 General

- 1 The header <ctype.h> declares several functions useful for classifying and mapping characters.<sup>239)</sup> In all cases the argument is an **int**, the value of which shall be representable as an **unsigned char** or shall equal the value of the macro **E0F**. If the argument has any other value, the behavior is undefined.
- 2 The behavior of these functions is affected by the current locale. Those functions that have locale-specific aspects only when not in the "C" locale are noted subsequently in this subclause.
- 3 The term *printing character* refers to a member of a locale-specific set of characters, each of which occupies one printing position on a display device; the term *control character* refers to a member of a locale-specific set of characters that are not printing characters.<sup>240)</sup> All letters and digits are printing characters.

#### Forward references:

**E0F** (7.23.1), localization (7.11).

### 7.4.2 Character classification functions

#### 7.4.2.1 General

- 1 The functions in this subclause return nonzero (true) if and only if the value of the argument **c** conforms to that in the description of the function.

#### 7.4.2.2 The **isalnum** function

##### Synopsis

```
1 #include <ctype.h>
 int isalnum(int c);
```

##### Description

- 2 The **isalnum** function tests for any character for which **isalpha** or **isdigit** is true.

#### 7.4.2.3 The **isalpha** function

##### Synopsis

```
1 #include <ctype.h>
 int isalpha(int c);
```

##### Description

- 2 The **isalpha** function tests for any character for which **isupper** or **islower** is true, or any character that is one of a locale-specific set of alphabetic characters for which none of **iscntrl**, **isdigit**, **ispunct**, or **isspace** is true.<sup>241)</sup> In the "C" locale, **isalpha** returns true only for the characters for which **isupper** or **islower** is true.

#### 7.4.2.4 The **isblank** function

##### Synopsis

```
1 #include <ctype.h>
 int isblank(int c);
```

<sup>239)</sup>See "future library directions" (7.34.3).

<sup>240)</sup>In an implementation that uses the seven-bit US ASCII character set, the printing characters are those whose values lie from 0x20 (space) through 0x7E (tilde); the control characters are those whose values lie from 0 (NUL) through 0x1F (US), and the character 0x7F (DEL).

<sup>241)</sup>The functions **islower** and **isupper** test true or false separately for each of these additional characters; all four combinations are possible.

**Description**

- 2 The **isblank** function tests for any character that is a standard blank character or is one of a locale-specific set of characters for which **isspace** is true and that is used to separate words within a line of text. The standard blank characters are the following: space (' '), and horizontal tab ('\t'). In the "C" locale, **isblank** returns true only for the standard blank characters.

**7.4.2.5 The iscntrl function****Synopsis**

```
1 #include <ctype.h>
 int iscntrl(int c);
```

**Description**

- 2 The **iscntrl** function tests for any control character.

**7.4.2.6 The isdigit function****Synopsis**

```
1 #include <ctype.h>
 int isdigit(int c);
```

**Description**

- 2 The **isdigit** function tests for any decimal-digit character (as defined in 5.3.1).

**7.4.2.7 The isgraph function****Synopsis**

```
1 #include <ctype.h>
 int isgraph(int c);
```

**Description**

- 2 The **isgraph** function tests for any printing character except space (' ').

**7.4.2.8 The islower function****Synopsis**

```
1 #include <ctype.h>
 int islower(int c);
```

**Description**

- 2 The **islower** function tests for any character that is a lowercase letter or is one of a locale-specific set of characters for which none of **iscntrl**, **isdigit**, **ispunct**, or **isspace** is true. In the "C" locale, **islower** returns true only for the lowercase letters (as defined in 5.3.1).

**7.4.2.9 The isprint function****Synopsis**

```
1 #include <ctype.h>
 int isprint(int c);
```

**Description**

- 2 The **isprint** function tests for any printing character including space (' ').

**7.4.2.10 The ispunct function****Synopsis**

```
1 #include <ctype.h>
 int ispunct(int c);
```

**Description**

- 2 The **ispunct** function tests for any printing character that is one of a locale-specific set of punctuation characters for which neither **isspace** nor **isalnum** is true. In the "C" locale, **ispunct** returns true for every printing character for which neither **isspace** nor **isalnum** is true.

**7.4.2.11 The **isspace** function****Synopsis**

```
1 #include <ctype.h>
 int isspace(int c);
```

**Description**

- 2 The **isspace** function tests for any character that is a standard white-space character or is one of a locale-specific set of characters for which **isalnum** is false. The standard white-space characters are the following: space (' '), form feed ('\f'), new-line ('\n'), carriage return ('\r'), horizontal tab ('\t'), and vertical tab ('\v'). In the "C" locale, **isspace** returns true only for the standard white-space characters.

**7.4.2.12 The **isupper** function****Synopsis**

```
1 #include <ctype.h>
 int isupper(int c);
```

**Description**

- 2 The **isupper** function tests for any character that is an uppercase letter or is one of a locale-specific set of characters for which none of **isctrl**, **isdigit**, **ispunct**, or **isspace** is true. In the "C" locale, **isupper** returns true only for the uppercase letters (as defined in 5.3.1).

**7.4.2.13 The **isxdigit** function****Synopsis**

```
1 #include <ctype.h>
 int isxdigit(int c);
```

**Description**

- 2 The **isxdigit** function tests for any hexadecimal-digit character (as defined in 6.4.5.2).

**7.4.3 Character case mapping functions****7.4.3.1 The **tolower** function****Synopsis**

```
1 #include <ctype.h>
 int tolower(int c);
```

**Description**

- 2 The **tolower** function converts an uppercase letter to a corresponding lowercase letter.

**Returns**

- 3 If the argument is a character for which **isupper** is true and there are one or more corresponding characters, as specified by the current locale, for which **islower** is true, the **tolower** function returns one of the corresponding characters (always the same one for any given locale); otherwise, the argument is returned unchanged.

**7.4.3.2 The **toupper** function**

## Synopsis

```
1 #include <ctype.h>
 int toupper(int c);
```

## Description

- 2 The **toupper** function converts a lowercase letter to a corresponding uppercase letter.

## Returns

- 3 If the argument is a character for which **islower** is true and there are one or more corresponding characters, as specified by the current locale, for which **isupper** is true, the **toupper** function returns one of the corresponding characters (always the same one for any given locale); otherwise, the argument is returned unchanged.

## 7.5 Errors <errno.h>

- 1 The header <errno.h> defines several macros, all relating to the reporting of error conditions.
- 2 The macros are

**EDOM**  
**EILSEQ**  
**ERANGE**

which expand to integer constant expressions with type **int**, distinct positive values, and which are suitable for use in conditional expression inclusion preprocessing directives; and

**errno**

which expands to a modifiable lvalue<sup>242)</sup> that has type **int** and thread storage duration, the value of which is set to a positive error number by several library functions. If a macro definition is suppressed to access an actual object, or a program defines an identifier with the name **errno**, the behavior is undefined.

- 3 The value of **errno** in the initial thread is zero at program startup (the initial representation of the object designated by **errno** in other threads is indeterminate), but is never set to zero by any library function.<sup>243)</sup> The value of **errno** may be set to nonzero by a library function call whether or not there is an error, provided the use of **errno** is not documented in the description of the function in this document.
- 4 Additional macro definitions, beginning with **E** and a digit or **E** and an uppercase letter,<sup>244)</sup> may also be specified by the implementation.

<sup>242)</sup>The macro **errno** is not expected to be the identifier of an object. Expansion to a modifiable lvalue resulting from a function call (for example, `(*errno())`) is a viable implementation strategy.

<sup>243)</sup>Thus, a program that uses **errno** for error checking would set it to zero before a library function call, then inspect it before a subsequent library function call. Of course, a library function can save the value of **errno** on entry and then set it to zero, as long as the original value is restored if **errno**'s value is still zero just before the return.

<sup>244)</sup>See "future library directions" (7.34.4).

## 7.6 Floating-point environment <fenv.h>

### 7.6.1 General

- 1 The header <fenv.h> defines several macros, and declares types and functions that provide access to the floating-point environment. The *floating-point environment* refers collectively to any floating-point status flags and control modes supported by the implementation.<sup>245)</sup>

A *floating-point status flag* is a system variable whose value is set (but never cleared) when a *floating-point exception* is raised, which occurs as a side effect of exceptional floating-point arithmetic to provide auxiliary information.<sup>246)</sup> A *floating-point control mode* is a system variable whose value may be set by the user to affect the subsequent behavior of floating-point arithmetic.

- 2 A floating-point control mode may be *constant* (7.6.3) or *dynamic*. The *dynamic floating-point environment* includes the dynamic floating-point control modes and the floating-point status flags.
- 3 The dynamic floating-point environment has thread storage duration. The initial state for a thread's dynamic floating-point environment is the current state of the dynamic floating-point environment of the thread that creates it. It is initialized at the time of the thread's creation.
- 4 Certain programming conventions support the intended model of use for the dynamic floating-point environment:<sup>247)</sup>

- a function call does not alter its caller's floating-point control modes, clear its caller's floating-point status flags, nor depend on the state of its caller's floating-point status flags unless the function is so documented;
- a function call is assumed to require default floating-point control modes, unless its documentation promises otherwise;
- a function call is assumed to have the potential for raising floating-point exceptions, unless its documentation promises otherwise.

- 5 The feature test macro `__STDC_VERSION_FENV_H__` expands to the token 202311L.

- 6 The type

`fenv_t`

represents the entire dynamic floating-point environment.

- 7 The type

`femode_t`

represents the collection of dynamic floating-point control modes supported by the implementation, including the dynamic rounding direction mode.

- 8 The type

`fexcept_t`

represents the floating-point status flags collectively, including any status the implementation associates with the flags.

- 9 Each of the macros

<sup>245)</sup>This header is designed to support the floating-point exception status flags and rounding-direction control modes required by ISO/IEC 60559, and other similar floating-point state information. It is also designed to facilitate code portability among all systems.

<sup>246)</sup>A floating-point status flag is not an object and can be set more than once within an expression.

<sup>247)</sup>With these conventions, a programmer can safely assume default floating-point control modes (or be unaware of them). The responsibilities associated with accessing the floating-point environment fall on the programmer or program that does so explicitly.

```
FE_DIVBYZERO
FE_INEXACT
FE_INVALID
FE_OVERFLOW
FE_UNDERFLOW
```

is defined if and only if the implementation supports the floating-point exception by means of the functions in 7.6.5.<sup>248)</sup> Additional implementation-defined floating-point exceptions, with macro definitions beginning with **FE\_** and an uppercase letter,<sup>249)</sup> may also be specified by the implementation. The defined macros expand to integer constant expressions with values such that bitwise ORs of all combinations of the macros result in distinct values, and furthermore, bitwise ANDs of all combinations of the macros result in zero.<sup>250)</sup>

- 10 Decimal floating-point operations and ISO/IEC 60559 binary floating-point operations (Annex F) access the same floating-point exception status flags.

- 11 The macro

```
FE_DFL_MODE
```

represents the default state for the collection of dynamic floating-point control modes supported by the implementation – and has type “pointer to const-qualified **femode\_t**”. Additional implementation-defined states for the dynamic mode collection, with macro definitions beginning with **FE\_** and an uppercase letter, and having type “pointer to const-qualified **femode\_t**”, may also be specified by the implementation.

- 12 The macro

```
FE_ALL_EXCEPT
```

is the bitwise OR of all floating-point exception macros defined by the implementation. If no such macros are defined, **FE\_ALL\_EXCEPT** shall be defined as 0.

- 13 Each of the macros

```
FE_DOWNWARD
FE_TONEAREST
FE_TONEARESTFROMZERO
FE_TOWARDZERO
FE_UPWARD
```

is defined if and only if the implementation supports getting and setting the represented rounding direction by means of the **fegetround** and **fesetround** functions. The defined macros expand to integer constant expressions whose values are distinct nonnegative values. Additional implementation-defined rounding directions, with macro definitions beginning with **FE\_** and an uppercase letter,<sup>251)</sup> may also be specified by the implementation.<sup>252)</sup>

- 14 If the implementation supports decimal floating types, each of the macros

```
FE_DEC_DOWNWARD
FE_DEC_TONEAREST
FE_DEC_TONEARESTFROMZERO
FE_DEC_TOWARDZERO
```

<sup>248)</sup> The implementation supports a floating-point exception if there are circumstances where a call to at least one of the functions in 7.6.5, using the macro as the appropriate argument, will succeed. It is not necessary for all the functions to succeed all the time.

<sup>249)</sup> See “future library directions” (7.34.5).

<sup>250)</sup> The macros are typically distinct powers of two.

<sup>251)</sup> See “future library directions” (7.34.5).

<sup>252)</sup> Even though the rounding direction macros can expand to constants corresponding to the values of **FLT\_ROUND**s, they are not required to do so.

**FE\_DEC\_UPWARD**

is defined for use with the **fe\_dec\_getround** and **fe\_dec\_setround** functions for getting and setting the dynamic rounding direction mode for decimal floating-point operations. The decimal rounding direction affects all (inexact) operations that produce a result of decimal floating type and all operations that produce an integer or character sequence result and have an operand of decimal floating type, unless stated otherwise. The macros expand to integer constant expressions whose values are distinct nonnegative values.

- 15 During translation, constant rounding direction modes for decimal floating-point arithmetic are in effect where specified. Elsewhere, during translation the decimal rounding direction mode is **FE\_DEC\_TONEAREST**.
- 16 At program startup the dynamic rounding direction mode for decimal floating-point arithmetic is initialized to **FE\_DEC\_TONEAREST**.
- 17 The macro

**FE\_DFL\_ENV**

represents the default dynamic floating-point environment — the one installed at program startup — and has type “pointer to const-qualified **fenv\_t**”. It can be used as an argument to `<fenv.h>` functions that manage the dynamic floating-point environment.

- 18 Additional implementation-defined environments, with macro definitions beginning with **FE\_** and an uppercase letter,<sup>253)</sup> and having type “pointer to const-qualified **fenv\_t**”, may also be specified by the implementation.

## 7.6.2 The **FENV\_ACCESS** pragma

### Synopsis

- 1 

```
#include <fenv.h>
#pragma STDC FENV_ACCESS on-off-switch
```

### Description

- 2 The **FENV\_ACCESS** pragma provides a means to inform the implementation when a program can access the floating-point environment to test floating-point status flags or run under non-default floating-point control modes.<sup>254)</sup> The pragma shall occur either outside external declarations or preceding all explicit declarations and statements inside a compound statement. When outside external declarations, the pragma takes effect from its occurrence until another **FENV\_ACCESS** pragma is encountered, or until the end of the translation unit. When inside a compound statement, the pragma takes effect from its occurrence until another **FENV\_ACCESS** pragma is encountered (including within a nested compound statement), or until the end of the compound statement. At the end of a compound statement, the state for the pragma is restored to its condition just before the compound statement. If this pragma is used in any other context, the behavior is undefined. If part of a program tests floating-point status flags or establishes or is executed with non-default floating-point mode settings using any means other than the **FENV\_ROUND** pragmas, but was translated with the state for the **FENV\_ACCESS** pragma “off”, the behavior is undefined. The default state (“on” or “off”) for the pragma is implementation-defined. (When execution passes from a part of the program translated with **FENV\_ACCESS** “off” to a part translated with **FENV\_ACCESS** “on”, the state of the floating-point status flags is unspecified and the floating-point control modes have their default settings.)
- 3 EXAMPLE

```
#include <fenv.h>
```

<sup>253)</sup>See “future library directions” (7.34.5).

<sup>254)</sup>The purpose of the **FENV\_ACCESS** pragma is to allow certain optimizations that can subvert flag tests and mode changes (e.g. global common subexpression elimination, code motion, and constant folding). In general, if the state of **FENV\_ACCESS** is “off”, the translator can assume that the flags are not tested, and that default modes are in effect, except where specified otherwise by an **FENV\_ROUND** pragma.



```

void f(double x)
{
 #pragma STDC FENV_ACCESS ON
 void g(double);
 void h(double);
 /* ... */
 g(x + 1);
 h(x + 1);
 /* ... */
}

```

- 4 If the function `g` can depend on status flags set as a side effect of the first `x + 1`, or if the second `x + 1` can depend on control modes set as a side effect of the call to function `g`, then the program has to contain an appropriately placed invocation of `#pragma STDC FENV_ACCESS ON` as shown.<sup>255)</sup>

### 7.6.3 The `FENV_ROUND` pragma

#### Synopsis

```

1 #include <fenv.h>
 #pragma STDC FENV_ROUND direction
 #pragma STDC FENV_ROUND FE_DYNAMIC

```

#### Description

- 2 The `FENV_ROUND` pragma provides a means to specify a constant rounding direction for floating-point operations for standard floating types within a translation unit or compound statement. The pragma shall occur either outside external declarations or before all explicit declarations and statements inside a compound statement. When outside external declarations, the pragma takes effect from its occurrence until another `FENV_ROUND` pragma is encountered, or until the end of the translation unit. When inside a compound statement, the pragma takes effect from its occurrence until another `FENV_ROUND` pragma is encountered (including within a nested compound statement), or until the end of the compound statement; at the end of a compound statement the static rounding mode is restored to its condition just before the compound statement. If this pragma is used in any other context, its behavior is undefined.
- 3 *direction* shall be: one of the names of the supported rounding direction macros for operations for standard floating types (7.6), to specify a constant rounding mode; or, `FE_DYNAMIC`, to specify dynamic rounding. If any other value is specified, the behavior is undefined. If no `FENV_ROUND` pragma is in effect, or the specified *direction* is `FE_DYNAMIC`, rounding is according to the mode specified by the dynamic floating-point environment, which is the dynamic rounding mode that was established either at thread creation or by a call to `fesetround`, `fesetmode`, `fesetenv`, or `feupdateenv`. If the *direction* `FE_DYNAMIC` is specified and `FENV_ACCESS` is “off”, the translator may assume that the default rounding mode is in effect.
- 4 The `FENV_ROUND` pragma affects operations for standard floating types. Within the scope of an `FENV_ROUND` pragma establishing a constant rounding mode, floating-point operators, implicit conversions (including the conversion of a value represented in a format wider than its semantic types to its semantic type, as done by classification macros), and invocations of functions indicated in Table 7.1, for which macro replacement has not been suppressed (7.1.4), shall be evaluated according to the specified constant rounding mode (as though no constant mode was specified and the corresponding dynamic rounding mode had been established by a call to `fesetround`). Invocations of functions for which macro replacement has been suppressed and invocations of functions other than those indicated in Table 7.1 shall not be affected by constant rounding modes – they are affected by (and affect) only the dynamic mode. Floating constants (6.4.5.3) of a standard floating type that occur in the scope of a constant rounding mode shall be interpreted according to that mode.

<sup>255)</sup>The side effects impose a temporal ordering that requires two evaluations of `x + 1`. On the other hand, without the `#pragma STDC FENV_ACCESS ON` pragma, and assuming the default state is “off”, just one evaluation of `x + 1` would suffice.

Table 7.1 — Functions affected by constant rounding modes – for standard floating types

| Header     | Function families                                                              |
|------------|--------------------------------------------------------------------------------|
| <math.h>   | <b>acos, acospi, asin, asinpi, atan, atan2, atan2pi, atanpi</b>                |
| <math.h>   | <b>cos, cospi, sin, sinpi, tan, tanpi</b>                                      |
| <math.h>   | <b>acosh, asinh, atanh</b>                                                     |
| <math.h>   | <b>cosh, sinh, tanh</b>                                                        |
| <math.h>   | <b>exp, exp10, exp10m1, exp2, exp2m1, expm1</b>                                |
| <math.h>   | <b>log, log10, log10p1, log1p, log2, log2p1, logp1</b>                         |
| <math.h>   | <b>scalbn, scalbln, ldexp</b>                                                  |
| <math.h>   | <b>cbrt, compoundn, hypot, pow, pown, powr, rootn, rsqrt, sqrt</b>             |
| <math.h>   | <b>erf, erfc</b>                                                               |
| <math.h>   | <b>lgamma, tgamma</b>                                                          |
| <math.h>   | <b>rint, nearbyint, lrint, llrint</b>                                          |
| <math.h>   | <b>fdim</b>                                                                    |
| <math.h>   | <b>fma</b>                                                                     |
| <math.h>   | <b>fadd, dadd, fsub, dsub, fmul, dmul, fdiv, ddiv, fma, dfma, fsqrt, dsqrt</b> |
| <stdlib.h> | <b>atof, strfrom, strt0</b>                                                    |
| <wchar.h>  | <b>wcsto</b>                                                                   |
| <stdio.h>  | <b>printf</b> and <b>scanf</b> families                                        |
| <wchar.h>  | <b>wprintf</b> and <b>wscanf</b> families                                      |

A function family listed in Table 7.1 indicates the functions for all standard floating types, where the function family is represented by the name of the functions without a suffix. For example, **acos** indicates the functions **acos**, **acosf**, and **acosl**.

- 5 NOTE Constant rounding modes can be implemented using dynamic rounding modes as illustrated in the following example, except that this method does not interpret inexact floating literals according to the constant rounding mode as required.

```

{
 #pragma STDC FENV_ROUND direction
 // compiler inserts:
 // #pragma STDC FENV_ACCESS ON
 // int __savedrnd;
 // __savedrnd = __swapround(direction);
 ... operations affected by constant rounding mode ...
 // compiler inserts:
 // __savedrnd = __swapround(__savedrnd);
 ... operations not affected by constant rounding mode ...
 // compiler inserts:
 // __savedrnd = __swapround(__savedrnd);
 ... operations affected by constant rounding mode ...
 // compiler inserts:
 // __swapround(__savedrnd);

```

```
}
```

where `__swapround` is defined by:

```
static inline int __swapround(const int new) {
 const int old = fegetround();
 fesetround(new);
 return old;
}
```

## 7.6.4 The `FENV_DEC_ROUND` pragma

### Synopsis

```
1 #include <fenv.h>
 #ifdef __STDC_IEC_60559_DFP__
 #pragma STDC FENV_DEC_ROUND dec-direction
 #endif
```

### Description

- 2 The `FENV_DEC_ROUND` pragma is a decimal floating-point analog of the `FENV_ROUND` pragma. If `FLT_RADIX` is not 10, the `FENV_DEC_ROUND` pragma affects operators, functions, and floating literals only for decimal floating types. The affected functions are listed in Table 7.2. If `FLT_RADIX` is 10, whether the `FENV_ROUND` and `FENV_DEC_ROUND` pragmas alter the rounding direction of both standard and decimal floating-point operations is implementation-defined. *dec-direction* shall be one of the decimal rounding direction macro names (`FE_DEC_DOWNWARD`, `FE_DEC_TONEAREST`, `FE_DEC_TONEARESTFROMZERO`, `FE_DEC_TOWARDZERO`, and `FE_DEC_UPWARD`) defined in 7.6, to specify a constant rounding mode, or `FE_DEC_DYNAMIC`, to specify dynamic rounding. The corresponding dynamic rounding mode can be established by a call to `fe_dec_setround`.

Table 7.2 — Functions affected by constant rounding modes – for decimal floating types

| Header   | Function families                                                                                                                                                                        |
|----------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <math.h> | <code>acos</code> , <code>acospi</code> , <code>asin</code> , <code>asinpi</code> , <code>atan</code> , <code>atan2</code> , <code>atan2pi</code> , <code>atanpi</code>                  |
| <math.h> | <code>cos</code> , <code>cospi</code> , <code>sin</code> , <code>sinpi</code> , <code>tan</code> , <code>tanpi</code>                                                                    |
| <math.h> | <code>acosh</code> , <code>asinh</code> , <code>atanh</code>                                                                                                                             |
| <math.h> | <code>cosh</code> , <code>sinh</code> , <code>tanh</code>                                                                                                                                |
| <math.h> | <code>exp</code> , <code>exp10</code> , <code>exp10m1</code> , <code>exp2</code> , <code>exp2m1</code> , <code>expm1</code>                                                              |
| <math.h> | <code>log</code> , <code>log10</code> , <code>log10p1</code> , <code>log1p</code> , <code>log2</code> , <code>log2p1</code> , <code>logp1</code>                                         |
| <math.h> | <code>scalbn</code> , <code>scalbln</code> , <code>ldexp</code>                                                                                                                          |
| <math.h> | <code>cbrt</code> , <code>compoundn</code> , <code>hypot</code> , <code>pow</code> , <code>pown</code> , <code>powr</code> , <code>rootn</code> , <code>rsqrt</code> , <code>sqrt</code> |
| <math.h> | <code>erf</code> , <code>erfc</code>                                                                                                                                                     |
| <math.h> | <code>lgamma</code> , <code>tgamma</code>                                                                                                                                                |
| <math.h> | <code>rint</code> , <code>nearbyint</code> , <code>lrint</code> , <code>llrint</code>                                                                                                    |
| <math.h> | <code>quantize</code>                                                                                                                                                                    |
| <math.h> | <code>fdim</code>                                                                                                                                                                        |
| <math.h> | <code>fma</code>                                                                                                                                                                         |

|            |                                                                                                         |
|------------|---------------------------------------------------------------------------------------------------------|
| <math.h>   | <b>d32add, d64add, d32sub, d64sub, d32mul, d64mul, d32div, d64div, d32fma, d64fma, d32sqrt, d64sqrt</b> |
| <stdlib.h> | <b>strfrom, strt</b>                                                                                    |
| <wchar.h>  | <b>wcsto</b>                                                                                            |
| <stdio.h>  | <b>printf</b> and <b>scanf</b> families                                                                 |
| <wchar.h>  | <b>wprintf</b> and <b>wscanf</b> families                                                               |

A function family listed in Table 7.2 indicates the functions for all decimal floating types, where the function family is represented by the name of the functions without a suffix. For example, **acos** indicates the functions **acosd32**, **acosd64**, and **acosd128**.

## 7.6.5 Floating-point exceptions

### 7.6.5.1 General

- The following functions provide access to the floating-point status flags.<sup>256)</sup> The **int** input argument for the functions represents a subset of floating-point exceptions, and can be zero or the bitwise OR of one or more floating-point exception macros, for example **FE\_OVERFLOW** | **FE\_INEXACT**. For other argument values, the behavior of these functions is undefined.

### 7.6.5.2 The **feclearexcept** function

#### Synopsis

```
1 #include <fenv.h>
 int feclearexcept(int excepts);
```

#### Description

- The **feclearexcept** function attempts to clear the supported floating-point exceptions represented by its argument.

#### Returns

- The **feclearexcept** function returns zero if the **excepts** argument is zero or if all the specified exceptions were successfully cleared. Otherwise, it returns a nonzero value.

### 7.6.5.3 The **fegetexceptflag** function

#### Synopsis

```
1 #include <fenv.h>
 int fegetexceptflag(fexcept_t *flagp, int excepts);
```

#### Description

- The **fegetexceptflag** function attempts to store an implementation-defined representation of the states of the floating-point status flags indicated by the argument **excepts** in the object pointed to by the argument **flagp**.

#### Returns

- The **fegetexceptflag** function returns zero if the representation was successfully stored. Otherwise, it returns a nonzero value.

### 7.6.5.4 The **feraiseexcept** function

<sup>256)</sup>The functions **fetestexcept**, **feraiseexcept**, and **feclearexcept** support the basic abstraction of flags that are either set or clear. An implementation can endow floating-point status flags with more information — for example, the address of the code which first raised the floating-point exception; the functions **fegetexceptflag** and **fesetexceptflag** deal with the full content of flags.

## Synopsis

```
1 #include <fenv.h>
 int feraiseexcept(int excepts);
```

## Description

- 2 The **feraiseexcept** function attempts to raise the supported floating-point exceptions represented by its argument.<sup>257)</sup> The order in which these floating-point exceptions are raised is unspecified, except as stated in F.8.7. Whether the **feraiseexcept** function additionally raises the “inexact” floating-point exception whenever it raises the “overflow” or “underflow” floating-point exception is implementation-defined.

## Returns

- 3 The **feraiseexcept** function returns zero if the **excepts** argument is zero or if all the specified exceptions were successfully raised. Otherwise, it returns a nonzero value.

## Recommended practice

Implementation extensions associated with raising a floating-point exception (for example, enabled traps or ISO/IEC 60559 alternate exception handling) should be honored by this function.

### 7.6.5.5 The **fesetexcept** function

#### Synopsis

```
1 #include <fenv.h>
 int fesetexcept(int excepts);
```

## Description

- 2 The **fesetexcept** function attempts to set the supported floating-point exception flags represented by its argument. This function does not clear any floating-point exception flags. This function changes the state of the floating-point exception flags, but does not cause any other side effects that can be associated with raising floating-point exceptions.<sup>258)</sup>

## Returns

- 3 The **fesetexcept** function returns zero if all the specified exceptions were successfully set or if the **excepts** argument is zero. Otherwise, it returns a nonzero value.

### 7.6.5.6 The **fesetexceptflag** function

#### Synopsis

```
1 #include <fenv.h>
 int fesetexceptflag(const fexcept_t *flagp, int excepts);
```

## Description

- 2 The **fesetexceptflag** function attempts to set the floating-point status flags indicated by the argument **excepts** to the states stored in the object pointed to by **flagp**. The value of **\*flagp** shall have been set by a previous call to **fegetexceptflag** whose second argument represented at least those floating-point exceptions represented by the argument **excepts**. Like **fesetexcept**, this function does not raise floating-point exceptions, but only sets the state of the flags.

## Returns

- 3 The **fesetexceptflag** function returns zero if the **excepts** argument is zero or if all the specified flags were successfully set to the appropriate state. Otherwise, it returns a nonzero value.

<sup>257)</sup>The effect is intended to be similar to that of floating-point exceptions raised by arithmetic operations. Hence, implementation extensions associated with raising a floating-point exception (for example, enabled traps or ISO/IEC 60559 alternate exception handling) should be honored. The specification in F.8.7 is in the same spirit.

<sup>258)</sup>Implementation extensions like traps for floating-point exceptions and ISO/IEC 60559 exception handling do not occur.

### 7.6.5.7 The **fetestexceptflag** function

#### Synopsis

```
1 #include <fenv.h>
 int fetestexceptflag(const fexcept_t *flagp, int excepts);
```

#### Description

- 2 The **fetestexceptflag** function determines which of a specified subset of the floating-point exception flags are set in the object pointed to by **flagp**. The value of **\*flagp** shall have been set by a previous call to **fegetexceptflag** whose second argument represented at least those floating-point exceptions represented by the argument **excepts**. The **excepts** argument specifies the floating-point status flags to be queried.

#### Returns

- 3 The **fetestexceptflag** function returns the value of the bitwise OR of the floating-point exception macros included in **excepts** corresponding to the floating-point exceptions set in **\*flagp**.

### 7.6.5.8 The **fetestexcept** function

#### Synopsis

```
1 #include <fenv.h>
 int fetestexcept(int excepts);
```

#### Description

- 2 The **fetestexcept** function determines which of a specified subset of the floating-point exception flags are currently set. The **excepts** argument specifies the floating-point status flags to be queried.<sup>259)</sup>

#### Returns

- 3 The **fetestexcept** function returns the value of the bitwise OR of the floating-point exception macros corresponding to the currently set floating-point exceptions included in **excepts**.
- 4 EXAMPLE Call **f** if “invalid” is set, then **g** if “overflow” is set:

```
#include <fenv.h>
/* ... */
{
 #pragma STDC FENV_ACCESS ON
 int set_excepts;
 feclearexcept(FE_INVALID | FE_OVERFLOW);
 // maybe raise exceptions
 set_excepts = fetestexcept(FE_INVALID | FE_OVERFLOW);
 if (set_excepts & FE_INVALID) f();
 if (set_excepts & FE_OVERFLOW) g();
 /* ... */
}
```

## 7.6.6 Rounding and other control modes

### 7.6.6.1 General

- 1 The **fegetround** and **fesetround** functions provide control of rounding direction modes. The **fegetmode** and **fesetmode** functions manage all the implementation’s dynamic floating-point control modes collectively.

### 7.6.6.2 The **fegetmode** function

#### Synopsis

- 1 <sup>259)</sup>This mechanism allows testing several floating-point exceptions with just one function call.

```
#include <fenv.h>
int fegetmode(femode_t *modep);
```

**Description**

- 2 The **fegetmode** function attempts to store all the dynamic floating-point control modes in the object pointed to by `modep`.

**Returns**

- 3 The **fegetmode** function returns zero if the modes were successfully stored. Otherwise, it returns a nonzero value.

**7.6.6.3 The fegetround function****Synopsis**

```
1 #include <fenv.h>
 int fegetround(void);
```

**Description**

- 2 The **fegetround** function gets the current value of the dynamic rounding direction mode.

**Returns**

- 3 The **fegetround** function returns the value of the rounding direction macro representing the current dynamic rounding direction or a negative value if there is no such rounding direction macro or the current dynamic rounding direction is not determinable.

**7.6.6.4 The fe\_dec\_getround function****Synopsis**

```
1 #include <fenv.h>
 #ifdef __STDC_IEC_60559_DFP__
 int fe_dec_getround(void);
 #endif
```

**Description**

- 2 The **fe\_dec\_getround** function gets the current value of the dynamic rounding direction mode for decimal floating-point operations.

**Returns**

- 3 The **fe\_dec\_getround** function returns the value of the rounding direction macro representing the current dynamic rounding direction for decimal floating-point operations, or a negative value if there is no such rounding macro or the current rounding direction is not determinable.

**7.6.6.5 The fesetmode function****Synopsis**

```
1 #include <fenv.h>
 int fesetmode(const femode_t *modep);
```

**Description**

- 2 The **fesetmode** function attempts to establish the dynamic floating-point modes represented by the object pointed to by `modep`. The argument `modep` shall point to an object set by a call to **fegetmode**, or equal **FE\_DFL\_MODE** or a dynamic floating-point mode state macro defined by the implementation.

**Returns**

The **fesetmode** function returns zero if the modes were successfully established. Otherwise, it returns a nonzero value.

### 7.6.6.6 The **fesetround** function

#### Synopsis

```
1 #include <fenv.h>
 int fesetround(int rnd);
```

#### Description

- 2 The **fesetround** function establishes the rounding direction represented by its argument **rnd**. If the argument is not equal to the value of a rounding direction macro, the rounding direction is not changed.

#### Returns

- 3 The **fesetround** function returns zero if and only if the dynamic rounding direction mode was set to the requested rounding direction.
- 4 **EXAMPLE** The following is a way to save, set, and restore the rounding direction, including reporting an error and abort if setting the rounding direction fails.

```
#include <fenv.h>
#include <assert.h>

void f(int rnd_dir)
{
 #pragma STDC FENV_ACCESS ON
 int save_round;
 int setround_ok;
 save_round = fegetround();
 setround_ok = fesetround(rnd_dir);
 assert(setround_ok == 0);
 /* ... */
 fesetround(save_round);
 /* ... */
}
```

### 7.6.6.7 The **fe\_dec\_setround** function

#### Synopsis

```
1 #include <fenv.h>
 #ifdef __STDC_IEC_60559_DFP__
 int fe_dec_setround(int rnd);
 #endif
```

#### Description

- 2 The **fe\_dec\_setround** function sets the dynamic rounding direction mode for decimal floating-point operations to be the rounding direction represented by its argument **rnd**. If the argument is not equal to the value of a decimal rounding direction macro, the rounding direction is not changed.
- 3 If **FLT\_RADIX** is not 10, the rounding direction altered by the **fesetround** function is independent of the rounding direction altered by the **fe\_dec\_setround** function; otherwise if **FLT\_RADIX** is 10, whether the **fesetround** and **fe\_dec\_setround** functions alter the rounding direction of both standard and decimal floating-point operations is implementation-defined.

#### Returns

- 4 The **fe\_dec\_setround** function returns a zero value if and only if the argument is equal to a decimal rounding direction macro (that is, if and only if the dynamic rounding direction mode for decimal floating-point operations was set to the requested rounding direction).



## 7.6.7 Environment

### 7.6.7.1 General

- 1 The functions in this subclause manage the floating-point environment — status flags and control modes — as one entity.

### 7.6.7.2 The **fegetenv** function

#### Synopsis

```
1 #include <fenv.h>
 int fegetenv(fenv_t *envp);
```

#### Description

- 2 The **fegetenv** function attempts to store the current dynamic floating-point environment in the object pointed to by **envp**.

#### Returns

- 3 The **fegetenv** function returns zero if the environment was successfully stored. Otherwise, it returns a nonzero value.

### 7.6.7.3 The **feholdexcept** function

#### Synopsis

```
1 #include <fenv.h>
 int feholdexcept(fenv_t *envp);
```

#### Description

- 2 The **feholdexcept** function saves the current dynamic floating-point environment in the object pointed to by **envp**, clears the floating-point status flags, and then installs a *non-stop* (continue on floating-point exceptions) mode, if available, for all floating-point exceptions.<sup>260)</sup>

#### Returns

- 3 The **feholdexcept** function returns zero if and only if non-stop floating-point exception handling was successfully installed.

### 7.6.7.4 The **fesetenv** function

#### Synopsis

```
1 #include <fenv.h>
 int fesetenv(const fenv_t *envp);
```

#### Description

- 2 The **fesetenv** function attempts to establish the dynamic floating-point environment represented by the object pointed to by **envp**. The argument **envp** shall point to an object set by a call to **fegetenv** or **feholdexcept**, or equal a dynamic floating-point environment macro. **fesetenv** merely installs the state of the floating-point status flags represented through its argument, and does not raise these floating-point exceptions.

#### Returns

- 3 The **fesetenv** function returns zero if the environment was successfully established. Otherwise, it returns a nonzero value.

### 7.6.7.5 The **feupdateenv** function

<sup>260)</sup>ISO/IEC 60559 systems have a default non-stop mode, and typically at least one other mode for trap handling or aborting; if the system provides only the non-stop mode then installing it is trivial. For such systems, the **feholdexcept** function can be used in conjunction with the **feupdateenv** function to write routines that hide spurious floating-point exceptions from their callers.

## Synopsis

```
1 #include <fenv.h>
 int feupdateenv(const fenv_t *envp);
```

## Description

- 2 The **feupdateenv** function attempts to save the currently raised floating-point exceptions in its automatic storage, install the dynamic floating-point environment represented by the object pointed to by **envp**, and then raise the saved floating-point exceptions. The argument **envp** shall point to an object set by a call to **feholdexcept** or **fegetenv**, or equal a dynamic floating-point environment macro.

## Returns

- 3 The **feupdateenv** function returns zero if all the actions were successfully carried out. Otherwise, it returns a nonzero value.
- 4 EXAMPLE Hide spurious underflow floating-point exceptions:

```
#include <fenv.h>
double f(double x)
{
 #pragma STDC FENV_ACCESS ON
 double result;
 fenv_t save_env;
 if (feholdexcept(&save_env))
 return /* indication of an environmental problem */;
 // compute result
 if (/* test spurious underflow */)
 if (feclearexcept(FE_UNDERFLOW))
 return /* indication of an environmental problem */;
 if (feupdateenv(&save_env))
 return /* indication of an environmental problem */;
 return result;
}
```

## 7.7 Characteristics of floating types <float.h>

- 1 The header <float.h> defines several macros that expand to various limits and parameters of the real floating types.
- 2 The macro

**\_\_STDC\_VERSION\_FLOAT\_H\_\_**

is an integer constant expression with a value equivalent to 202311L.

- 3 The rest of the macros, their meanings, and the constraints (or restrictions) on their values are listed in 5.3.5.3.3 and 5.3.5.3.4. A summary is given in Annex E.

## 7.8 Format conversion of integer types <inttypes.h>

### 7.8.1 General

- 1 The header <inttypes.h> includes the header <stdint.h> and extends it with additional facilities provided by hosted implementations.
- 2 It defines the macro

```
__STDC_VERSION_INTTYPES_H__
```

which is an integer constant expression with a value equivalent to 202311L.

- 3 It declares functions for manipulating greatest-width integers and converting numeric character strings to greatest-width integers, and it declares the type

```
imaxdiv_t
```

which is a structure type that is the type of the value returned by the `imaxdiv` function. For each type declared in <stdint.h>, it defines corresponding macros for conversion specifiers for use with the formatted input/output functions.<sup>261)</sup>

**Forward references:** integer types <stdint.h> (7.22), formatted input/output functions (7.23.6), formatted wide character input/output functions (7.32.2).

### 7.8.2 Macros for format specifiers

- 1 Each of the following object-like macros expands to an ordinary string literal containing a conversion specifier, possibly modified by a length modifier, suitable for use within the format argument of a formatted input/output function when converting the corresponding integer type. These macro names have the general form of **PRI** (ordinary string literals for the `fprintf` and `fwprintf` family) or **SCN** (ordinary string literals for the `scanf` and `fwscanf` family),<sup>262)</sup> followed by the conversion specifier, followed by a name corresponding to a similar type name in 7.22.2. In these names, *N* represents the width of the type as described in 7.22.2. For example, **PRIdFAST32** can be used in a format string to print the value of an integer of type `int_fast32_t`. The functions in the `fprintf` and `fwprintf` families shall behave as if they use `va_arg` with a type argument naming the type resulting from applying the default argument promotions to the type corresponding to the macro and then convert the result of the `va_arg` expansion to the type corresponding to the macro.
- 2 The `fprintf` macros for signed integers are:

```
PRIdN PRIdLEASTN PRIdFASTN PRIdMAX PRIdPTR
PRIiN PRIiLEASTN PRIiFASTN PRIiMAX PRIiPTR
```

- 3 The `fprintf` macros for unsigned integers are:

```
PRIBN PRIBLEASTN PRIBFASTN PRIBMAX PRIBPTR
PRIoN PRIoLEASTN PRIoFASTN PRIoMAX PRIoPTR
PRIuN PRIuLEASTN PRIuFASTN PRIuMAX PRIuPTR
PRIXN PRIXLEASTN PRIXFASTN PRIXMAX PRIXPTR
```

- 4 The following `fprintf` macros for unsigned integer types are optional:

```
PRIBN PRIBLEASTN PRIBFASTN PRIBMAX PRIBPTR
```

They shall be defined if the implementation supports the **B** specifier as indicated in 7.23.6.2 and 7.32.2.2; otherwise they shall not be defined.

- 5 The `scanf` macros for signed integers are:

```
SCNdN SCNdLEASTN SCNdFASTN SCNdMAX SCNdPTR
SCNiN SCNiLEASTN SCNiFASTN SCNiMAX SCNiPTR
```

<sup>261)</sup>See “future library directions” (7.34.7).

<sup>262)</sup>For any given type, the corresponding macros for `fprintf` and `scanf` functions can be distinct.

- 6 The **fscanf** macros for unsigned integers are:

```
SCNbN SCNbLEASTN SCNbFASTN SCNbMAX SCNbPTR
SCNoN SCNoLEASTN SCNoFASTN SCNoMAX SCNoPTR
SCNuN SCNuLEASTN SCNuFASTN SCNuMAX SCNuPTR
SCNxN SCNxLEASTN SCNxFASTN SCNxMAX SCNxPTR
```

- 7 For each type that the implementation provides in `<stdint.h>`, the corresponding **fprintf** macros shall be defined and the corresponding **fscanf** macros shall be defined unless the implementation does not have a suitable **fscanf** length modifier for the type.
- 8 EXAMPLE

```
#include <inttypes.h>
#include <wchar.h>
int main(void)
{
 uintmax_t i = UINTMAX_MAX; // this type always exists
 wprintf(L"The largest integer value is %020"
 PRIxMAX "\n", i);
 return 0;
}
```

## 7.8.3 Functions for greatest-width integer types

### 7.8.3.1 The **imaxabs** and **uimaxabs** functions

#### Synopsis

```
1 #include <inttypes.h>
 intmax_t imaxabs(intmax_t j);
 uintmax_t uimaxabs(intmax_t j);
```

#### Description

- 2 The **imaxabs** and **uimaxabs** functions compute the absolute value of an integer `j`. If the result cannot be represented, the behavior is undefined.<sup>263)</sup>

#### Returns

- 3 The **imaxabs** and **uimaxabs** functions return the absolute value.

### 7.8.3.2 The **imaxdiv** function

#### Synopsis

```
1 #include <inttypes.h>
 imaxdiv_t imaxdiv(intmax_t numer, intmax_t denom);
```

#### Description

- 2 The **imaxdiv** function computes `numer / denom` and `numer % denom` in a single operation.

#### Returns

- 3 The **imaxdiv** function returns a structure of type **imaxdiv\_t** comprising both the quotient and the remainder. The structure shall contain (in either order) the members **quot** (the quotient) and **rem** (the remainder), each of which has type **intmax\_t**. If either part of the result cannot be represented, the behavior is undefined.

### 7.8.3.3 The **strtoimax** and **strtoumax** functions

#### Synopsis

```
1 #include <inttypes.h>
 intmax_t strtoimax(const char * restrict nptr, char ** restrict endptr, int base);
```

<sup>263)</sup>Only the absolute value of the most negative number is not representable for **imaxabs**.

```
uintmax_t strtoumax(const char * restrict nptr, char ** restrict endptr, int base);
```

### Description

- 2 The **strtoumax** and **strtoumax** functions are equivalent to the **strtol**, **strtoll**, **strtoul**, and **strtoull** functions, except that the initial portion of the string is converted to **intmax\_t** and **uintmax\_t** representation, respectively.

### Returns

- 3 The **strtoumax** and **strtoumax** functions return the converted value, if any. If no conversion could be performed, zero is returned. If the correct value is outside the range of representable values, **INTMAX\_MAX**, **INTMAX\_MIN**, or **UINTMAX\_MAX** is returned (according to the return type and sign of the value, if any), and the value of the macro **ERANGE** is stored in **errno**.

Forward references: the **strtol**, **strtoll**, **strtoul**, and **strtoull** functions (7.24.2.8).

### 7.8.3.4 The **wcstoi** and **wcstou** functions

#### Synopsis

```
1 #include <stddef.h> // for wchar_t
 #include <inttypes.h>
 intmax_t wcstoi(const wchar_t *restrict nptr, wchar_t **restrict endptr, int base);
 uintmax_t wcstou(const wchar_t *restrict nptr, wchar_t **restrict endptr, int base);
```

### Description

- 2 The **wcstoi** and **wcstou** functions are equivalent to the **wcstol**, **wcstoll**, **wcstoul**, and **wcstoull** functions except that the initial portion of the wide string is converted to **intmax\_t** and **uintmax\_t** representation, respectively.

### Returns

- 3 The **wcstoi** function returns the converted value, if any. If no conversion could be performed, zero is returned. If the correct value is outside the range of representable values, **INTMAX\_MAX**, **INTMAX\_MIN**, or **UINTMAX\_MAX** is returned (according to the return type and sign of the value, if any), and the value of the macro **ERANGE** is stored in **errno**.

Forward references: the **wcstol**, **wcstoll**, **wcstoul**, and **wcstoull** functions (7.32.4.2.4).

## 7.9 Alternative spellings <iso646.h>

- 1 The header <iso646.h> defines the following eleven macros (on the left) that expand to the corresponding tokens (on the right):

|               |    |
|---------------|----|
| <b>and</b>    | && |
| <b>and_eq</b> | &= |
| <b>bitand</b> | &  |
| <b>bitor</b>  |    |
| <b>compl</b>  | ~  |
| <b>not</b>    | !  |
| <b>not_eq</b> | != |
| <b>or</b>     |    |
| <b>or_eq</b>  | =  |
| <b>xor</b>    | ^  |
| <b>xor_eq</b> | ^= |

## 7.10 Characteristics of integer types <limits.h>

- 1 The header <limits.h> defines several macros that expand to various limits and parameters of the standard integer types.
- 2 The macro

`__STDC_VERSION_LIMITS_H__`

is an integer constant expression with a value equivalent to 202311L.

- 3 The rest of the macros, their meanings, and the constraints (or restrictions) on their values are listed in 5.3.5.3.2. A summary is given in Annex E.



## 7.11 Localization <locale.h>

### 7.11.1 General

- 1 The header <locale.h> declares two functions, one type, and defines several macros.
- 2 The type is

```
struct lconv
```

which contains members related to the formatting of numeric values. The structure shall contain at least the following members, in any order. The semantics of the members and their normal ranges are explained in 7.11.3.1. In the "C" locale, the members shall have the values specified in the comments.

```
char *decimal_point; // "."
char *thousands_sep; // ""
char *grouping; // ""
char *mon_decimal_point; // ""
char *mon_thousands_sep; // ""
char *mon_grouping; // ""
char *positive_sign; // ""
char *negative_sign; // ""
char *currency_symbol; // ""
char frac_digits; // CHAR_MAX
char p_cs_precedes; // CHAR_MAX
char n_cs_precedes; // CHAR_MAX
char p_sep_by_space; // CHAR_MAX
char n_sep_by_space; // CHAR_MAX
char p_sign_posn; // CHAR_MAX
char n_sign_posn; // CHAR_MAX
char *int_curr_symbol; // ""
char int_frac_digits; // CHAR_MAX
char int_p_cs_precedes; // CHAR_MAX
char int_n_cs_precedes; // CHAR_MAX
char int_p_sep_by_space; // CHAR_MAX
char int_n_sep_by_space; // CHAR_MAX
char int_p_sign_posn; // CHAR_MAX
char int_n_sign_posn; // CHAR_MAX
```

- 3 The macros defined are **NULL** (described in 7.21); and

```
LC_ALL
LC_COLLATE
LC_CTYPE
LC_MONETARY
LC_NUMERIC
LC_TIME
```

which expand to integer constant expressions with distinct values, suitable for use as the first argument to the **setlocale** function.<sup>264)</sup> Additional macro definitions, beginning with the characters **LC\_** and an uppercase letter,<sup>265)</sup> may also be specified by the implementation.

### 7.11.2 The **setlocale** function

#### Synopsis

- 1 

```
#include <locale.h>
char *setlocale(int category, const char *locale);
```

<sup>264)</sup>ISO/IEC 9945 specifies locale and charmap formats that can be used to specify locales for C.

<sup>265)</sup>See "future library directions" (7.34.8).

## Description

- 2 The **setlocale** function selects the appropriate portion of the program's locale as specified by the **category** and **locale** arguments. The **setlocale** function may be used to change or query the program's entire current locale or portions thereof. The value **LC\_ALL** for **category** names the program's entire locale; the other values for **category** name only a portion of the program's locale. **LC\_COLLATE** affects the behavior of the **strcoll** and **strxfrm** functions. **LC\_CTYPE** affects the behavior of the character handling functions<sup>266)</sup> and the multibyte and wide character functions. **LC\_MONETARY** affects the monetary formatting information returned by the **localeconv** function. **LC\_NUMERIC** affects the decimal-point character for the formatted input/output functions and the string conversion functions, as well as the nonmonetary formatting information returned by the **localeconv** function. **LC\_TIME** affects the behavior of the **strftime** and **wcsftime** functions.
- 3 A value of "C" for **locale** specifies the minimal environment for C translation; a value of "" for **locale** specifies the locale-specific native environment. Other implementation-defined strings may be passed as the second argument to **setlocale**.
- 4 At program startup, the equivalent of

```
setlocale(LC_ALL, "C");
```

is executed.

- 5 A call to the **setlocale** function may introduce a data race with other calls to the **setlocale** function or with calls to functions that are affected by the current locale. The implementation shall behave as if no library function calls the **setlocale** function.

## Returns

- 6 If a pointer to a string is given for **locale** and the selection can be honored, the **setlocale** function returns a pointer to the string associated with the specified **category** for the new locale. If the selection cannot be honored, the **setlocale** function returns a null pointer and the program's locale is not changed.
- 7 A null pointer for **locale** causes the **setlocale** function to return a pointer to the string associated with the **category** for the program's current locale; the program's locale is not changed.<sup>267)</sup>
- 8 The pointer to string returned by the **setlocale** function is such that a subsequent call with that string value and its associated category will restore that part of the program's locale. The string pointed to shall not be modified by the program. The behavior is undefined if the returned value is used after a subsequent call to the **setlocale** function, or after the thread which called the **setlocale** function to obtain the returned value has exited.

**Forward references:** formatted input/output functions (7.23.6), multibyte/wide character conversion functions (7.24.8), multibyte/wide string conversion functions (7.24.9), numeric conversion functions (7.24.2), the **strcoll** function (7.27.4.4), the **strftime** function (7.30.3.6), the **strxfrm** function (7.27.4.6).

## 7.11.3 Numeric formatting convention inquiry

### 7.11.3.1 The **localeconv** function

#### Synopsis

```
1 #include <locale.h>
 struct lconv *localeconv(void);
```

## Description

- 2 The **localeconv** function sets the components of an object with type **struct lconv** with values appropriate for the formatting of numeric quantities (monetary and otherwise) according to the rules of the current locale.

<sup>266)</sup>The only functions in 7.4 whose behavior is not affected by the current locale are **isdigit** and **isxdigit**.

<sup>267)</sup>The implementation is thus required to encode in a string the various categories due to a heterogeneous locale when **category** has the value **LC\_ALL**.

- 3 The members of the structure with type **char \*** are pointers to strings, any of which (except **decimal\_point**) can point to "", to indicate that the value is not available in the current locale or is of zero length. Apart from **grouping** and **mon\_grouping**, the strings shall start and end in the initial shift state. The members with type **char** are nonnegative numbers, any of which can be **CHAR\_MAX** to indicate that the value is not available in the current locale. The members include the following:

**char \*decimal\_point**

The decimal-point character used to format nonmonetary quantities.

**char \*thousands\_sep**

The character used to separate groups of digits before the decimal-point character in formatted nonmonetary quantities.

**char \*grouping**

A string whose elements indicate the size of each group of digits in formatted nonmonetary quantities.

**char \*mon\_decimal\_point**

The decimal-point used to format monetary quantities.

**char \*mon\_thousands\_sep**

The separator for groups of digits before the decimal-point in formatted monetary quantities.

**char \*mon\_grouping**

A string whose elements indicate the size of each group of digits in formatted monetary quantities.

**char \*positive\_sign**

The string used to indicate a nonnegative-valued formatted monetary quantity.

**char \*negative\_sign**

The string used to indicate a negative-valued formatted monetary quantity.

**char \*currency\_symbol**

The local currency symbol applicable to the current locale.

**char frac\_digits**

The number of fractional digits (those after the decimal-point) to be displayed in a locally formatted monetary quantity.

**char p\_cs\_precedes**

Set to 1 or 0 if the **currency\_symbol** respectively precedes or succeeds the value for a nonnegative locally formatted monetary quantity.

**char n\_cs\_precedes**

Set to 1 or 0 if the **currency\_symbol** respectively precedes or succeeds the value for a negative locally formatted monetary quantity.

**char p\_sep\_by\_space**

Set to a value indicating the separation of the **currency\_symbol**, the sign string, and the value for a nonnegative locally formatted monetary quantity.

**char n\_sep\_by\_space**

Set to a value indicating the separation of the **currency\_symbol**, the sign string, and the value for a negative locally formatted monetary quantity.

**char p\_sign\_posn**

Set to a value indicating the positioning of the **positive\_sign** for a nonnegative locally formatted monetary quantity.

**char *n\_sign\_posn***

Set to a value indicating the positioning of the **negative\_sign** for a negative locally formatted monetary quantity.

**char \**int\_curr\_symbol***

The international currency symbol applicable to the current locale. The first three characters contain the alphabetic international currency symbol in accordance with those specified in ISO 4217. The fourth character (immediately preceding the null character) is the character used to separate the international currency symbol from the monetary quantity.

**char *int\_frac\_digits***

The number of fractional digits (those after the decimal-point) to be displayed in an internationally formatted monetary quantity.

**char *int\_p\_cs\_precedes***

Set to 1 or 0 if the **int\_curr\_symbol** respectively precedes or succeeds the value for a nonnegative internationally formatted monetary quantity.

**char *int\_n\_cs\_precedes***

Set to 1 or 0 if the **int\_curr\_symbol** respectively precedes or succeeds the value for a negative internationally formatted monetary quantity.

**char *int\_p\_sep\_by\_space***

Set to a value indicating the separation of the **int\_curr\_symbol**, the sign string, and the value for a nonnegative internationally formatted monetary quantity.

**char *int\_n\_sep\_by\_space***

Set to a value indicating the separation of the **int\_curr\_symbol**, the sign string, and the value for a negative internationally formatted monetary quantity.

**char *int\_p\_sign\_posn***

Set to a value indicating the positioning of the **positive\_sign** for a nonnegative internationally formatted monetary quantity.

**char *int\_n\_sign\_posn***

Set to a value indicating the positioning of the **negative\_sign** for a negative internationally formatted monetary quantity.

- 4 The elements of **grouping** and **mon\_grouping** are interpreted according to the following:

**CHAR\_MAX** No further grouping is to be performed.

0 The previous element is to be repeatedly used for the remainder of the digits.

*other* The integer value is the number of digits that compose the current group. The next element is examined to determine the size of the next group of digits before the current group.

- 5 The values of **p\_sep\_by\_space**, **n\_sep\_by\_space**, **int\_p\_sep\_by\_space**, and **int\_n\_sep\_by\_space** are interpreted according to the following:

0 No space separates the currency symbol and value.

1 If the currency symbol and sign string are adjacent, a space separates them from the value; otherwise, a space separates the currency symbol from the value.

2 If the currency symbol and sign string are adjacent, a space separates them; otherwise, a space separates the sign string from the value.

For **int\_p\_sep\_by\_space** and **int\_n\_sep\_by\_space**, the fourth character of **int\_curr\_symbol** is used instead of a space.

- 6 The values of **p\_sign\_posn**, **n\_sign\_posn**, **int\_p\_sign\_posn**, and **int\_n\_sign\_posn** are interpreted according to the following:
  - 0 Parentheses surround the quantity and currency symbol.
  - 1 The sign string precedes the quantity and currency symbol.
  - 2 The sign string succeeds the quantity and currency symbol.
  - 3 The sign string immediately precedes the currency symbol.
  - 4 The sign string immediately succeeds the currency symbol.
- 7 The implementation shall behave as if no library function calls the **localeconv** function.

#### Returns

- 8 The **localeconv** function returns a pointer to the filled-in object. The structure pointed to by the return value shall not be modified by the program, but may be overwritten by a subsequent call to the **localeconv** function. In addition, calls to the **setlocale** function with categories **LC\_ALL**, **LC\_MONETARY**, or **LC\_NUMERIC** may overwrite the contents of the structure.
- 9 EXAMPLE 1 Table 7.3 illustrates rules which can be used by four countries to format — potentially irrelevant, imaginary, and/or historical — monetary quantities.

**Table 7.3 — Formatting rule examples**

| Country  | Local format   |                 | International format |               |
|----------|----------------|-----------------|----------------------|---------------|
|          | Positive       | Negative        | Positive             | Negative      |
| Country1 | 1.234,56 mk    | -1.234,56 mk    | FIM 1.234,56         | FIM -1.234,56 |
| Country2 | L.1.234        | -L.1.234        | ITL 1.234            | -ITL 1.234    |
| Country3 | ƒ 1.234,56     | ƒ -1.234,56     | NLG 1.234,56         | NLG -1.234,56 |
| Country4 | SFr.s.1,234.56 | SFr.s.1,234.56C | CHF 1,234.56         | CHF 1,234.56C |

- 10 For these four countries, the respective values for the monetary members of the structure returned by **localeconv** can be as displayed in Table 7.4:

Table 7.4 — Formatting values for example countries

|                           | Country1 | Country2 | Country3 | Country4 |
|---------------------------|----------|----------|----------|----------|
| <b>mon_decimal_point</b>  | " , "    | " "      | " , "    | " . "    |
| <b>mon_thousands_sep</b>  | " . "    | " . "    | " . "    | " , "    |
| <b>mon_grouping</b>       | "\3"     | "\3"     | "\3"     | "\3"     |
| <b>positive_sign</b>      | " "      | " "      | " "      | " "      |
| <b>negative_sign</b>      | " _ "    | " _ "    | " _ "    | " ¢ "    |
| <b>currency_symbol</b>    | "mk"     | "L. "    | "\u0192" | "SFrs. " |
| <b>frac_digits</b>        | 2        | 0        | 2        | 2        |
| <b>p_cs_precedes</b>      | 0        | 1        | 1        | 1        |
| <b>n_cs_precedes</b>      | 0        | 1        | 1        | 1        |
| <b>p_sep_by_space</b>     | 1        | 0        | 1        | 0        |
| <b>n_sep_by_space</b>     | 1        | 0        | 2        | 0        |
| <b>p_sign_posn</b>        | 1        | 1        | 1        | 1        |
| <b>n_sign_posn</b>        | 1        | 1        | 4        | 2        |
| <b>int_curr_symbol</b>    | "FIM "   | "ITL "   | "NLG "   | "CHF "   |
| <b>int_frac_digits</b>    | 2        | 0        | 2        | 2        |
| <b>int_p_cs_precedes</b>  | 1        | 1        | 1        | 1        |
| <b>int_n_cs_precedes</b>  | 1        | 1        | 1        | 1        |
| <b>int_p_sep_by_space</b> | 1        | 1        | 1        | 1        |
| <b>int_n_sep_by_space</b> | 2        | 1        | 2        | 1        |
| <b>int_p_sign_posn</b>    | 1        | 1        | 1        | 1        |
| <b>int_n_sign_posn</b>    | 4        | 1        | 4        | 2        |

- 11 EXAMPLE 2 Table 7.5 illustrates how the **cs\_precedes**, **sep\_by\_space**, and **sign\_posn** members affect the formatted value.

Table 7.5 — Select member formatting affect

| <b>p_cs_precedes</b> | <b>p_sign_posn</b> | <b>p_sep_by_space</b> |           |          |
|----------------------|--------------------|-----------------------|-----------|----------|
|                      |                    | 0                     | 1         | 2        |
| 0                    | 0                  | (1.25\$)              | (1.25 \$) | (1.25\$) |
|                      | 1                  | +1.25\$               | +1.25 \$  | + 1.25\$ |
|                      | 2                  | 1.25\$+               | 1.25 \$+  | 1.25\$ + |
|                      | 3                  | 1.25+\$               | 1.25 +\$  | 1.25+ \$ |
|                      | 4                  | 1.25\$+               | 1.25 \$+  | 1.25\$ + |
| 1                    | 0                  | (\$1.25)              | (\$ 1.25) | (\$1.25) |
|                      | 1                  | +\$1.25               | +\$ 1.25  | + \$1.25 |
|                      | 2                  | \$1.25+               | \$ 1.25+  | \$1.25 + |
|                      | 3                  | +\$1.25               | +\$ 1.25  | + \$1.25 |
|                      | 4                  | \$+1.25               | \$+ 1.25  | \$ +1.25 |

## 7.12 Mathematics <math.h>

### 7.12.1 General

- 1 The header <math.h> declares two types and many mathematical functions and defines several macros. Most synopses specify a family of functions consisting of a principal function with one or more **double** parameters, a **double** return value, or both; and other functions with the same name but with **f** and **l** suffixes, which are corresponding functions with **float** and **long double** parameters, return values, or both.<sup>268)</sup> Integer arithmetic functions and conversion functions are discussed later.
- 2 The feature test macro `__STDC_VERSION_MATH_H__` expands to the token 202311L.
- 3 The types

```
float_t
double_t
```

are floating types such that the values of **float** and **double** are subsets of the values of **float\_t** and **double\_t**, respectively, and such that the values of **float\_t** are a subset of the values of **double\_t**. If `FLT_EVAL_METHOD` equals 0, **float\_t** and **double\_t** are **float** and **double**, respectively; if `FLT_EVAL_METHOD` equals 1, they are both **double**; if `FLT_EVAL_METHOD` equals 2, they are both **long double**; and for other values of `FLT_EVAL_METHOD`, they are otherwise implementation-defined.<sup>269)</sup> If they are not real floating types, the behavior is implementation-defined.

- 4 The types

```
_Decimal32_t
_Decimal64_t
```

are decimal floating types at least as wide as `_Decimal32` and `_Decimal64`, respectively, and such that `_Decimal64_t` is at least as wide as `_Decimal32_t`. They are present only if the implementation defines `__STDC_IEC_60559_DFP__`. If `DEC_EVAL_METHOD` equals 0, `_Decimal32_t` and `_Decimal64_t` are `_Decimal32` and `_Decimal64`, respectively; if `DEC_EVAL_METHOD` equals 1, they are both `_Decimal64`; if `DEC_EVAL_METHOD` equals 2, they are both `_Decimal128`; and for other values of `DEC_EVAL_METHOD`, they are otherwise implementation-defined.

- 5 The macro

```
HUGE_VAL
```

expands to a **double** arithmetic constant expression, not necessarily representable as a **float**, whose value is the maximum value returned by library functions when a floating result of type **double** overflows under the default rounding mode, either maximum finite number in the type or positive or unsigned infinity. The macros

```
HUGE_VALF
HUGE_VALL
```

are respectively **float** and **long double** analogs of `HUGE_VAL`.<sup>270)</sup>

- 6 The macros in this paragraph are only present if the implementation defines `__STDC_IEC_60559_DFP__`. The macro

```
HUGE_VAL_D32
```

<sup>268)</sup> Particularly on systems with wide expression evaluation, a <math.h> function can pass arguments and return values in wider format than the synopsis prototype indicates.

<sup>269)</sup> The types **float\_t** and **double\_t** are intended to be the implementation's most efficient types at least as wide as **float** and **double**, respectively. For `FLT_EVAL_METHOD` equal 0, 1, or 2, the type **float\_t** is the narrowest type used by the implementation to evaluate floating expressions.

<sup>270)</sup> `HUGE_VAL`, `HUGE_VALF`, and `HUGE_VALL` can be positive infinities in an implementation that supports infinities.

expands to an arithmetic constant expression of type **\_Decimal32** representing positive infinity. The macros

```
HUGE_VAL_D64
HUGE_VAL_D128
```

are respectively **\_Decimal64** and **\_Decimal128** analogs of **HUGE\_VAL\_D32**.

- 7 The macro

```
INFINITY
```

is defined if and only if the implementation supports an infinity for the type **float**. It expands to an arithmetic constant expression of type **float** representing positive or unsigned infinity.

- 8 The macro

```
DEC_INFINITY
```

expands to an arithmetic constant expression of type **\_Decimal32** representing positive infinity.

- 9 The macro

```
NAN
```

is defined if and only if the implementation supports quiet NaNs for the **float** type. It expands to an arithmetic constant expression of type **float** representing a quiet NaN.

- 10 The macro

```
DEC_NAN
```

expands to an arithmetic constant expression of type **\_Decimal32** representing a quiet NaN.

- 11 Use of the macros **INFINITY**, **DEC\_INFINITY**, **NAN**, and **DEC\_NAN** in `<math.h>` is an obsolescent feature. Instead, use the same macros in `<float.h>`.

- 12 The *number classification macros*

```
FP_INFINITE
FP_NAN
FP_NORMAL
FP_SUBNORMAL
FP_ZERO
```

represent mutually exclusive kinds of floating-point values. They expand to integer constant expressions with distinct values. Additional implementation-defined floating-point classifications, with macro definitions beginning with **FP\_** and an uppercase letter, may also be specified by the implementation.

- 13 The *math rounding direction macros*

```
FP_INT_UPWARD
FP_INT_DOWNWARD
FP_INT_TOWARDZERO
FP_INT_TONEARESTFROMZERO
FP_INT_TONEAREST
```

represent the rounding directions of the functions **ceil**, **floor**, **trunc**, **round**, and **roundeven**, respectively, that convert to integral values in floating-point formats. They expand to integer constant expressions with distinct values suitable for use as the second argument to the **fromfp**, **ufromfp**, **fromfpx**, and **ufromfpx** functions.



- 14 The macro

**FP\_FAST\_FMA**

is optionally defined. If defined, it indicates that the **fma** function generally executes about as fast as, or faster than, a multiply and an add of **double** operands.<sup>271)</sup> The macros

**FP\_FAST\_FMAF**  
**FP\_FAST\_FMAL**

are, respectively, **float** and **long double** analogs of **FP\_FAST\_FMA**. If defined, these macros expand to the integer literal 1.

- 15 The macros

**FP\_FAST\_FMAD32**  
**FP\_FAST\_FMAD64**  
**FP\_FAST\_FMAD128**

are, respectively, **\_Decimal32**, **\_Decimal64**, and **\_Decimal128** analogs of **FP\_FAST\_FMA**.

- 16 Each of the macros

|                      |                      |                       |                      |
|----------------------|----------------------|-----------------------|----------------------|
| <b>FP_FAST_FADD</b>  | <b>FP_FAST_DSUBL</b> | <b>FP_FAST_FDIVL</b>  | <b>FP_FAST_FFMA</b>  |
| <b>FP_FAST_FADDL</b> | <b>FP_FAST_FMUL</b>  | <b>FP_FAST_FDIVL</b>  | <b>FP_FAST_FFMAL</b> |
| <b>FP_FAST_DADDL</b> | <b>FP_FAST_FMULL</b> | <b>FP_FAST_FSQRT</b>  | <b>FP_FAST_DFMAL</b> |
| <b>FP_FAST_FSUB</b>  | <b>FP_FAST_DMULL</b> | <b>FP_FAST_FSQRTL</b> |                      |
| <b>FP_FAST_FSUBL</b> | <b>FP_FAST_FDIV</b>  | <b>FP_FAST_DSQRTL</b> |                      |

is optionally defined. If defined, it indicates that the corresponding function generally executes about as fast, or faster, than the corresponding operation or function of the argument type with result type the same as the argument type followed by conversion to the narrower type. For **FP\_FAST\_FFMA**, **FP\_FAST\_FFMAL**, and **FP\_FAST\_DFMAL**, the comparison is to a call to **fma** or **fmal** followed by a conversion, not to separate multiply, add, and conversion. If defined, these macros expand to the integer literal 1.

- 17 The macros

|                           |                           |                            |
|---------------------------|---------------------------|----------------------------|
| <b>FP_FAST_D32ADDD64</b>  | <b>FP_FAST_D32MULD64</b>  | <b>FP_FAST_D32FMAD64</b>   |
| <b>FP_FAST_D32ADDD128</b> | <b>FP_FAST_D32MULD128</b> | <b>FP_FAST_D32FMAD128</b>  |
| <b>FP_FAST_D64ADDD128</b> | <b>FP_FAST_D64MULD128</b> | <b>FP_FAST_D64FMAD128</b>  |
| <b>FP_FAST_D32SUBD64</b>  | <b>FP_FAST_D32DIVD64</b>  | <b>FP_FAST_D32SQRTD64</b>  |
| <b>FP_FAST_D32SUBD128</b> | <b>FP_FAST_D32DIVD128</b> | <b>FP_FAST_D32SQRTD128</b> |
| <b>FP_FAST_D64SUBD128</b> | <b>FP_FAST_D64DIVD128</b> | <b>FP_FAST_D64SQRTD128</b> |

are analogs of **FP\_FAST\_FADD**, **FP\_FAST\_FADDL**, **FP\_FAST\_DADDL**, etc., for decimal floating types.

- 18 The macros

**FP\_ILOGB0**  
**FP\_ILOGBNAN**

expand to integer constant expressions whose values are returned by **ilogb(x)** if **x** is zero or NaN, respectively. The value of **FP\_ILOGB0** shall be either **INT\_MIN** or **-INT\_MAX**. The value of **FP\_ILOGBNAN** shall be either **INT\_MAX** or **INT\_MIN**.

- 19 The macros

<sup>271)</sup>Typically, the **FP\_FAST\_FMA** macro is defined if and only if the **fma** function is implemented directly with a hardware multiply-add instruction. Software implementations are expected to be substantially slower.

**FP\_LLOGB0**  
**FP\_LLOGBNAN**

expand to integer constant expressions whose values are returned by **llogb(x)** if **x** is zero or NaN, respectively. The value of **FP\_LLOGB0** shall be **LONG\_MIN** if the value of **FP\_ILOGB0** is **INT\_MIN**, and shall be **-LONG\_MAX** if the value of **FP\_ILOGB0** is **-INT\_MAX**. The value of **FP\_LLOGBNAN** shall be **LONG\_MAX** if the value of **FP\_ILOGBNAN** is **INT\_MAX**, and shall be **LONG\_MIN** if the value of **FP\_ILOGBNAN** is **INT\_MIN**.

20 The macros

**MATH\_ERRNO**  
**MATH\_ERREXCEPT**

expand to the integer literals 1 and 2, respectively; the macro

**math\_errhandling**

expands to an expression that has type **int** and the value **MATH\_ERRNO**, **MATH\_ERREXCEPT**, the bitwise OR of both, or 0; the value shall not be 0 in a hosted implementation. The value of **math\_errhandling** is constant for the duration of the program. It is unspecified whether **math\_errhandling** is a macro or an identifier with external linkage. If a macro definition is suppressed or a program defines an identifier with the name **math\_errhandling**, the behavior is undefined. If the expression **math\_errhandling & MATH\_ERREXCEPT** can be nonzero, the implementation shall define the macros **FE\_DIVBYZERO**, **FE\_INVALID**, and **FE\_OVERFLOW** in `<fenv.h>`.

### 7.12.2 Treatment of error conditions

- 1 The behavior of each of the functions in `<math.h>` is specified for all representable values of its input arguments, except where explicitly stated otherwise. Each function shall execute as if it were a single operation without raising **SIGFPE** and without generating any of the floating-point exceptions “invalid”, “divide-by-zero”, or “overflow” except to reflect the result of the function.
- 2 For all functions, a *domain error* occurs if and only if an input argument is outside the domain over which the mathematical function is defined. The description of each function lists any required domain errors; an implementation may define additional domain errors, provided that such errors are consistent with the mathematical definition of the function.<sup>272)</sup> Whether a signaling NaN input causes a domain error is implementation-defined. On a domain error, the function returns an implementation-defined value; if the integer expression **math\_errhandling & MATH\_ERRNO** is nonzero, the integer expression **errno** acquires the value **EDOM**; if the integer expression **math\_errhandling & MATH\_ERREXCEPT** is nonzero, the “invalid” floating-point exception is raised.
- 3 Similarly, a *pole error* (also known as a singularity or infinitary) occurs in cases where the arguments are finite and where the mathematical function typically has an infinite limit (for example, a pole error occurs at **log(0.0)** because the mathematical log function has a right-hand limit of  $-\infty$  at 0). The description of each function lists any required pole errors; an implementation may define additional pole errors, provided that such errors are consistent with the mathematical definition of the function. On a pole error, the function returns an implementation-defined value; if the integer expression **math\_errhandling & MATH\_ERRNO** is nonzero, the integer expression **errno** acquires the value **ERANGE**; if the integer expression **math\_errhandling & MATH\_ERREXCEPT** is nonzero, the “divide-by-zero” floating-point exception is raised.
- 4 Likewise, a *range error* occurs if and only if the result overflows or underflows, as defined below. The description of each function lists any required range errors; an implementation may define additional range errors, provided that such errors are consistent with the mathematical definition of the function and are the result of either overflow or underflow. Range errors that are required or

<sup>272)</sup>In an implementation that supports infinities, this allows an infinity as an argument to be a domain error if the mathematical domain of the function does not include the infinity.

implementation-defined shall or may be reported, as specified in this subclause, respectively.

- 5 A floating result overflows if a finite result value with ordinary accuracy<sup>273)</sup> would have magnitude (absolute value) too large for the representation with full precision in the specified type. A result that is exactly an infinity does not overflow. If a floating result overflows and default rounding is in effect, then the function returns the value of the macro **HUGE\_VAL** suffixed according to the return type, with the same sign as the correct value of the function; however, for the types with reduced-precision representations of numbers beyond the overflow threshold, the function may return a representation of the result with less than full precision for the type. If a floating result overflows and default rounding is in effect and the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero, then the integer expression **errno** acquires the value **ERANGE**. If a floating result overflows, and the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, the “overflow” floating-point exception is raised (regardless of whether default rounding is in effect).
- 6 The result underflows if a nonzero result value with ordinary accuracy would have magnitude (absolute value) less than the minimum normalized number in the type; however a zero result that is specified to be an exact zero does not underflow. Also, a result with ordinary accuracy and the magnitude of the minimum normalized number may underflow.<sup>274)</sup> If the result underflows, the function returns an implementation-defined value whose magnitude is no greater than the smallest normalized positive number in the specified type; if the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero, whether **errno** acquires the value **ERANGE** is implementation-defined; if the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, whether the “underflow” floating-point exception is raised is implementation-defined.
- 7 If a domain, pole, or range error occurs and the integer expression **math\_errhandling** & **MATH\_ERRNO** is zero,<sup>275)</sup> then **errno** shall either be set to the value corresponding to the error or left unmodified. If no such error occurs, **errno** shall be left unmodified regardless of the setting of **math\_errhandling**.

### 7.12.3 The **FP\_CONTRACT** pragma

#### Synopsis

```
1 #include <math.h>
 #pragma STDC FP_CONTRACT on-off-switch
```

#### Description

- 2 The **FP\_CONTRACT** pragma can be used to allow (if the state is “on”) or disallow (if the state is “off”) the implementation to contract expressions (6.5.1). Each pragma can occur either outside external declarations or preceding all explicit declarations and statements inside a compound statement. When outside external declarations, the pragma takes effect from its occurrence until another **FP\_CONTRACT** pragma is encountered, or until the end of the translation unit. When inside a compound statement, the pragma takes effect from its occurrence until another **FP\_CONTRACT** pragma is encountered (including within a nested compound statement), or until the end of the compound statement; at the end of a compound statement the state for the pragma is restored to its condition just before the compound statement. If this pragma is used in any other context, the behavior is undefined. The default state (“on” or “off”) for the pragma is implementation-defined.

### 7.12.4 Classification macros

#### 7.12.4.1 General

- 1 Floating-point values can be classified as NaN, infinite, normal, subnormal, or zero, or into other implementation-defined categories. Numbers whose magnitude is at least  $b^{e_{\min}-1}$  (the minimum magnitude of normalized floating-point numbers in the type) and at most  $(1 - b^{-p})b^{e_{\max}}$  (the maximum magnitude of normalized floating-point numbers in the type), where  $b$ ,  $p$ ,  $e_{\min}$ , and  $e_{\max}$

<sup>273)</sup>Ordinary accuracy is determined by the implementation. It refers to the accuracy of the function where results are not compromised by extreme magnitude.

<sup>274)</sup>The term underflow here is intended to encompass both “gradual underflow” as in ISO/IEC 60559 and also “flush-to-zero” underflow. ISO/IEC 60559 underflow can occur in cases where the magnitude of the rounded result (accurate to the full precision of the type) equals the minimum normalized number in the format.

<sup>275)</sup>Math errors are being indicated by the floating-point exception flags rather than by **errno**.

are as in 5.3.5.3.3, are classified as normal. Larger magnitude finite numbers represented with full precision in the type may also be classified as normal. Nonzero numbers whose magnitude is less than  $b^{e_{\min}-1}$  are classified as subnormal.

- 2 In the synopses in this subclause, *real-floating* indicates that the argument shall be an expression of real floating type.

#### 7.12.4.2 The **fpclassify** macro

##### Synopsis

```
1 #include <math.h>
 int fpclassify(real-floating x);
```

##### Description

- 2 The **fpclassify** macro classifies its argument value as NaN, infinite, normal, subnormal, zero, or into another implementation-defined category. First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then classification is based on the type of the argument.<sup>276)</sup>

##### Returns

- 3 The **fpclassify** macro returns the value of the number classification macro appropriate to the value of its argument.

#### 7.12.4.3 The **iscanonical** macro

##### Synopsis

```
1 #include <math.h>
 int iscanonical(real-floating x);
```

##### Description

- 2 The **iscanonical** macro determines whether its argument value is canonical (5.3.5.3.3). First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then, determination is based on the type of the argument.

##### Returns

- 3 The **iscanonical** macro returns a nonzero value if and only if its argument is canonical.

#### 7.12.4.4 The **isfinite** macro

##### Synopsis

```
1 #include <math.h>
 int isfinite(real-floating x);
```

##### Description

- 2 The **isfinite** macro determines whether its argument has a finite value (zero, subnormal, or normal, and not infinite or NaN). First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then determination is based on the type of the argument.

##### Returns

- 3 The **isfinite** macro returns a nonzero value if and only if its argument has a finite value.

#### 7.12.4.5 The **isinf** macro

##### Synopsis

```
1 #include <math.h>
```

<sup>276)</sup>Since an expression can be evaluated with more range and precision than its type has, it is important to know the type that classification is based on. For example, a normal **long double** value can become subnormal when converted to **double**, and zero when converted to **float**.

```
int isinf(real-floating x);
```

#### Description

- 2 The **isinf** macro determines whether its argument value is (positive or negative) infinity. First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then determination is based on the type of the argument.

#### Returns

- 3 The **isinf** macro returns a nonzero value if and only if its argument has an infinite value.

#### 7.12.4.6 The **isnan** macro

##### Synopsis

```
1 #include <math.h>
 int isnan(real-floating x);
```

#### Description

- 2 The **isnan** macro determines whether its argument value is a NaN. First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then determination is based on the type of the argument.<sup>277)</sup>

#### Returns

- 3 The **isnan** macro returns a nonzero value if and only if its argument has a NaN value.

#### 7.12.4.7 The **isnormal** macro

##### Synopsis

```
1 #include <math.h>
 int isnormal(real-floating x);
```

#### Description

- 2 The **isnormal** macro determines whether its argument value is normal (neither zero, subnormal, infinite, nor NaN). First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then determination is based on the type of the argument.

#### Returns

- 3 The **isnormal** macro returns a nonzero value if and only if its argument has a normal value.

#### 7.12.4.8 The **signbit** macro

##### Synopsis

```
1 #include <math.h>
 int signbit(real-floating x);
```

#### Description

- 2 The **signbit** macro determines whether the sign of its argument value is negative.<sup>278)</sup> If the argument value is an unsigned zero, its sign is regarded as positive. Otherwise, if the argument value is unsigned, the result value (zero or nonzero) is implementation-defined.

#### Returns

- 3 The **signbit** macro returns a nonzero value if and only if the sign of its argument value is determined to be negative.

#### 7.12.4.9 The **issignaling** macro

<sup>277)</sup>For the **isnan** macro, the type for determination does not matter unless the implementation supports NaNs in the evaluation type but not in the semantic type.

<sup>278)</sup>The **signbit** macro determines the sign of all values, including infinities, zeros, and NaNs.

## Synopsis

```
1 #include <math.h>
 int issignaling(real-floating x);
```

## Description

2 The **issignaling** macro determines whether its argument value is a signaling NaN.

## Returns

3 The **issignaling** macro returns a nonzero value if and only if its argument is a signaling NaN.<sup>279)</sup>

### 7.12.4.10 The **issubnormal** macro

## Synopsis

```
1 #include <math.h>
 int issubnormal(real-floating x);
```

## Description

2 The **issubnormal** macro determines whether its argument value is subnormal. First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then determination is based on the type of the argument.

## Returns

3 The **issubnormal** macro returns a nonzero value if and only if its argument is subnormal.

### 7.12.4.11 The **iszero** macro

## Synopsis

```
1 #include <math.h>
 int iszero(real-floating x);
```

## Description

2 The **iszero** macro determines whether its argument value is (positive, negative, or unsigned) zero. First, an argument represented in a format wider than its semantic type is converted to its semantic type. Then, determination is based on the type of the argument.

## Returns

3 The **iszero** macro returns a nonzero value if and only if its argument is zero.

## 7.12.5 Trigonometric functions

### 7.12.5.1 The **acos** functions

## Synopsis

```
1 #include <math.h>
 double acos(double x);
 float acosf(float x);
 long double acosl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 acosd32(_Decimal32 x);
 _Decimal64 acosd64(_Decimal64 x);
 _Decimal128 acosd128(_Decimal128 x);
 #endif
```

<sup>279)</sup>F.3 specifies that **issignaling** (and all the other classification macros), raise no floating-point exception if the argument is a variable, or any other expression whose value is represented in the format of its semantic type, even if the value is a signaling NaN.

**Description**

- 2 The **acos** functions compute the principal value of the arc cosine of  $x$ . A domain error occurs for arguments not in the interval  $[-1, +1]$ .

**Returns**

- 3 The **acos** functions return  $\arccos x$  in the interval  $[0, \pi]$  radians.

**7.12.5.2 The **asin** functions****Synopsis**

```

1 #include <math.h>
 double asin(double x);
 float asinf(float x);
 long double asinl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 asind32(_Decimal32 x);
 _Decimal64 asind64(_Decimal64 x);
 _Decimal128 asind128(_Decimal128 x);
 #endif

```

**Description**

- 2 The **asin** functions compute the principal value of the arc sine of  $x$ . A domain error occurs for arguments not in the interval  $[-1, +1]$ . A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **asin** functions return  $\arcsin x$  in the interval  $[-\frac{\pi}{2}, +\frac{\pi}{2}]$  radians.

**7.12.5.3 The **atan** functions****Synopsis**

```

1 #include <math.h>
 double atan(double x);
 float atanf(float x);
 long double atanl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 atand32(_Decimal32 x);
 _Decimal64 atand64(_Decimal64 x);
 _Decimal128 atand128(_Decimal128 x);
 #endif

```

**Description**

- 2 The **atan** functions compute the principal value of the arc tangent of  $x$ . A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **atan** functions return  $\arctan x$  in the interval  $[-\frac{\pi}{2}, +\frac{\pi}{2}]$  radians.

**7.12.5.4 The **atan2** functions****Synopsis**

```

1 #include <math.h>
 double atan2(double y, double x);
 float atan2f(float y, float x);
 long double atan2l(long double y, long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 atan2d32(_Decimal32 y, _Decimal32 x);
 _Decimal64 atan2d64(_Decimal64 y, _Decimal64 x);
 _Decimal128 atan2d128(_Decimal128 y, _Decimal128 x);
 #endif

```

**Description**

- 2 The **atan2** functions compute the value of the arc tangent of  $y/x$ , using the signs of both arguments to determine the quadrant of the return value. A domain error may occur if both arguments are zero. A range error occurs if  $x$  is positive and nonzero  $\frac{y}{x}$  is too close to zero.

**Returns**

- 3 The **atan2** functions return  $\arctan(y/x)$  in the interval  $[-\pi, +\pi]$  radians.

**7.12.5.5 The **cos** functions****Synopsis**

```
1 #include <math.h>
 double cos(double x);
 float cosf(float x);
 long double cosl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 cosd32(_Decimal32 x);
 _Decimal64 cosd64(_Decimal64 x);
 _Decimal128 cosd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **cos** functions compute the cosine of  $x$  (measured in radians).

**Returns**

- 3 The **cos** functions return  $\cos x$ .

**7.12.5.6 The **sin** functions****Synopsis**

```
1 #include <math.h>
 double sin(double x);
 float sinf(float x);
 long double sinl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 sind32(_Decimal32 x);
 _Decimal64 sind64(_Decimal64 x);
 _Decimal128 sind128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **sin** functions compute the sine of  $x$  (measured in radians). A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **sin** functions return  $\sin x$ .

**7.12.5.7 The **tan** functions****Synopsis**

```
1 #include <math.h>
 double tan(double x);
 float tanf(float x);
 long double tanl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 tand32(_Decimal32 x);
 _Decimal64 tand64(_Decimal64 x);
 _Decimal128 tand128(_Decimal128 x);
 #endif
```



**Description**

- 2 The **tan** functions return the tangent of  $x$  (measured in radians). A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **tan** functions return  $\tan x$ .

**7.12.5.8 The **acospi** functions****Synopsis**

```
1 #include <math.h>
 double acospi(double x);
 float acospi(float x);
 long double acospil(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 acospid32(_Decimal32 x);
 _Decimal64 acospid64(_Decimal64 x);
 _Decimal128 acospid128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **acospi** functions compute the principal value of the arc cosine of  $x$ , divided by  $\pi$ , thus measuring the angle in half-revolutions. A domain error occurs for arguments not in the interval  $[-1, +1]$ .

**Returns**

- 3 The **acospi** functions return  $\arccos(x)/\pi$  in the interval  $[0, 1]$ .

**7.12.5.9 The **asinpi** functions****Synopsis**

```
1 #include <math.h>
 double asinpi(double x);
 float asinpi(float x);
 long double asinpil(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 asinpid32(_Decimal32 x);
 _Decimal64 asinpid64(_Decimal64 x);
 _Decimal128 asinpid128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **asinpi** functions compute the principal value of the arc sine of  $x$ , divided by  $\pi$ , thus measuring the angle in half-revolutions. A domain error occurs for arguments not in the interval  $[-1, +1]$ . A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **asinpi** functions return  $\arcsin(x)/\pi$  in the interval  $[-\frac{1}{2}, +\frac{1}{2}]$ .

**7.12.5.10 The **atanpi** functions****Synopsis**

```
1 #include <math.h>
 double atanpi(double x);
 float atanpi(float x);
 long double atanpil(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 atanpid32(_Decimal32 x);
 _Decimal64 atanpid64(_Decimal64 x);
 #endif
```

```
_Decimal128 atanpid128(_Decimal128 x);
#endif
```

### Description

- 2 The **atanpi** functions compute the principal value of the arc tangent of  $x$ , divided by  $\pi$ , thus measuring the angle in half-revolutions. A range error occurs if nonzero  $x$  is too close to zero.

### Returns

- 3 The **atanpi** functions return  $\arctan(x)/\pi$ . in the interval  $[-\frac{1}{2}, +\frac{1}{2}]$ .

#### 7.12.5.11 The **atan2pi** functions

##### Synopsis

```
1 #include <math.h>
 double atan2pi(double y, double x);
 float atan2pif(float y, float x);
 long double atan2pil(long double y, long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 atan2pid32(_Decimal32 y, _Decimal32 x);
 _Decimal64 atan2pid64(_Decimal64 y, _Decimal64 x);
 _Decimal128 atan2pid128(_Decimal128 y, _Decimal128 x);
 #endif
```

### Description

- 2 The **atan2pi** functions compute the angle, measured in half-revolutions, subtended at the origin by the point  $(x, y)$  and the positive x-axis. Thus, the **atan2pi** functions compute  $\arctan(\frac{y}{x})/\pi$ , in the range  $[-1, +1]$ . A domain error may occur if both arguments are zero. A range error occurs if  $x$  is positive and nonzero  $\frac{y}{x}$  is too close to zero.

### Returns

- 3 The **atan2pi** functions return the computed angle, in the interval  $[-1, +1]$ .

#### 7.12.5.12 The **cospi** functions

##### Synopsis

```
1 #include <math.h>
 double cospi(double x);
 float cospif(float x);
 long double cospil(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 cospid32(_Decimal32 x);
 _Decimal64 cospid64(_Decimal64 x);
 _Decimal128 cospid128(_Decimal128 x);
 #endif
```

### Description

- 2 The **cospi** functions compute the cosine of  $\pi \times x$ , thus regarding  $x$  as a measurement in half-revolutions.

### Returns

- 3 The **cospi** functions return  $\cos(\pi \times x)$ .

#### 7.12.5.13 The **sinpi** functions

##### Synopsis

```
1 #include <math.h>
 double sinpi(double x);
 float sinpif(float x);
```

```

long double sinpi(long double x);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 sinpid32(_Decimal32 x);
_Decimal64 sinpid64(_Decimal64 x);
_Decimal128 sinpid128(_Decimal128 x);
#endif

```

#### Description

- 2 The **sinpi** functions compute the sine of  $\pi \times x$ , thus regarding  $x$  as a measurement in half-revolutions. A range error occurs if nonzero  $x$  is too close to zero.

#### Returns

- 3 The **sinpi** functions return  $\sin(\pi \times x)$ .

### 7.12.5.14 The **tanpi** functions

#### Synopsis

```

1 #include <math.h>
 double tanpi(double x);
 float tanpif(float x);
 long double tanpil(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 tanpid32(_Decimal32 x);
 _Decimal64 tanpid64(_Decimal64 x);
 _Decimal128 tanpid128(_Decimal128 x);
 #endif

```

#### Description

- 2 The **tanpi** functions compute the tangent of  $\pi \times x$ , thus regarding  $x$  as a measurement in half-revolutions. A range error occurs if nonzero  $x$  is too close to zero. A pole error may occur if  $|x|$  is  $(n + 0.5)$  for integer  $n$ .

#### Returns

- 3 The **tanpi** functions return  $\tan(\pi \times x)$ .

## 7.12.6 Hyperbolic functions

### 7.12.6.1 The **acosh** functions

#### Synopsis

```

1 #include <math.h>
 double acosh(double x);
 float acoshf(float x);
 long double acoshl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 acoshd32(_Decimal32 x);
 _Decimal64 acoshd64(_Decimal64 x);
 _Decimal128 acoshd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The **acosh** functions compute the (nonnegative) arc hyperbolic cosine of  $x$ . A domain error occurs for arguments less than 1.

#### Returns

- 3 The **acosh** functions return  $\operatorname{arcosh} x$  in the interval  $[0, +\infty]$ .

### 7.12.6.2 The **asinh** functions

## Synopsis

```

1 #include <math.h>
 double asinh(double x);
 float asinhf(float x);
 long double asinhl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 asinhd32(_Decimal32 x);
 _Decimal64 asinhd64(_Decimal64 x);
 _Decimal128 asinhd128(_Decimal128 x);
 #endif

```

## Description

- 2 The **asinh** functions compute the arc hyperbolic sine of *x*. A range error occurs if nonzero *x* is too close to zero.

## Returns

- 3 The **asinh** functions return  $\operatorname{arsinh} x$ .

### 7.12.6.3 The **atanh** functions

## Synopsis

```

1 #include <math.h>
 double atanh(double x);
 float atanhf(float x);
 long double atanhl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 atanhd32(_Decimal32 x);
 _Decimal64 atanhd64(_Decimal64 x);
 _Decimal128 atanhd128(_Decimal128 x);
 #endif

```

## Description

- 2 The **atanh** functions compute the arc hyperbolic tangent of *x*. A domain error occurs for arguments not in the interval  $[-1, +1]$ . A pole error may occur if the argument equals  $-1$  or  $+1$ . A range error occurs if nonzero *x* is too close to zero.

## Returns

- 3 The **atanh** functions return  $\operatorname{artanh} x$ .

### 7.12.6.4 The **cosh** functions

## Synopsis

```

1 #include <math.h>
 double cosh(double x);
 float coshf(float x);
 long double coshl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 coshd32(_Decimal32 x);
 _Decimal64 coshd64(_Decimal64 x);
 _Decimal128 coshd128(_Decimal128 x);
 #endif

```

## Description

- 2 The **cosh** functions compute the hyperbolic cosine of *x*. A range error occurs if the magnitude of finite *x* is too large.

## Returns

- 3 The **cosh** functions return  $\cosh x$ .

### 7.12.6.5 The `sinh` functions

#### Synopsis

```

1 #include <math.h>
 double sinh(double x);
 float sinhf(float x);
 long double sinhl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 sinhd32(_Decimal32 x);
 _Decimal64 sinhd64(_Decimal64 x);
 _Decimal128 sinhd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The `sinh` functions compute the hyperbolic sine of `x`. A range error occurs if the magnitude of finite `x` is too large or if nonzero `x` is too close to zero.

#### Returns

- 3 The `sinh` functions return  $\sinh x$ .

### 7.12.6.6 The `tanh` functions

#### Synopsis

```

1 #include <math.h>
 double tanh(double x);
 float tanhf(float x);
 long double tanhl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 tanhd32(_Decimal32 x);
 _Decimal64 tanhd64(_Decimal64 x);
 _Decimal128 tanhd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The `tanh` functions compute the hyperbolic tangent of `x`. A range error occurs if nonzero `x` is too close to zero.

#### Returns

- 3 The `tanh` functions return  $\tanh x$ .

## 7.12.7 Exponential and logarithmic functions

### 7.12.7.1 The `exp` functions

#### Synopsis

```

1 #include <math.h>
 double exp(double x);
 float expf(float x);
 long double expl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 expd32(_Decimal32 x);
 _Decimal64 expd64(_Decimal64 x);
 _Decimal128 expd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The `exp` functions compute the base-*e* exponential of `x`. A range error occurs if the magnitude of finite `x` is too large.

**Returns**

- 3 The **exp** functions return  $e^x$ .

**7.12.7.2 The exp10 functions****Synopsis**

```
1 #include <math.h>
 double exp10(double x);
 float exp10f(float x);
 long double exp10l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 exp10d32(_Decimal32 x);
 _Decimal64 exp10d64(_Decimal64 x);
 _Decimal128 exp10d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **exp10** functions compute the base-10 exponential of  $x$ . A range error occurs if the magnitude of finite  $x$  is too large.

**Returns**

- 3 The **exp10** functions return  $10^x$ .

**7.12.7.3 The exp10m1 functions****Synopsis**

```
1 #include <math.h>
 double exp10m1(double x);
 float exp10m1f(float x);
 long double exp10m1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 exp10md32(_Decimal32 x);
 _Decimal64 exp10md64(_Decimal64 x);
 _Decimal128 exp10md128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **exp10m1** functions compute the base-10 exponential of the argument, minus 1. A range error occurs if positive finite  $x$  is too large or if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **exp10m1** functions return  $10^x - 1$ .

**7.12.7.4 The exp2 functions****Synopsis**

```
1 #include <math.h>
 double exp2(double x);
 float exp2f(float x);
 long double exp2l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 exp2d32(_Decimal32 x);
 _Decimal64 exp2d64(_Decimal64 x);
 _Decimal128 exp2d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **exp2** functions compute the base-2 exponential of  $x$ . A range error occurs if the magnitude of finite  $x$  is too large.

**Returns**

- 3 The **exp2** functions return  $2^x$ .

**7.12.7.5 The exp2m1 functions****Synopsis**

```
1 #include <math.h>
 double exp2m1(double x);
 float exp2m1f(float x);
 long double exp2m1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 exp2m1d32(_Decimal32 x);
 _Decimal64 exp2m1d64(_Decimal64 x);
 _Decimal128 exp2m1d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **exp2m1** functions compute the base-2 exponential of the argument, minus 1. A range error occurs if positive finite  $x$  is too large or if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **exp2m1** functions return  $2^x - 1$ .

**7.12.7.6 The expm1 functions****Synopsis**

```
1 #include <math.h>
 double expm1(double x);
 float expm1f(float x);
 long double expm1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 expm1d32(_Decimal32 x);
 _Decimal64 expm1d64(_Decimal64 x);
 _Decimal128 expm1d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **expm1** functions compute the base- $e$  exponential of the argument, minus 1. A range error occurs if positive finite  $x$  is too large or if nonzero  $x$  is too close to zero.<sup>280)</sup>

**Returns**

- 3 The **expm1** functions return  $e^x - 1$ .

**7.12.7.7 The frexp functions****Synopsis**

```
1 #include <math.h>
 double frexp(double value, int *p);
 float frexpf(float value, int *p);
 long double frexpl(long double value, int *p);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 frexpd32(_Decimal32 value, int *p);
 _Decimal64 frexpd64(_Decimal64 value, int *p);
 _Decimal128 frexpd128(_Decimal128 value, int *p);
 #endif
```

<sup>280)</sup>For small magnitude  $x$ , **expm1**( $x$ ) is expected to be more accurate than **exp**( $x$ ) - 1.

**Description**

- 2 The **frexp** functions break a floating-point number into a normalized fraction and an integer exponent. They store the integer in the **int** object pointed to by **p**. If the return type of the function is a standard floating type, the exponent is an integral power of 2. If the return type of the function is a decimal floating type, the exponent is an integral power of 10.

**Returns**

- 3 If **value** is not a floating-point number or if the integral power is outside the range of **int**, the results are unspecified. Otherwise, the **frexp** functions return the value **x**, such that **x** has a magnitude in the interval  $[\frac{1}{2}, 1)$  or zero, and **value** equals  $x \times 2^{*p}$ , when the return type of the function is a standard floating type; or **x** has a magnitude in the interval  $[1/10, 1)$  or zero, and **value** equals  $x \times 10^{*p}$ , when the return type of the function is a decimal floating type. If **value** is zero, both parts of the result are zero.

**7.12.7.8 The *ilogb* functions****Synopsis**

```
1 #include <math.h>
 int ilogb(double x);
 int ilogbf(float x);
 int ilogbl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 int ilogbd32(_Decimal32 x);
 int ilogbd64(_Decimal64 x);
 int ilogbd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **ilogb** functions extract the exponent of **x** as a signed **int** value. If **x** is zero they compute the value **FP\_ILOGB0**; if **x** is infinite they compute the value **INT\_MAX**; if **x** is a NaN they compute the value **FP\_ILOGBNAN**; otherwise, they are equivalent to calling the corresponding **logb** function and converting the returned value to type **int**. A domain error or range error may occur if **x** is zero, infinite, or NaN. If the correct value is outside the range of the return type, the numeric result is unspecified and a domain error or range error may occur.

**Returns**

- 3 The **ilogb** functions return the exponent of **x** as a signed **int** value.

**Forward references:** the **logb** functions (7.12.7.17).

**7.12.7.9 The *ldexp* functions****Synopsis**

```
1 #include <math.h>
 double ldexp(double x, int p);
 float ldexpf(float x, int p);
 long double ldexpl(long double x, int p);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 ldexpd32(_Decimal32 x, int p);
 _Decimal64 ldexpd64(_Decimal64 x, int p);
 _Decimal128 ldexpd128(_Decimal128 x, int p);
 #endif
```

**Description**

- 2 The **ldexp** functions multiply a floating-point number by an integral power of 2 when the return type of the function is a standard floating type, or by an integral power of 10 when the return type of the function is a decimal floating type. A range error occurs for some finite **x**, depending on **p**.



**Returns**

- 3 The **ldexp** functions return  $x \times 2^p$  when the return type of the function is a standard floating type, or return  $x \times 10^p$  when the return type of the function is a decimal floating type.

**7.12.7.10 The **llogb** functions****Synopsis**

```
1 #include <math.h>
 long int llogb(double x);
 long int llogbf(float x);
 long int llogbl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 long int llogbd32(_Decimal32 x);
 long int llogbd64(_Decimal64 x);
 long int llogbd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **llogb** functions extract the exponent of  $x$  as a signed **long int** value. If  $x$  is zero they compute the value **FP\_LLOGB0**; if  $x$  is infinite they compute the value **LONG\_MAX**; if  $x$  is a NaN they compute the value **FP\_LLOGBNAN**; otherwise, they are equivalent to calling the corresponding **logb** function and converting the returned value to type **long int**. A domain error or range error may occur if  $x$  is zero, infinite, or NaN. If the correct value is outside the range of the return type, the numeric result is unspecified.

**Returns**

- 3 The **llogb** functions return the exponent of  $x$  as a signed **long int** value.

**Forward references:** the **logb** functions (7.12.7.17).

**7.12.7.11 The **log** functions****Synopsis**

```
1 #include <math.h>
 double log(double x);
 float logf(float x);
 long double logl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 logd32(_Decimal32 x);
 _Decimal64 logd64(_Decimal64 x);
 _Decimal128 logd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **log** functions compute the base- $e$  (natural) logarithm of  $x$ . A domain error occurs if the argument is less than zero. A pole error may occur if the argument is zero.

**Returns**

- 3 The **log** functions return  $\log_e x$ .

**7.12.7.12 The **log10** functions****Synopsis**

```
1 #include <math.h>
 double log10(double x);
 float log10f(float x);
 long double log10l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 log10d32(_Decimal32 x);
 _Decimal64 log10d64(_Decimal64 x);
```

```
_Decimal128 log10d128(_Decimal128 x);
#endif
```

### Description

- 2 The **log10** functions compute the base-10 (common) logarithm of *x*. A domain error occurs if the argument is less than zero. A pole error may occur if the argument is zero.

### Returns

- 3 The **log10** functions return  $\log_{10} x$ .

### 7.12.7.13 The **log10p1** functions

#### Synopsis

```
1 #include <math.h>
 double log10p1(double x);
 float log10p1f(float x);
 long double log10p1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 log10pd32(_Decimal32 x);
 _Decimal64 log10pd64(_Decimal64 x);
 _Decimal128 log10pd128(_Decimal128 x);
 #endif
```

### Description

- 2 The **log10p1** functions compute the base-10 logarithm of 1 plus the argument. A domain error occurs if the argument is less than  $-1$ . A pole error may occur if the argument equals  $-1$ . A range error occurs if nonzero *x* is too close to zero.

### Returns

- 3 The **log10p1** functions return  $\log_{10}(1 + x)$ .

### 7.12.7.14 The **log1p** and **logp1** functions

#### Synopsis

```
1 #include <math.h>
 double log1p(double x);
 float log1pf(float x);
 long double log1pl(long double x);
 double logp1(double x);
 float logp1f(float x);
 long double logp1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 log1pd32(_Decimal32 x);
 _Decimal64 log1pd64(_Decimal64 x);
 _Decimal128 log1pd128(_Decimal128 x);
 _Decimal32 logp1d32(_Decimal32 x);
 _Decimal64 logp1d64(_Decimal64 x);
 _Decimal128 logp1d128(_Decimal128 x);
 #endif
```

### Description

- 2 The **log1p** functions are equivalent to the **logp1** functions.<sup>281)</sup> These functions compute the base-*e* (natural) logarithm of 1 plus the argument.<sup>282)</sup> A domain error occurs if the argument is less than  $-1$ . A pole error may occur if the argument equals  $-1$ . A range error occurs if nonzero *x* is too close to zero.

<sup>281)</sup>The **logp1** functions are preferred for name consistency with the **log10p1** and **log2p1** functions.

<sup>282)</sup>For small magnitude *x*, **logp1**(*x*) is expected to be more accurate than **log**(1 + *x*).

**Returns**

- 3 The **log1p** and **logp1** functions return  $\log_e(1 + x)$ .

**7.12.7.15 The log2 functions****Synopsis**

```
1 #include <math.h>
 double log2(double x);
 float log2f(float x);
 long double log2l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 log2d32(_Decimal32 x);
 _Decimal64 log2d64(_Decimal64 x);
 _Decimal128 log2d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **log2** functions compute the base-2 logarithm of  $x$ . A domain error occurs if the argument is less than zero. A pole error may occur if the argument is zero.

**Returns**

- 3 The **log2** functions return  $\log_2 x$ .

**7.12.7.16 The log2p1 functions****Synopsis**

```
1 #include <math.h>
 double log2p1(double x);
 float log2p1f(float x);
 long double log2p1l(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 log2p1d32(_Decimal32 x);
 _Decimal64 log2p1d64(_Decimal64 x);
 _Decimal128 log2p1d128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **log2p1** functions compute the base-2 logarithm of 1 plus the argument. A domain error occurs if the argument is less than  $-1$ . A pole error may occur if the argument equals  $-1$ . A range error occurs if nonzero  $x$  is too close to zero.

**Returns**

- 3 The **log2p1** functions return  $\log_2(1+x)$ .

**7.12.7.17 The logb functions****Synopsis**

```
1 #include <math.h>
 double logb(double x);
 float logbf(float x);
 long double logbl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 logbd32(_Decimal32 x);
 _Decimal64 logbd64(_Decimal64 x);
 _Decimal128 logbd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **logb** functions extract the exponent of  $x$ , as a signed integer value in floating-point format. If  $x$  is subnormal it is treated as though it were normalized; thus, for positive finite  $x$ ,

$$1 \leq x \times b^{-\text{logb}(x)} < b$$

where  $b = \text{FLT\_RADIX}$  if the return type of the function is a standard floating type, or  $b = 10$  if the return type of the function is a decimal floating type. A domain error or pole error may occur if the argument is zero.

**Returns**

- 3 The **logb** functions return the signed exponent of  $x$ .

**7.12.7.18 The modf functions****Synopsis**

```
1 #include <math.h>
 double modf(double value, double *iptr);
 float modff(float value, float *iptr);
 long double modfl(long double value, long double *iptr);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 modfd32(_Decimal32 x, _Decimal32 *iptr);
 _Decimal64 modfd64(_Decimal64 x, _Decimal64 *iptr);
 _Decimal128 modfd128(_Decimal128 x, _Decimal128 *iptr);
 #endif
```

**Description**

- 2 The **modf** functions break the argument  $value$  into integral and fractional parts, each of which has the same type and sign as the argument. They store the integral part (in floating-point format) in the object pointed to by  $iptr$ .

**Returns**

- 3 The **modf** functions return the signed fractional part of  $value$ .

**7.12.7.19 The scalbn and scalbln functions****Synopsis**

```
1 #include <math.h>
 double scalbn(double x, int n);
 float scalbnf(float x, int n);
 long double scalbnl(long double x, int n);
 double scalbln(double x, long int n);
 float scalblnf(float x, long int n);
 long double scalblnl(long double x, long int n);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 scalbnd32(_Decimal32 x, int n);
 _Decimal64 scalbnd64(_Decimal64 x, int n);
 _Decimal128 scalbnd128(_Decimal128 x, int n);
 _Decimal32 scalblnd32(_Decimal32 x, long int n);
 _Decimal64 scalblnd64(_Decimal64 x, long int n);
 _Decimal128 scalblnd128(_Decimal128 x, long int n);
 #endif
```

**Description**

- 2 The **scalbn** and **scalbln** functions compute  $x \times b^n$ , where  $b = \text{FLT\_RADIX}$  if the return type of the function is a standard floating type, or  $b = 10$  if the return type of the function is a decimal floating type. A range error occurs for some finite  $x$ , depending on  $n$ .

**Returns**

- 3 The **scalbn** and **scalbln** functions return  $x \times b^n$ .

## 7.12.8 Power and absolute-value functions

### 7.12.8.1 The **cbrt** functions

#### Synopsis

```

1 #include <math.h>
 double cbrt(double x);
 float cbrtf(float x);
 long double cbrtl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 cbrtd32(_Decimal32 x);
 _Decimal64 cbrtd64(_Decimal64 x);
 _Decimal128 cbrtd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The **cbrt** functions compute the real cube root of  $x$ .

#### Returns

- 3 The **cbrt** functions return  $x^{\frac{1}{3}}$ .

### 7.12.8.2 The **compoundn** functions

#### Synopsis

```

1 #include <math.h>
 double compoundn(double x, long long int n);
 float compoundnf(float x, long long int n);
 long double compoundnl(long double x, long long int n);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 compoundnd32(_Decimal32 x, long long int n);
 _Decimal64 compoundnd64(_Decimal64 x, long long int n);
 _Decimal128 compoundnd128(_Decimal128 x, long long int n);
 #endif

```

#### Description

- 2 The **compoundn** functions compute 1 plus  $x$ , raised to the power  $n$ . A domain error occurs if  $x < -1$ . Depending on  $n$ , a range error occurs if either positive finite  $x$  is too large or if  $x$  is too near but not equal to  $-1$ . A pole error may occur if  $x$  equals  $-1$  and  $n < 0$ .

#### Returns

- 3 The **compoundn** functions return  $(1 + x)^n$ .

### 7.12.8.3 The **fabs** functions

#### Synopsis

```

1 #include <math.h>
 double fabs(double x);
 float fabsf(float x);
 long double fabsl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fabsd32(_Decimal32 x);
 _Decimal64 fabsd64(_Decimal64 x);
 _Decimal128 fabsd128(_Decimal128 x);
 #endif

```

#### Description

- 2 The **fabs** functions compute the absolute value of  $x$ .

**Returns**

- 3 The **fabs** functions return  $|x|$ .

**7.12.8.4 The **hypot** functions****Synopsis**

```
1 #include <math.h>
 double hypot(double x, double y);
 float hypotf(float x, float y);
 long double hypotl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 hypotd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 hypotd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 hypotd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **hypot** functions compute the square root of the sum of the squares of  $x$  and  $y$ , without undue overflow or underflow. A range error occurs for some finite arguments.

**Returns**

- 3 The **hypot** functions return  $\sqrt{x^2 + y^2}$ .

**7.12.8.5 The **pow** functions****Synopsis**

```
1 #include <math.h>
 double pow(double x, double y);
 float powf(float x, float y);
 long double powl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 powd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 powd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 powd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **pow** functions compute  $x$  raised to the power  $y$ . A domain error occurs if  $x$  is finite and less than zero and  $y$  is finite and not an integer value. A domain error may occur if  $x$  is zero and  $y$  is zero. Depending on  $y$ , a range error occurs if either the magnitude of nonzero finite  $x$  is too large or too near zero. A domain error or pole error may occur if  $x$  is zero and  $y$  is less than zero.

**Returns**

- 3 The **pow** functions return  $x^y$ .

**7.12.8.6 The **pown** functions****Synopsis**

```
1 #include <math.h>
 double pown(double x, long long int n);
 float pownf(float x, long long int n);
 long double pownl(long double x, long long int n);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 pownd32(_Decimal32 x, long long int n);
 _Decimal64 pownd64(_Decimal64 x, long long int n);
 _Decimal128 pownd128(_Decimal128 x, long long int n);
 #endif
```

**Description**

- 2 The **pow** functions compute  $x$  raised to the  $n^{\text{th}}$  power. A pole error may occur if  $x$  equals 0 and  $n < 0$ . Depending on  $n$ , a range error occurs if either the magnitude of nonzero finite  $x$  is too large or too near zero.

**Returns**

- 3 The **pow** functions return  $x^n$ .

**7.12.8.7 The **pow** functions****Synopsis**

```
1 #include <math.h>
 double pow(double y, double x);
 float powf(float y, float x);
 long double powl(long double y, long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 powrd32(_Decimal32 y, _Decimal32 x);
 _Decimal64 powrd64(_Decimal64 y, _Decimal64 x);
 _Decimal128 powrd128(_Decimal128 y, _Decimal128 x);
 #endif
```

**Description**

- 2 The **pow** functions compute  $x$  raised to the power  $y$  as  $e^{y \log_e x}$ .<sup>283)</sup> A domain error occurs if  $x < 0$  or if  $x$  and  $y$  are both zero. Depending on  $y$ , a range error occurs if either positive nonzero finite  $x$  is too large or too near zero. A pole error may occur if  $x$  equals zero and finite  $y < 0$ .

**Returns**

- 3 The **pow** functions return  $e^{y \log_e x}$ .

**7.12.8.8 The **rootn** functions****Synopsis**

```
1 #include <math.h>
 double rootn(double x, long long int n);
 float rootnf(float x, long long int n);
 long double rootnl(long double x, long long int n);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 rootnd32(_Decimal32 x, long long int n);
 _Decimal64 rootnd64(_Decimal64 x, long long int n);
 _Decimal128 rootnd128(_Decimal128 x, long long int n);
 #endif
```

**Description**

- 2 The **rootn** functions compute the principal  $n^{\text{th}}$  root of  $x$ . A domain error occurs if  $n$  is 0 or if  $x < 0$  and  $n$  is even. If  $n$  is  $-1$ , a range error occurs if either the magnitude of nonzero finite  $x$  is too large or too near zero. A pole error may occur if  $x$  equals zero and  $n < 0$ .

**Returns**

- 3 The **rootn** functions return  $x^{\frac{1}{n}}$ .

**7.12.8.9 The **rsqrt** functions****Synopsis**

```
1 #include <math.h>
 double rsqrt(double x);
 float rsqrtf(float x);
```

<sup>283)</sup> Restricting the domain to that of the formula  $e^{y \log_e x}$  is intended to better meet expectations for a continuous power function and to allow implementations with fewer tests for special cases.

```

long double rsqrtl(long double x);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 rsqrt32(_Decimal32 x);
_Decimal64 rsqrt64(_Decimal64 x);
_Decimal128 rsqrt128(_Decimal128 x);
#endif

```

### Description

- 2 The **rsqrt** functions compute the reciprocal of the nonnegative square root of the argument. A domain error occurs if the argument is less than zero. A pole error may occur if the argument equals zero.

### Returns

- 3 The **rsqrt** functions return  $\frac{1}{\sqrt{x}}$ .

## 7.12.8.10 The **sqrt** functions

### Synopsis

```

1 #include <math.h>
double sqrt(double x);
float sqrtf(float x);
long double sqrtl(long double x);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 sqrt32(_Decimal32 x);
_Decimal64 sqrt64(_Decimal64 x);
_Decimal128 sqrt128(_Decimal128 x);
#endif

```

### Description

- 2 The **sqrt** functions compute the nonnegative square root of  $x$ . A domain error occurs if the argument is less than zero.

### Returns

- 3 The **sqrt** functions return  $\sqrt{x}$ .

## 7.12.9 Error and gamma functions

### 7.12.9.1 The **erf** functions

#### Synopsis

```

1 #include <math.h>
double erf(double x);
float erff(float x);
long double erfl(long double x);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 erfd32(_Decimal32 x);
_Decimal64 erfd64(_Decimal64 x);
_Decimal128 erfd128(_Decimal128 x);
#endif

```

### Description

- 2 The **erf** functions compute the error function of  $x$ . A range error occurs if nonzero  $x$  is too close to zero.

### Returns

- 3 The **erf** functions return  $\text{erf } x = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$ .



7.12.9.2 The **erfc** functions

## Synopsis

```

1 #include <math.h>
 double erfc(double x);
 float erfcf(float x);
 long double erfcl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 erfcd32(_Decimal32 x);
 _Decimal64 erfcd64(_Decimal64 x);
 _Decimal128 erfcd128(_Decimal128 x);
 #endif

```

## Description

- 2 The **erfc** functions compute the complementary error function of  $x$ . A range error occurs if positive finite  $x$  is too large.

## Returns

- 3 The **erfc** functions return  $\operatorname{erfc} x = 1 - \operatorname{erf} x = \frac{2}{\sqrt{\pi}} \int_x^{\infty} e^{-t^2} dt$ .

7.12.9.3 The **lgamma** functions

## Synopsis

```

1 #include <math.h>
 double lgamma(double x);
 float lgammaf(float x);
 long double lgammal(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 lgammad32(_Decimal32 x);
 _Decimal64 lgammad64(_Decimal64 x);
 _Decimal128 lgammad128(_Decimal128 x);
 #endif

```

## Description

- 2 The **lgamma** functions compute the natural logarithm of the absolute value of gamma of  $x$ . A range error occurs if positive finite  $x$  is too large. A pole error may occur if  $x$  is a negative integer or zero.

## Returns

- 3 The **lgamma** functions return  $\log_e |\Gamma(x)|$ .

7.12.9.4 The **tgamma** functions

## Synopsis

```

1 #include <math.h>
 double tgamma(double x);
 float tgammaf(float x);
 long double tgammal(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 tgammad32(_Decimal32 x);
 _Decimal64 tgammad64(_Decimal64 x);
 _Decimal128 tgammad128(_Decimal128 x);
 #endif

```

## Description

- 2 The **tgamma** functions compute the gamma function of  $x$ . A domain error or pole error may occur if  $x$  is a negative integer or zero. A range error occurs for some finite  $x$  less than zero, if positive finite  $x$  is too large, or nonzero  $x$  is too close to zero.

**Returns**

- 3 The **tgamma** functions return  $\Gamma(x)$ .

**7.12.10 Nearest integer functions****7.12.10.1 The **ceil** functions****Synopsis**

```
1 #include <math.h>
 double ceil(double x);
 float ceilf(float x);
 long double ceill(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 ceild32(_Decimal32 x);
 _Decimal64 ceild64(_Decimal64 x);
 _Decimal128 ceild128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **ceil** functions compute the smallest integer value not less than  $x$ .

**Returns**

- 3 The **ceil** functions return  $\lceil x \rceil$ , expressed as a floating-point number.

**7.12.10.2 The **floor** functions****Synopsis**

```
1 #include <math.h>
 double floor(double x);
 float floorf(float x);
 long double floorl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 floord32(_Decimal32 x);
 _Decimal64 floord64(_Decimal64 x);
 _Decimal128 floord128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **floor** functions compute the largest integer value not greater than  $x$ .

**Returns**

- 3 The **floor** functions return  $\lfloor x \rfloor$ , expressed as a floating-point number.

**7.12.10.3 The **nearbyint** functions****Synopsis**

```
1 #include <math.h>
 double nearbyint(double x);
 float nearbyintf(float x);
 long double nearbyintl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 nearbyintd32(_Decimal32 x);
 _Decimal64 nearbyintd64(_Decimal64 x);
 _Decimal128 nearbyintd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **nearbyint** functions round their argument to an integer value in floating-point format, using the current rounding direction and without raising the “inexact” floating-point exception.

**Returns**

- 3 The **nearbyint** functions return the rounded integer value.

**7.12.10.4 The rint functions****Synopsis**

```
1 #include <math.h>
 double rint(double x);
 float rintf(float x);
 long double rintl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 rintd32(_Decimal32 x);
 _Decimal64 rintd64(_Decimal64 x);
 _Decimal128 rintd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **rint** functions differ from the **nearbyint** functions (7.12.10.3) only in that the **rint** functions may raise the “inexact” floating-point exception if the result differs in value from the argument.

**Returns**

- 3 The **rint** functions return the rounded integer value.

**7.12.10.5 The lrint and llrint functions****Synopsis**

```
1 #include <math.h>
 long int lrint(double x);
 long int lrintf(float x);
 long int rintl(long double x);
 long long int llrint(double x);
 long long int llrintf(float x);
 long long int llrintl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 long int lrintd32(_Decimal32 x);
 long int lrintd64(_Decimal64 x);
 long int lrintd128(_Decimal128 x);
 long long int llrintd32(_Decimal32 x);
 long long int llrintd64(_Decimal64 x);
 long long int llrintd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **lrint** and **llrint** functions round their argument to the nearest integer value, rounding according to the current rounding direction. If the rounded value is outside the range of the return type, the numeric result is unspecified and a domain error or range error may occur.

**Returns**

- 3 The **lrint** and **llrint** functions return the rounded integer value.

**7.12.10.6 The round functions****Synopsis**

```
1 #include <math.h>
 double round(double x);
 float roundf(float x);
 long double roundl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 roundd32(_Decimal32 x);
```

```

_Decimal64 roundd64(_Decimal64 x);
_Decimal128 roundd128(_Decimal128 x);
#endif

```

### Description

- 2 The **round** functions round their argument to the nearest integer value in floating-point format, rounding halfway cases away from zero, regardless of the current rounding direction.

### Returns

- 3 The **round** functions return the rounded integer value.

#### 7.12.10.7 The **lround** and **llround** functions

##### Synopsis

```

1 #include <math.h>
 long int lround(double x);
 long int lroundf(float x);
 long int lroundl(long double x);
 long long int llround(double x);
 long long int llroundf(float x);
 long long int llroundl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 long int lroundd32(_Decimal32 x);
 long int lroundd64(_Decimal64 x);
 long int lroundd128(_Decimal128 x);
 long long int llroundd32(_Decimal32 x);
 long long int llroundd64(_Decimal64 x);
 long long int llroundd128(_Decimal128 x);
 #endif

```

### Description

- 2 The **lround** and **llround** functions round their argument to the nearest integer value, rounding halfway cases away from zero, regardless of the current rounding direction. If the rounded value is outside the range of the return type, the numeric result is unspecified and a domain error or range error may occur.

### Returns

- 3 The **lround** and **llround** functions return the rounded integer value.

#### 7.12.10.8 The **roundeven** functions

##### Synopsis

```

1 #include <math.h>
 double roundeven(double x);
 float roundevenf(float x);
 long double roundevenl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 roundevend32(_Decimal32 x);
 _Decimal64 roundevend64(_Decimal64 x);
 _Decimal128 roundevend128(_Decimal128 x);
 #endif

```

### Description

- 2 The **roundeven** functions round their argument to the nearest integer value in floating-point format, rounding halfway cases to even (that is, to the nearest value that is an even integer), regardless of the current rounding direction.

**Returns**

- 3 The **roundeven** functions return the rounded integer value.

**7.12.10.9 The **trunc** functions****Synopsis**

```
1 #include <math.h>
 double trunc(double x);
 float truncf(float x);
 long double truncfpl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 truncd32(_Decimal32 x);
 _Decimal64 truncd64(_Decimal64 x);
 _Decimal128 truncd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **trunc** functions round their argument to the integer value, in floating format, nearest to but no larger in magnitude than the argument.

**Returns**

- 3 The **trunc** functions return the truncated integer value.

**7.12.10.10 The **fromfp** and **ufromfp** functions****Synopsis**

```
1 #include <math.h>
 double fromfp(double x, int rnd, unsigned int width);
 float fromfpf(float x, int rnd, unsigned int width);
 long double fromfppl(long double x, int rnd, unsigned int width);
 double ufromfp(double x, int rnd, unsigned int width);
 float ufromfpf(float x, int rnd, unsigned int width);
 long double ufromfppl(long double x, int rnd, unsigned int width);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fromfpd32(_Decimal32 x, int rnd, unsigned int width);
 _Decimal64 fromfpd64(_Decimal64 x, int rnd, unsigned int width);
 _Decimal128 fromfpd128(_Decimal128 x, int rnd, unsigned int width);
 _Decimal32 ufromfpd32(_Decimal32 x, int rnd, unsigned int width);
 _Decimal64 ufromfpd64(_Decimal64 x, int rnd, unsigned int width);
 _Decimal128 ufromfpd128(_Decimal128 x, int rnd, unsigned int width);
 #endif
```

**Description**

- 2 The **fromfp** and **ufromfp** functions round *x*, using the math rounding direction indicated by *rnd*, to a signed or unsigned integer, respectively. If *width* is nonzero and the resulting integer is within the range

—  $[-2^{(\text{width}-1)}, 2^{(\text{width}-1)} - 1]$ , for signed

—  $[0, 2^{\text{width}} - 1]$ , for unsigned

then the functions return the integer value (represented in floating type). Otherwise, if *width* is zero or *x* does not round to an integer within the range, the functions return a NaN (of the type of the *x* argument), if available, else the value of *x*, and a domain error occurs. If the value of the *rnd* argument is not equal to the value of a math rounding direction macro (7.12), the direction of rounding is unspecified. The **fromfp** and **ufromfp** functions do not raise the “inexact” floating-point exception.

**Returns**

- 3 The **fromfp** and **ufromfp** functions return the rounded integer value.
- 4 EXAMPLE 1 Upward rounding of **double** *x* to type **int**, without raising the “inexact” floating-point exception, is achieved by

```
(int)fromfp(x, FP_INT_UPWARD, INT_WIDTH)
```

- 5 EXAMPLE 2 Unsigned integer wrapping is not performed in

```
ufromfp(-3.0, FP_INT_UPWARD, UINT_WIDTH) /* domain error */
```

**7.12.10.11 The fromfpx and ufromfpx functions****Synopsis**

```
1 #include <math.h>
 double fromfpx(double x, int rnd, unsigned int width);
 float fromfpxf(float x, int rnd, unsigned int width);
 long double fromfpxl(long double x, int rnd, unsigned int width);
 double ufromfpx(double x, int rnd, unsigned int width);
 float ufromfpxf(float x, int rnd, unsigned int width);
 long double ufromfpxl(long double x, int rnd, unsigned int width);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fromfpxd32(_Decimal32 x, int rnd, unsigned int width);
 _Decimal64 fromfpxd64(_Decimal64 x, int rnd, unsigned int width);
 _Decimal128 fromfpxd128(_Decimal128 x, int rnd, unsigned int width);
 _Decimal32 ufromfpxd32(_Decimal32 x, int rnd, unsigned int width);
 _Decimal64 ufromfpxd64(_Decimal64 x, int rnd, unsigned int width);
 _Decimal128 ufromfpxd128(_Decimal128 x, int rnd, unsigned int width);
 #endif
```

**Description**

- 2 The **fromfpx** and **ufromfpx** functions differ from the **fromfp** and **ufromfp** functions, respectively, only in that the **fromfpx** and **ufromfpx** functions raise the “inexact” floating-point exception if a rounded result not exceeding the specified width differs in value from the argument *x*.

**Returns**

- 3 The **fromfpx** and **ufromfpx** functions return the rounded integer value.
- 4 NOTE Conversions to integer types that are not required to raise the inexact exception can be done simply by rounding to integral value in floating type and then converting to the target integer type. For example, the conversion of **long double** *x* to **uint64\_t**, using upward rounding, is done by

```
(uint64_t)ceil(x)
```

**7.12.11 Remainder functions****7.12.11.1 The fmod functions****Synopsis**

```
1 #include <math.h>
 double fmod(double x, double y);
 float fmodf(float x, float y);
 long double fmodl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fmodd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fmodd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fmodd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **fmod** functions compute the floating-point remainder of  $x/y$ .

**Returns**

- 3 The **fmod** functions return the value  $x - ny$ , for some integer  $n$  such that, if  $y$  is nonzero, the result has the same sign as  $x$  and magnitude less than the magnitude of  $y$ . If  $y$  is zero, whether a domain error occurs or the **fmod** functions return zero is implementation-defined.

**7.12.11.2 The remainder functions****Synopsis**

```
1 #include <math.h>
 double remainder(double x, double y);
 float remainderf(float x, float y);
 long double remainderl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 remainderd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 remainderd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 remainderd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **remainder** functions compute the remainder  $x \text{ REM } y$  required by ISO/IEC 60559.<sup>284)</sup>

**Returns**

- 3 The **remainder** functions return  $x \text{ REM } y$ . If  $y$  is zero, whether a domain error occurs or the functions return zero is implementation-defined.

**7.12.11.3 The remquo functions****Synopsis**

```
1 #include <math.h>
 double remquo(double x, double y, int *quo);
 float remquof(float x, float y, int *quo);
 long double remquol(long double x, long double y, int *quo);
```

**Description**

- 2 The **remquo** functions compute the same remainder as the **remainder** functions. In the object pointed to by **quo** they store a value whose magnitude is congruent modulo  $2^n$  to the magnitude of the integral quotient of  $x/y$ , where  $n$  is an implementation-defined integer greater than or equal to 3. If the value stored is not zero, its sign is the sign of  $x/y$ .

**Returns**

- 3 The **remquo** functions return  $x \text{ REM } y$ . If  $y$  is zero, the value stored in the object pointed to by **quo** is unspecified and whether a domain error occurs or the functions return zero is implementation-defined.
- 4 NOTE There are no decimal floating-point versions of the **remquo** functions.

**7.12.12 Manipulation functions****7.12.12.1 The copysign functions****Synopsis**

```
1 #include <math.h>
 double copysign(double x, double y);
```

<sup>284)</sup>When  $y \neq 0$ , the remainder  $r = x \text{ REM } y$  is defined regardless of the rounding mode by the mathematical relation  $r = x - ny$ , where  $n$  is the integer nearest the exact value of  $\frac{x}{y}$ ; whenever  $|n - \frac{x}{y}| = \frac{1}{2}$ , then  $n$  is even. If  $r = 0$ , its sign shall be that of  $x$ . This definition is applicable for all implementations.

```

float copysignf(float x, float y);
long double copysignl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 copysign32(_Decimal32 x, _Decimal32 y);
_Decimal64 copysign64(_Decimal64 x, _Decimal64 y);
_Decimal128 copysign128(_Decimal128 x, _Decimal128 y);
#endif

```

### Description

- 2 The **copysign** functions produce a value with the magnitude of *x* and the sign of *y*. If *x* or *y* is an unsigned value, the sign (if any) of the result is implementation-defined. On implementations that represent a signed zero but do not treat negative zero consistently in arithmetic operations, the **copysign** functions should regard the sign of zero as positive.

### Returns

- 3 The **copysign** functions return a value with the magnitude of *x* and the sign of *y*.

#### 7.12.12.2 The **nan** functions

##### Synopsis

```

1 #include <math.h>
double nan(const char *tagp);
float nanf(const char *tagp);
long double nanl(const char *tagp);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 nand32(const char *tagp);
_Decimal64 nand64(const char *tagp);
_Decimal128 nand128(const char *tagp);
#endif

```

### Description

- 2 The **nan**, **nanf**, and **nanl** functions convert the string pointed to by *tagp* according to the following rules. The call **nan**("n-char-sequence") is equivalent to **strtod**("NAN(n-char-sequence)", **nullptr**); the call **nan**("") is equivalent to **strtod**("NAN()", **nullptr**). If *tagp* does not point to an empty string or an n-char sequence, the call is equivalent to **strtod**("NAN", **nullptr**). Calls to **nanf** and **nanl** are equivalent to the corresponding calls to **strtof** and **strtold**.

### Returns

- 3 The **nan** functions return a quiet NaN, if available, with content indicated through *tagp*. If the implementation does not support quiet NaNs, the functions return zero.

Forward references: the **strtod**, **strtof**, and **strtold** functions (7.24.2.6).

#### 7.12.12.3 The **nextafter** functions

##### Synopsis

```

1 #include <math.h>
double nextafter(double x, double y);
float nextafterf(float x, float y);
long double nextafterl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 nextafterd32(_Decimal32 x, _Decimal32 y);
_Decimal64 nextafterd64(_Decimal64 x, _Decimal64 y);
_Decimal128 nextafterd128(_Decimal128 x, _Decimal128 y);
#endif

```



**Description**

- 2 The **nextafter** functions determine the next representable value, in the return type of the function, after *x* in the direction of *y*, where *x* and *y* are first converted to the return type of the function.<sup>285)</sup> The **nextafter** functions return *y* if *x* equals *y*.

A range error occurs if the magnitude of *x* is the largest finite value representable in the type and the result is infinite or not representable in the type. If *x* != *y*, a range error occurs for either subnormal or zero results.

**Returns**

- 3 The **nextafter** functions return the next representable value in the specified format after *x* in the direction of *y*.

**7.12.12.4 The nexttoward functions****Synopsis**

```
1 #include <math.h>
 double nexttoward(double x, long double y);
 float nexttowardf(float x, long double y);
 long double nexttowardl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 nexttowarddd32(_Decimal32 x, _Decimal128 y);
 _Decimal64 nexttowarddd64(_Decimal64 x, _Decimal128 y);
 _Decimal128 nexttowarddd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **nexttoward** functions are equivalent to the **nextafter** functions except that the second parameter has type **long double** or **\_Decimal128** and the functions return *y* converted to the return type of the function if *x* equals *y*.<sup>286)</sup>

**Returns**

- 3 The **nexttoward** functions return the next representable value in the specified format after *x* in the direction of *y*.

**7.12.12.5 The nextup functions****Synopsis**

```
1 #include <math.h>
 double nextup(double x);
 float nextupf(float x);
 long double nextupl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 nextupdd32(_Decimal32 x);
 _Decimal64 nextupdd64(_Decimal64 x);
 _Decimal128 nextupdd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **nextup** functions determine the next representable value, in the return type of the function, greater than *x*. If *x* is the negative number of least magnitude in the type of *x*, **nextup**(*x*) is −0 if the type has signed zeros and is 0 otherwise. If *x* is zero, **nextup**(*x*) is the positive number of least magnitude in the type of *x*. If *x* is the positive number (finite or infinite) of maximum magnitude in the type, **nextup**(*x*) is *x*.

<sup>285)</sup>The argument values are converted to the return type of the function, even by a macro implementation of the function.

<sup>286)</sup>The result of the **nexttoward** functions is determined in the return type of the function, without loss of range or precision in a floating second argument.

**Returns**

- 3 The nextup functions return the next representable value in the specified type greater than *x*.

**7.12.12.6 The `nextdown` functions****Synopsis**

```
1 #include <math.h>
 double nextdown(double x);
 float nextdownf(float x);
 long double nextdownl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 nextdownd32(_Decimal32 x);
 _Decimal64 nextdownd64(_Decimal64 x);
 _Decimal128 nextdownd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The `nextdown` functions determine the next representable value, in the return type of the function, less than *x*. If *x* is the positive number of least magnitude in the type of *x*, `nextdown(x)` is +0 if the type has signed zeros and is 0 otherwise. If *x* is zero, `nextdown(x)` is the negative number of least magnitude in the type of *x*. If *x* is the negative number (finite or infinite) of maximum magnitude in the type, `nextdown(x)` is *x*.

**Returns**

- 3 The `nextdown` functions return the next representable value in the specified type less than *x*.

**7.12.12.7 The `canonicalize` functions****Synopsis**

```
1 #include <math.h>
 int canonicalize(double *cx, const double *x);
 int canonicalizef(float *cx, const float *x);
 int canonicalizel(long double *cx, const long double *x);
 #ifdef __STDC_IEC_60559_DFP__
 int canonicalized32(_Decimal32 *cx, const _Decimal32 *x);
 int canonicalized64(_Decimal64 *cx, const _Decimal64 *x);
 int canonicalized128(_Decimal128 *cx, const _Decimal128 *x);
 #endif
```

**Description**

- 2 The `canonicalize` functions attempt to produce a canonical version of the floating-point representation in the object pointed to by the argument *x*, as if to a temporary object of the specified type, and store the canonical result in the object pointed to by the argument *cx*.<sup>287)</sup> If the input *\*x* is a signaling NaN, the `canonicalize` functions are intended to store a canonical quiet NaN. If a canonical result is not produced the object pointed to by *cx* is unchanged.

**Returns**

- 3 The `canonicalize` functions return zero if a canonical result is stored in the object pointed to by *cx*. Otherwise they return a nonzero value.

**7.12.13 Maximum, minimum, and positive difference functions****7.12.13.1 The `fdim` functions****Synopsis**

```
1 #include <math.h>
 double fdim(double x, double y);
```

<sup>287)</sup>Arguments *x* and *cx* can point to the same object.

```
float fdimf(float x, float y);
long double fdiml(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fdimd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fdimd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fdimd128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- 2 The **fdim** functions determine the *positive difference* between their arguments:

$$\begin{cases} x - y & \text{if } x > y \\ +0 & \text{if } x \leq y \end{cases}$$

A range error occurs for some finite arguments.

### Returns

- 3 The **fdim** functions return the positive difference value.

### 7.12.13.2 The **fmax** functions

#### Synopsis

```
1 #include <math.h>
double fmax(double x, double y);
float fmaxf(float x, float y);
long double fmaxl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fmaxd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaxd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaxd128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- 2 The **fmax** functions determine the maximum numeric value of their arguments.<sup>288)</sup>

### Returns

- 3 The **fmax** functions return the maximum numeric value of their arguments.

### 7.12.13.3 The **fmin** functions

#### Synopsis

```
1 #include <math.h>
double fmin(double x, double y);
float fminf(float x, float y);
long double fminl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fmind32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmind64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmind128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- 2 The **fmin** functions determine the minimum numeric value of their arguments.<sup>289)</sup>

<sup>288)</sup>Quiet NaN arguments are treated as missing data: if one argument is a quiet NaN and the other numeric, then the **fmax** functions choose the numeric value. See F.10.10.2.

<sup>289)</sup>The **fmin** functions are analogous to the **fmax** functions in their treatment of quiet NaNs.

**Returns**

- 3 The **fmin** functions return the minimum numeric value of their arguments.
- 4 NOTE The **fmax** and **fmin** functions are similar to the **fmaximum\_num** and **fminimum\_num** functions, though can differ in which signed zero is returned when the arguments are differently signed zeros and in their treatment of signaling NaNs (see F.10.10.5).

**7.12.13.4 The fmaximum functions****Synopsis**

```
1 #include <math.h>
 double fmaximum(double x, double y);
 float fmaximumf(float x, float y);
 long double fmaximuml(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fmaximumd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fmaximumd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fmaximumd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **fmaximum** functions determine the maximum value of their arguments. For these functions, +0 is considered greater than −0. These functions differ from the **fmaximum\_num** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

**Returns**

- 3 The **fmaximum** functions return the maximum value of their arguments.

**7.12.13.5 The fminimum functions****Synopsis**

```
1 #include <math.h>
 double fminimum(double x, double y);
 float fminimumf(float x, float y);
 long double fminimuml(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fminimumd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fminimumd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fminimumd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **fminimum** functions determine the minimum value of their arguments. For these functions, −0 is considered less than +0. These functions differ from the **fminimum\_num** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

**Returns**

- 3 The **fminimum** functions return the minimum value of their arguments.

**7.12.13.6 The fmaximum\_mag functions****Synopsis**

```
1 #include <math.h>
 double fmaximum_mag(double x, double y);
 float fmaximum_magf(float x, float y);
 long double fmaximum_magl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fmaximum_magd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fmaximum_magd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fmaximum_magd128(_Decimal128 x, _Decimal128 y);
 #endif
```

```
#endif
```

### Description

- 2 The **fmaximum\_mag** functions determine the value of the argument of maximum magnitude:  $x$  if  $|x| > |y|$ ,  $y$  if  $|y| > |x|$ , and **fmaximum**( $x$ ,  $y$ ) otherwise. These functions differ from the **fmaximum\_mag\_num** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fmaximum\_mag** functions return the value of the argument of maximum magnitude.

### 7.12.13.7 The **fminimum\_mag** functions

#### Synopsis

```
1 #include <math.h>
 double fminimum_mag(double x, double y);
 float fminimum_magf(float x, float y);
 long double fminimum_magl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fminimum_magd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fminimum_magd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fminimum_magd128(_Decimal128 x, _Decimal128 y);
 #endif
```

### Description

- 2 The **fminimum\_mag** functions determine the value of the argument of minimum magnitude:  $x$  if  $|x| < |y|$ ,  $y$  if  $|y| < |x|$ , and **fminimum**( $x$ ,  $y$ ) otherwise. These functions differ from the **fminimum\_mag\_num** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fminimum\_mag** functions return the value of the argument of minimum magnitude.

### 7.12.13.8 The **fmaximum\_num** functions

#### Synopsis

```
1 #include <math.h>
 double fmaximum_num(double x, double y);
 float fmaximum_numf(float x, float y);
 long double fmaximum_numl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fmaximum_numd32(_Decimal32 x, _Decimal32 y);
 _Decimal64 fmaximum_numd64(_Decimal64 x, _Decimal64 y);
 _Decimal128 fmaximum_numd128(_Decimal128 x, _Decimal128 y);
 #endif
```

### Description

- 2 The **fmaximum\_num** functions determine the maximum value of their numeric arguments. They determine the number if one argument is a number and the other is a NaN. These functions differ from the **fmaximum** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fmaximum\_num** functions return the maximum value of their numeric arguments.

### 7.12.13.9 The **fminimum\_num** functions

#### Synopsis

```
1 #include <math.h>
 double fminimum_num(double x, double y);
 float fminimum_numf(float x, float y);
```

```

long double fminimum_numl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fminimum_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimum_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimum_numd128(_Decimal128 x, _Decimal128 y);
#endif

```

### Description

- 2 The **fminimum\_num** functions determine the minimum value of their numeric arguments. They determine the number if one argument is a number and the other is a NaN. These functions differ from the **fminimum** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fminimum\_num** functions return the minimum value of their numeric arguments.

### 7.12.13.10 The **fmaximum\_mag\_num** functions

#### Synopsis

```

1 #include <math.h>
double fmaximum_mag_num(double x, double y);
float fmaximum_mag_numf(float x, float y);
long double fmaximum_mag_numl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fmaximum_mag_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaximum_mag_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaximum_mag_numd128(_Decimal128 x, _Decimal128 y);
#endif

```

### Description

- 2 The **fmaximum\_mag\_num** functions determine the value of a numeric argument of maximum magnitude. They determine the number if one argument is a number and the other is a NaN. These functions differ from the **fmaximum\_mag** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fmaximum\_mag\_num** functions return the value of a numeric argument of maximum magnitude.

### 7.12.13.11 The **fminimum\_mag\_num** functions

#### Synopsis

```

1 #include <math.h>
double fminimum_mag_num(double x, double y);
float fminimum_mag_numf(float x, float y);
long double fminimum_mag_numl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 fminimum_mag_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimum_mag_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimum_mag_numd128(_Decimal128 x, _Decimal128 y);
#endif

```

### Description

- 2 The **fminimum\_mag\_num** functions determine the value of a numeric argument of minimum magnitude. They determine the number if one argument is a number and the other is a NaN. These functions differ from the **fminimum\_mag** functions only in their treatment of NaN arguments (see F.10.10.4, F.10.10.5).

### Returns

- 3 The **fminimum\_mag\_num** functions return the value of a numeric argument of minimum magnitude.

## 7.12.14 Fused multiply-add

### 7.12.14.1 The **fma** functions

#### Synopsis

```

1 #include <math.h>
 double fma(double x, double y, double z);
 float fmaf(float x, float y, float z);
 long double fmal(long double x, long double y, long double z);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 fmad32(_Decimal32 x, _Decimal32 y, _Decimal32 z);
 _Decimal64 fmad64(_Decimal64 x, _Decimal64 y, _Decimal64 z);
 _Decimal128 fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
 #endif

```

#### Description

- 2 The **fma** functions compute  $(x \times y) + z$ , rounded as one ternary operation: they compute the value (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error occurs for some infinite arguments.

#### Returns

- 3 The **fma** functions return  $(x \times y) + z$ , rounded as one ternary operation.

## 7.12.15 Functions that round result to narrower type

### 7.12.15.1 General

- 1 The functions in this subclause round their results to the return type, which is typically narrower<sup>290</sup> than the parameter types.

### 7.12.15.2 Add and round to narrower type

#### Synopsis

```

1 #include <math.h>
 float fadd(double x, double y);
 float faddl(long double x, long double y);
 double daddl(long double x, long double y);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 d32add64(_Decimal64 x, _Decimal64 y);
 _Decimal32 d32add128(_Decimal128 x, _Decimal128 y);
 _Decimal64 d64add128(_Decimal128 x, _Decimal128 y);
 #endif

```

#### Description

- 2 These functions compute the sum of  $x + y$ , rounded to the return type of the function. They compute the sum (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error may occur for infinite arguments.

#### Returns

- 3 These functions return the sum of  $x + y$ , rounded to the return type of the function.

### 7.12.15.3 Subtract and round to narrower type

#### Synopsis

```

1 #include <math.h>
 float fsub(double x, double y);

```

<sup>290</sup>In some cases the destination type can sometimes not be narrower than the parameter types. For example, **double** potentially is not narrower than **long double**.

```
float fsubl(long double x, long double y);
double dsubl(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 d32subd64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32subd128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64subd128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- These functions compute the difference of  $x - y$ , rounded to the return type of the function. They compute the difference (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error may occur for infinite arguments.

### Returns

- These functions return the difference of  $x - y$ , rounded to the return type of the function.

#### 7.12.15.4 Multiply and round to narrower type

##### Synopsis

```
1 #include <math.h>
float fmul(double x, double y);
float fmul_l(long double x, long double y);
double dmul(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 d32muld64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32muld128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64muld128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- These functions compute the product  $x \times y$ , rounded to the return type of the function. They compute the product (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error occurs for one infinite argument and one zero argument.

### Returns

- These functions return the product of  $x \times y$ , rounded to the return type of the function.

#### 7.12.15.5 Divide and round to narrower type

##### Synopsis

```
1 #include <math.h>
float fdiv(double x, double y);
float fdiv_l(long double x, long double y);
double ddiv(long double x, long double y);
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 d32divd64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32divd128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64divd128(_Decimal128 x, _Decimal128 y);
#endif
```

### Description

- These functions compute the quotient  $x \div y$ , rounded to the return type of the function. They compute the quotient (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error occurs for either both arguments infinite or both arguments zero. A pole error occurs for a finite  $x$  and a zero  $y$ .



**Returns**

- 3 These functions return the quotient  $x \div y$ , rounded to the return type of the function.

**7.12.15.6 Fused multiply-add and round to narrower type****Synopsis**

```
1 #include <math.h>
 float fmma(double x, double y, double z);
 float fmmal(long double x, long double y, long double z);
 double dfma(long double x, long double y, long double z);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 d32fmad64(_Decimal64 x, _Decimal64 y, _Decimal64 z);
 _Decimal32 d32fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
 _Decimal64 d64fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
 #endif
```

**Description**

- 2 These functions compute  $(x \times y) + z$  (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite arguments. A domain error may occur for an infinite argument.

**Returns**

- 3 These functions return  $(x \times y) + z$ , rounded to the return type of the function.

**7.12.15.7 Square root rounded to narrower type****Synopsis**

```
1 #include <math.h>
 float fsqrt(double x);
 float fsqrtl(long double x);
 double dsqrtl(long double x);
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 d32sqrtd64(_Decimal64 x);
 _Decimal32 d32sqrtd128(_Decimal128 x);
 _Decimal64 d64sqrtd128(_Decimal128 x);
 #endif
```

**Description**

- 2 These functions compute the square root of  $x$ , rounded to the return type of the function. They compute the square root (as if) to infinite precision and round once to the return type, according to the current rounding mode. A range error occurs for some finite positive arguments. A domain error occurs if the argument is less than zero.

**Returns**

- 3 These functions return the nonnegative square root of  $x$ , rounded to the return type of the function.

**7.12.16 Quantum and quantum exponent functions****7.12.16.1 The `quantizedN` functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 quantized32(_Decimal32 x, _Decimal32 y);
 _Decimal64 quantized64(_Decimal64 x, _Decimal64 y);
 _Decimal128 quantized128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **quantized***N* functions compute, if possible, a value with the numerical value of *x* and the quantum exponent of *y*. If the quantum exponent is being increased, the value shall be correctly rounded; if the result does not have the same value as *x*, the “inexact” floating-point exception shall be raised. If the quantum exponent is being decreased and the significand of the result has more digits than the type would allow, the result is NaN, the “invalid” floating-point exception is raised, and a domain error occurs. If one or both operands are NaN the result is NaN. Otherwise if only one operand is infinite, the result is NaN, the “invalid” floating-point exception is raised, and a domain error occurs. If both operands are infinite, the result is **DEC\_INFINITY** with the sign of *x*, converted to the return type of the function. The **quantized***N* functions do not raise the “overflow” and “underflow” floating-point exceptions.

**Returns**

- 3 The **quantized***N* functions return a value with the numerical value of *x* (except for any rounding) and the quantum exponent of *y*.

**7.12.16.2 The samequantumd*N* functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 bool samequantumd32(_Decimal32 x, _Decimal32 y);
 bool samequantumd64(_Decimal64 x, _Decimal64 y);
 bool samequantumd128(_Decimal128 x, _Decimal128 y);
 #endif
```

**Description**

- 2 The **samequantumd***N* functions determine if the quantum exponents of *x* and *y* are the same. If both *x* and *y* are NaN, or both infinite, they have the same quantum exponents; if exactly one operand is infinite or exactly one operand is NaN, they do not have the same quantum exponents. The **samequantumd***N* functions raise no floating-point exception.

**Returns**

- 3 The **samequantumd***N* functions return nonzero (**true**) when *x* and *y* have the same quantum exponents, zero (**false**) otherwise.

**7.12.16.3 The quantumd*N* functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 quantumd32(_Decimal32 x);
 _Decimal64 quantumd64(_Decimal64 x);
 _Decimal128 quantumd128(_Decimal128 x);
 #endif
```

**Description**

- 2 The **quantumd***N* functions compute the quantum (5.3.5.3.4) of a finite argument. If *x* is infinite, the result is  $+\infty$ .

**Returns**

- 3 The **quantumd***N* functions return the quantum of *x*.

**7.12.16.4 The llquantexpd*N* functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
```

```

long long int llquantexpd32(_Decimal32 x);
long long int llquantexpd64(_Decimal64 x);
long long int llquantexpd128(_Decimal128 x);
#endif

```

### Description

- The **llquantexpd***N* functions compute the quantum exponent (5.3.5.3.4) of a finite argument. If *x* is infinite or NaN, they compute **LLONG\_MIN**, the “invalid” floating-point exception is raised, and a domain error occurs.

### Returns

- The **llquantexpd***N* functions return the quantum exponent of *x*.

## 7.12.17 Decimal re-encoding functions

### 7.12.17.1 General

- ISO/IEC 60559 specifies two different schemes to encode significands in the object representation of a decimal floating-point object: one based on decimal encoding (which packs three decimal digits into 10 bits), the other based on binary encoding (as a binary integer). An implementation may use either of these encoding schemes for its decimal floating types. The re-encoding functions in this subclause provide conversions between external decimal data with a given encoding scheme and the implementation’s corresponding decimal floating type.

### 7.12.17.2 The **encodedecd***N* functions

#### Synopsis

```

1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 void encodedecd32(unsigned char encptr[restrict static 4],
 const _Decimal32 * restrict xptr);
 void encodedecd64(unsigned char encptr[restrict static 8],
 const _Decimal64 * restrict xptr);
 void encodedecd128(unsigned char encptr[restrict static 16],
 const _Decimal128 * restrict xptr);
 #endif

```

### Description

- The **encodedecd***N* functions convert *\*xptr* into an ISO/IEC 60559 decimal*N* encoding in the encoding scheme based on decimal encoding of the significand and store the resulting encoding as an *N*/8 element array, with 8 bits per array element, in the object pointed to by *encptr*. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions preserve the value of *\*xptr* and raise no floating-point exceptions. If *\*xptr* is non-canonical, these functions can possibly produce a canonical encoding.

### Returns

- The **encodedecd***N* functions return no value.

### 7.12.17.3 The **decodedecd***N* functions

#### Synopsis

```

1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 void decodedecd32(_Decimal32 * restrict xptr,
 const unsigned char encptr[restrict static 4]);
 void decodedecd64(_Decimal64 * restrict xptr,
 const unsigned char encptr[restrict static 8]);
 void decodedecd128(_Decimal128 * restrict xptr,
 const unsigned char encptr[restrict static 16]);
 #endif

```

**Description**

- 2 The **decodedecd** $N$  functions interpret the  $N/8$  element array pointed to by **encptr** as an ISO/IEC 60559 decimal $N$  encoding, with 8 bits per array element, in the encoding scheme based on decimal encoding of the significand. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions convert the given encoding into a value of the decimal floating type, and store the result in the object pointed to by **xptr**. These functions preserve the encoded value and raise no floating-point exceptions. If the encoding is non-canonical, these functions can possibly produce a canonical representation.

**Returns**

- 3 The **decodedecd** $N$  functions return no value.

**7.12.17.4 The **encodebind** $N$  functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 void encodebind32(unsigned char encptr[restrict static 4],
 const _Decimal32 * restrict xptr);
 void encodebind64(unsigned char encptr[restrict static 8],
 const _Decimal64 * restrict xptr);
 void encodebind128(unsigned char encptr[restrict static 16],
 const _Decimal128 * restrict xptr);
 #endif
```

**Description**

- 2 The **encodebind** $N$  functions convert **\*xptr** into an ISO/IEC 60559 decimal $N$  encoding in the encoding scheme based on binary encoding of the significand and store the resulting encoding as an  $N/8$  element array, with 8 bits per array element, in the object pointed to by **encptr**. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions preserve the value of **\*xptr** and raise no floating-point exceptions. If **\*xptr** is non-canonical, these functions can possibly produce a canonical encoding.

**Returns**

- 3 The **encodebind** $N$  functions return no value.

**7.12.17.5 The **decodebind** $N$  functions****Synopsis**

```
1 #include <math.h>
 #ifdef __STDC_IEC_60559_DFP__
 void decodebind32(_Decimal32 * restrict xptr,
 const unsigned char encptr[restrict static 4]);
 void decodebind64(_Decimal64 * restrict xptr,
 const unsigned char encptr[restrict static 8]);
 void decodebind128(_Decimal128 * restrict xptr,
 const unsigned char encptr[restrict static 16]);
 #endif
```

**Description**

- 2 The **decodebind** $N$  functions interpret the  $N/8$  element array pointed to by **encptr** as an ISO/IEC 60559 decimal $N$  encoding, with 8 bits per array element, in the encoding scheme based on binary encoding of the significand. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions convert the given encoding into a value of decimal floating type, and store the result in the object pointed to by **xptr**. These functions preserve the encoded value and raise no floating-point exceptions. If the encoding is non-canonical, these functions can produce a canonical representation.

**Returns**

- 3 The **decodebind***N* functions return no value.

**7.12.18 Comparison macros****7.12.18.1 General**

- 1 The relational and equality operators support the usual mathematical relationships between numeric values. For any ordered pair of numeric values exactly one of the relationships — *less*, *greater*, and *equal* — is true. Relational operators may raise the “invalid” floating-point exception when argument values are NaNs. For a NaN and a numeric value, or for two NaNs, just the *unordered* relationship is true.<sup>291)</sup> 7.12.18.2 through 7.12.18.7 provide macros that are quiet versions of the relational operators: the macros do not raise the “invalid” floating-point exception as an effect of quiet NaN arguments. The comparison macros facilitate writing efficient code that accounts for quiet NaNs without suffering the “invalid” floating-point exception. In the synopses in this subclause, *real-floating* indicates that the argument shall be an expression of real floating type<sup>292)</sup> (both arguments are not required to have the same type).<sup>293)</sup> If either argument has decimal floating type, the other argument shall have decimal floating type as well.

**7.12.18.2 The **isgreater** macro****Synopsis**

- ```
1  #include <math.h>
    int isgreater(real-floating x, real-floating y);
```

Description

- 2 The **isgreater** macro determines whether its first argument is greater than its second argument. The value of **isgreater**(*x*, *y*) is always equal to (*x*) > (*y*).

However, unlike (*x*) > (*y*), **isgreater**(*x*, *y*) does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN.

Returns

- 3 The **isgreater** macro returns the value of (*x*) > (*y*).

7.12.18.3 The **isgreaterequal macro****Synopsis**

- ```
1 #include <math.h>
 int isgreaterequal(real-floating x, real-floating y);
```

**Description**

- 2 The **isgreaterequal** macro determines whether its first argument is greater than or equal to its second argument. The value of **isgreaterequal**(*x*, *y*) is always equal to (*x*) >= (*y*).

However, unlike (*x*) >= (*y*), **isgreaterequal**(*x*, *y*) does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN.

**Returns**

- 3 The **isgreaterequal** macro returns the value of (*x*) >= (*y*).

**7.12.18.4 The **isless** macro****Synopsis**

- ```
1  #include <math.h>
```

²⁹¹⁾ISO/IEC 60559 requires that the built-in relational operators raise the “invalid” floating-point exception if the operands compare unordered, as an error indicator for programs written without consideration of NaNs; the result in these cases is false.

²⁹²⁾If any argument is of integer type, or any other type that is not a real floating type, the behavior is undefined.

²⁹³⁾Whether an argument represented in a format wider than its semantic type is converted to the semantic type is unspecified.

```
int isless(real-floating x, real-floating y);
```

Description

- 2 The **isless** macro determines whether its first argument is less than its second argument. The value of **isless**(*x*, *y*) is always equal to (*x*) < (*y*).

However, unlike (*x*) < (*y*), **isless**(*x*, *y*) does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN.

Returns

- 3 The **isless** macro returns the value of (*x*) < (*y*).

7.12.18.5 The **islessequal** macro

Synopsis

```
1 #include <math.h>
   int islessequal(real-floating x, real-floating y);
```

Description

- 2 The **islessequal** macro determines whether its first argument is less than or equal to its second argument. The value of **islessequal**(*x*, *y*) is always equal to (*x*) <= (*y*).

However, unlike (*x*) <= (*y*), **islessequal**(*x*, *y*) does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN.

Returns

- 3 The **islessequal** macro returns the value of (*x*) <= (*y*).

7.12.18.6 The **islessgreater** macro

Synopsis

```
1 #include <math.h>
   int islessgreater(real-floating x, real-floating y);
```

Description

- 2 The **islessgreater** macro determines whether its first argument is less than or greater than its second argument. The **islessgreater**(*x*, *y*) macro is similar to (*x*) < (*y*) || (*x*) > (*y*).

However, **islessgreater**(*x*, *y*) does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN (nor does it evaluate *x* and *y* twice).

Returns

- 3 The **islessgreater** macro returns the value of (*x*) < (*y*) || (*x*) > (*y*).

7.12.18.7 The **isunordered** macro

Synopsis

```
1 #include <math.h>
   int isunordered(real-floating x, real-floating y);
```

Description

- 2 The **isunordered** macro determines whether its arguments are unordered. It does not raise the “invalid” floating-point exception when *x* and *y* are unordered and neither is a signaling NaN.

Returns

- 3 The **isunordered** macro returns 1 if its arguments are unordered and 0 otherwise.

7.12.18.8 The **iseqsig** macro

Synopsis

```
1  #include <math.h>
    int iseqsig(real-floating x, real-floating y);
```

Description

- 2 The **iseqsig** macro determines whether its arguments are equal. If an argument is a NaN, a domain error occurs for the macro, as if a domain error occurred for a function (7.12.2).

Returns

- 3 The **iseqsig** macro returns 1 if its arguments are equal and 0 otherwise.

7.13 Non-local jumps <setjmp.h>

7.13.1 General

- 1 The header <setjmp.h> defines the macros **setjmp** and **__STDC_VERSION_SETJMP_H__**, and declares one function and one type, for bypassing the normal function call and return discipline.²⁹⁴⁾
- 2 The macro

```
__STDC_VERSION_SETJMP_H__
```

is an integer constant expression with a value equivalent to 202311L.

- 3 The type declared is

```
jmp_buf
```

which is an array type suitable for holding the information needed to restore a calling environment. The environment of an invocation of the **setjmp** macro consists of information sufficient for a call to the **longjmp** function to return execution to the correct block and invocation of that block, were it called recursively. It does not include the state of the floating-point environment, of open files, or of any other component of the abstract machine.

- 4 It is unspecified whether **setjmp** is a macro or an identifier declared with external linkage. If a macro definition is suppressed to access an actual function, or a program defines an external identifier with the name **setjmp**, the behavior is undefined.

7.13.2 Save calling environment

7.13.2.1 The **setjmp** macro

Synopsis

- 1

```
#include <setjmp.h>
int setjmp(jmp_buf env);
```

Description

- 2 The **setjmp** macro saves its calling environment in its **jmp_buf** argument for later use by the **longjmp** function.

Returns

- 3 If the return is from a direct invocation, the **setjmp** macro returns the value zero. If the return is from a call to the **longjmp** function, the **setjmp** macro returns a nonzero value.

Environmental limits

- 4 An invocation of the **setjmp** macro shall appear only in one of the following contexts:
 - the entire controlling expression of a selection or iteration statement;
 - one operand of a relational or equality operator with the other operand an integer constant expression, with the resulting expression being the entire controlling expression of a selection or iteration statement;
 - the operand of a unary ! operator with the resulting expression being the entire controlling expression of a selection or iteration statement; or
 - the entire expression of an expression statement (possibly cast to **void**).
- 5 If the invocation appears in any other context, the behavior is undefined.

7.13.3 Restore calling environment

7.13.3.1 The **longjmp** function

²⁹⁴⁾These functions are useful for dealing with unusual conditions encountered in a low-level function of a program.

Synopsis

```
1  #include <setjmp.h>
    [[noreturn]] void longjmp(jmp_buf env, int val);
```

Description

- 2 The **longjmp** function restores the environment saved by the most recent invocation of the **setjmp** macro in the same invocation of the program with the corresponding **jmp_buf** argument. If there has been no such invocation, or if the invocation was from another thread of execution, or if the function containing the invocation of the **setjmp** macro has terminated execution²⁹⁵⁾ in the interim, or if the invocation of the **setjmp** macro was within the scope of an identifier with variably modified type and execution has left that scope in the interim, the behavior is undefined.
- 3 All accessible objects have values, and all other components of the abstract machine²⁹⁶⁾ have state, as of the time the **longjmp** function was called, except that the representation of objects of automatic storage duration that are local to the function containing the invocation of the corresponding **setjmp** macro that do not have volatile-qualified type and have been changed between the **setjmp** invocation and **longjmp** call is indeterminate.

Returns

- 4 After **longjmp** is completed, thread execution continues as if the corresponding invocation of the **setjmp** macro had just returned the value specified by **val**. The **longjmp** function cannot cause the **setjmp** macro to return the value 0; if **val** is 0, the **setjmp** macro returns the value 1.
- 5 EXAMPLE The **longjmp** function that returns control back to the point of the **setjmp** invocation can cause memory associated with a variable length array object to be squandered.

```
#include <setjmp.h>
jmp_buf buf;
void g(int n);
void h(int n);
int n = 6;

void f(void)
{
    int x[n];           // valid: f is not terminated
    setjmp(buf);
    g(n);
}

void g(int n)
{
    int a[n];           // a can remain allocated
    h(n);
}

void h(int n)
{
    int b[n];           // b can remain allocated
    longjmp(buf, 2);    // can cause memory loss
}
```

²⁹⁵⁾For example, by executing a **return** statement or because another **longjmp** call has caused a transfer to a **setjmp** invocation in a function earlier in the set of nested calls.

²⁹⁶⁾This includes, but is not limited to, the floating-point environment and the state of open files.

7.14 Signal handling <signal.h>

7.14.1 General

- 1 The header <signal.h> declares a type and two functions and defines several macros, for handling various *signals* (conditions that may be reported during program execution).
- 2 The type defined is

```
sig_atomic_t
```

which is the (possibly volatile-qualified) integer type of an object that can be accessed as an atomic entity, even in the presence of asynchronous interrupts.

- 3 The macros defined are

```
SIG_DFL
SIG_ERR
SIG_IGN
```

which expand to constant expressions with distinct values that have type compatible with the second argument to, and the return value of, the **signal** function, and whose values compare unequal to the address of any declarable function; and the following, which expand to positive integer constant expressions with type **int** and distinct values that are the signal numbers, each corresponding to the specified condition:

- SIGABRT** abnormal termination, such as is initiated by the **abort** function
- SIGFPE** an erroneous arithmetic operation, such as zero divide or an operation resulting in overflow
- SIGILL** detection of an invalid function image, such as an invalid instruction
- SIGINT** receipt of an interactive attention signal
- SIGSEGV** an invalid access to storage
- SIGTERM** a termination request sent to the program

- 4 An implementation is not required to generate any of these signals, except as a result of explicit calls to the **raise** function. Additional signals and pointers to undeclarable functions, with macro definitions beginning, respectively, with the letters **SIG** and an uppercase letter or with **SIG_** and an uppercase letter,²⁹⁷⁾ may also be specified by the implementation. The complete set of signals, their semantics, and their default handling is implementation-defined; all signal numbers shall be positive.

7.14.2 Specify signal handling

7.14.2.1 The **signal** function

Synopsis

- 1

```
#include <signal.h>
void (*signal(int sig, void (*func)(int)))(int);
```

Description

- 2 The **signal** function chooses one of three ways in which receipt of the signal number **sig** is to be subsequently handled. If the value of **func** is **SIG_DFL**, default handling for that signal will occur. If the value of **func** is **SIG_IGN**, the signal will be ignored. Otherwise, **func** shall point to a function to be called when that signal occurs. An invocation of such a function because of a signal, or

²⁹⁷⁾See “future library directions” (7.34.10). The names of the signal numbers reflect the following terms (respectively): abort, floating-point exception, illegal instruction, interrupt, segmentation violation, and termination.

(recursively) of any further functions called by that invocation (other than functions in the standard library),²⁹⁸⁾ is called a *signal handler*.

- 3 When a signal occurs and `func` points to a function, it is implementation-defined whether the equivalent of `signal(sig, SIG_DFL)`; is executed or the implementation prevents some implementation-defined set of signals (at least including `sig`) from occurring until the current signal handling has completed; in the case of `SIGILL`, the implementation may alternatively define that no action is taken. Then the equivalent of `(*func)(sig)`; is executed. If and when the function returns, if the value of `sig` is `SIGFPE`, `SIGILL`, `SIGSEGV`, or any other implementation-defined value corresponding to a computational exception, the behavior is undefined; otherwise the program will resume execution at the point it was interrupted.
- 4 If the signal occurs as the result of calling the `abort` or `raise` function, the signal handler shall not call the `raise` function.
- 5 If the signal occurs other than as the result of calling the `abort` or `raise` function, the behavior is undefined if the signal handler refers to any object with static or thread storage duration that is not a lock-free atomic object and that is not declared with the `constexpr` storage-class specifier other than by assigning a value to an object declared as `volatile sig_atomic_t`, or the signal handler calls any function in the standard library other than

- the `abort` function,
- the `_Exit` function,
- the `quick_exit` function,
- the functions in `<stdatomic.h>` (except where explicitly stated otherwise) when the atomic arguments are lock-free,
- the `atomic_is_lock_free` function with any atomic argument, or
- the `signal` function with the first argument equal to the signal number corresponding to the signal that caused the invocation of the handler. Furthermore, if such a call to the `signal` function results in a `SIG_ERR` return, the object designated by `errno` has an indeterminate representation.²⁹⁹⁾

- 6 At program startup, the equivalent of

```
signal(sig, SIG_IGN);
```

may be executed for some signals selected in an implementation-defined manner; the equivalent of

```
signal(sig, SIG_DFL);
```

is executed for all other signals defined by the implementation.

- 7 Use of this function in a multi-threaded program results in undefined behavior. The implementation shall behave as if no library function calls the `signal` function.

Returns

- 8 If the request can be honored, the `signal` function returns the value of `func` for the most recent successful call to `signal` for the specified signal `sig`. Otherwise, a value of `SIG_ERR` is returned and a positive value is stored in `errno`.

Forward references: the `abort` function (7.24.5.1), the `exit` function (7.24.5.4), the `_Exit` function (7.24.5.5), the `quick_exit` function (7.24.5.7).

²⁹⁸⁾This includes functions called indirectly via standard library functions (e.g. a `SIGABRT` handler called via the `abort` function).

²⁹⁹⁾If any signal is generated by an asynchronous signal handler, the behavior is undefined.

7.14.3 Send signal

7.14.3.1 The **raise** function

Synopsis

1

```
#include <signal.h>
int raise(int sig);
```

Description

2

The **raise** function carries out the actions described in 7.14.2.1 for the signal **sig**. If a signal handler is called, the **raise** function shall not return until after the signal handler does.

Returns

3

The **raise** function returns zero if successful, nonzero if unsuccessful.

7.15 Alignment `<stdalign.h>`

- 1 The header `<stdalign.h>` provides no content.

7.16 Varying arguments <stdarg.h>

7.16.1 General

- 1 The header <stdarg.h> declares a type and defines five macros, for use with variadic functions and varying arguments.
- 2 The macro

```
__STDC_VERSION_STDARG_H__
```

is an integer constant expression with a value equivalent to 202311L.

- 3 A function may be called with a variable number of arguments of varying types if its parameter type list ends with an ellipsis.
- 4 The type declared is

```
va_list
```

which is a complete object type suitable for holding information needed by the macros **va_start**, **va_arg**, **va_end**, and **va_copy**. If access to the varying arguments is desired, the called function shall declare an object (generally referred to as **ap** in this subclause) having type **va_list**. The object **ap** may be passed as an argument to another function; if that function invokes the **va_arg** macro with parameter **ap**, the representation of **ap** in the calling function is indeterminate and shall be passed to the **va_end** macro prior to any further reference to **ap**.³⁰⁰⁾ Whether a byte copy of **va_list** can be used in place of the original is implementation-defined.

7.16.2 Varying argument list access macros

7.16.2.1 General

- 1 The **va_start** and **va_arg** macros described in this subclause shall be implemented as macros, not functions. It is unspecified whether **va_copy** and **va_end** are macros or identifiers declared with external linkage. If a macro definition is suppressed to access an actual function, or a program defines an external identifier with the same name, the behavior is undefined. Each invocation of the **va_start** and **va_copy** macros shall be matched by a corresponding invocation of the **va_end** macro in the same function.

7.16.2.2 The **va_arg** macro

Synopsis

- 1

```
#include <stdarg.h>
type va_arg(va_list ap, type);
```

Description

- 2 The **va_arg** macro expands to an expression that has the specified type and the value of the next argument in the call. The parameter **ap** shall have been initialized by the **va_start** or **va_copy** macro (without an intervening invocation of the **va_end** macro for the same **ap**). Each invocation of the **va_arg** macro modifies **ap** so that the values of successive arguments are returned in turn. The behavior is undefined if there is no actual next argument. The parameter *type* shall be an object type name. If *type* is not compatible with the type of the actual next argument (as promoted according to the default argument promotions), the behavior is undefined, except for the following cases:

- both types are pointers to qualified or unqualified versions of compatible types;
- one type is compatible with a signed integer type, the other type is compatible with the corresponding unsigned integer type, and the value is representable in both types;

³⁰⁰⁾ A pointer to a **va_list** can be created and passed to another function, in which case the original function can make further use of the original list after the other function returns.

- one type is pointer to qualified or unqualified **void** and the other is a pointer to a qualified or unqualified character type;
- or, the type of the next argument is **nullptr_t** and *type* is a pointer type that has the same representation and alignment requirements as a pointer to a character type.³⁰¹⁾

Returns

- 3 The first invocation of the **va_arg** macro after that of the **va_start** macro returns the value of the first argument without an explicit parameter, which matches the position of the `...` in the parameter list. Successive invocations return the values of the remaining arguments in succession.

7.16.2.3 The **va_copy** macro

Synopsis

```
1  #include <stdarg.h>
    void va_copy(va_list dest, va_list src);
```

Description

- 2 The **va_copy** macro initializes `dest` as a copy of `src`, as if the **va_start** macro had been applied to `dest` followed by the same sequence of uses of the **va_arg** macro as had previously been used to reach the present state of `src`. Neither the **va_copy** nor **va_start** macro shall be invoked to reinitialize `dest` without an intervening invocation of the **va_end** macro for the same `dest`.

Returns

- 3 The **va_copy** macro returns no value.

7.16.2.4 The **va_end** macro

Synopsis

```
1  #include <stdarg.h>
    void va_end(va_list ap);
```

Description

- 2 The **va_end** macro facilitates a normal return from the function whose varying argument list was referred to by the expansion of the **va_start** macro, or the function containing the expansion of the **va_copy** macro, that initialized the `va_list` `ap`. The **va_end** macro may modify `ap` so that it is no longer usable (without being reinitialized by the **va_start** or **va_copy** macro). If there is no corresponding invocation of the **va_start** or **va_copy** macro, or if the **va_end** macro is not invoked before the return, the behavior is undefined.

Returns

- 3 The **va_end** macro returns no value.

7.16.2.5 The **va_start** macro

Synopsis

```
1  #include <stdarg.h>
    void va_start(va_list ap, ...);
```

Description

- 2 The **va_start** macro shall be invoked before any access to the unnamed arguments.
- 3 The **va_start** macro initializes `ap` for subsequent use by the **va_arg** and **va_end** macros. Neither the **va_start** nor **va_copy** macro shall be invoked to reinitialize `ap` without an intervening invocation of the **va_end** macro for the same `ap`.

³⁰¹⁾Such types are in particular pointers to qualified or unqualified versions of **void**.

- 4 Only the first argument passed to **va_start** is evaluated. If any additional arguments expand to include unbalanced parentheses, or a preprocessing token that does not convert to a token, the behavior is undefined.
- 5 NOTE The macro allows additional arguments to be passed for **va_start** for compatibility with older versions of the library only.

Returns

- 6 The **va_start** macro returns no value.

Recommended practice

- 7 Additional arguments beyond the first given to the **va_start** macro may be expanded and used in unspecified contexts where they are unevaluated. For example, an implementation diagnoses potentially erroneous input for an invocation of **va_start** such as:

```
#include <stdarg.h>

void miaou (...) {
    va_list vl;
    va_start(vl, 1, 3.0, "12", xd); // diagnostic encouraged
    /* ... */
    va_end(vl);
}
```

Simultaneously, **va_start** usage consistent with older revisions of this document should not produce a diagnostic:

```
#include <stdarg.h>

void neigh (int last_arg, ...) {
    va_list vl;
    va_start(vl, last_arg); // no diagnostic
    /* ... */
    va_end(vl);
}
```

- 8 EXAMPLE 1 The function **f1** gathers into an array a list of arguments that are pointers to strings (but not more than **MAXARGS** arguments), then passes the array as a single argument to function **f2**. The number of pointers is specified by the first argument to **f1**.

```
#include <stdarg.h>
#define MAXARGS 31

void f1(int n_ptrs, ...)
{
    va_list ap;
    char *array[MAXARGS];
    int ptr_no = 0;

    if (n_ptrs > MAXARGS)
        n_ptrs = MAXARGS;
    va_start(ap);
    while (ptr_no < n_ptrs)
        array[ptr_no++] = va_arg(ap, char *);
    va_end(ap);
    f2(n_ptrs, array);
}
```

Each call to **f1** is required to have visible the definition of the function or a declaration such as

```
void f1(int, ...);
```


- 9 EXAMPLE 2 The function **f3** is similar, but saves the status of the varying argument list after the indicated number of arguments; after **f2** has been called once with the whole list, the trailing part of the list is gathered again and passed to function **f4**.

```
#include <stdarg.h>
#define MAXARGS 31

void f3(int n_ptrs, int f4_after, ...)
{
    va_list ap, ap_save;
    char *array[MAXARGS];
    int ptr_no = 0;
    if (n_ptrs > MAXARGS)
        n_ptrs = MAXARGS;
    va_start(ap);
    while (ptr_no < n_ptrs) {
        array[ptr_no++] = va_arg(ap, char *);
        if (ptr_no == f4_after)
            va_copy(ap_save, ap);
    }
    va_end(ap);
    f2(n_ptrs, array);

    // Now process the saved copy.

    n_ptrs -= f4_after;
    ptr_no = 0;
    while (ptr_no < n_ptrs)
        array[ptr_no++] = va_arg(ap_save, char *);
    va_end(ap_save);
    f4(n_ptrs, array);
}
```

- 10 EXAMPLE 3 The function **f5** is similar to **f1**, but instead of passing an explicit number of strings as the first argument, the argument list is terminated with a null pointer.

```
#include <stdarg.h>
#define MAXARGS 31

void f5(...)
{
    va_list ap;
    char *array[MAXARGS];
    int ptr_no = 0;
    va_start(ap);
    while (ptr_no < MAXARGS)
    {
        char *ptr = va_arg(ap, char *);
        if (!ptr)
            break;
        array[ptr_no++] = ptr;
    }
    va_end(ap);
    f6(ptr_no, array);
}
```

Each call to **f5** is required to have visible the definition of the function or a declaration such as

```
void f5(...);
```

and implicitly requires the last argument to be a null pointer.

7.17 Atomics <stdatomic.h>

7.17.1 Introduction

- 1 The header <stdatomic.h> defines several macros and declares several types and functions for performing atomic operations on data shared between threads.³⁰²⁾
- 2 Implementations that define the macro **__STDC_NO_ATOMICS__** may not provide this header nor support any of its facilities.
- 3 The macro

__STDC_VERSION_STDATOMIC_H__

is an integer constant expression with a value equivalent to 202311L.

- 4 The macros defined are the *atomic lock-free macros*

ATOMIC_BOOL_LOCK_FREE
ATOMIC_CHAR_LOCK_FREE
ATOMIC_CHAR8_T_LOCK_FREE
ATOMIC_CHAR16_T_LOCK_FREE
ATOMIC_CHAR32_T_LOCK_FREE
ATOMIC_WCHAR_T_LOCK_FREE
ATOMIC_SHORT_LOCK_FREE
ATOMIC_INT_LOCK_FREE
ATOMIC_LONG_LOCK_FREE
ATOMIC_LLONG_LOCK_FREE
ATOMIC_POINTER_LOCK_FREE

which expand to constant expressions suitable for use in conditional expression inclusion preprocessing directives and which indicate the lock-free property of the corresponding atomic types (both signed and unsigned); and

ATOMIC_FLAG_INIT

which expands to an initializer for an object of type **atomic_flag**.

- 5 The types include

memory_order

which is an enumerated type whose enumerators identify memory ordering constraints;

atomic_flag

which is a structure type representing a lock-free, primitive atomic flag; and several atomic analogs of integer types.

- 6 In the following synopses:
 - An *A* refers to an atomic type.
 - A *C* refers to its corresponding non-atomic type.
 - An *M* refers to the type of the other argument for arithmetic operations. For atomic integer types, *M* is *C*. For atomic pointer types, *M* is **ptrdiff_t**.
 - The functions not ending in **_explicit** have the same semantics as the corresponding **_explicit** function with **memory_order_seq_cst** for the **memory_order** argument.

³⁰²⁾See “future library directions” (7.34.11).

- 7 It is unspecified whether any generic function declared in `<stdatomic.h>` is a macro or an identifier declared with external linkage. If a macro definition is suppressed to access an actual function, or a program defines an external identifier with the name of a generic function, the behavior is undefined.
- 8 NOTE Many operations are volatile-qualified. The “volatile as device register” semantics have not changed in the standard. This qualification means that volatility is preserved when applying these operations to volatile objects.

7.17.2 Initialization

7.17.2.1 General

- 1 An atomic object with automatic storage duration that is not initialized or such an object with allocated storage duration initially has an indeterminate representation; equally, a non-atomic store to any byte of the representation (either directly or, for example, by calls to `memcpy` or `memset`) makes any atomic object have an indeterminate representation. Explicit or default initialization for atomic objects with static or thread storage duration that do not have the type `atomic_flag` is guaranteed to produce a valid state.³⁰³⁾
- 2 Concurrent access to an atomic object before it is set to a valid state, even via an atomic operation, constitutes a data race. If a signal occurs other than as the result of calling the `abort` or `raise` functions, the behavior is undefined if the signal handler reads or modifies an atomic object that has an indeterminate representation.
- 3 EXAMPLE The following definition ensure valid states for `guide` and `head` regardless if these are found in file scope or block scope. Thus any atomic operation that is performed on them after their initialization has been met is well defined.

```
_Atomic int guide = 42;
static _Atomic(void*) head;
```

7.17.2.2 The `atomic_init` generic function

Synopsis

- 1

```
#include <stdatomic.h>
void atomic_init(volatile A *obj, C value);
```

Description

- 2 The `atomic_init` generic function initializes the atomic object pointed to by `obj` to the value `value`, while also initializing any additional state that the implementation may need to carry for the atomic object. If the object is a byte array, after the call the effective type is the atomic type `A`.
- 3 Although this function initializes an atomic object, it does not avoid data races; concurrent access to the object being initialized, even via an atomic operation, constitutes a data race.
- 4 If a signal occurs other than as the result of calling the `abort` or `raise` functions, the behavior is undefined if the signal handler calls the `atomic_init` generic function.

Returns

- 5 The `atomic_init` generic function returns no value.
- 6 EXAMPLE

```
atomic_int guide;
atomic_init(&guide, 42);
```

³⁰³⁾See “future library directions” (7.34.11).

7.17.3 Order and consistency

7.17.3.1 General

- 1 The enumerated type **memory_order** specifies the detailed regular (non-atomic) memory synchronization operations as defined in 5.2.2.5 and may provide for operation ordering. Its enumeration constants are as follows:³⁰⁴⁾

```
memory_order_relaxed
memory_order_consume
memory_order_acquire
memory_order_release
memory_order_acq_rel
memory_order_seq_cst
```

- 2 For **memory_order_relaxed**, no operation orders memory.
- 3 For **memory_order_release**, **memory_order_acq_rel**, and **memory_order_seq_cst**, a store operation performs a release operation on the affected memory location.
- 4 For **memory_order_acquire**, **memory_order_acq_rel**, and **memory_order_seq_cst**, a load operation performs an acquire operation on the affected memory location.
- 5 For **memory_order_consume**, a load operation performs a consume operation on the affected memory location.
- 6 There shall be a single total order *S* on all **memory_order_seq_cst** operations, consistent with the “happens before” order and modification orders for all affected locations, such that each **memory_order_seq_cst** operation *B* that loads a value from an atomic object *M* observes one of the following values:
- the result of the last modification *A* of *M* that precedes *B* in *S*, if it exists, or
 - if *A* exists, the result of some modification of *M* that is not **memory_order_seq_cst** and that does not happen before *A*, or
 - if *A* does not exist, the result of some modification of *M* that is not **memory_order_seq_cst**.
- 7 NOTE 1 Although it is not explicitly required that *S* include lock operations, it can always be extended to an order that does include lock and unlock operations, since the ordering between those is already included in the “happens before” ordering.
- 8 NOTE 2 Atomic operations specifying **memory_order_relaxed** are relaxed only with respect to memory ordering. Implementations still guarantee that any given atomic access to a particular atomic object is indivisible with respect to all other atomic accesses to that object.
- 9 For an atomic operation *B* that reads the value of an atomic object *M*, if there is a **memory_order_seq_cst** fence *X* sequenced before *B*, then *B* observes either the last **memory_order_seq_cst** modification of *M* preceding *X* in the total order *S* or a later modification of *M* in its modification order.
- 10 For atomic operations *A* and *B* on an atomic object *M*, where *A* modifies *M* and *B* takes its value, if there is a **memory_order_seq_cst** fence *X* such that *A* is sequenced before *X* and *B* follows *X* in *S*, then *B* observes either the effects of *A* or a later modification of *M* in its modification order.
- 11 For atomic modifications *A* and *B* of an atomic object *M*, *B* occurs later than *A* in the modification order of *M* if:
- there is a **memory_order_seq_cst** fence *X* such that *A* is sequenced before *X*, and *X* precedes *B* in *S*, or
 - there is a **memory_order_seq_cst** fence *Y* such that *Y* is sequenced before *B*, and *A* precedes *Y* in *S*, or

³⁰⁴⁾See “future library directions” (7.34.11).

- there are **memory_order_seq_cst** fences X and Y such that A is sequenced before X , Y is sequenced before B , and X precedes Y in S .
- 12 Atomic read-modify-write operations shall always read the last value (in the modification order) stored before the write associated with the read-modify-write operation.
- 13 An atomic store shall only store a value that has been computed from constants and program input values by a finite sequence of program evaluations, such that each evaluation observes the stored values of the objects as computed by the last prior assignment in the sequence. The ordering of evaluations in this sequence shall be such that
- If an evaluation B observes a value computed by A in a different thread, then B does not happen before A .
 - If an evaluation A is included in the sequence, then all evaluations that assign to the same object and happen before A are also included.
- 14 NOTE 3 The second requirement disallows “out-of-thin-air”, or “speculative” stores of atomics when relaxed atomics are used. Since unordered operations are involved, evaluations can appear in this sequence out of thread order. For example, with x and y initially zero,

```
// Thread 1:
r1 = atomic_load_explicit(&y, memory_order_relaxed);
atomic_store_explicit(&x, r1, memory_order_relaxed);

// Thread 2:
r2 = atomic_load_explicit(&x, memory_order_relaxed);
atomic_store_explicit(&y, 42, memory_order_relaxed);
```

is allowed to produce $r1 == 42 \ \&\& \ r2 == 42$. The sequence of evaluations justifying this consists of:

```
atomic_store_explicit(&y, 42, memory_order_relaxed);
r1 = atomic_load_explicit(&y, memory_order_relaxed);
atomic_store_explicit(&x, r1, memory_order_relaxed);
r2 = atomic_load_explicit(&x, memory_order_relaxed);
```

On the other hand,

```
// Thread 1:
r1 = atomic_load_explicit(&y, memory_order_relaxed);
atomic_store_explicit(&x, r1, memory_order_relaxed);

// Thread 2:
r2 = atomic_load_explicit(&x, memory_order_relaxed);
atomic_store_explicit(&y, r2, memory_order_relaxed);
```

is not allowed to produce $r1 == 42 \ \&\& \ r2 == 42$, since there is no sequence of evaluations that results in the computation of 42. In the absence of “relaxed” operations and read-modify-write operations with weaker than **memory_order_acq_rel** ordering, the second requirement has no impact.

Recommended practice

- 15 The requirements do not forbid $r1 == 42 \ \&\& \ r2 == 42$ in the following example, with x and y initially zero:

```
// Thread 1:
r1 = atomic_load_explicit(&x, memory_order_relaxed);
if (r1 == 42)
    atomic_store_explicit(&y, r1, memory_order_relaxed);

// Thread 2:
```

```

r2 = atomic_load_explicit(&y, memory_order_relaxed);
if (r2 == 42)
    atomic_store_explicit(&x, 42, memory_order_relaxed);

```

However, this is not useful behavior, and implementations should not allow it.

- 16 Implementations should make atomic stores visible to atomic loads within a reasonable amount of time.

7.17.3.2 The **kill_dependency** macro

Synopsis

```

1  #include <stdatomic.h>
    type kill_dependency(type y);

```

Description

- 2 The **kill_dependency** macro terminates a dependency chain; the argument does not carry a dependency to the return value.

Returns

- 3 The **kill_dependency** macro returns the value of *y*.

7.17.4 Fences

7.17.4.1 General

- 1 This subclause introduces synchronization primitives called *fences*. Fences can have acquire semantics, release semantics, or both. A fence with acquire semantics is called an *acquire fence*; a fence with release semantics is called a *release fence*.
- 2 A release fence *A* synchronizes with an acquire fence *B* if there exist atomic operations *X* and *Y*, both operating on some atomic object *M*, such that *A* is sequenced before *X*, *X* modifies *M*, *Y* is sequenced before *B*, and *Y* reads the value written by *X* or a value written by any side effect in the hypothetical release sequence *X* would head if it were a release operation.
- 3 A release fence *A* synchronizes with an atomic operation *B* that performs an acquire operation on an atomic object *M* if there exists an atomic operation *X* such that *A* is sequenced before *X*, *X* modifies *M*, and *B* reads the value written by *X* or a value written by any side effect in the hypothetical release sequence *X* would head if it were a release operation.
- 4 An atomic operation *A* that is a release operation on an atomic object *M* synchronizes with an acquire fence *B* if there exists some atomic operation *X* on *M* such that *X* is sequenced before *B* and reads the value written by *A* or a value written by any side effect in the release sequence headed by *A*.

7.17.4.2 The **atomic_thread_fence** function

Synopsis

```

1  #include <stdatomic.h>
    void atomic_thread_fence(memory_order order);

```

Description

- 2 Depending on the value of *order*, this operation:
- has no effects, if *order* == **memory_order_relaxed**;
 - is an acquire fence, if *order* == **memory_order_acquire** or *order* == **memory_order_consume**;
 - is a release fence, if *order* == **memory_order_release**;
 - is both an acquire fence and a release fence, if *order* == **memory_order_acq_rel**;

— is a sequentially consistent acquire and release fence, if `order == memory_order_seq_cst`.

Returns

- 3 The `atomic_thread_fence` function returns no value.

7.17.4.3 The `atomic_signal_fence` function

Synopsis

```
1  #include <stdatomic.h>
    void atomic_signal_fence(memory_order order);
```

Description

- 2 Equivalent to `atomic_thread_fence(order)`, except that the resulting ordering constraints are established only between a thread and a signal handler executed in the same thread.
- 3 NOTE 1 The `atomic_signal_fence` function can be used to specify the order in which actions performed by the thread become visible to the signal handler.
- 4 NOTE 2 Compiler optimizations and reorderings of loads and stores are inhibited in the same way as with `atomic_thread_fence`, but the hardware fence instructions that `atomic_thread_fence` would have inserted are not emitted.

Returns

- 5 The `atomic_signal_fence` function returns no value.

7.17.5 Lock-free property

7.17.5.1 General

- 1 The atomic lock-free macros indicate the lock-free property of integer and address atomic types. A value of 0 indicates that the type is never lock-free; a value of 1 indicates that the type is sometimes lock-free; a value of 2 indicates that the type is always lock-free.

Recommended practice

- 2 Operations that are lock-free should also be *address-free*. That is, atomic operations on the same memory location via two different addresses will communicate atomically. The implementation should not depend on any per-process state. This restriction enables communication via memory mapped into a process more than once and memory shared between two processes.

7.17.5.2 The `atomic_is_lock_free` generic function

Synopsis

```
1  #include <stdatomic.h>
    bool atomic_is_lock_free(const volatile A *obj);
```

Description

- 2 The `atomic_is_lock_free` generic function indicates whether atomic operations on objects of the type pointed to by `obj` are lock-free.

Returns

- 3 The `atomic_is_lock_free` generic function returns nonzero (true) if and only if atomic operations on objects of the type pointed to by the argument are lock-free. In any given program execution, the result of the lock-free query shall be consistent for all pointers of the same type.³⁰⁵⁾

³⁰⁵⁾ `obj` can be a null pointer.

7.17.6 Atomic integer types

- 1 For each line in Table 7.6,³⁰⁶⁾ the atomic type name is declared as a type that has the same representation and alignment requirements as the corresponding direct type.³⁰⁷⁾

Table 7.6 — Type name equivalency

Atomic type name	Direct type
<code>atomic_bool</code>	<code>_Atomic bool</code>
<code>atomic_char</code>	<code>_Atomic char</code>
<code>atomic_schar</code>	<code>_Atomic signed char</code>
<code>atomic_uchar</code>	<code>_Atomic unsigned char</code>
<code>atomic_short</code>	<code>_Atomic short</code>
<code>atomic_ushort</code>	<code>_Atomic unsigned short</code>
<code>atomic_int</code>	<code>_Atomic int</code>
<code>atomic_uint</code>	<code>_Atomic unsigned int</code>
<code>atomic_long</code>	<code>_Atomic long</code>
<code>atomic_ulong</code>	<code>_Atomic unsigned long</code>
<code>atomic_llong</code>	<code>_Atomic long long</code>
<code>atomic_ullong</code>	<code>_Atomic unsigned long long</code>
<code>atomic_char8_t</code>	<code>_Atomic char8_t</code>
<code>atomic_char16_t</code>	<code>_Atomic char16_t</code>
<code>atomic_char32_t</code>	<code>_Atomic char32_t</code>
<code>atomic_wchar_t</code>	<code>_Atomic wchar_t</code>
<code>atomic_int_least8_t</code>	<code>_Atomic int_least8_t</code>
<code>atomic_uint_least8_t</code>	<code>_Atomic uint_least8_t</code>
<code>atomic_int_least16_t</code>	<code>_Atomic int_least16_t</code>
<code>atomic_uint_least16_t</code>	<code>_Atomic uint_least16_t</code>
<code>atomic_int_least32_t</code>	<code>_Atomic int_least32_t</code>
<code>atomic_uint_least32_t</code>	<code>_Atomic uint_least32_t</code>
<code>atomic_int_least64_t</code>	<code>_Atomic int_least64_t</code>
<code>atomic_uint_least64_t</code>	<code>_Atomic uint_least64_t</code>
<code>atomic_int_fast8_t</code>	<code>_Atomic int_fast8_t</code>
<code>atomic_uint_fast8_t</code>	<code>_Atomic uint_fast8_t</code>
<code>atomic_int_fast16_t</code>	<code>_Atomic int_fast16_t</code>
<code>atomic_uint_fast16_t</code>	<code>_Atomic uint_fast16_t</code>
<code>atomic_int_fast32_t</code>	<code>_Atomic int_fast32_t</code>
<code>atomic_uint_fast32_t</code>	<code>_Atomic uint_fast32_t</code>
<code>atomic_int_fast64_t</code>	<code>_Atomic int_fast64_t</code>
<code>atomic_uint_fast64_t</code>	<code>_Atomic uint_fast64_t</code>
<code>atomic_intptr_t</code>	<code>_Atomic intptr_t</code>
<code>atomic_uintptr_t</code>	<code>_Atomic uintptr_t</code>
<code>atomic_size_t</code>	<code>_Atomic size_t</code>
<code>atomic_ptrdiff_t</code>	<code>_Atomic ptrdiff_t</code>
<code>atomic_intmax_t</code>	<code>_Atomic intmax_t</code>
<code>atomic_uintmax_t</code>	<code>_Atomic uintmax_t</code>

- 2 Conversions to `atomic_bool` behave the same as conversions to `bool`.

Recommended practice

- 3 The representation of an atomic integer type is not required to have the same size as the corresponding regular type but it should have the same size whenever possible, as it eases effort required to port existing code.

³⁰⁶⁾See “future library directions” (7.34.11).

³⁰⁷⁾The same representation and alignment requirements are meant to imply interchangeability as arguments to functions, return values from functions, and members of unions.

7.17.7 Operations on atomic types

7.17.7.1 General

- 1 There are only a few kinds of operations on atomic types, though there are many instances of those kinds. This subclause specifies each general kind.

7.17.7.2 The `atomic_store` generic functions

Synopsis

```
1  #include <stdatomic.h>
    void atomic_store(volatile A *object, C desired);
    void atomic_store_explicit(volatile A *object, C desired, memory_order order);
```

Description

- 2 The `order` argument shall not be `memory_order_acquire`, `memory_order_consume`, nor `memory_order_acq_rel`. Atomically replace the value pointed to by `object` with the value of `desired`. Memory is affected according to the value of `order`.

Returns

- 3 The `atomic_store` generic functions return no value.

7.17.7.3 The `atomic_load` generic functions

Synopsis

```
1  #include <stdatomic.h>
    C atomic_load(const volatile A *object);
    C atomic_load_explicit(const volatile A *object, memory_order order);
```

Description

- 2 The `order` argument shall not be `memory_order_release` nor `memory_order_acq_rel`. Memory is affected according to the value of `order`.

Returns

- 3 Atomically returns the value pointed to by `object`.

7.17.7.4 The `atomic_exchange` generic functions

Synopsis

```
1  #include <stdatomic.h>
    C atomic_exchange(volatile A *object, C desired);
    C atomic_exchange_explicit(volatile A *object, C desired, memory_order order);
```

Description

- 2 Atomically replace the value pointed to by `object` with `desired`. Memory is affected according to the value of `order`. These operations are read-modify-write operations (5.2.2.5).

Returns

- 3 Atomically returns the value pointed to by `object` immediately before the effects.

7.17.7.5 The `atomic_compare_exchange` generic functions

Synopsis

```
1  #include <stdatomic.h>
    bool atomic_compare_exchange_strong(volatile A *object, C *expected, C desired);
    bool atomic_compare_exchange_strong_explicit(volatile A *object, C *expected,
        C desired, memory_order success, memory_order failure);
    bool atomic_compare_exchange_weak(volatile A *object, C *expected, C desired);
    bool atomic_compare_exchange_weak_explicit(volatile A *object, C *expected,
        C desired, memory_order success, memory_order failure);
```

Description

- 2 The **failure** argument shall not be **memory_order_release** nor **memory_order_acq_rel**. The **failure** argument shall be no stronger than the success argument.
- 3 Atomically, compares the contents of the memory pointed to by **object** for equality with that pointed to by **expected**, and if true, replaces the contents of the memory pointed to by **object** with **desired**, and if false, updates the contents of the memory pointed to by **expected** with that pointed to by **object**. Further, if the comparison is true, memory is affected according to the value of **success**, and if the comparison is false, memory is affected according to the value of **failure**. These operations are atomic read-modify-write operations (5.2.2.5).
- 4 NOTE 1 For example, the effect of **atomic_compare_exchange_strong** is

```
if (memcmp(object, expected, sizeof(*object)) == 0)
    memcpy(object, &desired, sizeof(*object));
else
    memcpy(expected, object, sizeof(*object));
```

- 5 A weak compare-and-exchange operation can fail spuriously. That is, even when the contents of memory referred to by **expected** and **object** are equal, it can return zero and store back to **expected** the same memory contents that were originally there.
- 6 NOTE 2 This spurious failure enables implementation of compare-and-exchange on a broader class of machines; e.g. load-locked store-conditional machines.
- 7 EXAMPLE A consequence of spurious failure is that nearly all uses of weak compare-and-exchange will be in a loop.

```
exp = atomic_load(&cur);
do {
    des = function(exp);
} while (!atomic_compare_exchange_weak(&cur, &exp, des));
```

When a compare-and-exchange is in a loop, the weak version will yield better performance on some platforms. When a weak compare-and-exchange would require a loop and a strong one would not, the strong one is preferable.

Returns

- 8 The result of the comparison.

7.17.7.6 The **atomic_fetch and modify generic functions**

- 1 The following operations perform arithmetic and bitwise computations. All these operations are applicable to an object of any atomic integer type other than **_Atomic bool**, **atomic_bool**, or the atomic version of an enumeration with underlying type **bool**. The key, operator, and computation correspondence is:

key	op	computation
add	+	addition
sub	-	subtraction
or		bitwise inclusive or
xor	^	bitwise exclusive or
and	&	bitwise and

Synopsis

- 2

```
#include <stdatomic.h>
C atomic_fetch_key(volatile A *object, M operand);
C atomic_fetch_key_explicit(volatile A *object, M operand, memory_order order);
```

Description

- 3 Atomically replaces the value pointed to by `object` with the result of the computation applied to the value pointed to by `object` and the given operand. Memory is affected according to the value of `order`. These operations are atomic read-modify-write operations (5.2.2.5). For signed integer types, arithmetic performs silent wraparound on integer overflow; there are no undefined results. For address types, the result may be an undefined address, but the operations otherwise have no undefined behavior.

Returns

- 4 Atomically, the value pointed to by `object` immediately before the effects.

7.17.8 Atomic flag type and operations**7.17.8.1 General**

- 1 The `atomic_flag` type provides the classic test-and-set functionality. It has two states, set and clear.
- 2 Operations on an object of type `atomic_flag` shall be lock free.
- 3 NOTE Hence, as per 7.17.5, the operations should also be address-free. No other type requires lock-free operations, so the `atomic_flag` type is the minimum hardware-implemented type needed to conform to this document. The remaining types can be emulated with `atomic_flag`, though with less than ideal properties.
- 4 The macro `ATOMIC_FLAG_INIT` can be used to initialize an `atomic_flag` to the clear state. An `atomic_flag` that is not explicitly initialized with `ATOMIC_FLAG_INIT` has initially an indeterminate representation.
- 5 EXAMPLE

```
atomic_flag guard = ATOMIC_FLAG_INIT;
```

7.17.8.2 The `atomic_flag_test_and_set` functions**Synopsis**

- ```
1 #include <stdatomic.h>
 bool atomic_flag_test_and_set(volatile atomic_flag *object);
 bool atomic_flag_test_and_set_explicit(volatile atomic_flag *object,
 memory_order order);
```

**Description**

- 2 Atomically places the atomic flag pointed to by `object` in the set state and returns the value corresponding to the immediately preceding state. Memory is affected according to the value of `order`. These operations are atomic read-modify-write operations (5.2.2.5).

**Returns**

- 3 The `atomic_flag_test_and_set` functions return the value that corresponds to the state of the atomic flag immediately before the effects. The return value true corresponds to the set state and the return value false corresponds to the clear state.

**7.17.8.3 The `atomic_flag_clear` functions****Synopsis**

- ```
1 #include <stdatomic.h>
    void atomic_flag_clear(volatile atomic_flag *object);
    void atomic_flag_clear_explicit(volatile atomic_flag *object,
        memory_order order);
```

Description

- 2 The `order` argument shall not be `memory_order_acquire` nor `memory_order_acq_rel`. Atomically places the atomic flag pointed to by `object` into the clear state. Memory is affected according to the value of `order`.

Returns

- 3 The **atomic_flag_clear** functions return no value.

7.18 Bit and byte utilities <stdbit.h>

7.18.1 General

- 1 The header <stdbit.h> defines the following macros, types, and functions, to work with the byte and bit representation of many types, typically integer types. This header makes available the `size_t` type name (7.21) and any `uintN_t`, `intN_t`, `uint_leastN_t`, or `int_leastN_t` type names defined by the implementation (7.22).
- 2 The macro

```
__STDC_VERSION_STDBIT_H__
```

is an integer constant expression with a value equivalent to 202ymmL.

- 3 The *most significant index* is the 0-based index counting from the most significant bit, 0, to the least significant bit, $w - 1$, where w is the width of the type that is having its most significant index computed.
- 4 The *least significant index* is the 0-based index counting from the least significant bit, 0, to the most significant bit, $w - 1$, where w is the width of the type that is having its least significant index computed.
- 5 It is unspecified whether any generic function declared in <stdbit.h> is a macro or an identifier declared with external linkage. If a macro definition is suppressed to access an actual function, or a program defines an external identifier with the name of a generic function, the behavior is unspecified.

7.18.2 Endian

- 1 Two common methods of byte ordering in multi-byte scalar types are *little-endian* and *big-endian*. Little-endian is a format for storage or transmission of binary data in which the least significant byte is placed first, with the rest in ascending order. Or, that the least significant byte is stored at the smallest memory address. Big-endian is a format for storage or transmission of binary data in which the most significant byte is placed first, with the rest in descending order. Or, that the most significant byte is stored at the smallest memory address. Other byte orderings are also possible.
- 2 The macros are:

```
__STDC_ENDIAN_LITTLE__
```

which represents a method of byte order storage in which the least significant byte is placed first and the rest are in ascending order, and is an integer constant expression;

```
__STDC_ENDIAN_BIG__
```

which represents a method of byte order storage in which the most significant byte is placed first and the rest are in descending order, and is an integer constant expression;

```
__STDC_ENDIAN_NATIVE__ /* see following description */
```

which represents the method of byte order storage for the execution environment and is an integer constant expression. `__STDC_ENDIAN_NATIVE__` describes the endianness of the execution environment with respect to bit-precise integer types, standard integer types, and extended integer types which do not have padding bits.

- 3 `__STDC_ENDIAN_NATIVE__` shall expand to an integer constant expression whose value is equivalent to the value of `__STDC_ENDIAN_LITTLE__` if the execution environment is little-endian. Otherwise, `__STDC_ENDIAN_NATIVE__` shall expand to an integer constant expression whose value is equivalent to the value of `__STDC_ENDIAN_BIG__` if the execution environment is big-endian. If the execution environment is neither little-endian nor big-endian, it then has some other implementation-defined byte order and the macro `__STDC_ENDIAN_NATIVE__` shall expand to an integer constant expression whose value is different from the values of

`__STDC_ENDIAN_LITTLE__` and `__STDC_ENDIAN_BIG__`. The values of the integer constant expressions for `__STDC_ENDIAN_LITTLE__` and `__STDC_ENDIAN_BIG__` are not equal.

7.18.3 Count Leading Zeros

Synopsis

```
1  #include <stdbit.h>
    unsigned int stdc_leading_zeros_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_leading_zeros_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_leading_zeros_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_leading_zeros_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_leading_zeros_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type stdc_leading_zeros(generic_value_type value) [[unsequenced]];
```

Returns

- 2 Returns the number of consecutive 0 bits in *value*, starting from the most significant bit.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.4 Count Leading Ones

Synopsis

```
1  #include <stdbit.h>
    unsigned int stdc_leading_ones_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_leading_ones_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_leading_ones_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_leading_ones_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_leading_ones_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type stdc_leading_ones(generic_value_type value) [[unsequenced]];
```

Returns

- 2 Returns the number of consecutive 1 bits in *value*, starting from the most significant bit.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.5 Count Trailing Zeros

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_trailing_zeros_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_trailing_zeros_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_trailing_zeros_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_trailing_zeros_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_trailing_zeros_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type
    stdc_trailing_zeros(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the number of consecutive 0 bits in *value*, starting from the least significant bit.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.6 Count Trailing Ones

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_trailing_ones_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_trailing_ones_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_trailing_ones_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_trailing_ones_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_trailing_ones_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type stdc_trailing_ones(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the number of consecutive 1 bits in *value*, starting from the least significant bit.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.7 First Leading Zero

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_first_leading_zero_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_first_leading_zero_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_first_leading_zero_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_first_leading_zero_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_first_leading_zero_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type
    stdc_first_leading_zero(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the most significant index of the first 0 bit in *value*, plus 1. If it is not found, this function returns 0.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.8 First Leading One

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_first_leading_one_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_first_leading_one_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_first_leading_one_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_first_leading_one_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_first_leading_one_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type
    stdc_first_leading_one(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the most significant index of the first 1 bit in *value*, plus 1. If it is not found, this function returns 0.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.9 First Trailing Zero

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_first_trailing_zero_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_first_trailing_zero_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_first_trailing_zero_ui(unsigned int value) [[unsequenced]];
    unsigned int
    stdc_first_trailing_zero_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_first_trailing_zero_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type
    stdc_first_trailing_zero(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the least significant index of the first 0 bit in *value*, plus 1. If it is not found, this function returns 0.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.10 First Trailing One

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_first_trailing_one_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_first_trailing_one_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_first_trailing_one_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_first_trailing_one_ul(unsigned long int value) [[unsequenced]];
    unsigned int
    stdc_first_trailing_one_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type
    stdc_first_trailing_one(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the least significant index of the first 1 bit in *value*, plus 1. If it is not found, this function returns 0.

The type-generic function (marked by its *generic_value_type* argument) returns the appropriate value based on the type of the input value, so long as it is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.11 Count Zeros

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_count_zeros_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_count_zeros_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_count_zeros_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_count_zeros_ul(unsigned long int value) [[unsequenced]];
    unsigned int
        stdc_count_zeros_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type stdc_count_zeros(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the total number of 0 bits within the given *value*.
- 3 The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:
 - standard unsigned integer type, excluding **bool**;
 - extended unsigned integer type;
 - or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.12 Count Ones

Synopsis

```

1  #include <stdbit.h>
    unsigned int stdc_count_ones_uc(unsigned char value) [[unsequenced]];
    unsigned int stdc_count_ones_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_count_ones_ui(unsigned int value) [[unsequenced]];
    unsigned int stdc_count_ones_ul(unsigned long int value) [[unsequenced]];
    unsigned int
        stdc_count_ones_ull(unsigned long long int value) [[unsequenced]];
    generic_return_type stdc_count_ones(generic_value_type value) [[unsequenced]];

```

Returns

- 2 Returns the total number of 1 bits within the given *value*.
- The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:
 - standard unsigned integer type, excluding **bool**;
 - extended unsigned integer type;
 - or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.13 Single-bit Check

Synopsis

1

```
#include <stdbit.h>
bool stdc_has_single_bit_uc(unsigned char value) [[unsequenced]];
bool stdc_has_single_bit_us(unsigned short value) [[unsequenced]];
bool stdc_has_single_bit_ui(unsigned int value) [[unsequenced]];
bool stdc_has_single_bit_ul(unsigned long int value) [[unsequenced]];
bool stdc_has_single_bit_ull(unsigned long long int value) [[unsequenced]];
bool stdc_has_single_bit(generic_value_type value) [[unsequenced]];
```

Returns

- 2 The **stdc_has_single_bit** functions return **true** if and only if there is a single 1 bit in **value**.

The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

7.18.14 Bit Width

Synopsis

1

```
#include <stdbit.h>
unsigned int stdc_bit_width_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_bit_width_us(unsigned short value) [[unsequenced]];
unsigned int stdc_bit_width_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_bit_width_ul(unsigned long int value) [[unsequenced]];
unsigned int
    stdc_bit_width_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_bit_width(generic_value_type value) [[unsequenced]];
```

Description

- 2 The **stdc_bit_width** functions compute the smallest number of bits needed to store **value**.

Returns

- 3 The **stdc_bit_width** functions return 0 if **value** is 0. Otherwise, they return $1 + \lfloor \log_2(\text{value}) \rfloor$.

The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

The *generic_return_type* type shall be a suitably large unsigned integer type capable of representing the computed result.

7.18.15 Bit Floor

Synopsis

```

1  #include <stdbit.h>
    unsigned char stdc_bit_floor_uc(unsigned char value) [[unsequenced]];
    unsigned short stdc_bit_floor_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_bit_floor_ui(unsigned int value) [[unsequenced]];
    unsigned long int stdc_bit_floor_ul(unsigned long int value) [[unsequenced]];
    unsigned long long int
        stdc_bit_floor_ull(unsigned long long int value) [[unsequenced]];
    generic_value_type stdc_bit_floor(generic_value_type value) [[unsequenced]];

```

Description

- 2 The **stdc_bit_floor** functions compute the largest integral power of 2 that is not greater than *value*.

Returns

- 3 The **stdc_bit_floor** functions return 0 if *value* is 0. Otherwise, they return the largest integral power of 2 that is not greater than *value*.

The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

7.18.16 Bit Ceiling

Synopsis

```

1  #include <stdbit.h>
    unsigned char stdc_bit_ceil_uc(unsigned char value) [[unsequenced]];
    unsigned short stdc_bit_ceil_us(unsigned short value) [[unsequenced]];
    unsigned int stdc_bit_ceil_ui(unsigned int value) [[unsequenced]];
    unsigned long int stdc_bit_ceil_ul(unsigned long int value) [[unsequenced]];
    unsigned long long int
        stdc_bit_ceil_ull(unsigned long long int value) [[unsequenced]];
    generic_value_type stdc_bit_ceil(generic_value_type value) [[unsequenced]];

```

Description

- 2 The **stdc_bit_ceil** functions compute the smallest integral power of 2 that is not less than *value*. If the computation does not fit in the given return type, they return 0.

Returns

- 3 The **stdc_bit_ceil** functions return the smallest integral power of 2 that is not less than *value* or 0 if such a value is not representable in the return type.

The type-generic function (marked by its *generic_value_type* argument) returns the previously described result for a given input value so long as the *generic_value_type* is a:

- standard unsigned integer type, excluding **bool**;
- extended unsigned integer type;
- or, bit-precise unsigned integer type whose width matches a standard or extended integer type, excluding **bool**.

7.18.17 Rotate Left

Synopsis

```

1  #include <stdbit.h>
    unsigned char stdc_rotate_left_uc(unsigned char value, unsigned int count);
    unsigned short stdc_rotate_left_us(unsigned short value, unsigned int count);
    unsigned int stdc_rotate_left_ui(unsigned int value, unsigned int count);
    unsigned long stdc_rotate_left_ul(unsigned long value, unsigned int count);
    unsigned long long stdc_rotate_left_ull(unsigned long long value, unsigned int
        count);
    generic_value_type stdc_rotate_left(
        generic_value_type value, generic_count_type count);

```

Description

- 2 The **stdc_rotate_left** functions perform a bitwise rotate left. This operation is typically known as a left circular shift.

Returns

- 3 Let N be the width corresponding to the type of the input value. Let `no_promote_value` be `(unsigned _BitInt(N))value`. Let r be `count % N`.
- If r is 0, returns `value`;
 - otherwise, returns `(no_promote_value << r) | (no_promote_value >> (N - r)).308)`
- 4 The type-generic function (marked by its `generic_value_type` argument) returns the above described result for a given input value so long as the `generic_value_type` is:
- a standard unsigned integer type, excluding `bool`;
 - an extended unsigned integer type;
 - or, a bit-precise unsigned integer type whose width matches any standard or extended integer type, excluding `bool`.

The `generic_count_type` count argument shall be a non-negative value of signed or unsigned integer type, or `char`.

7.18.18 Rotate Right

Synopsis

```

1  #include <stdbit.h>
    unsigned char stdc_rotate_right_uc(unsigned char value, unsigned int count);
    unsigned short stdc_rotate_right_us(unsigned short value, unsigned int count);
    unsigned int stdc_rotate_right_ui(unsigned int value, unsigned int count);
    unsigned long stdc_rotate_right_ul(unsigned long value, unsigned int count);
    unsigned long long
        stdc_rotate_right_ull(unsigned long long value, unsigned int count);
    generic_value_type
        stdc_rotate_right(generic_value_type value, generic_count_type count);

```

Description

- 2 The **stdc_rotate_right** functions perform a bitwise rotate right. This operation is typically known as a right circular shift.

³⁰⁸⁾ Bit-precise integer types do not perform integer promotion in these expressions and therefore prevent potential undefined behavior from a literal implementation which would use `value` in place of `no_promote_value` (e.g., from a function call such as `stdc_rotate_left_us(0xFFFF, 15)` with an `unsigned short` width of 16 with an `int` width of 24).

Returns

- 3 Let N be the width corresponding to the type of the input value. Let `no_promote_value` be `(unsigned _BitInt(N))value`. Let r be `count % N`.
 - If r is 0, returns value;
 - otherwise, returns `(no_promote_value >> r) | (no_promote_value << (N - r)).309)`
- 4 The type-generic function (marked by its `generic_value_type` argument) returns the above described result for a given input value so long as the `generic_value_type` is:
 - a standard unsigned integer type, excluding `bool`;
 - an extended unsigned integer type;
 - or, a bit-precise unsigned integer type whose width matches any standard or extended integer type, excluding `bool`.

The `generic_count_type` count argument shall be a non-negative value of signed or unsigned integer type, or `char`.

7.18.19 8-bit Memory Reversal**Synopsis**

```
1  #include <stdbit.h>
    #include <limits.h>
    #if CHAR_BIT == 8
    void stdc_memreverse8(size_t n, unsigned char ptr[static n]);
    #endif
```

Description

- 2 The `stdc_memreverse8` function reverses the order of the bytes pointed to by `ptr`. For every `unsigned char` at offset `index` in the range $0 \leq \text{index} \ \&\& \ \text{index} < n/2$, swap the values at `ptr[index]` and `ptr[n - index - 1]`.

7.18.20 Exact-width 8-bit Memory Reversal**Synopsis**

```
1  #include <stdbit.h>
    #include <limits.h>
    #include <stdint.h>
    #if CHAR_BIT == 8
    uintN_t stdc_memreverse8uN(uintN_t value);
    #endif
```

Description

- 2 The `stdc_memreverse8uN` functions provide an interface to swap the bytes of a corresponding `uintN_t` object, where N matches one of the exact-width integer types (7.22.2.2). If an implementation provides the corresponding `uintN_t` typedef, it shall define the corresponding exact-width memory reversal function for that value of N .

Returns

- 3 The `stdc_memreverse8uN` functions returns the 8-bit memory reversed `uintN_t` value, as if by invoking `stdc_memreverse8(sizeof(value), (unsigned char*)&value)` and then returning `value`.

³⁰⁹⁾Bit-precise integer types do not perform integer promotion in these expressions and therefore prevent potential undefined behavior from a literal implementation which would use `value` in place of `no_promote_value` (e.g., from a function call such as `stdc_rotate_right_us(0xFFFF, 15)` with an `unsigned short` width of 16 with an `int` width of 24).

7.18.21 Endian-Aware 8-bit Load

Synopsis

```

1  #include <stdbit.h>
    #include <limits.h>
    #if CHAR_BIT == 8
        uint_leastN_t stdc_load8_leuN(const unsigned char ptr[static (N / 8)]);
        uint_leastN_t stdc_load8_beuN(const unsigned char ptr[static (N / 8)]);
        uint_leastN_t stdc_load8_aligned_leuN(const unsigned char ptr[static (N / 8)]);
        uint_leastN_t stdc_load8_aligned_beuN(const unsigned char ptr[static (N / 8)]);
        int_leastN_t stdc_load8_lesN(const unsigned char ptr[static (N / 8)]);
        int_leastN_t stdc_load8_besN(const unsigned char ptr[static (N / 8)]);
        int_leastN_t stdc_load8_aligned_lesN(const unsigned char ptr[static (N / 8)]);
        int_leastN_t stdc_load8_aligned_besN(const unsigned char ptr[static (N / 8)]);
    #endif

```

Description

- 2 The 8-bit load family of functions read an `int_leastN_t` or `uint_leastN_t` object from the provided `ptr` in an endian-aware (7.18.2) manner, where N matches an existing minimum-width integer type (7.22.2.3). If an implementation provides the corresponding `uint_leastN_t` typedef, it shall define the corresponding endian-aware load function for that value of N . If this function is present, N shall be a multiple of 8 and `CHAR_BIT` shall be 8. The functions containing `_aligned` in the name shall assume that `ptr` is suitably aligned to access a signed or unsigned integer of width N for a signed or unsigned variant of the function, respectively. If the function name contains the `SN` suffix in the name, it is a signed variant. Otherwise, the function is an unsigned variant. If the function name contains the `lesN` or `leuN` suffix, it is a little-endian variant. Otherwise, if the function name contains the `besN` or `beuN` suffix, it is a big-endian variant.

Returns

- 3 Let the computed value *result* be:

$$\sum_{index=0}^{(N \div 8) - 1} b_{index} \times 2^{8 \times index}$$

where b_{index} is:

- `ptr[index]`, if the function is the little-endian variant;
- otherwise, `ptr[N - index - 1]`, if the function is the big-endian variant.

If the function is an unsigned variant, return *result*. Otherwise, if the function is a signed variant, return:

- *result*, if is less than 2^{N-1} ;
- otherwise, $result - 2^N$.

7.18.22 Endian-Aware 8-bit Store

Synopsis

```

1  #include <stdbit.h>
    #include <limits.h>
    #if CHAR_BIT == 8
        void stdc_store8_leuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
        void stdc_store8_beuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
        void stdc_store8_aligned_leuN(uint_leastN_t value, unsigned char ptr[static (N /
            8)]);
    #endif

```



```

void stdc_store8_aligned_beuN(uint_leastN_t value, unsigned char ptr[static (N /
8)]);
void stdc_store8_lesN(int_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_besN(int_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_aligned_lesN(int_leastN_t value, unsigned char ptr[static (N /
8)]);
void stdc_store8_aligned_besN(int_leastN_t value, unsigned char ptr[static (N /
8)]);
#endif

```

Description

- 2 The 8-bit store family of functions write a `int_leastN_t` or `uint_leastN_t` object into the provided `ptr` in an endian-aware (7.18.2) manner, where N matches an existing minimum-width integer type (7.22.2.3). If an implementation provides the corresponding `uint_leastN_t` typedef, it shall define the corresponding endian-aware store function for that value of N . If this function is present, N shall be a multiple of 8 and `CHAR_BIT` shall be 8. The functions containing `_aligned` in the name shall assume that `ptr` is suitably aligned to access a signed or unsigned integer of width N . If the function name contains the `sN` suffix in the name, it is a signed variant. Otherwise, the function is an unsigned variant. If the function name contains the `lesN` or `leuN` suffix, it is a little-endian variant. Otherwise, if the function name contains the `besN` or `beuN` suffix, it is a big-endian variant.
- 3 Let `value_unsigned` be `value` if the function is a unsigned variant. Otherwise, let `value_unsigned` be the conversion of `value` to its corresponding unsigned type, if the function is a signed variant.

Let `index` be an integer in a sequence that

- starts from 0 and increments by 8 in the range of $[0, N)$, if the function is a little-endian variant;
- starts from $N - 8$ and decrements by 8 in the range of $[0, N)$, if the function is a big-endian variant.

Let `ptr_index` be an integer that starts from 0. For each `index` in the order of the above-specified sequence:

- Sets the 8 bits in `ptr[ptr_index]` to `(value_unsigned >> index) & 0xFF`.
- Increments `ptr_index` by 1.

7.19 Boolean type and values `<stdbool.h>`

- 1 The header `<stdbool.h>` provides the obsolescent macro `__bool_true_false_are_defined` which expands to the integer literal 1.

7.20 Checked Integer Arithmetic <stdckdint.h>

7.20.1 General

- 1 The header <stdckdint.h> defines several macros for performing checked integer arithmetic.
- 2 The macro

```
__STDC_VERSION__STDCKDINT_H__
```

is an integer constant expression with a value equivalent to 202311L.

7.20.2 Checked Integer Operation Type-generic Macros

Synopsis

```
1 #include <stdckdint.h>
   bool ckd_add(type1 *result, type2 a, type3 b);
   bool ckd_sub(type1 *result, type2 a, type3 b);
   bool ckd_mul(type1 *result, type2 a, type3 b);
```

Description

- 2 These type-generic macros perform addition, subtraction, or multiplication of the mathematical values of **a** and **b**, storing the result of the operation in ***result**, (that is, ***result** is assigned the result of computing **a + b**, **a - b**, or **a * b**). Each operation is performed as if both operands were represented in a signed integer type with infinite range, and the result was then converted from this integer type to **type1**.
- 3 Both **type2** and **type3** shall be any integer type other than “plain” **char**, **bool**, a bit-precise integer type, or an enumerated type, and they can be the same. ***result** shall be a modifiable lvalue of any integer type other than “plain” **char**, **bool**, a bit-precise integer type, or an enumerated type.

Recommended practice

- 4 It is recommended to produce a diagnostic message if **type2** or **type3** are not suitable integer types, or if ***result** is not a modifiable lvalue of a suitable integer type.

Returns

- 5 If these type-generic macros return **false**, the value assigned to ***result** correctly represents the mathematical result of the operation. Otherwise, these type-generic macros return **true**. In this case, the value assigned to ***result** is the mathematical result of the operation wrapped around to the width of ***result**.
- 6 EXAMPLE If **a** and **b** are values of type **signed int**, and **result** is a **signed long**, then

```
ckd_sub(&result, a, b);
```

indicates if **a - b** can be expressed as a **signed long**. If **signed long** has a greater width than **signed int**, this is the case and this macro invocation returns **false**.

7.21 Common definitions <stddef.h>

7.21.1 General

- 1 The header <stddef.h> defines the following macros and declares the following types. Some are also defined in other headers, as noted in their respective subclauses.

- 2 The macro

```
__STDC_VERSION_STDDEF_H__
```

is an integer constant expression with a value equivalent to 202311L.

- 3 The types are

```
ptrdiff_t
```

which is the signed integer type of the result of subtracting two pointers;

```
size_t
```

which is the unsigned integer type of the result of the **sizeof** operator;

```
max_align_t
```

which is an object type whose alignment is the greatest fundamental alignment;

```
wchar_t
```

which is an integer type whose range of values can represent codes (or part of a sequence of codes) for all the members of the supported wide execution and wide literal encodings (6.2.9); the null character shall have the code value zero. Each member of the basic character set shall have a code value equal to its value when used as the lone character in an ordinary character literal if an implementation does not define **__STDC_MB_MIGHT_NEQ_WC__**; and,

```
nullptr_t
```

which is the type of the **nullptr** predefined constant, see the subsequent description in the following subclauses.

- 4 The macros are

```
NULL
```

which expands to an implementation-defined null pointer constant;

```
unreachable()
```

which expands to a void expression whose behavior is undefined if it is reached during execution; and

```
offsetof(type, member-designator)
```

which expands to an integer constant expression that has type **size_t**, the value of which is the offset in bytes, to the subobject (designated by *member-designator*), from the beginning of any object of type *type*. The type and member designator shall be such that given

```
static type t;
```

then the expression `&(t.member-designator)` evaluates to an address constant. If the specified type name contains a comma not between matching parentheses or if the specified member is a bit-field, the behavior is undefined.

Recommended practice

- 5 The types used for `size_t` and `ptrdiff_t` should not have an integer conversion rank greater than that of **signed long int** unless the implementation supports objects large enough to make this necessary.

7.21.2 The **unreachable** macro

Synopsis

```
1  #include <stddef.h>
    void unreachable(void);
```

Description

- 2 An invocation of the function-like macro **unreachable** indicates that the particular flow control that leads to the invocation will never be taken; it receives no arguments and expands to a void expression. The program execution shall not reach such an invocation.

Returns

- 3 If a macro invocation **unreachable**() is reached during execution, the behavior is undefined.
- 4 EXAMPLE 1 The following program assumes that each execution is provided with at least one command line argument. The behavior of an execution with no arguments is undefined.

```
#include <stddef.h>
#include <stdio.h>

int main (int argc, char* argv[static argc + 1]) {
    if (argc <= 2)
        unreachable();
    else
        return printf("%s: we see %s", argv[0], argv[1]);

    return puts("this should never be reached");
}
```

Here, the **static** array size expression and the annotation of the control flow with **unreachable** indicates that the pointed-to parameter array `argv` will hold at least three elements, regardless of the circumstances. A possible optimization is that the resulting executable never performs the comparison and unconditionally executes a tail call to `printf` that never returns to the `main` function. In particular, the entire call and reference to `puts` can be omitted from the executable. No diagnostic is expected.

Because `argv` and `argc`'s values are controlled by the implementation and cannot be deterministically accounted for by the program, this program runs a high risk of engaging in completely undefined behavior.

- 5 EXAMPLE 2 The following code expresses the expectation that the argument to the function will be one of the three enumerator values despite an enumerated type allowing other, non-enumerated values to be passed. Some implementations can diagnose the lack of a **return** statement after the **switch**, but use of the **unreachable** macro signals information to the implementation that this scenario should not be possible, allowing for better diagnostic and optimization properties.

```
enum Colors { Red, Green, Blue };
int get_channel_index(enum Colors c) {
    switch (c) {
        case Red: return 0;
        case Green: return 1;
        case Blue: return 2;
    }
    unreachable();
}
```

```
}

```

7.21.3 The `nullptr_t` type

Synopsis

```
1  #include <stddef.h>
    typedef typeof_unqual(nullptr) nullptr_t;
```

Description

- 2 The `nullptr_t` type is the type of the `nullptr` predefined constant. It has only a very limited use in contexts where this type is needed to distinguish `nullptr` from other expression types. It is an unqualified complete scalar type that is different from all pointer or arithmetic types and is neither an atomic or array type and has exactly one value, `nullptr`. Default or empty initialization of an object of this type is equivalent to an initialization by `nullptr`.
- 3 The size and alignment of `nullptr_t` is the same as for a pointer to character type. An object representation of the value `nullptr` is the same as the object representation of a null pointer value of type `void*`. An lvalue conversion of an object of type `nullptr_t` with such an object representation has the value `nullptr`; if the object representation is different, the behavior is undefined.³¹⁰⁾
- 4 NOTE Because it is considered to be a scalar type, `nullptr_t` can appear in many context where `(void*)0` would be valid, for example,
 - as the operand of `alignas`, `sizeof` or `typeof` operators,
 - as the operand of an implicit or explicit conversion to a pointer type,
 - as the assignment expression in an assignment or initialization of an object of type `nullptr_t`,
 - as an argument to a parameter of type `nullptr_t` or in a varying argument list,
 - as a void expression,
 - as the operand of an implicit or explicit conversion to `bool`,
 - as an operand of a `_Generic` primary expression,
 - as an operand of the `!`, `&&`, `||` or conditional operators, or
 - as the controlling expression of an `if` or iteration statement.

³¹⁰⁾Thus, during the whole program execution an object of type `nullptr_t` evaluates to the assumed value `nullptr`.

7.22 Integer types <stdint.h>

7.22.1 General

- 1 The header <stdint.h> declares sets of integer types having specified widths, and defines corresponding sets of macros.³¹¹⁾ It also defines macros that specify limits of integer types corresponding to types defined in other standard headers.
- 2 Types are defined in the following categories:
 - integer types having certain exact widths;
 - integer types having at least certain specified widths;
 - fastest integer types having at least certain specified widths;
 - integer types wide enough to hold pointers to objects;
 - integer types having greatest width.

(Some of these types may denote the same type.)

- 3 Corresponding macros specify limits of the declared types and construct suitable constants.
- 4 For each type described herein that the implementation provides,³¹²⁾ <stdint.h> shall declare that typedef name and define the associated macros. Conversely, for each type described herein that the implementation does not provide, <stdint.h> shall not declare that typedef name nor shall it define the associated macros. An implementation shall provide those types described as “required”, but may not provide any of the others (described as “optional”). None of the types shall be defined as a synonym for a bit-precise integer type.
- 5 The feature test macro `__STDC_VERSION_STDINT_H__` expands to the token 202311L.

7.22.2 Integer types

7.22.2.1 General

- 1 When typedef names differing only in the absence or presence of the initial `u` are defined, they shall denote corresponding signed and unsigned types as described in 6.2.5; an implementation providing one of these corresponding types shall also provide the other.
- 2 In the following descriptions, the symbol *N* represents an unsigned decimal integer with no leading zeros (e.g. 8 or 24, but not 04 or 048).

7.22.2.2 Exact-width integer types

- 1 The typedef name `intN_t` designates a signed integer type with width *N* and no padding bits. Thus, `int8_t` denotes such a signed integer type with a width of exactly 8 bits.
- 2 The typedef name `uintN_t` designates an unsigned integer type with width *N* and no padding bits. Thus, `uint24_t` denotes such an unsigned integer type with a width of exactly 24 bits.
- 3 If an implementation provides standard or extended integer types with a particular width and no padding bits, it shall define the corresponding typedef names.

7.22.2.3 Minimum-width integer types

- 1 The typedef name `int_leastN_t` designates a signed integer type with a width of at least *N*, such that no signed integer type with lesser size has at least the specified width. Thus, `int_least32_t` denotes a signed integer type with a width of at least 32 bits.
- 2 The typedef name `uint_leastN_t` designates an unsigned integer type with a width of at least *N*, such that no unsigned integer type with lesser size has at least the specified width. Thus, `uint_least16_t` denotes an unsigned integer type with a width of at least 16 bits.

³¹¹⁾See “future library directions” (7.34.15).

³¹²⁾Some of these types can denote implementation-defined extended integer types.

- 3 If the typedef name `intN_t` is defined, `int_leastN_t` designates the same type. If the typedef name `uintN_t` is defined, `uint_leastN_t` designates the same type.
- 4 The following types are required:

<code>int_least8_t</code>	<code>uint_least8_t</code>
<code>int_least16_t</code>	<code>uint_least16_t</code>
<code>int_least32_t</code>	<code>uint_least32_t</code>
<code>int_least64_t</code>	<code>uint_least64_t</code>

All other types of this form are optional.

7.22.2.4 Fastest minimum-width integer types

- 1 Each of the following types designates an integer type that is usually fastest³¹³⁾ to operate with among all integer types that have at least the specified width.
- 2 The typedef name `int_fastN_t` designates the fastest signed integer type with a width of at least *N*. The typedef name `uint_fastN_t` designates the fastest unsigned integer type with a width of at least *N*.
- 3 The following types are required:

<code>int_fast8_t</code>	<code>uint_fast8_t</code>
<code>int_fast16_t</code>	<code>uint_fast16_t</code>
<code>int_fast32_t</code>	<code>uint_fast32_t</code>
<code>int_fast64_t</code>	<code>uint_fast64_t</code>

All other types of this form are optional.

7.22.2.5 Integer types capable of holding object pointers

- 1 The following type designates a signed integer type, other than a bit-precise integer type, with the property that any valid pointer to `void` can be converted to this type, then converted back to pointer to `void`, and the result will compare equal to the original pointer:

`intptr_t`

The following type designates an unsigned integer type, other than a bit-precise integer type, with the property that any valid pointer to `void` can be converted to this type, then converted back to pointer to `void`, and the result will compare equal to the original pointer:

`uintptr_t`

These types are optional.

7.22.2.6 Greatest-width integer types

- 1 The following type designates a signed integer type, other than a bit-precise integer type, capable of representing any value of any signed integer type with the possible exceptions of signed bit-precise integer types and of signed extended integer types that are wider than `long long` and that are referred by the type definition for an exact width integer type:

`intmax_t`

The following type designates the unsigned integer type that corresponds to `intmax_t`.³¹⁴⁾

³¹³⁾The designated type is not guaranteed to be fastest for all purposes; if the implementation has no clear grounds for choosing one type over another, it will simply pick some integer type satisfying the signedness and width requirements.

³¹⁴⁾Thus this type is capable of representing any value of any unsigned integer type with the possible exception of bit-precise integer types and particular extended integer types that are wider than `unsigned long long`.

uintmax_t

These types are required.

7.22.3 Widths of specified-width integer types

7.22.3.1 General

- 1 The following object-like macros specify the width of the types declared in `<stdint.h>`. Each macro name corresponds to a similar type name in 7.22.2.
- 2 Each instance of any defined macro shall be replaced by a constant expression suitable for use in `#if` preprocessing directives. Its implementation-defined value shall be equal to or greater than the value given in the subsequent subclauses, except where stated to be exactly the given value. An implementation shall define only the macros corresponding to those typedef names it actually provides.³¹⁵⁾

7.22.3.2 Width of exact-width integer types

- | | | |
|---|--------------------|------------------|
| 1 | INTN_WIDTH | exactly <i>N</i> |
| | UINTN_WIDTH | exactly <i>N</i> |

7.22.3.3 Width of minimum-width integer types

- | | | |
|---|--------------------------|----------------------------------|
| 1 | INT_LEASTN_WIDTH | exactly UINT_LEASTN_WIDTH |
| | UINT_LEASTN_WIDTH | <i>N</i> |

7.22.3.4 Width of fastest minimum-width integer types

- | | | |
|---|-------------------------|---------------------------------|
| 1 | INT_FASTN_WIDTH | exactly UINT_FASTN_WIDTH |
| | UINT_FASTN_WIDTH | <i>N</i> |

7.22.3.5 Width of integer types capable of holding object pointers

- | | | |
|---|----------------------|------------------------------|
| 1 | INTPTR_WIDTH | exactly UINTPTR_WIDTH |
| | UINTPTR_WIDTH | 16 |

7.22.3.6 Width of greatest-width integer types

- | | | |
|---|----------------------|------------------------------|
| 1 | INTMAX_WIDTH | exactly UINTMAX_WIDTH |
| | UINTMAX_WIDTH | 64 |

7.22.4 Width of other integer types

7.22.4.1 General

- 1 The following object-like macros specify the width of integer types corresponding to types defined in other standard headers.
- 2 Each instance of these macros shall be replaced by a constant expression suitable for use in `#if` preprocessing directives. Its implementation-defined value shall be equal to or greater than the corresponding value given in the following subclauses. An implementation shall define only the macros corresponding to those typedef names it actually provides.³¹⁶⁾

7.22.4.2 Width of **ptrdiff_t**

- | | | |
|---|----------------------|----|
| 1 | PTRDIFF_WIDTH | 16 |
|---|----------------------|----|

³¹⁵⁾The exact-width and pointer-holding integer types are optional.

³¹⁶⁾A freestanding implementation is not expected to provide all these types.

7.22.4.3 Width of `sig_atomic_t`

1	<code>SIG_ATOMIC_WIDTH</code>	8
---	-------------------------------	---

7.22.4.4 Width of `size_t`

1	<code>SIZE_WIDTH</code>	16
---	-------------------------	----

7.22.4.5 Width of `wchar_t`

1	<code>WCHAR_WIDTH</code>	8
---	--------------------------	---

7.22.4.6 Width of `wint_t`

1	<code>WINT_WIDTH</code>	16
---	-------------------------	----

7.22.5 Macros for constants of integer type

7.22.5.1 General

- 1 The following function-like macros expand to integer constant expressions suitable for initializing objects that have integer types corresponding to types defined in `<stdint.h>`. Each macro name corresponds to a similar type name in 7.22.2.3 or 7.22.2.6.
- 2 The argument in any instance of these macros shall be an unsuffixed integer literal (as defined in 6.4.5.2) with a value that does not exceed the limits for the corresponding type.
- 3 Each invocation of one of these macros shall expand to an integer constant expression. The type of the expression shall have the same type as would an expression of the corresponding type converted according to the integer promotions. The value of the expression shall be that of the argument. If the value is in the range of the type `intmax_t` (for a signed type) or the type `uintmax_t` (for an unsigned type), see 7.22.2.6, the expression is suitable for use in conditional expression inclusion preprocessing directives.

7.22.5.2 Macros for minimum-width constants

- 1 The macro `INTN_C(value)` expands to an integer constant expression corresponding to the type `int_leastN_t`. The macro `UINTN_C(value)` expands to an integer constant expression corresponding to the type `uint_leastN_t`. For example, if `uint_least64_t` is a name for the type `unsigned long long int`, then `UINT64_C(0x123)` can expand to the integer literal `0x123ULL`.

7.22.5.3 Macros for greatest-width constants

- 1 The following macro expands to an integer constant expression having the value specified by its argument and the type `intmax_t`:

```
INTMAX_C(value)
```

The following macro expands to an integer constant expression having the value specified by its argument and the type `uintmax_t`:

```
UINTMAX_C(value)
```

7.22.6 Maximal and minimal values of integer types

- 1 For all integer types for which there is a macro with suffix `_WIDTH` holding the width, maximum macros with suffix `_MAX` and, for all signed types, minimum macros with suffix `_MIN` are defined as by 5.3.5.3. If it is unspecified if a type is signed or unsigned and the implementation has it as an unsigned type, a minimum macro with extension `_MIN` and value 0 of the corresponding type, converted according to the integer promotions, is defined.

7.23 Input/output <stdio.h>

7.23.1 Introduction

- 1 The header <stdio.h> defines several macros, and declares three types and many functions for performing input and output.
- 2 The macro

```
__STDC_VERSION_STDIO_H__
```

is an integer constant expression with a value equivalent to 202ymmL.

- 3 The types declared are `size_t` (described in 7.21);

```
FILE
```

which is an object type capable of recording all the information needed to control a stream, including its file position indicator, a pointer to its associated buffer (if any), an *error indicator* that records whether a read/write error has occurred, and an *end-of-file indicator* that records whether the end of the file has been reached; and

```
fpos_t
```

which is a complete object type other than an array type capable of recording all the information needed to specify uniquely every position within a file.

- 4 The macros are `NULL` (described in 7.21);

```
_IOFBF
_IOLBF
_IONBF
```

which expand to integer constant expressions with distinct values, suitable for use as the third argument to the `setvbuf` function;

```
BUFSIZ
```

which expands to an integer constant expression that is the size of the buffer used by the `setbuf` function;

```
EOF
```

which expands to an integer constant expression, with type `int` and a negative value, that is returned by several functions to indicate *end-of-file*, that is, no more input from a stream;

```
FOPEN_MAX
```

which expands to an integer constant expression that is the minimum number of files that the implementation guarantees can be open simultaneously;

```
FILENAME_MAX
```

which expands to an integer constant expression that is the size needed for an array of `char` large enough to hold the longest file name string that the implementation guarantees can be opened or, if the implementation imposes no practical limit on the length of file name strings, the recommended size of an array intended to hold a file name string;³¹⁷⁾

³¹⁷⁾Of course, file name string contents are subject to other system-specific constraints; therefore *all* possible strings of length `FILENAME_MAX` cannot be expected to be opened successfully.

_PRINTF_NAN_LEN_MAX

which expands to an integer constant expression (suitable for use in conditional expression inclusion preprocessing directives) that is the maximum number of characters output for any

[-]NAN(*n-char-sequence*)

sequence.³¹⁸⁾ If an implementation has no support for NaNs, it shall be 0. **_PRINTF_NAN_LEN_MAX** shall be less than 64;

L_tmpnam

which expands to an integer constant expression that is the size needed for an array of **char** large enough to hold a temporary file name string generated by the **tmpnam** function;

SEEK_CUR
SEEK_END
SEEK_SET

which expand to integer constant expressions with distinct values, suitable for use as the third argument to the **fseek** function;

TMP_MAX

which expands to an integer constant expression that is the minimum number of unique file names that can be generated by the **tmpnam** function;

stderr
stdin
stdout

which are expressions of type “pointer to **FILE**” that point to the **FILE** objects associated, respectively, with the standard error, input, and output streams.

- 5 The header `<wchar.h>` declares functions for wide character input and output. The wide character input/output functions described in that subclause provide operations analogous to most of those described here, except that the fundamental units internal to the program are wide characters. The external representation (in the file) is a sequence of generalized multibyte characters, as described further in 7.23.3.
- 6 The input/output functions are given the following collective terms:
 - The *wide character input functions* — those functions described in 7.32 that perform input into wide characters and wide strings: **fgetwc**, **fgetws**, **getwc**, **getwchar**, **fwscanf**, **wscanf**, **vfwscanf**, and **vwscanf**.
 - The *wide character output functions* — those functions described in 7.32 that perform output from wide characters and wide strings: **fputwc**, **fputws**, **putwc**, **putwchar**, **fwprintf**, **wprintf**, **vfwprintf**, and **vwprintf**.
 - The *wide character input/output functions* — the union of the **ungetwc** function, the wide character input functions, and the wide character output functions.
 - The *byte input/output functions* — those functions described in this subclause that perform input/output: **fgetc**, **fgets**, **fprintf**, **fputc**, **fputs**, **fread**, **fscanf**, **fwrite**, **getc**, **getchar**, **printf**, **putc**, **putchar**, **puts**, **scanf**, **ungetc**, **vfprintf**, **vscanf**, **vprintf**, and **vscanf**.

Forward references: files (7.23.3), the **fseek** function (7.23.9.2), streams (7.23.2), the **tmpnam** function (7.23.4.4), `<wchar.h>` (7.32).

³¹⁸⁾If the implementation only uses the [-]NAN style, then **_PRINTF_NAN_LEN_MAX** would have the value 4.

7.23.2 Streams

- 1 Input and output, whether to or from physical devices such as terminals and tape drives, or whether to or from files supported on structured storage devices, are mapped into logical data *streams*, whose properties are more uniform than their various inputs and outputs. Two forms of mapping are supported, for *text streams* and for *binary streams*.³¹⁹⁾
- 2 A text stream is an ordered sequence of characters composed into *lines*, each line consisting of zero or more characters plus a terminating new-line character. Whether the last line requires a terminating new-line character is implementation-defined. Characters can be added, altered, or deleted on input and output to conform to differing conventions for representing text in the host environment. Thus, there is not required to be a one-to-one correspondence between the characters in a stream and those in the external representation. Data read in from a text stream will necessarily compare equal to the data that were earlier written out to that stream only if:
 - the data consist only of printing characters and the control characters horizontal tab and new-line;
 - no new-line character is immediately preceded by space characters;
 - and, the last character is a new-line character.

Whether space characters that are written out immediately before a new-line character appear when read in is implementation-defined.

- 3 A binary stream is an ordered sequence of characters that can transparently record internal data. Data read in from a binary stream shall compare equal to the data that were earlier written out to that stream, under the same implementation. Such a stream may, however, have an implementation-defined number of null characters appended to the end of the stream.
- 4 Each stream has an *orientation*. After a stream is associated with an external file, but before any operations are performed on it, the stream is unoriented. Once a wide character input/output function has been applied to an unoriented stream, the stream becomes a *wide-oriented stream*. Similarly, once a byte input/output function has been applied to an unoriented stream, the stream becomes a *byte-oriented stream*. Only a call to the **freopen** function or the **fwide** function can otherwise alter the orientation of a stream. (A successful call to **freopen** removes any orientation.)³²⁰⁾
- 5 Byte input/output functions shall not be applied to a wide-oriented stream and wide character input/output functions shall not be applied to a byte-oriented stream. The remaining stream operations do not affect, and are not affected by, a stream's orientation, except for the following additional restrictions:
 - Binary wide-oriented streams have the file-positioning restrictions ascribed to both text and binary streams.
 - For any input and output stream with pre-existing multibyte characters beyond the file position indicator, any call to **fputc** and **fputwc** runs the risk of overwriting part of a multibyte character, leaving the file with bytes that no longer consist of valid multibyte characters.

- 6 Each wide-oriented stream has an associated **mbstate_t** object that stores the current parse state of the stream. A successful call to **fgetpos** stores a representation of the value of this **mbstate_t** object as part of the value of the **fpos_t** object. A later successful call to **fsetpos** using the same stored **fpos_t** value restores the value of the associated **mbstate_t** object as well as the position within the controlled stream.
- 7 Each stream has an associated lock that is used to prevent data races when multiple threads of execution access a stream, and to restrict the interleaving of stream operations performed by multiple

³¹⁹⁾An implementation does not need to distinguish between text streams and binary streams. In such an implementation, there need be no new-line characters in a text stream nor any limit to the length of a line.

³²⁰⁾The three predefined streams **stdin**, **stdout**, and **stderr** are unoriented at program startup.

threads. Only one thread can hold this lock at a time. The lock is reentrant: a single thread can hold the lock multiple times at a given time.

- 8 All functions that read, write, position, or query the position of a stream lock the stream before accessing it. They release the lock associated with the stream when the access is complete.

Environmental limits

- 9 An implementation shall support text files with lines containing at least 254 characters, including the terminating new-line character. The value of the macro **BUFSIZ** shall be at least 256.

Forward references: the **freopen** function (7.23.5.4), the **fwide** function (7.32.3.5), **mbstate_t** (7.32.1), the **fgetpos** function (7.23.9.1), the **fsetpos** function (7.23.9.3).

7.23.3 Files

- 1 A stream is associated with an external file (which can be a physical device) by *opening* a file, which can involve *creating* a new file. Creating an existing file causes its former contents to be discarded, if necessary. If a file can support positioning requests (such as a disk file, as opposed to a terminal), then a *file position indicator* associated with the stream is positioned at the start (character number zero) of the file, unless the file is opened with append mode in which case it is implementation-defined whether the file position indicator is initially positioned at the beginning or the end of the file. The file position indicator is maintained by subsequent reads, writes, and positioning requests, to facilitate an orderly progression through the file.
- 2 Binary files are not truncated, except as defined in 7.23.5.3. Whether a write on a text stream causes the associated file to be truncated beyond that point is implementation-defined.
- 3 When a stream is *unbuffered*, characters are intended to appear from the source or at the destination as soon as possible. Otherwise characters may be accumulated and transmitted to or from the host environment as a block. When a stream is *fully buffered*, characters are intended to be transmitted to or from the host environment as a block when a buffer is filled. When a stream is *line buffered*, characters are intended to be transmitted to or from the host environment as a block when a new-line character is encountered. Furthermore, characters are intended to be transmitted as a block to the host environment when a buffer is filled, when input is requested on an unbuffered stream, or when input is requested on a line buffered stream that requires the transmission of characters from the host environment. Support for these characteristics is implementation-defined, and can be affected via the **setbuf** and **setvbuf** functions.
- 4 A file may be disassociated from a controlling stream by *closing* the file. Output streams are flushed (any unwritten buffer contents are transmitted to the host environment) before the stream is disassociated from the file. The lifetime of a **FILE** object ends when the associated file is closed (including the standard text streams). Whether a file of zero length (on which no characters have been written by an output stream) actually exists is implementation-defined.
- 5 The file may be subsequently reopened, by the same or another program execution, and its contents reclaimed or modified (if it can be repositioned at its start). If the **main** function returns to its original caller, or if the **exit** function is called, all open files are closed (hence all output streams are flushed) before program termination. Other paths to program termination, such as calling the **abort** function, are not required to close all files properly.
- 6 The address of the **FILE** object used to control a stream can be significant; a copy of a **FILE** object is not required to serve in place of the original.
- 7 At program startup, three text streams are predefined and are already opened — *standard input* (for reading conventional input), *standard output* (for writing conventional output), and *standard error* (for writing diagnostic output). As initially opened, the standard error stream is not fully buffered; the standard input and standard output streams are fully buffered if and only if the stream can be determined not to refer to an interactive device.
- 8 Functions that open additional (nontemporary) files require a *file name*, which is a string. The rules for composing valid file names are implementation-defined. Whether the same file can be simultaneously open multiple times is also implementation-defined.

- 9 Although both text and binary wide-oriented streams are conceptually sequences of wide characters, the external file associated with a wide-oriented stream is a sequence of multibyte characters, generalized as follows:
- Multibyte encodings within files can contain embedded null bytes (unlike multibyte encodings valid for use internal to the program).
 - A file is not required to begin nor end in the initial shift state.³²¹⁾
- 10 Moreover, the encodings used for multibyte characters can differ among files. Both the nature and choice of such encodings are implementation-defined.
- 11 The wide character input functions read multibyte characters from the stream and convert them to wide characters as if they were read by successive calls to the **fgetc** function. Each conversion occurs as if by a call to the **mbtowc** function, with the conversion state described by the stream's own **mbstate_t** object. The byte input functions read characters from the stream as if by successive calls to the **fgetc** function.
- 12 The wide character output functions convert wide characters to multibyte characters and write them to the stream as if they were written by successive calls to the **fputc** function. Each conversion occurs as if by a call to the **wcrtomb** function, with the conversion state described by the stream's own **mbstate_t** object. The byte output functions write characters to the stream as if by successive calls to the **fputc** function.
- 13 In some cases, some of the byte input/output functions also perform conversions between multibyte characters and wide characters. These conversions also occur as if by calls to the **mbtowc** and **wcrtomb** functions.
- 14 An encoding error occurs if the character sequence presented to the underlying **mbtowc** function does not form a valid (generalized) multibyte character, or if the code value passed to the underlying **wcrtomb** does not correspond to a valid (generalized) multibyte character. The wide character input/output functions and the byte input/output functions store the value of the macro **EILSEQ** in **errno** if and only if an encoding error occurs.

Environmental limits

- 15 The value of **FOPEN_MAX** shall be at least eight, including the three standard text streams.

Forward references: the **exit** function (7.24.5.4), the **fgetc** function (7.23.7.1), the **fopen** function (7.23.5.3), the **fputc** function (7.23.7.3), the **setbuf** function (7.23.5.5), the **setvbuf** function (7.23.5.6), the **fgetc** function (7.32.3.1), the **fputc** function (7.32.3.3), conversion state (7.32.6), the **mbtowc** function (7.32.6.4.3), the **wcrtomb** function (7.32.6.4.4).

7.23.4 Operations on files

7.23.4.1 The **remove** function

Synopsis

```
1  #include <stdio.h>
    int remove(const char *filename);
```

Description

- 2 The **remove** function causes the file whose name is the string pointed to by **filename** to be no longer accessible by that name. A subsequent attempt to open that file using that name will fail, unless it is created anew. If the file is open, the behavior of the **remove** function is implementation-defined.

Returns

- 3 The **remove** function returns zero if the operation succeeds, nonzero if it fails.

³²¹⁾Setting the file position indicator to end-of-file, as with **fseek(file, 0, SEEK_END)**, has undefined behavior for a binary stream (because of possible trailing null characters) or for any stream with state-dependent encoding that does not assuredly end in the initial shift state.

7.23.4.2 The **rename** function

Synopsis

```
1  #include <stdio.h>
    int rename(const char *old, const char *new);
```

Description

- 2 The **rename** function causes the file whose name is the string pointed to by **old** to be henceforth known by the name given by the string pointed to by **new**. The file named **old** is no longer accessible by that name. If a file named by the string pointed to by **new** exists prior to the call to the **rename** function, the behavior is implementation-defined.

Returns

- 3 The **rename** function returns zero if the operation succeeds, nonzero if it fails,³²²⁾ in which case if the file existed previously it is still known by its original name.

7.23.4.3 The **tmpfile** function

Synopsis

```
1  #include <stdio.h>
    FILE *tmpfile(void);
```

Description

- 2 The **tmpfile** function creates a temporary binary file that is different from any other existing file and that will automatically be removed when it is closed or at program termination. If the program terminates abnormally, whether an open temporary file is removed is implementation-defined. The file is opened for update with "wb+" mode.

Recommended practice

- 3 It should be possible to open at least **TMP_MAX** temporary files during the lifetime of the program (this limit may be shared with **tmpnam**) and there should be no limit on the number simultaneously open other than this limit and any limit on the number of open files (**FOPEN_MAX**).

Returns

- 4 The **tmpfile** function returns a pointer to the stream of the file that it created. If the file cannot be created, the **tmpfile** function returns a null pointer.

Forward references: the **fopen** function (7.23.5.3).

7.23.4.4 The **tmpnam** function

Synopsis

```
1  #include <stdio.h>
    char *tmpnam(char *s);
```

Description

- 2 The **tmpnam** function generates a string that is a valid file name and that is not the same as the name of an existing file.³²³⁾ The function is potentially capable of generating at least **TMP_MAX** different strings, but any or all of them can already be in use by existing files and thus not be suitable return values.
- 3 The **tmpnam** function generates a different string each time it is called.

³²²⁾ Among the reasons the implementation can cause the **rename** function to fail are that the file is open or that it is necessary to copy its contents to effectuate its renaming.

³²³⁾ Files created using strings generated by the **tmpnam** function are temporary only in the sense that their names are not expected to collide with those generated by conventional naming rules for the implementation. It is still necessary to use the **remove** function to remove such files when their use is ended, and before program termination.

- 4 Calls to the **tmpnam** function with a null pointer argument may introduce data races with each other. The implementation shall behave as if no library function calls the **tmpnam** function.

Returns

- 5 If no suitable string can be generated, the **tmpnam** function returns a null pointer. Otherwise, if the argument is a null pointer, the **tmpnam** function leaves its result in an internal static object and returns a pointer to that object (subsequent calls to the **tmpnam** function may modify the same object). If the argument is not a null pointer, it is assumed to point to an array of at least **L_tmpnam** chars; the **tmpnam** function writes its result in that array and returns the argument as its value.

Environmental limits

- 6 The value of the macro **TMP_MAX** shall be at least 25.

7.23.5 File access functions

7.23.5.1 The **fclose** function

Synopsis

```
1  #include <stdio.h>
    int fclose(FILE *stream);
```

Description

- 2 A successful call to the **fclose** function causes the stream pointed to by **stream** to be flushed and the associated file to be closed. Any unwritten buffered data for the stream are delivered to the host environment to be written to the file; any unread buffered data are discarded. Whether the call succeeds or not, the stream is disassociated from the file and any buffer set by the **setbuf** or **setvbuf** function is disassociated from the stream (and deallocated if it was automatically allocated).

Returns

- 3 The **fclose** function returns zero if the stream was successfully closed, or **EOF** if any errors were detected.

7.23.5.2 The **fflush** function

Synopsis

```
1  #include <stdio.h>
    int fflush(FILE *stream);
```

Description

- 2 If **stream** points to an output stream or an update stream in which the most recent operation was not input, the **fflush** function causes any unwritten data for that stream to be delivered to the host environment to be written to the file; otherwise, the behavior is undefined.
- 3 If **stream** is a null pointer, the **fflush** function performs this flushing action on all streams for which the behavior is defined previously in this subclause.

Returns

- 4 The **fflush** function sets the error indicator for the stream and returns **EOF** if a write error occurs, otherwise it returns zero.

Forward references: the **fopen** function (7.23.5.3).

7.23.5.3 The **fopen** function

Synopsis

```
1  #include <stdio.h>
    FILE *fopen(const char * restrict filename, const char * restrict mode);
```

Description

- 2 The **fopen** function opens the file whose name is the string pointed to by **filename**, and associates a stream with it.
- 3 The argument **mode** points to a string. The behavior is unspecified if any character occurs more than once in the string. If the string begins with one of the following characters, then the file shall be opened in the mode indicated. Otherwise, the behavior is undefined.³²⁴⁾

r	Open the file for reading.
w	Truncate the file to zero length or create the file for writing.
a	Open or create the file for writing at end-of-file (append).

The remainder of the string can contain any of the following characters in any order, and further affect how the file is opened:

b	The contents of the file shall be treated as binary, otherwise they shall be treated as text.
x	The file shall be opened in exclusive mode.
p	The file shall be opened in private mode.
+	The file shall be opened for update (both reading and writing), rather than just reading or writing.

- 4 Opening a file with read mode (**'r'** as the first character in the **mode** argument) fails if the file does not exist or cannot be read.
- 5 Opening a file with append mode (**'a'** as the first character in the **mode** argument) causes all subsequent writes to the file to be forced to the then current end-of-file at the point of buffer flush or actual write, regardless of intervening calls to the **fseek**, **fsetpos**, or **rewind** functions. Incrementing the current end-of-file by the amount of data written is atomic with respect to other threads writing to the same file provided the file was also opened in append mode. If the implementation is not capable of incrementing the current end-of-file atomically, it shall fail instead of performing non-atomic end-of-file writes. In some implementations, opening a binary file with append mode (**'b'** as the second or third character in the previously described list of **mode** argument values) may initially position the file position indicator for the stream beyond the last data written, because of null character padding.
- 6 Opening a file with exclusive mode (**'x'** as a character in the **mode** argument) fails if the file already exists or cannot be created. The check for the existence of the file and the creation of the file if it does not exist is atomic with respect to other threads and other concurrent program executions. If the implementation is not capable of performing the check for the existence of the file and the creation of the file atomically, it shall fail instead of performing a non-atomic check and creation.
- 7 Opening a file with private mode (**'p'** as a character in the **mode** argument) creates a file whose contents cannot be accessed by other means at any point in time including after program end, and whose **filename** may be ignored by implementations except to identify which storage may hold the file. If the implementation is not capable of creating private files, it shall fail.
- 8 When a file is opened with update mode (**'+'** as the second or third character in the previously described list of **mode** argument values), both input and output can be performed on the associated stream. However, output shall not be directly followed by input without an intervening call to the **fflush** function or to a file positioning function (**fseek**, **fsetpos**, or **rewind**), and input shall not be directly followed by output without an intervening call to a file positioning function, unless the

³²⁴⁾If the string begins with one of the listed mode sequences, the implementation can choose to ignore the remaining characters, or it can use them to select different kinds of a file (some of which can potentially not conform to the properties in 7.23.2).

input operation encounters end-of-file. Opening (or creating) a text file with update mode may instead open (or create) a binary stream in some implementations.

- 9 When opened, a stream is fully buffered if and only if it can be determined not to refer to an interactive device. The error and end-of-file indicators for the stream are cleared.

Returns

- 10 The **fopen** function returns a pointer to the object controlling the stream. If the open operation fails, **fopen** returns a null pointer.

Forward references: file positioning functions (7.23.9).

7.23.5.4 The **freopen** function

Synopsis

```
1  #include <stdio.h>
    FILE *freopen(const char * restrict filename, const char * restrict mode,
                  FILE * restrict stream);
```

Description

- 2 The **freopen** function opens the file whose name is the string pointed to by **filename** and associates the stream pointed to by **stream** with it. The **mode** argument is used just as in the **fopen** function.³²⁵⁾
- 3 If **filename** is a null pointer, the **freopen** function attempts to change the mode of the stream to that specified by **mode**, as if the name of the file currently associated with the stream had been used. It is implementation-defined which changes of mode are permitted (if any), and under what circumstances.
- 4 The **freopen** function first attempts to close any file that is associated with the specified stream. Failure to close the file is ignored. The error and end-of-file indicators for the stream are cleared.

Returns

- 5 The **freopen** function returns a null pointer if the open operation fails. Otherwise, **freopen** returns the value of **stream**.

7.23.5.5 The **setbuf** function

Synopsis

```
1  #include <stdio.h>
    void setbuf(FILE * restrict stream, char * restrict buf);
```

Description

- 2 Except that it returns no value, the **setbuf** function is equivalent to the **setvbuf** function invoked with the values **_IOFBF** for **mode** and **BUFSIZ** for **size**, or (if **buf** is a null pointer), with the value **_IONBF** for **mode**.

Returns

- 3 The **setbuf** function returns no value.

Forward references: the **setvbuf** function (7.23.5.6).

7.23.5.6 The **setvbuf** function

Synopsis

```
1  #include <stdio.h>
    int setvbuf(FILE * restrict stream, char * restrict buf, int mode, size_t size);
```

³²⁵⁾The primary use of the **freopen** function is to change the file associated with a standard text stream (**stderr**, **stdin**, or **stdout**), as those identifiers are not required to be modifiable lvalues to which the value returned by the **fopen** function can be assigned.

Description

- 2 The **setvbuf** function may be used only after the stream pointed to by **stream** has been associated with an open file and before any other operation (other than an unsuccessful call to **setvbuf**) is performed on the stream. The argument **mode** determines how **stream** will be buffered, as follows:

_IOFBF causes input/output to be fully buffered;

_IOLBF causes input/output to be line buffered;

_IONBF causes input/output to be unbuffered.

If **buf** is not a null pointer, the array it points to can be used instead of a buffer allocated by the **setvbuf** function³²⁶⁾ and the argument **size** specifies the size of the array; otherwise, **size** may determine the size of a buffer allocated by the **setvbuf** function. The members of the array at any time have unspecified values.

Returns

- 3 The **setvbuf** function returns zero on success, or nonzero if an invalid value is given for **mode** or if the request cannot be honored.

7.23.6 Formatted input/output functions**7.23.6.1 General**

- 1 The formatted input/output functions shall behave as if there is a sequence point after the actions associated with each specifier.³²⁷⁾

7.23.6.2 The **fprintf function****Synopsis**

- 1

```
#include <stdio.h>
int fprintf(FILE * restrict stream, const char * restrict format, ...);
```

Description

- 2 The **fprintf** function writes output to the stream pointed to by **stream**, under control of the string pointed to by **format** that specifies how subsequent arguments are converted for output. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored. The **fprintf** function returns when the end of the format string is encountered.
- 3 The format shall be a multibyte character sequence, beginning and ending in its initial shift state. The format is composed of zero or more directives: ordinary multibyte characters (not %), which are copied unchanged to the output stream; and conversion specifications, each of which results in fetching zero or more subsequent arguments, converting them, if applicable, according to the corresponding conversion specifier, and then writing the result to the output stream.
- 4 Each conversion specification is introduced by the character %. After the %, the following appear in sequence:

- Zero or more *flags* (in any order) that modify the meaning of the conversion specification.
- An optional minimum *field width*. If the converted value has fewer characters than the field width, it is padded with spaces (by default) on the left (or right, if the left adjustment flag, described later, has been given) to the field width. The field width takes the form of an asterisk * (described later) or a nonnegative decimal integer.³²⁸⁾

³²⁶⁾The buffer has to have a lifetime at least as great as the open stream, so not closing the stream before a buffer that has automatic storage duration is deallocated upon block exit results in undefined behavior.

³²⁷⁾The **fprintf** functions perform writes to memory for the %n specifier.

³²⁸⁾0 is taken as a flag, not as the beginning of a field width.

- An optional *precision* that gives the minimum number of digits to appear for the **b**, **B**, **d**, **i**, **o**, **u**, **x**, and **X** conversions, the number of digits to appear after the decimal-point character for **a**, **A**, **e**, **E**, **f**, and **F** conversions, the maximum number of significant digits for the **g** and **G** conversions, or the maximum number of bytes to be written for **s** conversions. The precision takes the form of a period (.) followed either by an asterisk * (described later) or by an optional nonnegative decimal integer; if only the period is specified, the precision is taken as zero. If a precision appears with any other conversion specifier, the behavior is undefined.
 - An optional *length modifier* that specifies the size of the argument.
 - A *conversion specifier* character that specifies the type of conversion to be applied.
- 5 As noted previously, a field width, or precision, or both, may be indicated by an asterisk. In this case, an **int** argument supplies the field width or precision. The arguments specifying field width, or precision, or both, shall appear (in that order) before the argument (if any) to be converted. A negative field width argument is taken as a - flag followed by a positive field width. A negative precision argument is taken as if the precision were omitted.
- 6 The flag characters and their meanings are:
- The result of the conversion is left-justified within the field. (It is right-justified if this flag is not specified.)
 - + The result of a signed conversion always begins with a plus or minus sign. (It begins with a sign only when a value with a negative sign is converted if this flag is not specified.)³²⁹⁾
 - space* If the first character of a signed conversion is not a sign, or if a signed conversion results in no characters, a space is prefixed to the result. If the *space* and + flags both appear, the *space* flag is ignored.
 - # The result is converted to an “alternative form”. For **o** conversion, it increases the precision, if and only if necessary, to force the first digit of the result to be a zero (if the value and precision are both 0, a single 0 is printed). For **b** conversion, a nonzero result has **0b** prefixed to it. For the optional **B** conversion as described later in this subclause, a nonzero result has **0B** prefixed to it. For **x** (or **X**) conversion, a nonzero result has **0x** (or **0X**) prefixed to it. For **a**, **A**, **e**, **E**, **f**, **F**, **g**, and **G** conversions, the result of converting a floating-point number always contains a decimal-point character, even if no digits follow it. (Normally, a decimal-point character appears in the result of these conversions only if a digit follows it.) For **g** and **G** conversions, trailing zeros are *not* removed from the result. For other conversions, the behavior is undefined.
 - 0** For **b**, **B**, **d**, **i**, **o**, **u**, **x**, **X**, **a**, **A**, **e**, **E**, **f**, **F**, **g**, and **G** conversions, leading zeros (following any indication of sign or base) are used to pad to the field width rather than performing space padding, except when converting an infinity or NaN. If the **0** and - flags both appear, the **0** flag is ignored. For **b**, **B**, **d**, **i**, **o**, **u**, **x**, and **X** conversions, if a precision is specified, the **0** flag is ignored. For other conversions, the behavior is undefined.
- 7 The length modifiers and their meanings are:
- hh** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **signed char** or **unsigned char** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **signed char** or **unsigned char** before printing); or that a following **n** conversion specifier applies to a pointer to a **signed char** argument.
 - h** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **short int** or **unsigned short int** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **short int** or

³²⁹⁾The results of all floating conversions of a negative zero, and of negative values that round to zero, include a minus sign.

unsigned short int before printing); or that a following **n** conversion specifier applies to a pointer to a **short int** argument.

- l (ell) Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **long int** or **unsigned long int** argument; that a following **n** conversion specifier applies to a pointer to a **long int** argument; that a following **c** conversion specifier applies to a **wint_t** argument; that a following **s** conversion specifier applies to a pointer to a **wchar_t** argument; or has no effect on a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier.
- ll (ell-ell) Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **long long int** or **unsigned long long int** argument; or that a following **n** conversion specifier applies to a pointer to a **long long int** argument.
- j Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to an **intmax_t** or **uintmax_t** argument; or that a following **n** conversion specifier applies to a pointer to an **intmax_t** argument.
- z Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **size_t** or the corresponding signed integer type argument; or that a following **n** conversion specifier applies to a pointer to a signed integer type corresponding to **size_t** argument.
- t Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **ptrdiff_t** or the corresponding unsigned integer type argument; or that a following **n** conversion specifier applies to a pointer to a **ptrdiff_t** argument.
- wN Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to an integer argument with a specific width where **N** is a positive decimal integer with no leading zeros (the argument will have been promoted according to the integer promotions, but its value shall be converted to the unpromoted type); or that a following **n** conversion specifier applies to a pointer to an integer type argument with a width of **N** bits. All minimum-width integer types (7.22.2.3) and exact-width integer types (7.22.2.2) defined in the header `<stdint.h>` shall be supported. Other supported values of **N** are implementation-defined.
- wfN Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a fastest minimum-width integer argument with a specific width where **N** is a positive decimal integer with no leading zeros (the argument will have been promoted according to the integer promotions, but its value shall be converted to the unpromoted type); or that a following **n** conversion specifier applies to a pointer to a fastest minimum-width integer type argument with a width of **N** bits. All fastest minimum-width integer types (7.22.2.4) defined in the header `<stdint.h>` shall be supported. Other supported values of **N** are implementation-defined.
- L Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **long double** argument.
- H Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **_Decimal32** argument.
- D Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **_Decimal64** argument.
- DD Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **_Decimal128** argument.

If a length modifier appears with any conversion specifier other than as specified previously, the behavior is undefined.

- 8 The conversion specifiers and their meanings are:

- d, i** The **int** argument is converted to signed decimal in the style `[-]dddd`. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no characters.
- b, B, o, u, x, X** The **unsigned int** argument is converted to unsigned binary (**b** or **B**), unsigned octal (**o**), unsigned decimal (**u**), or unsigned hexadecimal notation (**x** or **X**) in the style `dddd`; the letters **abcdef** are used for **x** conversion and the letters **ABCDEF** for **X** conversion. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no characters. The specifier **B** is optional and provides the same functionality as **b**, except for the **##** flag as previously specified. The **PRIB** macros from `<inttypes.h>` shall only be defined if the implementation follows the specification as given here.
- f, F** A **double** argument representing a floating-point number is converted to decimal notation in the style `[-]ddd.ddd`, where the number of digits after the decimal-point character is equal to the precision specification. If the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point character appears. If a decimal-point character appears, at least one digit appears before it. The value is rounded to the appropriate number of digits.
- A **double** argument representing an infinity is converted in one of the styles `[-]inf` or `[-]infinity` — which style is implementation-defined. A **double** argument representing a NaN is converted in one of the styles `[-]nan` or `[-]nan(n-char-sequence)` — which style, and the meaning of any *n-char-sequence*, is implementation-defined. The **F** conversion specifier produces **INF**, **INFINITY**, or **NAN** instead of **inf**, **infinity**, or **nan**, respectively.³³⁰⁾
- e, E** A **double** argument representing a floating-point number is converted in the style `[-]d.ddde±dd`, where there is one digit (which is nonzero if the argument is nonzero) before the decimal-point character and the number of digits after it is equal to the precision; if the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point character appears. The value is rounded to the appropriate number of digits. The **E** conversion specifier produces a number with **E** instead of **e** introducing the exponent. The exponent always contains at least two digits, and only as many more digits as necessary to represent the exponent. If the value is zero, the exponent is zero.
- A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.
- g, G** A **double** argument representing a floating-point number is converted in style **f** or **e** (or in style **F** or **E** in the case of a **G** conversion specifier), depending on the value converted and the precision. Let *P* equal the precision if nonzero, 6 if the precision is omitted, or 1 if the precision is zero. Then, if a conversion with style **E** would have an exponent of *X*:
- if $P > X \geq -4$, the conversion is with style **f** (or **F**) and precision $P - (X + 1)$.
- otherwise, the conversion is with style **e** (or **E**) and precision $P - 1$.
- Finally, unless the **#** flag is used, any trailing zeros are removed from the fractional portion of the result and the decimal-point character is removed if there is no fractional portion remaining.
- A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.
- a, A** A **double** argument representing a floating-point number is converted in the style `[-]0xh.hhhhp±d`, where there is one hexadecimal digit (which is nonzero if the argument is a normalized floating-point number and is otherwise unspecified) before the decimal-point

³³⁰⁾When applied to infinite and NaN values, the **-**, **+**, and *space* flag characters have their usual meaning; the **#** and **0** flag characters have no effect.

character³³¹⁾ and the number of hexadecimal digits after it is equal to the precision; if the precision is missing and **FLT_RADIX** is a power of 2, then the precision is sufficient for an exact representation of the value; if the precision is missing and **FLT_RADIX** is not a power of 2, then the precision is sufficient to distinguish³³²⁾ values of type **double**, except that trailing zeros may be omitted; if the precision is zero and the **#** flag is not specified, no decimal-point character appears. The letters **abcdef** are used for **a** conversion and the letters **ABCDEF** for **A** conversion. The **A** conversion specifier produces a number with **X** and **P** instead of **x** and **p**. The exponent always contains at least one digit, and only as many more digits as necessary to represent the decimal exponent of 2. If the value is zero, the exponent is zero.

A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.

If an **H**, **D**, or **DD** modifier is present and the precision is missing, then for a decimal floating type argument represented by a triple of integers (s, c, q) , where n is the number of significant digits in the coefficient c ,

- if $-(n + 5) \leq q \leq 0$, use style **f** (or style **F** in the case of an **A** conversion specifier) with formatting precision equal to $-q$,
- otherwise, use style **e** (or style **E** in the case of an **A** conversion specifier) with formatting precision equal to $n - 1$, with the exceptions that if $c = 0$ then the digit-sequence in the exponent-part shall have the value q (rather than 0), and that the exponent is always expressed with the minimum number of digits required to represent its value (the exponent never contains a leading zero).

If the precision P is present (in the conversion specification) and is zero or at least as large as the precision p (5.3.5.3.3) of the decimal floating type, the conversion is as if the precision were missing. If the precision P is present (and nonzero) and less than the precision p of the decimal floating type, the conversion first obtains an intermediate result as follows, where n is the number of significant digits in the coefficient:

- If $n \leq P$, set the intermediate result to the input.
- If $n > P$, round the input value, according to the current rounding direction for decimal floating-point operations, to P decimal digits, with unbounded exponent range, representing the result with a P -digit integer coefficient when in the form (s, c, q) .

Convert the intermediate result in the manner described previously for the case where the precision is missing.

- c If no **l** length modifier is present, the **int** argument is converted to an **unsigned char**, and the resulting character is written.

If an **l** length modifier is present, the **wint_t** argument is converted as if by a call to the **wcrtomb** function with a pointer to storage of at least **MB_CUR_MAX** bytes, the **wint_t** argument converted to **wchar_t**, and an initial shift state.

³³¹⁾ Binary implementations can choose the hexadecimal digit to the left of the decimal-point character so that subsequent digits align to nibble (4-bit) boundaries. This implementation choice affects numerical values printed with a precision P that is insufficient to represent all values exactly. Implementations with different conventions about the most significant hexadecimal digit will round at different places, affecting the numerical value of the hexadecimal result. For example, possible printed output for the code

```
#include <stdio.h>
/* ... */
double x = 123.0;
printf("%.1a", x);
```

include **"0x1.fp+6"** and **"0xf.6p+3"** whose numerical values are 124 and 123, respectively. Portable code seeking identical numerical results on different platforms should avoid precisions P that require rounding.

³³²⁾ The formatting precision P is sufficient to distinguish values of the source type if $16^P > b^p$ where b (not a power of 2) and p are the base and precision of the source type (5.3.5.3.3). A smaller P potentially suffices depending on the implementation's scheme for determining the digit to the left of the decimal-point character.

- s If no `l` length modifier is present, the argument shall be a pointer to storage of character type.³³³⁾ Characters from the storage are written up to (but not including) the terminating null character. If the precision is specified, no more than that many bytes are written. If the precision is not specified or is greater than the size of the storage, the storage shall contain a null character.
- If an `l` length modifier is present, the argument shall be a pointer to storage of `wchar_t` type. Wide characters from the storage are converted to multibyte characters (each as if by a call to the `wcrtomb` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first wide character is converted) up to and including a terminating null wide character. The resulting multibyte characters are written up to (but not including) the terminating null character (byte). If no precision is specified, the storage shall contain a null wide character. If a precision is specified, no more than that many bytes are written (including shift sequences, if any), and the storage shall contain a null wide character if, to equal the multibyte character sequence length given by the precision, the function would need to access a wide character one past the end of the array. In no case is a partial multibyte character written.³³⁴⁾
- p The argument shall be a pointer to `void` or a pointer to a character type. The value of the pointer is converted to a sequence of printing characters, in an implementation-defined manner.
- n The argument shall be a pointer to signed integer whose type is specified by the length modifier, if any, for the conversion specification, or shall be `int` if no length modifier is specified for the conversion specification. The number of characters written to the output stream so far by this call to `fprintf` is stored into the integer object pointed to by the argument. No argument is converted, but one is consumed. If the conversion specification includes any flags, a field width, or a precision, the behavior is undefined.
- % A `%` character is written. No argument is converted. The complete conversion specification shall be `%%`.
- 9 If a conversion specification is invalid, the behavior is undefined.³³⁵⁾ `fprintf` shall behave as if it uses `va_arg` with a type argument naming the type resulting from applying the default argument promotions to the type corresponding to the conversion specification and then converting the result of the `va_arg` expansion to the type corresponding to the conversion specification.³³⁶⁾
- 10 In no case does a nonexistent or small field width cause truncation of a field; if the result of a conversion is wider than the field width, the field is expanded to contain the conversion result.
- 11 For `a` and `A` conversions, if `FLT_RADIX` is a power of 2, the value is correctly rounded to a hexadecimal floating number with the given precision.

Recommended practice

- 12 For `a` and `A` conversions, if `FLT_RADIX` is not a power of 2 and the result is not exactly representable in the given precision, the result should be one of the two adjacent numbers in hexadecimal floating style with the given precision, with the extra stipulation that the error should have a correct sign for the current rounding direction.
- 13 For `e`, `E`, `f`, `F`, `g`, and `G` conversions, if the number of significant decimal digits is at most the maximum value M of the `T_DECIMAL_DIG` macros (defined in `<float.h>`), then the result should be correctly rounded.³³⁷⁾ If the number of significant decimal digits is more than M but the source value is exactly representable with M digits, then the result should be an exact representation with trailing zeros. Otherwise, the source value is bounded by two adjacent decimal strings $L < U$, both having M significant digits; the value of the resultant decimal string D should satisfy $L \leq D \leq U$, with the

³³³⁾No special provisions are made for multibyte characters.

³³⁴⁾Redundant shift sequences can result if multibyte characters have a state-dependent encoding.

³³⁵⁾See “future library directions” (7.34.16).

³³⁶⁾The behavior is undefined when the types differ as specified for `va_arg` 7.16.2.2.

³³⁷⁾For binary-to-decimal conversion, the result format’s values are the numbers representable with the given format specifier. The number of significant digits is determined by the format specifier, and in the case of fixed-point conversion by the source value as well.

extra stipulations that the error should have a correct sign for the current rounding direction and that increasing the number of decimal digits for the result should not decrease the accuracy of the result.

- 14 The uppercase **B** format specifier is made optional by the previous description, because it used to be available for extensions in previous versions of this document. Implementations that did not use an uppercase **B** as their own extension before are encouraged to implement it as previously described.

Returns

- 15 The **fprintf** function returns the number of characters transmitted, or a negative value if an output or encoding error occurred or if the implementation does not support a specified width length modifier.

Environmental limits

- 16 The number of characters that can be produced by any single conversion shall be at least 4095.
- 17 EXAMPLE 1 To print a date and time in the form “Sunday, July 3, 10:02” followed by π to five decimal places:

```
#include <math.h>
#include <stdio.h>
/* ... */
char *weekday, *month;    // pointers to strings
int day, hour, min;
fprintf(stdout, "%s, %s %d, %.2d:%.2d\n",
        weekday, month, day, hour, min);
fprintf(stdout, "pi = %.5f\n", 4 * atan(1.0));
```

- 18 EXAMPLE 2 In this example, multibyte characters do not have a state-dependent encoding, and the members of the extended character set that consist of more than one byte each consist of exactly two bytes, the first of which is denoted here by a $\square$ and the second by an uppercase letter.
- 19 Given the following wide string with length seven,

```
static wchar_t wstr[] = L"\squareX\squareYabc\squareZ\squareW";
```

the seven calls

```
fprintf(stdout, "|1234567890123|\n");
fprintf(stdout, "|%13ls|\n", wstr);
fprintf(stdout, "|%-13.9ls|\n", wstr);
fprintf(stdout, "|%13.10ls|\n", wstr);
fprintf(stdout, "|%13.11ls|\n", wstr);
fprintf(stdout, "|%13.15ls|\n", &wstr[2]);
fprintf(stdout, "|%13lc|\n", (wint_t) wstr[5]);
```

will print the following seven lines:

```
|1234567890123|
| \squareX\squareYabc\squareZ\squareW|
|\squareX\squareYabc\squareZ\squareW|
| \squareX\squareYabc\squareZ\squareW|
| \squareX\squareYabc\squareZ\squareW|
| \squareX\squareYabc\squareZ\squareW|
| \squareX\squareYabc\squareZ\squareW|
```

- 20 EXAMPLE 3 Following are representations of **_Decimal64** arguments as triples (s, c, q) and the corresponding character sequences **fprintf** produces with “%Da”:

(+1, 123, 0)	123
(-1, 123, 0)	-123
(+1, 123, -2)	1.23
(+1, 123, 1)	1.23e+3
(-1, 123, 1)	-1.23e+3
(+1, 123, -8)	0.00000123
(+1, 123, -9)	1.23e-7
(+1, 120, -8)	0.00000120
(+1, 120, -9)	1.20e-7
(+1, 1234567890123456, 0)	1234567890123456
(+1, 1234567890123456, 1)	1.234567890123456e+16
(+1, 1234567890123456, -1)	123456789012345.6
(+1, 1234567890123456, -21)	0.000001234567890123456
(+1, 1234567890123456, -22)	1.234567890123456e-7
(+1, 0, 0)	0
(-1, 0, 0)	-0
(+1, 0, -6)	0.000000
(+1, 0, -7)	0e-7
(+1, 0, 2)	0e+2
(+1, 5, -6)	0.000005
(+1, 50, -7)	0.0000050
(+1, 5, -7)	5e-7

To illustrate the effects of a precision specification, the sequence:

```
_Decimal32 x = 6543.00DF;           // (+1, 654300, -2)
fprintf(stdout, "%Ha\n", x);
fprintf(stdout, "%.6Ha\n", x);
fprintf(stdout, "%.5Ha\n", x);
fprintf(stdout, "%.4Ha\n", x);
fprintf(stdout, "%.3Ha\n", x);
fprintf(stdout, "%.2Ha\n", x);
fprintf(stdout, "%.1Ha\n", x);
fprintf(stdout, "%.0Ha\n", x);
```

assuming default rounding, results in:

```
6543.00
6543.00
6543.0
6543
6.54e+3
6.5e+3
7e+3
6543.00
```

To illustrate the effects of the exponent range, the sequence:

```
_Decimal32 x = 9543210e87DF;        // (+1, 9543210, 87)
_Decimal32 y = 9500000e90DF;        // (+1, 9500000, 90)
fprintf(stdout, "%.6Ha\n", x);
fprintf(stdout, "%.5Ha\n", x);
fprintf(stdout, "%.4Ha\n", x);
fprintf(stdout, "%.3Ha\n", x);
fprintf(stdout, "%.2Ha\n", x);
fprintf(stdout, "%.1Ha\n", x);
fprintf(stdout, "%.1Ha\n", y);
```

assuming default rounding, results in:

```
9.54321e+93
9.5432e+93
9.543e+93
9.54e+93
```

```

9.5e+93
1e+94
1e+97

```

To further illustrate the effects of the exponent range, the sequence:

```

_Decimal32 x = 9512345e90DF;    // (+1, 9512345, 90)
_Decimal32 y = 9512345e86DF;    // (+1, 9512345, 86)
fprintf(stdout, "%.3Ha\n", x);
fprintf(stdout, "%.2Ha\n", x);
fprintf(stdout, "%.1Ha\n", x);
fprintf(stdout, "%.2Ha\n", y);

```

assuming default rounding, results in:

```

9.51e+96
9.5e+96
1e+97
9.5e+92

```

Forward references: conversion state (7.32.6), the `wcrtomb` function (7.32.6.4.4).

7.23.6.3 The `fscanf` function

Synopsis

```

1  #include <stdio.h>
    int fscanf(FILE * restrict stream, const char * restrict format, ...);

```

Description

- 2 The `fscanf` function reads input from the stream pointed to by `stream`, under control of the string pointed to by `format` that specifies the admissible input sequences and how they are to be converted for assignment, using subsequent arguments as pointers to the objects to receive the converted input. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored.
- 3 The format shall be a multibyte character sequence, beginning and ending in its initial shift state. The format is composed of zero or more directives: one or more white-space characters, an ordinary multibyte character (neither `%` nor a white-space character), or a conversion specification. Each conversion specification is introduced by the character `%`. After the `%`, the following appear in sequence:
 - An optional assignment-suppressing character `*`.
 - An optional decimal integer greater than zero that specifies the maximum field width (in characters).
 - An optional *length modifier* that specifies the size of the receiving object.
 - A *conversion specifier* character that specifies the type of conversion to be applied.
- 4 The `fscanf` function executes each directive of the format in turn. When all directives have been executed, or if a directive fails (as detailed later in this subclause), the function returns. Failures are described as input failures (due to the occurrence of an encoding error or the unavailability of input characters), or matching failures (due to inappropriate input).
- 5 A directive composed of white-space character(s) is executed by reading input up to the first non-white-space character (which remains unread), or until no more characters can be read. The directive never fails.
- 6 A directive that is an ordinary multibyte character is executed by reading the next characters of the stream. If any of those characters differ from the ones composing the directive, the directive fails

and the differing and subsequent characters remain unread. Similarly, if end-of-file, an encoding error, or a read error prevents a character from being read, the directive fails.

- 7 A directive that is a conversion specification defines a set of matching input sequences, as described further in this subclause for each specifier. A conversion specification is executed in the following steps:
- 8 Input white-space characters are skipped, unless the specification includes a `[`, `c`, or `n` specifier.³³⁸⁾
- 9 An input item is read from the stream, unless the specification includes an `n` specifier. An input item is defined as the longest sequence of input characters which does not exceed any specified field width and which is, or is a prefix of, a matching input sequence.³³⁹⁾ The first character, if any, after the input item remains unread. If the length of the input item is zero, the execution of the directive fails; this condition is a matching failure unless end-of-file, an encoding error, or a read error prevented input from the stream, in which case it is an input failure.
- 10 Except in the case of a `%` specifier, the input item (or, in the case of a `%n` directive, the count of input characters) is converted to a type appropriate to the conversion specifier. If the input item is not a matching sequence, the execution of the directive fails: this condition is a matching failure. Unless assignment suppression was indicated by a `*`, the result of the conversion is placed in the object pointed to by the first argument following the `format` argument that has not already received a conversion result. If this object does not have an appropriate type, or if the result of the conversion cannot be represented in the object, the behavior is undefined.
- 11 The length modifiers and their meanings are:

<code>hh</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to signed char or unsigned char .
<code>h</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to short int or unsigned short int .
<code>l</code> (ell)	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to long int or unsigned long int ; that a following <code>a</code> , <code>A</code> , <code>e</code> , <code>E</code> , <code>f</code> , <code>F</code> , <code>g</code> , or <code>G</code> conversion specifier applies to an argument with type pointer to double ; or that a following <code>c</code> , <code>s</code> , or <code>[</code> conversion specifier applies to an argument with type pointer to wchar_t .
<code>ll</code> (ell-ell)	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to long long int or unsigned long long int .
<code>j</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to intmax_t or uintmax_t .
<code>z</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to size_t or the corresponding signed integer type.
<code>t</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument with type pointer to ptrdiff_t or the corresponding unsigned integer type.
<code>wN</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument which is a pointer to an integer with a specific width where <i>N</i> is a positive decimal integer with no leading zeros. All minimum-width integer types (7.22.2.3) and exact-width integer types (7.22.2.2) defined in the header <code><stdint.h></code> shall be supported. Other supported values of <i>N</i> are implementation-defined.
<code>wfN</code>	Specifies that a following <code>b</code> , <code>d</code> , <code>i</code> , <code>o</code> , <code>u</code> , <code>x</code> , <code>X</code> , or <code>n</code> conversion specifier applies to an argument which is a pointer to a fastest minimum-width integer with a specific width where <i>N</i> is a positive decimal integer with no leading zeros. All fastest minimum-width integer types

³³⁸⁾These white-space characters are not counted against a specified field width.

³³⁹⁾`fscanf` pushes back at most one input character onto the input stream. Therefore, some sequences that are acceptable to `strtod`, `strtol`, etc., are unacceptable to `fscanf`.

(7.22.2.4) defined in the header `<stdint.h>` shall be supported. Other supported values of *N* are implementation-defined.

- L Specifies that a following *a*, *A*, *e*, *E*, *f*, *F*, *g*, or *G* conversion specifier applies to an argument with type pointer to **long double**.
- H Specifies that a following *a*, *A*, *e*, *E*, *f*, *F*, *g*, or *G* conversion specifier applies to an argument with type pointer to **_Decimal32**.
- D Specifies that a following *a*, *A*, *e*, *E*, *f*, *F*, *g*, or *G* conversion specifier applies to an argument with type pointer to **_Decimal64**.
- DD Specifies that a following *a*, *A*, *e*, *E*, *f*, *F*, *g*, or *G* conversion specifier applies to an argument with type pointer to **_Decimal128**.

If a length modifier appears with any conversion specifier other than as specified previously, the behavior is undefined.

- 12 In the following, the type of the corresponding argument for a conversion specifier shall be a pointer to a type determined by the length modifiers, if any, or specified by the conversion specifier. The conversion specifiers and their meanings are:

- d Matches an optionally signed decimal integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 10 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **int**.
- b Matches an optionally signed binary integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 2 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- i Matches an optionally signed integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 0 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **int**.
- o Matches an optionally signed octal integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 8 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- u Matches an optionally signed decimal integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 10 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- x Matches an optionally signed hexadecimal integer, whose format is the same as expected for the subject sequence of the **strtoul** function with the value 16 for the *base* argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- a, e, f, g* Matches an optionally signed floating-point number, infinity, or NaN, whose format is the same as expected for the subject sequence of the **strtod** function. Unless a length modifier is specified, the corresponding argument shall be a pointer to **float**.
- c Matches a sequence of characters of exactly the number specified by the field width (1 if no field width is present in the directive).³⁴⁰⁾

³⁴⁰⁾No special provisions are made for multibyte characters in the matching rules used by the *c*, *s*, and *[* conversion specifiers — the extent of the input field is determined on a byte-by-byte basis. The resulting field is nevertheless a sequence of multibyte characters that begins in the initial shift state.

If no `l` length modifier is present, the corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence. No null character is added.

If an `l` length modifier is present, the input shall be a sequence of multibyte characters that begins in the initial shift state. Each multibyte character in the sequence is converted to a wide character as if by a call to the `mbrtowc` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first multibyte character is converted. The corresponding argument shall be a pointer to storage of `wchar_t` large enough to accept the resulting sequence of wide characters. No null wide character is added.

s Matches a sequence of non-white-space characters.³⁴⁰⁾

If no `l` length modifier is present, the corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence and a terminating null character, which will be added automatically.

If an `l` length modifier is present, the input shall be a sequence of multibyte characters that begins in the initial shift state. Each multibyte character is converted to a wide character as if by a call to the `mbrtowc` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first multibyte character is converted. The corresponding argument shall be a pointer to storage of `wchar_t` large enough to accept the sequence and the terminating null wide character, which will be added automatically.

[Matches a nonempty sequence of characters from a set of expected characters (the *scanset*).³⁴⁰⁾

If no `l` length modifier is present, the corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence and a terminating null character, which will be added automatically.

If an `l` length modifier is present, the input shall be a sequence of multibyte characters that begins in the initial shift state. Each multibyte character is converted to a wide character as if by a call to the `mbrtowc` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first multibyte character is converted. The corresponding argument shall be a pointer that points to storage of `wchar_t` large enough to accept the sequence and the terminating null wide character, which will be added automatically.

The conversion specifier includes all subsequent characters in the `format` string, up to and including the matching right bracket (`]`). The characters between the brackets (the *scanlist*) compose the scanset, unless the character after the left bracket is a circumflex (`^`), in which case the scanset contains all characters that do not appear in the scanlist between the circumflex and the right bracket. If the conversion specifier begins with `[]` or `[^]`, the right bracket character is in the scanlist and the next following right bracket character is the matching right bracket that ends the specification; otherwise the first following right bracket character is the one that ends the specification. If a `-` character is in the scanlist and is not the first, nor the second where the first character is a `^`, nor the last character, the behavior is implementation-defined.

p Matches an implementation-defined set of sequences, which should be the same as the set of sequences that may be produced by the `%p` conversion of the `fprintf` function. The corresponding argument shall be a pointer to a pointer of **void**. The input item is converted to a pointer value in an implementation-defined manner. If the input item is a value converted earlier during the same program execution, the pointer that results shall compare equal to that value; otherwise the behavior of the `%p` conversion is undefined.

n No input is consumed. The corresponding argument shall be a pointer of a signed integer type. The number of characters read from the input stream so far by this call to the `fscanf` function is stored into the integer object pointed to by the argument. Execution of a `%n` directive does not increment the assignment count returned at the completion of execution of the `fscanf` function. No argument is converted, but one is consumed. If the conversion

specification includes an assignment-suppressing character or a field width, the behavior is undefined.

% Matches a single % character; no conversion or assignment occurs. The complete conversion specification shall be %%.

- 13 If a conversion specification is invalid, the behavior is undefined.³⁴¹⁾
- 14 The conversion specifiers **A**, **E**, **F**, **G**, and **X** are also valid and behave the same as, respectively, **a**, **e**, **f**, **g**, and **x**.
- 15 Trailing white-space characters (including new-line characters) are left unread unless matched by a directive. The success of literal matches and suppressed assignments is not directly determinable other than via the %n directive.

Returns

- 16 The **fscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure or if the implementation does not support a specific width length modifier.
- 17 EXAMPLE 1 The call:

```
#include <stdio.h>
/* ... */
int n, i; float x; char name[50];
n = fscanf(stdin, "%d%f%s", &i, &x, name);
```

with the input line:

```
25 54.32E-1 thompson
```

will assign to **n** the value 3, to **i** the value 25, to **x** the value 5.432, and to **name** the sequence thompson\0.

- 18 EXAMPLE 2 The call:

```
#include <stdio.h>
/* ... */
int i; float x; char name[50];
fscanf(stdin, "%2d%f*d %[0123456789]", &i, &x, name);
```

with input:

```
56789 0123 56a72
```

will assign to **i** the value 56 and to **x** the value 789.0, will skip 0123, and will assign to **name** the sequence 56\0. The next character read from the input stream will be a.

- 19 EXAMPLE 3 To accept repeatedly from **stdin** a quantity, a unit of measure, and an item name:

```
#include <stdio.h>
/* ... */
int count; float quant; char units[21], item[21];
do {
    count = fscanf(stdin, "%f%20s of %20s", &quant, units, item);
    fscanf(stdin, "%*[^\\n]");
} while (!feof(stdin) && !ferror(stdin));
```

- 20 If the **stdin** stream contains the following lines:

³⁴¹⁾See “future library directions” (7.34.16).


```

2 quarts of oil
-12.8degrees Celsius
lots of luck
10.0LBS      of
dirt
100ergs of energy

```

the execution of the preceding example will be analogous to the following assignments:

```

quant = 2; strcpy(units, "quarts"); strcpy(item, "oil");
count = 3;
quant = -12.8; strcpy(units, "degrees");
count = 2; // "C" fails to match "o"
count = 0; // "l" fails to match "%f"
quant = 10.0; strcpy(units, "LBS"); strcpy(item, "dirt");
count = 3;
count = 0; // "100e" fails to match "%f"
count = EOF;

```

21 EXAMPLE 4 In:

```

#include <stdio.h>
/* ... */
int d1, d2, n1, n2, i;
i = sscanf("123", "%d%n%n%d", &d1, &n1, &n2, &d2);

```

the value 123 is assigned to `d1` and the value 3 to `n1`. Because `%n` can never get an input failure, the value of 3 is also assigned to `n2`. The value of `d2` is not affected. The value 1 is assigned to `i`.

22 EXAMPLE 5 The call:

```

#include <stdio.h>
/* ... */
int n, i;
n = sscanf("foo %bar 42", "foo%%bar%d", &i);

```

will assign to `n` the value 1 and to `i` the value 42 because input white-space characters are skipped for both the `%` and `d` conversion specifiers.

23 EXAMPLE 6 In these examples, multibyte characters do have a state-dependent encoding, and the members of the extended character set that consist of more than one byte each consist of exactly two bytes, the first of which is denoted here by a `□` and the second by an uppercase letter, but are only recognized as such when in the alternate shift state. The shift sequences are denoted by `↑` and `↓`, in which the first causes entry into the alternate shift state.

24 After the call:

```

#include <stdio.h>
/* ... */
char str[50];
fscanf(stdin, "a%s", str);

```

with the input line:

```
a↑□X□Y↓ bc
```

`str` will contain `↑□X□Y↓\0` assuming that none of the bytes of the shift sequences (or of the multibyte characters, in the more general case) appears to be a single-byte white-space character.

25 In contrast, after the call:

```

#include <stdio.h>
#include <stddef.h>

```

```
/* ... */
wchar_t wstr[50];
fscanf(stdin, "a%ls", wstr);
```

with the same input line, `wstr` will contain the two wide characters that correspond to `␣` and `␣` and a terminating null wide character.

26 However, the call:

```
#include <stdio.h>
#include <stddef.h>
/* ... */
wchar_t wstr[50];
fscanf(stdin, "a↑␣↓%ls", wstr);
```

with the same input line will return zero due to a matching failure against the `↓` sequence in the format string.

27 Assuming that the first byte of the multibyte character `␣` is the same as the first byte of the multibyte character `␣`, after the call:

```
#include <stdio.h>
#include <stddef.h>
/* ... */
wchar_t wstr[50];
fscanf(stdin, "a↑␣↓%ls", wstr);
```

with the same input line, zero will again be returned, but `stdin` will be left with a partially consumed multibyte character.

Forward references: the `strtod`, `strtof`, and `strtold` functions (7.24.2.6), the `strtol`, `strtoll`, `strtoul`, and `strtoull` functions (7.24.2.8), conversion state (7.32.6), the `wcrtomb` function (7.32.6.4.4).

7.23.6.4 The `printf` function

Synopsis

```
1 #include <stdio.h>
   int printf(const char * restrict format, ...);
```

Description

2 The `printf` function is equivalent to `fprintf` with the argument `stdout` interposed before the arguments to `printf`.

Returns

3 The `printf` function returns the number of characters transmitted, or a negative value if an output or encoding error occurred.

7.23.6.5 The `scanf` function

Synopsis

```
1 #include <stdio.h>
   int scanf(const char * restrict format, ...);
```

Description

2 The `scanf` function is equivalent to `fscanf` with the argument `stdin` interposed before the arguments to `scanf`.

Returns

3 The `scanf` function returns the value of the macro `EOF` if an input failure occurs before the first conversion (if any) has completed. Otherwise, the `scanf` function returns the number of input items

assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

7.23.6.6 The **snprintf** function

Synopsis

```
1  #include <stdio.h>
    int snprintf(char * restrict s, size_t n, const char * restrict format, ...);
```

Description

- 2 The **snprintf** function is equivalent to **fprintf**, except that the output is written into an array (specified by argument **s**) rather than to a stream. If **n** is zero, nothing is written, and **s** can be a null pointer. Otherwise, output characters beyond the **n**-1st are discarded rather than being written to the array, and a null character is written at the end of the characters actually written into the array. If copying takes place between objects that overlap, the behavior is undefined.

Returns

- 3 The **snprintf** function returns the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character, or a negative value if an encoding error occurred. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than **n**.

7.23.6.7 The **sprintf** function

Synopsis

```
1  #include <stdio.h>
    int sprintf(char * restrict s, const char * restrict format, ...);
```

Description

- 2 The **sprintf** function is equivalent to **fprintf**, except that the output is written into an array (specified by the argument **s**) rather than to a stream. A null character is written at the end of the characters written; it is not counted as part of the returned value. If copying takes place between objects that overlap, the behavior is undefined.

Returns

- 3 The **sprintf** function returns the number of characters written in the array, not counting the terminating null character, or a negative value if an encoding error occurred.

7.23.6.8 The **sscanf** function

Synopsis

```
1  #include <stdio.h>
    int sscanf(const char * restrict s, const char * restrict format, ...);
```

Description

- 2 The **sscanf** function is equivalent to **fscanf**, except that input is obtained from a string (specified by the argument **s**) rather than from a stream. Reaching the end of the string is equivalent to encountering end-of-file for the **fscanf** function. If copying takes place between objects that overlap, the behavior is undefined.

Returns

- 3 The **sscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **sscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

7.23.6.9 The **vfprintf** function

Synopsis

```

1  #include <stdarg.h>
    #include <stdio.h>
    int vfprintf(FILE * restrict stream, const char * restrict format, va_list arg);

```

Description

2 The **vfprintf** function is equivalent to **fprintf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** invocations). The **vfprintf** function does not invoke the **va_end** macro.³⁴²⁾

Returns

3 The **vfprintf** function returns the number of characters transmitted, or a negative value if an output or encoding error occurred.

4 EXAMPLE The following shows the use of the **vfprintf** function in a general error-reporting routine.

```

#include <stdarg.h>
#include <stdio.h>

void error(char *function_name, char *format, ...)
{
    va_list args;

    va_start(args, format);
    // print out name of function causing error
    fprintf(stderr, "ERROR in %s: ", function_name);
    // print out remainder of message
    vfprintf(stderr, format, args);
    va_end(args);
}

```

7.23.6.10 The **vfscanf** function

Synopsis

```

1  #include <stdarg.h>
    #include <stdio.h>
    int vfscanf(FILE * restrict stream, const char * restrict format, va_list arg);

```

Description

2 The **vfscanf** function is equivalent to **fscanf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls). The **vfscanf** function does not invoke the **va_end** macro.³⁴²⁾

Returns

3 The **vfscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **vfscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

7.23.6.11 The **vprintf** function

Synopsis

```

1  #include <stdarg.h>
    #include <stdio.h>
    int vprintf(const char * restrict format, va_list arg);

```

³⁴²⁾As the functions **vfprintf**, **vfscanf**, **vprintf**, **vscanf**, **vsprintf**, **vsnprintf**, **vsprintf**, and **vsscanf** invoke the **va_arg** macro, **arg** after the return has an indeterminate representation.

Description

- 2 The **vprintf** function is equivalent to **printf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls). The **vprintf** function does not invoke the **va_end** macro.³⁴²⁾

Returns

- 3 The **vprintf** function returns the number of characters transmitted, or a negative value if an output or encoding error occurred.

7.23.6.12 The vscanf function**Synopsis**

```
1  #include <stdarg.h>
    #include <stdio.h>
    int vscanf(const char * restrict format, va_list arg);
```

Description

- 2 The **vscanf** function is equivalent to **scanf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls). The **vscanf** function does not invoke the **va_end** macro.³⁴²⁾

Returns

- 3 The **vscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **vscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

7.23.6.13 The vsnprintf function**Synopsis**

```
1  #include <stdarg.h>
    #include <stdio.h>
    int vsnprintf(char * restrict s, size_t n, const char * restrict format, va_list
        arg);
```

Description

- 2 The **vsnprintf** function is equivalent to **snprintf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls). The **vsnprintf** function does not invoke the **va_end** macro.³⁴²⁾ If copying takes place between objects that overlap, the behavior is undefined.

Returns

- 3 The **vsnprintf** function returns the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character, or a negative value if an encoding error occurred. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than **n**.

7.23.6.14 The vsprintf function**Synopsis**

```
1  #include <stdarg.h>
    #include <stdio.h>
    int vsprintf(char * restrict s, const char * restrict format, va_list arg);
```

Description

- 2 The **vsprintf** function is equivalent to **sprintf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls).

The **vsprintf** function does not invoke the **va_end** macro.³⁴²⁾ If copying takes place between objects that overlap, the behavior is undefined.

Returns

- 3 The **vsprintf** function returns the number of characters written in the array, not counting the terminating null character, or a negative value if an encoding error occurred.

7.23.6.15 The **vsscanf** function

Synopsis

```
1  #include <stdarg.h>
    #include <stdio.h>
    int vsscanf(const char * restrict s, const char * restrict format, va_list arg);
```

Description

- 2 The **vsscanf** function is equivalent to **sscanf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** calls). The **vsscanf** function does not invoke the **va_end** macro.³⁴²⁾

Returns

- 3 The **vsscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **vsscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

7.23.7 Character input/output functions

7.23.7.1 The **fgetc** function

Synopsis

```
1  #include <stdio.h>
    int fgetc(FILE *stream);
```

Description

- 2 If the end-of-file indicator for the input stream pointed to by **stream** is not set and a next character is present, the **fgetc** function obtains that character as an **unsigned char** converted to an **int** and advances the associated file position indicator for the stream (if defined).

Returns

- 3 If the end-of-file indicator for the stream is set, or if the stream is at end-of-file, the end-of-file indicator for the stream is set and the **fgetc** function returns **EOF**. Otherwise, the **fgetc** function returns the next character from the input stream pointed to by **stream**. If a read error occurs, the error indicator for the stream is set and the **fgetc** function returns **EOF**.³⁴³⁾

7.23.7.2 The **fgets** function

Synopsis

```
1  #include <stdio.h>
    char *fgets(char * restrict s, int n, FILE * restrict stream);
```

Description

- 2 The **fgets** function reads at most one less than the number of characters specified by **n** from the stream pointed to by **stream** into the array pointed to by **s**. No additional characters are read after a new-line character (which is retained) or after end-of-file. A null character is written immediately after the last character read into the array. If **n** is negative or zero, the behavior is undefined.

³⁴³⁾ An end-of-file and a read error can be distinguished by use of the **feof** and **ferror** functions.

Returns

- 3 The **fgetc** function returns **s** if successful. If end-of-file is encountered and no characters have been read into the array, the contents of the array remain unchanged and a null pointer is returned. If a read error occurs during the operation, the members of the array have unspecified values and a null pointer is returned.

7.23.7.3 The fputc function**Synopsis**

```
1  #include <stdio.h>
    int fputc(int c, FILE *stream);
```

Description

- 2 The **fputc** function writes the character specified by **c** (converted to an **unsigned char**) to the output stream pointed to by **stream**, at the position indicated by the associated file position indicator for the stream (if defined), and advances the indicator appropriately. If the file cannot support positioning requests, or if the stream was opened with append mode, the character is appended to the output stream.

Returns

- 3 The **fputc** function returns the character written. If a write error occurs, the error indicator for the stream is set and **fputc** returns **EOF**.

7.23.7.4 The fputs function**Synopsis**

```
1  #include <stdio.h>
    int fputs(const char * restrict s, FILE * restrict stream);
```

Description

- 2 The **fputs** function writes the string pointed to by **s** to the stream pointed to by **stream**. The terminating null character is not written.

Returns

- 3 The **fputs** function returns **EOF** if a write error occurs; otherwise it returns a nonnegative value.

7.23.7.5 The getc function**Synopsis**

```
1  #include <stdio.h>
    int getc(FILE *stream);
```

Description

- 2 The **getc** function is equivalent to **fgetc**, except that if it is implemented as a macro, it may evaluate **stream** more than once, so the argument should never be an expression with side effects.

Returns

- 3 The **getc** function returns the next character from the input stream pointed to by **stream**. If the stream is at end-of-file, the end-of-file indicator for the stream is set and **getc** returns **EOF**. If a read error occurs, the error indicator for the stream is set and **getc** returns **EOF**.

7.23.7.6 The getchar function**Synopsis**

```
1  #include <stdio.h>
    int getchar(void);
```

Description

- 2 The **getchar** function is equivalent to **getc** with the argument **stdin**.

Returns

- 3 The **getchar** function returns the next character from the input stream pointed to by **stdin**. If the stream is at end-of-file, the end-of-file indicator for the stream is set and **getchar** returns **EOF**. If a read error occurs, the error indicator for the stream is set and **getchar** returns **EOF**.

7.23.7.7 The **putc function****Synopsis**

```
1  #include <stdio.h>
    int putc(int c, FILE *stream);
```

Description

- 2 The **putc** function is equivalent to **fputc**, except that if it is implemented as a macro, it may evaluate **stream** more than once, so that argument should never be an expression with side effects.

Returns

- 3 The **putc** function returns the character written. If a write error occurs, the error indicator for the stream is set and **putc** returns **EOF**.

7.23.7.8 The **putchar function****Synopsis**

```
1  #include <stdio.h>
    int putchar(int c);
```

Description

- 2 The **putchar** function is equivalent to **putc** with the second argument **stdout**.

Returns

- 3 The **putchar** function returns the character written. If a write error occurs, the error indicator for the stream is set and **putchar** returns **EOF**.

7.23.7.9 The **puts function****Synopsis**

```
1  #include <stdio.h>
    int puts(const char *s);
```

Description

- 2 The **puts** function writes the string pointed to by **s** to the stream pointed to by **stdout**, and appends a new-line character to the output. The terminating null character is not written.

Returns

- 3 The **puts** function returns **EOF** if a write error occurs; otherwise it returns a nonnegative value.

7.23.7.10 The **ungetc function****Synopsis**

```
1  #include <stdio.h>
    int ungetc(int c, FILE *stream);
```

Description

- 2 The **ungetc** function pushes the character specified by **c** (converted to an **unsigned char**) back onto the input stream pointed to by **stream**. Pushed-back characters will be returned by subsequent

reads on that stream in the reverse order of their pushing. A successful intervening call (with the stream pointed to by `stream`) to a file positioning function (`fseek`, `fsetpos`, or `rewind`) discards any pushed-back characters for the stream. The external storage corresponding to the stream is unchanged.

- 3 One character of pushback is guaranteed. If the `ungetc` function is called too many times on the same stream without an intervening read or file positioning operation on that stream, the operation can fail.
- 4 If the value of `c` equals that of the macro `EOF`, the operation fails and the input stream is unchanged.
- 5 A successful call to the `ungetc` function clears the end-of-file indicator for the stream. The value of the file position indicator for the stream after reading or discarding all pushed-back characters shall be the same as it was before the characters were pushed back.³⁴⁴⁾ For a text stream, the value of its file position indicator after a successful call to the `ungetc` function is unspecified until all pushed-back characters are read or discarded. For a binary stream, its file position indicator is decremented by each successful call to the `ungetc` function; if its value was zero before a call, it has an indeterminate representation after the call.³⁴⁵⁾

Returns

- 6 The `ungetc` function returns the character pushed back after conversion, or `EOF` if the operation fails.

Forward references: file positioning functions (7.23.9).

7.23.8 Direct input/output functions

7.23.8.1 The `fread` function

Synopsis

```
1  #include <stdio.h>
    size_t fread(void * restrict ptr, size_t size, size_t nmemb,
                FILE * restrict stream);
```

Description

- 2 The `fread` function reads, into the array pointed to by `ptr`, up to `nmemb` elements whose size is specified by `size`, from the stream pointed to by `stream`. For each object, `size` calls are made to the `fgetc` function and the results stored, in the order read, in an array of `unsigned char` exactly overlaying the object. The file position indicator for the stream (if defined) is advanced by the number of characters successfully read. If an error occurs, the resulting representation of the file position indicator for the stream is indeterminate. If a partial element is read, its representation is indeterminate.

Returns

- 3 The `fread` function returns the number of elements successfully read, which can be less than `nmemb` if a read error or end-of-file is encountered. If `size` or `nmemb` is zero, `fread` returns zero and the contents of the array and the state of the stream remain unchanged.

7.23.8.2 The `fwrite` function

Synopsis

```
1  #include <stdio.h>
    size_t fwrite(const void * restrict ptr, size_t size, size_t nmemb,
                FILE * restrict stream);
```

Description

- 2 The `fwrite` function writes, from the array pointed to by `ptr`, up to `nmemb` elements whose size is specified by `size`, to the stream pointed to by `stream`. For each object, `size` calls are made to the

³⁴⁴⁾A file positioning function can further modify the file position indicator after discarding any pushed-back characters.

³⁴⁵⁾See “future library directions” (7.34.16).

fputc function, taking the values (in order) from an array of **unsigned char** exactly overlaying the object. The file position indicator for the stream (if defined) is advanced by the number of characters successfully written. If an error occurs, the resulting representation of the file position indicator for the stream is indeterminate.

Returns

- 3 The **fwrite** function returns the number of elements successfully written, which will be less than **nmemb** only if a write error is encountered. If **size** or **nmemb** is zero, **fwrite** returns zero and the state of the stream remains unchanged.

7.23.9 File positioning functions

7.23.9.1 The **fgetpos** function

Synopsis

```
1  #include <stdio.h>
    int fgetpos(FILE * restrict stream, fpos_t * restrict pos);
```

Description

- 2 The **fgetpos** function stores the current values of the parse state (if any) and file position indicator for the stream pointed to by **stream** in the object pointed to by **pos**. The values stored contain unspecified information usable by the **fsetpos** function for repositioning the stream to its position at the time of the call to the **fgetpos** function.

Returns

- 3 If successful, the **fgetpos** function returns zero; on failure, the **fgetpos** function returns nonzero and stores an implementation-defined positive value in **errno**.

Forward references: the **fsetpos** function (7.23.9.3).

7.23.9.2 The **fseek** function

Synopsis

```
1  #include <stdio.h>
    int fseek(FILE *stream, long int offset, int whence);
```

Description

- 2 The **fseek** function sets the file position indicator for the stream pointed to by **stream**. If a read or write error occurs, the error indicator for the stream is set and **fseek** fails.
- 3 For a binary stream, the new position, measured in characters from the beginning of the file, is obtained by adding **offset** to the position specified by **whence**. The specified position is the beginning of the file if **whence** is **SEEK_SET**, the current value of the file position indicator if **SEEK_CUR**, or end-of-file if **SEEK_END**. A binary stream can fail to meaningfully support **fseek** calls with a **whence** value of **SEEK_END**.
- 4 For a text stream, either **offset** shall be zero, or **offset** shall be a value returned by an earlier successful call to the **ftell** function on a stream associated with the same file and **whence** shall be **SEEK_SET**.
- 5 After determining the new position, a successful call to the **fseek** function undoes any effects of the **ungetc** function on the stream, clears the end-of-file indicator for the stream, and then establishes the new position. After a successful **fseek** call, the next operation on an update stream may be either input or output.

Returns

- 6 The **fseek** function returns nonzero only for a request that cannot be satisfied.

Forward references: the **ftell** function (7.23.9.4).

7.23.9.3 The **fsetpos** function

Synopsis

```
1  #include <stdio.h>
    int fsetpos(FILE *stream, const fpos_t *pos);
```

Description

- 2 The **fsetpos** function sets the **mbstate_t** object (if any) and file position indicator for the stream pointed to by **stream** according to the value of the object pointed to by **pos**, which shall be a value obtained from an earlier successful call to the **fgetpos** function on a stream associated with the same file. If a read or write error occurs, the error indicator for the stream is set and **fsetpos** fails.
- 3 A successful call to the **fsetpos** function undoes any effects of the **ungetc** function on the stream, clears the end-of-file indicator for the stream, and then establishes the new parse state and position. After a successful **fsetpos** call, the next operation on an update stream may be either input or output.

Returns

- 4 If successful, the **fsetpos** function returns zero; on failure, the **fsetpos** function returns nonzero and stores an implementation-defined positive value in **errno**.

7.23.9.4 The **ftell** function

Synopsis

```
1  #include <stdio.h>
    long int ftell(FILE *stream);
```

Description

- 2 The **ftell** function obtains the current value of the file position indicator for the stream pointed to by **stream**. For a binary stream, the value is the number of characters from the beginning of the file. For a text stream, its file position indicator contains unspecified information, usable by the **fseek** function for returning the file position indicator for the stream to its position at the time of the **ftell** call; the difference between two such return values is not necessarily a meaningful measure of the number of characters written or read.

Returns

- 3 If successful, the **ftell** function returns the current value of the file position indicator for the stream. On failure, the **ftell** function returns -1L and stores an implementation-defined positive value in **errno**.

7.23.9.5 The **rewind** function

Synopsis

```
1  #include <stdio.h>
    void rewind(FILE *stream);
```

Description

- 2 The **rewind** function sets the file position indicator for the stream pointed to by **stream** to the beginning of the file. It is equivalent to

```
(void)fseek(stream, 0L, SEEK_SET)
```

except that the error indicator for the stream is also cleared.

Returns

- 3 The **rewind** function returns no value.

7.23.10 Error-handling functions

7.23.10.1 The **clearerr** function

Synopsis

```
1  #include <stdio.h>
    void clearerr(FILE *stream);
```

Description

2 The **clearerr** function clears the end-of-file and error indicators for the stream pointed to by **stream**.

Returns

3 The **clearerr** function returns no value.

7.23.10.2 The feof function**Synopsis**

```
1  #include <stdio.h>
    int feof(FILE *stream);
```

Description

2 The **feof** function tests the end-of-file indicator for the stream pointed to by **stream**.

Returns

3 The **feof** function returns nonzero if and only if the end-of-file indicator is set for **stream**.

7.23.10.3 The ferror function**Synopsis**

```
1  #include <stdio.h>
    int ferror(FILE *stream);
```

Description

2 The **ferror** function tests the error indicator for the stream pointed to by **stream**.

Returns

3 The **ferror** function returns nonzero if and only if the error indicator is set for **stream**.

7.23.10.4 The perror function**Synopsis**

```
1  #include <stdio.h>
    void perror(const char *s);
```

Description

2 The **perror** function maps the error number in the integer expression **errno** to an error message. It writes a sequence of characters to the standard error stream thus: first (if **s** is not a null pointer and the character pointed to by **s** is not the null character), the string pointed to by **s** followed by a colon (:) and a space; then an appropriate error message string followed by a new-line character. The contents of the error message strings are the same as those returned by the **strerror** function with argument **errno**.

Returns

3 The **perror** function returns no value.

Forward references: the **strerror** function (7.27.6.3).

7.24 General utilities <stdlib.h>

7.24.1 General

- 1 The header <stdlib.h> declares several types and functions of general utility, and defines several macros.³⁴⁶⁾
- 2 The feature test macro `__STDC_VERSION_STDLIB_H` expands to the token 202ymmL.
- 3 The types declared are `size_t` and `wchar_t` (both described in 7.21), `once_flag` (described in 7.29),

```
div_t
```

which is a structure type that is the type of the value returned by the `div` function,

```
ldiv_t
```

which is a structure type that is the type of the value returned by the `ldiv` function, and

```
lldiv_t
```

which is a structure type that is the type of the value returned by the `lldiv` function.

- 4 The macros defined are `NULL` (described in 7.21); `ONCE_FLAG_INIT` (described in 7.29);

```
EXIT_FAILURE
```

and

```
EXIT_SUCCESS
```

which expand to integer constant expressions that can be used as the argument to the `exit` function to return unsuccessful or successful termination status, respectively, to the host environment;

```
RAND_MAX
```

which expands to an integer constant expression that is the maximum value returned by the `rand` function; and

```
MB_CUR_MAX
```

which expands to a positive integer expression with type `size_t` that is the maximum number of bytes in a multibyte character for the extended character set specified by the current locale (category `LC_CTYPE`), which is never greater than `MB_LEN_MAX`.

- 5 The function

```
#include <stdlib.h>
void call_once(once_flag *flag, void (*func)(void));
```

is described in 7.29.2.

7.24.2 Numeric conversion functions

7.24.2.1 General

- 1 The functions `atof`, `atoi`, `atol`, and `atoll` are not required to affect the value of the integer expression `errno` on an error. If the value of the result cannot be represented, the behavior is undefined.

7.24.2.2 The `atof` function

³⁴⁶⁾See “future library directions” (7.34.17).

Synopsis

```
1  #include <stdlib.h>
    double atof(const char *nptr);
```

Description

- 2 The **atof** function converts the initial portion of the string pointed to by **nptr** to **double** representation. Except for the behavior on error, it is equivalent to

```
    strtod(nptr, nullptr)
```

Returns

- 3 The **atof** function returns the converted value.

Forward references: the **strtod**, **strtodf**, and **strtold** functions (7.24.2.6).

7.24.2.3 The **atoi**, **atol**, and **atoll** functions

Synopsis

```
1  #include <stdlib.h>
    int atoi(const char *nptr);
    long int atol(const char *nptr);
    long long int atoll(const char *nptr);
```

Description

- 2 The **atoi**, **atol**, and **atoll** functions convert the initial portion of the string pointed to by **nptr** to **int**, **long int**, and **long long int** representation, respectively. Except for the behavior on error, they are equivalent to

```
    atoi: (int)strtoul(nptr, nullptr, 10)
    atol: strtoul(nptr, nullptr, 10)
    atoll: strtoll(nptr, nullptr, 10)
```

Returns

- 3 The **atoi**, **atol**, and **atoll** functions return the converted value.

Forward references: the **strtoul**, **strtoll**, **strtoul**, and **strtoull** functions (7.24.2.8).

7.24.2.4 The **strfromd**, **strfromf**, and **strfroml** functions

Synopsis

```
1  #include <stdlib.h>
    int strfromd(char * restrict s, size_t n, const char * restrict format,
        double fp);
    int strfromf(char * restrict s, size_t n, const char * restrict format,
        float fp);
    int strfroml(char * restrict s, size_t n, const char * restrict format,
        long double fp);
```

Description

- 2 The **strfromd**, **strfromf**, and **strfroml** functions are equivalent to **snprintf**(**s**, **n**, **format**, **fp**) (7.23.6.6), except that the default argument promotions are not applied and the format string shall only contain the character %, an optional precision that does not contain an asterisk *, and one of the conversion specifiers **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G**, which applies to the type (**double**, **float**, or **long double**) indicated by the function suffix (rather than by a length modifier). Use of these functions with any other format string results in undefined behavior.

Returns

- 3 The **strfromd**, **strfromf**, and **strfroml** functions return the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than **n**.

7.24.2.5 The **strfromdN functions****Synopsis**

```
1  #include <stdlib.h>
    #ifdef __STDC_IEC_60559_DFP__
    int strfromd32(char * restrict s, size_t n, const char * restrict format,
        _Decimal32 fp);
    int strfromd64(char * restrict s, size_t n, const char * restrict format,
        _Decimal64 fp);
    int strfromd128(char * restrict s, size_t n, const char * restrict format,
        _Decimal128 fp);
    #endif
```

Description

- 2 The **strfromdN** functions are equivalent to **snprintf(s, n, format, fp)** (7.23.6.6), except the format string contains only the character %, an optional precision that does not contain an asterisk *, and one of the conversion specifiers **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G**, which applies to the type (**_Decimal32**, **_Decimal64**, or **_Decimal128**) indicated by the function suffix (rather than by a length modifier). Use of these functions with any other format string results in undefined behavior.

Returns

- 3 The **strfromdN** functions return the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than **n**.

7.24.2.6 The **strtod, **strtof**, and **strtold** functions****Synopsis**

```
1  #include <stdlib.h>
    double strtod(const char * restrict nptr, char ** restrict endptr);
    float strtof(const char * restrict nptr, char ** restrict endptr);
    long double strtold(const char * restrict nptr, char ** restrict endptr);
```

Description

- 2 The **strtod**, **strtof**, and **strtold** functions convert the initial portion of the string pointed to by **nptr** to **double**, **float**, and **long double** representation, respectively. First, they decompose the input string into three parts: an initial, possibly empty, sequence of white-space characters, a subject sequence resembling a floating literal or representing an infinity or NaN; and a final string of one or more unrecognized characters, including the terminating null character of the input string. Then, they attempt to convert the subject sequence to a floating-point number, and return the result.
- 3 The expected form of the subject sequence is an optional plus or minus sign, then one of the following:
- decimal form: a nonempty sequence of decimal digits optionally containing a decimal-point character, then an optional exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2);
 - hexadecimal form: a **0x** or **0X**, then a nonempty sequence of hexadecimal digits optionally containing a decimal-point character, then an optional binary exponent part as defined in 6.4.5.3, excluding any digit separators;
 - **INF** or **INFINITY**, ignoring case

- **NAN** or **NAN**(*n-char-sequence_{opt}*), ignoring case in the **NAN** part, where:
n-char-sequence:

digit
nondigit
n-char-sequence digit
n-char-sequence nondigit

The subject sequence is defined as the longest initial subsequence of the input string, starting with the first non-white-space character, that is of the expected form. The subject sequence contains no characters if the input string is not of the expected form.

- 4 If the subject sequence has the decimal or hexadecimal form, the sequence of characters starting with the first digit or the decimal-point character (whichever occurs first) is interpreted as a floating literal according to the rules of 6.4.5.3, except that the decimal-point character is used in place of a period, and that if neither an exponent part nor a decimal-point character appears in a decimal form, or if a binary exponent part does not appear in a hexadecimal form, an exponent part of the appropriate type with value zero is assumed to follow the last digit in the string. If the subject sequence begins with a minus sign, the sequence is interpreted as arithmetically negated.³⁴⁷⁾
- 5 A character sequence **INF** or **INFINITY** is interpreted as an infinity, if representable in the return type, else like a floating literal that is too large for the range of the return type. A character sequence **NAN** or **NAN**(*n-char-sequence_{opt}*) is interpreted as a quiet NaN, if supported in the return type, else like a subject sequence part that does not have the expected form; the meaning of the n-char sequence is implementation-defined.³⁴⁸⁾ A pointer to the final string is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.
- 6 If the subject sequence has the hexadecimal form and **FLT_RADIX** is a power of 2, the value resulting from the conversion is correctly rounded.
- 7 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 8 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of **nptr** is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.

Recommended practice

- 9 If the subject sequence has the hexadecimal form, **FLT_RADIX** is not a power of 2, and the result is not exactly representable, the result should be one of the two numbers in the appropriate internal format that are adjacent to the hexadecimal floating source value, with the extra stipulation that the error should have a correct sign for the current rounding direction.
- 10 If the subject sequence has the decimal form and at most *M* significant digits, where *M* is the maximum value of the **T_DECIMAL_DIG** macros (defined in <float.h>), the result should be correctly rounded. If the subject sequence *D* has the decimal form and more than *M* significant digits, consider the two bounding, adjacent decimal strings *L* and *U*, both having *M* significant digits, such that the values of *L*, *D*, and *U* satisfy $L \leq D \leq U$. The result should be one of the (equal or adjacent) values that would be obtained by correctly rounding *L* and *U* according to the current rounding direction, with the extra stipulation that the error with respect to *D* should have a correct sign for the current rounding direction.³⁴⁹⁾

Returns

- 11 The functions return the converted value, if any. If no conversion could be performed, positive or unsigned zero is returned.
- 12 If the correct value overflows and default rounding is in effect (7.12.2), plus or minus **HUGE_VAL**,

³⁴⁷⁾ It is unspecified whether a minus-signed sequence is converted to a negative number directly or by arithmetically negating the value resulting from converting the corresponding unsigned sequence (see F.5); the two methods can yield different results if rounding is toward positive or negative infinity. In either case, the functions honor the sign of zero if floating-point arithmetic supports signed zeros.

³⁴⁸⁾ An implementation can use the n-char sequence to determine extra information to be represented in the NaN's significand.

³⁴⁹⁾ *M* is sufficiently large that *L* and *U* will usually correctly round to the same internal floating value, but if not will correctly round to adjacent values.

HUGE_VALF, or **HUGE_VALL** is returned (according to the return type and sign of the value); if the integer expression **math_errhandling** & **MATH_ERRNO** is nonzero, the integer expression **errno** acquires the value of **ERANGE**; if the integer expression **math_errhandling** & **MATH_ERREXCEPT** is nonzero, the “overflow” floating-point exception is raised.

- 13 If the result underflows (7.12.2), the functions return a value whose magnitude is no greater than the smallest normalized positive number in the return type; if the integer expression **math_errhandling** & **MATH_ERRNO** is nonzero, whether **errno** acquires the value **ERANGE** is implementation-defined; if the integer expression **math_errhandling** & **MATH_ERREXCEPT** is nonzero, whether the “underflow” floating-point exception is raised is implementation-defined.

7.24.2.7 The **strtodN** functions

Synopsis

```
1  #include <stdlib.h>
    #ifdef __STDC_IEC_60559_DFP__
    _Decimal32 strtod32(const char * restrict nptr, char ** restrict endptr);
    _Decimal64 strtod64(const char * restrict nptr, char ** restrict endptr);
    _Decimal128 strtod128(const char * restrict nptr, char ** restrict endptr);
    #endif
```

Description

- 2 The **strtodN** functions convert the initial portion of the string pointed to by **nptr** to decimal floating type representation. First, they decompose the input string into three parts: an initial, possibly empty, sequence of white-space characters; a subject sequence resembling a floating literal or representing an infinity or NaN; and a final string of one or more unrecognized characters, including the terminating null character of the input string. Then, they attempt to convert the subject sequence to a floating-point number, and return the result.
- 3 The expected form of the subject sequence is an optional plus or minus sign, then one of the following:
- decimal form: a nonempty sequence of decimal digits optionally containing a decimal-point character, then an optional exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2)
 - hexadecimal form: a **0x** or **0X**, then a nonempty sequence of hexadecimal digits optionally containing a decimal-point character, then an optional binary exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2)
 - **INF** or **INFINITY**, ignoring case
 - **NAN** or **NAN(*d-char-sequence_{opt}*)**, ignoring case in the **NAN** part, where:

d-char-sequence:

```
digit
nondigit
d-char-sequence digit
d-char-sequence nondigit
```

The subject sequence is defined as the longest initial subsequence of the input string, starting with the first non-white-space character, that is of the expected form. The subject sequence contains no characters if the input string is not of the expected form.

- 4 If the subject sequence has the decimal or hexadecimal form, the sequence of characters starting with the first digit or the decimal-point character (whichever occurs first) is interpreted as a floating literal according to the rules of 6.4.5.3, except that the decimal-point character is used in place of a period,

and that if neither an exponent part nor a decimal-point character appears in a decimal form, or if a binary exponent part does not appear in a hexadecimal form, an exponent part of the appropriate type with value zero is assumed to follow the last digit in the string. If the subject sequence begins with a minus sign, the sequence is interpreted as arithmetically negated before rounding and the sign s is set to -1 , else s is set to 1 .

- 5 If the subject sequence has the decimal form, the value resulting from the conversion is correctly rounded and the coefficient c and the quantum exponent q are determined by the rules in 6.4.5.3 for a decimal floating literal of decimal type.
- 6 If the subject sequence has the hexadecimal form, the value resulting from the conversion is correctly rounded provided the subject sequence has at most M significant hexadecimal digits, where $M \geq \lceil (P - 1)/4 \rceil + 1$ is implementation-defined, and P is the maximum precision of the supported radix-2 floating types and binary non-arithmetic interchange formats.³⁵⁰⁾ If the value resulting from the conversion is correctly rounded for all subject sequences of hexadecimal form, M may be regarded as infinite. If the subject sequence has more than M significant hexadecimal digits, the implementation may first round to M significant hexadecimal digits according to the applicable decimal rounding direction mode, signaling exceptions as though converting from a wider format, then correctly round the result of the shortened hexadecimal input to the result type. The preferred quantum exponent for the result is 0 if the hexadecimal number is exactly represented in the decimal type; the preferred quantum exponent for the result is the least possible if the hexadecimal number is not exactly represented in the decimal type.
- 7 A character sequence **INF** or **INFINITY** is interpreted as an infinity. A character sequence **NAN** or **NAN(*d-char-sequence_{opt}*)**, is interpreted as a quiet NaN; the meaning of the d-char sequence is implementation-defined.³⁵¹⁾ A pointer to the final string is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.
- 8 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 9 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of **nptr** is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.

Returns

- 10 The **strtodN** functions return the converted value, if any. If no conversion could be performed, the value of the triple $(+1, 0, 0)$ is returned. If the correct value overflows:
 - the value of the macro **ERANGE** is stored in **errno** if the integer expression **math_errhandling** & **MATH_ERRNO** is nonzero;
 - the "overflow" floating-point exception is raised if the integer expression **math_errhandling** & **MATH_ERREXCEPT** is nonzero.

If the result underflows (7.12.2), whether **errno** acquires the value **ERANGE** if the integer expression **math_errhandling** & **MATH_ERRNO** is nonzero is implementation-defined; if the integer expression **math_errhandling** & **MATH_ERREXCEPT** is nonzero, whether the "underflow" floating-point exception is raised is implementation-defined.

- 11 EXAMPLE Following are subject sequences of the decimal form and the resulting triples (s, c, q) produced by **strtod64**. For **_Decimal64**, the precision (maximum coefficient length) is 16 and the quantum exponent range is $-398 \leq q \leq 369$.

"0"	(+1, 0, 0)
"0.00"	(+1, 0, -2)
"123"	(+1, 123, 0)
"-123"	(-1, 123, 0)
"1.23E3"	(+1, 123, 1)
"1.23E+3"	(+1, 123, 1)

³⁵⁰⁾Non-arithmetic interchange formats are an optional feature in Annex H.

³⁵¹⁾An implementation can use the d-char sequence to determine extra information to be represented in the NaN's significand.

"12.3E+7"	(+1, 123, 6)
"12.0"	(+1, 120, -1)
"12.3"	(+1, 123, -1)
"0.00123"	(+1, 123, -5)
"-1.23E-12"	(-1, 123, -14)
"1234.5E-4"	(+1, 12345, -5)
"-0"	(-1, 0, 0)
"-0.00"	(-1, 0, -2)
"0E+7"	(+1, 0, 7)
"-0E-7"	(-1, 0, -7)
"12345678901234567890"	(+1, 1234567890123457, 4) or (+1, 1234567890123456, 4) depending on rounding mode
"1234E-400"	(+1, 12, -398) or (+1, 13, -398) depending on rounding mode
"1234E-402"	(+1, 0, -398) or (+1, 1, -398) depending on rounding mode
"1000."	(+1, 1000, 0)
".0001"	(+1, 1, -4)
"1000.e0"	(+1, 1000, 0)
".0001e0"	(+1, 1, -4)
"1000.0"	(+1, 10000, -1)
"0.0001"	(+1, 1, -4)
"1000.00"	(+1, 100000, -2)
"00.0001"	(+1, 1, -4)
"001000."	(+1, 1000, 0)
"001000.0"	(+1, 10000, -1)
"001000.00"	(+1, 100000, -2)
"00.00"	(+1, 0, -2)
"00."	(+1, 0, 0)
".00"	(+1, 0, -2)
"00.00e-5"	(+1, 0, -7)
"00.e-5"	(+1, 0, -5)
".00e-5"	(+1, 0, -7)
"0x1.8p+4"	(+1, 24, 0)
"infinite"	infinity, and a pointer to "infinite" is stored in the object pointed to by <code>endptr</code> , provided <code>endptr</code> is not a null pointer

7.24.2.8 The `strtol`, `strtoll`, `strtoul`, and `strtoull` functions

Synopsis

```

1  #include <stdlib.h>
    long int strtol(const char * restrict nptr, char ** restrict endptr, int base);
    long long int strtoll(const char * restrict nptr, char ** restrict endptr,
        int base);
    unsigned long int strtoul(const char * restrict nptr, char ** restrict endptr,
        int base);
    unsigned long long int strtoull(const char * restrict nptr,
        char ** restrict endptr, int base);

```

Description

- The `strtol`, `strtoll`, `strtoul`, and `strtoull` functions convert the initial portion of the string pointed to by `nptr` to `long int`, `long long int`, `unsigned long int`, and `unsigned long long int` representation, respectively. First, they decompose the input string into three parts: an initial, possibly empty, sequence of white-space characters, a subject sequence resembling an integer represented in some radix determined by the value of `base`, and a final string of one or more unrecognized characters, including the terminating null character of the input string. Then, they attempt to convert the subject sequence to an integer, and return the result.
- If the value of `base` is zero, the expected form of the subject sequence is that of an integer literal as described in 6.4.5.2, optionally preceded by a plus or minus sign, but not including an integer suffix or any optional digit separators. If the value of `base` is between 2 and 36 (inclusive), the expected form of the subject sequence is a sequence of letters and digits representing an integer with the radix

specified by **base**, optionally preceded by a plus or minus sign, but not including an integer suffix or any optional digit separators. The letters from **a** (or **A**) through **z** (or **Z**) are ascribed the values 10 through 35; only letters and digits whose ascribed values are less than that of **base** are permitted. If the value of **base** is 2, the characters **0b** or **0B** can optionally precede the sequence of letters and digits, following the sign if present. If the value of **base** is 8, the characters **0o** or **0O** can optionally precede the sequence of letters and digits, following the sign if present. If the value of **base** is 16, the characters **0x** or **0X** can optionally precede the sequence of letters and digits, following the sign if present.

- 4 The subject sequence is defined as the longest initial subsequence of the input string, starting with the first non-white-space character, that is of the expected form. The subject sequence contains no characters if the input string is empty or consists entirely of white-space characters, or if the first non-white-space character is other than a sign or a permissible letter or digit.
- 5 If the subject sequence has the expected form and the value of **base** is zero, the sequence of characters starting with the first digit is interpreted as an integer literal according to the rules of 6.4.5.2. If the subject sequence has the expected form and the value of **base** is between 2 and 36, it is used as the base for conversion, ascribing to each letter its value as previously given. If the subject sequence begins with a minus sign, the resulting value is the negative of the converted value; for functions whose return type is an unsigned integer type this action is performed in the return type. A pointer to the final string is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.
- 6 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 7 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of **nptr** is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.

Returns

- 8 The **strtoul**, **strtoll**, **strtoul**, and **strtoull** functions return the converted value, if any. If no conversion could be performed, zero is returned. If the correct value is outside the range of representable values, **LONG_MIN**, **LONG_MAX**, **LLONG_MIN**, **LLONG_MAX**, **ULONG_MAX**, or **ULLONG_MAX** is returned (according to the return type and sign of the value, if any), and the value of the macro **ERANGE** is stored in **errno**.

7.24.3 Pseudo-random sequence generation functions

7.24.3.1 The **rand** function

Synopsis

```
1  #include <stdlib.h>
   int rand(void);
```

Description

- 2 The **rand** function computes a sequence of pseudo-random integers in the range 0 to **RAND_MAX** inclusive.
- 3 The **rand** function is not required to avoid data races with other calls to pseudo-random sequence generation functions. The implementation shall behave as if no library function calls the **rand** function.
- 4 NOTE There are no guarantees as to the quality of the random sequence produced and some implementations are known to produce sequences with distressingly non-random low-order bits. Applications with particular requirements should use a generator that is known to be sufficient for their needs.

Returns

- 5 The **rand** function returns a pseudo-random integer.

Environmental limits

- 6 The value of the **RAND_MAX** macro shall be at least 32767.

7.24.3.2 The **srand** function

Synopsis

```
1  #include <stdlib.h>
    void srand(unsigned int seed);
```

Description

- 2 The **srand** function uses the argument as a seed for a new sequence of pseudo-random numbers to be returned by subsequent calls to **rand**. If **srand** is then called with the same seed value, the sequence of pseudo-random numbers shall be repeated. If **rand** is called before any calls to **srand** have been made, the same sequence shall be generated as when **srand** is first called with a seed value of 1.
- 3 The **srand** function is not required to avoid data races with other calls to pseudo-random sequence generation functions. The implementation shall behave as if no library function calls the **srand** function.

Returns

- 4 The **srand** function returns no value.
- 5 EXAMPLE The following functions define a portable implementation of **rand** and **srand**.

```
static unsigned long int next = 1;

int rand(void)    // RAND_MAX assumed to be 32767
{
    next = next * 1103515245 + 12345;
    return (unsigned int)(next/65536) % 32768;
}

void srand(unsigned int seed)
{
    next = seed;
}
```

7.24.4 Memory management functions

7.24.4.1 General

- 1 The order and contiguity of storage allocated by successive calls to the **aligned_alloc**, **calloc**, **malloc**, and **realloc** functions is unspecified. The pointer returned if the allocation succeeds is suitably aligned so that it can be assigned to a pointer to any type of object with a fundamental alignment requirement and size less than or equal to the size requested. It can then be used to access such an object or an array of such objects in the space allocated (until the space is explicitly deallocated). The lifetime of an allocated object extends from the allocation until the deallocation. Each such allocation shall yield a pointer to an object disjoint from any other object. The pointer returned points to the start (lowest byte address) of the allocated space. If the space cannot be allocated, a null pointer is returned. If the size of the space requested is zero, the behavior is implementation-defined: either a null pointer is returned to indicate an error, or the behavior is as if the size were some nonzero value, except that the returned pointer shall not be used to access an object.
- 2 For purposes of determining the existence of a data race, memory allocation functions behave as though they accessed only memory locations accessible through their arguments and not other static duration storage. These functions may, however, visibly modify the storage that they allocate or deallocate. Calls to these functions that allocate or deallocate a particular region of memory shall occur in a single total order, and each such deallocation call shall synchronize with the next allocation (if any) in this order.

7.24.4.2 The **aligned_alloc** function

Synopsis

```
1  #include <stdlib.h>
    void *aligned_alloc(size_t alignment, size_t size);
```

Description

- 2 The **aligned_alloc** function allocates space for an object whose alignment is specified by **alignment**,³⁵²⁾ whose size is specified by **size**, and whose representation is indeterminate. If the value of **alignment** is not a valid alignment supported by the implementation the function shall fail by returning a null pointer.

Returns

- 3 The **aligned_alloc** function returns either a null pointer or a pointer to the allocated space.

7.24.4.3 The **calloc** function

Synopsis

```
1  #include <stdlib.h>
    void *calloc(size_t nmemb, size_t size);
```

Description

- 2 The **calloc** function allocates space for an array of **nmemb** objects, each of whose size is **size**. The space is initialized to all bits zero.³⁵³⁾

Returns

- 3 The **calloc** function returns either a pointer to the allocated space or a null pointer if the space cannot be allocated or if the product **nmemb** * **size** would wraparound **size_t**.

7.24.4.4 The **free** function

Synopsis

```
1  #include <stdlib.h>
    void free(void *ptr);
```

Description

- 2 The **free** function causes the space pointed to by **ptr** to be deallocated, that is, made available for further allocation. If **ptr** is a null pointer, no action occurs. Otherwise, if the argument does not match a pointer earlier returned by a memory management function, or if the space has been deallocated by a call to **free** or **realloc**, the behavior is undefined.

Returns

- 3 The **free** function returns no value.

7.24.4.5 The **free_sized** function

Synopsis

```
1  #include <stdlib.h>
    void free_sized(void *ptr, size_t size);
```

Description

- 2 If **ptr** is a null pointer or the result obtained from a call to **malloc**, **realloc**, or **calloc**, where **size** size is equal to the requested allocation size, this function is equivalent to **free(ptr)**. Otherwise, the behavior is undefined. The result of an **aligned_alloc** call may not be passed to **free_sized**.
- 3 NOTE A conforming implementation can ignore **size** and call **free**.

³⁵²⁾The alignment requirements from 7.24.4 also apply even if the requested alignment is less strict.

³⁵³⁾This is not expected to be the same as the representation of floating-point zero or a null pointer constant.

Recommended practice

- 4 Implementations may provide extensions to query the usable size of an allocation, or to determine the usable size of the allocation that would result if a request for some other size were to succeed. Such implementations should allow passing the resulting usable size as the `size` parameter, and provide functionality equivalent to `free` in such cases.

Returns

- 5 The `free_sized` function returns no value.

7.24.4.6 The `free_aligned_sized` function**Synopsis**

```
1  #include <stdlib.h>
    void free_aligned_sized(void *ptr, size_t alignment, size_t size);
```

Description

- 2 If `ptr` is a null pointer or the result obtained from a call to `aligned_alloc`, where `alignment` is equal to the requested allocation alignment and `size` is equal to the requested allocation size, this function is equivalent to `free(ptr)`. Otherwise, the behavior is undefined. The result of an `malloc`, `calloc`, or `realloc` call may not be passed to `free_aligned_sized`.
- 3 NOTE A conforming implementation can ignore `alignment` and `size` and call `free`.

Recommended practice

- 4 Implementations may provide extensions to query the usable size of an allocation, or to determine the usable size of the allocation that would result if a request for some other size were to succeed. Such implementations should allow passing the resulting usable size as the `size` parameter, and provide functionality equivalent to `free` in such cases.

Returns

- 5 The `free_aligned_sized` function returns no value.

7.24.4.7 The `malloc` function**Synopsis**

```
1  #include <stdlib.h>
    void *malloc(size_t size);
```

Description

- 2 The `malloc` function allocates space for an object whose size is specified by `size` and whose representation is indeterminate.

Returns

- 3 The `malloc` function returns either a null pointer or a pointer to the allocated space.

7.24.4.8 The `realloc` function**Synopsis**

```
1  #include <stdlib.h>
    void *realloc(void *ptr, size_t size);
```

Description

- 2 The `realloc` function deallocates the old object pointed to by `ptr` and returns a pointer to a new object that has the size specified by `size`. The contents of the new object shall be the same as that of the old object prior to deallocation, up to the lesser of the new and old sizes. Any bytes in the new object beyond the size of the old object have unspecified values.
- 3 If `ptr` is a null pointer, the `realloc` function behaves like the `malloc` function for the specified size.

Otherwise, if `ptr` does not match a pointer earlier returned by a memory management function, or if the space has been deallocated by a call to the `free` or `realloc` function, or if the `size` is zero, the behavior is undefined. If memory for the new object is not allocated, the old object is not deallocated and its value is unchanged.

Returns

- 4 The `realloc` function returns a pointer to the new object (which can have the same value as a pointer to the old object), or a null pointer if the new object has not been allocated.

7.24.5 Communication with the environment

7.24.5.1 The `abort` function

Synopsis

```
1  #include <stdlib.h>
    [[noreturn]] void abort(void);
```

Description

- 2 The `abort` function causes abnormal program termination to occur, unless the signal `SIGABRT` is being caught and the signal handler does not return. Whether open streams with unwritten buffered data are flushed, open streams are closed, or temporary files are removed is implementation-defined. An implementation-defined form of the status *unsuccessful termination* is returned to the host environment by means of the function call `raise(SIGABRT)`.

Returns

- 3 The `abort` function does not return to its caller.

7.24.5.2 The `atexit` function

Synopsis

```
1  #include <stdlib.h>
    int atexit(void (*func)(void));
```

Description

- 2 The `atexit` function registers the function pointed to by `func`, to be called without arguments at normal program termination.³⁵⁴ It is unspecified whether a call to the `atexit` function that does not happen before the `exit` function is called will succeed.

Environmental limits

- 3 The implementation shall support the registration of at least 32 functions.

Returns

- 4 The `atexit` function returns zero if the registration succeeds, nonzero if it fails.

Forward references: the `at_quick_exit` function (7.24.5.3), the `exit` function (7.24.5.4).

7.24.5.3 The `at_quick_exit` function

Synopsis

```
1  #include <stdlib.h>
    int at_quick_exit(void (*func)(void));
```

Description

- 2 The `at_quick_exit` function registers the function pointed to by `func`, to be called without arguments should `quick_exit` be called.³⁵⁵ It is unspecified whether a call to the `at_quick_exit`

³⁵⁴The `atexit` function registrations are distinct from the `at_quick_exit` registrations, so applications potentially need to call both registration functions with the same argument.

³⁵⁵The `at_quick_exit` function registrations are distinct from the `atexit` registrations, so applications potentially need to call both registration functions with the same argument.

function that does not happen before the **quick_exit** function is called will succeed.

Environmental limits

- 3 The implementation shall support the registration of at least 32 functions.

Returns

- 4 The **at_quick_exit** function returns zero if the registration succeeds, nonzero if it fails.

Forward references: the **quick_exit** function (7.24.5.7).

7.24.5.4 The **exit** function

Synopsis

```
1      #include <stdlib.h>
      [[noreturn]] void exit(int status);
```

Description

- 2 The **exit** function causes normal program termination to occur. No functions registered by the **at_quick_exit** function are called. If a program calls the **exit** function more than once, or calls the **quick_exit** function in addition to the **exit** function, the behavior is undefined.
- 3 First, all functions registered by the **atexit** function are called, in the reverse order of their registration,³⁵⁶⁾ except that a function is called after any previously registered functions that had already been called at the time it was registered. If, during the call to any such function, a call to the **longjmp** function is made that would terminate the call to the registered function, the behavior is undefined.
- 4 Next, all open streams with unwritten buffered data are flushed, all open streams are closed, and all files created by the **tmpfile** function are removed.
- 5 Finally, control is returned to the host environment. If the value of **status** is zero or **EXIT_SUCCESS**, an implementation-defined form of the status *successful termination* is returned. If the value of **status** is **EXIT_FAILURE**, an implementation-defined form of the status *unsuccessful termination* is returned. Otherwise the status returned is implementation-defined.

Returns

- 6 The **exit** function cannot return to its caller.

7.24.5.5 The **_Exit** function

Synopsis

```
1      #include <stdlib.h>
      [[noreturn]] void _Exit(int status);
```

Description

- 2 The **_Exit** function causes normal program termination to occur and control to be returned to the host environment. No functions registered by the **atexit** function, the **at_quick_exit** function, or signal handlers registered by the **signal** function are called. The status returned to the host environment is determined in the same way as for the **exit** function (7.24.5.4). Whether open streams with unwritten buffered data are flushed, open streams are closed, or temporary files are removed is implementation-defined.

Returns

- 3 The **_Exit** function cannot return to its caller.

7.24.5.6 The **getenv** function

Synopsis

-
- 1 ³⁵⁶⁾Each function is called as many times as it was registered, and in the correct order with respect to other registered functions.

```
#include <stdlib.h>
char *getenv(const char *name);
```

Description

- 2 The **getenv** function searches an *environment list*, provided by the host environment, for a string that matches the string pointed to by **name**. The set of environment names and the method for altering the environment list are implementation-defined. The **getenv** function is not required to avoid data races with other threads of execution that modify the environment list.³⁵⁷⁾
- 3 The implementation shall behave as if no library function calls the **getenv** function.

Returns

- 4 The **getenv** function returns a pointer to a string associated with the matched list member. The string pointed to shall not be modified by the program, but can be overwritten by a subsequent call to the **getenv** function. If the specified **name** cannot be found, a null pointer is returned.

7.24.5.7 The **quick_exit** function

Synopsis

```
1 #include <stdlib.h>
   [[noreturn]] void quick_exit(int status);
```

Description

- 2 The **quick_exit** function causes normal program termination to occur. No functions registered by the **atexit** function or signal handlers registered by the **signal** function are called. If a program calls the **quick_exit** function more than once, or calls the **exit** function in addition to the **quick_exit** function, the behavior is undefined. If a signal is raised while the **quick_exit** function is executing, the behavior is undefined.
- 3 The **quick_exit** function first calls all functions registered by the **at_quick_exit** function, in the reverse order of their registration,³⁵⁸⁾ except that a function is called after any previously registered functions that had already been called at the time it was registered. If, during the call to any such function, a call to the **longjmp** function is made that would terminate the call to the registered function, the behavior is undefined.
- 4 Then control is returned to the host environment by means of the function call **_Exit(status)**.

Returns

- 5 The **quick_exit** function cannot return to its caller.

7.24.5.8 The **system** function

Synopsis

```
1 #include <stdlib.h>
   int system(const char *string);
```

Description

- 2 If **string** is a null pointer, the **system** function determines whether the host environment has a *command processor*. If **string** is not a null pointer, the **system** function passes the string pointed to by **string** to that command processor to be executed in a manner which the implementation shall document; this can then cause the program calling **system** to behave in a non-conforming manner or to terminate.

³⁵⁷⁾Many implementations provide non-standard functions that modify the environment list.

³⁵⁸⁾Each function is called as many times as it was registered, and in the correct order with respect to other registered functions.

Returns

- 3 If the argument is a null pointer, the **system** function returns nonzero only if a command processor is available. If the argument is not a null pointer, and the **system** function does return, it returns an implementation-defined value.

7.24.6 Searching and sorting utilities**7.24.6.1 General**

- 1 These utilities make use of a comparison function to search or sort arrays of unspecified type. Where an argument declared as **size_t nmemb** specifies the length of the array for a function, **nmemb** can have the value zero on a call to that function; the comparison function is not called, a search finds no matching element, and sorting performs no rearrangement. Pointer arguments on such a call may be null pointers.
- 2 The implementation shall ensure that the second argument of the comparison function (when called from **bsearch**), or both arguments (when called from **qsort**), are pointers to elements of the array.³⁵⁹⁾ The first argument when called from **bsearch** shall equal **key**.
- 3 The comparison function shall not alter the contents of the array. The implementation may reorder elements of the array between calls to the comparison function, but shall not alter the contents of any individual element.
- 4 When the same objects (consisting of **size** bytes, irrespective of their current positions in the array) are passed more than once to the comparison function, the results shall be consistent with one another. That is, for **qsort** they shall define a total ordering on the array, and for **bsearch** the same object shall always compare the same way with the key.
- 5 A sequence point occurs immediately before and immediately after each call to the comparison function, and also between any call to the comparison function and any movement of the objects passed as arguments to that call.

7.24.6.2 The **bsearch generic function****Synopsis**

```
1  #include <stdlib.h>
    QVoid *bsearch(const void *key, QVoid *base, size_t nmemb, size_t size,
                  int (*compar)(const void *, const void *));
```

Description

- 2 The **bsearch** generic function searches an array of **nmemb** objects, the initial element of which is pointed to by **base**, for an element that matches the object pointed to by **key**. The size of each element of the array is specified by **size**.
- 3 The comparison function pointed to by **compar** is called with two arguments that point to the **key** object and to an array element, in that order. The function shall return an integer less than, equal to, or greater than zero if the **key** object is considered, respectively, to be less than, to match, or to be greater than the array element. The array shall consist of: all the elements that compare less than, all the elements that compare equal to, and all the elements that compare greater than the **key** object, in that order.³⁶⁰⁾

Returns

- 4 The **bsearch** generic function returns a pointer to a matching element of the array, or a null pointer if no match is found. If two elements compare as equal, which element is matched is unspecified.

³⁵⁹⁾That is, if the value passed is **p**, then the following expressions are always nonzero:

```
((char *)p - (char *)base) % size == 0
(char *)p >= (char *)base
(char *)p < (char *)base + nmemb * size
```

³⁶⁰⁾In practice, the entire array is sorted according to the comparison function.

- 5 The **bsearch** function is generic in the qualification of the type pointed to by the argument **base**. If this argument is a pointer to a **const**-qualified object type, the returned pointer will be a pointer to **const**-qualified **void**. Otherwise, the argument shall be a pointer to an unqualified object type or a null pointer constant, and the returned pointer will be a pointer to unqualified **void**.

The external declaration of **bsearch** has the concrete type:

```
void * (const void *, const void *, size_t, size_t,
        int (*)(const void *, const void *))
```

which supports all correct uses. If a macro definition of this generic function is suppressed to access an actual function, the external declaration with this concrete type is visible.³⁶¹⁾

7.24.6.3 The **qsort** function

Synopsis

```
1  #include <stdlib.h>
    void qsort(void *base, size_t nmemb, size_t size,
               int (*compar)(const void *, const void *));
```

Description

- 2 The **qsort** function sorts an array of **nmemb** objects, the initial element of which is pointed to by **base**. The size of each object is specified by **size**.
- 3 The contents of the array are sorted into ascending order according to a comparison function pointed to by **compar**, which is called with two arguments that point to the objects being compared. The function shall return an integer less than, equal to, or greater than zero if the first argument is considered to be respectively less than, equal to, or greater than the second.
- 4 If two elements compare as equal, their order in the resulting sorted array is unspecified.

Returns

- 5 The **qsort** function returns no value.

7.24.7 Integer arithmetic functions

7.24.7.1 The **abs**, **uabs**, **labs**, **ulabs**, **llabs**, and **ullabs** functions

Synopsis

```
1  #include <stdlib.h>
    int abs(int j);
    unsigned int uabs(int j);
    long int labs(long int j);
    unsigned long int ulabs(long int j);
    long long int llabs(long long int j);
    unsigned long long int ullabs(long long int j);
```

Description

- 2 The **abs**, **uabs**, **labs**, **ulabs**, **llabs**, and **ullabs** functions compute the absolute value of an integer **j**. If the result cannot be represented, the behavior is undefined.³⁶²⁾

Returns

- 3 The **abs**, **uabs**, **labs**, **ulabs**, **llabs**, and **ullabs** functions return the absolute value.

7.24.7.2 The **div**, **ldiv**, and **lldiv** functions

Synopsis

```
1  #include <stdlib.h>
```

³⁶¹⁾This is an obsolescent feature.

³⁶²⁾Only the absolute value of the most negative number is not representable for **abs**, **ulabs**, **ullabs**.

```
div_t div(int numer, int denom);
ldiv_t ldiv(long int numer, long int denom);
lldiv_t lldiv(long long int numer, long long int denom);
```

Description

- The **div**, **ldiv**, and **lldiv** functions compute `numer/denom` and `numer%denom` in a single operation.

Returns

- The **div**, **ldiv**, and **lldiv** functions return a structure of type **div_t**, **ldiv_t**, and **lldiv_t**, respectively, comprising both the quotient and the remainder. The structures shall contain (in either order) the members **quot** (the quotient) and **rem** (the remainder), each of which has the same type as the arguments **numer** and **denom**. If either part of the result cannot be represented, the behavior is undefined.

7.24.8 Multibyte/wide character conversion functions

7.24.8.1 General

- The behavior of the multibyte character functions is affected by the **LC_CTYPE** category of the current locale. For a state-dependent encoding, each of the **mbtowc** and **wctomb** functions is placed into its initial conversion state prior to the first call to the function and can be returned to that state by a call for which its character pointer argument, **s**, is a null pointer. Subsequent calls with **s** as other than a null pointer cause the internal conversion state of the function to be altered as necessary. It is implementation-defined whether internal conversion state has thread storage duration, and whether a newly created thread has the same state as the current thread at the time of creation, or the initial conversion state. A call with **s** as a null pointer causes these functions to return a nonzero value if encodings have state dependency, and zero otherwise.³⁶³ It is implementation-defined whether these functions avoid data races with other calls to the same function.
- Changing the **LC_CTYPE** category causes the internal object describing the conversion state of the **mbtowc** and **wctomb** functions to have an indeterminate representation.

7.24.8.2 The **mblen** function

Synopsis

- ```
#include <stdlib.h>
int mblen(const char *s, size_t n);
```

### Description

- If **s** is not a null pointer, the **mblen** function determines the number of bytes contained in the multibyte character pointed to by **s**. Except that the conversion state of the **mbtowc** function is not affected, it is equivalent to

```
mbtowc((wchar_t *)0, (const char *)0, 0);
mbtowc((wchar_t *)0, s, n);
```

### Returns

- If **s** is a null pointer, the **mblen** function returns a nonzero or zero value, if multibyte character encodings, respectively, do or do not have state-dependent encodings. If **s** is not a null pointer, the **mblen** function either returns 0 (if **s** points to the null character), or returns the number of bytes that are contained in the multibyte character (if the next **n** or fewer bytes form a valid multibyte character), or returns -1 (if they do not form a valid multibyte character).

**Forward references:** the **mbtowc** function (7.24.8.3).

### 7.24.8.3 The **mbtowc** function

<sup>363</sup>If the locale employs special bytes to change the shift state, these bytes do not produce separate wide character codes, but are grouped with an adjacent multibyte character.

## Synopsis

```
1 #include <stdlib.h>
 int mbtowc(wchar_t * restrict pwc, const char * restrict s, size_t n);
```

## Description

- 2 If *s* is not a null pointer, the **mbtowc** function inspects at most *n* bytes beginning with the byte pointed to by *s* to determine the number of bytes needed to complete the next multibyte character (including any shift sequences). If the function determines that the next multibyte character is complete and valid, it determines the value of the corresponding wide character and then, if *pwc* is not a null pointer, stores that value in the object pointed to by *pwc*. If the corresponding wide character is the null wide character, the function is left in the initial conversion state.
- 3 The implementation shall behave as if no library function calls the **mbtowc** function.

## Returns

- 4 If *s* is a null pointer, the **mbtowc** function returns a nonzero or zero value, if multibyte character encodings, respectively, do or do not have state-dependent encodings. If *s* is not a null pointer, the **mbtowc** function either returns 0 (if *s* points to the null character), or returns the number of bytes that are contained in the converted multibyte character (if the next *n* or fewer bytes form a valid multibyte character), or returns -1 (if they do not form a valid multibyte character).
- 5 In no case will the value returned be greater than *n* or the value of the **MB\_CUR\_MAX** macro.

### 7.24.8.4 The **wctomb** function

## Synopsis

```
1 #include <stdlib.h>
 int wctomb(char *s, wchar_t wc);
```

## Description

- 2 The **wctomb** function determines the number of bytes needed to represent the multibyte character corresponding to the wide character given by *wc* (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by *s* (if *s* is not a null pointer). At most **MB\_CUR\_MAX** characters are stored. If *wc* is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state, and the function is left in the initial conversion state.
- 3 The implementation shall behave as if no library function calls the **wctomb** function.

## Returns

- 4 If *s* is a null pointer, the **wctomb** function returns a nonzero or zero value, if multibyte character encodings, respectively, do or do not have state-dependent encodings. If *s* is not a null pointer, the **wctomb** function returns -1 if the value of *wc* does not correspond to a valid multibyte character, or returns the number of bytes that are contained in the multibyte character corresponding to the value of *wc*.
- 5 In no case will the value returned be greater than the value of the **MB\_CUR\_MAX** macro.

## 7.24.9 Multibyte/wide string conversion functions

### 7.24.9.1 General

- 1 The behavior of the multibyte string functions is affected by the **LC\_CTYPE** category of the current locale.

### 7.24.9.2 The **mbstowcs** function

## Synopsis

```
1 #include <stdlib.h>
 size_t mbstowcs(wchar_t * restrict pwcs, const char * restrict s, size_t n);
```

**Description**

- 2 The **mbstowcs** function converts a sequence of multibyte characters that begins in the initial shift state from the array pointed to by **s** into a sequence of corresponding wide characters and stores not more than **n** wide characters into the array pointed to by **pwcs**. No multibyte characters that follow a null character (which is converted into a null wide character) will be examined or converted. Each multibyte character is converted as if by a call to the **mbtowc** function, except that the conversion state of the **mbtowc** function is not affected.
- 3 No more than **n** elements will be modified in the array pointed to by **pwcs**. If copying takes place between objects that overlap, the behavior is undefined.

**Returns**

- 4 If an invalid multibyte character is encountered, the **mbstowcs** function returns **(size\_t) (-1)**. Otherwise, the **mbstowcs** function returns the number of array elements modified, not including a terminating null wide character, if any.<sup>364)</sup>

**7.24.9.3 The wcstombs function****Synopsis**

```
1 #include <stdlib.h>
 size_t wcstombs(char * restrict s, const wchar_t * restrict pwcs, size_t n);
```

**Description**

- 2 The **wcstombs** function converts a sequence of wide characters from the array pointed to by **pwcs** into a sequence of corresponding multibyte characters that begins in the initial shift state, and stores these multibyte characters into the array pointed to by **s**, stopping if a multibyte character would exceed the limit of **n** total bytes or if a null character is stored. Each wide character is converted as if by a call to the **wctomb** function, except that the conversion state of the **wctomb** function is not affected.
- 3 No more than **n** bytes will be modified in the array pointed to by **s**. If copying takes place between objects that overlap, the behavior is undefined.

**Returns**

- 4 If a wide character is encountered that does not correspond to a valid multibyte character, the **wcstombs** function returns **(size\_t) (-1)**. Otherwise, the **wcstombs** function returns the number of bytes modified, not including a terminating null character, if any.<sup>364)</sup>

**7.24.10 Alignment of memory****7.24.10.1 The memalignment function****Synopsis**

```
1 #include <stdlib.h>
 size_t memalignment(const void *p);
```

**Description**

- 2 The **memalignment** function accepts a pointer to any object and returns the maximum alignment satisfied by its address value. The alignment may be an extended alignment and may also be beyond the range supported by the implementation for explicit use by **alignas**.<sup>365)</sup> If so, it will satisfy all alignments usable by the implementation. The value returned can be compared to the result of **alignof**, and if it is greater or equal, the alignment requirement for the type operand is satisfied.

<sup>364)</sup>The array will not be null-terminated if the value returned is **n**.

<sup>365)</sup>The actual alignment of an object can be stricter than the alignment requested for an object by **alignas** or (implicitly) by an allocation function, but will always satisfy it.

**Returns**

- 3 The alignment of the pointer `p`, which is a power of two. If `p` is a null pointer, an alignment of zero is returned.
- 4 NOTE An alignment of zero indicates that the tested pointer cannot be used to access an object of any type.



## 7.25 Text transcoding utilities <stdmchar.h>

### 7.25.1 General

- 1 The header <stdmchar.h> declares four status code enumerators, several macros, several types and several functions for transcoding encoded text safely and efficiently. It is meant to supersede conversion utilities from Unicode utilities (7.31) and Extended multibyte and wide character utilities (7.32). It is meant to represent “multi character” functions. These functions can be used to count the number of input that form a complete sequence, count the number of output characters required for a conversion with no additional allocation, validate an input sequence, or transcode text from one encoding to another encoding. Particularly, it provides single unit and multi unit transcoding functions for transcoding by working on code units and code points.
- 2 Inputs to the functions in this subclause are read until there is enough information taken in to perform an *indivisible unit of work*. An indivisible unit is the smallest possible input, as defined by the encoding, that can produce one or more outputs, perform a transformation of any state, or both. The conversion of these indivisible units is called an indivisible unit of work, and they are used to complete the transcoding operations specified in this subclause.
- 3 One or more of the following must hold for any given transcoding operation on an attempt to complete an indivisible unit of work:
  - enough input is consumed to perform an output or change the state;
  - output is written from consuming input, or output is written from the state which causes the state to change; or,
  - an error occurs and both the input and output do not change relative to the current indivisible unit of work.
- 4 For the multi unit functions, the process acts as if it completes one indivisible unit of work repeatedly. When an error occurs, only the input successfully consumed, the state successfully altered, and the output successfully written according to the last indivisible unit of work are reflected in the output values of the functions in this subclause: no other values are written.
- 5 Functions in <stdmchar.h> which use **char** and **wchar\_t**, or their qualified forms, derive their implementation-defined encodings from the narrow execution encoding or the wide execution encoding (6.2.9), respectively. The other encodings are UTF-8, associated with **char8\_t**, UTF-16, associated with **char16\_t**, and UTF-32, associated with **char32\_t**.
- 6 NOTE 1 Each value is treated as code units with an unsigned value and not as a container of octets. This means that the decision of, for example, UTF-16 in big or little endian encoding scheme is decided by the endianness of the code unit type. Only whole code unit values are used (i.e. a UTF-32 code point value of U+1F377 represents a value identical to how `U'\U0001F377'` is stored by the implementation).
- 7 For the UTF-8, UTF-16, and UTF-32 encodings, collectively referred to as the *Unicode encodings*, an indivisible unit of work for a read operation shall be the sequence of code units that corresponds to one Unicode code point. The value of each code unit is treated as a sequence of unsigned values. If input is exhausted before a sequence of code units corresponding to one Unicode code point can be reached, then **stdc\_mcerr\_incomplete\_input** shall be returned. If there is an illegal code unit sequence, then **stdc\_mcerr\_invalid** shall be returned. For the implementation-defined execution and wide execution encodings, they have the same aforementioned requirement if the implementation defines it to be one of the Unicode encodings.
- 8 NOTE 2 If an implementation chooses to provide, for example, an execution encoding as the input encoding for a transcoding function that is defined to be the same as the UTF-8 encoding, then it is required to read one full complete Unicode code point's worth of code units. If it cannot, then it returns **stdc\_mcerr\_incomplete\_input** (if the input sequence is not long enough but does not have any invalid code units in the sequence) or **stdc\_mcerr\_invalid** (if the input sequence is not a proper code unit sequence).
- 9 NOTE 3 The requirements for Unicode encodings do not apply to derivative encodings defined by the implementation. For example, an implementation may define a “partial UTF-8” execution encoding where it stores every read UTF-8 code unit in the state and, rather than returning **stdc\_mcerr\_incomplete\_input**, returns **stdc\_mcerr\_ok** and produces no output. It may accumulate code units and write out a code point

when it accumulates enough code units in its internal state. However, such an encoding is distinct and separate from the UTF-8 encoding used in the `c8` prefixed and suffixed functions described in this subclause.

- 10 NOTE 4 The implementation-defined execution, wide execution, literal, and wide literal encodings can also have different behaviors if they do not define themselves as one of the Unicode encodings. For example, if `__STDC_ENDIAN_NATIVE__` (7.18.2) is equivalent to `__STDC_ENDIAN_LITTLE__`, but the wide execution encoding is defined to be “UTF-16 Big Endian” (“UTF16-BE”), then it may be classified as not one of the three recognized Unicode encodings according to this subclause. As such, a sequence of `wchar_t` elements that is null-terminated produced by transcoding functions in this subclause can behave differently than expected; e.g. `L"U0001F377"`, if valid and defined to be UTF-16, can potentially not compare equal to a sequence of `wchar_t` objects produced by successful use of the transcoding functions from the code points `U+01F377` and `U+000000`.
- 11 For all functions in this subclause, when a code unit value of 0 (e.g. `'\0'`) is encountered in the input, the `mbstate_t` object in use for the transcoding operation is set to the initial shift state. The output associated with the indivisible unit of work consists of the appropriate null character preceded by any shift sequence necessary to cause the output to be in the initial shift state.
- 12 NOTE 5 As described in 7.32.6.1, an object of type `mbstate_t` can always be set to the initial conversion sequence by initializing it with `= {0}`; or `= {}`. An existing `mbstate_t` object can always be set to the initial conversion sequence by assigning to it from the expressions `(mbstate_t){0}` or `(mbstate_t){}`.
- 13 Changing the `LC_CTYPE` category causes any conversion state already in use with the functions in this subclause to be indeterminate.
- 14 The macro

```
__STDC_VERSION_STDMCHAR_H__
```

is an integer constant expression with a value equivalent to 202ymmL.

- 15 The types declared are `mbstate_t` (described in 7.32.1), `wchar_t` (described in 7.21), `char8_t` (described in 7.31), `char16_t` (described in 7.31), `char32_t` (described in 7.31), `size_t` (described in 7.21), and;

```
stdc_mcerr
```

which is both an enumerated type and a typedef whose enumerators identify the status codes from a function call described in this subclause.

- 16 The macros declared are `WCHAR_MIN`, `WCHAR_MAX`, and `WCHAR_WIDTH` (described in 7.22); `WCHAR_UTF8`, `WCHAR_UTF16`, `WCHAR_UTF32`, `MB_UTF8`, `MB_UTF16`, and `MB_UTF32` (described in 7.32.1); and,

```
STDC_C8_MAX
STDC_C16_MAX
STDC_C32_MAX
STDC_MC_MAX
STDC_MWC_MAX
```

which correspond to the maximum output for each single unit conversion function (7.25.2) and its corresponding output type. Each macro shall expand into an integer constant expression with minimum values, as described in Table 7.8.

- 17 There is an association of naming convention, types, encoding, and maximums, used to describe the functions in this subclause:

The maximum output macro values specified in Table 7.8 are related to the single unit conversion functions (7.25.2). These functions perform at most one indivisible unit of work, or return an error. The maximum output macro values shall be integer constant expressions large enough that conversions to the single unit conversion function’s specified encoding shall not overflow a buffer of the proper code unit type with that size. The maximum output macro values do not affect the multi unit conversion functions (7.25.3), which perform as many indivisible units of work as is possible until an error occurs, until the output space is exhausted, or until the input is exhausted.

- 18 Unlike the functions present in `<uchar.h>` and `<wchar.h>`, the functions present in this subclause

Table 7.8 — Transcoding function associations

| Prefix/Suffix    | Code Unit Type        | Encoding                                                           | Maximum Macro             | Minimum Value |
|------------------|-----------------------|--------------------------------------------------------------------|---------------------------|---------------|
| <code>mc</code>  | <code>char</code>     | The narrow execution encoding, influenced by <code>LC_CTYPE</code> | <code>STDC_MC_MAX</code>  | 1             |
| <code>mwc</code> | <code>wchar_t</code>  | The wide execution encoding, influenced by <code>LC_CTYPE</code>   | <code>STDC_MWC_MAX</code> | 1             |
| <code>c8</code>  | <code>char8_t</code>  | UTF-8                                                              | <code>STDC_C8_MAX</code>  | 4             |
| <code>c16</code> | <code>char16_t</code> | UTF-16                                                             | <code>STDC_C16_MAX</code> | 2             |
| <code>c32</code> | <code>char32_t</code> | UTF-32                                                             | <code>STDC_C32_MAX</code> | 1             |

can write more than one `wchar_t` value for conversions based on the wide execution encoding to accommodate a wider set of implementation-defined encodings, so long as the number of code units does not exceed the maximum output macro value of `STDC_MWC_MAX`.

- 19 The enumerators of the enumerated type `stdc_mcerr` are defined as follows:

```
stdc_mcerr_ok = 0
stdc_mcerr_invalid = -1,
stdc_mcerr_incomplete_input = -2,
stdc_mcerr_insufficient_output = -3
```

Each value represents a specific situation when calling the relevant transcoding functions in `<stdmchar.h>`:

- `stdc_mcerr_insufficient_output`, when the input is correct and an indivisible unit of work can be performed but there is not enough output space to write to;
- `stdc_mcerr_incomplete_input`, when input has been exhausted and the sequence is not incorrect but there are no more input values;
- `stdc_mcerr_invalid`, when a unrecorded encoding error occurred; and,
- `stdc_mcerr_ok`, when the operation was successful (none of the situations described for the other values of this enumerated type apply).

No other value shall be returned from the functions described in this subclause.

#### Recommended practice

- 20 The maximum output macro values are intended for use in making automatic storage duration array declarations. Implementations should choose values for the macros that are spacious enough to accommodate a variety of underlying implementation choices for the target encodings supported by the narrow execution encodings and wide execution encodings, which for some encodings can output more than one UTF-32 code point. A set of values which are most resilient to future additions and changes in implementations is as follows:

```
#define STDC_C8_MAX 64
#define STDC_C16_MAX 32
#define STDC_C32_MAX 16
#define STDC_MC_MAX 64
#define STDC_MWC_MAX 32
```

- 21 Beyond just the Unicode encodings specified previously, implementations are encouraged to not store partial reads or partial writes in the `mbstate_t` object with these functions unless as is strictly

necessary. Implementations providing additional encodings for use with these functions should, to the extent possible for a given encoding, always define an indivisible unit of work to transcode as complete a unit of information as is possible or produce an error. If a sequence of code units cannot form a complete shift sequence or produce output, then an implementation should return `stdc_mcerr_incomplete_input` if the input is exhausted, or `stdc_mcerr_invalid` if the input sequence is incorrect.

## 7.25.2 Single Unit Sized Conversion Functions

### Synopsis

1

```
#include <stdmchar.h>

stdc_mcerr stdc_mcnrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_mwcnrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c8nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
```

```

 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c16nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c32nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);

```

### Description

- Let transcoding function be one of the functions listed previously transcribed in the form

```

stdc_mcerr stdc_XnrtoYn(size_t*restrict output_size,
 charY*restrict *restrict output,

```

```
size_t*restrict input_size, const charX*restrict *restrict input,
mbstate_t*restrict state)
```

with the following properties:

- $X$  and  $Y$  be one of the prefixes/suffixes from Table 7.8;
- `charX` and `charY` be the associated code unit types for  $X$  and  $Y$  from Table 7.8;
- and, *encodingX* and *encodingY* be the associated encoding types for  $X$  and  $Y$  from Table 7.8.

The transcoding functions take an input buffer and possibly an output buffer of the associated code unit types, potentially with their sizes. The function consumes any number of code units of type `charX` to perform a single indivisible unit of work necessary to convert some amount of input from *encodingX* to *encodingY*, which results in zero or more output code units of type `charY`.

- 3 An `mbstate_t` object describes the conversion state for the current conversion if it is not in an unspecified state (as described further later in this subclause) and:

- the conversion is between Unicode encodings;
- or, the input fragment is the start of an encoded sequence in the input encoding and the `mbstate_t` object was initialized to the initial conversion state;
- or, the input fragment is a continuation of an encoded sequence and `mbstate_t` is the result of having advanced to the input and possibly output positions through the application of prior calls to the same transcoding function.

The behavior is undefined when a function described by this subclause is invoked and `*state` (or, if `state` is a null pointer, the `mbstate_t` object created that is unique to the current invocation) does not describe the conversion state for the current conversion.

- 4 The transcoding functions convert from code units of type `charX` interpreted according to *encodingX* to code units of type `charY` according to *encodingY* given a conversion state of value `*state`. This function only performs a single indivisible unit of work. It returns `stdc_mcerr_ok` if the input is empty. The input is considered empty if `input_size` is a null pointer or if `*input_size` is zero if `input_size` is not a null pointer, and `*state` is in the initial conversion sequence.
- 5 Any time *input code unit reads* in the following description is used:

- code units are read from `*input`, sequentially, and interpreted according to *encodingX*;
- if `*input_size` is smaller than the necessary amount of sequential reads that must be performed from input to complete an indivisible unit of work, the function does not modify any of `*input`, `*input_size`, `*output`, or `*output_size`. It returns `stdc_mcerr_incomplete_sequence`;
- if an unrecorded encoding error (6.2.9) occurs (e.g. the input read is invalid according to *encodingX* or the input is valid but cannot be converted to *encodingY*), then the function returns `stdc_mcerr_invalid`;
- if the function returns `stdc_mcerr_ok`, the function decrements `*input_size` by the number of input code units that were read and increments `*input` by the number of input code units that were read for the complete indivisible unit of work.

- 6 Any time *output code unit writes* in the following description is used:

- converted code units are potentially written into `*output`, sequentially, according to *encodingY*;



- if `output_size` is not a null pointer and if `*output_size` is smaller than the necessary amount of sequential writes that must be performed to complete an indivisible unit of work, the function does not modify any of `*input`, `*input_size`, `*output`, or `*output_size`. It returns `stdc_mcerr_insufficient_output`;
- if `output_size` is a null pointer, but `output` and `*output` are not null pointers, it is assumed `*output` has enough space to perform the necessary sequential writes and the behavior is undefined if the target output buffer is not large enough for this transcoding operation's indivisible unit of work;
- if the function returns `stdc_mcerr_ok`, the function decrements `*output_size` by the number of output code units that are, or could have been (if `output` or `*output` are null pointers), written. If `output` and `*output` are not null pointers, then `*output` is incremented by the number of output code units that are written to complete an indivisible unit of work.

7 The behavior of the transcoding functions is as follows:

- If `state` is a null pointer, then an automatic storage duration object of type `mbstate_t` is created which is unique to the current invocation. It is initialized to the initial conversion state and a pointer to this object is used wherever `state` is used in this paragraph.
- Then, if `input` is a null pointer or `*input` is a null pointer, then `*state` is set to the initial conversion state, the function returns `stdc_mcerr_ok`, and no other actions are taken.
- Otherwise, if `*state` is in an implementation-defined conversion state that requires it, any necessary output code units writes are performed to return `*state` to the initial conversion state.
- The function performs input code unit reads and subsequently performs the output code unit writes as is necessary to complete an indivisible unit of work.
- The function returns `stdc_mcerr_ok`.

8 NOTE If `state` is a null pointer, and the function uses e.g. a created automatic storage duration `mbstate_t` object that is discarded by the end of the invocation, then any potential conversion state contained in the created `mbstate_t` object and used during processing could become unrecoverable to the program.

### Returns

- 9 On success or failure, the transcoding functions shall return one of the above error codes (7.25.1). If `input` is a null pointer or `*input` is a null pointer, then `*state` is set to the initial conversion state and no other work is performed.
- 10 If the function returns `stdc_mcerr_ok`, then all of the following is true:
- if `input` and `*input` are not null pointers, `*input` is incremented by the number of code units read and successfully converted;
  - if `input_size` is not a null pointer, `*input_size` is decremented by the number of code units read and successfully converted from the input;
  - if `output` and `*output` are not null pointers, `*output` is incremented by the number of code units written to the output; and,
  - if `output_size` is not a null pointer, `*output_size` is decremented by the number of code units written to the output.

Otherwise, if an error is returned then none of the above occurs. If the return value is `stdc_mcerr_invalid`, then `*state` is in an unspecified state. If the return value is `stdc_mcerr_incomplete_input` or `stdc_mcerr_insufficient_output`, then `*state` is not changed.

**Recommended practice**

- 11 Implementations should take advantage of the information of null pointer values for the output size pointer, output data pointer, or both, to drastically improve performance characteristics for assumed unlimited write space, output counting scenarios, or input validation/counting, respectively.
- 12 Implementations should prefer returning an error for an incomplete input sequence over storing intermediate data within the state where possible for non-Unicode encodings. This can make it easier for functionality built on top of the functions in this subclause to report errors without skipping over potentially invalid input data, resulting in potentially more accurate reports. Error handling and recovery also greatly benefit from being able to examine invalid input; avoiding skipping over invalid data by consuming it into a state and reporting no errors means that functionality built on top can potentially discard what should be considered unneeded, already-processed data.

**7.25.3 Multi Unit Sized Conversion Functions****Synopsis**

1

```
#include <stdmchar.h>

stdc_mcerr stdc_mcsnrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrto8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrto16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrto32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_mwcnrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtomwcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrto8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrto16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrto32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c8snrtomwcn(size_t*restrict output_size,
```



```

 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c16snrtomwcsn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c32snrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtomwcsn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);

```

**Description**

- 2 Let *multi unit transcoding function* in this function be one of the functions listed above transcribed in the form

```

stdc_mcerr stdc_XsnrtoYsn(size_t*restrict output_size,
 charY*restrict *restrict output,
 size_t*restrict input_size, const charX*restrict *restrict input,
 mbstate_t*restrict state);

```

with the following properties:

- *X* and *Y* be one of the prefixes/suffixes from Table 7.8;
- **charX** and **charY** be the associated code unit types for *X* and *Y* from Table 7.8;
- and, *encodingX* and *encodingY* be the associated encoding types for *X* and *Y* from Table 7.8.

The multi unit transcoding functions take an input buffer and possibly an output buffer of the associated code unit types, potentially with their sizes. The functions consume any number of code units to perform a sequence of indivisible units of work, which results in zero or more output code units. The functions repeatedly perform an indivisible unit of work until either an error occurs or the input is exhausted.

- 3 An **mbstate\_t** object describes the conversion state for the current conversion if it is not in an unspecified state (as described later in this subclause) and:
- the conversion is between Unicode encodings;
  - or, the input fragment is the start of an encoded sequence in the input encoding and the **mbstate\_t** object was initialized to the initial conversion state;
  - or, the input fragment is a continuation of an encoded sequence and **mbstate\_t** is the result of having advanced to the input and possibly output positions through the application of prior calls to the same transcoding function.

The behavior is undefined when a function described by this subclause is invoked and **\*state** (or, if **state** is a null pointer, the **mbstate\_t** object created for that case) does not describe the conversion state for the current conversion.

- 4 If **input\_size** is a null pointer, **input** shall either be a null pointer or point to a null pointer. Otherwise, **input** shall be a pointer to a non-null pointer to an array of at least **\*input\_size** elements.
- 5 The multi unit transcoding functions convert from code units of type **charX** interpreted according to *encodingX* to code units of type **charY** according to *encodingY* given a conversion state of value **\*state**. The behavior of these functions is as-if the analogous single unit function **stdc\_XntoYn** was repeatedly called, with the same **output**, **output\_size**, **input**, **input\_size**, and **state** parameters, to perform multiple indivisible units of work. The function stops when an error occurs or the **input** is exhausted (only signified when **\*input\_size** is zero and the **mbsinit**(**\*state**) returns nonzero).
- 6 The multi unit transcoding functions behave as-if:
1. If **state** is a null pointer, then an automatic storage duration object of type **mbstate\_t** is created which is unique to the current invocation. It is initialized to the initial conversion state and a pointer to this object is used wherever **state** is used in this paragraph.
  2. **stdc\_XnrtoYn** is called with **output\_size**, **output**, **input\_size**, **input**, and **state** with its result stored in a temporary named **err**.
  3. If **input** is a null pointer or **\*input** is a null pointer, return **err**.

4. If `err` is not `stdc_mcerr_ok`, then return `err`.
5. Otherwise, if `*input_size` is greater than zero, go back to (2).
6. Otherwise, if `mbsinit(*state)` returns zero, go back to (2).
7. Otherwise, return `err`.

### Returns

- 7 On success or failure, the transcoding functions shall return one of the above error codes (7.25.1). If `state` is not a null pointer and `*state` is not initialized to the initial conversion state for the function on its first use, or is used after being input into a function whose result is not one of `stdc_mcerr_ok`, `stdc_mcerr_incomplete_input`, or `stdc_mcerr_insufficient_output`, the behavior of the functions is unspecified.
- 8 The following is true after the invocation:
  - `*input` is incremented by the number of code units read and successfully converted if `input` and `*input` are not null pointers. If `stdc_mcerr_ok` is returned, then all the input is consumed. Otherwise, `*input` points to the location just after the last successfully completed indivisible unit of work.
  - `*input_size` is decremented by the number of code units read from `*input` that were successfully converted. If no error occurred, then `*input_size` is 0.
  - if `output` and `*output` is not a null pointer, `*output` is incremented by the number of code units written from successfully completed indivisible unit of work.
  - if `output_size` is not a null pointer, `*output_size` is decremented by the number of code units written to the output or that would have been written to the output.

If the return value is `stdc_mcerr_invalid` and `state` is not a null pointer, then `*state` is in an unspecified state.

- 9 NOTE The object unique to the invocation is reused for every call in the second step of the multi unit sized conversion algorithm, and not recreated. If `state` is a null pointer, and the function uses e.g. a created automatic storage duration `mbstate_t` object that is discarded by the end of the invocation, then any potential conversion state contained in the created `mbstate_t` object, and used or accumulated during multi unit processing, could become unrecoverable to the program.
- 10 EXAMPLE 1 The following is an example of using a single indivisible unit sized conversion function `stdc_mcnrtoc8n` to implement a multi unit sized conversion algorithm:

```
#include <stdmchar.h>

stdc_mcerr sample_mcsnrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t* restrict state) {
 mbstate_t invocation_unique_internal_state;
 if (state == nullptr) {
 invocation_unique_internal_state = (mbstate_t){};
 state = &invocation_unique_internal_state;
 }
 if (input == nullptr || *input == nullptr) {
 return stdc_mcnrtoc8n(output_size, output,
 input_size, input, state);
 }
 for (;;) {
 stdc_mcerr err = stdc_mcnrtoc8n(output_size, output,
 input_size, input, state);

 if (err != stdc_mcerr_ok) {
 return err;
 }
 }
}
```

```

 }
 if (*input_size > 0) {
 continue;
 }
 // some execution encodings (6.2.9) may contain
 // additional output as input gets processed
 int state_finished = mbsinit(state);
 if (state_finished == 0) {
 continue;
 }
 return err;
}
}

```

- 11 EXAMPLE 2 The multi unit sized conversion functions can be used to perform other functionality, such as counting, validation, and more by using a null pointer value for specific arguments:

```

#include <stdmchar.h>

bool is_valid_utf16_from_utf8(size_t str_n, const char8_t str[restrict static
str_n]) {
 stdc_mcerr err = stdc_c8snrtoc16sn(nullptr, nullptr, &str_n, &str);
 return err == stdc_mcerr_ok;
}

size_t count_utf16_from_utf8(size_t str_n, const char8_t str[restrict static
str_n]) {
 const size_t utf16_before_n = SIZE_MAX;
 size_t utf16_after_n = utf16_before_n;
 stdc_mcerr err = stdc_c8snrtoc16sn(&utf16_after_n, nullptr, &str_n, &str);
 return err == stdc_mcerr_ok ? utf16_before_n - utf16_after_n : 0;
}

bool unbounded_conversion_utf16_from_utf8(size_t str_n,
const char8_t str[restrict static str_n],
char16_t* restrict dest_str) {
 stdc_mcerr err = stdc_c8snrtoc16sn(nullptr, &dest_str, &str_n, &str);
 return err == stdc_mcerr_ok;
}

int main () {
 const char8_t str[] = u8"\u201CSaw a \U0001F9DC \u2014"
 u8"didn't catch her\u2026 \U0001F61E\u201D\n"
 u8"\t- Sniff";
 // includes null terminator
 const size_t str_n = _Lengthof(str);

 if (!is_valid_utf16_from_utf8(str_n, str)) {
 // input not valid
 return 1;
 }

 size_t utf16_str_n = count_utf16_from_utf8(str_n, str);
 constexpr size_t utf16_str_max_size = STDC_C16_MAX * (sizeof(str) / sizeof
(*str));
 char16_t utf16_str[utf16_str_max_size] = {};
 if (utf16_str_max_size < utf16_str_n) {
 // buffer too small
 return 2;
 }
}

```

```

 if (!unbounded_conversion_utf16_from_utf8(str_n, str, utf16_str)) {
 // write failed
 return 3;
 }
 // At this point, utf16_str is a veritable UTF-16 string.
 // As noted above, null terminator from utf8_str was included:
 // utf16_str is a sequence of UTF-16 code units plus the null
 // terminator, in a suitable form at the end of the UTF-16 string.
 return 0;
}

```

The above program demonstrates validating, counting, and doing an unbounded (size unsafe) write using the provided functions. Caution should be taken when a program uses unbounded writes, as the size of the buffer is assumed to be large enough during the call to the multi unit sized conversion function when `output_size` is a null pointer. An implementation can detect the above cases where specific arguments or their pointed to values are a null pointer value, and provide improved implementations relying on properties from these assumptions.

#### Recommended practice

- 12 The multi unit transcoding functions are explicitly for the purpose of performing conversions on the largest contiguous section of valid data in the shortest amount of time possible. Implementations should take advantage of the information of null pointer values for the output size pointer, output data pointer, or both, to drastically improve performance characteristics for assumed unlimited write space, output counting scenarios, or input validation/counting, respectively.
- 13 Implementations should prefer returning an error for an incomplete input sequence over storing intermediate data within the state where possible for non-Unicode encodings. By leaving partial input unconsumed, it can be easier for functionality built on top of the functions in this subclause to report errors without skipping over potentially invalid input data.

## 7.26 **\_Noreturn** <stdnoreturn.h>

- 1 The header <stdnoreturn.h> defines the macro

**noreturn**

which expands to **\_Noreturn**.

- 2 The **noreturn** macro and the <stdnoreturn.h> header are obsolescent features.

## 7.27 String handling <string.h>

### 7.27.1 String function conventions

- 1 The header <string.h> declares one type, several functions, several type-generic functions, and defines two macros useful for manipulating arrays of character type and other objects treated as arrays of character type.<sup>366)</sup> The type is **size\_t** and one of the macros is **NULL** (both described in 7.21). Various methods are used for determining the lengths of the arrays, but in all cases a **char \*** or **void \*** argument points to the initial (lowest addressed) character of the array. If an array is accessed beyond the end of an object, the behavior is undefined.
- 2 The macro

```
__STDC_VERSION_STRING_H__
```

is an integer constant expression with a value equivalent to 202ymmL.

- 3 Where an argument declared as **size\_t n** specifies the length of the array for a function, **n** can have the value zero on a call to that function. Unless explicitly stated otherwise in the description of a particular function in this subclause, pointer arguments on such a call may be null pointers. On such a call, a function that locates a character finds no occurrence, a function that compares two character sequences returns zero, and a function that copies characters copies zero characters.
- 4 For all functions in this subclause, each character shall be interpreted as if it had the type **unsigned char** (and therefore every possible object representation is valid and has a different value).

### 7.27.2 Copying functions

#### 7.27.2.1 The **memcpy** function

##### Synopsis

```
1 #include <string.h>
 void *memcpy(void * restrict s1, const void * restrict s2, size_t n);
```

##### Description

- 2 The **memcpy** function copies **n** characters from the object pointed to by **s2** into the object pointed to by **s1**. If copying takes place between objects that overlap, the behavior is undefined.

##### Returns

- 3 The **memcpy** function returns the value of **s1**.

#### 7.27.2.2 The **memccpy** function

##### Synopsis

```
1 #include <string.h>
 void *memccpy(void * restrict s1, const void * restrict s2, int c, size_t n);
```

##### Description

- 2 The **memccpy** function copies characters from the object pointed to by **s2** into the object pointed to by **s1**, stopping after the first occurrence of character **c** (converted to an **unsigned char**) is copied, or after **n** characters are copied, whichever comes first. If copying takes place between objects that overlap, the behavior is undefined.

##### Returns

- 3 The **memccpy** function returns a pointer to the character after the copy of **c** in **s1**, or a null pointer if **c** was not found in the first **n** characters of **s2**.

#### 7.27.2.3 The **memmove** function

##### Synopsis

- 1 <sup>366)</sup>See “future library directions” (7.34.18).

```
#include <string.h>
void *memmove(void *s1, const void *s2, size_t n);
```

### Description

- 2 The **memmove** function copies *n* characters from the object pointed to by *s2* into the object pointed to by *s1*. Copying takes place as if the *n* characters from the object pointed to by *s2* are first copied into a temporary array of *n* characters that does not overlap the objects pointed to by *s1* and *s2*, and then the *n* characters from the temporary array are copied into the object pointed to by *s1*.

### Returns

- 3 The **memmove** function returns the value of *s1*.

#### 7.27.2.4 The **strcpy** function

##### Synopsis

```
1 #include <string.h>
char *strcpy(char * restrict s1, const char * restrict s2);
```

### Description

- 2 The **strcpy** function copies the string pointed to by *s2* (including the terminating null character) into the array pointed to by *s1*. If copying takes place between objects that overlap, the behavior is undefined.

### Returns

- 3 The **strcpy** function returns the value of *s1*.

#### 7.27.2.5 The **strncpy** function

##### Synopsis

```
1 #include <string.h>
char *strncpy(char * restrict s1, const char * restrict s2, size_t n);
```

### Description

- 2 The **strncpy** function copies not more than *n* characters (characters that follow a null character are not copied) from the array pointed to by *s2* to the array pointed to by *s1*.<sup>367)</sup> If copying takes place between objects that overlap, the behavior is undefined.
- 3 If the array pointed to by *s2* is a string that is shorter than *n* characters, null characters are appended to the copy in the array pointed to by *s1*, until *n* characters in all have been written.

### Returns

- 4 The **strncpy** function returns the value of *s1*.

#### 7.27.2.6 The **strdup** function

##### Synopsis

```
1 #include <string.h>
char *strdup(const char *s);
```

### Description

- 2 The **strdup** function creates a copy of the string pointed to by *s* in a space allocated as if by a call to **malloc**.

### Returns

- 3 The **strdup** function returns a pointer to the first character of the duplicate string. The returned pointer can be passed to **free**. If no space can be allocated the **strdup** function returns a null pointer.

<sup>367)</sup>Thus, if there is no null character in the first *n* characters of the array pointed to by *s2*, the result will not be null-terminated.



### 7.27.2.7 The **strndup** function

#### Synopsis

```
1 #include <string.h>
 char *strndup(const char *s, size_t n);
```

#### Description

- 2 The **strndup** function creates a string initialized with no more than **n** initial characters of the array pointed to by **s** and up to the first null character, whichever comes first, in a space allocated as if by a call to **malloc**. If the array pointed to by **s** does not contain a null within the first **n** characters, a null is appended to the copy of the array.

#### Returns

- 3 The **strndup** function returns a pointer to the first character of the created string. The returned pointer can be passed to **free**. If space cannot be allocated the **strndup** function returns a null pointer.

## 7.27.3 Concatenation functions

### 7.27.3.1 The **strcat** function

#### Synopsis

```
1 #include <string.h>
 char *strcat(char * restrict s1, const char * restrict s2);
```

#### Description

- 2 The **strcat** function appends a copy of the string pointed to by **s2** (including the terminating null character) to the end of the string pointed to by **s1**. The initial character of **s2** overwrites the null character at the end of **s1**. If copying takes place between objects that overlap, the behavior is undefined.

#### Returns

- 3 The **strcat** function returns the value of **s1**.

### 7.27.3.2 The **strncat** function

#### Synopsis

```
1 #include <string.h>
 char *strncat(char * restrict s1, const char * restrict s2, size_t n);
```

#### Description

- 2 The **strncat** function appends not more than **n** characters (a null character and characters that follow it are not appended) from the array pointed to by **s2** to the end of the string pointed to by **s1**. The initial character of **s2** overwrites the null character at the end of **s1**. A terminating null character is always appended to the result.<sup>368)</sup> If copying takes place between objects that overlap, the behavior is undefined.

#### Returns

- 3 The **strncat** function returns the value of **s1**.

Forward references: the **strlen** function (7.27.6.4).

## 7.27.4 Comparison functions

### 7.27.4.1 General

- 1 The sign of a nonzero value returned by the comparison functions **memcmp**, **strcmp**, and **strncmp** is determined by the sign of the difference between the values of the first pair of characters (both

<sup>368)</sup>Thus, the maximum number of characters that can end up in the array pointed to by **s1** is **strlen(s1)+n+1**.

interpreted as **unsigned char**) that differ in the objects being compared.

#### 7.27.4.2 The **memcmp** function

##### Synopsis

```
1 #include <string.h>
 int memcmp(const void *s1, const void *s2, size_t n);
```

##### Description

- 2 The **memcmp** function compares the first **n** characters of the object pointed to by **s1** to the first **n** characters of the object pointed to by **s2**.<sup>369)</sup>

##### Returns

- 3 The **memcmp** function returns an integer greater than, equal to, or less than zero, accordingly as the object pointed to by **s1** is greater than, equal to, or less than the object pointed to by **s2**.

#### 7.27.4.3 The **strcmp** function

##### Synopsis

```
1 #include <string.h>
 int strcmp(const char *s1, const char *s2);
```

##### Description

- 2 The **strcmp** function compares the string pointed to by **s1** to the string pointed to by **s2**.

##### Returns

- 3 The **strcmp** function returns an integer greater than, equal to, or less than zero, accordingly as the string pointed to by **s1** is greater than, equal to, or less than the string pointed to by **s2**.

#### 7.27.4.4 The **strcoll** function

##### Synopsis

```
1 #include <string.h>
 int strcoll(const char *s1, const char *s2);
```

##### Description

- 2 The **strcoll** function compares the string pointed to by **s1** to the string pointed to by **s2**, both interpreted as appropriate to the **LC\_COLLATE** category of the current locale.

##### Returns

- 3 The **strcoll** function returns an integer greater than, equal to, or less than zero, accordingly as the string pointed to by **s1** is greater than, equal to, or less than the string pointed to by **s2** when both are interpreted as appropriate to the current locale.

#### 7.27.4.5 The **strncmp** function

##### Synopsis

```
1 #include <string.h>
 int strncmp(const char *s1, const char *s2, size_t n);
```

##### Description

- 2 The **strncmp** function compares not more than **n** characters (characters that follow a null character are not compared) from the array pointed to by **s1** to the array pointed to by **s2**.

<sup>369)</sup>The unused bytes used as padding for purposes of alignment within structure objects take on unspecified values when a value is stored in the object (see 6.2.6.1). Strings shorter than their allocated space and unions can also cause problems in comparison.

**Returns**

- 3 The **strncmp** function returns an integer greater than, equal to, or less than zero, accordingly as the possibly null-terminated array pointed to by **s1** is greater than, equal to, or less than the possibly null-terminated array pointed to by **s2**.

**7.27.4.6 The **strxfrm** function****Synopsis**

```
1 #include <string.h>
 size_t strxfrm(char * restrict s1, const char * restrict s2, size_t n);
```

**Description**

- 2 The **strxfrm** function transforms the string pointed to by **s2** and places the resulting string into the array pointed to by **s1**. The transformation is such that if the **strcmp** function is applied to two transformed strings, it returns a value greater than, equal to, or less than zero, corresponding to the result of the **strcmp** function applied to the same two original strings. No more than **n** characters are placed into the resulting array pointed to by **s1**, including the terminating null character. If **n** is zero, **s1** is permitted to be a null pointer. If copying takes place between objects that overlap, the behavior is undefined.

**Returns**

- 3 The **strxfrm** function returns the length of the transformed string (not including the terminating null character). If the value returned is **n** or more, the members of the array pointed to by **s1** have an indeterminate representation.
- 4 EXAMPLE The value of the following expression is the size of the array needed to hold the transformation of the string pointed to by **s**.

```
1 + strxfrm(nullptr, s, 0)
```

**7.27.5 Search functions****7.27.5.1 Introduction**

- 1 The stateless search functions in this subclause (**memchr**, **strchr**, **strpbrk**, **strrchr**, **strstr**) are *generic functions*. These functions are generic in the qualification of the array to be searched and will return a result pointer to an element with the same qualification as the passed array. If the array to be searched is **const**-qualified, the result pointer will be to a **const**-qualified element. If the array to be searched is not **const**-qualified,<sup>370)</sup> the result pointer will be to an unqualified element.
- 2 The external declarations of these generic functions have a concrete function type that returns a pointer to an unqualified element (of type **char** when specified as *QChar*, and **void** when specified as *QVoid*), and accepts a pointer to a **const**-qualified array of the same type to search. This signature supports all correct uses. If a macro definition of any of these generic functions is suppressed to access an actual function, the external declaration with the corresponding concrete type is visible.<sup>371)</sup>
- 3 The **volatile** and **restrict** qualifiers are not accepted on the elements of the array to search.

**7.27.5.2 The **memchr** generic function****Synopsis**

```
1 #include <string.h>
 QVoid *memchr(QVoid *s, int c, size_t n);
```

**Description**

- 2 The **memchr** generic function locates the first occurrence of **c** (converted to an **unsigned char**) in the initial **n** characters (each interpreted as **unsigned char**) of the object pointed to by **s**. The

<sup>370)</sup>The null pointer constant is not a pointer to a **const**-qualified type, and therefore the result expression has the type of a pointer to an unqualified element.

<sup>371)</sup>This is an obsolescent feature.

implementation shall behave as if it reads the characters sequentially and stops as soon as a matching character is found.

#### Returns

- 3 The **memchr** generic function returns a pointer to the located character, or a null pointer if the character does not occur in the object.

#### 7.27.5.3 The **strchr** generic function

##### Synopsis

```
1 #include <string.h>
 QChar *strchr(QChar *s, int c);
```

##### Description

- 2 The **strchr** generic function locates the first occurrence of **c** (converted to a **char**) in the string pointed to by **s**. The terminating null character is considered to be part of the string.

#### Returns

- 3 The **strchr** generic function returns a pointer to the located character, or a null pointer if the character does not occur in the string.

#### 7.27.5.4 The **strcspn** function

##### Synopsis

```
1 #include <string.h>
 size_t strcspn(const char *s1, const char *s2);
```

##### Description

- 2 The **strcspn** function computes the length of the maximum initial segment of the string pointed to by **s1** which consists entirely of characters *not* from the string pointed to by **s2**.

#### Returns

- 3 The **strcspn** function returns the length of the segment.

#### 7.27.5.5 The **strpbrk** generic function

##### Synopsis

```
1 #include <string.h>
 QChar *strpbrk(QChar *s1, const char *s2);
```

##### Description

- 2 The **strpbrk** generic function locates the first occurrence in the string pointed to by **s1** of any character from the string pointed to by **s2**.

#### Returns

- 3 The **strpbrk** generic function returns a pointer to the character, or a null pointer if no character from **s2** occurs in **s1**.

#### 7.27.5.6 The **strrchr** generic function

##### Synopsis

```
1 #include <string.h>
 QChar *strrchr(QChar *s, int c);
```

##### Description

- 2 The **strrchr** generic function locates the last occurrence of **c** (converted to a **char**) in the string pointed to by **s**. The terminating null character is considered to be part of the string.

**Returns**

- 3 The **strchr** generic function returns a pointer to the character, or a null pointer if **c** does not occur in the string.

**7.27.5.7 The **strspn** function****Synopsis**

```
1 #include <string.h>
 size_t strspn(const char *s1, const char *s2);
```

**Description**

- 2 The **strspn** function computes the length of the maximum initial segment of the string pointed to by **s1** which consists entirely of characters from the string pointed to by **s2**.

**Returns**

- 3 The **strspn** function returns the length of the segment.

**7.27.5.8 The **strstr** generic function****Synopsis**

```
1 #include <string.h>
 QChar *strstr(QChar *s1, const char *s2);
```

**Description**

- 2 The **strstr** generic function locates the first occurrence in the string pointed to by **s1** of the sequence of characters (excluding the terminating null character) in the string pointed to by **s2**.

**Returns**

- 3 The **strstr** generic function returns a pointer to the located string, or a null pointer if the string is not found. If **s2** points to a string with zero length, the function returns **s1**.

**7.27.5.9 The **strtok** function****Synopsis**

```
1 #include <string.h>
 char *strtok(char * restrict s1, const char * restrict s2);
```

**Description**

- 2 A sequence of calls to the **strtok** function breaks the string pointed to by **s1** into a sequence of tokens, each of which is delimited by a character from the string pointed to by **s2**. The first call in the sequence has a non-null first argument; subsequent calls in the sequence have a null first argument. If any of the subsequent calls in the sequence is made by a different thread than the first, the behavior is undefined. The separator string pointed to by **s2** can be different from call to call.
- 3 The first call in the sequence searches the string pointed to by **s1** for the first character that is *not* contained in the current separator string pointed to by **s2**. If no such character is found, then there are no tokens in the string pointed to by **s1** and the **strtok** function returns a null pointer. If such a character is found, it is the start of the first token.
- 4 The **strtok** function then searches from there for a character that *is* contained in the current separator string. If no such character is found, the current token extends to the end of the string pointed to by **s1**, and subsequent searches for a token will return a null pointer. If such a character is found, it is overwritten by a null character, which terminates the current token. The **strtok** function saves a pointer to the following character, from which the next search for a token will start.
- 5 Each subsequent call, with a null pointer as the value of the first argument, starts searching from the saved pointer and behaves as described previously.

- 6 The **strtok** function is not required to avoid data races with other calls to the **strtok** function.<sup>372)</sup>  
The implementation shall behave as if no library function calls the **strtok** function.

#### Returns

- 7 The **strtok** function returns a pointer to the first character of a token, or a null pointer if there is no token.

- 8 EXAMPLE

```
#include <string.h>
static char str[] = "?a???b,,,#c";
char *t;

t = strtok(str, "?"); // t points to the token "a"
t = strtok(nullptr, ","); // t points to the token "??b"
t = strtok(nullptr, "#,"); // t points to the token "c"
t = strtok(nullptr, "?"); // t is a null pointer
```

Forward references: The **strtok\_s** function (K.3.7.4.1).

## 7.27.6 Miscellaneous functions

### 7.27.6.1 The **memset** function

#### Synopsis

```
1 #include <string.h>
 void *memset(void *s, int c, size_t n);
```

#### Description

- 2 The **memset** function copies the value of **c** (converted to an **unsigned char**) into each of the first **n** characters of the object pointed to by **s**.

#### Returns

- 3 The **memset** function returns the value of **s**.

### 7.27.6.2 The **memset\_explicit** function

#### Synopsis

```
1 #include <string.h>
 void *memset_explicit(void *s, int c, size_t n);
```

#### Description

- 2 The **memset\_explicit** function copies the value of **c** (converted to an **unsigned char**) into each of the first **n** characters of the object pointed to by **s**. The purpose of this function is to make sensitive information stored in the object inaccessible.<sup>373)</sup>

#### Returns

- 3 The **memset\_explicit** function returns the value of **s**.

### 7.27.6.3 The **strerror** function

#### Synopsis

```
1 #include <string.h>
 char *strerror(int errnum);
```

<sup>372)</sup>The **strtok\_s** function can be used instead to avoid data races.

<sup>373)</sup>The intention is that the memory store is always performed (i.e. never elided), regardless of optimizations. This is in contrast to calls to the **memset** function (7.27.6.1)

**Description**

- 2 The **strerror** function maps the number in **errnum** to a message string. Typically, the values for **errnum** come from **errno**, but **strerror** shall map any value of type **int** to a message.
- 3 The **strerror** function is not required to avoid data races with other calls to the **strerror** function.<sup>374)</sup> The implementation shall behave as if no library function calls the **strerror** function.

**Returns**

- 4 The **strerror** function returns a pointer to the string, the contents of which are locale-specific. The array pointed to shall not be modified by the program. The behavior is undefined if the returned value is used after a subsequent call to the **strerror** function, or after the thread which called the function to obtain the returned value has exited.

**Forward references:** The **strerror\_s** function (K.3.7.5.2).

**7.27.6.4 The **strlen** function****Synopsis**

```
1 #include <string.h>
 size_t strlen(const char *s);
```

**Description**

- 2 The **strlen** function computes the length of the string pointed to by **s**.

**Returns**

- 3 The **strlen** function returns the number of characters that precede the terminating null character.

**7.27.6.5 The **strnlen** function****Synopsis**

```
1 #include <string.h>
 size_t strnlen(const char *s, size_t n);
```

**Description**

- 2 The **strnlen** function counts not more than **n** characters (a null character and characters that follow it are not counted) in the array to which **s** points. At most the first **n** characters of **s** shall be accessed by **strnlen**. The implementation shall behave as if it reads the characters sequentially and stops as soon as a null character is found.

**Returns**

- 3 The **strnlen** function returns the number of characters that precede the terminating null character. If there is no null character in the first **n** characters of **s** then **strnlen** returns **n**.

<sup>374)</sup>The **strerror\_s** function can be used instead to avoid data races.

## 7.28 Type-generic math <tgmath.h>

- 1 The header <tgmath.h> includes the headers <math.h> and <complex.h> and defines several type-generic macros.
- 2 The feature test macro `__STDC_VERSION_TGMATH_H__` expands to the token 202311L.
- 3 This clause specifies a many-to-one correspondence of functions in <math.h> and <complex.h> with *type-generic macros*.<sup>375)</sup> Use of a type-generic macro invokes a corresponding function whose type is determined by the types of the arguments for particular parameters called the *generic parameters*.<sup>376)</sup>
- 4 Of the <math.h> and <complex.h> functions without an **f** (**float**) or **l** (**long double**) suffix, several have one or more parameters whose corresponding real type is **double**. For each such function, except the functions that round result to narrower type (7.12.15) (which are covered subsequently in this subclause) and **modf**, there is a corresponding type-generic macro. The parameters whose corresponding real type is **double** in the function synopsis are generic parameters.
- 5 Some of the <math.h> functions for decimal floating types have no unsuffixed counterpart. Of these functions with a **d64** suffix, some have one or more parameters whose type is **\_Decimal64**. For each such function, except **decodedecd64**, **encodedecd64**, **decodebind64**, and **encodebind64**, there is a corresponding type-generic macro. The parameters whose real type is **\_Decimal64** in the function synopsis are generic parameters.
- 6 If arguments for generic parameters of a type-generic macro are such that some argument has a corresponding real type that is of standard floating type and another argument is of decimal floating type, the behavior is undefined.
- 7 Except for the macros for functions that round result to a narrower type (7.12.15), use of a type-generic macro invokes a function whose generic parameters have the corresponding real type determined by the types of the arguments for the generic parameters as follows:
  - Arguments of integer type are regarded as having type **\_Decimal64** if any argument has decimal floating type, and as having type **double** otherwise.
  - If the function has exactly one generic parameter, the type determined is the corresponding real type of the argument for the generic parameter.
  - If the function has exactly two generic parameters, the type determined is the corresponding real type determined by the usual arithmetic conversions (6.3.2.8) applied to the arguments for the generic parameters.
  - If the function has more than two generic parameters, the type determined is the corresponding real type determined by repeatedly applying the usual arithmetic conversions, first to the first two arguments for generic parameters, then to that result type and the next argument for a generic parameter, and so forth until the usual arithmetic conversions have been applied to the last argument for a generic parameter.

If neither <math.h> and <complex.h> define a function whose generic parameters have the determined corresponding real type, the behavior is undefined.

- 8 For each unsuffixed function in <math.h> for which there is a function in <complex.h> with the same name except for a **c** prefix, the corresponding type-generic macro (for both functions) has the same name as the function in <math.h>. The corresponding type-generic macro for **fabs** and **cabs** is **fabs**. Table 7.9 shows the associations between functions and type-generic macros.

<sup>375)</sup> Like other function-like macros in standard libraries, each type-generic macro can be suppressed to make available the corresponding ordinary function.

<sup>376)</sup> If the type of the argument is not compatible with the type of the parameter for the selected function, the behavior is undefined.



Table 7.9 — Functions and type-generic macro relations

| <math.h><br>function | <complex.h><br>function | type-generic<br>macro |
|----------------------|-------------------------|-----------------------|
| <b>acos</b>          | <b>cacos</b>            | <b>acos</b>           |
| <b>asin</b>          | <b>casin</b>            | <b>asin</b>           |
| <b>atan</b>          | <b>catan</b>            | <b>atan</b>           |
| <b>acosh</b>         | <b>cacosh</b>           | <b>acosh</b>          |
| <b>asinh</b>         | <b>casinh</b>           | <b>asinh</b>          |
| <b>atanh</b>         | <b>catanh</b>           | <b>atanh</b>          |
| <b>cos</b>           | <b>ccos</b>             | <b>cos</b>            |
| <b>sin</b>           | <b>csin</b>             | <b>sin</b>            |
| <b>tan</b>           | <b>ctan</b>             | <b>tan</b>            |
| <b>cosh</b>          | <b>ccosh</b>            | <b>cosh</b>           |
| <b>sinh</b>          | <b>csinh</b>            | <b>sinh</b>           |
| <b>tanh</b>          | <b>ctanh</b>            | <b>tanh</b>           |
| <b>exp</b>           | <b>cexp</b>             | <b>exp</b>            |
| <b>log</b>           | <b>clog</b>             | <b>log</b>            |
| <b>pow</b>           | <b>cpow</b>             | <b>pow</b>            |
| <b>sqrt</b>          | <b>csqrt</b>            | <b>sqrt</b>           |
| <b>fabs</b>          | <b>cabs</b>             | <b>fabs</b>           |

If at least one argument for a generic parameter is complex, then use of the macro invokes a complex function; otherwise, use of the macro invokes a real function.

- 9 For each unsuffixed function in <math.h> without a **c**-prefixed counterpart in <complex.h> (except functions that round result to narrower type, **modf**, and **canonicalize**), the corresponding type-generic macro has the same name as the function. These type-generic macros are:

|                  |                         |                |                   |                  |
|------------------|-------------------------|----------------|-------------------|------------------|
| <b>acospi</b>    | <b>exp2</b>             | <b>fmod</b>    | <b>log2</b>       | <b>rootn</b>     |
| <b>asinpi</b>    | <b>expm1</b>            | <b>frexp</b>   | <b>logb</b>       | <b>roundeven</b> |
| <b>atan2pi</b>   | <b>fdim</b>             | <b>fromfpx</b> | <b>logp1</b>      | <b>round</b>     |
| <b>atan2</b>     | <b>floor</b>            | <b>fromfp</b>  | <b>lrint</b>      | <b>rsqrt</b>     |
| <b>atanpi</b>    | <b>fmax</b>             | <b>hypot</b>   | <b>lround</b>     | <b>scalbln</b>   |
| <b>cbrt</b>      | <b>fmaximum</b>         | <b>ilogb</b>   | <b>nearbyint</b>  | <b>scalbn</b>    |
| <b>ceil</b>      | <b>fmaximum_mag</b>     | <b>ldexp</b>   | <b>nextafter</b>  | <b>sinpi</b>     |
| <b>compoundn</b> | <b>fmaximum_num</b>     | <b>lgamma</b>  | <b>nextdown</b>   | <b>tanpi</b>     |
| <b>copysign</b>  | <b>fmaximum_mag_num</b> | <b>llogb</b>   | <b>nexttoward</b> | <b>tgamma</b>    |
| <b>cospi</b>     | <b>fma</b>              | <b>llrint</b>  | <b>nextup</b>     | <b>trunc</b>     |
| <b>erfc</b>      | <b>fmin</b>             | <b>llround</b> | <b>pown</b>       | <b>ufromfpx</b>  |
| <b>erf</b>       | <b>fminimum</b>         | <b>log10p1</b> | <b>powr</b>       | <b>ufromfp</b>   |
| <b>expl0m1</b>   | <b>fminimum_mag</b>     | <b>log10</b>   | <b>remainder</b>  |                  |
| <b>expl0</b>     | <b>fminimum_num</b>     | <b>log1p</b>   | <b>remquo</b>     |                  |
| <b>exp2m1</b>    | <b>fminimum_mag_num</b> | <b>log2p1</b>  | <b>rint</b>       |                  |

If all arguments for generic parameters are real, then use of the macro invokes a real function (provided <math.h> defines a function of the determined type); otherwise, use of the macro is undefined.

- 10 For each unsuffixed function in <complex.h> that is not a **c**-prefixed counterpart to a function in <math.h>, the corresponding type-generic macro has the same name as the function. These type-generic macros are:

**carg**                      **cimag**                      **conj**                      **cproj**                      **creal**

Use of the macro with any argument of standard floating or complex type invokes a complex function. Use of the macro with an argument of decimal floating type is undefined.

- 11 The functions that round result to a narrower type have type-generic macros whose names are obtained by omitting any suffix from the function names. Thus, the macros with **f** or **d** prefix are:

|             |             |             |             |             |              |
|-------------|-------------|-------------|-------------|-------------|--------------|
| <b>fadd</b> | <b>fsub</b> | <b>fmul</b> | <b>fdiv</b> | <b>ffma</b> | <b>fsqrt</b> |
| <b>dadd</b> | <b>dsub</b> | <b>dmul</b> | <b>ddiv</b> | <b>dfma</b> | <b>dsqrt</b> |

and the macros with **d32** or **d64** prefix are:

|               |               |               |               |               |                |
|---------------|---------------|---------------|---------------|---------------|----------------|
| <b>d32add</b> | <b>d32sub</b> | <b>d32mul</b> | <b>d32div</b> | <b>d32fma</b> | <b>d32sqrt</b> |
| <b>d64add</b> | <b>d64sub</b> | <b>d64mul</b> | <b>d64div</b> | <b>d64fma</b> | <b>d64sqrt</b> |

All arguments shall be real. If the macro prefix is **f** or **d**, use of an argument of decimal floating type is undefined. If the macro prefix is **d32** or **d64**, use of an argument of standard floating type is undefined. The function invoked is determined as follows:

- If any argument has type **\_Decimal128**, or if the macro prefix is **d64**, the function invoked has the name of the macro, with a **d128** suffix.
  - Otherwise, if the macro prefix is **d32**, the function invoked has the name of the macro, with a **d64** suffix.
  - Otherwise, if any argument has type **long double**, or if the macro prefix is **d**, the function invoked has the name of the macro, with an **l** suffix.
  - Otherwise, the function invoked has the name of the macro (with no suffix).
- 12 For each **d64**-suffixed function in `<math.h>`, except **decodedecd64**, **encodedecd64**, **decodebind64**, and **encodebind64**, that does not have an unsuffixed counterpart, the corresponding type-generic macro has the name of the function, but without the suffix. These type-generic macros are displayed in Table 7.10:

Table 7.10 — Decimal functions and type-generic macro relations

| <code>&lt;math.h&gt;</code> function | type-generic macro |
|--------------------------------------|--------------------|
| <b>quantized</b> <i>N</i>            | <b>quantize</b>    |
| <b>samequantum</b> <i>dN</i>         | <b>samequantum</b> |
| <b>quantum</b> <i>dN</i>             | <b>quantum</b>     |
| <b>llquantexp</b> <i>dN</i>          | <b>llquantexp</b>  |

Use of the macro with an argument of standard floating or complex type or with only integer type arguments is undefined.

- 13 A type-generic macro corresponding to a function indicated in Table 7.1 is affected by constant rounding modes (7.6.5).
- 14 NOTE The type-generic macro definition in the example in 6.5.2.1 does not conform to this specification. A conforming macro can be implemented as follows:

```
#define cbrt(X) \
_Generic((X), \
 long double: _Roundwise_cbrtl, \
 default: _Roundwise_cbrt, \
 float: _Roundwise_cbrtf \
)(X)
```

where **\_Roundwise\_cbrtl**, **\_Roundwise\_cbrt**, and **\_Roundwise\_cbrtf** are pointers to functions that are equivalent to **cbrtl**, **cbrt**, and **cbrtf**, respectively, but that are guaranteed to be affected by constant rounding modes (7.6.3).

## 15 EXAMPLE With the declarations

```

#include <tgmath.h>
int n;
float f;
double d;
long double ld;
float complex fc;
double complex dc;
long double complex ldc;
#ifdef __STDC_IEC_60559_DFP__
_Decimal32 d32;
_Decimal64 d64;
_Decimal128 d128;
#endif

```

functions invoked by use of type-generic macros are shown in the Table 7.11:

Table 7.11 — Generic macro use to underlying invocation

| macro use                             | invocation                                  |
|---------------------------------------|---------------------------------------------|
| <code>exp(n)</code>                   | <code>exp(n)</code> , the function          |
| <code>acosh(f)</code>                 | <code>acoshf(f)</code>                      |
| <code>sin(d)</code>                   | <code>sin(d)</code> , the function          |
| <code>atan(ld)</code>                 | <code>atanl(ld)</code>                      |
| <code>log(fc)</code>                  | <code>clogf(fc)</code>                      |
| <code>sqrt(dc)</code>                 | <code>csqrt(dc)</code>                      |
| <code>pow(ldc, f)</code>              | <code>cpowl(ldc, f)</code>                  |
| <code>remainder(n, n)</code>          | <code>remainder(n, n)</code> , the function |
| <code>nextafter(d, f)</code>          | <code>nextafter(d, f)</code> , the function |
| <code>nexttoward(f, ld)</code>        | <code>nexttowardf(f, ld)</code>             |
| <code>copysign(n, ld)</code>          | <code>copysignl(n, ld)</code>               |
| <code>ceil(fc)</code>                 | undefined                                   |
| <code>rint(dc)</code>                 | undefined                                   |
| <code>fmaximum(ldc, ld)</code>        | undefined                                   |
| <code>carg(n)</code>                  | <code>carg(n)</code> , the function         |
| <code>cproj(f)</code>                 | <code>cprojf(f)</code>                      |
| <code>creal(d)</code>                 | <code>creal(d)</code> , the function        |
| <code>cimag(ld)</code>                | <code>cimagl(ld)</code>                     |
| <code>fabs(fc)</code>                 | <code>cabsf(fc)</code>                      |
| <code>carg(dc)</code>                 | <code>carg(dc)</code> , the function        |
| <code>cproj(ldc)</code>               | <code>cprojl(ldc)</code>                    |
| <code>fsub(f, ld)</code>              | <code>fsubl(f, ld)</code>                   |
| <code>fdiv(d, n)</code>               | <code>fdiv(d, n)</code> , the function      |
| <code>dfma(f, d, ld)</code>           | <code>dfmal(f, d, ld)</code>                |
| <code>dadd(f, f)</code>               | <code>daddl(f, f)</code>                    |
| <code>dsqrt(dc)</code>                | undefined                                   |
| <code>exp(d64)</code>                 | <code>expd64(d64)</code>                    |
| <code>sqrt(d32)</code>                | <code>sqrtd32(d32)</code>                   |
| <code>fmaximum(d64, d128)</code>      | <code>fmaximumd128(d64, d128)</code>        |
| <code>pow(d32, n)</code>              | <code>powd64(d32, n)</code>                 |
| <code>remainder(d64, d)</code>        | undefined                                   |
| <code>creal(d64)</code>               | undefined                                   |
| <code>remquo(d32, d32, &amp;n)</code> | undefined                                   |
| <code>llquantexp(d)</code>            | undefined                                   |
| <code>quantize(dc)</code>             | undefined                                   |
| <code>samequantum(n, n)</code>        | undefined                                   |
| <code>d32sub(d32, d128)</code>        | <code>d32subd128(d32, d128)</code>          |
| <code>d32div(d64, n)</code>           | <code>d32divd64(d64, n)</code>              |
| <code>d64fma(d32, d64, d128)</code>   | <code>d64fmad128(d32, d64, d128)</code>     |

|                          |                              |
|--------------------------|------------------------------|
| <b>d64add</b> (d32, d32) | <b>d64addd128</b> (d32, d32) |
| <b>d64sqrt</b> (d)       | undefined                    |
| <b>dadd</b> (n, d64)     | undefined                    |

## 7.29 Threads <threads.h>

### 7.29.1 Introduction

- 1 The header <threads.h> includes the header <time.h>, defines macros, and declares types, enumeration constants, and functions that support multiple threads of execution.<sup>377)</sup>
- 2 Implementations that define the macro `__STDC_NO_THREADS__` may not provide this header nor support any of its facilities.
- 3 The macros are

`ONCE_FLAG_INIT`

which expands to a value that can be used to initialize an object of type `once_flag`; and

`TSS_DTOR_ITERATIONS`

which expands to an integer constant expression representing the maximum number of times that destructors will be called when a thread terminates.

- 4 The types are

`cnd_t`

which is a complete object type that holds an identifier for a condition variable;

`thrd_t`

which is a complete object type that holds an identifier for a thread;

`tss_t`

which is a complete object type that holds an identifier for a thread-specific storage pointer;

`mtx_t`

which is a complete object type that holds an identifier for a mutex;

`tss_dtor_t`

which is the function pointer type `void (*)(void*)`, used for a destructor for a thread-specific storage pointer;

`thrd_start_t`

which is the function pointer type `int (*)(void*)` that is passed to `thrd_create` to create a new thread; and

`once_flag`

which is a complete object type that holds a flag for use by `call_once`.

- 5 The enumeration constants are

`mtx_plain`

which is passed to `mtx_init` to create a mutex object that does not support timeout;

`mtx_recursive`

<sup>377)</sup>See “future library directions” (7.34.20).

which is passed to `mtx_init` to create a mutex object that supports recursive locking;

```
mtx_timed
```

which is passed to `mtx_init` to create a mutex object that supports timeout;

```
thrd_timedout
```

which is returned by a timed wait function to indicate that the time specified in the call was reached without acquiring the requested resource;

```
thrd_success
```

which is returned by a function to indicate that the requested operation succeeded;

```
thrd_busy
```

which is returned by a function to indicate that the requested operation failed because a resource requested by a test and return function is already in use;

```
thrd_error
```

which is returned by a function to indicate that the requested operation failed; and

```
thrd_nomem
```

which is returned by a function to indicate that the requested operation failed because it was unable to allocate memory.

**Forward references:** date and time (7.30).

## 7.29.2 Initialization functions

### 7.29.2.1 The `call_once` function

#### Synopsis

```
1 #include <threads.h>
 void call_once(once_flag *flag, void (*func)(void));
```

#### Description

- 2 The `call_once` function uses the `once_flag` pointed to by `flag` to ensure that `func` is called exactly once, the first time the `call_once` function is called with that value of `flag`. Completion of an effective call to the `call_once` function synchronizes with all subsequent calls to the `call_once` function with the same value of `flag`.

#### Returns

- 3 The `call_once` function returns no value.

## 7.29.3 Condition variable functions

### 7.29.3.1 The `cnd_broadcast` function

#### Synopsis

```
1 #include <threads.h>
 int cnd_broadcast(cnd_t *cnd);
```

**Description**

- 2 The **cnd\_broadcast** function unblocks all the threads that are blocked on the condition variable pointed to by **cond** at the time of the call. If no threads are blocked on the condition variable pointed to by **cond** at the time of the call, the function does nothing.

**Returns**

- 3 The **cnd\_broadcast** function returns **thrd\_success** on success, or **thrd\_error** if the request could not be honored.

**7.29.3.2 The **cnd\_destroy** function****Synopsis**

```
1 #include <threads.h>
 void cnd_destroy(cnd_t *cond);
```

**Description**

- 2 The **cnd\_destroy** function releases all resources used by the condition variable pointed to by **cond**. The **cnd\_destroy** function requires that no threads be blocked waiting for the condition variable pointed to by **cond**.

**Returns**

- 3 The **cnd\_destroy** function returns no value.

**7.29.3.3 The **cnd\_init** function****Synopsis**

```
1 #include <threads.h>
 int cnd_init(cnd_t *cond);
```

**Description**

- 2 The **cnd\_init** function creates a condition variable. If it succeeds it sets the object pointed to by **cond** to a value that uniquely identifies the newly created condition variable. A thread that calls **cnd\_wait** on a newly created condition variable will block.

**Returns**

- 3 The **cnd\_init** function returns **thrd\_success** on success, or **thrd\_nomem** if no memory could be allocated for the newly created condition, or **thrd\_error** if the request could not be honored.

**7.29.3.4 The **cnd\_signal** function****Synopsis**

```
1 #include <threads.h>
 int cnd_signal(cnd_t *cond);
```

**Description**

- 2 The **cnd\_signal** function unblocks one of the threads that are blocked on the condition variable pointed to by **cond** at the time of the call. If no threads are blocked on the condition variable at the time of the call, the function does nothing and returns success.

**Returns**

- 3 The **cnd\_signal** function returns **thrd\_success** on success or **thrd\_error** if the request could not be honored.

**7.29.3.5 The **cnd\_timedwait** function****Synopsis**

```
1 #include <threads.h>
 int cnd_timedwait(cnd_t * restrict cond, mtx_t * restrict mtx,
 const struct timespec * restrict ts);
```

**Description**

- 2 The **cnd\_timedwait** function atomically unlocks the mutex pointed to by **mtx** and blocks until the condition variable pointed to by **cond** is signaled by a call to **cnd\_signal** or to **cnd\_broadcast**, or until after the **TIME\_UTC**-based calendar time pointed to by **ts**, or until it is unblocked due to an unspecified reason. When the calling thread becomes unblocked it locks the object pointed to by **mtx** before it returns. The **cnd\_timedwait** function requires that the mutex pointed to by **mtx** be locked by the calling thread.

**Returns**

- 3 The **cnd\_timedwait** function returns **thrd\_success** upon success, or **thrd\_timedout** if the time specified in the call was reached without acquiring the requested resource, or **thrd\_error** if the request could not be honored.

**7.29.3.6 The cnd\_wait function****Synopsis**

```
1 #include <threads.h>
 int cnd_wait(cnd_t *cond, mtx_t *mtx);
```

**Description**

- 2 The **cnd\_wait** function atomically unlocks the mutex pointed to by **mtx** and blocks until the condition variable pointed to by **cond** is signaled by a call to **cnd\_signal** or to **cnd\_broadcast**, or until it is unblocked due to an unspecified reason. When the calling thread becomes unblocked it locks the mutex pointed to by **mtx** before it returns. The **cnd\_wait** function requires that the mutex pointed to by **mtx** be locked by the calling thread.

**Returns**

- 3 The **cnd\_wait** function returns **thrd\_success** on success or **thrd\_error** if the request could not be honored.

**7.29.4 Mutex functions****7.29.4.1 General**

- 1 For purposes of determining the existence of a data race, lock and unlock operations behave as atomic operations. All lock and unlock operations on a particular mutex occur in some particular total order.
- 2 NOTE This total order can be viewed as the modification order of the mutex.

**7.29.4.2 The mtx\_destroy function****Synopsis**

```
1 #include <threads.h>
 void mtx_destroy(mtx_t *mtx);
```

**Description**

- 2 The **mtx\_destroy** function releases any resources used by the mutex pointed to by **mtx**. No threads can be blocked waiting for the mutex pointed to by **mtx**.

**Returns**

- 3 The **mtx\_destroy** function returns no value.

**7.29.4.3 The mtx\_init function****Synopsis**

```
1 #include <threads.h>
 int mtx_init(mtx_t *mtx, int type);
```



**Description**

- 2 The `mtx_init` function creates a mutex object with properties indicated by `type`, which shall have one of these values:

`mtx_plain` for a simple non-recursive mutex,

`mtx_timed` for a non-recursive mutex that supports timeout,

`mtx_plain` | `mtx_recursive` for a simple recursive mutex, or

`mtx_timed` | `mtx_recursive` for a recursive mutex that supports timeout.

- 3 If the `mtx_init` function succeeds, it sets the mutex pointed to by `mtx` to a value that uniquely identifies the newly created mutex.

**Returns**

- 4 The `mtx_init` function returns `thrd_success` on success, or `thrd_error` if the request could not be honored.

**7.29.4.4 The `mtx_lock` function****Synopsis**

```
1 #include <threads.h>
 int mtx_lock(mtx_t *mtx);
```

**Description**

- 2 The `mtx_lock` function blocks until it locks the mutex pointed to by `mtx`. If the mutex is non-recursive, it shall not be locked by the calling thread. Prior calls to `mtx_unlock` on the same mutex synchronize with this operation.

**Returns**

- 3 The `mtx_lock` function returns `thrd_success` on success, or `thrd_error` if the request could not be honored.

**7.29.4.5 The `mtx_timedlock` function****Synopsis**

```
1 #include <threads.h>
 int mtx_timedlock(mtx_t * restrict mtx, const struct timespec * restrict ts);
```

**Description**

- 2 The `mtx_timedlock` function endeavors to block until it locks the mutex pointed to by `mtx` or until after the **TIME\_UTC**-based calendar time pointed to by `ts`. The specified mutex shall support timeout. If the operation succeeds, prior calls to `mtx_unlock` on the same mutex synchronize with this operation.

**Returns**

- 3 The `mtx_timedlock` function returns `thrd_success` on success, or `thrd_timedout` if the time specified was reached without acquiring the requested resource, or `thrd_error` if the request could not be honored.

**7.29.4.6 The `mtx_trylock` function****Synopsis**

```
1 #include <threads.h>
 int mtx_trylock(mtx_t *mtx);
```

**Description**

- 2 The **mtx\_trylock** function endeavors to lock the mutex pointed to by **mtx**. If the mutex is already locked, the function returns without blocking. If the operation succeeds, prior calls to **mtx\_unlock** on the same mutex synchronize with this operation.

**Returns**

- 3 The **mtx\_trylock** function returns **thrd\_success** on success, or **thrd\_busy** if the resource requested is already in use, or **thrd\_error** if the request could not be honored. **mtx\_trylock** can spuriously fail to lock an unused resource, in which case it returns **thrd\_busy**.

**7.29.4.7 The **mtx\_unlock** function****Synopsis**

```
1 #include <threads.h>
 int mtx_unlock(mtx_t *mtx);
```

**Description**

- 2 The **mtx\_unlock** function unlocks the mutex pointed to by **mtx**. The mutex pointed to by **mtx** shall be locked by the calling thread.

**Returns**

- 3 The **mtx\_unlock** function returns **thrd\_success** on success or **thrd\_error** if the request could not be honored.

**7.29.5 Thread functions****7.29.5.1 The **thrd\_create** function****Synopsis**

```
1 #include <threads.h>
 int thrd_create(thrd_t *thr, thrd_start_t func, void *arg);
```

**Description**

- 2 The **thrd\_create** function creates a new thread executing **func(arg)**. If the **thrd\_create** function succeeds, it sets the object pointed to by **thr** to the identifier of the newly created thread. (A thread's identifier can be reused for a different thread once the original thread has exited and either been detached or joined to another thread.) The completion of the **thrd\_create** function synchronizes with the beginning of the execution of the new thread.
- 3 Returning from **func** has the same behavior as invoking **thrd\_exit** with the value returned from **func**.

**Returns**

- 4 The **thrd\_create** function returns **thrd\_success** on success, or **thrd\_nomem** if no memory could be allocated for the thread requested, or **thrd\_error** if the request could not be honored.

**7.29.5.2 The **thrd\_current** function****Synopsis**

```
1 #include <threads.h>
 thrd_t thrd_current(void);
```

**Description**

- 2 The **thrd\_current** function identifies the thread that called it.

**Returns**

- 3 The **thrd\_current** function returns the identifier of the thread that called it.

**7.29.5.3 The `thrd_detach` function****Synopsis**

```
1 #include <threads.h>
 int thrd_detach(thrd_t thr);
```

**Description**

- 2 The `thrd_detach` function tells the operating system to dispose of any resources allocated to the thread identified by `thr` when that thread terminates. The thread identified by `thr` shall not have been previously detached or joined with another thread.

**Returns**

- 3 The `thrd_detach` function returns `thrd_success` on success or `thrd_error` if the request could not be honored.

**7.29.5.4 The `thrd_equal` function****Synopsis**

```
1 #include <threads.h>
 int thrd_equal(thrd_t thr0, thrd_t thr1);
```

**Description**

- 2 The `thrd_equal` function will determine whether the thread identified by `thr0` refers to the thread identified by `thr1`.

**Returns**

- 3 The `thrd_equal` function returns zero if the thread `thr0` and the thread `thr1` refer to different threads. Otherwise the `thrd_equal` function returns a nonzero value.

**7.29.5.5 The `thrd_exit` function****Synopsis**

```
1 #include <threads.h>
 [[noreturn]] void thrd_exit(int res);
```

**Description**

- 2 For every thread-specific storage key which was created with a non-null destructor and for which the value is non-null, `thrd_exit` sets the value associated with the key to a null pointer value and then invokes the destructor with its previous value. The order in which destructors are invoked is unspecified.
- 3 If after this process there remain keys with both non-null destructors and values, the implementation repeats this process up to `TSS_DTOR_ITERATIONS` times.
- 4 Following this, the `thrd_exit` function terminates execution of the calling thread and sets its result code to `res`.
- 5 The program terminates normally after the last thread has been terminated. The behavior is as if the program called the `exit` function with the status `EXIT_SUCCESS` at thread termination time.

**Returns**

- 6 The `thrd_exit` function returns no value.

**7.29.5.6 The `thrd_join` function****Synopsis**

```
1 #include <threads.h>
 int thrd_join(thrd_t thr, int *res);
```

**Description**

- 2 The **thrd\_join** function joins the thread identified by **thr** with the current thread by blocking until the other thread has terminated. If the parameter **res** is not a null pointer, it stores the thread's result code in the integer pointed to by **res**. The termination of the other thread synchronizes with the completion of the **thrd\_join** function. The thread identified by **thr** shall not have been previously detached or joined with another thread.

**Returns**

- 3 The **thrd\_join** function returns **thrd\_success** on success or **thrd\_error** if the request could not be honored.

**7.29.5.7 The **thrd\_sleep** function****Synopsis**

```
1 #include <threads.h>
 int thrd_sleep(const struct timespec *duration, struct timespec *remaining);
```

**Description**

- 2 The **thrd\_sleep** function suspends execution of the calling thread until either the interval specified by **duration** has elapsed or a signal which is not being ignored is received. If interrupted by a signal and the **remaining** argument is not null, the amount of time remaining (the requested interval minus the time actually slept) is stored in the interval it points to. The **duration** and **remaining** arguments can point to the same object.
- 3 The suspension time can be longer than requested because the interval is rounded up to an integer multiple of the sleep resolution or because of the scheduling of other activity by the system. But, except for the case of being interrupted by a signal, the suspension time will not be less than that specified, as measured by the system clock **TIME\_UTC**.

**Returns**

- 4 The **thrd\_sleep** function returns zero if the requested time has elapsed,  $-1$  if it has been interrupted by a signal, or a negative value (which can also be  $-1$ ) if it fails.

**7.29.5.8 The **thrd\_yield** function****Synopsis**

```
1 #include <threads.h>
 void thrd_yield(void);
```

**Description**

- 2 The **thrd\_yield** function endeavors to permit other threads to run, even if the current thread would ordinarily continue to run.

**Returns**

- 3 The **thrd\_yield** function returns no value.

**7.29.6 Thread-specific storage functions****7.29.6.1 The **tss\_create** function****Synopsis**

```
1 #include <threads.h>
 int tss_create(tss_t *key, tss_dtor_t dtor);
```

**Description**

- 2 The **tss\_create** function creates a thread-specific storage pointer with destructor **dtor**, which can be null.

- 3 A null pointer value is associated with the newly created key in all existing threads. Upon subsequent thread creation, the value associated with all keys is initialized to a null pointer value in the new thread.
- 4 Destructors associated with thread-specific storage are not invoked at program termination.
- 5 The **tss\_create** function shall not be called from within a destructor.

#### Returns

- 6 If the **tss\_create** function is successful, it sets the thread-specific storage pointed to by **key** to a value that uniquely identifies the newly created pointer and returns **thrd\_success**; otherwise, **thrd\_error** is returned and the thread-specific storage pointed to by **key** is set to an indeterminate representation.

#### 7.29.6.2 The **tss\_delete** function

##### Synopsis

```
1 #include <threads.h>
 void tss_delete(tss_t key);
```

##### Description

- 2 The **tss\_delete** function releases any resources used by the thread-specific storage identified by **key**. The **tss\_delete** function shall only be called with a value for **key** that was returned by a call to **tss\_create** before the thread commenced executing destructors.
- 3 If **tss\_delete** is called while another thread is executing destructors, whether this will affect the number of invocations of the destructor associated with **key** on that thread is unspecified.
- 4 Calling **tss\_delete** will not result in the invocation of any destructors.

#### Returns

- 5 The **tss\_delete** function returns no value.

#### 7.29.6.3 The **tss\_get** function

##### Synopsis

```
1 #include <threads.h>
 void *tss_get(tss_t key);
```

##### Description

- 2 The **tss\_get** function returns the value for the current thread held in the thread-specific storage identified by **key**. The **tss\_get** function shall only be called with a value for **key** that was returned by a call to **tss\_create** before the thread commenced executing destructors.

#### Returns

- 3 The **tss\_get** function returns the value for the current thread if successful, or zero if unsuccessful.

#### 7.29.6.4 The **tss\_set** function

##### Synopsis

```
1 #include <threads.h>
 int tss_set(tss_t key, void *val);
```

##### Description

- 2 The **tss\_set** function sets the value for the current thread held in the thread-specific storage identified by **key** to **val**. The **tss\_set** function shall only be called with a value for **key** that was returned by a call to **tss\_create** before the thread commenced executing destructors.
- 3 This action will not invoke the destructor associated with the key on the value being replaced.

#### Returns

- 4 The **tss\_set** function returns **thrd\_success** on success or **thrd\_error** if the request could not be honored.

## 7.30 Date and time <time.h>

### 7.30.1 Components of time

- 1 The header <time.h> defines several macros, and declares types and functions for manipulating time. Many functions deal with a *calendar time* that represents the current date (according to the Gregorian calendar) and time. Some functions deal with *local time*, which is the calendar time expressed for some specific time zone, and with *Daylight Saving Time*, which is a temporary change in the algorithm for determining local time. The local time zone and Daylight Saving Time are implementation-defined.
- 2 The feature test macro `__STDC_VERSION_TIME_H__` expands to the token 202ymmL. The other macros defined are `NULL` (described in 7.21);

```
CLOCKS_PER_SEC
```

which expands to an expression with type `clock_t` (described later in this subclause) that is the number per second of the value returned by the `clock` function;

```
TIME_UTC
TIME_MONOTONIC
```

which expand to integer constant expressions greater than 0 designating the calendar time and monotonic time bases, respectively. Additional time base macro definitions, beginning with `TIME_` and an uppercase letter, may also be specified by the implementation;<sup>378)</sup> and,

```
TIME_ACTIVE
TIME_THREAD_ACTIVE
```

which, if defined, expand to integer values, designating overall execution and thread-specific active processing time bases, respectively.

- 3 The definition of macros for time bases other than `TIME_UTC` are optional. If defined, the corresponding time bases are supported by `timespec_get` and `timespec_getres`, and their values are positive. If defined, the value of the optional macro `TIME_ACTIVE` shall be different from the constants `TIME_UTC` and `TIME_MONOTONIC` and shall not change during the same program invocation. The optional macro `TIME_THREAD_ACTIVE` shall not be defined if the implementation does not support threads; its value shall be different from `TIME_UTC`, `TIME_MONOTONIC`, and `TIME_ACTIVE`, it shall be the same for all expansions of the macro for the same thread, and the value provided for one thread shall not be used by a different thread as the `base` argument of `timespec_get` or `timespec_getres`.
- 4 The types declared are `size_t` (described in 7.21);

```
clock_t
```

and

```
time_t
```

which are real types capable of representing times;

```
struct timespec
```

which holds an interval specified in seconds and nanoseconds (which may represent a calendar time based on a particular epoch); and

```
struct tm
```

<sup>378)</sup>See future library directions (7.34). Implementations can define additional time bases, but are only required to support a real time clock based on UTC.

which holds the components of a calendar time, called the *broken-down time*.

- 5 The range and precision of times representable in `clock_t` and `time_t` are implementation-defined. The `timespec` structure shall contain at least the following members, in any order. The semantics of the members and their normal ranges are expressed in the comments.<sup>379)</sup>

```
time_t tv_sec; // whole seconds -- ≥ 0
/* see the following */ tv_nsec; // nanoseconds -- [0, 999999999]
```

The `tv_nsec` member shall be an implementation-defined signed integer type capable of representing the range [0, 999999999].

- 6 The `tm` structure shall contain at least the following members, in any order.<sup>380)</sup> The semantics of the members and their normal ranges are expressed in the comments.

```
int tm_sec; // seconds after the minute -- [0, 60]
int tm_min; // minutes after the hour -- [0, 59]
int tm_hour; // hours since midnight -- [0, 23]
int tm_mday; // day of the month -- [1, 31]
int tm_mon; // months since January -- [0, 11]
int tm_year; // years since 1900
int tm_wday; // days since Sunday -- [0, 6]
int tm_yday; // days since January 1 -- [0, 365]
int tm_isdst; // Daylight Saving Time flag
```

The value of `tm_isdst` is positive if Daylight Saving Time is in effect, zero if Daylight Saving Time is not in effect, and negative if the information is not available.

## 7.30.2 Time manipulation functions

### 7.30.2.1 The `clock` function

#### Synopsis

```
1 #include <time.h>
 clock_t clock(void);
```

#### Description

- 2 The `clock` function determines the processor time used.

#### Returns

- 3 The `clock` function returns the implementation's best approximation of the active processing time associated with the program execution since the beginning of an implementation-defined era related only to the program invocation. To determine the time in seconds, the value returned by the `clock` function should be divided by the value of the macro `CLOCKS_PER_SEC`. If the processor time used is not available, the function returns the value `(clock_t)(-1)`. If the value cannot be represented, the function returns an unspecified value.<sup>381)</sup>

### 7.30.2.2 The `difftime` function

#### Synopsis

```
1 #include <time.h>
 double difftime(time_t time1, time_t time0);
```

#### Description

- 2 The `difftime` function computes the difference between two calendar times: `time1 - time0`.

<sup>379)</sup>The `tv_sec` member is a linear count of seconds and potentially does not have the normal semantics of a `time_t`.

<sup>380)</sup>The range [0, 60] for `tm_sec` allows for a positive leap second.

<sup>381)</sup>This can be due to overflow of the `clock_t` type.



**Returns**

- 3 The **difftime** function returns the difference expressed in seconds as a **double**.

**7.30.2.3 The mktime function****Synopsis**

```
1 #include <time.h>
 time_t mktime(struct tm *timeptr);
```

**Description**

- 2 The **mktime** function converts the broken-down time, expressed as local time, in the structure pointed to by **timeptr** into a calendar time value with the same encoding as that of the values returned by the **time** function. The original values of the **tm\_wday** and **tm\_yday** components of the structure are ignored, and the original values of the other components are not restricted to the ranges indicated previously. If the local time to be used for the conversion is one that includes Daylight Saving Time adjustments, a positive or zero value for **tm\_isdst** causes the **mktime** function to perform the conversion as if Daylight Saving Time, respectively, is or is not in effect for the specified time. A negative value causes it to attempt to determine whether Daylight Saving Time is in effect for the specified time; if it determines that Daylight Saving Time is in effect it produces the same result as an equivalent call with a positive **tm\_isdst** value, otherwise it produces the same result as an equivalent call with a **tm\_isdst** value of zero.<sup>382)</sup> On successful completion, the components of the structure are set to the same values that would be returned by a call to the **localtime** function with the calculated calendar time as its argument.

**Returns**

- 3 The **mktime** function returns the specified calendar time encoded as a value of type **time\_t**. If the calendar time cannot be represented in the **time\_t** encoding used for the return value or the value to be returned in the **tm\_year** component of the structure pointed to by **timeptr** cannot be represented as an **int**, the function returns the value (**time\_t**) (-1) and does not change the value of the **tm\_wday** component of the structure.
- 4 EXAMPLE What day of the week is July 4, 2001?

```
#include <stdio.h>
#include <time.h>
static const char *const wday[] = {
 "Sunday", "Monday", "Tuesday", "Wednesday",
 "Thursday", "Friday", "Saturday", "-unknown-"
};
struct tm time_str;
/* ... */

time_str.tm_year = 2001 - 1900;
time_str.tm_mon = 7 - 1;
time_str.tm_mday = 4;
time_str.tm_hour = 0;
time_str.tm_min = 0;
time_str.tm_sec = 1;
time_str.tm_isdst = -1;
time_str.tm_wday = 7;
mktime(&time_str);
printf("%s\n", wday[time_str.tm_wday]);
```

**7.30.2.4 The timegm function****Synopsis**

- 1 <sup>382)</sup>If the broken-down time specifies a time that is either skipped over or repeated when a transition to or from Daylight Saving Time occurs, it is unspecified whether the **mktime** function produces the same result as an equivalent call with a positive **tm\_isdst** value or as an equivalent call with a **tm\_isdst** value of zero.

```
#include <time.h>
time_t timegm(struct tm *timeptr);
```

### Description

- 2 The **timegm** function converts the broken-down time, expressed as UTC time, in the structure pointed to by **timeptr** into a calendar time value with the same encoding as that of the values returned by the **time** function. The original values of the **tm\_wday** and **tm\_yday** components of the structure are ignored, and the original values of the other components are not restricted to the ranges indicated previously. On successful completion, the values of the **tm\_wday** and **tm\_yday** components of the structure are set appropriately, and the other components are set to represent the specified calendar time, but with their values forced to the ranges indicated previously; the final value of **tm\_mday** is not set until **tm\_mon** and **tm\_year** are determined.

### Returns

- 3 The **timegm** function returns the specified calendar time encoded as a value of type **time\_t**. If the calendar time cannot be represented in the **time\_t** encoding used for the return value or the value to be returned in the **tm\_year** component of the structure pointed to by **timeptr** cannot be represented as an **int**, the function returns the value **(time\_t) (-1)** and does not change the value of the **tm\_wday** component of the structure.

#### 7.30.2.5 The **time** function

##### Synopsis

```
1 #include <time.h>
 time_t time(time_t *timer);
```

### Description

- 2 The **time** function determines the current calendar time. The encoding of the value is unspecified.

### Returns

- 3 The **time** function returns the implementation's best approximation to the current calendar time. The value **(time\_t) (-1)** is returned if the calendar time is not available. If **timer** is not a null pointer, the return value is also assigned to the object it points to.

#### 7.30.2.6 The **timespec\_get** function

##### Synopsis

```
1 #include <time.h>
 int timespec_get(struct timespec *ts, int base);
```

### Description

- 2 The **timespec\_get** function sets the interval pointed to by **ts** to hold the current calendar time based on the specified time base.
- 3 If **base** is **TIME\_UTC**, the **tv\_sec** member is set to the number of seconds since an implementation-defined *epoch*, truncated to a whole value and the **tv\_nsec** member is set to the integral number of nanoseconds, rounded to the resolution of the system clock.<sup>383)</sup> The optional time base **TIME\_MONOTONIC** is the same, but the reference point is an implementation-defined time point; different program invocations can refer to the same or different reference points.<sup>384)</sup> For the same program invocation, the results of two calls to **timespec\_get** with **TIME\_MONOTONIC** such that the first happens before the second shall not be decreasing. It is implementation-defined if **TIME\_MONOTONIC** accounts for time during which the execution environment is suspended.<sup>385)</sup> For the optional time

<sup>383)</sup> Although a **struct timespec** object describes times with nanosecond resolution, the available resolution is system dependent and can even be greater than 1 second.

<sup>384)</sup> Commonly, this reference point is the boot time of the execution environment or the start of the execution.

<sup>385)</sup> The execution environment can, for example, lack the ability to track physical time that elapsed during suspension in a low power consumption mode.

bases **TIME\_ACTIVE** and **TIME\_THREAD\_ACTIVE** the result is similar, but the call measures the amount of active processing time associated with the whole program invocation or with the calling thread, respectively.

#### Returns

- 4 If the **timespec\_get** function is successful it returns the nonzero value **base**; otherwise, it returns zero.

#### Recommended practice

- 5 It is recommended practice that timing results of calls to **timespec\_get** with **TIME\_ACTIVE**, if defined, and of calls to **clock** are as close to each other as their types, value ranges, and resolutions (obtained with **timespec\_getres** and **CLOCKS\_PER\_SEC**, respectively) allow. Because of its wider value range and improved indications on error, **timespec\_get** with time base **TIME\_ACTIVE** should be used instead of **clock** by new code whenever possible.

### 7.30.2.7 The **timespec\_getres** function

#### Synopsis

```
1 #include <time.h>
 int timespec_getres(struct timespec *ts, int base);
```

#### Description

- 2 If **ts** is non-null and **base** is supported by the **timespec\_get** function, the **timespec\_getres** function returns the resolution of the time provided by the **timespec\_get** function for **base** in the **timespec** structure pointed to by **ts**. For each supported **base**, multiple calls to the **timespec\_getres** function during the same program execution shall have identical results.

#### Returns

- 3 If the value **base** is supported by the **timespec\_get** function, the **timespec\_getres** function returns the nonzero value **base**; otherwise, it returns zero.

## 7.30.3 Time conversion functions

### 7.30.3.1 General

- 1 Functions with a **\_r** suffix place the result of the conversion into the buffer referred by **buf** and return that pointer. These functions and the function **strftime** shall not be subject to data races, unless the time or calendar state is changed in a multi-thread execution.<sup>386)</sup>
- 2 Functions **gmtime** and **localtime** are the same as their counterparts suffixed with **\_r**. In place of the parameter **buf**, they use a pointer to one or two broken-down time structures. Similarly, an array of **char** is commonly used by **asctime** and **ctime**. Execution of any of the functions that return a pointer to one of these objects may overwrite the information returned from any previous call to one of these functions that uses the same object. These functions are not reentrant and are not required to avoid data races with each other. Accessing the returned pointer after the thread that called the function that returned it has exited results in undefined behavior. The implementation shall behave as if no other library functions call these functions.

### 7.30.3.2 The **asctime** function

#### Synopsis

```
1 #include <time.h>
 [[deprecated]] char *asctime(const struct tm *timeptr);
```

#### Description

- 2 This function is obsolescent and should be avoided in new code.

<sup>386)</sup>This does not mean that these functions are forbidden to read global state that describes the time and calendar settings of the execution, such as the **LC\_TIME** locale or the implementation-defined specification of the local time zone. Only the setting of that state by **setlocale** or by means of implementation-defined functions can constitute races.

- 3 The **asctime** function converts the broken-down time in the structure pointed to by **timeptr** into a string in the form

```
Sun Sep 16 01:03:52 1973\n\0
```

using the equivalent of the following algorithm.

```
[[deprecated]] char *asctime(const struct tm *timeptr)
{
 static const char wday_name[7][3] = {
 "Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"
 };
 static const char mon_name[12][3] = {
 "Jan", "Feb", "Mar", "Apr", "May", "Jun",
 "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"
 };
 static char result[26];

 snprintf(result, 26, "%.3s %.3s%3d %.2d:%.2d:%.2d %d\n",
 wday_name[timeptr->tm_wday],
 mon_name[timeptr->tm_mon],
 timeptr->tm_mday, timeptr->tm_hour,
 timeptr->tm_min, timeptr->tm_sec,
 1900 + timeptr->tm_year);
 return result;
}
```

- 4 If any of the members of the broken-down time contain values that are outside their normal ranges,<sup>387)</sup> the behavior of the **asctime** function is undefined. Likewise, if the calculated year exceeds four digits or is less than the year 1000, the behavior is undefined.

#### Returns

- 5 The **asctime** function returns a pointer to the string.

#### 7.30.3.3 The **ctime** function

##### Synopsis

```
1 #include <time.h>
 [[deprecated]] char *ctime(const time_t *timer);
```

#### Description

- 2 This function is obsolescent and should be avoided in new code.
- 3 The **ctime** function converts the calendar time pointed to by **timer** to local time in the form of a string. They are equivalent to:

```
asctime(localtime(timer))
```

#### Returns

- 4 The **ctime** function returns the pointer returned by the **asctime** function with that broken-down time as argument.

Forward references: the **localtime** functions (7.30.3.5).

#### 7.30.3.4 The **gmtime** functions

##### Synopsis

```
1 #include <time.h>
```

<sup>387)</sup>See 7.30.1.

```
struct tm *gmtime(const time_t *timer);
struct tm *gmtime_r(const time_t *timer, struct tm *buf);
```

#### Description

- 2 The **gmtime** functions convert the calendar time pointed to by **timer** into a broken-down time, expressed as UTC.

#### Returns

- 3 The **gmtime** functions return a pointer to the broken-down time, or a null pointer if the specified time cannot be converted to UTC.

### 7.30.3.5 The **localtime** functions

#### Synopsis

```
1 #include <time.h>
 struct tm *localtime(const time_t *timer);
 struct tm *localtime_r(const time_t *timer, struct tm *buf);
```

#### Description

- 2 The **localtime** functions convert the calendar time pointed to by **timer** into a broken-down time, expressed as local time.

#### Returns

- 3 The **localtime** functions return a pointer to the broken-down time, or a null pointer if the specified time cannot be converted to local time.

### 7.30.3.6 The **strftime** function

#### Synopsis

```
1 #include <time.h>
 size_t strftime(char * restrict s, size_t maxsize, const char * restrict format,
 const struct tm * restrict timeptr);
```

#### Description

- 2 The **strftime** function places characters into the array pointed to by **s** as controlled by the string pointed to by **format**. The format shall be a multibyte character sequence, beginning and ending in its initial shift state. The **format** string consists of zero or more conversion specifiers and ordinary multibyte characters. A conversion specifier consists of a % character, possibly followed by an E or O modifier character (described later), followed by a character that determines the behavior of the conversion specifier. All ordinary multibyte characters (including the terminating null character) are copied unchanged into the array. If copying takes place between objects that overlap, the behavior is undefined. No more than **maxsize** characters are placed into the array.
- 3 Each conversion specifier shall be replaced by appropriate characters as described in the following list. The appropriate characters shall be determined using the **LC\_TIME** category of the current locale and by the values of zero or more members of the broken-down time structure pointed to by **timeptr**, as specified in brackets in the description. If any of the specified values is outside the normal range, the characters stored are unspecified.

%a is replaced by the locale's abbreviated weekday name. [**tm\_wday**]

%A is replaced by the locale's full weekday name. [**tm\_wday**]

%b is replaced by the locale's abbreviated month name. [**tm\_mon**]

%B is replaced by the locale's full month name. [**tm\_mon**]

%c is replaced by the locale's appropriate date and time representation. [all specified in 7.30.1]

%C is replaced by the year divided by 100 and truncated to an integer, as a decimal number (00–99). [**tm\_year**]

- %d is replaced by the day of the month as a decimal number (01–31). [**tm\_mday**]
- %D is equivalent to "%m/%d/%y". [**tm\_mon, tm\_mday, tm\_year**]
- %e is replaced by the day of the month as a decimal number (1–31); a single digit is preceded by a space. [**tm\_mday**]
- %F is equivalent to "%Y-%m-%d" (the ISO 8601 date format, when the year is between 1000 and 9999 inclusive). [**tm\_year, tm\_mon, tm\_mday**]
- %g is replaced by the last 2 digits of the week-based year (see further in this subclause) as a decimal number (00–99). [**tm\_year, tm\_wday, tm\_yday**]
- %G is replaced by the week-based year (see further in this subclause) as a decimal number (e.g. 1997). [**tm\_year, tm\_wday, tm\_yday**]
- %h is equivalent to "%b". [**tm\_mon**]
- %H is replaced by the hour (24-hour clock) as a decimal number (00–23). [**tm\_hour**]
- %I is replaced by the hour (12-hour clock) as a decimal number (01–12). [**tm\_hour**]
- %j is replaced by the day of the year as a decimal number (001–366). [**tm\_yday**]
- %m is replaced by the month as a decimal number (01–12). [**tm\_mon**]
- %M is replaced by the minute as a decimal number (00–59). [**tm\_min**]
- %n is replaced by a new-line character.
- %p is replaced by the locale's equivalent of the AM/PM designations associated with a 12-hour clock. [**tm\_hour**]
- %r is replaced by the locale's 12-hour clock time. [**tm\_hour, tm\_min, tm\_sec**]
- %R is equivalent to "%H:%M". [**tm\_hour, tm\_min**]
- %S is replaced by the second as a decimal number (00–60). [**tm\_sec**]
- %t is replaced by a horizontal-tab character.
- %T is equivalent to "%H:%M:%S" (the ISO 8601 time format). [**tm\_hour, tm\_min, tm\_sec**]
- %u is replaced by the ISO 8601 weekday as a decimal number (1–7), where Monday is 1. [**tm\_wday**]
- %U is replaced by the week number of the year (the first Sunday as the first day of week 1) as a decimal number (00–53). [**tm\_year, tm\_wday, tm\_yday**]
- %V is replaced by the ISO 8601 week number (see further in this subclause) as a decimal number (01–53). [**tm\_year, tm\_wday, tm\_yday**]
- %w is replaced by the weekday as a decimal number (0–6), where Sunday is 0. [**tm\_wday**]
- %W is replaced by the week number of the year (the first Monday as the first day of week 1) as a decimal number (00–53). [**tm\_year, tm\_wday, tm\_yday**]
- %x is replaced by the locale's appropriate date representation. [all specified in 7.30.1]
- %X is replaced by the locale's appropriate time representation. [all specified in 7.30.1]
- %y is replaced by the last 2 digits of the year as a decimal number (00–99). [**tm\_year**]
- %Y is replaced by the year as a decimal number (e.g. 1997). [**tm\_year**]
- %z is replaced by the offset from UTC in the ISO 8601 format "-0430" (meaning 4 hours 30 minutes behind UTC, west of Greenwich), or by no characters if no time zone is determinable. Behavior is implementation-defined if the broken-down time structure does not have a value that could be returned by **localtime** or **gmtime**. [all members, including any non-standard additional members]
- %Z is replaced by the locale's time zone name or abbreviation, or by no characters if no time zone is determinable. Behavior is implementation-defined if the broken-down time structure does not have a value that could be returned by **localtime** or **gmtime**. [all members, including any non-standard additional members]
- %% is replaced by %.

- 4 Some conversion specifiers can be modified by the inclusion of an E or O modifier character to indicate an alternative format or specification. If the alternative format or specification does not exist for the current locale, the modifier is ignored.

|     |                                                                                                                                                                                               |
|-----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| %Ec | is replaced by the locale's alternative date and time representation.                                                                                                                         |
| %EC | is replaced by the name of the base year (period) in the locale's alternative representation.                                                                                                 |
| %Ex | is replaced by the locale's alternative date representation.                                                                                                                                  |
| %EX | is replaced by the locale's alternative time representation.                                                                                                                                  |
| %Ey | is replaced by the offset from %EC (year only) in the locale's alternative representation.                                                                                                    |
| %EY | is replaced by the locale's full alternative year representation.                                                                                                                             |
| %Ob | is replaced by the locale's abbreviated alternative month name.                                                                                                                               |
| %OB | is replaced by the locale's alternative appropriate full month name.                                                                                                                          |
| %Od | is replaced by the day of the month, using the locale's alternative numeric symbols (filled as needed with leading zeros, or with leading spaces if there is no alternative symbol for zero). |
| %Oe | is replaced by the day of the month, using the locale's alternative numeric symbols (filled as needed with leading spaces).                                                                   |
| %OH | is replaced by the hour (24-hour clock), using the locale's alternative numeric symbols.                                                                                                      |
| %OI | is replaced by the hour (12-hour clock), using the locale's alternative numeric symbols.                                                                                                      |
| %Om | is replaced by the month, using the locale's alternative numeric symbols.                                                                                                                     |
| %OM | is replaced by the minutes, using the locale's alternative numeric symbols.                                                                                                                   |
| %OS | is replaced by the seconds, using the locale's alternative numeric symbols.                                                                                                                   |
| %Ou | is replaced by the ISO 8601 weekday as a number in the locale's alternative representation, where Monday is 1.                                                                                |
| %OU | is replaced by the week number, using the locale's alternative numeric symbols.                                                                                                               |
| %OV | is replaced by the ISO 8601 week number, using the locale's alternative numeric symbols.                                                                                                      |
| %Ow | is replaced by the weekday as a number, using the locale's alternative numeric symbols.                                                                                                       |
| %OW | is replaced by the week number of the year, using the locale's alternative numeric symbols.                                                                                                   |
| %Oy | is replaced by the last 2 digits of the year, using the locale's alternative numeric symbols.                                                                                                 |

- 5 %g, %G, and %V give values according to the ISO 8601 week-based year. In this system, weeks begin on a Monday and week 1 of the year is the week that includes January 4th, which is also the week that includes the first Thursday of the year, and is also the first week that contains at least four days in the year. If the first Monday of January is the 2nd, 3rd, or 4th, the preceding days are part of the last week of the preceding year; thus, for Saturday 2nd January 1999, %G is replaced by 1998 and %V is replaced by 53. If December 29th, 30th, or 31st is a Monday, it and any following days are part of week 1 of the following year. Thus, for Tuesday 30th December 1997, %G is replaced by 1998 and %V is replaced by 01.
- 6 If a conversion specifier is not one of the ones previously specified, the behavior is undefined.
- 7 In the "C" locale, the E and O modifiers are ignored and the replacement strings for the following specifiers are:

|    |                                                |
|----|------------------------------------------------|
| %a | the first three characters of %A.              |
| %A | one of "Sunday", "Monday", ..., "Saturday".    |
| %b | the first three characters of %B.              |
| %B | one of "January", "February", ..., "December". |
| %c | equivalent to "%a %b %e %T %Y".                |
| %p | one of "AM" or "PM".                           |



%r equivalent to “%I:%M:%S %p”.

%x equivalent to “%m/%d/%y”.

%X equivalent to %T.

%Z implementation-defined.

#### Returns

- 8 If the total number of resulting characters including the terminating null character is not more than `maxsize`, the `strftime` function returns the number of characters placed into the array pointed to by `s` not including the terminating null character. Otherwise, zero is returned and the members of the array have an indeterminate representation.



## 7.31 Unicode utilities <uchar.h>

### 7.31.1 General

- 1 The header <uchar.h> declares one macro, a few types, and several functions for manipulating Unicode characters.
- 2 The macro

```
__STDC_VERSION_UCHAR_H__
```

is an integer constant expression with a value equivalent to 202311L.

- 3 The types declared are **mbstate\_t** (described in 7.32.1) and **size\_t** (described in 7.21);

```
char8_t
```

which is an unsigned integer type used for 8-bit characters and is the same type as **unsigned char**;

```
char16_t
```

which is an unsigned integer type used for 16-bit characters and is the same type as **uint\_least16\_t** (described in 7.22.2.3); and

```
char32_t
```

which is an unsigned integer type used for 32-bit characters and is the same type as **uint\_least32\_t** (also described in 7.22.2.3).

### 7.31.2 Restartable multibyte/wide character conversion functions

#### 7.31.2.1 General

- 1 These functions have a parameter, **ps**, of type pointer to **mbstate\_t** that points to an object that can completely describe the current conversion state of the associated multibyte character sequence, which the functions alter as necessary. If **ps** is a null pointer, each function uses its own internal **mbstate\_t** object instead, which is initialized prior to the first call to the function to the initial conversion state; the functions are not required to avoid data races with other calls to the same function in this case. It is implementation-defined whether the internal **mbstate\_t** object has thread storage duration; if it has thread storage duration, it is initialized to the initial conversion state prior to the first call to the function on the new thread. The implementation behaves as if no library function calls these functions with a null pointer for **ps**.
- 2 When used in the functions in this subclause, the encoding of **char8\_t**, **char16\_t**, and **char32\_t** objects, and sequences of such objects, is UTF-8, UTF-16, and UTF-32, respectively. Similarly, the encoding of **char** and **wchar\_t**, and sequences of such objects, is the execution and wide execution encodings (6.2.9), respectively.

#### 7.31.2.2 The **mbrtoc8** function

##### Synopsis

- 1

```
#include <uchar.h>
size_t mbrtoc8(char8_t * restrict pc8, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
```

##### Description

- 2 If **s** is a null pointer, the **mbrtoc8** function is equivalent to the call:

```
mbrtoc8(nullptr, "", 1, ps)
```

In this case, the values of the parameters **pc8** and **n** are ignored.

- 3 If `s` is not a null pointer, the `mbrtoc8` function function inspects at most `n` bytes beginning with the byte pointed to by `s` to determine the number of bytes needed to complete the next multibyte character (including any shift sequences). If the function determines that the next multibyte character is complete and valid, it determines the values of the corresponding characters and then, if `pc8` is not a null pointer, stores the value of the first (or only) such character in the object pointed to by `pc8`. Subsequent calls will store successive characters without consuming any additional input until all the characters have been stored. If the corresponding character is the null character, the resulting state described is the initial conversion state.

#### Returns

- 4 The `mbrtoc8` function returns the first of the following that applies (given the current conversion state):

0 if the next `n` or fewer bytes complete the multibyte character that corresponds to the null character (which is the value stored).

*between 1 and n inclusive* if the next `n` or fewer bytes complete a valid multibyte character (which is the value stored); the value returned is the number of bytes that complete the multibyte character.

(`size_t`) (-3) if the next character resulting from a previous call has been stored (no bytes from the input have been consumed by this call).

(`size_t`) (-2) if the next `n` bytes contribute to an incomplete (but potentially valid) multibyte character, and all `n` bytes have been processed (no value is stored).<sup>388)</sup>

(`size_t`) (-1) if an encoding error occurs, in which case the next `n` or fewer bytes do not contribute to a complete and valid multibyte character (no value is stored); the value of the macro `EILSEQ` is stored in `errno`, and the conversion state is unspecified.

#### 7.31.2.3 The `c8rtomb` function

##### Synopsis

```
1 #include <uchar.h>
 size_t c8rtomb(char * restrict s, char8_t c8, mbstate_t * restrict ps);
```

##### Description

- 2 If `s` is a null pointer, the `c8rtomb` function is equivalent to the call

```
c8rtomb(buf, u8'\0', ps)
```

where `buf` is an internal buffer.

- 3 If `s` is not a null pointer, the `c8rtomb` function determines the number of bytes needed to represent the multibyte character that corresponds to the character given or completed by `c8` (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by `s`, or stores nothing if `c8` does not represent a complete character. At most `MB_CUR_MAX` bytes are stored. If `c8` is a null character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state; the resulting state described is the initial conversion state.

#### Returns

- 4 The `c8rtomb` function returns the number of bytes stored in the array object (including any shift sequences). When `c8` is not a valid character, an encoding error occurs: the function stores the value of the macro `EILSEQ` in `errno` and returns (`size_t`) (-1); the conversion state is unspecified.

#### 7.31.2.4 The `mbrtoc16` function

<sup>388)</sup>When `n` has at least the value of the `MB_CUR_MAX` macro, this case can only occur if `s` points at a sequence of redundant shift sequences (for implementations with state-dependent encodings).

## Synopsis

```
1 #include <uchar.h>
 size_t mbrtoc16(char16_t * restrict pc16, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
```

## Description

- 2 If *s* is a null pointer, the **mbrtoc16** function is equivalent to the call:

```
mbrtoc16(nullptr, "", 1, ps)
```

In this case, the values of the parameters *pc16* and *n* are ignored.

- 3 If *s* is not a null pointer, the **mbrtoc16** function inspects at most *n* bytes beginning with the byte pointed to by *s* to determine the number of bytes needed to complete the next multibyte character (including any shift sequences). If the function determines that the next multibyte character is complete and valid, it determines the values of the corresponding wide characters and then, if *pc16* is not a null pointer, stores the value of the first (or only) such character in the object pointed to by *pc16*. Subsequent calls will store successive wide characters without consuming any additional input until all the characters have been stored. If the corresponding wide character is the null wide character, the resulting state described is the initial conversion state.

## Returns

- 4 The **mbrtoc16** function returns the first of the following that applies (given the current conversion state):

0 if the next *n* or fewer bytes complete the multibyte character that corresponds to the null wide character (which is the value stored).

*between 1 and n inclusive* if the next *n* or fewer bytes complete a valid multibyte character (which is the value stored); the value returned is the number of bytes that complete the multibyte character.

(*size\_t*) (-3) if the next character resulting from a previous call has been stored (no bytes from the input have been consumed by this call).

(*size\_t*) (-2) if the next *n* bytes contribute to an incomplete (but potentially valid) multibyte character, and all *n* bytes have been processed (no value is stored).<sup>389)</sup>

(*size\_t*) (-1) if an encoding error occurs, in which case the next *n* or fewer bytes do not contribute to a complete and valid multibyte character (no value is stored); the value of the macro **EILSEQ** is stored in **errno**, and the conversion state is unspecified.

### 7.31.2.5 The **c16rtomb** function

#### Synopsis

```
1 #include <uchar.h>
 size_t c16rtomb(char * restrict s, char16_t c16, mbstate_t * restrict ps);
```

## Description

- 2 If *s* is a null pointer, the **c16rtomb** function is equivalent to the call

```
c16rtomb(buf, u'\0', ps)
```

where *buf* is an internal buffer.

<sup>389)</sup>When *n* has at least the value of the **MB\_CUR\_MAX** macro, this case can only occur if *s* points at a sequence of redundant shift sequences (for implementations with state-dependent encodings).

- 3 If `s` is not a null pointer, the `c16rtomb` function determines the number of bytes needed to represent the multibyte character that corresponds to the wide character given or completed by `c16` (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by `s`, or stores nothing if `c16` does not represent a complete character. At most `MB_CUR_MAX` bytes are stored. If `c16` is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state; the resulting state described is the initial conversion state.

#### Returns

- 4 The `c16rtomb` function returns the number of bytes stored in the array object (including any shift sequences). When `c16` is not a valid wide character, an encoding error occurs: the function stores the value of the macro `EILSEQ` in `errno` and returns `(size_t) (-1)`; the conversion state is unspecified.

#### 7.31.2.6 The `mbrtoc32` function

##### Synopsis

```
1 #include <uchar.h>
 size_t mbrtoc32(char32_t * restrict pc32, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
```

##### Description

- 2 If `s` is a null pointer, the `mbrtoc32` function is equivalent to the call:

```
mbrtoc32(nullptr, "", 1, ps)
```

In this case, the values of the parameters `pc32` and `n` are ignored.

- 3 If `s` is not a null pointer, the `mbrtoc32` function inspects at most `n` bytes beginning with the byte pointed to by `s` to determine the number of bytes needed to complete the next multibyte character (including any shift sequences). If the function determines that the next multibyte character is complete and valid, it determines the values of the corresponding wide characters and then, if `pc32` is not a null pointer, stores the value of the first (or only) such character in the object pointed to by `pc32`. Subsequent calls will store successive wide characters without consuming any additional input until all the characters have been stored. If the corresponding wide character is the null wide character, the resulting state described is the initial conversion state.

#### Returns

- 4 The `mbrtoc32` function returns the first of the following that applies (given the current conversion state):

0 if the next `n` or fewer bytes complete the multibyte character that corresponds to the null wide character (which is the value stored).

*between 1 and n inclusive* if the next `n` or fewer bytes complete a valid multibyte character (which is the value stored); the value returned is the number of bytes that complete the multibyte character.

`(size_t) (-3)` if the next character resulting from a previous call has been stored (no bytes from the input have been consumed by this call).

`(size_t) (-2)` if the next `n` bytes contribute to an incomplete (but potentially valid) multibyte character, and all `n` bytes have been processed (no value is stored).<sup>390)</sup>

`(size_t) (-1)` if an encoding error occurs, in which case the next `n` or fewer bytes do not contribute to a complete and valid multibyte character (no value is stored); the value of the macro `EILSEQ` is stored in `errno`, and the conversion state is unspecified.

<sup>390)</sup>When `n` has at least the value of the `MB_CUR_MAX` macro, this case can only occur if `s` points at a sequence of redundant shift sequences (for implementations with state-dependent encodings).

7.31.2.7 The **c32rtomb** function

## Synopsis

```
1 #include <uchar.h>
 size_t c32rtomb(char * restrict s, char32_t c32, mbstate_t * restrict ps);
```

## Description

- 2 If **s** is a null pointer, the **c32rtomb** function is equivalent to the call

```
c32rtomb(buf, U'\0', ps)
```

where **buf** is an internal buffer.

- 3 If **s** is not a null pointer, the **c32rtomb** function determines the number of bytes needed to represent the multibyte character that corresponds to the wide character given by **c32** (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by **s**. At most **MB\_CUR\_MAX** bytes are stored. If **c32** is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state; the resulting state described is the initial conversion state.

## Returns

- 4 The **c32rtomb** function returns the number of bytes stored in the array object (including any shift sequences). When **c32** is not a valid wide character, an encoding error occurs: the function stores the value of the macro **EILSEQ** in **errno** and returns **(size\_t) (-1)**; the conversion state is unspecified.

## 7.32 Extended multibyte and wide character utilities <wchar.h>

### 7.32.1 Introduction

- 1 The header <wchar.h> defines several macros, and declares four data types, one tag, and many functions.<sup>391)</sup>
- 2 The macro

```
__STDC_VERSION_WCHAR_H__
```

is an integer constant expression with a value equivalent to 202ymmL.

- 3 The types declared are **wchar\_t** and **size\_t** (both described in 7.21);

```
mbstate_t
```

which is a complete object type other than an array type that can hold the conversion state information necessary to convert between sequences of multibyte characters and wide characters;

```
wint_t
```

which is an integer type unchanged by default argument promotions that can hold any value corresponding to members of the extended character set, as well as at least one value that does not correspond to any member of the extended character set (see subsequent **WEOF** description);<sup>392)</sup> and

```
struct tm
```

which is declared as an incomplete structure type (the contents are described in 7.30.1).

- 4 The macros defined are **NULL** (described in 7.21); **WCHAR\_MIN**, **WCHAR\_MAX**, and **WCHAR\_WIDTH** (described in 7.22);

```
WCHAR_UTF8
WCHAR_UTF16
WCHAR_UTF32
```

which expand to an expression of signed or unsigned integer type (that is potentially not an integer constant expression) whose value is nonzero if:

- the wide execution encoding (6.2.9) is capable of representing every character in the required Unicode set;
- the width of **wchar\_t** is at least 8, 16, or 32 for UTF-8, UTF-16, or UTF-32, respectively;
- and, the values of a sequence of **wchar\_t** objects consumed and produced by related character functions have a values consistent with a sequence of code units of the UTF-8, UTF-16, or UTF-32 encodings, respectively;

```
MB_UTF8
MB_UTF16
MB_UTF32
```

which expand to an expression of signed or unsigned integer type (that is potentially not an integer constant expression) whose value is nonzero if:

- the execution encoding (6.2.9) is capable of representing every character in the required Unicode set;

<sup>391)</sup>See “future library directions” (7.34.21).

<sup>392)</sup>**wchar\_t** and **wint\_t** can be the same integer type.

- the width of **char** is at least 8, 16, or 32 for UTF-8, UTF-16, or UTF-32, respectively;
- and, the values of a sequence of **char** objects (treated a sequence of unsigned char objects) consumed and produced by related character functions have a values consistent with a sequence of code units of the UTF-8, UTF-16, or UTF-32 encodings, respectively;

and

**WEOF**

which expands to a constant expression of type **wint\_t** whose value does not correspond to any member of the extended character set.<sup>393)</sup> It is accepted (and returned) by several functions in this subclause to indicate *end-of-file*, that is, no more input from a stream. It is also used as a wide character value that does not correspond to any member of the extended character set.

- 5 The functions declared are grouped as follows:
  - Functions that perform input and output of wide characters, or multibyte characters, or both;
  - Functions that provide wide string numeric conversion;
  - Functions that perform general wide string manipulation;
  - Functions for wide string date and time conversion; and
  - Functions that provide extended capabilities for conversion between multibyte and wide character sequences.
- 6 Arguments to the functions in this subclause may point to arrays containing **wchar\_t** values that do not correspond to members of the extended character set. Such values shall be processed according to the specified semantics, except that it is unspecified whether an encoding error occurs if such a value appears in the format string for a function in 7.32.2 or 7.32.5 and the specified semantics do not require that value to be processed by **wcrtomb**.
- 7 Unless explicitly stated otherwise, if the execution of a function described in this subclause causes copying to take place between objects that overlap, the behavior is undefined.

## 7.32.2 Formatted wide character input/output functions

### 7.32.2.1 General

- 1 The formatted wide character input/output functions shall behave as if there is a sequence point after the actions associated with each specifier.<sup>394)</sup>

### 7.32.2.2 The **fwprintf** function

#### Synopsis

```
1 #include <stdio.h>
 #include <wchar.h>
 int fwprintf(FILE * restrict stream, const wchar_t * restrict format, ...);
```

#### Description

- 2 The **fwprintf** function writes output to the stream pointed to by **stream**, under control of the wide string pointed to by **format** that specifies how subsequent arguments are converted for output. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored. The **fwprintf** function returns when the end of the format string is encountered.
- 3 The format is composed of zero or more directives: ordinary wide characters (not %), which are copied unchanged to the output stream; and conversion specifications, each of which results in

<sup>393)</sup>The value of the macro **WEOF** can differ from that of **EOF** and the value can be positive.

<sup>394)</sup>The **fwprintf** functions perform writes to memory for the **%n** specifier.



fetching zero or more subsequent arguments, converting them, if applicable, according to the corresponding conversion specifier, and then writing the result to the output stream.

- 4 Each conversion specification is introduced by the wide character `%`. After the `%`, the following appear in sequence:
  - Zero or more *flags* (in any order) that modify the meaning of the conversion specification.
  - An optional minimum *field width*. If the converted value has fewer wide characters than the field width, it is padded with spaces (by default) on the left (or right, if the left adjustment flag, described later, has been given) to the field width. The field width takes the form of an asterisk `*` (described later) or a nonnegative decimal integer.<sup>395)</sup>
  - An optional *precision* that gives the minimum number of digits to appear for the `b`, `B`, `d`, `i`, `o`, `u`, `x`, and `X` conversions, the number of digits to appear after the decimal-point wide character for `a`, `A`, `e`, `E`, `f`, and `F` conversions, the maximum number of significant digits for the `g` and `G` conversions, or the maximum number of wide characters to be written for `s` conversions. The precision takes the form of a period (`.`) followed either by an asterisk `*` (described later) or by an optional nonnegative decimal integer; if only the period is specified, the precision is taken as zero. If a precision appears with any other conversion specifier, the behavior is undefined.
  - An optional *length modifier* that specifies the size of the argument.
  - A *conversion specifier* wide character that specifies the type of conversion to be applied.
- 5 As noted previously, a field width, or precision, or both, may be indicated by an asterisk. In this case, an `int` argument supplies the field width or precision. The arguments specifying field width, or precision, or both, shall appear (in that order) before the argument (if any) to be converted. A negative field width argument is taken as a `-` flag followed by a positive field width. A negative precision argument is taken as if the precision were omitted.
- 6 The flag wide characters and their meanings are:
  - The result of the conversion is left-justified within the field. (It is right-justified if this flag is not specified.)
  - + The result of a signed conversion always begins with a plus or minus sign. (It begins with a sign only when a value with a negative sign is converted if this flag is not specified.)<sup>396)</sup>
  - space* If the first wide character of a signed conversion is not a sign, or if a signed conversion results in no wide characters, a space is prefixed to the result. If the *space* and `+` flags both appear, the *space* flag is ignored.
  - # The result is converted to an “alternative form”. For `o` conversion, it increases the precision, if and only if necessary, to force the first digit of the result to be a zero (if the value and precision are both 0, a single 0 is printed). For `b` conversion, a nonzero result has `0b` prefixed to it. For the optional `B` conversion as described later in this subclause, a nonzero result has `0B` prefixed to it. For `x` (or `X`) conversion, a nonzero result has `0x` (or `0X`) prefixed to it. For `a`, `A`, `e`, `E`, `f`, `F`, `g`, and `G` conversions, the result of converting a floating-point number always contains a decimal-point wide character, even if no digits follow it. (Normally, a decimal-point wide character appears in the result of these conversions only if a digit follows it.) For `g` and `G` conversions, trailing zeros are *not* removed from the result. For other conversions, the behavior is undefined.
  - 0 For `b`, `B`, `d`, `i`, `o`, `u`, `x`, `X`, `a`, `A`, `e`, `E`, `f`, `F`, `g`, and `G` conversions, leading zeros (following any indication of sign or base) are used to pad to the field width rather than performing space padding, except when converting an infinity or NaN. If the `0` and `-` flags both appear, the `0` flag is ignored. For `b`, `B`, `d`, `i`, `o`, `u`, `x`, and `X` conversions, if a precision is specified, the `0` flag is ignored. For other conversions, the behavior is undefined.

<sup>395)</sup> `0` is taken as a flag, not as the beginning of a field width.

<sup>396)</sup> The results of all floating conversions of a negative zero, and of negative values that round to zero, include a minus sign.



## 7 The length modifiers and their meanings are:

- hh** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **signed char** or **unsigned char** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **signed char** or **unsigned char** before printing); or that a following **n** conversion specifier applies to a pointer to a **signed char** argument.
- h** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **short int** or **unsigned short int** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **short int** or **unsigned short int** before printing); or that a following **n** conversion specifier applies to a pointer to a **short int** argument.
- l** (ell) Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **long int** or **unsigned long int** argument; that a following **n** conversion specifier applies to a pointer to a **long int** argument; that a following **c** conversion specifier applies to a **wint\_t** argument; that a following **s** conversion specifier applies to a pointer to a **wchar\_t** argument; or has no effect on a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier.
- ll** (ell-ell) Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **long long int** or **unsigned long long int** argument; or that a following **n** conversion specifier applies to a pointer to a **long long int** argument.
- j** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to an **intmax\_t** or **uintmax\_t** argument; or that a following **n** conversion specifier applies to a pointer to an **intmax\_t** argument.
- z** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **size\_t** or the corresponding signed integer type argument; or that a following **n** conversion specifier applies to a pointer to a signed integer type corresponding to **size\_t** argument.
- t** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a **ptrdiff\_t** or the corresponding unsigned integer type argument; or that a following **n** conversion specifier applies to a pointer to a **ptrdiff\_t** argument.
- wN** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to an integer argument with a specific width where *N* is a positive decimal integer with no leading zeros (the argument will have been promoted according to the integer promotions, but its value shall be converted to the unpromoted type); or that a following **n** conversion specifier applies to a pointer to an integer type argument with a width of *N* bits. All minimum-width integer types (7.22.2.3) and exact-width integer types (7.22.2.2) defined in the header `<stdint.h>` shall be supported. Other supported values of *N* are implementation-defined.
- wfN** Specifies that a following **b**, **B**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a fastest minimum-width integer argument with a specific width where *N* is a positive decimal integer with no leading zeros (the argument will have been promoted according to the integer promotions, but its value shall be converted to the unpromoted type); or that a following **n** conversion specifier applies to a pointer to a fastest minimum-width integer type argument with a width of *N* bits. All fastest minimum-width integer types (7.22.2.4) defined in the header `<stdint.h>` shall be supported. Other supported values of *N* are implementation-defined.
- L** Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **long double** argument.
- H** Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **\_Decimal32** argument.

- D** Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **\_Decimal64** argument.
- DD** Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, or **G** conversion specifier applies to a **\_Decimal128** argument.

If a length modifier appears with any conversion specifier other than as specified previously, the behavior is undefined.

8 The conversion specifiers and their meanings are:

- d, i** The **int** argument is converted to signed decimal in the style *[-]dddd*. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no wide characters.
- b, B, o, u, x, X** The **unsigned int** argument is converted to unsigned binary (**b** or **B**), unsigned octal (**o**), unsigned decimal (**u**), or unsigned hexadecimal notation (**x** or **X**) in the style *dddd*; the letters **abcdef** are used for **x** conversion and the letters **ABCDEF** for **X** conversion. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no wide characters. The specifier **B** is optional and provides the same functionality as **b**, except for the **#** flag as previously specified. The **PRIB** macros from `<inttypes.h>` shall only be defined if the implementation follows the specification as given here.
- f, F** A **double** argument representing a floating-point number is converted to decimal notation in the style *[-]ddd.ddd*, where the number of digits after the decimal-point wide character is equal to the precision specification. If the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point wide character appears. If a decimal-point wide character appears, at least one digit appears before it. The value is rounded to the appropriate number of digits.
- A **double** argument representing an infinity is converted in one of the styles *[-]inf* or *[-]infinity* — which style is implementation-defined. A **double** argument representing a NaN is converted in one of the styles *[-]nan* or *[-]nan(*n-wchar-sequence*)* — which style, and the meaning of any *n-wchar-sequence*, is implementation-defined. The **F** conversion specifier produces **INF**, **INFINITY**, or **NAN** instead of **inf**, **infinity**, or **nan**, respectively.<sup>397)</sup>
- e, E** A **double** argument representing a floating-point number is converted in the style *[-]d.ddde±dd*, where there is one digit (which is nonzero if the argument is nonzero) before the decimal-point wide character and the number of digits after it is equal to the precision; if the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point wide character appears. The value is rounded to the appropriate number of digits. The **E** conversion specifier produces a number with **E** instead of **e** introducing the exponent. The exponent always contains at least two digits, and only as many more digits as necessary to represent the exponent. If the value is zero, the exponent is zero.
- A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.
- g, G** A **double** argument representing a floating-point number is converted in style **f** or **e** (or in style **F** or **E** in the case of a **G** conversion specifier), depending on the value converted and the precision. Let *P* equal the precision if nonzero, 6 if the precision is omitted, or 1 if the precision is zero. Then, if a conversion with style **E** would have an exponent of *X*:
- if  $P > X \geq -4$ , the conversion is with style **f** (or **F**) and precision  $P - (X + 1)$ .
- otherwise, the conversion is with style **e** (or **E**) and precision  $P - 1$ .

<sup>397)</sup>When applied to infinite and NaN values, the **-**, **+**, and *space* flag wide characters have their usual meaning; the **#** and **0** flag wide characters have no effect.

Finally, unless the **#** flag is used, any trailing zeros are removed from the fractional portion of the result and the decimal-point wide character is removed if there is no fractional portion remaining.

A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.

**a, A** A **double** argument representing a floating-point number is converted in the style `[-]0xh.hhhhp±d`, where there is one hexadecimal digit (which is nonzero if the argument is a normalized floating-point number and is otherwise unspecified) before the decimal-point wide character<sup>398)</sup> and the number of hexadecimal digits after it is equal to the precision; if the precision is missing and **FLT\_RADIX** is a power of 2, then the precision is sufficient for an exact representation of the value; if the precision is missing and **FLT\_RADIX** is not a power of 2, then the precision is sufficient to distinguish<sup>399)</sup> values of type **double**, except that trailing zeros may be omitted; if the precision is zero and the **#** flag is not specified, no decimal-point wide character appears. The letters **abcdef** are used for **a** conversion and the letters **ABCDEF** for **A** conversion. The **A** conversion specifier produces a number with **X** and **P** instead of **x** and **p**. The exponent always contains at least one digit, and only as many more digits as necessary to represent the decimal exponent of 2. If the value is zero, the exponent is zero.

A **double** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.

If an **H**, **D**, or **DD** modifier is present and the precision is missing, then for a decimal floating type argument represented by a triple of integers  $(s, c, q)$ , where  $n$  is the number of significant digits in the coefficient  $c$ ,

- if  $-(n + 5) \leq q \leq 0$ , use style **f** (or style **F** in the case of an **A** conversion specifier) with formatting precision equal to  $-q$ ,
- otherwise, use style **e** (or style **E** in the case of an **A** conversion specifier) with formatting precision equal to  $n - 1$ , with the exceptions that if  $c = 0$  then the digit-sequence in the exponent-part shall have the value  $q$  (rather than 0), and that the exponent is always expressed with the minimum number of digits required to represent its value (the exponent never contains a leading zero).

If the precision  $P$  is present (in the conversion specification) and is zero or at least as large as the precision  $p$  (5.3.5.3.3) of the decimal floating type, the conversion is as if the precision were missing. If the precision  $P$  is present (and nonzero) and less than the precision  $p$  of the decimal floating type, the conversion first obtains an intermediate result as follows, where  $n$  is the number of significant digits in the coefficient:

- If  $n \leq P$ , set the intermediate result to the input.
- If  $n > P$ , round the input value, according to the current rounding direction for decimal floating-point operations, to  $P$  decimal digits, with unbounded exponent

<sup>398)</sup> Binary implementations can choose the hexadecimal digit to the left of the decimal-point wide character so that subsequent digits align to nibble (4-bit) boundaries. This implementation choice affects numerical values printed with a precision  $P$  that is insufficient to represent all values exactly. Implementations with different conventions about the most significant hexadecimal digit will round at different places, affecting the numerical value of the hexadecimal result. For example, possible printed output for the code

```
#include <stdio.h>
/* ... */
double x = 123.0;
printf("%.1a", x);
```

include `"0x1.fp+6"` and `"0xf.6p+3"` whose numerical values are 124 and 123, respectively. Portable code seeking identical numerical results on different platforms should avoid precisions  $P$  that require rounding.

<sup>399)</sup> The formatting precision  $P$  is sufficient to distinguish values of the source type if  $16^P > b^p$  where  $b$  (not a power of 2) and  $p$  are the base and precision of the source type (5.3.5.3.3). A smaller  $P$  potentially suffices depending on the implementation's scheme for determining the digit to the left of the decimal-point wide character.

range, representing the result with a  $P$ -digit integer coefficient when in the form  $(s, c, q)$ .

Convert the intermediate result in the manner described previously for the case where the precision is missing.

- c** If no **l** length modifier is present, the **int** argument is converted to a wide character as if by calling **btowc** and the resulting wide character is written.  
If an **l** length modifier is present, the **wint\_t** argument is converted to **wchar\_t** and written.
- s** If no **l** length modifier is present, the argument shall be a pointer to storage of character type containing a multibyte character sequence beginning in the initial shift state. Characters from the storage are converted as if by repeated calls to the **mbrtowc** function, with the conversion state described by an **mbstate\_t** object initialized to zero before the first multibyte character is converted, and written up to (but not including) the terminating null wide character. If the precision is specified, no more than that many wide characters are written. If the precision is not specified or is greater than the size of the converted storage, the converted storage shall contain a null wide character.  
If an **l** length modifier is present, the argument shall be a pointer to storage of **wchar\_t** type. Wide characters from the storage are written up to (but not including) a terminating null wide character. If the precision is specified, no more than that many wide characters are written. If the precision is not specified or is greater than the size of the array, the storage shall contain a null wide character.
- p** The argument shall be a pointer to **void** or a pointer to a character type. The value of the pointer is converted to a sequence of printing wide characters, in an implementation-defined manner.
- n** The argument shall be a pointer to signed integer whose type is specified by the length modifier, if any, for the conversion specification, or shall be **int** if no length modifier is specified for the conversion specification. The number of wide characters written to the output stream so far by this call to **fwprintf** is stored into the integer object pointed to by the argument. No argument is converted, but one is consumed. If the conversion specification includes any flags, a field width, or a precision, the behavior is undefined.
- %** A % wide character is written. No argument is converted. The complete conversion specification shall be **%%**.

- 9 If a conversion specification is invalid, the behavior is undefined.<sup>400)</sup> **fwprintf** shall behave as if it uses **va\_arg** with a type argument naming the type resulting from applying the default argument promotions to the type corresponding to the conversion specification and then converting the result of the **va\_arg** expansion to the type corresponding to the conversion specification.<sup>401)</sup>
- 10 In no case does a nonexistent or small field width cause truncation of a field; if the result of a conversion is wider than the field width, the field is expanded to contain the conversion result.
- 11 For **a** and **A** conversions, if **FLT\_RADIX** is a power of 2, the value is correctly rounded to a hexadecimal floating number with the given precision.

#### Recommended practice

- 12 For **a** and **A** conversions, if **FLT\_RADIX** is not a power of 2 and the result is not exactly representable in the given precision, the result should be one of the two adjacent numbers in hexadecimal floating style with the given precision, with the extra stipulation that the error should have a correct sign for the current rounding direction.
- 13 For **e**, **E**, **f**, **F**, **g**, and **G** conversions, if the number of significant decimal digits is at most the maximum value  $M$  of the **T\_DECIMAL\_DIG** macros (defined in `<float.h>`), then the result should be correctly

<sup>400)</sup>See “future library directions” (7.34.21).

<sup>401)</sup>The behavior is undefined when the types differ as specified for **va\_arg** 7.16.2.2.

rounded.<sup>402)</sup> If the number of significant decimal digits is more than  $M$  but the source value is exactly representable with  $M$  digits, then the result should be an exact representation with trailing zeros. Otherwise, the source value is bounded by two adjacent decimal strings  $L < U$ , both having  $M$  significant digits; the value of the resultant decimal string  $D$  should satisfy  $L \leq D \leq U$ , with the extra stipulations that the error should have a correct sign for the current rounding direction and that increasing the number of decimal digits for the result should not decrease the accuracy of the result.

- 14 The uppercase **B** format specifier is made optional by the previous description, because it used to be available for extensions in previous versions of this document. Implementations that did not use an uppercase **B** as their own extension before are encouraged to implement it as previously described.

#### Returns

- 15 The **fwprintf** function returns the number of wide characters transmitted, or a negative value if an output or encoding error occurred or if the implementation does not support a specified width length modifier.

#### Environmental limits

- 16 The number of wide characters that can be produced by any single conversion shall be at least 4095.
- 17 EXAMPLE To print a date and time in the form “Sunday, July 3, 10:02” followed by  $\pi$  to five decimal places:

```
#include <math.h>
#include <stdio.h>
#include <wchar.h>
/* ... */
wchar_t *weekday, *month; // pointers to wide strings
int day, hour, min;
fwprintf(stdout, L"%ls, %ls %d, %.2d:%.2d\n",
 weekday, month, day, hour, min);
fwprintf(stdout, L"pi = %.5f\n", 4 * atan(1.0));
```

Forward references: the **btowc** function (7.32.6.2.1), the **mbrtowc** function (7.32.6.4.3).

### 7.32.2.3 The **fwscanf** function

#### Synopsis

```
1 #include <stdio.h>
 #include <wchar.h>
 int fwscanf(FILE * restrict stream, const wchar_t * restrict format, ...);
```

#### Description

- 2 The **fwscanf** function reads input from the stream pointed to by **stream**, under control of the wide string pointed to by **format** that specifies the admissible input sequences and how they are to be converted for assignment, using subsequent arguments as pointers to the objects to receive the converted input. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored.
- 3 The format is composed of zero or more directives: one or more white-space wide characters, an ordinary wide character (neither **%** nor a white-space wide character), or a conversion specification. Each conversion specification is introduced by the wide character **%**. After the **%**, the following appear in sequence:

— An optional assignment-suppressing wide character **\***.

<sup>402)</sup>For binary-to-decimal conversion, the result format’s values are the numbers representable with the given format specifier. The number of significant digits is determined by the format specifier, and in the case of fixed-point conversion by the source value as well.



- An optional decimal integer greater than zero that specifies the maximum field width (in wide characters).
- An optional *length modifier* that specifies the size of the receiving object.
- A *conversion specifier* wide character that specifies the type of conversion to be applied.

- 4 The **fwscanf** function executes each directive of the format in turn. When all directives have been executed, or if a directive fails (as detailed later in this subclause), the function returns. Failures are described as input failures (due to the occurrence of an encoding error or the unavailability of input characters), or matching failures (due to inappropriate input).
- 5 A directive composed of white-space wide character(s) is executed by reading input up to the first non-white-space wide character (which remains unread), or until no more wide characters can be read. The directive never fails.
- 6 A directive that is an ordinary wide character is executed by reading the next wide character of the stream. If that wide character differs from the directive, the directive fails and the differing and subsequent wide characters remain unread. Similarly, if end-of-file, an encoding error, or a read error prevents a wide character from being read, the directive fails.
- 7 A directive that is a conversion specification defines a set of matching input sequences, as described further in this subclause for each specifier. A conversion specification is executed in the following steps:
- 8 Input white-space wide characters are skipped, unless the specification includes a **[**, **c**, or **n** specifier.<sup>403)</sup>
- 9 An input item is read from the stream, unless the specification includes an **n** specifier. An input item is defined as the longest sequence of input wide characters which does not exceed any specified field width and which is, or is a prefix of, a matching input sequence.<sup>404)</sup> The first wide character, if any, after the input item remains unread. If the length of the input item is zero, the execution of the directive fails; this condition is a matching failure unless end-of-file, an encoding error, or a read error prevented input from the stream, in which case it is an input failure.
- 10 Except in the case of a **%** specifier, the input item (or, in the case of a **%n** directive, the count of input wide characters) is converted to a type appropriate to the conversion specifier. If the input item is not a matching sequence, the execution of the directive fails: this condition is a matching failure. Unless assignment suppression was indicated by a **\***, the result of the conversion is placed in the object pointed to by the first argument following the **format** argument that has not already received a conversion result. If this object does not have an appropriate type, or if the result of the conversion cannot be represented in the object, the behavior is undefined.
- 11 The length modifiers and their meanings are:

- |                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>hh</b>           | Specifies that a following <b>b</b> , <b>d</b> , <b>i</b> , <b>o</b> , <b>u</b> , <b>x</b> , <b>X</b> , or <b>n</b> conversion specifier applies to an argument with type pointer to <b>signed char</b> or <b>unsigned char</b> .                                                                                                                                                                                                                                                                                                                                     |
| <b>h</b>            | Specifies that a following <b>b</b> , <b>d</b> , <b>i</b> , <b>o</b> , <b>u</b> , <b>x</b> , <b>X</b> , or <b>n</b> conversion specifier applies to an argument with type pointer to <b>short int</b> or <b>unsigned short int</b> .                                                                                                                                                                                                                                                                                                                                  |
| <b>l</b> (ell)      | Specifies that a following <b>b</b> , <b>d</b> , <b>i</b> , <b>o</b> , <b>u</b> , <b>x</b> , <b>X</b> , or <b>n</b> conversion specifier applies to an argument with type pointer to <b>long int</b> or <b>unsigned long int</b> ; that a following <b>a</b> , <b>A</b> , <b>e</b> , <b>E</b> , <b>f</b> , <b>F</b> , <b>g</b> , or <b>G</b> conversion specifier applies to an argument with type pointer to <b>double</b> ; or that a following <b>c</b> , <b>s</b> , or <b>[</b> conversion specifier applies to an argument with type pointer to <b>wchar_t</b> . |
| <b>ll</b> (ell-ell) | Specifies that a following <b>b</b> , <b>d</b> , <b>i</b> , <b>o</b> , <b>u</b> , <b>x</b> , <b>X</b> , or <b>n</b> conversion specifier applies to an argument with type pointer to <b>long long int</b> or <b>unsigned long long int</b> .                                                                                                                                                                                                                                                                                                                          |

<sup>403)</sup> These white-space wide characters are not counted against a specified field width.

<sup>404)</sup> **fwscanf** pushes back at most one input wide character onto the input stream. Therefore, some sequences that are acceptable to **wcstod**, **wcstol**, etc., are unacceptable to **fwscanf**.

- j** Specifies that a following **b, d, i, o, u, x, X, or n** conversion specifier applies to an argument with type pointer to **intmax\_t** or **uintmax\_t**.
- z** Specifies that a following **b, d, i, o, u, x, X, or n** conversion specifier applies to an argument with type pointer to **size\_t** or the corresponding signed integer type.
- t** Specifies that a following **b, d, i, o, u, x, X, or n** conversion specifier applies to an argument with type pointer to **ptrdiff\_t** or the corresponding unsigned integer type.
- wN** Specifies that a following **b, d, i, o, u, x, X, or n** conversion specifier applies to an argument which is a pointer to an integer with a specific width where *N* is a positive decimal integer with no leading zeros. All minimum-width integer types (7.22.2.3) and exact-width integer types (7.22.2.2) defined in the header `<stdint.h>` shall be supported. Other supported values of *N* are implementation-defined.
- wfN** Specifies that a following **b, d, i, o, u, x, X, or n** conversion specifier applies to an argument which is a pointer to a fastest minimum-width integer with a specific width where *N* is a positive decimal integer with no leading zeros. All fastest minimum-width integer types (7.22.2.4) defined in the header `<stdint.h>` shall be supported. Other supported values of *N* are implementation-defined.
- L** Specifies that a following **a, A, e, E, f, F, g, or G** conversion specifier applies to an argument with type pointer to **long double**.
- H** Specifies that a following **a, A, e, E, f, F, g, or G** conversion specifier applies to an argument with type pointer to **\_Decimal32**.
- D** Specifies that a following **a, A, e, E, f, F, g, or G** conversion specifier applies to an argument with type pointer to **\_Decimal64**.
- DD** Specifies that a following **a, A, e, E, f, F, g, or G** conversion specifier applies to an argument with type pointer to **\_Decimal128**.

If a length modifier appears with any conversion specifier other than as specified previously, the behavior is undefined.

- 12 In the following, the type of the corresponding argument for a conversion specifier shall be a pointer to a type determined by the length modifiers, if any, or specified by the conversion specifier. The conversion specifiers and their meanings are:

- d** Matches an optionally signed decimal integer, whose format is the same as expected for the subject sequence of the **wcstol** function with the value 10 for the **base** argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **int**.
- b** Matches an optionally signed binary integer, whose format is the same as expected for the subject sequence of the **wcstoul** function with the value 2 for the **base** argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- i** Matches an optionally signed integer, whose format is the same as expected for the subject sequence of the **wcstol** function with the value 0 for the **base** argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **int**.
- o** Matches an optionally signed octal integer, whose format is the same as expected for the subject sequence of the **wcstoul** function with the value 8 for the **base** argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.

- u Matches an optionally signed decimal integer, whose format is the same as expected for the subject sequence of the `wcstoul` function with the value 10 for the `base` argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- x Matches an optionally signed hexadecimal integer, whose format is the same as expected for the subject sequence of the `wcstoul` function with the value 16 for the `base` argument. Unless a length modifier is specified, the corresponding argument shall be a pointer to **unsigned int**.
- a, e, f, g Matches an optionally signed floating-point number, infinity, or NaN, whose format is the same as expected for the subject sequence of the `wcstod` function. Unless a length modifier is specified, the corresponding argument shall be a pointer to **float**.
- c Matches a sequence of wide characters of exactly the number specified by the field width (1 if no field width is present in the directive).  
 If no `l` length modifier is present, characters from the input field are converted as if by repeated calls to the `wcrtomb` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first wide character is converted. The corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence. No null character is added.  
 If an `l` length modifier is present, the corresponding argument shall be a pointer to storage of `wchar_t` large enough to accept the sequence. No null wide character is added.
- s Matches a sequence of non-white-space wide characters.  
 If no `l` length modifier is present, characters from the input field are converted as if by repeated calls to the `wcrtomb` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first wide character is converted. The corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence and a terminating null character, which will be added automatically.  
 If an `l` length modifier is present, the corresponding argument shall be a pointer to storage of `wchar_t` large enough to accept the sequence and the terminating null wide character, which will be added automatically.
- [ Matches a nonempty sequence of wide characters from a set of expected characters (the *scanset*).  
 If no `l` length modifier is present, characters from the input field are converted as if by repeated calls to the `wcrtomb` function, with the conversion state described by an `mbstate_t` object initialized to zero before the first wide character is converted. The corresponding argument shall be a pointer to **char**, **signed char**, **unsigned char**, or **void** that points to storage large enough to accept the sequence and a terminating null character, which will be added automatically.  
 If an `l` length modifier is present, the corresponding argument shall be a pointer that points to storage of `wchar_t` large enough to accept the sequence and the terminating null wide character, which will be added automatically.  
 The conversion specifier includes all subsequent wide characters in the `format` string, up to and including the matching right bracket (]). The wide characters between the brackets (the *scanlist*) compose the scanset, unless the wide character after the left bracket is a circumflex (^), in which case the scanset contains all wide characters that do not appear in the scanlist between the circumflex and the right bracket. If the conversion specifier begins with [ ] or [ ^ ], the right bracket wide character is in the scanlist and the next following right bracket wide character is the matching right bracket that ends the specification; otherwise the first following right bracket wide character is the one that ends the specification. If a - wide character is in the scanlist and is not the first, nor



the second where the first wide character is a `^`, nor the last character, the behavior is implementation-defined.

- p** Matches an implementation-defined set of sequences, which should be the same as the set of sequences that may be produced by the `%p` conversion of the `fwprintf` function. The corresponding argument shall be a pointer to a pointer of **void**. The input item is converted to a pointer value in an implementation-defined manner. If the input item is a value converted earlier during the same program execution, the pointer that results shall compare equal to that value; otherwise the behavior of the `%p` conversion is undefined.
- n** No input is consumed. The corresponding argument shall be a pointer of a signed integer type. The number of wide characters read from the input stream so far by this call to the `fwscanf` function is stored into the integer object pointed to by the argument. Execution of a `%n` directive does not increment the assignment count returned at the completion of execution of the `fwscanf` function. No argument is converted, but one is consumed. If the conversion specification includes an assignment-suppressing wide character or a field width, the behavior is undefined.
- %** Matches a single `%` wide character; no conversion or assignment occurs. The complete conversion specification shall be `%%`.

- 13 If a conversion specification is invalid, the behavior is undefined.<sup>405)</sup>
- 14 The conversion specifiers **A**, **E**, **F**, **G**, and **X** are also valid and behave the same as, respectively, **a**, **e**, **f**, **g**, and **x**.
- 15 Trailing white-space wide characters (including new-line wide characters) are left unread unless matched by a directive. The success of literal matches and suppressed assignments is not directly determinable other than via the `%n` directive.

#### Returns

- 16 The `fwscanf` function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure or if the implementation does not support a specific width length modifier.
- 17 EXAMPLE 1 The call:

```
#include <stdio.h>
#include <wchar.h>
/* ... */
int n, i; float x; wchar_t name[50];
n = fwscanf(stdin, L"%d%f%ls", &i, &x, name);
```

with the input line:

```
25 54.32E-1 thompson
```

will assign to `n` the value 3, to `i` the value 25, to `x` the value 5.432, and to `name` the sequence `thompson\0`.

- 18 EXAMPLE 2 The call:

```
#include <stdio.h>
#include <wchar.h>
/* ... */
int i; float x; double y;
fwscanf(stdin, L"%2d%f%d %lf", &i, &x, &y);
```

with input:

```
56789 0123 56a72
```

<sup>405)</sup>See “future library directions” (7.34.21).

will assign to `i` the value 56 and to `x` the value 789.0, will skip past `0123`, and will assign to `y` the value 56.0. The next wide character read from the input stream will be `a`.

**Forward references:** the `wcstod`, `wcstof`, and `wcstold` functions (7.32.4.2.2), the `wcstol`, `wcstoll`, `wcstoul`, and `wcstoull` functions (7.32.4.2.4), the `wcrtomb` function (7.32.6.4.4).

#### 7.32.2.4 The `swprintf` function

##### Synopsis

```
1 #include <wchar.h>
 int swprintf(wchar_t * restrict s, size_t n, const wchar_t * restrict format,
 ...);
```

##### Description

- 2 The `swprintf` function is equivalent to `fwprintf`, except that the argument `s` specifies an array of wide characters into which the generated output is to be written, rather than written to a stream. No more than `n` wide characters are written, including a terminating null wide character, which is always added (unless `n` is zero).

##### Returns

- 3 The `swprintf` function returns the number of wide characters written in the array, not counting the terminating null wide character, or a negative value if an encoding error occurred or if `n` or more wide characters were requested to be written.

#### 7.32.2.5 The `swscanf` function

##### Synopsis

```
1 #include <wchar.h>
 int swscanf(const wchar_t * restrict s, const wchar_t * restrict format, ...);
```

##### Description

- 2 The `swscanf` function is equivalent to `fwscanf`, except that the argument `s` specifies a wide string from which the input is to be obtained, rather than from a stream. Reaching the end of the wide string is equivalent to encountering end-of-file for the `fwscanf` function.

##### Returns

- 3 The `swscanf` function returns the value of the macro `EOF` if an input failure occurs before the first conversion (if any) has completed. Otherwise, the `swscanf` function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

#### 7.32.2.6 The `vfwprintf` function

##### Synopsis

```
1 #include <stdarg.h>
 #include <stdio.h>
 #include <wchar.h>
 int vfwprintf(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);
```

##### Description

- 2 The `vfwprintf` function is equivalent to `fwprintf`, with the varying argument list replaced by `arg`, which shall have been initialized by the `va_start` macro (and possibly subsequent `va_arg` invocations). The `vfwprintf` function does not invoke the `va_end` macro.<sup>406)</sup>

<sup>406)</sup>As the functions `vfwprintf`, `vswprintf`, `vfwscanf`, `vwprintf`, `vwscanf`, and `vswscanf` invoke the `va_arg` macro, the representation of `arg` after the return is indeterminate.

**Returns**

- 3 The **vfwprintf** function returns the number of wide characters transmitted, or a negative value if an output or encoding error occurred.
- 4 EXAMPLE The following shows the use of the **vfwprintf** function in a general error-reporting routine.

```

#include <stdarg.h>
#include <stdio.h>
#include <wchar.h>

void error(char *function_name, wchar_t *format, ...)
{
 va_list args;

 va_start(args, format);
 // print out name of function causing error
 fprintf(stderr, L"ERROR in %s: ", function_name);
 // print out remainder of message
 vfwprintf(stderr, format, args);
 va_end(args);
}

```

**7.32.2.7 The vfwscanf function****Synopsis**

```

1 #include <stdarg.h>
 #include <stdio.h>
 #include <wchar.h>
 int vfwscanf(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);

```

**Description**

- 2 The **vfwscanf** function is equivalent to **fscanf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va\_start** macro (and possibly subsequent **va\_arg** calls). The **vfwscanf** function does not invoke the **va\_end** macro.<sup>406)</sup>

**Returns**

- 3 The **vfwscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **vfwscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**7.32.2.8 The vswprintf function****Synopsis**

```

1 #include <stdarg.h>
 #include <wchar.h>
 int vswprintf(wchar_t * restrict s, size_t n, const wchar_t * restrict format,
 va_list arg);

```

**Description**

- 2 The **vswprintf** function is equivalent to **swprintf**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va\_start** macro (and possibly subsequent **va\_arg** calls). The **vswprintf** function does not invoke the **va\_end** macro.<sup>406)</sup>

**Returns**

- 3 The **vswprintf** function returns the number of wide characters written in the array, not counting the terminating null wide character, or a negative value if an encoding error occurred or if **n** or more wide characters were requested to be generated.

**7.32.2.9 The `vswscanf` function****Synopsis**

```

1 #include <stdarg.h>
 #include <wchar.h>
 int vswscanf(const wchar_t * restrict s, const wchar_t * restrict format,
 va_list arg);

```

**Description**

- 2 The **`vswscanf`** function is equivalent to **`swscanf`**, with the varying argument list replaced by **`arg`**, which shall have been initialized by the **`va_start`** macro (and possibly subsequent **`va_arg`** calls). The **`vswscanf`** function does not invoke the **`va_end`** macro.<sup>406)</sup>

**Returns**

- 3 The **`vswscanf`** function returns the value of the macro **`EOF`** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **`vswscanf`** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**7.32.2.10 The `vwprintf` function****Synopsis**

```

1 #include <stdarg.h>
 #include <wchar.h>
 int vwprintf(const wchar_t * restrict format, va_list arg);

```

**Description**

- 2 The **`vwprintf`** function is equivalent to **`wprintf`**, with the varying argument list replaced by **`arg`**, which shall have been initialized by the **`va_start`** macro (and possibly subsequent **`va_arg`** calls). The **`vwprintf`** function does not invoke the **`va_end`** macro.<sup>406)</sup>

**Returns**

- 3 The **`vwprintf`** function returns the number of wide characters transmitted, or a negative value if an output or encoding error occurred.

**7.32.2.11 The `vwscanf` function****Synopsis**

```

1 #include <stdarg.h>
 #include <wchar.h>
 int vwscanf(const wchar_t * restrict format, va_list arg);

```

**Description**

- 2 The **`vwscanf`** function is equivalent to **`wscanf`**, with the varying argument list replaced by **`arg`**, which shall have been initialized by the **`va_start`** macro (and possibly subsequent **`va_arg`** calls). The **`vwscanf`** function does not invoke the **`va_end`** macro.<sup>406)</sup>

**Returns**

- 3 The **`vwscanf`** function returns the value of the macro **`EOF`** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **`vwscanf`** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**7.32.2.12 The `wprintf` function****Synopsis**

```

1 #include <wchar.h>
 int wprintf(const wchar_t * restrict format, ...);

```

**Description**

- 2 The **wprintf** function is equivalent to **fwprintf** with the argument **stdout** interposed before the arguments to **wprintf**.

**Returns**

- 3 The **wprintf** function returns the number of wide characters transmitted, or a negative value if an output or encoding error occurred.

**7.32.2.13 The wscanf function****Synopsis**

```
1 #include <wchar.h>
 int wscanf(const wchar_t * restrict format, ...);
```

**Description**

- 2 The **wscanf** function is equivalent to **fwscanf** with the argument **stdin** interposed before the arguments to **wscanf**.

**Returns**

- 3 The **wscanf** function returns the value of the macro **EOF** if an input failure occurs before the first conversion (if any) has completed. Otherwise, the **wscanf** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**7.32.3 Wide character input/output functions****7.32.3.1 The fgetwc function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 wint_t fgetwc(FILE *stream);
```

**Description**

- 2 If the end-of-file indicator for the input stream pointed to by **stream** is not set and a next wide character is present, the **fgetwc** function obtains that wide character as a **wchar\_t** converted to a **wint\_t** and advances the associated file position indicator for the stream (if defined).

**Returns**

- 3 If the end-of-file indicator for the stream is set, or if the stream is at end-of-file, the end-of-file indicator for the stream is set and the **fgetwc** function returns **WEOF**. Otherwise, the **fgetwc** function returns the next wide character from the input stream pointed to by **stream**. If a read error occurs, the error indicator for the stream is set and the **fgetwc** function returns **WEOF**. If an encoding error occurs (including too few bytes), the error indicator for the stream is set and the value of the macro **EILSEQ** is stored in **errno** and the **fgetwc** function returns **WEOF**.<sup>407)</sup>

**7.32.3.2 The fgetws function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 wchar_t *fgetws(wchar_t * restrict s, int n, FILE * restrict stream);
```

<sup>407)</sup>An end-of-file and a read error can be distinguished by use of the **feof** and **ferror** functions. Also, **errno** will be set to **EILSEQ** by input/output functions only if an encoding error occurs.

**Description**

- 2 The **fgetws** function reads at most one less than the number of wide characters specified by **n** from the stream pointed to by **stream** into the array pointed to by **s**. No additional wide characters are read after a new-line wide character (which is retained) or after end-of-file. A null wide character is written immediately after the last wide character read into the array. If **n** is negative or zero, the behavior is undefined.

**Returns**

- 3 The **fgetws** function returns **s** if successful. If end-of-file is encountered and no characters have been read into the array, the contents of the array remain unchanged and a null pointer is returned. If a read or encoding error occurs during the operation, the array members have an indeterminate representation and a null pointer is returned.

**7.32.3.3 The fputwc function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 wint_t fputwc(wchar_t c, FILE *stream);
```

**Description**

- 2 The **fputwc** function writes the wide character specified by **c** to the output stream pointed to by **stream**, at the position indicated by the associated file position indicator for the stream (if defined), and advances the indicator appropriately. If the file cannot support positioning requests, or if the stream was opened with append mode, the character is appended to the output stream.

**Returns**

- 3 The **fputwc** function returns the wide character written. If a write error occurs, the error indicator for the stream is set and **fputwc** returns **WEOF**. If an encoding error occurs, the error indicator for the stream is set and the value of the macro **EILSEQ** is stored in **errno** and **fputwc** returns **WEOF**.

**7.32.3.4 The fputws function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 int fputws(const wchar_t * restrict s, FILE * restrict stream);
```

**Description**

- 2 The **fputws** function writes the wide string pointed to by **s** to the stream pointed to by **stream**. The terminating null wide character is not written.

**Returns**

- 3 The **fputws** function returns **EOF** if a write or encoding error occurs; otherwise, it returns a nonnegative value.

**7.32.3.5 The fwide function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 int fwide(FILE *stream, int mode);
```

**Description**

- 2 The **fwide** function determines the orientation of the stream pointed to by **stream**. If **mode** is greater than zero, the function first attempts to make the stream wide oriented. If **mode** is less than zero,

the function first attempts to make the stream byte oriented.<sup>408)</sup> Otherwise, `mode` is zero and the function does not alter the orientation of the stream.

#### Returns

- 3 The `fwide` function returns a value greater than zero if, after the call, the stream has wide orientation, a value less than zero if the stream has byte orientation, or zero if the stream has no orientation.

#### 7.32.3.6 The `getwc` function

##### Synopsis

```
1 #include <stdio.h>
 #include <wchar.h>
 wint_t getwc(FILE *stream);
```

#### Description

- 2 The `getwc` function is equivalent to `fgetwc`, except that if it is implemented as a macro, it may evaluate `stream` more than once, so the argument should never be an expression with side effects.

#### Returns

- 3 The `getwc` function returns the next wide character from the input stream pointed to by `stream`, or `WEOF`.

#### 7.32.3.7 The `getwchar` function

##### Synopsis

```
1 #include <wchar.h>
 wint_t getwchar(void);
```

#### Description

- 2 The `getwchar` function is equivalent to `getwc` with the argument `stdin`.

#### Returns

- 3 The `getwchar` function returns the next wide character from the input stream pointed to by `stdin`, or `WEOF`.

#### 7.32.3.8 The `putwc` function

##### Synopsis

```
1 #include <stdio.h>
 #include <wchar.h>
 wint_t putwc(wchar_t c, FILE *stream);
```

#### Description

- 2 The `putwc` function is equivalent to `fputwc`, except that if it is implemented as a macro, it may evaluate `stream` more than once, so that argument should never be an expression with side effects.

#### Returns

- 3 The `putwc` function returns the wide character written, or `WEOF`.

#### 7.32.3.9 The `putwchar` function

##### Synopsis

```
1 #include <wchar.h>
 wint_t putwchar(wchar_t c);
```

<sup>408)</sup>If the orientation of the stream has already been determined, `fwide` does not change it.

**Description**

- 2 The **putwchar** function is equivalent to **putc** with the second argument **stdout**.

**Returns**

- 3 The **putwchar** function returns the character written, or **WEOF**.

**7.32.3.10 The **ungetwc** function****Synopsis**

```
1 #include <stdio.h>
 #include <wchar.h>
 wint_t ungetwc(wint_t c, FILE *stream);
```

**Description**

- 2 The **ungetwc** function pushes the wide character specified by **c** back onto the input stream pointed to by **stream**. Pushed-back wide characters will be returned by subsequent reads on that stream in the reverse order of their pushing. A successful intervening call (with the stream pointed to by **stream**) to a file positioning function (**fseek**, **fsetpos**, or **rewind**) discards any pushed-back wide characters for the stream. The external storage corresponding to the stream is unchanged.
- 3 One wide character of pushback is guaranteed, even if the call to the **ungetwc** function follows just after a call to a formatted wide character input function **fwscanf**, **vfwscanf**, **wscanf**, or **wscanf**. If the **ungetwc** function is called too many times on the same stream without an intervening read or file positioning operation on that stream, the operation may fail.
- 4 If the value of **c** equals that of the macro **WEOF**, the operation fails and the input stream is unchanged.
- 5 A successful call to the **ungetwc** function clears the end-of-file indicator for the stream. The value of the file position indicator for the stream after reading or discarding all pushed-back wide characters is the same as it was before the wide characters were pushed back.<sup>409)</sup> For a text or binary stream, the value of its file position indicator after a successful call to the **ungetwc** function is unspecified until all pushed-back wide characters are read or discarded.

**Returns**

- 6 The **ungetwc** function returns the wide character pushed back, or **WEOF** if the operation fails.

**7.32.4 General wide string utilities****7.32.4.1 General**

- 1 The header `<wchar.h>` declares functions for wide string manipulation. Various methods are used for determining the lengths of the arrays, but in all cases a **wchar\_t\*** argument points to the initial (lowest addressed) element of the array. If an array is accessed beyond the end of an object, the behavior is undefined.
- 2 Where an argument declared as **size\_t n** determines the length of the array for a function, **n** can have the value zero on a call to that function. Unless explicitly stated otherwise in the description of a particular function in this subclause, pointer arguments on such a call may be null pointers. On such a call, a function that locates a wide character finds no occurrence, a function that compares two wide character sequences returns zero, and a function that copies wide characters copies zero wide characters.

**7.32.4.2 Wide string numeric conversion functions****7.32.4.2.1 General**

- 1 This subclause describes wide string analogs of the **strtod** family of functions (7.24.2.6, 7.24.2.7).<sup>410)</sup>

<sup>409)</sup> A file positioning function can further modify the file position indicator after discarding any pushed-back wide characters.

<sup>410)</sup> Wide string analogs of the **strfromd** family of functions (7.24.2.6, 7.24.2.7) are not provided because those conversions can be done by using **mbstowcs** (7.24.9.2) to convert the result of **strfromd**, **strfromf**, and similar to wide string. For example, the following converts **double d** to wide string **ws** with at most **n-1** non-null wide characters, using style **g** formatting, and computes the number **nc** of wide characters that would have been written had **n** been sufficiently large, not counting the



7.32.4.2.2 The **wcstod**, **wcstof**, and **wcstold** functions

## Synopsis

```

1 #include <wchar.h>
 double wcstod(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 float wcstof(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 long double wcstold(const wchar_t * restrict nptr, wchar_t ** restrict endptr);

```

## Description

- 2 The **wcstod**, **wcstof**, and **wcstold** functions convert the initial portion of the wide string pointed to by **nptr** to **double**, **float**, and **long double** representation, respectively. First, they decompose the input string into three parts: an initial, possibly empty, sequence of white-space wide characters, a subject sequence resembling a floating literal or representing an infinity or NaN; and a final wide string of one or more unrecognized wide characters, including the terminating null wide character of the input wide string. Then, they attempt to convert the subject sequence to a floating-point number, and return the result.
- 3 The expected form of the subject sequence is an optional plus or minus sign, then one of the following:

- decimal form: a nonempty sequence of decimal digits optionally containing a decimal-point wide character, then an optional exponent part as defined for the corresponding single-byte characters in 6.4.5.3, excluding any digit separators (6.4.5.2);
- hexadecimal form: a **0x** or **0X**, then a nonempty sequence of hexadecimal digits optionally containing a decimal-point wide character, then an optional binary exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2);
- **INF** or **INFINITY**, or any other wide string equivalent except for case
- **NAN** or **NAN**(*n-wchar-sequence<sub>opt</sub>*), or any other wide string equivalent except for case in the **NAN** part, where:

*n-wchar-sequence*:

```

 digit
 nondigit
 n-wchar-sequence digit
 n-wchar-sequence nondigit

```

The subject sequence is defined as the longest initial subsequence of the input wide string, starting with the first non-white-space wide character, that is of the expected form. The subject sequence contains no wide characters if the input wide string is not of the expected form.

- 4 If the subject sequence has the decimal or hexadecimal form, the sequence of wide characters starting with the first digit or the decimal-point wide character (whichever occurs first) is interpreted as a floating literal according to the rules of 6.4.5.3, except that the decimal-point wide character is used

---

terminating null wide character.

```

#include <stdlib.h>
const size_t n = 20;
double d;
//...
// convert d to single-byte character string s
char s[n];
int nc = strfromd(s, n, "%g", d);
// convert s (regarded as a multi-byte character
// string) to wide string ws
wchar_t ws[n];
(void)mbstowcs(ws, s, n);

```

in place of a period, and that if neither an exponent part nor a decimal-point wide character appears in a decimal form, or if a binary exponent part does not appear in a hexadecimal form, an exponent part of the appropriate type with value zero is assumed to follow the last digit in the string.

- 5 If the subject sequence begins with a minus sign, the sequence is interpreted as arithmetically negated.<sup>411)</sup>
- 6 A wide character sequence **INF** or **INFINITY** is interpreted as an infinity, if representable in the return type, else like a floating literal that is too large for the range of the return type. A wide character sequence **NAN** or **NAN**(*n-wchar-sequence<sub>opt</sub>*) is interpreted as a quiet NaN, if supported in the return type, else like a subject sequence part that does not have the expected form; the meaning of the n-wchar sequence is implementation-defined.<sup>412)</sup>
- 7 A pointer to the final wide string is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.
- 8 If the subject sequence has the hexadecimal form and **FLT\_RADIX** is a power of 2, the value resulting from the conversion is correctly rounded.
- 9 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 10 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of **nptr** is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.

#### Recommended practice

- 11 If the subject sequence has the hexadecimal form, **FLT\_RADIX** is not a power of 2, and the result is not exactly representable, the result should be one of the two numbers in the appropriate internal format that are adjacent to the hexadecimal floating source value, with the extra stipulation that the error should have a correct sign for the current rounding direction.
- 12 If the subject sequence has the decimal form and at most *M* significant digits, where *M* is the maximum value of the **T\_DECIMAL\_DIG** macros (defined in <float.h>), the result should be correctly rounded. If the subject sequence *D* has the decimal form and more than *M* significant digits, consider the two bounding, adjacent decimal strings *L* and *U*, both having *M* significant digits, such that the values of *L*, *D*, and *U* satisfy  $L \leq D \leq U$ . The result should be one of the (equal or adjacent) values that would be obtained by correctly rounding *L* and *U* according to the current rounding direction, with the extra stipulation that the error with respect to *D* should have a correct sign for the current rounding direction.<sup>413)</sup>

#### Returns

- 13 The functions return the converted value, if any. If no conversion could be performed, positive or unsigned zero is returned.
- 14 If the correct value overflows and default rounding is in effect (7.12.2), plus or minus **HUGE\_VAL**, **HUGE\_VALF**, or **HUGE\_VALL** is returned (according to the return type and sign of the value); if the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero, the integer expression **errno** acquires the value of **ERANGE**; if the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, the "overflow" floating-point exception is raised.
- 15 If the result underflows (7.12.2), the functions return a value whose magnitude is no greater than the smallest normalized positive number in the return type; if the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero, whether **errno** acquires the value **ERANGE** is implementation-defined; if the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, whether the "underflow" floating-point exception is raised is implementation-defined.

<sup>411)</sup>It is unspecified whether a minus-signed sequence is converted to a negative number directly or by arithmetically negating the value resulting from converting the corresponding unsigned sequence (see F.5); the two methods can yield different results if rounding is toward positive or negative infinity. In either case, the functions honor the sign of zero if floating-point arithmetic supports signed zeros.

<sup>412)</sup>An implementation can use the n-wchar sequence to determine extra information to be represented in the NaN's significand.

<sup>413)</sup>*M* is sufficiently large that *L* and *U* will usually correctly round to the same internal floating value, but if not will correctly round to adjacent values.

7.32.4.2.3 The **wcstodN** functions

## Synopsis

```

1 #include <wchar.h>
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 wcstod32(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 _Decimal64 wcstod64(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 _Decimal128 wcstod128(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 #endif

```

## Description

- 2 The **wcstodN** functions convert the initial portion of the wide string pointed to by **nptr** to decimal floating type representation. First, they decompose the input wide string into three parts: an initial, possibly empty, sequence of white-space wide characters; a subject sequence resembling a floating literal or representing an infinity or NaN; and a final wide string of one or more unrecognized wide characters, including the terminating null wide character of the input wide string. Then, they attempt to convert the subject sequence to a floating-point number, and return the result.
- 3 The expected form of the subject sequence is an optional plus or minus sign, then one of the following:

- decimal form: a nonempty sequence of decimal digits optionally containing a decimal-point wide character, then an optional exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2)
- hexadecimal form: a **0x** or **0X**, then a nonempty sequence of hexadecimal digits optionally containing a decimal-point wide character, then an optional binary exponent part as defined in 6.4.5.3, excluding any digit separators (6.4.5.2)
- **INF** or **INFINITY**, ignoring case
- **NAN** or **NAN(*d-wchar-sequence<sub>opt</sub>*)**, ignoring case in the **NAN** part, where:

*d-wchar-sequence*:

```

digit
nondigit
d-wchar-sequence digit
d-wchar-sequence nondigit

```

The subject sequence is defined as the longest initial subsequence of the input wide string, starting with the first non-white-space wide character, that is of the expected form. The subject sequence contains no wide characters if the input wide string is not of the expected form.

- 4 If the subject sequence has the decimal or hexadecimal form, the sequence of wide characters starting with the first digit or the decimal-point wide character (whichever occurs first) is interpreted as a floating literal according to the rules of 6.4.5.3, except that the decimal-point wide character is used in place of a period, and that if neither an exponent part nor a decimal-point wide character appears in a decimal form, or if a binary exponent part does not appear in a hexadecimal form, an exponent part of the appropriate type with value zero is assumed to follow the last digit in the wide string. If the subject sequence begins with a minus sign, the sequence is interpreted as arithmetically negated before rounding and the sign *s* is set to  $-1$ , else *s* is set to  $1$ .
- 5 If the subject sequence has the decimal form, the value resulting from the conversion is correctly rounded and the coefficient *c* and the quantum exponent *q* are determined by the rules in 6.4.5.3 for a decimal floating literal of decimal type.
- 6 If the subject sequence has the hexadecimal form, the value resulting from the conversion is correctly rounded provided the subject sequence has at most *M* significant hexadecimal digits, where  $M \geq$

$\lceil (P-1)/4 \rceil + 1$  is implementation-defined, and  $P$  is the maximum precision of the supported radix-2 floating types and binary non-arithmetic interchange formats.<sup>414)</sup> If the value resulting from the conversion is correctly rounded for all subject sequences of hexadecimal form,  $M$  may be regarded as infinite. If the subject sequence has more than  $M$  significant hexadecimal digits, the implementation may first round to  $M$  significant hexadecimal digits according to the applicable decimal rounding direction mode, signaling exceptions as though converting from a wider format, then correctly round the result of the shortened hexadecimal input to the result type. The preferred quantum exponent for the result is 0 if the hexadecimal number is exactly represented in the decimal type; the preferred quantum exponent for the result is the least possible if the hexadecimal number is not exactly represented in the decimal type.

- 7 A wide character sequence **INF** or **INFINITY** is interpreted as an infinity. A wide character sequence **NAN** or **NAN(*d-wchar-sequence<sub>opt</sub>*)**, is interpreted as a quiet NaN; the meaning of the *d-wchar* sequence is implementation-defined.<sup>415)</sup> A pointer to the final wide string is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.
- 8 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 9 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of **nptr** is stored in the object pointed to by **endptr**, provided that **endptr** is not a null pointer.

#### Returns

- 10 The **wcstodN** functions return the converted value, if any. If no conversion could be performed, the value of the triple (+1, 0, 0) is returned. If the correct value overflows:
  - the value of the macro **ERANGE** is stored in **errno** if the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero;
  - the "overflow" floating-point exception is raised if the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero.

If the result underflows (7.12.2), whether **errno** acquires the value **ERANGE** if the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero is implementation-defined; if the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, whether the "underflow" floating-point exception is raised is implementation-defined.

#### 7.32.4.2.4 The **wcstol**, **wcstoll**, **wcstoul**, and **wcstoull** functions

##### Synopsis

```
1 #include <wchar.h>
 long int wcstol(const wchar_t * restrict nptr, wchar_t ** restrict endptr,
 int base);
 long long int wcstoll(const wchar_t * restrict nptr, wchar_t ** restrict endptr,
 int base);
 unsigned long int wcstoul(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr, int base);
 unsigned long long int wcstoull(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr, int base);
```

##### Description

- 2 The **wcstol**, **wcstoll**, **wcstoul**, and **wcstoull** functions convert the initial portion of the wide string pointed to by **nptr** to **long int**, **long long int**, **unsigned long int**, and **unsigned long long int** representation, respectively. First, they decompose the input string into three parts: an initial, possibly empty, sequence of white-space wide characters, a subject sequence resembling an integer represented in some radix determined by the value of **base**, and a final wide string of one

<sup>414)</sup>Non-arithmetic interchange formats are an optional feature in Annex H.

<sup>415)</sup>An implementation can use the *d-wchar* sequence to determine extra information to be represented in the NaN's significand.

or more unrecognized wide characters, including the terminating null wide character of the input wide string. Then, they attempt to convert the subject sequence to an integer, and return the result.

- 3 If the value of `base` is zero, the expected form of the subject sequence is that of an integer literal as described for the corresponding single-byte characters in 6.4.5.2, optionally preceded by a plus or minus sign, but not including an integer suffix or any optional digit separators (6.4.5.2). If the value of `base` is between 2 and 36 (inclusive), the expected form of the subject sequence is a sequence of letters and digits representing an integer with the radix specified by `base`, optionally preceded by a plus or minus sign, but not including an integer suffix or any optional digit separators. The letters from `a` (or `A`) through `z` (or `Z`) are ascribed the values 10 through 35; only letters and digits whose ascribed values are less than that of `base` are permitted. If the value of `base` is 2, the wide characters `0b` or `0B` can optionally precede the sequence of letters and digits, following the sign if present. If the value of `base` is 8, the wide characters `0o` or `0O` can optionally precede the sequence of letters and digits, following the sign if present. If the value of `base` is 16, the wide characters `0x` or `0X` can optionally precede the sequence of letters and digits, following the sign if present.
- 4 The subject sequence is defined as the longest initial subsequence of the input wide string, starting with the first non-white-space wide character, that is of the expected form. The subject sequence contains no wide characters if the input wide string is empty or consists entirely of white-space wide characters, or if the first non-white-space wide character is other than a sign or a permissible letter or digit.
- 5 If the subject sequence has the expected form and the value of `base` is zero, the sequence of wide characters starting with the first digit is interpreted as an integer literal according to the rules of 6.4.5.2. If the subject sequence has the expected form and the value of `base` is between 2 and 36, it is used as the base for conversion, ascribing to each letter its value as previously given. If the subject sequence begins with a minus sign, the resulting value is the negative of the converted value; for functions whose return type is an unsigned integer type this action is performed in the return type. A pointer to the final wide string is stored in the object pointed to by `endptr`, provided that `endptr` is not a null pointer.
- 6 In other than the "C" locale, additional locale-specific subject sequence forms may be accepted.
- 7 If the subject sequence is empty or does not have the expected form, no conversion is performed; the value of `nptr` is stored in the object pointed to by `endptr`, provided that `endptr` is not a null pointer.

#### Returns

- 8 The `wcstol`, `wcstoll`, `wcstoul`, and `wcstoull` functions return the converted value, if any. If no conversion could be performed, zero is returned. If the correct value is outside the range of representable values, `LONG_MIN`, `LONG_MAX`, `LLONG_MIN`, `LLONG_MAX`, `ULONG_MAX`, or `ULLONG_MAX` is returned (according to the return type sign of the value, if any), and the value of the macro `ERANGE` is stored in `errno`.

### 7.32.4.3 Wide string copying functions

#### 7.32.4.3.1 The `wcscpy` function

##### Synopsis

```
1 #include <wchar.h>
 wchar_t *wcscpy(wchar_t * restrict s1, const wchar_t * restrict s2);
```

##### Description

- 2 The `wcscpy` function copies the wide string pointed to by `s2` (including the terminating null wide character) into the array pointed to by `s1`.

##### Returns

- 3 The `wcscpy` function returns the value of `s1`.

#### 7.32.4.3.2 The `wcsncpy` function

##### Synopsis

```
1
```

```
#include <wchar.h>
wchar_t *wcsncpy(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
```

**Description**

- 2 The **wcsncpy** function copies not more than *n* wide characters (those that follow a null wide character are not copied) from the array pointed to by *s2* to the array pointed to by *s1*.<sup>416)</sup>
- 3 If the array pointed to by *s2* is a wide string that is shorter than *n* wide characters, null wide characters are appended to the copy in the array pointed to by *s1*, until *n* wide characters in all have been written.

**Returns**

- 4 The **wcsncpy** function returns the value of *s1*.

**7.32.4.3.3 The wmemcpy function****Synopsis**

```
1 #include <wchar.h>
 wchar_t *wmemcpy(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
```

**Description**

- 2 The **wmemcpy** function copies *n* wide characters from the object pointed to by *s2* to the object pointed to by *s1*.

**Returns**

- 3 The **wmemcpy** function returns the value of *s1*.

**7.32.4.3.4 The wmemmove function****Synopsis**

```
1 #include <wchar.h>
 wchar_t *wmemmove(wchar_t *s1, const wchar_t *s2, size_t n);
```

**Description**

- 2 The **wmemmove** function copies *n* wide characters from the object pointed to by *s2* to the object pointed to by *s1*. Copying takes place as if the *n* wide characters from the object pointed to by *s2* are first copied into a temporary array of *n* wide characters that does not overlap the objects pointed to by *s1* or *s2*, and then the *n* wide characters from the temporary array are copied into the object pointed to by *s1*.

**Returns**

- 3 The **wmemmove** function returns the value of *s1*.

**7.32.4.4 Wide string concatenation functions****7.32.4.4.1 The wcscat function****Synopsis**

```
1 #include <wchar.h>
 wchar_t *wcscat(wchar_t * restrict s1, const wchar_t * restrict s2);
```

**Description**

- 2 The **wcscat** function appends a copy of the wide string pointed to by *s2* (including the terminating null wide character) to the end of the wide string pointed to by *s1*. The initial wide character of *s2* overwrites the null wide character at the end of *s1*.

<sup>416)</sup>Thus, if there is no null wide character in the first *n* wide characters of the array pointed to by *s2*, the result will not be null-terminated.



**Returns**

- 3 The **wscat** function returns the value of **s1**.

**7.32.4.4.2 The wcsncat function****Synopsis**

```
1 #include <wchar.h>
 wchar_t *wcsncat(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
```

**Description**

- 2 The **wcsncat** function appends not more than **n** wide characters (a null wide character and wide characters that follow it are not appended) from the array pointed to by **s2** to the end of the wide string pointed to by **s1**. The initial wide character of **s2** overwrites the null wide character at the end of **s1**. A terminating null wide character is always appended to the result.<sup>417)</sup>

**Returns**

- 3 The **wcsncat** function returns the value of **s1**.

**7.32.4.5 Wide string comparison functions****7.32.4.5.1 General**

- 1 Unless explicitly stated otherwise, the functions described in this subclause order two wide characters the same way as two integers of the underlying integer type designated by **wchar\_t**.

**7.32.4.5.2 The wscmp function****Synopsis**

```
1 #include <wchar.h>
 int wscmp(const wchar_t *s1, const wchar_t *s2);
```

**Description**

- 2 The **wscmp** function compares the wide string pointed to by **s1** to the wide string pointed to by **s2**.

**Returns**

- 3 The **wscmp** function returns an integer greater than, equal to, or less than zero, accordingly as the wide string pointed to by **s1** is greater than, equal to, or less than the wide string pointed to by **s2**.

**7.32.4.5.3 The wscoll function****Synopsis**

```
1 #include <wchar.h>
 int wscoll(const wchar_t *s1, const wchar_t *s2);
```

**Description**

- 2 The **wscoll** function compares the wide string pointed to by **s1** to the wide string pointed to by **s2**, both interpreted as appropriate to the **LC\_COLLATE** category of the current locale.

**Returns**

- 3 The **wscoll** function returns an integer greater than, equal to, or less than zero, accordingly as the wide string pointed to by **s1** is greater than, equal to, or less than the wide string pointed to by **s2** when both are interpreted as appropriate to the current locale.

**7.32.4.5.4 The wcsncmp function****Synopsis**

```
1 #include <wchar.h>
 int wcsncmp(const wchar_t *s1, const wchar_t *s2, size_t n);
```

<sup>417)</sup>Thus, the maximum number of wide characters that can end up in the array pointed to by **s1** is **wcslen(s1)+n+1**.

**Description**

- 2 The **wcsncmp** function compares not more than *n* wide characters (those that follow a null wide character are not compared) from the array pointed to by *s1* to the array pointed to by *s2*.

**Returns**

- 3 The **wcsncmp** function returns an integer greater than, equal to, or less than zero, accordingly as the possibly null-terminated array pointed to by *s1* is greater than, equal to, or less than the possibly null-terminated array pointed to by *s2*.

**7.32.4.5.5 The wcsxfrm function****Synopsis**

```
1 #include <wchar.h>
 size_t wcsxfrm(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
```

**Description**

- 2 The **wcsxfrm** function transforms the wide string pointed to by *s2* and places the resulting wide string into the array pointed to by *s1*. The transformation is such that if the **wscmp** function is applied to two transformed wide strings, it returns a value greater than, equal to, or less than zero, corresponding to the result of the **wscoll** function applied to the same two original wide strings. No more than *n* wide characters are placed into the resulting array pointed to by *s1*, including the terminating null wide character. If *n* is zero, *s1* is permitted to be a null pointer.

**Returns**

- 3 The **wcsxfrm** function returns the length of the transformed wide string (not including the terminating null wide character). If the value returned is *n* or greater, the members of the array pointed to by *s1* have an indeterminate representation.
- 4 **EXAMPLE** The value of the following expression is the length of the array needed to hold the transformation of the wide string pointed to by *s*:

```
1 + wcsxfrm(nullptr, s, 0)
```

**7.32.4.5.6 The wmemcmp function****Synopsis**

```
1 #include <wchar.h>
 int wmemcmp(const wchar_t *s1, const wchar_t *s2, size_t n);
```

**Description**

- 2 The **wmemcmp** function compares the first *n* wide characters of the object pointed to by *s1* to the first *n* wide characters of the object pointed to by *s2*.

**Returns**

- 3 The **wmemcmp** function returns an integer greater than, equal to, or less than zero, accordingly as the object pointed to by *s1* is greater than, equal to, or less than the object pointed to by *s2*.

**7.32.4.6 Wide string search functions****7.32.4.6.1 Introduction**

- 1 The stateless search functions in this subclause (**wcschr**, **wcspbrk**, **wcsrchr**, **wmemchr**, **wcsstr**) are *generic functions*. These functions are generic in the qualification of the array to be searched and will return a result pointer to an element with the same qualification as the passed array. If the array to be searched is **const**-qualified, the result pointer will be to a **const**-qualified element. If the array to be searched is not **const**-qualified,<sup>418)</sup> the result pointer will be to an unqualified element.

<sup>418)</sup>The null pointer constant is not a pointer to a **const**-qualified type, and therefore the result expression has the type of a pointer to an unqualified element.



- 2 The external declarations of these generic functions have a concrete function type that returns a pointer to an unqualified element of type `wchar_t` (indicated by `QWchar_t`), and accepts a pointer to a **const**-qualified array of the same type to search. This signature supports all correct uses. If a macro definition of any of these generic functions is suppressed to access an actual function, the external declaration with this concrete type is visible.<sup>419)</sup>
- 3 The **volatile** and **restrict** qualifiers are not accepted on the elements of the array to search.

#### 7.32.4.6.2 The `wcschr` generic function

##### Synopsis

```
1 #include <wchar.h>
 QWchar_t *wcschr(QWchar_t *s, wchar_t c);
```

##### Description

- 2 The `wcschr` generic function locates the first occurrence of `c` in the wide string pointed to by `s`. The terminating null wide character is considered to be part of the wide string.

##### Returns

- 3 The `wcschr` generic function returns a pointer to the located wide character, or a null pointer if the wide character does not occur in the wide string.

#### 7.32.4.6.3 The `wscspn` function

##### Synopsis

```
1 #include <wchar.h>
 size_t wscspn(const wchar_t *s1, const wchar_t *s2);
```

##### Description

- 2 The `wscspn` function computes the length of the maximum initial segment of the wide string pointed to by `s1` which consists entirely of wide characters *not* from the wide string pointed to by `s2`.

##### Returns

- 3 The `wscspn` function returns the length of the segment.

#### 7.32.4.6.4 The `wcspbrk` generic function

##### Synopsis

```
1 #include <wchar.h>
 QWchar_t *wcspbrk(QWchar_t *s1, const wchar_t *s2);
```

##### Description

- 2 The `wcspbrk` generic function locates the first occurrence in the wide string pointed to by `s1` of any wide character from the wide string pointed to by `s2`.

##### Returns

- 3 The `wcspbrk` generic function returns a pointer to the wide character in `s1`, or a null pointer if no wide character from `s2` occurs in `s1`.

#### 7.32.4.6.5 The `wcsrchr` generic function

##### Synopsis

```
1 #include <wchar.h>
 QWchar_t *wcsrchr(QWchar_t *s, wchar_t c);
```

<sup>419)</sup>This is an obsolescent feature.

**Description**

- 2 The **wcsrchr** generic function locates the last occurrence of **c** in the wide string pointed to by **s**. The terminating null wide character is considered to be part of the wide string.

**Returns**

- 3 The **wcsrchr** generic function returns a pointer to the wide character, or a null pointer if **c** does not occur in the wide string.

**7.32.4.6.6 The wcsspfn function****Synopsis**

```
1 #include <wchar.h>
 size_t wcsspfn(const wchar_t *s1, const wchar_t *s2);
```

**Description**

- 2 The **wcsspfn** function computes the length of the maximum initial segment of the wide string pointed to by **s1** which consists entirely of wide characters from the wide string pointed to by **s2**.

**Returns**

- 3 The **wcsspfn** function returns the length of the segment.

**7.32.4.6.7 The wcsstr generic function****Synopsis**

```
1 #include <wchar.h>
 QWchar_t *wcsstr(QWchar_t *s1, const wchar_t *s2);
```

**Description**

- 2 The **wcsstr** generic function locates the first occurrence in the wide string pointed to by **s1** of the sequence of wide characters (excluding the terminating null wide character) in the wide string pointed to by **s2**.

**Returns**

- 3 The **wcsstr** generic function returns a pointer to the located wide string, or a null pointer if the wide string is not found. If **s2** points to a wide string with zero length, the function returns **s1**.

**7.32.4.6.8 The wcstok function****Synopsis**

```
1 #include <wchar.h>
 wchar_t *wcstok(wchar_t * restrict s1, const wchar_t * restrict s2,
 wchar_t ** restrict ptr);
```

**Description**

- 2 A sequence of calls to the **wcstok** function breaks the wide string pointed to by **s1** into a sequence of tokens, each of which is delimited by a wide character from the wide string pointed to by **s2**. The third argument points to a caller-provided **wchar\_t** pointer into which the **wcstok** function stores information necessary for it to continue scanning the same wide string.
- 3 The first call in a sequence has a non-null first argument and stores an initial value in the object pointed to by **ptr**. Subsequent calls in the sequence have a null first argument and the object pointed to by **ptr** is required to have the value stored by the previous call in the sequence, which is then updated. The separator wide string pointed to by **s2** may be different from call to call.
- 4 The first call in the sequence searches the wide string pointed to by **s1** for the first wide character that is *not* contained in the current separator wide string pointed to by **s2**. If no such wide character is found, then there are no tokens in the wide string pointed to by **s1** and the **wcstok** function returns a null pointer. If such a wide character is found, it is the start of the first token.

- 5 The **wcstok** function then searches from there for a wide character that is contained in the current separator wide string. If no such wide character is found, the current token extends to the end of the wide string pointed to by **s1**, and subsequent searches in the same wide string for a token return a null pointer. If such a wide character is found, it is overwritten by a null wide character, which terminates the current token.
- 6 In all cases, the **wcstok** function stores sufficient information in the pointer pointed to by **ptr** so that subsequent calls, with a null pointer for **s1** and the unmodified pointer value for **ptr**, shall start searching just past the element overwritten by a null wide character (if any).

#### Returns

- 7 The **wcstok** function returns a pointer to the first wide character of a token, or a null pointer if there is no token.
- 8 EXAMPLE

```
#include <wchar.h>
static wchar_t str1[] = L"?a???b,,,#c";
static wchar_t str2[] = L"\t \t";
wchar_t *t, *ptr1, *ptr2;

t = wcstok(str1, L"?", &ptr1); // t points to the token L"a"
t = wcstok(nullptr, L",", &ptr1); // t points to the token L"???b"
t = wcstok(str2, L" \t", &ptr2); // t is a null pointer
t = wcstok(nullptr, L"#", &ptr1); // t points to the token L"c"
t = wcstok(nullptr, L"?", &ptr1); // t is a null pointer
```

#### 7.32.4.6.9 The **wmemchr** generic function

##### Synopsis

- 1 

```
#include <wchar.h>
QWchar_t *wmemchr(QWchar_t *s, wchar_t c, size_t n);
```

##### Description

- 2 The **wmemchr** generic function locates the first occurrence of **c** in the initial **n** wide characters of the object pointed to by **s**.

##### Returns

- 3 The **wmemchr** generic function returns a pointer to the located wide character, or a null pointer if the wide character does not occur in the object.

#### 7.32.4.7 Miscellaneous functions

##### 7.32.4.7.1 The **wcslen** function

##### Synopsis

- 1 

```
#include <wchar.h>
size_t wcslen(const wchar_t *s);
```

##### Description

- 2 The **wcslen** function computes the length of the wide string pointed to by **s**.

##### Returns

- 3 The **wcslen** function returns the number of wide characters that precede the terminating null wide character.

##### 7.32.4.7.2 The **wcsnlen** function

##### Synopsis

- 1 

```
#include <wchar.h>
```

```
size_t wcsnlen(const wchar_t *s, size_t n);
```

### Description

- 2 The **wcsnlen** function counts not more than **n** wide characters (a null wide character and wide characters that follow it are not counted) in the array to which **s** points. At most the first **n** wide characters of **s** shall be accessed by **wcsnlen**. The implementation shall behave as if it reads the wide characters sequentially and stops as soon as a null wide character is found.

### Returns

- 3 The **wcsnlen** function returns the number of wide characters that precede the terminating null wide character. If there is no null wide character in the first **n** wide characters of **s** then **wcsnlen** returns **n**.

#### 7.32.4.7.3 The **wmemset** function

##### Synopsis

```
1 #include <wchar.h>
 wchar_t *wmemset(wchar_t *s, wchar_t c, size_t n);
```

### Description

- 2 The **wmemset** function copies the value of **c** into each of the first **n** wide characters of the object pointed to by **s**.

### Returns

- 3 The **wmemset** function returns the value of **s**.

## 7.32.5 Wide character time conversion functions

### 7.32.5.1 The **wcsftime** function

##### Synopsis

```
1 #include <time.h>
 #include <wchar.h>
 size_t wcsftime(wchar_t * restrict s, size_t maxsize,
 const wchar_t * restrict format, const struct tm * restrict timeptr);
```

### Description

- 2 The **wcsftime** function is equivalent to the **strftime** function, except that:
- The argument **s** points to the initial element of an array of wide characters into which the generated output is to be placed.
  - The argument **maxsize** indicates the limiting number of wide characters.
  - The argument **format** is a wide string and the conversion specifiers are replaced by corresponding sequences of wide characters.
  - The return value indicates the number of wide characters.

### Returns

- 3 If the total number of resulting wide characters including the terminating null wide character is not more than **maxsize**, the **wcsftime** function returns the number of wide characters placed into the array pointed to by **s** not including the terminating null wide character. Otherwise, zero is returned and the members of the array have an indeterminate representation.

## 7.32.6 Extended multibyte/wide character conversion utilities

### 7.32.6.1 General

- 1 The header `<wchar.h>` declares an extended set of functions useful for conversion between multibyte characters and wide characters.

- 2 Most of the following functions — those that are listed as “restartable”, 7.32.6.4 and 7.32.6.5 — take as a last argument a pointer to an object of type `mbstate_t` that is used to describe the current *conversion state* from a particular multibyte character sequence to a wide character sequence (or the reverse) under the rules of a particular setting for the `LC_CTYPE` category of the current locale.
- 3 The initial conversion state corresponds, for a conversion in either direction, to the beginning of a new multibyte character in the initial shift state. A zero-valued `mbstate_t` object is (at least) one way to describe an initial conversion state. A zero-valued `mbstate_t` object can be used to initiate conversion involving any multibyte character sequence, in any `LC_CTYPE` category setting. If an `mbstate_t` object has been altered by any of the functions described in this subclause, and is then used with a different multibyte character sequence, or in the other conversion direction, or with a different `LC_CTYPE` category setting than on earlier function calls, the behavior is undefined.<sup>420)</sup>
- 4 On entry, each function takes the described conversion state (either internal or pointed to by an argument) as current. The conversion state described by the referenced object is altered as needed to track the shift state, and the position within a multibyte character, for the associated multibyte character sequence.

### 7.32.6.2 Single-byte/wide character conversion functions

#### 7.32.6.2.1 The `btowc` function

##### Synopsis

```
1 #include <wchar.h>
 wint_t btowc(int c);
```

##### Description

- 2 The `btowc` function determines whether `c` constitutes a valid single-byte character in the initial shift state.

##### Returns

- 3 The `btowc` function returns `WEOF` if `c` has the value `EOF` or if `(unsigned char)c` does not constitute a valid single-byte character in the initial shift state. Otherwise, it returns the wide character representation of that character.

#### 7.32.6.2.2 The `wctob` function

##### Synopsis

```
1 #include <wchar.h>
 int wctob(wint_t c);
```

##### Description

- 2 The `wctob` function determines whether `c` corresponds to a member of the extended character set whose multibyte character representation is a single byte when in the initial shift state.

##### Returns

- 3 The `wctob` function returns `E0F` if `c` does not correspond to a multibyte character with length one in the initial shift state. Otherwise, it returns the single-byte representation of that character as an `unsigned char` converted to an `int`.

### 7.32.6.3 Conversion state functions

#### 7.32.6.3.1 The `mbsinit` function

##### Synopsis

```
1 #include <wchar.h>
 int mbsinit(const mbstate_t *ps);
```

<sup>420)</sup> Thus, a particular `mbstate_t` object can be used, for example, with both the `mbrtowc` and `mbsrtowcs` functions as long as they are used to step sequentially through the same multibyte character string.

**Description**

- 2 If `ps` is not a null pointer, the `mbstate_t` object describes an initial conversion state.

**Returns**

- 3 The `mbstate_t` function returns nonzero if `ps` is a null pointer or if the referenced object describes an initial conversion state; otherwise, it returns zero.

**7.32.6.4 Restartable multibyte/wide character conversion functions****7.32.6.4.1 General**

- 1 These functions differ from the corresponding multibyte character functions of 7.24.8 (`mblen`, `mbtowc`, and `wctomb`) in that they have an extra parameter, `ps`, of type pointer to `mbstate_t` that points to an object that can completely describe the current conversion state of the associated multibyte character sequence. If `ps` is a null pointer, each function uses its own internal `mbstate_t` object instead, which is initialized prior to the first call to the function to the initial conversion state; the functions are not required to avoid data races with other calls to the same function in this case. It is implementation-defined whether the internal `mbstate_t` object has thread storage duration; if it has thread storage duration, it is initialized to the initial conversion state prior to the first call to the function on the new thread. The implementation behaves as if no library function calls these functions with a null pointer for `ps`.
- 2 Also unlike their corresponding functions, the return value does not represent whether the encoding is state-dependent.

**7.32.6.4.2 The `mbrlen` function****Synopsis**

```
1 #include <wchar.h>
 size_t mbrlen(const char * restrict s, size_t n, mbstate_t * restrict ps);
```

**Description**

- 2 The `mbrlen` function is equivalent to the call:

```
mbrtowc(nullptr, s, n, ps != nullptr ? ps: &internal)
```

where `internal` is the `mbstate_t` object for the `mbrlen` function, except that the expression designated by `ps` is evaluated only once.

**Returns**

- 3 The `mbrlen` function returns a value between zero and `n`, inclusive, `(size_t)(-2)`, or `(size_t)(-1)`.

**Forward references:** the `mbrtowc` function (7.32.6.4.3).

**7.32.6.4.3 The `mbrtowc` function****Synopsis**

```
1 #include <wchar.h>
 size_t mbrtowc(wchar_t * restrict pwc, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
```

**Description**

- 2 If `s` is a null pointer, the `mbrtowc` function is equivalent to the call:

```
mbrtowc(nullptr, "", 1, ps)
```

In this case, the values of the parameters `pwc` and `n` are ignored.

- 3 If `s` is not a null pointer, the `mbrtowc` function inspects at most `n` bytes beginning with the byte pointed to by `s` to determine the number of bytes needed to complete the next multibyte character (including any shift sequences). If the function determines that the next multibyte character is complete and valid, it determines the value of the corresponding wide character and then, if `pwc` is not a null pointer, stores that value in the object pointed to by `pwc`. If the corresponding wide character is the null wide character, the resulting state described is the initial conversion state.

#### Returns

- 4 The `mbrtowc` function returns the first of the following that applies (given the current conversion state):
- 0 if the next `n` or fewer bytes complete the multibyte character that corresponds to the null wide character (which is the value stored).
  - between 1 and `n` inclusive* if the next `n` or fewer bytes complete a valid multibyte character (which is the value stored); the value returned is the number of bytes that complete the multibyte character.
  - (`size_t`) (-2) if the next `n` bytes contribute to an incomplete (but potentially valid) multibyte character, and all `n` bytes have been processed (no value is stored).<sup>421)</sup>
  - (`size_t`) (-1) if an encoding error occurs, in which case the next `n` or fewer bytes do not contribute to a complete and valid multibyte character (no value is stored); the value of the macro `EILSEQ` is stored in `errno`, and the conversion state is unspecified.

#### 7.32.6.4.4 The `wrtomb` function

##### Synopsis

```
1 #include <wchar.h>
 size_t wrtomb(char * restrict s, wchar_t wc, mbstate_t * restrict ps);
```

##### Description

- 2 If `s` is a null pointer, the `wrtomb` function is equivalent to the call

```
wrtomb(buf, L'\0', ps)
```

where `buf` is an internal buffer.

- 3 If `s` is not a null pointer, the `wrtomb` function determines the number of bytes needed to represent the multibyte character that corresponds to the wide character given by `wc` (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by `s`. At most `MB_CUR_MAX` bytes are stored. If `wc` is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state; the resulting state described is the initial conversion state.

#### Returns

- 4 The `wrtomb` function returns the number of bytes stored in the array object (including any shift sequences). When `wc` is not a valid wide character, an encoding error occurs: the function stores the value of the macro `EILSEQ` in `errno` and returns (`size_t`) (-1); the conversion state is unspecified.

#### 7.32.6.5 Restartable multibyte/wide string conversion functions

##### 7.32.6.5.1 General

- 1 These functions differ from the corresponding multibyte string functions of 7.24.9 (`mbstowcs` and `wcstombs`) in that they have an extra parameter, `ps`, of type pointer to `mbstate_t` that points to an object that can completely describe the current conversion state of the associated multibyte

<sup>421)</sup>When `n` has at least the value of the `MB_CUR_MAX` macro, this case can only occur if `s` points at a sequence of redundant shift sequences (for implementations with state-dependent encodings).



character sequence. If `ps` is a null pointer, each function uses its own internal `mbstate_t` object instead, which is initialized prior to the first call to the function to the initial conversion state; the functions are not required to avoid data races with other calls to the same function in this case. It is implementation-defined whether the internal `mbstate_t` object has thread storage duration; if it has thread storage duration, it is initialized to the initial conversion state prior to the first call to the function on the new thread. The implementation behaves as if no library function calls these functions with a null pointer for `ps`.

- 2 Also unlike their corresponding functions, the conversion source parameter, `src`, has a pointer-to-pointer type. When the function is storing the results of conversions (that is, when `dst` is not a null pointer), the pointer object pointed to by this parameter is updated to reflect the amount of the source processed by that invocation.

#### 7.32.6.5.2 The `mbsrtowcs` function

##### Synopsis

```
1 #include <wchar.h>
 size_t mbsrtowcs(wchar_t * restrict dst, const char ** restrict src, size_t len,
 mbstate_t * restrict ps);
```

##### Description

- 2 The `mbsrtowcs` function converts a sequence of multibyte characters that begins in the conversion state described by the object pointed to by `ps`, from the array indirectly pointed to by `src` into a sequence of corresponding wide characters. If `dst` is not a null pointer, the converted characters are stored into the array pointed to by `dst`. Conversion continues up to and including a terminating null character, which is also stored. Conversion stops earlier in two cases: when a sequence of bytes is encountered that does not form a valid multibyte character, or (if `dst` is not a null pointer) when `len` wide characters have been stored into the array pointed to by `dst`.<sup>422)</sup> Each conversion takes place as if by a call to the `mbrtowc` function.
- 3 If `dst` is not a null pointer, the pointer object pointed to by `src` is assigned either a null pointer (if conversion stopped due to reaching a terminating null character) or the address just past the last multibyte character converted (if any). If conversion stopped due to reaching a terminating null character and if `dst` is not a null pointer, the resulting state described is the initial conversion state.

##### Returns

- 4 If the input conversion encounters a sequence of bytes that do not form a valid multibyte character, an encoding error occurs: the `mbsrtowcs` function stores the value of the macro `EILSEQ` in `errno` and returns `(size_t)(-1)`; the conversion state is unspecified. Otherwise, it returns the number of multibyte characters successfully converted, not including the terminating null character (if any).

#### 7.32.6.5.3 The `wcsrtombs` function

##### Synopsis

```
1 #include <wchar.h>
 size_t wcsrtombs(char * restrict dst, const wchar_t ** restrict src, size_t len,
 mbstate_t * restrict ps);
```

##### Description

- 2 The `wcsrtombs` function converts a sequence of wide characters from the array indirectly pointed to by `src` into a sequence of corresponding multibyte characters that begins in the conversion state described by the object pointed to by `ps`. If `dst` is not a null pointer, the converted characters are then stored into the array pointed to by `dst`. Conversion continues up to and including a terminating null wide character, which is also stored. Conversion stops earlier in two cases: when a wide character is reached that does not correspond to a valid multibyte character, or (if `dst` is not a null pointer) when the next multibyte character would exceed the limit of `len` total bytes to be stored into the

<sup>422)</sup>Thus, the value of `len` is ignored if `dst` is a null pointer.



array pointed to by `dst`. Each conversion takes place as if by a call to the `wcrtomb` function.<sup>423)</sup>

- 3 If `dst` is not a null pointer, the pointer object pointed to by `src` is assigned either a null pointer (if conversion stopped due to reaching a terminating null wide character) or the address just past the last wide character converted (if any). If conversion stopped due to reaching a terminating null wide character, the resulting state described is the initial conversion state.

#### Returns

- 4 If conversion stops because a wide character is reached that does not correspond to a valid multibyte character, an encoding error occurs: the `wcsrtombs` function stores the value of the macro `EILSEQ` in `errno` and returns `(size_t) (-1)`; the conversion state is unspecified. Otherwise, it returns the number of bytes in the resulting multibyte character sequence, not including the terminating null character (if any).

---

<sup>423)</sup> If conversion stops because a terminating null wide character has been reached, the bytes stored include those necessary to reach the initial shift state immediately before the null byte.

### 7.33 Wide character classification and mapping utilities <wctype.h>

#### 7.33.1 Introduction

- 1 The header <wctype.h> defines one macro, and declares three data types and many functions.<sup>424)</sup>
- 2 The types declared are `wint_t` described in 7.32.1;

`wctrans_t`

which is a scalar type that can hold values which represent locale-specific character mappings; and

`wctype_t`

which is a scalar type that can hold values which represent locale-specific character classifications.

- 3 The macro defined is `WEOF` (described in 7.32.1).
- 4 The functions declared are grouped as follows:
  - Functions that provide wide character classification;
  - Extensible functions that provide wide character classification;
  - Functions that provide wide character case mapping;
  - Extensible functions that provide wide character mapping.
- 5 For all functions described in this subclause that accept an argument of type `wint_t`, the value shall be representable as a `wchar_t` or shall equal the value of the macro `WEOF`. If this argument has any other value, the behavior is undefined.
- 6 The behavior of these functions is affected by the `LC_CTYPE` category of the current locale.

#### 7.33.2 Wide character classification utilities

##### 7.33.2.1 General

- 1 The header <wctype.h> declares several functions useful for classifying wide characters.
- 2 The term *printing wide character* refers to a member of a locale-specific set of wide characters, each of which occupies at least one printing position on a display device. The term *control wide character* refers to a member of a locale-specific set of wide characters that are not printing wide characters.

##### 7.33.2.2 Wide character classification functions

###### 7.33.2.2.1 General

- 1 The functions in this subclause return nonzero (true) if and only if the value of the argument `wc` conforms to that in the description of the function.
- 2 Each of the following functions returns true for each wide character that corresponds (as if by a call to the `wctob` function) to a single-byte character for which the corresponding character classification function from 7.4.2 returns true, except that the `iswgraph` and `iswpunct` functions may differ with respect to wide characters other than `L' '` that are both printing and white-space wide characters.<sup>425)</sup>

**Forward references:** the `wctob` function (7.32.6.2.2).

<sup>424)</sup>See “future library directions” (7.34.22).

<sup>425)</sup>For example, if the expression `isalpha(wctob(wc))` evaluates to true, then the call `iswalphabet(wc)` also returns true. But, if the expression `isgraph(wctob(wc))` evaluates to true (which cannot occur for `wc == L' '` of course), then either `iswgraph(wc)` or `iswprint(wc)&& iswspace(wc)` is true, but not both.

7.33.2.2.2 The **iswalnum** function

## Synopsis

```
1 #include <wctype.h>
 int iswalnum(wint_t wc);
```

## Description

- 2 The **iswalnum** function tests for any wide character for which **iswalpha** or **iswdigit** is true.

7.33.2.2.3 The **iswalpha** function

## Synopsis

```
1 #include <wctype.h>
 int iswalpha(wint_t wc);
```

## Description

- 2 The **iswalpha** function tests for any wide character for which **iswupper** or **iswlower** is true, or any wide character that is one of a locale-specific set of alphabetic wide characters for which none of **iswcntrl**, **iswdigit**, **iswpunct**, or **iswspace** is true.<sup>426)</sup>

7.33.2.2.4 The **iswblank** function

## Synopsis

```
1 #include <wctype.h>
 int iswblank(wint_t wc);
```

## Description

- 2 The **iswblank** function tests for any wide character that is a standard blank wide character or is one of a locale-specific set of wide characters for which **iswspace** is true and that is used to separate words within a line of text. The standard blank wide characters are the following: space (L' '), and horizontal tab (L'\t'). In the "C" locale, **iswblank** returns true only for the standard blank characters.

7.33.2.2.5 The **iswcntrl** function

## Synopsis

```
1 #include <wctype.h>
 int iswcntrl(wint_t wc);
```

## Description

- 2 The **iswcntrl** function tests for any control wide character.

7.33.2.2.6 The **iswdigit** function

## Synopsis

```
1 #include <wctype.h>
 int iswdigit(wint_t wc);
```

## Description

- 2 The **iswdigit** function tests for any wide character that corresponds to a decimal-digit character (as defined in 5.3.1).

<sup>426)</sup>The functions **iswlower** and **iswupper** test true or false separately for each of these additional wide characters; all four combinations are possible.

7.33.2.2.7 The **iswgraph** function

## Synopsis

```
1 #include <wctype.h>
 int iswgraph(wint_t wc);
```

## Description

2 The **iswgraph** function tests for any wide character for which **iswprint** is true and **iswspace** is false.<sup>427)</sup>

7.33.2.2.8 The **iswlower** function

## Synopsis

```
1 #include <wctype.h>
 int iswlower(wint_t wc);
```

## Description

2 The **iswlower** function tests for any wide character that corresponds to a lowercase letter or is one of a locale-specific set of wide characters for which none of **iswcntrl**, **iswdigit**, **iswpunct**, or **iswspace** is true.

7.33.2.2.9 The **iswprint** function

## Synopsis

```
1 #include <wctype.h>
 int iswprint(wint_t wc);
```

## Description

2 The **iswprint** function tests for any printing wide character.

7.33.2.2.10 The **iswpunct** function

## Synopsis

```
1 #include <wctype.h>
 int iswpunct(wint_t wc);
```

## Description

2 The **iswpunct** function tests for any printing wide character that is one of a locale-specific set of punctuation wide characters for which neither **iswspace** nor **iswalnum** is true.<sup>427)</sup>

7.33.2.2.11 The **iswspace** function

## Synopsis

```
1 #include <wctype.h>
 int iswspace(wint_t wc);
```

## Description

2 The **iswspace** function tests for any wide character that corresponds to a locale-specific set of white-space wide characters for which none of **iswalnum**, **iswgraph**, or **iswpunct** is true.

7.33.2.2.12 The **iswupper** function

## Synopsis

```
1 #include <wctype.h>
 int iswupper(wint_t wc);
```

<sup>427)</sup>The behavior of the **iswgraph** and **iswpunct** functions can differ from their corresponding functions in 7.4.2 with respect to printing, white-space, single-byte execution characters other than ' '.

**Description**

- 2 The **iswupper** function tests for any wide character that corresponds to an uppercase letter or is one of a locale-specific set of wide characters for which none of **iswcntrl**, **iswdigit**, **iswpunct**, or **iswspace** is true.

**7.33.2.2.13 The **iswxdigit** function****Synopsis**

```
1 #include <wctype.h>
 int iswxdigit(wint_t wc);
```

**Description**

- 2 The **iswxdigit** function tests for any wide character that corresponds to a hexadecimal-digit character (as defined in 6.4.5.2).

**7.33.2.3 Extensible wide character classification functions****7.33.2.3.1 General**

- 1 The functions **wctype** and **iswctype** provide extensible wide character classification as well as testing equivalent to that performed by the functions described in the previous subclause (7.33.2.2).

**7.33.2.3.2 The **iswctype** function****Synopsis**

```
1 #include <wctype.h>
 int iswctype(wint_t wc, wctype_t desc);
```

**Description**

- 2 The **iswctype** function determines whether the wide character **wc** has the property described by **desc**. The current setting of the **LC\_CTYPE** category shall be the same as during the call to **wctype** that returned the value **desc**.
- 3 Each of the following expressions has a truth-value equivalent to the call to the wide character classification function (7.33.2.2) in the comment that follows the expression:

```
iswctype(wc, wctype("alnum")) // iswalnum(wc)
iswctype(wc, wctype("alpha")) // iswalpha(wc)
iswctype(wc, wctype("blank")) // iswblank(wc)
iswctype(wc, wctype("cntrl")) // iswcntrl(wc)
iswctype(wc, wctype("digit")) // iswdigit(wc)
iswctype(wc, wctype("graph")) // iswgraph(wc)
iswctype(wc, wctype("lower")) // iswlower(wc)
iswctype(wc, wctype("print")) // iswprint(wc)
iswctype(wc, wctype("punct")) // iswpunct(wc)
iswctype(wc, wctype("space")) // iswspace(wc)
iswctype(wc, wctype("upper")) // iswupper(wc)
iswctype(wc, wctype("xdigit")) // iswxdigit(wc)
```

**Returns**

- 4 The **iswctype** function returns nonzero (true) if and only if the value of the wide character **wc** has the property described by **desc**. If **desc** is zero, the **iswctype** function returns zero (false).

**Forward references:** the **wctype** function (7.33.2.3.3).

**7.33.2.3.3 The **wctype** function****Synopsis**

```
1 #include <wctype.h>
 wctype_t wctype(const char *property);
```

**Description**

- 2 The **wctype** function constructs a value with type **wctype\_t** that describes a class of wide characters identified by the string argument **property**.
- 3 The strings listed in the description of the **iswctype** function shall be valid in all locales as **property** arguments to the **wctype** function.

**Returns**

- 4 If **property** identifies a valid class of wide characters according to the **LC\_CTYPE** category of the current locale, the **wctype** function returns a nonzero value that is valid as the second argument to the **iswctype** function; otherwise, it returns zero.

**7.33.3 Wide character case mapping utilities****7.33.3.1 Wide character case mapping functions****7.33.3.1.1 The **towlower** function****Synopsis**

```
1 #include <wctype.h>
 wint_t tolower(wint_t wc);
```

**Description**

- 2 The **towlower** function converts an uppercase letter to a corresponding lowercase letter.

**Returns**

- 3 If the argument is a wide character for which **iswupper** is true and there are one or more corresponding wide characters, as specified by the current locale, for which **iswlower** is true, the **towlower** function returns one of the corresponding wide characters (always the same one for any given locale); otherwise, the argument is returned unchanged.

**7.33.3.1.2 The **towupper** function****Synopsis**

```
1 #include <wctype.h>
 wint_t towupper(wint_t wc);
```

**Description**

- 2 The **towupper** function converts a lowercase letter to a corresponding uppercase letter.

**Returns**

- 3 If the argument is a wide character for which **iswlower** is true and there are one or more corresponding wide characters, as specified by the current locale, for which **iswupper** is true, the **towupper** function returns one of the corresponding wide characters (always the same one for any given locale); otherwise, the argument is returned unchanged.

**7.33.3.2 Extensible wide character case mapping functions****7.33.3.2.1 General**

- 1 The functions **wctrans** and **towctrans** provide extensible wide character mapping as well as case mapping equivalent to that performed by the functions described in the previous subclause (7.33.3.1).

**7.33.3.2.2 The **towctrans** function****Synopsis**

```
1 #include <wctype.h>
 wint_t towctrans(wint_t wc, wctrans_t desc);
```

**Description**

- 2 The **towctrans** function maps the wide character **wc** using the mapping described by **desc**. The current setting of the **LC\_CTYPE** category shall be the same as during the call to **wctrans** that returned the value **desc**.
- 3 Each of the following expressions behaves the same as the call to the wide character case mapping function (7.33.3.1) in the comment that follows the expression:

```
towctrans(wc, wctrans("tolower")) // tolower(wc)
towctrans(wc, wctrans("toupper")) // toupper(wc)
```

**Returns**

- 4 The **towctrans** function returns the mapped value of **wc** using the mapping described by **desc**. If **desc** is zero, the **towctrans** function returns the value of **wc**.

**7.33.3.2.3 The wctrans function****Synopsis**

```
1 #include <wctype.h>
 wctrans_t wctrans(const char *property);
```

**Description**

- 2 The **wctrans** function constructs a value with type **wctrans\_t** that describes a mapping between wide characters identified by the string argument **property**.
- 3 The strings listed in the description of the **towctrans** function shall be valid in all locales as **property** arguments to the **wctrans** function.

**Returns**

- 4 If **property** identifies a valid mapping of wide characters according to the **LC\_CTYPE** category of the current locale, the **wctrans** function returns a nonzero value that is valid as the second argument to the **towctrans** function; otherwise, it returns zero.

## 7.34 Future library directions

### 7.34.1 General

- 1 Although grouped under individual headers, all the external names identified as reserved identifiers or potentially reserved identifiers in this subclause remain so regardless of which headers are included in the program.

### 7.34.2 Complex arithmetic `<complex.h>`

- 1 The function names

|                         |                       |                      |                      |
|-------------------------|-----------------------|----------------------|----------------------|
| <code>cacospi</code>    | <code>cexp10m1</code> | <code>clog10</code>  | <code>crootn</code>  |
| <code>casinpi</code>    | <code>cexp10</code>   | <code>clog1p</code>  | <code>crsqr</code>   |
| <code>catanpi</code>    | <code>cexp2m1</code>  | <code>clog2p1</code> | <code>csinpi</code>  |
| <code>ccompoundn</code> | <code>cexp2</code>    | <code>clog2</code>   | <code>ctanpi</code>  |
| <code>ccospi</code>     | <code>cexpm1</code>   | <code>clogp1</code>  | <code>ctgamma</code> |
| <code>cerfc</code>      | <code>clgamma</code>  | <code>cpow</code>    |                      |
| <code>cerf</code>       | <code>clog10p1</code> | <code>cpowr</code>   |                      |

and the same names suffixed with **f** or **l** are potentially reserved identifiers and may be added to the declarations in the `<complex.h>` header.

### 7.34.3 Character handling `<ctype.h>`

- 1 Function names that begin with either **is** or **to**, and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the `<ctype.h>` header.

### 7.34.4 Errors `<errno.h>`

- 1 Macros that begin with **E** and a digit or **E** and an uppercase letter may be added to the macros defined in the `<errno.h>` header by a future edition of this document or by an implementation.

### 7.34.5 Floating-point environment `<fenv.h>`

- 1 Macros that begin with **FE\_** and an uppercase letter may be added to the macros defined in the `<fenv.h>` header by a future edition of this document or by an implementation.

### 7.34.6 Characteristics of floating types `<float.h>`

- 1 Macros that begin with **DBL\_**, **DEC32\_**, **DEC64\_**, **DEC128\_**, **DEC\_**, **FLT\_**, or **LDBL\_** and an uppercase letter are potentially reserved identifiers and may be added to the macros defined in the `<float.h>` header.
- 2 Use of the **DECIMAL\_DIG** macro is an obsolescent feature. A similar type-specific macro, such as **LDBL\_DECIMAL\_DIG**, can be used instead.
- 3 The use of **FLT\_HAS\_SUBNORM**, **DBL\_HAS\_SUBNORM**, and **LDBL\_HAS\_SUBNORM** macros is an obsolescent feature.

### 7.34.7 Format conversion of integer types `<inttypes.h>`

- 1 Macros that begin with either **PRI** or **SCN**, and either a lowercase letter, **B**, or **X** are potentially reserved identifiers and may be added to the macros defined in the `<inttypes.h>` header.
- 2 Function names that begin with **str**, or **wcs** and a lowercase letter are potentially reserved identifiers may be added to the declarations in the `<inttypes.h>` header.

### 7.34.8 Localization `<locale.h>`

- 1 Macros that begin with **LC\_** and an uppercase letter may be added to the macros defined in the `<locale.h>` header by a future edition of this document or by an implementation.

### 7.34.9 Mathematics `<math.h>`

- 1 Macros that begin with **FP\_** and an uppercase letter may be added to the macros defined in the `<math.h>` header by a future edition of this document or by an implementation.



- 2 Macros that begin with **MATH\_** and an uppercase letter are potentially reserved identifiers and may be added to the macros in the `<math.h>` header.
- 3 Function names that begin with **is** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the `<math.h>` header.
- 4 Function names that begin with **cr\_** are potentially reserved identifiers and may be added to the `<math.h>` header. The **cr\_** prefix is intended to indicate a correctly rounded version of the function.
- 5 Use of the macros **INFINITY**, **DEC\_INFINITY**, **NAN**, and **DEC\_NAN** in `<math.h>` is an obsolescent feature. Instead, use the same macros in `<float.h>`.

#### 7.34.10 Signal handling `<signal.h>`

- 1 Macros that begin with either **SIG** and an uppercase letter or **SIG\_** and an uppercase letter may be added to the macros defined in the `<signal.h>` header by a future edition of this document or by an implementation.

#### 7.34.11 Atomics `<stdatomic.h>`

- 1 Macros that begin with **ATOMIC\_** and an uppercase letter are potentially reserved identifiers and may be added to the macros defined in the `<stdatomic.h>` header. Typedef names that begin with either **atomic\_** or **memory\_**, and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the `<stdatomic.h>` header. Enumeration constants that begin with **memory\_order\_** and a lowercase letter are potentially reserved identifiers and may be added to the definition of the **memory\_order** type in the `<stdatomic.h>` header. Function names that begin with **atomic\_** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the `<stdatomic.h>` header.

#### 7.34.12 Boolean type and values `<stdbool.h>`

- 1 The macro **\_\_bool\_true\_false\_are\_defined** is an obsolescent feature.

#### 7.34.13 Bit and byte utilities `<stdbit.h>`

- 1 Type and function names that begin with **stdc\_** are potentially reserved identifiers and may be added to the declarations in the `<stdbit.h>` header.

#### 7.34.14 Checked Arithmetic Functions `<stdckdint.h>`

- 1 Type and function names that begin with **ckd\_** are potentially reserved identifiers and may be added to the declarations in the `<stdckdint.h>` header.

#### 7.34.15 Integer types `<stdint.h>`

- 1 Typedef names beginning with **int** or **uint** and ending with **\_t** are potentially reserved identifiers and may be added to the types defined in the `<stdint.h>` header. Macro names beginning with **INT** or **UINT** and ending with **\_MAX**, **\_MIN**, **\_WIDTH**, or **\_C** are potentially reserved identifiers and may be added to the macros defined in the `<stdint.h>` header.

#### 7.34.16 Input/output `<stdio.h>`

- 1 Lowercase letters may be added to the conversion specifiers and length modifiers in **fprintf** and **fscanf**. Other characters may be used in extensions. The specifier **B** for **printf** may become mandatory in future versions of this document.
- 2 The use of **ungetc** on a binary stream where the file position indicator is zero prior to the call is an obsolescent feature.

#### 7.34.17 General utilities `<stdlib.h>`

- 1 Function names that begin with **str** or **wcs** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the `<stdlib.h>` header.
- 2 Suppressing the macro definition of **bsearch** to access the actual function is an obsolescent feature.

**7.34.18 String handling <string.h>**

- 1 Function names that begin with **str**, **mem**, or **wcs** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the <string.h> header.
- 2 Suppressing the macro definitions of **memchr**, **strchr**, **strpbrk**, **strrchr**, or **strstr** to access the corresponding actual function is an obsolescent feature.

**7.34.19 Date and time <time.h>**

- 1 Macros beginning with **TIME\_** and an uppercase letter may be added to the macros in the <time.h> header by a future edition of this document or by an implementation.
- 2 The time bases **TIME\_MONOTONIC**, **TIME\_ACTIVE** and **TIME\_THREAD\_ACTIVE** may become mandatory in future versions of this standard.

**7.34.20 Threads <threads.h>**

- 1 Function names, type names, and enumeration constants that begin with either **cnd\_**, **mtx\_**, **thrd\_**, or **tss\_**, and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the <threads.h> header.

**7.34.21 Extended multibyte and wide character utilities <wchar.h>**

- 1 Function names that begin with **wcs** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the <wchar.h> header.
- 2 Lowercase letters may be added to the conversion specifiers and length modifiers in **fwprintf** and **fwscanf**. Other characters may be used in extensions.
- 3 Suppressing the macro definitions of **wcschr**, **wcspbrk**, **wcsrchr**, **wmemchr**, or **wcsstr** to access the corresponding actual function is an obsolescent feature.

**7.34.22 Wide character classification and mapping utilities <wctype.h>**

- 1 Function names that begin with **is** or **to** and a lowercase letter are potentially reserved identifiers and may be added to the declarations in the <wctype.h> header.

## Annex A

(informative)

### Language syntax summary

#### A.1 Notation

- 1 The notation is described in 6.1.

#### A.2 Lexical grammar

##### A.2.1 Lexical elements

(6.4.1) *token*:

*keyword*  
*identifier*  
*constant*  
*string-literal*  
*punctuator*

(6.4.1) *preprocessing-token*:

*header-name*  
*identifier*  
*pp-number*  
*character-literal*  
*string-literal*  
*punctuator*

each universal character name that cannot be one of the above  
 each non-white-space character that cannot be one of the above

##### A.2.2 Keywords

(6.4.2) *keyword*: one of

|                  |               |                      |                      |                    |
|------------------|---------------|----------------------|----------------------|--------------------|
| <b>alignas</b>   | <b>do</b>     | <b>int</b>           | <b>struct</b>        | <b>while</b>       |
| <b>alignof</b>   | <b>double</b> | <b>long</b>          | <b>switch</b>        | <b>_Atomic</b>     |
| <b>auto</b>      | <b>else</b>   | <b>nullptr</b>       | <b>thread_local</b>  | <b>_BitInt</b>     |
| <b>bool</b>      | <b>enum</b>   | <b>register</b>      | <b>true</b>          | <b>_Complex</b>    |
| <b>break</b>     | <b>extern</b> | <b>restrict</b>      | <b>typedef</b>       | <b>_Decimal128</b> |
| <b>case</b>      | <b>false</b>  | <b>return</b>        | <b>typeof</b>        | <b>_Decimal32</b>  |
| <b>char</b>      | <b>float</b>  | <b>short</b>         | <b>typeof_unqual</b> | <b>_Decimal64</b>  |
| <b>const</b>     | <b>for</b>    | <b>signed</b>        | <b>union</b>         | <b>_Generic</b>    |
| <b>constexpr</b> | <b>goto</b>   | <b>sizeof</b>        | <b>unsigned</b>      | <b>_Lengthof</b>   |
| <b>continue</b>  | <b>if</b>     | <b>static</b>        | <b>void</b>          | <b>_Noreturn</b>   |
| <b>default</b>   | <b>inline</b> | <b>static_assert</b> | <b>volatile</b>      |                    |

##### A.2.3 Identifiers

(6.4.3.1) *identifier*:

*identifier-start*  
*identifier* *identifier-continue*

(6.4.3.1) *identifier-start*:

*nondigit*  
 XID\_Start character  
 universal character name of class XID\_Start

(6.4.3.1) *identifier-continue*:

*digit*  
*nondigit*  
 XID\_Continue character  
 universal character name of class XID\_Continue

(6.4.3.1) *nondigit*: one of

– a b c d e f g h i j k l m  
 n o p q r s t u v w x y z  
 A B C D E F G H I J K L M  
 N O P Q R S T U V W X Y Z

(6.4.3.1) *digit*: one of

0 1 2 3 4 5 6 7 8 9

## A.2.4 Universal character names

(6.4.4) *universal-character-name*:

\u *hex-quad*  
 \U *hex-quad hex-quad*  
 \u{ *simple-hexadecimal-digit-sequence* }  
 \U{ *simple-hexadecimal-digit-sequence* }

(6.4.4) *hex-quad*:

*hexadecimal-digit hexadecimal-digit hexadecimal-digit hexadecimal-digit*

(6.4.4) *simple-hexadecimal-digit-sequence*:

*hexadecimal-digit*  
*simple-hexadecimal-digit-sequence hexadecimal-digit*

## A.2.5 Constants

(6.4.5.1) *constant*:

*integer-literal*  
*floating-literal*  
*enumeration-constant*  
*character-literal*  
*predefined-constant*

(6.4.5.2) *integer-literal*:

*decimal-literal integer-suffix<sub>opt</sub>*  
*octal-literal integer-suffix<sub>opt</sub>*  
*hexadecimal-literal integer-suffix<sub>opt</sub>*  
*binary-literal integer-suffix<sub>opt</sub>*

(6.4.5.2) *decimal-literal*:

*nonzero-digit*  
*decimal-literal ' <sub>opt</sub> digit*

(6.4.5.2) *hexadecimal-literal*:

*hexadecimal-prefix hexadecimal-digit-sequence*

(6.4.5.2) *octal-literal*:

*prefixed-octal-literal*  
*unprefixed-octal-literal*

(6.4.5.2) *unprefixed-octal-literal*:

0  
 0 ' <sub>opt</sub> *octal-digit-sequence*

(6.4.5.2) *prefixed-octal-literal*:

*octal-prefix octal-digit-sequence*

(6.4.5.2) *binary-literal*:

*binary-prefix binary-digit*  
*binary-literal ' <sub>opt</sub> binary-digit*

(6.4.5.2) *hexadecimal-prefix*: one of

0x 0X

(6.4.5.2) *octal-prefix*: one of

**0o 0O**

(6.4.5.2) *binary-prefix*: one of

**0b 0B**

(6.4.5.2) *nonzero-digit*: one of

**1 2 3 4 5 6 7 8 9**

*hexadecimal-digit-sequence*:

*hexadecimal-digit*

*hexadecimal-digit-sequence* ' <sub>opt</sub> *hexadecimal-digit*

*octal-digit-sequence*:

*octal-digit*

*octal-digit-sequence* ' <sub>opt</sub> *octal-digit*

(6.4.5.2) *hexadecimal-digit*: one of

**0 1 2 3 4 5 6 7 8 9**

**a b c d e f**

**A B C D E F**

(6.4.5.2) *octal-digit*: one of

**0 1 2 3 4 5 6 7**

(6.4.5.2) *binary-digit*: one of

**0 1**

(6.4.5.2) *integer-suffix*:

*unsigned-suffix* *long-suffix*<sub>opt</sub>

*unsigned-suffix* *long-long-suffix*

*unsigned-suffix* *bit-precise-int-suffix*

*long-suffix* *unsigned-suffix*<sub>opt</sub>

*long-long-suffix* *unsigned-suffix*<sub>opt</sub>

*bit-precise-int-suffix* *unsigned-suffix*<sub>opt</sub>

(6.4.5.2) *bit-precise-int-suffix*: one of

**wb WB**

(6.4.5.2) *unsigned-suffix*: one of

**u U**

(6.4.5.2) *long-suffix*: one of

**l L**

(6.4.5.2) *long-long-suffix*: one of

**ll LL**

(6.4.5.3) *floating-literal*:

*decimal-floating-literal*

*hexadecimal-floating-literal*

(6.4.5.3) *decimal-floating-literal*:

*fractional-literal* *exponent-part*<sub>opt</sub> *floating-suffix*<sub>opt</sub>

*digit-sequence* *exponent-part* *floating-suffix*<sub>opt</sub>

(6.4.5.3) *hexadecimal-floating-literal*:

*hexadecimal-prefix* *hexadecimal-fractional-literal*

*binary-exponent-part* *floating-suffix*<sub>opt</sub>

*hexadecimal-prefix* *hexadecimal-digit-sequence*

*binary-exponent-part* *floating-suffix*<sub>opt</sub>

(6.4.5.3) *fractional-literal*:

*digit-sequence*<sub>opt</sub> . *digit-sequence*

*digit-sequence* .

(6.4.5.3) *exponent-part*:

**e** *sign*<sub>opt</sub> *digit-sequence*  
**E** *sign*<sub>opt</sub> *digit-sequence*

(6.4.5.3) *sign*: one of

**+** **-**

(6.4.5.3) *digit-sequence*:

*digit*  
*digit-sequence* ' <sub>opt</sub> *digit*

(6.4.5.3) *hexadecimal-fractional-literal*:

*hexadecimal-digit-sequence*<sub>opt</sub> . *hexadecimal-digit-sequence*  
*hexadecimal-digit-sequence* .

(6.4.5.3) *binary-exponent-part*:

**p** *sign*<sub>opt</sub> *digit-sequence*  
**P** *sign*<sub>opt</sub> *digit-sequence*

(6.4.5.3) *real-floating-suffix*: one of

**f** **l** **F** **L** **df** **dd** **dL** **DF** **DD** **DL**

(6.4.5.3) *floating-suffix*:

*real-floating-suffix* *complex-suffix*<sub>opt</sub>  
*complex-suffix* *real-floating-suffix*<sub>opt</sub>

(6.4.5.3) *complex-suffix*: one of

**i** **I** **j** **J**

(6.4.5.4) *enumeration-constant*:

*identifier*

(6.4.5.5) *character-literal*:

*encoding-prefix*<sub>opt</sub> ' *c-char-sequence* '

(6.4.5.5) *encoding-prefix*: one of

**u8** **u** **U** **L**

(6.4.5.5) *c-char-sequence*:

*c-char*  
*c-char-sequence* *c-char*

(6.4.5.5) *c-char*:

any member of the source character set except  
the single-quote ' , backslash \ , or new-line character  
*escape-sequence*

(6.4.5.5) *escape-sequence*:

*simple-escape-sequence*  
*octal-escape-sequence*  
*hexadecimal-escape-sequence*  
*universal-character-name*

(6.4.5.5) *simple-escape-sequence*: one of

\ ' \ " \ ? \ \  
\ a \ b \ f \ n \ r \ t \ v

(6.4.5.5) *octal-escape-sequence*:

\ *octal-digit*  
\ *octal-digit* *octal-digit*  
\ *octal-digit* *octal-digit* *octal-digit*  
\ o { *simple-octal-digit-sequence* }

(6.4.5.5) *simple-octal-digit-sequence*:

*octal-digit*  
*simple-octal-digit-sequence* *octal-digit*

(6.4.5.5) *hexadecimal-escape-sequence*:

$\backslash x$  *simple-hexadecimal-digit-sequence*  
 $\backslash x\{$  *simple-hexadecimal-digit-sequence*  $\}$

(6.4.5.6) *predefined-constant*:

**false**  
**true**  
**nullptr**

## A.2.6 String literals

(6.4.6) *string-literal*:

*encoding-prefix*<sub>opt</sub> " *s-char-sequence*<sub>opt</sub> "

(6.4.6) *s-char-sequence*:

*s-char*  
*s-char-sequence* *s-char*

(6.4.6) *s-char*:

any member of the source character set except  
the double-quote " , backslash \ , or new-line character  
*escape-sequence*

## A.2.7 Punctuators

(6.4.7) *punctuator*: one of

[ ] ( ) { } . ->  
++ -- & \* + - ~ !  
/ % << >> < > <= >= == != ^ | && ||  
? : :: ; ...  
= \*= /= %= += -= <<= >>= &= ^= |=  
, # ##  
<: :> <% %> %: %: %:

## A.2.8 Header names

(6.4.8) *header-name*:

< *h-char-sequence* >  
" *q-char-sequence* "

(6.4.8) *h-char-sequence*:

*h-char*  
*h-char-sequence* *h-char*

(6.4.8) *h-char*:

any member of the source character set except  
the new-line character and >

(6.4.8) *q-char-sequence*:

*q-char*  
*q-char-sequence* *q-char*

(6.4.8) *q-char*:

any member of the source character set except  
the new-line character and "

## A.2.9 Preprocessing numbers

(6.4.9) *pp-number*:

*digit*  
*.* *digit*  
*pp-number* *identifier-continue*  
*pp-number* ' *digit*  
*pp-number* ' *nondigit*  
*pp-number* **e** *sign*  
*pp-number* **E** *sign*  
*pp-number* **p** *sign*  
*pp-number* **P** *sign*  
*pp-number* **.**

## A.3 Phrase structure grammar

### A.3.1 Expressions

(6.5.2) *primary-expression*:

*identifier*  
*constant*  
*string-literal*  
*( expression )*  
*generic-selection*

(6.5.2.1) *generic-selection*:

**\_Generic** ( *generic-controlling-operand* , *generic-assoc-list* )

(6.5.2.1) *generic-controlling-operand*:

*assignment-expression*  
*type-name*

(6.5.2.1) *generic-assoc-list*:

*generic-association*  
*generic-assoc-list* , *generic-association*

(6.5.2.1) *generic-association*:

*type-name* : *assignment-expression*  
**default** : *assignment-expression*

(6.5.3.1) *postfix-expression*:

*primary-expression*  
*postfix-expression* [ *expression* ]  
*postfix-expression* ( *argument-expression-list*<sub>opt</sub> )  
*postfix-expression* . *identifier*  
*postfix-expression* -> *identifier*  
*postfix-expression* ++  
*postfix-expression* --  
*compound-literal*

(6.5.3.1) *argument-expression-list*:

*assignment-expression*  
*argument-expression-list* , *assignment-expression*

(6.5.3.6) *compound-literal*:

( *storage-class-specifiers*<sub>opt</sub> *type-name* ) *braced-initializer*

(6.5.3.6) *storage-class-specifiers*:

*storage-class-specifier*  
*storage-class-specifiers* *storage-class-specifier*



(6.5.4.1) *unary-expression*:

*postfix-expression*  
**++** *unary-expression*  
**--** *unary-expression*  
*unary-operator* *cast-expression*  
**\_Lengthof** *unary-expression*  
**\_Lengthof** ( *type-name* )  
**sizeof** *unary-expression*  
**sizeof** ( *type-name* )  
**alignof** ( *type-name* )

(6.5.4.1) *unary-operator*: one of

**& \* + - ~ !**

(6.5.5) *cast-expression*:

*unary-expression*  
 ( *type-name* ) *cast-expression*

(6.5.6) *multiplicative-expression*:

*cast-expression*  
*multiplicative-expression* \* *cast-expression*  
*multiplicative-expression* / *cast-expression*  
*multiplicative-expression* % *cast-expression*

(6.5.7) *additive-expression*:

*multiplicative-expression*  
*additive-expression* + *multiplicative-expression*  
*additive-expression* - *multiplicative-expression*

(6.5.8) *shift-expression*:

*additive-expression*  
*shift-expression* << *additive-expression*  
*shift-expression* >> *additive-expression*

(6.5.9) *relational-expression*:

*shift-expression*  
*relational-expression* < *shift-expression*  
*relational-expression* > *shift-expression*  
*relational-expression* <= *shift-expression*  
*relational-expression* >= *shift-expression*

(6.5.10) *equality-expression*:

*relational-expression*  
*equality-expression* == *relational-expression*  
*equality-expression* != *relational-expression*

(6.5.11) *AND-expression*:

*equality-expression*  
*AND-expression* & *equality-expression*

(6.5.12) *exclusive-OR-expression*:

*AND-expression*  
*exclusive-OR-expression* ^ *AND-expression*

(6.5.13) *inclusive-OR-expression*:

*exclusive-OR-expression*  
*inclusive-OR-expression* | *exclusive-OR-expression*

(6.5.14) *logical-AND-expression*:

*inclusive-OR-expression*  
*logical-AND-expression* && *inclusive-OR-expression*

(6.5.15) *logical-OR-expression*:

*logical-AND-expression*  
*logical-OR-expression* **||** *logical-AND-expression*

(6.5.16) *conditional-expression*:

*logical-OR-expression*  
*logical-OR-expression* **?** *expression* **:** *conditional-expression*

(6.5.17.1) *assignment-expression*:

*conditional-expression*  
*unary-expression* *assignment-operator* *assignment-expression*

(6.5.17.1) *assignment-operator*: one of

**=** **\*=** **/=** **%=** **+=** **-=** **<<=** **>>=** **&=** **^=** **|=**

(6.5.18) *expression*:

*assignment-expression*  
*expression* **,** *assignment-expression*

(6.6.1) *constant-expression*:

*conditional-expression*

(6.6.2) *constant-range-expression*:

*constant-expression* **...** *constant-expression*

### A.3.2 Declarations

(6.7.1) *declaration*:

*declaration-specifiers* *init-declarator-list*<sub>opt</sub> **;**  
*attribute-specifier-sequence* *declaration-specifiers* *init-declarator-list* **;**  
*static\_assert-declaration*  
*attribute-declaration*

(6.7.1) *declaration-specifiers*:

*declaration-specifier* *attribute-specifier-sequence*<sub>opt</sub>  
*declaration-specifier* *declaration-specifiers*

(6.7.1) *declaration-specifier*:

*storage-class-specifier*  
*type-specifier-qualifier*  
*function-specifier*

(6.7.1) *init-declarator-list*:

*init-declarator*  
*init-declarator-list* **,** *init-declarator*

(6.7.1) *init-declarator*:

*declarator*  
*declarator* **=** *initializer*

(6.7.1) *attribute-declaration*:

*attribute-specifier-sequence* **;**

(6.7.1) *simple-declaration*:

*attribute-specifier-sequence*<sub>opt</sub> *declaration-specifiers* *declarator* **=** *initializer*

(6.7.2) *storage-class-specifier*:

**auto**  
**constexpr**  
**extern**  
**register**  
**static**  
**thread\_local**  
**typedef**



(6.7.3.3) *enumerator-list*:

*enumerator*  
*enumerator-list* , *enumerator*

(6.7.3.3) *enumerator*:

*enumeration-constant* *attribute-specifier-sequence*<sub>opt</sub>  
*enumeration-constant* *attribute-specifier-sequence*<sub>opt</sub> = *constant-expression*

(6.7.3.3) *enum-type-specifier*:

: *specifier-qualifier-list*

(6.7.3.5) *atomic-type-specifier*:

**\_Atomic** ( *type-name* )

(6.7.3.6) *typeof-specifier*:

**typeof** ( *typeof-specifier-argument* )  
**typeof\_unqual** ( *typeof-specifier-argument* )

(6.7.3.6) *typeof-specifier-argument*:

*expression*  
*type-name*

(6.7.4.1) *type-qualifier*:

**const**  
**restrict**  
**volatile**  
**\_Atomic**

(6.7.5) *function-specifier*:

**inline**  
**\_Noreturn**

(6.7.6) *alignment-specifier*:

**alignas** ( *type-name* )  
**alignas** ( *constant-expression* )

(6.7.7.1) *declarator*:

*pointer*<sub>opt</sub> *direct-declarator*

(6.7.7.1) *direct-declarator*:

*identifier* *attribute-specifier-sequence*<sub>opt</sub>  
( *declarator* )  
*array-declarator* *attribute-specifier-sequence*<sub>opt</sub>  
*function-declarator* *attribute-specifier-sequence*<sub>opt</sub>

(6.7.7.1) *array-declarator*:

*direct-declarator* [ *type-qualifier-list*<sub>opt</sub> *assignment-expression*<sub>opt</sub> ]  
*direct-declarator* [ **static** *type-qualifier-list*<sub>opt</sub> *assignment-expression* ]  
*direct-declarator* [ *type-qualifier-list* **static** *assignment-expression* ]  
*direct-declarator* [ *type-qualifier-list*<sub>opt</sub> \* ]

(6.7.7.1) *function-declarator*:

*direct-declarator* ( *parameter-type-list*<sub>opt</sub> )

(6.7.7.1) *pointer*:

\* *attribute-specifier-sequence*<sub>opt</sub> *type-qualifier-list*<sub>opt</sub>  
\* *attribute-specifier-sequence*<sub>opt</sub> *type-qualifier-list*<sub>opt</sub> *pointer*

(6.7.7.1) *type-qualifier-list*:

*type-qualifier*  
*type-qualifier-list* *type-qualifier*

(6.7.7.1) *parameter-type-list*:

*parameter-list*  
*parameter-list* , ...  
...

(6.7.7.1) *parameter-list*:

*parameter-declaration*  
*parameter-list* , *parameter-declaration*

(6.7.7.1) *parameter-declaration*:

*attribute-specifier-sequence*<sub>opt</sub> *declaration-specifiers* *declarator*  
*attribute-specifier-sequence*<sub>opt</sub> *declaration-specifiers* *abstract-declarator*<sub>opt</sub>

(6.7.8) *type-name*:

*specifier-qualifier-list* *abstract-declarator*<sub>opt</sub>

(6.7.8) *abstract-declarator*:

*pointer*  
*pointer*<sub>opt</sub> *direct-abstract-declarator*

(6.7.8) *direct-abstract-declarator*:

( *abstract-declarator* )  
*array-abstract-declarator* *attribute-specifier-sequence*<sub>opt</sub>  
*function-abstract-declarator* *attribute-specifier-sequence*<sub>opt</sub>

(6.7.8) *array-abstract-declarator*:

*direct-abstract-declarator*<sub>opt</sub> [ *type-qualifier-list*<sub>opt</sub> *assignment-expression*<sub>opt</sub> ]  
*direct-abstract-declarator*<sub>opt</sub> [ **static** *type-qualifier-list*<sub>opt</sub> *assignment-expression* ]  
*direct-abstract-declarator*<sub>opt</sub> [ *type-qualifier-list* **static** *assignment-expression* ]  
*direct-abstract-declarator*<sub>opt</sub> [ \* ]

(6.7.8) *function-abstract-declarator*:

*direct-abstract-declarator*<sub>opt</sub> ( *parameter-type-list*<sub>opt</sub> )

(6.7.9) *typedef-name*:

*identifier*

(6.7.11) *braced-initializer*:

{ }  
{ *initializer-list* }  
{ *initializer-list* , }

(6.7.11) *initializer*:

*assignment-expression*  
*braced-initializer*

(6.7.11) *initializer-list*:

*designation*<sub>opt</sub> *initializer*  
*initializer-list* , *designation*<sub>opt</sub> *initializer*

(6.7.11) *designation*:

*designator-list* =

(6.7.11) *designator-list*:

*designator*  
*designator-list* *designator*

(6.7.11) *designator*:

[ *constant-expression* ]  
. *identifier*

(6.7.12) *static\_assert-declaration*:

**static\_assert** ( *constant-expression* , *string-literal* ) ;  
**static\_assert** ( *constant-expression* ) ;

(6.7.13.2) *attribute-specifier-sequence*:

*attribute-specifier-sequence*<sub>opt</sub> *attribute-specifier*

(6.7.13.2) *attribute-specifier*:

[ [ *attribute-list* ] ]

(6.7.13.2) *attribute-list*:

*attribute*<sub>opt</sub>  
*attribute-list* , *attribute*<sub>opt</sub>

(6.7.13.2) *attribute*:

*attribute-token attribute-argument-clause*<sub>opt</sub>

(6.7.13.2) *attribute-token*:

*standard-attribute*  
*attribute-prefixed-token*

(6.7.13.2) *standard-attribute*:

*identifier*

(6.7.13.2) *attribute-prefixed-token*:

*attribute-prefix* :: *identifier*

(6.7.13.2) *attribute-prefix*:

*identifier*

(6.7.13.2) *attribute-argument-clause*:

( *balanced-token-sequence*<sub>opt</sub> )

(6.7.13.2) *balanced-token-sequence*:

*balanced-token*  
*balanced-token-sequence balanced-token*

(6.7.13.2) *balanced-token*:

( *balanced-token-sequence*<sub>opt</sub> )  
 [ *balanced-token-sequence*<sub>opt</sub> ]  
 { *balanced-token-sequence*<sub>opt</sub> }  
 any token other than a parenthesis, a bracket, or a brace

### A.3.3 Statements

(6.8.1) *statement*:

*labeled-statement*  
*unlabeled-statement*

(6.8.1) *unlabeled-statement*:

*expression-statement*  
*attribute-specifier-sequence*<sub>opt</sub> *primary-block*  
*attribute-specifier-sequence*<sub>opt</sub> *jump-statement*

(6.8.1) *primary-block*:

*compound-statement*  
*selection-statement*  
*iteration-statement*

(6.8.1) *secondary-block*:

*statement*

(6.8.2) *label*:

*attribute-specifier-sequence*<sub>opt</sub> *identifier* :  
*attribute-specifier-sequence*<sub>opt</sub> **case** *constant-range-expression* :  
*attribute-specifier-sequence*<sub>opt</sub> **case** *constant-expression* :  
*attribute-specifier-sequence*<sub>opt</sub> **default** :

(6.8.2) *labeled-statement*:

*label statement*

(6.8.3) *compound-statement*:

{ *block-item-list*<sub>opt</sub> }

(6.8.3) *block-item-list*:

*block-item*  
*block-item-list block-item*

(6.8.3) *block-item*:

*declaration*  
*unlabeled-statement*  
*label*

(6.8.4) *expression-statement*:

*expression*<sub>opt</sub> ;  
*attribute-specifier-sequence expression* ;

(6.8.5.1) *selection-statement*:

**if** ( *selection-header* ) *secondary-block*  
**if** ( *selection-header* ) *secondary-block* **else** *secondary-block*  
**switch** ( *selection-header* ) *secondary-block*

(6.8.5.1) *selection-header*:

*expression*  
*declaration expression*  
*simple-declaration*

(6.8.6.1) *iteration-statement*:

**while** ( *expression* ) *secondary-block*  
**do** *secondary-block* **while** ( *expression* ) ;  
**for** ( *expression*<sub>opt</sub> ; *expression*<sub>opt</sub> ; *expression*<sub>opt</sub> ) *secondary-block*  
**for** ( *declaration expression*<sub>opt</sub> ; *expression*<sub>opt</sub> ) *secondary-block*

(6.8.7.1) *jump-statement*:

**goto** *identifier* ;  
**continue** *identifier*<sub>opt</sub> ;  
**break** *identifier*<sub>opt</sub> ;  
**return** *expression*<sub>opt</sub> ;

### A.3.4 External definitions

(6.9.1) *translation-unit*:

*external-declaration*  
*translation-unit external-declaration*

(6.9.1) *external-declaration*:

*function-definition*  
*declaration*

(6.9.2) *function-definition*:

*attribute-specifier-sequence*<sub>opt</sub> *declaration-specifiers declarator function-body*

(6.9.2) *function-body*:

*compound-statement*

## A.4 Preprocessing directives

(6.10.1) *preprocessing-file*:

*group*<sub>opt</sub>

(6.10.1) *group*:

*group-part*  
*group group-part*

(6.10.1) *group-part*:

*if-section*  
*control-line*  
*text-line*  
**#** *non-directive*

(6.10.1) *if-section*:

*if-group elif-groups<sub>opt</sub> else-group<sub>opt</sub> endif-line*

(6.10.1) *if-group*:

**# if** *constant-expression new-line group<sub>opt</sub>*  
**# ifdef** *identifier new-line group<sub>opt</sub>*  
**# ifndef** *identifier new-line group<sub>opt</sub>*

(6.10.1) *elif-groups*:

*elif-group*  
*elif-groups elif-group*

(6.10.1) *elif-group*:

**# elif** *constant-expression new-line group<sub>opt</sub>*  
**# elifdef** *identifier new-line group<sub>opt</sub>*  
**# elifndef** *identifier new-line group<sub>opt</sub>*

(6.10.1) *else-group*:

**# else** *new-line group<sub>opt</sub>*

(6.10.1) *endif-line*:

**# endif** *new-line*

(6.10.1) *control-line*:

**# include** *pp-tokens new-line*  
**# embed** *pp-tokens new-line*  
**# define** *identifier replacement-list new-line*  
**# define** *identifier lparen identifier-list<sub>opt</sub> ) replacement-list new-line*  
**# define** *identifier lparen ... ) replacement-list new-line*  
**# define** *identifier lparen identifier-list , ... ) replacement-list new-line*  
**# undef** *identifier new-line*  
**# line** *pp-tokens new-line*  
**# error** *pp-tokens<sub>opt</sub> new-line*  
**# warning** *pp-tokens<sub>opt</sub> new-line*  
**# pragma** *pp-tokens<sub>opt</sub> new-line*  
**#** *new-line*

(6.10.1) *text-line*:

*pp-tokens<sub>opt</sub> new-line*

(6.10.1) *non-directive*:

*pp-tokens new-line*

(6.10.1) *lparen*:

a ( character not immediately preceded by white space

(6.10.1) *replacement-list*:

*pp-tokens<sub>opt</sub>*

(6.10.1) *pp-tokens*:

*preprocessing-token*  
*pp-tokens preprocessing-token*

(6.10.1) *new-line*:

the new-line character

(6.10.1) *identifier-list*:

*identifier*  
*identifier-list , identifier*

(6.10.1) *pp-parameter*:

*pp-parameter-name pp-parameter-clause<sub>opt</sub>*

(6.10.1) *pp-parameter-name*:

*pp-standard-parameter*  
*pp-prefixed-parameter*



(6.10.1) *pp-standard-parameter*:

*identifier*

(6.10.1) *pp-prefixed-parameter*:

*identifier* :: *identifier*

(6.10.1) *pp-parameter-clause*:

( *pp-balanced-token-sequence*<sub>opt</sub> )

(6.10.1) *pp-balanced-token-sequence*:

*pp-balanced-token*

*pp-balanced-token-sequence* *pp-balanced-token*

(6.10.1) *pp-balanced-token*:

( *pp-balanced-token-sequence*<sub>opt</sub> )

[ *pp-balanced-token-sequence*<sub>opt</sub> ]

{ *pp-balanced-token-sequence*<sub>opt</sub> }

any *pp-token* other than a parenthesis, a bracket, or a brace

(6.10.1) *embed-parameter-sequence*:

*pp-parameter*

*embed-parameter-sequence* *pp-parameter*

*defined-macro-expression*:

**defined** *identifier*

**defined** ( *identifier* )

*h-preprocessing-token*:

any *preprocessing-token* other than >

*h-pp-tokens*:

*h-preprocessing-token*

*h-pp-tokens* *h-preprocessing-token*

*header-name-tokens*:

*string-literal*

< *h-pp-tokens* >

*has-include-expression*:

**\_\_has\_include** ( *header-name* )

**\_\_has\_include** ( *header-name-tokens* )

*has-embed-expression*:

**\_\_has\_embed** ( *header-name* *embed-parameter-sequence*<sub>opt</sub> )

**\_\_has\_embed** ( *header-name-tokens* *pp-balanced-token-sequence*<sub>opt</sub> )

*has-c-attribute-express*:

**\_\_has\_c\_attribute** ( *pp-tokens* )

*va-opt-replacement*:

**\_\_VA\_OPT\_\_** ( *pp-tokens*<sub>opt</sub> )

(6.10.8) *standard-pragma*:

**# pragma STDC FP\_CONTRACT** *on-off-switch*

**# pragma STDC FENV\_ACCESS** *on-off-switch*

**# pragma STDC FENV\_DEC\_ROUND** *dec-direction*

**# pragma STDC FENV\_ROUND** *direction*

**# pragma STDC CX\_LIMITED\_RANGE** *on-off-switch*

(6.10.8) *on-off-switch*: one of

**ON** **OFF** **DEFAULT**

(6.10.8) *direction*: one of

**FE\_DOWNWARD** **FE\_TONEAREST** **FE\_TONEARESTFROMZERO**  
**FE\_TOWARDZERO** **FE\_UPWARD** **FE\_DYNAMIC**

(6.10.8) *dec-direction*: one of

**FE\_DEC\_DOWNWARD** **FE\_DEC\_TONEAREST** **FE\_DEC\_TONEARESTFROMZERO**  
**FE\_DEC\_TOWARDZERO** **FE\_DEC\_UPWARD** **FE\_DEC\_DYNAMIC**

## A.5 Floating-point subject sequence

### A.5.1 NaN char sequence

(7.24.2.6) *n-char-sequence*:

*digit*  
*nondigit*  
*n-char-sequence digit*  
*n-char-sequence nondigit*

### A.5.2 NaN **wchar\_t** sequence

(7.32.4.2.2) *n-wchar-sequence*:

*digit*  
*nondigit*  
*n-wchar-sequence digit*  
*n-wchar-sequence nondigit*

## A.6 Decimal floating-point subject sequence

### A.6.1 NaN decimal char sequence

(7.24.2.7) *d-char-sequence*:

*digit*  
*nondigit*  
*d-char-sequence digit*  
*d-char-sequence nondigit*

### A.6.2 NaN decimal **wchar\_t** sequence

(7.32.4.2.3) *d-wchar-sequence*:

*digit*  
*nondigit*  
*d-wchar-sequence digit*  
*d-wchar-sequence nondigit*

## Annex B

(informative)

### Library summary

#### B.1 Diagnostics <assert.h>

**\_\_STDC\_VERSION\_ASSERT\_H\_\_****NDEBUG**

```
void assert(scalar expression);
```

#### B.2 Complex <complex.h>

**\_\_STDC\_NO\_COMPLEX\_\_****\_Complex\_I****\_\_STDC\_VERSION\_COMPLEX\_H\_\_****I****complex**

```
#pragma STDC CX_LIMITED_RANGE on-off-switch
double complex cacos(double complex z);
float complex cacosf(float complex z);
long double complex cacosl(long double complex z);
double complex casin(double complex z);
float complex casinf(float complex z);
long double complex casinl(long double complex z);
double complex catan(double complex z);
float complex catanf(float complex z);
long double complex catanl(long double complex z);
double complex ccos(double complex z);
float complex ccosf(float complex z);
long double complex ccosl(long double complex z);
double complex csin(double complex z);
float complex csinf(float complex z);
long double complex csinl(long double complex z);
double complex ctan(double complex z);
float complex ctanf(float complex z);
long double complex ctanl(long double complex z);
double complex cacosh(double complex z);
float complex cacoshf(float complex z);
long double complex cacoshl(long double complex z);
double complex casinh(double complex z);
float complex casinhf(float complex z);
long double complex casinhl(long double complex z);
double complex catanh(double complex z);
float complex catanhf(float complex z);
long double complex catanhl(long double complex z);
double complex ccosh(double complex z);
float complex ccoshf(float complex z);
long double complex ccoshl(long double complex z);
double complex csinh(double complex z);
float complex csinhf(float complex z);
long double complex csinhl(long double complex z);
double complex ctanh(double complex z);
float complex ctanhf(float complex z);
long double complex ctanhl(long double complex z);
double complex cexp(double complex z);
float complex cexpf(float complex z);
long double complex cexpl(long double complex z);
double complex clog(double complex z);
float complex clogf(float complex z);
```

```

long double complex clogl(long double complex z);
double cabs(double complex z);
float cabsf(float complex z);
long double cabsl(long double complex z);
double complex cpow(double complex x, double complex y);
float complex cpowf(float complex x, float complex y);
long double complex cpowl(long double complex x, long double complex y);
double complex csqrt(double complex z);
float complex csqrtf(float complex z);
long double complex csqrtl(long double complex z);
double carg(double complex z);
float cargf(float complex z);
long double cargl(long double complex z);
double cimag(double complex z);
float cimagf(float complex z);
long double cimagl(long double complex z);
double complex CMPLX(double x, double y);
float complex CMPLXF(float x, float y);
long double complex CMPLXL(long double x, long double y);
double complex conj(double complex z);
float complex conjf(float complex z);
long double complex conjl(long double complex z);
double complex cproj(double complex z);
float complex cprojf(float complex z);
long double complex cprojl(long double complex z);
double creal(double complex z);
float crealf(float complex z);
long double creall(long double complex z);

```

Only if the implementation defines `__STDC_IEC_60559_TYPES__` and additionally the user code defines `__STDC_WANT_IEC_60559_TYPES_EXT__` before any inclusion of `<complex.h>`:

```

_FloatNx complex cacosfNx(_FloatNx complex z);
_FloatNx complex cacosfNx(_FloatNx complex z);
_FloatNx complex casinfNx(_FloatNx complex z);
_FloatNx complex casinfNx(_FloatNx complex z);
_FloatNx complex catanfNx(_FloatNx complex z);
_FloatNx complex catanfNx(_FloatNx complex z);
_FloatNx complex ccosfNx(_FloatNx complex z);
_FloatNx complex ccosfNx(_FloatNx complex z);
_FloatNx complex csinfNx(_FloatNx complex z);
_FloatNx complex csinfNx(_FloatNx complex z);
_FloatNx complex ctanfNx(_FloatNx complex z);
_FloatNx complex ctanfNx(_FloatNx complex z);
_FloatNx complex cacoshfNx(_FloatNx complex z);
_FloatNx complex cacoshfNx(_FloatNx complex z);
_FloatNx complex casinhfNx(_FloatNx complex z);
_FloatNx complex casinhfNx(_FloatNx complex z);
_FloatNx complex catanhfNx(_FloatNx complex z);
_FloatNx complex catanhfNx(_FloatNx complex z);
_FloatNx complex ccoshfNx(_FloatNx complex z);
_FloatNx complex ccoshfNx(_FloatNx complex z);
_FloatNx complex csinhfNx(_FloatNx complex z);
_FloatNx complex csinhfNx(_FloatNx complex z);
_FloatNx complex ctanhfNx(_FloatNx complex z);
_FloatNx complex ctanhfNx(_FloatNx complex z);
_FloatNx complex cexpfNx(_FloatNx complex z);
_FloatNx complex cexpfNx(_FloatNx complex z);
_FloatNx complex clogfNx(_FloatNx complex z);
_FloatNx complex clogfNx(_FloatNx complex z);
_FloatNx cabsfNx(_FloatNx complex z);

```

```

_FloatNx cabsfNx(_FloatNx complex z);
_FloatN complex cpowfN(_FloatN complex x, _FloatN complex y);
_FloatNx complex cpowfNx(_FloatNx complex x, _FloatNx complex y);
_FloatN complex csqrtfN(_FloatN complex z);
_FloatNx complex csqrtfNx(_FloatNx complex z);
_FloatN cargfN(_FloatN complex z);
_FloatNx cargfNx(_FloatNx complex z);
_FloatN cimagfN(_FloatN complex z);
_FloatNx cimagfNx(_FloatNx complex z);
_FloatN complex CMPLXFN(_FloatN x, _FloatN y);
_FloatNx complex CMPLXFNX(_FloatNx x, _FloatNx y);
_FloatN complex conjfN(_FloatN complex z);
_FloatNx complex conjfNx(_FloatNx complex z);
_FloatN complex cprojfN(_FloatN complex z);
_FloatNx complex cprojfNx(_FloatNx complex z);
_FloatN crealfN(_FloatN complex z);
_FloatNx crealfNx(_FloatNx complex z);

```

### B.3 Character handling <ctype.h>

```

int isalnum(int c);
int isalpha(int c);
int isblank(int c);
int iscntrl(int c);
int isdigit(int c);
int isgraph(int c);
int islower(int c);
int isprint(int c);
int ispunct(int c);
int isspace(int c);
int isupper(int c);
int isxdigit(int c);
int tolower(int c);
int toupper(int c);

```

### B.4 Errors <errno.h>

EDOM

EILSEQ

ERANGE

errno

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of <errno.h>:

`errno_t`

### B.5 Floating-point environment <fenv.h>

|                           |                            |                                      |
|---------------------------|----------------------------|--------------------------------------|
| <code>fenv_t</code>       | <code>FE_INVALID</code>    | <code>FE_TONEARESTFROMZERO</code>    |
| <code>fexcept_t</code>    | <code>FE_OVERFLOW</code>   | <code>FE_TOWARDZERO</code>           |
| <code>femode_t</code>     | <code>FE_UNDERFLOW</code>  | <code>FE_UPWARD</code>               |
| <code>FE_DFL_MODE</code>  | <code>FE_ALL_EXCEPT</code> | <code>FE_DFL_ENV</code>              |
| <code>FE_DIVBYZERO</code> | <code>FE_DOWNWARD</code>   | <code>__STDC_VERSION_FENV_H__</code> |
| <code>FE_INEXACT</code>   | <code>FE_TONEAREST</code>  |                                      |

```

#pragma STDC FENV_ACCESS on-off-switch
#pragma STDC FENV_ROUND direction
#pragma STDC FENV_ROUND FE_DYNAMIC
int feclearexcept(int excepts);

```

```

int fegetexceptflag(fexcept_t *flagp, int excepts);
int feraiseexcept(int excepts);
int fesetexcept(int excepts);
int fesetexceptflag(const fexcept_t *flagp, int excepts);
int fetestexceptflag(const fexcept_t *flagp, int excepts);
int fetestexcept(int excepts);
int fegetmode(femode_t *modep);
int fegetround(void);
int fesetmode(const femode_t *modep);
int fesetround(int rnd);
int fegetenv(fenv_t *envp);
int feholdexcept(fenv_t *envp);
int fesetenv(const fenv_t *envp);
int feupdateenv(const fenv_t *envp);

```

Only if the implementation defines `__STDC_IEC_60559_DFP__`:

```

FE_DEC_DOWNWARD FE_DEC_TONEARESTFROMZERO FE_DEC_UPWARD
FE_DEC_TONEAREST FE_DEC_TOWARDZERO

```

```

#pragma STDC FENV_DEC_ROUND dec-direction
int fe_dec_getround(void);
int fe_dec_setround(int rnd);

```

Only if the implementation follows the recommended practice from F.2.2:

```
FE_SNANS_ALWAYS_SIGNAL
```

## B.6 Characteristics of floating types <float.h>

### B.6.1 Macros

|                                       |                              |                            |
|---------------------------------------|------------------------------|----------------------------|
| <code>__STDC_VERSION_FLOAT_H__</code> | <code>DBL_DIG</code>         | <code>DBL_NORM_MAX</code>  |
| <code>FLT_ROUNDS</code>               | <code>LDBL_DIG</code>        | <code>LDBL_NORM_MAX</code> |
| <code>FLT_EVAL_METHOD</code>          | <code>FLT_MIN_EXP</code>     | <code>FLT_EPSILON</code>   |
| <code>FLT_HAS_SUBNORM</code>          | <code>DBL_MIN_EXP</code>     | <code>DBL_EPSILON</code>   |
| <code>DBL_HAS_SUBNORM</code>          | <code>LDBL_MIN_EXP</code>    | <code>LDBL_EPSILON</code>  |
| <code>LDBL_HAS_SUBNORM</code>         | <code>FLT_MIN_10_EXP</code>  | <code>FLT_MIN</code>       |
| <code>FLT_RADIX</code>                | <code>DBL_MIN_10_EXP</code>  | <code>DBL_MIN</code>       |
| <code>FLT_MANT_DIG</code>             | <code>LDBL_MIN_10_EXP</code> | <code>LDBL_MIN</code>      |
| <code>DBL_MANT_DIG</code>             | <code>FLT_MAX_EXP</code>     | <code>FLT_SNAN</code>      |
| <code>LDBL_MANT_DIG</code>            | <code>DBL_MAX_EXP</code>     | <code>DBL_SNAN</code>      |
| <code>FLT_DECIMAL_DIG</code>          | <code>LDBL_MAX_EXP</code>    | <code>LDBL_SNAN</code>     |
| <code>DBL_DECIMAL_DIG</code>          | <code>FLT_MAX_10_EXP</code>  | <code>FLT_TRUE_MIN</code>  |
| <code>LDBL_DECIMAL_DIG</code>         | <code>DBL_MAX_10_EXP</code>  | <code>DBL_TRUE_MIN</code>  |
| <code>DECIMAL_DIG</code>              | <code>LDBL_MAX_10_EXP</code> | <code>LDBL_TRUE_MIN</code> |
| <code>FLT_IS_IEC_60559</code>         | <code>FLT_MAX</code>         | <code>INFINITY</code>      |
| <code>DBL_IS_IEC_60559</code>         | <code>DBL_MAX</code>         | <code>NAN</code>           |
| <code>LDBL_IS_IEC_60559</code>        | <code>LDBL_MAX</code>        |                            |
| <code>FLT_DIG</code>                  | <code>FLT_NORM_MAX</code>    |                            |

The following macro is provided only if the program defines `__STDC_WANT_IEC_60559_EXT__` before inclusion of the header <float.h>:

## CR\_DECIMAL\_DIG

### B.6.2 Characteristics of decimal floating types

- 1 The following macros are provided only if the implementation defines `__STDC_IEC_60559_DFP__`.  
 $N$  is 32, 64 and 128.

DEC\_EVAL\_METHOD  
DEC\_INFINITY  
DEC\_NAN

DEC<sub>N</sub>\_EPSILON  
DEC<sub>N</sub>\_MANT\_DIG  
DEC<sub>N</sub>\_MAX\_EXP

DEC<sub>N</sub>\_MAX  
DEC<sub>N</sub>\_MIN\_EXP  
DEC<sub>N</sub>\_MIN

DEC<sub>N</sub>\_TRUE\_MIN  
DEC<sub>N</sub>\_SNAN

### B.6.3 Characteristics of ISO/IEC 60559 interchange and extended types

The following macros are provided only if the implementation defines `__STDC_IEC_60559_TYPES__` and additionally the user code defines `__STDC_WANT_IEC_60559_TYPES_EXT__` before any inclusion of `<float.h>`.

```

FLT_N_DECIMAL_DIG
FLT_N_DIG
FLT_N_EPSILON
FLT_N_MANT_DIG
FLT_N_MAX_10_EXP
FLT_N_MAX_EXP
FLT_N_MAX
FLT_N_MIN_10_EXP
FLT_N_MIN_EXP
FLT_N_MIN

```

```

FLT_N_SNAN
FLT_TRUE_MIN
FLT_DECIMAL_DIG
FLT_DIG
FLT_EPSILON
FLT_MANT_DIG
FLT_MAX_10_EXP
FLT_MAX_EXP
FLT_MAX
FLT_MIN_10_EXP

```

```
FLT_NX_MIN_EXP
FLT_NX_MIN
FLT_NX_SNAN
FLT_NX_TRUE_MIN
DEC_N_EPSILON
DEC_N_MANT_DIG
DEC_N_MAX_EXP
DEC_N_MAX
DEC_N_MIN_EXP
DEC_N_MIN
```

```
DEC_N_SNaN
DEC_TRUE_MIN
DEC_EPSILON
DEC_MANT_DIG
DEC_MAX_EXP
DEC_MAX
DEC_MIN_EXP
DEC_MIN
DEC_N_SNaN
DEC_TRUE_MIN
```

## B.7 Format conversion of integer types <inttypes.h>

\_\_STDC\_VERSION\_INTTYPES\_H\_\_

**imaxdiv\_t**

|                      |                                   |                                  |                        |                        |
|----------------------|-----------------------------------|----------------------------------|------------------------|------------------------|
| <b>PRIB</b> <i>N</i> | <b>PRIB</b> <i>LEAST</i> <i>N</i> | <b>PRIB</b> <i>FAST</i> <i>N</i> | <b>PRIB</b> <i>MAX</i> | <b>PRIB</b> <i>PTR</i> |
| <b>PRIB</b> <i>N</i> | <b>PRIB</b> <i>LEAST</i> <i>N</i> | <b>PRIB</b> <i>FAST</i> <i>N</i> | <b>PRIB</b> <i>MAX</i> | <b>PRIB</b> <i>PTR</i> |
| <b>PRID</b> <i>N</i> | <b>PRID</b> <i>LEAST</i> <i>N</i> | <b>PRID</b> <i>FAST</i> <i>N</i> | <b>PRID</b> <i>MAX</i> | <b>PRID</b> <i>PTR</i> |
| <b>PRII</b> <i>N</i> | <b>PRII</b> <i>LEAST</i> <i>N</i> | <b>PRII</b> <i>FAST</i> <i>N</i> | <b>PRII</b> <i>MAX</i> | <b>PRII</b> <i>PTR</i> |
| <b>PRIO</b> <i>N</i> | <b>PRIO</b> <i>LEAST</i> <i>N</i> | <b>PRIO</b> <i>FAST</i> <i>N</i> | <b>PRIO</b> <i>MAX</i> | <b>PRIO</b> <i>PTR</i> |
| <b>PRIU</b> <i>N</i> | <b>PRIU</b> <i>LEAST</i> <i>N</i> | <b>PRIU</b> <i>FAST</i> <i>N</i> | <b>PRIU</b> <i>MAX</i> | <b>PRIU</b> <i>PTR</i> |
| <b>PRIX</b> <i>N</i> | <b>PRIX</b> <i>LEAST</i> <i>N</i> | <b>PRIX</b> <i>FAST</i> <i>N</i> | <b>PRIX</b> <i>MAX</i> | <b>PRIX</b> <i>PTR</i> |
| <b>PRIX</b> <i>N</i> | <b>PRIX</b> <i>LEAST</i> <i>N</i> | <b>PRIX</b> <i>FAST</i> <i>N</i> | <b>PRIX</b> <i>MAX</i> | <b>PRIX</b> <i>PTR</i> |
| <b>SCNb</b> <i>N</i> | <b>SCNb</b> <i>LEAST</i> <i>N</i> | <b>SCNb</b> <i>FAST</i> <i>N</i> | <b>SCNb</b> <i>MAX</i> | <b>SCNb</b> <i>PTR</i> |
| <b>SCNd</b> <i>N</i> | <b>SCNd</b> <i>LEAST</i> <i>N</i> | <b>SCNd</b> <i>FAST</i> <i>N</i> | <b>SCNd</b> <i>MAX</i> | <b>SCNd</b> <i>PTR</i> |
| <b>SCNi</b> <i>N</i> | <b>SCNi</b> <i>LEAST</i> <i>N</i> | <b>SCNi</b> <i>FAST</i> <i>N</i> | <b>SCNi</b> <i>MAX</i> | <b>SCNi</b> <i>PTR</i> |
| <b>SCNo</b> <i>N</i> | <b>SCNo</b> <i>LEAST</i> <i>N</i> | <b>SCNo</b> <i>FAST</i> <i>N</i> | <b>SCNo</b> <i>MAX</i> | <b>SCNo</b> <i>PTR</i> |
| <b>SCNu</b> <i>N</i> | <b>SCNu</b> <i>LEAST</i> <i>N</i> | <b>SCNu</b> <i>FAST</i> <i>N</i> | <b>SCNu</b> <i>MAX</i> | <b>SCNu</b> <i>PTR</i> |
| <b>SCNx</b> <i>N</i> | <b>SCNx</b> <i>LEAST</i> <i>N</i> | <b>SCNx</b> <i>FAST</i> <i>N</i> | <b>SCNx</b> <i>MAX</i> | <b>SCNx</b> <i>PTR</i> |

```
intmax_t imaxabs(intmax_t j);
uintmax_t uimaxabs(intmax_t j);
imaxdiv_t imaxdiv(intmax_t numer, intmax_t denom);
intmax_t strtointmax(const char * restrict nptr, char ** restrict endptr, int base);
uintmax_t strtouintmax(const char * restrict nptr, char ** restrict endptr, int base);
intmax_t wcstointmax(const wchar_t * restrict nptr, wchar_t ** restrict endptr, int base);
uintmax_t wcstouintmax(const wchar_t * restrict nptr, wchar_t ** restrict endptr, int base);
```

## B.8 Alternative spellings `<iso646.h>`

|               |              |               |               |
|---------------|--------------|---------------|---------------|
| <b>and</b>    | <b>bitor</b> | <b>not_eq</b> | <b>xor</b>    |
| <b>and_eq</b> | <b>compl</b> | <b>or</b>     | <b>xor_eq</b> |
| <b>bitand</b> | <b>not</b>   | <b>or_eq</b>  |               |

## B.9 Sizes of integer types <limits.h>

|                                  |                     |                   |
|----------------------------------|---------------------|-------------------|
| <b>__STDC_VERSION_LIMITS_H__</b> | <b>ULONG_WIDTH</b>  | <b>USHRT_MAX</b>  |
| <b>BITINT_MAXWIDTH</b>           | <b>LLONG_WIDTH</b>  | <b>INT_MIN</b>    |
| <b>BOOL_WIDTH</b>                | <b>ULLONG_WIDTH</b> | <b>INT_MAX</b>    |
| <b>CHAR_BIT</b>                  | <b>BOOL_MAX</b>     | <b>UINT_MAX</b>   |
| <b>CHAR_WIDTH</b>                | <b>SCHAR_MIN</b>    | <b>LONG_MIN</b>   |
| <b>SCHAR_WIDTH</b>               | <b>SCHAR_MAX</b>    | <b>LONG_MAX</b>   |
| <b>UCHAR_WIDTH</b>               | <b>UCHAR_MAX</b>    | <b>ULONG_MAX</b>  |
| <b>SHRT_WIDTH</b>                | <b>CHAR_MIN</b>     | <b>LLONG_MIN</b>  |
| <b>USHRT_WIDTH</b>               | <b>CHAR_MAX</b>     | <b>LLONG_MAX</b>  |
| <b>INT_WIDTH</b>                 | <b>MB_LEN_MAX</b>   | <b>ULLONG_MAX</b> |
| <b>UINT_WIDTH</b>                | <b>SHRT_MIN</b>     |                   |
| <b>LONG_WIDTH</b>                | <b>SHRT_MAX</b>     |                   |

## B.10 Localization <locale.h>

|                     |                   |                    |                   |
|---------------------|-------------------|--------------------|-------------------|
| <b>struct lconv</b> | <b>LC_ALL</b>     | <b>LC_CTYPE</b>    | <b>LC_NUMERIC</b> |
| <b>NULL</b>         | <b>LC_COLLATE</b> | <b>LC_MONETARY</b> | <b>LC_TIME</b>    |

```
char *setlocale(int category, const char *locale);
struct lconv *localeconv(void);
```

## B.11 Mathematics <math.h>

|                                |                                 |                         |
|--------------------------------|---------------------------------|-------------------------|
| <b>__STDC_VERSION_MATH_H__</b> | <b>FP_INT_TONEARESTFROMZERO</b> | <b>FP_FAST_DDIVL</b>    |
| <b>float_t</b>                 | <b>FP_INT_TONEAREST</b>         | <b>FP_FAST_FSQRT</b>    |
| <b>double_t</b>                | <b>FP_FAST_FMA</b>              | <b>FP_FAST_FSQRTL</b>   |
| <b>HUGE_VAL</b>                | <b>FP_FAST_FMAF</b>             | <b>FP_FAST_DSQRTL</b>   |
| <b>HUGE_VALF</b>               | <b>FP_FAST_FMAL</b>             | <b>FP_FAST_FFMA</b>     |
| <b>HUGE_VALL</b>               | <b>FP_FAST_FADD</b>             | <b>FP_FAST_FFMAL</b>    |
| <b>INFINITY</b>                | <b>FP_FAST_FADDL</b>            | <b>FP_FAST_DFMAL</b>    |
| <b>NAN</b>                     | <b>FP_FAST_DADDL</b>            | <b>FP_ILOGB0</b>        |
| <b>FP_INFINITE</b>             | <b>FP_FAST_FSUB</b>             | <b>FP_ILOGBNAN</b>      |
| <b>FP_NAN</b>                  | <b>FP_FAST_FSUBL</b>            | <b>FP_LLOGB0</b>        |
| <b>FP_NORMAL</b>               | <b>FP_FAST_DSUBL</b>            | <b>FP_LLOGBNAN</b>      |
| <b>FP_SUBNORMAL</b>            | <b>FP_FAST_FMUL</b>             | <b>MATH_ERRNO</b>       |
| <b>FP_ZERO</b>                 | <b>FP_FAST_FMULL</b>            | <b>MATH_ERREXCEPT</b>   |
| <b>FP_INT_UPWARD</b>           | <b>FP_FAST_DMULL</b>            | <b>math_errhandling</b> |
| <b>FP_INT_DOWNWARD</b>         | <b>FP_FAST_FDIV</b>             |                         |
| <b>FP_INT_TOWARDZERO</b>       | <b>FP_FAST_FDIVL</b>            |                         |

```
#pragma STDC FP_CONTRACT on-off-switch
int fpclassify(real-floating x);
int iscanonical(real-floating x);
int isfinite(real-floating x);
int isinf(real-floating x);
int isnan(real-floating x);
int isnormal(real-floating x);
int signbit(real-floating x);
int issignaling(real-floating x);
```



```

int issubnormal(real-floating x);
int iszero(real-floating x);
double acos(double x);
float acosf(float x);
long double acosl(long double x);
double asin(double x);
float asinf(float x);
long double asinl(long double x);
double atan(double x);
float atanf(float x);
long double atanl(long double x);
double atan2(double y, double x);
float atan2f(float y, float x);
long double atan2l(long double y, long double x);
double cos(double x);
float cosf(float x);
long double cosl(long double x);
double sin(double x);
float sinf(float x);
long double sinl(long double x);
double tan(double x);
float tanf(float x);
long double tanl(long double x);
double acospi(double x);
float acospif(float x);
long double acospil(long double x);
double asinpi(double x);
float asinpif(float x);
long double asinpil(long double x);
double atanpi(double x);
float atanpif(float x);
long double atanpil(long double x);
double atan2pi(double y, double x);
float atan2pif(float y, float x);
long double atan2pil(long double y, long double x);
double cospi(double x);
float cospif(float x);
long double cospil(long double x);
double sinpi(double x);
float sinpif(float x);
long double sinpil(long double x);
double tanpi(double x);
float tanpif(float x);
long double tanpil(long double x);
double acosh(double x);
float acoshf(float x);
long double acoshl(long double x);
double asinh(double x);
float asinhf(float x);
long double asinhl(long double x);
double atanh(double x);
float atanhf(float x);
long double atanhhl(long double x);
double cosh(double x);
float coshf(float x);
long double coshl(long double x);
double sinh(double x);
float sinhlf(float x);
long double sinhl(long double x);
double tanh(double x);
float tanhf(float x);

```

```

long double tanhl(long double x);
double exp(double x);
float expf(float x);
long double expl(long double x);
double exp10(double x);
float exp10f(float x);
long double exp10l(long double x);
double exp10m1(double x);
float exp10m1f(float x);
long double exp10m1l(long double x);
double exp2(double x);
float exp2f(float x);
long double exp2l(long double x);
double exp2m1(double x);
float exp2m1f(float x);
long double exp2m1l(long double x);
double expm1(double x);
float expm1f(float x);
long double expm1l(long double x);
double frexp(double value, int *p);
float frexpf(float value, int *p);
long double frexpl(long double value, int *p);
int ilogb(double x);
int ilogbf(float x);
int ilogbl(long double x);
double ldexp(double x, int p);
float ldexpf(float x, int p);
long double ldexpl(long double x, int p);
long int llogb(double x);
long int llogbf(float x);
long int llogbl(long double x);
double log(double x);
float logf(float x);
long double logl(long double x);
double log10(double x);
float log10f(float x);
long double log10l(long double x);
double log10p1(double x);
float log10p1f(float x);
long double log10p1l(long double x);
double log1p(double x);
float log1pf(float x);
long double log1pl(long double x);
double logp1(double x);
float logp1f(float x);
long double logp1l(long double x);
double log2(double x);
float log2f(float x);
long double log2l(long double x);
double log2p1(double x);
float log2p1f(float x);
long double log2p1l(long double x);
double logb(double x);
float logbf(float x);
long double logbl(long double x);
double modf(double value, double *iptr);
float modff(float value, float *iptr);
long double modfl(long double value, long double *iptr);
double scalbn(double x, int n);
float scalbnf(float x, int n);
long double scalbnl(long double x, int n);

```

```

double scalbln(double x, long int n);
float scalblnf(float x, long int n);
long double scalblnl(long double x, long int n);
double cbrt(double x);
float cbrtf(float x);
long double cbrtl(long double x);
double compoundn(double x, long long int n);
float compoundnf(float x, long long int n);
long double compoundnl(long double x, long long int n);
double fabs(double x);
float fabsf(float x);
long double fabsl(long double x);
double hypot(double x, double y);
float hypotf(float x, float y);
long double hypotl(long double x, long double y);
double pow(double x, double y);
float powf(float x, float y);
long double powl(long double x, long double y);
double pown(double x, long long int n);
float pownf(float x, long long int n);
long double pownl(long double x, long long int n);
double powr(double y, double x);
float powrf(float y, float x);
long double powrl(long double y, long double x);
double rootn(double x, long long int n);
float rootnf(float x, long long int n);
long double rootnl(long double x, long long int n);
double rsqrt(double x);
float rsqrtf(float x);
long double rsqrtl(long double x);
double sqrt(double x);
float sqrtf(float x);
long double sqrtl(long double x);
double erf(double x);
float erff(float x);
long double erfl(long double x);
double erfc(double x);
float erfcf(float x);
long double erfcl(long double x);
double lgamma(double x);
float lgammaf(float x);
long double lgammal(long double x);
double tgamma(double x);
float tgammaf(float x);
long double tgammaL(long double x);
double ceil(double x);
float ceilf(float x);
long double ceill(long double x);
double floor(double x);
float floorf(float x);
long double floorl(long double x);
double nearbyint(double x);
float nearbyintf(float x);
long double nearbyintl(long double x);
double rint(double x);
float rintf(float x);
long double rintl(long double x);
long int lrint(double x);
long int lrintf(float x);
long int lrintl(long double x);
long long int llrint(double x);

```

```

long long int llrintf(float x);
long long int llrintl(long double x);
double round(double x);
float roundf(float x);
long double roundl(long double x);
long int lround(double x);
long int lroundf(float x);
long int lroundl(long double x);
long long int llround(double x);
long long int llroundf(float x);
long long int llroundl(long double x);
double roundeven(double x);
float roundevenf(float x);
long double roundevenl(long double x);
double trunc(double x);
float truncf(float x);
long double truncf(long double x);
double fromfp(double x, int rnd, unsigned int width);
float fromfpf(float x, int rnd, unsigned int width);
long double fromfpl(long double x, int rnd, unsigned int width);
double ufromfp(double x, int rnd, unsigned int width);
float ufromfpf(float x, int rnd, unsigned int width);
long double ufromfpl(long double x, int rnd, unsigned int width);
double fromfpx(double x, int rnd, unsigned int width);
float fromfpxf(float x, int rnd, unsigned int width);
long double fromfpxl(long double x, int rnd, unsigned int width);
double ufromfpx(double x, int rnd, unsigned int width);
float ufromfpxf(float x, int rnd, unsigned int width);
long double ufromfpxl(long double x, int rnd, unsigned int width);
double fmod(double x, double y);
float fmodf(float x, float y);
long double fmodl(long double x, long double y);
double remainder(double x, double y);
float remainderf(float x, float y);
long double remainderl(long double x, long double y);
double remquo(double x, double y, int *quo);
float remquof(float x, float y, int *quo);
long double remquol(long double x, long double y, int *quo);
double copysign(double x, double y);
float copysignf(float x, float y);
long double copysignl(long double x, long double y);
double nan(const char *tagp);
float nanf(const char *tagp);
long double nanl(const char *tagp);
double nextafter(double x, double y);
float nextafterf(float x, float y);
long double nextafterl(long double x, long double y);
double nexttoward(double x, long double y);
float nexttowardf(float x, long double y);
long double nexttowardl(long double x, long double y);
double nextup(double x);
float nextupf(float x);
long double nextupl(long double x);
double nextdown(double x);
float nextdownf(float x);
long double nextdownl(long double x);
int canonicalize(double *cx, const double *x);
int canonicalizef(float *cx, const float *x);
int canonicalizel(long double *cx, const long double *x);
double fdim(double x, double y);
float fdimf(float x, float y);

```

```

long double fdiml(long double x, long double y);
double fmax(double x, double y);
float fmaxf(float x, float y);
long double fmaxl(long double x, long double y);
double fmin(double x, double y);
float fminf(float x, float y);
long double fminl(long double x, long double y);
double fmaximum(double x, double y);
float fmaximumf(float x, float y);
long double fmaximuml(long double x, long double y);
double fminimum(double x, double y);
float fminimumf(float x, float y);
long double fminimuml(long double x, long double y);
double fmaximum_mag(double x, double y);
float fmaximum_magf(float x, float y);
long double fmaximum_magl(long double x, long double y);
double fminimum_mag(double x, double y);
float fminimum_magf(float x, float y);
long double fminimum_magl(long double x, long double y);
double fmaximum_num(double x, double y);
float fmaximum_numf(float x, float y);
long double fmaximum_numl(long double x, long double y);
double fminimum_num(double x, double y);
float fminimum_numf(float x, float y);
long double fminimum_numl(long double x, long double y);
double fmaximum_mag_num(double x, double y);
float fmaximum_mag_numf(float x, float y);
long double fmaximum_mag_numl(long double x, long double y);
double fminimum_mag_num(double x, double y);
float fminimum_mag_numf(float x, float y);
long double fminimum_mag_numl(long double x, long double y);
double fma(double x, double y, double z);
float fmaf(float x, float y, float z);
long double fmal(long double x, long double y, long double z);
float fadd(double x, double y);
float faddl(long double x, long double y);
double daddl(long double x, long double y);
float fsub(double x, double y);
float fsubl(long double x, long double y);
double dsubl(long double x, long double y);
float fmul(double x, double y);
float fmul_l(long double x, long double y);
double dmul_l(long double x, long double y);
float fdiv(double x, double y);
float fdivl(long double x, long double y);
double ddivl(long double x, long double y);
float ffma(double x, double y, double z);
float ffmal(long double x, long double y, long double z);
double dfmal(long double x, long double y, long double z);
float fsqrt(double x);
float fsqrtl(long double x);
double dsqrtl(long double x);
int isgreater(real-floating x, real-floating y);
int isgreaterequal(real-floating x, real-floating y);
int isless(real-floating x, real-floating y);
int islessequal(real-floating x, real-floating y);
int islessgreater(real-floating x, real-floating y);
int isunordered(real-floating x, real-floating y);
int iseql(real-floating x, real-floating y);

```

Only if the implementation defines `__STDC_IEC_60559_DFP__`:

|                                 |                                 |                                  |
|---------------------------------|---------------------------------|----------------------------------|
| <code>_Decimal32_t</code>       | <code>FP_FAST_D32SUBD64</code>  | <code>FP_FAST_D32FMAD128</code>  |
| <code>_Decimal64_t</code>       | <code>FP_FAST_D32SUBD128</code> | <code>FP_FAST_D64FMAD128</code>  |
| <code>DEC_INFINITY</code>       | <code>FP_FAST_D64SUBD128</code> | <code>FP_FAST_D32SQRTD64</code>  |
| <code>DEC_NAN</code>            | <code>FP_FAST_D32MULD64</code>  | <code>FP_FAST_D32SQRTD128</code> |
| <code>FP_FAST_FMAD32</code>     | <code>FP_FAST_D32MULD128</code> | <code>FP_FAST_D64SQRTD128</code> |
| <code>FP_FAST_FMAD64</code>     | <code>FP_FAST_D64MULD128</code> | <code>HUGE_VAL_D32</code>        |
| <code>FP_FAST_FMAD128</code>    | <code>FP_FAST_D32DIVD64</code>  | <code>HUGE_VAL_D64</code>        |
| <code>FP_FAST_D32ADDD64</code>  | <code>FP_FAST_D32DIVD128</code> | <code>HUGE_VAL_D128</code>       |
| <code>FP_FAST_D32ADDD128</code> | <code>FP_FAST_D64DIVD128</code> |                                  |
| <code>FP_FAST_D64ADDD128</code> | <code>FP_FAST_D32FMAD64</code>  |                                  |

```

_Decimal32 acosd32(_Decimal32 x);
_Decimal64 acosd64(_Decimal64 x);
_Decimal128 acosd128(_Decimal128 x);
_Decimal32 asind32(_Decimal32 x);
_Decimal64 asind64(_Decimal64 x);
_Decimal128 asind128(_Decimal128 x);
_Decimal32 atand32(_Decimal32 x);
_Decimal64 atand64(_Decimal64 x);
_Decimal128 atand128(_Decimal128 x);
_Decimal32 atan2d32(_Decimal32 y, _Decimal32 x);
_Decimal64 atan2d64(_Decimal64 y, _Decimal64 x);
_Decimal128 atan2d128(_Decimal128 y, _Decimal128 x);
_Decimal32 cosd32(_Decimal32 x);
_Decimal64 cosd64(_Decimal64 x);
_Decimal128 cosd128(_Decimal128 x);
_Decimal32 sind32(_Decimal32 x);
_Decimal64 sind64(_Decimal64 x);
_Decimal128 sind128(_Decimal128 x);
_Decimal32 tand32(_Decimal32 x);
_Decimal64 tand64(_Decimal64 x);
_Decimal128 tand128(_Decimal128 x);
_Decimal32 acospid32(_Decimal32 x);
_Decimal64 acospid64(_Decimal64 x);
_Decimal128 acospid128(_Decimal128 x);
_Decimal32 asinpid32(_Decimal32 x);
_Decimal64 asinpid64(_Decimal64 x);
_Decimal128 asinpid128(_Decimal128 x);
_Decimal32 atanpid32(_Decimal32 x);
_Decimal64 atanpid64(_Decimal64 x);
_Decimal128 atanpid128(_Decimal128 x);
_Decimal32 atan2pid32(_Decimal32 y, _Decimal32 x);
_Decimal64 atan2pid64(_Decimal64 y, _Decimal64 x);
_Decimal128 atan2pid128(_Decimal128 y, _Decimal128 x);
_Decimal32 cospid32(_Decimal32 x);
_Decimal64 cospid64(_Decimal64 x);
_Decimal128 cospid128(_Decimal128 x);
_Decimal32 sinpid32(_Decimal32 x);
_Decimal64 sinpid64(_Decimal64 x);
_Decimal128 sinpid128(_Decimal128 x);
_Decimal32 tanpid32(_Decimal32 x);
_Decimal64 tanpid64(_Decimal64 x);
_Decimal128 tanpid128(_Decimal128 x);
_Decimal32 acoshd32(_Decimal32 x);
_Decimal64 acoshd64(_Decimal64 x);
_Decimal128 acoshd128(_Decimal128 x);
_Decimal32 asinhd32(_Decimal32 x);
_Decimal64 asinhd64(_Decimal64 x);
_Decimal128 asinhd128(_Decimal128 x);
_Decimal32 atanh32(_Decimal32 x);

```

```

_Decimal64 atanh64(_Decimal64 x);
_Decimal128 atanh128(_Decimal128 x);
_Decimal32 coshd32(_Decimal32 x);
_Decimal64 coshd64(_Decimal64 x);
_Decimal128 coshd128(_Decimal128 x);
_Decimal32 sinhd32(_Decimal32 x);
_Decimal64 sinhd64(_Decimal64 x);
_Decimal128 sinhd128(_Decimal128 x);
_Decimal32 tanhd32(_Decimal32 x);
_Decimal64 tanhd64(_Decimal64 x);
_Decimal128 tanhd128(_Decimal128 x);
_Decimal32 expd32(_Decimal32 x);
_Decimal64 expd64(_Decimal64 x);
_Decimal128 expd128(_Decimal128 x);
_Decimal32 exp10d32(_Decimal32 x);
_Decimal64 exp10d64(_Decimal64 x);
_Decimal128 exp10d128(_Decimal128 x);
_Decimal32 exp10md32(_Decimal32 x);
_Decimal64 exp10md64(_Decimal64 x);
_Decimal128 exp10md128(_Decimal128 x);
_Decimal32 exp2d32(_Decimal32 x);
_Decimal64 exp2d64(_Decimal64 x);
_Decimal128 exp2d128(_Decimal128 x);
_Decimal32 exp2md32(_Decimal32 x);
_Decimal64 exp2md64(_Decimal64 x);
_Decimal128 exp2md128(_Decimal128 x);
_Decimal32 expm1d32(_Decimal32 x);
_Decimal64 expm1d64(_Decimal64 x);
_Decimal128 expm1d128(_Decimal128 x);
_Decimal32 frexpd32(_Decimal32 value, int *p);
_Decimal64 frexpd64(_Decimal64 value, int *p);
_Decimal128 frexpd128(_Decimal128 value, int *p);
int ilogbd32(_Decimal32 x);
int ilogbd64(_Decimal64 x);
int ilogbd128(_Decimal128 x);
_Decimal32 ldexpd32(_Decimal32 x, int p);
_Decimal64 ldexpd64(_Decimal64 x, int p);
_Decimal128 ldexpd128(_Decimal128 x, int p);
long int llogbd32(_Decimal32 x);
long int llogbd64(_Decimal64 x);
long int llogbd128(_Decimal128 x);
_Decimal32 logd32(_Decimal32 x);
_Decimal64 logd64(_Decimal64 x);
_Decimal128 logd128(_Decimal128 x);
_Decimal32 log10d32(_Decimal32 x);
_Decimal64 log10d64(_Decimal64 x);
_Decimal128 log10d128(_Decimal128 x);
_Decimal32 log10pd32(_Decimal32 x);
_Decimal64 log10pd64(_Decimal64 x);
_Decimal128 log10pd128(_Decimal128 x);
_Decimal32 log1pd32(_Decimal32 x);
_Decimal64 log1pd64(_Decimal64 x);
_Decimal128 log1pd128(_Decimal128 x);
_Decimal32 logp1d32(_Decimal32 x);
_Decimal64 logp1d64(_Decimal64 x);
_Decimal128 logp1d128(_Decimal128 x);
_Decimal32 log2d32(_Decimal32 x);
_Decimal64 log2d64(_Decimal64 x);
_Decimal128 log2d128(_Decimal128 x);
_Decimal32 log2pd32(_Decimal32 x);
_Decimal64 log2pd64(_Decimal64 x);

```



```

_Decimal128 log2p1d128(_Decimal128 x);
_Decimal32 logbd32(_Decimal32 x);
_Decimal64 logbd64(_Decimal64 x);
_Decimal128 logbd128(_Decimal128 x);
_Decimal32 modfd32(_Decimal32 x, _Decimal32 *iptr);
_Decimal64 modfd64(_Decimal64 x, _Decimal64 *iptr);
_Decimal128 modfd128(_Decimal128 x, _Decimal128 *iptr);
_Decimal32 scalbnd32(_Decimal32 x, int n);
_Decimal64 scalbnd64(_Decimal64 x, int n);
_Decimal128 scalbnd128(_Decimal128 x, int n);
_Decimal32 scalblnd32(_Decimal32 x, long int n);
_Decimal64 scalblnd64(_Decimal64 x, long int n);
_Decimal128 scalblnd128(_Decimal128 x, long int n);
_Decimal32 cbrtd32(_Decimal32 x);
_Decimal64 cbrtd64(_Decimal64 x);
_Decimal128 cbrtd128(_Decimal128 x);
_Decimal32 compoundnd32(_Decimal32 x, long long int n);
_Decimal64 compoundnd64(_Decimal64 x, long long int n);
_Decimal128 compoundnd128(_Decimal128 x, long long int n);
_Decimal32 fabsd32(_Decimal32 x);
_Decimal64 fabsd64(_Decimal64 x);
_Decimal128 fabsd128(_Decimal128 x);
_Decimal32 hypotd32(_Decimal32 x, _Decimal32 y);
_Decimal64 hypotd64(_Decimal64 x, _Decimal64 y);
_Decimal128 hypotd128(_Decimal128 x, _Decimal128 y);
_Decimal32 powd32(_Decimal32 x, _Decimal32 y);
_Decimal64 powd64(_Decimal64 x, _Decimal64 y);
_Decimal128 powd128(_Decimal128 x, _Decimal128 y);
_Decimal32 pownd32(_Decimal32 x, long long int n);
_Decimal64 pownd64(_Decimal64 x, long long int n);
_Decimal128 pownd128(_Decimal128 x, long long int n);
_Decimal32 powrd32(_Decimal32 y, _Decimal32 x);
_Decimal64 powrd64(_Decimal64 y, _Decimal64 x);
_Decimal128 powrd128(_Decimal128 y, _Decimal128 x);
_Decimal32 rootnd32(_Decimal32 x, long long int n);
_Decimal64 rootnd64(_Decimal64 x, long long int n);
_Decimal128 rootnd128(_Decimal128 x, long long int n);
_Decimal32 rsqrt32(_Decimal32 x);
_Decimal64 rsqrt64(_Decimal64 x);
_Decimal128 rsqrt128(_Decimal128 x);
_Decimal32 sqrt32(_Decimal32 x);
_Decimal64 sqrt64(_Decimal64 x);
_Decimal128 sqrt128(_Decimal128 x);
_Decimal32 erfd32(_Decimal32 x);
_Decimal64 erfd64(_Decimal64 x);
_Decimal128 erfd128(_Decimal128 x);
_Decimal32 erfcd32(_Decimal32 x);
_Decimal64 erfcd64(_Decimal64 x);
_Decimal128 erfcd128(_Decimal128 x);
_Decimal32 lgammad32(_Decimal32 x);
_Decimal64 lgammad64(_Decimal64 x);
_Decimal128 lgammad128(_Decimal128 x);
_Decimal32 tgammd32(_Decimal32 x);
_Decimal64 tgammd64(_Decimal64 x);
_Decimal128 tgammd128(_Decimal128 x);
_Decimal32 ceild32(_Decimal32 x);
_Decimal64 ceild64(_Decimal64 x);
_Decimal128 ceild128(_Decimal128 x);
_Decimal32 floord32(_Decimal32 x);
_Decimal64 floord64(_Decimal64 x);
_Decimal128 floord128(_Decimal128 x);

```



```

_Decimal32 nearbyintd32(_Decimal32 x);
_Decimal64 nearbyintd64(_Decimal64 x);
_Decimal128 nearbyintd128(_Decimal128 x);
_Decimal32 rintd32(_Decimal32 x);
_Decimal64 rintd64(_Decimal64 x);
_Decimal128 rintd128(_Decimal128 x);
long int lrintd32(_Decimal32 x);
long int lrintd64(_Decimal64 x);
long int lrintd128(_Decimal128 x);
long long int llrintd32(_Decimal32 x);
long long int llrintd64(_Decimal64 x);
long long int llrintd128(_Decimal128 x);
_Decimal32 roundd32(_Decimal32 x);
_Decimal64 roundd64(_Decimal64 x);
_Decimal128 roundd128(_Decimal128 x);
long int lroundd32(_Decimal32 x);
long int lroundd64(_Decimal64 x);
long int lroundd128(_Decimal128 x);
long long int llroundd32(_Decimal32 x);
long long int llroundd64(_Decimal64 x);
long long int llroundd128(_Decimal128 x);
_Decimal32 roundevend32(_Decimal32 x);
_Decimal64 roundevend64(_Decimal64 x);
_Decimal128 roundevend128(_Decimal128 x);
_Decimal32 truncd32(_Decimal32 x);
_Decimal64 truncd64(_Decimal64 x);
_Decimal128 truncd128(_Decimal128 x);
_Decimal32 fromfpd32(_Decimal32 x, int rnd, unsigned int width);
_Decimal64 fromfpd64(_Decimal64 x, int rnd, unsigned int width);
_Decimal128 fromfpd128(_Decimal128 x, int rnd, unsigned int width);
_Decimal32 ufromfpd32(_Decimal32 x, int rnd, unsigned int width);
_Decimal64 ufromfpd64(_Decimal64 x, int rnd, unsigned int width);
_Decimal128 ufromfpd128(_Decimal128 x, int rnd, unsigned int width);
_Decimal32 fromfpdx32(_Decimal32 x, int rnd, unsigned int width);
_Decimal64 fromfpdx64(_Decimal64 x, int rnd, unsigned int width);
_Decimal128 fromfpdx128(_Decimal128 x, int rnd, unsigned int width);
_Decimal32 ufromfpdx32(_Decimal32 x, int rnd, unsigned int width);
_Decimal64 ufromfpdx64(_Decimal64 x, int rnd, unsigned int width);
_Decimal128 ufromfpdx128(_Decimal128 x, int rnd, unsigned int width);
_Decimal32 fmodd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmodd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmodd128(_Decimal128 x, _Decimal128 y);
_Decimal32 remainderd32(_Decimal32 x, _Decimal32 y);
_Decimal64 remainderd64(_Decimal64 x, _Decimal64 y);
_Decimal128 remainderd128(_Decimal128 x, _Decimal128 y);
_Decimal32 copysignd32(_Decimal32 x, _Decimal32 y);
_Decimal64 copysignd64(_Decimal64 x, _Decimal64 y);
_Decimal128 copysignd128(_Decimal128 x, _Decimal128 y);
_Decimal32 nand32(const char *tagp);
_Decimal64 nand64(const char *tagp);
_Decimal128 nand128(const char *tagp);
_Decimal32 nextafterd32(_Decimal32 x, _Decimal32 y);
_Decimal64 nextafterd64(_Decimal64 x, _Decimal64 y);
_Decimal128 nextafterd128(_Decimal128 x, _Decimal128 y);
_Decimal32 nexttowardd32(_Decimal32 x, _Decimal128 y);
_Decimal64 nexttowardd64(_Decimal64 x, _Decimal128 y);
_Decimal128 nexttowardd128(_Decimal128 x, _Decimal128 y);
_Decimal32 nexttupd32(_Decimal32 x);
_Decimal64 nexttupd64(_Decimal64 x);
_Decimal128 nexttupd128(_Decimal128 x);
_Decimal32 nexttdownd32(_Decimal32 x);

```

```

_Decimal64 nextdownd64(_Decimal64 x);
_Decimal128 nextdownd128(_Decimal128 x);
int canonicalized32(_Decimal32 *cx, const _Decimal32 *x);
int canonicalized64(_Decimal64 *cx, const _Decimal64 *x);
int canonicalized128(_Decimal128 *cx, const _Decimal128 *x);
_Decimal32 fdimd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fdimd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fdimd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmaxd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaxd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaxd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmind32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmind64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmind128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmaximumd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaximumd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaximumd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fminimumd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimumd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimumd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmaximum_magd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaximum_magd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaximum_magd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fminimum_magd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimum_magd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimum_magd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmaximum_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaximum_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaximum_numd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fminimum_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimum_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimum_numd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmaximum_mag_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fmaximum_mag_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fmaximum_mag_numd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fminimum_mag_numd32(_Decimal32 x, _Decimal32 y);
_Decimal64 fminimum_mag_numd64(_Decimal64 x, _Decimal64 y);
_Decimal128 fminimum_mag_numd128(_Decimal128 x, _Decimal128 y);
_Decimal32 fmad32(_Decimal32 x, _Decimal32 y, _Decimal32 z);
_Decimal64 fmad64(_Decimal64 x, _Decimal64 y, _Decimal64 z);
_Decimal128 fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
_Decimal32 d32addd64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32addd128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64addd128(_Decimal128 x, _Decimal128 y);
_Decimal32 d32subd64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32subd128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64subd128(_Decimal128 x, _Decimal128 y);
_Decimal32 d32muld64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32muld128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64muld128(_Decimal128 x, _Decimal128 y);
_Decimal32 d32divd64(_Decimal64 x, _Decimal64 y);
_Decimal32 d32divd128(_Decimal128 x, _Decimal128 y);
_Decimal64 d64divd128(_Decimal128 x, _Decimal128 y);
_Decimal32 d32fmad64(_Decimal64 x, _Decimal64 y, _Decimal64 z);
_Decimal32 d32fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
_Decimal64 d64fmad128(_Decimal128 x, _Decimal128 y, _Decimal128 z);
_Decimal32 d32sqrtd64(_Decimal64 x);
_Decimal32 d32sqrtd128(_Decimal128 x);
_Decimal64 d64sqrtd128(_Decimal128 x);
_Decimal32 quantized32(_Decimal32 x, _Decimal32 y);
_Decimal64 quantized64(_Decimal64 x, _Decimal64 y);

```

```

_Decimal128 quantized128(_Decimal128 x, _Decimal128 y);
bool samequantumd32(_Decimal32 x, _Decimal32 y);
bool samequantumd64(_Decimal64 x, _Decimal64 y);
bool samequantumd128(_Decimal128 x, _Decimal128 y);
_Decimal32 quantumd32(_Decimal32 x);
_Decimal64 quantumd64(_Decimal64 x);
_Decimal128 quantumd128(_Decimal128 x);
long long int llquantexpd32(_Decimal32 x);
long long int llquantexpd64(_Decimal64 x);
long long int llquantexpd128(_Decimal128 x);
void encodedecd32(unsigned char encptr[restrict static 4],
 const _Decimal32 * restrict xptr);
void encodedecd64(unsigned char encptr[restrict static 8],
 const _Decimal64 * restrict xptr);
void encodedecd128(unsigned char encptr[restrict static 16],
 const _Decimal128 * restrict xptr);
void decodedecd32(_Decimal32 * restrict xptr,
 const unsigned char encptr[restrict static 4]);
void decodedecd64(_Decimal64 * restrict xptr,
 const unsigned char encptr[restrict static 8]);
void decodedecd128(_Decimal128 * restrict xptr,
 const unsigned char encptr[restrict static 16]);
void encodebind32(unsigned char encptr[restrict static 4],
 const _Decimal32 * restrict xptr);
void encodebind64(unsigned char encptr[restrict static 8],
 const _Decimal64 * restrict xptr);
void encodebind128(unsigned char encptr[restrict static 16],
 const _Decimal128 * restrict xptr);
void decodebind32(_Decimal32 * restrict xptr,
 const unsigned char encptr[restrict static 4]);
void decodebind64(_Decimal64 * restrict xptr,
 const unsigned char encptr[restrict static 8]);
void decodebind128(_Decimal128 * restrict xptr,
 const unsigned char encptr[restrict static 16]);

```

Only if the implementation defines `__STDC_IEC_60559_BFP__` or `__STDC_IEC_559__` and additionally the user code defines `__STDC_WANT_IEC_60559_EXT__` before any inclusion of `<math.h>`:

```

int totalorder(const double *x, const double *y);
int totalorderf(const float *x, const float *y);
int totalorderl(const long double *x, const long double *y);
int totalordermag(const double *x, const double *y);
int totalordermagf(const float *x, const float *y);
int totalordermagl(const long double *x, const long double *y);
double getpayload(const double *x);
float getpayloadf(const float *x);
long double getpayloadl(const long double *x);
int setpayload(double *res, double pl);
int setpayloadf(float *res, float pl);
int setpayloadl(long double *res, long double pl);
int setpayloadsig(double *res, double pl);
int setpayloadsigf(float *res, float pl);
int setpayloadsigl(long double *res, long double pl);

```

Only if the implementation defines `__STDC_IEC_60559_DFP__` and additionally the user code defines `__STDC_WANT_IEC_60559_EXT__` before any inclusion of `<math.h>`:

```

int totalorderd32(const _Decimal32 *x, const _Decimal32 *y);
int totalorderd64(const _Decimal64 *x, const _Decimal64 *y);
int totalorderd128(const _Decimal128 *x, const _Decimal128 *y);
int totalordermagd32(const _Decimal32 *x, const _Decimal32 *y);

```

```

int totalordermagd64(const _Decimal64 *x, const _Decimal64 *y);
int totalordermagd128(const _Decimal128 *x, const _Decimal128 *y);
_Decimal32 getpayloadd32(const _Decimal32 *x);
_Decimal64 getpayloadd64(const _Decimal64 *x);
_Decimal128 getpayloadd128(const _Decimal128 *x);
int setpayloadd32(_Decimal32 *res, _Decimal32 pl);
int setpayloadd64(_Decimal64 *res, _Decimal64 pl);
int setpayloadd128(_Decimal128 *res, _Decimal128 pl);
int setpayloadsigd32(_Decimal32 *res, _Decimal32 pl);
int setpayloadsigd64(_Decimal64 *res, _Decimal64 pl);
int setpayloadsigd128(_Decimal128 *res, _Decimal128 pl);

```

Only if the implementation defines `__STDC_IEC_60559_TYPES__` and additionally the user code defines `__STDC_WANT_IEC_60559_TYPES_EXT__` before any inclusion of `<math.h>`:

|                              |                                |                                |
|------------------------------|--------------------------------|--------------------------------|
| <code>long_double_t</code>   | <code>FP_FAST_FMSUBFN</code>   | <code>FP_FAST_DMDDIVDN</code>  |
| <code>_FloatN_t</code>       | <code>FP_FAST_FMXSUBFN</code>  | <code>FP_FAST_DMXXDIVDN</code> |
| <code>_DecimalN_t</code>     | <code>FP_FAST_FMXSUBFN</code>  | <code>FP_FAST_DMXXDIVDN</code> |
| <code>HUGE_VAL_FN</code>     | <code>FP_FAST_DMSUBDN</code>   | <code>FP_FAST_FMFMAFN</code>   |
| <code>HUGE_VAL_DN</code>     | <code>FP_FAST_DMSUBDN</code>   | <code>FP_FAST_FMFMAFN</code>   |
| <code>HUGE_VAL_FNX</code>    | <code>FP_FAST_DMXXSUBDN</code> | <code>FP_FAST_FMXFMAFN</code>  |
| <code>HUGE_VAL_DNX</code>    | <code>FP_FAST_DMXXSUBDN</code> | <code>FP_FAST_FMXFMAFN</code>  |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMMULFN</code>   | <code>FP_FAST_DMFMADN</code>   |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMMULFN</code>   | <code>FP_FAST_DMFMADN</code>   |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMXMULFN</code>  | <code>FP_FAST_DMXXFMADN</code> |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMXMULFN</code>  | <code>FP_FAST_DMXXFMADN</code> |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMMULDN</code>   | <code>FP_FAST_FMSQRTFN</code>  |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMMULDN</code>   | <code>FP_FAST_FMSQRTFN</code>  |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_DMXXMULDN</code> | <code>FP_FAST_FMXSQRTFN</code> |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_DMXXMULDN</code> | <code>FP_FAST_FMXSQRTFN</code> |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMDIVFN</code>   | <code>FP_FAST_DMSQRTDN</code>  |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMDIVFN</code>   | <code>FP_FAST_DMSQRTDN</code>  |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMXDIVFN</code>  | <code>FP_FAST_DMXSQRTDN</code> |
| <code>FP_FAST_FMAFN</code>   | <code>FP_FAST_FMXDIVFN</code>  | <code>FP_FAST_DMXSQRTDN</code> |
| <code>FP_FAST_FMSUBFN</code> | <code>FP_FAST_DMDDIVDN</code>  |                                |

```

_FloatN acosfN(_FloatN x);
_FloatNx acosfNx(_FloatNx x);
_DecimalN acosdN(_DecimalN x);
_DecimalNx acosdNx(_DecimalNx x);
_FloatN asinfN(_FloatN x);
_FloatNx asinfNx(_FloatNx x);
_DecimalN asindN(_DecimalN x);
_DecimalNx asindNx(_DecimalNx x);
_FloatN atanfN(_FloatN x);
_FloatNx atanfNx(_FloatNx x);
_DecimalN atandN(_DecimalN x);
_DecimalNx atandNx(_DecimalNx x);
_FloatN atan2fN(_FloatN y, _FloatN x);
_FloatNx atan2fNx(_FloatNx y, _FloatNx x);
_DecimalN atan2dN(_DecimalN y, _DecimalN x);
_DecimalNx atan2dNx(_DecimalNx y, _DecimalNx x);
_FloatN cosfN(_FloatN x);
_FloatNx cosfNx(_FloatNx x);
_DecimalN cosdN(_DecimalN x);
_DecimalNx cosdNx(_DecimalNx x);
_FloatN sinfN(_FloatN x);
_FloatNx sinfNx(_FloatNx x);

```

```

_DecimalN sindN(_DecimalN x);
_DecimalNx sindNx(_DecimalNx x);
_FloatN tanfN(_FloatN x);
_FloatNx tanfNx(_FloatNx x);
_DecimalN tandN(_DecimalN x);
_DecimalNx tandNx(_DecimalNx x);
_FloatN acospifN(_FloatN x);
_FloatNx acospifNx(_FloatNx x);
_DecimalN acospidN(_DecimalN x);
_DecimalNx acospidNx(_DecimalNx x);
_FloatN asinpifN(_FloatN x);
_FloatNx asinpifNx(_FloatNx x);
_DecimalN asinpidN(_DecimalN x);
_DecimalNx asinpidNx(_DecimalNx x);
_FloatN atanpifN(_FloatN x);
_FloatNx atanpifNx(_FloatNx x);
_DecimalN atanpidN(_DecimalN x);
_DecimalNx atanpidNx(_DecimalNx x);
_FloatN atan2pifN(_FloatN y, _FloatN x);
_FloatNx atan2pifNx(_FloatNx y, _FloatNx x);
_DecimalN atan2pidN(_DecimalN y, _DecimalN x);
_DecimalNx atan2pidNx(_DecimalNx y, _DecimalNx x);
_FloatN cospifN(_FloatN x);
_FloatNx cospifNx(_FloatNx x);
_DecimalN cospidN(_DecimalN x);
_DecimalNx cospidNx(_DecimalNx x);
_FloatN sinpifN(_FloatN x);
_FloatNx sinpifNx(_FloatNx x);
_DecimalN sinpidN(_DecimalN x);
_DecimalNx sinpidNx(_DecimalNx x);
_FloatN tanpifN(_FloatN x);
_FloatNx tanpifNx(_FloatNx x);
_DecimalN tanpidN(_DecimalN x);
_DecimalNx tanpidNx(_DecimalNx x);
_FloatN acoshfN(_FloatN x);
_FloatNx acoshfNx(_FloatNx x);
_DecimalN acoshdN(_DecimalN x);
_DecimalNx acoshdNx(_DecimalNx x);
_FloatN asinhfN(_FloatN x);
_FloatNx asinhfNx(_FloatNx x);
_DecimalN asinhdN(_DecimalN x);
_DecimalNx asinhdNx(_DecimalNx x);
_FloatN atanhfN(_FloatN x);
_FloatNx atanhfNx(_FloatNx x);
_DecimalN atanhdN(_DecimalN x);
_DecimalNx atanhdNx(_DecimalNx x);
_FloatN coshfN(_FloatN x);
_FloatNx coshfNx(_FloatNx x);
_DecimalN coshdN(_DecimalN x);
_DecimalNx coshdNx(_DecimalNx x);
_FloatN sinhfn(_FloatN x);
_FloatNx sinhfnNx(_FloatNx x);
_DecimalN sinhndN(_DecimalN x);
_DecimalNx sinhndNx(_DecimalNx x);
_FloatN tanhfN(_FloatN x);
_FloatNx tanhfNx(_FloatNx x);
_DecimalN tanhdN(_DecimalN x);
_DecimalNx tanhdNx(_DecimalNx x);
_FloatN expfN(_FloatN x);
_FloatNx expfNx(_FloatNx x);
_DecimalN expdN(_DecimalN x);

```

```

_DecimalNx expdNx(_DecimalNx x);
_FloatN exp10fN(_FloatN x);
_FloatNx exp10fNx(_FloatNx x);
_DecimalN exp10dN(_DecimalN x);
_DecimalNx exp10dNx(_DecimalNx x);
_FloatN exp10m1fN(_FloatN x);
_FloatNx exp10m1fNx(_FloatNx x);
_DecimalN exp10m1dN(_DecimalN x);
_DecimalNx exp10m1dNx(_DecimalNx x);
_FloatN exp2fN(_FloatN x);
_FloatNx exp2fNx(_FloatNx x);
_DecimalN exp2dN(_DecimalN x);
_DecimalNx exp2dNx(_DecimalNx x);
_FloatN exp2m1fN(_FloatN x);
_FloatNx exp2m1fNx(_FloatNx x);
_DecimalN exp2m1dN(_DecimalN x);
_DecimalNx exp2m1dNx(_DecimalNx x);
_FloatN expm1fN(_FloatN x);
_FloatNx expm1fNx(_FloatNx x);
_DecimalN expm1dN(_DecimalN x);
_DecimalNx expm1dNx(_DecimalNx x);
_FloatN frexpfN(_FloatN value, int *exp);
_FloatNx frexpfNx(_FloatNx value, int *exp);
_DecimalN frexpdN(_DecimalN value, int *exp);
_DecimalNx frexpdNx(_DecimalNx value, int *exp);
int ilogbfN(_FloatN x);
int ilogbfNx(_FloatNx x);
int ilogbdN(_DecimalN x);
int ilogbdNx(_DecimalNx x);
_FloatN ldexpfN(_FloatN value, int exp);
_FloatNx ldexpfNx(_FloatNx value, int exp);
_DecimalN ldexpdN(_DecimalN value, int exp);
_DecimalNx ldexpdNx(_DecimalNx value, int exp);
long int llogbfN(_FloatN x);
long int llogbfNx(_FloatNx x);
long int llogbdN(_DecimalN x);
long int llogbdNx(_DecimalNx x);
_FloatN logfN(_FloatN x);
_FloatNx logfNx(_FloatNx x);
_DecimalN logdN(_DecimalN x);
_DecimalNx logdNx(_DecimalNx x);
_FloatN log10fN(_FloatN x);
_FloatNx log10fNx(_FloatNx x);
_DecimalN log10dN(_DecimalN x);
_DecimalNx log10dNx(_DecimalNx x);
_FloatN log10p1fN(_FloatN x);
_FloatNx log10p1fNx(_FloatNx x);
_DecimalN log10p1dN(_DecimalN x);
_DecimalNx log10p1dNx(_DecimalNx x);
_FloatN log1pfN(_FloatN x);
_FloatNx log1pfNx(_FloatNx x);
_FloatN logp1fN(_FloatN x);
_FloatNx logp1fNx(_FloatNx x);
_DecimalN log1pdN(_DecimalN x);
_DecimalNx log1pdNx(_DecimalNx x);
_DecimalN logp1dN(_DecimalN x);
_DecimalNx logp1dNx(_DecimalNx x);
_FloatN log2fN(_FloatN x);
_FloatNx log2fNx(_FloatNx x);
_DecimalN log2dN(_DecimalN x);
_DecimalNx log2dNx(_DecimalNx x);

```



```

_FloatN log2p1fN(_FloatN x);
_FloatNx log2p1fNx(_FloatNx x);
_DecimalN log2p1dN(_DecimalN x);
_DecimalNx log2p1dNx(_DecimalNx x);
_FloatN logbfN(_FloatN x);
_FloatNx logbfNx(_FloatNx x);
_DecimalN logbdN(_DecimalN x);
_DecimalNx logbdNx(_DecimalNx x);
_FloatN modffN(_FloatN x, _FloatN *iptr);
_FloatNx modffNx(_FloatNx x, _FloatNx *iptr);
_DecimalN modfdN(_DecimalN x, _DecimalN *iptr);
_DecimalNx modfdNx(_DecimalNx x, _DecimalNx *iptr);
_FloatN scalbnfN(_FloatN value, int exp);
_FloatNx scalbnfNx(_FloatNx value, int exp);
_DecimalN scalbndN(_DecimalN value, int exp);
_DecimalNx scalbndNx(_DecimalNx value, int exp);
_FloatN scalblnfN(_FloatN value, long int exp);
_FloatNx scalblnfNx(_FloatNx value, long int exp);
_DecimalN scalblndN(_DecimalN value, long int exp);
_DecimalNx scalblndNx(_DecimalNx value, long int exp);
_FloatN cbrtfN(_FloatN x);
_FloatNx cbrtfNx(_FloatNx x);
_DecimalN cbrtdN(_DecimalN x);
_DecimalNx cbrtdNx(_DecimalNx x);
_FloatN compoundnfN(_FloatN x, long long int n);
_FloatNx compoundnfNx(_FloatNx x, long long int n);
_DecimalN compoundndN(_DecimalN x, long long int n);
_DecimalNx compoundndNx(_DecimalNx x, long long int n);
_FloatN fabsfN(_FloatN x);
_FloatNx fabsfNx(_FloatNx x);
_DecimalN fabsdN(_DecimalN x);
_DecimalNx fabsdNx(_DecimalNx x);
_FloatN hypotfN(_FloatN x, _FloatN y);
_FloatNx hypotfNx(_FloatNx x, _FloatNx y);
_DecimalN hypotdN(_DecimalN x, _DecimalN y);
_DecimalNx hypotdNx(_DecimalNx x, _DecimalNx y);
_FloatN powfN(_FloatN x, _FloatN y);
_FloatNx powfNx(_FloatNx x, _FloatNx y);
_DecimalN powdN(_DecimalN x, _DecimalN y);
_DecimalNx powdNx(_DecimalNx x, _DecimalNx y);
_FloatN pownfN(_FloatN x, long long int n);
_FloatNx pownfNx(_FloatNx x, long long int n);
_DecimalN powndN(_DecimalN x, long long int n);
_DecimalNx powndNx(_DecimalNx x, long long int n);
_FloatN powrfN(_FloatN x, _FloatN y);
_FloatNx powrfNx(_FloatNx x, _FloatNx y);
_DecimalN powrdN(_DecimalN x, _DecimalN y);
_DecimalNx powrdNx(_DecimalNx x, _DecimalNx y);
_FloatN rootnfN(_FloatN x, long long int n);
_FloatNx rootnfNx(_FloatNx x, long long int n);
_DecimalN rootndN(_DecimalN x, long long int n);
_DecimalNx rootndNx(_DecimalNx x, long long int n);
_FloatN rsqrtfN(_FloatN x);
_FloatNx rsqrtfNx(_FloatNx x);
_DecimalN rsqrtndN(_DecimalN x);
_DecimalNx rsqrtndNx(_DecimalNx x);
_FloatN sqrtfN(_FloatN x);
_FloatNx sqrtfNx(_FloatNx x);
_DecimalN sqrtdN(_DecimalN x);
_DecimalNx sqrtdNx(_DecimalNx x);
_FloatN erffN(_FloatN x);

```

```

_FloatNx erffNx(_FloatNx x);
_DecimalN erfdN(_DecimalN x);
_DecimalNx erfdNx(_DecimalNx x);
_FloatN erfcfN(_FloatN x);
_FloatNx erfcfNx(_FloatNx x);
_DecimalN erfcdN(_DecimalN x);
_DecimalNx erfcdNx(_DecimalNx x);
_FloatN lgammafN(_FloatN x);
_FloatNx lgammafNx(_FloatNx x);
_DecimalN lgammadN(_DecimalN x);
_DecimalNx lgammadNx(_DecimalNx x);
_FloatN tgammafN(_FloatN x);
_FloatNx tgammafNx(_FloatNx x);
_DecimalN tgammaN(_DecimalN x);
_DecimalNx tgammaNx(_DecimalNx x);
_FloatN ceilfN(_FloatN x);
_FloatNx ceilfNx(_FloatNx x);
_DecimalN ceildN(_DecimalN x);
_DecimalNx ceildNx(_DecimalNx x);
_FloatN floorfN(_FloatN x);
_FloatNx floorfNx(_FloatNx x);
_DecimalN floordN(_DecimalN x);
_DecimalNx floordNx(_DecimalNx x);
_FloatN nearbyintfN(_FloatN x);
_FloatNx nearbyintfNx(_FloatNx x);
_DecimalN nearbyintdN(_DecimalN x);
_DecimalNx nearbyintdNx(_DecimalNx x);
_FloatN rintfN(_FloatN x);
_FloatNx rintfNx(_FloatNx x);
_DecimalN rintdN(_DecimalN x);
_DecimalNx rintdNx(_DecimalNx x);
long int lrintfN(_FloatN x);
long int lrintfNx(_FloatNx x);
long int lrintdN(_DecimalN x);
long int lrintdNx(_DecimalNx x);
long long int llrintfN(_FloatN x);
long long int llrintfNx(_FloatNx x);
long long int llrintdN(_DecimalN x);
long long int llrintdNx(_DecimalNx x);
_FloatN roundfN(_FloatN x);
_FloatNx roundfNx(_FloatNx x);
_DecimalN rounddN(_DecimalN x);
_DecimalNx rounddNx(_DecimalNx x);
long int lroundfN(_FloatN x);
long int lroundfNx(_FloatNx x);
long int lrounddN(_DecimalN x);
long int lrounddNx(_DecimalNx x);
long long int llroundfN(_FloatN x);
long long int llroundfNx(_FloatNx x);
long long int llrounddN(_DecimalN x);
long long int llrounddNx(_DecimalNx x);
_FloatN roundevenfN(_FloatN x);
_FloatNx roundevenfNx(_FloatNx x);
_DecimalN roundevendN(_DecimalN x);
_DecimalNx roundevendNx(_DecimalNx x);
_FloatN truncfN(_FloatN x);
_FloatNx truncfNx(_FloatNx x);
_DecimalN truncdN(_DecimalN x);
_DecimalNx truncdNx(_DecimalNx x);
_FloatN fromfpfN(_FloatN x, int rnd, unsigned int width);
_FloatNx fromfpfNx(_FloatNx x, int rnd, unsigned int width);

```



```

_DecimalN fromfpdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx fromfpdNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN ufromfpfN(_FloatN x, int rnd, unsigned int width);
_FloatNx ufromfpfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN ufromfpdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx ufromfpdNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN fromfpxfN(_FloatN x, int rnd, unsigned int width);
_FloatNx fromfpxfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN fromfpdxN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx fromfpdxNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN ufromfpxfN(_FloatN x, int rnd, unsigned int width);
_FloatNx ufromfpxfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN ufromfpdxN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx ufromfpdxNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN fmodfN(_FloatN x, _FloatN y);
_FloatNx fmodfNx(_FloatNx x, _FloatNx y);
_DecimalN fmoddN(_DecimalN x, _DecimalN y);
_DecimalNx fmoddNx(_DecimalNx x, _DecimalNx y);
_FloatN remainderfN(_FloatN x, _FloatN y);
_FloatNx remainderfNx(_FloatNx x, _FloatNx y);
_DecimalN remainderdN(_DecimalN x, _DecimalN y);
_DecimalNx remainderdNx(_DecimalNx x, _DecimalNx y);
_FloatN remquofN(_FloatN x, _FloatN y, int *quo);
_FloatNx remquofNx(_FloatNx x, _FloatNx y, int *quo);
_FloatN copysignfN(_FloatN x, _FloatN y);
_FloatNx copysignfNx(_FloatNx x, _FloatNx y);
_DecimalN copysigndN(_DecimalN x, _DecimalN y);
_DecimalNx copysigndNx(_DecimalNx x, _DecimalNx y);
_FloatN nanfN(const char *tagp);
_FloatNx nanfNx(const char *tagp);
_DecimalN nandN(const char *tagp);
_DecimalNx nandNx(const char *tagp);
_FloatN nextafterfN(_FloatN x, _FloatN y);
_FloatNx nextafterfNx(_FloatNx x, _FloatNx y);
_DecimalN nextafterdN(_DecimalN x, _DecimalN y);
_DecimalNx nextafterdNx(_DecimalNx x, _DecimalNx y);
_FloatN nextupfN(_FloatN x);
_FloatNx nextupfNx(_FloatNx x);
_DecimalN nextupdN(_DecimalN x);
_DecimalNx nextupdNx(_DecimalNx x);
_FloatN nextdownfN(_FloatN x);
_FloatNx nextdownfNx(_FloatNx x);
_DecimalN nextdowndN(_DecimalN x);
_DecimalNx nextdowndNx(_DecimalNx x);
int canonicalizefN(_FloatN *cx, const _FloatN *x);
int canonicalizefNx(_FloatNx *cx, const _FloatNx *x);
int canonicalizedN(_DecimalN *cx, const _DecimalN *x);
int canonicalizedNx(_DecimalNx *cx, const _DecimalNx *x);
_FloatN fdimfN(_FloatN x, _FloatN y);
_FloatNx fdimfNx(_FloatNx x, _FloatNx y);
_DecimalN fdimdN(_DecimalN x, _DecimalN y);
_DecimalNx fdimdNx(_DecimalNx x, _DecimalNx y);
_FloatN fmaximumfN(_FloatN x, _FloatN y);
_FloatNx fmaximumfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximumdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximumdNx(_DecimalNx x, _DecimalNx y);
_FloatN fminimumfN(_FloatN x, _FloatN y);
_FloatNx fminimumfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimumdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimumdNx(_DecimalNx x, _DecimalNx y);
_FloatN fmaximum_magfN(_FloatN x, _FloatN y);

```

```

_FloatNx fmaximum_magfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_magdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_magdNx(_DecimalNx x, _DecimalNx y);
_FloatN fminimum_magfN(_FloatN x, _FloatN y);
_FloatNx fminimum_magfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_magdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_magdNx(_DecimalNx x, _DecimalNx y);
_FloatN fmaximum_numfN(_FloatN x, _FloatN y);
_FloatNx fmaximum_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_numdNx(_DecimalNx x, _DecimalNx y);
_FloatN fminimum_numfN(_FloatN x, _FloatN y);
_FloatNx fminimum_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_numdNx(_DecimalNx x, _DecimalNx y);
_FloatN fmaximum_mag_numfN(_FloatN x, _FloatN y);
_FloatNx fmaximum_mag_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_mag_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_mag_numdNx(_DecimalNx x, _DecimalNx y);
_FloatN fminimum_mag_numfN(_FloatN x, _FloatN y);
_FloatNx fminimum_mag_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_mag_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_mag_numdNx(_DecimalNx x, _DecimalNx y);
_FloatN fmafN(_FloatN x, _FloatN y, _FloatN z);
_FloatNx fmafNx(_FloatNx x, _FloatNx y, _FloatNx z);
_DecimalN fmadN(_DecimalN x, _DecimalN y, _DecimalN z);
_DecimalNx fmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z);
_FloatM fMaddfN(_FloatN x, _FloatN y); // M < N
_FloatM fMaddfNx(_FloatNx x, _FloatNx y); // M ≤ N
_FloatMx fMxaddfN(_FloatN x, _FloatN y); // M < N
_FloatMx fMxaddfNx(_FloatNx x, _FloatNx y); // M < N
_DecimalM dMadddN(_DecimalN x, _DecimalN y); // M < N
_DecimalM dMadddNx(_DecimalNx x, _DecimalNx y); // M ≤ N
_DecimalMx dMxadddN(_DecimalN x, _DecimalN y); // M < N
_DecimalMx dMxadddNx(_DecimalNx x, _DecimalNx y); // M < N
_FloatM fMsubfN(_FloatN x, _FloatN y); // M < N
_FloatM fMsubfNx(_FloatNx x, _FloatNx y); // M ≤ N
_FloatMx fMxsubfN(_FloatN x, _FloatN y); // M < N
_FloatMx fMxsubfNx(_FloatNx x, _FloatNx y); // M < N
_DecimalM dMsubdN(_DecimalN x, _DecimalN y); // M < N
_DecimalM dMsubdNx(_DecimalNx x, _DecimalNx y); // M ≤ N
_DecimalMx dMxsubdN(_DecimalN x, _DecimalN y); // M < N
_DecimalMx dMxsubdNx(_DecimalNx x, _DecimalNx y); // M < N
_FloatM fMmulfN(_FloatN x, _FloatN y); // M < N
_FloatM fMmulfNx(_FloatNx x, _FloatNx y); // M ≤ N
_FloatMx fMxmulfN(_FloatN x, _FloatN y); // M < N
_FloatMx fMxmulfNx(_FloatNx x, _FloatNx y); // M < N
_DecimalM dMmuldN(_DecimalN x, _DecimalN y); // M < N
_DecimalM dMmuldNx(_DecimalNx x, _DecimalNx y); // M ≤ N
_DecimalMx dMxmuldN(_DecimalN x, _DecimalN y); // M < N
_DecimalMx dMxmuldNx(_DecimalNx x, _DecimalNx y); // M < N
_FloatM fMdivfN(_FloatN x, _FloatN y); // M < N
_FloatM fMdivfNx(_FloatNx x, _FloatNx y); // M ≤ N
_FloatMx fMxdivfN(_FloatN x, _FloatN y); // M < N
_FloatMx fMxdivfNx(_FloatNx x, _FloatNx y); // M < N
_DecimalM dMdivdN(_DecimalN x, _DecimalN y); // M < N
_DecimalM dMdivdNx(_DecimalNx x, _DecimalNx y); // M ≤ N
_DecimalMx dMxdivdN(_DecimalN x, _DecimalN y); // M < N
_DecimalMx dMxdivdNx(_DecimalNx x, _DecimalNx y); // M < N
_FloatM fMfmfN(_FloatN x, _FloatN y, _FloatN z); // M < N
_FloatM fMfmfNx(_FloatNx x, _FloatNx y, _FloatNx z); // M ≤ N

```

```

_FloatMx fMxfmafN(_FloatN x, _FloatN y, _FloatN z); // M < N
_FloatMx fMxfmafNx(_FloatNx x, _FloatNx y, _FloatNx z); // M < N
_DecimalM dMfmadN(_DecimalN x, _DecimalN y, _DecimalN z); // M < N
_DecimalM dMfmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z); // M ≤ N
_DecimalMx dMxfmadN(_DecimalN x, _DecimalN y, _DecimalN z); // M < N
_DecimalMx dMxfmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z); // M < N
_FloatM fMsqrtfN(_FloatN x); // M < N
_FloatM fMsqrtfNx(_FloatNx x); // M ≤ N
_FloatMx fMxsqrtfN(_FloatN x); // M < N
_FloatMx fMxsqrtfNx(_FloatNx x); // M < N
_DecimalM dMsqrtN(_DecimalN x); // M < N
_DecimalM dMsqrtNx(_DecimalNx x); // M ≤ N
_DecimalMx dMxsqrtN(_DecimalN x); // M < N
_DecimalMx dMxsqrtNx(_DecimalNx x); // M < N
_DecimalN quantizedN(_DecimalN x, _DecimalN y);
_DecimalNx quantizedNx(_DecimalNx x, _DecimalNx y);
bool samequantumdN(_DecimalN x, _DecimalN y);
bool samequantumdNx(_DecimalNx x, _DecimalNx y);
_DecimalN quantumdN(_DecimalN x);
_DecimalNx quantumdNx(_DecimalNx x);
long long int llquantexpdN(_DecimalN x);
long long int llquantexpdNx(_DecimalNx x);
void encodedecdN(unsigned char * restrict encptr,
 const _DecimalN * restrict xptr);
void decodedecdN(_DecimalN * restrict xptr,
 const unsigned char * restrict encptr);
void encodebindN(unsigned char * restrict encptr,
 const _DecimalN * restrict xptr);
void decodebindN(_DecimalN * restrict xptr,
 const unsigned char * restrict encptr);
int totalorderfN(const _FloatN *x, const _FloatN *y);
int totalorderfNx(const _FloatNx *x, const _FloatNx *y);
int totalorderdN(const _DecimalN *x, const _DecimalN *y);
int totalorderdNx(const _DecimalNx *x, const _DecimalNx *y);
int totalordermagfN(const _FloatN *x, const _FloatN *y);
int totalordermagfNx(const _FloatNx *x, const _FloatNx *y);
int totalordermagdN(const _DecimalN *x, const _DecimalN *y);
int totalordermagdNx(const _DecimalNx *x, const _DecimalNx *y);
_FloatN getpayloadfN(const _FloatN *x);
_FloatNx getpayloadfNx(const _FloatNx *x);
_DecimalN getpayloaddN(const _DecimalN *x);
_DecimalNx getpayloaddNx(const _DecimalNx *x);
int setpayloadfN(_FloatN *res, _FloatN pl);
int setpayloadfNx(_FloatNx *res, _FloatNx pl);
int setpayloaddN(_DecimalN *res, _DecimalN pl);
int setpayloaddNx(_DecimalNx *res, _DecimalNx pl);
int setpayloadsigfN(_FloatN *res, _FloatN pl);
int setpayloadsigfNx(_FloatNx *res, _FloatNx pl);
int setpayloadsigdN(_DecimalN *res, _DecimalN pl);
int setpayloadsigdNx(_DecimalNx *res, _DecimalNx pl);
void encodefN(unsigned char encptr[restrict static N/8],
 const _FloatN * restrict xptr);
void decodefN(_FloatN * restrict xptr,
 const unsigned char encptr[restrict static N/8]);
void fMencfN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);
void dMencdecN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);
void dMencbindN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);

```

**B.12 Non-local jumps <setjmp.h>****\_\_STDC\_VERSION\_SETJMP\_H\_\_****jmp\_buf**

```
int setjmp(jmp_buf env);
[[noreturn]] void longjmp(jmp_buf env, int val);
```

**B.13 Signal handling <signal.h>****sig\_atomic\_t****SIG\_IGN****SIGILL****SIGTERM****SIG\_DFL****SIGABRT****SIGINT****SIG\_ERR****SIGFPE****SIGSEGV**

```
void (*signal(int sig, void (*func)(int)))(int);
int raise(int sig);
```

**B.14 Alignment <stdalign.h>**

The header &lt;stdalign.h&gt; provides no content.

**B.15 Varying arguments <stdarg.h>****va\_list****\_\_STDC\_VERSION\_STDARG\_H\_\_**

```
type va_arg(va_list ap, type);
void va_copy(va_list dest, va_list src);
void va_end(va_list ap);
void va_start(va_list ap, ...);
```

**B.16 Atomics <stdatomic.h>****\_\_STDC\_NO\_ATOMICS\_\_****\_\_STDC\_VERSION\_STDATOMIC\_H\_\_****ATOMIC\_BOOL\_LOCK\_FREE****ATOMIC\_CHAR\_LOCK\_FREE****ATOMIC\_CHAR8\_T\_LOCK\_FREE****ATOMIC\_CHAR16\_T\_LOCK\_FREE****ATOMIC\_CHAR32\_T\_LOCK\_FREE****ATOMIC\_WCHAR\_T\_LOCK\_FREE****ATOMIC\_SHORT\_LOCK\_FREE****ATOMIC\_INT\_LOCK\_FREE****ATOMIC\_LONG\_LOCK\_FREE****ATOMIC\_LLONG\_LOCK\_FREE****ATOMIC\_POINTER\_LOCK\_FREE****ATOMIC\_FLAG\_INIT****memory\_order****atomic\_flag****memory\_order\_relaxed****memory\_order\_consume****memory\_order\_acquire****memory\_order\_release****memory\_order\_acq\_rel****memory\_order\_seq\_cst****atomic\_bool****atomic\_char****atomic\_schar****atomic\_uchar****atomic\_short****atomic\_ushort****atomic\_int****atomic\_uint****atomic\_long****atomic\_ulong****atomic\_llong****atomic\_ullong****atomic\_char8\_t****atomic\_char16\_t****atomic\_char32\_t****atomic\_wchar\_t****atomic\_int\_least8\_t****atomic\_uint\_least8\_t****atomic\_int\_least16\_t****atomic\_uint\_least16\_t****atomic\_int\_least32\_t****atomic\_uint\_least32\_t****atomic\_int\_least64\_t****atomic\_uint\_least64\_t**

|                                   |                                   |
|-----------------------------------|-----------------------------------|
| <code>atomic_int_fast8_t</code>   | <code>atomic_uint_fast64_t</code> |
| <code>atomic_uint_fast8_t</code>  | <code>atomic_intptr_t</code>      |
| <code>atomic_int_fast16_t</code>  | <code>atomic_uintptr_t</code>     |
| <code>atomic_uint_fast16_t</code> | <code>atomic_size_t</code>        |
| <code>atomic_int_fast32_t</code>  | <code>atomic_ptrdiff_t</code>     |
| <code>atomic_uint_fast32_t</code> | <code>atomic_intmax_t</code>      |
| <code>atomic_int_fast64_t</code>  | <code>atomic_uintmax_t</code>     |

```

void atomic_init(volatile A *obj, C value);
type kill_dependency(type y);
void atomic_thread_fence(memory_order order);
void atomic_signal_fence(memory_order order);
bool atomic_is_lock_free(const volatile A *obj);
void atomic_store(volatile A *object, C desired);
void atomic_store_explicit(volatile A *object, C desired, memory_order order);
C atomic_load(const volatile A *object);
C atomic_load_explicit(const volatile A *object, memory_order order);
C atomic_exchange(volatile A *object, C desired);
C atomic_exchange_explicit(volatile A *object, C desired, memory_order order);
bool atomic_compare_exchange_strong(volatile A *object, C *expected, C desired);
bool atomic_compare_exchange_strong_explicit(volatile A *object, C *expected,
 C desired, memory_order success, memory_order failure);
bool atomic_compare_exchange_weak(volatile A *object, C *expected, C desired);
bool atomic_compare_exchange_weak_explicit(volatile A *object, C *expected,
 C desired, memory_order success, memory_order failure);
C atomic_fetch_key(volatile A *object, M operand);
C atomic_fetch_key_explicit(volatile A *object, M operand, memory_order order);
bool atomic_flag_test_and_set(volatile atomic_flag *object);
bool atomic_flag_test_and_set_explicit(volatile atomic_flag *object,
 memory_order order);
void atomic_flag_clear(volatile atomic_flag *object);
void atomic_flag_clear_explicit(volatile atomic_flag *object,
 memory_order order);

```

## B.17 Bit and byte utilities <stdbit.h>

|                                     |                                        |
|-------------------------------------|----------------------------------------|
| <code>__STDC_ENDIAN_BIG__</code>    | <code>__STDC_ENDIAN_NATIVE__</code>    |
| <code>__STDC_ENDIAN_LITTLE__</code> | <code>__STDC_VERSION_STDBIT_H__</code> |

```

unsigned int stdc_leading_zeros_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_leading_zeros_us(unsigned short value) [[unsequenced]];
unsigned int stdc_leading_zeros_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_leading_zeros_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_leading_zeros_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_leading_zeros(generic_value_type value) [[unsequenced]];
unsigned int stdc_leading_ones_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_leading_ones_us(unsigned short value) [[unsequenced]];
unsigned int stdc_leading_ones_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_leading_ones_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_leading_ones_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_leading_ones(generic_value_type value) [[unsequenced]];
unsigned int stdc_trailing_zeros_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_trailing_zeros_us(unsigned short value) [[unsequenced]];
unsigned int stdc_trailing_zeros_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_trailing_zeros_ul(unsigned long int value) [[unsequenced]];
unsigned int

```



```

stdc_trailing_zeros_ull(unsigned long long int value) [[unsequenced]];
generic_return_type
stdc_trailing_zeros(generic_value_type value) [[unsequenced]];
unsigned int stdc_trailing_ones_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_trailing_ones_us(unsigned short value) [[unsequenced]];
unsigned int stdc_trailing_ones_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_trailing_ones_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_trailing_ones_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_trailing_ones(generic_value_type value) [[unsequenced]];
unsigned int stdc_first_leading_zero_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_first_leading_zero_us(unsigned short value) [[unsequenced]];
unsigned int stdc_first_leading_zero_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_first_leading_zero_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_first_leading_zero_ull(unsigned long long int value) [[unsequenced]];
generic_return_type
stdc_first_leading_zero(generic_value_type value) [[unsequenced]];
unsigned int stdc_first_leading_one_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_first_leading_one_us(unsigned short value) [[unsequenced]];
unsigned int stdc_first_leading_one_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_first_leading_one_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_first_leading_one_ull(unsigned long long int value) [[unsequenced]];
generic_return_type
stdc_first_leading_one(generic_value_type value) [[unsequenced]];
unsigned int stdc_first_trailing_zero_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_first_trailing_zero_us(unsigned short value) [[unsequenced]];
unsigned int stdc_first_trailing_zero_ui(unsigned int value) [[unsequenced]];
unsigned int
stdc_first_trailing_zero_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_first_trailing_zero_ull(unsigned long long int value) [[unsequenced]];
generic_return_type
stdc_first_trailing_zero(generic_value_type value) [[unsequenced]];
unsigned int stdc_first_trailing_one_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_first_trailing_one_us(unsigned short value) [[unsequenced]];
unsigned int stdc_first_trailing_one_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_first_trailing_one_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_first_trailing_one_ull(unsigned long long int value) [[unsequenced]];
generic_return_type
stdc_first_trailing_one(generic_value_type value) [[unsequenced]];
unsigned int stdc_count_zeros_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_count_zeros_us(unsigned short value) [[unsequenced]];
unsigned int stdc_count_zeros_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_count_zeros_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_count_zeros_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_count_zeros(generic_value_type value) [[unsequenced]];
unsigned int stdc_count_ones_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_count_ones_us(unsigned short value) [[unsequenced]];
unsigned int stdc_count_ones_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_count_ones_ul(unsigned long int value) [[unsequenced]];
unsigned int
stdc_count_ones_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_count_ones(generic_value_type value) [[unsequenced]];
bool stdc_has_single_bit_uc(unsigned char value) [[unsequenced]];
bool stdc_has_single_bit_us(unsigned short value) [[unsequenced]];
bool stdc_has_single_bit_ui(unsigned int value) [[unsequenced]];
bool stdc_has_single_bit_ul(unsigned long int value) [[unsequenced]];

```

```

bool stdc_has_single_bit_ull(unsigned long long int value) [[unsequenced]];
bool stdc_has_single_bit(generic_value_type value) [[unsequenced]];
unsigned int stdc_bit_width_uc(unsigned char value) [[unsequenced]];
unsigned int stdc_bit_width_us(unsigned short value) [[unsequenced]];
unsigned int stdc_bit_width_ui(unsigned int value) [[unsequenced]];
unsigned int stdc_bit_width_ul(unsigned long int value) [[unsequenced]];
unsigned int
 stdc_bit_width_ull(unsigned long long int value) [[unsequenced]];
generic_return_type stdc_bit_width(generic_value_type value) [[unsequenced]];
unsigned char stdc_bit_floor_uc(unsigned char value) [[unsequenced]];
unsigned short stdc_bit_floor_us(unsigned short value) [[unsequenced]];
unsigned int stdc_bit_floor_ui(unsigned int value) [[unsequenced]];
unsigned long int stdc_bit_floor_ul(unsigned long int value) [[unsequenced]];
unsigned long long int
 stdc_bit_floor_ull(unsigned long long int value) [[unsequenced]];
generic_value_type stdc_bit_floor(generic_value_type value) [[unsequenced]];
unsigned char stdc_bit_ceil_uc(unsigned char value) [[unsequenced]];
unsigned short stdc_bit_ceil_us(unsigned short value) [[unsequenced]];
unsigned int stdc_bit_ceil_ui(unsigned int value) [[unsequenced]];
unsigned long int stdc_bit_ceil_ul(unsigned long int value) [[unsequenced]];
unsigned long long int
 stdc_bit_ceil_ull(unsigned long long int value) [[unsequenced]];
generic_value_type stdc_bit_ceil(generic_value_type value) [[unsequenced]];
unsigned char stdc_rotate_left_uc(unsigned char value, unsigned int count);
unsigned short stdc_rotate_left_us(unsigned short value, unsigned int count);
unsigned int stdc_rotate_left_ui(unsigned int value, unsigned int count);
unsigned long stdc_rotate_left_ul(unsigned long value, unsigned int count);
unsigned long long stdc_rotate_left_ull(unsigned long long value, unsigned int count);
generic_value_type stdc_rotate_left(
 generic_value_type value, generic_count_type count);
unsigned char stdc_rotate_right_uc(unsigned char value, unsigned int count);
unsigned short stdc_rotate_right_us(unsigned short value, unsigned int count);
unsigned int stdc_rotate_right_ui(unsigned int value, unsigned int count);
unsigned long stdc_rotate_right_ul(unsigned long value, unsigned int count);
unsigned long long
 stdc_rotate_right_ull(unsigned long long value, unsigned int count);
generic_value_type
 stdc_rotate_right(generic_value_type value, generic_count_type count);

```

Only if the implementation has **CHAR\_BIT** equal to 8:

```

void stdc_memreverse8(size_t n, unsigned char ptr[static n]);
uintN_t stdc_memreverse8uN(uintN_t value);
uint_leastN_t stdc_load8_leuN(const unsigned char ptr[static (N / 8)]);
uint_leastN_t stdc_load8_beuN(const unsigned char ptr[static (N / 8)]);
uint_leastN_t stdc_load8_aligned_leuN(const unsigned char ptr[static (N / 8)]);
uint_leastN_t stdc_load8_aligned_beuN(const unsigned char ptr[static (N / 8)]);
int_leastN_t stdc_load8_lesN(const unsigned char ptr[static (N / 8)]);
int_leastN_t stdc_load8_besN(const unsigned char ptr[static (N / 8)]);
int_leastN_t stdc_load8_aligned_lesN(const unsigned char ptr[static (N / 8)]);
int_leastN_t stdc_load8_aligned_besN(const unsigned char ptr[static (N / 8)]);
void stdc_store8_leuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_beuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_aligned_leuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_aligned_beuN(uint_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_lesN(int_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_besN(int_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_aligned_lesN(int_leastN_t value, unsigned char ptr[static (N / 8)]);
void stdc_store8_aligned_besN(int_leastN_t value, unsigned char ptr[static (N / 8)]);

```

**B.18 Boolean type and values <stdbool.h>**`__bool_true_false_are_defined`**B.19 Checked Integer Operations <stdckdint.h>**`__STDC_VERSION__` `STDCKDINT_H__`

```
bool ckd_add(type1 *result, type2 a, type3 b);
bool ckd_sub(type1 *result, type2 a, type3 b);
bool ckd_mul(type1 *result, type2 a, type3 b);
```

**B.20 Common definitions <stddef.h>**

|                        |                                                       |                   |
|------------------------|-------------------------------------------------------|-------------------|
| <code>ptrdiff_t</code> | <code>max_align_t</code>                              | <code>NULL</code> |
| <code>nullptr_t</code> | <code>wchar_t</code>                                  |                   |
| <code>size_t</code>    | <code>__STDC_VERSION__</code> <code>STDDEF_H__</code> |                   |

`offsetof(type, member-designator)``unreachable()`

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of <stddef.h>:

`rsizet`**B.21 Integer types <stdint.h>**

|                                                       |                                |                                 |
|-------------------------------------------------------|--------------------------------|---------------------------------|
| <code>intN_t</code>                                   | <code>INT_LEASTN_WIDTH</code>  | <code>PTRDIFF_MIN</code>        |
| <code>uintN_t</code>                                  | <code>UINT_LEASTN_MAX</code>   | <code>PTRDIFF_MAX</code>        |
| <code>int_leastN_t</code>                             | <code>UINT_LEASTN_WIDTH</code> | <code>SIG_ATOMIC_MIN</code>     |
| <code>uint_leastN_t</code>                            | <code>INT_FASTN_MIN</code>     | <code>SIG_ATOMIC_MAX</code>     |
| <code>int_fastN_t</code>                              | <code>INT_FASTN_MAX</code>     | <code>SIG_ATOMIC_WIDTH</code>   |
| <code>uint_fastN_t</code>                             | <code>INT_FASTN_WIDTH</code>   | <code>SIZE_MAX</code>           |
| <code>intptr_t</code>                                 | <code>UINT_FASTN_MAX</code>    | <code>SIZE_WIDTH</code>         |
| <code>uintptr_t</code>                                | <code>UINT_FASTN_WIDTH</code>  | <code>WCHAR_MIN</code>          |
| <code>intmax_t</code>                                 | <code>INTPTR_MIN</code>        | <code>WCHAR_MAX</code>          |
| <code>uintmax_t</code>                                | <code>INTPTR_MAX</code>        | <code>WCHAR_WIDTH</code>        |
| <code>__STDC_VERSION__</code> <code>STDINT_H__</code> | <code>INTPTR_WIDTH</code>      | <code>WINT_MIN</code>           |
| <code>INTN_MIN</code>                                 | <code>UINTPTR_MAX</code>       | <code>WINT_MAX</code>           |
| <code>INTN_MAX</code>                                 | <code>UINTPTR_WIDTH</code>     | <code>WINT_WIDTH</code>         |
| <code>INTN_WIDTH</code>                               | <code>INTMAX_MIN</code>        | <code>INTN_C( value )</code>    |
| <code>UINTN_MAX</code>                                | <code>INTMAX_MAX</code>        | <code>UINTN_C( value )</code>   |
| <code>UINTN_WIDTH</code>                              | <code>INTMAX_WIDTH</code>      | <code>INTMAX_C( value )</code>  |
| <code>INT_LEASTN_MIN</code>                           | <code>UINTMAX_MAX</code>       | <code>UINTMAX_C( value )</code> |
| <code>INT_LEASTN_MAX</code>                           | <code>UINTMAX_WIDTH</code>     |                                 |

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of <stdint.h>:



**RSIZE\_MAX****B.22 Input/output <stdio.h>**

|               |                     |                 |                                 |
|---------------|---------------------|-----------------|---------------------------------|
| <b>size_t</b> | <b>_IONBF</b>       | <b>SEEK_CUR</b> | <b>stdout</b>                   |
| <b>FILE</b>   | <b>BUFSIZ</b>       | <b>SEEK_END</b> | <b>_PRINTF_NAN_LEN_MAX</b>      |
| <b>fpos_t</b> | <b>EOF</b>          | <b>SEEK_SET</b> |                                 |
| <b>NULL</b>   | <b>FOPEN_MAX</b>    | <b>TMP_MAX</b>  | <b>__STDC_VERSION_STDIO_H__</b> |
| <b>_IOFBF</b> | <b>FILENAME_MAX</b> | <b>stderr</b>   |                                 |
| <b>_IOLBF</b> | <b>L_tmpnam</b>     | <b>stdin</b>    |                                 |

```

int remove(const char *filename);
int rename(const char *old, const char *new);
FILE *tmpfile(void);
char *tmpnam(char *s);
int fclose(FILE *stream);
int fflush(FILE *stream);
FILE *fopen(const char * restrict filename, const char * restrict mode);
FILE *freopen(const char * restrict filename, const char * restrict mode,
FILE * restrict stream);
void setbuf(FILE * restrict stream, char * restrict buf);
int setvbuf(FILE * restrict stream, char * restrict buf, int mode, size_t size);
int printf(const char * restrict format, ...);
int scanf(const char * restrict format, ...);
int snprintf(char * restrict s, size_t n, const char * restrict format, ...);
int sprintf(char * restrict s, const char * restrict format, ...);
int sscanf(const char * restrict s, const char * restrict format, ...);
int vfprintf(FILE * restrict stream, const char * restrict format, va_list arg);
int vfprintf(FILE * restrict stream, const char * restrict format, va_list arg);
int vprintf(const char * restrict format, va_list arg);
int vscanf(const char * restrict format, va_list arg);
int vsnprintf(char * restrict s, size_t n, const char * restrict format, va_list arg);
int vsprintf(char * restrict s, const char * restrict format, va_list arg);
int vsscanf(const char * restrict s, const char * restrict format, va_list arg);
int fgetc(FILE *stream);
char *fgets(char * restrict s, int n, FILE * restrict stream);
int fputc(int c, FILE *stream);
int fputs(const char * restrict s, FILE * restrict stream);
int getc(FILE *stream);
int getchar(void);
int putc(int c, FILE *stream);
int putchar(int c);
int puts(const char *s);
int ungetc(int c, FILE *stream);
size_t fread(void * restrict ptr, size_t size, size_t nmemb,
FILE * restrict stream);
size_t fwrite(const void * restrict ptr, size_t size, size_t nmemb,
FILE * restrict stream);
int fgetpos(FILE * restrict stream, fpos_t * restrict pos);
int fseek(FILE *stream, long int offset, int whence);
int fsetpos(FILE *stream, const fpos_t *pos);
long int ftell(FILE *stream);
void rewind(FILE *stream);
void clearerr(FILE *stream);
int feof(FILE *stream);
int ferror(FILE *stream);
void perror(const char *s);
int fprintf(FILE * restrict stream, const char * restrict format, ...);
int fscanf(FILE * restrict stream, const char * restrict format, ...);

```

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of `<stdio.h>`:

`L_tmpnam_s`                      `TMP_MAX_S`                      `errno_t`                      `rsizet`

```

errno_t tmpfile_s(FILE * restrict * restrict streamptr);
errno_t tmpnam_s(char *s, rsizet maxsize);
errno_t fopen_s(FILE * restrict * restrict streamptr,
 const char * restrict filename, const char * restrict mode);
errno_t freopen_s(FILE * restrict * restrict newstreamptr,
 const char * restrict filename, const char * restrict mode,
 FILE * restrict stream);
int fprintf_s(FILE * restrict stream, const char * restrict format, ...);
int fscanf_s(FILE * restrict stream, const char * restrict format, ...);
int printf_s(const char * restrict format, ...);
int scanf_s(const char * restrict format, ...);
int snprintf_s(char * restrict s, rsizet n, const char * restrict format, ...);
int sprintf_s(char * restrict s, rsizet n, const char * restrict format, ...);
int sscanf_s(const char * restrict s, const char * restrict format, ...);
int vfprintf_s(FILE * restrict stream, const char * restrict format, va_list arg);
int vfscanf_s(FILE * restrict stream, const char * restrict format, va_list arg);
int vprintf_s(const char * restrict format, va_list arg);
int vscanf_s(const char * restrict format, va_list arg);
int vsnprintf_s(char * restrict s, rsizet n, const char * restrict format,
 va_list arg);
int vsprintf_s(char * restrict s, rsizet n, const char * restrict format,
 va_list arg);
int vsscanf_s(const char * restrict s, const char * restrict format, va_list arg);
char *gets_s(char *s, rsizet n);

```

## B.23 General utilities `<stdlib.h>`

`size_t`                                      `once_flag`                                      `NULL`  
`wchar_t`                                      `__STDC_VERSION_STDLIB_H__`                                      `ONCE_FLAG_INIT`  
`div_t`                                      `EXIT_FAILURE`                                      `RAND_MAX`  
`ldiv_t`                                      `EXIT_SUCCESS`  
`lldiv_t`                                      `MB_CUR_MAX`

```

void call_once(once_flag *flag, void (*func)(void));
double atof(const char *nptr);
int atoi(const char *nptr);
long int atol(const char *nptr);
long long int atoll(const char *nptr);
int strfromd(char * restrict s, size_t n, const char * restrict format,
 double fp);
int strfromf(char * restrict s, size_t n, const char * restrict format,
 float fp);
int strfroml(char * restrict s, size_t n, const char * restrict format,
 long double fp);
double strtod(const char * restrict nptr, char ** restrict endptr);
float strtod(const char * restrict nptr, char ** restrict endptr);
long double strtold(const char * restrict nptr, char ** restrict endptr);
long int strtol(const char * restrict nptr, char ** restrict endptr, int base);
long long int strtoll(const char * restrict nptr, char ** restrict endptr,
 int base);
unsigned long int strtoul(const char * restrict nptr, char ** restrict endptr,
 int base);
unsigned long long int strtoull(const char * restrict nptr,

```

```

 char ** restrict endptr, int base);
int rand(void);
void srand(unsigned int seed);
void *aligned_alloc(size_t alignment, size_t size);
void *calloc(size_t nmemb, size_t size);
void free(void *ptr);
void free_sized(void *ptr, size_t size);
void free_aligned_sized(void *ptr, size_t alignment, size_t size);
void *malloc(size_t size);
void *realloc(void *ptr, size_t size);
[[noreturn]] void abort(void);
int atexit(void (*func)(void));
int at_quick_exit(void (*func)(void));
[[noreturn]] void exit(int status);
[[noreturn]] void _Exit(int status);
char *getenv(const char *name);
[[noreturn]] void quick_exit(int status);
int system(const char *string);
QVoid *bsearch(const void *key, QVoid *base, size_t nmemb, size_t size,
 int (*compar)(const void *, const void *));
void qsort(void *base, size_t nmemb, size_t size,
 int (*compar)(const void *, const void *));
int abs(int j);
unsigned int uabs(int j);
long int labs(long int j);
unsigned long int ulabs(long int j);
long long int llabs(long long int j);
unsigned long long int ullabs(long long int j);
div_t div(int numer, int denom);
ldiv_t ldiv(long int numer, long int denom);
lldiv_t lldiv(long long int numer, long long int denom);
int mblen(const char *s, size_t n);
int mbtowc(wchar_t * restrict pwc, const char * restrict s, size_t n);
int wctomb(char *s, wchar_t wc);
size_t mbstowcs(wchar_t * restrict pwcs, const char * restrict s, size_t n);
size_t wcstombs(char * restrict s, const wchar_t * restrict pwcs, size_t n);
size_t memalignment(const void *p);

```

Only if the implementation defines `__STDC_IEC_60559_DFP__`:

```

int strfromd32(char * restrict s, size_t n, const char * restrict format,
 _Decimal32 fp);
int strfromd64(char * restrict s, size_t n, const char * restrict format,
 _Decimal64 fp);
int strfromd128(char * restrict s, size_t n, const char * restrict format,
 _Decimal128 fp);
_Decimal32 strtod32(const char * restrict nptr, char ** restrict endptr);
_Decimal64 strtod64(const char * restrict nptr, char ** restrict endptr);
_Decimal128 strtod128(const char * restrict nptr, char ** restrict endptr);

```

Only if the implementation defines `__STDC_IEC_60559_TYPES__` and additionally the user code defines `__STDC_WANT_IEC_60559_TYPES_EXT__` before any inclusion of `<stdlib.h>`:

```

int strfromfN(char * restrict s, size_t n,
 const char * restrict format, _FloatN fp);
int strfromfNx(char * restrict s, size_t n,
 const char * restrict format, _FloatNx fp);
int strfromdN(char * restrict s, size_t n,
 const char * restrict format, _DecimalN fp);
int strfromdNx(char * restrict s, size_t n,
 const char * restrict format, _DecimalNx fp);

```

```

_FloatN strtOfN(const char * restrict nptr,
 char ** restrict endptr);
_FloatNx strtOfNx(const char * restrict nptr,
 char ** restrict endptr);
_DecimalN strtodN(const char * restrict nptr,
 char ** restrict endptr);
_DecimalNx strtodNx(const char * restrict nptr,
 char ** restrict endptr);
int strfromencfN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);
int strfromencdecN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);
int strfromencbindN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);
void strtOencfN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);
void strtOencdecN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);
void strtOencbindN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);

```

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of `<stdlib.h>`:

`errno_t` `rsizet` `constraint_handler_t`

```

constraint_handler_t set_constraint_handler_s(constraint_handler_t handler);
void abort_handler_s(const char * restrict msg, void * restrict ptr,
 errno_t error);
void ignore_handler_s(const char * restrict msg, void * restrict ptr,
 errno_t error);
errno_t getenv_s(size_t * restrict len, char * restrict value, rsizet maxsize,
 const char * restrict name);
QVoid *bsearch_s(const void *key, QVoid *base, rsizet nmemb, rsizet size,
 int (*compar)(const void *k, const void *y, void *context),
 void *context);
errno_t qsort_s(void *base, rsizet nmemb, rsizet size,
 int (*compar)(const void *x, const void *y, void *context),
 void *context);
errno_t wctomb_s(int * restrict status, char * restrict s, rsizet smax,
 wchar_t wc);
errno_t mbstowcs_s(size_t * restrict retval, wchar_t * restrict dst,
 rsizet dstmax, const char * restrict src, rsizet len);
errno_t wcstombs_s(size_t * restrict retval, char * restrict dst, rsizet dstmax,
 const wchar_t * restrict src, rsizet len);

```

## B.24 Text transcoding utilities `<stdmchar.h>`

|                                          |                           |
|------------------------------------------|---------------------------|
| <code>__STDC_VERSION_STDMCHAR_H__</code> | <code>WCHAR_WIDTH</code>  |
| <code>mbstate_t</code>                   | <code>WCHAR_UTF8</code>   |
| <code>wchar_t</code>                     | <code>WCHAR_UTF16</code>  |
| <code>char8_t</code>                     | <code>WCHAR_UTF32</code>  |
| <code>char16_t</code>                    | <code>MB_UTF8</code>      |
| <code>char32_t</code>                    | <code>MB_UTF16</code>     |
| <code>size_t</code>                      | <code>MB_UTF32</code>     |
| <code>stdc_mcerr</code>                  | <code>STDC_C8_MAX</code>  |
| <code>WCHAR_MAX</code>                   | <code>STDC_C16_MAX</code> |
| <code>WCHAR_MIN</code>                   | <code>STDC_C32_MAX</code> |

STDC\_MC\_MAX  
STDC\_MWC\_MAX  
stdc\_mcerr\_ok

stdc\_mcerr\_invalid  
stdc\_mcerr\_incomplete\_input  
stdc\_mcerr\_insufficient\_output

```
stdc_mcerr stdc_mcnrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcnrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_mwcnrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcnrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c8nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
```

```

stdc_mcerr stdc_c8nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c16nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c32nrtomcn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc8n(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc16n(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32nrtoc32n(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_mcsnrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,

```

```

 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mcsnrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_mwcsnrtoombsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcsnrtoombsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcsnrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcsnrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_mwcsnrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, wchar_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c8snrtombsn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtombsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c8snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char8_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c16snrtombsn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtombsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);

```



```

stdc_mcerr stdc_c16snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c16snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char16_t*restrict *restrict input,
 mbstate_t*restrict state);

stdc_mcerr stdc_c32snrtomcsn(size_t*restrict output_size,
 char*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtomwcn(size_t*restrict output_size,
 wchar_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc8sn(size_t*restrict output_size,
 char8_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc16sn(size_t*restrict output_size,
 char16_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);
stdc_mcerr stdc_c32snrtoc32sn(size_t*restrict output_size,
 char32_t*restrict *restrict output,
 size_t*restrict input_size, const char32_t*restrict *restrict input,
 mbstate_t*restrict state);

```

## B.25 `_Noreturn <stdnoreturn.h>`

`noreturn`

## B.26 String handling <string.h>

`size_t` `NULL`  
`__STDC_VERSION_STRING_H__`

```

void *memcpy(void * restrict s1, const void * restrict s2, size_t n);
void *memccpy(void * restrict s1, const void * restrict s2, int c, size_t n);
void *memmove(void *s1, const void *s2, size_t n);
char *strcpy(char * restrict s1, const char * restrict s2);
char *strncpy(char * restrict s1, const char * restrict s2, size_t n);
char *strdup(const char *s);
char *strndup(const char *s, size_t n);
char *strcat(char * restrict s1, const char * restrict s2);
char *strncat(char * restrict s1, const char * restrict s2, size_t n);
int memcmp(const void *s1, const void *s2, size_t n);
int strcmp(const char *s1, const char *s2);
int strcoll(const char *s1, const char *s2);
int strncmp(const char *s1, const char *s2, size_t n);
size_t strxfrm(char * restrict s1, const char * restrict s2, size_t n);
QVoid *memchr(QVoid *s, int c, size_t n);
QChar *strchr(QChar *s, int c);
size_t strcspn(const char *s1, const char *s2);

```



```

QChar *strpbrk(QChar *s1, const char *s2);
QChar *strrchr(QChar *s, int c);
size_t strspn(const char *s1, const char *s2);
QChar *strstr(QChar *s1, const char *s2);
char *strtok(char * restrict s1, const char * restrict s2);
void *memset(void *s, int c, size_t n);
void *memset_explicit(void *s, int c, size_t n);
char *strerror(int errnum);
size_t strlen(const char *s);
size_t strnlen(const char *s, size_t n);

```

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of `<string.h>`:

`errno_t`

`rsize_t`

```

errno_t memcpy_s(void * restrict s1, rsize_t slmax, const void * restrict s2,
 rsize_t n);
errno_t memmove_s(void *s1, rsize_t slmax, const void *s2, rsize_t n);
errno_t strcpy_s(char * restrict s1, rsize_t slmax, const char * restrict s2);
errno_t strncpy_s(char * restrict s1, rsize_t slmax, const char * restrict s2,
 rsize_t n);
errno_t strcat_s(char * restrict s1, rsize_t slmax, const char * restrict s2);
errno_t strncat_s(char * restrict s1, rsize_t slmax, const char * restrict s2,
 rsize_t n);
char *strtok_s(char * restrict s1, rsize_t * restrict slmax,
 const char * restrict s2, char ** restrict ptr);
errno_t memset_s(void *s, rsize_t smax, int c, rsize_t n)
errno_t strerror_s(char *s, rsize_t maxsize, errno_t errnum);
size_t strerrorlen_s(errno_t errnum);
size_t strnlen_s(const char *s, size_t maxsize);

```

## B.27 Type-generic math `<tgmath.h>`

|                                        |                               |                               |
|----------------------------------------|-------------------------------|-------------------------------|
| <code>__STDC_VERSION_TGMATH_H__</code> | <code>cbirt</code>            | <code>fminimum_num</code>     |
| <code>acos</code>                      | <code>ceil</code>             | <code>fminimum_mag_num</code> |
| <code>asin</code>                      | <code>compoundn</code>        | <code>fmod</code>             |
| <code>atan</code>                      | <code>copysign</code>         | <code>frexp</code>            |
| <code>acosh</code>                     | <code>cospi</code>            | <code>fromfpx</code>          |
| <code>asinh</code>                     | <code>erfc</code>             | <code>fromfp</code>           |
| <code>atanh</code>                     | <code>erf</code>              | <code>hypot</code>            |
| <code>cos</code>                       | <code>expl0m1</code>          | <code>ilogb</code>            |
| <code>sin</code>                       | <code>expl0</code>            | <code>ldexp</code>            |
| <code>tan</code>                       | <code>exp2m1</code>           | <code>lgamma</code>           |
| <code>cosh</code>                      | <code>exp2</code>             | <code>llogb</code>            |
| <code>sinh</code>                      | <code>expm1</code>            | <code>llrint</code>           |
| <code>tanh</code>                      | <code>fdim</code>             | <code>llround</code>          |
| <code>exp</code>                       | <code>floor</code>            | <code>log10p1</code>          |
| <code>log</code>                       | <code>fmax</code>             | <code>log10</code>            |
| <code>pow</code>                       | <code>fmaximum</code>         | <code>log1p</code>            |
| <code>sqrt</code>                      | <code>fmaximum_mag</code>     | <code>log2p1</code>           |
| <code>fabs</code>                      | <code>fmaximum_num</code>     | <code>log2</code>             |
| <code>acospi</code>                    | <code>fmaximum_mag_num</code> | <code>logb</code>             |
| <code>asimpi</code>                    | <code>fma</code>              | <code>logp1</code>            |
| <code>atan2pi</code>                   | <code>fmin</code>             | <code>lrint</code>            |
| <code>atan2</code>                     | <code>fminimum</code>         | <code>lround</code>           |
| <code>atanpi</code>                    | <code>fminimum_mag</code>     | <code>nearbyint</code>        |

|            |          |       |
|------------|----------|-------|
| nextafter  | round    | dadd  |
| nextdown   | rsqrt    | fsub  |
| nexttoward | scalbln  | dsub  |
| nextup     | scalbn   | fmul  |
| pown       | sinpi    | dmul  |
| powr       | tanpi    | fdiv  |
| remainder  | tgamma   | ddiv  |
| remquo     | trunc    | ffma  |
| rint       | ufromfpx | dfma  |
| rootn      | ufromfp  | fsqrt |
| roundeven  | fadd     | dsqrt |

Only if the implementation does not define `___STDC_NO_COMPLEX___`:

|      |       |      |       |       |
|------|-------|------|-------|-------|
| carg | cimag | conj | cproj | creal |
|------|-------|------|-------|-------|

Only if the implementation defines `___STDC_IEC_60559_DFP___`:

|        |        |        |         |             |            |
|--------|--------|--------|---------|-------------|------------|
| d32add | d64sub | d32div | d64fma  | quantize    | llquantexp |
| d64add | d32mul | d64div | d32sqrt | samequantum |            |
| d32sub | d64mul | d32fma | d64sqrt | quantum     |            |

Only if the implementation defines `___STDC_IEC_60559_TYPES___` and additionally the user code defines `___STDC_WANT_IEC_60559_TYPES_EXT___` before any inclusion of `<tgmath.h>`:

|        |        |         |        |        |         |
|--------|--------|---------|--------|--------|---------|
| fMadd  | fMsub  | fMmul   | fMdiv  | fMfma  | fMsqrt  |
| fMxadd | fMxsub | fMxmMul | fMxdiv | fMxfma | fMxsqrt |
| dMadd  | dMsub  | dMmul   | dMdiv  | dMfma  | dMsqrt  |
| dMxadd | dMxsub | dMxmMul | dMxdiv | dMxfma | dMxsqrt |

## B.28 Threads `<threads.h>`

|                                    |                            |                            |
|------------------------------------|----------------------------|----------------------------|
| <code>___STDC_NO_THREADS___</code> | <code>mtx_t</code>         | <code>mtx_timed</code>     |
| <code>ONCE_FLAG_INIT</code>        | <code>tss_dtor_t</code>    | <code>thrd_timedout</code> |
| <code>TSS_DTOR_ITERATIONS</code>   | <code>thrd_start_t</code>  | <code>thrd_success</code>  |
| <code>cnd_t</code>                 | <code>once_flag</code>     | <code>thrd_busy</code>     |
| <code>thrd_t</code>                | <code>mtx_plain</code>     | <code>thrd_error</code>    |
| <code>tss_t</code>                 | <code>mtx_recursive</code> | <code>thrd_nomem</code>    |

```
void call_once(once_flag *flag, void (*func)(void));
int cnd_broadcast(cnd_t *cond);
void cnd_destroy(cnd_t *cond);
int cnd_init(cnd_t *cond);
int cnd_signal(cnd_t *cond);
int cnd_timedwait(cnd_t * restrict cond, mtx_t * restrict mtx,
 const struct timespec * restrict ts);
int cnd_wait(cnd_t *cond, mtx_t *mtx);
void mtx_destroy(mtx_t *mtx);
int mtx_init(mtx_t *mtx, int type);
int mtx_lock(mtx_t *mtx);
int mtx_timedlock(mtx_t * restrict mtx, const struct timespec * restrict ts);
int mtx_trylock(mtx_t *mtx);
int mtx_unlock(mtx_t *mtx);
int thrd_create(thrd_t *thr, thrd_start_t func, void *arg);
thrd_t thrd_current(void);
int thrd_detach(thrd_t thr);
```

```

int thrd_equal(thrd_t thr0, thrd_t thr1);
[[noreturn]] void thrd_exit(int res);
int thrd_join(thrd_t thr, int *res);
int thrd_sleep(const struct timespec *duration, struct timespec *remaining);
void thrd_yield(void);
int tss_create(tss_t *key, tss_dtor_t dtor);
void tss_delete(tss_t key);
void *tss_get(tss_t key);
int tss_set(tss_t key, void *val);

```

## B.29 Date and time <time.h>

|                                      |                       |                              |
|--------------------------------------|-----------------------|------------------------------|
| <code>__STDC_VERSION_TIME_H__</code> | <code>TIME_UTC</code> | <code>time_t</code>          |
| <code>NULL</code>                    | <code>size_t</code>   | <code>struct timespec</code> |
| <code>CLOCKS_PER_SEC</code>          | <code>clock_t</code>  | <code>struct tm</code>       |

```

clock_t clock(void);
double difftime(time_t time1, time_t time0);
time_t mktime(struct tm *timeptr);
time_t timegm(struct tm *timeptr);
time_t time(time_t *timer);
int timespec_get(struct timespec *ts, int base);
int timespec_getres(struct timespec *ts, int base);
[[deprecated]] char *asctime(const struct tm *timeptr);
[[deprecated]] char *ctime(const time_t *timer);
struct tm *gmtime(const time_t *timer);
struct tm *gmtime_r(const time_t *timer, struct tm *buf);
struct tm *localtime(const time_t *timer);
struct tm *localtime_r(const time_t *timer, struct tm *buf);
size_t strftime(char * restrict s, size_t maxsize, const char * restrict format,
 const struct tm * restrict timeptr);

```

Only if supported by the implementation:

|                             |                          |
|-----------------------------|--------------------------|
| <code>TIME_MONOTONIC</code> | <code>TIME_ACTIVE</code> |
|-----------------------------|--------------------------|

Only if threads are supported and it is supported by the implementation:

`TIME_THREAD_ACTIVE`

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of <time.h>:

|                      |                      |
|----------------------|----------------------|
| <code>errno_t</code> | <code>rsize_t</code> |
|----------------------|----------------------|

```

errno_t asctime_s(char *s, rsize_t maxsize, const struct tm *timeptr);
errno_t ctime_s(char *s, rsize_t maxsize, const time_t *timer);
struct tm *gmtime_s(const time_t * restrict timer, struct tm * restrict result);
struct tm *localtime_s(const time_t * restrict timer, struct tm * restrict result);

```

## B.30 Unicode utilities <uchar.h>

|                                       |                      |                       |
|---------------------------------------|----------------------|-----------------------|
| <code>__STDC_VERSION_UCHAR_H__</code> | <code>size_t</code>  | <code>char16_t</code> |
| <code>mbstate_t</code>                | <code>char8_t</code> | <code>char32_t</code> |

```

size_t mbrtoc8(char8_t * restrict pc8, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
size_t c8rtomb(char * restrict s, char8_t c8, mbstate_t * restrict ps);
size_t mbrtoc16(char16_t * restrict pc16, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
size_t c16rtomb(char * restrict s, char16_t c16, mbstate_t * restrict ps);
size_t mbrtoc32(char32_t * restrict pc32, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
size_t c32rtomb(char * restrict s, char32_t c32, mbstate_t * restrict ps);

```

### B.31 Extended multibyte/wide character utilities <wchar.h>

|                                       |                          |                          |
|---------------------------------------|--------------------------|--------------------------|
| <code>wchar_t</code>                  | <code>NULL</code>        | <code>WCHAR_UTF32</code> |
| <code>size_t</code>                   | <code>WCHAR_MAX</code>   | <code>MB_UTF8</code>     |
| <code>mbstate_t</code>                | <code>WCHAR_MIN</code>   | <code>MB_UTF16</code>    |
| <code>wint_t</code>                   | <code>WCHAR_WIDTH</code> | <code>MB_UTF32</code>    |
| <code>struct tm</code>                | <code>WCHAR_UTF8</code>  | <code>WEOF</code>        |
| <code>__STDC_VERSION_WCHAR_H__</code> | <code>WCHAR_UTF16</code> |                          |

```

int fwprintf(FILE * restrict stream, const wchar_t * restrict format, ...);
int fwscanf(FILE * restrict stream, const wchar_t * restrict format, ...);
int swprintf(wchar_t * restrict s, size_t n, const wchar_t * restrict format,
 ...);
int swscanf(const wchar_t * restrict s, const wchar_t * restrict format, ...);
int vfwprintf(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);
int vfwscanf(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);
int vswprintf(wchar_t * restrict s, size_t n, const wchar_t * restrict format,
 va_list arg);
int vswscanf(const wchar_t * restrict s, const wchar_t * restrict format,
 va_list arg);
int vwprintf(const wchar_t * restrict format, va_list arg);
int vwscanf(const wchar_t * restrict format, va_list arg);
int wprintf(const wchar_t * restrict format, ...);
int wscanf(const wchar_t * restrict format, ...);
wint_t fgetwc(FILE *stream);
wchar_t *fgetws(wchar_t * restrict s, int n, FILE * restrict stream);
wint_t fputwc(wchar_t c, FILE *stream);
int fputws(const wchar_t * restrict s, FILE * restrict stream);
int fwide(FILE *stream, int mode);
wint_t getwc(FILE *stream);
wint_t getwchar(void);
wint_t putwc(wchar_t c, FILE *stream);
wint_t putwchar(wchar_t c);
wint_t ungetwc(wint_t c, FILE *stream);
double wcstod(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
float wcstof(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
long double wcstold(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
long int wcstol(const wchar_t * restrict nptr, wchar_t ** restrict endptr,
 int base);
long long int wcstoll(const wchar_t * restrict nptr, wchar_t ** restrict endptr,
 int base);
unsigned long int wcstoul(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr, int base);

```

```

unsigned long long int wcstoull(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr, int base);
wchar_t *wcscpy(wchar_t * restrict s1, const wchar_t * restrict s2);
wchar_t *wcsncpy(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
wchar_t *wmemcpy(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
wchar_t *wmemmove(wchar_t *s1, const wchar_t *s2, size_t n);
wchar_t *wcscat(wchar_t * restrict s1, const wchar_t * restrict s2);
wchar_t *wcsncat(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
int wcscmp(const wchar_t *s1, const wchar_t *s2);
int wscoll(const wchar_t *s1, const wchar_t *s2);
int wcsncmp(const wchar_t *s1, const wchar_t *s2, size_t n);
size_t wcsxfrm(wchar_t * restrict s1, const wchar_t * restrict s2, size_t n);
int wmemcmp(const wchar_t *s1, const wchar_t *s2, size_t n);
QWchar_t *wcschr(QWchar_t *s, wchar_t c);
size_t wcslen(const wchar_t *s);
QWchar_t *wcpbrk(QWchar_t *s1, const wchar_t *s2);
QWchar_t *wcsrchr(QWchar_t *s, wchar_t c);
size_t wcsspn(const wchar_t *s1, const wchar_t *s2);
QWchar_t *wcsstr(QWchar_t *s1, const wchar_t *s2);
wchar_t *wcstok(wchar_t * restrict s1, const wchar_t * restrict s2,
 wchar_t ** restrict ptr);
QWchar_t *wmemchr(QWchar_t *s, wchar_t c, size_t n);
size_t wcslen(const wchar_t *s);
size_t wcsnlen(const wchar_t *s, size_t n);
wchar_t *wmemset(wchar_t *s, wchar_t c, size_t n);
size_t wcsftime(wchar_t * restrict s, size_t maxsize,
 const wchar_t * restrict format, const struct tm * restrict timeptr);
wint_t btowc(int c);
int wctob(wint_t c);
int mbsinit(const mbstate_t *ps);
size_t mbrlen(const char * restrict s, size_t n, mbstate_t * restrict ps);
size_t mbrtowc(wchar_t * restrict pwc, const char * restrict s, size_t n,
 mbstate_t * restrict ps);
size_t wctrmb(char * restrict s, wchar_t wc, mbstate_t * restrict ps);
size_t mbsrtowcs(wchar_t * restrict dst, const char ** restrict src, size_t len,
 mbstate_t * restrict ps);
size_t wcsrtombs(char * restrict dst, const wchar_t ** restrict src, size_t len,
 mbstate_t * restrict ps);

```

Only if the implementation defines `__STDC_IEC_60559_DFP__`:

```

_Decimal32 wcstod32(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
_Decimal64 wcstod64(const wchar_t * restrict nptr, wchar_t ** restrict endptr);
_Decimal128 wcstod128(const wchar_t * restrict nptr, wchar_t ** restrict endptr);

```

Only if the implementation defines `__STDC_IEC_60559_TYPES__` and additionally the user code defines `__STDC_WANT_IEC_60559_TYPES_EXT__` before any inclusion of `<wchar.h>`:

```

_FloatN wcstofN(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
_FloatNx wcstofNx(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
_DecimalN wcstodN(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
_DecimalNx wcstodNx(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
void wcstoencfN(unsigned char encptr[restrict static N/8],
 const wchar_t * restrict nptr, wchar_t ** restrict endptr);
void wcstoencdecN(unsigned char encptr[restrict static N/8],
 const wchar_t * restrict nptr, wchar_t ** restrict endptr);
void wcstoencbindN(unsigned char encptr[restrict static N/8],

```

```
const wchar_t * restrict nptr, wchar_t ** restrict endptr);
```

Only if the implementation defines `__STDC_LIB_EXT1__` and additionally the user code defines `__STDC_WANT_LIB_EXT1__` before any inclusion of `<wchar.h>`:

**errno\_t****rsize\_t**

```
int fwprintf_s(FILE * restrict stream, const wchar_t * restrict format, ...);
int fwsscanf_s(FILE * restrict stream, const wchar_t * restrict format, ...);
int snwprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
...);
int swprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
...);
int swscanf_s(const wchar_t * restrict s, const wchar_t * restrict format, ...);
int vfwprintf_s(FILE * restrict stream, const wchar_t * restrict format,
va_list arg);
int vfwscanf_s(FILE * restrict stream, const wchar_t * restrict format,
va_list arg);
int vsnprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
va_list arg);
int vswprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
va_list arg);
int vswscanf_s(const wchar_t * restrict s, const wchar_t * restrict format,
va_list arg);
int vwprintf_s(const wchar_t * restrict format, va_list arg);
int vwsscanf_s(const wchar_t * restrict format, va_list arg);
int wprintf_s(const wchar_t * restrict format, ...);
int wscanf_s(const wchar_t * restrict format, ...);
errno_t wcsncpy_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2);
errno_t wcsncpy_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2, rsize_t n);
errno_t wmemcpy_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2, rsize_t n);
errno_t wmemmove_s(wchar_t *s1, rsize_t slmax, const wchar_t *s2, rsize_t n);
errno_t wscat_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2);
errno_t wcsncat_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2, rsize_t n);
wchar_t *wcstok_s(wchar_t * restrict s1, rsize_t slmax,
const wchar_t * restrict s2, wchar_t ** restrict ptr);
size_t wcsnlen_s(const wchar_t *s, size_t maxsize);
errno_t wcrctomb_s(size_t * restrict retval, char * restrict s, rsize_t smax,
wchar_t wc, mbstate_t * restrict ps);
errno_t mbsrtowcs_s(size_t * restrict retval, wchar_t * restrict dst,
rsize_t dstmax, const char ** restrict src, rsize_t len,
mbstate_t * restrict ps);
errno_t wcsrtombs_s(size_t * restrict retval, char * restrict dst,
rsize_t dstmax, const wchar_t ** restrict src, rsize_t len,
mbstate_t * restrict ps);
```

## B.32 Wide character classification and mapping utilities `<wctype.h>`

**wint\_t****wctrans\_t****wctype\_t****WEOF**

```
int iswalnum(wint_t wc);
int iswalpna(wint_t wc);
int iswblank(wint_t wc);
```

```
int iswcntrl(wint_t wc);
int iswdigit(wint_t wc);
int iswgraph(wint_t wc);
int iswlower(wint_t wc);
int iswprint(wint_t wc);
int iswpunct(wint_t wc);
int iswspace(wint_t wc);
int iswupper(wint_t wc);
int iswxdigit(wint_t wc);
int iswctype(wint_t wc, wctype_t desc);
wctype_t wctype(const char *property);
wint_t towlower(wint_t wc);
wint_t towupper(wint_t wc);
wint_t towctrans(wint_t wc, wctrans_t desc);
wctrans_t wctrans(const char *property);
```

## Annex C

(informative)

### Sequence points

#### C.1 Known Sequence Points

- 1 The following are the sequence points described in 5.2.2.4:
  - Between the evaluations of the function designator and actual arguments in a function call and the actual call. (6.5.3.3).
  - Between the evaluations of the first and second operands of the following operators: logical AND && (6.5.14); logical OR || (6.5.15); comma , (6.5.18).
  - Between the evaluations of the first operand of the conditional ?: operator and whichever of the second and third operands is evaluated (6.5.16).
  - Between the evaluation of a full expression and the next full expression to be evaluated. The following are full expressions: a full declarator for a variably modified type; an initializer that is not part of a compound literal (6.7.11); the expression in an expression statement (6.8.4); the controlling expression of a selection statement (**if** or **switch**) (6.8.5); the controlling expression of a **while** or **do** statement (6.8.6); each of the (optional) expressions of a **for** statement (6.8.6.4); the (optional) expression in a **return** statement (6.8.7.5).
  - Immediately before a library function returns (7.1.4).
  - After the actions associated with each formatted input/output function conversion specifier (7.23.6, 7.32.2).
  - Immediately before and immediately after each call to a comparison function, and also between any call to a comparison function and any movement of the objects passed as arguments to that call (7.24.6).



## Annex D

(informative)

### Universal character names for identifiers

#### D.1 Introduction

- 1 This subclause describes the choices made in application of UAX #31 (“Unicode Identifier and Pattern Syntax”) to C of the requirements from UAX #31 and how they do or do not apply to C. For UAX #31, C conforms by meeting the requirements “Default Identifiers” (D.2) and “Equivalent Normalized Identifiers” (D.2). The other requirements, also listed in the following subclauses, are either alternatives not taken or do not apply to C.

#### D.2 Default Identifiers

##### D.2.1 General

- 1 UAX #31 specifies a default syntax for identifiers based on properties from the Unicode Character Database, UAX #44. The general syntax is

```
<Identifier> := <Start> <Continue>* (<Medial> <Continue>+)*
```

where `<Start>` has the `XID_Start` property, `<Continue>` has the `XID_Continue` property, and `<Medial>` is a list of characters permitted between continue characters. For C we add the character `_` (U+005F, LOW LINE) to the set of permitted Start characters, the Medial set is empty, and the Continue characters are unmodified. In the grammar used in UAX #31, this is

```
<Identifier> := <Start> <Continue>*
<Start> := XID_Start + U+005F
<Continue> := <Start> + XID_Continue
```

Additionally, implementations may add the character `$` (U+0024, DOLLAR SIGN) to the set of permitted Start and Continue characters. This is described in the C grammar (6.4.3.1), where *identifier* is formed from *identifier-start* or *identifier* followed by *identifier-continue*.

##### D.2.2 Restricted Format Characters

- 1 If an implementation of UAX #31 wishes to allow format characters such as ZERO WIDTH JOINER or ZERO WIDTH NON-JOINER it shall define a profile allowing them, or describe precisely which combinations are permitted.
- 2 C does not allow format characters in identifiers, so this does not apply.

##### D.2.3 Stable Identifiers

- 1 An implementation of UAX #31 may choose to guarantee that identifiers are stable across versions of the Unicode Standard. Once a string qualifies as an identifier it does so in all future versions. C does not make this guarantee, except to the extent that UAX #31 guarantees the stability of the `XID_Start` and `XID_Continue` properties.

#### D.3 Immutable Identifiers

- 1 An implementation may choose to guarantee that the set of identifiers will never change by fixing the set of code points allowed in identifiers forever.
- 2 C does not choose to make this guarantee. As scripts are added to Unicode, additional characters in those scripts may become available for use in identifiers.

#### D.4 Pattern\_White\_Space and Pattern\_Syntax Characters

- 1 UAX #31 describes how languages that use or interpret patterns of characters, such as regular expressions or number formats, may describe that syntax with Unicode properties.
- 2 C does not do this as part of the language, deferring to library components for such usage of patterns. This requirement does not apply to C.

### **D.5 Equivalent Normalized Identifiers**

- 1 UAX #31 requires that implementations describe how identifiers are compared and considered equivalent.
- 2 C requires that identifiers be in Normalization Form C and therefore identifiers that compare the same under NFC are equivalent. This is described in 6.4.3.

### **D.6 Equivalent Case-Insensitive Identifiers**

- 1 C considers case to be significant in identifier comparison, and does not do any case folding. This requirement does not apply to C

### **D.7 Filtered Normalized Identifiers**

- 1 If any characters are excluded from normalization, UAX #31 requires a precise specification of those exclusions.
- 2 C does not make any such exclusions.

### **D.8 Filtered Case-Insensitive Identifiers**

- 1 C identifiers are case sensitive, and therefore this requirement does not apply.

### **D.9 Hashtag Identifiers**

- 1 There are no hashtags in C, so this requirement does not apply.

## Annex E

(informative)

### Implementation limits

#### E.1 Introduction

- 1 The contents of the header `<limits.h>` are given in the following subclauses. The values shall all be constant expressions suitable for use in conditional expression inclusion preprocessing directives. The components are described further in 5.3.5.3.2.

#### E.2 Minimum values

- 1 For the following macros, the minimum values shown shall be replaced by implementation-defined values.

|                |                        |                     |                                                    |
|----------------|------------------------|---------------------|----------------------------------------------------|
| <b>#define</b> | <b>BOOL_WIDTH</b>      | 1                   | // exact value                                     |
| <b>#define</b> | <b>CHAR_BIT</b>        | 8                   |                                                    |
| <b>#define</b> | <b>USHRT_WIDTH</b>     | 16                  |                                                    |
| <b>#define</b> | <b>UINT_WIDTH</b>      | 16                  |                                                    |
| <b>#define</b> | <b>ULONG_WIDTH</b>     | 32                  |                                                    |
| <b>#define</b> | <b>ULLONG_WIDTH</b>    | 64                  |                                                    |
| <b>#define</b> | <b>BITINT_MAXWIDTH</b> | <b>ULLONG_WIDTH</b> | // at minimum as large<br>// as unsigned long long |
| <b>#define</b> | <b>MB_LEN_MAX</b>      | 1                   |                                                    |

- 2 For the following macros, the minimum magnitudes shown shall be replaced by implementation-defined magnitudes with the same sign that are deduced from the prior macros as indicated.<sup>428)</sup>

|                |                    |                                      |                                    |
|----------------|--------------------|--------------------------------------|------------------------------------|
| <b>#define</b> | <b>BOOL_MAX</b>    | 1                                    | // $2^{\text{BOOL\_WIDTH}} - 1$    |
| <b>#define</b> | <b>CHAR_MAX</b>    | <b>UCHAR_MAX</b> or <b>SCHAR_MAX</b> |                                    |
| <b>#define</b> | <b>CHAR_MIN</b>    | 0 or <b>SCHAR_MIN</b>                |                                    |
| <b>#define</b> | <b>CHAR_WIDTH</b>  | 8                                    | // CHAR_BIT                        |
| <b>#define</b> | <b>UCHAR_MAX</b>   | 255                                  | // $2^{\text{UCHAR\_WIDTH}} - 1$   |
| <b>#define</b> | <b>UCHAR_WIDTH</b> | 8                                    | // CHAR_BIT                        |
| <b>#define</b> | <b>USHRT_MAX</b>   | 65535                                | // $2^{\text{USHRT\_WIDTH}} - 1$   |
| <b>#define</b> | <b>SCHAR_MAX</b>   | +127                                 | // $2^{\text{SCHAR\_WIDTH}-1} - 1$ |
| <b>#define</b> | <b>SCHAR_MIN</b>   | -128                                 | // $-2^{\text{SCHAR\_WIDTH}-1}$    |
| <b>#define</b> | <b>SCHAR_WIDTH</b> | 8                                    | // CHAR_BIT                        |
| <b>#define</b> | <b>SHRT_MAX</b>    | +32767                               | // $2^{\text{SHRT\_WIDTH}-1} - 1$  |
| <b>#define</b> | <b>SHRT_MIN</b>    | -32768                               | // $-2^{\text{SHRT\_WIDTH}-1}$     |
| <b>#define</b> | <b>SHRT_WIDTH</b>  | 16                                   | // USHRT_WIDTH                     |
| <b>#define</b> | <b>INT_MAX</b>     | +32767                               | // $2^{\text{INT\_WIDTH}-1} - 1$   |
| <b>#define</b> | <b>INT_MIN</b>     | -32768                               | // $-2^{\text{INT\_WIDTH}-1}$      |
| <b>#define</b> | <b>INT_WIDTH</b>   | 16                                   | // UINT_WIDTH                      |
| <b>#define</b> | <b>UINT_MAX</b>    | 65535                                | // $2^{\text{UINT\_WIDTH}} - 1$    |
| <b>#define</b> | <b>LONG_MAX</b>    | +2147483647                          | // $2^{\text{LONG\_WIDTH}-1} - 1$  |
| <b>#define</b> | <b>LONG_MIN</b>    | -2147483648                          | // $-2^{\text{LONG\_WIDTH}-1}$     |
| <b>#define</b> | <b>LONG_WIDTH</b>  | 32                                   | // ULONG_WIDTH                     |
| <b>#define</b> | <b>LLONG_MAX</b>   | +9223372036854775807                 | // $2^{\text{LLONG\_WIDTH}-1} - 1$ |
| <b>#define</b> | <b>LLONG_MIN</b>   | -9223372036854775808                 | // $-2^{\text{LLONG\_WIDTH}-1}$    |
| <b>#define</b> | <b>LLONG_WIDTH</b> | 64                                   | // ULLONG_WIDTH                    |
| <b>#define</b> | <b>ULONG_MAX</b>   | 4294967295                           | // $2^{\text{ULONG\_WIDTH}} - 1$   |
| <b>#define</b> | <b>ULLONG_MAX</b>  | 18446744073709551615                 | // $2^{\text{ULLONG\_WIDTH}} - 1$  |

- 3 The contents of the header `<float.h>` are given in the following subclauses. All integer values, except **FLT\_ROUNDS**, shall be constant expressions suitable for use in **#if** preprocessing directives;

<sup>428)</sup>For the minimum value of a signed integer type there is no expression consisting of a minus sign and a decimal literal of that same type. The numbers in the table are only given as indications for the values and do not represent suitable expressions to be used for these macros.

all floating values shall be arithmetic constant expressions. The components are described further in 5.3.5.3.3 and 5.3.5.3.4.

- 4 The values given in the following list shall be replaced by implementation-defined expressions:

```
#define FLT_EVAL_METHOD
#define FLT_ROUND5
#ifdef __STDC_IEC_60559_DFP__
#define DEC_EVAL_METHOD
#endif
```

- 5 The values given in the following list shall be replaced by implementation-defined constant expressions that are greater or equal in magnitude (absolute value) to those shown, with the same sign:

```
#define DBL_DECIMAL_DIG 10
#define DBL_DIG 10
#define DBL_MANT_DIG
#define DBL_MAX_10_EXP +37
#define DBL_MAX_EXP
#define DBL_MIN_10_EXP -37
#define DBL_MIN_EXP
#define DECIMAL_DIG 10
#define FLT_DECIMAL_DIG 6
#define FLT_DIG 6
#define FLT_MANT_DIG
#define FLT_MAX_10_EXP +37
#define FLT_MAX_EXP
#define FLT_MIN_10_EXP -37
#define FLT_MIN_EXP
#define FLT_RADIX 2
#define LDBL_DECIMAL_DIG 10
#define LDBL_DIG 10
#define LDBL_MANT_DIG
#define LDBL_MAX_10_EXP +37
#define LDBL_MAX_EXP
#define LDBL_MIN_10_EXP -37
#define LDBL_MIN_EXP
```

- 6 The values given in the following list shall be replaced by implementation-defined constant expressions with values that are greater than or equal to those shown:

```
#define DBL_MAX 1E+37
#define DBL_NORM_MAX 1E+37
#define FLT_MAX 1E+37
#define FLT_NORM_MAX 1E+37
#define LDBL_MAX 1E+37
#define LDBL_NORM_MAX 1E+37
```

- 7 The values given in the following list shall be replaced by implementation-defined constant expressions with (positive) values that are less than or equal to those shown:

```
#define DBL_EPSILON 1E-9
#define DBL_MIN 1E-37
#define FLT_EPSILON 1E-5
#define FLT_MIN 1E-37
#define LDBL_EPSILON 1E-9
#define LDBL_MIN 1E-37
```

- 8 If the implementation supports decimal floating types, the following macros provide the parameters of these types as exact values.

[illegible]

## Annex F

(normative)

# ISO/IEC 60559 floating-point arithmetic

### F.1 Introduction

- 1 This annex specifies C language support for the *ISO/IEC 60559 floating-point standard*. The *ISO/IEC 60559 floating-point standard* is specifically Floating-point arithmetic (ISO/IEC 60559:2020), also designated as *IEEE Standard for Floating-Point Arithmetic* (IEEE 754–2019). ISO/IEC 60559 generally refers to the floating-point standard, as in ISO/IEC 60559 operation, ISO/IEC 60559 format, etc.
- 2 The ISO/IEC 60559 floating-point standard is a minor upgrade to ISO/IEC/IEEE 60559:2011 (IEEE 754–2008). ISO/IEC/IEEE 60559:2011 was a major upgrade to IEC 60559:1989 (IEEE 754–1985), specifying decimal as well as binary floating-point arithmetic.
- 3 An implementation that defines `__STDC_IEC_60559_BFP__` to 202ymmL shall conform to the specifications in this annex for binary floating-point arithmetic and shall also define `__STDC_IEC_559__` to 1.<sup>429)</sup>
- 4 An implementation that defines `__STDC_IEC_60559_DFP__` to 202ymmL shall conform to the specifications for decimal floating-point arithmetic in the following subclauses of this annex:
  - F.2.2 Infinities and NaNs
  - F.3 Operations
  - F.4 Floating to integer conversions
  - F.6 The **return** statement
  - F.7 Contracted expressions
  - F.8 Floating-point environment
  - F.9 Optimization
  - F.10 Mathematics `<math.h>` and `<tgmath.h>`

For the purpose of specifying these conformance requirements, the macros, functions, and values mentioned in the subclauses listed prior are understood to refer to the corresponding macros, functions, and values for decimal floating types. Likewise, the “rounding direction mode” is understood to refer to the rounding direction mode for decimal floating-point arithmetic.

- 5 Where a binding between the C language and ISO/IEC 60559 is indicated, the ISO/IEC 60559-specified behavior is adopted by reference, unless stated otherwise.

### F.2 Types

#### F.2.1 General

- 1 The C floating types match the ISO/IEC 60559 formats as follows:
  - The **float** type matches the ISO/IEC 60559 binary32 format.
  - The **double** type matches the ISO/IEC 60559 binary64 format.
  - The **long double** type matches the ISO/IEC 60559 binary128 format, else an ISO/IEC 60559 binary64-extended format,<sup>430)</sup> else a non-ISO/IEC 60559 extended format, else the ISO/IEC 60559 binary64 format.

Any non-ISO/IEC 60559 extended format used for the **long double** type shall have more precision than ISO/IEC 60559 binary64 and at least the range of ISO/IEC 60559 binary64.<sup>431)</sup> The value

<sup>429)</sup>Implementations that do not define either of `__STDC_IEC_60559_BFP__` and `__STDC_IEC_559__` are not required to conform to these specifications. New code should not use the obsolescent macro `__STDC_IEC_559__` to test for conformance to this annex.

<sup>430)</sup>ISO/IEC 60559 binary64-extended formats include the common 80-bit ISO/IEC 60559 format.

<sup>431)</sup>A non-ISO/IEC 60559 **long double** type provides signed infinities, signed zeros, and NaNs, as its values include all **double** values.

of **FLT\_ROUNDS** applies to all ISO/IEC 60559 types supported by the implementation, but is not required to apply to non-ISO/IEC 60559 types.

#### Recommended practice

- 2 The **long double** type should match the ISO/IEC 60559 binary128 format, else an ISO/IEC 60559 binary64-extended format.

### F.2.2 Infinities and NaNs

- 1 Since negative and positive infinity are representable in ISO/IEC 60559 formats, all real numbers lie within the range of representable values (5.3.5.3.3).
- 2 The **NAN** and **INFINITY** macros in `<float.h>` and the **nan** functions in `<math.h>` provide designations for ISO/IEC 60559 quiet NaNs and infinities. The **FLT\_SNAN**, **DBL\_SNAN**, and **LDBL\_SNAN** macros in `<float.h>` provide designations for ISO/IEC 60559 signaling NaNs.
- 3 This annex does not require the full support for signaling NaNs specified in ISO/IEC 60559. This annex uses the term NaN, unless explicitly qualified, to denote quiet NaNs. Where specification of signaling NaNs is not provided, the behavior of signaling NaNs is implementation-defined (either treated as an ISO/IEC 60559 quiet NaN or treated as an ISO/IEC 60559 signaling NaN).<sup>432)</sup>
- 4 Any operator or `<math.h>` function that raises an “invalid” floating-point exception, if delivering a floating type result, shall return a quiet NaN, unless explicitly specified otherwise.
- 5 To support signaling NaNs as specified in ISO/IEC 60559, an implementation should adhere to the following recommended practice.

#### Recommended practice

- 6 Any floating-point operator or `<math.h>` function or macro with a signaling NaN input, unless explicitly specified otherwise, raises an “invalid” floating-point exception.
- 7 NOTE Some functions do not propagate quiet NaN arguments. For example, **hypot**(*x*, *y*) returns infinity if *x* or *y* is infinite and the other is a quiet NaN. The recommended practice in this subclause specifies that such functions (and others) raise the “invalid” floating-point exception if an argument is a signaling NaN, which also implies they return a quiet NaN in these cases.
- 8 The `<fenv.h>` header defines the macro **FE\_SNANS\_ALWAYS\_SIGNAL** if and only if the implementation follows the recommended practice in this subclause. If defined, **FE\_SNANS\_ALWAYS\_SIGNAL** expands to the integer literal 1.

### F.3 Operations

- 1 C operators, functions, and function-like macros provide operations specified by ISO/IEC 60559 as shown in Table F.2. In the table, C functions are represented by the function name without a type suffix. Specifications for the C facilities are provided in the listed clauses. The C specifications are intended to match ISO/IEC 60559, unless stated otherwise.

Table F.2 — Operation binding — mathematical operations

| ISO/IEC 60559 operation       | C operation      | Clause              |
|-------------------------------|------------------|---------------------|
| roundToIntegralTiesToEven     | <b>roundeven</b> | 7.12.10.8, F.10.7.8 |
| roundToIntegralTiesAway       | <b>round</b>     | 7.12.10.6, F.10.7.6 |
| roundToIntegralTowardZero     | <b>trunc</b>     | 7.12.10.9, F.10.7.9 |
| roundToIntegralTowardPositive | <b>ceil</b>      | 7.12.10.1, F.10.7.1 |
| roundToIntegralTowardNegative | <b>floor</b>     | 7.12.10.2, F.10.7.2 |

<sup>432)</sup>Since NaNs created by ISO/IEC 60559 arithmetic operations are always quiet, quiet NaNs (along with infinities) are sufficient for closure of the arithmetic.

|                        |                                    |                                                                      |
|------------------------|------------------------------------|----------------------------------------------------------------------|
| roundToIntegralExact   | <b>rint</b>                        | 7.12.10.4, F.10.7.4                                                  |
| nextUp                 | <b>nextup</b>                      | 7.12.12.5, F.10.9.5                                                  |
| nextDown               | <b>nextdown</b>                    | 7.12.12.6, F.10.9.6                                                  |
| getPayload             | <b>getpayload</b>                  | F.10.14.2                                                            |
| setPayload             | <b>setpayload</b>                  | F.10.14.3                                                            |
| setPayloadSignaling    | <b>setpayloadsig</b>               | F.10.14.4                                                            |
| quantize               | <b>quantize</b>                    | 7.12.16.1                                                            |
| sameQuantum            | <b>samequantum</b>                 | 7.12.16.2                                                            |
| quantum                | <b>quantum</b>                     | 7.12.16.3                                                            |
| encodeDecimal          | <b>encodedec</b>                   | 7.12.17.2                                                            |
| decodeDecimal          | <b>decodedec</b>                   | 7.12.17.3                                                            |
| encodeBinary           | <b>encodebin</b>                   | 7.12.17.4                                                            |
| decodeBinary           | <b>decodebin</b>                   | 7.12.17.5                                                            |
| remainder              | <b>remainder, remquo</b>           | 7.12.11.2, F.10.8.2,<br>7.12.11.3, F.10.8.3                          |
| maximum                | <b>fmaximum</b>                    | 7.12.13.4, F.10.10.4                                                 |
| minimum                | <b>fminimum</b>                    | 7.12.13.5, F.10.10.4                                                 |
| maximumMagnitude       | <b>fmaximum_mag</b>                | 7.12.13.6, F.10.10.4                                                 |
| minimumMagnitude       | <b>fminimum_mag</b>                | 7.12.13.7, F.10.10.4                                                 |
| maximumNumber          | <b>fmaximum_num</b>                | 7.12.13.8, F.10.10.5                                                 |
| minimumNumber          | <b>fminimum_num</b>                | 7.12.13.9, F.10.10.5                                                 |
| maximumMagnitudeNumber | <b>fmaximum_mag_num</b>            | 7.12.13.10, F.10.10.5                                                |
| minimumMagnitudeNumber | <b>fminimum_mag_num</b>            | 7.12.13.11, F.10.10.5                                                |
| scaleB                 | <b>scalbn, scalbln</b>             | 7.12.7.19, F.10.4.19                                                 |
| logB                   | <b>logb, ilogb, llogb</b>          | 7.12.7.17, F.10.4.17,<br>7.12.7.8, F.10.4.8,<br>7.12.7.10, F.10.4.10 |
| addition               | <b>+, fadd, faddl, daddl</b>       | 6.5.7, 7.12.15.2,<br>F.10.12                                         |
| subtraction            | <b>-, fsub, fsubl, dsubl</b>       | 6.5.7, 7.12.15.3,<br>F.10.12                                         |
| multiplication         | <b>*, fmul, fmul, dmull</b>        | 6.5.6, 7.12.15.4,<br>F.10.12                                         |
| division               | <b>/, fdiv, fdivl, ddivl</b>       | 6.5.6, 7.12.15.5,<br>F.10.12                                         |
| squareRoot             | <b>sqrt, fsqrt, fsqrtl, dsqrtl</b> | 7.12.8.10, F.10.5.10,<br>7.12.15.7, F.10.12                          |
| fusedMultiplyAdd       | <b>fma, ffma, fmal, dfmal</b>      | 7.12.14.1, F.10.11.1,<br>7.12.15.6, F.10.12                          |



|                                                                                                                                                                                     |                                                                         |                                                |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------|------------------------------------------------|
| convertFromInt                                                                                                                                                                      | cast and implicit conversion                                            | 6.3.2.4, 6.5.5                                 |
| convertToIntegerTiesToEven<br>convertToIntegerTowardZero<br>convertToIntegerTowardPositive<br>convertToIntegerTowardNegative                                                        | <b>fromfp, ufromfp</b>                                                  | 7.12.10.10, F.10.7.10                          |
| convertToIntegerTiesToAway                                                                                                                                                          | <b>fromfp, ufromfp, lround, llround</b>                                 | 7.12.10.10, F.10.7.10, 7.12.10.7, F.10.7.7     |
| convertToIntegerExactTiesToEven<br>convertToIntegerExactTowardZero<br>convertToIntegerExactTowardPositive<br>convertToIntegerExactTowardNegative<br>convertToIntegerExactTiesToAway | <b>fromfpx, ufromfpx</b>                                                | 7.12.10.11, F.10.7.11                          |
| convertFormat - different formats                                                                                                                                                   | cast and implicit conversions                                           | 6.3.2.5, 6.5.5                                 |
| convertFormat - same format                                                                                                                                                         | <b>canonicalize</b>                                                     | 7.12.12.7, F.10.9.7                            |
| convertFromDecimalCharacter                                                                                                                                                         | <b>strtod, wcstod, scanf, wscanf</b> ,<br>decimal floating literals     | 7.24.2.6, 7.32.4.2.2, 7.23.6.5, 7.32.2.13, F.5 |
| convertToDecimalCharacter                                                                                                                                                           | <b>printf, wprintf, strfromd</b>                                        | 7.23.6.4, 7.32.2.12, 7.24.2.4, F.5             |
| convertFromHexCharacter                                                                                                                                                             | <b>strtod, wcstod, scanf, wscanf</b> ,<br>hexadecimal floating literals | 7.24.2.6, 7.32.4.2.2, 7.23.6.5, 7.32.2.13, F.5 |
| convertToHexCharacter                                                                                                                                                               | <b>printf, wprintf, strfromd</b>                                        | 7.23.6.4, 7.32.2.12, 7.24.2.4, F.5             |
| copy                                                                                                                                                                                | <b>memcpy, memmove, + (x)</b>                                           | 7.27.2.1, 7.27.2.3                             |
| negate                                                                                                                                                                              | - (x)                                                                   | 6.5.4.4                                        |
| abs                                                                                                                                                                                 | <b>fabs</b>                                                             | 7.12.8.3, F.10.5.3                             |
| copySign                                                                                                                                                                            | <b>copysign</b>                                                         | 7.12.12.1, F.10.9.1                            |
| compareQuietEqual                                                                                                                                                                   | ==                                                                      | 6.5.10, F.9.4                                  |
| compareQuietNotEqual                                                                                                                                                                | !=                                                                      | 6.5.10, F.9.4                                  |
| compareSignalingEqual                                                                                                                                                               | <b>iseqsig</b>                                                          | 7.12.18.8, F.10.15.2                           |
| compareSignalingGreater                                                                                                                                                             | >                                                                       | 6.5.9, F.9.4                                   |
| compareSignalingGreaterEqual                                                                                                                                                        | >=                                                                      | 6.5.9, F.9.4                                   |
| compareSignalingLess                                                                                                                                                                | <                                                                       | 6.5.9, F.9.4                                   |
| compareSignalingLessEqual                                                                                                                                                           | <=                                                                      | 6.5.9, F.9.4                                   |
| compareSignalingNotEqual                                                                                                                                                            | <b>! iseqsig(x)</b>                                                     | 7.12.18.8, F.10.15.2                           |
| compareSignalingNotGreater                                                                                                                                                          | <b>! (x &gt; y)</b>                                                     | 6.5.9, F.9.4                                   |
| compareSignalingLessUnordered                                                                                                                                                       | <b>! (x &gt;= y)</b>                                                    | 6.5.9, F.9.4                                   |
| compareSignalingNotLess                                                                                                                                                             | <b>! (x &lt; y)</b>                                                     | 6.5.9, F.9.4                                   |
| compareSignalingGreaterUnordered                                                                                                                                                    | <b>! (x &lt;= y)</b>                                                    | 6.5.9, F.9.4                                   |
| compareQuietGreater                                                                                                                                                                 | <b>isgreater</b>                                                        | 7.12.18.2                                      |

|                              |                                         |                              |
|------------------------------|-----------------------------------------|------------------------------|
| compareQuietGreaterEqual     | <b>isgreaterequal</b>                   | 7.12.18.3                    |
| compareQuietLess             | <b>isless</b>                           | 7.12.18.4                    |
| compareQuietLessEqual        | <b>islessequal</b>                      | 7.12.18.5                    |
| compareQuietUnordered        | <b>isunordered</b>                      | 7.12.18.7                    |
| compareQuietNotGreater       | <b>! isgreater</b> (x, y)               | 7.12.18.2                    |
| compareQuietLessUnordered    | <b>! isgreaterequal</b> (x, y)          | 7.12.18.3                    |
| compareQuietNotLess          | <b>! isless</b> (x, y)                  | 7.12.18.4                    |
| compareQuietGreaterUnordered | <b>! islessequal</b> (x, y)             | 7.12.18.5                    |
| compareQuietOrdered          | <b>! isunordered</b> (x, y)             | 7.12.18.7                    |
| class                        | <b>fpclassify, signbit, issignaling</b> | 7.12.4.2, 7.12.4.8, 7.12.4.9 |
| isSignMinus                  | <b>signbit</b>                          | 7.12.4.8                     |
| isNormal                     | <b>isnormal</b>                         | 7.12.4.7                     |
| isFinite                     | <b>isfinite</b>                         | 7.12.4.4                     |
| isZero                       | <b>iszero</b>                           | 7.12.4.11                    |
| isSubnormal                  | <b>issubnormal</b>                      | 7.12.4.10                    |
| isInfinite                   | <b>isinf</b>                            | 7.12.4.5                     |
| isNaN                        | <b>isnan</b>                            | 7.12.4.6                     |
| isSignaling                  | <b>issignaling</b>                      | 7.12.4.9                     |
| isCanonical                  | <b>iscanonical</b>                      | 7.12.4.3                     |
| radix                        | <b>FLT_RADIX</b>                        | 5.3.5.3.3                    |
| totalOrder                   | <b>totalorder</b>                       | F.10.13.2                    |
| totalOrderMag                | <b>totalordermag</b>                    | F.10.13.3                    |
| lowerFlags                   | <b>feclearexcept</b>                    | 7.6.5.2                      |
| raiseFlags                   | <b>fesetexcept</b>                      | 7.6.5.5                      |
| testFlags                    | <b>fetestexcept</b>                     | 7.6.5.8                      |
| testSavedFlags               | <b>fetestexceptflag</b>                 | 7.6.5.7                      |
| restoreFlags                 | <b>fesetexceptflag</b>                  | 7.6.5.6                      |
| saveAllFlags                 | <b>fegetexceptflag</b>                  | 7.6.5.3                      |
| getBinaryRoundingDirection   | <b>fegetround</b>                       | 7.6.6.3                      |
| setBinaryRoundingDirection   | <b>fesetround</b>                       | 7.6.6.6                      |
| saveModes                    | <b>fegetmode</b>                        | 7.6.6.2                      |

|              |                                         |              |
|--------------|-----------------------------------------|--------------|
| restoreModes | <b>fesetmode</b>                        | 7.6.6.5      |
| defaultModes | <b>fesetmode</b> ( <b>FE_DFL_MODE</b> ) | 7.6.6.5, 7.6 |

- 2 The ISO/IEC 60559 requirement that certain of its operations be provided for operands of different formats (of the same radix) is satisfied by C's usual arithmetic conversions (6.3.2.8) and function-call argument conversions (6.5.3.3). For example, the following operations take **float** **f** and **double** **d** inputs and produce a **long double** result:

```
(long double)f * d
powl(f, d)
```

- 3 The functions **fmin** and **fmax** have been superseded by **fminimum\_num** and **fmaximum\_num**. The **fmin** and **fmax** functions provide the minNum and maxNum operations specified in (the superseded) ISO/IEC/IEEE 60559:2011.
- 4 Whether C assignment (6.5.17) (and conversion as if by assignment) to the same format is an ISO/IEC 60559 convertFormat or copy operation<sup>433</sup> is implementation-defined, even if <fenv.h> defines the macro **FE\_SNANS\_ALWAYS\_SIGNAL** (F.2.2). If the return expression of a **return** statement is evaluated to the floating-point format of the return type, it is implementation-defined whether a convertFormat operation is applied to the result of the return expression.
- 5 The unary + and - operators raise no floating-point exceptions, even if the operand is a signaling NaN.
- 6 The C classification macros **fpclassify**, **iscanonical**, **isfinite**, **isinf**, **isnan**, **isnormal**, **issignaling**, **issubnormal**, **iszero**, and **signbit** provide the ISO/IEC 60559 operations indicated in Table F.2 provided their arguments are in the format of their semantic type. Then these macros raise no floating-point exceptions, even if an argument is a signaling NaN.
- 7 The **signbit** macro, providing the ISO/IEC 60559 **isSignMinus** operation, determines the sign of its argument value as the sign bit of the value's representation. This applies to all values, including NaNs whose sign bit is not generally interpreted by ISO/IEC 60559.
- 8 The C **nearbyint** functions (7.12.10.3, F.10.7.3) provide the nearbyinteger function recommended in the Appendix to (superseded) ANSI/IEEE 854–1987.
- 9 The C **nextafter** (7.12.12.3, F.10.9.3) and **nexttoward** (7.12.12.4, F.10.9.4) functions provide the **nextafter** function recommended in the Appendix to (superseded) IEC 60559:1989 (but with a minor change to better handle signed zeros).
- 10 The macros (7.6) **FE\_DOWNWARD**, **FE\_TONEAREST**, **FE\_TONEARESTFROMZERO**, **FE\_TOWARDZERO**, and **FE\_UPWARD**, which are used in conjunction with the **fegetround** and **fesetround** functions and the **FENV\_ROUND** pragma, represent the ISO/IEC 60559 rounding-direction attributes roundTowardNegative, roundTiesToEven, roundTiesToAway, roundTowardZero, and roundTowardPositive, respectively, for binary floating-point arithmetic. Support for the roundTiesToAway attribute for binary floating-point arithmetic, and hence for the **FE\_TONEARESTFROMZERO** macro, is optional.
- 11 The C **fegetenv** (7.6.7.2), **feholdexcept** (7.6.7.3), **fesetenv** (7.6.7.4) and **feupdateenv** (7.6.7.5) functions provide a facility to manage the dynamic floating-point environment, comprising the ISO/IEC 60559 status flags and dynamic control modes.
- 12 ISO/IEC 60559 requires operations with specified operand and result formats. Therefore, math functions that are bound to ISO/IEC 60559 operations (see Table F.2) shall remove any extra range and precision from arguments or results.
- 13 ISO/IEC 60559 requires operations that round their result to formats the same as and wider than the operands, in addition to the operations that round their result to narrower formats (see 7.12.15).

<sup>433</sup>Where the source and destination formats are the same, convertFormat operations differ from copy operations in that convertFormat operations raise the “invalid” floating-point exception on signaling NaN inputs and do not propagate non-canonical encodings.

Operators (+, -, \*, and /) whose evaluation formats are wider than the semantic type (5.3.5.3.3) may not support some of the ISO/IEC 60559 operations, because getting a result in a given format can require a cast that can introduce an extra rounding error. The functions that round result to narrower type (7.12.15) provide the ISO/IEC 60559 operations that round result to same and wider (as well as narrower) formats, in those cases where built-in operators and casts do not. For example, `ddivl(x, y)` computes a correctly rounded **double** divide of **float** `x` by **float** `y`, regardless of the evaluation method.

- 14 Decimal versions of the **remquo** library function are not provided. (The decimal **remainder** functions provide the remainder operation defined by ISO/IEC 60559.)
- 15 The binding for the `convertFormat` operation applies to all conversions among ISO/IEC 60559 formats. Therefore, for implementations that conform to this annex, conversions between decimal floating types and standard floating types with ISO/IEC 60559 formats are correctly rounded and raise floating-point exceptions as specified in ISO/IEC 60559.
- 16 ISO/IEC 60559 specifies the `convertFromHexCharacter` and `convertToHexCharacter` operations only for binary floating-point arithmetic.
- 17 The integer literal `10` provides the radix operation defined in ISO/IEC 60559 for decimal floating-point arithmetic.
- 18 The **fe\_dec\_getround** (7.6.6.4) and **fe\_dec\_setround** (7.6.6.7) functions provide the `getDecimalRoundingDirection` and `setDecimalRoundingDirection` operations defined in ISO/IEC 60559 for decimal floating-point arithmetic. The macros (7.6) **FE\_DEC\_DOWNWARD**, **FE\_DEC\_TONEAREST**, **FE\_DEC\_TONEARESTFROMZERO**, **FE\_DEC\_TOWARDZERO**, and **FE\_DEC\_UPWARD**, which are used in conjunction with the **fe\_dec\_getround** and **fe\_dec\_setround** functions and the **FENV\_DEC\_ROUND** pragma, represent the ISO/IEC 60559 rounding-direction attributes `roundTowardNegative`, `roundTiesToEven`, `roundTiesToAway`, `roundTowardZero`, and `roundTowardPositive`, respectively, for decimal floating-point arithmetic.
- 19 The **llquantexpdN** (7.12.16.4) functions compute the (quantum) exponent  $q$  defined in ISO/IEC 60559 for decimal numbers viewed as having integer significands.
- 20 The C functions in Table F.3 correspond to mathematical operations recommended by ISO/IEC 60559. However, correct rounding, which ISO/IEC 60559 specifies for its operations, is not required for the C functions in the table. 7.34.9 (potentially) reserves **cr\_** prefixed names for functions fully matching the ISO/IEC 60559 mathematical operations. In the table, the C functions are represented by the function name without a type suffix.

Table F.3 — ISO/IEC 60559 operation to C function

| ISO/IEC 60559 operation | C function                  | Clause               |
|-------------------------|-----------------------------|----------------------|
| <code>exp</code>        | <b>exp</b>                  | 7.12.7.1, F.10.4.1   |
| <code>expm1</code>      | <b>expm1</b>                | 7.12.7.6, F.10.4.6   |
| <code>exp2</code>       | <b>exp2</b>                 | 7.12.7.4, F.10.4.4   |
| <code>exp2m1</code>     | <b>exp2m1</b>               | 7.12.7.5, F.10.4.5   |
| <code>exp10</code>      | <b>exp10</b>                | 7.12.7.2, F.10.4.2   |
| <code>exp10m1</code>    | <b>exp10m1</b>              | 7.12.7.3, F.10.4.3   |
| <code>log</code>        | <b>log</b>                  | 7.12.7.11, F.10.4.11 |
| <code>log2</code>       | <b>log2</b>                 | 7.12.7.15, F.10.4.15 |
| <code>log10</code>      | <b>log10</b>                | 7.12.7.12, F.10.4.12 |
| <code>logp1</code>      | <b>log1p</b> , <b>logp1</b> | 7.12.7.14, F.10.4.14 |
| <code>log2p1</code>     | <b>log2p1</b>               | 7.12.7.16, F.10.4.16 |
| <code>log10p1</code>    | <b>log10p1</b>              | 7.12.7.13, F.10.4.13 |
| <code>hypot</code>      | <b>hypot</b>                | 7.12.8.4, F.10.5.4   |
| <code>rSqrt</code>      | <b>rsqrt</b>                | 7.12.8.9, F.10.5.9   |
| <code>compound</code>   | <b>compoundn</b>            | 7.12.8.2, F.10.5.2   |
| <code>rootn</code>      | <b>rootn</b>                | 7.12.8.8, F.10.5.8   |

Table F.3 — ISO/IEC 60559 operation to C function

| ISO/IEC 60559 operation | C function     | Clause               |
|-------------------------|----------------|----------------------|
| pown                    | <b>pown</b>    | 7.12.8.6, F.10.5.6   |
| pow                     | <b>pow</b>     | 7.12.8.5, F.10.5.5   |
| powr                    | <b>powr</b>    | 7.12.8.7, F.10.5.7   |
| sin                     | <b>sin</b>     | 7.12.5.6, F.10.2.6   |
| cos                     | <b>cos</b>     | 7.12.5.5, F.10.2.5   |
| tan                     | <b>tan</b>     | 7.12.5.7, F.10.2.7   |
| sinPi                   | <b>sinpi</b>   | 7.12.5.13, F.10.2.13 |
| cosPi                   | <b>cospi</b>   | 7.12.5.12, F.10.2.12 |
| tanPi                   | <b>tanpi</b>   | 7.12.5.14, F.10.2.14 |
| asinPi                  | <b>asinpi</b>  | 7.12.5.9, F.10.2.9   |
| acosPi                  | <b>acospi</b>  | 7.12.5.8, F.10.2.8   |
| atanPi                  | <b>atanpi</b>  | 7.12.5.10, F.10.2.10 |
| atan2Pi                 | <b>atan2pi</b> | 7.12.5.11, F.10.2.11 |
| asin                    | <b>asin</b>    | 7.12.5.2, F.10.2.2   |
| acos                    | <b>acos</b>    | 7.12.5.1, F.10.2.1   |
| atan                    | <b>atan</b>    | 7.12.5.3, F.10.2.3   |
| atan2                   | <b>atan2</b>   | 7.12.5.4, F.10.2.4   |
| sinh                    | <b>sinh</b>    | 7.12.6.5, F.10.3.5   |
| cosh                    | <b>cosh</b>    | 7.12.6.4, F.10.3.4   |
| tanh                    | <b>tanh</b>    | 7.12.6.6, F.10.3.6   |
| asinh                   | <b>asinh</b>   | 7.12.6.2, F.10.3.2   |
| acosh                   | <b>acosh</b>   | 7.12.6.1, F.10.3.1   |
| atanh                   | <b>atanh</b>   | 7.12.6.3, F.10.3.3   |

#### F.4 Floating to integer conversion

- 1 If the integer type is **bool**, 6.3.2.2 applies and the conversion raises no floating-point exceptions if the floating-point value is not a signaling NaN. Otherwise, if the floating value is infinite or NaN or if the integral part of the floating value exceeds the range of the integer type, then the “invalid” floating-point exception is raised and the resulting value is unspecified. Otherwise, the resulting value is determined by 6.3.2.4. Conversion of an integral floating value that does not exceed the range of the integer type raises no floating-point exceptions; whether conversion of a non-integral floating value raises the “inexact” floating-point exception is unspecified.<sup>434)</sup>

#### F.5 Conversions between binary floating types and decimal character sequences

- 1 The `<float.h>` header defines the macro

**CR\_DECIMAL\_DIG**

if and only if **\_\_STDC\_WANT\_IEC\_60559\_EXT\_\_** is defined as a macro at the point in the source file where `<float.h>` is first included. If defined, **CR\_DECIMAL\_DIG** expands to an integer constant expression suitable for use in conditional expression inclusion preprocessing directives whose value is a number such that conversions between all supported ISO/IEC 60559 binary formats and character sequences with at most **CR\_DECIMAL\_DIG** significant decimal digits are correctly rounded. The value of **CR\_DECIMAL\_DIG** shall be at least  $M + 3$ , where  $M$  is the maximum value of the **T\_DECIMAL\_DIG** macros for ISO/IEC 60559 binary formats. If the implementation correctly rounds for all numbers of significant decimal digits, then **CR\_DECIMAL\_DIG** shall have the value of the macro **UINTMAX\_MAX**.

- 2 Conversions of types with ISO/IEC 60559 binary formats to character sequences with more than

<sup>434)</sup>ISO/IEC 60559 recommends that implicit floating-to-integer conversions raise the “inexact” floating-point exception for non-integer in-range values. In those cases where it matters, library functions can be used to effect such conversions with or without raising the “inexact” floating-point exception. See **fromfp**, **ufromfp**, **fromfpx**, **ufromfpx**, **rint**, **lrint**, **llrint**, and **nearbyint** in `<math.h>`.

**CR\_DECIMAL\_DIG** significant decimal digits shall correctly round to **CR\_DECIMAL\_DIG** significant digits and pad zeros on the right.

- 3 Conversions from character sequences with more than **CR\_DECIMAL\_DIG** significant decimal digits to types with ISO/IEC 60559 binary formats shall correctly round to an intermediate character sequence with **CR\_DECIMAL\_DIG** significant decimal digits, according to the applicable rounding direction, and correctly round the intermediate result (having **CR\_DECIMAL\_DIG** significant decimal digits) to the destination type. The “inexact” floating-point exception is raised (once) if either conversion is inexact.<sup>435)</sup> (The second conversion can raise the “overflow” or “underflow” floating-point exception.)
- 4 The specification in this subclause assures conversion between ISO/IEC 60559 binary format and decimal character sequence follows all pertinent recommended practice. It also assures conversion from ISO/IEC 60559 format to decimal character sequence with at least **T\_DECIMAL\_DIG** digits and back, using to-nearest rounding, is the identity function, where *T* is the macro prefix for the format.
- 5 Functions such as **strtod** that convert character sequences to floating types honor the rounding direction. Hence, if the rounding direction can be upward or downward, the implementation cannot convert a minus-signed sequence by arithmetically negating the converted unsigned sequence.
- 6 NOTE ISO/IEC 60559 specifies that conversion to one-digit character strings using **roundTiesToEven**, when both choices have an odd least significant digit, produce the value with the larger magnitude. This can happen with **9.5e2** whose nearest neighbors are **9.e2** and **1.e3**, both of which have a single odd digit in the significand part.

## F.6 The return statement

If the return expression is evaluated in a floating-point format different from the return type, the expression is converted as if by assignment<sup>436)</sup> to the return type of the function and the resulting value is returned to the caller.

## F.7 Contracted expressions

- 1 A contracted expression is correctly rounded (once) and treats infinities, NaNs, signed zeros, subnormals, and the rounding directions in a manner consistent with the basic arithmetic operations covered by ISO/IEC 60559.

### Recommended practice

- 2 A contracted expression should raise floating-point exceptions in a manner generally consistent with the basic arithmetic operations.

## F.8 Floating-point environment

### F.8.1 General

- 1 The floating-point environment defined in `<fenv.h>` includes the ISO/IEC 60559 floating-point exception status flags and rounding-direction control modes. It may also include other floating-point status or modes that the implementation provides as extensions.<sup>437)</sup>
- 2 This annex does not include support for ISO/IEC 60559’s optional alternate exception handling. The specification in this annex assumes ISO/IEC 60559 default exception handling: the flag is set, a default result is delivered, and execution continues. Implementations may provide alternate exception handling as an extension.

### F.8.2 Environment management

- 1 ISO/IEC 60559 requires that floating-point operations implicitly raise floating-point exception status flags, and that rounding control modes can be set explicitly to affect result values of floating-point operations. These changes to the floating-point state are treated as side effects which respect

<sup>435)</sup> The intermediate conversion is exact only if all input digits after the first **CR\_DECIMAL\_DIG** digits are 0.

<sup>436)</sup> Assignment removes any extra range and precision.

<sup>437)</sup> Dynamic rounding precision and trap enablement modes are examples of such extensions.



sequence points.<sup>438)</sup>

### F.8.3 Translation

- 1 During translation, constant rounding direction modes (7.6.3) are in effect where specified. Elsewhere, during translation the ISO/IEC 60559 default modes are in effect:
  - The rounding direction mode is rounding to nearest.
  - The rounding precision mode (if supported) is set so that results are not shortened.
  - Trapping or stopping (if supported) is disabled on all floating-point exceptions.

#### Recommended practice

- 2 The implementation should produce a diagnostic message for each translation-time floating-point exception, other than “inexact”;<sup>439)</sup> the implementation should then proceed with the translation of the program.

### F.8.4 Execution

- 1 At program startup the dynamic floating-point environment is initialized as prescribed by ISO/IEC 60559:
  - All floating-point exception status flags are cleared.
  - The dynamic rounding direction mode is rounding to nearest.
  - The dynamic rounding precision mode (if supported) is set so that results are not shortened.
  - Trapping or stopping (if supported) is disabled on all floating-point exceptions.

### F.8.5 Constant expressions

- 1 An arithmetic constant expression of floating type, other than one in an initializer for an object that has static or thread storage duration or that is declared with storage-class specifier **constexpr**, is evaluated (as if) during execution; thus, it is affected by any operative floating-point control modes and raises floating-point exceptions as required by ISO/IEC 60559 (provided the state for the **FENV\_ACCESS** pragma is “on”).<sup>440)</sup>
- 2 EXAMPLE

```
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
void f(void)
{
 constexpr double v = 0.0/0.0; // does not raise an exception
 float w[] = { 0.0/0.0 }; // raises an exception
 static float x = 0.0/0.0; // does not raise an exception
 float y = 0.0/0.0; // raises an exception
 double z = 0.0/0.0; // raises an exception
 /* ... */
}
```

<sup>438)</sup> If the state for the **FENV\_ACCESS** pragma is “off”, the implementation is free to assume the dynamic floating-point control modes will be the default ones and the floating-point status flags will not be tested, which allows certain optimizations (see F.9).

<sup>439)</sup> As floating literals are converted to appropriate internal representations at translation time, their conversion is subject to constant or default rounding modes and raises no execution-time floating-point exceptions (even where the state of the **FENV\_ACCESS** pragma is “on”). Library functions, for example **strtod**, provide execution-time conversion of numeric strings.

<sup>440)</sup> Where the state for the **FENV\_ACCESS** pragma is “on”, results of inexact expressions like  $1.0/3.0$  are affected by rounding modes set at execution time, and expressions such as  $0.0/0.0$  and  $1.0/0.0$  generate execution-time floating-point exceptions. The programmer can achieve the efficiency of translation-time evaluation through static initialization, such as

```
const static double one_third = 1.0/3.0;
```

- 3 For the **static** and **constexpr** initializations, the division is done at translation time, raising no (execution-time) floating-point exceptions. On the other hand, for the three automatic initializations the invalid division occurs at execution time.

## F.8.6 Initialization

- 1 All computation for automatic initialization is done (as if) at execution time; thus, it is affected by any operative modes and raises floating-point exceptions as required by ISO/IEC 60559 (provided the state for the **FENV\_ACCESS** pragma is “on”). All computation for initialization of objects that have static or thread storage duration, or that are declared with storage-class specifier **constexpr**, is done (as if) at translation time.
- 2 EXAMPLE

```
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
void f(void)
{
 constexpr float t = (float)1.1e75; // does not raise exceptions
 float u[] = { 1.1e75 }; // raises exceptions
 static float v = 1.1e75; // does not raise exceptions
 float w = 1.1e75; // raises exceptions
 double x = 1.1e75; // may raise exceptions
 float y = 1.1e75f; // may raise exceptions
 long double z = 1.1e75; // does not raise exceptions
 /* ... */
}
```

- 3 The **constexpr** initialization of **t** and the static initialization of **v** raise no (execution-time) floating-point exceptions because their computation is done at translation time. The automatic initialization of **u** and **w** require an execution-time conversion to **float** of the wider value **1.1e75**, which raises floating-point exceptions. The automatic initializations of **x** and **y** entail execution-time conversion; however, in some expression evaluation methods, the conversions are not to a narrower format, in which case no floating-point exception is raised.<sup>441)</sup> The automatic initialization of **z** entails execution-time conversion, but not to a narrower format, so no floating-point exception is raised. The conversions of the floating literals **1.1e75** and **1.1e75f** to their internal representations occur at translation time in all cases.

## F.8.7 Changing the environment

- 1 Operations defined in 6.5.1 and functions and macros defined for the standard libraries change floating-point status flags and control modes just as indicated by their specifications (including conformance to ISO/IEC 60559). They do not change flags or modes (so as to be detectable by the user) in any other cases.
- 2 If the floating-point exceptions represented by the argument to the **feraiseexcept** function in **<fenv.h>** include both “overflow” and “inexact”, then “overflow” is raised before “inexact”. Similarly, if the represented exceptions include both “underflow” and “inexact”, then “underflow” is raised before “inexact”.

## F.9 Optimization

### F.9.1 General

- 1 This subclause identifies code transformations that can subvert ISO/IEC 60559-specified behavior, and others that do not.

<sup>441)</sup>Use of **float\_t** and **double\_t** variables increases the likelihood of translation-time computation. For example, the automatic initialization

```
double_t x = 1.1e75;
```

can be done at translation time, regardless of the expression evaluation method.



## F.9.2 Global transformations

- 1 Floating-point arithmetic operations and external function calls can entail side effects which optimization shall honor, at least where the state of the **FENV\_ACCESS** pragma is “on”. The flags and modes in the floating-point environment can be regarded as global variables; floating-point operations (+, \*, etc.) implicitly read the modes and write the flags.
- 2 Concern about side effects can inhibit code motion and removal of seemingly useless code. For example, in

```
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
void f(double x)
{
 /* ... */
 for (i = 0; i < n; i++) x + 1;
 /* ... */
}
```

$x+1$  may raise floating-point exceptions, so cannot be removed. And since the loop body may not execute (maybe  $0 \geq n$ ),  $x+1$  cannot be moved out of the loop. (Of course these optimizations are valid if the implementation can rule out the nettlesome cases.)

- 3 This specification does not require support for trap handlers that maintain information about the order or count of floating-point exceptions. Therefore, between function calls, the side effects due to floating-point exceptions are not required be precise: the actual order and number of occurrences of floating-point exceptions ( $> 1$ ) may vary from what the source code expresses. Thus, the preceding loop could be treated as

```
if (0 < n) x + 1;
```

## F.9.3 Expression transformations

- 1 Valid expression transformations shall preserve numerical values.
- 2 The equivalences noted in the following description apply to expressions of standard floating types.

$x/2 \leftrightarrow x \times 0.5$  Although similar transformations involving inexact constants generally do not yield equivalent expressions, if the constants are exact then such transformations can be made on ISO/IEC 60559 machines and others that round perfectly.

$1 \times x$  and  $x/1 \rightarrow x$  The expressions  $1 \times x$ ,  $x/1$ , and  $x$  may be regarded as equivalent (on ISO/IEC 60559 machines, among others).<sup>442)</sup>

$x/x \rightarrow 1.0$  The expressions  $x/x$  and 1.0 are not equivalent if  $x$  can be zero, infinite, or NaN.

$x - y \leftrightarrow x + (-y)$  The expressions  $x - y$ ,  $x + (-y)$ , and  $(-y) + x$  are equivalent (on ISO/IEC 60559 machines, among others).

$x - y \leftrightarrow -(y - x)$  The expressions  $x - y$  and  $-(y - x)$  are not equivalent because  $1 - 1$  is  $+0$  but  $-(1 - 1)$  is  $-0$  (in the default rounding direction).<sup>443)</sup>

$x - x \rightarrow 0.0$  The expressions  $x - x$  and 0.0 are not equivalent if  $x$  is a NaN or infinite.

<sup>442)</sup>Implementations can have non-required features that invalidate these and other transformations that remove arithmetic operators. Examples include strict support for signaling NaNs (an optional feature) and alternate exception handling (not included in this specification).

<sup>443)</sup>ISO/IEC 60559 prescribes a signed zero to preserve mathematical identities across certain discontinuities. Examples include:

$1/(1/\pm\infty)$  is  $\pm\infty$

and

$\text{conj}(\text{csqrt}(z))$  is  $\text{csqrt}(\text{conj}(z))$ ,

for complex  $z$ .

|                              |                                                                                                                                                                                                                                                                                                           |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $0 \times x \rightarrow 0.0$ | The expressions $0 \times x$ and $0.0$ are not equivalent if $x$ is a NaN, infinite, or $-0$ .                                                                                                                                                                                                            |
| $x + 0 \rightarrow x$        | The expressions $x + 0$ and $x$ are not equivalent if $x$ is $-0$ , because $(-0) + (+0)$ yields $+0$ (in the default rounding direction), not $-0$ .                                                                                                                                                     |
| $x - 0 \rightarrow x$        | $(+0) - (+0)$ yields $-0$ when rounding is downward (toward $-\infty$ ), but $+0$ otherwise, and $(-0) - (+0)$ always yields $-0$ ; so, if the state of the <b>FENV_ACCESS</b> pragma is “off”, promising default rounding, then the implementation can replace $x - 0$ by $x$ , even if $x$ may be zero. |
| $-x \leftrightarrow 0 - x$   | The expressions $-x$ and $0 - x$ are not equivalent if $x$ is $+0$ , because $-(+0)$ yields $-0$ , but $0 - (+0)$ yields $+0$ (unless rounding is downward).                                                                                                                                              |

- 3 For expressions of decimal floating types, transformations shall preserve quantum exponents, as well as numerical values (5.3.5.3.4).
- 4 EXAMPLE The computation  $1. \times x \rightarrow x$  is valid for decimal floating-point expressions  $x$ , but  $1.0 \times x \rightarrow x$  is not:

$$\begin{array}{llll} 1. \times 12.34 & = & (+1, 1, 0) \times (+1, 1234, -2) & \text{yields } (+1, 1234, -2) = 12.34 \\ 1.0 \times 12.34 & = & (+1, 10, -1) \times (+1, 1234, -2) & \text{yields } (+1, 12340, -3) = 12.340 \end{array}$$

In the second case, the factor 12.34 and the result 12.340 have different quantum exponents, demonstrating that  $1.0 \times x$  and  $x$  are not equivalent expressions.

## F.9.4 Relational operators

- 1  $x \neq x \rightarrow \text{false}$  The expression  $x \neq x$  is true if  $x$  is a NaN.
- $x = x \rightarrow \text{true}$  The expression  $x = x$  is false if  $x$  is a NaN.
- $x < y \rightarrow \text{isless}(x, y)$  (and similarly for  $\leq$ ,  $>$ ,  $\geq$ ) Though equal, these expressions are not equivalent because of side effects when  $x$  or  $y$  is a NaN and the state of the **FENV\_ACCESS** pragma is “on”. This transformation, which would be desirable if extra code were required to cause the “invalid” floating-point exception for unordered cases, could be performed provided the state of the **FENV\_ACCESS** pragma is “off”.

The sense of relational operators shall be maintained. This includes handling unordered cases as expressed by the source code.

- 2 EXAMPLE

```
// calls g and raises “invalid” if a and b are unordered
if (a < b)
 f();
else
 g();
```

is not equivalent to

```
// calls f and raises “invalid” if a and b are unordered
if (a >= b)
 g();
else
 f();
```

nor to

```
// calls f without raising “invalid” if a and b are unordered
if (isgreaterqual(a,b))
 g();
else
 f();
```

nor, unless the state of the **FENV\_ACCESS** pragma is “off”, to

```
// calls g without raising “invalid” if a and b are unordered
if (isless(a,b))
 f();
else
 g();
```

but is equivalent to

```
if (!(a < b))
 g();
else
 f();
```

## F.9.5 Constant arithmetic

- 1 The implementation shall honor floating-point exceptions raised by execution-time constant arithmetic wherever the state of the **FENV\_ACCESS** pragma is “on”. (See F.8.5 and F.8.6.) An operation on literals that raises no floating-point exception can be folded during translation, except, if the state of the **FENV\_ACCESS** pragma is “on”, a further check is required to assure that changing the rounding direction to downward does not alter the sign of the result,<sup>444</sup> and implementations that support dynamic rounding precision modes shall assure further that the result of the operation raises no floating-point exception when converted to the semantic type of the operation.

## F.10 Mathematics <math.h> and <tgmath.h>

### F.10.1 General

- 1 This subclause contains specifications of <math.h> and <tgmath.h> facilities that are particularly suited for ISO/IEC 60559 implementations.
- 2 The Standard C macro **HUGE\_VAL** and its **float** and **long double** analogs, **HUGE\_VALF** and **HUGE\_VALL**, expand to expressions whose values are positive infinities.
- 3 For each single-argument function **f** in <math.h> whose mathematical counterpart is symmetric (even), **f(-x)** is **f(x)** for all rounding modes and for all **x** in the (valid) domain of the function. For each single-argument function **f** in <math.h> whose mathematical counterpart is antisymmetric (odd), **f(-x)** is **-f(x)** for the ISO/IEC 60559 rounding modes **roundTiesToEven**, **roundTiesToAway**, and **roundTowardZero**, and for all **x** in the (valid) domain of the function. The **atan2** and **atan2pi** functions are odd in their first argument.
- 4 Special cases for functions in <math.h> are covered directly or indirectly by ISO/IEC 60559. The functions that ISO/IEC 60559 specifies directly are identified in F.3. The other functions in <math.h> treat infinities, NaNs, signed zeros, subnormals, and (provided the state of the **FENV\_ACCESS** pragma is “on”) the floating-point status flags in a manner consistent with ISO/IEC 60559 operations.
- 5 The expression **math\_errhandling** & **MATH\_ERREXCEPT** shall evaluate to a nonzero value.
- 6 The functions bound to operations in ISO/IEC 60559 (F.3) are fully specified by ISO/IEC 60559, including rounding behaviors and floating-point exceptions.
- 7 The “invalid” and “divide-by-zero” floating-point exceptions are raised as specified in subsequent subclauses of this annex.
- 8 The “overflow” floating-point exception is raised whenever an infinity — or, because of rounding direction, a maximal-magnitude finite number — is returned in lieu of a finite value whose magnitude is too large.
- 9 The “underflow” floating-point exception is raised whenever a computed result is tiny<sup>445</sup> and the returned result is inexact.

<sup>444</sup>) **0-0** yields **-0** instead of **+0** just when the rounding direction is downward.

<sup>445</sup>) Tiny generally indicates having a magnitude in the subnormal range. See ISO/IEC 60559 for details about detecting tininess.

- 10 Whether or when library functions not listed in Table F.2 raise the “inexact” floating-point exception is unspecified, unless stated otherwise.
- 11 Whether or when library functions not listed in Table F.2 raise a spurious “underflow” floating-point exception is not specified by this annex.<sup>446)</sup>
- 12 As implied by F.8.7, library functions do not raise spurious “invalid”, “overflow”, or “divide-by-zero” floating-point exceptions (detectable by the user).
- 13 Whether the functions not listed in Table F.2 honor the rounding direction mode is implementation-defined, unless explicitly specified otherwise.
- 14 Functions with a NaN argument return a NaN result and raise no floating-point exception, except where explicitly stated otherwise.
- 15 The specifications in the following subclauses append to the definitions in `<math.h>`. For families of functions, the specifications apply to all the functions even though only the principal function is shown. Unless otherwise specified, where the symbol “ $\pm$ ” occurs in both an argument and the result, the result has the same sign as the argument.

### Recommended practice

- 16 ISO/IEC 60559 specifies correct rounding for the operations in Table F.2 recommended by ISO/IEC 60559, and thereby preserves useful mathematical properties such as symmetry, monotonicity, and periodicity. The corresponding functions with (potentially) reserved **cr\_**-prefixed names (7.34.9) do the same. The C functions in the table, however, are not required to be correctly rounded, but implementations should still preserve as many of these useful mathematical properties as possible.
- 17 If a function with one or more NaN arguments returns a NaN result, the result should be the same as one of the NaN arguments (after possible type conversion), except perhaps for the sign.

## F.10.2 Trigonometric functions

### F.10.2.1 The **acos** functions

- 1 — **acos**(1) returns +0.
- **acos**( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $|x| > 1$ .

### F.10.2.2 The **asin** functions

- 1 — **asin**( $\pm 0$ ) returns  $\pm 0$ .
- **asin**( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $|x| > 1$ .

### F.10.2.3 The **atan** functions

- 1 — **atan**( $\pm 0$ ) returns  $\pm 0$ .
- **atan**( $\pm \infty$ ) returns  $\pm \frac{\pi}{2}$ .

### F.10.2.4 The **atan2** functions

- 1 — **atan2**( $\pm 0, -0$ ) returns  $\pm \pi$ .<sup>447)</sup>
- **atan2**( $\pm 0, +0$ ) returns  $\pm 0$ .
- **atan2**( $\pm 0, x$ ) returns  $\pm \pi$  for  $x < 0$ .
- **atan2**( $\pm 0, x$ ) returns  $\pm 0$  for  $x > 0$ .

<sup>446)</sup>It is intended that spurious “underflow” and “inexact” floating-point exceptions are raised only if avoiding them would be too costly. 7.12.2 specifies that if **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero, then an “underflow” floating-point exception shall not be raised unless an underflow range error occurs.

<sup>447)</sup>**atan2**(0, 0) does not raise the “invalid” floating-point exception, nor does **atan2**( $y, 0$ ) raise the “divide-by-zero” floating-point exception.

- **atan2**( $y, \pm 0$ ) returns  $-\frac{\pi}{2}$  for  $y < 0$ .
- **atan2**( $y, \pm 0$ ) returns  $\frac{\pi}{2}$  for  $y > 0$ .
- **atan2**( $\pm y, -\infty$ ) returns  $\pm\pi$  for finite  $y > 0$ .
- **atan2**( $\pm y, +\infty$ ) returns  $\pm 0$  for finite  $y > 0$ .
- **atan2**( $\pm\infty, x$ ) returns  $\pm\frac{\pi}{2}$  for finite  $x$ .
- **atan2**( $\pm\infty, -\infty$ ) returns  $\pm\frac{3\pi}{4}$ .
- **atan2**( $\pm\infty, +\infty$ ) returns  $\pm\frac{\pi}{4}$ .

#### F.10.2.5 The **cos** functions

- 1 — **cos**( $\pm 0$ ) returns 1.
- **cos**( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

#### F.10.2.6 The **sin** functions

- 1 — **sin**( $\pm 0$ ) returns  $\pm 0$ .
- **sin**( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

#### F.10.2.7 The **tan** functions

- 1 — **tan**( $\pm 0$ ) returns  $\pm 0$ .
- **tan**( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

#### F.10.2.8 The **acospi** functions

- 1 — **acospi**(+1) returns +0.
- **acospi**( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $|x| > 1$ .

#### F.10.2.9 The **asinpi** functions

- 1 — **asinpi**( $\pm 0$ ) returns  $\pm 0$ .
- **asinpi**( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $|x| > 1$ .

#### F.10.2.10 The **atanpi** functions

- 1 — **atanpi**( $\pm 0$ ) returns  $\pm 0$ .
- **atanpi**( $\pm\infty$ ) returns  $\pm\frac{1}{2}$ .

#### F.10.2.11 The **atan2pi** functions

- 1 — **atan2pi**( $\pm 0, -0$ ) returns  $\pm 1$ .<sup>448)</sup>
- **atan2pi**( $\pm 0, +0$ ) returns  $\pm 0$ .
- **atan2pi**( $\pm 0, x$ ) returns  $\pm 1$  for  $x < 0$ .
- **atan2pi**( $\pm 0, x$ ) returns  $\pm 0$  for  $x > 0$ .
- **atan2pi**( $y, \pm 0$ ) returns  $-\frac{1}{2}$  for  $y < 0$ .
- **atan2pi**( $y, \pm 0$ ) returns  $+\frac{1}{2}$  for  $y > 0$ .
- **atan2pi**( $\pm y, -\infty$ ) returns  $\pm 1$  for finite  $y > 0$ .
- **atan2pi**( $\pm y, +\infty$ ) returns  $\pm 0$  for finite  $y > 0$ .
- **atan2pi**( $\pm\infty, x$ ) returns  $\pm\frac{1}{2}$  for finite  $x$ .
- **atan2pi**( $\pm\infty, -\infty$ ) returns  $\pm\frac{3}{4}$ .
- **atan2pi**( $\pm\infty, +\infty$ ) returns  $\pm\frac{1}{4}$ .

<sup>448)</sup> **atan2pi**(0, 0) does not raise the “invalid” floating-point exception, nor does **atan2pi**( $y, 0$ ) raise the “divide-by-zero” floating-point exception.

**F.10.2.12 The `cospi` functions**

- 1 — `cospi`( $\pm 0$ ) returns 1.
- `cospi`( $n + \frac{1}{2}$ ) returns  $+0$ , for integers  $n$ .
- `cospi`( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

**F.10.2.13 The `sinpi` functions**

- 1 — `sinpi`( $\pm 0$ ) returns  $\pm 0$ .
- `sinpi`( $\pm n$ ) returns  $\pm 0$ , for positive integers  $n$ .
- `sinpi`( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

**F.10.2.14 The `tanpi` functions**

- 1 — `tanpi`( $\pm 0$ ) returns  $\pm 0$ .
- `tanpi`( $n$ ) returns  $+0$ , for positive even and negative odd integers  $n$ .
- `tanpi`( $n$ ) returns  $-0$ , for positive odd and negative even integers  $n$ .
- `tanpi`( $n + \frac{1}{2}$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception, for even integers  $n$ .
- `tanpi`( $n + \frac{1}{2}$ ) returns  $-\infty$  and raises the “divide-by-zero” floating-point exception, for odd integers  $n$ .
- `tanpi`( $\pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.

**F.10.3 Hyperbolic functions****F.10.3.1 The `acosh` functions**

- 1 — `acosh`(1) returns  $+0$ .
- `acosh`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < 1$ .
- `acosh`( $+\infty$ ) returns  $+\infty$ .

**F.10.3.2 The `asinh` functions**

- 1 — `asinh`( $\pm 0$ ) returns  $\pm 0$ .
- `asinh`( $\pm\infty$ ) returns  $\pm\infty$ .

**F.10.3.3 The `atanh` functions**

- 1 — `atanh`( $\pm 0$ ) returns  $\pm 0$ .
- `atanh`( $\pm 1$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception.
- `atanh`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $|x| > 1$ .

**F.10.3.4 The `cosh` functions**

- 1 — `cosh`( $\pm 0$ ) returns 1.
- `cosh`( $\pm\infty$ ) returns  $+\infty$ .

**F.10.3.5 The `sinh` functions**

- 1 — `sinh`( $\pm 0$ ) returns  $\pm 0$ .
- `sinh`( $\pm\infty$ ) returns  $\pm\infty$ .

**F.10.3.6 The `tanh` functions**

- 1 — `tanh`( $\pm 0$ ) returns  $\pm 0$ .
- `tanh`( $\pm \infty$ ) returns  $\pm 1$ .

**F.10.4 Exponential and logarithmic functions****F.10.4.1 The `exp` functions**

- 1 — `exp`( $\pm 0$ ) returns 1.
- `exp`( $-\infty$ ) returns  $+0$ .
- `exp`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.2 The `exp10` functions**

- 1 — `exp10`( $\pm 0$ ) returns 1.
- `exp10`( $-\infty$ ) returns  $+0$ .
- `exp10`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.3 The `exp10m1` functions**

- 1 — `exp10m1`( $\pm 0$ ) returns  $\pm 0$ .
- `exp10m1`( $-\infty$ ) returns  $-1$ .
- `exp10m1`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.4 The `exp2` functions**

- 1 — `exp2`( $\pm 0$ ) returns 1.
- `exp2`( $-\infty$ ) returns  $+0$ .
- `exp2`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.5 The `exp2m1` functions**

- 1 — `exp2m1`( $\pm 0$ ) returns  $\pm 0$ .
- `exp2m1`( $-\infty$ ) returns  $-1$ .
- `exp2m1`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.6 The `expm1` functions**

- 1 — `expm1`( $\pm 0$ ) returns  $\pm 0$ .
- `expm1`( $-\infty$ ) returns  $-1$ .
- `expm1`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.7 The `frex` functions**

- 1 — `frex`( $\pm 0, p$ ) returns  $\pm 0$ , and stores 0 in the object pointed to by  $p$ .
- `frex`( $\pm \infty, p$ ) returns  $\pm \infty$ , and stores an unspecified value in the object pointed to by  $p$ .
- `frex`(NaN,  $p$ ) stores an unspecified value in the object pointed to by  $p$  (and returns a NaN).
- 2 `frex` raises no floating-point exceptions if `value` is not a signaling NaN.
- 3 The returned value is independent of the current rounding direction mode.
- 4 On a binary system, the body of the `frex` function may be

```

{
 *p = (value == 0 || !isfinite(value)) ? 0: (int)(1 + logb(value));
 return scalbn(value, -(*p));
}

```

**F.10.4.8 The `ilogb` functions**

- 1 When the correct result is representable in the range of the return type, the returned value is exact and is independent of the current rounding direction mode.
- 2 If the correct result is outside the range of the return type, the numeric result is unspecified and the “invalid” floating-point exception is raised.
- 3 `ilogb(x)`, for  $x$  zero, infinite, or NaN, raises the “invalid” floating-point exception and returns the value specified in 7.12.7.8.

**F.10.4.9 The `ldexp` functions**

- 1 On a binary system, `ldexp(x, exp)` is equivalent to `scalbn(x, exp)`.

**F.10.4.10 The `llogb` functions**

- 1 The `llogb` functions are equivalent to the `ilogb` functions, except that the `llogb` functions determine a result in the `long int` type.

**F.10.4.11 The `log` functions**

- 1 — `log(±0)` returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log(1)` returns  $+0$ .
- `log(x)` returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- `log(+∞)` returns  $+\infty$ .

**F.10.4.12 The `log10` functions**

- 1 — `log10(±0)` returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log10(1)` returns  $+0$ .
- `log10(x)` returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- `log10(+∞)` returns  $+\infty$ .

**F.10.4.13 The `log10p1` functions**

- 1 — `log10p1(±0)` returns  $\pm 0$ .
- `log10p1(−1)` returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log10p1(x)` returns a NaN and raises the “invalid” floating-point exception for  $x < -1$ .
- `log10p1(+∞)` returns  $+\infty$ .

**F.10.4.14 The `log1p` and `logp1` functions**

- 1 — `log1p(±0)` returns  $\pm 0$ .
- `log1p(−1)` returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log1p(x)` returns a NaN and raises the “invalid” floating-point exception for  $x < -1$ .
- `log1p(+∞)` returns  $+\infty$ .

The `log1p` functions are equivalent to the `logp1` functions.

**F.10.4.15 The `log2` functions**

- 1 — `log2(±0)` returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log2(1)` returns  $+0$ .
- `log2(x)` returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- `log2(+∞)` returns  $+\infty$ .



**F.10.4.16 The `log2p1` functions**

- 1 — `log2p1`( $\pm 0$ ) returns  $\pm 0$ .
- `log2p1`( $-1$ ) returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `log2p1`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < -1$ .
- `log2p1`( $+\infty$ ) returns  $+\infty$ .

**F.10.4.17 The `logb` functions**

- 1 — `logb`( $\pm 0$ ) returns  $-\infty$  and raises the “divide-by-zero” floating-point exception.
- `logb`( $\pm\infty$ ) returns  $+\infty$ .
- 2 The returned value is exact and is independent of the current rounding direction mode.

**F.10.4.18 The `modf` functions**

- 1 — `modf`( $\pm x$ , *iptr*) returns a result with the same sign as  $x$ .
- `modf`( $\pm\infty$ , *iptr*) returns  $\pm 0$  and stores  $\pm\infty$  in the object pointed to by *iptr*.
- `modf`(NaN, *iptr*) stores a NaN in the object pointed to by *iptr* (and returns a NaN).
- 2 The returned values are exact and are independent of the current rounding direction mode.
- 3 `modf` behaves as though implemented by

```
#include <math.h>
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
double modf(double value, double *iptr)
{
 int save_round = fegetround();
 fesetround(FE_TOWARDZERO);
 *iptr = nearbyint(value);
 fesetround(save_round);
 return copysign(
 isinf(value) ? 0.0:
 value - (*iptr), value);
}
```

**F.10.4.19 The `scalbn` and `scalbln` functions**

- 1 — `scalbn`( $\pm 0$ ,  $n$ ) returns  $\pm 0$ .
- `scalbn`( $x$ ,  $0$ ) returns  $x$ .
- `scalbn`( $\pm\infty$ ,  $n$ ) returns  $\pm\infty$ .
- 2 If the calculation does not overflow or underflow, the returned value is exact and independent of the current rounding direction mode.

**F.10.5 Power and absolute value functions****F.10.5.1 The `cbrt` functions**

- 1 — `cbrt`( $\pm 0$ ) returns  $\pm 0$ .
- `cbrt`( $\pm\infty$ ) returns  $\pm\infty$ .

**F.10.5.2 The compoundn functions**

- 1 — **compoundn**( $x, 0$ ) returns 1 for  $x \geq -1$  or  $x$  a NaN.
- **compoundn**( $x, n$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < -1$ .
- **compoundn**( $-1, n$ ) returns  $+\infty$  and raises the divide-by-zero floating-point exception for  $n < 0$ .
- **compoundn**( $-1, n$ ) returns  $+0$  for  $n > 0$ .
- **compoundn**( $+\infty, n$ ) returns  $+\infty$  for  $n > 0$ .
- **compoundn**( $+\infty, n$ ) returns  $+0$  for  $n < 0$ .

**F.10.5.3 The fabs functions**

- 1 — **fabs**( $\pm 0$ ) returns  $+0$ .
- **fabs**( $\pm\infty$ ) returns  $+\infty$ .
- 2 **fabs**( $x$ ) raises no floating-point exceptions, even if  $x$  is a signaling NaN. The returned value is independent of the current rounding direction mode.

**F.10.5.4 The hypot functions**

- 1 — **hypot**( $x, y$ ), **hypot**( $y, x$ ), and **hypot**( $x, -y$ ) are equivalent.
- **hypot**( $x, \pm 0$ ) returns the absolute value of  $x$ , if  $x$  is not a NaN.
- **hypot**( $\pm\infty, y$ ) returns  $+\infty$ , even if  $y$  is a NaN.
- **hypot**( $x, \text{NaN}$ ) returns a NaN, if  $x$  is not  $\pm\infty$ .

**F.10.5.5 The pow functions**

- 1 — **pow**( $\pm 0, y$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception for  $y$  an odd integer  $< 0$ .
- **pow**( $\pm 0, y$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception for  $y < 0$ , finite, and not an odd integer.
- **pow**( $\pm 0, -\infty$ ) returns  $+\infty$ .
- **pow**( $\pm 0, y$ ) returns  $\pm 0$  for  $y$  an odd integer  $> 0$ .
- **pow**( $\pm 0, y$ ) returns  $+0$  for  $y > 0$  and not an odd integer.
- **pow**( $-1, \pm\infty$ ) returns 1.
- **pow**( $+1, y$ ) returns 1 for any  $y$ , even a NaN.
- **pow**( $x, \pm 0$ ) returns 1 for any  $x$ , even a NaN.
- **pow**( $x, y$ ) returns a NaN and raises the “invalid” floating-point exception for finite  $x < 0$  and finite non-integer  $y$ .
- **pow**( $x, -\infty$ ) returns  $+\infty$  for  $|x| < 1$ .
- **pow**( $x, -\infty$ ) returns  $+0$  for  $|x| > 1$ .
- **pow**( $x, +\infty$ ) returns  $+0$  for  $|x| < 1$ .
- **pow**( $x, +\infty$ ) returns  $+\infty$  for  $|x| > 1$ .
- **pow**( $-\infty, y$ ) returns  $-0$  for  $y$  an odd integer  $< 0$ .
- **pow**( $-\infty, y$ ) returns  $+0$  for  $y < 0$  and not an odd integer.
- **pow**( $-\infty, y$ ) returns  $-\infty$  for  $y$  an odd integer  $> 0$ .
- **pow**( $-\infty, y$ ) returns  $+\infty$  for  $y > 0$  and not an odd integer.
- **pow**( $+\infty, y$ ) returns  $+0$  for  $y < 0$ .
- **pow**( $+\infty, y$ ) returns  $+\infty$  for  $y > 0$ .

**F.10.5.6 The `pown` functions**

- 1 — `pown`( $x, 0$ ) returns 1 for all  $x$  not a signaling NaN.
- `pown`( $\pm 0, n$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception for odd  $n < 0$ .
- `pown`( $\pm 0, n$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception for even  $n < 0$ .
- `pown`( $\pm 0, n$ ) returns  $+0$  for even  $n > 0$ .
- `pown`( $\pm 0, n$ ) returns  $\pm 0$  for odd  $n > 0$ .
- `pown`( $\pm\infty, n$ ) is equivalent to `pown`( $\pm 0, -n$ ) for  $n$  not 0, except that the “divide-by-zero” floating-point exception is not raised.

**F.10.5.7 The `powr` functions**

- 1 — `powr`( $x, \pm 0$ ) returns 1 for finite  $x > 0$ .
- `powr`( $\pm 0, y$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception for finite  $y < 0$ .
- `powr`( $\pm 0, -\infty$ ) returns  $+\infty$ .
- `powr`( $\pm 0, y$ ) returns  $+0$  for  $y > 0$ .
- `powr`( $+1, y$ ) returns 1 for finite  $y$ .
- `powr`( $+1, \pm\infty$ ) returns a NaN and raises the “invalid” floating-point exception.
- `powr`( $x, y$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- `powr`( $\pm 0, \pm 0$ ) returns a NaN and raises the “invalid” floating-point exception.
- `powr`( $+\infty, \pm 0$ ) returns a NaN and raises the “invalid” floating-point exception.

**F.10.5.8 The `rootn` functions**

- 1 — `rootn`( $\pm 0, n$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception for odd  $n < 0$ .
- `rootn`( $\pm 0, n$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception for even  $n < 0$ .
- `rootn`( $\pm 0, n$ ) returns  $+0$  for even  $n > 0$ .
- `rootn`( $\pm 0, n$ ) returns  $\pm 0$  for odd  $n > 0$ .
- `rootn`( $+\infty, n$ ) returns  $+\infty$  for  $n > 0$ .
- `rootn`( $-\infty, n$ ) returns  $-\infty$  for odd  $n > 0$ .
- `rootn`( $-\infty, n$ ) returns a NaN and raises the “invalid” floating-point exception for even  $n > 0$ .
- `rootn`( $+\infty, n$ ) returns  $+0$  for  $n < 0$ .
- `rootn`( $-\infty, n$ ) returns  $-0$  for odd  $n < 0$ .
- `rootn`( $-\infty, n$ ) returns a NaN and raises the “invalid” floating-point exception for even  $n < 0$ .
- `rootn`( $x, 0$ ) returns a NaN and raises the “invalid” floating-point exception for all  $x$  (including NaN).
- `rootn`( $x, n$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < 0$  and  $n$  even.

**F.10.5.9 The `rsqrt` functions**

- 1 — `rsqrt`( $\pm 0$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception.
- `rsqrt`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- `rsqrt`( $+\infty$ ) returns  $+0$ .

**F.10.5.10 The `sqrt` functions**

- 1 — `sqrt`( $\pm 0$ ) returns  $\pm 0$ .
  - `sqrt`( $+\infty$ ) returns  $+\infty$ .
  - `sqrt`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $x < 0$ .
- 2 The returned value is dependent on the current rounding direction mode.

**F.10.6 Error and gamma functions****F.10.6.1 The `erf` functions**

- 1 — `erf`( $\pm 0$ ) returns  $\pm 0$ .
- `erf`( $\pm\infty$ ) returns  $\pm 1$ .

**F.10.6.2 The `erfc` functions**

- 1 — `erfc`( $-\infty$ ) returns 2.
- `erfc`( $+\infty$ ) returns  $+0$ .

**F.10.6.3 The `lgamma` functions**

- 1 — `lgamma`(1) returns  $+0$ .
- `lgamma`(2) returns  $+0$ .
- `lgamma`( $x$ ) returns  $+\infty$  and raises the “divide-by-zero” floating-point exception for  $x$  a negative integer or zero.
- `lgamma`( $-\infty$ ) returns  $+\infty$ .
- `lgamma`( $+\infty$ ) returns  $+\infty$ .

**F.10.6.4 The `tgamma` functions**

- 1 — `tgamma`( $\pm 0$ ) returns  $\pm\infty$  and raises the “divide-by-zero” floating-point exception.
- `tgamma`( $x$ ) returns a NaN and raises the “invalid” floating-point exception for  $x$  a negative integer.
- `tgamma`( $-\infty$ ) returns a NaN and raises the “invalid” floating-point exception.
- `tgamma`( $+\infty$ ) returns  $+\infty$ .

**F.10.7 Nearest integer functions****F.10.7.1 The `ceil` functions**

- 1 — `ceil`( $\pm 0$ ) returns  $\pm 0$ .
  - `ceil`( $\pm\infty$ ) returns  $\pm\infty$ .
- 2 The returned value is exact and is independent of the current rounding direction mode.
- 3 The **double** version of `ceil` behaves as though implemented by

```

#include <math.h>
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
double ceil(double x)
{
 double result;
 int save_round = fegetround();
 fesetround(FE_UPWARD);
 result = nearbyint(x);
 fesetround(save_round);
 return result;
}

```

#### F.10.7.2 The **floor** functions

- 1 — **floor**( $\pm 0$ ) returns  $\pm 0$ .  
— **floor**( $\pm \infty$ ) returns  $\pm \infty$ .
- 2 The returned value is exact and is independent of the current rounding direction mode.
- 3 See the sample implementation for **ceil** in F.10.7.1.

#### F.10.7.3 The **nearbyint** functions

- 1 The **nearbyint** functions use ISO/IEC 60559 rounding according to the current rounding direction. They do not raise the “inexact” floating-point exception if the result differs in value from the argument.  
— **nearbyint**( $\pm 0$ ) returns  $\pm 0$  (for all rounding directions).  
— **nearbyint**( $\pm \infty$ ) returns  $\pm \infty$  (for all rounding directions).

#### F.10.7.4 The **rint** functions

- 1 The **rint** functions differ from the **nearbyint** functions only in that they do raise the “inexact” floating-point exception if the result differs in value from the argument.

#### F.10.7.5 The **lrint** and **llrint** functions

- 1 The **lrint** and **llrint** functions provide floating-to-integer conversion as prescribed by ISO/IEC 60559. They round according to the current rounding direction. If the rounded value is outside the range of the return type, the numeric result is unspecified and the “invalid” floating-point exception is raised. When they raise no other floating-point exception and the result differs from the argument, they raise the “inexact” floating-point exception.

#### F.10.7.6 The **round** functions

- 1 — **round**( $\pm 0$ ) returns  $\pm 0$ .  
— **round**( $\pm \infty$ ) returns  $\pm \infty$ .
- 2 The returned value is independent of the current rounding direction mode.
- 3 The **double** version of **round** behaves as though implemented by:<sup>449)</sup>

```

#include <math.h>
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
double round(double x)
{
 double result;

```

<sup>449)</sup>This code does not handle signaling NaNs as required of implementations that define **FE\_SNANS\_ALWAYS\_SIGNAL**.

```

 fenv_t save_env;
 feholdexcept(&save_env);
 result = rint(x);
 if (fetestexcept(FE_INEXACT)) {
 fesetround(FE_TOWARDZERO);
 result = rint(copysign(0.5 + fabs(x), x));
 feclearexcept(FE_INEXACT);
 }
 feupdateenv(&save_env);
 return result;
}

```

#### F.10.7.7 The **lround** and **llround** functions

- 1 The **lround** and **llround** functions differ from the **lrint** and **llrint** functions with the default rounding direction just in that the **lround** and **llround** functions round halfway cases away from zero and are not required to raise the “inexact” floating-point exception for non-integer arguments that round to within the range of the return type.

#### F.10.7.8 The **roundeven** functions

- 1
  - **roundeven**( $\pm 0$ ) returns  $\pm 0$ .
  - **roundeven**( $\pm \infty$ ) returns  $\pm \infty$ .
- 2 The returned value is exact and is independent of the current rounding direction mode.
- 3 See the sample implementation for **ceil** in F.10.7.1.

#### F.10.7.9 The **trunc** functions

- 1 The **trunc** functions use ISO/IEC 60559 rounding toward zero (regardless of the current rounding direction).
  - **trunc**( $\pm 0$ ) returns  $\pm 0$ .
  - **trunc**( $\pm \infty$ ) returns  $\pm \infty$ .
- 2 The returned value is exact and is independent of the current rounding direction mode.

#### F.10.7.10 The **fromfp** and **ufromfp** functions

- 1 The **fromfp** and **ufromfp** functions raise the “invalid” floating-point exception and return a NaN if the argument **width** is zero or if the floating-point argument **x** is infinite or NaN or rounds to an integral value that is outside the range determined by the argument **width** (see 7.12.10.10).
- 2 These functions do not raise the “inexact” floating-point exception.

#### F.10.7.11 The **fromfpx** and **ufromfpx** functions

- 1 The **fromfpx** and **ufromfpx** functions raise the “invalid” floating-point exception and return a NaN if the argument **width** is zero or if the floating-point argument **x** is infinite or NaN or rounds to an integral value that is outside the range determined by the argument **width** (see 7.12.10.11).
- 2 These functions raise the “inexact” floating-point exception if a valid result differs in value from the floating-point argument **x**.

### F.10.8 Remainder functions

#### F.10.8.1 The **fmod** functions

- 1
  - **fmod**( $\pm 0, y$ ) returns  $\pm 0$  for  $y$  not zero.
  - **fmod**( $x, y$ ) returns a NaN and raises the “invalid” floating-point exception for  $x$  infinite or  $y$  zero (and neither is a NaN).

- **fmod**( $x, \pm\infty$ ) returns  $x$  for finite  $x$ .
- 2 When subnormal results are supported, the returned value is exact and is independent of the current rounding direction mode.
- 3 The **double** version of **fmod** behaves as though implemented by

```
#include <math.h>
#include <fenv.h>
#pragma STDC FENV_ACCESS ON
double fmod(double x, double y)
{
 double result;
 result = remainder(fabs(x), (y = fabs(y)));
 if (signbit(result)) result += y;
 return copysign(result, x);
}
```

### F.10.8.2 The **remainder** functions

- 1 — **remainder**( $\pm 0, y$ ) returns  $\pm 0$  for  $y$  not zero.
- **remainder**( $x, y$ ) returns a NaN and raises the “invalid” floating-point exception for  $x$  infinite or  $y$  zero (and neither is a NaN).
- **remainder**( $x, \pm\infty$ ) returns  $x$  for finite  $x$ .

- 2 When subnormal results are supported, the returned value is exact and is independent of the current rounding direction mode.

### F.10.8.3 The **remquo** functions

- 1 The **remquo** functions follow the specifications for the **remainder** functions.
- 2 If a NaN is returned, the value stored in the object pointed to by **quo** is unspecified.
- 3 When subnormal results are supported, the returned value is exact and is independent of the current rounding direction mode.

## F.10.9 Manipulation functions

### F.10.9.1 The **copysign** functions

- 1 **copysign**( $x, y$ ) raises no floating-point exceptions, even if  $x$  or  $y$  is a signaling NaN. The returned value is independent of the current rounding direction mode.

### F.10.9.2 The **nan** functions

- 1 All ISO/IEC 60559 implementations support quiet NaNs, in all floating formats.
- 2 The returned value is exact and is independent of the current rounding direction mode.

### F.10.9.3 The **nextafter** functions

- 1 — **nextafter**( $x, y$ ) raises the “overflow” and “inexact” floating-point exceptions for  $x$  finite and the function value infinite.
- **nextafter**( $x, y$ ) raises the “underflow” and “inexact” floating-point exceptions for the function value subnormal or zero and  $x \neq y$ .
- 2 Even though underflow or overflow can occur, the returned value is independent of the current rounding direction mode.

### F.10.9.4 The **nexttoward** functions

- 1 No additional requirements beyond those on **nextafter**.
- 2 Even though underflow or overflow can occur, the returned value is independent of the current rounding direction mode.

**F.10.9.5 The `nextup` functions**

- 1 — `nextup`( $+\infty$ ) returns  $+\infty$ .
- `nextup`( $-\infty$ ) returns the largest-magnitude negative finite number in the return type of the function.
- 2 `nextup`(*x*) raises no floating-point exceptions if *x* is not a signaling NaN. The returned value is independent of the current rounding direction mode.

**F.10.9.6 The `nextdown` functions**

- 1 — `nextdown`( $-\infty$ ) returns  $-\infty$ .
- `nextdown`( $+\infty$ ) returns the largest-magnitude positive finite number in the type of the function.
- 2 `nextdown`(*x*) raises no floating-point exceptions if *x* is not a signaling NaN. The returned value is independent of the current rounding direction mode.

**F.10.9.7 The `canonicalize` functions**

- 1 The `canonicalize` functions produce<sup>450)</sup> the canonical version of the representation in the object pointed to by the argument *x*. If the input *\*x* is a signaling NaN, the “invalid” floating-point exception is raised and a (canonical) quiet NaN (which should be the canonical version of that signaling NaN made quiet) is produced. For quiet NaN, infinity, and finite inputs, the functions raise no floating-point exceptions.

**F.10.10 Maximum, minimum, and positive difference functions****F.10.10.1 The `fdim` functions**

- 1 No additional requirements.

**F.10.10.2 The `fmax` functions**

- 1 If just one argument is a NaN, the `fmax` functions return the other argument (if both arguments are NaNs, the functions return a NaN).
- 2 The returned value is exact and is independent of the current rounding direction mode.
- 3 The body of the `fmax` function may be:<sup>451)</sup>

```

{
 double r = (isgreater(x, y) || isnan(y)) ? x : y;
 (void) canonicalize(&r, &r);
 return r;
}

```

**F.10.10.3 The `fmin` functions**

- 1 The `fmin` functions are analogous to the `fmax` functions (see F.10.10.2).
- 2 The returned value is exact and is independent of the current rounding direction mode.

**F.10.10.4 The `fmaximum`, `fminimum`, `fmaximum_mag`, and `fminimum_mag` functions**

- 1 These functions treat NaNs like other functions in `<math.h>` (see F.10). They differ from the corresponding `fmaximum_num`, `fminimum_num`, `fmaximum_mag_num`, and `fminimum_mag_num` functions only in their treatment of NaNs.

<sup>450)</sup>As if `*x * 1e0` were computed.

<sup>451)</sup>If possible, `fmax` is sensitive to the sign of zero, for example `fmax(-0.0, +0.0)` ideally returns `+0`. Note also that this implementation does not handle signaling NaNs as required of implementations that define `FE_SNANS_ALWAYS_SIGNAL`.



**F.10.10.5 The `fmaximum_num`, `fminimum_num`, `fmaximum_mag_num`, and `fminimum_mag_num` functions**

- 1 These functions return the number if one argument is a number and the other is a quiet or signaling NaN. If both arguments are NaNs, a quiet NaN is returned. If an argument is a signaling NaN, the “invalid” floating-point exception is raised (even though the function returns the number when the other argument is a number).

**F.10.11 Fused multiply-add****F.10.11.1 The `fma` functions**

- 1 — `fma`( $x, y, z$ ) computes  $xy + z$ , correctly rounded once.
- `fma`( $x, y, z$ ) returns a NaN and optionally raises the “invalid” floating-point exception if one of  $x$  and  $y$  is infinite, the other is zero, and  $z$  is a NaN.
- `fma`( $x, y, z$ ) returns a NaN and raises the “invalid” floating-point exception if one of  $x$  and  $y$  is infinite, the other is zero, and  $z$  is not a NaN.
- `fma`( $x, y, z$ ) returns a NaN and raises the “invalid” floating-point exception if  $x$  times  $y$  is an exact infinity and  $z$  is also an infinity but with the opposite sign.

**F.10.12 Functions that round result to narrower type**

- 1 The functions that round their result to narrower type (7.12.15) are fully specified in ISO/IEC 60559. The returned value is dependent on the current rounding direction mode.
- 2 These functions treat zero and infinite arguments like the corresponding operation or function: `+`, `-`, `*`, `/`, `fma`, or `sqrt`.

**F.10.13 Total order functions****F.10.13.1 General**

- 1 This subclause specifies the total order functions required by ISO/IEC 60559.
- 2 NOTE These functions are specified only in this annex because the functions for standard floating types depend on details of ISO/IEC 60559 formats that are conditionally supported based on if the relevant feature test macro, `__STDC_IEC_60559_BFP__` or `__STDC_IEC_60559_DFP__`, is defined or not.

**F.10.13.2 The `totalorder` functions****Synopsis**

```

1 #define __STDC_WANT_IEC_60559_EXT__
 #include <math.h>
 #ifdef __STDC_IEC_60559_BFP__
 int totalorder(const double *x, const double *y);
 int totalorderf(const float *x, const float *y);
 int totalorderl(const long double *x, const long double *y);
 #endif
 #ifdef __STDC_IEC_60559_DFP__
 int totalorderd32(const _Decimal32 *x, const _Decimal32 *y);
 int totalorderd64(const _Decimal64 *x, const _Decimal64 *y);
 int totalorderd128(const _Decimal128 *x, const _Decimal128 *y);
 #endif

```

**Description**

- 2 The `totalorder` functions determine whether the total order relationship, defined by ISO/IEC 60559, is true for the ordered pair of  $*x, *y$ . These functions are fully specified in ISO/IEC 60559. These functions are independent of the current rounding direction mode and raise no floating-point exceptions, even if  $*x$  or  $*y$  is a signaling NaN.

**Returns**

- 3 The **totalorder** functions return nonzero if and only if the total order relation is true for the ordered pair of *\*x*, *\*y*.

**F.10.13.3 The **totalordermag** functions****Synopsis**

```

1 #define __STDC_WANT_IEC_60559_EXT__
 #include <math.h>
 #ifdef __STDC_IEC_60559_BFP__
 int totalordermag(const double *x, const double *y);
 int totalordermagf(const float *x, const float *y);
 int totalordermagl(const long double *x, const long double *y);
 #endif
 #ifdef __STDC_IEC_60559_DFP__
 int totalordermagd32(const _Decimal32 *x, const _Decimal32 *y);
 int totalordermagd64(const _Decimal64 *x, const _Decimal64 *y);
 int totalordermagd128(const _Decimal128 *x, const _Decimal128 *y);
 #endif

```

**Description**

- 2 The **totalordermag** functions determine whether the total order relationship, defined by ISO/IEC 60559, is true for the ordered pair of the magnitudes of *\*x*, *\*y*. These functions are fully specified in ISO/IEC 60559. These functions are independent of the current rounding direction mode and raise no floating-point exceptions, even if *\*x* or *\*y* is a signaling NaN.

**Returns**

- 3 The **totalordermag** functions return nonzero if and only if the total order relation is true for the ordered pair of the magnitudes of *\*x*, *\*y*.

**F.10.14 Payload functions****F.10.14.1 General**

- 1 ISO/IEC 60559 defines the payload to be information contained in a quiet or signaling NaN. The payload is intended for implementation-defined diagnostic information about the NaN, such as where or how the NaN was created.<sup>452)</sup> The implementation interprets the payload as a nonnegative integer suitable for use with the functions in this subclause, which get and set payloads. The implementation may restrict which payloads are admissible for the user to set.
- 2 NOTE These functions are specified only in this annex because the functions for standard floating types depend on details of ISO/IEC 60559 formats that are conditionally supported based on if the relevant feature test macro, **\_\_STDC\_IEC\_60559\_BFP\_\_** or **\_\_STDC\_IEC\_60559\_DFP\_\_**, is defined or not.

**F.10.14.2 The **getpayload** functions****Synopsis**

```

1 #define __STDC_WANT_IEC_60559_EXT__
 #include <math.h>
 #ifdef __STDC_IEC_60559_BFP__
 double getpayload(const double *x);
 float getpayloadf(const float *x);
 long double getpayloadl(const long double *x);
 #endif
 #ifdef __STDC_IEC_60559_DFP__
 _Decimal32 getpayloadd32(const _Decimal32 *x);
 _Decimal64 getpayloadd64(const _Decimal64 *x);
 _Decimal128 getpayloadd128(const _Decimal128 *x);
 #endif

```

<sup>452)</sup> For the purpose of determining value inclusion (as in 6.2.5, 7.12, and H.11), quiet NaN representations can be regarded as having the same value, regardless of payloads.

**Description**

- The **getpayload** functions extract the payload of a quiet or signaling NaN input and return it as a positive-signed floating-point integer. If *\*x* is not a NaN, the return result is  $-1$ . These functions raise no floating-point exceptions, even if *\*x* is a signaling NaN.

**Returns**

- The **getpayload** functions return the payload of the NaN input as a positive-signed floating-point integer.

**F.10.14.3 The **setpayload** functions****Synopsis**

```

1 #define __STDC_WANT_IEC_60559_EXT__
 #include <math.h>
 #ifdef __STDC_IEC_60559_BFP__
 int setpayload(double *res, double pl);
 int setpayloadf(float *res, float pl);
 int setpayloadl(long double *res, long double pl);
 #endif
 #ifdef __STDC_IEC_60559_DFP__
 int setpayloadd32(_Decimal32 *res, _Decimal32 pl);
 int setpayloadd64(_Decimal64 *res, _Decimal64 pl);
 int setpayloadd128(_Decimal128 *res, _Decimal128 pl);
 #endif

```

**Description**

- The **setpayload** functions create a quiet NaN with the payload specified by *pl* and a zero sign bit and store that NaN in the object pointed to by *\*res*. If *pl* is not a floating-point integer representing an admissible payload, *\*res* is set to  $+0$ .

**Returns**

- If the **setpayload** functions stored the specified NaN, they return a zero value, otherwise a nonzero value (and *\*res* is set to  $+0$ ).

**F.10.14.4 The **setpayloadsig** functions****Synopsis**

```

1 #define __STDC_WANT_IEC_60559_EXT__
 #include <math.h>
 #ifdef __STDC_IEC_60559_BFP__
 int setpayloadsig(double *res, double pl);
 int setpayloadsigf(float *res, float pl);
 int setpayloadsigl(long double *res, long double pl);
 #endif
 #ifdef __STDC_IEC_60559_DFP__
 int setpayloadsigd32(_Decimal32 *res, _Decimal32 pl);
 int setpayloadsigd64(_Decimal64 *res, _Decimal64 pl);
 int setpayloadsigd128(_Decimal128 *res, _Decimal128 pl);
 #endif

```

**Description**

- The **setpayloadsig** functions create a signaling NaN with the payload specified by *pl* and a zero sign bit and store that NaN in the object pointed to by *\*res*. If *pl* is not a floating-point integer representing an admissible payload, *\*res* is set to  $+0$ .

**Returns**

- If the **setpayloadsig** functions stored the specified NaN, they return a zero value, otherwise a nonzero value (and *\*res* is set to  $+0$ ).

## F.10.15 Comparison macros

### F.10.15.1 General

- 1 Relational operators and their corresponding comparison macros (7.12.18) produce equivalent result values, even if argument values are represented in wider formats. Thus, comparison macro arguments represented in formats wider than their semantic types are not converted to the semantic types, unless the wide evaluation method converts operands of relational operators to their semantic types. The standard wide evaluation methods characterized by **FLT\_EVAL\_METHOD** and **DEC\_EVAL\_METHOD** equal to 1 or 2 (5.3.5.3.3, 5.3.5.3.4), do not convert operands of relational operators to their semantic types.

### F.10.15.2 The **iseqsig** macro

- 1 The equality operator `==` and the **iseqsig** macro produce equivalent results, except that the **iseqsig** macro raises the “invalid” floating-point exception if an argument is a NaN.

## Annex G

(normative)

### ISO/IEC 60559-compatible complex arithmetic

#### G.1 Introduction

- 1 This annex supplements Annex F to specify complex arithmetic for compatibility with ISO/IEC 60559 real floating-point arithmetic. An implementation that defines **\_\_STDC\_IEC\_60559\_COMPLEX\_\_** shall conform to the specifications in this annex.<sup>453)</sup>

#### G.2 Conventions

- 1 A complex value with at least one infinite part is regarded as an *infinity* (even if its other part is a quiet NaN). A complex value is a *finite number* if each of its parts is a finite number (neither infinite nor NaN). A complex value is a *zero* if each of its parts is a zero.

#### G.3 Binary operators

##### G.3.1 General

- 1 For most operand types, the value of the result of a binary operator with a complex operand is completely determined, with reference to real arithmetic, by the usual mathematical formula. For some operand types, the usual mathematical formula is problematic because of its treatment of infinities and because of undue overflow or underflow; in these cases the result satisfies certain properties (specified in G.3.2), but is not completely determined.

##### G.3.2 Multiplicative operators

###### Semantics

- 1 If the operands are not both complex, then the result and floating-point exception behavior of the `*` operator is defined by the usual mathematical formula shown in Table G.1:

**Table G.1 — Results of multiplication operations**

| <code>*</code> | $u$            | $u + iv$       |
|----------------|----------------|----------------|
| $x$            | $xu$           | $(xu) + i(xv)$ |
| $x + iy$       | $(xu) + i(yu)$ |                |

- 2 If the second operand is not complex, then the result and floating-point exception behavior of the `/` operator is defined by the usual mathematical formula as in Table G.2:

**Table G.2 — Results of division operations**

| <code>/</code> | $u$              |
|----------------|------------------|
| $x$            | $x/u$            |
| $x + iy$       | $(x/u) + i(y/u)$ |

- 3 The `*` and `/` operators satisfy the following infinity properties for all real and complex operands:<sup>454)</sup>
- if one operand is an infinity and the other operand is a nonzero finite number or an infinity, then the result of the `*` operator is an infinity;
  - if the first operand is an infinity and the second operand is a finite number, then the result of the `/` operator is an infinity;
  - if the first operand is a finite number and the second operand is an infinity, then the result of the `/` operator is a zero;

<sup>453)</sup>Implementations that do not define **\_\_STDC\_IEC\_60559\_COMPLEX\_\_** are not required to conform to these specifications.

<sup>454)</sup>These properties are already implied for those cases covered in Table G.1 and Table G.2, but are required for all cases (at least where the state for **CX\_LIMITED\_RANGE** is “off”).

— if the first operand is a nonzero finite number or an infinity and the second operand is a zero, then the result of the / operator is an infinity.

- 4 If both operands of the \* operator are complex or if the second operand of the / operator is complex, the operator raises floating-point exceptions if appropriate for the calculation of the parts of the result, and may raise spurious floating-point exceptions.
- 5 EXAMPLE 1 Multiplication of **double \_Complex** operands could be implemented as follows.

```
#include <math.h>
#include <complex.h>

/* Multiply z * w ... */
double complex _Cm ultd(double complex z, double complex w)
{
 #pragma STDC FP_CONTRACT OFF
 double a, b, c, d, ac, bd, ad, bc, x, y;
 a = creal(z); b = cimag(z);
 c = creal(w); d = cimag(w);
 ac = a * c; bd = b * d;
 ad = a * d; bc = b * c;
 x = ac - bd; y = ad + bc;
 if (isnan(x) && isnan(y)) {
 /* Recover infinities that computed as NaN+iNaN ... */
 int recalc = 0;
 if (isinf(a) || isinf(b)) { // z is infinite
 /* "Box" the infinity and change NaNs in the other factor to 0 */
 a = copysign(isinf(a) ? 1.0 : 0.0, a);
 b = copysign(isinf(b) ? 1.0 : 0.0, b);
 if (isnan(c)) c = copysign(0.0, c);
 if (isnan(d)) d = copysign(0.0, d);
 recalc = 1;
 }
 if (isinf(c) || isinf(d)) { // w is infinite
 /* "Box" the infinity and change NaNs in the other factor to 0 */
 c = copysign(isinf(c) ? 1.0 : 0.0, c);
 d = copysign(isinf(d) ? 1.0 : 0.0, d);
 if (isnan(a)) a = copysign(0.0, a);
 if (isnan(b)) b = copysign(0.0, b);
 recalc = 1;
 }
 if (!recalc && (isinf(ac) || isinf(bd) ||
 isinf(ad) || isinf(bc))) {
 /* Recover infinities from overflow by changing NaNs to 0 ... */
 if (isnan(a)) a = copysign(0.0, a);
 if (isnan(b)) b = copysign(0.0, b);
 if (isnan(c)) c = copysign(0.0, c);
 if (isnan(d)) d = copysign(0.0, d);
 recalc = 1;
 }
 if (recalc) {
 x = INFINITY * (a * c - b * d);
 y = INFINITY * (a * d + b * c);
 }
 }
 return CMLX(x, y);
}
```

- 6 This implementation achieves the required treatment of infinities at the cost of only one **isnan** test in ordinary (finite) cases. It is less than ideal in that undue overflow and underflow can occur.
- 7 EXAMPLE 2 Division of two **double \_Complex** operands can be implemented as follows.

```

#include <math.h>
#include <complex.h>

/* Divide z / w ... */
double complex _Cdivd(double complex z, double complex w)
{
 #pragma STDC FP_CONTRACT OFF
 double a, b, c, d, logbw, denom, x, y;
 int ilogbw = 0;
 a = creal(z); b = cimag(z);
 c = creal(w); d = cimag(w);
 logbw = logb(fmaximum_num(fabs(c), fabs(d)));
 if (isfinite(logbw)) {
 ilogbw = (int)logbw;
 c = scalbn(c, -ilogbw); d = scalbn(d, -ilogbw);
 }
 denom = c * c + d * d;
 x = scalbn((a * c + b * d) / denom, -ilogbw);
 y = scalbn((b * c - a * d) / denom, -ilogbw);

 /* Recover infinities and zeros that computed as NaN+iNaN; */
 /* the only cases are nonzero/zero, infinite/finite, and finite/infinite, ... */

 if (isnan(x) && isnan(y)) {
 if ((denom == 0.0) &&
 (!isnan(a) || !isnan(b))) {
 x = copysign(INFINITY, c) * a;
 y = copysign(INFINITY, c) * b;
 }
 else if ((isinf(a) || isinf(b)) &&
 isfinite(c) && isfinite(d)) {
 a = copysign(isinf(a) ? 1.0 : 0.0, a);
 b = copysign(isinf(b) ? 1.0 : 0.0, b);
 x = INFINITY * (a * c + b * d);
 y = INFINITY * (b * c - a * d);
 }
 else if ((logbw == INFINITY) &&
 isfinite(a) && isfinite(b)) {
 c = copysign(isinf(c) ? 1.0 : 0.0, c);
 d = copysign(isinf(d) ? 1.0 : 0.0, d);
 x = 0.0 * (a * c + b * d);
 y = 0.0 * (b * c - a * d);
 }
 }
 return CMPLX(x, y);
}

```

- 8 Scaling the denominator alleviates the main overflow and underflow problem, which is more serious than for multiplication. In the spirit of the preceding multiplication example, this code does not defend against overflow and underflow in the calculation of the numerator. Scaling with the `scalbn` function, instead of with division, provides better roundoff characteristics.

### G.3.3 Additive operators

#### Semantics

- 1 In all cases the result and floating-point exception behavior of a `+` or `-` operator is defined by the usual mathematical formula in Table G.3:

Table G.3 — Results of addition or subtraction operations

| + or -   | $u$              | $u + iv$                 |
|----------|------------------|--------------------------|
| $x$      | $x \pm u$        | $(x \pm u) \pm iv$       |
| $x + iy$ | $(x \pm u) + iy$ | $(x \pm u) + i(y \pm v)$ |

## G.4 Complex arithmetic <complex.h>

### G.4.1 General

- 1 This subclause contains specifications for the <complex.h> functions that are particularly suited to ISO/IEC 60559 implementations. For families of functions, the specifications apply to all of the functions even though only the principal function is shown. Unless otherwise specified, where the symbol “ $\pm$ ” occurs in both an argument and the result, the result has the same sign as the argument.
- 2 The functions are continuous onto both sides of their branch cuts, taking into account the sign of zero. For example, `csqrt`( $-2 \pm i0$ ) =  $\pm i\sqrt{2}$ .
- 3 Since complex values are composed of real values, each function may be regarded as computing real values from real values. Except as noted, the functions treat real infinities, NaNs, signed zeros, subnormals, and the floating-point exception flags in a manner consistent with the specifications for real functions in F.10.<sup>455)</sup>
- 4 In subsequent subclauses in G.4 “NaN” refers to a quiet NaN. The behavior of signaling NaNs in this annex is implementation-defined.
- 5 The functions `cimag`, `conj`, `cproj`, and `creal` are fully specified for all implementations, including ISO/IEC 60559 ones, in 7.3.9. These functions raise no floating-point exceptions.
- 6 Each of the functions `cabs` and `carg` is specified by a formula in terms of a real function (whose special cases are covered in Annex F):

```
cabs(x + iy) = hypot(x, y)
carg(x + iy) = atan2(y, x)
```

- 7 Each of the functions `casin`, `catan`, `ccos`, `csin`, and `ctan` is specified implicitly by a formula in terms of other complex functions (whose special cases are specified below):

```
casin(z) = -i casinh(iz)
catan(z) = -i catanh(iz)
ccos(z) = ccosh(iz)
csin(z) = -i csinh(iz)
ctan(z) = -i ctanh(iz)
```

- 8 For the other functions, the following subclauses specify behavior for special cases, including treatment of the “invalid” and “divide-by-zero” floating-point exceptions. For families of functions, the specifications apply to all of the functions even though only the principal function is shown. For a function  $f$  satisfying  $f(\text{conj}(z)) = \text{conj}(f(z))$ , the specifications for the upper half-plane imply the specifications for the lower half-plane; if the function  $f$  is also either even,  $f(-z) = f(z)$ , or odd,  $f(-z) = -f(z)$ , then the specifications for the first quadrant imply the specifications for the other three quadrants.
- 9 In the following subclauses,  $\text{cis}(y)$  is defined as  $\cos(y) + i \sin(y)$ .

### G.4.2 Trigonometric functions

#### G.4.2.1 The `cacos` functions

- 1 — `cacos`(`conj`( $z$ )) = `conj`(`cacos`( $z$ )).
- `cacos`( $\pm 0 + i0$ ) returns  $\frac{\pi}{2} - i0$ .
- `cacos`( $\pm 0 + i\text{NaN}$ ) returns  $\frac{\pi}{2} + i\text{NaN}$ .

<sup>455)</sup>As noted in G.2, a complex value with at least one infinite part is regarded as an infinity even if its other part is a quiet NaN.



- **cacos**( $x + i\infty$ ) returns  $\frac{\pi}{2} - i\infty$ , for finite  $x$ .
- **cacos**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for nonzero finite  $x$ .
- **cacos**( $-\infty + iy$ ) returns  $\pi - i\infty$ , for positive-signed finite  $y$ .
- **cacos**( $+\infty + iy$ ) returns  $+0 - i\infty$ , for positive-signed finite  $y$ .
- **cacos**( $-\infty + i\infty$ ) returns  $3\frac{\pi}{4} - i\infty$ .
- **cacos**( $+\infty + i\infty$ ) returns  $\frac{\pi}{4} - i\infty$ .
- **cacos**( $\pm\infty + i\text{NaN}$ ) returns  $\text{NaN} \pm i\infty$  (where the sign of the imaginary part of the result is unspecified).
- **cacos**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $y$ .
- **cacos**( $\text{NaN} + i\infty$ ) returns  $\text{NaN} - i\infty$ .
- **cacos**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

### G.4.3 Hyperbolic functions

#### G.4.3.1 The **cacosh** functions

1

- **cacosh**(**conj**( $z$ )) = **conj**(**cacosh**( $z$ )).
- **cacosh**( $\pm 0 + i0$ ) returns  $+0 + \frac{i\pi}{2}$ .
- **cacosh**( $x + i\infty$ ) returns  $+\infty + \frac{i\pi}{2}$ , for finite  $x$ .
- **cacosh**( $0 + i\text{NaN}$ ) returns  $\text{NaN} \pm \frac{i\pi}{2}$  (where the sign of the imaginary part of the result is unspecified).
- **cacosh**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **cacosh**( $-\infty + iy$ ) returns  $+\infty + i\pi$ , for positive-signed finite  $y$ .
- **cacosh**( $+\infty + iy$ ) returns  $+\infty + i0$ , for positive-signed finite  $y$ .
- **cacosh**( $-\infty + i\infty$ ) returns  $+\infty + i\frac{3\pi}{4}$ .
- **cacosh**( $+\infty + i\infty$ ) returns  $+\infty + \frac{i\pi}{4}$ .
- **cacosh**( $\pm\infty + i\text{NaN}$ ) returns  $+\infty + i\text{NaN}$ .
- **cacosh**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $y$ .
- **cacosh**( $\text{NaN} + i\infty$ ) returns  $+\infty + i\text{NaN}$ .
- **cacosh**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

**G.4.3.2 The `casinh` functions**

- 1 — **`casinh(conj(z)) = conj(casinh(z))`** and **`casinh`** is odd.
- **`casinh(+0 + i0)`** returns  $0 + i0$ .
- **`casinh(x + i∞)`** returns  $+\infty + \frac{i\pi}{2}$  for positive-signed finite  $x$ .
- **`casinh(x + iNaN)`** returns NaN +  $iNaN$  and optionally raises the “invalid” floating-point exception, for finite  $x$ .
- **`casinh(+∞ + iy)`** returns  $+\infty + i0$  for positive-signed finite  $y$ .
- **`casinh(+∞ + i∞)`** returns  $+\infty + \frac{i\pi}{4}$ .
- **`casinh(+∞ + iNaN)`** returns  $+\infty + iNaN$ .
- **`casinh(NaN + i0)`** returns NaN +  $i0$ .
- **`casinh(NaN + iy)`** returns NaN +  $iNaN$  and optionally raises the “invalid” floating-point exception, for finite nonzero  $y$ .
- **`casinh(NaN + i∞)`** returns  $\pm\infty + iNaN$  (where the sign of the real part of the result is unspecified).
- **`casinh(NaN + iNaN)`** returns NaN +  $iNaN$ .

**G.4.3.3 The `catanh` functions**

- 1 — **`catanh(conj(z)) = conj(catanh(z))`** and **`catanh`** is odd.
- **`catanh(+0 + i0)`** returns  $+0 + i0$ .
- **`catanh(+0 + iNaN)`** returns  $+0 + iNaN$ .
- **`catanh(+1 + i0)`** returns  $+\infty + i0$  and raises the “divide-by-zero” floating-point exception.
- **`catanh(x + i∞)`** returns  $+0 + \frac{i\pi}{2}$ , for finite positive-signed  $x$ .
- **`catanh(x + iNaN)`** returns NaN +  $iNaN$  and optionally raises the “invalid” floating-point exception, for nonzero finite  $x$ .
- **`catanh(+∞ + iy)`** returns  $+0 + \frac{i\pi}{2}$ , for finite positive-signed  $y$ .
- **`catanh(+∞ + i∞)`** returns  $+0 + \frac{i\pi}{2}$ .
- **`catanh(+∞ + iNaN)`** returns  $+0 + iNaN$ .
- **`catanh(NaN + iy)`** returns NaN +  $iNaN$  and optionally raises the “invalid” floating-point exception, for finite  $y$ .
- **`catanh(NaN + i∞)`** returns  $\pm 0 + \frac{i\pi}{2}$  (where the sign of the real part of the result is unspecified).
- **`catanh(NaN + iNaN)`** returns NaN +  $iNaN$ .

**G.4.3.4 The `ccosh` functions**

- 1 — **`ccosh(conj(z)) = conj(ccosh(z))`** and **`ccosh`** is even.
- **`ccosh(+0 + i0)`** returns  $1 + i0$ .
- **`ccosh(+0 + i∞)`** returns NaN ±  $i0$  (where the sign of the imaginary part of the result is unspecified) and raises the “invalid” floating-point exception.
- **`ccosh(+0 + iNaN)`** returns NaN ±  $i0$  (where the sign of the imaginary part of the result is unspecified).

- **ccosh**( $x + i\infty$ ) returns  $\text{NaN} + i\text{NaN}$  and raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **ccosh**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **ccosh**( $+\infty + i0$ ) returns  $+\infty + i0$ .
- **ccosh**( $+\infty + iy$ ) returns  $+\infty \text{ cis}(y)$ , for finite nonzero  $y$ .
- **ccosh**( $+\infty + i\infty$ ) returns  $\pm\infty + i\text{NaN}$  (where the sign of the real part of the result is unspecified) and raises the “invalid” floating-point exception.
- **ccosh**( $+\infty + i\text{NaN}$ ) returns  $+\infty + i\text{NaN}$ .
- **ccosh**( $\text{NaN} + i0$ ) returns  $\text{NaN} \pm i0$  (where the sign of the imaginary part of the result is unspecified).
- **ccosh**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for all nonzero numbers  $y$ .
- **ccosh**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

#### G.4.3.5 The **csinh** functions

1

- **csinh**(**conj**( $z$ )) = **conj**(**csinh**( $z$ )) and **csinh** is odd.
- **csinh**( $+0 + i0$ ) returns  $+0 + i0$ .
- **csinh**( $+0 + i\infty$ ) returns  $\pm 0 + i\text{NaN}$  (where the sign of the real part of the result is unspecified) and raises the “invalid” floating-point exception.
- **csinh**( $+0 + i\text{NaN}$ ) returns  $\pm 0 + i\text{NaN}$  (where the sign of the real part of the result is unspecified).
- **csinh**( $x + i\infty$ ) returns  $\text{NaN} + i\text{NaN}$  and raises the “invalid” floating-point exception, for positive finite  $x$ .
- **csinh**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **csinh**( $+\infty + i0$ ) returns  $+\infty + i0$ .
- **csinh**( $+\infty + iy$ ) returns  $+\infty \text{ cis}(y)$ , for positive finite  $y$ .
- **csinh**( $+\infty + i\infty$ ) returns  $\pm\infty + i\text{NaN}$  (where the sign of the real part of the result is unspecified) and raises the “invalid” floating-point exception.
- **csinh**( $+\infty + i\text{NaN}$ ) returns  $\pm\infty + i\text{NaN}$  (where the sign of the real part of the result is unspecified).
- **csinh**( $\text{NaN} + i0$ ) returns  $\text{NaN} + i0$ .
- **csinh**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for all nonzero numbers  $y$ .
- **csinh**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

#### G.4.3.6 The **ctanh** functions

- 1 — **ctanh**(**conj**( $z$ )) = **conj**(**ctanh**( $z$ )) and **ctanh** is odd.
- **ctanh**( $+0 + i0$ ) returns  $+0 + i0$ .
- **ctanh**( $0 + i\infty$ ) returns  $0 + i\text{NaN}$  and raises the “invalid” floating-point exception.
- **ctanh**( $x + i\infty$ ) returns  $\text{NaN} + i\text{NaN}$  and raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **ctanh**( $0 + i\text{NaN}$ ) returns  $0 + i\text{NaN}$ .
- **ctanh**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite nonzero  $x$ .
- **ctanh**( $+\infty + iy$ ) returns  $1 + i0 \sin(2y)$ , for positive-signed finite  $y$ .
- **ctanh**( $+\infty + i\infty$ ) returns  $1 \pm i0$  (where the sign of the imaginary part of the result is unspecified).
- **ctanh**( $+\infty + i\text{NaN}$ ) returns  $1 \pm i0$  (where the sign of the imaginary part of the result is unspecified).
- **ctanh**( $\text{NaN} + i0$ ) returns  $\text{NaN} + i0$ .
- **ctanh**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for all nonzero numbers  $y$ .
- **ctanh**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

### G.4.4 Exponential and logarithmic functions

#### G.4.4.1 The **cexp** functions

- 1 — **cexp**(**conj**( $z$ )) = **conj**(**cexp**( $z$ )).
- **cexp**( $\pm 0 + i0$ ) returns  $1 + i0$ .
- **cexp**( $x + i\infty$ ) returns  $\text{NaN} + i\text{NaN}$  and raises the “invalid” floating-point exception, for finite  $x$ .
- **cexp**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $x$ .
- **cexp**( $+\infty + i0$ ) returns  $+\infty + i0$ .
- **cexp**( $-\infty + iy$ ) returns  $+0 \text{ cis}(y)$ , for finite  $y$ .
- **cexp**( $+\infty + iy$ ) returns  $+\infty \text{ cis}(y)$ , for finite nonzero  $y$ .
- **cexp**( $-\infty + i\infty$ ) returns  $\pm 0 \pm i0$  (where the signs of the real and imaginary parts of the result are unspecified).
- **cexp**( $+\infty + i\infty$ ) returns  $\pm \infty + i\text{NaN}$  and raises the “invalid” floating-point exception (where the sign of the real part of the result is unspecified).
- **cexp**( $-\infty + i\text{NaN}$ ) returns  $\pm 0 \pm i0$  (where the signs of the real and imaginary parts of the result are unspecified).
- **cexp**( $+\infty + i\text{NaN}$ ) returns  $\pm \infty + i\text{NaN}$  (where the sign of the real part of the result is unspecified).
- **cexp**( $\text{NaN} + i0$ ) returns  $\text{NaN} + i0$ .
- **cexp**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for all nonzero numbers  $y$ .
- **cexp**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

### G.4.4.2 The **clog** functions

- 1 — **clog**(**conj**( $z$ )) = **conj**(**clog**( $z$ )).
- **clog**( $-0 + i0$ ) returns  $-\infty + i\pi$  and raises the “divide-by-zero” floating-point exception.
- **clog**( $+0 + i0$ ) returns  $-\infty + i0$  and raises the “divide-by-zero” floating-point exception.
- **clog**( $x + i\infty$ ) returns  $+\infty + \frac{i\pi}{2}$ , for finite  $x$ .
- **clog**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $x$ .
- **clog**( $-\infty + iy$ ) returns  $+\infty + i\pi$ , for finite positive-signed  $y$ .
- **clog**( $+\infty + iy$ ) returns  $+\infty + i0$ , for finite positive-signed  $y$ .
- **clog**( $-\infty + i\infty$ ) returns  $+\infty + i\frac{3\pi}{4}$ .
- **clog**( $+\infty + i\infty$ ) returns  $+\infty + i\frac{\pi}{4}$ .
- **clog**( $\pm\infty + i\text{NaN}$ ) returns  $+\infty + i\text{NaN}$ .
- **clog**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $y$ .
- **clog**( $\text{NaN} + i\infty$ ) returns  $+\infty + i\text{NaN}$ .
- **clog**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

## G.4.5 Power and absolute-value functions

### G.4.5.1 The **cpow** functions

- 1 The **cpow** functions raise floating-point exceptions if appropriate for the calculation of the parts of the result, and may also raise spurious floating-point exceptions.<sup>456)</sup>

### G.4.5.2 The **csqrt** functions

- 1 — **csqrt**(**conj**( $z$ )) = **conj**(**csqrt**( $z$ )).
- **csqrt**( $\pm 0 + i0$ ) returns  $+0 + i0$ .
- **csqrt**( $x + i\infty$ ) returns  $+\infty + i\infty$ , for all  $x$  (including  $\text{NaN}$ ).
- **csqrt**( $x + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $x$ .
- **csqrt**( $-\infty + iy$ ) returns  $+0 + i\infty$ , for finite positive-signed  $y$ .
- **csqrt**( $+\infty + iy$ ) returns  $+\infty + i0$ , for finite positive-signed  $y$ .
- **csqrt**( $-\infty + i\text{NaN}$ ) returns  $\text{NaN} \pm i\infty$  (where the sign of the imaginary part of the result is unspecified).
- **csqrt**( $+\infty + i\text{NaN}$ ) returns  $+\infty + i\text{NaN}$ .
- **csqrt**( $\text{NaN} + iy$ ) returns  $\text{NaN} + i\text{NaN}$  and optionally raises the “invalid” floating-point exception, for finite  $y$ .
- **csqrt**( $\text{NaN} + i\text{NaN}$ ) returns  $\text{NaN} + i\text{NaN}$ .

<sup>456)</sup>This allows **cpow**( $z, c$ ) to be implemented as **cexp**(**cclog**( $z$ )) without precluding implementations that treat special cases more carefully.

## Annex H

(normative)

### ISO/IEC 60559 interchange and extended types

#### H.1 Introduction

- 1 This annex specifies extension types for programming language C that have the arithmetic interchange and extended floating-point formats specified in ISO/IEC 60559. This annex also includes functions that support the non-arithmetic interchange formats in that standard. This annex was adapted from ISO/IEC TS 18661-3:2015, Floating-point extensions for C —Interchange and extended types.
- 2 An implementation that defines `__STDC_IEC_60559_TYPES__` to 202ymmL shall conform to the specifications in this annex. An implementation may define `__STDC_IEC_60559_TYPES__` only if it defines `__STDC_IEC_60559_BFP__`, indicating support for ISO/IEC 60559 binary floating-point arithmetic, or defines `__STDC_IEC_60559_DFP__`, indicating support for ISO/IEC 60559 decimal floating-point arithmetic (or defines both). Where a binding between the C language and ISO/IEC 60559 is indicated, the ISO/IEC 60559-specified behavior is adopted by reference, unless stated otherwise.

#### H.2 Types

##### H.2.1 General

- 1 This clause specifies types that support ISO/IEC 60559 arithmetic interchange and extended formats. The encoding conversion functions (H.11.4) and numeric conversion functions for encodings (H.12.4, H.12.5) support the non-arithmetic interchange formats specified in ISO/IEC 60559.

##### H.2.2 Interchange floating types

- 1 ISO/IEC 60559 specifies interchange formats, and their encodings, which can be used for the exchange of floating-point data between implementations. These formats are identified by their radix (binary or decimal) and their storage width  $N$ . Table H.1, Table H.2, Table H.3, and Table H.4 give the C floating-point model parameters<sup>457)</sup> (5.3.5.3.3) for the ISO/IEC 60559 interchange formats, where the function `round()` rounds to the nearest integer.

**Table H.1 — Binary interchange format parameters**

| Parameter                               | binary16 | binary32 | binary64 | binary128 |
|-----------------------------------------|----------|----------|----------|-----------|
| $N$ , storage width in bits             | 16       | 32       | 64       | 128       |
| $p$ , precision in binary digits (bits) | 11       | 24       | 53       | 113       |
| $e_{\max}$ , maximum exponent $e$       | 16       | 128      | 1024     | 16384     |
| $e_{\min}$ , minimum exponent $e$       | −13      | −125     | −1021    | −16381    |

**Table H.2 — Binary interchange format parameters for arbitrary  $N$**

| Parameter                               | binary $N$ ( $N \geq 128$ )                 |
|-----------------------------------------|---------------------------------------------|
| $N$ , storage width in bits             | $N$ , a multiple of 32                      |
| $p$ , precision in binary digits (bits) | $N - \text{round}(4 \times \log_2(N)) + 13$ |
| $e_{\max}$ , maximum exponent $e$       | $2^{(N-p-1)}$                               |
| $e_{\min}$ , minimum exponent $e$       | $3 - e_{\max}$                              |

<sup>457)</sup>In ISO/IEC 60559, normal floating-point numbers are expressed with the first significant digit to the left of the radix point. Hence the exponent in the C model (shown in Table H.1, Table H.2 Table H.3, and Table H.4) is 1 more than the exponent of the same number in the ISO/IEC 60559 model.

Table H.3 — Decimal interchange format parameters

| Parameter                         | decimal32 | decimal64 | decimal128 |
|-----------------------------------|-----------|-----------|------------|
| $N$ , storage width in bits       | 32        | 64        | 128        |
| $p$ , precision in decimal digits | 7         | 16        | 34         |
| $e_{\max}$ , maximum exponent $e$ | 97        | 385       | 6145       |
| $e_{\min}$ , minimum exponent $e$ | −94       | −382      | −6142      |

Table H.4 — Decimal interchange format parameters for arbitrary  $N$ 

| Parameter                         | decimal $N$ ( $N \geq 32$ )          |
|-----------------------------------|--------------------------------------|
| $N$ , storage width in bits       | $N$ , a multiple of 32               |
| $p$ , precision in decimal digits | $9 \times (N \div 32) - 2$           |
| $e_{\max}$ , maximum exponent $e$ | $3 \times 2^{((N \div 16) + 3)} + 1$ |
| $e_{\min}$ , minimum exponent $e$ | $3 - e_{\max}$                       |

- 2 EXAMPLE For the binary160 format,  $p = 144$ ,  $e_{\max} = 32678$  and  $e_{\min} = -32765$ . For the decimal160 format,  $p = 43$ ,  $e_{\max} = 24577$  and  $e_{\min} = -24574$ .
- 3 Types designated:

**\_Float $N$**

where  $N$  is 16, 32, 64, or  $\geq 128$  and a multiple of 32; and, types designated

**\_Decimal $N$**

where  $N \geq 32$  and a multiple of 32, are collectively called the *interchange floating types*. Each interchange floating type has the ISO/IEC 60559 interchange format corresponding to its width ( $N$ ) and radix (2 for **\_Float $N$** , 10 for **\_Decimal $N$** ). Each interchange floating type is not compatible with any other type.

- 4 An implementation that defines **\_\_STDC\_IEC\_60559\_BFP\_\_** and **\_\_STDC\_IEC\_60559\_TYPES\_\_** shall provide **\_Float32** and **\_Float64** as interchange floating types with the same representation and alignment requirements as **float** and **double**, respectively. If the implementation's **long double** type supports an ISO/IEC 60559 interchange format of width  $N > 64$ , then the implementation shall also provide the type **\_Float $N$**  as an interchange floating type with the same representation and alignment requirements as **long double**. The implementation may provide other radix-2 interchange floating types **\_Float $N$** ; the set of such types supported is implementation-defined.
- 5 An implementation that defines **\_\_STDC\_IEC\_60559\_DFP\_\_** provides the decimal floating types **\_Decimal32**, **\_Decimal64**, and **\_Decimal128** (6.2.5). If the implementation also defines **\_\_STDC\_IEC\_60559\_TYPES\_\_**, it may provide other radix-10 interchange floating types **\_Decimal $N$** ; the set of such types supported is implementation-defined.

### H.2.3 Non-arithmetic interchange formats

- 1 An implementation supports ISO/IEC 60559 non-arithmetic interchange formats by providing the associated encoding-to-encoding conversion functions (H.11.4.3) in `<math.h>` and the string-from-encoding functions (H.12.4) and string-to-encoding functions (H.12.5) in `<stdlib.h>`.
- 2 An implementation that defines **\_\_STDC\_IEC\_60559\_BFP\_\_** and **\_\_STDC\_IEC\_60559\_TYPES\_\_** supports some ISO/IEC 60559 radix-2 interchange formats as arithmetic formats by providing types **\_Float $N$**  (as well as **float** and **double**) with those formats. The implementation may support other ISO/IEC 60559 radix-2 interchange formats as non-arithmetic formats; the set of such formats supported is implementation-defined.
- 3 An implementation that defines **\_\_STDC\_IEC\_60559\_DFP\_\_** and **\_\_STDC\_IEC\_60559\_TYPES\_\_** supports some ISO/IEC 60559 radix-10 interchange formats as arithmetic formats by providing types



**\_DecimalN** with those formats. The implementations may support other ISO/IEC 60559 radix-10 interchange formats as non-arithmetic formats; the set of such formats supported is implementation-defined.

## H.2.4 Extended floating types

- For each of its basic formats, ISO/IEC 60559 specifies an extended format whose maximum exponent and precision exceed those of the basic format it is associated with. Extended formats are intended for arithmetic with more precision and exponent range than is available in the basic formats used for the input data. The extra precision and range often mitigate round-off error and eliminate overflow and underflow in intermediate computations. Table H.5 gives the minimum values of these parameters, as defined for the C floating-point model (5.3.5.3.3). For all ISO/IEC 60559 extended (and interchange) formats,  $e_{\min} = 3 - e_{\max}$ .

Table H.5 — Extended format parameters for floating-point numbers

|                   | Extended formats associated with: |          |           |           |            |
|-------------------|-----------------------------------|----------|-----------|-----------|------------|
| Parameter         | binary32                          | binary64 | binary128 | decimal64 | decimal128 |
| $p$ digits $\geq$ | 32                                | 64       | 128       | 22        | 40         |
| $e_{\max} \geq$   | 1024                              | 16384    | 65536     | 6145      | 24577      |

- Types designated **\_Float32x**, **\_Float64x**, **\_Float128x**, **\_Decimal64x**, and **\_Decimal128x** support the corresponding ISO/IEC 60559 extended formats and are collectively called the *extended floating types*. The set of values of **\_Float32x** is a subset of the set of values of **\_Float64x**; the set of values of **\_Float64x** is a subset of the set of values of **\_Float128x**. The set of values of **\_Decimal64x** is a subset of the set of values of **\_Decimal128x**. Each extended floating type is not compatible with any other type. An implementation that defines **\_\_STDC\_IEC\_60559\_BFP\_\_** and **\_\_STDC\_IEC\_60559\_TYPES\_\_** shall provide **\_Float32x**, and may provide one or both of the types **\_Float64x** and **\_Float128x**. An implementation that defines **\_\_STDC\_IEC\_60559\_DFP\_\_** and **\_\_STDC\_IEC\_60559\_TYPES\_\_** shall provide **\_Decimal64x**, and may provide **\_Decimal128x**. Which (if any) of the optional extended floating types are provided is implementation-defined.
- NOTE 1 ISO/IEC 60559 does not specify an extended format associated with the decimal32 format, nor does this annex specify an extended type associated with the **\_Decimal32** type.
- NOTE 2 The **\_Float32x** type can have the same format as **double**. The **\_Decimal64x** type can have the same format as **\_Decimal128**.

## H.2.5 Classification of real floating types

- 6.2.5 defines standard floating types as a collective name for the types **float**, **double** and **long double** and it defines decimal floating types as a collective name for the types **\_Decimal32**, **\_Decimal64**, and **\_Decimal128**.
- H.2.2 defines interchange floating types and H.2.4 defines extended floating types.
- The types **\_FloatN** and **\_FloatNx** are collectively called *binary floating types*.
- This subclause broadens *decimal floating types* to include the types **\_DecimalN** and **\_DecimalNx**, introduced in this annex, as well as **\_Decimal32**, **\_Decimal64**, and **\_Decimal128**.
- This subclause broadens *real floating types* to include all interchange floating types and extended floating types, as well as standard floating types.
- Thus, in this annex, real floating types are classified as follows:

- standard floating types, composed of **float**, **double**, **long double**;
- decimal floating types, composed of **\_DecimalN**, **\_DecimalNx**;



- binary floating types, composed of `_FloatN`, `_FloatNx`;
- interchange floating types, composed of `_FloatN`, `_DecimalN`; and,
- extended floating types, composed of `_FloatNx`, `_DecimalNx`.

- 7 NOTE Standard floating types (which have an implementation-defined radix) are not included in either binary floating types (which always have radix 2) or decimal floating types (which always have radix 10).

## H.2.6 Complex types

- 1 This subclause broadens the C complex types (6.2.5) to also include similar types whose corresponding real parts have binary floating types. For the types `_FloatN` and `_FloatNx`, there are complex types designated respectively as `_FloatN_Complex` and `_FloatNx_Complex`. (Complex types are a conditional feature that implementations are not required to support; see 6.10.10.4.)

## H.3 Characteristics in `<float.h>`

- 1 This subclause enhances the `FLT_EVAL_METHOD` and `DEC_EVAL_METHOD` macros to apply to the types introduced in this annex.
- 2 If `FLT_RADIX` is 2, the value of `FLT_EVAL_METHOD` (5.3.5.3.3) characterizes the use of evaluation formats for standard floating types and for binary floating types:

- |              |                                                                                                                                                                                                                                                                                                                                                                       |
|--------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| -1           | indeterminable;                                                                                                                                                                                                                                                                                                                                                       |
| 0            | evaluate all operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>float</code> , to the range and precision of <code>float</code> ; evaluate all other operations and literals to the range and precision of the semantic type;                                                                      |
| 1            | evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>double</code> , to the range and precision of <code>double</code> ; evaluate all other operations and literals to the range and precision of the semantic type;                                                                        |
| 2            | evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>long double</code> , to the range and precision of <code>long double</code> ; evaluate all other operations and literals to the range and precision of the semantic type;                                                              |
| <i>N</i>     | where <code>_FloatN</code> is a supported interchange floating type, evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>_FloatN</code> , to the range and precision of <code>_FloatN</code> ; evaluate all other operations and literals to the range and precision of the semantic type; |
| <i>N</i> + 1 | where <code>_FloatNx</code> is a supported extended floating type, evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>_FloatNx</code> , to the range and precision of <code>_FloatNx</code> ; evaluate all other operations and literals to the range and precision of the semantic type. |

If `FLT_RADIX` is not 2, the use of evaluation formats for operations and literals of binary floating types is implementation-defined.

- 3 The implementation-defined value of `DEC_EVAL_METHOD` (5.3.5.3.4) characterizes the use of evaluation formats for decimal floating types:

- |    |                                                                                                                                                                                                                                                                                                        |
|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| -1 | indeterminable;                                                                                                                                                                                                                                                                                        |
| 0  | evaluate all operations and literals just to the range and precision of the type;                                                                                                                                                                                                                      |
| 1  | evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of <code>_Decimal64</code> , to the range and precision of <code>_Decimal64</code> ; evaluate all other operations and literals to the range and precision of the semantic type; |

- 2 evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of **\_Decimal128**, to the range and precision of **\_Decimal128**; evaluate all other operations and literals to the range and precision of the semantic type;
- $N$  where **\_Decimal $N$**  is a supported interchange floating type, evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of **\_Decimal $N$** , to the range and precision of **\_Decimal $N$** ; evaluate all other operations and literals to the range and precision of the semantic type;
- $N + 1$  where **\_Decimal $N$ x** is a supported extended floating type, evaluate operations and literals, whose semantic type comprises a set of values that is a strict subset of the values of **\_Decimal $N$ x**, to the range and precision of **\_Decimal $N$ x**; evaluate all other operations and literals to the range and precision of the semantic type.
- 4 This subclause also specifies `<float.h>` macros, analogous to the macros for standard floating types, that characterize binary floating types in terms of the model presented in 5.3.5.3.3. This subclause generalizes the specification of characteristics in 5.3.5.3.4 to include the decimal floating types introduced in this annex. The prefix **FLT $N$ \_** indicates the type **\_Float $N$**  or the non-arithmetic binary interchange format of width  $N$ . The prefix **FLT $N$ x\_** indicates the type **\_Float $N$ x**. The prefix **DEC $N$ \_** indicates the type **\_Decimal $N$**  or the non-arithmetic decimal interchange format of width  $N$ . The prefix **DEC $N$ x\_** indicates the type **\_Decimal $N$ x**. The type parameters  $p$ ,  $e_{\max}$ , and  $e_{\min}$  for extended floating types are for the extended floating type itself, not for the basic format that it extends.
- 5 If **\_\_STDC\_WANT\_IEC\_60559\_TYPES\_EXT\_\_** is defined (by the user) at the point in the code where `<float.h>` is first included, the following applies (H.8). For each interchange or extended floating type that the implementation provides, `<float.h>` shall define the associated macros in the following lists. Conversely, for each such type that the implementation does not provide, `<float.h>` shall not define the associated macros in the following list, except, the implementation shall define the macros **FLT $N$ \_DECIMAL\_DIG** and **FLT $N$ \_DIG** if it supports the ISO/IEC 60559 non-arithmetic binary interchange format of width  $N$  (H.2.3).
- 6 The signaling NaN macros

```

FLT N _SNAN
DEC N _SNAN
FLT N x_SNAN
DEC N x_SNAN

```

expand to constant expressions of types **\_Float $N$** , **\_Decimal $N$** , **\_Float $N$ x**, and **\_Decimal $N$ x** respectively, representing a signaling NaN. If a signaling NaN macro, optionally preceded by the unary + or - operator, is used as an initializer that is evaluated at translation time to initialize an object of the same type, the object is initialized with a signaling NaN value.

- 7 The integer values given in the following lists shall be replaced by integer constant expressions:

- radix of exponent representation,  $b$  (2 for binary, 10 for decimal)

For the standard floating types, this value is implementation-defined and is specified by the macro **FLT\_RADIX**. For the interchange and extended floating types there is no corresponding macro; the radix is an inherent property of the types.

- The number of bits in the floating-point significand,  $p$

```

FLT N _MANT_DIG
FLT N x_MANT_DIG

```

- The number of digits in the coefficient,  $p$

**DECN\_MANT\_DIG**  
**DECNX\_MANT\_DIG**

- number of decimal digits,  $n$ , such that any floating-point number with  $p$  bits can be rounded to a floating-point number with  $n$  decimal digits and back again, using to-nearest rounding for both roundings, without change to the value,  $\lceil 1 + p \log_{10}(2) \rceil$

**FLTN\_DECIMAL\_DIG**  
**FLTNX\_DECIMAL\_DIG**

- number of decimal digits,  $q$ , such that any floating-point number with  $q$  decimal digits can be rounded to a floating-point number with  $p$  bits and back again, using to-nearest rounding for both roundings, without a change to the  $q$  decimal digits,  $\lfloor (p - 1) \log_{10}(2) \rfloor$

**FLTN\_DIG**  
**FLTNX\_DIG**

- minimum negative integer such that the radix raised to one less than that power is a normalized floating-point number,  $e_{\min}$

**FLTN\_MIN\_EXP**  
**FLTNX\_MIN\_EXP**  
**DECN\_MIN\_EXP**  
**DECNX\_MIN\_EXP**

- minimum negative integer such that 10 raised to that power is in the range of normalized floating-point numbers,  $\lceil \log_{10}(2)^{e_{\min} - 1} \rceil$

**FLTN\_MIN\_10\_EXP**  
**FLTNX\_MIN\_10\_EXP**

- maximum positive integer such that the radix raised to one less than that power is a representable finite floating-point number,  $e_{\max}$

**FLTN\_MAX\_EXP**  
**FLTNX\_MAX\_EXP**  
**DECN\_MAX\_EXP**  
**DECNX\_MAX\_EXP**

- maximum integer such that 10 raised to that power is in the range of representable finite floating-point numbers,  $\lfloor \log_{10}((1 - 2^{-p})2^{e_{\max}}) \rfloor$

**FLTN\_MAX\_10\_EXP**  
**FLTNX\_MAX\_10\_EXP**

- maximum representable finite floating-point number,  $(1 - b^{-p})b^{e_{\max}}$

**FLTN\_MAX**  
**FLTNX\_MAX**  
**DECN\_MAX**  
**DECNX\_MAX**

- the difference between 1 and the least normalized value greater than 1 that is representable in the given floating type,  $b^{1-p}$

```

FLT_N_EPSILON
FLT_X_EPSILON
DEC_N_EPSILON
DEC_X_EPSILON

```

- minimum normalized positive floating-point number,  $b^{e_{\min} - 1}$

```

FLT_N_MIN
FLT_X_MIN
DEC_N_MIN
DEC_X_MIN

```

- minimum positive floating-point number,  $b^{e_{\min} - p}$

```

FLT_N_TRUE_MIN
FLT_X_TRUE_MIN
DEC_N_TRUE_MIN
DEC_X_TRUE_MIN

```

## H.4 Conversions

### H.4.1 General

- 1 This subclause enhances the usual arithmetic conversions (6.3.2.8) to handle interchange and extended floating types. It supports the ISO/IEC 60559 recommendation against allowing implicit conversions of operands to obtain a common type where the conversion is between types where neither is a subset of (or equivalent to) the other.
- 2 This subclause also broadens the operation binding in F.3 for the ISO/IEC 60559 `convertFormat` operation to apply to ISO/IEC 60559 arithmetic and non-arithmetic formats.

### H.4.2 Real floating and integer

- 1 When a finite value of interchange or extended floating type is converted to an integer type other than **bool**, the fractional part is discarded (i.e. the value is truncated toward zero). If the value of the integral part cannot be represented by the integer type, the “invalid” floating-point exception shall be raised and the result of the conversion is unspecified.
- 2 When a value of integer type is converted to an interchange or extended floating type, if the value being converted can be represented exactly in the new type, it is unchanged. If the value being converted cannot be represented exactly, the result shall be correctly rounded with exceptions raised as specified in ISO/IEC 60559.

### H.4.3 Usual arithmetic conversions

- 1 If either operand is of floating type, the common real type is determined as follows:
  - If one operand has decimal floating type, the other operand shall not have standard floating type, binary floating type, or complex type.
  - If only one operand has a floating type, the other operand is converted to the corresponding real type of the operand of floating type.
  - If both operands have the same corresponding real type, no further conversion is needed.
  - If both operands have floating types and neither of the sets of values of their corresponding real types is a subset of (or equivalent to) the other, the behavior is undefined.
  - Otherwise, if both operands are floating types and the sets of values of their corresponding real types are not equivalent, the operand whose set of values of its corresponding real type is a strict subset of the set of values of the corresponding real type of the other operand is

converted, without change of type domain, to a type with the corresponding real type of that other operand.

- Otherwise, if both operands are floating types and the sets of values of their corresponding real types are equivalent, then the following rules are applied:
  - If the corresponding real type of either operand is an interchange floating type, the other operand is converted, without change of type domain, to a type whose corresponding real type is that same interchange floating type.
  - Otherwise, if the corresponding real type of either operand is **long double**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **long double**.
  - Otherwise, if the corresponding real type of either operand is **double**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **double**.<sup>458)</sup>
  - Otherwise, if the corresponding real type of either operand is **\_Float128x** or **\_Decimal128x**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **\_Float128x** or **\_Decimal128x**, respectively.
  - Otherwise, if the corresponding real type of either operand is **\_Float64x** or **\_Decimal64x**, the other operand is converted, without change of type domain, to a type whose corresponding real type is **\_Float64x** or **\_Decimal64x**, respectively.

#### H.4.4 Arithmetic and non-arithmetic formats

- 1 The operation binding in F.3 for the ISO/IEC 60559 `convertFormat` operation applies to ISO/IEC 60559 arithmetic and non-arithmetic formats as follows:

- For conversions between arithmetic formats supported by floating types (same or different radix) – casts and implicit conversions.
- For same-radix conversions between non-arithmetic interchange formats – encoding-to-encoding conversion functions (H.11.4.3).
- For conversions between non-arithmetic interchange formats (same or different radix) – compositions of string-from-encoding functions (H.12.4) (converting exactly) and string-to-encoding functions (H.12.5).
- For same-radix conversions from interchange formats supported by interchange floating types to non-arithmetic interchange formats – compositions of encode functions (H.11.4.2.2, 7.12.17.2, 7.12.17.4) and encoding-to-encoding functions (H.11.4.3).
- For same radix conversions from non-arithmetic interchange formats to interchange formats supported by interchange floating types – compositions of encoding-to-encoding conversion functions (H.11.4.3) and decode functions (H.11.4.2.3, 7.12.17.3, 7.12.17.5). See the example in H.11.4.3.2.
- For conversions from non-arithmetic interchange formats to arithmetic formats supported by floating types (same or different radix) – compositions of string-from-encoding functions (H.12.4) (converting exactly) and numeric conversion functions **strtod**, etc. (7.24.2.6, 7.24.2.7). See the example in H.12.3.
- For conversions from arithmetic formats supported by floating types to non-arithmetic interchange formats (same or different radix) – compositions of numeric conversion functions **strfromd**, etc. (7.24.2.4, 7.24.2.5) (converting exactly) and string-to-encoding functions (H.12.5).

<sup>458)</sup>All cases where **float** is expected to have the same format as another type are covered in the preceding paragraphs.

## H.5 Lexical Elements

### H.5.1 Keywords

- 1 This subclause expands the list of keywords (6.4.2) to also include:

- **\_Float***N*, where *N* is 16, 32, 64, or  $\geq 128$  and a multiple of 32
- **\_Float32x**
- **\_Float64x**
- **\_Float128x**
- **\_Decimal***N*, where *N* is 96 or  $> 128$  and a multiple of 32
- **\_Decimal64x**
- **\_Decimal128x**

### H.5.2 Literals

- 1 This subclause specifies literals of interchange and extended floating types.
- 2 This subclause expands real-floating-suffix (6.4.5.3) to also include: **f***N*, **FN**, **fNx**, **FNx**, **d***N*, **DN**, **dNx**, or **DNx**.
- 3 A floating suffix **d***N*, **DN**, **dNx**, or **DNx** shall not be used in a hexadecimal-floating-literal.
- 4 A floating suffix shall not designate a type that the implementation does not provide.
- 5 If a floating literal is suffixed by **f***N* or **FN**, it has type **\_Float***N*. If suffixed by **fNx** or **FNx**, it has type **\_FloatNx**. If suffixed by **d***N* or **DN**, it has type **\_Decimal***N*. If suffixed by **dNx** or **DNx**, it has type **\_DecimalNx**.
- 6 The quantum exponent of a floating literal of decimal floating type is the same as for the result value of the corresponding **strtod***N* or **strtodNx** function (H.12.3) for the same numeric string.
- 7 NOTE 1 For *N* = 32, 64, and 128, the suffixes **d***N* and **DN** in this subclause for literals of type **\_Decimal***N* are equivalent alternatives to the suffixes **df**, **dd**, **dL**, **DF**, **DD**, and **DL** in 6.4.5.3 for the same types.
- 8 NOTE 2 It follows from here and from 6.4.5.3 that the valid combinations of such an expanded real floating suffix with a complex suffix correspond to the supported complex types as described in H.2.5.

## H.6 Expressions

- 1 This subclause expands the specification of expressions to also cover interchange and extended floating types.
- 2 Operators involving operands of interchange or extended floating type are evaluated according to the semantics of ISO/IEC 60559, including production of decimal floating-point results with the preferred quantum exponent as specified in ISO/IEC 60559 (see 5.3.5.3.4).
- 3 For multiplicative operators (6.5.6), additive operators (6.5.7), relational operators (6.5.9), equality operators (6.5.10), and compound assignment operators (6.5.17.3), if either operand has decimal floating type, the other operand shall not have standard floating type, binary floating type, or complex type.
- 4 For conditional operators (6.5.16), if the second or third operand has decimal floating type, the other of those operands shall not have standard floating type, binary floating type, or complex type.
- 5 The equivalence of expressions noted in F.9.3 apply to expressions of binary floating types, as well as standard floating types.

## H.7 Declarations

- 1 This subclause expands the list of type specifiers (6.7.3) to also include:

- **\_Float***N*, where *N* is 16, 32, 64, or  $\geq 128$  and a multiple of 32

- **\_Float32x**
  - **\_Float64x**
  - **\_Float128x**
  - **\_DecimalN**, where  $N$  is 96 or  $> 128$  and a multiple of 32
  - **\_Decimal64x**
  - **\_Decimal128x**
- 2 The type specifiers **\_FloatN** (where  $N$  is 16, 32, 64, or  $\geq 128$  and a multiple of 32), **\_Float32x**, **\_Float64x**, **\_Float128x**, **\_DecimalN** (where  $N$  is 96 or  $> 128$  and a multiple of 32), **\_Decimal64x**, and **\_Decimal128x** shall not be used if the implementation does not support the corresponding types (see 6.10.10.4 and H.2).
- 3 This subclause also expands the list under Constraints in 6.7.3 to also include:
- **\_FloatN**, where  $N$  is 16, 32, 64, or  $\geq 128$  and a multiple of 32
  - **\_Float32x**
  - **\_Float64x**
  - **\_Float128x**
  - **\_DecimalN**, where  $N$  is 96 or  $> 128$  and a multiple of 32
  - **\_Decimal64x**
  - **\_Decimal128x**
  - **\_FloatN \_Complex**, where  $N$  is 16, 32, 64, or  $\geq 128$  and a multiple of 32
  - **\_Float32x \_Complex**
  - **\_Float64x \_Complex**
  - **\_Float128x \_Complex**

## H.8 Identifiers in standard headers

- 1 The identifiers added to library headers by this annex are defined or declared by their respective headers only if the macro **\_\_STDC\_WANT\_IEC\_60559\_TYPES\_EXT\_\_** is defined (by the user) at the point in the code where the appropriate header is first included.

## H.9 Complex arithmetic <complex.h>

- 1 This subclause specifies complex functions for corresponding real types that are binary floating types.
- 2 Each function synopsis in 7.3 specifies a family of functions including a principal function with one or more **double complex** parameters and a **double complex** or **double** return value. This subclause expands the synopsis to also include other functions, with the same name as the principal function but with **fN** and **fNx** suffixes, which are corresponding functions whose parameters and return values have corresponding real types **\_FloatN** and **\_FloatNx**.
- 3 The following function prototypes are added to the synopses of the respective subclauses in 7.3. For each binary floating type that the implementation provides, <complex.h> shall declare the associated functions (see H.8). Conversely, for each such type that the implementation does not provide, <complex.h> shall not declare the associated functions.

### 7.3.5 Trigonometric functions



```

_FloatN complex cacosfN(_FloatN complex z);
_FloatNx complex cacosfNx(_FloatNx complex z);
_FloatN complex casinfN(_FloatN complex z);
_FloatNx complex casinfNx(_FloatNx complex z);
_FloatN complex catanfN(_FloatN complex z);
_FloatNx complex catanfNx(_FloatNx complex z);
_FloatN complex ccosfN(_FloatN complex z);
_FloatNx complex ccosfNx(_FloatNx complex z);
_FloatN complex csinfN(_FloatN complex z);
_FloatNx complex csinfNx(_FloatNx complex z);
_FloatN complex ctanfN(_FloatN complex z);
_FloatNx complex ctanfNx(_FloatNx complex z);

```

### 7.3.6 Hyperbolic functions

```

_FloatN complex cacoshfN(_FloatN complex z);
_FloatNx complex cacoshfNx(_FloatNx complex z);
_FloatN complex casinhfN(_FloatN complex z);
_FloatNx complex casinhfNx(_FloatNx complex z);
_FloatN complex catanhfN(_FloatN complex z);
_FloatNx complex catanhfNx(_FloatNx complex z);
_FloatN complex ccoshfN(_FloatN complex z);
_FloatNx complex ccoshfNx(_FloatNx complex z);
_FloatN complex csinhfN(_FloatN complex z);
_FloatNx complex csinhfNx(_FloatNx complex z);
_FloatN complex ctanhfN(_FloatN complex z);
_FloatNx complex ctanhfNx(_FloatNx complex z);

```

### 7.3.7 Exponential and logarithmic functions

```

_FloatN complex cexpfN(_FloatN complex z);
_FloatNx complex cexpfNx(_FloatNx complex z);
_FloatN complex clogfN(_FloatN complex z);
_FloatNx complex clogfNx(_FloatNx complex z);

```

### 7.3.8 Power and absolute value functions

```

_FloatN cabsfN(_FloatN complex z);
_FloatNx cabsfNx(_FloatNx complex z);
_FloatN complex cpowfN(_FloatN complex x, _FloatN complex y);
_FloatNx complex cpowfNx(_FloatNx complex x, _FloatNx complex y);
_FloatN complex csqrtfN(_FloatN complex z);
_FloatNx complex csqrtfNx(_FloatNx complex z);

```

### 7.3.9 Manipulation functions

```

_FloatN cargfN(_FloatN complex z);
_FloatNx cargfNx(_FloatNx complex z);
_FloatN cimagfN(_FloatN complex z);
_FloatNx cimagfNx(_FloatNx complex z);
_FloatN complex CMPLXFN(_FloatN x, _FloatN y);
_FloatNx complex CMPLXFNx(_FloatNx x, _FloatNx y);
_FloatN complex conjfN(_FloatN complex z);
_FloatNx complex conjfNx(_FloatNx complex z);
_FloatN complex cprojfN(_FloatN complex z);
_FloatNx complex cprojfNx(_FloatNx complex z);
_FloatN crealfN(_FloatN complex z);
_FloatNx crealfNx(_FloatNx complex z);

```



- 4 For the functions listed in “future library directions” for `<complex.h>` (7.34.2), the possible suffixes are expanded to also include `fN` and `fNx`.

## H.10 Floating-point environment

- 1 This subclause broadens the effects of the floating-point environment (7.6) to apply to types and formats specified in this annex.
- 2 The same floating-point status flags are used by floating-point operations for all floating types, including those types introduced in this annex, and by conversions for ISO/IEC 60559 non-arithmetic interchange formats.
- 3 Both the dynamic rounding direction mode accessed by `fegetround` and `fesetround` and the `FENV_ROUND` rounding control pragma apply to operations for binary floating types, as well as for standard floating types, and also to conversions for radix-2 non-arithmetic interchange formats. Likewise, both the dynamic rounding direction mode accessed by `fe_dec_getround` and `fe_dec_setround` and the `FENV_DEC_ROUND` rounding control pragmas apply to operations for all the decimal floating types, including those decimal floating types introduced in this annex, and to conversions for radix-10 non-arithmetic interchange formats.
- 4 Table 7.1, which describes functions affected by constant rounding modes for standard floating types, applies also for binary floating types. Each `<math.h>` function family listed in the table indicates the family of functions of all standard and binary floating types (for example, the `acos` family includes `acosf`, `acosl`, `acosfN`, and `acosfNx` as well as `acos`). The `fMencfN`, `strfromencfN`, and `strtoencfN` functions are also affected by these constant rounding modes.
- 5 Table 7.2, which describes functions affected by constant rounding modes for decimal floating types, each `<math.h>` function family indicates the family of functions of all decimal floating types (for example, the `acos` family includes `acosdN` and `acosdNx`). The `dMencbindN`, `dMencdecN`, `strfromencbindN`, `strfromencdecN`, `strtoencbindN`, and `strtoencdecN` functions are also affected by these constant rounding modes.

## H.11 Mathematics `<math.h>`

### H.11.1 General

- 1 This subclause specifies types, functions, and macros for interchange and extended floating types, generally corresponding to those specified in 7.12 and F.10.
- 2 All classification macros (7.12.4) and comparison macros (7.12.18) naturally extend to handle interchange and extended floating types. For comparison macros, if neither of the sets of values of the argument formats is a subset of (or equivalent to) the other, the behavior is undefined.
- 3 This subclause also specifies encoding conversion functions that are part of support for the non-arithmetic interchange formats in ISO/IEC 60559 (see H.2.3).
- 4 Most function synopses in 7.12 specify a family of functions including a principal function with one or more `double` parameters, a `double` return value, or both. The synopses are expanded to also include functions with the same name as the principal function but with `fN`, `fNx`, `dN`, and `dNx` suffixes, which are corresponding functions whose parameters, return values, or both are of types `_FloatN`, `_FloatNx`, `_DecimalN`, and `_DecimalNx`, respectively.
- 5 For each interchange or extended floating type that the implementation provides, `<math.h>` shall define the associated types and macros and declare the associated functions (see H.8). Conversely, for each such type that the implementation does not provide, `<math.h>` shall not define the associated types and macros or declare the associated functions unless explicitly specified otherwise.
- 6 With the types

```
float_t
double_t
```

in 7.12 are included the type

**long\_double\_t**

and for each supported type **\_FloatN**, the type

**\_FloatN\_t**

and for each supported type **\_DecimalN**, the type

**\_DecimalN\_t**

These are floating types, such that:

- the values of **long double** are a subset of the values of **long\_double\_t**;
- the values of **\_FloatN** are a subset of the values of **\_FloatN\_t**;
- the values of **\_DecimalN** are a subset of the values of **\_DecimalN\_t**;
- the values of **double\_t** are a subset of the values of **long\_double\_t**;
- the values of **\_FloatM\_t** are a subset of the values of **\_FloatN\_t** if  $M < N$ ;
- the values of **\_DecimalM\_t** are a subset of the values of **\_DecimalN\_t** if  $M < N$ .

If **FLT\_RADIX** is 2 and **FLT\_EVAL\_METHOD** (H.3) is nonnegative, then each of the types corresponding to a standard or binary floating type is the type whose range and precision are specified by **FLT\_EVAL\_METHOD** to be used for evaluating operations and literals of that standard or binary floating type. If **DEC\_EVAL\_METHOD** (H.3) is nonnegative, then each of the types corresponding to a decimal floating type is the type whose range and precision are specified by **DEC\_EVAL\_METHOD** to be used for evaluating operations and literals of that decimal floating type.

7 EXAMPLE If the supported standard and binary floating types are as shown in Table H.6:

Table H.6 — Example supported types

| Type                               | ISO/IEC 60559 format     |
|------------------------------------|--------------------------|
| <b>_Float16</b>                    | binary16                 |
| <b>float, _Float32</b>             | binary32                 |
| <b>double, _Float64, _Float32x</b> | binary64                 |
| <b>long double, _Float64x</b>      | 80-bit binary64-extended |
| <b>_Float128</b>                   | binary128                |

then Table H.7 gives the types with **\_t** suffixes for various values for a **FLT\_EVAL\_METHOD** of a given value  $m$ :

Table H.7 — **\_t** type (vertical) vs.  $m$  (horizontal) relation

| <b>_t</b> type       | $m$              |                  |                  |                  |
|----------------------|------------------|------------------|------------------|------------------|
|                      | 0                | 1                | 2                | 32               |
| <b>_Float16_t</b>    | float            | double           | long double      | <b>_Float32</b>  |
| <b>float_t</b>       | float            | double           | long double      | float            |
| <b>_Float32_t</b>    | <b>_Float32</b>  | double           | long double      | <b>_Float32</b>  |
| <b>double_t</b>      | double           | double           | long double      | double           |
| <b>_Float64_t</b>    | <b>_Float64</b>  | <b>_Float64</b>  | long double      | <b>_Float64</b>  |
| <b>long_double_t</b> | long double      | long double      | long double      | long double      |
| <b>_Float128_t</b>   | <b>_Float128</b> | <b>_Float128</b> | <b>_Float128</b> | <b>_Float128</b> |

| <b>_t</b> type       | <i>m</i>           |                  |                    |                    |
|----------------------|--------------------|------------------|--------------------|--------------------|
|                      | 64                 | 128              | 33                 | 65                 |
| <b>_Float16_t</b>    | <b>_Float64</b>    | <b>_Float128</b> | <b>_Float32x</b>   | <b>_Float64x</b>   |
| <b>float_t</b>       | <b>_Float64</b>    | <b>_Float128</b> | <b>_Float32x</b>   | <b>_Float64x</b>   |
| <b>_Float32_t</b>    | <b>_Float64</b>    | <b>_Float128</b> | <b>_Float32x</b>   | <b>_Float64x</b>   |
| <b>double_t</b>      | <b>double</b>      | <b>_Float128</b> | <b>double</b>      | <b>_Float64x</b>   |
| <b>_Float64_t</b>    | <b>_Float64</b>    | <b>_Float128</b> | <b>_Float64</b>    | <b>_Float64x</b>   |
| <b>long_double_t</b> | <b>long double</b> | <b>_Float128</b> | <b>long double</b> | <b>long double</b> |
| <b>_Float128_t</b>   | <b>_Float128</b>   | <b>_Float128</b> | <b>_Float128</b>   | <b>_Float128</b>   |

## H.11.2 Macros

- 1 This subclause adds macros in 7.12 as follows.
- 2 The macros

```
HUGE_VAL_FN
HUGE_VAL_DN
HUGE_VAL_FNX
HUGE_VAL_DNX
```

expand to constant expressions of types **\_FloatN**, **\_DecimalN**, **\_FloatNx**, and **\_DecimalNx**, respectively, representing positive infinity.

- 3 The macros

```
FP_FAST_FMAFN
FP_FAST_FMADN
FP_FAST_FMAFNX
FP_FAST_FMAFNX
```

are, respectively, **\_FloatN**, **\_DecimalN**, **\_FloatNx**, and **\_DecimalNx** analogues of **FP\_FAST\_FMA**.

- 4 The macros in the following lists are interchange and extended floating type analogues of **FP\_FAST\_FADD**, **FP\_FAST\_FADDL**, **FP\_FAST\_DADDL**, etc.
- 5 For  $M < N$ , the macros

```
FP_FAST_FMADDFN
FP_FAST_FMSUBFN
FP_FAST_FMMULFN
FP_FAST_FMDIVFN
FP_FAST_FMFMAFN
FP_FAST_FMSQRTFN
FP_FAST_DMADDDN
FP_FAST_DMSUBDN
FP_FAST_DMMULDN
FP_FAST_DMDIVDN
FP_FAST_DMFADN
FP_FAST_DMSQRTDN
```

characterize the corresponding functions whose arguments are of an interchange floating type of width  $N$  and whose return type is an interchange floating type of width  $M$ .

- 6 For  $M \leq N$ , the macros

```
FP_FAST_FMADDFNX
FP_FAST_FMSUBFNX
FP_FAST_FMMULFNX
FP_FAST_FMDIVFNX
FP_FAST_FMFMAFNX
FP_FAST_FMSQRTFNX
```

```

FP_FAST_DMADDDNX
FP_FAST_DMSUBDNX
FP_FAST_DMMULDNX
FP_FAST_DMDIVDNX
FP_FAST_DMFMADNX
FP_FAST_DMSQRDNX

```

characterize the corresponding functions whose arguments are of an extended floating type that extends a format of width  $N$  and whose return type is an interchange floating type of width  $M$ .

- 7 For  $M < N$ , the macros

```

FP_FAST_FMXADDFN
FP_FAST_FMXSUBFN
FP_FAST_FMXMULFN
FP_FAST_FMXDIVFN
FP_FAST_FMXFMADN
FP_FAST_FMXSQRTFN
FP_FAST_DMxADDDN
FP_FAST_DMXSUBDN
FP_FAST_DMXMULDN
FP_FAST_DMXDIVDN
FP_FAST_DMXFMADN
FP_FAST_DMXSQRTDN

```

characterize the corresponding functions whose arguments are of an interchange floating type of width  $N$  and whose return type is an extended floating type that extends a format of width  $M$ .

- 8 For  $M < N$ , the macros

```

FP_FAST_FMXADDFNX
FP_FAST_FMXSUBFNX
FP_FAST_FMXMULFNX
FP_FAST_FMXDIVFNX
FP_FAST_FMXFMADNX
FP_FAST_FMXSQRTFNX
FP_FAST_DMxADDDNX
FP_FAST_DMXSUBDNX
FP_FAST_DMXMULDNX
FP_FAST_DMXDIVDNX
FP_FAST_DMXFMADNX
FP_FAST_DMXSQRTDNX

```

characterize the corresponding functions whose arguments are of an extended floating type that extends a format of width  $N$  and whose return type is an extended floating type that extends a format of width  $M$ .

### H.11.3 Functions

- 1 This subclause adds the following functions to the synopses of the respective subclauses in 7.12.

#### 7.12.5 Trigonometric functions

```

_FloatN acosfN(_FloatN x);
_FloatNx acosfNx(_FloatNx x);
_DecimalN acosdN(_DecimalN x);
_DecimalNx acosdNx(_DecimalNx x);

_FloatN asinfN(_FloatN x);
_FloatNx asinfNx(_FloatNx x);
_DecimalN asindN(_DecimalN x);
_DecimalNx asindNx(_DecimalNx x);

```

```

_FloatN atanfN(_FloatN x);
_FloatNx atanfNx(_FloatNx x);
_DecimalN atandN(_DecimalN x);
_DecimalNx atandNx(_DecimalNx x);

_FloatN atan2fN(_FloatN y, _FloatN x);
_FloatNx atan2fNx(_FloatNx y, _FloatNx x);
_DecimalN atan2dN(_DecimalN y, _DecimalN x);
_DecimalNx atan2dNx(_DecimalNx y, _DecimalNx x);

_FloatN cosfN(_FloatN x);
_FloatNx cosfNx(_FloatNx x);
_DecimalN cosdN(_DecimalN x);
_DecimalNx cosdNx(_DecimalNx x);

_FloatN sinfN(_FloatN x);
_FloatNx sinfNx(_FloatNx x);
_DecimalN sindN(_DecimalN x);
_DecimalNx sindNx(_DecimalNx x);

_FloatN tanfN(_FloatN x);
_FloatNx tanfNx(_FloatNx x);
_DecimalN tandN(_DecimalN x);
_DecimalNx tandNx(_DecimalNx x);

_FloatN acospiN(_FloatN x);
_FloatNx acospiNx(_FloatNx x);
_DecimalN acospidN(_DecimalN x);
_DecimalNx acospidNx(_DecimalNx x);

_FloatN asinpiN(_FloatN x);
_FloatNx asinpiNx(_FloatNx x);
_DecimalN asinpidN(_DecimalN x);
_DecimalNx asinpidNx(_DecimalNx x);

_FloatN atanpiN(_FloatN x);
_FloatNx atanpiNx(_FloatNx x);
_DecimalN atanpidN(_DecimalN x);
_DecimalNx atanpidNx(_DecimalNx x);

_FloatN atan2piN(_FloatN y, _FloatN x);
_FloatNx atan2piNx(_FloatNx y, _FloatNx x);
_DecimalN atan2pidN(_DecimalN y, _DecimalN x);
_DecimalNx atan2pidNx(_DecimalNx y, _DecimalNx x);

_FloatN cospifN(_FloatN x);
_FloatNx cospifNx(_FloatNx x);
_DecimalN cospidN(_DecimalN x);
_DecimalNx cospidNx(_DecimalNx x);

_FloatN sinpiN(_FloatN x);
_FloatNx sinpiNx(_FloatNx x);
_DecimalN sinpidN(_DecimalN x);
_DecimalNx sinpidNx(_DecimalNx x);

_FloatN tanpiN(_FloatN x);
_FloatNx tanpiNx(_FloatNx x);
_DecimalN tanpidN(_DecimalN x);
_DecimalNx tanpidNx(_DecimalNx x);

```

## 7.12.6 Hyperbolic functions

```

_FloatN acoshfN(_FloatN x);
_FloatNx acoshfNx(_FloatNx x);
_DecimalN acoshdN(_DecimalN x);
_DecimalNx acoshdNx(_DecimalNx x);

_FloatN asinhfN(_FloatN x);
_FloatNx asinhfNx(_FloatNx x);
_DecimalN asinhdN(_DecimalN x);
_DecimalNx asinhdNx(_DecimalNx x);

_FloatN atanhfN(_FloatN x);
_FloatNx atanhfNx(_FloatNx x);
_DecimalN atanhdN(_DecimalN x);
_DecimalNx atanhdNx(_DecimalNx x);

_FloatN coshfN(_FloatN x);
_FloatNx coshfNx(_FloatNx x);
_DecimalN coshdN(_DecimalN x);
_DecimalNx coshdNx(_DecimalNx x);

_FloatN sinhN(_FloatN x);
_FloatNx sinhNx(_FloatNx x);
_DecimalN sinhN(_DecimalN x);
_DecimalNx sinhNx(_DecimalNx x);

_FloatN tanhfN(_FloatN x);
_FloatNx tanhfNx(_FloatNx x);
_DecimalN tanhdN(_DecimalN x);
_DecimalNx tanhdNx(_DecimalNx x);

```

## 7.12.7 Exponential and logarithmic functions

```

_FloatN expfN(_FloatN x);
_FloatNx expfNx(_FloatNx x);
_DecimalN expdN(_DecimalN x);
_DecimalNx expdNx(_DecimalNx x);

_FloatN expl0fN(_FloatN x);
_FloatNx expl0fNx(_FloatNx x);
_DecimalN expl0dN(_DecimalN x);
_DecimalNx expl0dNx(_DecimalNx x);

_FloatN expl0m1fN(_FloatN x);
_FloatNx expl0m1fNx(_FloatNx x);
_DecimalN expl0m1dN(_DecimalN x);
_DecimalNx expl0m1dNx(_DecimalNx x);

_FloatN exp2fN(_FloatN x);
_FloatNx exp2fNx(_FloatNx x);
_DecimalN exp2dN(_DecimalN x);
_DecimalNx exp2dNx(_DecimalNx x);

_FloatN exp2m1fN(_FloatN x);
_FloatNx exp2m1fNx(_FloatNx x);
_DecimalN exp2m1dN(_DecimalN x);
_DecimalNx exp2m1dNx(_DecimalNx x);

_FloatN expm1fN(_FloatN x);
_FloatNx expm1fNx(_FloatNx x);

```

```

_DecimalN expm1dN(_DecimalN x);
_DecimalNx expm1dNx(_DecimalNx x);

_FloatN frexpfN(_FloatN value, int *exp);
_FloatNx frexpfNx(_FloatNx value, int *exp);
_DecimalN frexpdN(_DecimalN value, int *exp);
_DecimalNx frexpdNx(_DecimalNx value, int *exp);

int ilogbfN(_FloatN x);
int ilogbfNx(_FloatNx x);
int ilogbdN(_DecimalN x);
int ilogbdNx(_DecimalNx x);

_FloatN ldexpfN(_FloatN value, int exp);
_FloatNx ldexpfNx(_FloatNx value, int exp);
_DecimalN ldexpdN(_DecimalN value, int exp);
_DecimalNx ldexpdNx(_DecimalNx value, int exp);

long int llogbfN(_FloatN x);
long int llogbfNx(_FloatNx x);
long int llogbdN(_DecimalN x);
long int llogbdNx(_DecimalNx x);

_FloatN logfN(_FloatN x);
_FloatNx logfNx(_FloatNx x);
_DecimalN logdN(_DecimalN x);
_DecimalNx logdNx(_DecimalNx x);

_FloatN log10fN(_FloatN x);
_FloatNx log10fNx(_FloatNx x);
_DecimalN log10dN(_DecimalN x);
_DecimalNx log10dNx(_DecimalNx x);

_FloatN log10p1fN(_FloatN x);
_FloatNx log10p1fNx(_FloatNx x);
_DecimalN log10p1dN(_DecimalN x);
_DecimalNx log10p1dNx(_DecimalNx x);

_FloatN log1pfN(_FloatN x);
_FloatNx log1pfNx(_FloatNx x);
_FloatN logp1fN(_FloatN x);
_FloatNx logp1fNx(_FloatNx x);
_DecimalN log1pdN(_DecimalN x);
_DecimalNx log1pdNx(_DecimalNx x);
_DecimalN logp1dN(_DecimalN x);
_DecimalNx logp1dNx(_DecimalNx x);

_FloatN log2fN(_FloatN x);
_FloatNx log2fNx(_FloatNx x);
_DecimalN log2dN(_DecimalN x);
_DecimalNx log2dNx(_DecimalNx x);

_FloatN log2p1fN(_FloatN x);
_FloatNx log2p1fNx(_FloatNx x);
_DecimalN log2p1dN(_DecimalN x);
_DecimalNx log2p1dNx(_DecimalNx x);

_FloatN logbfN(_FloatN x);
_FloatNx logbfNx(_FloatNx x);
_DecimalN logbdN(_DecimalN x);
_DecimalNx logbdNx(_DecimalNx x);

```

```

_FloatN modffN(_FloatN x, _FloatN *iptr);
_FloatNx modffNx(_FloatNx x, _FloatNx *iptr);
_DecimalN modfdN(_DecimalN x, _DecimalN *iptr);
_DecimalNx modfdNx(_DecimalNx x, _DecimalNx *iptr);

_FloatN scalbnfN(_FloatN value, int exp);
_FloatNx scalbnfNx(_FloatNx value, int exp);
_DecimalN scalbndN(_DecimalN value, int exp);
_DecimalNx scalbndNx(_DecimalNx value, int exp);

_FloatN scalblnfN(_FloatN value, long int exp);
_FloatNx scalblnfNx(_FloatNx value, long int exp);
_DecimalN scalblndN(_DecimalN value, long int exp);
_DecimalNx scalblndNx(_DecimalNx value, long int exp);

```

### 7.12.8 Power and absolute-value functions

```

_FloatN cbrtfN(_FloatN x);
_FloatNx cbrtfNx(_FloatNx x);
_DecimalN cbrtdN(_DecimalN x);
_DecimalNx cbrtdNx(_DecimalNx x);

_FloatN compoundnfN(_FloatN x, long long int n);
_FloatNx compoundnfNx(_FloatNx x, long long int n);
_DecimalN compoundndN(_DecimalN x, long long int n);
_DecimalNx compoundndNx(_DecimalNx x, long long int n);

_FloatN fabsfN(_FloatN x);
_FloatNx fabsfNx(_FloatNx x);
_DecimalN fabsdN(_DecimalN x);
_DecimalNx fabsdNx(_DecimalNx x);

_FloatN hypotfN(_FloatN x, _FloatN y);
_FloatNx hypotfNx(_FloatNx x, _FloatNx y);
_DecimalN hypotdN(_DecimalN x, _DecimalN y);
_DecimalNx hypotdNx(_DecimalNx x, _DecimalNx y);

_FloatN powfN(_FloatN x, _FloatN y);
_FloatNx powfNx(_FloatNx x, _FloatNx y);
_DecimalN powdN(_DecimalN x, _DecimalN y);
_DecimalNx powdNx(_DecimalNx x, _DecimalNx y);

_FloatN pownfN(_FloatN x, long long int n);
_FloatNx pownfNx(_FloatNx x, long long int n);
_DecimalN powndN(_DecimalN x, long long int n);
_DecimalNx powndNx(_DecimalNx x, long long int n);

_FloatN powrfN(_FloatN x, _FloatN y);
_FloatNx powrfNx(_FloatNx x, _FloatNx y);
_DecimalN powrdN(_DecimalN x, _DecimalN y);
_DecimalNx powrdNx(_DecimalNx x, _DecimalNx y);

_FloatN rootnfN(_FloatN x, long long int n);
_FloatNx rootnfNx(_FloatNx x, long long int n);
_DecimalN rootndN(_DecimalN x, long long int n);
_DecimalNx rootndNx(_DecimalNx x, long long int n);

_FloatN rsqrtfN(_FloatN x);
_FloatNx rsqrtfNx(_FloatNx x);
_DecimalN rsqrtndN(_DecimalN x);

```



```

_DecimalNx rsqrtfN(_DecimalNx x);

_FloatN sqrtfN(_FloatN x);
_FloatNx sqrtfNx(_FloatNx x);
_DecimalN sqrtfN(_DecimalN x);
_DecimalNx sqrtfNx(_DecimalNx x);

```

#### 7.12.9 Error and gamma functions

```

_FloatN erffN(_FloatN x);
_FloatNx erffNx(_FloatNx x);
_DecimalN erfdN(_DecimalN x);
_DecimalNx erfdNx(_DecimalNx x);

_FloatN erfcfN(_FloatN x);
_FloatNx erfcfNx(_FloatNx x);
_DecimalN erfcdN(_DecimalN x);
_DecimalNx erfcdNx(_DecimalNx x);

_FloatN lgammafN(_FloatN x);
_FloatNx lgammafNx(_FloatNx x);
_DecimalN lgammadN(_DecimalN x);
_DecimalNx lgammadNx(_DecimalNx x);

_FloatN tgammafN(_FloatN x);
_FloatNx tgammafNx(_FloatNx x);
_DecimalN tgammafNx(_DecimalN x);
_DecimalNx tgammafNx(_DecimalNx x);

```

#### 7.12.10 Nearest integer functions

```

_FloatN ceilfN(_FloatN x);
_FloatNx ceilfNx(_FloatNx x);
_DecimalN ceildN(_DecimalN x);
_DecimalNx ceildNx(_DecimalNx x);

_FloatN floorfN(_FloatN x);
_FloatNx floorfNx(_FloatNx x);
_DecimalN floordN(_DecimalN x);
_DecimalNx floordNx(_DecimalNx x);

_FloatN nearbyintfN(_FloatN x);
_FloatNx nearbyintfNx(_FloatNx x);
_DecimalN nearbyintdN(_DecimalN x);
_DecimalNx nearbyintdNx(_DecimalNx x);

_FloatN rintfN(_FloatN x);
_FloatNx rintfNx(_FloatNx x);
_DecimalN rintdN(_DecimalN x);
_DecimalNx rintdNx(_DecimalNx x);

long int lrintfN(_FloatN x);
long int lrintfNx(_FloatNx x);
long int lrintdN(_DecimalN x);
long int lrintdNx(_DecimalNx x);

long long int llrintfN(_FloatN x);
long long int llrintfNx(_FloatNx x);
long long int llrintdN(_DecimalN x);
long long int llrintdNx(_DecimalNx x);

```

```

_FloatN roundfN(_FloatN x);
_FloatNx roundfNx(_FloatNx x);
_DecimalN rounddN(_DecimalN x);
_DecimalNx rounddNx(_DecimalNx x);

long int lroundfN(_FloatN x);
long int lroundfNx(_FloatNx x);
long int lrounddN(_DecimalN x);
long int lrounddNx(_DecimalNx x);

long long int llroundfN(_FloatN x);
long long int llroundfNx(_FloatNx x);
long long int llrounddN(_DecimalN x);
long long int llrounddNx(_DecimalNx x);

_FloatN roundevenfN(_FloatN x);
_FloatNx roundevenfNx(_FloatNx x);
_DecimalN roundevendN(_DecimalN x);
_DecimalNx roundevendNx(_DecimalNx x);

_FloatN truncfN(_FloatN x);
_FloatNx truncfNx(_FloatNx x);
_DecimalN truncdN(_DecimalN x);
_DecimalNx truncdNx(_DecimalNx x);

_FloatN fromfpfN(_FloatN x, int rnd, unsigned int width);
_FloatNx fromfpfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN fromfpdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx fromfpdNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN ufromfpfN(_FloatN x, int rnd, unsigned int width);
_FloatNx ufromfpfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN ufromfpdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx ufromfpdNx(_DecimalNx x, int rnd, unsigned int width);

_FloatN fromfpxfN(_FloatN x, int rnd, unsigned int width);
_FloatNx fromfpxfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN fromfpxdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx fromfpxdNx(_DecimalNx x, int rnd, unsigned int width);
_FloatN ufromfpxfN(_FloatN x, int rnd, unsigned int width);
_FloatNx ufromfpxfNx(_FloatNx x, int rnd, unsigned int width);
_DecimalN ufromfpxdN(_DecimalN x, int rnd, unsigned int width);
_DecimalNx ufromfpxdNx(_DecimalNx x, int rnd, unsigned int width);

```

#### 7.12.11.2 Remainder functions

```

_FloatN fmodfN(_FloatN x, _FloatN y);
_FloatNx fmodfNx(_FloatNx x, _FloatNx y);
_DecimalN fmoddN(_DecimalN x, _DecimalN y);
_DecimalNx fmoddNx(_DecimalNx x, _DecimalNx y);

_FloatN remainderfN(_FloatN x, _FloatN y);
_FloatNx remainderfNx(_FloatNx x, _FloatNx y);
_DecimalN remainderdN(_DecimalN x, _DecimalN y);
_DecimalNx remainderdNx(_DecimalNx x, _DecimalNx y);

_FloatN remquofN(_FloatN x, _FloatN y, int *quo);
_FloatNx remquofNx(_FloatNx x, _FloatNx y, int *quo);

```

#### 7.12.12 Manipulation functions

```

_FloatN copysignfN(_FloatN x, _FloatN y);

```

```

_FloatNx copysignfNx(_FloatNx x, _FloatNx y);
_DecimalN copysigndN(_DecimalN x, _DecimalN y);
_DecimalNx copysigndNx(_DecimalNx x, _DecimalNx y);

_FloatN nanfN(const char *tagp);
_FloatNx nanfNx(const char *tagp);
_DecimalN nandN(const char *tagp);
_DecimalNx nandNx(const char *tagp);

_FloatN nextafterfN(_FloatN x, _FloatN y);
_FloatNx nextafterfNx(_FloatNx x, _FloatNx y);
_DecimalN nextafterdN(_DecimalN x, _DecimalN y);
_DecimalNx nextafterdNx(_DecimalNx x, _DecimalNx y);

_FloatN nextupfN(_FloatN x);
_FloatNx nextupfNx(_FloatNx x);
_DecimalN nextupdN(_DecimalN x);
_DecimalNx nextupdNx(_DecimalNx x);

_FloatN nextdownfN(_FloatN x);
_FloatNx nextdownfNx(_FloatNx x);
_DecimalN nextdowndN(_DecimalN x);
_DecimalNx nextdowndNx(_DecimalNx x);

int canonicalizefN(_FloatN *cx, const _FloatN *x);
int canonicalizefNx(_FloatNx *cx, const _FloatNx *x);
int canonicalizedN(_DecimalN *cx, const _DecimalN *x);
int canonicalizedNx(_DecimalNx *cx, const _DecimalNx *x);

```

### 7.12.13 Maximum, minimum, and positive difference functions

```

_FloatN fdimfN(_FloatN x, _FloatN y);
_FloatNx fdimfNx(_FloatNx x, _FloatNx y);
_DecimalN fdimdN(_DecimalN x, _DecimalN y);
_DecimalNx fdimdNx(_DecimalNx x, _DecimalNx y);

_FloatN fmaximumfN(_FloatN x, _FloatN y);
_FloatNx fmaximumfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximumdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximumdNx(_DecimalNx x, _DecimalNx y);

_FloatN fminimumfN(_FloatN x, _FloatN y);
_FloatNx fminimumfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimumdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimumdNx(_DecimalNx x, _DecimalNx y);

_FloatN fmaximum_magfN(_FloatN x, _FloatN y);
_FloatNx fmaximum_magfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_magdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_magdNx(_DecimalNx x, _DecimalNx y);

_FloatN fminimum_magfN(_FloatN x, _FloatN y);
_FloatNx fminimum_magfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_magdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_magdNx(_DecimalNx x, _DecimalNx y);

_FloatN fmaximum_numfN(_FloatN x, _FloatN y);
_FloatNx fmaximum_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_numdNx(_DecimalNx x, _DecimalNx y);

```

```

_FloatN fminimum_numfN(_FloatN x, _FloatN y);
_FloatNx fminimum_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_numdNx(_DecimalNx x, _DecimalNx y);

_FloatN fmaximum_mag_numfN(_FloatN x, _FloatN y);
_FloatNx fmaximum_mag_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fmaximum_mag_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fmaximum_mag_numdNx(_DecimalNx x, _DecimalNx y);

_FloatN fminimum_mag_numfN(_FloatN x, _FloatN y);
_FloatNx fminimum_mag_numfNx(_FloatNx x, _FloatNx y);
_DecimalN fminimum_mag_numdN(_DecimalN x, _DecimalN y);
_DecimalNx fminimum_mag_numdNx(_DecimalNx x, _DecimalNx y);

```

## 7.12.14.1 Fused multiply-add

```

_FloatN fmafN(_FloatN x, _FloatN y, _FloatN z);
_FloatNx fmafNx(_FloatNx x, _FloatNx y, _FloatNx z);
_DecimalN fmadN(_DecimalN x, _DecimalN y, _DecimalN z);
_DecimalNx fmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z);

```

## 7.12.15 Functions that round result to narrower type

```

_FloatM fMaddfN(_FloatN x, _FloatN y); // $M < N$
_FloatM fMaddfNx(_FloatNx x, _FloatNx y); // $M \leq N$
_FloatMx fMxaddfN(_FloatN x, _FloatN y); // $M < N$
_FloatMx fMxaddfNx(_FloatNx x, _FloatNx y); // $M < N$
_DecimalM dMadddN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalM dMadddNx(_DecimalNx x, _DecimalNx y); // $M \leq N$
_DecimalMx dMxadddN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalMx dMxadddNx(_DecimalNx x, _DecimalNx y); // $M < N$

_FloatM fMsubfN(_FloatN x, _FloatN y); // $M < N$
_FloatM fMsubfNx(_FloatNx x, _FloatNx y); // $M \leq N$
_FloatMx fMxsubfN(_FloatN x, _FloatN y); // $M < N$
_FloatMx fMxsubfNx(_FloatNx x, _FloatNx y); // $M < N$
_DecimalM dMsubdN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalM dMsubdNx(_DecimalNx x, _DecimalNx y); // $M \leq N$
_DecimalMx dMxsubdN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalMx dMxsubdNx(_DecimalNx x, _DecimalNx y); // $M < N$

_FloatM fMmulfN(_FloatN x, _FloatN y); // $M < N$
_FloatM fMmulfNx(_FloatNx x, _FloatNx y); // $M \leq N$
_FloatMx fMxmulfN(_FloatN x, _FloatN y); // $M < N$
_FloatMx fMxmulfNx(_FloatNx x, _FloatNx y); // $M < N$
_DecimalM dMmuldN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalM dMmuldNx(_DecimalNx x, _DecimalNx y); // $M \leq N$
_DecimalMx dMxmuldN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalMx dMxmuldNx(_DecimalNx x, _DecimalNx y); // $M < N$

_FloatM fMdivfN(_FloatN x, _FloatN y); // $M < N$
_FloatM fMdivfNx(_FloatNx x, _FloatNx y); // $M \leq N$
_FloatMx fMxdivfN(_FloatN x, _FloatN y); // $M < N$
_FloatMx fMxdivfNx(_FloatNx x, _FloatNx y); // $M < N$
_DecimalM dMdivdN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalM dMdivdNx(_DecimalNx x, _DecimalNx y); // $M \leq N$
_DecimalMx dMxdivdN(_DecimalN x, _DecimalN y); // $M < N$
_DecimalMx dMxdivdNx(_DecimalNx x, _DecimalNx y); // $M < N$

_FloatM fMfmafN(_FloatN x, _FloatN y, _FloatN z); // $M < N$

```

```

_FloatM fMfmafNx(_FloatNx x, _FloatNx y, _FloatNx z); // M ≤ N
_FloatMx fMxfmafN(_FloatN x, _FloatN y, _FloatN z); // M < N
_FloatMx fMxfmafNx(_FloatNx x, _FloatNx y, _FloatNx z); // M < N
_DecimalM dMfmadN(_DecimalN x, _DecimalN y, _DecimalN z); // M < N
_DecimalM dMfmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z); // M ≤ N
_DecimalMx dMxfmadN(_DecimalN x, _DecimalN y, _DecimalN z); // M < N
_DecimalMx dMxfmadNx(_DecimalNx x, _DecimalNx y, _DecimalNx z); // M < N

_FloatM fMsqrtfN(_FloatN x); // M < N
_FloatM fMsqrtfNx(_FloatNx x); // M ≤ N
_FloatMx fMxsqrtfN(_FloatN x); // M < N
_FloatMx fMxsqrtfNx(_FloatNx x); // M < N
_DecimalM dMsqrtN(_DecimalN x); // M < N
_DecimalM dMsqrtNx(_DecimalNx x); // M ≤ N
_DecimalMx dMxsqrtN(_DecimalN x); // M < N
_DecimalMx dMxsqrtNx(_DecimalNx x); // M < N

```

#### 7.12.16 Quantum and quantum exponent functions

```

_DecimalN quantizedN(_DecimalN x, _DecimalN y);
_DecimalNx quantizedNx(_DecimalNx x, _DecimalNx y);

bool samequantumdN(_DecimalN x, _DecimalN y);
bool samequantumdNx(_DecimalNx x, _DecimalNx y);

_DecimalN quantumdN(_DecimalN x);
_DecimalNx quantumdNx(_DecimalNx x);

long long int llquantexpdN(_DecimalN x);
long long int llquantexpdNx(_DecimalNx x);

```

#### 7.12.17 Decimal re-encoding functions

```

void encodedecdN(unsigned char * restrict encptr,
 const _DecimalN * restrict xptr);
void decodedecdN(_DecimalN * restrict xptr,
 const unsigned char * restrict encptr);
void encodebindN(unsigned char * restrict encptr,
 const _DecimalN * restrict xptr);
void decodebindN(_DecimalN * restrict xptr,
 const unsigned char * restrict encptr);

```

#### F.10.13 Total order functions

```

int totalorderfN(const _FloatN *x, const _FloatN *y);
int totalorderfNx(const _FloatNx *x, const _FloatNx *y);
int totalorderdN(const _DecimalN *x, const _DecimalN *y);
int totalorderdNx(const _DecimalNx *x, const _DecimalNx *y);

int totalordermagfN(const _FloatN *x, const _FloatN *y);
int totalordermagfNx(const _FloatNx *x, const _FloatNx *y);
int totalordermagdN(const _DecimalN *x, const _DecimalN *y);
int totalordermagdNx(const _DecimalNx *x, const _DecimalNx *y);

```

#### F.10.14 Payload functions

```

_FloatN getpayloadfN(const _FloatN *x);
_FloatNx getpayloadfNx(const _FloatNx *x);
_DecimalN getpayloaddN(const _DecimalN *x);
_DecimalNx getpayloaddNx(const _DecimalNx *x);

```

```

int setpayloadfN(_FloatN *res, _FloatN pl);
int setpayloadfNx(_FloatNx *res, _FloatNx pl);
int setpayloaddN(_DecimalN *res, _DecimalN pl);
int setpayloaddNx(_DecimalNx *res, _DecimalNx pl);

int setpayloadsigfN(_FloatN *res, _FloatN pl);
int setpayloadsigfNx(_FloatNx *res, _FloatNx pl);
int setpayloadsigdN(_DecimalN *res, _DecimalN pl);
int setpayloadsigdNx(_DecimalNx *res, _DecimalNx pl);

```

- 2 The specification of the **frexp** functions (7.12.7.7) applies to the functions for binary floating types like those for standard floating types: the exponent is an integral power of 2 and, when applicable, **value** equals  $x \times 2^{\text{exp}}$ .
- 3 The specification of the **ldexp** functions (7.12.7.9) applies to the functions for binary floating types like those for standard floating types: they return  $x2^{\text{exp}}$ .
- 4 The specification of the **logb** functions (7.12.7.17) applies to binary floating types, with  $b = 2$ .
- 5 The specification of the **scalbn** and **scalbln** functions (7.12.7.19) applies to binary floating types, with  $b = 2$ .

## H.11.4 Encoding conversion functions

### H.11.4.1 General

- 1 This subclause introduces `<math.h>` functions that, together with the numerical conversion functions for encodings in H.12, support the non-arithmetic interchange formats specified by ISO/IEC 60559. Support for these formats is an optional feature of this annex. Implementations that do not support non-arithmetic interchange formats are not required to declare the functions in this subclause.
- 2 Non-arithmetic interchange formats are not associated with floating types. Arrays of element type **unsigned char** are used as parameters for conversion functions, to represent encodings in interchange formats that may be non-arithmetic formats.

### H.11.4.2 Encode and decode functions

#### H.11.4.2.1 General

- 1 This subclause specifies functions to map representations in binary floating types to and from encodings in **unsigned char** arrays.

#### H.11.4.2.2 The **encodefN** functions

##### Synopsis

- 1

```

#define __STDC_WANT_IEC_60559_TYPES_EXT__
#include <math.h>

void encodefN(unsigned char encptr[restrict static N/8],
 const _FloatN * restrict xptr);

```

##### Description

- 2 The **encodefN** functions convert *\*xptr* into an ISO/IEC 60559 binary*N* encoding and store the resulting encoding as an *N/8* element array, with 8 bits per array element, in the object pointed to by *encptr*. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions preserve the value of *\*xptr* and raise no floating-point exceptions. If *\*xptr* is non-canonical, these functions can produce a canonical encoding.

##### Returns

- 3 The **encodefN** functions return no value.

#### H.11.4.2.3 The **decodefN** functions

## Synopsis

```

1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <math.h>

 void decodefN(_FloatN * restrict xptr,
 const unsigned char encptr[restrict static N/8]);

```

## Description

- 2 The **decodefN** functions interpret the  $N/8$  element array pointed to by **encptr** as an ISO/IEC 60559 binary $N$  encoding, with 8 bits per array element. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2). These functions convert the given encoding into a representation in the type **\_FloatN**, and store the result in the object pointed to by **xptr**. These functions preserve the encoded value and raise no floating-point exceptions. If the encoding is non-canonical, these functions can produce a canonical representation.

## Returns

- 3 The **decodefN** functions return no value.
- 4 See Example in H.11.4.3.2.

## H.11.4.3 Encoding-to-encoding conversion functions

### H.11.4.3.1 General

- 1 An implementation shall declare an **fMencfN** function for each  $M$  and  $N$  equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic binary interchange format,  $M \neq N$ . An implementation shall provide both **dMencdecN** and **dMencbindN** functions for each  $M$  and  $N$  equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic decimal interchange format,  $M \neq N$ .

### H.11.4.3.2 The **fMencfN** functions

## Synopsis

```

1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <math.h>

 void fMencfN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);

```

## Description

- 2 The **fMencfN** functions convert between ISO/IEC 60559 binary interchange formats. These functions interpret the  $N/8$  element array pointed to by **encNptr** as an encoding of width  $N$  bits. They convert the encoding to an encoding of width  $M$  bits and store the resulting encoding as an  $M/8$  element array in the object pointed to by **encMptr**. The conversion rounds and raises floating-point exceptions as specified in ISO/IEC 60559. The order of bytes in the arrays follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

## Returns

- 3 These functions return no value.
- 4 EXAMPLE If the ISO/IEC 60559 binary16 format is supported as a non-arithmetic format, data in binary16 format can be converted to type **float** as follows:

```

 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <math.h>
 unsigned char b16[2]; // for input binary16 datum
 float f; // for result
 unsigned char b32[4];
 _Float32 f32;

```



```

// store input binary16 datum in array b16
...
f32encf16(b32, b16);
decodef32(&f32, b32);
f = f32;
...

```

#### H.11.4.3.3 The **dMencdecdN** and **dMencbindN** functions

##### Synopsis

```

1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <math.h>

 void dMencdecdN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);
 void dMencbindN(unsigned char encMptr[restrict static M/8],
 const unsigned char encNptr[restrict static N/8]);

```

##### Description

- The **dMencdecdN** and **dMencbindN** functions convert between ISO/IEC 60559 decimal interchange formats that use the same encoding scheme. The **dMencdecdN** functions convert between formats using the encoding scheme based on decimal encoding of the significand. The **dMencbindN** functions convert between formats using the encoding scheme based on binary encoding of the significand. These functions interpret the  $N/8$  element array pointed to by **encNptr** as an encoding of width  $N$  bits. They convert the encoding to an encoding of width  $M$  bits and store the resulting encoding as an  $M/8$  element array in the object pointed to by **encMptr**. The conversion rounds and raises floating-point exceptions as specified in ISO/IEC 60559. The order of bytes in the arrays follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

##### Returns

- These functions return no value.

## H.12 Numeric conversion functions <stdlib.h>

### H.12.1 General

- This clause expands the specification of numeric conversion functions in <stdlib.h> (7.24.2) to also include conversions of strings from and to interchange and extended floating types. The conversions from floating are provided by functions analogous to the **strfromd** function. The conversions to floating are provided by functions analogous to the **strtod** function.
- This clause also specifies functions to convert strings from and to ISO/IEC 60559 interchange format encodings.
- For each interchange or extended floating type that the implementation provides, <stdlib.h> shall declare the associated functions specified in the following subclauses in H.12.2 and H.12.3 (see H.8). Conversely, for each such type that the implementation does not provide, <stdlib.h> shall not declare the associated functions.
- For each ISO/IEC 60559 arithmetic or non-arithmetic format that the implementation supports, <stdlib.h> shall declare the associated functions specified the following subclauses in H.12.4 and H.12.5 (see H.8). Conversely, for each such format that the implementation does not provide, <stdlib.h> shall not declare the associated functions.

### H.12.2 String from floating

- This subclause expands 7.24.2.4 and 7.24.2.5 to also include functions for the interchange and extended floating types. It adds to the synopsis in 7.24.2.4 the prototypes

```

int strfromfN(char * restrict s, size_t n,
 const char * restrict format, _FloatN fp);

```



```
int strfromfNx(char * restrict s, size_t n,
 const char * restrict format, _FloatNx fp);
```

It encompasses the prototypes in 7.24.2.5 by replacing them with

```
int strfromdN(char * restrict s, size_t n,
 const char * restrict format, _DecimalW fp);
int strfromdNx(char * restrict s, size_t n,
 const char * restrict format, _DecimalNx fp);
```

- 2 The descriptions and returns for the added functions are analogous to the ones in 7.24.2.4 and 7.24.2.5.

### H.12.3 String to floating

- 1 This subclause expands 7.24.2.6, 7.32.4.2.2, 7.24.2.7, and 7.32.4.2.3 to also include functions for the interchange and extended floating types.
- 2 It adds to the synopsis in 7.24.2.6 the prototypes

```
_FloatN strtofN(const char * restrict nptr,
 char ** restrict endptr);
_FloatNx strtofNx(const char * restrict nptr,
 char ** restrict endptr);
```

It adds to the synopsis in 7.32.4.2.2 the prototypes

```
_FloatW wcstofN(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
_FloatNx wcstofNx(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
```

It encompasses the prototypes in 7.24.2.7 by replacing them with

```
_DecimalW strtodN(const char * restrict nptr,
 char ** restrict endptr);
_DecimalNx strtodNx(const char * restrict nptr,
 char ** restrict endptr);
```

It encompasses the prototypes in 7.32.4.2.3 by replacing them with

```
_DecimalW wcstodN(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
_DecimalNx wcstodNx(const wchar_t * restrict nptr,
 wchar_t ** restrict endptr);
```

- 3 The descriptions and returns for the added functions are analogous to the ones in 7.24.2.6, 7.32.4.2.2, 7.24.2.7, and 7.32.4.2.3.
- 4 EXAMPLE If the ISO/IEC 60559 binary128 format is supported as a non-arithmetic format, data in binary128 format can be converted to type `_Decimal128` as follows:

```
#define __STDC_WANT_IEC_60559_TYPES_EXT__
#include <stdlib.h>
#define MAXSIZE 41 // > intermediate hex string length
 // for the "C" locale
unsigned char b128[16]; // for input binary128 datum
_Decimal128 d128; // for result
char s[MAXSIZE];
// store input binary128 datum in array b128
...
```

```

strfromencf128(s, MAXSIZE, "%a", b128);
d128 = strtod128(s, nullptr);
...

```

Use of "%a" for formatting assures an exact conversion of the value in binary format to character sequence. The value of that character sequence will be correctly rounded to **Decimal128**, as specified previously in this subclause. Assuming a single-byte decimal-point character as in the "C" locale, the array *s* for the output of **strfromencf128** need have no greater size than 41, which is the maximum length of strings of the form

$[-]0xh.h \dots hp \pm d$

where there are up to 29 hexadecimal digits *h* and *d* has 5 digits plus 1 for the null character.

## H.12.4 String from encoding

### H.12.4.1 General

- 1 An implementation shall declare the **strfromencfN** function for each *N* equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic binary interchange format. An implementation shall declare both the **strfromencdecN** and **strfromencbindN** functions for each *N* equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic decimal interchange format.

### H.12.4.2 The **strfromencfN** functions

#### Synopsis

```

1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <stdlib.h>

 int strfromencfN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);

```

#### Description

- 2 The **strfromencfN** functions are similar to the **strfromfN** functions, except the input is the value of the *N*/8 element array pointed to by *encptr*, interpreted as an ISO/IEC 60559 binary *N* encoding. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

#### Returns

- 3 The **strfromencfN** functions return the same values as corresponding **strfromfN** functions.

### H.12.4.3 The **strfromencdecN** and **strfromencbindN** functions

#### Synopsis

```

1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <stdlib.h>

 int strfromencdecN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);
 int strfromencbindN(char * restrict s, size_t n, const char * restrict format,
 const unsigned char encptr[restrict static N/8]);

```

#### Description

- 2 The **strfromencdecN** functions are similar to the **strfromdN** functions except the input is the value of the *N*/8 element array pointed to by *encptr*, interpreted as an ISO/IEC 60559 decimal *N* encoding in the coding scheme based on decimal encoding of the significand. The **strfromencbindN** functions are similar to the **strfromdN** functions except the input is the value of the *N*/8 element array pointed to by *encptr*, interpreted as an ISO/IEC 60559 decimal *N* encoding in the coding scheme based on binary encoding of the significand. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

**Returns**

- 3 The **strfromencdec***N* and **strfromencbind***N* functions return the same values as corresponding **strfrom***N* functions.

**H.12.5 String to encoding****H.12.5.1 General**

- 1 An implementation shall declare the **strtoencf***N* and **wcstoencf***N* functions for each *N* equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic binary interchange format. An implementation shall declare the **strtoencdec***N*, **strtoencbind***N*, **wcstoencdec***N*, and **wcstoencbind***N* functions for each *N* equal to the width of a supported ISO/IEC 60559 arithmetic or non-arithmetic decimal interchange format.

**H.12.5.2 The **strtoencf***N* functions****Synopsis**

```
1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <stdlib.h>

 void strtoencfN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);
```

**Description**

- 2 The **strtoencf***N* functions are similar to the **strtof***N* functions, except they store an ISO/IEC 60559 encoding of the result as an *N*/8 element array in the object pointed to by **encptr**. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

**Returns**

- 3 These functions return no value.

**H.12.5.3 The **wcstoencf***N* functions****Synopsis**

```
1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <wchar.h>

 void wcstoencfN(unsigned char encptr[restrict static N/8],
 const wchar_t * restrict nptr, wchar_t ** restrict endptr);
```

**Description**

- 2 The **wcstoencf***N* functions are similar to the **wcstof***N* functions, except they store an ISO/IEC 60559 encoding of the result as an *N*/8 element array in the object pointed to by **encptr**. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

**Returns**

- 3 These functions return no value.

**H.12.5.4 The **strtoencdec***N* and **strtoencbind***N* functions****Synopsis**

```
1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <stdlib.h>

 void strtoencdecN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);
 void strtoencbindN(unsigned char encptr[restrict static N/8],
 const char * restrict nptr, char ** restrict endptr);
```

**Description**

- 2 The **strtoencdecN** and **strtoencbindN** functions are similar to the **strtodN** functions, except they store an ISO/IEC 60559 encoding of the result as an  $N/8$  element array in the object pointed to by **encptr**. The **strtoencdecN** functions produce an encoding in the encoding scheme based on decimal encoding of the significand. The **strtoencbindN** functions produce an encoding in the encoding scheme based on binary encoding of the significand. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

**Returns**

- 3 These functions return no value.

**H.12.5.5 The **wcstoencdecN** and **wcstoencbindN** functions****Synopsis**

```
1 #define __STDC_WANT_IEC_60559_TYPES_EXT__
 #include <wchar.h>

 void wcstoencdecN(unsigned char encptr[restrict static N/8],
 const wchar_t * restrict nptr, wchar_t ** restrict endptr);
 void wcstoencbindN(unsigned char encptr[restrict static N/8],
 const wchar_t * restrict nptr, wchar_t ** restrict endptr);
```

**Description**

- 2 The **wcstoencdecN** and **wcstoencbindN** functions are similar to the **wcstodN** functions, except they store an ISO/IEC 60559 encoding of the result as an  $N/8$  element array in the object pointed to by **encptr**. The **wcstoencdecN** functions produce an encoding in the encoding scheme based on decimal encoding of the significand. The **wcstoencbindN** functions produce an encoding in the encoding scheme based on binary encoding of the significand. The order of bytes in the array follows the endianness specified with **\_\_STDC\_ENDIAN\_NATIVE\_\_** (7.18.2).

**Returns**

- 3 These functions return no value.

**H.13 Type-generic macros <tgmath.h>**

- 1 This clause enhances the specification of type-generic macros in <tgmath.h> (7.28) to apply to interchange and extended floating types, as well as standard floating types.
- 2 If arguments for generic parameters of a type-generic macro are such that some argument has a corresponding real type that is a standard floating type or a binary floating type and another argument is of decimal floating type, the behavior is undefined.
- 3 The treatment of arguments of integer type in 7.28 is expanded to cases where another argument has extended type. Arguments of integer type are regarded as having type:

- **\_Decimal64x**, if any argument has a decimal extended type; otherwise
- **\_Float32x**, if any argument has a binary extended type; otherwise
- **\_Decimal64**, if any argument has decimal type; otherwise
- **double**

- 4 Use of the macros **carg**, **cimag**, **conj**, **cproj**, or **creal** with any argument of standard floating type, binary floating type, or complex type invokes a complex function. Use of the macro with an argument of a decimal floating type results in undefined behavior.
- 5 The functions that round results to a narrower type have type-generic macros whose names are obtained by omitting any suffix from the function names. Thus, the macros with **f** or **d** prefix are (as in 7.28):

|             |             |              |
|-------------|-------------|--------------|
| <b>fadd</b> | <b>fmul</b> | <b>ffma</b>  |
| <b>dadd</b> | <b>dmul</b> | <b>dfma</b>  |
| <b>fsub</b> | <b>fdiv</b> | <b>fsqrt</b> |
| <b>dsub</b> | <b>ddiv</b> | <b>dsqrt</b> |

and the macros with **fM**, **fMx**, **dM**, or **dMx** prefix are:

|               |                |                |
|---------------|----------------|----------------|
| <b>fMadd</b>  | <b>fMxmUL</b>  | <b>dMfma</b>   |
| <b>fMsub</b>  | <b>fMxdiv</b>  | <b>dMsqrt</b>  |
| <b>fMmul</b>  | <b>fMxfma</b>  | <b>dMxadd</b>  |
| <b>fMdiv</b>  | <b>fMxsqrt</b> | <b>dMxsub</b>  |
| <b>fMfma</b>  | <b>dMadd</b>   | <b>dMxmUL</b>  |
| <b>fMsqrt</b> | <b>dMsub</b>   | <b>dMxdiv</b>  |
| <b>fMxadd</b> | <b>dMmul</b>   | <b>dMxfma</b>  |
| <b>fMxsub</b> | <b>dMdiv</b>   | <b>dMxsqrt</b> |

All arguments are generic. If any argument is not real, use of the macro results in undefined behavior. The following specification uses the notation  $type1 \subseteq type2$  to mean the values of  $type1$  are a subset of (or the same as) the values of  $type2$ . The generic parameter type  $T$  for the function invoked by the macro is determined as follows:

- First, obtain a preliminary type  $P$  for the generic parameters: if all arguments are of integer type, then  $P$  is **double** if the macro prefix is **f**, **d**, **fN**, or **fNx** and  $P$  is **\_Decimal64** if the macro prefix is **dN** or **dNx**; otherwise (if some argument is not of integer type), apply the rules (for determining the corresponding real type of the generic parameters) in 7.28 for macros that do not round result to narrower type, using the usual arithmetic conversion rules in H.4.3, to obtain  $P$ .
- If there exists a corresponding function whose generic parameters have type  $P$ , then  $T$  is  $P$ .
- Otherwise,  $T$  is determined from  $P$  and the macro prefix as follows:
  - For prefix **f**: if  $P$  is a standard or binary floating type, then  $T$  is the first standard floating type of either **double** or **long double**, such that  $P \subseteq T$ , if such a type  $T$  exists. Otherwise (if no such type  $T$  exists or  $P$  is a decimal floating type), the behavior is undefined.
  - For prefix **d**: if  $P$  is a standard or binary floating type, then  $T$  is **long double** if  $P \subseteq \text{long double}$ . Otherwise (if  $P \subseteq \text{long double}$  is **false** or  $P$  is a decimal floating type), the behavior is undefined.
  - For prefix **fM**: if  $P$  is a standard or binary floating type, then  $T$  is **\_FloatN** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported; otherwise  $T$  is **\_FloatNx** for minimum  $N \geq M$  such that  $P \subseteq T$ , if such a type  $T$  is supported. Otherwise (if no such **\_FloatN** or **\_FloatNx** is supported or  $P$  is a decimal floating type), the behavior is undefined.
  - For prefix **fMx**: if  $P$  is a standard or binary floating type, then  $T$  is **\_FloatNx** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported; otherwise  $T$  is **\_FloatN** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported. Otherwise (if no such **\_FloatNx** or **\_FloatN** is supported or  $P$  is a decimal floating type), the behavior is undefined.
  - For prefix **dM**: if  $P$  is a decimal floating type, then  $T$  is **\_DecimalN** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported; otherwise  $T$  is **\_DecimalNx** for minimum  $N \geq M$  such that  $P \subseteq T$ . Otherwise ( $P$  is a standard or binary floating type), the behavior is undefined.
  - For prefix **dMx**: if  $P$  is a decimal floating type, then  $T$  is **\_DecimalNx** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported; otherwise  $T$  is **\_DecimalN** for minimum  $N > M$  such that  $P \subseteq T$ , if such a type  $T$  is supported. Otherwise ( $P$  is a standard or binary floating type), the behavior is undefined.

## 6 EXAMPLE With the declarations

```

#define __STDC_WANT_IEC_60559_TYPES_EXT__

#include <tgmath.h>

int n;
double d;
long double ld;
double complex dc;
_Float32x f32x;
_Float64 f64;
_Float64x f64x;
_Float128 f128;
_Float64x complex f64xc;

```

functions invoked by use of type-generic macros are shown in the following Table H.11 and Table H.12.  $type1 \subseteq type2$  means the values of  $type1$  are a subset of (or the same as) the values of  $type2$ , and  $type1 \subset type2$  means the values of  $type1$  are a strict subset of the values of  $type2$ :

Table H.11 — Type-generic macro resolution

| macro use                  | invokes                                                                                                                                                                                         |
|----------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>cos(f64xc)</code>    | <code>ccosf64x</code>                                                                                                                                                                           |
| <code>pow(dc, f128)</code> | <code>cpowf128</code>                                                                                                                                                                           |
| <code>pow(f64, d)</code>   | <code>powf64</code>                                                                                                                                                                             |
| <code>pow(d, f32x)</code>  | <code>pow</code> , the function, if <code>_Float32x</code> $\subseteq$ <code>double</code> , else <code>powf32x</code> if <code>double</code> $\subset$ <code>_Float32x</code> , else undefined |
| <code>pow(f32, n)</code>   | <code>pow</code> , the function                                                                                                                                                                 |
| <code>pow(f32x, n)</code>  | <code>powf32x</code>                                                                                                                                                                            |

Some type-generic macros that round the result to a narrower type function behave as shown in Table H.12.

Table H.12 — Type-generic narrow rounding examples

| macro use                          | invokes                                                                                                                                                                                         |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>fsub(d, ld)</code>           | <code>fsubl</code>                                                                                                                                                                              |
| <code>dsub(d, f32)</code>          | <code>dsubl</code>                                                                                                                                                                              |
| <code>fmul(dc, d)</code>           | undefined                                                                                                                                                                                       |
| <code>ddiv(ld, f128)</code>        | <code>ddivl</code> if <code>_Float128</code> $\subseteq$ <code>long double</code> , else undefined                                                                                              |
| <code>f32add(f64x, f64)</code>     | <code>f32addf64x</code>                                                                                                                                                                         |
| <code>f32xsqrt(n)</code>           | <code>f32xsqrtf64x</code>                                                                                                                                                                       |
| <code>f32mul(f128, f32x)</code>    | <code>f32mulf128</code> if <code>_Float32x</code> $\subseteq$ <code>_Float128</code> , else <code>f32mulf32x</code> if <code>_Float128</code> $\subset$ <code>_Float32x</code> , else undefined |
| <code>f32fma(f32x, n, f32x)</code> | <code>f32fmaf32x</code>                                                                                                                                                                         |
| <code>f32add(f32, f32)</code>      | <code>f32addf64</code>                                                                                                                                                                          |
| <code>f32xsqrt(f32)</code>         | <code>f32xsqrtf64x</code> , as prior declarations show <code>_Float64x</code> is supported                                                                                                      |
| <code>f64div(f32x, f32x)</code>    | <code>f64divf128</code> if <code>_Float32x</code> $\subseteq$ <code>_Float128</code> , else <code>f64divf64x</code>                                                                             |

## Annex I

(informative)

### Common warnings

#### I.1 Introduction

- 1 An implementation may generate warnings in many situations to help find a source of unintended behavior during the translation or execution of a program. Many such situations are not specified as part of this document.

#### I.2 Common situations

- 1 The following are a few of the common situations where an implementation may generate a warning:
  - A new **struct** or **union** type appears in a function prototype (6.2.1, 6.7.3.4).
  - A block with initialization of an object that has automatic storage duration is jumped into (6.2.4).
  - An implicit narrowing conversion is encountered, such as the assignment of a **long int** or a **double** to an **int**, or a pointer to **void** to a pointer to any type other than a character type (6.3).
  - A hexadecimal floating literal cannot be represented exactly in its evaluation format (6.4.5.3).
  - An ordinary character literal includes more than one character or a wide character literal includes more than one multibyte character (6.4.5.5).
  - The characters `/*` are found in a comment (6.4.8).
  - An “unordered” binary operator (not comma, `&&`, or `| |`) contains a side effect to an lvalue in one operand, and a side effect to, or an access to the value of, the identical lvalue in the other operand (6.5.1).
  - An object is defined but not used (6.7).
  - A value is given to an object of an enumerated type other than by assignment of an enumeration constant that is a member of that type, or an enumeration object that has the same type, or the value of a function that returns the same enumerated type (6.7.3.3).
  - An aggregate has a partly bracketed initialization (6.7.9).
  - A statement cannot be reached (6.8).
  - A statement with no apparent effect is encountered (6.8).
  - A constant expression is used as the controlling expression of a selection statement (6.8.5).
  - An incorrectly formed preprocessing group is encountered while skipping a preprocessing group (6.10.2).
  - An unrecognized **#pragma** directive is encountered (6.10.8).



## Annex J

(informative)

### Portability issues

#### J.1 Unspecified behavior

- 1 The following are unspecified:
  - (1) The manner and timing of static initialization (5.2.2).
  - (2) The termination status returned to the hosted environment if the return type of `main` is not compatible with `int` (5.2.2.3.4).
  - (3) The values of objects that are neither lock-free atomic objects nor of type `volatile sig_atomic_t` and the state of the floating-point environment, when the processing of the abstract machine is interrupted by receipt of a signal (5.2.2.4).
  - (4) The behavior of the display device if a printing character is written when the active position is at the final position of a line (5.3.3).
  - (5) The behavior of the display device if a backspace character is written when the active position is at the initial position of a line (5.3.3).
  - (6) The behavior of the display device if a horizontal tab character is written when the active position is at or past the last defined horizontal tabulation position (5.3.3).
  - (7) The behavior of the display device if a vertical tab character is written when the active position is at or past the last defined vertical tabulation position (5.3.3).
  - (8) How an extended source character that does not correspond to a universal character name counts toward the significant initial characters in an external identifier (5.3.5.2).
  - (9) Many aspects of the representations of types (6.2.6).
  - (10) The value of padding bytes when storing values in structures or unions (6.2.6.1).
  - (11) The values of bytes that correspond to union members other than the one last stored into (6.2.6.1).
  - (12) The representation used when storing a value in an object that has more than one object representation for that value (6.2.6.1).
  - (13) The values of any padding bits in integer representations (6.2.6.2).
  - (14) Whether two string literals result in distinct arrays (6.4.6).
  - (15) The order in which subexpressions are evaluated and the order in which side effects take place, except as specified for the function-call `()`, `&&`, `||`, `?:`, and comma operators (6.5.1).
  - (16) The order in which the function designator, arguments, and subexpressions within the arguments are evaluated in a function call (6.5.3.3).
  - (17) The order of side effects among compound literal initialization list expressions (6.5.3.6).
  - (18) The order in which the operands of an assignment operator are evaluated (6.5.17).
  - (19) The alignment of the addressable storage unit allocated to hold a bit-field (6.7.3.2).
  - (20) Whether a call to an inline function uses the inline definition or the external definition of the function (6.7.5).
  - (21) Whether a size expression is evaluated when it is part of the operand of a `sizeof` operator and changing the value of the size expression would not affect the result of the operator (6.7.7.3).



- (22) The order in which any side effects occur among the initialization list expressions in an initializer (6.7.11).
- (23) The layout of storage for function parameters (6.9.2).
- (24) When a fully expanded macro replacement list contains a function-like macro name as its last preprocessing token and the next preprocessing token from the source file is a (, and the fully expanded replacement of that macro ends with the name of the first macro and the next preprocessing token from the source file is again a (, whether that is considered a nested replacement (6.10.5).
- (25) The order in which # and ## operations are evaluated during macro substitution (6.10.5.3, 6.10.5.4).
- (26) The line number of a preprocessing token, in particular `__LINE__`, that spans multiple physical lines (6.10.6).
- (27) The line number of a preprocessing directive that spans multiple physical lines (6.10.6).
- (28) The line number of a macro invocation that spans multiple physical or logical lines (6.10.6).
- (29) The line number following a directive of the form `#line __LINE__ new-line` (6.10.6).
- (30) The state of the floating-point status flags when execution passes from a part of the program translated with `FENV_ACCESS` “off” to a part translated with `FENV_ACCESS` “on” (7.6.2).
- (31) The order in which `feraiseexcept` raises floating-point exceptions, except as stated in F.8.7 (7.6.5.4).
- (32) Whether `math_errhandling` is a macro or an identifier with external linkage (7.12).
- (33) The results of the `frexp` functions when the specified value is not a floating-point number (7.12.7.7).
- (34) The numeric result of the `ilogb` functions when the correct value is outside the range of the return type (7.12.7.8, F.10.4.8).
- (35) The result of rounding when the value is out of range (7.12.10.5, 7.12.10.7, F.10.7.5).
- (36) The value stored by the `remquo` functions in the object pointed to by `quo` when `y` is zero (7.12.11.3).
- (37) Whether a comparison macro argument that is represented in a format wider than its semantic type is converted to the semantic type (7.12.18).
- (38) Whether `setjmp` is a macro or an identifier with external linkage (7.13).
- (39) Whether `va_copy` and `va_end` are macros or identifiers with external linkage (7.16.2).
- (40) The hexadecimal digit before the decimal point when a non-normalized floating-point number is printed with an `a` or `A` conversion specifier (7.23.6.2, 7.32.2.2).
- (41) The value of the file position indicator after a successful call to the `ungetc` function for a text stream, or the `ungetc` function for any stream, until all pushed-back characters are read or discarded (7.23.7.10, 7.32.3.10).
- (42) The details of the value stored by the `fgetpos` function (7.23.9.1).
- (43) The details of the value returned by the `ftell` function for a text stream (7.23.9.4).
- (44) Whether the `strtod`, `strtof`, `strtold`, `wctod`, `wctof`, and `wctold` functions convert a minus-signed sequence to a negative number directly or by arithmetically negating the value resulting from converting the corresponding unsigned sequence (7.24.2.6, 7.32.4.2.2).

- (45) The order and contiguity of storage allocated by successive calls to the **calloc**, **malloc**, **realloc**, and **aligned\_alloc** functions (7.24.4).
- (46) The amount of storage allocated by a successful call to the **calloc**, **malloc**, **realloc**, or **aligned\_alloc** function when 0 bytes was requested (7.24.4).
- (47) Whether a call to the **atexit** function that does not happen before the **exit** function is called will succeed (7.24.5.2).
- (48) Whether a call to the **at\_quick\_exit** function that does not happen before the **quick\_exit** function is called will succeed (7.24.5.3).
- (49) Which of two elements that compare as equal is matched by the **bsearch** function (7.24.6.2).
- (50) The order of two elements that compare as equal in an array sorted by the **qsort** function (7.24.6.3).
- (51) The order in which destructors are invoked by **thrd\_exit** (7.29.5.5).
- (52) Whether calling **tss\_delete** on a key while another thread is executing destructors affects the number of invocations of the destructors associated with the key on that thread (7.29.6.2).
- (53) The encoding of the calendar time returned by the **time** function (7.30.2.5).
- (54) The characters stored by the **strftime** or **wcsftime** function if any of the time values being converted is outside the normal range (7.30.3.6, 7.32.5.1).
- (55) Whether an encoding error occurs if a **wchar\_t** value that does not correspond to a member of the extended character set appears in the format string for a function in 7.32.2 or 7.32.5 and the specified semantics do not require that value to be processed by **wcrtomb** (7.32.1).
- (56) The conversion state after an encoding error (6.2.9) occurs (7.32.6.4.3, 7.32.6.4.4, 7.32.6.5.2, 7.32.6.5.3, 7.31.2.2, 7.31.2.3, 7.31.2.4, 7.31.2.5, 7.31.2.6, 7.31.2.7).
- (57) The conversion state after a unrecorded encoding error (6.2.9) occurs (7.25).
- (58) The use of an **mbstate\_t** object that contains conversion state from an unrelated conversion (7.25).
- (59) The resulting value when the “invalid” floating-point exception is raised during ISO/IEC 60559 floating to integer conversion (F.4).
- (60) Whether conversion of non-integer ISO/IEC 60559 floating values to integer raises the “inexact” floating-point exception (F.4).
- (61) Whether or when library functions in `<math.h>` raise the “inexact” floating-point exception in an ISO/IEC 60559 conformant implementation (F.10).
- (62) Whether or when library functions in `<math.h>` raise an undeserved “underflow” floating-point exception in an ISO/IEC 60559 conformant implementation (F.10).
- (63) The exponent value stored by **frexp** for a NaN or infinity (F.10.4.7).
- (64) The numeric result returned by the **lrint**, **llrint**, **lround**, and **llround** functions if the rounded value is outside the range of the return type (F.10.7.5, F.10.7.7).
- (65) The sign of one part of the **complex** result of several math functions for certain special cases in ISO/IEC 60559 compatible implementations (G.4.2.1, G.4.3.2, G.4.3.3, G.4.3.4, G.4.3.5, G.4.3.6, G.4.4.1, G.4.5.2).

## J.2 Undefined behavior

1 The behavior is undefined in the following circumstances:

- (1) A “shall” or “shall not” requirement that appears outside of a constraint is violated (Clause 4).
- (2) A nonempty source file does not end in a new-line character which is not immediately preceded by a backslash character or ends in a partial preprocessing token or comment (5.2.1.2).
- (3) Token concatenation produces a character sequence matching the syntax of a universal character name (5.2.1.2).
- (4) A program in a hosted environment does not define a function named **main** using one of the specified forms (5.2.2.3.2).
- (5) The execution of a program contains a data race (5.2.2.5).
- (6) A character not in the basic source character set is encountered in a source file, except in an identifier, a character literal, a string literal, a header name, a comment, or a preprocessing token that is never converted to a token (5.3.1).
- (7) An identifier, comment, string literal, character literal, or header name contains an invalid multibyte character or does not begin and end in the initial shift state (5.3.2).
- (8) The same identifier has both internal and external linkage in the same translation unit (6.2.2).
- (9) An object is referred to outside of its lifetime (6.2.4).
- (10) The value of a pointer to an object whose lifetime has ended is used (6.2.4).
- (11) The value of an object with automatic storage duration is used while the object has an indeterminate representation (6.2.4, 6.7.11, 6.8).
- (12) A non-value representation is read by an lvalue expression that does not have character type (6.2.6.1).
- (13) A non-value representation is produced by a side effect that modifies any part of the object using an lvalue expression that does not have character type (6.2.6.1).
- (14) Two declarations of the same object or function specify types that are not compatible (6.2.7).
- (15) A program requires the formation of a composite type from a variable length array type whose size is specified by an expression that is not evaluated (6.2.7).
- (16) Conversion to or from an integer type produces a value outside the range that can be represented (6.3.2.4).
- (17) Demotion of one real floating type to another produces a value outside the range that can be represented (6.3.2.5).
- (18) An lvalue does not designate an object when evaluated (6.3.3.1).
- (19) A non-array lvalue with an incomplete type is used in a context that requires the value of the designated object (6.3.3.1).
- (20) An lvalue designating an object of automatic storage duration that could have been declared with the **register** storage class is used in a context that requires the value of the designated object, but the object is uninitialized. (6.3.3.1).
- (21) An attempt is made to use the value of a void expression, or an implicit or explicit conversion (except to **void**) is applied to a void expression (6.3.3.2).
- (22) Conversion of a pointer to an integer type produces a value outside the range that can be represented (6.3.3.3).

- (23) Conversion between two pointer types produces a result that is incorrectly aligned (6.3.3.3).
- (24) A pointer is used to call a function whose type is not compatible with the referenced type (6.3.3.3).
- (25) An unmatched ' or " character is encountered on a logical source line during tokenization (6.4).
- (26) A reserved keyword token is used in translation phase 7 or 8 (5.2.1.2) for some purpose other than as a keyword (6.4.2).
- (27) A universal character name in an identifier does not designate a character whose encoding falls into one of the specified ranges (6.4.3.1).
- (28) The initial character of an identifier is a universal character name designating a digit (6.4.3.1).
- (29) Two identifiers differ only in nonsignificant characters (6.4.3.1).
- (30) The identifier `__func__` is explicitly declared (6.4.3.2).
- (31) The program attempts to modify a string literal (6.4.6).
- (32) The characters ' , \ , " , // , or /\* occur in the sequence between the < and > delimiters, or the characters ' , \ , // , or /\* occur in the sequence between the " delimiters, in a header name preprocessing token (6.4.8).
- (33) A side effect on a scalar object is unsequenced relative to either a different side effect on the same scalar object or a value computation using the value of the same scalar object (6.5.1).
- (34) An exceptional condition occurs during the evaluation of an expression (6.5.1).
- (35) An object has its stored value accessed other than by an lvalue of an allowable type (6.5.1).
- (36) A function is defined with a type that is not compatible with the type (of the expression) pointed to by the expression that denotes the called function (6.5.3.3).
- (37) A member of an atomic structure or union is accessed (6.5.3.4).
- (38) The operand of the unary \* operator has an invalid value (6.5.4.3).
- (39) The value of the second operand of the / or % operator is zero (6.5.6).
- (40) If the quotient `a/b` is not representable, the behavior of both `a/b` and `a%b` (6.5.6).
- (41) Addition or subtraction of a pointer into, or just beyond, an array object and an integer type produces a result that does not point into, or just beyond, the same array object (6.5.7).
- (42) Addition or subtraction of a pointer into, or just beyond, an array object and an integer type produces a result that points just beyond the array object and is used as the operand of a unary \* operator that is evaluated (6.5.7).
- (43) Pointers that do not point into, or just beyond, the same array object are subtracted (6.5.7).
- (44) An array subscript is out of range, even if an object is apparently accessible with the given subscript (as in the lvalue expression `a[1][7]` given the declaration `int a[4][5]`) (6.5.7).
- (45) The result of subtracting two pointers is not representable in an object of type `ptrdiff_t` (6.5.7).
- (46) An expression is shifted by a negative number or by an amount greater than or equal to the width of the promoted expression (6.5.8).
- (47) An expression having signed promoted type is left-shifted and either the value of the expression is negative or the result of shifting would not be representable in the promoted type (6.5.8).

- (48) Pointers that do not point to the same aggregate or union (nor just beyond the same array object) are compared using relational operators (6.5.9).
- (49) An object is assigned to an inexactly overlapping object or to an exactly overlapping object with incompatible type (6.5.17.2).
- (50) An expression that is required to be an integer constant expression does not have an integer type; has operands that are not integer literals, named constants, compound literal constants, enumeration constants, character literals, predefined constants, **sizeof** or **\_Lengthof** expressions whose results are integer constant expressions, **alignof** expressions, or immediately-cast floating literals; or contains casts (outside operands to **sizeof**, **\_Lengthof** and **alignof** operators) other than conversions of arithmetic types to integer types (6.6.1).
- (51) A constant expression in an initializer is not, or does not evaluate to, one of the following: a named constant, a compound literal constant, an arithmetic constant expression, a null pointer constant, an address constant, or an address constant for a complete object type plus or minus an integer constant expression (6.6.1).
- (52) An arithmetic constant expression does not have arithmetic type; has operands that are not integer literals, floating literals, named and compound literal constants of arithmetic type, character literals, predefined constants, **sizeof** or **\_Lengthof** expressions whose results are integer constant expressions, or **alignof** expressions; or contains casts (outside operands to **sizeof** or **alignof** operators) other than conversions of arithmetic types to arithmetic types (6.6.1).
- (53) The value of an object is accessed by an array-subscript `[]`, member-access `.` or `->`, address `&`, or indirection `*` operator or a pointer cast in creating an address constant (6.6.1).
- (54) An attempt is made to access, or generate a pointer to just past, a flexible array member of a structure when the referenced object provides no elements for that array (6.7.3.2).
- (55) An attempt is made to modify an object defined with a `const`-qualified type through use of an lvalue with non-`const`-qualified type (6.7.4).
- (56) An attempt is made to refer to an object defined with a `volatile`-qualified type through use of an lvalue with non-`volatile`-qualified type (6.7.4).
- (57) Two qualified types that are required to be compatible do not have the identically qualified version of a compatible type (6.7.4).
- (58) An object which has been modified is accessed through a `restrict`-qualified pointer to a `const`-qualified type, or through a `restrict`-qualified pointer and another pointer that are not both based on the same object (6.7.4.2).
- (59) A `restrict`-qualified pointer is assigned a value based on another restricted pointer whose associated block neither began execution before the block associated with this pointer, nor ended before the assignment (6.7.4.2).
- (60) A function declared with a **\_Noreturn** function specifier returns to its caller (6.7.5).
- (61) The definition of an object has an alignment specifier and another declaration of that object has a stricter alignment specifier in a different translation unit (6.7.6).
- (62) Declarations of an object in different translation units have different alignment specifiers (6.7.6).
- (63) Two pointer types that are required to be compatible are not identically qualified, or are not pointers to compatible types (6.7.7.2).
- (64) The size expression in an array declaration is not a constant expression and evaluates at program execution time to a nonpositive value (6.7.7.3).

- (65) In a context requiring two array types to be compatible, they do not have compatible element types, or their size specifiers evaluate to unequal values (6.7.7.3).
- (66) A declaration of an array parameter includes the keyword **static** within the [ and ] and the corresponding argument does not provide access to the first element of an array with at least the specified number of elements (6.7.7.4).
- (67) In a context requiring two function types to be compatible, they do not have compatible return types, or their parameters disagree in use of the ellipsis terminator or the number and type of parameters (6.7.7.4).
- (68) A function definition that does not have the asserted property is called by a function declaration or a function pointer with a type that has the **unsequenced** or **reproducible** attribute (6.7.13.8).
- (69) An identifier with external linkage is used, but in the program there does not exist exactly one external definition for the identifier, or the identifier is not used and there exist multiple external definitions for the identifier (6.9).
- (70) A function that accepts a variable number of arguments is defined without a parameter type list that ends with the ellipsis notation (6.9.2).
- (71) The } that terminates a function is reached, and the value of the function call is used by the caller (6.9.2).
- (72) A non-directive preprocessing directive is executed (6.10).
- (73) The token **defined** is generated during the expansion of a **#if** or **#elif** preprocessing directive, or the use of the **defined** unary operator does not match one of the two specified forms prior to macro replacement (6.10.2).
- (74) The **#include** preprocessing directive that results after expansion does not match one of the two header name forms (6.10.3).
- (75) The **#embed** preprocessing directive that results after expansion does not match any of the name forms (6.10.4).
- (76) The character sequence in an **#include** preprocessing directive does not start with a letter (6.10.3).
- (77) The character sequence in an **#embed** preprocessing directive does not start with a letter (6.10.4).
- (78) There are sequences of preprocessing tokens within the list of macro arguments that would otherwise act as preprocessing directives (6.10.5).
- (79) The result of the preprocessing operator **#** is not a valid ordinary string literal (6.10.5.3).
- (80) The result of the preprocessing operator **##** is not a valid preprocessing token (6.10.5.4).
- (81) The **#line** preprocessing directive that results after expansion does not match one of the two well-defined forms, or its digit sequence specifies zero or a number greater than 2147483647 (6.10.6).
- (82) A non-**STDC** **#pragma** preprocessing directive that is documented as causing translation failure or some other form of undefined behavior is encountered (6.10.8).
- (83) A **#pragma STDC** preprocessing directive does not match one of the well-defined forms (6.10.8).
- (84) The name of a predefined macro, or the identifier **defined**, is the subject of a **#define** or **#undef** preprocessing directive (6.10.10).
- (85) An attempt is made to copy an object to an overlapping object by use of a library function, other than as explicitly allowed (e.g. **memmove**) (Clause 7).

- (86) A file with the same name as one of the standard headers, not provided as part of the implementation, is placed in any of the standard places that are searched for included source files (7.1.2).
- (87) A header is included within an external declaration or definition (7.1.2).
- (88) A function, object, type, or macro that is specified as being declared or defined by some standard header is used before any header that declares or defines it is included (7.1.2).
- (89) A standard header is included while a macro is defined with the same name as a keyword (7.1.2).
- (90) The program attempts to declare a library function itself, rather than via a standard header, but the declaration does not have external linkage (7.1.2).
- (91) The program declares or defines a reserved identifier, other than as allowed by 7.1.4 (7.1.3).
- (92) The program removes the definition of a macro whose name begins with an underscore and either an uppercase letter or another underscore (7.1.3).
- (93) An argument to a library function has an invalid value or a type not expected by a variadic function (7.1.4).
- (94) The pointer passed to a library function array parameter does not have a value such that all address computations and object accesses are valid (7.1.4).
- (95) The macro definition of **assert** is suppressed to access an actual function (7.2).
- (96) The argument to the **assert** macro does not have a scalar type (7.2).
- (97) The **CX\_LIMITED\_RANGE**, **FENV\_ACCESS**, or **FP\_CONTRACT** pragma is used in any context other than outside all external declarations or preceding all explicit declarations and statements inside a compound statement (7.3.4, 7.6.2, 7.12.3).
- (98) The value of an argument to a character handling function is neither equal to the value of **EOF** nor representable as an **unsigned char** (7.4).
- (99) A macro definition of **errno** is suppressed to access an actual object, or the program defines an identifier with the name **errno** (7.5).
- (100) Part of the program tests floating-point status flags, sets floating-point control modes, or runs under non-default mode settings, but was translated with the state for the **FENV\_ACCESS** pragma “off” (7.6.2).
- (101) The exception-mask argument for one of the functions that provide access to the floating-point status flags has a nonzero value not obtained by bitwise OR of the floating-point exception macros (7.6.5).
- (102) The **fesetexceptflag** function is used to set floating-point status flags that were not specified in the call to the **fegetexceptflag** function that provided the value of the corresponding **fexcept\_t** object (7.6.5.6).
- (103) The argument to **fesetenv** or **feupdateenv** is neither an object set by a call to **fegetenv** or **fehldexcept**, nor is it an environment macro (7.6.7.4, 7.6.7.5).
- (104) The value of the result of an integer arithmetic or conversion function cannot be represented (7.8.3.1, 7.8.3.2, 7.8.3.3, 7.8.3.4, 7.24.7.1, 7.24.7.2, 7.24.2).
- (105) The program modifies the string pointed to by the value returned by the **setlocale** function (7.11.2).
- (106) A pointer returned by the **setlocale** function is used after a subsequent call to the function, or after the calling thread has exited (7.11.2).



- (107) The program modifies the structure pointed to by the value returned by the **localeconv** function (7.11.3.1).
- (108) A macro definition of **math\_errhandling** is suppressed or the program defines an identifier with the name **math\_errhandling** (7.12).
- (109) An argument to a floating-point classification or comparison macro is not of real floating type (7.12.4, 7.12.18).
- (110) A macro definition of **setjmp** is suppressed to access an actual function, or the program defines an external identifier with the name **setjmp** (7.13).
- (111) An invocation of the **setjmp** macro occurs other than in an allowed context (7.13.3.1).
- (112) The **longjmp** function is invoked to restore a nonexistent environment (7.13.3.1).
- (113) After a **longjmp**, there is an attempt to access the value of an object of automatic storage duration that does not have volatile-qualified type, local to the function containing the invocation of the corresponding **setjmp** macro, that was changed between the **setjmp** invocation and **longjmp** call (7.13.3.1).
- (114) The program specifies an invalid pointer to a signal handler function (7.14.2.1).
- (115) A signal handler returns when the signal corresponded to a computational exception (7.14.2.1).
- (116) A signal handler called in response to **SIGFPE**, **SIGILL**, **SIGSEGV**, or any other implementation-defined value corresponding to a computational exception returns (7.14.2.1).
- (117) A signal occurs as the result of calling the **abort** or **raise** function, and the signal handler calls the **raise** function (7.14.2.1).
- (118) A signal occurs other than as the result of calling the **abort** or **raise** function, and the signal handler refers to an object with static or thread storage duration that is not a lock-free atomic object other than by assigning a value to an object declared as **volatile sig\_atomic\_t**, or calls any function in the standard library other than the **abort** function, the **\_Exit** function, the **quick\_exit** function, the functions in `<stdatomic.h>` (except where explicitly stated otherwise) when the atomic arguments are lock-free, the **atomic\_is\_lock\_free** function with any atomic argument, or the **signal** function (for the same signal number) (7.14.2.1).
- (119) The value of **errno** is referred to after a signal occurred other than as the result of calling the **abort** or **raise** function and the corresponding signal handler obtained a **SIG\_ERR** return from a call to the **signal** function (7.14.2.1).
- (120) A signal is generated by an asynchronous signal handler (7.14.2.1).
- (121) The **signal** function is used in a multi-threaded program (7.14.2.1).
- (122) A variadic function attempts to access its varying arguments other than through a properly declared and initialized **va\_list** object, or before the **va\_start** macro is invoked (7.16, 7.16.2.2, 7.16.2.5).
- (123) The macro **va\_arg** is invoked using the parameter **ap** that was passed to a function that invoked the macro **va\_arg** with the same parameter (7.16).
- (124) A macro definition of **va\_start**, **va\_arg**, **va\_copy**, or **va\_end** is suppressed to access an actual function, or the program defines an external identifier with the name **va\_copy** or **va\_end** (7.16.2).
- (125) The **va\_start** or **va\_copy** macro is invoked without a corresponding invocation of the **va\_end** macro in the same function, or vice versa (7.16.2, 7.16.2.3, 7.16.2.4, 7.16.2.5).
- (126) The **va\_arg** macro is invoked when there is no actual next argument, or with a specified type that is not compatible with the promoted type of the actual next argument, with certain exceptions (7.16.2.2).



- (127) The *type* parameter to the **va\_arg** macro does not name an object type (7.16.2.2).
- (128) Using a null pointer constant in form of an integer expression as an argument to a . . . function and then interpreting it as a **void\*** or **char\*** (7.16.2.2).
- (129) The **va\_copy** or **va\_start** macro is invoked to initialize a **va\_list** that was previously initialized by either macro without an intervening invocation of the **va\_end** macro for the same **va\_list** (7.16.2.3, 7.16.2.5).
- (130) The **va\_start** macro is invoked with additional arguments that include unbalanced parentheses, or unrecognized preprocessing tokens (7.16.2.5).
- (131) The macro definition of a generic function is suppressed to access an actual function (7.17.1, 7.18).
- (132) The *type* parameter of an **offsetof** macro contains a comma not between matching parentheses (7.21).
- (133) When program execution reaches an **unreachable**( ) macro invocation (7.21.2).
- (134) Arbitrarily copying or changing the bytes of or copying from a non-null pointer into a **nullptr\_t** object and then reading that object (7.21.3).
- (135) The *member-designator* parameter of an **offsetof** macro is an invalid right operand of the . operator for the *type* parameter, or designates a bit-field (7.21).
- (136) The argument in an instance of one of the macros for constants of integer type is not a decimal, octal, or hexadecimal literal, or it has a value that exceeds the limits for the corresponding type (7.22.5).
- (137) A byte input/output function is applied to a wide-oriented stream, or a wide character input/output function is applied to a byte-oriented stream (7.23.2).
- (138) The value of a pointer to a **FILE** object is used after the associated file is closed (7.23.3).
- (139) The stream for the **fflush** function points to an input stream or to an update stream in which the most recent operation was input (7.23.5.2).
- (140) The string pointed to by the **mode** argument in a call to the **fopen** function does not exactly match one of the specified character sequences (7.23.5.3).
- (141) An output operation on an update stream is followed by an input operation without an intervening call to the **fflush** function or a file positioning function, or an input operation on an update stream is followed by an output operation without an intervening call to a file positioning function (7.23.5.3).
- (142) An attempt is made to use the contents of the array that was supplied in a call to the **setvbuf** function (7.23.5.6).
- (143) There are insufficient arguments for the format in a call to one of the formatted input/output functions, or an argument does not have an appropriate type (7.23.6.2, 7.23.6.3, 7.32.2.2, 7.32.2.3).
- (144) The format in a call to one of the formatted input/output functions or to the **strftime** or **wcsftime** function is not a valid multibyte character sequence that begins and ends in its initial shift state (7.23.6.2, 7.23.6.3, 7.30.3.6, 7.32.2.2, 7.32.2.3, 7.32.5.1).
- (145) In a call to one of the formatted output functions, a precision appears with a conversion specifier other than those described (7.23.6.2, 7.32.2.2).
- (146) A conversion specification for a formatted output function uses an asterisk to denote an argument-supplied field width or precision, but the corresponding argument is not provided (7.23.6.2, 7.32.2.2).

- (147) A conversion specification for a formatted output function uses a **#** or **0** flag with a conversion specifier other than those described (7.23.6.2, 7.32.2.2).
- (148) A conversion specification for one of the formatted input/output functions uses a length modifier with a conversion specifier other than those described (7.23.6.2, 7.23.6.3, 7.32.2.2, 7.32.2.3).
- (149) An **s** conversion specifier is encountered by one of the formatted output functions, and the argument is missing the null terminator (unless a precision is specified that does not require null termination) (7.23.6.2, 7.32.2.2).
- (150) An **n** conversion specification for one of the formatted input/output functions includes any flags, an assignment-suppressing character, a field width, or a precision (7.23.6.2, 7.23.6.3, 7.32.2.2, 7.32.2.3).
- (151) A **%** conversion specifier is encountered by one of the formatted input/output functions, but the complete conversion specification is not exactly **%%** (7.23.6.2, 7.23.6.3, 7.32.2.2, 7.32.2.3).
- (152) An invalid conversion specification is found in the format for one of the formatted input/output functions, or the **strftime** or **wcsftime** function (7.23.6.2, 7.23.6.3, 7.30.3.6, 7.32.2.2, 7.32.2.3, 7.32.5.1).
- (153) The number of characters or wide characters transmitted by a formatted output function (or written to an array, or that would have been written to an array) is greater than **INT\_MAX** (7.23.6.2, 7.32.2.2).
- (154) The number of input items assigned by a formatted input function is greater than **INT\_MAX** (7.23.6.3, 7.32.2.3).
- (155) The result of a conversion by one of the formatted input functions cannot be represented in the corresponding object, or the receiving object does not have an appropriate type (7.23.6.3, 7.32.2.3).
- (156) A **c**, **s**, or **[]** conversion specifier is encountered by one of the formatted input functions, and the array pointed to by the corresponding argument is not large enough to accept the input sequence (and a null terminator if the conversion specifier is **s** or **[]**) (7.23.6.3, 7.32.2.3).
- (157) A **c**, **s**, or **[]** conversion specifier with an **l** qualifier is encountered by one of the formatted input functions, but the input is not a valid multibyte character sequence that begins in the initial shift state (7.23.6.3, 7.32.2.3).
- (158) The input item for a **%p** conversion by one of the formatted input functions is not a value converted earlier during the same program execution (7.23.6.3, 7.32.2.3).
- (159) The **vfprintf**, **vfprintf**, **vprintf**, **vscanf**, **vsprintf**, **vsscanf**, **vfwprintf**, **vfwscanf**, **vswprintf**, **vswscanf**, **wprintf**, or **wscanf** function is called with an improperly initialized **va\_list** argument, or the argument is used (other than in an invocation of **va\_end**) after the function returns (7.23.6.9, 7.23.6.10, 7.23.6.11, 7.23.6.12, 7.23.6.13, 7.23.6.14, 7.23.6.15, 7.32.2.6, 7.32.2.7, 7.32.2.8, 7.32.2.9, 7.32.2.10, 7.32.2.11).
- (160) The contents of the array supplied in a call to the **fgets** or **fgetws** function are used after a read error occurred (7.23.7.2, 7.32.3.2).
- (161) The **n** parameter is negative or zero for a call to **fgets** or **fgetws**. (7.23.7.2, 7.32.3.2).
- (162) The file position indicator for a binary stream is used after a call to the **ungetc** function where its value was zero before the call (7.23.7.10).
- (163) The file position indicator for a stream is used after an error occurred during a call to the **fread** or **fwrite** function (7.23.8.1, 7.23.8.2).
- (164) A partial element read by a call to the **fread** function is used (7.23.8.1).

- (165) The **fseek** function is called for a text stream with a nonzero offset and either the offset was not returned by a previous successful call to the **ftell** function on a stream associated with the same file or **whence** is not **SEEK\_SET** (7.23.9.2).
- (166) The **fsetpos** function is called to set a position that was not returned by a previous successful call to the **fgetpos** function on a stream associated with the same file (7.23.9.3).
- (167) A non-null pointer returned by a call to the **calloc**, **malloc**, **realloc**, or **aligned\_alloc** function with a zero requested size is used to access an object (7.24.4).
- (168) The value of a pointer that refers to space deallocated by a call to the **free**, **free\_sized**, **free\_aligned\_sized** or **realloc** function is used (7.24.4).
- (169) The pointer argument to the **free** or **realloc** function is unequal to a null pointer and does not match a pointer earlier returned by a memory management function, or the space has been deallocated by a call to **free**, **free\_sized**, **free\_aligned\_sized** or **realloc** (7.24.4.4, 7.24.4.8).
- (170) The value of the object allocated by the **malloc** function is used (7.24.4.7).
- (171) The values of any bytes in a new object allocated by the **realloc** function beyond the size of the old object are used (7.24.4.8).
- (172) The program calls the **exit** or **quick\_exit** function more than once, or calls both functions (7.24.5.4, 7.24.5.7).
- (173) During the call to a function registered with the **atexit** or **at\_quick\_exit** function, a call is made to the **longjmp** function that would terminate the call to the registered function (7.24.5.4, 7.24.5.7).
- (174) The string set up by the **getenv** or **strerror** function is modified by the program (7.24.5.6, 7.27.6.3).
- (175) A signal is raised while the **quick\_exit** function is executing (7.24.5.7).
- (176) A command is executed through the **system** function in a way that is documented as causing termination or some other form of undefined behavior (7.24.5.8).
- (177) A searching or sorting utility function is called with an invalid pointer argument, even if the number of elements is zero (7.24.6).
- (178) The comparison function called by a searching or sorting utility function alters the contents of the array being searched or sorted, or returns ordering values inconsistently (7.24.6).
- (179) The array being searched by the **bsearch** function does not have its elements in proper order (7.24.6.2).
- (180) The current conversion state is used by a multibyte/wide character conversion function after changing the **LC\_CTYPE** category (7.24.8).
- (181) A string or wide string utility function is instructed to access an array beyond the end of an object (7.27.1, 7.32.4).
- (182) A string or wide string utility function is called with an invalid pointer argument, even if the length is zero (7.27.1, 7.32.4).
- (183) The contents of the destination array are used after a call to the **strxfrm**, **strftime**, **wcsxfrm**, or **wcsftime** function in which the specified length was too small to hold the entire null-terminated result (7.27.4.6, 7.30.3.6, 7.32.4.5.5, 7.32.5.1).
- (184) A sequence of calls of the **strtok** function is made from different threads (7.27.5.9).
- (185) The first argument in the very first call to the **strtok** or **wcstok** is a null pointer (7.27.5.9, 7.32.4.6.8).

- (186) A pointer returned by the **strerror** function is used after a subsequent call to the function, or after the calling thread has exited (7.27.6.3).
- (187) The type of an argument to a type-generic macro is not compatible with the type of the corresponding parameter of the selected function (7.28).
- (188) Arguments for generic parameters of a type-generic macro are such that some argument has a corresponding real type that is of standard floating type and another argument is of decimal floating type (7.28).
- (189) Arguments for generic parameters of a type-generic macro are such that neither `<math.h>` and `<complex.h>` define a function whose generic parameters have the determined corresponding real type (7.28).
- (190) A complex argument is supplied for a generic parameter of a type-generic macro that has no corresponding complex function (7.28).
- (191) A decimal floating argument is supplied for a generic parameter of a type-generic macro that expects a complex argument (7.28).
- (192) A standard floating or complex argument is supplied for a generic parameter of a type-generic macro that expects a decimal floating type argument (7.28).
- (193) A non-recursive mutex passed to **mtx\_lock** is locked by the calling thread (7.29.4.4).
- (194) The mutex passed to **mtx\_timedlock** does not support timeout (7.29.4.5).
- (195) The mutex passed to **mtx\_unlock** is not locked by the calling thread (7.29.4.7).
- (196) The thread passed to **thrd\_detach** or **thrd\_join** was previously detached or joined with another thread (7.29.5.3, 7.29.5.6).
- (197) The **tss\_create** function is called from within a destructor (7.29.6.1).
- (198) The key passed to **tss\_delete**, **tss\_get**, or **tss\_set** was not returned by a call to **tss\_create** before the thread commenced executing destructors (7.29.6.2, 7.29.6.3, 7.29.6.4).
- (199) An attempt is made to access the pointer returned by the time conversion functions after the thread that originally called the function to obtain it has exited (7.30.3).
- (200) At least one member of the broken-down time passed to **asctime** contains a value outside its normal range, or the calculated year exceeds four digits or is less than the year 1000 (7.30.3.2).
- (201) The argument corresponding to an **s** specifier without an **l** qualifier in a call to the **fwprintf** function does not point to a valid multibyte character sequence that begins in the initial shift state (7.32.2.12).
- (202) In a call to the **wcstok** function, the object pointed to by **ptr** does not have the value stored by the previous call for the same wide string (7.32.4.6.8).
- (203) An **mbstate\_t** object is used inappropriately (7.32.6).
- (204) The value of an argument of type **wint\_t** to a wide character classification or case mapping function is neither equal to the value of **WEOF** nor representable as a **wchar\_t** (7.33.1).
- (205) The **iswctype** function is called using a different **LC\_CTYPE** category from the one in effect for the call to the **wctype** function that returned the description (7.33.2.3.2).
- (206) The **towctrans** function is called using a different **LC\_CTYPE** category from the one in effect for the call to the **wctrans** function that returned the description (7.33.3.2.2).

### J.3 Implementation-defined behavior

#### J.3.1 General

- 1 A conforming implementation is required to document its choice of behavior in each of the areas listed in this subclause. The following are implementation-defined:

#### J.3.2 Translation

- 1
  - (1) How a diagnostic is identified (3.15, 5.2.1.3).
  - (2) Whether each nonempty sequence of white-space characters other than new-line is retained or replaced by one space character in translation phase 3 (5.2.1.2).

#### J.3.3 Environment

- 1
  - (1) The mapping between physical source file multibyte characters and the source character set in translation phase 1 (5.2.1.2).
  - (2) The name and type of the function called at program startup in a freestanding environment (5.2.2.2).
  - (3) The effect of program termination in a freestanding environment (5.2.2.2).
  - (4) An alternative manner in which the **main** function may be defined (5.2.2.3.2).
  - (5) The values given to the strings pointed to by the **argv** argument to **main** (5.2.2.3.2).
  - (6) What constitutes an interactive device (5.2.2.4).
  - (7) Whether a program can have more than one thread of execution in a freestanding environment (5.2.2.5).
  - (8) The set of signals, their semantics, and their default handling (7.14).
  - (9) Signal values other than **SIGFPE**, **SIGILL**, and **SIGSEGV** that correspond to a computational exception (7.14.2.1).
  - (10) Signals for which the equivalent of **signal(sig, SIG\_IGN)**; is executed at program startup (7.14.2.1).
  - (11) The set of environment names and the method for altering the environment list used by the **getenv** function (7.24.5.6).
  - (12) The manner of execution of the string by the **system** function (7.24.5.8).

#### J.3.4 Identifiers

- 1
  - (1) Which additional multibyte characters may appear in identifiers and their correspondence to universal character names (6.4.3).
  - (2) The number of significant initial characters in an identifier (5.3.5.2, 6.4.3).

#### J.3.5 Characters

- 1
  - (1) The number of bits in a byte (3.7).
  - (2) The values of the members of the execution character set (5.3.1).
  - (3) The unique value of the member of the execution character set produced for each of the standard alphabetic escape sequences (5.3.3).
  - (4) The value of a **char** object into which has been stored any character other than a member of the basic execution character set (6.2.5).
  - (5) Which of **signed char** or **unsigned char** has the same range, representation, and behavior as “plain” **char** (6.2.5, 6.3.2.1).

- (6) The literal encoding, which maps of the characters of the execution character set to the values in a character literal or string literal (6.2.9, 6.4.5.5).
- (7) The wide literal encoding, of the characters of the execution character set to the values in a `wchar_t` character literal or `wchar_t` string literal (6.2.9, 6.4.5.5).
- (8) The mapping of members of the source character set (in character literals and string literals) to members of the execution character set (6.4.5.5, 5.2.1.2).
- (9) The value of an ordinary character literal containing more than one character or containing a character or escape sequence that does not map to a single-byte execution character (6.4.5.5).
- (10) The value of a wide character literal containing more than one multibyte character or a single multibyte character that maps to multiple members of the extended execution character set, or containing a multibyte character or escape sequence not represented in the extended execution character set (6.4.5.5).
- (11) The current locale used to convert a wide character literal consisting of a single multibyte character that maps to a member of the extended execution character set into a corresponding wide character code (6.4.5.5).
- (12) The current locale used to convert a wide string literal into corresponding wide character codes (6.4.6).
- (13) The value of a string literal containing a multibyte character or escape sequence not represented in the execution character set (6.4.6).
- (14) The encoding of `wchar_t` where the macro `__STDC_ISO_10646__` is not defined (6.10.10.3).

### J.3.6 Integers

- <sup>1</sup> (1) Any extended integer types that exist in the implementation (6.2.5).
- (2) The rank of any extended integer type relative to another extended integer type with the same precision (6.3.2.1).
- (3) The result of, or the signal raised by, converting an integer to a signed integer type when the value cannot be represented in an object of that type (6.3.2.3).
- (4) The results of some bitwise operations on signed integers (6.5.1).

### J.3.7 Floating-point

- <sup>1</sup> (1) The accuracy of the floating-point operations and of the library functions in `<math.h>` and `<complex.h>` that return floating-point results (5.3.5.3.3).
- (2) The accuracy of the conversions between floating-point internal representations and string representations performed by the library functions in `<stdio.h>`, `<stdlib.h>`, and `<wchar.h>` (5.3.5.3.3).
- (3) The rounding behaviors characterized by non-standard values of `FLT_ROUNDS` (5.3.5.3.3).
- (4) The evaluation methods characterized by non-standard negative values of `FLT_EVAL_METHOD` (5.3.5.3.3).
- (5) The evaluation methods characterized by non-standard negative values of `DEC_EVAL_METHOD` (5.3.5.3.4).
- (6) If decimal floating types are supported (6.2.5).
- (7) The direction of rounding when an integer is converted to a floating-point number that cannot exactly represent the original value (6.3.2.4).



- (8) The direction of rounding when a floating-point number is converted to a narrower floating-point number (6.3.2.5).
- (9) How the nearest representable value or the larger or smaller representable value immediately adjacent to the nearest representable value is chosen for certain floating literals (6.4.5.3).
- (10) Whether and how floating expressions are contracted when not disallowed by the **FP\_CONTRACT** pragma (6.5.1).
- (11) The default state for the **FENV\_ACCESS** pragma (7.6.2).
- (12) Additional floating-point exceptions, rounding modes, environments, and classifications, and their macro names (7.6, 7.12).
- (13) The default state for the **FP\_CONTRACT** pragma (7.12.3).

### J.3.8 Constant expressions

- <sup>1</sup> (1) Whether or not an expression not explicitly sanctioned by this document is an extended constant expression, whether or not such extended constant expressions can be used in the same contexts as this document, and whether or not such extended constant expressions can affect potentially detectable semantic changes in the program (6.6.1).

### J.3.9 Arrays and pointers

- <sup>1</sup> (1) The result of converting a pointer to an integer or vice versa (6.3.3.3).
- (2) The size of the result of subtracting two pointers to elements of the same array (6.5.7).
- (3) An lvalue having array type is converted to a pointer to the initial element of the array, and the array object has register storage class (6.3.3.1).

### J.3.10 Hints

- <sup>1</sup> (1) The extent to which suggestions made by using the **register** storage-class specifier are effective (6.7.2).
- (2) The extent to which suggestions made by using the **inline** function specifier are effective (6.7.5).

### J.3.11 Structures, unions, enumerations, and bit-fields

- <sup>1</sup> (1) Whether a “plain” **int** bit-field is treated as a **signed int** bit-field or as an **unsigned int** bit-field (6.7.3, 6.7.3.2).
- (2) Allowable bit-field types other than **bool**, **signed int**, **unsigned int**, and bit-precise integer types (6.7.3.2).
- (3) Whether atomic types are permitted for bit-fields (6.7.3.2).
- (4) Whether a bit-field can straddle a storage-unit boundary (6.7.3.2).
- (5) The order of allocation of bit-fields within a unit (6.7.3.2).
- (6) The alignment of non-bit-field members of structures (6.7.3.2). This should present no problem unless binary data written by one implementation is read by another.
- (7) The integer type compatible with each enumerated type without fixed underlying type (6.7.3.3).
- (8) A structure or union is defined without any named members (including those specified indirectly via anonymous structures and unions) (6.7.3.2).

### J.3.12 Qualifiers

- <sup>1</sup> (1) What constitutes an access to an object that has volatile-qualified type (6.7.4).
- (2) The specification of a function type includes any type qualifiers (6.7.4).

### J.3.13 Types

- <sup>1</sup> (1) A program forms the composite type of an enumerated type and a non-enumeration integer type (6.2.7).
- (2) Whether or not it is supported for a declaration for which a type is inferred to contain a pointer, array, or function declarator (6.7.10).
- (3) Whether or not it is supported for a declaration for which a type is inferred to contain no or more than one declarators (6.7.10).

### J.3.14 Preprocessing directives

- <sup>1</sup> (1) The locations within **#pragma** directives where header name preprocessing tokens are recognized (6.4, 6.4.8).
- (2) How sequences in both forms of header names are mapped to headers or external source file names (6.4.8).
- (3) Whether the value of a character literal in a constant expression that controls conditional inclusion matches the value of the same character literal in the execution character set (6.10.2).
- (4) Whether the value of a single-character character literal in a constant expression that controls conditional inclusion may have a negative value (6.10.2).
- (5) The places that are searched for an included < > delimited header, and how the places are specified or the header is identified (6.10.3).
- (6) How the named source file is searched for in an included " " delimited header name (6.10.3).
- (7) How the named resource file is searched for in an embedded " " delimited resource name (6.10.4).
- (8) The method by which preprocessing tokens (possibly resulting from macro expansion) in a **#include** directive are combined into a header name (6.10.3).
- (9) The nesting limit for **#include** processing (6.10.3).
- (10) The method by which preprocessing tokens (possibly resulting from macro expansion) in a **#embed** directive are combined into a resource name (6.10.4).
- (11) The mapping between a resource's data and the values of the integer constant expressions, if any, in the replacement of a **#embed** directive (6.10.4).
- (12) The width of a resource located by the **#embed** directive (6.10.4).
- (13) Whether the **#** operator inserts a \ character before the \ character that begins a universal character name in a character literal or string literal (6.10.5.3).
- (14) The behavior on each recognized non-**STDC** **#pragma** directive (6.10.8).
- (15) The definitions for **\_\_DATE\_\_** and **\_\_TIME\_\_** when respectively, the date and time of translation are not available (6.10.10.2).



### J.3.15 Library functions

- 1 (1) Any library facilities available to a freestanding program, other than the minimal set required by Clause 4 (5.2.2.2).
- (2) The format of the diagnostic printed by the **assert** macro (7.2.2.1).
- (3) The representation of the floating-point status flags stored by the **fegetexceptflag** function (7.6.5.3).
- (4) Whether the **feraiseexcept** function raises the “inexact” floating-point exception in addition to the “overflow” or “underflow” floating-point exception (7.6.5.4).
- (5) Strings other than “C” and “” that may be passed as the second argument to the **setlocale** function (7.11.2).
- (6) The types defined for **float\_t** and **double\_t** when the value of the **FLT\_EVAL\_METHOD** macro is less than 0 (7.12).
- (7) The types defined for **\_Decimal32\_t** and **\_Decimal64\_t** when the value of the **DEC\_EVAL\_METHOD** macro is less than 0 (7.12).
- (8) Domain errors for the mathematics functions, other than those required by this document (7.12.2).
- (9) The values returned by the mathematics functions on domain errors or pole errors (7.12.2).
- (10) The values returned by the mathematics functions on underflow range errors, whether **errno** is set to the value of the macro **ERANGE** when the integer expression **math\_errhandling** & **MATH\_ERRNO** is nonzero, and whether the “underflow” floating-point exception is raised when the integer expression **math\_errhandling** & **MATH\_ERREXCEPT** is nonzero. (7.12.2).
- (11) Whether a domain error occurs or zero is returned when an **fmod** function has a second argument of zero (7.12.11.1).
- (12) Whether a domain error occurs or zero is returned when a **remainder** function has a second argument of zero (7.12.11.2).
- (13) The base-2 logarithm of the modulus used by the **remquo** functions in reducing the quotient (7.12.11.3).
- (14) Whether a domain error occurs or zero is returned when a **remquo** function has a second argument of zero (7.12.11.3).
- (15) Whether a byte copy of **va\_list** can be used in place of the original (7.16.1).
- (16) Whether the equivalent of **signal(sig, SIG\_DFL)**; is executed prior to the call of a signal handler, and, if not, the blocking of signals that is performed (7.14.2.1).
- (17) The value of **\_\_STDC\_ENDIAN\_NATIVE\_\_** if the execution environment is not big-endian or little-endian (7.18.2).
- (18) The null pointer constant to which the macro **NULL** expands (7.21).
- (19) Whether the last line of a text stream requires a terminating new-line character (7.23.2).
- (20) Whether space characters that are written out to a text stream immediately before a new-line character appear when read in (7.23.2).
- (21) The number of null characters that may be appended to data written to a binary stream (7.23.2).
- (22) Whether the file position indicator of an append-mode stream is initially positioned at the beginning or end of the file (7.23.3).

- (23) Whether a write on a text stream causes the associated file to be truncated beyond that point (7.23.3).
- (24) The characteristics of file buffering (7.23.3).
- (25) Whether a zero-length file actually exists (7.23.3).
- (26) The rules for composing valid file names (7.23.3).
- (27) Whether the same file can be simultaneously open multiple times (7.23.3).
- (28) The nature and choice of encodings used for multibyte characters in files (7.23.3).
- (29) The effect of the **remove** function on an open file (7.23.4.1).
- (30) The effect if a file with the new name exists prior to a call to the **rename** function (7.23.4.2).
- (31) Whether an open temporary file is removed upon abnormal program termination (7.23.4.3).
- (32) Which changes of mode are permitted (if any), and under what circumstances (7.23.5.4).
- (33) The style used to print an infinity or NaN, and the meaning of any n-char or n-wchar sequence printed for a NaN (7.23.6.2, 7.32.2.2).
- (34) The output for %p conversion in the **fprintf** or **fwprintf** function (7.23.6.2, 7.32.2.2).
- (35) The interpretation of a - character that is neither the first nor the last character, nor the second where a ^ character is the first, in the scanlist for %[ conversion in the **fscanf** or **fwscanf** function (7.23.6.3, 7.32.2.2).
- (36) The set of sequences matched by a %p conversion and the interpretation of the corresponding input item in the **fscanf** or **fwscanf** function (7.23.6.3, 7.32.2.3).
- (37) The value to which the macro **errno** is set by the **fgetpos**, **fsetpos**, or **ftell** functions on failure (7.23.9.1, 7.23.9.3, 7.23.9.4).
- (38) The meaning of any n-char or n-wchar sequence in a string representing a NaN that is converted by the **strtod**, **strtof**, **strtold**, **wctod**, **wctof**, or **wctold** function (7.24.2.6, 7.32.4.2.2).
- (39) Whether the **strtod**, **strtof**, **strtold**, **wctod**, **wctof**, or **wctold** function sets **errno** to **ERANGE** when underflow occurs (7.24.2.6, 7.32.4.2.2).
- (40) The meaning of any d-char or d-wchar sequence in a string representing a NaN that is converted by the **strtod32**, **strtod64**, **strtod128**, **wctod32**, **wctod64**, or **wctod128** function (7.24.2.7, 7.32.4.2.3).
- (41) Whether the **strtod32**, **strtod64**, **strtod128**, **wctod32**, **wctod64**, or **wctod128** function sets **errno** to **ERANGE** when underflow occurs (7.24.2.7, 7.32.4.2.3).
- (42) Whether the **calloc**, **malloc**, **realloc**, and **aligned\_alloc** functions return a null pointer or a pointer to an allocated object when the size requested is zero (7.24.4).
- (43) Whether open streams with unwritten buffered data are flushed, open streams are closed, or temporary files are removed when the **abort** or **\_Exit** function is called (7.24.5.1, 7.24.5.5).
- (44) The termination status returned to the host environment by the **abort**, **exit**, **\_Exit**, or **quick\_exit** function (7.24.5.1, 7.24.5.4, 7.24.5.5, 7.24.5.7).
- (45) The value returned by the **system** function when its argument is not a null pointer (7.24.5.8).
- (46) Whether the internal state of multibyte/wide character conversion functions has thread-storage duration, and its initial value in newly created threads (7.24.8).

- (47) Whether the multibyte/wide character conversion functions avoid data races with other calls to the same function (7.24.8).
- (48) The range and precision of times representable in `clock_t` and `time_t` (7.30).
- (49) The local time zone and Daylight Saving Time (7.30.1).
- (50) Whether `TIME_MONOTONIC` or `TIME_ACTIVE` are supported time bases (7.30.1).
- (51) Whether `TIME_THREAD_ACTIVE` is a supported time bases (7.30.1, 7.29.1).
- (52) The era for the `clock` function (7.30.2.1).
- (53) The `TIME_UTC` epoch (7.30.2.6).
- (54) The replacement string for the `%Z` specifier to the `strftime`, and `wcsftime` functions in the "C" locale (7.30.3.6, 7.32.5.1).
- (55) The behavior of the `strftime` and `wcsftime` functions if the `%z` or `%Z` specifier is used and the broken-down time structure does not have a value that could be returned by `localtime` or `gmtime`.
- (56) Whether internal `mbstate_t` objects have thread storage duration (7.31.2, 7.32.6.4, 7.32.6.5).
- (57) Whether the functions in `<math.h>` honor the rounding direction mode in an ISO/IEC 60559 conformant implementation, unless explicitly specified otherwise (F.10).

### J.3.16 Architecture

- 1
  - (1) The values or expressions assigned to the macros specified in the headers `<float.h>`, `<limits.h>`, and `<stdint.h>` (5.3.5.3, 7.22).
  - (2) The result of attempting to indirectly access an object with automatic or thread storage duration from a thread other than the one with which it is associated (6.2.4).
  - (3) The number, order, and encoding of bytes in any object (when not explicitly specified in this document) (6.2.6.1).
  - (4) Whether any extended alignments are supported and the contexts in which they are supported (6.2.8).
  - (5) Valid alignment values other than those returned by an `alignof` expression for fundamental types, if any (6.2.8).
  - (6) Whether any atomic types have fundamental alignment (6.2.8).
  - (7) The value of the result of the `sizeof` and `alignof` operators (6.5.4.5).

### J.4 Locale-specific behavior

- 1 The following characteristics of a hosted environment are locale-specific and are required to be documented by the implementation:
  - (1) Additional members of the source and execution character sets beyond the basic character set (5.3.1).
  - (2) The presence, meaning, and representation of additional multibyte characters in the execution character set beyond the basic character set (5.3.2).
  - (3) The shift states used for the encoding of multibyte characters (5.3.2).
  - (4) The direction of writing of successive printing characters (5.3.3).
  - (5) The decimal-point character (7.1.1).

- (6) The set of printing characters (7.4, 7.33.2).
- (7) The set of control characters (7.4, 7.33.2).
- (8) The sets of characters tested for by the **isalpha**, **isblank**, **islower**, **ispunct**, **isspace**, **isupper**, **iswalph**, **iswblank**, **iswlower**, **iswpunct**, **iswspace**, or **iswupper** functions (7.4.2.3, 7.4.2.4, 7.4.2.8, 7.4.2.10, 7.4.2.11, 7.4.2.12, 7.33.2.2.3, 7.33.2.2.4, 7.33.2.2.8, 7.33.2.2.10, 7.33.2.2.11, 7.33.2.2.12).
- (9) The native environment (7.11.2).
- (10) Additional subject sequences accepted by the numeric conversion functions (7.24.2, 7.32.4.2).
- (11) The collation sequence of the execution character set (7.27.4.4, 7.32.4.5.3).
- (12) The contents of the error message strings set up by the **strerror** function (7.27.6.3).
- (13) The formats for time and date (7.30.3.6, 7.32.5.1).
- (14) Character mappings that are supported by the **towctrans** function (7.33.1).
- (15) Character classifications that are supported by the **iswctype** function (7.33.1).

## J.5 Common extensions

### J.5.1 General

- 1 The following extensions are widely used in many systems, but are not portable to all implementations. The inclusion of any extension that may cause a strictly conforming program to become invalid renders an implementation nonconforming. Examples of such extensions are new keywords, extra library functions declared in standard headers, or predefined macros with names that do not begin with an underscore.

### J.5.2 Environment arguments

- 1 In a hosted environment, the **main** function receives a third argument, **char \*envp[]**, that points to a null-terminated array of pointers to **char**, each of which points to a string that provides information about the environment for this execution of the program (5.2.2.3.2).

### J.5.3 Specialized identifiers

- 1 Characters other than the underscore **\_**, letters, and digits, that are not part of the basic source character set (such as the dollar sign **\$**, or characters in national character sets) may appear in an identifier (6.4.3).

### J.5.4 Lengths and cases of identifiers

- 1 All characters in identifiers (with or without external linkage) are significant (6.4.3).

### J.5.5 Scopes of identifiers

- 1 A function identifier, or the identifier of an object the declaration of which contains the keyword **extern**, has file scope (6.2.1).

### J.5.6 Writable string literals

- 1 String literals are modifiable (in which case, identical string literals should denote distinct objects) (6.4.6).

### J.5.7 Other arithmetic types

- 1 Additional arithmetic types, such as **\_\_int128** or **double double**, and their appropriate conversions are defined (6.2.5, 6.3.2). Additional floating types may have more range or precision than **long double**, may be used for evaluating expressions of other floating types, and may be used to define **float\_t** or **double\_t**. Additional floating types may also have less range or precision than **float**.

### J.5.8 Function pointer casts

- 1 A pointer to an object or to **void** may be cast to a pointer to a function, allowing data to be invoked as a function (6.5.5).
- 2 A pointer to a function may be cast to a pointer to an object or to **void**, allowing a function to be inspected or modified (for example, by a debugger) (6.5.5).

### J.5.9 Extended bit-field types

- 1 A bit-field may be declared with a type other than **bool**, **unsigned int**, **signed int**, or a bit-precise integer type, with an appropriate maximum width (6.7.3.2).

### J.5.10 The **fortran** keyword

- 1 The **fortran** function specifier may be used in a function declaration to indicate that calls suitable for FORTRAN should be generated, or that a different representation for the external name is to be generated (6.7.5).

### J.5.11 The **asm** keyword

- 1 The **asm** keyword may be used to insert assembly language directly into the translator output (6.8). The most common implementation is via a statement of the form:

```
asm (character-string-literal);
```

### J.5.12 Type inference

- 1 A declaration for which a type is inferred (6.7.10) may additionally accept pointer declarators, function declarators, and may have more than one declarator.

### J.5.13 Multiple external definitions

- 1 There may be more than one external definition for the identifier of an object, with or without the explicit use of the keyword **extern**; if the definitions disagree, or more than one is initialized, the behavior is undefined (6.9.3).

### J.5.14 Predefined macro names

- 1 Macro names that do not begin with an underscore, describing the translation and execution environments, are defined by the implementation before translation begins (6.10.10).

### J.5.15 Floating-point status flags

- 1 If any floating-point status flags are set on normal termination after all calls to functions registered by the **atexit** function have been made (see 7.24.5.4), the implementation writes some diagnostics indicating the fact to the **stderr** stream, if it is still open,

### J.5.16 Extra arguments for signal handlers

- 1 Handlers for specific signals are called with extra arguments in addition to the signal number (7.14.2.1).

### J.5.17 Additional stream types and file-opening modes

- 1 Additional mappings from files to streams are supported (7.23.2).
- 2 Additional file-opening modes may be specified by characters appended to the **mode** argument of the **fopen** function (7.23.5.3).

### J.5.18 Defined file position indicator

- 1 The file position indicator is decremented by each successful call to the **ungetc** or **ungetwc** function for a text stream, except if its value was zero before a call (7.23.7.10, 7.32.3.10).

### J.5.19 Math error reporting

- 1 Functions declared in `<complex.h>` and `<math.h>` raise **SIGFPE** to report errors instead of, or in addition to, setting **errno** or raising floating-point exceptions (7.3, 7.12).

## J.6 Reserved identifiers and keywords

### J.6.1 General

- 1 A lot of identifier preprocessing tokens are used for specific purposes in regular clauses or appendices from translation phase 3 (5.2.1.2) onwards. Using any of these for a purpose different from their description in this document, even if the use is in a context where they are normatively permitted, may have an impact on the portability of code and should thus be avoided.

### J.6.2 Rule based identifiers

- 1 The following 53 regular expressions characterize identifiers that are systematically reserved by some clause in this document.

|                                   |                                         |
|-----------------------------------|-----------------------------------------|
| <b>atomic_</b> [a-z][a-zA-Z0-9_]* | <b>INT</b> [a-zA-Z0-9_]*_ <b>C</b>      |
| <b>ATOMIC_</b> [A-Z][a-zA-Z0-9_]* | <b>INT</b> [a-zA-Z0-9_]*_ <b>MAX</b>    |
| [a-zA-Z0-9_]*_ <b>DECIMAL_DIG</b> | <b>INT</b> [a-zA-Z0-9_]*_ <b>MIN</b>    |
| [a-zA-Z0-9_]*_ <b>DIG</b>         | <b>int</b> [a-zA-Z0-9_]*_ <b>t</b>      |
| [a-zA-Z0-9_]*_ <b>EPSILON</b>     | <b>INT</b> [a-zA-Z0-9_]*_ <b>WIDTH</b>  |
| [a-zA-Z0-9_]*_ <b>MANT_DIG</b>    | <b>is</b> [a-z][a-zA-Z0-9_]*            |
| [a-zA-Z0-9_]*_ <b>MAX</b>         | <b>LC_</b> [A-Z][a-zA-Z0-9_]*           |
| [a-zA-Z0-9_]*_ <b>MAX_10_EXP</b>  | <b>LDBL_</b> [A-Z][a-zA-Z0-9_]*         |
| [a-zA-Z0-9_]*_ <b>MAX_EXP</b>     | <b>MATH_</b> [A-Z][a-zA-Z0-9_]*         |
| [a-zA-Z0-9_]*_ <b>MIN</b>         | <b>mem</b> [a-z][a-zA-Z0-9_]*           |
| [a-zA-Z0-9_]*_ <b>MIN_10_EXP</b>  | <b>mtx_</b> [a-z][a-zA-Z0-9_]*          |
| [a-zA-Z0-9_]*_ <b>MIN_EXP</b>     | <b>PRI</b> [a-zA-Z0-9_]*_ <b>zBX</b>    |
| [a-zA-Z0-9_]*_ <b>SNAN</b>        | <b>SCN</b> [a-zA-Z0-9_]*_ <b>zBX</b>    |
| [a-zA-Z0-9_]*_ <b>TRUE_MIN</b>    | <b>SIG</b> [A-Z][a-zA-Z0-9_]*           |
| <b>_</b> [a-zA-Z_][a-zA-Z0-9_]*   | <b>SIG_</b> [A-Z][a-zA-Z0-9_]*          |
| <b>ckd_</b> [a-z][a-zA-Z0-9_]*    | <b>stdc_</b> [a-z][a-zA-Z0-9_]*         |
| <b>cnd_</b> [a-z][a-zA-Z0-9_]*    | <b>str</b> [a-z][a-zA-Z0-9_]*           |
| <b>cr_</b> [a-z][a-zA-Z0-9_]*     | <b>thrd_</b> [a-z][a-zA-Z0-9_]*         |
| <b>DBL_</b> [A-Z][a-zA-Z0-9_]*    | <b>TIME_</b> [A-Z][a-zA-Z0-9_]*         |
| <b>DEC128_</b> [A-Z][a-zA-Z0-9_]* | <b>to</b> [a-z][a-zA-Z0-9_]*            |
| <b>DEC32_</b> [A-Z][a-zA-Z0-9_]*  | <b>tss_</b> [a-z][a-zA-Z0-9_]*          |
| <b>DEC64_</b> [A-Z][a-zA-Z0-9_]*  | <b>UINT</b> [a-zA-Z0-9_]*_ <b>C</b>     |
| <b>DEC_</b> [A-Z][a-zA-Z0-9_]*    | <b>UINT</b> [a-zA-Z0-9_]*_ <b>MAX</b>   |
| <b>E</b> [0-9A-Z][a-zA-Z0-9_]*    | <b>uint</b> [a-zA-Z0-9_]*_ <b>t</b>     |
| <b>FE_</b> [A-Z][a-zA-Z0-9_]*     | <b>UINT</b> [a-zA-Z0-9_]*_ <b>WIDTH</b> |
| <b>FLT_</b> [A-Z][a-zA-Z0-9_]*    | <b>wcs</b> [a-z][a-zA-Z0-9_]*           |
| <b>FP_</b> [A-Z][a-zA-Z0-9_]*     |                                         |

- 2 The following 922 identifiers or keywords match these patterns and have particular semantics provided by this document.

|                                  |                                                |
|----------------------------------|------------------------------------------------|
| <b>_Alignas</b>                  | <b>ATOMIC_CHAR32_T_LOCK_FREE</b>               |
| <b>_Alignof</b>                  | <b>atomic_char8_t</b>                          |
| <b>_Atomic</b>                   | <b>ATOMIC_CHAR8_T_LOCK_FREE</b>                |
| <b>atomic_bool</b>               | <b>ATOMIC_CHAR_LOCK_FREE</b>                   |
| <b>ATOMIC_BOOL_LOCK_FREE</b>     | <b>atomic_compare_exchange_strong</b>          |
| <b>atomic_char</b>               | <b>atomic_compare_exchange_strong_explicit</b> |
| <b>atomic_char16_t</b>           | <b>atomic_compare_exchange_weak</b>            |
| <b>ATOMIC_CHAR16_T_LOCK_FREE</b> | <b>atomic_compare_exchange_weak_explicit</b>   |
| <b>atomic_char32_t</b>           | <b>atomic_exchange</b>                         |

|                                   |                               |
|-----------------------------------|-------------------------------|
| atomic_exchange_explicit          | atomic_uint_least64_t         |
| atomic_fetch_                     | atomic_uint_least8_t          |
| atomic_fetch_add                  | atomic_uintmax_t              |
| atomic_fetch_add_explicit         | atomic_uintptr_t              |
| atomic_fetch_and                  | atomic_ullong                 |
| atomic_fetch_and_explicit         | atomic_ulong                  |
| atomic_fetch_or                   | atomic_ushort                 |
| atomic_fetch_or_explicit          | atomic_wchar_t                |
| atomic_fetch_sub                  | ATOMIC_WCHAR_T_LOCK_FREE      |
| atomic_fetch_sub_explicit         | _BitInt                       |
| atomic_fetch_xor                  | _Bool                         |
| atomic_fetch_xor_explicit         | BOOL_MAX                      |
| atomic_flag                       | __bool_true_false_are_defined |
| atomic_flag_clear                 | CHAR_MAX                      |
| atomic_flag_clear_explicit        | CHAR_MIN                      |
| ATOMIC_FLAG_INIT                  | ckd_add                       |
| atomic_flag_test_and_set          | ckd_div                       |
| atomic_flag_test_and_set_explicit | ckd_mul                       |
| atomic_init                       | ckd_sub                       |
| atomic_int                        | cnd_broadcast                 |
| atomic_int_fast16_t               | cnd_destroy                   |
| atomic_int_fast32_t               | cnd_init                      |
| atomic_int_fast64_t               | cnd_signal                    |
| atomic_int_fast8_t                | cnd_t                         |
| atomic_int_least16_t              | cnd_timedwait                 |
| atomic_int_least32_t              | cnd_wait                      |
| atomic_int_least64_t              | _Complex                      |
| atomic_int_least8_t               | _Complex_I                    |
| ATOMIC_INT_LOCK_FREE              | __cplusplus                   |
| atomic_intmax_t                   | CR_DECIMAL_DIG                |
| atomic_intptr_t                   | __DATE__                      |
| atomic_is_lock_free               | DBL_DECIMAL_DIG               |
| atomic_llong                      | DBL_DIG                       |
| ATOMIC_LLONG_LOCK_FREE            | DBL_EPSILON                   |
| atomic_load                       | DBL_HAS_SUBNORM               |
| atomic_load_explicit              | DBL_IS_IEC_60559              |
| atomic_long                       | DBL_MANT_DIG                  |
| ATOMIC_LONG_LOCK_FREE             | DBL_MAX                       |
| ATOMIC_POINTER_LOCK_FREE          | DBL_MAX_10_EXP                |
| atomic_ptrdiff_t                  | DBL_MAX_EXP                   |
| atomic_schar                      | DBL_MIN                       |
| atomic_short                      | DBL_MIN_10_EXP                |
| ATOMIC_SHORT_LOCK_FREE            | DBL_MIN_EXP                   |
| atomic_signal_fence               | DBL_NORM_MAX                  |
| atomic_size_t                     | DBL_SNAN                      |
| atomic_store                      | DBL_TRUE_MIN                  |
| atomic_store_explicit             | DEC128_EPSILON                |
| atomic_thread_fence               | DEC128_MANT_DIG               |
| atomic_uchar                      | DEC128_MAX                    |
| atomic_uint                       | DEC128_MAX_EXP                |
| atomic_uint_fast16_t              | DEC128_MIN                    |
| atomic_uint_fast32_t              | DEC128_MIN_EXP                |
| atomic_uint_fast64_t              | DEC128_SNAN                   |
| atomic_uint_fast8_t               | DEC128_TRUE_MIN               |
| atomic_uint_least16_t             | DEC32_EPSILON                 |
| atomic_uint_least32_t             | DEC32_MANT_DIG                |



|                          |                     |
|--------------------------|---------------------|
| DEC32_MAX                | __FILE__            |
| DEC32_MAX_EXP            | FILENAME_MAX        |
| DEC32_MIN                | _Float128           |
| DEC32_MIN_EXP            | _Float128_t         |
| DEC32_SNAN               | _Float128x          |
| DEC32_TRUE_MIN           | _Float16            |
| DEC64_EPSILON            | _Float16_t          |
| DEC64_MANT_DIG           | _Float32            |
| DEC64_MAX                | _Float32_t          |
| DEC64_MAX_EXP            | _Float32x           |
| DEC64_MIN                | _Float64            |
| DEC64_MIN_EXP            | _Float64_t          |
| DEC64_SNAN               | _Float64x           |
| DEC64_TRUE_MIN           | FLT_DECIMAL_DIG     |
| DEC_EVAL_METHOD          | FLT_DIG             |
| _Decimal128              | FLT_EPSILON         |
| _Decimal128x             | FLT_EVAL_METHOD     |
| _Decimal32               | FLT_HAS_SUBNORM     |
| _Decimal32_t             | FLT_IS_IEC_60559    |
| _Decimal64               | FLT_MANT_DIG        |
| _Decimal64_t             | FLT_MAX             |
| _Decimal64x              | FLT_MAX_10_EXP      |
| DECIMAL_DIG              | FLT_MAX_EXP         |
| DEC_INFINITY             | FLT_MIN             |
| DEC_NAN                  | FLT_MIN_10_EXP      |
| __deprecated__           | FLT_MIN_EXP         |
| EDOM                     | FLT_NORM_MAX        |
| EILSEQ                   | FLT_RADIX           |
| EOF                      | FLT_ROUNDS          |
| EOL                      | FLT_SNAN            |
| ERANGE                   | FLT_TRUE_MIN        |
| _Exit                    | FOPEN_MAX           |
| EXIT_FAILURE             | FP_CONTRACT         |
| EXIT_SUCCESS             | FP_FAST_D32ADDD128  |
| __fallthrough__          | FP_FAST_D32ADDD64   |
| FE_ALL_EXCEPT            | FP_FAST_D32DIVD128  |
| FE_DEC_DOWNWARD          | FP_FAST_D32DIVD64   |
| FE_DEC_DYNAMIC           | FP_FAST_D32FMAD128  |
| FE_DEC_TONEAREST         | FP_FAST_D32FMAD64   |
| FE_DEC_TONEARESTFROMZERO | FP_FAST_D32MULD128  |
| FE_DEC_TOWARDZERO        | FP_FAST_D32MULD64   |
| FE_DEC_UPWARD            | FP_FAST_D32SQRTD128 |
| FE_DFL_ENV               | FP_FAST_D32SQRTD64  |
| FE_DFL_MODE              | FP_FAST_D32SUBD128  |
| FE_DIVBYZERO             | FP_FAST_D32SUBD64   |
| FE_DOWNWARD              | FP_FAST_D64ADDD128  |
| FE_DYNAMIC               | FP_FAST_D64DIVD128  |
| FE_INEXACT               | FP_FAST_D64FMAD128  |
| FE_INVALID               | FP_FAST_D64MULD128  |
| FE_OVERFLOW              | FP_FAST_D64SQRTD128 |
| FE_SNANS_ALWAYS_SIGNAL   | FP_FAST_D64SUBD128  |
| FE_TONEAREST             | FP_FAST_DADDL       |
| FE_TONEARESTFROMZERO     | FP_FAST_DDIVL       |
| FE_TOWARDZERO            | FP_FAST_DFMAL       |
| FE_UNDERFLOW             | FP_FAST_DMULL       |
| FE_UPWARD                | FP_FAST_DSQRTL      |



|                          |                |
|--------------------------|----------------|
| FP_FAST_DSUBL            | INT8_C         |
| FP_FAST_FADD             | INT8_MAX       |
| FP_FAST_FADDL            | INT8_MIN       |
| FP_FAST_FDIV             | int8_t         |
| FP_FAST_FDIVL            | INT8_WIDTH     |
| FP_FAST_FFMA             | int_fast16_t   |
| FP_FAST_FFMAL            | int_fast32_t   |
| FP_FAST_FMA              | int_fast64_t   |
| FP_FAST_FMA128           | int_fast8_t    |
| FP_FAST_FMA32            | int_least16_t  |
| FP_FAST_FMA64            | int_least32_t  |
| FP_FAST_FMAF             | int_least64_t  |
| FP_FAST_FMAL             | int_least8_t   |
| FP_FAST_FMUL             | INT_MAX        |
| FP_FAST_FMULL            | INTMAX_C       |
| FP_FAST_FSQRT            | INTMAX_MAX     |
| FP_FAST_FSQRTL           | INTMAX_MIN     |
| FP_FAST_FSUB             | intmax_t       |
| FP_FAST_FSUBL            | INTMAX_WIDTH   |
| FP_ILOGB0                | INT_MIN        |
| FP_ILOGBNAN              | INTPTR_MAX     |
| FP_INFINITE              | INTPTR_MIN     |
| FP_INT_DOWNWARD          | intptr_t       |
| FP_INT_TONEAREST         | INTPTR_WIDTH   |
| FP_INT_TONEARESTFROMZERO | INT_WIDTH      |
| FP_INT_TOWARDZERO        | _IOFBF         |
| FP_INT_UPWARD            | _IOLBF         |
| FP_LLOGB0                | _IONBF         |
| FP_LLOGBNAN              | isalnum        |
| FP_NAN                   | isalpha        |
| FP_NORMAL                | isblank        |
| FP_SUBNORMAL             | iscanonical    |
| FP_ZERO                  | iscntrl        |
| __func__                 | isdigit        |
| _Generic                 | iseqsig        |
| __has_c_attribute        | isfinite       |
| __has_embed              | isgraph        |
| __has_include            | isgreater      |
| __if_empty__             | isgreaterequal |
| _Imaginary               | isinf          |
| _Imaginary_I             | isless         |
| INT16_C                  | islessequal    |
| INT16_MAX                | islessgreater  |
| INT16_MIN                | islower        |
| int16_t                  | isnan          |
| INT16_WIDTH              | isnormal       |
| INT32_C                  | isprint        |
| INT32_MAX                | ispunct        |
| INT32_MIN                | issignaling    |
| int32_t                  | isspace        |
| INT32_WIDTH              | issubnormal    |
| INT64_C                  | isunordered    |
| INT64_MAX                | isupper        |
| INT64_MIN                | iswalnum       |
| int64_t                  | iswalpha       |
| INT64_WIDTH              | iswblank       |

|                      |                      |
|----------------------|----------------------|
| iswcntrl             | memory_order_consume |
| iswctype             | memory_order_relaxed |
| iswdigit             | memory_order_release |
| iswgraph             | memory_order_seq_cst |
| iswlower             | memset               |
| iswprint             | memset_explicit      |
| iswpunct             | memset_s             |
| iswspace             | mtx_destroy          |
| iswupper             | mtx_init             |
| iswxdigit            | mtx_lock             |
| isxdigit             | mtx_plain            |
| iszero               | mtx_recursive        |
| LC_ALL               | mtx_t                |
| LC_COLLATE           | mtx_timed            |
| LC_CTYPE             | mtx_timedlock        |
| LC_MONETARY          | mtx_trylock          |
| LC_NUMERIC           | mtx_unlock           |
| LC_TIME              | __nodiscard__        |
| LDBL_DECIMAL_DIG     | __Noreturn__         |
| LDBL_DIG             | __noreturn__         |
| LDBL_EPSILON         | _Noreturn            |
| LDBL_HAS_SUBNORM     | _Pragma              |
| LDBL_IS_IEC_60559    | PRId32               |
| LDBL_MANT_DIG        | PRId64               |
| LDBL_MAX             | PRIdFAST32           |
| LDBL_MAX_10_EXP      | PRIdFAST64           |
| LDBL_MAX_EXP         | PRIdLEAST32          |
| LDBL_MIN             | PRIdLEAST64          |
| LDBL_MIN_10_EXP      | PRIdMAX              |
| LDBL_MIN_EXP         | PRIdPTR              |
| LDBL_NORM_MAX        | PRId32               |
| LDBL_SNAN            | PRId64               |
| LDBL_TRUE_MIN        | PRIdFAST32           |
| _Lengthof            | PRIdFAST64           |
| __limit__            | PRIdLEAST32          |
| __LINE__             | PRIdLEAST64          |
| LLONG_MAX            | PRIdMAX              |
| LLONG_MIN            | PRIdPTR              |
| LONG_MAX             | PRId32               |
| LONG_MIN             | PRId64               |
| MATH_ERREXCEPT       | PRIdFAST32           |
| MATH_ERRNO           | PRIdFAST64           |
| __maybe_unused__     | PRIdLEAST32          |
| MB_CUR_MAX           | PRIdLEAST64          |
| MB_LEN_MAX           | PRIdMAX              |
| memalignment         | PRIdPTR              |
| memcpy               | _PRINTF_NAN_LEN_MAX  |
| memchr               | PRId32               |
| memcmp               | PRId64               |
| memcpy               | PRIdFAST32           |
| memcpy_s             | PRIdFAST64           |
| memmove              | PRIdLEAST32          |
| memmove_s            | PRIdLEAST64          |
| memory_order         | PRIdMAX              |
| memory_order_acq_rel | PRIdPTR              |
| memory_order_acquire | PRId32               |

|                     |                     |
|---------------------|---------------------|
| PRIu64              | stdc_bit_ceil_ull   |
| PRIuFAST32          | stdc_bit_ceil_us    |
| PRIuFAST64          | stdc_bit_floor      |
| PRIuLEAST32         | stdc_bit_floor_uc   |
| PRIuLEAST64         | stdc_bit_floor_ui   |
| PRIuMAX             | stdc_bit_floor_ul   |
| PRIuPTR             | stdc_bit_floor_ull  |
| PRIX32              | stdc_bit_floor_us   |
| PRIX64              | stdc_bit_width      |
| PRIXFAST32          | stdc_bit_width_uc   |
| PRIXFAST64          | stdc_bit_width_ui   |
| PRIXLEAST32         | stdc_bit_width_ul   |
| PRIXLEAST64         | stdc_bit_width_ull  |
| PRIXMAX             | stdc_bit_width_us   |
| PRIXPTR             | STDC_C16_MAX        |
| PTRDIFF_MAX         | stdc_c16nrtoc16n    |
| PTRDIFF_MIN         | stdc_c16nrtoc32n    |
| RAND_MAX            | stdc_c16nrtoc8n     |
| __reproducible__    | stdc_c16nrtomcn     |
| RSIZE_MAX           | stdc_c16nrtomwcn    |
| SCHAR_MAX           | stdc_c16snrtoc16sn  |
| SCHAR_MIN           | stdc_c16snrtoc32sn  |
| SCNbMAX             | stdc_c16snrtoc8sn   |
| SCNbPTR             | stdc_c16snrtomcsn   |
| SCNdMAX             | stdc_c16snrtomwcn   |
| SCNdPTR             | STDC_C32_MAX        |
| SCNiMAX             | stdc_c32nrtoc16n    |
| SCNiPTR             | stdc_c32nrtoc32n    |
| SCNoMAX             | stdc_c32nrtoc8n     |
| SCNoPTR             | stdc_c32nrtomcn     |
| SCNuMAX             | stdc_c32nrtomwcn    |
| SCNuPTR             | stdc_c32snrtoc16sn  |
| SCNxMAX             | stdc_c32snrtoc32sn  |
| SCNxPTR             | stdc_c32snrtoc8sn   |
| SHRT_MAX            | stdc_c32snrtomcsn   |
| SHRT_MIN            | stdc_c32snrtomwcn   |
| SIGABRT             | STDC_C8_MAX         |
| SIG_ATOMIC_MAX      | stdc_c8nrtoc16n     |
| SIG_ATOMIC_MIN      | stdc_c8nrtoc32n     |
| SIG_ATOMIC_WIDTH    | stdc_c8nrtoc8n      |
| SIG_DFL             | stdc_c8nrtomcn      |
| SIG_ERR             | stdc_c8nrtomwcn     |
| SIGFPE              | stdc_c8snrtoc16sn   |
| SIG_IGN             | stdc_c8snrtoc32sn   |
| SIGILL              | stdc_c8snrtoc8sn    |
| SIGINT              | stdc_c8snrtomcsn    |
| SIGSEGV             | stdc_c8snrtomwcn    |
| SIGTERM             | stdc_count_ones     |
| SIZE_MAX            | stdc_count_ones_uc  |
| _Static_assert      | stdc_count_ones_ui  |
| __STDC__            | stdc_count_ones_ul  |
| __STDC_ANALYZABLE__ | stdc_count_ones_ull |
| stdc_bit_ceil       | stdc_count_ones_us  |
| stdc_bit_ceil_uc    | stdc_count_zeros    |
| stdc_bit_ceil_ui    | stdc_count_zeros_uc |
| stdc_bit_ceil_ul    | stdc_count_zeros_ui |

|                              |                                |
|------------------------------|--------------------------------|
| stdc_count_zeros_ul          | stdc_leading_zeros_ull         |
| stdc_count_zeros_ull         | stdc_leading_zeros_us          |
| stdc_count_zeros_us          | __STDC_LIB_EXT1__              |
| __STDC_EMBED_EMPTY__         | __STDC_LITERAL_UTF16__         |
| __STDC_EMBED_FOUND__         | __STDC_LITERAL_UTF32__         |
| __STDC_EMBED_NOT_FOUND__     | __STDC_LITERAL_UTF8__          |
| __STDC_ENDIAN_BIG__          | stdc_load8_aligned_bes         |
| __STDC_ENDIAN_LITTLE__       | stdc_load8_aligned_beu         |
| __STDC_ENDIAN_NATIVE__       | stdc_load8_aligned_les         |
| stdc_first_leading_one       | stdc_load8_aligned_leu         |
| stdc_first_leading_one_uc    | stdc_load8_bes                 |
| stdc_first_leading_one_ui    | stdc_load8_beu                 |
| stdc_first_leading_one_ul    | stdc_load8_les                 |
| stdc_first_leading_one_ull   | stdc_load8_leu                 |
| stdc_first_leading_one_us    | __STDC_MB_MIGHT_NEQ_WC__       |
| stdc_first_leading_zero      | stdc_mcerr                     |
| stdc_first_leading_zero_uc   | stdc_mcerr_incomplete_input    |
| stdc_first_leading_zero_ui   | stdc_mcerr_insufficient_output |
| stdc_first_leading_zero_ul   | stdc_mcerr_invalid             |
| stdc_first_leading_zero_ull  | stdc_mcerr_ok                  |
| stdc_first_leading_zero_us   | STDC_MC_MAX                    |
| stdc_first_trailing_one      | stdc_mcnrtoc16n                |
| stdc_first_trailing_one_uc   | stdc_mcnrtoc32n                |
| stdc_first_trailing_one_ui   | stdc_mcnrtoc8n                 |
| stdc_first_trailing_one_ul   | stdc_mcnrtomcn                 |
| stdc_first_trailing_one_ull  | stdc_mcnrtomwcn                |
| stdc_first_trailing_one_us   | stdc_mcsnrtoc16sn              |
| stdc_first_trailing_zero     | stdc_mcsnrtoc32sn              |
| stdc_first_trailing_zero_uc  | stdc_mcsnrtoc8sn               |
| stdc_first_trailing_zero_ui  | stdc_mcsnrtomcsn               |
| stdc_first_trailing_zero_ul  | stdc_mcsnrtomwcn               |
| stdc_first_trailing_zero_ull | stdc_memreverse                |
| stdc_first_trailing_zero_us  | stdc_memreverse8u              |
| stdc_has_single_bit          | STDC_MWC_MAX                   |
| stdc_has_single_bit_uc       | stdc_mwcnrtoc16n               |
| stdc_has_single_bit_ui       | stdc_mwcnrtoc32n               |
| stdc_has_single_bit_ul       | stdc_mwcnrtoc8n                |
| stdc_has_single_bit_ull      | stdc_mwcnrtomcn                |
| stdc_has_single_bit_us       | stdc_mwcnrtomwcn               |
| __STDC_HOSTED__              | stdc_mwcnrtoc16sn              |
| __STDC_IEC_559__             | stdc_mwcnrtoc32sn              |
| __STDC_IEC_60559_BFP__       | stdc_mwcnrtoc8sn               |
| __STDC_IEC_60559_COMPLEX__   | stdc_mwcnrtomcsn               |
| __STDC_IEC_60559_DFP__       | stdc_mwcnrtomwcn               |
| __STDC_IEC_60559_TYPES__     | stdc_mwcnrtoc16sn              |
| __STDC_ISO_10646__           | stdc_mwcnrtoc32sn              |
| stdc_leading_ones            | stdc_mwcnrtoc8sn               |
| stdc_leading_ones_uc         | stdc_mwcnrtomcsn               |
| stdc_leading_ones_ui         | stdc_mwcnrtomwcn               |
| stdc_leading_ones_ul         | __STDC_NO_ATOMICS__            |
| stdc_leading_ones_ull        | __STDC_NO_COMPLEX__            |
| stdc_leading_ones_us         | __STDC_NO_THREADS__            |
| stdc_leading_zeros           | __STDC_NO_VLA__                |
| stdc_leading_zeros_uc        | stdc_rotate_left               |
| stdc_leading_zeros_ui        | stdc_rotate_left_uc            |
| stdc_leading_zeros_ul        | stdc_rotate_left_ui            |
|                              | stdc_rotate_left_ul            |
|                              | stdc_rotate_left_ull           |
|                              | stdc_rotate_left_us            |
|                              | stdc_rotate_right              |
|                              | stdc_rotate_right_uc           |

|                                   |                |
|-----------------------------------|----------------|
| stdc_rotate_right_ui              | strchr         |
| stdc_rotate_right_ul              | strcmp         |
| stdc_rotate_right_ull             | strcoll        |
| stdc_rotate_right_us              | strcpy         |
| stdc_store8_aligned_bes           | strcpy_s       |
| stdc_store8_aligned_beu           | strcspn        |
| stdc_store8_aligned_les           | strdup         |
| stdc_store8_aligned_leu           | strerror       |
| stdc_store8_bes                   | strerrorlen_s  |
| stdc_store8_beu                   | strerror_s     |
| stdc_store8_les                   | strfromd       |
| stdc_store8_leu                   | strfromd128    |
| stdc_trailing_ones                | strfromd32     |
| stdc_trailing_ones_uc             | strfromd64     |
| stdc_trailing_ones_ui             | strfromencf128 |
| stdc_trailing_ones_ul             | strfromf       |
| stdc_trailing_ones_ull            | strfroml       |
| stdc_trailing_ones_us             | strftime       |
| stdc_trailing_zeros               | strlen         |
| stdc_trailing_zeros_uc            | strncat        |
| stdc_trailing_zeros_ui            | strncat_s      |
| stdc_trailing_zeros_ul            | strncmp        |
| stdc_trailing_zeros_ull           | strncpy        |
| stdc_trailing_zeros_us            | strncpy_s      |
| __STDC_UTF_16__                   | strndup        |
| __STDC_UTF_32__                   | strnlen        |
| __STDC_VERSION__                  | strnlen_s      |
| __STDC_VERSION_ASSERT_H__         | strpbrk        |
| __STDC_VERSION_COMPLEX_H__        | strrchr        |
| __STDC_VERSION_FENV_H__           | strspn         |
| __STDC_VERSION_FLOAT_H__          | strstr         |
| __STDC_VERSION_INTTYPES_H__       | strtod         |
| __STDC_VERSION_LIMITS_H__         | strtod128      |
| __STDC_VERSION_MATH_H__           | strtod32       |
| __STDC_VERSION_SETJMP_H__         | strtod64       |
| __STDC_VERSION_STDARG_H__         | strtof         |
| __STDC_VERSION_STDATOMIC_H__      | strtoimax      |
| __STDC_VERSION_STDBIT_H__         | strtok         |
| __STDC_VERSION_STDCKDINT_H__      | strtok_s       |
| __STDC_VERSION_STDDEF_H__         | strtol         |
| __STDC_VERSION_STDINT_H__         | strtold        |
| __STDC_VERSION_STDIO_H__          | strtoll        |
| __STDC_VERSION_STDLIB_H__         | strtoul        |
| __STDC_VERSION_STRING_H__         | strtoull       |
| __STDC_VERSION_TGMATH_H__         | strtoumax      |
| __STDC_VERSION_TIME_H__           | struct         |
| __STDC_VERSION_UCHAR_H__          | strxfrm        |
| __STDC_VERSION_WCHAR_H__          | thrd_busy      |
| __STDC_WANT_IEC_60559_EXT__       | thrd_create    |
| __STDC_WANT_IEC_60559_TYPES_EXT__ | thrd_current   |
| __STDC_WANT_LIB_EXT1__            | thrd_detach    |
| __STDC_WIDE_LITERAL_UTF16__       | thrd_equal     |
| __STDC_WIDE_LITERAL_UTF32__       | thrd_error     |
| __STDC_WIDE_LITERAL_UTF8__        | thrd_exit      |
| strcat                            | thrd_join      |
| strcat_s                          | thrd_nomem     |

|                    |                 |
|--------------------|-----------------|
| thrd_sleep         | uint_fast8_t    |
| thrd_start_t       | uint_least16_t  |
| thrd_success       | uint_least32_t  |
| thrd_t             | uint_least64_t  |
| thrd_timedout      | uint_least8_t   |
| thrd_yield         | UINT_MAX        |
| _Thread_local      | UINTMAX_C       |
| __TIME__           | UINTMAX_MAX     |
| TIME_ACTIVE        | uintmax_t       |
| TIME_MONOTONIC     | UINTMAX_WIDTH   |
| TIME_THREAD_ACTIVE | UINTPTR_MAX     |
| TIME_UTC           | uintptr_t       |
| TMP_MAX            | UINTPTR_WIDTH   |
| tolower            | UINT_WIDTH      |
| totalorder         | ULLONG_MAX      |
| totalorderd128     | ULONG_MAX       |
| totalorderd32      | __unsequenced__ |
| totalorderd64      | USHRT_MAX       |
| totalorderf        | __VA_ARGS__     |
| totalorderl        | __VA_OPT__      |
| totalordermag      | WCHAR_MAX       |
| totalordermagd128  | WCHAR_MIN       |
| totalordermagd32   | wcscat          |
| totalordermagd64   | wcscat_s        |
| totalordermagf     | wcschr          |
| totalordermagl     | wcscmp          |
| toupper            | wcscoll         |
| towctrans          | wcscpy          |
| tolower            | wcscpy_s        |
| toupper            | wcscspn         |
| tss_create         | wcsftime        |
| tss_delete         | wcslen          |
| tss_dtor_t         | wcsncat         |
| tss_get            | wcsncat_s       |
| tss_set            | wcsncmp         |
| tss_t              | wcsncpy         |
| UCHAR_MAX          | wcsncpy_s       |
| UINT16_C           | wcsnlen         |
| UINT16_MAX         | wcsnlen_s       |
| uint16_t           | wcspbrk         |
| UINT16_WIDTH       | wcsrchr         |
| UINT32_C           | wcsrtombs       |
| UINT32_MAX         | wcsrtombs_s     |
| uint32_t           | wcsspn          |
| UINT32_WIDTH       | wcssstr         |
| UINT64_C           | wcstod          |
| UINT64_MAX         | wcstod128       |
| uint64_t           | wcstod32        |
| UINT64_WIDTH       | wcstod64        |
| UINT8_C            | wcstof          |
| UINT8_MAX          | wcstoimax       |
| uint8_t            | wcstok          |
| UINT8_WIDTH        | wcstok_s        |
| uint_fast16_t      | wcstol          |
| uint_fast32_t      | wcstold         |
| uint_fast64_t      | wcstoll         |

wcstombs  
wcstombs\_s  
wcstoul  
wcstoull

wcstoumax  
wcsxfrm  
WINT\_MAX  
WINT\_MIN

### J.6.3 Particular identifiers or keywords

- 1 The following 1246 identifiers or keywords are not covered by the previously listed matching patterns and have particular semantics provided by this document.

|                 |                 |                  |
|-----------------|-----------------|------------------|
| abort           | atan            | c16rtomb         |
| abort_handler_s | atan2           | c32rtomb         |
| abs             | atan2d128       | c8rtomb          |
| acos            | atan2d32        | cabs             |
| acosd128        | atan2d64        | cabsf            |
| acosd32         | atan2f          | cabsl            |
| acosd64         | atan2l          | cacos            |
| acosf           | atan2pi         | cacosf           |
| acosh           | atan2pid128     | cacosh           |
| acoshd128       | atan2pid32      | cacoshf          |
| acoshd32        | atan2pid64      | cacoshl          |
| acoshd64        | atan2pif        | cacosl           |
| acoshf          | atan2pil        | cacospi          |
| acoshl          | atand128        | cacospif         |
| acosl           | atand32         | cacospil         |
| acospi          | atand64         | calloc           |
| acospid128      | atanf           | call_once        |
| acospid32       | atanh           | canonicalize     |
| acospid64       | atanhd128       | canonicalized128 |
| acospif         | atanhd32        | canonicalized32  |
| acospil         | atanhd64        | canonicalized64  |
| alignas         | atanhf          | canonicalizef    |
| aligned_alloc   | atanhl          | canonicalizel    |
| alignof         | atanl           | carg             |
| and             | atanpi          | cargf            |
| and_eq          | atanpid128      | cargl            |
| asctime         | atanpid32       | case             |
| asctime_s       | atanpid64       | casin            |
| asin            | atanpif         | casinf           |
| asind128        | atanpil         | casinh           |
| asind32         | atexit          | casinhf          |
| asind64         | atof            | casinhl          |
| asinf           | atoi            | casinl           |
| asinh           | atol            | casinpi          |
| asinhd128       | atoll           | casinpif         |
| asinhd32        | at_quick_exit   | casinpil         |
| asinhd64        | auto            | catan            |
| asinhf          | bitand          | catanf           |
| asinhhl         | BITINT_MAXWIDTH | catanh           |
| asinl           | bitor           | catanhf          |
| asinpi          | bool            | catanhhl         |
| asinpil         | BOOL_WIDTH      | catanl           |
| asinpil128      | break           | catanpi          |
| asinpil32       | bsearch         | catanpif         |
| asinpil64       | bsearch_s       | catanpil         |
| asinpif         | btowc           | cbrt             |
| asinpil         | BUFSIZ          | cbrtd128         |
| assert          |                 |                  |

|             |                      |           |
|-------------|----------------------|-----------|
| cbrtd32     | clearerr             | cosf      |
| cbrtd64     | clgamma              | cosh      |
| cbrtf       | clgammaf             | coshd128  |
| cbrtl       | clgamma1             | coshd32   |
| ccompoundn  | clock                | coshd64   |
| ccompoundnf | CLOCKS_PER_SEC       | coshf     |
| ccompoundnl | clock_t              | coshl     |
| ccos        | clog                 | cosl      |
| ccosf       | clog10               | cospi     |
| ccosf64x    | clog10f              | cospid128 |
| ccosh       | clog10l              | cospid32  |
| ccoshf      | clog10p1             | cospid64  |
| ccoshl      | clog10p1f            | cospif    |
| ccosl       | clog10p1l            | cospil    |
| ccospi      | clog1p               | cpow      |
| ccospif     | clog1pf              | cpowf     |
| ccospil     | clog1pl              | cpowf128  |
| ceil        | clog2                | cpowl     |
| ceild128    | clog2f               | cpown     |
| ceild32     | clog2l               | cpownf    |
| ceild64     | clog2p1              | cpownl    |
| ceilf       | clog2p1f             | cpowr     |
| ceill       | clog2p1l             | cpowrf    |
| cerf        | clogf                | cpowrl    |
| cerfc       | clogl                | cproj     |
| cerfcf      | clogp1               | cprojf    |
| cerfcl      | clogp1f              | cprojl    |
| cerff       | clogp1l              | creal     |
| cerfl       | CMPLX                | crealf    |
| cexp        | CMPLXF               | creall    |
| cexp10      | CMPLXL               | crootn    |
| cexp10f     | compl                | crootnf   |
| cexp10l     | complex              | crootnl   |
| cexp10m1    | compoundn            | crsqr     |
| cexp10m1f   | compoundnd128        | crsqr     |
| cexp10m1l   | compoundnd32         | crsqr     |
| cexp2       | compoundnd64         | csin      |
| cexp2f      | compoundnf           | csinf     |
| cexp2l      | compoundnl           | csinh     |
| cexp2m1     | conj                 | csinhf    |
| cexp2m1f    | conjf                | csinh     |
| cexp2m1l    | conjl                | csinl     |
| cexpf       | const                | csinpi    |
| cexpl       | constexpr            | csinpif   |
| cexpm1      | constraint_handler_t | csinpil   |
| cexpm1f     | continue             | csqrt     |
| cexpm1l     | copysign             | csqr      |
| char        | copysign             | csqr      |
| char16_t    | copysign             | ctan      |
| char32_t    | copysign             | ctanf     |
| char8_t     | copysign             | ctanh     |
| CHAR_BIT    | copysign             | ctanhf    |
| CHAR_WIDTH  | cos                  | ctanh     |
| cimag       | cosd128              | ctanl     |
| cimagf      | cosd32               | ctanpi    |
| cimagl      | cosd64               | ctanpif   |



|                  |               |                 |
|------------------|---------------|-----------------|
| ctanpil          | div           | exp2d64         |
| ctgamma          | div_t         | exp2f           |
| ctgammaf         | dmul          | exp2l           |
| ctgammal         | dmull         | exp2m1          |
| ctime            | do            | exp2m1d128      |
| ctime_s          | double        | exp2m1d32       |
| currency_symbol  | double_t      | exp2m1d64       |
| CX_LIMITED_RANGE | dsqrt         | exp2m1f         |
| d32add           | dsqrtl        | exp2m1l         |
| d32addd128       | dsub          | expd128         |
| d32addd64        | dsubl         | expd32          |
| d32div           | elif          | expd64          |
| d32divd128       | elifdef       | expf            |
| d32divd64        | elifndef      | expl            |
| d32fma           | else          | expm1           |
| d32fmad128       | embed         | expm1d128       |
| d32fmad64        | encodebind128 | expm1d32        |
| d32mul           | encodebind32  | expm1d64        |
| d32muld128       | encodebind64  | expm1f          |
| d32muld64        | encodedecd128 | expm1l          |
| d32sqrt          | encodedecd32  | extern          |
| d32sqrtd128      | encodedecd64  | f32add          |
| d32sqrtd64       | endif         | f32addf64       |
| d32sub           | enum          | f32addf64x      |
| d32subd128       | erf           | f32fma          |
| d32subd64        | erfc          | f32fmaf32x      |
| d64add           | erfcd128      | f32mul          |
| d64addd128       | erfcd32       | f32mulf128      |
| d64div           | erfcd64       | f32mulf32x      |
| d64divd128       | erfcf         | f32xsqrt        |
| d64fma           | erfcl         | f32xsqrtf54x    |
| d64fmad128       | erfd128       | f32xsqrtf64x    |
| d64mul           | erfd32        | f64div          |
| d64muld128       | erfd64        | f64divf128      |
| d64sqrt          | erff          | f64divf64x      |
| d64sqrtd128      | erfl          | fabs            |
| d64sub           | errno         | fabsd128        |
| d64subd128       | errno_t       | fabsd32         |
| dadd             | error         | fabsd64         |
| daddl            | exit          | fabsf           |
| ddiv             | exp           | fabsl           |
| ddivl            | exp10         | fadd            |
| decimal_point    | exp10d128     | faddl           |
| decodebind128    | exp10d32      | fallthrough     |
| decodebind32     | exp10d64      | false           |
| decodebind64     | exp10f        | fclose          |
| decodedecd128    | exp10l        | fdim            |
| decodedecd32     | exp10m1       | fdimd128        |
| decodedecd64     | exp10m1d128   | fdimd32         |
| DEFAULT          | exp10m1d32    | fdimd64         |
| define           | exp10m1d64    | fdimf           |
| defined          | exp10m1f      | fdiml           |
| deprecated       | exp10m1l      | fdiv            |
| dfma             | exp2          | fdivl           |
| dfmal            | exp2d128      | feclearexcept   |
| difftime         | exp2d32       | fe_dec_getround |

|                  |                      |                    |
|------------------|----------------------|--------------------|
| fe_dec_setround  | fmaximuml            | fmul               |
| fegetenv         | fmaximum_mag         | fmull              |
| fegetexceptflag  | fmaximum_magd128     | fopen              |
| fegetmode        | fmaximum_magd32      | fopen_s            |
| fegetround       | fmaximum_magd64      | for                |
| feholdexcept     | fmaximum_magf        | fpclassify         |
| femode_t         | fmaximum_magl        | fpos_t             |
| FENV_ACCESS      | fmaximum_mag_num     | fprintf            |
| FENV_DEC_ROUND   | fmaximum_mag_numd128 | fprintf_s          |
| FENV_ROUND       | fmaximum_mag_numd32  | fputc              |
| fenv_t           | fmaximum_mag_numd64  | fputs              |
| feof             | fmaximum_mag_numf    | fputwc             |
| feraiseexcept    | fmaximum_mag_numl    | fputws             |
| ferror           | fmaximum_num         | frac_digits        |
| fesetenv         | fmaximum_numd128     | fread              |
| fesetexcept      | fmaximum_numd32      | free               |
| fesetexceptflag  | fmaximum_numd64      | free_aligned_sized |
| fesetmode        | fmaximum_numf        | free_sized         |
| fesetround       | fmaximum_numl        | freopen            |
| fetestexcept     | fmaxl                | freopen_s          |
| fetestexceptflag | fmin                 | frexp              |
| feupdateenv      | fmind128             | frexpd128          |
| fexcept_t        | fmind32              | frexpd32           |
| fflush           | fmind64              | frexpd64           |
| ffma             | fminf                | frexpf             |
| ffmal            | fminimum             | frexpl             |
| fgetc            | fminimumd128         | fromfp             |
| fgetpos          | fminimumd32          | fromfpd128         |
| fgets            | fminimumd64          | fromfpd32          |
| fgetwc           | fminimumf            | fromfpd64          |
| fgetws           | fminimuml            | fromfpf            |
| FILE             | fminimum_mag         | fromfpl            |
| float            | fminimum_magd128     | fromfpx            |
| float_t          | fminimum_magd32      | fromfpxd128        |
| floor            | fminimum_magd64      | fromfpxd32         |
| floord128        | fminimum_magf        | fromfpxd64         |
| floord32         | fminimum_magl        | fromfpxf           |
| floord64         | fminimum_mag_num     | fromfpxl           |
| floorf           | fminimum_mag_numd128 | fscanf             |
| floorl           | fminimum_mag_numd32  | fscanf_s           |
| fma              | fminimum_mag_numd64  | fseek              |
| fmad128          | fminimum_mag_numf    | fsetpos            |
| fmad32           | fminimum_mag_numl    | fsqrt              |
| fmad64           | fminimum_num         | fsqrtl             |
| fmaf             | fminimum_numd128     | fsub               |
| fmal             | fminimum_numd32      | fsubl              |
| fmax             | fminimum_numd64      | ftell              |
| fmaxd128         | fminimum_numf        | fwide              |
| fmaxd32          | fminimum_numl        | fwprintf           |
| fmaxd64          | fminl                | fwprintf_s         |
| fmaxf            | fmod                 | fwrite             |
| fmaximum         | fmodd128             | fwscanf            |
| fmaximumd128     | fmodd32              | fwscanf_s          |
| fmaximumd32      | fmodd64              | getc               |
| fmaximumd64      | fmodf                | getchar            |
| fmaximumf        | fmodl                | getenv             |

|                    |                 |               |
|--------------------|-----------------|---------------|
| getenv_s           | kill_dependency | log10p1       |
| getpayload         | labs            | log10p1d128   |
| getpayloadadd128   | lconv           | log10p1d32    |
| getpayloadadd32    | ldexp           | log10p1d64    |
| getpayloadadd64    | ldexpd128       | log10p1f      |
| getpayloadf        | ldexpd32        | log10p1l      |
| getpayloadl        | ldexpd64        | log1p         |
| gets               | ldexpf          | log1pd128     |
| gets_s             | ldexpl          | log1pd32      |
| getwc              | ldiv            | log1pd64      |
| getwchar           | ldiv_t          | log1pf        |
| gmtime             | lgamma          | log1pl        |
| gmtime_r           | lgammad128      | log2          |
| gmtime_s           | lgammad32       | log2d128      |
| goto               | lgammad64       | log2d32       |
| grouping           | lgammaf         | log2d64       |
| HUGE_VAL           | lgammal         | log2f         |
| HUGE_VAL_D128      | limit           | log2l         |
| HUGE_VAL_D32       | line            | log2p1        |
| HUGE_VAL_D64       | llabs           | log2p1d128    |
| HUGE_VALF          | lldiv           | log2p1d32     |
| HUGE_VALL          | lldiv_t         | log2p1d64     |
| hypot              | llogb           | log2p1f       |
| hypotd128          | llogbd128       | log2p1l       |
| hypotd32           | llogbd32        | logb          |
| hypotd64           | llogbd64        | logbd128      |
| hypotf             | llogbf          | logbd32       |
| hypotl             | llogbl          | logbd64       |
| I                  | LLONG_WIDTH     | logbf         |
| if                 | llquantexp      | logbl         |
| ifdef              | llquantexpd128  | logd128       |
| if_empty           | llquantexpd32   | logd32        |
| ifndef             | llquantexpd64   | logd64        |
| ignore_handler_s   | llrint          | logf          |
| ilogb              | llrintd128      | logl          |
| ilogbd128          | llrintd32       | logp1         |
| ilogbd32           | llrintd64       | logp1d128     |
| ilogbd64           | llrintf         | logp1d32      |
| ilogbf             | llrintl         | logp1d64      |
| ilogbl             | llround         | logp1f        |
| imaginary          | llroundd128     | logp1l        |
| imaxabs            | llroundd32      | long          |
| imaxdiv            | llroundd64      | long_double_t |
| imaxdiv_t          | llroundf        | longjmp       |
| include            | llroundl        | LONG_WIDTH    |
| INFINITY           | localeconv      | lrint         |
| inline             | localtime       | lrintd128     |
| int_curr_symbol    | localtime_r     | lrintd32      |
| int_frac_digits    | localtime_s     | lrintd64      |
| int_n_cs_precedes  | log             | lrintf        |
| int_n_sep_by_space | log10           | lrintl        |
| int_n_sign_posn    | log10d128       | lround        |
| int_p_cs_precedes  | log10d32        | lroundd128    |
| int_p_sep_by_space | log10d64        | lroundd32     |
| int_p_sign_posn    | log10f          | lroundd64     |
| jmp_buf            | log10l          | lroundf       |

|                   |                |                |
|-------------------|----------------|----------------|
| lroundl           | nextdownd128   | powrl          |
| L_tmpnam          | nextdownd32    | pragma         |
| L_tmpnam_s        | nextdownd64    | prefix         |
| main              | nextdownf      | printf         |
| malloc            | nextdownl      | printf_s       |
| math_errhandling  | nexttoward     | p_sep_by_space |
| max_align_t       | nexttowardd128 | p_sign_posn    |
| maybe_unused      | nexttowardd32  | ptrdiff_t      |
| mblen             | nexttowardd64  | PTRDIFF_WIDTH  |
| mbrlen            | nexttowardf    | putc           |
| mbrotoc16         | nexttowardl    | putchar        |
| mbrotoc32         | nextup         | puts           |
| mbrotoc8          | nextupd128     | putwc          |
| mbrotowc          | nextupd32      | putwchar       |
| mbsinit           | nextupd64      | qsort          |
| mbsrtowcs         | nextupf        | qsort_s        |
| mbsrtowcs_s       | nextupl        | quantize       |
| mbstate_t         | nodiscard      | quantized128   |
| mbstowcs          | noreturn       | quantized32    |
| mbstowcs_s        | not            | quantized64    |
| mbtowc            | not_eq         | quantum        |
| MB_UTF16          | n_sep_by_space | quantumd128    |
| MB_UTF32          | n_sign_posn    | quantumd32     |
| MB_UTF8           | NULL           | quantumd64     |
| mktime            | nullptr        | quick_exit     |
| modf              | nullptr_t      | raise          |
| modfd128          | OFF            | rand           |
| modfd32           | offsetof       | realloc        |
| modfd64           | ON             | register       |
| modff             | once_flag      | remainder      |
| modfl             | ONCE_FLAG_INIT | remainderd128  |
| mon_decimal_point | or             | remainderd32   |
| mon_grouping      | or_eq          | remainderd64   |
| mon_thousands_sep | p_cs_precedes  | remainderf     |
| nan               | perror         | remainderl     |
| nand128           | positive_sign  | remove         |
| nand32            | pow            | remquo         |
| nand64            | powd128        | remquoof       |
| nanf              | powd32         | remquoof       |
| nanl              | powd64         | rename         |
| n_cs_precedes     | powf           | reproducible   |
| NDEBUG            | powf32         | restrict       |
| nearbyint         | powf32x        | return         |
| nearbyintd128     | powf64         | rewind         |
| nearbyintd32      | powl           | rint           |
| nearbyintd64      | pown           | rintd128       |
| nearbyintf        | pownd128       | rintd32        |
| nearbyintl        | pownd32        | rintd64        |
| negative_sign     | pownd64        | rintf          |
| nextafter         | pownf          | rintl          |
| nextafterd128     | pownl          | rootn          |
| nextafterd32      | powr           | rootnd128      |
| nextafterd64      | powrd128       | rootnd32       |
| nextafterf        | powrd32        | rootnd64       |
| nextafterl        | powrd64        | rootnf         |
| nextdown          | powrf          | rootnl         |

|                          |                |                     |
|--------------------------|----------------|---------------------|
| round                    | setpayloadsigl | tan                 |
| roundd128                | setvbuf        | tand128             |
| roundd32                 | short          | tand32              |
| roundd64                 | SHRT_WIDTH     | tand64              |
| roundeven                | sig_atomic_t   | tanf                |
| roundevend128            | signal         | tanh                |
| roundevend32             | signbit        | tanhhd128           |
| roundevend64             | signed         | tanhhd32            |
| roundevenf               | sin            | tanhhd64            |
| roundevenl               | sind128        | tanhf               |
| roundf                   | sind32         | tanhld              |
| roundl                   | sind64         | tanl                |
| rsizet                   | sinf           | tanpi               |
| rsqrt                    | sinh           | tanpid128           |
| rsqrtd128                | sinhd128       | tanpid32            |
| rsqrtd32                 | sinhd32        | tanpid64            |
| rsqrtd64                 | sinhd64        | tanpif              |
| rsqrtf                   | sinhf          | tanpil              |
| rsqrtl                   | sinhl          | tgamma              |
| samequantum              | sinl           | tgamma128           |
| samequantumd128          | sinpi          | tgamma32            |
| samequantumd32           | sinpid128      | tgamma64            |
| samequantumd64           | sinpid32       | tgammaf             |
| scalbln                  | sinpid64       | tgamma1             |
| scalblnd128              | sinpif         | thousands_sep       |
| scalblnd32               | sinpil         | thread_local        |
| scalblnd64               | sizeof         | time                |
| scalblnf                 | size_t         | timegm              |
| scalblnl                 | SIZE_WIDTH     | timespec            |
| scalbn                   | snprintf       | timespec_get        |
| scalbnd128               | snprintf_s     | timespec_getres     |
| scalbnd32                | snwprintf_s    | time_t              |
| scalbnd64                | sprintf        | tm                  |
| scalbnf                  | sprintf_s      | tm_hour             |
| scalbnl                  | sqrt           | tm_isdst            |
| scanf                    | sqrtd128       | tm_mday             |
| scanf_s                  | sqrtd32        | tm_min              |
| SCHAR_WIDTH              | sqrtd64        | tm_mon              |
| SEEK_CUR                 | sqrtf          | tmpfile             |
| SEEK_END                 | sqrtl          | tmpfile_s           |
| SEEK_SET                 | srand          | TMP_MAX_S           |
| setbuf                   | sscanf         | tmpnam              |
| set_constraint_handler_s | sscanf_s       | tmpnam_s            |
| setjmp                   | static         | tm_sec              |
| setlocale                | static_assert  | tm_wday             |
| setpayload               | STDC           | tm_yday             |
| setpayloadadd128         | stderr         | tm_year             |
| setpayloadadd32          | stdin          | true                |
| setpayloadadd64          | stdout         | trunc               |
| setpayloadf              | suffix         | truncd128           |
| setpayloadl              | switch         | truncd32            |
| setpayloadsig            | swprintf       | truncd64            |
| setpayloadsigd128        | swprintf_s     | truncf              |
| setpayloadsigd32         | swscanf        | truncld             |
| setpayloadsigd64         | swscanf_s      | TSS_DTOR_ITERATIONS |
| setpayloadsigf           | system         | tv_nsec             |

|               |             |             |
|---------------|-------------|-------------|
| tv_sec        | va_end      | wchar_t     |
| typedef       | va_list     | WCHAR_UTF16 |
| typeof        | va_start    | WCHAR_UTF32 |
| typeof_unqual | vfprintf    | WCHAR_UTF8  |
| uabs          | vfprintf_s  | WCHAR_WIDTH |
| UCHAR_WIDTH   | vfscanf     | wctomb      |
| ufromfp       | vfscanf_s   | wctomb_s    |
| ufromfpd128   | vfwprintf   | wctob       |
| ufromfpd32    | vfwprintf_s | wctomb      |
| ufromfpd64    | vfwscanf    | wctomb_s    |
| ufromfpx      | vfwscanf_s  | wctrans     |
| ufromfpl      | void        | wctrans_t   |
| ufromfpx      | volatile    | wctype      |
| ufromfpxd128  | vprintf     | wctype_t    |
| ufromfpxd32   | vprintf_s   | WEOF        |
| ufromfpxd64   | vscanf      | while       |
| ufromfpxf     | vscanf_s    | wint_t      |
| ufromfpxl     | vsprintf    | WINT_WIDTH  |
| uimaxabs      | vsprintf_s  | wmemchr     |
| ulabs         | vsprintf_s  | wmemcmp     |
| ullabs        | vsscanf     | wmemcpy     |
| ULLONG_WIDTH  | vsscanf_s   | wmemcpy_s   |
| ULONG_WIDTH   | vswprintf   | wmemmove    |
| undef         | vswprintf_s | wmemmove_s  |
| ungetc        | vswscanf    | wmemset     |
| ungetwc       | vswscanf_s  | wprintf     |
| union         | vwprintf    | wprintf_s   |
| unreachable   | vwprintf_s  | wscanf      |
| unsequenced   | vwscanf     | wscanf_s    |
| unsigned      | warning     | xor         |
| USHRT_WIDTH   |             | xor_eq      |
| va_arg        |             |             |
| va_copy       |             |             |

## Annex K

(normative)

### Bounds-checking interfaces

#### K.1 Background

- 1 Traditionally, the C Library has contained many functions that trust the programmer to provide output character arrays big enough to hold the result being produced. Not only do these functions not check that the arrays are big enough, they frequently lack the information needed to perform such checks. While it is possible to write safe, robust, and error-free code using the existing library, the library tends to promote programming styles that lead to mysterious failures if a result is too big for the provided array.
- 2 A common programming style is to declare character arrays large enough to handle most practical cases. However, if these arrays are not large enough to handle the resulting strings, data can be written past the end of the array overwriting other data and program structures. The program never gets any indication that a problem exists, and so never has a chance to recover or to fail gracefully.
- 3 Worse, this style of programming has compromised the security of computers and networks. Buffer overflows can often be exploited to run arbitrary code with the permissions of the vulnerable (defective) program.
- 4 If the programmer writes runtime checks to verify lengths before calling library functions, then those runtime checks frequently duplicate work done inside the library functions, which discover string lengths as a side effect of doing their job.
- 5 This annex provides alternative library functions that promote safer, more secure programming. The alternative functions verify that output buffers are large enough for the intended result and return a failure indicator if they are not. Data is never written past the end of an array. All string results are null terminated.
- 6 This annex also addresses another problem that complicates writing robust code: functions that are not reentrant because they return pointers to static objects owned by the function. Such functions can be troublesome since a previously returned result can change if the function is called again, perhaps by another thread.

#### K.2 Scope

- 1 This annex specifies a series of optional extensions that can be useful in the mitigation of security vulnerabilities in programs, and comprise new functions, macros, and types declared or defined in existing standard headers.
- 2 An implementation that defines **\_\_STDC\_LIB\_EXT1\_\_** shall conform to the specifications in this annex.<sup>459)</sup>
- 3 This annex should be read as if it were merged into the parallel structure of named subclauses of Clause 7.

#### K.3 Library

##### K.3.1 Introduction

###### K.3.1.1 Standard headers

- 1 The functions, macros, and types declared or defined in this annex and its subclauses are not declared or defined by their respective headers if **\_\_STDC\_WANT\_LIB\_EXT1\_\_** is defined as a macro which expands to the integer literal 0 at the point in the source file where the appropriate header is first included.
- 2 The functions, macros, and types declared or defined in this annex and its subclauses are declared and defined by their respective headers if **\_\_STDC\_WANT\_LIB\_EXT1\_\_** is defined as a macro which expands to the integer literal 1 at the point in the source file where the appropriate header is first

<sup>459)</sup>Implementations that do not define **\_\_STDC\_LIB\_EXT1\_\_** are not required to conform to these specifications.

included.<sup>460)</sup>

- 3 It is implementation-defined whether the functions, macros, and types declared or defined in this annex and its subclauses are declared or defined by their respective headers if `__STDC_WANT_LIB_EXT1__` is not defined as a macro at the point in the source file where the appropriate header is first included.<sup>461)</sup>
- 4 Within a preprocessing translation unit, `__STDC_WANT_LIB_EXT1__` shall be defined identically for all inclusions of any headers from this annex. If `__STDC_WANT_LIB_EXT1__` is defined differently for any such inclusion, the implementation shall issue a diagnostic as if a preprocessor error directive were used.

#### K.3.1.2 Reserved identifiers

- 1 Each macro name in any of the following subclauses is reserved for use as specified if it is defined by any of its associated headers when included; unless explicitly stated otherwise (see 7.1.4).
- 2 All identifiers with external linkage in any of the following subclauses are reserved for use as identifiers with external linkage if any of them are used by the program. None of them are reserved if none of them are used.
- 3 Each identifier with file scope listed in any of the following subclauses is reserved for use as a macro name and as an identifier with file scope in the same name space if it is defined by any of its associated headers when included.

#### K.3.1.3 Use of `errno`

- 1 An implementation can set `errno` for the functions defined in this annex, but is not required to.

#### K.3.1.4 Runtime-constraint violations

- 1 Most functions in this annex include as part of their specification a list of runtime-constraints. These runtime-constraints are requirements on the program using the library.<sup>462)</sup>
- 2 Implementations shall verify that the runtime-constraints for a function are not violated by the program. If a runtime-constraint is violated, the implementation shall call the currently registered runtime-constraint handler (see `set_constraint_handler_s` in `<stdlib.h>`). Multiple runtime-constraint violations in the same call to a library function result in only one call to the runtime-constraint handler. It is unspecified which one of the multiple runtime-constraint violations cause the handler to be called.
- 3 If the runtime-constraints section for a function states an action to be performed when a runtime-constraint violation occurs, the function shall perform the action before calling the runtime-constraint handler. If the runtime-constraints section lists actions that are prohibited when a runtime-constraint violation occurs, then such actions are prohibited to the function both before calling the handler and after the handler returns.
- 4 The runtime-constraint handler is permitted not to return. If the handler does return, the library function whose runtime-constraint was violated shall return some indication of failure as given by the returns section in the function's specification.

### K.3.2 Errors `<errno.h>`

- 1 The header `<errno.h>` defines a type.
- 2 The type is

|                      |
|----------------------|
| <code>errno_t</code> |
|----------------------|

<sup>460)</sup> Future revisions of this document can define meanings for other values of `__STDC_WANT_LIB_EXT1__`.

<sup>461)</sup> 7.1.3 reserves certain names and patterns of names that an implementation can use in headers. All other names are not reserved, and a conforming implementation is not permitted to use them. While some of the names defined in this annex and its subclauses are (potentially) reserved, others are not. If an unreserved name is defined in a header when `__STDC_WANT_LIB_EXT1__` is defined as 0, the implementation is not conforming.

<sup>462)</sup> Although runtime-constraints replace many cases of undefined behavior, undefined behavior still exists in this annex. Implementations are free to detect any case of undefined behavior and treat it as a runtime-constraint violation by calling the runtime-constraint handler. This license comes directly from the definition of undefined behavior.



which is type `int`.<sup>463)</sup>

### K.3.3 Common definitions <stddef.h>

- 1 The header <stddef.h> defines a type.
- 2 The type is

`rsize_t`

which is the type `size_t`.<sup>464)</sup>

### K.3.4 Integer types <stdint.h>

- 1 The header <stdint.h> defines a macro.
- 2 The macro is

`RSIZE_MAX`

which expands to a value of type `size_t`. It can be an expression that is not constant. Functions that have parameters of type `rsize_t` consider it a runtime-constraint violation if the values of those parameters are greater than `RSIZE_MAX`.

#### Recommended practice

- 3 Extremely large object sizes are frequently a sign that an object's size was calculated incorrectly. For example, negative numbers appear as very large positive numbers when converted to an unsigned type like `size_t`. Also, some implementations do not support objects as large as the maximum value that can be represented by type `size_t`.
- 4 For those reasons, it is sometimes beneficial to restrict the range of object sizes to detect programming errors. For implementations targeting machines with large address spaces, it is recommended that `RSIZE_MAX` be defined as the smaller of the size of the largest object supported or `(SIZE_MAX >> 1)`, even if this limit is smaller than the size of some legitimate, but very large, objects. Implementations targeting machines with small address spaces may wish to define `RSIZE_MAX` as `SIZE_MAX`, which means that there is no object size that is considered a runtime-constraint violation.

### K.3.5 Input/output <stdio.h>

#### K.3.5.1 General

- 1 The header <stdio.h> defines several macros and two types.
- 2 The macros are

`L_tmpnam_s`

which expands to an integer constant expression that is the size needed for an array of `char` large enough to hold a temporary file name string generated by the `tmpnam_s` function;

`TMP_MAX_S`

which expands to an integer constant expression that is the maximum number of unique file names that can be generated by the `tmpnam_s` function.

- 3 The types are

`errno_t`

which is type `int`; and

<sup>463)</sup>As a matter of programming style, `errno_t` can be used as the type of something that deals only with the values that can be found in `errno`. For example, a function which returns the value of `errno` can be declared as having the return type `errno_t`.

<sup>464)</sup>See the description of the `RSIZE_MAX` macro in <stdint.h>.

**rsizet**

which is the type **size\_t**.

### K.3.5.2 Operations on files

#### K.3.5.2.1 The **tmpfile\_s** function

##### Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 errno_t tmpfile_s(FILE * restrict, * restrict streamptr);

```

##### Runtime-constraints

- 2 **streamptr** shall not be a null pointer.
- 3 If there is a runtime-constraint violation, **tmpfile\_s** does not attempt to create a file.

##### Description

- 4 The **tmpfile\_s** function creates a temporary binary file that is different from any other existing file and that will automatically be removed when it is closed or at program termination. If the program terminates abnormally, whether an open temporary file is removed is implementation-defined. The file is opened for update with "wb+" mode with the meaning that mode has in the **fopen\_s** function (including the mode's effect on exclusive access and file permissions).
- 5 If the file was created successfully, then the pointer to **FILE** pointed to by **streamptr** will be set to the pointer to the object controlling the opened file. Otherwise, the pointer to **FILE** pointed to by **streamptr** will be set to a null pointer.

##### Recommended practice

It should be possible to open at least **TMP\_MAX\_S** temporary files during the lifetime of the program (this limit can be shared with **tmpnam\_s**) and there should be no limit on the number simultaneously open other than this limit and any limit on the number of open files (**FOPEN\_MAX**).

##### Returns

- 6 The **tmpfile\_s** function returns zero if it created the file. If it did not create the file or there was a runtime-constraint violation, **tmpfile\_s** returns a nonzero value.

#### K.3.5.2.2 The **tmpnam\_s** function

##### Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 errno_t tmpnam_s(char *s, rsizet maxsize);

```

##### Runtime-constraints

- 2 **s** shall not be a null pointer. **maxsize** shall be less than or equal to **RSIZE\_MAX**. **maxsize** shall be greater than the length of the generated file name string.

##### Description

- 3 The **tmpnam\_s** function generates a string that is a valid file name and that is not the same as the name of an existing file.<sup>465</sup> The function is potentially capable of generating **TMP\_MAX\_S** different strings, but any or all of them can already be in use by existing files and thus not be suitable return values. The lengths of these strings shall be less than the value of the **L\_tmpnam\_s** macro.
- 4 The **tmpnam\_s** function generates a different string each time it is called.

<sup>465</sup>Files created using strings generated by the **tmpnam\_s** function are temporary only in the sense that their names are not expected to collide with those generated by conventional naming rules for the implementation. It is still necessary to use the **remove** function to remove such files when their use is ended, and before program termination.

- 5 It is assumed that `s` points to an array of at least `maxsize` characters. This array will be set to the generated string, as specified in the rest of this subclause.
- 6 The implementation shall behave as if no library function except `tmpnam` calls the `tmpnam_s` function.<sup>466)</sup>

#### Recommended practice

- 7 After a program obtains a file name using the `tmpnam_s` function and before the program creates a file with that name, the possibility exists that someone else can create a file with that same name. To avoid this race condition, the `tmpfile_s` function should be used instead of `tmpnam_s` when possible. One situation that requires the use of the `tmpnam_s` function is when the program needs to create a temporary directory rather than a temporary file.
- 8 Implementations should take care in choosing the patterns used for names returned by `tmpnam_s`. For example, making a thread ID part of the names avoids the race condition and possible conflict when multiple programs run simultaneously by the same user generate the same temporary file names.

#### Returns

- 9 If no suitable string can be generated, or if there is a runtime-constraint violation, the `tmpnam_s` function:
- if `s` is not null and `maxsize` is both greater than zero and not greater than `RSIZE_MAX`, writes a null character to `s[0]`
  - returns a nonzero value.
- 10 Otherwise, the `tmpnam_s` function writes the string in the array pointed to by `s` and returns zero.

#### Environmental limits

- 11 The value of the macro `TMP_MAX_S` shall be at least 25.

### K.3.5.3 File access functions

#### K.3.5.3.1 The `fopen_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 errno_t fopen_s(FILE * restrict * restrict streamptr,
 const char * restrict filename, const char * restrict mode);
```

#### Runtime-constraints

- 2 None of `streamptr`, `filename`, or `mode` shall be a null pointer.
- 3 If there is a runtime-constraint violation, `fopen_s` does not attempt to open a file. Furthermore, if `streamptr` is not a null pointer, `fopen_s` sets `*streamptr` to the null pointer.

#### Description

- 4 The `fopen_s` function opens the file whose name is the string pointed to by `filename`, and associates a stream with it.
- 5 The `mode` string shall be as described for `fopen`, with the addition that modes starting with the character 'w' or 'a' can be preceded by the character 'u', see the following:

`uw` truncate to zero length or create file for writing, default permissions

`ua` append; open or create file for writing at end-of-file, default permissions

<sup>466)</sup> An implementation can have `tmpnam` call `tmpnam_s` (perhaps so there is only one naming convention for temporary files), but this is not required.

The remainder of the string may contain any of the characters permitted by **fopen**.

- 6 Opening a file with exclusive mode ('x' as a character in the mode argument) fails if the file already exists or cannot be created. The check for the existence of the file and the creation of the file if it does not exist is atomic with respect to other threads and processes. If the implementation is not capable of performing the check for the existence of the file and the creation of the file atomically, it shall fail instead of performing a non-atomic check and creation.
- 7 Opening a file with private mode ('p' as a character in the mode argument) creates a file whose contents cannot be accessed by other means at any point in time including after program end, and whose **filename** may be ignored by implementations except to identify which storage ought to hold the file. If the implementation is not capable of creating private files, it shall fail.
- 8 If the file was opened successfully, then the pointer to **FILE** pointed to by **streamptr** will be set to the pointer to the object controlling the opened file. Otherwise, the pointer to **FILE** pointed to by **streamptr** will be set to a null pointer.

#### Returns

- 9 The **fopen\_s** function returns zero if it opened the file. If it did not open the file or if there was a runtime-constraint violation, **fopen\_s** returns a nonzero value.

#### K.3.5.3.2 The **freopen\_s** function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 errno_t freopen_s(FILE * restrict * restrict newstreamptr,
 const char * restrict filename, const char * restrict mode,
 FILE * restrict stream);
```

#### Runtime-constraints

- 2 None of **newstreamptr**, **mode**, and **stream** shall be a null pointer.
- 3 If there is a runtime-constraint violation, **freopen\_s** neither attempts to close any file associated with **stream** nor attempts to open a file. Furthermore, if **newstreamptr** is not a null pointer, **fopen\_s** sets **\*newstreamptr** to the null pointer.

#### Description

- 4 The **freopen\_s** function opens the file whose name is the string pointed to by **filename** and associates the stream pointed to by **stream** with it. The **mode** argument has the same meaning as in the **fopen\_s** function (including the mode's effect on exclusive access and file permissions).
- 5 If **filename** is a null pointer, the **freopen\_s** function attempts to change the mode of the stream to that specified by **mode**, as if the name of the file currently associated with the stream had been used. It is implementation-defined which changes of mode are permitted (if any), and under what circumstances.
- 6 The **freopen\_s** function first attempts to close any file that is associated with **stream**. Failure to close the file is ignored. The error and end-of-file indicators for the stream are cleared.
- 7 If the file was opened successfully, then the pointer to **FILE** pointed to by **newstreamptr** will be set to the value of **stream**. Otherwise, the pointer to **FILE** pointed to by **newstreamptr** will be set to a null pointer.

#### Returns

- 8 The **freopen\_s** function returns zero if it opened the file. If it did not open the file or there was a runtime-constraint violation, **freopen\_s** returns a nonzero value.

### K.3.5.4 Formatted input/output functions

#### K.3.5.4.1 General

- 1 Unless explicitly stated otherwise, if the execution of a function described in this subclause causes copying to take place between objects that overlap, the objects take on unspecified values.

#### K.3.5.4.2 The `fprintf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 int fprintf_s(FILE * restrict stream, const char * restrict format, ...);
```

##### Runtime-constraints

- 2 Neither `stream` nor `format` shall be a null pointer. The `%n` specifier<sup>467)</sup> (modified or not by flags, field width, or precision) shall not appear in the string pointed to by `format`. Any argument to `fprintf_s` corresponding to a `%s` specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the<sup>468)</sup> `fprintf_s` function does not attempt to produce further output, and it is unspecified to what extent `fprintf_s` produced output before discovering the runtime-constraint violation.

##### Description

- 4 The `fprintf_s` function is equivalent to the `fprintf` function except for the previously listed explicit runtime-constraints.

##### Returns

- 5 The `fprintf_s` function returns the number of characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

#### K.3.5.4.3 The `fscanf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 int fscanf_s(FILE * restrict stream, const char * restrict format, ...);
```

##### Runtime-constraints

- 2 Neither `stream` nor `format` shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the<sup>469)</sup> `fscanf_s` function does not attempt to perform further input, and it is unspecified to what extent `fscanf_s` performed input before discovering the runtime-constraint violation.

##### Description

- 4 The `fscanf_s` function is equivalent to `fscanf` except that the `c`, `s`, and `[` conversion specifiers apply to a pair of arguments (unless assignment suppression is indicated by a `*`). The first of these arguments is the same as for `fscanf`. That argument is immediately followed in the argument list by the second argument, which has type `rsize_t` and gives the number of elements in the array pointed to by the first argument of the pair. If the first argument points to a scalar object, it is considered to be an array of one element.<sup>470)</sup>

<sup>467)</sup>It is not a runtime-constraint violation for the characters `%n` to appear in sequence in the string pointed at by `format` when those characters are not interpreted as a `%n` specifier. For example, if the entire format string was `%%n`.

<sup>468)</sup>Because an implementation can treat any undefined behavior as a runtime-constraint violation, an implementation can treat any unsupported specifiers in the string pointed to by `format` as a runtime-constraint violation.

<sup>469)</sup>Because an implementation can treat any undefined behavior as a runtime-constraint violation, an implementation can treat any unsupported specifiers in the string pointed to by `format` as a runtime-constraint violation.

<sup>470)</sup>If the format is known at translation time, an implementation can issue a diagnostic for any argument used to store the result from a `c`, `s`, or `[` conversion specifier if that argument is not followed by an argument of a type compatible with

- 5 A matching failure occurs if the number of elements in a receiving object is insufficient to hold the converted input (including any trailing null character).

### Returns

- 6 The **fscanf\_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **fscanf\_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.
- 7 EXAMPLE 1 The call:

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
/* ... */
int n, i; float x; char name[50];
n = fscanf_s(stdin, "%d%f%S", &i, &x, name, (rsize_t) 50);
```

with the input line:

```
25 54.32E-1 thompson
```

will assign to **n** the value 3, to **i** the value 25, to **x** the value 5.432, and to **name** the sequence **thompson\0**.

- 8 EXAMPLE 2 The call:

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
/* ... */
int n; char s[5];
n = fscanf_s(stdin, "%s", s, sizeof s);
```

with the input line:

```
hello
```

will assign to **n** the value 0 since a matching failure occurred because the sequence **hello\0** requires an array of six characters to store it.

### K.3.5.4.4 The **printf\_s** function

#### Synopsis

- 1
- ```
#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
int printf_s(const char * restrict format, ...);
```

Runtime-constraints

- 2 **format** shall not be a null pointer. The **%n** specifier⁴⁷¹⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by **format**. Any argument to **printf_s** corresponding to a **%s** specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **printf_s** function does not attempt to produce further output, and it is unspecified to what extent **printf_s** produced output before discovering the runtime-constraint violation.

rsize_t. A limited amount of checking can be done if even if the format is not known at translation time. For example, an implementation can issue a diagnostic for each argument after **format** that has of type pointer to one of **char**, **signed char**, **unsigned char**, or **void** that is not followed by an argument of a type compatible with **rsize_t**. The diagnostic can warn that unless the pointer is being used with a conversion specifier using the **hh** length modifier, a length argument is expected to follow the pointer argument. Another useful diagnostic can flag any non-pointer argument following **format** that did not have a type compatible with **rsize_t**.

⁴⁷¹⁾It is not a runtime-constraint violation for the characters **%n** to appear in sequence in the string pointed to by **format** when those characters are not interpreted as a **%n** specifier. For example, if the entire format string was **%n**.

Description

- 4 The **printf_s** function is equivalent to the **printf** function except for the previously listed explicit runtime-constraints.

Returns

- 5 The **printf_s** function returns the number of characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

K.3.5.4.5 The scanf_s function**Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdio.h>
    int scanf_s(const char * restrict format, ...);
```

Runtime-constraints

- 2 **format** shall not be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **scanf_s** function does not attempt to perform further input, and it is unspecified to what extent **scanf_s** performed input before discovering the runtime-constraint violation.

Description

- 4 The **scanf_s** function is equivalent to **fscanf_s** with the argument **stdin** interposed before the arguments to **scanf_s**.

Returns

- 5 The **scanf_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **scanf_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

K.3.5.4.6 The snprintf_s function**Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdio.h>
    int snprintf_s(char * restrict s, rsize_t n, const char * restrict format, ...);
```

Runtime-constraints

- 2 Neither **s** nor **format** shall be a null pointer. **n** shall neither equal zero nor be greater than **RSIZE_MAX**. The **%n** specifier⁴⁷²⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by **format**. Any argument to **snprintf_s** corresponding to a **%s** specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if **s** is not a null pointer and **n** is greater than zero and not greater than **RSIZE_MAX**, then the **snprintf_s** function sets **s[0]** to the null character.

Description

- 4 The **snprintf_s** function is equivalent to the **snprintf** function except for the previously listed explicit runtime-constraints.
- 5 The **snprintf_s** function, unlike **sprintf_s**, will truncate the result to fit within the array pointed to by **s**.

⁴⁷²⁾It is not a runtime-constraint violation for the characters **%n** to appear in sequence in the string pointed at by **format** when those characters are not interpreted as a **%n** specifier. For example, if the entire format string was **%%n**.

Returns

- 6 The **snprintf_s** function returns the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character, or a negative value if a runtime-constraint violation occurred. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than **n**.

K.3.5.4.7 The **sprintf_s function****Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdio.h>
    int sprintf_s(char * restrict s, rsize_t n, const char * restrict format, ...);
```

Runtime-constraints

- 2 Neither **s** nor **format** shall be a null pointer. **n** shall neither equal zero nor be greater than **RSIZE_MAX**. The number of characters (including the trailing null) required for the result to be written to the array pointed to by **s** shall not be greater than **n**. The **%n** specifier⁴⁷³⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by **format**. Any argument to **sprintf_s** corresponding to a **%s** specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if **s** is not a null pointer and **n** is greater than zero and not greater than **RSIZE_MAX**, then the **sprintf_s** function sets **s[0]** to the null character.

Description

- 4 The **sprintf_s** function is equivalent to the **sprintf** function except for the parameter **n** and the previously listed explicit runtime-constraints.
- 5 The **sprintf_s** function, unlike **snprintf_s**, treats a result too big for the array pointed to by **s** as a runtime-constraint violation.

Returns

- 6 If no runtime-constraint violation occurred, the **sprintf_s** function returns the number of characters written in the array, not counting the terminating null character. If an encoding error occurred, **sprintf_s** returns a negative value. If any other runtime-constraint violation occurred, **sprintf_s** returns zero.

K.3.5.4.8 The **sscanf_s function****Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdio.h>
    int sscanf_s(const char * restrict s, const char * restrict format, ...);
```

Runtime-constraints

- 2 Neither **s** nor **format** shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **sscanf_s** function does not attempt to perform further input, and it is unspecified to what extent **sscanf_s** performed input before discovering the runtime-constraint violation.

Description

- 4 The **sscanf_s** function is equivalent to **fscanf_s**, except that input is obtained from a string (specified by the argument **s**) rather than from a stream. Reaching the end of the string is equivalent to encountering end-of-file for the **fscanf_s** function. If copying takes place between objects that overlap, the objects take on unspecified values.

⁴⁷³⁾ It is not a runtime-constraint violation for the characters **%n** to appear in sequence in the string pointed at by **format** when those characters are not interpreted as a **%n** specifier. For example, if the entire format string was **%n**.

Returns

- 5 The **sscanf_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **sscanf_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

K.3.5.4.9 The **vfprintf_s function****Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vfprintf_s(FILE * restrict stream, const char * restrict format, va_list arg)
        ;
```

Runtime-constraints

- 2 Neither **stream** nor **format** shall be a null pointer. The **%n** specifier⁴⁷⁴ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by **format**. Any argument to **vfprintf_s** corresponding to a **%s** specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **vfprintf_s** function does not attempt to produce further output, and it is unspecified to what extent **vfprintf_s** produced output before discovering the runtime-constraint violation.

Description

- 4 The **vfprintf_s** function is equivalent to the **vfprintf** function except for the previously listed explicit runtime-constraints.

Returns

- 5 The **vfprintf_s** function returns the number of characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

⁴⁷⁴) It is not a runtime-constraint violation for the characters **%n** to appear in sequence in the string pointed at by **format** when those characters are not interpreted as a **%n** specifier. For example, if the entire format string was **%%n**.

K.3.5.4.10 The `vfscanf_s` function**Synopsis**

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vfscanf_s(FILE * restrict stream, const char * restrict format, va_list arg);

```

Runtime-constraints

- 2 Neither `stream` nor `format` shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `vfscanf_s` function does not attempt to perform further input, and it is unspecified to what extent `vfscanf_s` performed input before discovering the runtime-constraint violation.

Description

- 4 The `vfscanf_s` function is equivalent to `fscanf_s`, with the varying argument list replaced by `arg`, which shall have been initialized by the `va_start` macro (and possibly subsequent `va_arg` invocations). The `vfscanf_s` function does not invoke the `va_end` macro.⁴⁷⁵⁾

Returns

- 5 The `vfscanf_s` function returns the value of the macro `EOF` if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the `vfscanf_s` function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

K.3.5.4.11 The `vprintf_s` function**Synopsis**

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vprintf_s(const char * restrict format, va_list arg);

```

Runtime-constraints

- 2 `format` shall not be a null pointer. The `%n` specifier⁴⁷⁶⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by `format`. Any argument to `vprintf_s` corresponding to a `%s` specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `vprintf_s` function does not attempt to produce further output, and it is unspecified to what extent `vprintf_s` produced output before discovering the runtime-constraint violation.

Description

- 4 The `vprintf_s` function is equivalent to the `vprintf` function except for the previously listed explicit runtime-constraints.

Returns

- 5 The `vprintf_s` function returns the number of characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

K.3.5.4.12 The `vscanf_s` function

⁴⁷⁵⁾As the functions `vfprintf_s`, `vfscanf_s`, `vprintf_s`, `vscanf_s`, `vsprintf_s`, `vsnprintf_s`, `vsprintf_s`, and `vsscanf_s` invoke the `va_arg` macro, the representation of `arg` after the return is indeterminate.

⁴⁷⁶⁾It is not a runtime-constraint violation for the characters `%n` to appear in sequence in the string pointed at by `format` when those characters are not interpreted as a `%n` specifier. For example, if the entire format string was `%%n`.

Synopsis

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vscanf_s(const char * restrict format, va_list arg);

```

Runtime-constraints

- 2 **format** shall not be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **vsscanf_s** function does not attempt to perform further input, and it is unspecified to what extent **vsscanf_s** performed input before discovering the runtime-constraint violation.

Description

- 4 The **vsscanf_s** function is equivalent to **scanf_s**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va_start** macro (and possibly subsequent **va_arg** invocations). The **vsscanf_s** function does not invoke the **va_end** macro.⁴⁷⁷⁾

Returns

- 5 The **vsscanf_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **vsscanf_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

K.3.5.4.13 The **vsnprintf_s** function

Synopsis

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vsnprintf_s(char * restrict s, rsize_t n, const char * restrict format,
        va_list arg);

```

Runtime-constraints

- 2 Neither **s** nor **format** shall be a null pointer. **n** shall neither equal zero nor be greater than **RSIZE_MAX**. The **%n** specifier⁴⁷⁸⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by **format**. Any argument to **vsnprintf_s** corresponding to a **%s** specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if **s** is not a null pointer and **n** is greater than zero and not greater than **RSIZE_MAX**, then the **vsnprintf_s** function sets **s[0]** to the null character.

Description

- 4 The **vsnprintf_s** function is equivalent to the **vsnprintf** function except for the previously listed explicit runtime-constraints.
- 5 The **vsnprintf_s** function, unlike **vsprintf_s**, will truncate the result to fit within the array pointed to by **s**.

Returns

- 6 The **vsnprintf_s** function returns the number of characters that would have been written had **n** been sufficiently large, not counting the terminating null character, or a negative value if a runtime-constraint violation occurred. Thus, the null-terminated output has been completely written if and

⁴⁷⁷⁾As the functions **vfprintf_s**, **vfscanf_s**, **vprintf_s**, **vsprintf_s**, **vsnprintf_s**, **vsprintf_s**, and **vsscanf_s** invoke the **va_arg** macro, the representation of **arg** after the return is indeterminate.

⁴⁷⁸⁾It is not a runtime-constraint violation for the characters **%n** to appear in sequence in the string pointed at by **format** when those characters are not interpreted as a **%n** specifier. For example, if the entire format string was **%%n**.

only if the returned value is both nonnegative and less than `n`.

K.3.5.4.14 The `vsprintf_s` function

Synopsis

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vsprintf_s(char * restrict s, rsize_t n, const char * restrict format,
        va_list arg);
```

Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. `n` shall neither equal zero nor be greater than `RSIZE_MAX`. The number of characters (including the trailing null) required for the result to be written to the array pointed to by `s` shall not be greater than `n`. The `%n` specifier⁴⁷⁹⁾ (modified or not by flags, field width, or precision) shall not appear in the string pointed to by `format`. Any argument to `vsprintf_s` corresponding to a `%s` specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `n` is greater than zero and not greater than `RSIZE_MAX`, then the `vsprintf_s` function sets `s[0]` to the null character.

Description

- 4 The `vsprintf_s` function is equivalent to the `vsprintf` function except for the parameter `n` and the previously listed explicit runtime-constraints.
- 5 The `vsprintf_s` function, unlike `vsprintf_s`, treats a result too big for the array pointed to by `s` as a runtime-constraint violation.

Returns

- 6 If no runtime-constraint violation occurred, the `vsprintf_s` function returns the number of characters written in the array, not counting the terminating null character. If an encoding error occurred, `vsprintf_s` returns a negative value. If any other runtime-constraint violation occurred, `vsprintf_s` returns zero.

K.3.5.4.15 The `vsscanf_s` function

Synopsis

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdarg.h>
    #include <stdio.h>
    int vsscanf_s(const char * restrict s, const char * restrict format, va_list arg)
        ;
```

Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `vsscanf_s` function does not attempt to perform further input, and it is unspecified to what extent `vsscanf_s` performed input before discovering the runtime-constraint violation.

Description

- 4 The `vsscanf_s` function is equivalent to `sscanf_s`, with the varying argument list replaced by `arg`, which shall have been initialized by the `va_start` macro (and possibly subsequent `va_arg` calls). The `vsscanf_s` function does not invoke the `va_end` macro.⁴⁸⁰⁾

⁴⁷⁹⁾It is not a runtime-constraint violation for the characters `%n` to appear in sequence in the string pointed at by `format` when those characters are not interpreted as a `%n` specifier. For example, if the entire format string was `%n`.

⁴⁸⁰⁾As the functions `vfprintf_s`, `vfscanf_s`, `vprintf_s`, `vscanf_s`, `vsprintf_s`, `vscanf_s`, and `vsscanf_s` invoke the `va_arg` macro, the value of `arg` after the return is indeterminate.

Returns

- 5 The **vsscanf_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **vscanf_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

K.3.5.5 Character input/output functions**K.3.5.5.1 The **gets_s** function****Synopsis**

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdio.h>
    char *gets_s(char *s, rsize_t n);
```

Runtime-constraints

- 2 **s** shall not be a null pointer. **n** shall neither be equal to zero nor be greater than **RSIZE_MAX**. A new-line character, end-of-file, or read error shall occur within reading **n** - 1 characters from **stdin**.⁴⁸¹⁾
- 3 If there is a runtime-constraint violation, characters are read and discarded from **stdin** until a new-line character is read, or end-of-file or a read error occurs, and if **s** is not a null pointer, **s[0]** is set to the null character.

Description

- 4 The **gets_s** function reads at most one less than the number of characters specified by **n** from the stream pointed to by **stdin**, into the array pointed to by **s**. No additional characters are read after a new-line character (which is discarded) or after end-of-file. The discarded new-line character does not count towards number of characters read. A null character is written immediately after the last character read into the array.
- 5 If end-of-file is encountered and no characters have been read into the array, or if a read error occurs during the operation, then **s[0]** is set to the null character, and the other elements of **s** take unspecified values.

Recommended practice

- 6 The **fgets** function allows properly-written programs to safely process input lines too long to store in the result array. In general this requires that callers of **fgets** pay attention to the presence or absence of a new-line character in the result array. It is recommended to use **fgets** (along with any needed processing based on new-line characters) instead of **gets_s**.

Returns

- 7 The **gets_s** function returns **s** if successful. If there was a runtime-constraint violation, or if end-of-file is encountered and no characters have been read into the array, or if a read error occurs during the operation, then a null pointer is returned.

⁴⁸¹⁾The **gets_s** function, unlike the historical **gets** function, makes it a runtime-constraint violation for a line of input to overflow the buffer to store it. Unlike the **fgets** function, **gets_s** maintains a one-to-one relationship between input lines and successful calls to **gets_s**. Programs that use **gets** expect such a relationship.

K.3.6 General utilities <stdlib.h>

K.3.6.1 General

- 1 The header <stdlib.h> defines three types.
- 2 The types are

```
errno_t
```

which is type `int`; and

```
rsize_t
```

which is the type `size_t`; and

```
constraint_handler_t
```

which has the following definition

```
typedef void (*constraint_handler_t)(
    const char * restrict msg,
    void * restrict ptr,
    errno_t error);
```

K.3.6.2 Runtime-constraint handling

K.3.6.2.1 The `set_constraint_handler_s` function

Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
   #include <stdlib.h>
   constraint_handler_t set_constraint_handler_s(constraint_handler_t handler);
```

Description

- 2 The `set_constraint_handler_s` function sets the runtime-constraint handler to be `handler`. The runtime-constraint handler is the function to be called when a library function detects a runtime-constraint violation. Only the most recent handler registered with `set_constraint_handler_s` is called when a runtime-constraint violation occurs.
- 3 When the handler is called, it is passed the following arguments in the following order:
 1. A pointer to a character string describing the runtime-constraint violation.
 2. A null pointer or a pointer to an implementation-defined object.
 3. If the function calling the handler has a return type declared as `errno_t`, the return value of the function is passed. Otherwise, a positive value of type `errno_t` is passed.
- 4 The implementation has a default constraint handler that is used if no calls to the `set_constraint_handler_s` function have been made. The behavior of the default handler is implementation-defined, and it can cause the program to exit or abort.
- 5 If the `handler` argument to `set_constraint_handler_s` is a null pointer, the implementation default handler becomes the current constraint handler.

Returns

- 6 The `set_constraint_handler_s` function returns a pointer to the previously registered handler.⁴⁸²⁾

⁴⁸²⁾If the previous handler was registered by calling `set_constraint_handler_s` with a null pointer argument, a pointer to the implementation default handler is returned (not null).

K.3.6.2.2 The **abort_handler_s** function

Synopsis

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdlib.h>
    void abort_handler_s(const char * restrict msg, void * restrict ptr,
                        errno_t error);

```

Description

- 2 A pointer to the **abort_handler_s** function shall be a suitable argument to the **set_constraint_handler_s** function.
- 3 The **abort_handler_s** function writes a message on the standard error stream in an implementation-defined format. The message shall include the string pointed to by **msg**. The **abort_handler_s** function then calls the **abort** function.⁴⁸³⁾

Returns

- 4 The **abort_handler_s** function does not return to its caller.

K.3.6.2.3 The **ignore_handler_s** function

Synopsis

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdlib.h>
    void ignore_handler_s(const char * restrict msg, void * restrict ptr,
                        errno_t error);

```

Description

- 2 A pointer to the **ignore_handler_s** function shall be a suitable argument to the **set_constraint_handler_s** function.
- 3 The **ignore_handler_s** function simply returns to its caller.⁴⁸⁴⁾

Returns

- 4 The **ignore_handler_s** function returns no value.

K.3.6.3 Communication with the environment

K.3.6.3.1 The **getenv_s** function

Synopsis

```

1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdlib.h>
    errno_t getenv_s(size_t * restrict len, char * restrict value, rsize_t maxsize,
                    const char * restrict name);

```

Runtime-constraints

- 2 **name** shall not be a null pointer. **maxsize** shall not be greater than **RSIZE_MAX**. If **maxsize** is not equal to zero, then **value** shall not be a null pointer.

⁴⁸³⁾Many implementations invoke a debugger when the **abort** function is called.

⁴⁸⁴⁾If the runtime-constraint handler is set to the **ignore_handler_s** function, any library function in which a runtime-constraint violation occurs will return to its caller. The caller can determine whether a runtime-constraint violation occurred based on the library function's specification (usually, the library function returns a nonzero **errno_t**).

- 3 If there is a runtime-constraint violation, the integer pointed to by `len` is set to 0 (if `len` is not null), and the environment list is not searched.

Description

- 4 The `getenv_s` function searches an *environment list*, provided by the host environment, for a string that matches the string pointed to by `name`.
- 5 If that name is found then `getenv_s` performs the following actions. If `len` is not a null pointer, the length of the string associated with the matched list member is stored in the integer pointed to by `len`. If the length of the associated string is less than `maxsize`, then the associated string is copied to the array pointed to by `value`.
- 6 If that name is not found then `getenv_s` performs the following actions. If `len` is not a null pointer, zero is stored in the integer pointed to by `len`. If `maxsize` is greater than zero, then `value[0]` is set to the null character.
- 7 The set of environment names and the method for altering the environment list are implementation-defined. The `getenv_s` function is not required to avoid data races with other threads of execution that modify the environment list.⁴⁸⁵⁾

Returns

- 8 The `getenv_s` function returns zero if the specified `name` is found and the associated string was successfully stored in `value`. Otherwise, a nonzero value is returned.

K.3.6.4 Searching and sorting utilities

K.3.6.4.1 General

- 1 These utilities make use of a comparison function to search or sort arrays of unspecified type. Where an argument declared as `size_t nmemb` specifies the length of the array for a function, if `nmemb` has the value zero on a call to that function, then the comparison function is not called, a search finds no matching element, sorting performs no rearrangement, and the pointer to the array can be null.
- 2 The implementation shall ensure that the second argument of the comparison function (when called from `bsearch_s`), or both arguments (when called from `qsort_s`), are pointers to elements of the array.⁴⁸⁶⁾ The first argument when called from `bsearch_s` shall equal `key`.
- 3 The comparison function shall not alter the contents of either the array or search key. The implementation may reorder elements of the array between calls to the comparison function, but shall not otherwise alter the contents of any individual element.
- 4 When the same objects (consisting of `size` bytes, irrespective of their current positions in the array) are passed more than once to the comparison function, the results shall be consistent with one another. That is, for `qsort_s` they shall define a total ordering on the array, and for `bsearch_s` the same object shall always compare the same way with the key.
- 5 A sequence point occurs immediately before and immediately after each call to the comparison function, and also between any call to the comparison function and any movement of the objects passed as arguments to that call.

K.3.6.4.2 The `bsearch_s` generic function

Synopsis

```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <stdlib.h>
    QVoid *bsearch_s(const void *key, QVoid *base, rsize_t nmemb, rsize_t size,
                     int (*compar)(const void *k, const void *y, void *context),
                     void *context);
```

⁴⁸⁵⁾Many implementations provide non-standard functions that modify the environment list.

⁴⁸⁶⁾That is, if the value passed is `p`, then the following expressions are always valid and nonzero:

```
( (char *)p - (char *)base ) % size == 0
(char *)p >= (char *)base
(char *)p < (char *)base + nmemb * size
```


Runtime-constraints

- 2 Neither `nmemb` nor `size` shall be greater than `RSIZE_MAX`. If `nmemb` is not equal to zero, then none of `key`, `base`, or `compar` shall be a null pointer.
- 3 If there is a runtime-constraint violation, the `bsearch_s` generic function does not search the array.

Description

- 4 The `bsearch_s` generic function searches an array of `nmemb` objects, the initial element of which is pointed to by `base`, for an element that matches the object pointed to by `key`. The size of each element of the array is specified by `size`.
- 5 The comparison function pointed to by `compar` is called with three arguments. The first two point to the `key` object and to an array element, in that order. The function shall return an integer less than, equal to, or greater than zero if the `key` object is considered, respectively, to be less than, to match, or to be greater than the array element. The array shall consist of: all the elements that compare less than, all the elements that compare equal to, and all the elements that compare greater than the `key` object, in that order.⁴⁸⁷⁾ The third argument to the comparison function is the `context` argument passed to `bsearch_s`. The sole use of `context` by `bsearch_s` is to pass it to the comparison function.⁴⁸⁸⁾

Returns

- 6 The `bsearch_s` generic function returns a pointer to a matching element of the array, or a null pointer if no match is found or there is a runtime-constraint violation. If two elements compare as equal, which element is matched is unspecified.
- 7 The `bsearch_s` generic function is generic in the qualification of the type pointed to by the argument `base`. If this argument is a pointer to a `const`-qualified object type, the returned pointer will be a pointer to `const`-qualified `void`. Otherwise, the argument shall be a pointer to an unqualified object type or a null pointer constant,⁴⁸⁹⁾ and the returned pointer will be a pointer to unqualified `void`.
- 8 The external declaration of `bsearch_s` has the concrete type:

```
void * (const void *, const void *, rsize_t, rsize_t,
       int (*)(const void *, const void *), void *)
```

which supports all correct uses. If a macro definition of the generic function is suppressed to access an actual function, the external declaration with this concrete type is visible.⁴⁹⁰⁾

K.3.6.4.3 The `qsort_s` function**Synopsis**

- 1

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <stdlib.h>
errno_t qsort_s(void *base, rsize_t nmemb, rsize_t size,
               int (*compar)(const void *x, const void *y, void *context),
               void *context);
```

Runtime-constraints

- 2 Neither `nmemb` nor `size` shall be greater than `RSIZE_MAX`. If `nmemb` is not equal to zero, then neither `base` nor `compar` shall be a null pointer.
- 3 If there is a runtime-constraint violation, the `qsort_s` function does not sort the array.

⁴⁸⁷⁾In practice, this means that the entire array has been sorted according to the comparison function.

⁴⁸⁸⁾The `context` argument is for the use of the comparison function in performing its duties. For example, it can specify a collating sequence used by the comparison function.

⁴⁸⁹⁾If the argument is a null pointer and the call is executed, the behavior is undefined.

⁴⁹⁰⁾This is an obsolescent feature.

Description

- 4 The **qsort_s** function sorts an array of **nmemb** objects, the initial element of which is pointed to by **base**. The size of each object is specified by **size**.
- 5 The contents of the array are sorted into ascending order according to a comparison function pointed to by **compar**, which is called with three arguments. The first two point to the objects being compared. The function shall return an integer less than, equal to, or greater than zero if the first argument is considered to be respectively less than, equal to, or greater than the second. The third argument to the comparison function is the **context** argument passed to **qsort_s**. The sole use of **context** by **qsort_s** is to pass it to the comparison function.⁴⁹¹⁾
- 6 If two elements compare as equal, their relative order in the resulting sorted array is unspecified.

Returns

- 7 The **qsort_s** function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

K.3.6.5 Multibyte/wide character conversion functions**K.3.6.5.1 General**

- 1 The behavior of the multibyte character functions is affected by the **LC_CTYPE** category of the current locale. For a state-dependent encoding, each function is placed into its initial conversion state by a call for which its character pointer argument, **s**, is a null pointer. Subsequent calls with **s** as other than a null pointer cause the internal conversion state of the function to be altered as necessary. A call with **s** as a null pointer causes these functions to set the **int** pointed to by their **status** argument to a nonzero value if encodings have state dependency, and zero otherwise.⁴⁹²⁾

Changing the **LC_CTYPE** category causes the internal object describing the conversion state of these functions to have an indeterminate representation.

K.3.6.5.2 The **wctomb_s function****Synopsis**

- ```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdlib.h>
 errno_t wctomb_s(int * restrict status, char * restrict s, rsize_t smax,
 wchar_t wc);
```

**Runtime-constraints**

- 2 Let  $n$  denote the number of bytes needed to represent the multibyte character corresponding to the wide character given by **wc** (including any shift sequences).
- 3 If **s** is not a null pointer, then **smax** shall not be less than  $n$ , and **smax** shall not be greater than **RSIZE\_MAX**. If **s** is a null pointer, then **smax** shall equal zero.
- 4 If there is a runtime-constraint violation, **wctomb\_s** does not modify the **int** pointed to by **status**, and if **s** is not a null pointer, no more than **smax** elements in the array pointed to by **s** will be accessed.

**Description**

- 5 The **wctomb\_s** function determines  $n$  and stores the multibyte character representation of **wc** in the array whose first element is pointed to by **s** (if **s** is not a null pointer). The number of characters stored never exceeds **MB\_CUR\_MAX** or **smax**. If **wc** is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state, and the function is left in the initial conversion state.
- 6 The implementation shall behave as if no library function calls the **wctomb\_s** function.

<sup>491)</sup>The **context** argument is for the use of the comparison function in performing its duties. For example, it can specify a collating sequence used by the comparison function.

<sup>492)</sup>If the locale employs special bytes to change the shift state, these bytes do not produce separate wide character codes, but are grouped with an adjacent multibyte character.

- 7 If `s` is a null pointer, the `wctomb_s` function stores into the `int` pointed to by `status` a nonzero or zero value, if multibyte character encodings, respectively, do or do not have state-dependent encodings.
- 8 If `s` is not a null pointer, the `wctomb_s` function stores into the `int` pointed to by `status` either `n` or `-1` if `wc`, respectively, does or does not correspond to a valid multibyte character.
- 9 In no case will the `int` pointed to by `status` be set to a value greater than the `MB_CUR_MAX` macro.

#### Returns

- 10 The `wctomb_s` function returns zero if successful, and a nonzero value if there was a runtime-constraint violation or `wc` did not correspond to a valid multibyte character.

### K.3.6.6 Multibyte/wide string conversion functions

#### K.3.6.6.1 General

- 1 The behavior of the multibyte string functions is affected by the `LC_CTYPE` category of the current locale.

#### K.3.6.6.2 The `mbstowcs_s` function

##### Synopsis

```
1 #include <stdlib.h>
 errno_t mbstowcs_s(size_t * restrict retval, wchar_t * restrict dst,
 rsize_t dstmax, const char * restrict src, rsize_t len);
```

##### Runtime-constraints

- 2 Neither `retval` nor `src` shall be a null pointer. If `dst` is not a null pointer, then neither `len` nor `dstmax` shall be greater than `RSIZE_MAX/sizeof(wchar_t)`. If `dst` is a null pointer, then `dstmax` shall equal zero. If `dst` is not a null pointer, then `dstmax` shall not equal zero. If `dst` is not a null pointer and `len` is not less than `dstmax`, then a null character shall occur within the first `dstmax` multibyte characters of the array pointed to by `src`.
- 3 If there is a runtime-constraint violation, then `mbstowcs_s` does the following. If `retval` is not a null pointer, then `mbstowcs_s` sets `*retval` to `(size_t) (-1)`. If `dst` is not a null pointer and `dstmax` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then `mbstowcs_s` sets `dst[0]` to the null wide character.

##### Description

- 4 The `mbstowcs_s` function converts a sequence of multibyte characters that begins in the initial shift state from the array pointed to by `src` into a sequence of corresponding wide characters. If `dst` is not a null pointer, the converted characters are stored into the array pointed to by `dst`. Conversion continues up to and including a terminating null character, which is also stored. Conversion stops earlier in two cases: when a sequence of bytes is encountered that does not form a valid multibyte character, or (if `dst` is not a null pointer) when `len` wide characters have been stored into the array pointed to by `dst`.<sup>493</sup> If `dst` is not a null pointer and no null wide character was stored into the array pointed to by `dst`, then `dst[len]` is set to the null wide character. Each conversion takes place as if by a call to the `mbrtowc` function.
- 5 Regardless of whether `dst` is or is not a null pointer, if the input conversion encounters a sequence of bytes that do not form a valid multibyte character, an encoding error occurs: the `mbstowcs_s` function stores the value `(size_t) (-1)` into `*retval`. Otherwise, the `mbstowcs_s` function stores into `*retval` the number of multibyte characters successfully converted, not including the terminating null character (if any).
- 6 All elements following the terminating null wide character (if any) written by `mbstowcs_s` in the array of `dstmax` wide characters pointed to by `dst` take unspecified values when `mbstowcs_s`

<sup>493</sup>Thus, the value of `len` is ignored if `dst` is a null pointer.

returns.<sup>494)</sup>

- 7 If copying takes place between objects that overlap, the objects take on unspecified values.

#### Returns

- 8 The **mbstowcs\_s** function returns zero if no runtime-constraint violation and no encoding error occurred. Otherwise, a nonzero value is returned.

#### K.3.6.6.3 The **wcstombs\_s** function

##### Synopsis

```
1 #include <stdlib.h>
 errno_t wcstombs_s(size_t * restrict retval, char * restrict dst, rsize_t dstmax,
 const wchar_t * restrict src, rsize_t len);
```

##### Runtime-constraints

- 2 Neither **retval** nor **src** shall be a null pointer. If **dst** is not a null pointer, then **len** shall not be greater than **RSIZE\_MAX/sizeof(wchar\_t)** and **dstmax** shall be nonzero and not greater than **RSIZE\_MAX**. If **dst** is a null pointer, then **dstmax** shall equal zero. If **dst** is not a null pointer and **len** is not less than **dstmax**, then the conversion shall have been stopped (see the following) because a terminating null wide character was reached or because an encoding error occurred.
- 3 If there is a runtime-constraint violation, then **wcstombs\_s** does the following. If **retval** is not a null pointer, then **wcstombs\_s** sets **\*retval** to **(size\_t)(-1)**. If **dst** is not a null pointer and **dstmax** is greater than zero and not greater than **RSIZE\_MAX**, then **wcstombs\_s** sets **dst[0]** to the null character.

##### Description

- 4 The **wcstombs\_s** function converts a sequence of wide characters from the array pointed to by **src** into a sequence of corresponding multibyte characters that begins in the initial shift state. If **dst** is not a null pointer, the converted characters are then stored into the array pointed to by **dst**. Conversion continues up to and including a terminating null wide character, which is also stored. Conversion stops earlier in two cases:

- when a wide character is reached that does not correspond to a valid multibyte character;
- (if **dst** is not a null pointer) when the next multibyte character would exceed the limit of *n* total bytes to be stored into the array pointed to by **dst**. If the wide character being converted is the null wide character, then *n* is the lesser of **len** or **dstmax**. Otherwise, *n* is the lesser of **len** or **dstmax-1**.

If the conversion stops without converting a null wide character and **dst** is not a null pointer, then a null character is stored into the array pointed to by **dst** immediately following any multibyte characters already stored. Each conversion takes place as if by a call to the **wcrtomb** function.<sup>495)</sup>

- 5 Regardless of whether **dst** is or is not a null pointer, if the input conversion encounters a wide character that does not correspond to a valid multibyte character, an encoding error occurs: the **wcstombs\_s** function stores the value **(size\_t)(-1)** into **\*retval**. Otherwise, the **wcstombs\_s** function stores into **\*retval** the number of bytes in the resulting multibyte character sequence, not including the terminating null character (if any).
- 6 All elements following the terminating null character (if any) written by **wcstombs\_s** in the array of **dstmax** elements pointed to by **dst** take unspecified values when **wcstombs\_s** returns.<sup>496)</sup>
- 7 If copying takes place between objects that overlap, the objects take on unspecified values.

<sup>494)</sup>This allows an implementation to attempt converting the multibyte string before discovering a terminating null character did not occur where required.

<sup>495)</sup>If conversion stops because a terminating null wide character has been reached, the bytes stored include those necessary to reach the initial shift state immediately before the null byte. However, if the conversion stops before a terminating null wide character has been reached, the result will be null terminated, but potentially not end in the initial shift state.

<sup>496)</sup>When **len** is not less than **dstmax**, the implementation can fill the array before discovering a runtime-constraint violation.

**Returns**

- 8 The `wcstombs_s` function returns zero if no runtime-constraint violation and no encoding error occurred. Otherwise, a nonzero value is returned.

**K.3.7 String handling <string.h>****K.3.7.1 General**

- 1 The header <string.h> defines two types.
- 2 The types are

```
errno_t
```

which is type `int`; and

```
rsize_t
```

which is the type `size_t`.

**K.3.7.2 Copying functions****K.3.7.2.1 The `memcpy_s` function****Synopsis**

- ```
1  #define __STDC_WANT_LIB_EXT1__ 1
    #include <string.h>
    errno_t memcpy_s(void * restrict s1, rsize_t slmax, const void * restrict s2,
                    rsize_t n);
```

Runtime-constraints

- 2 Neither `s1` nor `s2` shall be a null pointer. Neither `slmax` nor `n` shall be greater than `RSIZE_MAX`. `n` shall not be greater than `slmax`. Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, the `memcpy_s` function stores zeros in the first `slmax` characters of the object pointed to by `s1` if `s1` is not a null pointer and `slmax` is not greater than `RSIZE_MAX`.

Description

- 4 The `memcpy_s` function copies `n` characters from the object pointed to by `s2` into the object pointed to by `s1`.

Returns

- 5 The `memcpy_s` function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

K.3.7.2.2 The `memmove_s` function**Synopsis**

- ```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t memmove_s(void *s1, rsize_t slmax, const void *s2, rsize_t n);
```

**Runtime-constraints**

- 2 Neither `s1` nor `s2` shall be a null pointer. Neither `slmax` nor `n` shall be greater than `RSIZE_MAX`. `n` shall not be greater than `slmax`.
- 3 If there is a runtime-constraint violation, the `memmove_s` function stores zeros in the first `slmax` characters of the object pointed to by `s1` if `s1` is not a null pointer and `slmax` is not greater than `RSIZE_MAX`.

**Description**

- 4 The **memmove\_s** function copies *n* characters from the object pointed to by *s2* into the object pointed to by *s1*. This copying takes place as if the *n* characters from the object pointed to by *s2* are first copied into a temporary array of *n* characters that does not overlap the objects pointed to by *s1* or *s2*, and then the *n* characters from the temporary array are copied into the object pointed to by *s1*.

**Returns**

- 5 The **memmove\_s** function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

**K.3.7.2.3 The **strcpy\_s** function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t strcpy_s(char * restrict s1, rsize_t s1max, const char * restrict s2);
```

**Runtime-constraints**

- 2 Neither *s1* nor *s2* shall be a null pointer. *s1max* shall not be greater than **RSIZE\_MAX**. *s1max* shall not equal zero. *s1max* shall be greater than **strnlen\_s**(*s2*, *s1max*). Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, then if *s1* is not a null pointer and *s1max* is greater than zero and not greater than **RSIZE\_MAX**, then **strcpy\_s** sets *s1*[0] to the null character.

**Description**

- 4 The **strcpy\_s** function copies the string pointed to by *s2* (including the terminating null character) into the array pointed to by *s1*.
- 5 All elements following the terminating null character (if any) written by **strcpy\_s** in the array of *s1max* characters pointed to by *s1* take unspecified values when **strcpy\_s** returns.<sup>497)</sup>

**Returns**

- 6 The **strcpy\_s** function returns zero<sup>498)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

**K.3.7.2.4 The **strncpy\_s** function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t strncpy_s(char * restrict s1, rsize_t s1max, const char * restrict s2,
 rsize_t n);
```

**Runtime-constraints**

- 2 Neither *s1* nor *s2* shall be a null pointer. Neither *s1max* nor *n* shall be greater than **RSIZE\_MAX**. *s1max* shall not equal zero. If *n* is not less than *s1max*, then *s1max* shall be greater than **strnlen\_s**(*s2*, *s1max*). Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, then if *s1* is not a null pointer and *s1max* is greater than zero and not greater than **RSIZE\_MAX**, then **strncpy\_s** sets *s1*[0] to the null character.

<sup>497)</sup>This allows an implementation to copy characters from *s2* to *s1* while simultaneously checking if any of those characters are null. Such an approach can write a character to every element of *s1* before discovering that the first element was set to the null character.

<sup>498)</sup>A zero return value implies that all the requested characters from the string pointed to by *s2* fit within the array pointed to by *s1* and that the result in *s1* is null terminated.

## Description

- 4 The **strncpy\_s** function copies not more than *n* successive characters (characters that follow a null character are not copied) from the array pointed to by *s2* to the array pointed to by *s1*. If no null character was copied from *s2*, then *s1[n]* is set to a null character.
- 5 All elements following the terminating null character (if any) written by **strncpy\_s** in the array of *s1max* characters pointed to by *s1* take unspecified values when **strncpy\_s** returns a nonzero value.<sup>499)</sup>

## Returns

- 6 The **strncpy\_s** function returns zero<sup>500)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.
- 7 EXAMPLE The **strncpy\_s** function can be used to copy a string without the danger that the result will not be null terminated or that characters will be written past the end of the destination array.

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <string.h>
/* ... */
char src1[100] = "hello";
char src2[7] = {'g', 'o', 'o', 'd', ' ', 'y', 'e'};
char dst1[6], dst2[5], dst3[5];
int r1, r2, r3;
r1 = strncpy_s(dst1, 6, src1, 100);
r2 = strncpy_s(dst2, 5, src2, 7);
r3 = strncpy_s(dst3, 5, src2, 4);
```

The first call will assign to *r1* the value zero and to *dst1* the sequence `hello\0`.

The second call will assign to *r2* a nonzero value and to *dst2* the sequence `\0`.

The third call will assign to *r3* the value zero and to *dst3* the sequence `good\0`.

## K.3.7.3 Concatenation functions

### K.3.7.3.1 The **strcat\_s** function

#### Synopsis

- 1 

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <string.h>
errno_t strcat_s(char * restrict s1, rsize_t s1max, const char * restrict s2);
```

## Runtime-constraints

- 2 Let *m* denote the value *s1max* - **strlen\_s**(*s1*, *s1max*) upon entry to **strcat\_s**.
- 3 Neither *s1* nor *s2* shall be a null pointer. *s1max* shall not be greater than **RSIZE\_MAX**. *s1max* shall not equal zero. *m* shall not equal zero.<sup>501)</sup> *m* shall be greater than **strlen\_s**(*s2*, *m*). Copying shall not take place between objects that overlap.
- 4 If there is a runtime-constraint violation, then if *s1* is not a null pointer and *s1max* is greater than zero and not greater than **RSIZE\_MAX**, then **strcat\_s** sets *s1*[0] to the null character.

## Description

- 5 The **strcat\_s** function appends a copy of the string pointed to by *s2* (including the terminating null character) to the end of the string pointed to by *s1*. The initial character from *s2* overwrites the null character at the end of *s1*.

<sup>499)</sup>This allows an implementation to copy characters from *s2* to *s1* while simultaneously checking if any of those characters are null. Such an approach can write a character to every element of *s1* before discovering that the first element was set to the null character.

<sup>500)</sup>A zero return value implies that all of the requested characters from the string pointed to by *s2* fit within the array pointed to by *s1* and that the result in *s1* is null terminated.

<sup>501)</sup>Zero means that *s1* was not null terminated upon entry to **strcat\_s**.



- 6 All elements following the terminating null character (if any) written by **strcat\_s** in the array of **s1max** characters pointed to by **s1** take unspecified values when **strcat\_s** returns.<sup>502)</sup>

#### Returns

- 7 The **strcat\_s** function returns zero<sup>503)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

#### K.3.7.3.2 The **strncat\_s** function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t strncat_s(char * restrict s1, rsize_t s1max, const char * restrict s2,
 rsize_t n);
```

##### Runtime-constraints

- 2 Let  $m$  denote the value **s1max** - **strlen\_s(s1, s1max)** upon entry to **strncat\_s**.
- 3 Neither **s1** nor **s2** shall be a null pointer. Neither **s1max** nor **n** shall be greater than **RSIZE\_MAX**. **s1max** shall not equal zero.  $m$  shall not equal zero.<sup>504)</sup> If **n** is not less than  $m$ , then  $m$  shall be greater than **strlen\_s(s2, m)**. Copying shall not take place between objects that overlap.
- 4 If there is a runtime-constraint violation, then if **s1** is not a null pointer and **s1max** is greater than zero and not greater than **RSIZE\_MAX**, then **strncat\_s** sets **s1[0]** to the null character.

##### Description

- 5 The **strncat\_s** function appends not more than **n** successive characters (characters that follow a null character are not copied) from the array pointed to by **s2** to the end of the string pointed to by **s1**. The initial character from **s2** overwrites the null character at the end of **s1**. If no null character was copied from **s2**, then **s1[s1max - m + n]** is set to a null character.
- 6 All elements following the terminating null character (if any) written by **strncat\_s** in the array of **s1max** characters pointed to by **s1** take unspecified values when **strncat\_s** returns.<sup>505)</sup>

#### Returns

- 7 The **strncat\_s** function returns zero<sup>506)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.
- 8 EXAMPLE The **strncat\_s** function can be used to copy a string without the danger that the result will not be null terminated or that characters will be written past the end of the destination array.

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <string.h>
/* ... */
char s1[100] = "good";
char s2[6] = "hello";
char s3[6] = "hello";
char s4[7] = "abc";
char s5[1000] = "bye";
int r1, r2, r3, r4;
r1 = strncat_s(s1, 100, s5, 1000);
```

<sup>502)</sup>This allows an implementation to append characters from **s2** to **s1** while simultaneously checking if any of those characters are null. Such an approach can write a character to every element of **s1** before discovering that the first element was set to the null character.

<sup>503)</sup>A zero return value implies that all the requested characters from the string pointed to by **s2** were appended to the string pointed to by **s1** and that the result in **s1** is null terminated.

<sup>504)</sup>Zero means that **s1** was not null terminated upon entry to **strncat\_s**.

<sup>505)</sup>This allows an implementation to append characters from **s2** to **s1** while simultaneously checking if any of those characters are null. Such an approach can write a character to every element of **s1** before discovering that the first element was set to the null character.

<sup>506)</sup>A zero return value implies that all the requested characters from the string pointed to by **s2** were appended to the string pointed to by **s1** and that the result in **s1** is null terminated.



```

r2 = strncat_s(s2, 6, "", 1);
r3 = strncat_s(s3, 6, "X", 2);
r4 = strncat_s(s4, 7, "defghijklmn", 3);

```

After the first call `r1` will have the value zero and `s1` will contain the sequence `goodbye\0`.

After the second call `r2` will have the value zero and `s2` will contain the sequence `hello\0`.

After the third call `r3` will have a nonzero value and `s3` will contain the sequence `\0`.

After the fourth call `r4` will have the value zero and `s4` will contain the sequence `abcdef\0`.

### K.3.7.4 Search functions

#### K.3.7.4.1 The `strtok_s` function

##### Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 char *strtok_s(char * restrict s1, rsize_t * restrict s1max,
 const char * restrict s2, char ** restrict ptr);

```

##### Runtime-constraints

- 2 None of `s1max`, `s2`, or `ptr` shall be a null pointer. If `s1` is a null pointer, then `*ptr` shall not be a null pointer. The value of `*s1max` shall not be greater than `RSIZE_MAX`. The end of the token found shall occur within the first `*s1max` characters of `s1` for the first call, and shall occur within the first `*s1max` characters of where searching resumes on subsequent calls.
- 3 If there is a runtime-constraint violation, the `strtok_s` function does not indirect through the `s1` or `s2` pointers, and does not store a value in the object pointed to by `ptr`.

##### Description

- 4 A sequence of calls to the `strtok_s` function breaks the string pointed to by `s1` into a sequence of tokens, each of which is delimited by a character from the string pointed to by `s2`. The fourth argument points to a caller-provided `char` pointer into which the `strtok_s` function stores information necessary for it to continue scanning the same string.
- 5 The first call in a sequence has a non-null first argument and `s1max` points to an object whose value is the number of elements in the character array pointed to by the first argument. The first call stores an initial value in the object pointed to by `ptr` and updates the value pointed to by `s1max` to reflect the number of elements that remain in relation to `ptr`. Subsequent calls in the sequence have a null first argument and the objects pointed to by `s1max` and `ptr` are required to have the values stored by the previous call in the sequence, which are then updated. The separator string pointed to by `s2` can be different from call to call.
- 6 The first call in the sequence searches the string pointed to by `s1` for the first character that is *not* contained in the current separator string pointed to by `s2`. If no such character is found, then there are no tokens in the string pointed to by `s1` and the `strtok_s` function returns a null pointer. If such a character is found, it is the start of the first token.
- 7 The `strtok_s` function then searches from there for the first character in `s1` that *is* contained in the current separator string. If no such character is found, the current token extends to the end of the string pointed to by `s1`, and subsequent searches in the same string for a token return a null pointer. If such a character is found, it is overwritten by a null character, which terminates the current token.
- 8 In all cases, the `strtok_s` function stores sufficient information in the pointer pointed to by `ptr` so that subsequent calls, with a null pointer for `s1` and the unmodified pointer value for `ptr`, shall start searching just past the element overwritten by a null character (if any).

##### Returns

- 9 The `strtok_s` function returns a pointer to the first character of a token, or a null pointer if there is no token or there is a runtime-constraint violation.

## 10 EXAMPLE

```

#define __STDC_WANT_LIB_EXT1__ 1
#include <string.h>
static char str1[] = "?a???b,,,#c";
static char str2[] = "\t \t";
char *t, *ptr1, *ptr2;
rsize_t max1 = _Lengthof(str1);
rsize_t max2 = _Lengthof(str2);

t = strtok_s(str1, &max1, "?", &ptr1); // t points to the token "a"
t = strtok_s(nullptr, &max1, ",", &ptr1); // t points to the token "???b"
t = strtok_s(str2, &max2, " \t", &ptr2); // t is a null pointer
t = strtok_s(nullptr, &max1, "#", &ptr1); // t points to the token "c"
t = strtok_s(nullptr, &max1, "?", &ptr1); // t is a null pointer

```

## K.3.7.5 Miscellaneous functions

K.3.7.5.1 The `memset_s` function

## Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t memset_s(void *s, rsize_t smax, int c, rsize_t n)

```

## Runtime-constraints

- 2 `s` shall not be a null pointer. Neither `smax` nor `n` shall be greater than `RSIZE_MAX`. `n` shall not be greater than `smax`.
- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `smax` is not greater than `RSIZE_MAX`, the `memset_s` function stores the value of `c` (converted to an **unsigned char**) into each of the first `smax` characters of the object pointed to by `s`.

## Description

- 4 The `memset_s` function copies the value of `c` (converted to an **unsigned char**) into each of the first `n` characters of the object pointed to by `s`. Unlike `memset`, any call to the `memset_s` function shall be evaluated strictly according to the rules of the abstract machine as described in 5.2.2.4. That is, any call to the `memset_s` function shall assume that the memory indicated by `s` and `n` may be accessible in the future and thus contains the values indicated by `c`.

## Returns

- 5 The `memset_s` function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

K.3.7.5.2 The `strerror_s` function

## Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 errno_t strerror_s(char *s, rsize_t maxsize, errno_t errnum);

```

## Runtime-constraints

- 2 `s` shall not be a null pointer. `maxsize` shall not be greater than `RSIZE_MAX`. `maxsize` shall not equal zero.
- 3 If there is a runtime-constraint violation, then the array (if any) pointed to by `s` is not modified.

## Description

- 4 The `strerror_s` function maps the number in `errnum` to a locale-specific message string. Typically, the values for `errnum` come from `errno`, but `strerror_s` shall map any value of type `int` to a

message.

- 5 If the length of the desired string is less than `maxsize`, then the string is copied to the array pointed to by `s`.
- 6 Otherwise, if `maxsize` is greater than zero, then `maxsize-1` characters are copied from the string to the array pointed to by `s` and then `s[maxsize-1]` is set to the null character. Then, if `maxsize` is greater than 3, then `s[maxsize-2]`, `s[maxsize-3]`, and `s[maxsize-4]` are set to the character period (`.`).

#### Returns

- 7 The `strerror_s` function returns zero if the length of the desired string was less than `maxsize` and there was no runtime-constraint violation. Otherwise, the `strerror_s` function returns a nonzero value.

### K.3.7.5.3 The `strerrorlen_s` function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 size_t strerrorlen_s(errno_t errnum);
```

#### Description

- 2 The `strerrorlen_s` function calculates the length of the (untruncated) locale-specific message string that the `strerror_s` function maps to `errnum`.

#### Returns

- 3 The `strerrorlen_s` function returns the number of characters (not including the null character) in the full message string.

### K.3.7.5.4 The `strnlen_s` function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <string.h>
 size_t strnlen_s(const char *s, size_t maxsize);
```

#### Description

- 2 The `strnlen_s` function counts not more than `maxsize` characters (a null character and characters that follow it are not counted) in the array to which `s` points. At most the first `maxsize` characters of `s` shall be accessed by `strnlen_s`. The implementation shall behave as if it reads the characters sequentially and stops as soon as a null character is found.

#### Returns

- 3 If `s` is a null pointer,<sup>507</sup> then the `strnlen_s` function returns zero.
- 4 Otherwise, the `strnlen_s` function returns the number of characters that precede the terminating null character. If there is no null character in the first `maxsize` characters of `s` then `strnlen_s` returns `maxsize`.

## K.3.8 Date and time `<time.h>`

### K.3.8.1 General

- 1 The header `<time.h>` defines two types.
- 2 The types are

<sup>507</sup>The `strnlen_s` function has no runtime-constraints. This lack of runtime-constraints along with the values returned for a null pointer or an unterminated string argument make `strnlen_s` useful in algorithms that gracefully handle such exceptional data.

**errno\_t**

which is type **int**; and

**rsizet**

which is the type **size\_t**.

### K.3.8.2 Components of time

- 1 A broken-down time is *normalized* if the values of the members of the **tm** structure are in their normal ranges.<sup>508)</sup>

### K.3.8.3 Time conversion functions

#### K.3.8.3.1 General

- 1 Like the **strftime** function, the **asctime\_s** and **ctime\_s** functions do not return a pointer to a static object, and other library functions are permitted to call them.

#### K.3.8.3.2 The **asctime\_s** function

##### Synopsis

- 1
 

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <time.h>
errno_t asctime_s(char *s, rsizet maxsize, const struct tm *timeptr);
```

##### Runtime-constraints

- 2 Neither **s** nor **timeptr** shall be a null pointer. **maxsize** shall not be less than 26 and shall not be greater than **RSIZE\_MAX**. The broken-down time pointed to by **timeptr** shall be normalized. The calendar year represented by the broken-down time pointed to by **timeptr** shall not be less than calendar year 0 and shall not be greater than calendar year 9999.
- 3 If there is a runtime-constraint violation, there is no attempt to convert the time, and **s[0]** is set to a null character if **s** is not a null pointer and **maxsize** is not zero and is not greater than **RSIZE\_MAX**.

##### Description

- 4 The **asctime\_s** function converts the normalized broken-down time in the structure pointed to by **timeptr** into a 26 character (including the null character) string in the form

Sun Sep 16 01:03:52 1973\n\n0

The fields making up this string are (in order):

1. The name of the day of the week represented by **timeptr->tm\_wday** using the following three character weekday names: Sun, Mon, Tue, Wed, Thu, Fri, and Sat.
2. The character space.
3. The name of the month represented by **timeptr->tm\_mon** using the following three character month names: Jan, Feb, Mar, Apr, May, Jun, Jul, Aug, Sep, Oct, Nov, and Dec.
4. The character space.
5. The value of **timeptr->tm\_mday** as if printed using the **fprintf** format "%2d".
6. The character space.
7. The value of **timeptr->tm\_hour** as if printed using the **fprintf** format "%.2d".
8. The character colon.

<sup>508)</sup>The normal ranges are defined in 7.30.1.

9. The value of `timeptr->tm_min` as if printed using the `fprintf` format `"%.2d"`.
10. The character colon.
11. The value of `timeptr->tm_sec` as if printed using the `fprintf` format `"%.2d"`.
12. The character space.
13. The value of `timeptr->tm_year` + 1900 as if printed using the `fprintf` format `"%4d"`.
14. The character new line.
15. The null character.

### Recommended practice

The `strftime` function allows more flexible formatting and supports locale-specific behavior. If you do not require the exact form of the result string produced by the `asctime_s` function, consider using the `strftime` function instead.

### Returns

- 5 The `asctime_s` function returns zero if the time was successfully converted and stored into the array pointed to by `s`. Otherwise, it returns a nonzero value.

### K.3.8.3.3 The `ctime_s` function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <time.h>
 errno_t ctime_s(char *s, rsize_t maxsize, const time_t *timer);
```

### Runtime-constraints

- 2 Neither `s` nor `timer` shall be a null pointer. `maxsize` shall not be less than 26 and shall not be greater than `RSIZE_MAX`.
- 3 If there is a runtime-constraint violation, `s[0]` is set to a null character if `s` is not a null pointer and `maxsize` is not equal zero and is not greater than `RSIZE_MAX`.

### Description

- 4 The `ctime_s` function converts the calendar time pointed to by `timer` to local time in the form of a string. It is equivalent to

```
asctime_s(s, maxsize, localtime_s(timer, &(struct tm){ 0 }))
```

### Recommended practice

The `strftime` function allows more flexible formatting and supports locale-specific behavior. If you do not require the exact form of the result string produced by the `ctime_s` function, consider using the `strftime` function instead.

### Returns

- 5 The `ctime_s` function returns zero if the time was successfully converted and stored into the array pointed to by `s`. Otherwise, it returns a nonzero value.

### K.3.8.3.4 The `gmtime_s` function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <time.h>
 struct tm *gmtime_s(const time_t * restrict timer, struct tm * restrict result);
```

**Runtime-constraints**

- 2 Neither `timer` nor `result` shall be a null pointer.
- 3 If there is a runtime-constraint violation, there is no attempt to convert the time.

**Description**

- 4 The `gmtime_s` function converts the calendar time pointed to by `timer` into a broken-down time, expressed as UTC. The broken-down time is stored in the structure pointed to by `result`.

**Returns**

- 5 The `gmtime_s` function returns `result`, or a null pointer if the specified time cannot be converted to UTC or there is a runtime-constraint violation.

**K.3.8.3.5 The `localtime_s` function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <time.h>
 struct tm *localtime_s(const time_t * restrict timer, struct tm * restrict result
);
```

**Runtime-constraints**

- 2 Neither `timer` nor `result` shall be a null pointer.
- 3 If there is a runtime-constraint violation, there is no attempt to convert the time.

**Description**

- 4 The `localtime_s` function converts the calendar time pointed to by `timer` into a broken-down time, expressed as local time. The broken-down time is stored in the structure pointed to by `result`.

**Returns**

- 5 The `localtime_s` function returns `result`, or a null pointer if the specified time cannot be converted to local time or there is a runtime-constraint violation.

**K.3.9 Extended multibyte and wide character utilities `<wchar.h>`****K.3.9.1 General**

- 1 The header `<wchar.h>` defines two types.
- 2 The types are

```
errno_t
```

which is type `int`; and

```
rsize_t
```

which is the type `size_t`.

- 3 Unless explicitly stated otherwise, if the execution of a function described in this subclause causes copying to take place between objects that overlap, the objects take on unspecified values.

**K.3.9.2 Formatted wide character input/output functions****K.3.9.2.1 The `fwprintf_s` function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int fwprintf_s(FILE * restrict stream, const wchar_t * restrict format, ...);
```

**Runtime-constraints**

- 2 Neither `stream` nor `format` shall be a null pointer. The `%n` specifier<sup>509)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by `format`. Any argument to `fwprintf_s` corresponding to a `%s` specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `fwprintf_s` function does not attempt to produce further output, and it is unspecified to what extent `fwprintf_s` produced output before discovering the runtime-constraint violation.

**Description**

- 4 The `fwprintf_s` function is equivalent to the `fwprintf` function except for the previously listed explicit runtime-constraints.

**Returns**

- 5 The `fwprintf_s` function returns the number of wide characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

**K.3.9.2.2 The `fwscanf_s` function****Synopsis**

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdio.h>
 #include <wchar.h>
 int fwscanf_s(FILE * restrict stream, const wchar_t * restrict format, ...);

```

**Runtime-constraints**

- 2 Neither `stream` nor `format` shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `fwscanf_s` function does not attempt to perform further input, and it is unspecified to what extent `fwscanf_s` performed input before discovering the runtime-constraint violation.

**Description**

- 4 The `fwscanf_s` function is equivalent to `fwscanf` except that the `c`, `s`, and `[` conversion specifiers apply to a pair of arguments (unless assignment suppression is indicated by a `*`). The first of these arguments is the same as for `fwscanf`. That argument is immediately followed in the argument list by the second argument, which has type `size_t` and gives the number of elements in the array pointed to by the first argument of the pair. If the first argument points to a scalar object, it is considered to be an array of one element.<sup>510)</sup>
- 5 A matching failure occurs if the number of elements in a receiving object is insufficient to hold the converted input (including any trailing null character).

**Returns**

- 6 The `fwscanf_s` function returns the value of the macro `EOF` if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the `fwscanf_s` function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

<sup>509)</sup>It is not a runtime-constraint violation for the wide characters `%n` to appear in sequence in the wide string pointed at by `format` when those wide characters are not interpreted as a `%n` specifier. For example, if the entire format string was `L"%n%n"`.

<sup>510)</sup>If the format is known at translation time, an implementation can issue a diagnostic for any argument used to store the result from a `c`, `s`, or `[` conversion specifier if that argument is not followed by an argument of a type compatible with `rsize_t`. A limited amount of checking can be done if even if the format is not known at translation time. For example, an implementation can issue a diagnostic for each argument after `format` that has of type pointer to one of `char`, `signed char`, `unsigned char`, or `void` that is not followed by an argument of a type compatible with `rsize_t`. The diagnostic can warn that unless the pointer is being used with a conversion specifier using the `hh` length modifier, a length argument is expected to follow the pointer argument. Another useful diagnostic can flag any non-pointer argument following `format` that did not have a type compatible with `rsize_t`.



K.3.9.2.3 The `snwprintf_s` function

## Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int snwprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
 ...);

```

## Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. `n` shall neither equal zero nor be greater than `RSIZE_MAX/sizeof(wchar_t)`. The `%n` specifier<sup>511)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by `format`. Any argument to `snwprintf_s` corresponding to a `%s` specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `n` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then the `snwprintf_s` function sets `s[0]` to the null wide character.

## Description

- 4 The `snwprintf_s` function is equivalent to the `swprintf` function except for the previously listed explicit runtime-constraints.
- 5 The `snwprintf_s` function, unlike `swprintf_s`, will truncate the result to fit within the array pointed to by `s`.

## Returns

- 6 The `snwprintf_s` function returns the number of wide characters that would have been written had `n` been sufficiently large, not counting the terminating wide null character, or a negative value if a runtime-constraint violation occurred. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than `n`.

K.3.9.2.4 The `swprintf_s` function

## Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int swprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
 ...);

```

## Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. `n` shall neither equal zero nor be greater than `RSIZE_MAX/sizeof(wchar_t)`. The number of wide characters (including the trailing null) required for the result to be written to the array pointed to by `s` shall not be greater than `n`. The `%n` specifier<sup>512)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by `format`. Any argument to `swprintf_s` corresponding to a `%s` specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `n` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then the `swprintf_s` function sets `s[0]` to the null wide character.

<sup>511)</sup>It is not a runtime-constraint violation for the wide characters `%n` to appear in sequence in the wide string pointed at by `format` when those wide characters are not a interpreted as a `%n` specifier. For example, if the entire format string was `L"%n%n"`.

<sup>512)</sup>It is not a runtime-constraint violation for the wide characters `%n` to appear in sequence in the wide string pointed at by `format` when those wide characters are not a interpreted as a `%n` specifier. For example, if the entire format string was `L"%n%n"`.



**Description**

- 4 The **swprintf\_s** function is equivalent to the **swprintf** function except for the previously listed explicit runtime-constraints.
- 5 The **swprintf\_s** function, unlike **snwprintf\_s**, treats a result too big for the array pointed to by **s** as a runtime-constraint violation.

**Returns**

- 6 If no runtime-constraint violation occurred, the **swprintf\_s** function returns the number of wide characters written in the array, not counting the terminating null wide character. If an encoding error occurred or if **n** or more wide characters are requested to be written, **swprintf\_s** returns a negative value. If any other runtime-constraint violation occurred, **swprintf\_s** returns zero.

**K.3.9.2.5 The swscanf\_s function****Synopsis**

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int swscanf_s(const wchar_t * restrict s, const wchar_t * restrict format, ...);

```

**Runtime-constraints**

- 2 Neither **s** nor **format** shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **swscanf\_s** function does not attempt to perform further input, and it is unspecified to what extent **swscanf\_s** performed input before discovering the runtime-constraint violation.

**Description**

- 4 The **swscanf\_s** function is equivalent to **fwscanf\_s**, except that the argument **s** specifies a wide string from which the input is to be obtained, rather than from a stream. Reaching the end of the wide string is equivalent to encountering end-of-file for the **fwscanf\_s** function.

**Returns**

- 5 The **swscanf\_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **swscanf\_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**K.3.9.2.6 The vfwprintf\_s function****Synopsis**

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <stdio.h>
 #include <wchar.h>
 int vfwprintf_s(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);

```

**Runtime-constraints**

- 2 Neither **stream** nor **format** shall be a null pointer. The **%n** specifier<sup>513</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by **format**. Any argument to **vfwprintf\_s** corresponding to a **%s** specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **vfwprintf\_s** function does not attempt to produce further output, and it is unspecified to what extent **vfwprintf\_s** produced output before discovering

<sup>513</sup>) It is not a runtime-constraint violation for the wide characters **%n** to appear in sequence in the wide string pointed at by **format** when those wide characters are not interpreted as a **%n** specifier. For example, if the entire format string was **L"%n"**.

the runtime-constraint violation.

### Description

- 4 The **vfwprintf\_s** function is equivalent to the **vfwprintf** function except for the previously listed explicit runtime-constraints.

### Returns

- 5 The **vfwprintf\_s** function returns the number of wide characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

#### K.3.9.2.7 The **vfwscanf\_s** function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <stdio.h>
 #include <wchar.h>
 int vfwscanf_s(FILE * restrict stream, const wchar_t * restrict format,
 va_list arg);
```

##### Runtime-constraints

- 2 Neither **stream** nor **format** shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **vfwscanf\_s** function does not attempt to perform further input, and it is unspecified to what extent **vfwscanf\_s** performed input before discovering the runtime-constraint violation.

##### Description

- 4 The **vfwscanf\_s** function is equivalent to **fwscanf\_s**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va\_start** macro (and possibly subsequent **va\_arg** invocations). The **vfwscanf\_s** function does not invoke the **va\_end** macro.<sup>514)</sup>

##### Returns

- 5 The **vfwscanf\_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **vfwscanf\_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

#### K.3.9.2.8 The **vsnwprintf\_s** function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <wchar.h>
 int vsnwprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format
 , va_list arg);
```

##### Runtime-constraints

- 2 Neither **s** nor **format** shall be a null pointer. **n** shall neither equal zero nor be greater than **RSIZE\_MAX** / **sizeof(wchar\_t)**. The **%n** specifier<sup>515)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by **format**. Any argument to **vsnwprintf\_s** corresponding to a **%s** specifier shall not be a null pointer. No encoding error shall occur.

<sup>514)</sup>As the functions **vfwscanf\_s**, **vwscanf\_s**, and **vswscanf\_s** invoke the **va\_arg** macro, the representation of **arg** after the return is indeterminate.

<sup>515)</sup>It is not a runtime-constraint violation for the wide characters **%n** to appear in sequence in the wide string pointed at by **format** when those wide characters are not interpreted as a **%n** specifier. For example, if the entire format string was **L"%n"**.

- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `n` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then the `vsnwprintf_s` function sets `s[0]` to the null wide character.

#### Description

- 4 The `vsnwprintf_s` function is equivalent to the `vswprintf` function except for the previously listed explicit runtime-constraints.
- 5 The `vsnwprintf_s` function, unlike `vswprintf_s`, will truncate the result to fit within the array pointed to by `s`.

#### Returns

- 6 The `vsnwprintf_s` function returns the number of wide characters that would have been written had `n` been sufficiently large, not counting the terminating null character, or a negative value if a runtime-constraint violation occurred. Thus, the null-terminated output has been completely written if and only if the returned value is both nonnegative and less than `n`.

#### K.3.9.2.9 The `vswprintf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <wchar.h>
 int vswprintf_s(wchar_t * restrict s, rsize_t n, const wchar_t * restrict format,
 va_list arg);
```

#### Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. `n` shall neither equal zero nor be greater than `RSIZE_MAX/sizeof(wchar_t)`. The number of wide characters (including the trailing null) required for the result to be written to the array pointed to by `s` shall not be greater than `n`. The `%n` specifier<sup>516</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by `format`. Any argument to `vswprintf_s` corresponding to a `%s` specifier shall not be a null pointer. No encoding error shall occur.
- 3 If there is a runtime-constraint violation, then if `s` is not a null pointer and `n` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then the `vswprintf_s` function sets `s[0]` to the null wide character.

#### Description

- 4 The `vswprintf_s` function is equivalent to the `vswprintf` function except for the previously listed explicit runtime-constraints.
- 5 The `vswprintf_s` function, unlike `vsnwprintf_s`, treats a result too big for the array pointed to by `s` as a runtime-constraint violation.

#### Returns

- 6 If no runtime-constraint violation occurred, the `vswprintf_s` function returns the number of wide characters written in the array, not counting the terminating null wide character. If an encoding error occurred or if `n` or more wide characters are requested to be written, `vswprintf_s` returns a negative value. If any other runtime-constraint violation occurred, `vswprintf_s` returns zero.

#### K.3.9.2.10 The `vswscanf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
```

<sup>516</sup>It is not a runtime-constraint violation for the wide characters `%n` to appear in sequence in the wide string pointed at by `format` when those wide characters are not interpreted as a `%n` specifier. For example, if the entire format string was `L"%n%n"`.

```
#include <wchar.h>
int vswscanf_s(const wchar_t * restrict s, const wchar_t * restrict format,
 va_list arg);
```

#### Runtime-constraints

- 2 Neither `s` nor `format` shall be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `vswscanf_s` function does not attempt to perform further input, and it is unspecified to what extent `vswscanf_s` performed input before discovering the runtime-constraint violation.

#### Description

- 4 The `vswscanf_s` function is equivalent to `swscanf_s`, with the varying argument list replaced by `arg`, which shall have been initialized by the `va_start` macro (and possibly subsequent `va_arg` invocations). The `vswscanf_s` function does not invoke the `va_end` macro.<sup>517)</sup>

#### Returns

- 5 The `vswscanf_s` function returns the value of the macro `EOF` if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the `vswscanf_s` function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

#### K.3.9.2.11 The `vwprintf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <wchar.h>
 int vwprintf_s(const wchar_t * restrict format, va_list arg);
```

#### Runtime-constraints

- 2 `format` shall not be a null pointer. The `%n` specifier<sup>518)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by `format`. Any argument to `vwprintf_s` corresponding to a `%s` specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the `vwprintf_s` function does not attempt to produce further output, and it is unspecified to what extent `vwprintf_s` produced output before discovering the runtime-constraint violation.

#### Description

- 4 The `vwprintf_s` function is equivalent to the `vwprintf` function except for the previously listed explicit runtime-constraints.

#### Returns

- 5 The `vwprintf_s` function returns the number of wide characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

#### K.3.9.2.12 The `vwscanf_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <stdarg.h>
 #include <wchar.h>
```

<sup>517)</sup>As the functions `vfwscanf_s`, `vwscanf_s`, and `vswscanf_s` invoke the `va_arg` macro, the representation of `arg` after the return is indeterminate.

<sup>518)</sup>It is not a runtime-constraint violation for the wide characters `%n` to appear in sequence in the wide string pointed at by `format` when those wide characters are not interpreted as a `%n` specifier. For example, if the entire format string was `L"%n%n"`.

```
int vwscanf_s(const wchar_t * restrict format, va_list arg);
```

#### Runtime-constraints

- 2 **format** shall not be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **vwscanf\_s** function does not attempt to perform further input, and it is unspecified to what extent **vwscanf\_s** performed input before discovering the runtime-constraint violation.

#### Description

- 4 The **vwscanf\_s** function is equivalent to **wscanf\_s**, with the varying argument list replaced by **arg**, which shall have been initialized by the **va\_start** macro (and possibly subsequent **va\_arg** invocations). The **vwscanf\_s** function does not invoke the **va\_end** macro.<sup>519)</sup>

#### Returns

- 5 The **vwscanf\_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **vwscanf\_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

### K.3.9.2.13 The **wprintf\_s** function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int wprintf_s(const wchar_t * restrict format, ...);
```

#### Runtime-constraints

- 2 **format** shall not be a null pointer. The **%n** specifier<sup>520)</sup> (modified or not by flags, field width, or precision) shall not appear in the wide string pointed to by **format**. Any argument to **wprintf\_s** corresponding to a **%s** specifier shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **wprintf\_s** function does not attempt to produce further output, and it is unspecified to what extent **wprintf\_s** produced output before discovering the runtime-constraint violation.

#### Description

- 4 The **wprintf\_s** function is equivalent to the **wprintf** function except for the previously listed explicit runtime-constraints.

#### Returns

- 5 The **wprintf\_s** function returns the number of wide characters transmitted, or a negative value if an output error, encoding error, or runtime-constraint violation occurred.

### K.3.9.2.14 The **wscanf\_s** function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 int wscanf_s(const wchar_t * restrict format, ...);
```

<sup>519)</sup>As the functions **vfwscanf\_s**, **vwscanf\_s**, and **vswscanf\_s** invoke the **va\_arg** macro, the representation of **arg** after the return is indeterminate.

<sup>520)</sup>It is not a runtime-constraint violation for the wide characters **%n** to appear in sequence in the wide string pointed at by **format** when those wide characters are not interpreted as a **%n** specifier. For example, if the entire format string was **L"%n%n"**.

**Runtime-constraints**

- 2 **format** shall not be a null pointer. Any argument indirected though to store converted input shall not be a null pointer.
- 3 If there is a runtime-constraint violation, the **wscanf\_s** function does not attempt to perform further input, and it is unspecified to what extent **wscanf\_s** performed input before discovering the runtime-constraint violation.

**Description**

- 4 The **wscanf\_s** function is equivalent to **fwscanf\_s** with the argument **stdin** interposed before the arguments to **wscanf\_s**.

**Returns**

- 5 The **wscanf\_s** function returns the value of the macro **EOF** if an input failure occurs before any conversion or if there is a runtime-constraint violation. Otherwise, the **wscanf\_s** function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure.

**K.3.9.3 General wide string utilities****K.3.9.3.1 Wide string copying functions****K.3.9.3.1.1 The **wscpy\_s** function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 errno_t wscpy_s(wchar_t * restrict s1, rsize_t slmax,
 const wchar_t * restrict s2);
```

**Runtime-constraints**

- 2 Neither **s1** nor **s2** shall be a null pointer. **slmax** shall not be greater than **RSIZE\_MAX/sizeof(wchar\_t)**. **slmax** shall not equal zero. **slmax** shall be greater than **wcsnlen\_s(s2, slmax)**. Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, then if **s1** is not a null pointer and **slmax** is greater than zero and not greater than **RSIZE\_MAX/sizeof(wchar\_t)**, then **wscpy\_s** sets **s1[0]** to the null wide character.

**Description**

- 4 The **wscpy\_s** function copies the wide string pointed to by **s2** (including the terminating null wide character) into the array pointed to by **s1**.
- 5 All elements following the terminating null wide character (if any) written by **wscpy\_s** in the array of **slmax** wide characters pointed to by **s1** take unspecified values when **wscpy\_s** returns.<sup>521)</sup>

**Returns**

- 6 The **wscpy\_s** function returns zero<sup>522)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

**K.3.9.3.1.2 The **wcsncpy\_s** function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 errno_t wcsncpy_s(wchar_t * restrict s1, rsize_t slmax,
 const wchar_t * restrict s2, rsize_t n);
```

<sup>521)</sup>This allows an implementation to copy wide characters from **s2** to **s1** while simultaneously checking if any of those wide characters are null. Such an approach can write a wide character to every element of **s1** before discovering that the first element was set to the null wide character.

<sup>522)</sup>A zero return value implies that all the requested wide characters from the string pointed to by **s2** fit within the array pointed to by **s1** and that the result in **s1** is null terminated.

**Runtime-constraints**

- 2 Neither `s1` nor `s2` shall be a null pointer. Neither `s1max` nor `n` shall be greater than `RSIZE_MAX/sizeof(wchar_t)`. `s1max` shall not equal zero. If `n` is not less than `s1max`, then `s1max` shall be greater than `wcsnlen_s(s2, s1max)`. Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, then if `s1` is not a null pointer and `s1max` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then `wcsncpy_s` sets `s1[0]` to the null wide character.

**Description**

- 4 The `wcsncpy_s` function copies not more than `n` successive wide characters (wide characters that follow a null wide character are not copied) from the array pointed to by `s2` to the array pointed to by `s1`. If no null wide character was copied from `s2`, then `s1[n]` is set to a null wide character.
- 5 All elements following the terminating null wide character (if any) written by `wcsncpy_s` in the array of `s1max` wide characters pointed to by `s1` take unspecified values when `wcsncpy_s` returns.<sup>523)</sup>

**Returns**

- 6 The `wcsncpy_s` function returns zero<sup>524)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.
- 7 EXAMPLE The `wcsncpy_s` function can be used to copy a wide string without the danger that the result will not be null terminated or that wide characters will be written past the end of the destination array.

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <wchar.h>
/* ... */
wchar_t src1[100] = L"hello";
wchar_t src2[7] = {L'g', L'o', L'o', L'd', L'b', L'y', L'e'};
wchar_t dst1[6], dst2[5], dst3[5];
int r1, r2, r3;
r1 = wcsncpy_s(dst1, 6, src1, 100);
r2 = wcsncpy_s(dst2, 5, src2, 7);
r3 = wcsncpy_s(dst3, 5, src2, 4);
```

The first call will assign to `r1` the value zero and to `dst1` the sequence of wide characters `hello\0`.

The second call will assign to `r2` a nonzero value and to `dst2` the sequence of wide characters `\0`.

The third call will assign to `r3` the value zero and to `dst3` the sequence of wide characters `good\0`.

**K.3.9.3.1.3 The `wmemcpy_s` function****Synopsis**

- 1 

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <wchar.h>
errno_t wmemcpy_s(wchar_t * restrict s1, rsize_t s1max,
 const wchar_t * restrict s2, rsize_t n);
```

**Runtime-constraints**

- 2 Neither `s1` nor `s2` shall be a null pointer. Neither `s1max` nor `n` shall be greater than `RSIZE_MAX/sizeof(wchar_t)`. `n` shall not be greater than `s1max`. Copying shall not take place between objects that overlap.
- 3 If there is a runtime-constraint violation, the `wmemcpy_s` function stores zeros in the first `s1max` wide characters of the object pointed to by `s1` if `s1` is not a null pointer and `s1max` is not greater than `RSIZE_MAX/sizeof(wchar_t)`.

<sup>523)</sup> This allows an implementation to copy wide characters from `s2` to `s1` while simultaneously checking if any of those wide characters are null. Such an approach can write a wide character to every element of `s1` before discovering that the first element was set to the null wide character.

<sup>524)</sup> A zero return value implies that all the requested wide characters from the string pointed to by `s2` fit within the array pointed to by `s1` and that the result in `s1` is null terminated.



**Description**

- 4 The **wmemcpy\_s** function copies *n* successive wide characters from the object pointed to by *s2* into the object pointed to by *s1*.

**Returns**

- 5 The **wmemcpy\_s** function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

**K.3.9.3.1.4 The wmemmove\_s function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 errno_t wmemmove_s(wchar_t *s1, rsize_t slmax, const wchar_t *s2, rsize_t n);
```

**Runtime-constraints**

- 2 Neither *s1* nor *s2* shall be a null pointer. Neither *slmax* nor *n* shall be greater than **RSIZE\_MAX/sizeof(wchar\_t)**. *n* shall not be greater than *slmax*.
- 3 If there is a runtime-constraint violation, the **wmemmove\_s** function stores zeros in the first *slmax* wide characters of the object pointed to by *s1* if *s1* is not a null pointer and *slmax* is not greater than **RSIZE\_MAX/sizeof(wchar\_t)**.

**Description**

- 4 The **wmemmove\_s** function copies *n* successive wide characters from the object pointed to by *s2* into the object pointed to by *s1*. This copying takes place as if the *n* wide characters from the object pointed to by *s2* are first copied into a temporary array of *n* wide characters that does not overlap the objects pointed to by *s1* or *s2*, and then the *n* wide characters from the temporary array are copied into the object pointed to by *s1*.

**Returns**

- 5 The **wmemmove\_s** function returns zero if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

**K.3.9.3.2 Wide string concatenation functions****K.3.9.3.2.1 The wcscat\_s function****Synopsis**

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 errno_t wcscat_s(wchar_t * restrict s1, rsize_t slmax,
 const wchar_t * restrict s2);
```

**Runtime-constraints**

- 2 Let *m* denote the value *slmax* - **wcsnlen\_s**(*s1*, *slmax*) upon entry to **wcscat\_s**.
- 3 Neither *s1* nor *s2* shall be a null pointer. *slmax* shall not be greater than **RSIZE\_MAX/sizeof(wchar\_t)**. *slmax* shall not equal zero. *m* shall not equal zero.<sup>525)</sup> *m* shall be greater than **wcsnlen\_s**(*s2*, *m*). Copying shall not take place between objects that overlap.
- 4 If there is a runtime-constraint violation, then if *s1* is not a null pointer and *slmax* is greater than zero and not greater than **RSIZE\_MAX/sizeof(wchar\_t)**, then **wcscat\_s** sets *s1*[0] to the null wide character.

**Description**

- 5 The **wcscat\_s** function appends a copy of the wide string pointed to by *s2* (including the terminating null wide character) to the end of the wide string pointed to by *s1*. The initial wide character from

<sup>525)</sup>Zero means that *s1* was not null terminated upon entry to **wcscat\_s**.



`s2` overwrites the null wide character at the end of `s1`.

- 6 All elements following the terminating null wide character (if any) written by `wscat_s` in the array of `s1max` wide characters pointed to by `s1` take unspecified values when `wscat_s` returns.<sup>526)</sup>

#### Returns

- 7 The `wscat_s` function returns zero<sup>527)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.

#### K.3.9.3.2.2 The `wcsncat_s` function

##### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 errno_t wcsncat_s(wchar_t * restrict s1, rsize_t s1max,
 const wchar_t * restrict s2, rsize_t n);
```

##### Runtime-constraints

- 2 Let  $m$  denote the value `s1max - wcsnlen_s(s1, s1max)` upon entry to `wcsncat_s`.
- 3 Neither `s1` nor `s2` shall be a null pointer. Neither `s1max` nor `n` shall be greater than `RSIZE_MAX/sizeof(wchar_t)`. `s1max` shall not equal zero.  $m$  shall not equal zero.<sup>528)</sup> If `n` is not less than  $m$ , then  $m$  shall be greater than `wcsnlen_s(s2, m)`. Copying shall not take place between objects that overlap.
- 4 If there is a runtime-constraint violation, then if `s1` is not a null pointer and `s1max` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then `wcsncat_s` sets `s1[0]` to the null wide character.

##### Description

- 5 The `wcsncat_s` function appends not more than `n` successive wide characters (wide characters that follow a null wide character are not copied) from the array pointed to by `s2` to the end of the wide string pointed to by `s1`. The initial wide character from `s2` overwrites the null wide character at the end of `s1`. If no null wide character was copied from `s2`, then `s1[s1max - m + n]` is set to a null wide character.
- 6 All elements following the terminating null wide character (if any) written by `wcsncat_s` in the array of `s1max` wide characters pointed to by `s1` take unspecified values when `wcsncat_s` returns.<sup>529)</sup>

#### Returns

- 7 The `wcsncat_s` function returns zero<sup>530)</sup> if there was no runtime-constraint violation. Otherwise, a nonzero value is returned.
- 8 EXAMPLE The `wcsncat_s` function can be used to copy a wide string without the danger that the result will not be null terminated or that wide characters will be written past the end of the destination array.

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <wchar.h>
/* ... */
wchar_t s1[100] = L"good";
wchar_t s2[6] = L"hello";
```

<sup>526)</sup>This allows an implementation to append wide characters from `s2` to `s1` while simultaneously checking if any of those wide characters are null. Such an approach can write a wide character to every element of `s1` before discovering that the first element was set to the null wide character.

<sup>527)</sup>A zero return value implies that all the requested wide characters from the wide string pointed to by `s2` were appended to the wide string pointed to by `s1` and that the result in `s1` is null terminated.

<sup>528)</sup>Zero means that `s1` was not null terminated upon entry to `wcsncat_s`.

<sup>529)</sup>This allows an implementation to append wide characters from `s2` to `s1` while simultaneously checking if any of those wide characters are null. Such an approach can write a wide character to every element of `s1` before discovering that the first element was set to the null wide character.

<sup>530)</sup>A zero return value implies that all the requested wide characters from the wide string pointed to by `s2` were appended to the wide string pointed to by `s1` and that the result in `s1` is null terminated.

```

wchar_t s3[6] = L"hello";
wchar_t s4[7] = L"abc";
wchar_t s5[1000] = L"bye";
int r1, r2, r3, r4;
r1 = wcsncat_s(s1, 100, s5, 1000);
r2 = wcsncat_s(s2, 6, L"", 1);
r3 = wcsncat_s(s3, 6, L"X", 2);
r4 = wcsncat_s(s4, 7, L"defghijklmn", 3);

```

After the first call `r1` will have the value zero and `s1` will be the wide character sequence `goodbye\0`.

After the second call `r2` will have the value zero and `s2` will be the wide character sequence `hello\0`.

After the third call `r3` will have a nonzero value and `s3` will be the wide character sequence `\0`.

After the fourth call `r4` will have the value zero and `s4` will be the wide character sequence `abcdef\0`.

### K.3.9.3.3 Wide string search functions

#### K.3.9.3.3.1 The `wcstok_s` function

##### Synopsis

```

1 #define __STDC_WANT_LIB_EXT1__ 1
 #include <wchar.h>
 wchar_t *wcstok_s(wchar_t * restrict s1, rsize_t * restrict s1max,
 const wchar_t * restrict s2, wchar_t ** restrict ptr);

```

##### Runtime-constraints

- 2 None of `s1max`, `s2`, or `ptr` shall be a null pointer. If `s1` is a null pointer, then `*ptr` shall not be a null pointer. The value of `*s1max` shall not be greater than `RSIZE_MAX/sizeof(wchar_t)`. The end of the token found shall occur within the first `*s1max` wide characters of `s1` for the first call, and shall occur within the first `*s1max` wide characters of where searching resumes on subsequent calls.
- 3 If there is a runtime-constraint violation, the `wcstok_s` function does not indirect through the `s1` or `s2` pointers, and does not store a value in the object pointed to by `ptr`.

##### Description

- 4 A sequence of calls to the `wcstok_s` function breaks the wide string pointed to by `s1` into a sequence of tokens, each of which is delimited by a wide character from the wide string pointed to by `s2`. The fourth argument points to a caller-provided `wchar_t` pointer into which the `wcstok_s` function stores information necessary for it to continue scanning the same wide string.
- 5 The first call in a sequence has a non-null first argument and `s1max` points to an object whose value is the number of elements in the wide character array pointed to by the first argument. The first call stores an initial value in the object pointed to by `ptr` and updates the value pointed to by `s1max` to reflect the number of elements that remain in relation to `ptr`. Subsequent calls in the sequence have a null first argument and the objects pointed to by `s1max` and `ptr` are required to have the values stored by the previous call in the sequence, which are then updated. The separator wide string pointed to by `s2` can be different from call to call.
- 6 The first call in the sequence searches the wide string pointed to by `s1` for the first wide character that is *not* contained in the current separator wide string pointed to by `s2`. If no such wide character is found, then there are no tokens in the wide string pointed to by `s1` and the `wcstok_s` function returns a null pointer. If such a wide character is found, it is the start of the first token.
- 7 The `wcstok_s` function then searches from there for the first wide character in `s1` that is contained in the current separator wide string. If no such wide character is found, the current token extends to the end of the wide string pointed to by `s1`, and subsequent searches in the same wide string for a token return a null pointer. If such a wide character is found, it is overwritten by a null wide character, which terminates the current token.
- 8 In all cases, the `wcstok_s` function stores sufficient information in the pointer pointed to by `ptr` so that subsequent calls, with a null pointer for `s1` and the unmodified pointer value for `ptr`, shall start

searching just past the element overwritten by a null wide character (if any).

### Returns

- 9 The **wcstok\_s** function returns a pointer to the first wide character of a token, or a null pointer if there is no token or there is a runtime-constraint violation.

### EXAMPLE

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <wchar.h>
static wchar_t str1[] = L"?a???b,,,#c";
static wchar_t str2[] = L"\t \t";
wchar_t *t, *ptr1, *ptr2;
rsize_t max1 = wcslen(str1)+1;
rsize_t max2 = wcslen(str2)+1;

t = wcstok_s(str1, &max1, "?", &ptr1); // t points to the token "a"
t = wcstok_s(nullptr, &max1, ",", &ptr1); // t points to the token "??b"
t = wcstok_s(str2, &max2, " \t", &ptr2); // t is a null pointer
t = wcstok_s(nullptr, &max1, "#", &ptr1); // t points to the token "c"
t = wcstok_s(nullptr, &max1, "?", &ptr1); // t is a null pointer
```

## K.3.9.3.4 Miscellaneous functions

### K.3.9.3.4.1 The **wcsnlen\_s** function

#### Synopsis

```
1 #define __STDC_WANT_LIB_EXT1__ 1
#include <wchar.h>
size_t wcsnlen_s(const wchar_t *s, size_t maxsize);
```

#### Description

- 2 The **wcsnlen\_s** function counts not more than **maxsize** wide characters (a null wide character and wide characters that follow it are not counted) in the array to which **s** points. At most the first **maxsize** wide characters of **s** shall be accessed by **wcsnlen\_s**. The implementation shall behave as if it reads the wide characters sequentially and stops as soon as a null wide character is found.

#### Returns

- 3 If **s** is a null pointer,<sup>531)</sup> then the **wcsnlen\_s** function returns zero.
- 4 Otherwise, the **wcsnlen\_s** function returns the number of wide characters that precede the terminating null wide character. If there is no null wide character in the first **maxsize** wide characters of **s** then **wcsnlen\_s** returns **maxsize**.

## K.3.9.4 Extended multibyte/wide character conversion utilities

### K.3.9.4.1 Restartable multibyte/wide character conversion functions

#### K.3.9.4.1.1 General

- 1 Unlike **wcrtomb**, **wcrtomb\_s** does not permit the **ps** parameter (the pointer to the conversion state) to be a null pointer.

#### K.3.9.4.1.2 The **wcrtomb\_s** function

##### Synopsis

```
1 #include <wchar.h>
errno_t wcrtomb_s(size_t * restrict retval, char * restrict s, rsize_t smax,
 wchar_t wc, mbstate_t * restrict ps);
```

<sup>531)</sup>The **wcsnlen\_s** function has no runtime-constraints. This lack of runtime-constraints along with the values returned for a null pointer or an unterminated wide string argument make **wcsnlen\_s** useful in algorithms that gracefully handle such exceptional data.

**Runtime-constraints**

- 2 Neither `retval` nor `ps` shall be a null pointer. If `s` is not a null pointer, then `smax` shall not equal zero and shall not be greater than `RSIZE_MAX`. If `s` is not a null pointer, then `smax` shall not be less than the number of bytes to be stored in the array pointed to by `s`. If `s` is a null pointer, then `smax` shall equal zero.
- 3 If there is a runtime-constraint violation, then `wcrtomb_s` does the following. If `s` is not a null pointer and `smax` is greater than zero and not greater than `RSIZE_MAX`, then `wcrtomb_s` sets `s[0]` to the null character. If `retval` is not a null pointer, then `wcrtomb_s` sets `*retval` to `(size_t)(-1)`.

**Description**

- 4 If `s` is a null pointer, the `wcrtomb_s` function is equivalent to the call

```
wcrtomb_s(&retval, buf, _Lengthof buf, L'\0', ps)
```

where `retval` and `buf` are internal objects of the appropriate types, and the size of `buf` is greater than `MB_CUR_MAX`.

- 5 If `s` is not a null pointer, the `wcrtomb_s` function determines the number of bytes needed to represent the multibyte character that corresponds to the wide character given by `wc` (including any shift sequences), and stores the multibyte character representation in the array whose first element is pointed to by `s`. At most `MB_CUR_MAX` bytes are stored. If `wc` is a null wide character, a null byte is stored, preceded by any shift sequence needed to restore the initial shift state; the resulting state described is the initial conversion state.
- 6 If `wc` does not correspond to a valid multibyte character, an encoding error occurs: the `wcrtomb_s` function stores the value `(size_t)(-1)` into `*retval` and the conversion state is unspecified. Otherwise, the `wcrtomb_s` function stores into `*retval` the number of bytes (including any shift sequences) stored in the array pointed to by `s`.

**Returns**

- 7 The `wcrtomb_s` function returns zero if no runtime-constraint violation and no encoding error occurred. Otherwise, a nonzero value is returned.

**K.3.9.4.2 Restartable multibyte/wide string conversion functions****K.3.9.4.2.1 General**

- 1 Unlike `mbsrtowcs` and `wcsrtombs`, `mbsrtowcs_s` and `wcsrtombs_s` do not permit the `ps` parameter (the pointer to the conversion state) to be a null pointer.

**K.3.9.4.2.2 The `mbsrtowcs_s` function****Synopsis**

- 1 

```
#include <wchar.h>
errno_t mbsrtowcs_s(size_t * restrict retval, wchar_t * restrict dst,
 rsize_t dstmax, const char ** restrict src, rsize_t len,
 mbstate_t * restrict ps);
```

**Runtime-constraints**

- 2 None of `retval`, `src`, `*src`, or `ps` shall be null pointers. If `dst` is not a null pointer, then neither `len` nor `dstmax` shall be greater than `RSIZE_MAX/sizeof(wchar_t)`. If `dst` is a null pointer, then `dstmax` shall equal zero. If `dst` is not a null pointer, then `dstmax` shall not equal zero. If `dst` is not a null pointer and `len` is not less than `dstmax`, then a null character shall occur within the first `dstmax` multibyte characters of the array pointed to by `*src`.
- 3 If there is a runtime-constraint violation, then `mbsrtowcs_s` does the following. If `retval` is not a null pointer, then `mbsrtowcs_s` sets `*retval` to `(size_t)(-1)`. If `dst` is not a null pointer and `dstmax` is greater than zero and not greater than `RSIZE_MAX/sizeof(wchar_t)`, then `mbsrtowcs_s` sets `dst[0]` to the null wide character.

### Description

- 4 The **mbstowcs\_s** function converts a sequence of multibyte characters that begins in the conversion state described by the object pointed to by **ps**, from the array indirectly pointed to by **src** into a sequence of corresponding wide characters. If **dst** is not a null pointer, the converted characters are stored into the array pointed to by **dst**. Conversion continues up to and including a terminating null character, which is also stored. Conversion stops earlier in two cases: when a sequence of bytes is encountered that does not form a valid multibyte character, or (if **dst** is not a null pointer) when **len** wide characters have been stored into the array pointed to by **dst**.<sup>532)</sup> If **dst** is not a null pointer and no null wide character was stored into the array pointed to by **dst**, then **dst[*len*]** is set to the null wide character. Each conversion takes place as if by a call to the **mbrtowc** function.
- 5 If **dst** is not a null pointer, the pointer object pointed to by **src** is assigned either a null pointer (if conversion stopped due to reaching a terminating null character) or the address just past the last multibyte character converted (if any). If conversion stopped due to reaching a terminating null character and if **dst** is not a null pointer, the resulting state described is the initial conversion state.
- 6 Regardless of whether **dst** is or is not a null pointer, if the input conversion encounters a sequence of bytes that do not form a valid multibyte character, an encoding error occurs: the **mbstowcs\_s** function stores the value (**size\_t**) (-1) into **\*retval** and the conversion state is unspecified. Otherwise, the **mbstowcs\_s** function stores into **\*retval** the number of multibyte characters successfully converted, not including the terminating null character (if any).
- 7 All elements following the terminating null wide character (if any) written by **mbstowcs\_s** in the array of **dstmax** wide characters pointed to by **dst** take unspecified values when **mbstowcs\_s** returns.<sup>533)</sup>
- 8 If copying takes place between objects that overlap, the objects take on unspecified values.

### Returns

- 9 The **mbstowcs\_s** function returns zero if no runtime-constraint violation and no encoding error occurred. Otherwise, a nonzero value is returned.

#### K.3.9.4.2.3 The **wcsrtombs\_s** function

##### Synopsis

```
1 #include <wchar.h>
 errno_t wcsrtombs_s(size_t * restrict retval, char * restrict dst,
 rsize_t dstmax, const wchar_t ** restrict src, rsize_t len,
 mbstate_t * restrict ps);
```

### Runtime-constraints

- 2 None of **retval**, **src**, **\*src**, or **ps** shall be null pointers. If **dst** is not a null pointer, then neither **len** shall be greater than **RSIZE\_MAX/sizeof(wchar\_t)** nor **dstmax** shall be greater than **RSIZE\_MAX**. If **dst** is a null pointer, then **dstmax** shall equal zero. If **dst** is not a null pointer, then **dstmax** shall not equal zero. If **dst** is not a null pointer and **len** is not less than **dstmax**, then the conversion shall have been stopped (see the following) because a terminating null wide character was reached or because an encoding error occurred.
- 3 If there is a runtime-constraint violation, then **wcsrtombs\_s** does the following. If **retval** is not a null pointer, then **wcsrtombs\_s** sets **\*retval** to (**size\_t**) (-1). If **dst** is not a null pointer and **dstmax** is greater than zero and not greater than **RSIZE\_MAX**, then **wcsrtombs\_s** sets **dst[0]** to the null character.

### Description

- 4 The **wcsrtombs\_s** function converts a sequence of wide characters from the array indirectly pointed to by **src** into a sequence of corresponding multibyte characters that begins in the conversion state

<sup>532)</sup>Thus, the value of **len** is ignored if **dst** is a null pointer.

<sup>533)</sup>This allows an implementation to attempt converting the multibyte string before discovering a terminating null character did not occur where required.

described by the object pointed to by `ps`. If `dst` is not a null pointer, the converted characters are then stored into the array pointed to by `dst`. Conversion continues up to and including a terminating null wide character, which is also stored. Conversion stops earlier in two cases:

- when a wide character is reached that does not correspond to a valid multibyte character;
- (if `dst` is not a null pointer) when the next multibyte character would exceed the limit of  $n$  total bytes to be stored into the array pointed to by `dst`. If the wide character being converted is the null wide character, then  $n$  is the lesser of `len` or `dstmax`. Otherwise,  $n$  is the lesser of `len` or `dstmax` - 1.

If the conversion stops without converting a null wide character and `dst` is not a null pointer, then a null character is stored into the array pointed to by `dst` immediately following any multibyte characters already stored. Each conversion takes place as if by a call to the `wcrtomb` function.<sup>534)</sup>

- 5 If `dst` is not a null pointer, the pointer object pointed to by `src` is assigned either a null pointer (if conversion stopped due to reaching a terminating null wide character) or the address just past the last wide character converted (if any). If conversion stopped due to reaching a terminating null wide character, the resulting state described is the initial conversion state.
- 6 Regardless of whether `dst` is or is not a null pointer, if the input conversion encounters a wide character that does not correspond to a valid multibyte character, an encoding error occurs: the `wcsrtombs_s` function stores the value (`size_t`) (-1) into `*retval` and the conversion state is unspecified. Otherwise, the `wcsrtombs_s` function stores into `*retval` the number of bytes in the resulting multibyte character sequence, not including the terminating null character (if any).
- 7 All elements following the terminating null character (if any) written by `wcsrtombs_s` in the array of `dstmax` elements pointed to by `dst` take unspecified values when `wcsrtombs_s` returns.<sup>535)</sup>
- 8 If copying takes place between objects that overlap, the objects take on unspecified values.

#### Returns

- 9 The `wcsrtombs_s` function returns zero if no runtime-constraint violation and no encoding error occurred. Otherwise, a nonzero value is returned.

<sup>534)</sup>If conversion stops because a terminating null wide character has been reached, the bytes stored include those necessary to reach the initial shift state immediately before the null byte. However, if the conversion stops before a terminating null wide character has been reached, the result will be null terminated, but does not necessarily end in the initial shift state.

<sup>535)</sup>When `len` is not less than `dstmax`, the implementation can fill the array before discovering a runtime-constraint violation.



## Annex L

(normative)

## Analyzability

### L.1 Scope

- 1 This Annex specifies optional behavior that can aid in the analyzability of C programs.
- 2 An implementation that defines **\_\_STDC\_ANALYZABLE\_\_** shall conform to the specifications in this annex (see also 6.10.10.4).<sup>536)</sup>

### L.2 Requirements

- 1 If the program performs a trap (3.27), the implementation is permitted to invoke a runtime-constraint handler. Any such semantics are implementation-defined.
- 2 All undefined behavior shall be limited to bounded undefined behavior, except for the following which are permitted to result in critical undefined behavior:
  - An object is referred to outside of its lifetime (6.2.4).
  - A store is performed to an object that has two incompatible declarations (6.2.7),
  - A pointer is used to call a function whose type is not compatible with the referenced type (6.2.7, 6.3.3.3, 6.5.3.3).
  - An lvalue does not designate an object when evaluated (6.3.3.1).
  - The program attempts to modify a string literal (6.4.6).
  - The operand of the unary \* operator has an invalid value (6.5.4.3).
  - Addition or subtraction of a pointer into, or just beyond, an array object and an integer type produces a result that points just beyond the array object and is used as the operand of a unary \* operator that is evaluated (6.5.7).
  - An attempt is made to modify an object defined with a const-qualified type through use of an lvalue with non-const-qualified type (6.7.4).
  - An argument to a function or macro defined in the standard library has an invalid value or a type not expected by a variadic function (7.1.4).
  - The **longjmp** function is called with a **jmp\_buf** argument where the most recent invocation of the **setjmp** macro in the same invocation of the program with the corresponding **jmp\_buf** argument is nonexistent, or the invocation was from another thread of execution, or the function containing the invocation has terminated execution in the interim, or the invocation was within the scope of an identifier with variably modified type and execution has left that scope in the interim (7.13.3.1).
  - The value of a pointer that refers to space deallocated by a call to the **free**, **free\_sized**, **free\_aligned\_sized** or **realloc** function is used (7.24.4).
  - A string or wide string utility function accesses an array beyond the end of an object (7.27.1, 7.32.4).

<sup>536)</sup>Implementations that do not define **\_\_STDC\_ANALYZABLE\_\_** are not required to conform to these specifications.

## Annex M

(informative)

### Change History

#### M.1 Attribute Changes

- 1 The attribute feature was introduced in the fifth edition of this document, as detailed in the change list in M.3. Table M.1 illustrates what usage of the `__has_c_attribute` preprocessor conditional expression should return in the latest edition, previous values, and the change associated with that value.
- 2 Programs and implementations can use Table M.1 to know what values were being used at any specific point in time, leading up to the publication of this document. The value at the bottom of a particular row in Table M.1 is the latest value and corresponds with the behavior for the given attribute described in this document.

Table M.1: `__has_c_attribute` values and associated changes

| attribute tokens    | value   | semantic and/or syntactic changes                                                                                    |
|---------------------|---------|----------------------------------------------------------------------------------------------------------------------|
| <b>deprecated</b>   | 201904L | Initial introduction.                                                                                                |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>fallthrough</b>  | 201904L | Initial introduction.                                                                                                |
|                     | 201910L | Expanded locations where <b>fallthrough</b> provides diagnostics due to improvements in specification of blocks.     |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>maybe_unused</b> | 201904L | Initial introduction.                                                                                                |
|                     | 202106L | <b>maybe_unused</b> may appertain to labels.                                                                         |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>nodiscard</b>    | 201904L | Initial introduction.                                                                                                |
|                     | 202003L | Added a form which accepts a string literal for diagnostic purposes, e.g. <b>nodiscard</b> ("should have a reason"). |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>noreturn</b>     | 202202L | Initial introduction.                                                                                                |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>reproducible</b> | 202207L | Initial introduction.                                                                                                |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |
| <b>unsequenced</b>  | 202207L | Initial introduction.                                                                                                |
|                     | 202311L | Harmonized for fifth edition.                                                                                        |

#### M.2 Sixth Edition

- 1 Major changes in this sixth edition (`__STDC_VERSION__` 202ymmL) include:
  - Uppercase characters “A” through “F” (‘A’ through ‘F’) and lowercase characters “a” through “f” (‘a’ through ‘f’), colloquially known as the “hexadecimal digits”, are now guaranteed to be sequential.
  - Added a new mode character ‘p’ to **foopen**.
  - Clarified accessing and utilizing typed and untyped byte arrays.
  - **\_Generic** selection allows for a type to be used directly.
  - The removal of explicit and independent imaginary types.
  - ++ and -- being well-defined for complex values.
  - Clarified and improved recommendations for rounding in **printf**.



- **alignof** can be used on an incomplete array type.
- Added literals for complex types.
- New **strlen** and **wcsnlen** functions.
- New **u(1)(1)abs** and **uimaxabs** functions, to compute the absolute value without undefined behavior.
- New **0o** and **00** prefixes for octal literals.
- Delimited escape sequences.
- Declarations in the header of **if** and **switch** statements.
- Allowed atomic types not to have fundamental alignments.
- Allowed zero-length operations on null pointers.
- Named loops.
- Ranges in **case** labels.
- **\_Lengthof** operator for number of elements in an array.
- Memory reversal, endian-aware load and store and rotate operations in `<stdbit.h>`.
- New header `<stdmchar.h>` for text transcoding utilities.

### M.3 Fifth Edition

- 1 Major changes in this fifth edition (**\_\_STDC\_VERSION\_\_** 202311L) include:
  - add new keywords such as **bool**, **static\_assert**, **true**, **false**, **thread\_local** and others, and allowed implementations to provide macros for the older spelling with a leading under-score followed by a capital letter as well as defining old and new keywords as macros to enable transition of programs easily;
  - removed integer width constraints and obsolete sign representations (so-called “1’s complement” and “sign-magnitude”);
  - added a one-argument version of **static\_assert**;
  - removed support for function definitions with identifier lists;
  - mandated function declarations whose parameter list is empty be treated the same as a parameter list which only contain a single **void**;
  - harmonization with ISO/IEC 9945 (POSIX):
    - extended month name formats for **strftime**
    - integration of functions: **gmtime\_r**, **localtime\_r**, **memccpy**, **strdup**, **strndup**
  - harmonization with floating-point standard ISO/IEC 60559:
    - integration of binary floating-point technical specification ISO/IEC TS 18661-1:2014
    - integration of decimal floating-point technical specification ISO/IEC TS 18661-2:2015
    - integration of floating-point types technical specification ISO/IEC TS 18661-3:2015
    - integration of mathematical functions technical specification ISO/IEC TS 18661-4:2015
    - new maximum and minimum functions for operations recommended in ISO/IEC 60559:2020
  - made the **DECIMAL\_DIG** macro obsolescent;

- added version test macros to library headers that contained changes to aid in upgrading and portability to be used alongside the `__STDC_VERSION__` macro;
- allowed placement of labels in front of declarations and at the end of compound statement;
- added the attributes feature, which includes the attributes:
  - **deprecated**, for marking entities as discouraged for future use;
  - **fallthrough**, for explicitly marking cases where falling through in switches or labels is intended rather than accidental;
  - **maybe\_unused**, for marking entities which can end up not being used;
  - **nodiscard**, for marking entities which, when used, should have their value handled in some way by a program;
  - **noreturn**, for indicating a function shall never return;
  - **reproducible**, for marking function types for which inputs always produce predictable output if given the same input (e.g. cached data) but for which the order of such calls still matter; and,
  - **unsequenced**, for marking function types which always produce predictable output and have no dependencies upon other data (and other relevant caveats);
- added the `u8` character prefix to match the `u8` string prefix;
- mandated all `u8`, `u`, and `U` strings be UTF-8, UTF-16, and UTF-32, respectively, as defined by ISO/IEC 10646;
- separated the literal, wide literal, and UTF-8 literal, UTF-16 literal, and UTF-32 literal encodings for strings and characters and now have a solely execution-based version of these, particularly execution and wide execution encodings;
- added `mbrtoc8` and `c8rtomb` functions missing from `<uchar.h>`;
- compound literals may also include storage-class specifiers as part of the type to change the lifetime of the compound literal (and possibly turn it into a constant expression)
- added the **constexpr** specifier for object definitions and improved what is recognized as a constant expression in conjunction with the **constexpr** storage-class specifier;
- added support for initialization of objects with empty braces;
- added the **typeof** and **typeof\_unqual** operations for deducing the type of an expression;
- improved tag compatibility rules, enabling more types to be compatible with other types;
- added bit-precise integer types `_BitInt( N )` and `unsigned _BitInt( N )`, where `N` can be an integer constant expression whose value is from one to `BITINT_MAXWIDTH`, inclusive.
- improved rules for handling enumerations without underlying types, in particular allowing for enumerations without fixed underlying type to have value representations that have a greater range than `int`;
- added a new colon-delimited type specifier for enumerations to specify a fixed underlying type (and which, subject to an implementation's definitions governing such constructs, adopt the fixed underlying type's rules for padding, alignment, and sizing within structures and unions as well as with bit-fields);
- added a new header `<stdbit.h>` and a suite of bit and byte-handling utilities for portable access to many implementations' most efficient functionality;
- modified existing functions to preserve the **const**-ness of the type placed into the function;

- added a feature to embed binary data as faithfully as possible with a new preprocessor directive **#embed**;
- added a **nullptr** constant and a **nullptr\_t** type with a well-defined underlying representation identical to a pointer to **void**;
- added the **\_\_VA\_OPT\_\_** specifier and clarified language in the handling of macro invocation and arguments;
- mandated support for variably modified types (but not variable length arrays themselves);
- parameter names may be omitted in function definitions;
- ellipses on functions may appear without a preceding parameter in the parameter list of functions and **va\_start** no longer requires such an argument to be passed to it;
- Unicode identifiers allowed in syntax following Unicode Standard Annex, UAX #31;
- added the **memset\_explicit** function for making sensitive information inaccessible;
- added **memalignment** function to query the alignment of a pointer;
- certain type definitions (i.e. exact-width integer types such as **int128\_t**), bit-precise integer types, and extended integer types may exceed the normal boundaries of **intmax\_t** and **uintmax\_t** for signed and unsigned integer types, respectively;
- names of functions, macros, and variables in this document, where clarified, are potentially reserved rather than reserved to avoid undefined behavior for a large class of identifiers used by programs existing and to be created;
- mandated support for **call\_once**;
- allowed **ptrdiff\_t** to be an integer type with a width of at least 16, rather than requiring an integer type with a width of at least 17;
- added the **\_\_has\_include** feature for conditional inclusion expression preprocessor directives to check if a header is available for inclusion;
- changed the type qualifiers of the **\_Imaginary\_I** and **\_Complex\_I** macros;
- added qualifier preserving macros for **bsearch**, **memchr**, **strchr**, **strpbrk**, **strrchr**, **strstr**, **wcschr**, **wcspbrk**, **wcsrchr**, **wmemchr**, and **wcsstr**;
- added (U+0040, COMMERCIAL AT), \$ (U+0024, DOLLAR SIGN), and ` (U+0060, GRAVE ACCENT, “Backtick”) into the source and execution character set;
- enhanced the **auto** type specifier for single object definitions using type inference;
- added the **#elifdef** and **#elifndef** conditional inclusion preprocessor directives;
- added the **#warning** preprocessing directive;
- binary integer literals and appropriate formatting for input/output of binary integer numbers;
- digit separators with **'** (single quotation mark);
- removed conditional support for mixed wide and narrow string literal concatenation;
- added support for additional time bases, as well as **timespec\_getres**, in **<time.h>**;
- added support for new interface **timegm** to retrieve the broken-down time, in **<time.h>**;
- zero-sized reallocations with **realloc** are undefined behavior;
- added **free\_sized** and **free\_aligned\_sized** functions;

- added an **unreachable** feature which has undefined behavior if reached during program execution;
- added **printf** and **scanf** length modifiers for **intN\_t**, **int\_fastN\_t**, **uintN\_t**, and **uint\_fastN\_t**.

## M.4 Fourth Edition

- 1 There were no major changes in the fourth edition (**\_\_STDC\_VERSION\_\_** 201710L), only technical corrections and clarifications.

## M.5 Third Edition

- 1 Major changes in the third edition (**\_\_STDC\_VERSION\_\_** 201112L) included:
  - conditional (optional) features (including some that were previously mandatory)
  - support for multiple threads of execution including an improved memory sequencing model, atomic objects, and thread storage (`<stdatomic.h>` and `<threads.h>`)
  - additional floating-point characteristic macros (`<float.h>`)
  - querying and specifying alignment of objects (`<stdalign.h>`, `<stdlib.h>`)
  - Unicode characters and strings (`<uchar.h>`) (originally specified in ISO/IEC TR 19769:2004)
  - type-generic expressions
  - static assertions
  - anonymous structures and unions
  - no-return functions
  - macros to create complex numbers (`<complex.h>`)
  - support for opening files for exclusive access
  - removed the **gets** function (`<stdio.h>`)
  - added the **aligned\_alloc**, **at\_quick\_exit**, and **quick\_exit** functions (`<stdlib.h>`)
  - (conditional) support for bounds-checking interfaces (originally specified in ISO/IEC TR 24731-1:2007)
  - (conditional) support for analyzability

## M.6 Second Edition

- 1 Major changes in the second edition (**\_\_STDC\_VERSION\_\_** 199901L) included:
  - restricted character set support via digraphs and `<iso646.h>` (originally specified in ISO/IEC 9899:1990/Amd 1:1995)
  - wide character library support in `<wchar.h>` and `<wctype.h>` (originally specified in ISO/IEC 9899:1990/Amd 1:1995)
  - more precise aliasing rules via effective type
  - restricted pointers
  - variable length arrays
  - flexible array members
  - **static** and type qualifiers in parameter array declarators

- complex (and imaginary) support in `<complex.h>`
- type-generic math macros in `<tgmath.h>`
- the **long long int** type and library functions
- extended integer types
- increased minimum translation limits
- additional floating-point characteristics in `<float.h>`
- remove implicit **int**
- reliable integer division
- universal character names (`\u` and `\U`)
- extended identifiers
- hexadecimal floating literals and `%a` and `%A` **printf/scanf** conversion specifiers
- compound literals
- designated initializers
- `//` comments
- specified width integer types and corresponding library functions in `<inttypes.h>` and `<stdint.h>`
- remove implicit function declaration
- preprocessor arithmetic done in **intmax\_t/uintmax\_t**
- mixed declarations and statements
- new block scopes for selection and iteration statements
- integer constant type rules
- integer promotion rules
- macros with a variable number of arguments (**\_\_VA\_ARGS\_\_**)
- the **vscanf** family of functions in `<stdio.h>` and `<wchar.h>`
- additional math library functions in `<math.h>`
- treatment of error conditions by math library functions (**math\_errhandling**)
- floating-point environment access in `<fenv.h>`
- ISO/IEC 60559 (also known as IEC 559 or IEEE 754 arithmetic) support
- trailing comma allowed in **enum** declaration
- `%lf` conversion specifier allowed in **printf**
- inline functions
- the **snprintf** family of functions in `<stdio.h>`
- boolean type in `<stdbool.h>`
- idempotent type qualifiers
- empty macro arguments

- new structure type compatibility rules (tag compatibility)
- additional predefined macro names
- **\_Pragma** preprocessing operator
- standard pragmas
- **\_\_func\_\_** predefined identifier
- **va\_copy** macro
- additional **strftime** conversion specifiers
- LIA compatibility annex
- deprecate **ungetc** at the beginning of a binary file
- remove deprecation of aliased array parameters
- conversion of array to pointer not limited to lvalues
- relaxed constraints on aggregate and union initialization
- relaxed restrictions on portable header names
- **return** without expression not permitted in function that returns a value (and vice versa)

## M.7 First Edition, Amendment 1

- 1 Major changes in the amendment to the first edition (**\_\_STDC\_VERSION\_\_** 199409L) included:

- addition of the predefined **\_\_STDC\_VERSION\_\_** macro
- restricted character set support via digraphs and `<iso646.h>`
- wide character library support in `<wchar.h>` and `<wctype.h>`

## Bibliography

- [1] ISO/IEC 14882 Programming languages — C++.
- [2] ISO/IEC 646:1991, *Information technology — ISO 7-bit coded character set for information interchange*.
- [3] ISO/IEC 9945–2:1993, *Information technology — Portable Operating System Interface (POSIX) — Part 2: Shell and Utilities*.
- [4] ISO/IEC TR 10176:2003, *Information technology — Guidelines for the preparation of programming language standards*.
- [5] ISO/IEC 10967–1:2012, *Information technology — Language independent arithmetic — Part 1: Integer and floating point arithmetic*.
- [6] ISO/IEC TR 19769:2004, *Information technology — Programming languages, their environments and system software interfaces — Extensions for the programming language C to support new character data types*.
- [7] ISO/IEC TR 24731–1:2007, *Information technology — Programming languages, their environments and system software interfaces — Extensions to the C library — Part 1: Bounds-checking interfaces*.
- [8] ISO 80000–3, *Quantities and units — Part 3: Space and time*.
- [9] IEC 60559:1989, *Binary floating-point arithmetic for microprocessor systems*.
- [10] ISO/IEC/IEEE 60559:2011, *Floating-point arithmetic*.
- [11] IEEE 754–1985 *IEEE Standard for Binary Floating-Point Arithmetic*.
- [12] IEEE 754–2019 *IEEE Standard for Floating-Point Arithmetic*.
- [13] ANSI/IEEE 854–1987, *American National Standard for Radix-Independent Floating-Point Arithmetic*.
- [14] ANSI X3/TR–1–82 (1982), *American National Dictionary for Information Processing Systems, Information Processing Systems Technical Report*.
- [15] “The C Reference Manual” by Dennis M. Ritchie, a version of which was published in *The C Programming Language* by Brian W. Kernighan and Dennis M. Ritchie, Prentice-Hall, Inc., (1978). Copyright owned by AT&T.
- [16] 1984 */usr/group Standard* by the */usr/group Standards Committee*, Santa Clara, California, USA, November 1984.
- [17] The Unicode Consortium. *Unicode Standard Annex, UAX #31, Unicode Identifier and Pattern Syntax [online]*. Edited by Mark Davis and Robin Leroy. Available at <https://www.unicode.org/reports/tr31>.

## Index

- ! (exclamation-mark punctuation), 72
- ! (logical negation operator), 87
- != (inequality operator), 93
- != (not-equal punctuation), 72
- # (hash punctuation), 72
- # preprocessing directive, 199
- # punctuation, 175
- ## (hash-hash punctuation), 72
- #define preprocessing directive, 191
- #elif preprocessing directive, 179
- #elifdef preprocessing directive, 180
- #elifndef preprocessing directive, 180
- #else preprocessing directive, 180
- #embed preprocessing directive, 183
- #endif preprocessing directive, 180
- #error preprocessing directive, 9, 198
- #if preprocessing directive, 179
- #ifdef preprocessing directive, 180
- #ifndef preprocessing directive, 180
- #include preprocessing directive, 11, 182
- #line preprocessing directive, 197
- #pragma preprocessing directive, 198
- #undef preprocessing directive, 195, 206
- [*x*], 7
- [*x*], 8
- & (bitwise AND operator), 93
- && (logical AND operator), 17, 94
- &= (bitwise AND assignment operator), 98
- ' ' (space character), 11, 20, 54, 218, 219, 483
- ( (opening parenthesis punctuation), 72
- () (cast operator), 89
- () (function-call operator), 80
- () (parentheses punctuation), 137, 163, 165
- (){} (compound-literal operator), 83
- ) (closing parenthesis punctuation), 72
- \* (asterisk punctuation), 72, 134, 135
- \* (indirection operator), 79, 86
- \* (multiplication operator), 89, 589
- \*= (asterisk-equal punctuation), 72
- \*= (multiplication assignment operator), 98
- + (addition operator), 79, 86, 90, 591
- + (plus punctuation), 72
- + (unary plus operator), 87
- + format flag, 349, 448
- ++ (plus-plus punctuation), 72
- ++ (postfix increment operator), 51, 82
- ++ (prefix increment operator), 51, 86
- += (addition assignment operator), 98
- += (plus-equal punctuation), 72
- , (comma operator), 17, 99
- , (comma punctuation), 72, 79, 104, 112, 116, 123, 144
- (minus punctuation), 72
- (subtraction operator), 90, 591
- (unary minus operator), 87, 563
- format flag, 349, 448
- (minus-minus punctuation), 72
- (postfix decrement operator), 51, 83
- (prefix decrement operator), 51, 86
- = (minus-equal punctuation), 72
- = (subtraction assignment operator), 98
- > (minus-greater punctuation), 72
- > (structure/union pointer operator), 81
- . (dot punctuation), 72, 144
- . (structure/union member operator), 51, 81
- ... (ellipsis punctuation), 72, 137, 191
- / (division operator), 89, 589
- / (slash punctuation), 72
- /\* \*/ (comment delimiters), 74
- // (comment delimiter), 74
- /= (division assignment operator), 98
- /= (slash-equal punctuation), 72
- : (colon punctuation), 72, 112
- :> (alternative spelling of `]`), 73
- :> (colon greater punctuation), 72
- ; (semicolon punctuation), 72, 104, 112, 162, 165, 166
- < (less punctuation), 72
- < (less-than operator), 92
- <: (alternative spelling of `[]`), 73
- <: (less-colon punctuation), 72
- << (left-shift operator), 91
- << (less-less punctuation), 72
- <<= (left-shift assignment operator), 98
- <<= (less-less equal punctuation), 72
- <= (less-equal punctuation), 72
- <= (less-than-or-equal-to operator), 92
- <% (alternative spelling of `{}`), 73
- <% (less-percent punctuation), 72
- <assert.h> header, 186–188, 205, 208, 232, 507
- <complex.h> header, 25, 30, 108, 146, 201, 204, 205, 209–216, 416, 417, 488, 507, 508, 590, 591, 592, 607, 609, 644, 646, 654, 724, 725
- <ctype.h> header, 205, 217, 218–220, 488, 509
- <errno.h> header, 157, 205, 221, 488, 509, 672
- <fcntl.h> header, 9, 14, 16, 25, 30, 35, 98, 157, 205, 222, 224, 225, 227–234, 250, 488, 509, 559, 563, 566–569, 577, 581, 583, 725
- <float.h> header, 9, 10, 22, 24, 25, 29, 30, 108, 205, 235, 248, 353, 376, 452, 466, 488, 489, 510, 511, 555, 559, 565, 601, 602,



- 651, 724, 725
- <inttypes.h> header, 205, **236**, 237, 238, 351, 450, **488**, **511**, 725
- <iso646.h> header, 9, 10, 205, **239**, **511**, 724, 726
- <limits.h> header, 9, 10, 22, **23**, 24, 40, 41, 205, **240**, 327, 328, **512**, 555, 651
- <locale.h> header, 157, 205, **241**, 242, **488**, **512**
- <math.h> header, ii, 9, 25, 30, 35, 77, 157, 205, 226–228, **247**, 248, 250–253, **254–287**, 288, **289–294**, 295, 354, 416–418, 453, **488**, 489, **512**, 523, 524, 558, 559, 565, **571**, 572, 577, 581, 583–587, 590, 591, 599, **609**, **622–624**, 634, 644, 646, 651, 654, 725
- <setjmp.h> header, 205, **296**, 297, **532**
- <signal.h> header, 205, **298**, 300, **489**, **532**
- <stdalign.h> header, 9, 10, 205, **301**, **532**, 724
- <stdarg.h> header, 9, 10, 137, 205, **302**, 303–305, 364–366, 458–460, **532**, 681–684, 705–708
- <stdatomic.h> header, 201, 204, 205, 299, **307**, 308, 311, 312, 314–316, **489**, **532**, 640, 724
- <stdbit.h> header, 9, 205, **318**, 319–328, **489**, **533**, 721, 722
- <stdbool.h> header, 9, 10, 205, **330**, **489**, **536**, 725
- <stdckdint.h> header, 205, **331**, **489**, **536**
- <stddef.h> header, 9, 10, 52, 70, 72, 87, 88, 91, 149, 185, 186, 205, 238, **332**, 333, 334, 361, 362, **536**, **673**
- <stdint.h> header, 9, 10, 22, 24, 179, 205, 236, 237, 327, **335**, 337, 338, 350, 357, 358, 449, 455, **489**, **536**, 651, **673**, 725
- <stdio.h> header, 16, 25, 30, 35, 58, 157, 182, 186, 205, 226, 228, **339**, 343–345, 347, 348, 354, 356, 360–372, 433, 447, 453, 457–459, 461–464, **489**, **537**, 538, 646, **673**, 674–685, 703, 705, 706, 724, 725
- <stdlib.h> header, 9, 25, 30, 35, 205, 207, 226, 228, **373**, 374, 375, 377, 379–391, **489**, **538**, 539, 540, 599, **624**, 625, **626**, **627**, 646, 672, **686**, 687–692, 724
- <stdmchar.h> header, **393**, 395, 396, 400, 403, 404, **540**, 721
- <stdnoreturn.h> header, 9, 10, 155, 205, **406**, **544**
- <string.h> header, 9, 185, 188, 189, 205, **407**, 408–415, **490**, **544**, 545, **693**, 694–699
- <tgmath.h> header, 35, 205, **416**, 419, **545**, 558, **571**, **628**, 630, 725
- <threads.h> header, 157, 201, 204, 205, **421**, 422–429, **490**, **546**, 724
- <time.h> header, 205, 421, **431**, 432–437, 476, **490**, **547**, **699**, 700–702, 723
- <uchar.h> header, 70, 72, 205, 394, **441**, 442–445, **547**, 722, 724
- <wchar.h> header, 25, 30, 35, 157, 205, 226, 228, 237, 340, 394, **446**, 447, 453, 457–465, 467–480, **490**, **548**, 549, 550, **627**, **628**, 646, **702**, 703–717, 724–726
- <wctype.h> header, 205, **482**, 483–487, **490**, **550**, 724, 726
- = (equal-sign punctuator), **72**, 104, 116, 144
- = (simple assignment operator), **97**
- == (equal-equal punctuator), **72**
- == (equality operator), **93**
- > (greater punctuator), **72**
- > (greater-than operator), **92**
- >= (greater-equal punctuator), **72**
- >= (greater-than-or-equal-to operator), **92**
- >> (greater greater punctuator), **72**
- >> (right-shift operator), **91**
- >>= (greater-greater-equal punctuator), **72**
- >>= (right-shift assignment operator), **98**
- ? (question-mark punctuator), **72**
- ?: (conditional operator), 17, **95**
- [ (opening bracket punctuator), **72**
- [ ] (array subscript operator), **79**, 86
- [ ] (brackets punctuator), 135, 144
- # format flag, 349, 448
- % (percent punctuator), **72**
- % (remainder operator), **89**
- %: (alternative spelling of #), **73**
- %: (percent-colon punctuator), **72**
- %: (alternative spelling of ##), **73**
- %: (percent-percent punctuator), **72**
- %= (percent-equal punctuator), **72**
- %= (remainder assignment operator), **98**
- %> (alternative spelling of }), **73**
- %> (percent-greater punctuator), **72**
- %A conversion specifier, 351, 437, 451
- %B conversion specifier, 351, 437, 450
- %C conversion specifier, 437
- %D conversion specifier, 438
- %E conversion specifier, 351, 450
- %F conversion specifier, 351, 438, 450
- %G conversion specifier, 351, 438, 450
- %H conversion specifier, 438
- %I conversion specifier, 438
- %M conversion specifier, 438
- %R conversion specifier, 438
- %S conversion specifier, 438
- %T conversion specifier, 438
- %U conversion specifier, 438
- %V conversion specifier, 438
- %W conversion specifier, 438
- %X conversion specifier, 351, 438, 450
- %Y conversion specifier, 438
- %Z conversion specifier, 438
- %[ conversion specifier, 359, 456

- %% conversion specifier, 353, 360, 452, 457
- %a conversion specifier, 351, 358, 437, 451, 456
- %b conversion specifier, 351, 437, 450
- %c conversion specifier, 352, 358, 437, 452, 456
- %d conversion specifier, 351, 358, 438, 450, 455
- %e conversion specifier, 351, 358, 438, 450, 456
- %f conversion specifier, 351, 358, 450, 456
- %g conversion specifier, 351, 358, 438, 450, 456
- %h conversion specifier, 438
- %i conversion specifier, 351, 450
- %j conversion specifier, 438
- %m conversion specifier, 438
- %n conversion specifier, 353, 359, 438, 452, 457
- %o conversion specifier, 351, 358, 450, 455
- %p conversion specifier, 353, 359, 438, 452, 457
- %r conversion specifier, 438
- %s conversion specifier, 353, 359, 452, 456
- %t conversion specifier, 438
- %u conversion specifier, 351, 358, 438, 450, 456
- %w conversion specifier, 438
- %x conversion specifier, 351, 358, 438, 450, 456
- %y conversion specifier, 438
- %z conversion specifier, 438
- & (address operator), 51, 86
- & (ampersand punctuator), 72
- &= (ampersand-equal punctuator), 72
- && (ampersand-ampersand punctuator), 72
- \ (backslash character), 11, 19, 68
- \ (backslash escape sequence), 69, 202
- \ " (double-quote escape sequence), 69, 71, 202
- \ ' (single-quote escape sequence), 69, 71
- \0 (null character), 19, 70, 71
  - padding of binary stream, 341
- \? (question-mark escape sequence), 69
- \U (universal character names), 58
- \a (alert escape sequence), 21, 69
- \b (backspace escape sequence), 21, 69
- \ (escape character), 68
- \f (form-feed escape sequence), 21, 69, 219
- \n (new-line escape sequence), 21, 69, 219
- \octal digit (octal-character escape sequence), 69
- \r (carriage-return escape sequence), 21, 69, 219
- \t (horizontal-tab escape sequence), 21, 69, 218, 219, 483
- \u (universal character names), 58
- \v (vertical-tab escape sequence), 21, 69, 219
- \_Alignas**, 55
- \_Alignof**, 55
- \_Atomic**, 42
- \_Atomic** keyword, 55
- \_Atomic** type qualifier, 127
- \_Atomic** type specifier, 124
- \_BitInt**, 23, 40, 48, 50, 51, 55, 62, 63, 110, 111, 326, 327, 722
- \_BitInt** keyword, 55
- \_BitInt** type, 110
- \_Bool**, 55
- \_C** identifier suffix, 338, 489, 536
- \_Complex**, 26, 41, 50, 55, 108–111, 209, 590, 601, 607
- \_Complex** keyword, 55
- \_Complex** type, 41, 110, 589
- \_Complex\_I** macro, 209, 507, 723
- \_DECIMAL\_DIG** identifier suffix, 29, 353, 376, 452, 466, 565, 566
- \_Decimal128**, 40
- \_Decimal128** keyword, 55
- \_Decimal128** type, 110
- \_Decimal128x** type, 600
- \_Decimal32**, 40
- \_Decimal32** keyword, 55
- \_Decimal32** type, 110
- \_Decimal32\_t** type, 247, 518, 649
- \_Decimal64**, 40
- \_Decimal64** keyword, 55
- \_Decimal64** type, 110
- \_Decimal64\_t** type, 247, 518, 649
- \_Decimal64x** type, 600
- \_DecimalN** type, 599
- \_DecimalN\_t** type, 610
- \_DecimalNx** type, 601
- \_Exit** function, 299, 385, 386, 539, 640, 650
- \_Float128\_t** type, 610, 611
- \_Float128x** type, 600
- \_Float16\_t** type, 610, 611
- \_Float32\_t** type, 610, 611
- \_Float32x** type, 600
- \_Float64\_t** type, 610, 611
- \_Float64x** type, 600
- \_FloatN** type, 599
- \_FloatN\_t** type, 610
- \_FloatNx** type, 601
- \_Generic**, i, 55, 78, 120, 128, 143, 334, 418, 720
- \_Generic** keyword, 55
- \_H\_** identifier suffix, 205
- \_IOFBF** macro, 339, 347, 348, 537
- \_IOLBF** macro, 339, 348, 537
- \_IONBF** macro, 339, 347, 348, 537
- \_Imaginary\_I** (obsolete), 723
- \_Lengthof** keyword, 55
- \_Lengthof** operator, 86, 87
- \_MAX** identifier suffix, 24, 49, 337, 338, 489, 511, 536, 603
- \_MIN** identifier suffix, 24, 337, 338, 489, 511, 536, 604
- \_Noreturn**, 131
- \_Noreturn** attribute, 150, 154, 155
- \_Noreturn** keyword, 55
- \_PRINTF\_NAN\_LEN\_MAX** macro, 340, 537
- \_Pragma** operator, 201

- `__Static_assert`, 55
- `__Thread_local`, 55
- `__WIDTH` identifier suffix, 24, 337, 338, 489, 536
- `__STDC` identifier prefix, 203
- `__STDC_VERSION` identifier prefix, 205
- `__DATE__` macro, 199, 648
- `__FILE__` macro, 181, 199, 208
- `__LINE__` macro, 197, 198, 199, 208, 633
- `__STDC_ANALYZABLE__` macro, 201, 719
- `__STDC_EMBED_EMPTY__` macro, 179, 199
- `__STDC_EMBED_FOUND__` macro, 179, 199
- `__STDC_EMBED_NOT_FOUND__` macro, 178, 199
- `__STDC_ENDIAN_BIG__` macro, 318, 319, 533
- `__STDC_ENDIAN_LITTLE__` macro, 318, 319, 394, 533
- `__STDC_ENDIAN_NATIVE__` macro, 291, 292, 318, 394, 533, 622–624, 626–628, 649
- `__STDC_HOSTED__` macro, 199
- `__STDC_IEC_559__` (obsolete), 24, 201, 203, 204, 523, 558
- `__STDC_IEC_60559_BFP__` macro, 9, 24, 201, 204, 523, 558, 585–587, 598–600
- `__STDC_IEC_60559_COMPLEX__` macro, 24, 201, 204, 589
- `__STDC_IEC_60559_DFP__` macro, 9, 30, 201, 204, 227, 231, 232, 247, 254–281, 282, 283–287, 288, 289–292, 375, 377, 419, 467, 510, 511, 517, 523, 539, 549, 556–558, 585–587, 598–600
- `__STDC_IEC_60559_TYPES__` macro, 201, 204, 508, 511, 524, 539, 549, 598–600
- `__STDC_ISO_10646__` macro, 200, 646
- `__STDC_LIB_EXT1__` macro, 201, 204, 509, 536, 538, 540, 545, 547, 550, 671
- `__STDC_LITERAL_UTF16__` macro, 200
- `__STDC_LITERAL_UTF32__` macro, 200
- `__STDC_LITERAL_UTF8__` macro, 200
- `__STDC_MB_MIGHT_NEQ_WC__` macro, 47, 200, 332
- `__STDC_NO_ATOMICS__` macro, 201, 307, 532
- `__STDC_NO_COMPLEX__` macro, 65, 201, 209, 507
- `__STDC_NO_THREADS__` macro, 16, 201, 421, 546
- `__STDC_NO_VLA__` macro, 201
- `__STDC_UTF_16__` macro, 200
- `__STDC_UTF_32__` macro, 200
- `__STDC_VERSION_ASSERT_H__` macro, 208, 507
- `__STDC_VERSION_COMPLEX_H__` macro, 209, 507
- `__STDC_VERSION_FENV_H__` macro, 222, 509
- `__STDC_VERSION_FLOAT_H__` macro, 235, 510
- `__STDC_VERSION_INTTYPES_H__` macro, 236, 511
- `__STDC_VERSION_LIMITS_H__` macro, 240, 512
- `__STDC_VERSION_MATH_H__` macro, 247, 512
- `__STDC_VERSION_SETJMP_H__` macro, 296, 532
- `__STDC_VERSION_STDARG_H__` macro, 302, 532
- `__STDC_VERSION_STDATOMIC_H__` macro, 307, 532
- `__STDC_VERSION_STDBIT_H__` macro, 318, 533
- `__STDC_VERSION_STDCCKDINT_H__` macro, 331, 536
- `__STDC_VERSION_STDEDEF_H__` macro, 332, 536
- `__STDC_VERSION_STDINT_H__` macro, 335, 536
- `__STDC_VERSION_STDIO_H__` macro, 339, 537
- `__STDC_VERSION_STDLIB_H__` macro, 373, 538
- `__STDC_VERSION_STRING_H__` macro, 407, 544
- `__STDC_VERSION_TGMATH_H__` macro, 416
- `__STDC_VERSION_TIME_H__` macro, 431, 547
- `__STDC_VERSION_UCHAR_H__` macro, 441, 548
- `__STDC_VERSION_WCHAR_H__` macro, 446, 548
- `__STDC_VERSION__` macro, 200, 720–722, 724, 726
- `__STDC_WANT_IEC_60559_EXT__` macro, 510, 523, 565, 585–587
- `__STDC_WANT_IEC_60559_TYPES_EXT__` macro, 508, 511, 524, 539, 549, 602, 607, 622–624, 625, 626–628, 630
- `__STDC_WANT_LIB_EXT1__` macro, 509, 536, 538, 540, 545, 547, 550, 671, 672, 674–690, 693–715
- `__STDC_WIDE_LITERAL_UTF16__` macro, 200
- `__STDC_WIDE_LITERAL_UTF32__` macro, 201
- `__STDC_WIDE_LITERAL_UTF8__` macro, 200
- `__STDC__` macro, 180, 199
- `__TIME__` macro, 200, 648
- `__VA_ARGS__` identifier, 190, 191, 192, 725
- `__VA_OPT__` identifier, 190, 191–193, 723
- `__bool_true_false_are_defined` (obsolete), 330, 489, 536
- `__cplusplus` macro, 180, 199
- `__deprecated__` attribute, 151
- `__func__` identifier, 57, 58, 208, 636, 726
- `__has_c_attribute`, 149, 151–155, 157, 179, 180, 199, 720
- `__has_c_attribute` operator, 199
- `__has_embed`, 54, 73, 179, 181, 199
- `__has_embed` operator, 199
- `__has_include`, 54, 73, 179, 180, 199, 723
- `__has_include` operator, 199
- `__limit__` embed parameter, 177
- `__nodiscard__` attribute, 149

- \_explicit** identifier suffix, 307, 315, 533
- \_r** identifier suffix, 435
- \_t** identifier suffix, 327–329, 335, 336, 338, 489, 524, 535, 536, 610, 611, 724
- wchar\_t** character literal, 68
- wchar\_t** string literal, 71
- { } (braces punctuator), 116, 123, 144, 161
- { } (compound-literal operator), 83
- { (opening brace punctuator), 72
- } (closing brace punctuator), 72
- ] (closing bracket punctuator), 72
- ^ (bitwise exclusive OR operator), 94
- ^ (caret punctuator), 72
- ^= (bitwise exclusive OR assignment operator), 98
- ^= (caret-equal punctuator), 72
- constant expression, 161
- integer constant expression, 164, 178
- address constant, 101
- | (bitwise inclusive OR operator), 94
- | (vertical-line punctuator), 72
- |= (bitwise inclusive OR assignment operator), 98
- |= (vertical-line-equal punctuator), 72
- || (logical OR operator), 17, 95
- || (vertical-vertical punctuator), 72
- ~ (bitwise complement operator), 87
- ~ (tilde punctuator), 72
- 0 format flag, 349, 448
  
- abort** function, 155, 208, 298, 299, 308, 342, 384, 539, 640, 650, 687
- abort\_handler\_s** function, 540, 687
- abs** function, ii, 206, 207, 388, 539, 721
- abs** macro, 206, 207
- absolute-value function
  - complex, 214, 597
  - integer, 237, 388
  - real, 269, 577
- abstract declarator, 139
- abstract machine, 14
- access, 8, 128
- access (verb), 3
- acos** function, 226, 254, 255, 417, 513, 565, 572, 609
- acos** type-generic macro, 417
- acosd128** function, 228, 254, 518
- acosd32** function, 228, 254, 518
- acosd64** function, 228, 254, 518
- acosdN** function, 613
- acosdNx** function, 613
- acosf** function, 226, 254, 513
- acosfN** function, 613
- acosfNx** function, 613
- acosh** function, 259, 417, 513, 565, 574
- acosh** type-generic macro, 417
- acoshd128** function, 259, 518
- acoshd32** function, 259, 518
- acoshd64** function, 259, 518
- acoshdN** function, 614
- acoshdNx** function, 614
- acoshf** function, 259, 419, 513
- acoshfN** function, 614
- acoshfNx** function, 614
- acoshl** function, 259, 513
- acosl** function, 226, 254, 513, 609
- acospi** function, 257, 513, 565, 573
- acospi** type-generic macro, 417
- acospid128** function, 257, 518
- acospid32** function, 257, 518
- acospid64** function, 257, 518
- acospidN** function, 613
- acospidNx** function, 613
- acospiif** function, 257, 513
- acospiifN** function, 613
- acospiifNx** function, 613
- acospil** function, 257, 513
- acquire fence, 311
- acquire operation, 17
- active position, 21
- add and round to narrower type, 287
- addition assignment operator (+=), 98
- addition operator (+), 79, 86, 90, 591
- additive expression, 90, 591
- address operator (&), 51, 86
- address-free, 312
- adjustment, 82
- aggregate initialization, 145
- aggregate type, 42
- alert, 21
- alert escape sequence (\a), 21, 69
- aliasing, 76
- alignas**, 132
- alignas** keyword, 55
- aligned\_alloc** function, 381, 382, 383, 539, 634, 643, 650, 724
- alignment, 3, 46, 382
  - pointer, 42, 52
  - structure/union member, 114
- alignment header, 301
- alignment of memory, 391
- alignment specifier, 132
- alignof** keyword, 55
- alignof** operator, 86, 87
- allocated storage
  - order and contiguity, 381
- alternative spellings header, 239
- and** macro, 239, 512
- AND operator
  - bitwise (&), 93
  - bitwise assignment (&=), 98
  - logical (&&), 94
- AND operator

- logical (&&), 17
- and\_eq** macro, 239, 512
- anonymous structure, 113
- anonymous union, 113
- argc** (**main** function parameter), 13
- argument, 3
  - array, 172
  - complex, 215
  - default promotion, 80
  - function, 80, 172
  - macro, substitution, 191
  - variable, 191
- argv** (**main** function parameter), 13
- arithmetic
  - pointer, 90
- arithmetic constant expression, 101
- arithmetic conversions
  - usual, *see* usual arithmetic conversion
- arithmetic operator
  - additive, 90, 591
  - bitwise, 87, 93, 94
  - increment and decrement, 82, 86
  - multiplicative, 89, 589
  - shift, 91
  - unary, 87
- arithmetic type, 41
- arithmetically negate, 3
- array
  - argument, 172
  - declarator, 135
  - initialization, 145
  - multidimensional, 79
  - parameter, 172
  - storage order, 79
  - subscript operator ([  ]), 79, 86
  - subscripting, 79
  - type, 41
  - type conversion, 52
  - variable length, 134, 135, 201
- arrow operator (->), 81
- as-if rule, 14
- asctime** function, 199, 200, 435, 436, 547, 644
- asctime\_s** function, 547, 700, 701
- asin** function, 255, 417, 513, 565, 572
- asin** type-generic macro, 417
- asind128** function, 255, 518
- asind32** function, 255, 518
- asind64** function, 255, 518
- asindN** function, 613
- asindNx** function, 613
- asinf** function, 255, 513
- asinfN** function, 613
- asinfNx** function, 613
- asinh** function, 259, 260, 417, 513, 565, 574
- asinh** type-generic macro, 417
- asinh128** function, 260, 518
- asinh32** function, 260, 518
- asinhN** function, 614
- asinhNx** function, 614
- asinhf** function, 260, 513
- asinhfN** function, 614
- asinhfNx** function, 614
- asinh1** function, 260, 513
- asinl** function, 255, 513
- asinp** function, 257, 513, 565, 573
- asinp** type-generic macro, 417
- asinp128** function, 257, 518
- asinp32** function, 257, 518
- asinp64** function, 257, 518
- asinpN** function, 613
- asinpNx** function, 613
- asnpif** function, 257, 513
- asnpifN** function, 613
- asnpifNx** function, 613
- asnpil** function, 257, 513
- assert** macro, 152, 187, 189, 208, 232, 507, 639, 649
- assignment
  - compound, 98
  - conversion, 97
  - expression, 96
  - operator, 51, 96
  - simple, 97
- associativity of operator, 76
- asterisk punctuator (\*), 134, 135
- at\_quick\_exit** function, 384, 385, 386, 539, 634, 643, 724
- atan** function, 255, 354, 417, 453, 513, 565, 572
- atan** type-generic macro, 417
- atan2** function, 255, 256, 513, 565, 571, 572, 573, 592
- atan2** type-generic macro, 417
- atan2d128** function, 255, 518
- atan2d32** function, 255, 518
- atan2d64** function, 255, 518
- atan2dN** function, 613
- atan2dNx** function, 613
- atan2f** function, 255, 513
- atan2fN** function, 613
- atan2fNx** function, 613
- atan2l** function, 255, 513
- atan2pi** function, 258, 513, 565, 571, 573
- atan2pi** type-generic macro, 417
- atan2pid128** function, 258, 518
- atan2pid32** function, 258, 518
- atan2pid64** function, 258, 518
- atan2pidN** function, 613
- atan2pidNx** function, 613
- atan2pif** function, 258, 513
- atan2pifN** function, 613
- atan2pifNx** function, 613



- atan2pil** function, 258, 513
- atand128** function, 255, 518
- atand32** function, 255, 518
- atand64** function, 255, 518
- atandN** function, 613
- atandNx** function, 613
- atanf** function, 255, 513
- atanfN** function, 613
- atanfNx** function, 613
- atanh** function, 260, 417, 513, 565, 574
- atanh** type-generic macro, 417
- atanhd128** function, 260, 519
- atanhd32** function, 260, 518
- atanhd64** function, 260, 519
- atanhdN** function, 614
- atanhdNx** function, 614
- atanhf** function, 260, 513
- atanhfN** function, 614
- atanhfNx** function, 614
- atanhl** function, 260, 513
- atanl** function, 255, 419, 513
- atanpi** function, 257, 258, 513, 565, 573
- atanpi** type-generic macro, 417
- atanpid128** function, 258, 518
- atanpid32** function, 257, 518
- atanpid64** function, 257, 518
- atanpidN** function, 613
- atanpidNx** function, 613
- atanpif** function, 257, 513
- atanpifN** function, 613
- atanpifNx** function, 613
- atanpil** function, 257, 513
- atexit** function, 384, 385, 386, 539, 634, 643, 653
- atof** function, 226, 373, 374, 538
- atoi** function, 207, 373, 374, 538
- atol** function, 373, 374, 538
- atoll** function, 373, 374, 538
- atomic lock-free macro, 307, 312
- atomic operation, 17
- atomic type, 14, 42, 43, 51, 81, 82, 98, 124, 201, 313
- ATOMIC\_** identifier prefix, 489
- atomic\_** identifier prefix, 489
- atomic\_bool** type, 313, 315, 532
- ATOMIC\_BOOL\_LOCK\_FREE** macro, 307, 532
- atomic\_char** type, 313, 532
- atomic\_char16\_t** type, 313, 532
- ATOMIC\_CHAR16\_T\_LOCK\_FREE** macro, 307, 532
- atomic\_char32\_t** type, 313, 532
- ATOMIC\_CHAR32\_T\_LOCK\_FREE** macro, 307, 532
- atomic\_char8\_t** type, 313, 532
- ATOMIC\_CHAR8\_T\_LOCK\_FREE** macro, 307, 532
- ATOMIC\_CHAR\_LOCK\_FREE** macro, 307, 532
- atomic\_compare\_exchange\_strong** function, 82, 98, 99, 314, 315, 533
- atomic\_compare\_exchange\_strong\_explicit** function, 314, 533
- atomic\_compare\_exchange\_weak** function, 314, 315, 533
- atomic\_compare\_exchange\_weak\_explicit** function, 314, 533
- atomic\_exchange** function, 314, 533
- atomic\_exchange\_explicit** function, 314, 533
- atomic\_fetch\_** function, 315, 533
- atomic\_flag** type, 307, 308, 316, 532, 533
- atomic\_flag\_clear** function, 316, 317, 533
- atomic\_flag\_clear\_explicit** function, 316, 533
- ATOMIC\_FLAG\_INIT** macro, 307, 316, 532
- atomic\_flag\_test\_and\_set** function, 316, 533
- atomic\_flag\_test\_and\_set\_explicit** function, 316, 533
- atomic\_init** function, 308, 533
- atomic\_int** type, 308, 313, 532
- atomic\_int\_fast16\_t** type, 313, 533
- atomic\_int\_fast32\_t** type, 313, 533
- atomic\_int\_fast64\_t** type, 313, 533
- atomic\_int\_fast8\_t** type, 313, 533
- atomic\_int\_least16\_t** type, 313, 532
- atomic\_int\_least32\_t** type, 313, 532
- atomic\_int\_least64\_t** type, 313, 532
- atomic\_int\_least8\_t** type, 313, 532
- ATOMIC\_INT\_LOCK\_FREE** macro, 307, 532
- atomic\_intmax\_t** type, 313, 533
- atomic\_intptr\_t** type, 313, 533
- atomic\_is\_lock\_free** function, 299, 312, 533, 640
- atomic\_llong** type, 313, 532
- ATOMIC\_LLONG\_LOCK\_FREE** macro, 307, 532
- atomic\_load** function, 314, 315, 533
- atomic\_load\_explicit** function, 310, 311, 314, 533
- atomic\_long** type, 313, 532
- ATOMIC\_LONG\_LOCK\_FREE** macro, 307, 532
- ATOMIC\_POINTER\_LOCK\_FREE** macro, 307, 532
- atomic\_ptrdiff\_t** type, 313, 533
- atomic\_schar** type, 313, 532
- atomic\_short** type, 313, 532
- ATOMIC\_SHORT\_LOCK\_FREE** macro, 307, 532
- atomic\_signal\_fence** function, 312, 533
- atomic\_size\_t** type, 313, 533
- atomic\_store** function, 314, 533
- atomic\_store\_explicit** function, 310, 311, 314, 533
- atomic\_thread\_fence** function, 157, 311, 312, 533

- `atomic_uchar` type, 313, 532
- `atomic_uint` type, 313, 532
- `atomic_uint_fast16_t` type, 313, 533
- `atomic_uint_fast32_t` type, 313, 533
- `atomic_uint_fast64_t` type, 313, 533
- `atomic_uint_fast8_t` type, 313, 533
- `atomic_uint_least16_t` type, 313, 532
- `atomic_uint_least32_t` type, 313, 532
- `atomic_uint_least64_t` type, 313, 532
- `atomic_uint_least8_t` type, 313, 532
- `atomic_uintmax_t` type, 313, 533
- `atomic_uintptr_t` type, 313, 533
- `atomic_ullong` type, 313, 532
- `atomic_ulong` type, 313, 532
- `atomic_ushort` type, 313, 532
- `atomic_wchar_t` type, 313, 532
- `ATOMIC_WCHAR_T_LOCK_FREE` macro, 307, 532
- atomics header, 307, 489
- attribute, 149
  - `_Noreturn`, 150, 154, 155
  - `__deprecated__`, 151
  - `__nodiscard__`, 149
  - `deprecated`, 105, 114, 150, 151, 152, 153, 435, 436, 547, 720, 722
  - `fallthrough`, 150, 154, 180, 720, 722
  - for function type, 155
  - `maybe_unused`, 150, 152, 720, 722
  - `nodiscard`, 149, 150, 151, 152, 720, 722
  - `noreturn`, 132, 150, 154, 155, 297, 384–386, 427, 532, 539, 544, 547, 720, 722
  - `reproducible`, 150, 155, 157, 158, 638, 720, 722
  - `unsequenced`, 150, 155, 157, 158, 319–325, 533–535, 638, 720, 722
- attribute declaration, 104
- attribute prefixed token, 149
- attribute token, 149
- `auto`, 55, 106, 107, 136, 137, 142, 143, 171, 203, 723
- `auto` storage-class specifier, 55, 203
- `auto` type specifier, 203
- automatic storage duration, 21, 39
- backslash character (`\`), 11, 19, 68
- backslash escape sequence (`\`), 69, 202
- backspace, 21
- backspace escape sequence (`\b`), 21, 69
- basic character set, 4, 5, 19
- basic type, 41
- behavior, 3
- big-endian, 318
- binary digit, 61
- binary floating type, 600
- binary literal, 61
- binary prefix, 61
- binary resource inclusion, 183
- binary stream, 341, 369–371
- bit, 4
  - high-order, 4, 5
  - low-order, 5
- bit and byte utilities header, 489
- bit and byte utility, 318
- bit-field, 113
- bit-precise integer suffix, `wb` or `WB`, 62
- bit-precise integer type, 40
- bit-precise signed integer type, 40
- bit-precise unsigned integer type, 40
- `bitand` macro, 239, 512
- `BITINT_MAXWIDTH` macro, 23, 111, 512, 555, 722
- `bitor` macro, 239, 512
- bitwise operator, 76
  - AND, 93
  - AND assignment (`&=`), 98
  - complement (`~`), 87
  - exclusive OR, 94
  - exclusive OR assignment (`^=`), 98
  - inclusive OR, 94
  - inclusive OR assignment (`|=`), 98
  - shift, 91
- blank character, 218
- block, 160, 162
  - primary, 160
  - secondary, 160
- block scope, 36
- block structure, 36
- bold type** convention, 36
- `bool` keyword, 55
- `bool` type, 48, 110
- `bool` type conversion, 48
- `BOOL_MAX` macro, 512, 555
- `BOOL_WIDTH` macro, 23, 512, 555
- boolean type, 48
- boolean type and values header, 330, 489
- boolean type conversion, 47, 48
- bounded undefined behavior, 8
- braces punctuator (`{}`), 116, 123, 144, 161
- brackets operator (`[ ]`), 79, 86
- brackets punctuator (`[ ]`), 135, 144
- branch cut, 209
- `break` keyword, 55
- `break` statement, 168
- broken-down time, 432, 433–437, 700, 702
- `bsearch` macro, 387, 388, 489, 539, 634, 643, 723
- `bsearch_s` macro, 540, 688, 689
- `btowc` function, 452, 453, 477, 549
- `BUFSIZ` macro, 339, 342, 347, 537
- byte, 4, 87
- byte array, 4
- byte input/output function, 340
- byte type, 5
- byte-oriented stream, 341

- c** identifier prefix, 416, 417
- C** program, 11
- c16rtomb** function, 443, 444, 548
- c32rtomb** function, 445, 548
- c8rtomb** function, 442, 548, 722
- cabs** function, 214, 416, 417, 508, 592
- cabs** function
  - type-generic macro for, 417
- cabsf** function, 214, 419, 508
- cabsfN** function, 608
- cabsfNx** function, 608
- cabsl** function, 214, 508
- cacos** function, 210, 417, 507, 592, 593
- cacos** function
  - type-generic macro for, 417
- cacosf** function, 210, 507
- cacosfN** function, 608
- cacosfNx** function, 608
- cacosh** function, 211, 212, 417, 507, 593
- cacosh** function
  - type-generic macro for, 417
- cacoshf** function, 211, 507
- cacoshfN** function, 608
- cacoshfNx** function, 608
- cacoshl** function, 211, 507
- cacosl** function, 210, 507
- acospi** function, 488
- calendar time, 431, 432–434, 436, 437, 701, 702
- call by value, 80
- call\_once** function, 373, 421, 422, 538, 546, 723
- calloc** function, 381, 382, 383, 539, 634, 643, 650
- canonical representation, 25
- canonicalize** family, 34, 282
- canonicalize** function, 25, 282, 417, 516, 561, 584
- canonicalized128** function, 282, 522
- canonicalized32** function, 282, 522
- canonicalized64** function, 282, 522
- canonicalizedN** function, 619
- canonicalizedNx** function, 619
- canonicalizef** function, 282, 516
- canonicalizefN** function, 619
- canonicalizefNx** function, 619
- canonicalizel** function, 282, 516
- carg** function, 214, 215, 419, 508, 592, 628
- carg** type-generic macro, 417
- cargf** function, 214, 508
- cargfN** function, 608
- cargfNx** function, 608
- cargl** function, 214, 508
- carriage return, 21
- carriage-return escape sequence (`\r`), 21, 69, 219
- carries a dependency, 17
- case** keyword, 55
- case** label, 161, 164
- case mapping function
  - character, 219
  - extensible wide character, 486
  - wide character, 486
- casin** function, 210, 417, 507, 592
- casin** function
  - type-generic macro for, 417
- casinf** function, 210, 507
- casinfN** function, 608
- casinfNx** function, 608
- casinh** function, 212, 417, 507, 592, 594
- casinh** function
  - type-generic macro for, 417
- casinhf** function, 212, 507
- casinhfN** function, 608
- casinhfNx** function, 608
- casinhl** function, 212, 507
- casinl** function, 210, 507
- casinpi** function, 488
- cast, 89
- cast expression, 88
- cast operator (`()`), 89
- catan** function, 210, 211, 417, 507, 592
- catan** function
  - type-generic macro for, 417
- catanf** function, 210, 507
- catanfN** function, 608
- catanfNx** function, 608
- catanh** function, 212, 417, 507, 592, 594
- catanh** function
  - type-generic macro for, 417
- catanhf** function, 212, 507
- catanhfN** function, 608
- catanhfNx** function, 608
- catanhl** function, 212, 507
- catanl** function, 210, 507
- catanpi** function, 488
- cbrt** function, 78, 269, 418, 515, 577
- cbrt** type-generic macro, 417
- cbrtd128** function, 269, 520
- cbrtd32** function, 269, 520
- cbrtd64** function, 269, 520
- cbrtdN** function, 617
- cbrtdNx** function, 617
- cbrtf** function, 78, 269, 418, 515
- cbrtfN** function, 617
- cbrtfNx** function, 617
- cbrtl** function, 78, 269, 418, 515
- ccompoundn** function, 488
- ccos** function, 211, 417, 507, 592
- ccos** function
  - type-generic macro for, 417
- ccosf** function, 211, 507
- ccosf64x** function, 630



- ccosfN** function, 608
- ccosfNx** function, 608
- ccosh** function, 212, 417, 507, 592, 594, 595
- ccosh** function
  - type-generic macro for, 417
- ccoshf** function, 212, 507
- ccoshfN** function, 608
- ccoshfNx** function, 608
- ccoshl** function, 212, 507
- ccosl** function, 211, 507
- ccospi** function, 488
- ceil** function, 248, 274, 515, 559, 580, 581, 582
- ceil** type-generic macro, 417
- ceild128** function, 35, 274, 520
- ceild32** function, 35, 274, 520
- ceild64** function, 35, 274, 520
- ceildN** function, 618
- ceildNx** function, 618
- ceilf** function, 274, 515
- ceilfN** function, 618
- ceilfNx** function, 618
- ceiling, 7
- ceil** function, 274, 278, 515
- cerf** function, 488
- cerfc** function, 488
- cexp** function, 213, 417, 507, 596, 597
- cexp** function
  - type-generic macro for, 417
- cexp10** function, 488
- cexp10m1** function, 488
- cexp2** function, 488
- cexp2m1** function, 488
- cexpf** function, 213, 507
- cexpfN** function, 608
- cexpfNx** function, 608
- cexpl** function, 213, 507
- cexpm1** function, 488
- change history, 720
- char** keyword, 55
- char** type, 110
- char** type conversion, 48, 49
- char16\_t** type, 47, 69–72, 144, 200, 313, 393–397, 400, 401, 404, 441, 443, 540–544, 548
- char32\_t** type, 47, 69–72, 144, 200, 313, 393–397, 400, 401, 404, 441, 444, 445, 540–544, 548
- char8\_t** type, 47, 69–72, 108, 313, 393–397, 400, 401, 403, 404, 441, 442, 540–544, 548
- CHAR\_BIT** macro, 23, 43, 112, 183–185, 187, 327–329, 512, 535, 555
- CHAR\_MAX** macro, 24, 243, 244, 512, 555
- CHAR\_MIN** macro, 24, 41, 512, 555
- CHAR\_WIDTH** macro, 23, 512, 555
- character, 5
  - extended, 19
  - graphic, 20
- character array initialization, 145
- character case mapping function, 219
  - extensible wide character, 486
  - wide character, 486
- character classification function, 217
  - extensible wide character, 485
  - wide character, 482
- character display semantics, 21
- character handling header, 217, 242, 488
- character input/output function, 366, 685
  - wide character, 461
- character literal, 12, 19, 67
- character set, 19
- character type, 41, 145
- character type conversion, 47
- characteristics of floating types header, 235, 488
- characteristics of integer types header, 240
- checked arithmetic function, 489
- checked integer arithmetic header, 331
- cimag** function, 215, 216, 508, 590, 591, 592, 628
- cimag** type-generic macro, 417
- cimagf** function, 215, 508
- cimagfN** function, 608
- cimagfNx** function, 608
- cimagl** function, 215, 419, 508
- cis function, 592
- ckd\_** identifier prefix, 489
- ckd\_add** macro, 331, 536
- ckd\_mul** macro, 331, 536
- ckd\_sub** macro, 331, 536
- classification function
  - character, 217
  - extensible wide character, 485
  - floating-point, 251
  - wide character, 482
- clearerr** function, 371, 372, 537
- clgamma** function, 488
- clock** function, 431, 432, 435, 547, 651
- clock\_t** type, 431, 432, 547, 651
- CLOCKS\_PER\_SEC** macro, 431, 432, 435, 547
- clog** function, 213, 417, 507, 597
- clog** function
  - type-generic macro for, 417
- clog10** function, 488
- clog10p1** function, 488
- clog1p** function, 488
- clog2** function, 488
- clog2p1** function, 488
- clogf** function, 213, 419, 507
- clogfN** function, 608
- clogfNx** function, 608
- clogl** function, 213, 508

- clogp1** function, 488
- closed range, 102
- closing, 342
- CMPLX** macro, 109, 209, 215, 216, 508, 590, 591
- CMPLXF** macro, 215, 508, 509, 608
- CMPLXFN** macro, 608
- CMPLXFNx** macro, 608
- CMPLXL** macro, 215, 508
- cnd\_** identifier prefix, 490
- cnd\_broadcast** function, 422, 423, 424, 546
- cnd\_destroy** function, 423, 546
- cnd\_init** function, 423, 546
- cnd\_signal** function, 423, 424, 546
- cnd\_t** type, 421, 422–424, 546
- cnd\_timedwait** function, 423, 424, 546
- cnd\_wait** function, 423, 424, 546
- code point, 47, 553
- code unit, 47
- coefficient, 32
- collating sequence, 19
- colon punctuator (:), 112
- comma operator (,), 17, 99
- comma punctuator (,), 79, 104, 112, 116, 123, 144
- command processor, 386
- comment, 11, 54, 74
- comment delimiters (/ \* \*/ and //), 74
- common definitions header, 332
- common extension, 652
- common initial sequence, 81
- common real type, 49
- common warning, 12, 631
- comparison
  - pointer, 92
- comparison function, 387, 388, 688, 689, 690
  - string, 409
  - wide string, 471
- comparison macro, 293
- compatible type, 44, 111, 128, 134
- compl** macro, 239, 512
- complement operator (~), 87
- complete, 39
- complete type, 39
- complex arithmetic header, 209, 488
- complex** macro, 146, 209–216, 419, 507–509, 590, 591, 607, 608, 630, 634
- complex number, 41, 589
- complex type, 41, 110, 201, 589
- complex type conversion, 49
- complex type domain, 41
- components of time, 431, 700
- composite type, 45
- compound assignment, 98
- compound literal, 83
- compound literal constant, 100
- compound statement, 162
- compound-literal operator (( {} ), 83
- compoundn** function, 269, 515, 564, 578
- compoundn** type-generic macro, 417
- compoundnd128** function, 269, 520
- compoundnd32** function, 269, 520
- compoundnd64** function, 269, 520
- compoundndN** function, 617
- compoundndNx** function, 617
- compoundndnf** function, 269, 515
- compoundndnfN** function, 617
- compoundndnfNx** function, 617
- compoundnrl** function, 269, 515
- concatenation function
  - string, 409, 695
  - wide string, 470, 712
- conceptual model, 11
- conditional expression inclusion preprocess-  
ing directive, 23, 25, 178, 207
- conditional feature, 9, 40, 41, 42, 135, 201, 204, 558, 589, 671, 719
- conditional inclusion, 178
- conditional inclusion preprocessing directive, 178
- conditional operator (? :), 17, 95
- conflict, 16
- conformance, 9
- conforming freestanding implementation, 9
- conforming hosted implementation, 9
- conforming implementation, 9
- conforming program, 9
- conj** function, 215, 508, 592, 593–597, 628
- conj** type-generic macro, 417
- conjf** function, 215, 508
- conjfN** function, 608
- conjfNx** function, 608
- conjl** function, 215, 508
- const**, 42
- const** keyword, 55
- const** type qualifier, 127
- const-qualified type, 42, 51, 128
- constant, 60
  - enumeration, 36, 67
- constant expression, 100, 567
- constant range expression, 101, 102, 161, 164
- constants
  - as primary expression, 77
- constexpr**, 37–39, 55, 84, 100, 101, 102, 105–110, 144, 203, 299, 404, 567, 568, 722
- constexpr** storage-class specifier, 55
- constraint, 5, 9
- constraint\_handler\_t** type, 540, 686
- consume operation, 17
- content of structure/union/enumeration, 123
- contiguity of allocated storage, 381
- continue**, 55, 166, 167, 168, 404
- continue** keyword, 55

- continue** statement, 167
- contracted, 77
- contracted expression, 77, 251, 566
- control character, 20, 217
- control wide character, 482
- controlling type, 78
- conversion, 47
  - arithmetic operand, 47
  - array, 51, 52
  - array argument, 172
  - array parameter, 172
  - boolean, 48
  - boolean, characters, and integer, 47
  - by assignment, 97
  - by **return** statement, 169
  - complex type, 49
  - explicit, 47
  - function, 52
  - function argument, 80, 172
  - function designator, 51
  - function parameter, 172
  - implicit, 47
  - lvalue, 51
  - nullptr\_t**, 53
  - pointer, 52
  - real and complex, 49
  - real floating and integer, 49, 565
  - real floating type, 49
  - signed and unsigned integer, 48
  - usual arithmetic, *see* usual arithmetic conversion
  - void** type, 52
- conversion function
  - multibyte/wide character, 389, 690
    - extended, 476, 715
    - restartable, 441, 478, 715
  - multibyte/wide string, 390, 691
    - restartable, 479, 716
  - numeric, 237, 373
    - wide string, 238, 464
  - single byte/wide character, 477
  - time, 435
    - wide character, 476
- conversion functions
  - time, 700
- conversion specifier, 349, 356, 448, 454
  - %A, 351, 437, 451
  - %B, 351, 437, 450
  - %C, 437
  - %D, 438
  - %E, 351, 450
  - %F, 351, 438, 450
  - %G, 351, 438, 450
  - %H, 438
  - %I, 438
  - %M, 438
  - %R, 438
  - %S, 438
  - %T, 438
  - %U, 438
  - %V, 438
  - %W, 438
  - %X, 351, 438, 450
  - %Y, 438
  - %Z, 438
  - %, 359, 456
  - %%, 353, 360, 452, 457
  - %a, 351, 358, 437, 451, 456
  - %b, 351, 437, 450
  - %c, 352, 358, 437, 452, 456
  - %d, 351, 358, 438, 450, 455
  - %e, 351, 358, 438, 450, 456
  - %f, 351, 358, 450, 456
  - %g, 351, 358, 438, 450, 456
  - %h, 438
  - %i, 351, 450
  - %j, 438
  - %m, 438
  - %n, 353, 359, 438, 452, 457
  - %o, 351, 358, 450, 455
  - %p, 353, 359, 438, 452, 457
  - %r, 438
  - %s, 353, 359, 452, 456
  - %t, 438
  - %u, 351, 358, 438, 450, 456
  - %w, 438
  - %x, 351, 358, 438, 450, 456
  - %y, 438
  - %z, 438
- conversion state, 389, 441–445, 477, 478–480, 690, 715–718
- conversion state function, 477
- copying function
  - string, 407, 693
  - wide string, 469
- copying functions
  - wide string, 710
- copysign** function, 216, 279, 280, 516, 561, 577, 582, 583, 590, 591
- copysign** type-generic macro, 417
- copysign128** function, 280, 521
- copysign32** function, 280, 521
- copysign64** function, 280, 521
- copysignN** function, 619
- copysignNx** function, 619
- copysignf** function, 280, 516
- copysignfN** function, 619
- copysignfNx** function, 619
- copysignl** function, 280, 419, 516
- correctly rounded result, 5
- corresponding real type, 41
- corresponding unsigned integer type, 40

- cos** function, 143, 256, 417, 513, 565, 573
- cos** type-generic macro, 417
- cosd128** function, 256, 518
- cosd32** function, 256, 518
- cosd64** function, 256, 518
- cosdN** function, 613
- cosdNx** function, 613
- cosf** function, 256, 513
- cosfN** function, 613
- cosfNx** function, 613
- cosh** function, 260, 417, 513, 565, 574
- cosh** type-generic macro, 417
- coshd128** function, 260, 519
- coshd32** function, 260, 519
- coshd64** function, 260, 519
- coshdN** function, 614
- coshdNx** function, 614
- coshf** function, 260, 513
- coshfN** function, 614
- coshfNx** function, 614
- coshl** function, 260, 513
- cosl** function, 256, 513
- cosp** function, 258, 513, 565, 574
- cosp** type-generic macro, 417
- cospid128** function, 258, 518
- cospid32** function, 258, 518
- cospid64** function, 258, 518
- cospidN** function, 613
- cospidNx** function, 613
- cospif** function, 258, 513
- cospifN** function, 613
- cospifNx** function, 613
- cospil** function, 258, 513
- cpow** function, 214, 417, 508, 597
- cpow** function
  - type-generic macro for, 417
- cpowf** function, 214, 508
- cpowf128** function, 630
- cpowfN** function, 608
- cpowfNx** function, 608
- cpowl** function, 214, 419, 508
- cpown** function, 488
- cpowr** function, 488
- cproj** function, 215, 216, 508, 592, 628
- cproj** type-generic macro, 417
- cprojf** function, 216, 419, 508
- cprojfN** function, 608
- cprojfNx** function, 608
- cprojl** function, 216, 419, 508
- cr\_** identifier prefix, 489, 564, 572
- CR\_DECIMAL\_DIG** macro, 25, 511, 565, 566
- creal** function, 215, 216, 419, 508, 590, 591, 592, 628
- creal** type-generic macro, 417
- crealf** function, 216, 508
- crealfN** function, 608
- crealfNx** function, 608
- creall** function, 216, 508
- creating, 342
- critical undefined behavior, 8
- crootn** function, 488
- crsqrt** function, 488
- csin** function, 211, 417, 507, 592
- csin** function
  - type-generic macro for, 417
- csinf** function, 211, 507
- csinfN** function, 608
- csinfNx** function, 608
- csinh** function, 212, 213, 417, 507, 592, 595
- csinh** function
  - type-generic macro for, 417
- csinhf** function, 213, 507
- csinhfN** function, 608
- csinhfNx** function, 608
- csinhl** function, 213, 507
- csinl** function, 211, 507
- csinpi** function, 488
- csqrt** function, 214, 417, 419, 508, 592, 597
- csqrt** function
  - type-generic macro for, 417
- csqrtf** function, 214, 508
- csqrtfN** function, 608
- csqrtfNx** function, 608
- csqrtl** function, 214, 508
- ctan** function, 211, 417, 507, 592
- ctan** function
  - type-generic macro for, 417
- ctanf** function, 211, 507
- ctanfN** function, 608
- ctanfNx** function, 608
- ctanh** function, 213, 417, 507, 592, 596
- ctanh** function
  - type-generic macro for, 417
- ctanhf** function, 213, 507
- ctanhfN** function, 608
- ctanhfNx** function, 608
- ctanhl** function, 213, 507
- ctanl** function, 211, 507
- ctanpi** function, 488
- ctgamma** function, 488
- ctime** function, 435, 436, 547
- ctime\_s** function, 547, 700, 701
- currency\_symbol** structure member, 241, 243, 246
- current object, 145
- CX\_LIMITED\_RANGE** pragma, vii, 198, 199, 209, 210, 507, 589, 639
- D format modifier, 350, 358, 450, 455
- d** identifier prefix, 418
- d-wchar sequence, 377, 467
- d128** identifier prefix, 418
- d32** identifier prefix, 418

- d32add** macro, 33, 228, 418
- d32add** type-generic macro, 418
- d32addd128** function, 287, 522
- d32addd64** function, 287, 522
- d32div** macro, 33, 228, 418
- d32div** type-generic macro, 418
- d32divd128** function, 288, 522
- d32divd64** function, 288, 419, 522
- d32fma** macro, 34, 228, 418
- d32fma** type-generic macro, 418
- d32fmad128** function, 289, 522
- d32fmad64** function, 289, 522
- d32mul** macro, 33, 228, 418
- d32mul** type-generic macro, 418
- d32muld128** function, 288, 522
- d32muld64** function, 288, 522
- d32sqrt** macro, 33, 228, 418
- d32sqrt** type-generic macro, 418
- d32sqrtd128** function, 289, 522
- d32sqrtd64** function, 289, 522
- d32sub** macro, 33, 228, 418
- d32sub** type-generic macro, 418
- d32subd128** function, 288, 419, 522
- d32subd64** function, 288, 522
- d64** identifier prefix, 416, 418
- d64add** macro, 33, 228, 418
- d64add** type-generic macro, 418
- d64addd128** function, 287, 420, 522
- d64div** macro, 33, 228, 418
- d64div** type-generic macro, 418
- d64divd128** function, 288, 522
- d64fma** macro, 34, 228, 418
- d64fma** type-generic macro, 418
- d64fmad128** function, 289, 419, 522
- d64mul** macro, 33, 228, 418
- d64mul** type-generic macro, 418
- d64muld128** function, 288, 522
- d64sqrt** macro, 33, 228, 418
- d64sqrt** type-generic macro, 418
- d64sqrtd128** function, 289, 522
- d64sub** macro, 33, 228, 418
- d64sub** type-generic macro, 418
- d64subd128** function, 288, 522
- dadd** macro, 418, 629
- dadd** type-generic macro, 418
- daddl** function, 287, 419, 517, 560
- data race, 19, 207, 380, 381, 386, 414, 415, 435, 441, 478, 480, 688
- date and time header, 421, 431, 490, 699
- Daylight Saving Time, 431
- DBL\_** identifier prefix, 488
- DBL\_DECIMAL\_DIG** macro, 27, 30, 510, 556
- DBL\_DIG** macro, 28, 30, 510, 556
- DBL\_EPSILON** macro, 29, 30, 510, 556
- DBL\_HAS\_SUBNORM** macro, 26, 30, 488, 510
- DBL\_IS\_IEC\_60559** macro, 26, 30, 510
- DBL\_MANT\_DIG** macro, 27, 30, 109, 510, 556
- DBL\_MAX** macro, 28, 30, 510, 556
- DBL\_MAX\_10\_EXP** macro, 28, 30, 510, 556
- DBL\_MAX\_EXP** macro, 28, 30, 510, 556
- DBL\_MIN** macro, 29, 30, 510, 556
- DBL\_MIN\_10\_EXP** macro, 28, 30, 510, 556
- DBL\_MIN\_EXP** macro, 28, 30, 510, 556
- DBL\_NORM\_MAX** macro, 29, 510, 556
- DBL\_SNAN** macro, 27, 109, 164, 510, 559
- DBL\_TRUE\_MIN** macro, 29, 30, 510
- DD format modifier, 350, 358, 450, 455
- ddiv** macro, 418, 629, 630
- ddiv** type-generic macro, 418
- ddivl** function, 288, 517, 560, 564, 630
- DEC** identifier prefix, 511
- DEC128\_** identifier prefix, 30, 488
- DEC128\_EPSILON** macro, 32, 557
- DEC128\_MANT\_DIG** macro, 31, 557
- DEC128\_MAX** macro, 32, 557
- DEC128\_MAX\_EXP** macro, 32, 557
- DEC128\_MIN** macro, 32, 557
- DEC128\_MIN\_EXP** macro, 31, 557
- DEC128\_SNAN** macro, 31
- DEC128\_TRUE\_MIN** macro, 32, 557
- DEC32\_** identifier prefix, 30, 488
- DEC32\_EPSILON** macro, 32, 557
- DEC32\_MANT\_DIG** macro, 31, 557
- DEC32\_MAX** macro, 32, 557
- DEC32\_MAX\_EXP** macro, 32, 557
- DEC32\_MIN** macro, 32, 557
- DEC32\_MIN\_EXP** macro, 31, 557
- DEC32\_SNAN** macro, 31, 109
- DEC32\_TRUE\_MIN** macro, 32, 557
- DEC64\_** identifier prefix, 30, 488
- DEC64\_EPSILON** macro, 32, 557
- DEC64\_MANT\_DIG** macro, 31, 557
- DEC64\_MAX** macro, 32, 557
- DEC64\_MAX\_EXP** macro, 32, 557
- DEC64\_MIN** macro, 32, 557
- DEC64\_MIN\_EXP** macro, 31, 557
- DEC64\_SNAN** macro, 31, 109
- DEC64\_TRUE\_MIN** macro, 32, 108, 109, 557
- DEC\_** identifier prefix, 488
- DEC\_EVAL\_METHOD** macro, 25, 30, 66, 100, 247, 511, 556, 588, 601, 610, 646, 649
- DEC\_INFINITY** macro, 31, 109, 248, 290, 489, 511, 518
- DEC\_NAN** macro, 31, 109, 248, 489, 511, 518
- decimal digit, 20
- decimal floating type, 40, 600
- decimal literal, 60
- decimal re-encoding function, 291
- decimal rounding control pragma, 227
- decimal-point character, 204, 243
- decimal128 suffix, **dl** or **DL**, 65
- decimal32 suffix, **df** or **DF**, 65



- decimal64 suffix, **dd** or **DD**, 65
- DECIMAL\_DIG** macro, 25, 28, 488, 510, 556, 721
- decimal\_point** structure member, 241, 243
- declaration, 104
  - \_Static\_assert**, 55
  - function, 137
  - pointer, 134
  - static\_assert**, 55, 149
  - structure/union, 112
  - typedef**, 140
- declaration specifier, 104
- declarator, 133
  - abstract, 139
- declarator type derivation, 42, 134
- DECN\_EPSILON** macro, 604
- DECN\_MANT\_DIG** macro, 603
- DECN\_MAX** macro, 603
- DECN\_MAX\_EXP** macro, 603
- DECN\_MIN** macro, 604
- DECN\_MIN\_EXP** macro, 603
- DECN\_SNAN** macro, 602
- DECN\_TRUE\_MIN** macro, 604
- DECNX\_EPSILON** macro, 604
- DECNX\_MANT\_DIG** macro, 603
- DECNX\_MAX** macro, 603
- DECNX\_MAX\_EXP** macro, 603
- DECNX\_MIN** macro, 604
- DECNX\_MIN\_EXP** macro, 603
- DECNX\_SNAN** macro, 602
- DECNX\_TRUE\_MIN** macro, 604
- decodebin** family, 34, 292
- decodebind** family, 292, 293
- decodebind128** function, 292, 523
- decodebind32** function, 292, 523
- decodebind64** function, 292, 416, 418, 523
- decodebindN** function, 621
- decodedec** family, 34, 292
- decodedecd** family, 291, 292
- decodedecd128** function, 291, 523
- decodedecd32** function, 291, 523
- decodedecd64** function, 291, 416, 418, 523
- decodedecdN** function, 621
- decodedefN** function, 623
- default argument promotion, 80
- default initialization, 145
- default** keyword, 55
- default** label, 161, 164
- define** preprocessing directive, 191
- defined**, 179, 180, 186, 199, 638
- defined** operator, 199
- definition, 105
  - escaping, 156
  - function, 171
- dependency-ordered before, 17
- deprecated** attribute, 105, 114, 150, 151, 152, 153, 435, 436, 547, 720, 722
- derived declarator type, 42
- derived type, 41
- designated initializer, 145
- destringizing, 202
- device input/output, 14
- dfma** macro, 418, 629
- dfma** type-generic macro, 418
- dfmal** function, 289, 419, 517, 560
- diagnostic, 12
- diagnostic message, 5, 12
- diagnostics header, 208
- difftime** function, 432, 433, 547
- digit, 20, 217
- digraph, 73
- direct input/output function, 369
- display device, 21
- div** function, 143, 373, 388, 389, 539
- div\_t** type, 148, 373, 388, 389, 538, 539
- divide and round to narrower type, 288
- division assignment operator (/=), 98
- division operator (/), 89, 589
- dMadddN** function, 621
- dMadddNx** function, 621
- dMdivdN** function, 621
- dMdivdNx** function, 621
- dMencbindN** function, 624
- dMencdecN** function, 624
- dMfmadN** function, 621
- dMfmadNx** function, 621
- dMmuldN** function, 621
- dMmuldNx** function, 621
- dMsqrtdN** function, 621
- dMsqrtdNx** function, 621
- dMsubdN** function, 621
- dMsubdNx** function, 621
- dmul** macro, 418, 629
- dmul** type-generic macro, 418
- dnull** function, 288, 517, 560
- dMxadddN** function, 621
- dMxadddNx** function, 621
- dMxdivdN** function, 621
- dMxdivdNx** function, 621
- dMxfmadN** function, 621
- dMxfmadNx** function, 621
- dMxmuldN** function, 621
- dMxmuldNx** function, 621
- dMxsqrtdN** function, 621
- dMxsqrtdNx** function, 621
- dMxsubdN** function, 621
- dMxsubdNx** function, 621
- do** keyword, 55
- do** statement, 166
- documentation of implementation, 10
- domain error, 250, 255–260, 264–273, 275, 276, 279
- dot operator (.), 81

- double** **\_Complex** type, 41
- double** **\_Complex** type conversion, 49
- double** arithmetic, 15
- double** keyword, 55
- double** type, 40, 110
- double** type conversion, 49
- double-quote escape sequence (`\`"), 69, 71, 202
- double\_t** type, 247, 512, 568, 609, 610, 611, 649, 652
- dsqrt** macro, 418, 629
- dsqrt** type-generic macro, 418
- dsqrtl** function, 289, 517, 560
- dsub** macro, 418, 629, 630
- dsub** type-generic macro, 418
- dsubl** function, 288, 517, 560, 630
- dynamic floating-point environment, 222
- E format modifier, 439
- E** identifier prefix, 221, 488
- EDOM** macro, 221, 250, 509
- effective type, 76
- effectless function call, 156
- EILSEQ** macro, 221, 343, 442–445, 461, 462, 479–481, 509
- element type, 41
- elif** preprocessing directive, 179
- elifdef** preprocessing directive, 180
- elifndef** preprocessing directive, 180
- ellipsis punctuator (`...`), 137, 191
- else** keyword, 55
- else** preprocessing directive, 180
- else** statement, 163
- embed element width, 183
- embed parameter, 183
  - \_\_limit\_\_**, 177
  - if\_empty**, 184, 189, 190
  - limit**, 74, 177, 181, 183, 184, 186, 187–190
  - prefix**, 184, 188, 189
  - suffix**, 184, 188, 189, 190
- embed parameter sequence, 183
- empty initialization, 144, 145
- empty initializer, 144
- empty resource, 183
- empty statement, 162
- encodebin** family, 34, 292
- encodebind** family, 292
- encodebind128** function, 292, 523
- encodebind32** function, 292, 523
- encodebind64** function, 292, 416, 418, 523
- encodebindN** function, 621
- encodedec** family, 34, 291
- encodedecd** family, 291
- encodedecd128** function, 291, 523
- encodedecd32** function, 291, 523
- encodedecd64** function, 291, 416, 418, 523
- encodedecdN** function, 621
- encoddefN** function, 622
- encoding, 46
- encoding error, 47, 354, 356, 357, 362–366, 442–445, 447, 453, 454, 458–462, 479–481, 691, 692, 716–718
- end-of-file, 447
- end-of-file indicator, 339, 347, 366–372, 461, 464
- end-of-line indicator, 20
- endif** preprocessing directive, 180
- enum** keyword, 55
- enum** type, 41, 110, 116
- enumerated type, 41
- enumeration, 41, 116
- enumeration constant, 36, 67
- enumeration content, 123
- enumeration member, 116
- enumeration member type, 117
- enumeration specifier, 116
- enumeration tag, 38, 123
- enumerator, 116
- environment, 11
- environment function, 384, 687
- environment list, 386, 688
- environmental consideration, 19
- environmental limit, 21, 296, 342, 343, 345, 354, 380, 384, 385, 453, 675
- EOF** macro, 217, 339, 345, 360, 362–369, 447, 457–462, 477, 537, 639, 678, 679, 681–683, 685, 703, 705, 706, 708–710
- epoch, 434
- equal-sign punctuator (`=`), 104, 116, 144
- equality expression, 92
- equality operator (`==`), 93
- ERANGE** macro, 221, 238, 250, 251, 377, 378, 380, 466, 468, 469, 509, 649, 650
- erf** function, 272, 515, 580
- erf** type-generic macro, 417
- erfc** function, 273, 515, 580
- erfc** type-generic macro, 417
- erfcd128** function, 273, 520
- erfcd32** function, 273, 520
- erfcd64** function, 273, 520
- erfcdN** function, 617
- erfcdNx** function, 617
- erfcf** function, 273, 515
- erfcfN** function, 617
- erfcfNx** function, 617
- erfcl** function, 273, 515
- erfd128** function, 272, 520
- erfd32** function, 272, 520
- erfd64** function, 272, 520
- erfdN** function, 617
- erfdNx** function, 617
- erff** function, 272, 515
- erffN** function, 617
- erffNx** function, 617

**erfl** function, 272, 515  
**errno** identifier, 47, 156, 157, 206, 209, 221, 238, 250, 251, 299, 343, 370–373, 377, 378, 380, 415, 442–445, 461, 462, 466, 468, 469, 479–481, 509, 639, 640, 649, 650, 654, 672, 673, 698  
**errno\_t** type, 509, 538, 540, 545, 547, 550, 672, 673, 674–676, 686, 687, 689–692, 693, 694–696, 698, 699, 700, 701, 702, 710–713, 715–717  
error condition, 250  
error function, 272, 580  
error indicator, 339, 347, 366–368, 370, 372, 461, 462  
**error** preprocessing directive, 9, 198  
error-handling function, 371, 414, 698, 699  
errors header, 221, 488  
escape character (\), 68  
escape sequence, 19, 21, 68, 203  
escaping definition, 156  
evaluation format, 26, 66, 247  
evaluation method, 26, 77, 568  
evaluation of expression, 14  
exceptional condition, 76  
excess precision, 26, 50, 169  
excess range, 26, 50, 169  
exclusive OR operator  
    bitwise (^), 94  
    bitwise assignment (^=), 98  
executable program, 11  
execution character set, 19  
execution encodings, 47  
execution environment, 11, 12  
execution sequence, 14, 160  
**EXIT\_FAILURE** macro, 373, 385, 538  
**EXIT\_SUCCESS** macro, 373, 385, 427, 538  
**exp** function, 261, 262, 263, 315, 417, 419, 514, 564, 575  
**exp** type-generic macro, 417  
**exp10** function, 262, 514, 564, 575  
**exp10** type-generic macro, 417  
**exp10d128** function, 262, 519  
**exp10d32** function, 262, 519  
**exp10d64** function, 262, 519  
**exp10dN** function, 616  
**exp10dNx** function, 616  
**exp10f** function, 262, 514  
**exp10fN** function, 616  
**exp10fNx** function, 616  
**exp10l** function, 262, 514  
**exp10m1** function, 262, 514, 564, 575  
**exp10m1** type-generic macro, 417  
**exp10m1d128** function, 262, 519  
**exp10m1d32** function, 262, 519  
**exp10m1d64** function, 262, 519  
**exp10m1dN** function, 616  
**exp10m1dNx** function, 616  
**exp10mf** function, 262, 514  
**exp10mfN** function, 616  
**exp10mfNx** function, 616  
**exp10ml** function, 262, 514  
**exp10ml** type-generic macro, 417  
**exp10ml128** function, 263, 519  
**exp10ml32** function, 263, 519  
**exp10ml64** function, 263, 519  
**exp10mlN** function, 616  
**exp10mlNx** function, 616  
**exp10mlf** function, 263, 514  
**exp10mlfN** function, 616  
**exp10mlfNx** function, 616  
**exp10ml1** function, 263, 514  
**exp10ml1** type-generic macro, 417  
**exp10ml1d128** function, 263, 519  
**exp10ml1d32** function, 263, 519  
**exp10ml1d64** function, 263, 519  
**exp10ml1dN** function, 616  
**exp10ml1dNx** function, 616  
**exp10ml1f** function, 263, 514  
**exp10ml1fN** function, 616  
**exp10ml1fNx** function, 616  
**exp10ml1l** function, 263, 514  
**exp10ml1l** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function, 263, 514  
**exp10ml1lfN** function, 616  
**exp10ml1lfNx** function, 616  
**exp10ml1ll** function, 263, 514  
**exp10ml1ll** type-generic macro, 417  
**exp10ml1ld128** function, 263, 519  
**exp10ml1ld32** function, 263, 519  
**exp10ml1ld64** function, 263, 519  
**exp10ml1ldN** function, 616  
**exp10ml1ldNx** function, 616  
**exp10ml1lf** function,



- constant range, 101
- evaluation, 14
- full, 160
- parenthesized, 77
- primary, 77
- unary, 85
- void, 52
- expression statement, 162
- extended alignment, 46
- extended character, 19
- extended character set, 5, 19, 20
- extended constant expressions, 101
- extended floating type, 600
- extended integer type, 40, 48, 62, 335
- extended multibyte and wide character utilities header, 446, 490
- extended multibyte/wide character conversion utility, 476, 715
- extended signed integer types, 40
- extended unsigned integer type, 40
- extensible wide character case mapping function, 486
- extensible wide character classification function, 485
- extern**, 37, 38, 55, 87, 102, 103, 106, 107, 121, 127–132, 136, 137, 141, 158, 165, 172–174, 188, 189, 207, 652, 653
- extern** storage-class specifier, 37, 55
- external definition, 171
- external identifiers
  - underscore, 206
- external linkage, 37
- external name, 57
- external object definition, 173
- f** identifier suffix, 209, 247, 416, 418, 488
- f32add** macro, 630
- f32addf64** function, 630
- f32addf64x** function, 630
- f32fma** macro, 630
- f32fmaf32x** function, 630
- f32mul** macro, 630
- f32mulf128** function, 630
- f32mulf32x** function, 630
- f32xsqrt** macro, 630
- f32xsqrtf64x** function, 630
- f64div** macro, 630
- f64divf128** function, 630
- f64divf64x** function, 630
- fabs** function, 269, 270, 416, 417, 515, 561, 578, 582, 583, 591
- fabs** type-generic macro, 417
- fabsd128** function, 269, 520
- fabsd32** function, 269, 520
- fabsd64** function, 269, 520
- fabsdN** function, 617
- fabsdNx** function, 617
- fabsf** function, 269, 515
- fabsfN** function, 617
- fabsfNx** function, 617
- fabsl** function, 269, 515
- fadd** function, 287, 517, 560, 629
- fadd** type-generic macro, 418
- faddl** function, 287, 517, 560
- fallthrough** attribute, 150, 154, 180, 720, 722
- fallthrough declaration, 154
- false** keyword, 55
- false** predefined constant, 70
- family
  - canonicalize**, 34, 282
  - decodebin**, 34, 292
  - decodedec**, 34, 292
  - encodebin**, 34, 292
  - encodedec**, 34, 291
  - modf**, 34, 268, 416
  - strto**, 34, 377
  - wcsto**, 34, 467
- fclose** function, 186, 345, 537
- fdim** function, 282, 283, 516, 584
- fdim** type-generic macro, 417
- fdimd128** function, 283, 522
- fdimd32** function, 283, 522
- fdimd64** function, 283, 522
- fdimdN** function, 620
- fdimdNx** function, 620
- fdimf** function, 283, 516
- fdimfN** function, 620
- fdimfNx** function, 620
- fdiml** function, 283, 517
- fdiv** function, 288, 419, 517, 560, 629
- fdiv** type-generic macro, 418
- fdivl** function, 288, 517, 560
- FE\_** identifier prefix, 223, 224, 488
- FE\_ALL\_EXCEPT** macro, 99, 223, 509
- FE\_DEC\_DOWNWARD** macro, 223, 227, 510, 564
- FE\_DEC\_DOWNWARD** pragma, 199
- FE\_DEC\_DYNAMIC** pragma, 199, 227
- fe\_dec\_getround** function, 224, 231, 510, 564, 609
- fe\_dec\_setround** function, 224, 227, 232, 510, 564, 609
- FE\_DEC\_TONEAREST** macro, 223, 224, 227, 510, 564
- FE\_DEC\_TONEAREST** pragma, 199
- FE\_DEC\_TONEARESTFROMZERO** macro, 223, 227, 510, 564
- FE\_DEC\_TONEARESTFROMZERO** pragma, 199
- FE\_DEC\_TOWARDZERO** macro, 223, 227, 510, 564
- FE\_DEC\_TOWARDZERO** pragma, 199
- FE\_DEC\_UPWARD** macro, 223, 227, 510, 564
- FE\_DEC\_UPWARD** pragma, 199
- FE\_DFL\_ENV** macro, 224, 509
- FE\_DFL\_MODE** macro, 223, 231, 509, 563

- FE\_DIVBYZERO** macro, 223, 250, 509
- FE\_DOWNWARD** macro, 223, 509, 563
- FE\_DOWNWARD** pragma, 199
- FE\_DYNAMIC** pragma, 199, 225, 509
- FE\_INEXACT** macro, 223, 228, 509, 582
- FE\_INVALID** macro, 223, 230, 250, 509
- FE\_OVERFLOW** macro, 223, 228, 230, 250, 509
- FE\_SNANS\_ALWAYS\_SIGNAL** macro, 510, 559, 563, 581, 584
- FE\_TONEAREST** macro, 158, 223, 509, 563
- FE\_TONEAREST** pragma, 199
- FE\_TONEARESTFROMZERO** macro, 223, 509, 563
- FE\_TONEARESTFROMZERO** pragma, 199
- FE\_TOWARDZERO** macro, 223, 509, 563, 577, 582
- FE\_TOWARDZERO** pragma, 199
- FE\_UNDERFLOW** macro, 223, 234, 509
- FE\_UPWARD** macro, 9, 223, 509, 563, 581
- FE\_UPWARD** pragma, 199
- feature test macro, 205, 222, 247, 335, 373, 416, 431
- feclearexcept** function, 99, 228, 230, 234, 509, 562, 582
- fegetenv** function, 233, 234, 510, 563, 639
- fegetexceptflag** function, 228, 229, 230, 510, 562, 639, 649
- fegetmode** function, 230, 231, 510, 562
- fegetround** function, 223, 227, 230, 231, 232, 510, 562, 563, 577, 581, 609
- fehohldexcept** function, 98, 233, 234, 510, 563, 582, 639
- femode\_t** type, 222, 223, 231, 509, 510
- fence, 17, 311
  - acquire, 311
  - release, 311
- FENV\_ACCESS** pragma, vii, 9, 98, 198, 199, 224, 225, 230, 232, 234, 509, 566–571, 577, 581, 583, 633, 639, 647
- FENV\_DEC\_ROUND** pragma, vii, 66, 198, 199, 227, 510, 564, 609
- FENV\_ROUND** pragma, vii, 30, 66, 158, 198, 199, 224, 225, 226, 227, 509, 563, 609
- fenv\_t** type, 98, 222, 224, 233, 234, 509, 510, 581
- feof** function, 360, 366, 372, 461, 537
- feraiseexcept** function, 228, 229, 510, 568, 633, 649
- ferror** function, 360, 366, 372, 461, 537
- fesetenv** function, 225, 233, 510, 563, 639
- fesetexcept** function, 229, 510, 562
- fesetexceptflag** function, 228, 229, 510, 562, 639
- fesetmode** function, 225, 230, 231, 510, 563
- fesetround** function, 9, 25, 107, 223, 225, 227, 230, 232, 510, 562, 563, 577, 581, 582, 609
- fetestexcept** function, 228, 230, 510, 562, 582
- fetestexceptflag** function, 230, 510, 562
- feupdateenv** function, 99, 225, 233, 234, 510, 563, 582, 639
- fexcept\_t** type, 222, 228–230, 509, 510, 639
- fflush** function, 345, 346, 537, 641
- ffma** function, 289, 517, 560, 629
- ffma** type-generic macro, 418
- ffmal** function, 289, 517, 560
- fgetc** function, 340, 343, 366, 367, 369, 537
- fgetpos** function, 341, 342, 370, 371, 537, 633, 643, 650
- fgets** function, 340, 366, 367, 537, 642, 685
- fgetwc** function, 340, 343, 461, 463, 548
- fgetws** function, 340, 461, 462, 548, 642
- field width, 348, 448
- file, 342
  - access function, 345, 675
  - name, 342
  - operation, 343, 674
  - position indicator, 339, 341, 342, 346, 366, 367, 369–371, 461, 462, 464
  - positioning function, 370
- file name, 342
- file position indicator, 342
- file scope, 36, 171
- FILE** type, 76, 186, 339, 340, 342, 344, 345, 347, 348, 356, 364, 366–372, 447, 453, 458, 459, 461–464, 537, 538, 548, 550, 641, 674–677, 681, 682, 702, 703, 705, 706
- FILENAME\_MAX** macro, 339, 537
- finite number, 589
- fixed underlying type, 116
- flag, 348, 448
- flexible array member, 114
- float \_Complex** type, 41
- float \_Complex** type conversion, 49
- float** arithmetic, 15
- float** keyword, 55
- float** type, 40, 110
- float** type conversion, 49
- float\_t** type, 247, 512, 568, 609, 610, 611, 649, 652
- floating literal, 63
- floating suffix, **f** or **F**, 65
- floating type, 41, 203
- floating type conversion, 49, 565
- floating-point accuracy, 25, 66, 77, 376, 565
- floating-point arithmetic function, 247, 571
- floating-point classification function, 251
- floating-point control mode, 222, 568
- floating-point environment, 222, 566, 568
  - dynamic, 222
- floating-point environment header, 222, 488
- floating-point exception, 222, 228, 571
- floating-point number, 24, 40
- floating-point rounding mode, 25

- floating-point status flag, 222, 568
- floor, 8
- floor** function, 248, 274, 515, 559, 581
- floor** type-generic macro, 417
- floord128** function, 274, 520
- floord32** function, 274, 520
- floord64** function, 274, 520
- floordN** function, 618
- floordNx** function, 618
- floorf** function, 274, 515
- floorfN** function, 618
- floorfNx** function, 618
- floorl** function, 274, 515
- FLT** identifier prefix, 511
- FLT\_** identifier prefix, 488
- FLT\_DECIMAL\_DIG** macro, 27, 29, 30, 510, 556
- FLT\_DIG** macro, 28, 29, 30, 510, 556
- FLT\_EPSILON** macro, 29, 30, 510, 556
- FLT\_EVAL\_METHOD** macro, 25, 26, 29, 30, 98–100, 109, 247, 510, 556, 588, 601, 610, 646, 649
- FLT\_HAS\_SUBNORM** macro, 26, 30, 488, 510
- FLT\_IS\_IEC\_60559** macro, 26, 30, 510
- FLT\_MANT\_DIG** macro, 27, 29, 30, 109, 510, 556
- FLT\_MAX** macro, 28, 29, 30, 510, 556
- FLT\_MAX\_10\_EXP** macro, 28, 29, 30, 510, 556
- FLT\_MAX\_EXP** macro, 28, 29, 30, 510, 556
- FLT\_MIN** macro, 29, 30, 510, 556
- FLT\_MIN\_10\_EXP** macro, 28, 29, 30, 510, 556
- FLT\_MIN\_EXP** macro, 28, 29, 30, 510, 556
- FLT\_NORM\_MAX** macro, 29, 510, 556
- FLT\_RADIX** macro, 25, 27, 28–31, 65, 227, 232, 268, 352, 353, 376, 451, 452, 466, 510, 556, 562, 601, 602, 610
- FLT\_ROUND** macro, 25, 26, 223, 510, 555, 556, 559, 646
- FLT\_SNAN** macro, 27, 109, 510, 559
- FLT\_TRUE\_MIN** macro, 29, 30, 510
- FLTN\_DECIMAL\_DIG** macro, 603
- FLTN\_DIG** macro, 603
- FLTN\_EPSILON** macro, 604
- FLTN\_MANT\_DIG** macro, 602
- FLTN\_MAX** macro, 603
- FLTN\_MAX\_10\_EXP** macro, 603
- FLTN\_MAX\_EXP** macro, 603
- FLTN\_MIN** macro, 604
- FLTN\_MIN\_10\_EXP** macro, 603
- FLTN\_MIN\_EXP** macro, 603
- FLTN\_SNAN** macro, 602
- FLTN\_TRUE\_MIN** macro, 604
- FLTNX\_DECIMAL\_DIG** macro, 603
- FLTNX\_DIG** macro, 603
- FLTNX\_EPSILON** macro, 604
- FLTNX\_MANT\_DIG** macro, 602
- FLTNX\_MAX** macro, 603
- FLTNX\_MAX\_10\_EXP** macro, 603
- FLTNX\_MAX\_EXP** macro, 603
- FLTNX\_MIN** macro, 604
- FLTNX\_MIN\_10\_EXP** macro, 603
- FLTNX\_MIN\_EXP** macro, 603
- FLTNX\_SNAN** macro, 602
- FLTNX\_TRUE\_MIN** macro, 604
- fma** function, 249, 287, 517, 560, 585
- fma** type-generic macro, 417
- fmad128** function, 287, 522
- fmad32** function, 287, 522
- fmad64** function, 287, 522
- fMaddfN** function, 621
- fMaddfNx** function, 621
- fmadN** function, 620
- fmadNx** function, 620
- fmaf** function, 287, 517
- fmafN** function, 620
- fmafNx** function, 620
- fmall** function, 249, 287, 517
- fmax** function, 283, 284, 517, 563, 584
- fmax** type-generic macro, 417
- fmaxd128** function, 283, 522
- fmaxd32** function, 283, 522
- fmaxd64** function, 283, 522
- fmaxf** function, 283, 517
- fmaximum** function, 284, 285, 517, 560, 584
- fmaximum** type-generic macro, 417
- fmaximum\_mag** type-generic macro, 417
- fmaximum\_mag\_num** type-generic macro, 417
- fmaximum\_num** type-generic macro, 417
- fmaximum\_mag** function, 284, 285, 286, 517, 560, 584
- fmaximum\_mag\_num** function, 285, 286, 517, 560, 584, 585
- fmaximum\_mag\_numd128** function, 286, 522
- fmaximum\_mag\_numd32** function, 286, 522
- fmaximum\_mag\_numd64** function, 286, 522
- fmaximum\_mag\_numdN** function, 620
- fmaximum\_mag\_numdNx** function, 620
- fmaximum\_mag\_numf** function, 286, 517
- fmaximum\_mag\_numfN** function, 620
- fmaximum\_mag\_numfNx** function, 620
- fmaximum\_mag\_numl** function, 286, 517
- fmaximum\_magd128** function, 284, 522
- fmaximum\_magd32** function, 284, 522
- fmaximum\_magd64** function, 284, 522
- fmaximum\_magdN** function, 620
- fmaximum\_magdNx** function, 620
- fmaximum\_magf** function, 284, 517
- fmaximum\_magfN** function, 620
- fmaximum\_magfNx** function, 620
- fmaximum\_magl** function, 284, 517
- fmaximum\_num** function, 284, 285, 517, 560, 563, 584, 585, 591
- fmaximum\_numd128** function, 285, 522
- fmaximum\_numd32** function, 285, 522

**fmaximum\_numd64** function, 285, 522  
**fmaximum\_numdN** function, 620  
**fmaximum\_numdNx** function, 620  
**fmaximum\_numf** function, 285, 517  
**fmaximum\_numfN** function, 620  
**fmaximum\_numfNx** function, 620  
**fmaximum\_numl** function, 285, 517  
**fmaximumd128** function, 284, 419, 522  
**fmaximumd32** function, 284, 522  
**fmaximumd64** function, 284, 522  
**fmaximumdN** function, 620  
**fmaximumdNx** function, 620  
**fmaximumf** function, 284, 517  
**fmaximumfN** function, 620  
**fmaximumfNx** function, 620  
**fmaximuml** function, 284, 517  
**fmaxl** function, 283, 517  
**fMdivfN** function, 621  
**fMdivfNx** function, 621  
**fMencfN** function, 623  
**fMfafN** function, 621  
**fMfafNx** function, 621  
**fmin** function, 283, 284, 517, 563, 584  
**fmin** type-generic macro, 417  
**fmind128** function, 283, 522  
**fmind32** function, 283, 522  
**fmind64** function, 283, 522  
**fminf** function, 283, 517  
**fminimum** function, 284, 285, 286, 517, 560, 584  
**fminimum** type-generic macro, 417  
**fminimum\_mag** type-generic macro, 417  
**fminimum\_mag\_num** type-generic macro, 417  
**fminimum\_mag** function, 285, 286, 517, 560, 584  
**fminimum\_mag\_num** function, 285, 286, 517, 560, 584, 585  
**fminimum\_mag\_numd128** function, 286, 522  
**fminimum\_mag\_numd32** function, 286, 522  
**fminimum\_mag\_numd64** function, 286, 522  
**fminimum\_mag\_numdN** function, 620  
**fminimum\_mag\_numdNx** function, 620  
**fminimum\_mag\_numf** function, 286, 517  
**fminimum\_mag\_numfN** function, 620  
**fminimum\_mag\_numfNx** function, 620  
**fminimum\_mag\_numl** function, 286, 517  
**fminimum\_magd128** function, 285, 522  
**fminimum\_magd32** function, 285, 522  
**fminimum\_magd64** function, 285, 522  
**fminimum\_magdN** function, 620  
**fminimum\_magdNx** function, 620  
**fminimum\_magf** function, 285, 517  
**fminimum\_magfN** function, 620  
**fminimum\_magfNx** function, 620  
**fminimum\_magl** function, 285, 517  
**fminimum\_num** function, 284, 285, 286, 517, 560, 563, 584, 585  
**fminimum\_numd128** function, 286, 522  
**fminimum\_numd32** function, 286, 522  
**fminimum\_numd64** function, 286, 522  
**fminimum\_numdN** function, 620  
**fminimum\_numdNx** function, 620  
**fminimum\_numf** function, 285, 517  
**fminimum\_numfN** function, 620  
**fminimum\_numfNx** function, 620  
**fminimum\_numl** function, 286, 517  
**fminimumd128** function, 284, 522  
**fminimumd32** function, 284, 522  
**fminimumd64** function, 284, 522  
**fminimumdN** function, 620  
**fminimumdNx** function, 620  
**fminimumf** function, 284, 517  
**fminimumfN** function, 620  
**fminimumfNx** function, 620  
**fminimuml** function, 284, 517  
**fminl** function, 283, 517  
**fMmulfN** function, 621  
**fMmulfNx** function, 621  
**fmod** function, 278, 279, 516, 582, 583, 649  
**fmod** type-generic macro, 417  
**fmodd128** function, 278, 521  
**fmodd32** function, 278, 521  
**fmodd64** function, 278, 521  
**fmoddN** function, 618  
**fmoddNx** function, 618  
**fmodf** function, 278, 516  
**fmodfN** function, 618  
**fmodfNx** function, 618  
**fmodl** function, 278, 516  
**fMsqrtfN** function, 621  
**fMsqrtfNx** function, 621  
**fMsubfN** function, 621  
**fMsubfNx** function, 621  
**fmul** function, 288, 517, 560, 629  
**fmul** type-generic macro, 418  
**fmull** function, 288, 517, 560  
**fMxaddfN** function, 621  
**fMxaddfNx** function, 621  
**fMxdivfN** function, 621  
**fMxdivfNx** function, 621  
**fMxfmafN** function, 621  
**fMxfmafNx** function, 621  
**fMxmulfN** function, 621  
**fMxmulfNx** function, 621  
**fMxsqrtfN** function, 621  
**fMxsqrtfNx** function, 621  
**fMxsubfN** function, 621  
**fMxsubfNx** function, 621  
**fopen** function, i, 186, 343, 344, 345, 346, 347, 537, 641, 653, 675, 676, 720  
**FOPEN\_MAX** macro, 339, 343, 344, 537, 674  
**open\_s** function, 538, 674, 675, 676  
**for** keyword, 55



- for**, 166
- form feed, 21
- form-feed character, 20, 54
- form-feed escape sequence (`\f`), 21, 69, 219
- format conversion of integer types header, 488
- format flag
  - `+`, 349, 448
  - `-`, 349, 448
  - `#`, 349, 448
  - `0`, 349, 448
  - `space`, 349, 448
- format modifier
  - `D`, 350, 358, 450, 455
  - `DD`, 350, 358, 450, 455
  - `E`, 439
  - `H`, 350, 358, 449, 455
  - `h`, 349, 357, 449, 454
  - `hh`, 349, 357, 449, 454
  - `j`, 350, 357, 449, 455
  - `L`, 350, 358, 449, 455
  - `l`, 350, 357, 449, 454
  - `ll`, 350, 357, 449, 454
  - `O`, 439
  - `t`, 350, 357, 449, 455
  - `wfN`, 350, 357, 449, 455
  - `wN`, 350, 357, 449, 455
  - `z`, 350, 357, 449, 455
- formatted input/output function, 242, 348, 677
  - wide character, 447, 702
- forward reference, 6
- FP\_** identifier prefix, 248, 488
- FP\_CONTRACT** pragma, viii, 77, 158, 198, 199, 251, 512, 590, 591, 639, 647
- FP\_FAST\_D32ADDD128** macro, 249, 518
- FP\_FAST\_D32ADDD64** macro, 249, 518
- FP\_FAST\_D32DIVD128** macro, 249, 518
- FP\_FAST\_D32DIVD64** macro, 249, 518
- FP\_FAST\_D32FMAD128** macro, 249, 518
- FP\_FAST\_D32FMAD64** macro, 249, 518
- FP\_FAST\_D32MULD128** macro, 249, 518
- FP\_FAST\_D32MULD64** macro, 249, 518
- FP\_FAST\_D32SQRTD128** macro, 249, 518
- FP\_FAST\_D32SQRTD64** macro, 249, 518
- FP\_FAST\_D32SUBD128** macro, 249, 518
- FP\_FAST\_D32SUBD64** macro, 249, 518
- FP\_FAST\_D64ADDD128** macro, 249, 518
- FP\_FAST\_D64DIVD128** macro, 249, 518
- FP\_FAST\_D64FMAD128** macro, 249, 518
- FP\_FAST\_D64MULD128** macro, 249, 518
- FP\_FAST\_D64SQRTD128** macro, 249, 518
- FP\_FAST\_D64SUBD128** macro, 249, 518
- FP\_FAST\_DADDL** macro, 249, 512, 611
- FP\_FAST\_DDIVL** macro, 249, 512
- FP\_FAST\_DFMA** macro, 249, 512
- FP\_FAST\_DMADDDN** macro, 612
- FP\_FAST\_DMDDIVN** macro, 611
- FP\_FAST\_FMDIVN** macro, 612
- FP\_FAST\_DMFADN** macro, 611
- FP\_FAST\_FMFADN** macro, 612
- FP\_FAST\_DMMULDN** macro, 611
- FP\_FAST\_FMMULDN** macro, 612
- FP\_FAST\_DMSQRTDN** macro, 611
- FP\_FAST\_FMSQRTDN** macro, 612
- FP\_FAST\_DMSUBDN** macro, 611
- FP\_FAST\_FMSUBDN** macro, 612
- FP\_FAST\_DMULL** macro, 249, 512
- FP\_FAST\_DMADDDN** macro, 612
- FP\_FAST\_FMXADDDN** macro, 612
- FP\_FAST\_DMDDIVN** macro, 612
- FP\_FAST\_FMXDIVN** macro, 612
- FP\_FAST\_DMXFADN** macro, 612
- FP\_FAST\_FMXFADN** macro, 612
- FP\_FAST\_DMXMULDN** macro, 612
- FP\_FAST\_FMXMULDN** macro, 612
- FP\_FAST\_DMXSQRTDN** macro, 612
- FP\_FAST\_FMXSQRTDN** macro, 612
- FP\_FAST\_DMXSUBDN** macro, 612
- FP\_FAST\_FMXSUBDN** macro, 612
- FP\_FAST\_DSQRTL** macro, 249, 512
- FP\_FAST\_DSUBL** macro, 249, 512
- FP\_FAST\_FADD** macro, 249, 512, 611
- FP\_FAST\_FADDL** macro, 249, 512, 611
- FP\_FAST\_FDIV** macro, 249, 512
- FP\_FAST\_FDIVL** macro, 249, 512
- FP\_FAST\_FFMA** macro, 249, 512
- FP\_FAST\_FFMAL** macro, 249, 512
- FP\_FAST\_FMA** macro, 249, 512, 611
- FP\_FAST\_FMAD128** macro, 249, 518
- FP\_FAST\_FMAD32** macro, 249, 518
- FP\_FAST\_FMAD64** macro, 249, 518
- FP\_FAST\_FMADDFN** macro, 611
- FP\_FAST\_FMADDFN** macro, 612
- FP\_FAST\_FMADN** macro, 611
- FP\_FAST\_FMAFN** macro, 611
- FP\_FAST\_FMAF** macro, 249, 512, 524, 611
- FP\_FAST\_FMAFN** macro, 611
- FP\_FAST\_FMAFN** macro, 611
- FP\_FAST\_FMAL** macro, 249, 512
- FP\_FAST\_FMDIVN** macro, 611
- FP\_FAST\_FMDIVN** macro, 612
- FP\_FAST\_FMFMAFN** macro, 611
- FP\_FAST\_FMFMAFN** macro, 612
- FP\_FAST\_FMMULFN** macro, 611
- FP\_FAST\_FMMULFN** macro, 612
- FP\_FAST\_FMSQRTFN** macro, 611
- FP\_FAST\_FMSQRTFN** macro, 612
- FP\_FAST\_FMSUBFN** macro, 611
- FP\_FAST\_FMSUBFN** macro, 612
- FP\_FAST\_FMUL** macro, 249, 512
- FP\_FAST\_FMULL** macro, 249, 512

- FP\_FAST\_FMXADDFN** macro, 612  
**FP\_FAST\_FMXADDFNX** macro, 612  
**FP\_FAST\_FMXDIVFN** macro, 612  
**FP\_FAST\_FMXDIVFNX** macro, 612  
**FP\_FAST\_FMXFMFNF** macro, 612  
**FP\_FAST\_FMXFMFNFNX** macro, 612  
**FP\_FAST\_FMXMULFN** macro, 612  
**FP\_FAST\_FMXMULFNX** macro, 612  
**FP\_FAST\_FMXSQRTFN** macro, 612  
**FP\_FAST\_FMXSQRTFNX** macro, 612  
**FP\_FAST\_FMXSUBFN** macro, 612  
**FP\_FAST\_FMXSUBFNX** macro, 612  
**FP\_FAST\_FSQRT** macro, 249, 512  
**FP\_FAST\_FSQRTL** macro, 249, 512  
**FP\_FAST\_FSUB** macro, 249, 512  
**FP\_FAST\_FSUBL** macro, 249, 512  
**FP\_ILOGB0** macro, 249, 250, 264, 512  
**FP\_ILOGBNAN** macro, 249, 250, 264, 512  
**FP\_INFINITE** macro, 248, 512  
**FP\_INT\_DOWNWARD** macro, 248, 512  
**FP\_INT\_TONEAREST** macro, 248, 512  
**FP\_INT\_TONEARESTFROMZERO** macro, 248, 512  
**FP\_INT\_TOWARDZERO** macro, 248, 512  
**FP\_INT\_UPWARD** macro, 248, 278, 512  
**FP\_LLOGB0** macro, 250, 265, 512  
**FP\_LLOGBNAN** macro, 250, 265, 512  
**FP\_NAN** macro, 248, 512  
**FP\_NORMAL** macro, 248, 512  
**FP\_SUBNORMAL** macro, 248, 512  
**FP\_ZERO** macro, 248, 512  
**fpclassify** macro, 252, 512, 562, 563  
**fpos\_t** type, 339, 341, 370, 371, 537  
**fprintf\_s** function, 538, 677  
**fputc** function, 21, 340, 341, 343, 367, 368, 370, 537  
**fputs** function, 196, 340, 367, 537  
**fputwc** function, 340, 341, 343, 462, 463, 548  
**fputws** function, 340, 462, 548  
**frac\_digits** structure member, 241, 243, 246  
**fread** function, 184, 186, 340, 369, 537, 642  
**free** function, 382, 383, 384, 408, 409, 539, 643, 719  
**free\_aligned\_sized** function, 383, 539, 643, 719, 723  
**free\_sized** function, 382, 383, 539, 643, 719, 723  
 freestanding execution environment, 9, 12  
**freopen** function, 341, 342, 347, 537  
**freopen\_s** function, 538, 676  
**frexp** function, 263, 264, 514, 575, 622, 633, 634  
**frexp** type-generic macro, 417  
**frexpd128** function, 263, 519  
**frexpd32** function, 263, 519  
**frexpd64** function, 263, 519  
**frexpdN** function, 616  
**frexpdNx** function, 616  
**frexpf** function, 263, 514  
**frexpfN** function, 616  
**frexpfNx** function, 616  
**frexpl** function, 263, 514  
**fromfp** function, 248, 277, 278, 516, 561, 565, 582  
**fromfp** function, 277  
**fromfp** type-generic macro, 417  
**fromfpd128** function, 277, 521  
**fromfpd32** function, 277, 521  
**fromfpd64** function, 277, 521  
**fromfpdN** function, 618  
**fromfpdNx** function, 618  
**fromfpf** function, 277, 516  
**fromfpfN** function, 618  
**fromfpfNx** function, 618  
**fromfpl** function, 277, 516  
**fromfpx** function, 248, 278, 516, 561, 565, 582  
**fromfpx** function, 278  
**fromfpx** type-generic macro, 417  
**fromfpxd128** function, 278, 521  
**fromfpxd32** function, 278, 521  
**fromfpxd64** function, 278, 521  
**fromfpxdN** function, 618  
**fromfpxdNx** function, 618  
**fromfpxf** function, 278, 516  
**fromfpxfN** function, 618  
**fromfpxfNx** function, 618  
**fromfpxl** function, 278, 516  
**fscanf** function, 236, 237, 340, 356, 359–364, 489, 537, 650, 677  
**fscanf\_s** function, 538, 677, 678–680, 682  
**fseek** function, 340, 343, 346, 369, 370, 371, 464, 537, 643  
**fsetpos** function, 341, 342, 346, 369, 370, 371, 464, 537, 643, 650  
**fsqrt** function, 289, 517, 560, 629  
**fsqrt** type-generic macro, 418  
**fsqrtl** function, 289, 517, 560  
**fsub** function, 287, 517, 560, 629  
**fsub** type-generic macro, 418  
**fsubl** function, 288, 419, 517, 560, 630  
**ftell** function, 370, 371, 537, 633, 643, 650  
 full declarator, 134  
 full expression, 160  
 fully buffered, 342  
 fully buffered stream, 342  
 function  
     argument, 80, 172  
     body, 171  
     byte input/output, 340  
     call, 80  
         library, 206  
     declarator, 137  
     definition, 137, 171

- designator, 52
- image, 21
- inline, 131
- library, 11, 206
- name length, 22, 57, 203
- no-return, 132
- parameter, 13, 80, 172
- prototype, 13, 36, 137, 172, 247
- prototype scope, 37, 135, 136
- recursive call, 80
- return, 169, 566
- scope, 36
- stateless, 156
- type, 41
- type conversion, 52
- function call
  - effectless, 156
- function prototype, 36
- function prototype scope, 37
- function scope, 36
- function specifier, 131
  - \_Noreturn**, 55
  - inline**, 55
- function type, 39
- function type attribute, 155
- function-call operator `(( ))`, 80
- function-like macro, 191
- fundamental alignment, 46
- Fused multiply-add and round to narrower type, 289
- future direction
  - language, 203
  - library, 488
- fwide** function, 341, 342, 462, 463, 548
- fwprintf** function, 236, 340, 447, 453, 457–459, 461, 490, 548, 644, 650, 703
- fwprintf\_s** function, 550, 702, 703
- fwrite** function, 340, 369, 370, 537, 642
- fscanf** function, 236, 340, 453, 454, 457–459, 461, 464, 490, 548, 650, 703
- fscanf\_s** function, 550, 703, 705, 706, 710
- gamma function, 272, 580
- general utilities header, 373, 489
- general utility, 686
  - wide string, 464, 710
- general wide string utility, 464, 710
- generic association, 78
- generic parameter, 416
- generic selection, 77, 78
- getc** function, 340, 367, 368, 537
- getchar** function, 16, 340, 367, 368, 537
- getenv** function, 385, 386, 539, 643, 645
- getenv\_s** function, 540, 687, 688
- getpayload** function, 523, 560, 586, 587
- getpayload128** function, 524, 586
- getpayload32** function, 524, 586
- getpayload64** function, 524, 586
- getpayloadN** function, 622
- getpayloadNx** function, 622
- getpayloadf** function, 523, 586
- getpayloadfN** function, 622
- getpayloadfNx** function, 622
- getpayloadl** function, 523, 586
- gets** (obsolete), 685, 724
- gets\_s** function, 538, 685
- getwc** function, 340, 463, 548
- getwchar** function, 340, 463, 548
- gmtime** function, 435, 436, 437, 438, 547, 651
- gmtime\_r** function, 437, 547, 721
- gmtime\_s** function, 547, 701, 702
- goto** keyword, 55
- goto** statement, 36, 166
- graphic character, 20
- greater-than operator `(>)`, 92
- greater-than-or-equal-to operator `(>=)`, 92
- grouping** structure member, 241, 243, 244
- H format modifier, 350, 358, 449, 455
- h format modifier, 349, 357, 449, 454
- happens before, 18
- header, 11, 204
- header name, 54, 73, 182
- hexadecimal digit, 61, 65, 69
- `\xhexadecimal digit` (hexadecimal-character escape sequence), 69
- hexadecimal digit sequence, 61
- hexadecimal literal, 60
- hexadecimal prefix, 61
- hexadecimal-character escape sequence (`\xhexadecimal digit`), 69
- hh format modifier, 349, 357, 449, 454
- hidden, 37
- high-order bit, 4, 5
- horizontal tab, 21
- horizontal-tab character, 20, 54
- horizontal-tab escape sequence (`\t`), 21, 69, 218, 219, 483
- hosted execution environment, 9, 12, 13
- HUGE\_VAL** macro, ii, 247, 251, 376, 466, 512, 571
- HUGE\_VAL\_D128** macro, 248, 518
- HUGE\_VAL\_D32** macro, 247, 248, 518
- HUGE\_VAL\_D64** macro, 248, 518
- HUGE\_VAL\_DN** macro, 611
- HUGE\_VAL\_DNx** macro, 611
- HUGE\_VAL\_FN** macro, 611
- HUGE\_VAL\_FNx** macro, 611
- HUGE\_VALF** macro, 247, 377, 466, 512, 571
- HUGE\_VALL** macro, 247, 377, 466, 512, 571
- hyperbolic function
  - complex, 211, 593
  - real, 259, 574
- hypot** function, 270, 515, 559, 564, 578, 592
- hypot** type-generic macro, 417

- hypotd128** function, 270, 520
- hypotd32** function, 270, 520
- hypotd64** function, 270, 520
- hypotdN** function, 617
- hypotdNx** function, 617
- hypotf** function, 270, 515
- hypotfN** function, 617
- hypotfNx** function, 617
- hypotl** function, 270, 515
- I** macro, 209, 216
- idempotent, 156
- identifier, 56, 77
  - continue, 56, 553
  - maximum length, 56
  - name space, 38
  - reserved, 55, 205, 654, 663, 672
  - rule, 654
  - scope, 36
  - start, 56, 553
  - type, 39
- identifier continue, 56, 553
- identifier list, 176
- identifier start, 56, 553
- if** keyword, 55
- if** preprocessing directive, 179
- if** statement, 163
- if\_empty** embed parameter, 184, 189, 190
- ifdef**, 9, 175, 178, 179, 180, 227, 231, 232, 254–281, 282, 283–287, 288, 289–292, 375, 377, 419, 467, 556, 557, 585–587
- ifdef** preprocessing directive, 180
- ifndef** preprocessing directive, 180
- ignore\_handler\_s** function, 540, 687
- ilogb** function, 249, 264, 514, 560, 576, 633
- ilogb** type-generic macro, 417
- ilogbd128** function, 264, 519
- ilogbd32** function, 264, 519
- ilogbd64** function, 264, 519
- ilogbdN** function, 616
- ilogbdNx** function, 616
- ilogbf** function, 264, 514
- ilogbfN** function, 616
- ilogbfNx** function, 616
- ilogbl** function, 264, 514
- imaginary** (obsolete), 209
- imaxabs** function, 237, 511
- imaxdiv** function, 236, 237, 511
- imaxdiv\_t** type, 236, 237, 511
- implementation, 6
- implementation limit, 6, 9, 22, 57, 134, 164, 555
- implementation resource width, 183
- implementation-defined behavior, 3, 9, 645
- implementation-defined value, 7
- implicit conversion, 47
- implicit initialization, 145
- include** preprocessing directive, 11, 182
- inclusive OR operator
  - bitwise (**|**), 94
  - bitwise assignment (**|=**), 98
- incomplete, 39
- incomplete type, 39
- independent, 156
- indeterminate representation, 7
- indeterminately sequenced, 14, 80, 82, 98
- indirection operator (**\***), 79, 86
- indivisible unit of work, 393
- inequality operator (**!=**), 93
- infinitary, 250
- infinity, 589
- INFINITY** macro, 27, 109, 216, 248, 489, 510, 512, 559, 590, 591
- initial position, 21
- initial shift state, 20
- initialization, 12, 38, 51, 83, 143, 568
  - in block, 160
- initialized, 12
- initializer, 144
  - permitted form, 100
  - string literal, 51
- inline**, 131
- inline definition, 132
- inline function, 131
- inline** keyword, 55
- inner scope, 37
- input code unit reads, 398
- input failure, 459, 460, 678, 679, 681–683, 685, 703, 705, 706, 708–710
- input/output
  - device, 14
- input/output function
  - character, 366, 685
  - direct, 369
  - formatted, 348, 677
  - wide character, 447, 702
  - wide character, 461
  - formatted, 447, 702
- input/output header, 339, 489, 673
- INT** identifier prefix, 337, 338, 489, 536
- int** identifier prefix, 335, 336, 489, 536, 724
- int** keyword, 55
- int** type, 40, 48, 62, 110
- int** type conversion, 47–49
- intN\_t** type, 335
- INTN\_C** macro, 338
- INTN\_MAX** macro, 337
- INTN\_MIN** macro, 337, 338
- INTN\_WIDTH** macro, 337
- int8\_t** type, 335
- INT\_FAST** identifier prefix, 337, 536
- int\_fast** identifier prefix, 336, 536, 724
- INT\_LEAST** identifier prefix, 337, 536



- int\_least** identifier prefix, 328, 329, 335, 336, 338, 535, 536
- int\_curr\_symbol** structure member, 241, 244–246
- INT\_FASTN\_MAX** macro, 337
- INT\_FASTN\_MIN** macro, 337, 338
- INT\_FASTN\_WIDTH** macro, 337
- int\_fast16\_t** type, 313, 336
- int\_fast32\_t** type, 236, 313, 336
- int\_fast64\_t** type, 313, 336
- int\_fast8\_t** type, 313, 336
- int\_fastN\_t** type, 336
- int\_frac\_digits** structure member, 241, 244, 246
- INT\_LEASTN\_MAX** macro, 337
- INT\_LEASTN\_MIN** macro, 337, 338
- int\_leastN\_t** type, 335
- INT\_LEASTN\_WIDTH** macro, 337
- int\_least16\_t** type, 313, 336
- int\_least32\_t** type, 313, 335, 336
- int\_least64\_t** type, 313, 336
- int\_least8\_t** type, 313, 336
- INT\_MAX** macro, 24, 40, 110, 179, 249, 250, 264, 512, 555, 642
- INT\_MIN** macro, 24, 40, 249, 250, 512, 555
- int\_n\_cs\_precedes** structure member, 241, 244, 246
- int\_n\_sep\_by\_space** structure member, 241, 244–246
- int\_n\_sign\_posn** structure member, 241, 244–246
- int\_p\_cs\_precedes** structure member, 241, 244, 246
- int\_p\_sep\_by\_space** structure member, 241, 244–246
- int\_p\_sign\_posn** structure member, 241, 244–246
- INT\_WIDTH** macro, 23, 278, 512, 555
- integer arithmetic function, 237, 388
- integer constant expression, 52, 100, 112, 116, 135, 144, 149, 207
- integer conversion rank, 47
- integer literal, 60
- integer promotion, 15, 48, 87, 91, 164, 349, 449
- integer suffix, 62
- integer type, 41, 335
  - extended, 40, 48, 62, 335
- integer type conversion, 47, 48, 49, 565
- integer types header, 335, 489
- inter-thread happens before, 18
- interactive device, 14, 342, 347
- interchange floating type, 599
- internal linkage, 37
- internal name, 57
- interrupt, 21
- INTMAX\_C** macro, 338, 536
- INTMAX\_MAX** macro, 238, 337, 536
- INTMAX\_MIN** macro, 238, 337, 338, 536
- intmax\_t** type, 24, 179, 237, 238, 313, 336, 338, 350, 357, 449, 455, 511, 536, 723, 725
- INTMAX\_WIDTH** macro, 337, 536
- INTPTR\_MAX** macro, 337, 536
- INTPTR\_MIN** macro, 337, 338, 536
- intptr\_t** type, 313, 336, 536
- INTPTR\_WIDTH** macro, 337, 536
- is** identifier prefix, 488–490
- isalnum** function, 217, 219, 509
- isalpha** function, 217, 482, 509, 652
- isblank** function, 217, 218, 509, 652
- iscanonical** macro, 25, 252, 512, 562, 563
- iscntrl** function, 217, 218, 219, 509
- isdigit** function, 217, 218, 219, 242, 509
- iseqsig** macro, 294, 295, 517, 561, 588
- isfinite** macro, 252, 512, 562, 563, 575, 591
- isgraph** function, 218, 482, 509
- isgreater** macro, 293, 517, 561, 562
- isgreaterequal** macro, 293, 517, 562, 570, 584
- isinf** macro, 252, 253, 512, 562, 563, 577, 590, 591
- isless** macro, 293, 294, 517, 562, 570, 571
- islessequal** macro, 294, 517, 562
- islessgreater** macro, 294, 517
- islower** function, 4, 217, 218, 219, 220, 509, 652
- isnan** macro, 253, 512, 562, 563, 584, 590, 591
- isnormal** macro, 253, 512, 562, 563
- ISO/IEC 10646, 2, 200
- ISO/IEC 2382, 2, 3
- ISO/IEC 60559, 2, 201, 209, 222, 233, 251, 279, 293, 558, 589, 598
- ISO/IEC 9945–2, 241
- ISO 4217, 2, 244
- ISO 80000–2, 2, 3
- ISO 80000–3, 727
- ISO 8601 series, 2, 438
- isprint** function, 21, 218, 509
- ispunct** function, 217, 218, 219, 509, 652
- issignaling** macro, 253, 254, 512, 562, 563
- isspace** function, 204, 217, 218, 219, 509, 652
- issubnormal** macro, 254, 513, 562, 563
- isunordered** macro, 294, 517, 562
- isupper** function, 217, 219, 220, 509, 652
- iswalnum** function, 483, 484, 485, 550
- iswalpha** function, 482, 483, 485, 550, 652
- iswblank** function, 483, 485, 550, 652
- iswcntrl** function, 483, 484, 485, 551
- iswctype** function, 485, 486, 551, 644, 652
- iswdigit** function, 483, 484, 485, 551
- iswgraph** function, 482, 484, 485, 551
- iswlower** function, 483, 484, 485, 486, 551, 652
- iswprint** function, 482, 484, 485, 551

- iswpunct** function, 482, 483, **484**, 485, 551, 652
- iswspace** function, 204, 482, 483, **484**, 485, 551, 652
- iswupper** function, 483, **484**, 485, 486, 551, 652
- iswxdigit** function, 485, 551
- isxdigit** function, 219, 242, 509
- iszero** macro, **254**, 513, 562, 563
- italic type convention*, 3, **36**
- iteration statement, **165**
  
- j format modifier, 350, 357, 449, 455
- jmp\_buf** type, **296**, 297, 532, 719
- jump statement, **166**
  
- keyword, 55, 653
  - \_Lengthof**, 55
- kill\_dependency** macro, 17, **311**, 533
- known constant size, **42**
  
- L** encoding prefix, 67–69, 71
- L format modifier, 350, 358, 449, 455
- l format modifier, 350, 357, 449, 454
- l** identifier suffix, 209, **247**, 416, 418, 488
- L\_tmpnam** macro, **340**, 345, 537
- L\_tmpnam\_s** macro, 538, **673**, 674
- label
  - case**, 55
  - default**, 55
- label name, 36, **38**
- labeled statement, **161**
- labs** function, **388**, 539
- language, **36**
  - encoding prefix
    - L**, 67–69, 71
    - U**, 67–69, 71, 72
    - u**, 67–69, 71, 72
    - u8**, 67–69, 71, 72
  - future direction, **203**
  - syntax summary, **491**
- Latin alphabet, **19**
- LC\_** identifier prefix, **241**, 488
- LC\_ALL** macro, **241**, **242**, 245, 512
- LC\_COLLATE** macro, **241**, **242**, 410, 471, 512
- LC\_CTYPE** macro, 47, **241**, **242**, 373, 389, 390, 394, 395, 477, 482, 485–487, 512, 643, 644, 690, 691
- LC\_MONETARY** macro, **241**, **242**, 245, 512
- LC\_NUMERIC** macro, **241**, **242**, 245, 512
- LC\_TIME** macro, **241**, **242**, 435, 437, 512
- lconv** structure type, **241**, 242, 512
- LDBL\_** identifier prefix, 488
- LDBL\_DECIMAL\_DIG** macro, 27, 488, 510, 556
- LDBL\_DIG** macro, 28, 510, 556
- LDBL\_EPSILON** macro, 29, 510, 556
- LDBL\_HAS\_SUBNORM** macro, 26, 488, 510
- LDBL\_IS\_IEC\_60559** macro, 26, 510
- LDBL\_MANT\_DIG** macro, 27, 510, 556
- LDBL\_MAX** macro, 28, 510, 556
- LDBL\_MAX\_10\_EXP** macro, 28, 510, 556
- LDBL\_MAX\_EXP** macro, 28, 510, 556
- LDBL\_MIN** macro, 29, 510, 556
- LDBL\_MIN\_10\_EXP** macro, 28, 510, 556
- LDBL\_MIN\_EXP** macro, 28, 510, 556
- LDBL\_NORM\_MAX** macro, 29, 510, 556
- LDBL\_SNAN** macro, 27, 510, 559
- LDBL\_TRUE\_MIN** macro, 29, 510
- ldexp** function, **264**, 265, 514, **576**, 622
- ldexp** type-generic macro, **417**
- ldexpd128** function, **264**, 519
- ldexpd32** function, **264**, 519
- ldexpd64** function, **264**, 519
- ldexpdN** function, **616**
- ldexpdNx** function, **616**
- ldexpf** function, **264**, 514
- ldexpfN** function, **616**
- ldexpfNx** function, **616**
- ldexpl** function, **264**, 514
- ldiv** function, 143, 373, **388**, 389, 539
- ldiv\_t** type, 373, 389, 538, 539
- leading underscore in identifier, **206**
- least significant index, **318**
- left-shift assignment operator (<=<), **98**
- left-shift operator (<<), **91**
- length
  - external name, **22**, **57**, **203**
  - function name, **22**, **57**, **203**
  - identifier, **56**
  - internal name, **22**, **57**
- length function, 389, 415, 475, 478, 699, 715
- length modifier, **349**, **356**, **448**, **454**
- length of a string, **204**
- length of a wide string, **204**
- less-than operator (<), **92**
- less-than-or-equal-to operator (<=), **92**
- letter, **20**, 217
  - lowercase, **20**
- lexical element, 11, **54**
- lgamma** function, **273**, 515, **580**
- lgamma** type-generic macro, **417**
- lgammad128** function, **273**, 520
- lgammad32** function, **273**, 520
- lgammad64** function, **273**, 520
- lgammadN** function, **617**
- lgammadNx** function, **617**
- lgammaf** function, **273**, 515
- lgammafN** function, **617**
- lgammafNx** function, **617**
- lgammal** function, **273**, 515
- library, 11, **204**, **671**
  - constant
    - memory\_order\_acq\_rel**, **309**, 310, 311, 314–316, 532

- memory\_order\_acquire**, 309, 311, 314, 316, 532
- memory\_order\_consume**, 309, 311, 314, 532
- memory\_order\_relaxed**, 157, 309, 310, 311, 532
- memory\_order\_release**, 309, 311, 314, 315, 532
- memory\_order\_seq\_cst**, 19, 43, 82, 97, 98, 307, 309, 310, 312, 532
- mtx\_plain**, 421, 425, 546
- mtx\_recursive**, 421, 425, 546
- mtx\_timed**, 422, 425, 546
- thrd\_busy**, 422, 426, 546
- thrd\_error**, 422, 423–430, 546
- thrd\_nomem**, 422, 423, 426, 546
- thrd\_success**, 422, 423–430, 546
- thrd\_timedout**, 422, 424, 425, 546
- family
  - canonicalize**, 34
  - decodebin**, 34
  - decodebind**, 292, 293
  - decodedec**, 34
  - decodedecd**, 291, 292
  - encodebin**, 34
  - encodebind**, 292
  - encodedec**, 34
  - encodedecd**, 291
  - modf**, 34
  - strfrom**, 226, 228
  - strfromd**, 375
  - strto**, 34, 226, 228
  - strtod**, 377, 378
  - wcsto**, 34, 226, 228
  - wcstod**, 467, 468
- function
  - \_Exit**, 299, 385, 386, 539, 640, 650
  - abort**, 155, 208, 298, 299, 308, 342, 384, 539, 640, 650, 687
  - abort\_handler\_s**, 540, 687
  - abs**, ii, 206, 207, 388, 539, 721
  - acos**, 226, 254, 255, 417, 513, 565, 572, 609
  - acosd128**, 228, 254, 518
  - acosd32**, 228, 254, 518
  - acosd64**, 228, 254, 518
  - acosf**, 226, 254, 513
  - acosh**, 259, 417, 513, 565, 574
  - acoshd128**, 259, 518
  - acoshd32**, 259, 518
  - acoshd64**, 259, 518
  - acoshf**, 259, 419, 513
  - acoshl**, 259, 513
  - acosl**, 226, 254, 513, 609
  - acospi**, 257, 513, 565, 573
  - acospid128**, 257, 518
  - acospid32**, 257, 518
  - acospid64**, 257, 518
  - acospif**, 257, 513
  - acospil**, 257, 513
  - aligned\_alloc**, 381, 382, 383, 539, 634, 643, 650, 724
  - asctime**, 199, 200, 435, 436, 547, 644
  - asctime\_s**, 547, 700, 701
  - asin**, 255, 417, 513, 565, 572
  - asind128**, 255, 518
  - asind32**, 255, 518
  - asind64**, 255, 518
  - asinf**, 255, 513
  - asinh**, 259, 260, 417, 513, 565, 574
  - asinhd128**, 260, 518
  - asinhd32**, 260, 518
  - asinhd64**, 260, 518
  - asinhf**, 260, 513
  - asinhL**, 260, 513
  - asinl**, 255, 513
  - asinpi**, 257, 513, 565, 573
  - asinp128**, 257, 518
  - asinp32**, 257, 518
  - asinp64**, 257, 518
  - asinpif**, 257, 513
  - asinpil**, 257, 513
  - at\_quick\_exit**, 384, 385, 386, 539, 634, 643, 724
  - atan**, 255, 354, 417, 453, 513, 565, 572
  - atan2**, 255, 256, 513, 565, 571, 572, 573, 592
  - atan2d128**, 255, 518
  - atan2d32**, 255, 518
  - atan2d64**, 255, 518
  - atan2f**, 255, 513
  - atan2L**, 255, 513
  - atan2pi**, 258, 513, 565, 571, 573
  - atan2pid128**, 258, 518
  - atan2pid32**, 258, 518
  - atan2pid64**, 258, 518
  - atan2pif**, 258, 513
  - atan2pil**, 258, 513
  - atand128**, 255, 518
  - atand32**, 255, 518
  - atand64**, 255, 518
  - atanf**, 255, 513
  - atanh**, 260, 417, 513, 565, 574
  - atanhd128**, 260, 519
  - atanhd32**, 260, 518
  - atanhd64**, 260, 519
  - atanhf**, 260, 513
  - atanhL**, 260, 513
  - atanl**, 255, 419, 513
  - atanpi**, 257, 258, 513, 565, 573
  - atanpid128**, 258, 518
  - atanpid32**, 257, 518

**atanpid64**, 257, 518  
**atanpif**, 257, 513  
**atanpil**, 257, 513  
**atexit**, 384, 385, 386, 539, 634, 643, 653  
**atof**, 226, 373, 374, 538  
**atoi**, 207, 373, 374, 538  
**atol**, 373, 374, 538  
**atoll**, 373, 374, 538  
**atomic\_compare\_exchange\_strong**,  
82, 98, 99, 314, 315, 533  
**atomic\_compare\_exchange\_strong\_explicit**,  
314, 533  
**atomic\_compare\_exchange\_weak**,  
314, 315, 533  
**atomic\_compare\_exchange\_weak\_explicit**,  
314, 533  
**atomic\_exchange**, 314, 533  
**atomic\_exchange\_explicit**, 314, 533  
**atomic\_fetch\_**, 315, 533  
**atomic\_flag\_clear**, 316, 317, 533  
**atomic\_flag\_clear\_explicit**, 316,  
533  
**atomic\_flag\_test\_and\_set**, 316, 533  
**atomic\_flag\_test\_and\_set\_explicit**,  
316, 533  
**atomic\_init**, 308, 533  
**atomic\_is\_lock\_free**, 299, 312, 533,  
640  
**atomic\_load**, 314, 315, 533  
**atomic\_load\_explicit**, 310, 311, 314,  
533  
**atomic\_signal\_fence**, 312, 533  
**atomic\_store**, 314, 533  
**atomic\_store\_explicit**, 310, 311,  
314, 533  
**atomic\_thread\_fence**, 157, 311, 312,  
533  
**btowc**, 452, 453, 477, 549  
**c16rtomb**, 443, 444, 548  
**c32rtomb**, 445, 548  
**c8rtomb**, 442, 548, 722  
**cabs**, 214, 416, 417, 508, 592  
**cabsf**, 214, 419, 508  
**cabsl**, 214, 508  
**cacos**, 210, 417, 507, 592, 593  
**cacosf**, 210, 507  
**cacosh**, 211, 212, 417, 507, 593  
**cacoshf**, 211, 507  
**cacoshl**, 211, 507  
**cacosl**, 210, 507  
**cacospi**, 488  
**call\_once**, 373, 421, 422, 538, 546, 723  
**calloc**, 381, 382, 383, 539, 634, 643, 650  
**canonicalize**, 25, 282, 417, 516, 561,  
584  
**canonicalized128**, 282, 522  
**canonicalized32**, 282, 522  
**canonicalized64**, 282, 522  
**canonicalizelf**, 282, 516  
**canonicalizel**, 282, 516  
**carg**, 214, 215, 419, 508, 592, 628  
**cargf**, 214, 508  
**cargl**, 214, 508  
**casin**, 210, 417, 507, 592  
**casinf**, 210, 507  
**casinh**, 212, 417, 507, 592, 594  
**casinhf**, 212, 507  
**casinhl**, 212, 507  
**casinl**, 210, 507  
**casinpi**, 488  
**catan**, 210, 211, 417, 507, 592  
**catanf**, 210, 507  
**catanh**, 212, 417, 507, 592, 594  
**catanhf**, 212, 507  
**catanhl**, 212, 507  
**catanl**, 210, 507  
**catanpi**, 488  
**cbrt**, 78, 269, 418, 515, 577  
**cbrtd128**, 269, 520  
**cbrtd32**, 269, 520  
**cbrtd64**, 269, 520  
**cbrtf**, 78, 269, 418, 515  
**cbrtl**, 78, 269, 418, 515  
**ccompoundn**, 488  
**ccos**, 211, 417, 507, 592  
**ccosf**, 211, 507  
**ccosf64x**, 630  
**ccosh**, 212, 417, 507, 592, 594, 595  
**ccoshf**, 212, 507  
**ccoshl**, 212, 507  
**ccosl**, 211, 507  
**ccospi**, 488  
**ceil**, 248, 274, 515, 559, 580, 581, 582  
**ceild128**, 35, 274, 520  
**ceild32**, 35, 274, 520  
**ceild64**, 35, 274, 520  
**ceilf**, 274, 515  
**ceilL**, 274, 278, 515  
**cerf**, 488  
**cerfc**, 488  
**cexp**, 213, 417, 507, 596, 597  
**cexp10**, 488  
**cexp10m1**, 488  
**cexp2**, 488  
**cexp2m1**, 488  
**cexpf**, 213, 507  
**cexpl**, 213, 507  
**cexpm1**, 488  
**cimag**, 215, 216, 508, 590–592, 628  
**cimagf**, 215, 508  
**cimagl**, 215, 419, 508

**clearerr**, 371, 372, 537  
**clgamma**, 488  
**clock**, 431, 432, 435, 547, 651  
**clog**, 213, 417, 507, 597  
**clog10**, 488  
**clog10p1**, 488  
**cloglp**, 488  
**clog2**, 488  
**clog2p1**, 488  
**clogf**, 213, 419, 507  
**clogl**, 213, 508  
**clogp1**, 488  
**cnd\_broadcast**, 422, 423, 424, 546  
**cnd\_destroy**, 423, 546  
**cnd\_init**, 423, 546  
**cnd\_signal**, 423, 424, 546  
**cnd\_timedwait**, 423, 424, 546  
**cnd\_wait**, 423, 424, 546  
**compoundn**, 269, 515, 564, 578  
**compoundnd128**, 269, 520  
**compoundnd32**, 269, 520  
**compoundnd64**, 269, 520  
**compoundnf**, 269, 515  
**compoundnl**, 269, 515  
**conj**, 215, 508, 592–597, 628  
**conjf**, 215, 508  
**conjl**, 215, 508  
**copysign**, 216, 279, 280, 516, 561, 577, 582, 583, 590, 591  
**copysignd128**, 280, 521  
**copysignd32**, 280, 521  
**copysignd64**, 280, 521  
**copysignf**, 280, 516  
**copysignl**, 280, 419, 516  
**cos**, 143, 256, 417, 513, 565, 573  
**cosd128**, 256, 518  
**cosd32**, 256, 518  
**cosd64**, 256, 518  
**cosf**, 256, 513  
**cosh**, 260, 417, 513, 565, 574  
**coshd128**, 260, 519  
**coshd32**, 260, 519  
**coshd64**, 260, 519  
**coshf**, 260, 513  
**coshl**, 260, 513  
**cosl**, 256, 513  
**cospi**, 258, 513, 565, 574  
**cospid128**, 258, 518  
**cospid32**, 258, 518  
**cospid64**, 258, 518  
**cospif**, 258, 513  
**cospil**, 258, 513  
**cpow**, 214, 417, 508, 597  
**cpowf**, 214, 508  
**cpowf128**, 630  
**cpowl**, 214, 419, 508  
**cpown**, 488  
**cpowr**, 488  
**cproj**, 215, 216, 508, 592, 628  
**cprojf**, 216, 419, 508  
**cprojl**, 216, 419, 508  
**creal**, 215, 216, 419, 508, 590–592, 628  
**crealf**, 216, 508  
**creall**, 216, 508  
**crootn**, 488  
**crsqrtr**, 488  
**csin**, 211, 417, 507, 592  
**csinf**, 211, 507  
**csinh**, 212, 213, 417, 507, 592, 595  
**csinhf**, 213, 507  
**csinhl**, 213, 507  
**csinl**, 211, 507  
**csinpi**, 488  
**csqrt**, 214, 417, 419, 508, 592, 597  
**csqrtrf**, 214, 508  
**csqrtrl**, 214, 508  
**ctan**, 211, 417, 507, 592  
**ctanf**, 211, 507  
**ctanh**, 213, 417, 507, 592, 596  
**ctanhf**, 213, 507  
**ctanhl**, 213, 507  
**ctanl**, 211, 507  
**ctanpi**, 488  
**ctgamma**, 488  
**ctime**, 435, 436, 547  
**ctime\_s**, 547, 700, 701  
**d32addd128**, 287, 522  
**d32addd64**, 287, 522  
**d32divd128**, 288, 522  
**d32divd64**, 288, 419, 522  
**d32fmad128**, 289, 522  
**d32fmad64**, 289, 522  
**d32muld128**, 288, 522  
**d32muld64**, 288, 522  
**d32sqrtd128**, 289, 522  
**d32sqrtd64**, 289, 522  
**d32subd128**, 288, 419, 522  
**d32subd64**, 288, 522  
**d64addd128**, 287, 420, 522  
**d64divd128**, 288, 522  
**d64fmad128**, 289, 419, 522  
**d64muld128**, 288, 522  
**d64sqrtd128**, 289, 522  
**d64subd128**, 288, 522  
**daddl**, 287, 419, 517, 560  
**ddivl**, 288, 517, 560, 564, 630  
**decodebind128**, 292, 523  
**decodebind32**, 292, 523  
**decodebind64**, 292, 416, 418, 523  
**decodedecd128**, 291, 523  
**decodedecd32**, 291, 523  
**decodedecd64**, 291, 416, 418, 523



**dfmal**, 289, 419, 517, 560  
**difftime**, 432, 433, 547  
**div**, 143, 373, 388, 389, 539  
**dmull**, 288, 517, 560  
**dsqrtl**, 289, 517, 560  
**dsubl**, 288, 517, 560, 630  
**encodebind128**, 292, 523  
**encodebind32**, 292, 523  
**encodebind64**, 292, 416, 418, 523  
**encodedecd128**, 291, 523  
**encodedecd32**, 291, 523  
**encodedecd64**, 291, 416, 418, 523  
**erf**, 272, 515, 580  
**erfc**, 273, 515, 580  
**erfcd128**, 273, 520  
**erfcd32**, 273, 520  
**erfcd64**, 273, 520  
**erfcf**, 273, 515  
**erfcl**, 273, 515  
**erfd128**, 272, 520  
**erfd32**, 272, 520  
**erfd64**, 272, 520  
**erff**, 272, 515  
**erfl**, 272, 515  
**exp**, 261, 262, 263, 315, 417, 419, 514, 564, 575  
**exp10**, 262, 514, 564, 575  
**exp10d128**, 262, 519  
**exp10d32**, 262, 519  
**exp10d64**, 262, 519  
**exp10f**, 262, 514  
**exp10l**, 262, 514  
**exp10m1**, 262, 514, 564, 575  
**exp10m1d128**, 262, 519  
**exp10m1d32**, 262, 519  
**exp10m1d64**, 262, 519  
**exp10m1f**, 262, 514  
**exp10m1l**, 262, 514  
**exp2**, 262, 263, 514, 564, 575  
**exp2d128**, 262, 519  
**exp2d32**, 262, 519  
**exp2d64**, 262, 519  
**exp2f**, 262, 514  
**exp2l**, 262, 514  
**exp2m1**, 263, 514, 564, 575  
**exp2m1d128**, 263, 519  
**exp2m1d32**, 263, 519  
**exp2m1d64**, 263, 519  
**exp2m1f**, 263, 514  
**exp2m1l**, 263, 514  
**expd128**, 261, 519  
**expd32**, 261, 519  
**expd64**, 261, 419, 519  
**expf**, 261, 514  
**expl**, 261, 514  
**expm1**, 263, 514, 564, 575  
**expm1d128**, 263, 519  
**expm1d32**, 263, 519  
**expm1d64**, 263, 519  
**expm1f**, 263, 514  
**expm1l**, 263, 514  
**f32addf64**, 630  
**f32addf64x**, 630  
**f32fmaf32x**, 630  
**f32mulf128**, 630  
**f32mulf32x**, 630  
**f32xsqrtf64x**, 630  
**f64divf128**, 630  
**f64divf64x**, 630  
**fabs**, 269, 270, 416, 417, 515, 561, 578, 582, 583, 591  
**fabsd128**, 269, 520  
**fabsd32**, 269, 520  
**fabsd64**, 269, 520  
**fabsf**, 269, 515  
**fabsl**, 269, 515  
**fadd**, 287, 517, 560, 629  
**faddl**, 287, 517, 560  
**fclose**, 186, 345, 537  
**fdim**, 282, 283, 516, 584  
**fdimd128**, 283, 522  
**fdimd32**, 283, 522  
**fdimd64**, 283, 522  
**fdimf**, 283, 516  
**fdiml**, 283, 517  
**fdiv**, 288, 419, 517, 560, 629  
**fdivl**, 288, 517, 560  
**fe\_dec\_getround**, 224, 231, 510, 564, 609  
**fe\_dec\_setround**, 224, 227, 232, 510, 564, 609  
**feclearexcept**, 99, 228, 230, 234, 509, 562, 582  
**fegetenv**, 233, 234, 510, 563, 639  
**fegetexceptflag**, 228, 229, 230, 510, 562, 639, 649  
**fegetmode**, 230, 231, 510, 562  
**fegetround**, 223, 227, 230, 231, 232, 510, 562, 563, 577, 581, 609  
**feholdexcept**, 98, 233, 234, 510, 563, 582, 639  
**feof**, 360, 366, 372, 461, 537  
**feraiseexcept**, 228, 229, 510, 568, 633, 649  
**ferror**, 360, 366, 372, 461, 537  
**fesetenv**, 225, 233, 510, 563, 639  
**fesetexcept**, 229, 510, 562  
**fesetexceptflag**, 228, 229, 510, 562, 639  
**fesetmode**, 225, 230, 231, 510, 563  
**fesetround**, 9, 25, 107, 223, 225, 227, 230, 232, 510, 562, 563, 577, 581, 582,

609

**fetestexcept**, 228, 230, 510, 562, 582

**fetestexceptflag**, 230, 510, 562

**feupdateenv**, 99, 225, 233, 234, 510, 563, 582, 639

**fflush**, 345, 346, 537, 641

**ffma**, 289, 517, 560, 629

**ffmal**, 289, 517, 560

**fgetc**, 340, 343, 366, 367, 369, 537

**fgetpos**, 341, 342, 370, 371, 537, 633, 643, 650

**fgets**, 340, 366, 367, 537, 642, 685

**fgetwc**, 340, 343, 461, 463, 548

**fgetws**, 340, 461, 462, 548, 642

**floor**, 248, 274, 515, 559, 581

**floord128**, 274, 520

**floord32**, 274, 520

**floord64**, 274, 520

**floorf**, 274, 515

**floorl**, 274, 515

**fma**, 249, 287, 517, 560, 585

**fmad128**, 287, 522

**fmad32**, 287, 522

**fmad64**, 287, 522

**fmaf**, 287, 517

**fmal**, 249, 287, 517

**fmax**, 283, 284, 517, 563, 584

**fmaxd128**, 283, 522

**fmaxd32**, 283, 522

**fmaxd64**, 283, 522

**fmaxf**, 283, 517

**fmaximum**, 284, 285, 517, 560, 584

**fmaximum\_mag**, 284, 285, 286, 517, 560, 584

**fmaximum\_mag\_num**, 285, 286, 517, 560, 584, 585

**fmaximum\_mag\_numd128**, 286, 522

**fmaximum\_mag\_numd32**, 286, 522

**fmaximum\_mag\_numd64**, 286, 522

**fmaximum\_mag\_numf**, 286, 517

**fmaximum\_mag\_numl**, 286, 517

**fmaximum\_magd128**, 284, 522

**fmaximum\_magd32**, 284, 522

**fmaximum\_magd64**, 284, 522

**fmaximum\_magf**, 284, 517

**fmaximum\_magl**, 284, 517

**fmaximum\_num**, 284, 285, 517, 560, 563, 584, 585, 591

**fmaximum\_numd128**, 285, 522

**fmaximum\_numd32**, 285, 522

**fmaximum\_numd64**, 285, 522

**fmaximum\_numf**, 285, 517

**fmaximum\_numl**, 285, 517

**fmaximumd128**, 284, 419, 522

**fmaximumd32**, 284, 522

**fmaximumd64**, 284, 522

**fmaximumf**, 284, 517

**fmaximuml**, 284, 517

**fmin**, 283, 284, 517, 563, 584

**fmind128**, 283, 522

**fmind32**, 283, 522

**fmind64**, 283, 522

**fminf**, 283, 517

**fminimum**, 284, 285, 286, 517, 560, 584

**fminimum\_mag**, 285, 286, 517, 560, 584

**fminimum\_mag\_num**, 285, 286, 517, 560, 584, 585

**fminimum\_mag\_numd128**, 286, 522

**fminimum\_mag\_numd32**, 286, 522

**fminimum\_mag\_numd64**, 286, 522

**fminimum\_mag\_numf**, 286, 517

**fminimum\_mag\_numl**, 286, 517

**fminimum\_magd128**, 285, 522

**fminimum\_magd32**, 285, 522

**fminimum\_magd64**, 285, 522

**fminimum\_magf**, 285, 517

**fminimum\_magl**, 285, 517

**fminimum\_num**, 284, 285, 286, 517, 560, 563, 584, 585

**fminimum\_numd128**, 286, 522

**fminimum\_numd32**, 286, 522

**fminimum\_numd64**, 286, 522

**fminimum\_numf**, 285, 517

**fminimum\_numl**, 286, 517

**fminimumd128**, 284, 522

**fminimumd32**, 284, 522

**fminimumd64**, 284, 522

**fminimumf**, 284, 517

**fminimuml**, 284, 517

**fminl**, 283, 517

**fmod**, 278, 279, 516, 582, 583, 649

**fmodd128**, 278, 521

**fmodd32**, 278, 521

**fmodd64**, 278, 521

**fmodf**, 278, 516

**fmodl**, 278, 516

**fmul**, 288, 517, 560, 629

**fmull**, 288, 517, 560

**fopen**, i, 186, 343, 344, 345, 346, 347, 537, 641, 653, 675, 676, 720

**fopen\_s**, 538, 674, 675, 676

**fprintf\_s**, 538, 677

**fputc**, 21, 340, 341, 343, 367, 368, 370, 537

**fputs**, 196, 340, 367, 537

**fputwc**, 340, 341, 343, 462, 463, 548

**fputws**, 340, 462, 548

**fread**, 184, 186, 340, 369, 537, 642

**free**, 382, 383, 384, 408, 409, 539, 643, 719

**free\_aligned\_sized**, 383, 539, 643, 719, 723  
**free\_sized**, 382, 383, 539, 643, 719, 723  
**freopen**, 341, 342, 347, 537  
**freopen\_s**, 538, 676  
**frexp**, 263, 264, 514, 575, 622, 633, 634  
**frexpd128**, 263, 519  
**frexpd32**, 263, 519  
**frexpd64**, 263, 519  
**frexpf**, 263, 514  
**frexpl**, 263, 514  
**fromfp**, 248, 277, 278, 516, 561, 565, 582  
**fromfpd128**, 277, 521  
**fromfpd32**, 277, 521  
**fromfpd64**, 277, 521  
**fromfpx**, 277, 516  
**fromfpl**, 277, 516  
**fromfpx**, 248, 278, 516, 561, 565, 582  
**fromfpxd128**, 278, 521  
**fromfpxd32**, 278, 521  
**fromfpxd64**, 278, 521  
**fromfpxf**, 278, 516  
**fromfpxl**, 278, 516  
**fscanf**, 236, 237, 340, 356, 359–364, 489, 537, 650, 677  
**fscanf\_s**, 538, 677, 678–680, 682  
**fseek**, 340, 343, 346, 369, 370, 371, 464, 537, 643  
**fsetpos**, 341, 342, 346, 369, 370, 371, 464, 537, 643, 650  
**fsqrt**, 289, 517, 560, 629  
**fsqrtl**, 289, 517, 560  
**fsub**, 287, 517, 560, 629  
**fsubl**, 288, 419, 517, 560, 630  
**ftell**, 370, 371, 537, 633, 643, 650  
**fwide**, 341, 342, 462, 463, 548  
**fwprintf**, 236, 340, 447, 453, 457–459, 461, 490, 548, 644, 650, 703  
**fwprintf\_s**, 550, 702, 703  
**fwrite**, 340, 369, 370, 537, 642  
**fwscanf**, 236, 340, 453, 454, 457–459, 461, 464, 490, 548, 650, 703  
**fwscanf\_s**, 550, 703, 705, 706, 710  
**getc**, 340, 367, 368, 537  
**getchar**, 16, 340, 367, 368, 537  
**getenv**, 385, 386, 539, 643, 645  
**getenv\_s**, 540, 687, 688  
**getpayload**, 523, 560, 586, 587  
**getpayloadd128**, 524, 586  
**getpayloadd32**, 524, 586  
**getpayloadd64**, 524, 586  
**getpayloadf**, 523, 586  
**getpayloadl**, 523, 586  
**gets\_s**, 538, 685  
**getwc**, 340, 463, 548  
**getwchar**, 340, 463, 548  
**gmtime**, 435, 436, 437, 438, 547, 651  
**gmtime\_r**, 437, 547, 721  
**gmtime\_s**, 547, 701, 702  
**hypot**, 270, 515, 559, 564, 578, 592  
**hypotd128**, 270, 520  
**hypotd32**, 270, 520  
**hypotd64**, 270, 520  
**hypotf**, 270, 515  
**hypotl**, 270, 515  
**ignore\_handler\_s**, 540, 687  
**ilogb**, 249, 264, 514, 560, 576, 633  
**ilogbd128**, 264, 519  
**ilogbd32**, 264, 519  
**ilogbd64**, 264, 519  
**ilogbf**, 264, 514  
**ilogbl**, 264, 514  
**imaxabs**, 237, 511  
**imaxdiv**, 236, 237, 511  
**isalnum**, 217, 219, 509  
**isalpha**, 217, 482, 509, 652  
**isblank**, 217, 218, 509, 652  
**iscntrl**, 217, 218, 219, 509  
**isdigit**, 217, 218, 219, 242, 509  
**isgraph**, 218, 482, 509  
**islower**, 4, 217, 218, 219, 220, 509, 652  
**isprint**, 21, 218, 509  
**ispunct**, 217, 218, 219, 509, 652  
**isspace**, 204, 217, 218, 219, 509, 652  
**isupper**, 217, 219, 220, 509, 652  
**iswalnum**, 483, 484, 485, 550  
**iswalpha**, 482, 483, 485, 550, 652  
**iswblank**, 483, 485, 550, 652  
**iswcntrl**, 483, 484, 485, 551  
**iswctype**, 485, 486, 551, 644, 652  
**iswdigit**, 483, 484, 485, 551  
**iswgraph**, 482, 484, 485, 551  
**iswlower**, 483, 484, 485, 486, 551, 652  
**iswprint**, 482, 484, 485, 551  
**iswpunct**, 482, 483, 484, 485, 551, 652  
**iswspace**, 204, 482, 483, 484, 485, 551, 652  
**iswupper**, 483, 484, 485, 486, 551, 652  
**iswxdigit**, 485, 551  
**isxdigit**, 219, 242, 509  
**labs**, 388, 539  
**ldexp**, 264, 265, 514, 576, 622  
**ldexpd128**, 264, 519  
**ldexpd32**, 264, 519  
**ldexpd64**, 264, 519  
**ldexpf**, 264, 514  
**ldexpl**, 264, 514  
**ldiv**, 143, 373, 388, 389, 539  
**lgamma**, 273, 515, 580  
**lgammad128**, 273, 520  
**lgammad32**, 273, 520  
**lgammad64**, 273, 520



**lgammaf**, 273, 515  
**lgammal**, 273, 515  
**llabs**, 388, 539  
**lldiv**, 143, 373, 388, 389, 539  
**llogb**, 250, 265, 514, 560, 576  
**llogbd128**, 265, 519  
**llogbd32**, 265, 519  
**llogbd64**, 265, 519  
**llogbf**, 265, 514  
**llogbl**, 265, 514  
**llquantexpd128**, 291, 523  
**llquantexpd32**, 291, 523  
**llquantexpd64**, 291, 523  
**llrint**, 275, 515, 565, 581, 582, 634  
**llrintd128**, 275, 521  
**llrintd32**, 275, 521  
**llrintd64**, 275, 521  
**llrintf**, 275, 516  
**llrintl**, 275, 516  
**llround**, 276, 516, 561, 582, 634  
**llroundd128**, 276, 521  
**llroundd32**, 276, 521  
**llroundd64**, 276, 521  
**llroundf**, 276, 516  
**llroundl**, 276, 516  
**localeconv**, 242, 245, 512, 640  
**localtime**, 433, 435, 436, 437, 438, 547, 651  
**localtime\_r**, 437, 547, 721  
**localtime\_s**, 547, 701, 702  
**log**, 250, 265, 266, 417, 514, 564, 576  
**log10**, 265, 266, 514, 564, 576  
**log10d128**, 266, 519  
**log10d32**, 265, 519  
**log10d64**, 265, 519  
**log10f**, 265, 514  
**log10l**, 265, 514  
**log10p1**, 266, 514, 564, 576  
**log10pd128**, 266, 519  
**log10pd32**, 266, 519  
**log10pd64**, 266, 519  
**log10p1f**, 266, 514  
**log10p1l**, 266, 514  
**log1p**, 266, 267, 514, 564, 576  
**log1pd128**, 266, 519  
**log1pd32**, 266, 519  
**log1pd64**, 266, 519  
**log1pf**, 266, 514  
**log1pl**, 266, 514  
**log2**, 267, 514, 564, 576  
**log2d128**, 267, 519  
**log2d32**, 267, 519  
**log2d64**, 267, 519  
**log2f**, 267, 514  
**log2l**, 267, 514  
**log2p1**, 266, 267, 514, 564, 577  
**log2pd128**, 267, 520  
**log2pd32**, 267, 519  
**log2pd64**, 267, 519  
**log2p1f**, 267, 514  
**log2p1l**, 267, 514  
**logb**, 264, 265, 267, 268, 514, 560, 575, 577, 591, 622  
**logbd128**, 267, 520  
**logbd32**, 267, 520  
**logbd64**, 267, 520  
**logbf**, 267, 514  
**logbl**, 267, 514  
**logd128**, 265, 519  
**logd32**, 265, 519  
**logd64**, 265, 519  
**logf**, 265, 514  
**logl**, 265, 514  
**logp1**, 266, 267, 514, 564, 576  
**logpd128**, 266, 519  
**logpd32**, 266, 519  
**logpd64**, 266, 519  
**logp1f**, 266, 514  
**logp1l**, 266, 514  
**longjmp**, 296, 297, 385, 386, 532, 640, 643, 719  
**lrint**, 275, 515, 565, 581, 582, 634  
**lrintd128**, 275, 521  
**lrintd32**, 275, 521  
**lrintd64**, 275, 521  
**lrintf**, 275, 515  
**lrintl**, 275, 515  
**lround**, 276, 516, 561, 582, 634  
**lroundd128**, 276, 521  
**lroundd32**, 276, 521  
**lroundd64**, 276, 521  
**lroundf**, 276, 516  
**lroundl**, 276, 516  
**mblen**, 389, 478, 539  
**mbrlen**, 478, 549  
**mbrtoc16**, 442, 443, 548  
**mbrtoc32**, 444, 548  
**mbrtoc8**, 441, 442, 548, 722  
**mbrtowc**, 343, 359, 452, 453, 477, 478, 479, 480, 549, 691, 717  
**mbsinit**, 402–404, 477, 478, 549  
**mbsrtowcs**, 477, 480, 549, 716  
**mbsrtowcs\_s**, 550, 716, 717  
**mbstowcs**, 72, 390, 391, 464, 465, 479, 539  
**mbstowcs\_s**, 540, 691, 692  
**mbtowc**, 70, 389, 390, 391, 478, 539  
**memalignment**, 9, 391, 539, 723  
**memccpy**, 407, 544, 721  
**memcmp**, 43, 186, 315, 409, 410, 544  
**memcpy**, 43, 76, 183, 207, 308, 315, 407, 544, 561

**memcpy\_s**, 545, 693  
**memmove**, 76, 407, 408, 544, 561, 638  
**memmove\_s**, 545, 693, 694  
**memset**, 308, 414, 545, 698  
**memset\_explicit**, 414, 545, 723  
**memset\_s**, 545, 698  
**mktime**, 433, 547  
**modf**, 268, 416, 417, 514, 577  
**modfd128**, 268, 520  
**modfd32**, 268, 520  
**modfd64**, 268, 520  
**modff**, 268, 514  
**modfl**, 268, 514  
**mtx\_destroy**, 424, 546  
**mtx\_init**, 421, 422, 424, 425, 546  
**mtx\_lock**, 425, 546, 644  
**mtx\_timedlock**, 425, 546, 644  
**mtx\_trylock**, 425, 426, 546  
**mtx\_unlock**, 425, 426, 546, 644  
**nan**, 280, 516, 559, 583  
**nand128**, 280, 521  
**nand32**, 280, 521  
**nand64**, 280, 521  
**nanf**, 280, 516  
**nanl**, 280, 516  
**nearbyint**, 274, 275, 515, 563, 565, 577, 581  
**nearbyintd128**, 274, 521  
**nearbyintd32**, 274, 521  
**nearbyintd64**, 274, 521  
**nearbyintf**, 274, 515  
**nearbyintl**, 274, 515  
**nextafter**, 280, 281, 419, 516, 563, 583  
**nextafterd128**, 280, 521  
**nextafterd32**, 280, 521  
**nextafterd64**, 280, 521  
**nextafterf**, 280, 516  
**nextafterl**, 280, 516  
**nextdown**, 282, 516, 560, 584  
**nextdownd128**, 282, 522  
**nextdownd32**, 282, 521  
**nextdownd64**, 282, 522  
**nextdownf**, 282, 516  
**nextdownl**, 282, 516  
**nexttoward**, 281, 516, 563, 583  
**nexttowardd128**, 281, 521  
**nexttowardd32**, 281, 521  
**nexttowardd64**, 281, 521  
**nexttowardf**, 281, 419, 516  
**nexttowardl**, 281, 516  
**nextup**, 281, 516, 560, 584  
**nextupd128**, 281, 521  
**nextupd32**, 281, 521  
**nextupd64**, 281, 521  
**nextupf**, 281, 516  
**nextupl**, 281, 516  
**perror**, 372, 537  
**pow**, 270, 417, 515, 565, 578, 630  
**powd128**, 270, 520  
**powd32**, 270, 520  
**powd64**, 270, 419, 520  
**powf**, 270, 515  
**powf32x**, 630  
**powf64**, 630  
**powl**, 270, 515, 563  
**pown**, 270, 271, 515, 565, 579  
**pownd128**, 270, 520  
**pownd32**, 270, 520  
**pownd64**, 270, 520  
**pownf**, 270, 515  
**pownl**, 270, 515  
**powr**, 271, 515, 565, 579  
**powrd128**, 271, 520  
**powrd32**, 271, 520  
**powrd64**, 271, 520  
**powrf**, 271, 515  
**powrl**, 271, 515  
**printf\_s**, 538, 678, 679  
**puts**, 197, 333, 340, 368, 537  
**putwc**, 340, 463, 464, 548  
**putwchar**, 340, 463, 464, 548  
**qsort**, 387, 388, 539, 634  
**qsort\_s**, 540, 688, 689, 690  
**quantized128**, 289, 523  
**quantized32**, 289, 522  
**quantized64**, 289, 522  
**quantumd128**, 290, 523  
**quantumd32**, 290, 523  
**quantumd64**, 290, 523  
**quick\_exit**, 299, 384, 385, 386, 539, 634, 640, 643, 650, 724  
**raise**, 298, 299, 300, 308, 384, 532, 640  
**rand**, 373, 380, 381, 539  
**realloc**, 381, 382, 383, 384, 539, 634, 643, 650, 719, 723  
**remainder**, 279, 419, 516, 560, 564, 583, 649  
**remainderd128**, 279, 521  
**remainderd32**, 279, 521  
**remainderd64**, 279, 521  
**remainderf**, 279, 516  
**remainderl**, 279, 516  
**remove**, 343, 344, 537, 650, 674  
**remquo**, 279, 516, 560, 564, 583, 633, 649  
**remquoof**, 279, 516  
**remquol**, 279, 516  
**rename**, 344, 537, 650  
**rewind**, 346, 369, 371, 464, 537  
**rint**, 275, 515, 560, 565, 581, 582  
**rintd128**, 275, 521  
**rintd32**, 275, 521  
**rintd64**, 275, 521

**rintf**, 275, 515  
**rintl**, 275, 515  
**rootn**, 271, 515, 564, 579  
**rootnd128**, 271, 520  
**rootnd32**, 271, 520  
**rootnd64**, 271, 520  
**rootnf**, 271, 515  
**rootnl**, 271, 515  
**round**, 248, 275, 276, 516, 559, 581  
**roundd128**, 276, 521  
**roundd32**, 275, 521  
**roundd64**, 276, 521  
**roundeven**, 248, 276, 277, 516, 559, 582  
**roundevend128**, 276, 521  
**roundevend32**, 276, 521  
**roundevend64**, 276, 521  
**roundevenf**, 276, 516  
**roundevenl**, 276, 516  
**roundf**, 275, 516  
**roundl**, 275, 516  
**rsqrt**, 271, 272, 515, 564, 580  
**rsqrt128**, 272, 520  
**rsqrt32**, 272, 520  
**rsqrt64**, 272, 520  
**rsqrtf**, 271, 515  
**rsqrtl**, 272, 515  
**samequantumd128**, 290, 523  
**samequantumd32**, 290, 523  
**samequantumd64**, 290, 523  
**scalbln**, 268, 515, 560, 577, 622  
**scalblnd128**, 268, 520  
**scalblnd32**, 268, 520  
**scalblnd64**, 268, 520  
**scalblnf**, 268, 515  
**scalblnl**, 268, 515  
**scalbn**, 268, 514, 560, 575, 576, 577, 591, 622  
**scalbnd128**, 268, 520  
**scalbnd32**, 268, 520  
**scalbnd64**, 268, 520  
**scalbnf**, 268, 514  
**scalbnl**, 268, 514  
**scanf**, 34, 226, 228, 340, 362, 365, 537, 561, 724, 725  
**scanf\_s**, 538, 679, 683  
**set\_constraint\_handler\_s**, 540, 672, 686, 687  
**setbuf**, 339, 342, 343, 345, 347, 537  
**setjmp**, 206, 296, 297, 532, 633, 640, 719  
**setlocale**, viii, 204, 241, 242, 245, 435, 512, 639, 649  
**setpayload**, 523, 560, 587  
**setpayloadd128**, 524, 587  
**setpayloadd32**, 524, 587  
**setpayloadd64**, 524, 587  
**setpayloadf**, 523, 587  
**setpayloadl**, 523, 587  
**setpayloadsig**, 523, 560, 587  
**setpayloadsigd128**, 524, 587  
**setpayloadsigd32**, 524, 587  
**setpayloadsigd64**, 524, 587  
**setpayloadsigf**, 523, 587  
**setpayloadsigl**, 523, 587  
**setvbuf**, 339, 342, 343, 345, 347, 348, 537, 641  
**signal**, 15, 16, 141, 298, 299, 385, 386, 532, 640, 645, 649  
**sin**, 82, 256, 417, 419, 513, 565, 573  
**sind128**, 256, 518  
**sind32**, 256, 518  
**sind64**, 256, 518  
**sinf**, 256, 513  
**sinh**, 261, 417, 513, 565, 574  
**sinhd128**, 261, 519  
**sinhd32**, 261, 519  
**sinhd64**, 261, 519  
**sinhf**, 261, 513  
**sinhl**, 261, 513  
**sinl**, 256, 513  
**sinpi**, 258, 259, 513, 565, 574  
**sinpid128**, 259, 518  
**sinpid32**, 259, 518  
**sinpid64**, 259, 518  
**sinpif**, 258, 513  
**sinpil**, 259, 513  
**snprintf**, 363, 365, 374, 375, 436, 537, 679, 725  
**snprintf\_s**, 538, 679, 680  
**snwprintf\_s**, 550, 704, 705  
**sprintf**, 363, 365, 537, 680  
**sprintf\_s**, 538, 679, 680  
**sqrt**, 158, 272, 417, 515, 560, 580, 585  
**sqrt128**, 272, 520  
**sqrt32**, 272, 419, 520  
**sqrt64**, 272, 520  
**sqrtf**, 272, 515  
**sqrtl**, 272, 515  
**srand**, 381, 539  
**sscanf**, 361, 363, 366, 537  
**sscanf\_s**, 538, 680, 681, 684  
**stdc\_bit\_ceil\_uc**, 325, 535  
**stdc\_bit\_ceil\_ui**, 325, 535  
**stdc\_bit\_ceil\_ul**, 325, 535  
**stdc\_bit\_ceil\_ull**, 325, 535  
**stdc\_bit\_ceil\_us**, 325, 535  
**stdc\_bit\_floor\_uc**, 325, 535  
**stdc\_bit\_floor\_ui**, 325, 535  
**stdc\_bit\_floor\_ul**, 325, 535  
**stdc\_bit\_floor\_ull**, 325, 535  
**stdc\_bit\_floor\_us**, 325, 535  
**stdc\_bit\_width\_uc**, 324, 535  
**stdc\_bit\_width\_ui**, 324, 535

|                                                    |                                                      |
|----------------------------------------------------|------------------------------------------------------|
| <code>stdc_bit_width_ul</code> , 324, 535          | 534                                                  |
| <code>stdc_bit_width_ull</code> , 324, 535         | <code>stdc_first_leading_zero_ul</code> , 321, 534   |
| <code>stdc_bit_width_us</code> , 324, 535          | <code>stdc_first_leading_zero_ull</code> , 321, 534  |
| <code>stdc_c16nrtoc16n</code> , 397, 542           | <code>stdc_first_leading_zero_us</code> , 321, 534   |
| <code>stdc_c16nrtoc32n</code> , 397, 542           | <code>stdc_first_trailing_one_uc</code> , 322, 534   |
| <code>stdc_c16nrtoc8n</code> , 397, 542            | <code>stdc_first_trailing_one_ui</code> , 322, 534   |
| <code>stdc_c16nrtomcn</code> , 397, 542            | <code>stdc_first_trailing_one_ul</code> , 322, 534   |
| <code>stdc_c16nrtomwcn</code> , 397, 542           | <code>stdc_first_trailing_one_ull</code> , 322, 534  |
| <code>stdc_c16snrtoc16sn</code> , 401, 544         | <code>stdc_first_trailing_one_us</code> , 322, 534   |
| <code>stdc_c16snrtoc32sn</code> , 401, 544         | <code>stdc_first_trailing_zero_uc</code> , 322, 534  |
| <code>stdc_c16snrtoc8sn</code> , 401, 544          | <code>stdc_first_trailing_zero_ui</code> , 322, 534  |
| <code>stdc_c16snrtomcsn</code> , 401, 543          | <code>stdc_first_trailing_zero_ul</code> , 322, 534  |
| <code>stdc_c16snrtomwcn</code> , 401, 543          | <code>stdc_first_trailing_zero_ull</code> , 322, 534 |
| <code>stdc_c32nrtoc16n</code> , 397, 542           | <code>stdc_first_trailing_zero_us</code> , 322, 534  |
| <code>stdc_c32nrtoc32n</code> , 397, 542           | <code>stdc_has_single_bit_uc</code> , 324, 534       |
| <code>stdc_c32nrtoc8n</code> , 397, 542            | <code>stdc_has_single_bit_ui</code> , 324, 534       |
| <code>stdc_c32nrtomcn</code> , 397, 542            | <code>stdc_has_single_bit_ul</code> , 324, 534       |
| <code>stdc_c32nrtomwcn</code> , 397, 542           | <code>stdc_has_single_bit_us</code> , 324, 534       |
| <code>stdc_c32snrtoc16sn</code> , 401, 544         | <code>stdc_leading_ones_uc</code> , 319, 533         |
| <code>stdc_c32snrtoc32sn</code> , 401, 544         | <code>stdc_leading_ones_ui</code> , 319, 533         |
| <code>stdc_c32snrtoc8sn</code> , 401, 544          | <code>stdc_leading_ones_ul</code> , 319, 533         |
| <code>stdc_c32snrtomcsn</code> , 401, 544          | <code>stdc_leading_ones_ull</code> , 319, 533        |
| <code>stdc_c32snrtomwcn</code> , 401, 544          | <code>stdc_leading_ones_us</code> , 319, 533         |
| <code>stdc_c8nrtoc16n</code> , 397, 542            | <code>stdc_leading_zeros_uc</code> , 319, 533        |
| <code>stdc_c8nrtoc32n</code> , 397, 542            | <code>stdc_leading_zeros_ui</code> , 319, 533        |
| <code>stdc_c8nrtoc8n</code> , 397, 541             | <code>stdc_leading_zeros_ul</code> , 319, 533        |
| <code>stdc_c8nrtomcn</code> , 396, 541             | <code>stdc_leading_zeros_ull</code> , 319, 533       |
| <code>stdc_c8nrtomwcn</code> , 396, 541            | <code>stdc_leading_zeros_us</code> , 319, 533        |
| <code>stdc_c8snrtoc16sn</code> , 401, 404, 543     | <code>stdc_load8_aligned_bes</code> , 328, 535       |
| <code>stdc_c8snrtoc32sn</code> , 401, 543          | <code>stdc_load8_aligned_beu</code> , 328, 535       |
| <code>stdc_c8snrtoc8sn</code> , 401, 543           | <code>stdc_load8_aligned_les</code> , 328, 535       |
| <code>stdc_c8snrtomcsn</code> , 401, 543           | <code>stdc_load8_aligned_leu</code> , 328, 535       |
| <code>stdc_c8snrtomwcn</code> , 400, 543           | <code>stdc_load8_bes</code> , 328, 535               |
| <code>stdc_count_ones_uc</code> , 323, 534         | <code>stdc_load8_beu</code> , 328, 535               |
| <code>stdc_count_ones_ui</code> , 323, 534         | <code>stdc_load8_les</code> , 328, 535               |
| <code>stdc_count_ones_ul</code> , 323, 534         | <code>stdc_load8_leu</code> , 328, 535               |
| <code>stdc_count_ones_ull</code> , 323, 534        | <code>stdc_mcnrtoc16n</code> , 396, 541              |
| <code>stdc_count_ones_us</code> , 323, 534         | <code>stdc_mcnrtoc32n</code> , 396, 541              |
| <code>stdc_count_zeros_uc</code> , 323, 534        | <code>stdc_mcnrtoc8n</code> , 396, 403, 541          |
| <code>stdc_count_zeros_ui</code> , 323, 534        | <code>stdc_mcnrtomcn</code> , 396, 541               |
| <code>stdc_count_zeros_ul</code> , 323, 534        | <code>stdc_mcnrtomwcn</code> , 396, 541              |
| <code>stdc_count_zeros_ull</code> , 323, 534       | <code>stdc_mcsnrtoc16sn</code> , 400, 543            |
| <code>stdc_count_zeros_us</code> , 323, 534        |                                                      |
| <code>stdc_first_leading_one_uc</code> , 321, 534  |                                                      |
| <code>stdc_first_leading_one_ui</code> , 321, 534  |                                                      |
| <code>stdc_first_leading_one_ul</code> , 321, 534  |                                                      |
| <code>stdc_first_leading_one_ull</code> , 321, 534 |                                                      |
| <code>stdc_first_leading_one_us</code> , 321, 534  |                                                      |
| <code>stdc_first_leading_zero_uc</code> , 321, 534 |                                                      |
| <code>stdc_first_leading_zero_ui</code> , 321,     |                                                      |

**stdc\_mcsnrtoc32sn**, 400, 543  
**stdc\_mcsnrtoc8sn**, 400, 542  
**stdc\_mcsnrtomcsn**, 400, 542  
**stdc\_mcsnrtomwcn**, 400, 542  
**stdc\_memreverse8u**, 327, 535  
**stdc\_mwcnrtoc16n**, 396, 541  
**stdc\_mwcnrtoc32n**, 396, 541  
**stdc\_mwcnrtoc8n**, 396, 541  
**stdc\_mwcnrtomcn**, 396, 541  
**stdc\_mwcnrtomwcn**, 396, 541  
**stdc\_mwcsnrtoc16sn**, 400, 543  
**stdc\_mwcsnrtoc32sn**, 400, 543  
**stdc\_mwcsnrtoc8sn**, 400, 543  
**stdc\_mwcsnrtomcsn**, 400, 543  
**stdc\_mwcsnrtomwcn**, 400, 543  
**stdc\_rotate\_left\_uc**, 326, 535  
**stdc\_rotate\_left\_ui**, 326, 535  
**stdc\_rotate\_left\_ull**, 326, 535  
**stdc\_rotate\_left\_us**, 326, 535  
**stdc\_rotate\_right\_uc**, 326, 535  
**stdc\_rotate\_right\_ui**, 326, 535  
**stdc\_rotate\_right\_ull**, 326, 535  
**stdc\_rotate\_right\_us**, 326, 535  
**stdc\_store8\_aligned\_bes**, 329, 535  
**stdc\_store8\_aligned\_beu**, 329, 535  
**stdc\_store8\_aligned\_les**, 329, 535  
**stdc\_store8\_aligned\_leu**, 328, 535  
**stdc\_store8\_bes**, 329, 535  
**stdc\_store8\_beu**, 328, 535  
**stdc\_store8\_les**, 329, 535  
**stdc\_store8\_leu**, 328, 535  
**stdc\_trailing\_ones\_uc**, 320, 534  
**stdc\_trailing\_ones\_ui**, 320, 534  
**stdc\_trailing\_ones\_ull**, 320, 534  
**stdc\_trailing\_ones\_us**, 320, 534  
**stdc\_trailing\_zeros\_uc**, 320, 533  
**stdc\_trailing\_zeros\_ui**, 320, 533  
**stdc\_trailing\_zeros\_ull**, 320, 533  
**stdc\_trailing\_zeros\_us**, 320, 533  
**strcat**, 409, 544  
**strcat\_s**, 545, 695, 696  
**strcmp**, 409, 410, 411, 544  
**strcoll**, 9, 242, 410, 411, 544  
**strcpy**, 188–190, 408, 544  
**strcpy\_s**, 545, 694  
**strcspn**, 412, 544  
**strdup**, 9, 408, 544, 721  
**strerror**, 9, 372, 414, 415, 545, 643, 644, 652  
**strerror\_s**, 415, 545, 698, 699  
**strerrorlen\_s**, 545, 699  
**strfromd**, 374, 375, 464, 465, 538, 561  
**strfromd128**, 375, 539  
**strfromd32**, 375, 539  
**strfromd64**, 375, 539  
**strfromencf128**, 625, 626  
**strfromf**, 374, 375, 464, 538  
**strfroml**, 374, 375, 538  
**strftime**, ii, 242, 435, 437, 440, 476, 547, 634, 641–643, 651, 700, 701, 721, 726  
**strlen**, 409, 415, 545  
**strncat**, 409, 544  
**strncat\_s**, 545, 696, 697  
**strncmp**, 196, 409, 410, 411, 544  
**strncpy**, 408, 544  
**strncpy\_s**, 545, 694, 695  
**strndup**, 9, 409, 544, 721  
**strnlen**, ii, 415, 545, 721  
**strnlen\_s**, 545, 694–696, 699  
**strspn**, 413, 545  
**strtod**, 66, 280, 357, 358, 362, 374, 375, 464, 538, 561, 566, 567, 633, 650  
**strtod128**, 377, 539, 626, 650  
**strtod32**, 377, 539, 650  
**strtod64**, 377, 378, 539, 650  
**strtof**, 280, 362, 374, 375, 538, 633, 650  
**strtoimax**, 237, 238, 511  
**strtok**, 9, 413, 414, 545, 643  
**strtok\_s**, 414, 545, 697, 698  
**strtoul**, 238, 357, 358, 362, 374, 379, 380, 538  
**strtold**, 280, 362, 374, 375, 538, 633, 650  
**strtoll**, 238, 362, 374, 379, 380, 538  
**strtol**, 238, 358, 362, 374, 379, 380, 538  
**strtoull**, 238, 362, 374, 379, 380, 538  
**strtoimax**, 237, 238, 511  
**strxfrm**, 9, 242, 411, 544, 643  
**swprintf**, 458, 459, 548, 704, 705  
**swprintf\_s**, 550, 704, 705  
**swscanf**, 458, 460, 548  
**swscanf\_s**, 550, 705, 708  
**system**, 386, 387, 539, 643, 645, 650  
**tan**, 256, 257, 417, 513, 565, 573  
**tand128**, 256, 518  
**tand32**, 256, 518  
**tand64**, 256, 518  
**tanf**, 256, 513  
**tanh**, 261, 417, 513, 565, 575  
**tanhd128**, 261, 519  
**tanhd32**, 261, 519  
**tanhd64**, 261, 519  
**tanhf**, 261, 513  
**tanh1**, 261, 514  
**tanl**, 256, 513  
**tanpi**, 259, 513, 565, 574



**tanpid128**, 259, 518  
**tanpid32**, 259, 518  
**tanpid64**, 259, 518  
**tanpif**, 259, 513  
**tanpil**, 259, 513  
**tgamma**, 273, 274, 515, 580  
**tgamma128**, 273, 520  
**tgamma32**, 273, 520  
**tgamma64**, 273, 520  
**tgammaf**, 273, 515  
**tgammaL**, 273, 515  
**thrd\_create**, 421, 426, 546  
**thrd\_current**, 426, 546  
**thrd\_detach**, 427, 546, 644  
**thrd\_equal**, 427, 547  
**thrd\_exit**, 426, 427, 547, 634  
**thrd\_join**, 427, 428, 547, 644  
**thrd\_sleep**, 428, 547  
**thrd\_yield**, 428, 547  
**time**, 433, 434, 547, 634  
**timegm**, 433, 434, 547, 723  
**timespec\_get**, 431, 434, 435, 547  
**timespec\_getres**, 431, 435, 547, 723  
**tmpfile**, 344, 385, 537  
**tmpfile\_s**, 538, 674, 675  
**tmpnam**, 340, 344, 345, 537, 675  
**tmpnam\_s**, 538, 673, 674, 675  
**tolower**, 219, 509  
**totalorder**, 523, 562, 585, 586  
**totalorderd128**, 523, 585  
**totalorderd32**, 523, 585  
**totalorderd64**, 523, 585  
**totalorderf**, 523, 585  
**totalorderL**, 523, 585  
**totalordermag**, 523, 562, 586  
**totalordermagd128**, 524, 586  
**totalordermagd32**, 523, 586  
**totalordermagd64**, 524, 586  
**totalordermagf**, 523, 586  
**totalordermagL**, 523, 586  
**toupper**, 219, 220, 509  
**towctrans**, 486, 487, 551, 644, 652  
**tolower**, 486, 487, 551  
**towupper**, 486, 487, 551  
**trunc**, 248, 277, 516, 559, 582  
**truncd128**, 277, 521  
**truncd32**, 277, 521  
**truncd64**, 277, 521  
**truncf**, 277, 516  
**truncL**, 277, 516  
**tss\_create**, 428, 429, 547, 644  
**tss\_delete**, 429, 547, 634, 644  
**tss\_get**, 429, 547, 644  
**tss\_set**, 429, 430, 547, 644  
**uabs**, 388, 539  
**ufromfp**, 248, 277, 278, 516, 561, 565, 582  
**ufromfpd128**, 277, 521  
**ufromfpd32**, 277, 521  
**ufromfpd64**, 277, 521  
**ufromfpf**, 277, 516  
**ufromfpL**, 277, 516  
**ufromfpx**, 248, 278, 516, 561, 565, 582  
**ufromfpxd128**, 278, 521  
**ufromfpxd32**, 278, 521  
**ufromfpxd64**, 278, 521  
**ufromfpxf**, 278, 516  
**ufromfpxL**, 278, 516  
**uimaxabs**, 237, 511, 721  
**ulabs**, 388, 539  
**ullabs**, 388, 539  
**ungetc**, 340, 368, 369–371, 489, 537, 633, 642, 653, 726  
**ungetwc**, 340, 464, 548, 633, 653  
**va\_arg**, 236, 302, 303–305, 353, 364–366, 452, 458–460, 532, 640, 641, 682–684, 706, 708, 709  
**va\_copy**, 206, 302, 303, 305, 532, 633, 640, 641, 726  
**va\_end**, 206, 302, 303, 304, 305, 364–366, 458–460, 532, 633, 640–642, 682–684, 706, 708, 709  
**va\_start**, 302, 303, 304, 305, 364–366, 458–460, 532, 640, 641, 682–684, 706, 708, 709, 723  
**vfprintf**, 340, 363, 364, 537, 642, 681  
**vfprintf\_s**, 538, 681, 682–684  
**vfscanf**, 340, 364, 537, 642  
**vfscanf\_s**, 538, 682, 683, 684  
**vfwprintf**, 340, 458, 459, 548, 642, 706  
**vfwprintf\_s**, 550, 705, 706  
**vfwscanf**, 340, 458, 459, 464, 548, 642  
**vfwscanf\_s**, 550, 706, 708, 709  
**vprintf**, 340, 364, 365, 537, 642, 682  
**vprintf\_s**, 538, 682, 683, 684  
**vscanf**, 340, 364, 365, 537, 642, 725  
**vscanf\_s**, 538, 682, 683–685  
**vsprintf**, 364, 365, 537, 642, 683  
**vsprintf\_s**, 538, 682, 683, 684  
**vsnprintf**, 364, 365, 537, 642, 683  
**vsnprintf\_s**, 538, 682, 683, 684  
**vsnwprintf**, 550, 706, 707  
**vsprintf**, 364, 365, 366, 537, 642, 684  
**vsprintf\_s**, 538, 682, 683, 684  
**vsscanf**, 364, 366, 537, 642  
**vsscanf\_s**, 538, 682, 683, 684, 685  
**vswprintf**, 458, 459, 548, 642, 707  
**vswprintf\_s**, 550, 707  
**vswscanf**, 458, 460, 548, 642  
**vswscanf\_s**, 550, 706, 707, 708, 709  
**vwprintf**, 340, 458, 460, 548, 642, 708  
**vwprintf\_s**, 550, 708  
**vwscanf**, 340, 458, 460, 464, 548, 642

**vwscanf\_s**, 550, 706, **708**, 709  
**wcrtomb**, 343, 352, 353, 356, 362, 447, 456, 458, **479**, 481, 549, 634, 692, 715, 718  
**wcrtomb\_s**, 550, **715**, 716  
**wcscat**, **470**, 471, 549  
**wcscat\_s**, 550, **712**, 713  
**wscmp**, **471**, 472, 549  
**wscoll**, **471**, 472, 549  
**wscpy**, **469**, 549  
**wscpy\_s**, 550, **710**  
**wscspn**, **473**, 549  
**wcsftime**, 242, **476**, 549, 634, 641–643, 651  
**wcslen**, 471, **475**, 549, 715  
**wcsncat**, **471**, 549  
**wcsncat\_s**, 550, **713**, 714  
**wcsncmp**, **471**, 472, 549  
**wcsncpy**, **469**, 470, 549  
**wcsncpy\_s**, 550, **710**, 711  
**wcsnlen**, ii, **475**, 476, 549, 721  
**wcsnlen\_s**, 550, 710–713, **715**  
**wcsrtombs**, **480**, 481, 549, 716  
**wcsrtombs\_s**, 550, 716, **717**, 718  
**wcsspn**, **474**, 549  
**wcstod**, 454, 456, 458, **465**, 548, 561, 633, 650  
**wcstod128**, **465**, 467, 549, 650  
**wcstod32**, **465**, 467, 549, 650  
**wcstod64**, **465**, 467, 549, 650  
**wcstof**, 458, **465**, 548, 633, 650  
**wcstoimax**, **238**, 511  
**wcstok**, **474**, 475, 549, 643, 644  
**wcstok\_s**, 550, **714**, 715  
**wcstol**, 238, 454, 455, 458, **468**, 469, 548  
**wcstold**, 458, **465**, 548, 633, 650  
**wcstoll**, 238, 458, **468**, 469, 548  
**wcstombs**, **391**, 479, 539  
**wcstombs\_s**, 540, **692**, 693  
**wcstoul**, 238, 455, 456, 458, **468**, 469, 548  
**wcstoull**, 238, 458, **468**, 469, 549  
**wcstoumax**, **238**, 511  
**wcsxfrm**, 472, 549, 643  
**wctob**, 477, 482, 549  
**wctomb**, 389, **390**, 391, 478, 539  
**wctomb\_s**, 540, **690**, 691  
**wctrans**, 486, **487**, 551, 644  
**wctype**, 485, 486, 551, 644  
**wmemcmp**, **472**, 549  
**wmemcpy**, **470**, 549  
**wmemcpy\_s**, 550, **711**, 712  
**wmemmove**, **470**, 549  
**wmemmove\_s**, 550, **712**  
**wmemset**, **476**, 549  
**wprintf**, 226, 228, 237, 340, **460**, 461, 548, 561, 709  
**wprintf\_s**, 550, **709**  
**wscanf**, 226, 228, 340, 460, **461**, 464, 548, 561  
**wscanf\_s**, 550, **709**, 710  
future direction, **488**  
identifier  
  **\_\_VA\_ARGS\_\_**, **190**, 191, **192**, 725  
  **\_\_VA\_OPT\_\_**, **190**, 191–193, 723  
  **\_\_func\_\_**, **57**, 58, 208, 636, 726  
**errno**, 47, 156, 157, 206, 209, **221**, 238, 250, 251, 299, 343, 370–373, 377, 378, 380, 415, 442–445, 461, 462, 466, 468, 469, 479–481, 509, 639, 640, 649, 650, 654, 672, 673, 698  
identifier prefix  
  **\_\_STDC\_\_**, **203**  
  **\_\_STDC\_VERSION\_\_**, **205**  
**ATOMIC\_\_**, 489  
**atomic\_\_**, 489  
**c**, 416, 417  
**ckd\_\_**, 489  
**cnd\_\_**, 490  
**cr\_\_**, 489, 564, 572  
**d**, 418  
**d128**, 418  
**d32**, 418  
**d64**, 416, 418  
**DBL\_\_**, 488  
**DEC**, 511  
**DEC128\_\_**, 30, 488  
**DEC32\_\_**, 30, 488  
**DEC64\_\_**, 30, 488  
**DEC\_\_**, 488  
**E**, 221, 488  
**FE\_\_**, 223, **224**, 488  
**FLT**, 511  
**FLT\_\_**, 488  
**FP\_\_**, **248**, 488  
**INT**, **337**, **338**, 489, 536  
**int**, 335, 336, 489, 536, 724  
**INT\_FAST**, **337**, 536  
**int\_fast**, **336**, 536, 724  
**INT\_LEAST**, **337**, 536  
**int\_least**, 328, 329, **335**, **336**, 338, 535, 536  
**is**, 488–490  
**LC\_\_**, **241**, 488  
**LDBL\_\_**, 488  
**llquantexpd**, **290**, 291, 418  
**MATH\_\_**, 489  
**mem**, 490  
**memory\_\_**, 489  
**memory\_order\_\_**, 489  
**mtx\_\_**, 490

- PRI**, 236, 488  
**quantized**, 289, 290, 418  
**quantumd**, 290, 418  
**samequantumd**, 290, 418  
**SCN**, 236, 488  
**SIG**, 298, 489  
**SIG\_**, 298, 489  
**stdc\_**, 489  
**str**, 488–490  
**thrd\_**, 490  
**TIME\_**, 431, 490  
**to**, 488, 490  
**tss\_**, 490  
**UINT**, 337, 338, 489, 536  
**uint**, 327, 335, 336, 489, 535, 536, 724  
**UINT\_FAST**, 337, 536  
**uint\_fast**, 336, 536, 724  
**UINT\_LEAST**, 337, 536  
**uint\_least**, 328, 329, 335, 336, 338, 535, 536  
**wcs**, 488–490  
 identifier suffix  
**\_DECIMAL\_DIG**, 29, 353, 376, 452, 466, 565, 566  
**\_r**, 435  
**\_C**, 338, 489, 536  
**\_H\_**, 205  
**\_MAX**, 24, 49, 337, 338, 489, 511, 536, 603  
**\_MIN**, 24, 337, 338, 489, 511, 536, 604  
**\_WIDTH**, 24, 337, 338, 489, 536  
**\_explicit**, 307, 315, 533  
**\_t**, 327–329, 335, 336, 338, 489, 524, 535, 536, 610, 611, 724  
**f**, 209, 247, 416, 418, 488  
**l**, 209, 247, 416, 418, 488  
 macro  
**\_Complex\_I**, 209, 507, 723  
**\_IOFBF**, 339, 347, 348, 537  
**\_IOLBF**, 339, 348, 537  
**\_IONBF**, 339, 347, 348, 537  
**\_PRINTF\_NAN\_LEN\_MAX**, 340, 537  
**\_DATE\_**, 199, 648  
**\_FILE\_**, 181, 199, 208  
**\_LINE\_**, 197, 198, 199, 208, 633  
**\_STDC\_ANALYZABLE\_**, 201, 719  
**\_STDC\_EMBED\_EMPTY\_**, 179, 199  
**\_STDC\_EMBED\_FOUND\_**, 179, 199  
**\_STDC\_EMBED\_NOT\_FOUND\_**, 178, 199  
**\_STDC\_ENDIAN\_BIG\_**, 318, 319, 533  
**\_STDC\_ENDIAN\_LITTLE\_**, 318, 319, 394, 533  
**\_STDC\_ENDIAN\_NATIVE\_**, 291, 292, 318, 394, 533, 622–624, 626–628, 649  
**\_STDC\_HOSTED\_**, 199  
**\_STDC\_IEC\_60559\_BFP\_**, 9, 24, 201, 204, 523, 558, 585–587, 598–600  
**\_STDC\_IEC\_60559\_COMPLEX\_**, 24, 201, 204, 589  
**\_STDC\_IEC\_60559\_DFP\_**, 9, 30, 201, 204, 227, 231, 232, 247, 254–281, 282, 283–287, 288, 289–292, 375, 377, 419, 467, 510, 511, 517, 523, 539, 549, 556–558, 585–587, 598–600  
**\_STDC\_IEC\_60559\_TYPES\_**, 201, 204, 508, 511, 524, 539, 549, 598–600  
**\_STDC\_ISO\_10646\_**, 200, 646  
**\_STDC\_LIB\_EXT1\_**, 201, 204, 509, 536, 538, 540, 545, 547, 550, 671  
**\_STDC\_LITERAL\_UTF16\_**, 200  
**\_STDC\_LITERAL\_UTF32\_**, 200  
**\_STDC\_LITERAL\_UTF8\_**, 200  
**\_STDC\_MB\_MIGHT\_NEQ\_WC\_**, 47, 200, 332  
**\_STDC\_NO\_ATOMICS\_**, 201, 307, 532  
**\_STDC\_NO\_COMPLEX\_**, 65, 201, 209, 507  
**\_STDC\_NO\_THREADS\_**, 16, 201, 421, 546  
**\_STDC\_NO\_VLA\_**, 201  
**\_STDC\_UTF\_16\_**, 200  
**\_STDC\_UTF\_32\_**, 200  
**\_STDC\_VERSION\_ASSERT\_H\_**, 208, 507  
**\_STDC\_VERSION\_COMPLEX\_H\_**, 209, 507  
**\_STDC\_VERSION\_FENV\_H\_**, 222, 509  
**\_STDC\_VERSION\_FLOAT\_H\_**, 235, 510  
**\_STDC\_VERSION\_INTTYPES\_H\_**, 236, 511  
**\_STDC\_VERSION\_LIMITS\_H\_**, 240, 512  
**\_STDC\_VERSION\_MATH\_H\_**, 247, 512  
**\_STDC\_VERSION\_SETJMP\_H\_**, 296, 532  
**\_STDC\_VERSION\_STDARG\_H\_**, 302, 532  
**\_STDC\_VERSION\_STDATOMIC\_H\_**, 307, 532  
**\_STDC\_VERSION\_STDBIT\_H\_**, 318, 533  
**\_STDC\_VERSION\_STDCCKDINT\_H\_**, 331, 536  
**\_STDC\_VERSION\_STDDEF\_H\_**, 332, 536  
**\_STDC\_VERSION\_STDINT\_H\_**, 335, 536  
**\_STDC\_VERSION\_STDIO\_H\_**, 339, 537  
**\_STDC\_VERSION\_STDLIB\_H\_**, 373, 538  
**\_STDC\_VERSION\_STRING\_H\_**, 407, 544  
**\_STDC\_VERSION\_TGMATH\_H\_**, 416



**\_\_STDC\_VERSION\_TIME\_H\_\_**, 431, 547  
**\_\_STDC\_VERSION\_UCHAR\_H\_\_**, 441, 548  
**\_\_STDC\_VERSION\_WCHAR\_H\_\_**, 446, 548  
**\_\_STDC\_VERSION\_\_**, 200, 720–722, 724, 726  
**\_\_STDC\_WANT\_IEC\_60559\_EXT\_\_**, 510, 523, 565, 585–587  
**\_\_STDC\_WANT\_IEC\_60559\_TYPES\_EXT\_\_**, 508, 511, 524, 539, 549, 602, 607, 622–624, 625, 626–628, 630  
**\_\_STDC\_WANT\_LIB\_EXT1\_\_**, 509, 536, 538, 540, 545, 547, 550, 671, 672, 674–690, 693–715  
**\_\_STDC\_WIDE\_LITERAL\_UTF16\_\_**, 200  
**\_\_STDC\_WIDE\_LITERAL\_UTF32\_\_**, 201  
**\_\_STDC\_WIDE\_LITERAL\_UTF8\_\_**, 200  
**\_\_STDC\_\_**, 180, 199  
**\_\_TIME\_\_**, 200, 648  
**\_\_cplusplus**, 180, 199  
**and**, 239, 512  
**and\_eq**, 239, 512  
**assert**, 152, 187, 189, 208, 232, 507, 639, 649  
**ATOMIC\_BOOL\_LOCK\_FREE**, 307, 532  
**ATOMIC\_CHAR16\_T\_LOCK\_FREE**, 307, 532  
**ATOMIC\_CHAR32\_T\_LOCK\_FREE**, 307, 532  
**ATOMIC\_CHAR8\_T\_LOCK\_FREE**, 307, 532  
**ATOMIC\_CHAR\_LOCK\_FREE**, 307, 532  
**ATOMIC\_FLAG\_INIT**, 307, 316, 532  
**ATOMIC\_INT\_LOCK\_FREE**, 307, 532  
**ATOMIC\_LLONG\_LOCK\_FREE**, 307, 532  
**ATOMIC\_LONG\_LOCK\_FREE**, 307, 532  
**ATOMIC\_POINTER\_LOCK\_FREE**, 307, 532  
**ATOMIC\_SHORT\_LOCK\_FREE**, 307, 532  
**ATOMIC\_WCHAR\_T\_LOCK\_FREE**, 307, 532  
**bitand**, 239, 512  
**BITINT\_MAXWIDTH**, 23, 111, 512, 555, 722  
**bitor**, 239, 512  
**BOOL\_MAX**, 512, 555  
**BOOL\_WIDTH**, 23, 512, 555  
**bsearch**, 387, 388, 489, 539, 634, 643, 723  
**bsearch\_s**, 540, 688, 689  
**BUFSIZ**, 339, 342, 347, 537  
**CHAR\_BIT**, 23, 43, 112, 183–185, 187, 327–329, 512, 535, 555  
**CHAR\_MAX**, 24, 243, 244, 512, 555  
**CHAR\_MIN**, 24, 41, 512, 555  
**CHAR\_WIDTH**, 23, 512, 555  
**ckd\_add**, 331, 536  
**ckd\_mul**, 331, 536  
**ckd\_sub**, 331, 536  
**CLOCKS\_PER\_SEC**, 431, 432, 435, 547  
**CMPLX**, 109, 209, 215, 216, 508, 590, 591  
**CMPLXF**, 215, 508, 509, 608  
**CMPLXL**, 215, 508  
**compl**, 239, 512  
**complex**, 146, 209–216, 419, 507–509, 590, 591, 607, 608, 630, 634  
**CR\_DECIMAL\_DIG**, 25, 511, 565, 566  
**d32add**, 33, 228, 418  
**d32div**, 33, 228, 418  
**d32fma**, 34, 228, 418  
**d32mul**, 33, 228, 418  
**d32sqrt**, 33, 228, 418  
**d32sub**, 33, 228, 418  
**d64add**, 33, 228, 418  
**d64div**, 33, 228, 418  
**d64fma**, 34, 228, 418  
**d64mul**, 33, 228, 418  
**d64sqrt**, 33, 228, 418  
**d64sub**, 33, 228, 418  
**dadd**, 418, 629  
**DBL\_DECIMAL\_DIG**, 27, 30, 510, 556  
**DBL\_DIG**, 28, 30, 510, 556  
**DBL\_EPSILON**, 29, 30, 510, 556  
**DBL\_HAS\_SUBNORM**, 26, 30, 488, 510  
**DBL\_IS\_IEC\_60559**, 26, 30, 510  
**DBL\_MANT\_DIG**, 27, 30, 109, 510, 556  
**DBL\_MAX**, 28, 30, 510, 556  
**DBL\_MAX\_10\_EXP**, 28, 30, 510, 556  
**DBL\_MAX\_EXP**, 28, 30, 510, 556  
**DBL\_MIN**, 29, 30, 510, 556  
**DBL\_MIN\_10\_EXP**, 28, 30, 510, 556  
**DBL\_MIN\_EXP**, 28, 30, 510, 556  
**DBL\_NORM\_MAX**, 29, 510, 556  
**DBL\_SNAN**, 27, 109, 164, 510, 559  
**DBL\_TRUE\_MIN**, 29, 30, 510  
**ddiv**, 418, 629, 630  
**DEC128\_EPSILON**, 32, 557  
**DEC128\_MANT\_DIG**, 31, 557  
**DEC128\_MAX**, 32, 557  
**DEC128\_MAX\_EXP**, 32, 557  
**DEC128\_MIN**, 32, 557  
**DEC128\_MIN\_EXP**, 31, 557  
**DEC128\_SNAN**, 31  
**DEC128\_TRUE\_MIN**, 32, 557  
**DEC32\_EPSILON**, 32, 557  
**DEC32\_MANT\_DIG**, 31, 557  
**DEC32\_MAX**, 32, 557  
**DEC32\_MAX\_EXP**, 32, 557  
**DEC32\_MIN**, 32, 557  
**DEC32\_MIN\_EXP**, 31, 557  
**DEC32\_SNAN**, 31, 109  
**DEC32\_TRUE\_MIN**, 32, 557  
**DEC64\_EPSILON**, 32, 557  
**DEC64\_MANT\_DIG**, 31, 557  
**DEC64\_MAX**, 32, 557  
**DEC64\_MAX\_EXP**, 32, 557

**DEC64\_MIN**, 32, 557  
**DEC64\_MIN\_EXP**, 31, 557  
**DEC64\_SNAN**, 31, 109  
**DEC64\_TRUE\_MIN**, 32, 108, 109, 557  
**DEC\_EVAL\_METHOD**, 25, 30, 66, 100, 247, 511, 556, 588, 601, 610, 646, 649  
**DEC\_INFINITY**, 31, 109, 248, 290, 489, 511, 518  
**DEC\_NAN**, 31, 109, 248, 489, 511, 518  
**DECIMAL\_DIG**, 25, 28, 488, 510, 556, 721  
**dfma**, 418, 629  
**dmul**, 418, 629  
**dsqrt**, 418, 629  
**dsub**, 418, 629, 630  
**EDOM**, 221, 250, 509  
**EILSEQ**, 221, 343, 442–445, 461, 462, 479–481, 509  
**EOF**, 217, 339, 345, 360, 362–369, 447, 457–462, 477, 537, 639, 678, 679, 681–683, 685, 703, 705, 706, 708–710  
**ERANGE**, 221, 238, 250, 251, 377, 378, 380, 466, 468, 469, 509, 649, 650  
**EXIT\_FAILURE**, 373, 385, 538  
**EXIT\_SUCCESS**, 373, 385, 427, 538  
**f32add**, 630  
**f32fma**, 630  
**f32mul**, 630  
**f32xsqrt**, 630  
**f64div**, 630  
**FE\_ALL\_EXCEPT**, 99, 223, 509  
**FE\_DEC\_DOWNWARD**, 223, 227, 510, 564  
**FE\_DEC\_TONEAREST**, 223, 224, 227, 510, 564  
**FE\_DEC\_TONEARESTFROMZERO**, 223, 227, 510, 564  
**FE\_DEC\_TOWARDZERO**, 223, 227, 510, 564  
**FE\_DEC\_UPWARD**, 223, 227, 510, 564  
**FE\_DFL\_ENV**, 224, 509  
**FE\_DFL\_MODE**, 223, 231, 509, 563  
**FE\_DIVBYZERO**, 223, 250, 509  
**FE\_DOWNWARD**, 223, 509, 563  
**FE\_INEXACT**, 223, 228, 509, 582  
**FE\_INVALID**, 223, 230, 250, 509  
**FE\_OVERFLOW**, 223, 228, 230, 250, 509  
**FE\_SNANS\_ALWAYS\_SIGNAL**, 510, 559, 563, 581, 584  
**FE\_TONEAREST**, 158, 223, 509, 563  
**FE\_TONEARESTFROMZERO**, 223, 509, 563  
**FE\_TOWARDZERO**, 223, 509, 563, 577, 582  
**FE\_UNDERFLOW**, 223, 234, 509  
**FE\_UPWARD**, 9, 223, 509, 563, 581  
**FILENAME\_MAX**, 339, 537  
**FLT\_DECIMAL\_DIG**, 27, 29, 30, 510, 556  
**FLT\_DIG**, 28, 29, 30, 510, 556  
**FLT\_EPSILON**, 29, 30, 510, 556  
**FLT\_EVAL\_METHOD**, 25, 26, 29, 30, 98–100, 109, 247, 510, 556, 588, 601, 610, 646, 649  
**FLT\_HAS\_SUBNORM**, 26, 30, 488, 510  
**FLT\_IS\_IEC\_60559**, 26, 30, 510  
**FLT\_MANT\_DIG**, 27, 29, 30, 109, 510, 556  
**FLT\_MAX**, 28, 29, 30, 510, 556  
**FLT\_MAX\_10\_EXP**, 28, 29, 30, 510, 556  
**FLT\_MAX\_EXP**, 28, 29, 30, 510, 556  
**FLT\_MIN**, 29, 30, 510, 556  
**FLT\_MIN\_10\_EXP**, 28, 29, 30, 510, 556  
**FLT\_MIN\_EXP**, 28, 29, 30, 510, 556  
**FLT\_NORM\_MAX**, 29, 510, 556  
**FLT\_RADIX**, 25, 27, 28–31, 65, 227, 232, 268, 352, 353, 376, 451, 452, 466, 510, 556, 562, 601, 602, 610  
**FLT\_ROUNDS**, 25, 26, 223, 510, 555, 556, 559, 646  
**FLT\_SNAN**, 27, 109, 510, 559  
**FLT\_TRUE\_MIN**, 29, 30, 510  
**FOPEN\_MAX**, 339, 343, 344, 537, 674  
**FP\_FAST\_D32ADDD128**, 249, 518  
**FP\_FAST\_D32ADDD64**, 249, 518  
**FP\_FAST\_D32DIVD128**, 249, 518  
**FP\_FAST\_D32DIVD64**, 249, 518  
**FP\_FAST\_D32FMAD128**, 249, 518  
**FP\_FAST\_D32FMAD64**, 249, 518  
**FP\_FAST\_D32MULD128**, 249, 518  
**FP\_FAST\_D32MULD64**, 249, 518  
**FP\_FAST\_D32SQRTD128**, 249, 518  
**FP\_FAST\_D32SQRTD64**, 249, 518  
**FP\_FAST\_D32SUBD128**, 249, 518  
**FP\_FAST\_D32SUBD64**, 249, 518  
**FP\_FAST\_D64ADDD128**, 249, 518  
**FP\_FAST\_D64DIVD128**, 249, 518  
**FP\_FAST\_D64FMAD128**, 249, 518  
**FP\_FAST\_D64MULD128**, 249, 518  
**FP\_FAST\_D64SQRTD128**, 249, 518  
**FP\_FAST\_D64SUBD128**, 249, 518  
**FP\_FAST\_DADDL**, 249, 512, 611  
**FP\_FAST\_DDIVL**, 249, 512  
**FP\_FAST\_DFMAL**, 249, 512  
**FP\_FAST\_DMULL**, 249, 512  
**FP\_FAST\_DSQRTL**, 249, 512  
**FP\_FAST\_DSUBL**, 249, 512  
**FP\_FAST\_FADD**, 249, 512, 611  
**FP\_FAST\_FADDL**, 249, 512, 611  
**FP\_FAST\_FDIV**, 249, 512  
**FP\_FAST\_FDIVL**, 249, 512  
**FP\_FAST\_FFMA**, 249, 512  
**FP\_FAST\_FFMAL**, 249, 512  
**FP\_FAST\_FMA**, 249, 512, 611  
**FP\_FAST\_FMA128**, 249, 518  
**FP\_FAST\_FMA32**, 249, 518  
**FP\_FAST\_FMA64**, 249, 518  
**FP\_FAST\_FMAF**, 249, 512, 524, 611  
**FP\_FAST\_FMAL**, 249, 512

**FP\_FAST\_FMUL**, 249, 512  
**FP\_FAST\_FMULL**, 249, 512  
**FP\_FAST\_FSQRT**, 249, 512  
**FP\_FAST\_FSQRTL**, 249, 512  
**FP\_FAST\_FSUB**, 249, 512  
**FP\_FAST\_FSUBL**, 249, 512  
**FP\_ILOGB0**, 249, 250, 264, 512  
**FP\_ILOGBNAN**, 249, 250, 264, 512  
**FP\_INFINITE**, 248, 512  
**FP\_INT\_DOWNWARD**, 248, 512  
**FP\_INT\_TONEAREST**, 248, 512  
**FP\_INT\_TONEARESTFROMZERO**, 248, 512  
**FP\_INT\_TOWARDZERO**, 248, 512  
**FP\_INT\_UPWARD**, 248, 278, 512  
**FP\_LLOGB0**, 250, 265, 512  
**FP\_LLOGBNAN**, 250, 265, 512  
**FP\_NAN**, 248, 512  
**FP\_NORMAL**, 248, 512  
**FP\_SUBNORMAL**, 248, 512  
**FP\_ZERO**, 248, 512  
**fpclassify**, 252, 512, 562, 563  
**HUGE\_VAL**, ii, 247, 251, 376, 466, 512, 571  
**HUGE\_VAL\_D128**, 248, 518  
**HUGE\_VAL\_D32**, 247, 248, 518  
**HUGE\_VAL\_D64**, 248, 518  
**HUGE\_VALF**, 247, 377, 466, 512, 571  
**HUGE\_VALL**, 247, 377, 466, 512, 571  
**I**, 209  
**INFINITY**, 27, 109, 216, 248, 489, 510, 512, 559, 590, 591  
**INT\_MAX**, 24, 40, 110, 179, 249, 250, 264, 512, 555, 642  
**INT\_MIN**, 24, 40, 249, 250, 512, 555  
**INT\_WIDTH**, 23, 278, 512, 555  
**INTMAX\_C**, 338, 536  
**INTMAX\_MAX**, 238, 337, 536  
**INTMAX\_MIN**, 238, 337, 536  
**INTMAX\_WIDTH**, 337, 536  
**INTPTR\_MAX**, 337, 536  
**INTPTR\_MIN**, 337, 536  
**INTPTR\_WIDTH**, 337, 536  
**iscanonical**, 25, 252, 512, 562, 563  
**iseqsig**, 294, 295, 517, 561, 588  
**isfinite**, 252, 512, 562, 563, 575, 591  
**isgreater**, 293, 517, 561, 562  
**isgreaterequal**, 293, 517, 562, 570, 584  
**isinf**, 252, 253, 512, 562, 563, 577, 590, 591  
**isless**, 293, 294, 517, 562, 570, 571  
**islessequal**, 294, 517, 562  
**islessgreater**, 294, 517  
**isnan**, 253, 512, 562, 563, 584, 590, 591  
**isnormal**, 253, 512, 562, 563  
**issignaling**, 253, 254, 512, 562, 563  
**issubnormal**, 254, 513, 562, 563  
**isunordered**, 294, 517, 562  
**iszero**, 254, 513, 562, 563  
**kill\_dependency**, 17, 311, 533  
**L\_tmpnam**, 340, 345, 537  
**L\_tmpnam\_s**, 538, 673, 674  
**LC\_ALL**, 241, 242, 245, 512  
**LC\_COLLATE**, 241, 242, 410, 471, 512  
**LC\_CTYPE**, 47, 241, 242, 373, 389, 390, 394, 395, 477, 482, 485–487, 512, 643, 644, 690, 691  
**LC\_MONETARY**, 241, 242, 245, 512  
**LC\_NUMERIC**, 241, 242, 245, 512  
**LC\_TIME**, 241, 242, 435, 437, 512  
**LDBL\_DECIMAL\_DIG**, 27, 488, 510, 556  
**LDBL\_DIG**, 28, 510, 556  
**LDBL\_EPSILON**, 29, 510, 556  
**LDBL\_HAS\_SUBNORM**, 26, 488, 510  
**LDBL\_IS\_IEC\_60559**, 26, 510  
**LDBL\_MANT\_DIG**, 27, 510, 556  
**LDBL\_MAX**, 28, 510, 556  
**LDBL\_MAX\_10\_EXP**, 28, 510, 556  
**LDBL\_MAX\_EXP**, 28, 510, 556  
**LDBL\_MIN**, 29, 510, 556  
**LDBL\_MIN\_10\_EXP**, 28, 510, 556  
**LDBL\_MIN\_EXP**, 28, 510, 556  
**LDBL\_NORM\_MAX**, 29, 510, 556  
**LDBL\_SNAN**, 27, 510, 559  
**LDBL\_TRUE\_MIN**, 29, 510  
**LLONG\_MAX**, 24, 380, 469, 512, 555  
**LLONG\_MIN**, 24, 291, 380, 469, 512, 555  
**LLONG\_WIDTH**, 23, 512, 555  
**LONG\_MAX**, 24, 250, 265, 380, 469, 512, 555  
**LONG\_MIN**, 24, 250, 380, 469, 512, 555  
**LONG\_WIDTH**, 23, 512, 555  
**MATH\_ERREXCEPT**, 250, 251, 377, 378, 466, 468, 512, 571, 572, 649  
**math\_errhandling**, 206, 250, 251, 377, 378, 466, 468, 512, 571, 572, 633, 640, 649, 725  
**MATH\_ERRNO**, 250, 251, 377, 378, 466, 468, 512, 649  
**MB\_CUR\_MAX**, 204, 352, 373, 390, 442–445, 479, 538, 690, 691, 716  
**MB\_LEN\_MAX**, 23, 204, 373, 512, 555  
**MB\_UTF16**, 394, 446, 540, 548  
**MB\_UTF32**, 394, 446, 540, 548  
**MB\_UTF8**, 394, 446, 540, 548  
**memchr**, 411, 412, 490, 544, 723  
**NAN**, 27, 109, 248, 489, 510, 512, 559  
**NDEBUG**, 152, 205, 208, 507  
**noreturn**, 406  
**not**, 239, 512  
**not\_eq**, 239, 512  
**NULL**, 52, 241, 332, 339, 373, 407, 431, 446, 512, 536–538, 544, 547, 548, 649

**offsetof**, 115, 332, 536, 641  
**ONCE\_FLAG\_INIT**, 373, 421, 538, 546  
**or**, 239, 512  
**or\_eq**, 239, 512  
**PRIBMAX**, 236, 511  
**PRiBMAX**, 236, 511  
**PRIBPTR**, 236, 511  
**PRiBPTR**, 236, 511  
**PRIdFAST32**, 236  
**PRIdMAX**, 236, 511  
**PRIdPTR**, 236, 511  
**PRiMAX**, 236, 511  
**PRiPTR**, 236, 511  
**PRiOMAX**, 236, 511  
**PRiOPTR**, 236, 511  
**PRiUMAX**, 236, 511  
**PRiUPTR**, 236, 511  
**PRIXMAX**, 236, 511  
**PRIXMAX**, 236, 237, 511  
**PRIXPTR**, 236, 511  
**PRIXPTR**, 236, 511  
**PTRDIFF\_MAX**, 536  
**PTRDIFF\_MIN**, 536  
**PTRDIFF\_WIDTH**, 337  
**putc**, 340, 368, 537  
**putchar**, 340, 368, 537  
**quantize**, 34, 227, 560  
**quantum**, 34, 560  
**RAND\_MAX**, 373, 380, 381, 538  
**RSIZE\_MAX**, 537, 673, 674, 675, 679, 680, 683–685, 687, 689–698, 700, 701, 704, 706, 707, 710–714, 716, 717  
**samequantum**, 560  
**SCHAR\_MAX**, 24, 512, 555  
**SCHAR\_MIN**, 24, 41, 512, 555  
**SCHAR\_WIDTH**, 23, 512, 555  
**SCNbMAX**, 237, 511  
**SCNbPTR**, 237, 511  
**SCNdMAX**, 236, 511  
**SCNdPTR**, 236, 511  
**SCNiMAX**, 236, 511  
**SCNiPTR**, 236, 511  
**SCNoMAX**, 237, 511  
**SCNoPTR**, 237, 511  
**SCNuMAX**, 237, 511  
**SCNuPTR**, 237, 511  
**SCNxMAX**, 237, 511  
**SCNxPTR**, 237, 511  
**SEEK\_CUR**, 340, 370, 537  
**SEEK\_END**, 340, 343, 370, 537  
**SEEK\_SET**, 340, 370, 371, 537, 643  
**SHRT\_MAX**, 24, 512, 555  
**SHRT\_MIN**, 24, 512, 555  
**SHRT\_WIDTH**, 23, 512, 555  
**SIG\_ATOMIC\_MAX**, 536  
**SIG\_ATOMIC\_MIN**, 536  
**SIG\_ATOMIC\_WIDTH**, 338, 536  
**SIG\_DFL**, 298, 299, 532, 649  
**SIG\_ERR**, 298, 299, 532, 640  
**SIG\_IGN**, 298, 299, 532, 645  
**SIGABRT**, 298, 299, 384, 532  
**SIGFPE**, 250, 298, 299, 532, 640, 645, 654  
**SIGILL**, 298, 299, 532, 640, 645  
**SIGINT**, 298, 532  
**signbit**, 253, 512, 562, 563, 583  
**SIGSEGV**, 298, 299, 532, 640, 645  
**SIGTERM**, 298, 532  
**SIZE\_MAX**, 42, 404, 536, 673  
**SIZE\_WIDTH**, 338, 536  
**stdc\_bit\_ceil**, 325, 535  
**stdc\_bit\_floor**, 325, 535  
**stdc\_bit\_width**, 324, 535  
**STDC\_C16\_MAX**, 394, 395, 404, 540  
**STDC\_C32\_MAX**, 394, 395, 540  
**STDC\_C8\_MAX**, 394, 395, 540  
**stdc\_count\_ones**, 323, 534  
**stdc\_count\_zeros**, 323, 534  
**stdc\_first\_leading\_one**, 321, 534  
**stdc\_first\_leading\_zero**, 321, 534  
**stdc\_first\_trailing\_one**, 322, 534  
**stdc\_first\_trailing\_zero**, 322, 534  
**stdc\_has\_single\_bit**, 324, 535  
**stdc\_leading\_ones**, 319, 533  
**stdc\_leading\_zeros**, 319, 533  
**STDC\_MC\_MAX**, 394, 395, 541  
**STDC\_MWC\_MAX**, 394, 395, 541  
**stdc\_rotate\_left**, 326, 535  
**stdc\_rotate\_right**, 326, 535  
**stdc\_trailing\_ones**, 320, 534  
**stdc\_trailing\_zeros**, 320, 534  
**strchr**, 411, 412, 490, 544, 723  
**strpbrk**, 411, 412, 490, 545, 723  
**strrchr**, 411, 412, 413, 490, 545, 723  
**strstr**, 411, 413, 490, 545, 723  
**TIME\_ACTIVE**, 431, 435, 490, 547, 651  
**TIME\_MONOTONIC**, 431, 434, 490, 547, 651  
**TIME\_THREAD\_ACTIVE**, 431, 435, 490, 547, 651  
**TIME\_UTC**, 424, 425, 428, 431, 434, 547, 651  
**TMP\_MAX**, 340, 344, 345, 537  
**TMP\_MAX\_S**, 538, 673, 674, 675  
**TSS\_DTOR\_ITERATIONS**, 421, 427, 546  
**UCHAR\_MAX**, 24, 512, 555  
**UCHAR\_WIDTH**, 23, 512, 555  
**UINT64\_C**, 338  
**UINT\_MAX**, 24, 110, 119, 179, 512, 555  
**UINT\_WIDTH**, 23, 278, 512, 555  
**UINTMAX\_C**, 338, 536  
**UINTMAX\_MAX**, 237, 238, 337, 536, 565  
**UINTMAX\_WIDTH**, 337, 536

- UINTPTR\_MAX**, 337, 536
- UINTPTR\_WIDTH**, 337, 536
- ULLONG\_MAX**, 24, 109, 119, 380, 469, 512, 555
- ULLONG\_WIDTH**, 23, 512, 555
- ULONG\_MAX**, 24, 380, 469, 512, 555
- ULONG\_WIDTH**, 23, 512, 555
- unreachable**, ix, 332, 333, 536, 641, 724
- USHRT\_MAX**, 24, 118, 119, 512, 555
- USHRT\_WIDTH**, 23, 512, 555
- WCHAR\_MAX**, 394, 446, 536, 540, 548
- WCHAR\_MIN**, 394, 446, 536, 540, 548
- WCHAR\_UTF16**, 394, 446, 540, 548
- WCHAR\_UTF32**, 394, 446, 540, 548
- WCHAR\_UTF8**, 394, 446, 540, 548
- WCHAR\_WIDTH**, 338, 394, 446, 536, 540, 548
- wcschr**, 472, 473, 490, 549, 723
- wcspbrk**, 472, 473, 490, 549, 723
- wcsrchr**, 472, 473, 474, 490, 549, 723
- wcsstr**, 472, 474, 490, 549, 723
- WEOF**, 446, 447, 461–464, 477, 482, 548, 550, 644
- WINT\_MAX**, 536
- WINT\_MIN**, 536
- WINT\_WIDTH**, 338, 536
- wmemchr**, 472, 475, 490, 549, 723
- xor**, 239, 512
- xor\_eq**, 239, 512
- obsolete
  - \_Imaginary\_I**, 723
  - \_\_STDC\_IEC\_559\_\_**, 24, 201, 203, 204, 523, 558
  - \_\_bool\_true\_false\_are\_defined**, 330, 489, 536
  - gets**, 685, 724
  - imaginary**, 209
- stream
  - stderr**, 197, 340, 341, 342, 347, 364, 459, 537, 653
  - stdin**, 340, 341, 342, 347, 360–362, 368, 457, 461, 463, 537, 678, 679, 685, 710
  - stdout**, 340, 341, 342, 347, 354, 362, 368, 453, 461, 464, 537
- structure member
  - currency\_symbol**, 241, 243, 246
  - decimal\_point**, 241, 243
  - frac\_digits**, 241, 243, 246
  - grouping**, 241, 243, 244
  - int\_curr\_symbol**, 241, 244–246
  - int\_frac\_digits**, 241, 244, 246
  - int\_n\_cs\_precedes**, 241, 244, 246
  - int\_n\_sep\_by\_space**, 241, 244–246
  - int\_n\_sign\_posn**, 241, 244–246
  - int\_p\_cs\_precedes**, 241, 244, 246
  - int\_p\_sep\_by\_space**, 241, 244–246
  - int\_p\_sign\_posn**, 241, 244–246
  - mon\_decimal\_point**, 241, 243, 246
  - mon\_grouping**, 241, 243, 244, 246
  - mon\_thousands\_sep**, 241, 243, 246
  - n\_cs\_precedes**, 241, 243, 246
  - n\_sep\_by\_space**, 241, 243, 244, 246
  - n\_sign\_posn**, 241, 244–246
  - negative\_sign**, 241, 243, 244, 246
  - p\_cs\_precedes**, 241, 243, 246
  - p\_sep\_by\_space**, 241, 243, 244, 246
  - p\_sign\_posn**, 241, 243, 245, 246
  - positive\_sign**, 241, 243, 244, 246
  - stdc\_mcerr\_incomplete\_input**, 393, 395, 396, 399, 403, 541
  - stdc\_mcerr\_insufficient\_output**, 395, 399, 403, 541
  - stdc\_mcerr\_invalid**, 393, 395, 396, 398, 399, 403, 541
  - stdc\_mcerr\_ok**, 393, 395, 398, 399, 403, 404, 541
  - thousands\_sep**, 241, 243
  - tm\_hour**, 432, 433, 436, 438, 700
  - tm\_isdst**, 432, 433
  - tm\_mday**, 432, 433, 434, 436, 438, 700
  - tm\_min**, 432, 433, 436, 438, 701
  - tm\_mon**, 432, 433, 434, 436–438, 700
  - tm\_sec**, 432, 433, 436, 438, 701
  - tm\_wday**, 432, 433, 434, 436–438, 700
  - tm\_yday**, 432, 433, 434, 438
  - tm\_year**, 432, 433, 434, 436–438, 701
  - tv\_nsec**, 432, 434
  - tv\_sec**, 432, 434
- structure type
  - lconv**, 241, 242, 512
  - timespec**, 423, 425, 428, 431, 432, 434, 435, 546, 547
  - tm**, 431, 432, 433–437, 446, 476, 547–549, 700, 701, 702
- summary, 507
- term, 204
- type
  - \_BitInt**, 110
  - \_Complex**, 110
  - \_Decimal128**, 110
  - \_Decimal128x**, 600
  - \_Decimal32**, 110
  - \_Decimal32\_t**, 247, 518, 649
  - \_Decimal64**, 110
  - \_Decimal64\_t**, 247, 518, 649
  - \_Decimal64x**, 600
  - \_Float128\_t**, 610, 611
  - \_Float128x**, 600
  - \_Float16\_t**, 610, 611
  - \_Float32\_t**, 610, 611
  - \_Float32x**, 600
  - \_Float64\_t**, 610, 611



**\_Float64x**, 600  
**atomic\_bool**, 313, 315, 532  
**atomic\_char**, 313, 532  
**atomic\_char16\_t**, 313, 532  
**atomic\_char32\_t**, 313, 532  
**atomic\_char8\_t**, 313, 532  
**atomic\_flag**, 307, 308, 316, 532, 533  
**atomic\_int**, 308, 313, 532  
**atomic\_int\_fast16\_t**, 313, 533  
**atomic\_int\_fast32\_t**, 313, 533  
**atomic\_int\_fast64\_t**, 313, 533  
**atomic\_int\_fast8\_t**, 313, 533  
**atomic\_int\_least16\_t**, 313, 532  
**atomic\_int\_least32\_t**, 313, 532  
**atomic\_int\_least64\_t**, 313, 532  
**atomic\_int\_least8\_t**, 313, 532  
**atomic\_intmax\_t**, 313, 533  
**atomic\_intptr\_t**, 313, 533  
**atomic\_llong**, 313, 532  
**atomic\_long**, 313, 532  
**atomic\_ptrdiff\_t**, 313, 533  
**atomic\_schar**, 313, 532  
**atomic\_short**, 313, 532  
**atomic\_size\_t**, 313, 533  
**atomic\_uchar**, 313, 532  
**atomic\_uint**, 313, 532  
**atomic\_uint\_fast16\_t**, 313, 533  
**atomic\_uint\_fast32\_t**, 313, 533  
**atomic\_uint\_fast64\_t**, 313, 533  
**atomic\_uint\_fast8\_t**, 313, 533  
**atomic\_uint\_least16\_t**, 313, 532  
**atomic\_uint\_least32\_t**, 313, 532  
**atomic\_uint\_least64\_t**, 313, 532  
**atomic\_uint\_least8\_t**, 313, 532  
**atomic\_uintmax\_t**, 313, 533  
**atomic\_uintptr\_t**, 313, 533  
**atomic\_ullong**, 313, 532  
**atomic\_ulong**, 313, 532  
**atomic\_ushort**, 313, 532  
**atomic\_wchar\_t**, 313, 532  
**bool**, 48, 110  
**char**, 110  
**char16\_t**, 47, 69–72, 144, 200, 313, 393–397, 400, 401, 404, **441**, 443, 540–544, 548  
**char32\_t**, 47, 69–72, 144, 200, 313, 393–397, 400, 401, **441**, 444, 445, 540–544, 548  
**char8\_t**, 47, 69–72, 108, 313, 393–397, 400, 401, 403, 404, **441**, 442, 540–544, 548  
**clock\_t**, 431, 432, 547, 651  
**cnd\_t**, 421, 422–424, 546  
**constraint\_handler\_t**, 540, **686**  
**div\_t**, 148, 373, 388, 389, 538, 539  
**double**, 110  
**double\_t**, 247, 512, 568, **609**, 610, 611, 649, 652  
**enum**, 110  
**errno\_t**, 509, 538, 540, 545, 547, 550, 672, **673**, 674–676, **686**, 687, 689–692, 693, 694–696, 698, 699, **700**, 701, **702**, 710–713, 715–717  
**femode\_t**, 222, 223, 231, 509, 510  
**fenv\_t**, 98, 222, 224, 233, 234, 509, 510, 581  
**fexcept\_t**, 222, 228–230, 509, 510, 639  
**FILE**, 76, 186, **339**, 340, 342, 344, 345, 347, 348, 356, 364, 366–372, 447, 453, 458, 459, 461–464, 537, 538, 548, 550, 641, 674–677, 681, 682, 702, 703, 705, 706  
**float**, 110  
**float\_t**, 247, 512, 568, **609**, 610, 611, 649, 652  
**fpos\_t**, 339, 341, 370, 371, 537  
**imaxdiv\_t**, 236, 237, 511  
**int**, 48, 62, 110  
**int8\_t**, 335  
**int\_fast16\_t**, 313, 336  
**int\_fast32\_t**, 236, 313, 336  
**int\_fast64\_t**, 313, 336  
**int\_fast8\_t**, 313, 336  
**int\_least16\_t**, 313, 336  
**int\_least32\_t**, 313, 335, 336  
**int\_least64\_t**, 313, 336  
**int\_least8\_t**, 313, 336  
**intmax\_t**, 24, 179, 237, 238, 313, **336**, 338, 350, 357, 449, 455, 511, 536, 723, 725  
**intptr\_t**, 313, 336, 536  
**jmp\_buf**, 296, 297, 532, 719  
**ldiv\_t**, 373, 389, 538, 539  
**lldiv\_t**, 373, 389, 538, 539  
**long\_double\_t**, 524, **609**, 610, 611  
**max\_align\_t**, 46, 332, 536  
**mbstate\_t**, 341–343, 353, 359, 371, 394–403, **441**, 442–445, **446**, 452, 456, 477–480, 540–544, 548–550, 634, 644, 651, 715–717  
**memory\_order**, 307, **309**, 311, 312, 314–316, 489, 532, 533  
**mtx\_t**, 421, 423–426, 546  
**nullptr\_t**, ix, 42, 48, 52, **53**, 89, 93, 95–97, 303, **332**, **334**, 536, 641, 723  
**once\_flag**, 373, **421**, 422, 538, 546  
**ptrdiff\_t**, 82, 91, 307, 313, **332**, 333, 337, 350, 357, 449, 455, 536, 636, 723  
**rsize\_t**, 536, 538, 540, 545, 547, 550, 673, 674, 677–680, 683–685, **686**, 687–692, **693**, 694–698, **700**, 701, **702**, 703, 704, 706, 707, 710–717

- sig\_atomic\_t**, 14, 298, 299, 338, 532, 632, 640
- signed**, 110
- stdc\_mcerr**, 394, 395–397, 400–404, 540–544
- thrd\_start\_t**, 421, 426, 546
- thrd\_t**, 421, 426, 427, 546, 547
- time\_t**, 431, 432, 433, 434, 436, 437, 547, 651, 701, 702
- tss\_dtor\_t**, 421, 428, 546, 547
- tss\_t**, 421, 428, 429, 546, 547
- uint64\_t**, 278
- uint\_fast16\_t**, 313, 336
- uint\_fast32\_t**, 313, 336
- uint\_fast64\_t**, 313, 336
- uint\_fast8\_t**, 313, 336
- uint\_least16\_t**, 313, 335, 336, 441
- uint\_least32\_t**, 313, 336, 441
- uint\_least64\_t**, 313, 336, 338
- uint\_least8\_t**, 313, 336
- uintmax\_t**, 24, 179, 237, 238, 313, 337, 338, 350, 357, 449, 455, 511, 536, 723, 725
- uintptr\_t**, 313, 336, 536
- unsigned**, 110, 349, 357, 449, 454
- va\_list**, 302, 303–305, 364–366, 458–460, 532, 537, 538, 548, 550, 640–642, 649, 681–684, 705–709
- void**, 52, 110
- wchar\_t**, 5, 46, 47, 68–71, 144, 200, 201, 238, 313, 332, 338, 350, 352, 353, 357, 359, 362, 373, 389–391, 393–397, 400, 401, 441, 446, 447, 449, 452–454, 456–465, 467–476, 478–480, 482, 511, 536, 538–544, 548–550, 625, 627, 628, 634, 644, 646, 690–692, 702–717
- wctrans\_t**, 482, 486, 487, 550, 551
- wctype\_t**, 482, 485, 486, 550, 551
- wint\_t**, 338, 350, 352, 446, 447, 449, 452, 461–464, 477, 482, 483–486, 548–551, 644
  - use of function, 206
- lifetime, 38
- limit** embed parameter, 74, 177, 181, 183, 184, 186, 187–190
- line, 11
  - preprocessing directive, 176
- line buffered, 342
- line buffered stream, 342
- line number, 197, 199
- line** preprocessing directive, 197
- lines, 341
- linkage, 37, 104, 131, 135, 171, 173, 203
- literal
  - binary, 61
  - character, 67
  - floating, 63
  - hexadecimal, 60
  - integer, 60
  - octal, 60
  - prefixed octal, 60
  - unprefixed octal, 60
- literal encoding, 46
- little-endian, 318
- ll format modifier, 350, 357, 449, 454
- llabs** function, 388, 539
- lldiv** function, 143, 373, 388, 389, 539
- lldiv\_t** type, 373, 389, 538, 539
- llogb** function, 250, 265, 514, 560, 576
- llogb** type-generic macro, 417
- llogbd128** function, 265, 519
- llogbd32** function, 265, 519
- llogbd64** function, 265, 519
- llogbdN** function, 616
- llogbdNx** function, 616
- llogbf** function, 265, 514
- llogbfN** function, 616
- llogbfNx** function, 616
- llogbl** function, 265, 514
- LLONG\_MAX** macro, 24, 380, 469, 512, 555
- LLONG\_MIN** macro, 24, 291, 380, 469, 512, 555
- LLONG\_WIDTH** macro, 23, 512, 555
- llquantexpd** identifier prefix, 290, 291, 418
- llquantexpd128** function, 291, 523
- llquantexpd32** function, 291, 523
- llquantexpd64** function, 291, 523
- llquantexpdN** function, 621
- llquantexpdNx** function, 621
- llrint** function, 275, 515, 565, 581, 582, 634
- llrint** type-generic macro, 417
- llrintd128** function, 275, 521
- llrintd32** function, 275, 521
- llrintd64** function, 275, 521
- llrintdN** function, 618
- llrintdNx** function, 618
- llrintf** function, 275, 516
- llrintfN** function, 618
- llrintfNx** function, 618
- llrintl** function, 275, 516
- llround** function, 276, 516, 561, 582, 634
- llround** function, 276
- llround** type-generic macro, 417
- llroundd128** function, 276, 521
- llroundd32** function, 276, 521
- llroundd64** function, 276, 521
- llrounddN** function, 618
- llrounddNx** function, 618
- llroundf** function, 276, 516
- llroundfN** function, 618
- llroundfNx** function, 618
- llroundl** function, 276, 516
- local, 156

- local time, 431
- locale, 4
- locale-specific behavior, 4, 651
- localeconv** function, 242, 245, 512, 640
- localization header, 241, 488
- localtime** function, 433, 435, 436, 437, 438, 547, 651
- localtime\_r** function, 437, 547, 721
- localtime\_s** function, 547, 701, 702
- log** function, 250, 265, 266, 417, 514, 564, 576
- log** type-generic macro, 417
- log10** function, 265, 266, 514, 564, 576
- log10** type-generic macro, 417
- log10d128** function, 266, 519
- log10d32** function, 265, 519
- log10d64** function, 265, 519
- log10dN** function, 616
- log10dNx** function, 616
- log10f** function, 265, 514
- log10fN** function, 616
- log10fNx** function, 616
- log10l** function, 265, 514
- log10p1** function, 266, 514, 564, 576
- log10p1** type-generic macro, 417
- log10p1d128** function, 266, 519
- log10p1d32** function, 266, 519
- log10p1d64** function, 266, 519
- log10p1dN** function, 616
- log10p1dNx** function, 616
- log10p1f** function, 266, 514
- log10p1fN** function, 616
- log10p1fNx** function, 616
- log10p1l** function, 266, 514
- log1p** function, 266, 267, 514, 564, 576
- log1p** type-generic macro, 417
- log1pd128** function, 266, 519
- log1pd32** function, 266, 519
- log1pd64** function, 266, 519
- log1pdN** function, 616
- log1pdNx** function, 616
- log1pf** function, 266, 514
- log1pfN** function, 616
- log1pfNx** function, 616
- log1pl** function, 266, 514
- log2** function, 267, 514, 564, 576
- log2** type-generic macro, 417
- log2d128** function, 267, 519
- log2d32** function, 267, 519
- log2d64** function, 267, 519
- log2dN** function, 616
- log2dNx** function, 616
- log2f** function, 267, 514
- log2fN** function, 616
- log2fNx** function, 616
- log2l** function, 267, 514
- log2p1** function, 266, 267, 514, 564, 577
- log2p1** type-generic macro, 417
- log2p1d128** function, 267, 520
- log2p1d32** function, 267, 519
- log2p1d64** function, 267, 519
- log2p1dN** function, 616
- log2p1dNx** function, 616
- log2p1f** function, 267, 514
- log2p1fN** function, 616
- log2p1fNx** function, 616
- log2p1l** function, 267, 514
- logarithmic function
  - complex, 213, 596
  - real, 261, 575
- logb** function, 264, 265, 267, 268, 514, 560, 575, 577, 591, 622
- logb** type-generic macro, 417
- logbd128** function, 267, 520
- logbd32** function, 267, 520
- logbd64** function, 267, 520
- logbdN** function, 616
- logbdNx** function, 616
- logbf** function, 267, 514
- logbfN** function, 616
- logbfNx** function, 616
- logbl** function, 267, 514
- logd128** function, 265, 519
- logd32** function, 265, 519
- logd64** function, 265, 519
- logdN** function, 616
- logdNx** function, 616
- logf** function, 265, 514
- logfN** function, 616
- logfNx** function, 616
- logical operator
  - AND (&&), 17, 94
  - negation (!), 87
  - OR (| |), 95
  - OR (| |), 17
- logical source line, 11
- logl** function, 265, 514
- logp1** function, 266, 267, 514, 564, 576
- logp1** type-generic macro, 417
- logp1d128** function, 266, 519
- logp1d32** function, 266, 519
- logp1d64** function, 266, 519
- logp1dN** function, 616
- logp1dNx** function, 616
- logp1f** function, 266, 514
- logp1fN** function, 616
- logp1fNx** function, 616
- logp1l** function, 266, 514
- long double \_Complex** type, 41
- long double \_Complex** type conversion, 49
- long double suffix, l or L, 65
- long double** type, 40, 110
- long double** type conversion, 49



- `long int` type, 40, 110
- `long int` type conversion, 47–49
- long integer suffix, `l` or `L`, 62
- `long` keyword, 55
- `long long int` type, 40, 110
- `long long int` type conversion, 47–49
- long long integer suffix, `ll` or `LL`, 62
- `long_double_t` type, 524, 609, 610, 611
- `LONG_MAX` macro, 24, 250, 265, 380, 469, 512, 555
- `LONG_MIN` macro, 24, 250, 380, 469, 512, 555
- `LONG_WIDTH` macro, 23, 512, 555
- `longjmp` function, 296, 297, 385, 386, 532, 640, 643, 719
- loop body, 165
- low-order bit, 5
- lowercase letter, 20
- `lrint` function, 275, 515, 565, 581, 582, 634
- `lrint` type-generic macro, 417
- `lrintd128` function, 275, 521
- `lrintd32` function, 275, 521
- `lrintd64` function, 275, 521
- `lrintdN` function, 618
- `lrintdNx` function, 618
- `lrintf` function, 275, 515
- `lrintfN` function, 618
- `lrintfNx` function, 618
- `lrintl` function, 275, 515
- `lround` function, 276, 516, 561, 582, 634
- `lround` function, 276
- `lround` type-generic macro, 417
- `lroundd128` function, 276, 521
- `lroundd32` function, 276, 521
- `lroundd64` function, 276, 521
- `lrounddN` function, 618
- `lrounddNx` function, 618
- `lroundf` function, 276, 516
- `lroundfN` function, 618
- `lroundfNx` function, 618
- `lroundl` function, 276, 516
- lvalue, 51, 82, 86, 96, 124
- lvalue conversion, 51, 96–98
- macro
  - atomic lock-free, 307
  - math rounding direction, 248
  - number classification, 248
  - type-generic, 416
- macro argument substitution, 191
- macro definition
  - library function, 206
- macro invocation, 191
- macro name, 190
  - length, 22
  - predefined, 199, 203
  - redefinition, 190
  - scope, 195
- macro parameter, 191
- macro preprocessor, 175
- macro replacement, 190
- magnitude
  - complex, 214
- manipulation function
  - complex, 214
  - real, 279, 583
- matching failure, 459, 460, 706, 708, 709
- math rounding direction macro, 248
- `MATH_` identifier prefix, 489
- `MATH_ERREXCEPT` macro, 250, 251, 377, 378, 466, 468, 512, 571, 572, 649
- `math_errhandling` macro, 206, 250, 251, 377, 378, 466, 468, 512, 571, 572, 633, 640, 649, 725
- `MATH_ERRNO` macro, 250, 251, 377, 378, 466, 468, 512, 649
- mathematics header, 247, 488
- `max_align_t` type, 46, 332, 536
- maximal munch, 54
- maximum function, 282, 584
- `maybe_unused` attribute, 150, 152, 720, 722
- `MB_CUR_MAX` macro, 204, 352, 373, 390, 442–445, 479, 538, 690, 691, 716
- `MB_LEN_MAX` macro, 23, 204, 373, 512, 555
- `MB_UTF16` macro, 394, 446, 540, 548
- `MB_UTF32` macro, 394, 446, 540, 548
- `MB_UTF8` macro, 394, 446, 540, 548
- `mblen` function, 389, 478, 539
- `mbrlen` function, 478, 549
- `mbrtoc16` function, 442, 443, 548
- `mbrtoc32` function, 444, 548
- `mbrtoc8` function, 441, 442, 548, 722
- `mbrtowc` function, 343, 359, 452, 453, 477, 478, 479, 480, 549, 691, 717
- `mbsinit` function, 402–404, 477, 478, 549
- `mbsrtowcs` function, 477, 480, 549, 716
- `mbsrtowcs_s` function, 550, 716, 717
- `mbstate_t` type, 341–343, 353, 359, 371, 394–403, 441, 442–445, 446, 452, 456, 477–480, 540–544, 548–550, 634, 644, 651, 715–717
- `mbstowcs` function, 72, 390, 391, 464, 465, 479, 539
- `mbstowcs_s` function, 540, 691, 692
- `mbtowc` function, 70, 389, 390, 391, 478, 539
- `mem` identifier prefix, 490
- `memalign` function, 9, 391, 539, 723
- member access operators (`.` and `->`), 80
- member alignment, 114
- members, 38
- `memccpy` function, 407, 544, 721
- `memchr` macro, 411, 412, 490, 544, 723
- `memcpy` function, 43, 186, 315, 409, 410, 544
- `memcpy` function, 43, 76, 183, 207, 308, 315, 407,

- 544, 561
- memcpy\_s** function, 545, 693
- memmove** function, 76, 407, 408, 544, 561, 638
- memmove\_s** function, 545, 693, 694
- memory location, 6
- memory management function, 381
- memory\_** identifier prefix, 489
- memory\_order\_** identifier prefix, 489
- memory\_order** type, 307, 309, 311, 312, 314–316, 489, 532, 533
- memory\_order\_acq\_rel** constant, 309, 310, 311, 314–316, 532
- memory\_order\_acquire** constant, 309, 311, 314, 316, 532
- memory\_order\_consume** constant, 309, 311, 314, 532
- memory\_order\_relaxed** constant, 157, 309, 310, 311, 532
- memory\_order\_release** constant, 309, 311, 314, 315, 532
- memory\_order\_seq\_cst** constant, 19, 43, 82, 97, 98, 307, 309, 310, 312, 532
- memset**
  - explicit, 414
- memset** function, 308, 414, 545, 698
- memset\_explicit** function, 414, 545, 723
- memset\_s** function, 545, 698
- minimum function, 282, 584
- minus operator
  - unary, 563
- minus operator, unary, 87
- miscellaneous function
  - string, 414
  - wide string, 475
- miscellaneous functions
  - string, 698
  - wide string, 715
- mktime** function, 433, 547
- modf** family, 34, 268, 416
- modf** function, 268, 416, 417, 514, 577
- modfd128** function, 268, 520
- modfd32** function, 268, 520
- modfd64** function, 268, 520
- modfdN** function, 616
- modfdNx** function, 616
- modfff** function, 268, 514
- modfffN** function, 616
- modfffNx** function, 616
- modfl** function, 268, 514
- modifiable lvalue, 51
- modification order, 17
- modulus
  - complex, 214
- modulus function, 268
- mon\_decimal\_point** structure member, 241, 243, 246
- mon\_grouping** structure member, 241, 243, 244, 246
- mon\_thousands\_sep** structure member, 241, 243, 246
- most significant index, 318
- mtx\_** identifier prefix, 490
- mtx\_destroy** function, 424, 546
- mtx\_init** function, 421, 422, 424, 425, 546
- mtx\_lock** function, 425, 546, 644
- mtx\_plain** constant, 421, 425, 546
- mtx\_recursive** constant, 421, 425, 546
- mtx\_t** type, 421, 423–426, 546
- mtx\_timed** constant, 422, 425, 546
- mtx\_timedlock** function, 425, 546, 644
- mtx\_trylock** function, 425, 426, 546
- mtx\_unlock** function, 425, 426, 546, 644
- multi unit transcoding function, 402
- multibyte character, 5, 20, 68
- multibyte conversion function
  - wide character, 389, 690
  - extended, 476, 715
  - restartable, 441, 478, 715
  - wide string, 390, 691
  - restartable, 479, 716
- multibyte string, 204
- multibyte/wide character conversion function, 389, 690
  - extended, 476, 715
  - restartable, 441, 478, 715
- multibyte/wide string conversion function, 390, 691
  - restartable, 479, 716
- multidimensional array, 79
- multiplication assignment operator (\*=), 98
- multiplication operator (\*), 89, 589
- multiplicative expression, 89, 589
- multiply and round to narrower type, 288
- n-char sequence, 376
- n-wchar sequence, 465
- n\_cs\_precedes** structure member, 241, 243, 246
- n\_sep\_by\_space** structure member, 241, 243, 244, 246
- n\_sign\_posn** structure member, 241, 244–246
- name
  - external, 22, 57, 203
  - file, 342
  - internal, 22, 57
  - label, 38
  - structure/union member, 38
- name space, 38
- named, 161
- named constant, 100
- named label, 161
- NaN, 25
- nan** function, 280, 516, 559, 583

- NAN** macro, 27, 109, 248, 489, 510, 512, 559
- nand128** function, 280, 521
- nand32** function, 280, 521
- nand64** function, 280, 521
- nandN** function, 619
- nandNx** function, 619
- nanf** function, 280, 516
- nanfN** function, 619
- nanfNx** function, 619
- nanl** function, 280, 516
- narrow character literals, 68
- narrow execution encoding, 47
- NDEBUG** macro, 152, 205, 208, 507
- nearbyint** function, 274, 275, 515, 563, 565, 577, 581
- nearbyint** type-generic macro, 417
- nearbyintd128** function, 274, 521
- nearbyintd32** function, 274, 521
- nearbyintd64** function, 274, 521
- nearbyintdN** function, 618
- nearbyintdNx** function, 618
- nearbyintf** function, 274, 515
- nearbyintfN** function, 618
- nearbyintfNx** function, 618
- nearbyintl** function, 274, 515
- nearest integer function, 274, 580
- negation operator (!), 87
- negative zero, 280
- negative\_sign** structure member, 241, 243, 244, 246
- new line, 21
- new-line character, 11, 20, 54, 177, 197
- new-line escape sequence (\n), 21, 69, 219
- nextafter** function, 280, 281, 419, 516, 563, 583
- nextafter** type-generic macro, 417
- nextafterd128** function, 280, 521
- nextafterd32** function, 280, 521
- nextafterd64** function, 280, 521
- nextafterdN** function, 619
- nextafterdNx** function, 619
- nextafterf** function, 280, 516
- nextafterfN** function, 619
- nextafterfNx** function, 619
- nextafterl** function, 280, 516
- nextdown** function, 282, 516, 560, 584
- nextdown** type-generic macro, 417
- nextdownd128** function, 282, 522
- nextdownd32** function, 282, 521
- nextdownd64** function, 282, 522
- nextdowndN** function, 619
- nextdowndNx** function, 619
- nextdownf** function, 282, 516
- nextdownfN** function, 619
- nextdownfNx** function, 619
- nextdownl** function, 282, 516
- nexttoward** function, 281, 516, 563, 583
- nexttoward** type-generic macro, 417
- nexttowarddd128** function, 281, 521
- nexttowarddd32** function, 281, 521
- nexttowarddd64** function, 281, 521
- nexttowardf** function, 281, 419, 516
- nexttowardl** function, 281, 516
- nextup** function, 281, 516, 560, 584
- nextup** type-generic macro, 417
- nextupd128** function, 281, 521
- nextupd32** function, 281, 521
- nextupd64** function, 281, 521
- nextupdN** function, 619
- nextupdNx** function, 619
- nextupf** function, 281, 516
- nextupfN** function, 619
- nextupfNx** function, 619
- nextupl** function, 281, 516
- no linkage, 37
- no-return function, 132
- nodiscard** attribute, 149, 150, 151, 152, 720, 722
- non-canonical, 25
- non-canonical representation, 25
- non-graphic character, 21, 69
- non-local jumps header, 296
- non-stop floating-point control mode, 233
- non-value representation, 7, 43, 52, 81
- noreturn** attribute, 132, 150, 154, 155, 297, 384–386, 427, 532, 539, 544, 547, 720, 722
- noreturn** macro, 406
- norm
  - complex, 214
- normalized, 700
- normalized broken-down time, 700
- normalized floating-point number, 24
- not** macro, 239, 512
- not\_eq** macro, 239, 512
- null character (\0), 19, 70, 71
  - padding of binary stream, 341
- NULL** macro, 52, 241, 332, 339, 373, 407, 431, 446, 512, 536–538, 544, 547, 548, 649
- null pointer, 52
- null pointer constant, 52
- null preprocessing directive, 199
- null statement, 162
- null wide character, 204
- nullptr** keyword, 55
- nullptr** predefined constant, 70
- nullptr\_t** type, ix, 42, 48, 52, 53, 89, 93, 95–97, 303, 332, 334, 536, 641, 723
- nullptr\_t** type conversion, 53
- number classification macro, 248, 252
- numeric conversion function, 237, 373
  - wide string, 238, 464

numerical limit, 22

O format modifier, 439

object, 6

object representation, 43

object type, 39

object types, 39

object-like macro, 191

observable, 156

observable behavior, 14

observed, 156

obsolescence, xv, 203, 488

octal digit, 61, 69

octal digit sequence, 61

octal literal, 60

octal literals, 203

octal prefix, 61

octal-character escape sequence (*\octal digit*), 69

**OFF** pragma, 158, 590, 591

**offsetof** macro, 115, 332, 536, 641

**ON** pragma, 98, 225, 230, 232, 234, 567–569, 577, 581, 583

on-off switch, 198

**once\_flag** type, 373, 421, 422, 538, 546

**ONCE\_FLAG\_INIT** macro, 373, 421, 538, 546

opening, 342

operand, 72, 76

operating system, 12, 386

operations on file, 343

operations on files, 674

operator, 72, 76

**\_Alignas**, 55

**\_Alignof**, 55

**\_Lengthof**, 86, 87

**\_\_has\_c\_attribute**, 149, 151–155, 157, 179, 180, 199, 720

**\_\_has\_embed**, 54, 73, 179, 181, 199

**\_\_has\_include**, 54, 73, 179, 180, 199, 723

additive, 90

**alignas**, 55

**alignof**, 55, 86, 87

assignment, 96

associativity, 76

**defined**, 179, 180, 186, 199, 638

equality, 92

multiplicative, 89, 589

postfix, 79

precedence, 76

preprocessing, 193, 201

relational, 92

shift, 91

**sizeof**, 55, 86, 87

**typeof**, 124

**typeof**, 55

**typeof\_unqual**, 55

unary, 85

unary arithmetic, 87

**or** macro, 239, 512

OR operator

bitwise exclusive (^), 94

bitwise exclusive assignment (^=), 98

bitwise inclusive (|), 94

bitwise inclusive assignment (|=), 98

logical (||), 95

logical (||), 17

**or\_eq** macro, 239, 512

order of allocated storage, 381

order of evaluation, 76, 96, 193, 194

ordinary character literal, 68

ordinary identifier name space, 38

ordinary string literal, *see* string literal

orientation, 341

orientation of stream, 341, 462

out-of-bounds store, 8

outer scope, 37

output code unit writes, 398

over-aligned, 46

**p\_cs\_precedes** structure member, 241, 243, 246

**p\_sep\_by\_space** structure member, 241, 243, 244, 246

**p\_sign\_posn** structure member, 241, 243, 245, 246

padding

binary stream, 341

bit, 43, 335

structure/union, 43, 114

parameter, 6

array, 172

ellipsis, 137, 191

embed, 183

function, 80, 172

generic, 416

macro, 191

**main** function, 13

program, 13

parameter type list, 137

parentheses punctuator (( )), 137, 163, 165

parenthesized expression, 77

parse state, 341

perform a trap, 7

permitted form of initializer, 100

**pererror** function, 372, 537

phase angle

complex, 215

physical source line, 11

placemaker, 193

plus operator

unary, 87

pointer

null, 52

pointer arithmetic, 90

- pointer comparison, 92
- pointer declarator, 134
- pointer operator (->), 81
- pointer to a string, 204
- pointer to a wide string, 204
- pointer to function, 80
- pointer type, 42
- pointer type conversion, 52
- pole error, 250, 259, 260, 265–273
- portability, 9
- positive difference, 283
- positive difference function, 282, 584
- positive\_sign** structure member, 241, 243, 244, 246
- postfix decrement operator (--), 51, 83
- postfix expression, 79
- postfix increment operator (++), 51, 82
- pow** function, 270, 417, 515, 565, 578, 630
- pow** type-generic macro, 417
- powd128** function, 270, 520
- powd32** function, 270, 520
- powd64** function, 270, 419, 520
- powdN** function, 617
- powdNx** function, 617
- power function
  - complex, 214, 597
  - real, 269, 577
- powf** function, 270, 515
- powf32x** function, 630
- powf64** function, 630
- powfN** function, 617
- powfNx** function, 617
- powl** function, 270, 515, 563
- pown** function, 270, 271, 515, 565, 579
- pown** type-generic macro, 417
- pownd128** function, 270, 520
- pownd32** function, 270, 520
- pownd64** function, 270, 520
- powndN** function, 617
- powndNx** function, 617
- pownf** function, 270, 515
- pownfN** function, 617
- pownfNx** function, 617
- pownl** function, 270, 515
- powr** function, 271, 515, 565, 579
- powr** type-generic macro, 417
- powrd128** function, 271, 520
- powrd32** function, 271, 520
- powrd64** function, 271, 520
- powrdN** function, 617
- powrdNx** function, 617
- powrf** function, 271, 515
- powrfN** function, 617
- powrfNx** function, 617
- powrl** function, 271, 515
- pp-number, 74
- pragma
  - CX\_LIMITED\_RANGE**, vii, 198, 199, 209, 210, 507, 589, 639
  - FE\_DEC\_DOWNWARD**, 199
  - FE\_DEC\_DYNAMIC**, 199, 227
  - FE\_DEC\_TONEAREST**, 199
  - FE\_DEC\_TONEARESTFROMZERO**, 199
  - FE\_DEC\_TOWARDZERO**, 199
  - FE\_DEC\_UPWARD**, 199
  - FE\_DOWNWARD**, 199
  - FE\_DYNAMIC**, 199, 225, 509
  - FE\_TONEAREST**, 199
  - FE\_TONEARESTFROMZERO**, 199
  - FE\_TOWARDZERO**, 199
  - FE\_UPWARD**, 199
  - FENV\_ACCESS**, vii, 9, 98, 198, 199, 224, 225, 230, 232, 234, 509, 566–571, 577, 581, 583, 633, 639, 647
  - FENV\_DEC\_ROUND**, vii, 66, 198, 199, 227, 510, 564, 609
  - FENV\_ROUND**, vii, 30, 66, 158, 198, 199, 224, 225, 226, 227, 509, 563, 609
  - FP\_CONTRACT**, viii, 77, 158, 198, 199, 251, 512, 590, 591, 639, 647
  - OFF**, 158, 590, 591
  - ON**, 98, 225, 230, 232, 234, 567–569, 577, 581, 583
  - STDC**, 98, 158, 198, 203, 209, 224–227, 230, 232, 234, 251, 507, 509, 510, 512, 567–569, 577, 581, 583, 590, 591, 638, 648
- pragma**, 54, 73, 98, 158, 175, 198, 202, 209, 224–227, 230, 232, 234, 251, 507, 509, 510, 512, 567–569, 577, 581, 583, 590, 591, 631, 638, 648
- pragma operator, 201
- pragma** preprocessing directive, 198, 203
- precedence of operator, 76
- precedence of syntax rule, 11
- precision, 44, 47, 349, 448
  - excess, 26, 50, 169
- predefined constant
  - false**, 55, 70
  - nullptr**, 55, 70
  - true**, 55, 70
- predefined macro name, 199, 203
- preferred quantum exponent, 33
- prefix decrement operator (--), 51, 86
- prefix** embed parameter, 184, 188, 189
- prefix increment operator (++), 51, 86
- prefixed octal literal, 60
- preprocessing, 177
- preprocessing concatenation, 193
- preprocessing directive, 11, 175, 176
  - ifdef**, 9, 175, 178, 179, 180, 227, 231, 232, 254–281, 282, 283–287, 288, 289–292, 375, 377, 419, 467, 556, 557, 585–587



- pragma**, 54, 73, 98, 158, 175, 198, 202, **209**, 224–227, 230, 232, 234, 251, 507, 509, 510, 512, 567–569, 577, 581, 583, 590, 591, 631, 638, 648
- undef**, 57, 175, 178, **195**, 199, 206, 207, 638
- preprocessing file, **11**, **175**
- preprocessing number, 54, **74**
- preprocessing operator
  - #**, **193**
  - ##**, **193**
  - \_Pragma**, **201**
- preprocessing parameter, **177**
- preprocessing token, 11, **54**, 176
- preprocessing translation unit, **11**
- preprocessor, **175**
- preprocessor parameter, **177**
  - prefixed, **177**
  - standard, **177**
- preprocessor prefixed parameter, **177**
- preprocessor standard parameter, **177**
- PRI** identifier prefix, **236**, 488
- PRIBMAX** macro, **236**, 511
- PRiBMAX** macro, **236**, 511
- PRIBPTR** macro, **236**, 511
- PRiBPTR** macro, **236**, 511
- PRiCFASTN** macro, **236**
- PRiCLEASTN** macro, **236**
- PRiCN** macro, **236**
- PRiDFAST32** macro, **236**
- PRiDMAX** macro, **236**, 511
- PRiDPTR** macro, **236**, 511
- PRiIMAX** macro, **236**, 511
- PRiIPTR** macro, **236**, 511
- primary block, **160**
- primary expression, **77**
- printf\_s** function, 538, **678**, 679
- printing character, 21, **217**, 218
- printing wide character, **482**
- PRiOMAX** macro, **236**, 511
- PRiOPTR** macro, **236**, 511
- PRiUMAX** macro, **236**, 511
- PRiUPTR** macro, **236**, 511
- PRiXMAX** macro, **236**, 511
- PRiXMAX** macro, **236**, 237, 511
- PRiXPTR** macro, **236**, 511
- PRiXPTR** macro, **236**, 511
- program
  - conforming, **9**
  - strictly conforming, **9**
- program diagnostic, **208**
- program execution, **13**
- program file, **11**
- program image, **12**
- program name, **13**
- program name (**argv**[0]), **13**
- program parameter, **13**
- program parameters, **13**
- Program semantics, **14**
- program startup, **12**, **13**
- program structure, **11**
- program termination, **12**, **13**, **14**
- promotion
  - integer, **15**, **48**
- promotions
  - default argument, **80**
- pseudo-random sequence function, **380**
- PTRDIFF\_MAX** macro, **337**, 536
- PTRDIFF\_MIN** macro, **337**, **338**, 536
- ptrdiff\_t** type, 82, 91, 307, 313, **332**, **333**, **337**, 350, 357, 449, 455, 536, 636, 723
- PTRDIFF\_WIDTH** macro, **337**
- punctuator, **72**
- putc** macro, 340, **368**, 537
- putchar** macro, 340, **368**, 537
- puts** function, 197, 333, 340, **368**, 537
- putwc** function, 340, **463**, 464, 548
- putwchar** function, 340, **463**, 464, 548
- qsort** function, 387, **388**, 539, 634
- qsort\_s** function, 540, 688, **689**, 690
- qualified type, **42**
- qualified version of type, **42**
- qualifier
  - \_Atomic**, **42**, 55
  - const**, **42**, 55
  - restrict**, **42**, 55
  - volatile**, **42**, 55
- quantize** macro, 34, 227, 560
- quantized** identifier prefix, **289**, 290, 418
- quantized128** function, **289**, 523
- quantized32** function, **289**, 522
- quantized64** function, **289**, 522
- quantizedN** function, **621**
- quantizedNx** function, **621**
- quantum, **32**
- quantum exponent, **32**
- quantum exponent function, **289**
- quantum function, **289**
- quantum** macro, 34, 560
- quantumd** identifier prefix, **290**, 418
- quantumd128** function, **290**, 523
- quantumd32** function, **290**, 523
- quantumd64** function, **290**, 523
- quantumdN** function, **621**
- quantumdNx** function, **621**
- question-mark escape sequence (**\?**), **69**
- quick\_exit** function, 299, 384, 385, **386**, 539, 634, 640, 643, 650, 724
- quiet NaN, **25**
- raise** function, 298, 299, **300**, 308, 384, 532, 640
- rand** function, 373, **380**, 381, 539

- RAND\_MAX** macro, 373, 380, 381, 538
- range
  - excess, 26, 50, 169
- range error, 250, 255–264, 266–273, 275, 276, 281, 283, 287
- read-modify-write operation, 17
- read-read coherence, 19
- read-write coherence, 19
- real floating type, 41
- real floating type conversion, 49, 565
- real floating types, 600
- real type, 41
- real type domain, 41
- real-floating, 252
- realloc** function, 381, 382, 383, 384, 539, 634, 643, 650, 719, 723
- recommended practice, 7
- recursion, 80
- recursive function call, 80
- redefinition of macro, 190
- reentrancy, 14, 21
  - library function, 207
- referenced type, 42
- register**, 106
- register** storage-class specifier, 55, 171
- relational expression, 92
- relaxed atomic operation, 17
- release fence, 311
- release operation, 17
- release sequence, 17
- reliability of data
  - interrupted, 14
- remainder assignment operator (%=), 98
- remainder function, 278, 582
- remainder** function, 279, 419, 516, 560, 564, 583, 649
- remainder operator (%), 89
- remainder** type-generic macro, 417
- remainderd128** function, 279, 521
- remainderd32** function, 279, 521
- remainderd64** function, 279, 521
- remainderdN** function, 618
- remainderdNx** function, 618
- remainderf** function, 279, 516
- remainderfN** function, 618
- remainderfNx** function, 618
- remainderl** function, 279, 516
- remove** function, 343, 344, 537, 650, 674
- remquo** function, 279, 516, 560, 564, 583, 633, 649
- remquo** type-generic macro, 417
- remquof** function, 279, 516
- remquofN** function, 618
- remquofNx** function, 618
- remquol** function, 279, 516
- rename** function, 344, 537, 650
- representation
  - canonical, 25
  - non-canonical, 25
- representations of type, 43
  - pointer, 42
- reproducible, 156
- reproducible** attribute, 150, 155, 157, 158, 638, 720, 722
- rescanning and replacement, 194
- reserved identifier, 55, 205, 654, 663, 672
- resource, 178, 183
  - #embed** preprocessing directive, 183
  - empty, 183
- resource width, 183
- restartable multibyte/wide character conversion function, 441, 478, 715
- restartable multibyte/wide string conversion function, 479, 716
- restore calling environment function, 296
- restrict**, 42
- restrict** keyword, 55
- restrict** type qualifier, 127, 129
- restrict-qualified type, 42, 128
- return** keyword, 55
- return** statement, 169
- rewind** function, 346, 369, 371, 464, 537
- right-shift assignment operator (>>=), 98
- right-shift operator (>>), 91
- rint** function, 275, 515, 560, 565, 581, 582
- rint** type-generic macro, 417
- rintd128** function, 275, 521
- rintd32** function, 275, 521
- rintd64** function, 275, 521
- rintdN** function, 618
- rintdNx** function, 618
- rintf** function, 275, 515
- rintfN** function, 618
- rintfNx** function, 618
- rintl** function, 275, 515
- rootn** function, 271, 515, 564, 579
- rootn** type-generic macro, 417
- rootnd128** function, 271, 520
- rootnd32** function, 271, 520
- rootnd64** function, 271, 520
- rootndN** function, 617
- rootndNx** function, 617
- rootnf** function, 271, 515
- rootnfN** function, 617
- rootnfNx** function, 617
- rootnl** function, 271, 515
- round** function, 248, 275, 276, 516, 559, 581
- round to narrower type, 287
- round** type-generic macro, 417
- roundd128** function, 276, 521
- roundd32** function, 275, 521
- roundd64** function, 276, 521

- rounddN** function, 618
- rounddNx** function, 618
- roundeven** function, 248, 276, 277, 516, 559, 582
- roundeven** type-generic macro, 417
- roundevend128** function, 276, 521
- roundevend32** function, 276, 521
- roundevend64** function, 276, 521
- roundevendN** function, 618
- roundevendNx** function, 618
- roundevenf** function, 276, 516
- roundevenfN** function, 618
- roundevenfNx** function, 618
- roundevenl** function, 276, 516
- roundf** function, 275, 516
- roundfN** function, 618
- roundfNx** function, 618
- rounding, 276
- rounding control pragma, 225
- rounding direction, 248
- rounding mode
  - floating-point, 25
- roundl** function, 275, 516
- RSIZE\_MAX** macro, 537, 673, 674, 675, 679, 680, 683–685, 687, 689–698, 700, 701, 704, 706, 707, 710–714, 716, 717
- rsizet** type, 536, 538, 540, 545, 547, 550, 673, 674, 677–680, 683–685, 686, 687–692, 693, 694–698, 700, 701, 702, 703, 704, 706, 707, 710–717
- rsqrt** function, 271, 272, 515, 564, 580
- rsqrt** type-generic macro, 417
- rsqrtd128** function, 272, 520
- rsqrtd32** function, 272, 520
- rsqrtd64** function, 272, 520
- rsqrtdN** function, 617
- rsqrtdNx** function, 617
- rsqrtf** function, 271, 515
- rsqrtfN** function, 617
- rsqrtfNx** function, 617
- rsqrtl** function, 272, 515
- runtime-constraint, 7
- Runtime-constraint handling function, 686
- rvalue, 51
- same scope, 37
- samequantum** macro, 560
- samequantumd** identifier prefix, 290, 418
- samequantumd128** function, 290, 523
- samequantumd32** function, 290, 523
- samequantumd64** function, 290, 523
- samequantumdN** function, 621
- samequantumdNx** function, 621
- save calling environment function, 296
- scalar type, 42
- scalbln** function, 268, 515, 560, 577, 622
- scalbln** type-generic macro, 417
- scalblnd128** function, 268, 520
- scalblnd32** function, 268, 520
- scalblnd64** function, 268, 520
- scalblndN** function, 616
- scalblndNx** function, 616
- scalblnf** function, 268, 515
- scalblnfN** function, 616
- scalblnfNx** function, 616
- scalblnl** function, 268, 515
- scalbn** function, 268, 514, 560, 575, 576, 577, 591, 622
- scalbn** type-generic macro, 417
- scalbnd128** function, 268, 520
- scalbnd32** function, 268, 520
- scalbnd64** function, 268, 520
- scalbndN** function, 616
- scalbndNx** function, 616
- scalbnf** function, 268, 514
- scalbnfN** function, 616
- scalbnfNx** function, 616
- scalbnl** function, 268, 514
- scanf** function, 34, 226, 228, 340, 362, 365, 537, 561, 724, 725
- scanf\_s** function, 538, 679, 683
- scanlist, 359, 456
- scanset, 359, 456
- SCHAR\_MAX** macro, 24, 512, 555
- SCHAR\_MIN** macro, 24, 41, 512, 555
- SCHAR\_WIDTH** macro, 23, 512, 555
- SCN** identifier prefix, 236, 488
- SCNbMAX** macro, 237, 511
- SCNbPTR** macro, 237, 511
- SCNcFASTN** macro, 236
- SCNcLEASTN** macro, 236
- SCNcN** macro, 236
- SCNdMAX** macro, 236, 511
- SCNdPTR** macro, 236, 511
- SCNiMAX** macro, 236, 511
- SCNiPTR** macro, 236, 511
- SCNoMAX** macro, 237, 511
- SCNoPTR** macro, 237, 511
- SCNuMAX** macro, 237, 511
- SCNuPTR** macro, 237, 511
- SCNxMAX** macro, 237, 511
- SCNxPTR** macro, 237, 511
- scope, 36
- scope of identifier, 36, 173
- search function
  - string, 411
  - utility, 387, 688
  - wide string, 472
- search functions
  - string, 697
  - wide string, 714
- secondary block, 160
- SEEK\_CUR** macro, 340, 370, 537



- SEEK\_END** macro, 340, 343, 370, 537
- SEEK\_SET** macro, 340, 370, 371, 537, 643
- selection
  - \_Generic**, i, 55, 78, 120, 128, 143, 334, 418, 720
- selection statement, 162
- self-referential structure, 123
- semicolon punctuator (;), 104, 112, 162, 165, 166
- separate compilation, 11
- separate translation, 11
- sequence point, 14, 80, 94, 95, 99, 128, 129, 160, 206, 207, 348, 387, 447, 552, 688
- sequenced before, 14, 76, 80, 82, 96
- sequenced during a function call, 156
- sequencing of statement, 160
- sequential consistency, 19
- set\_constraint\_handler\_s** function, 540, 672, 686, 687
- setbuf** function, 339, 342, 343, 345, 347, 537
- setjmp** function, 206, 296, 297, 532, 633, 640, 719
- setlocale** function, viii, 204, 241, 242, 245, 435, 512, 639, 649
- setpayload** function, 523, 560, 587
- setpayloadd128** function, 524, 587
- setpayloadd32** function, 524, 587
- setpayloadd64** function, 524, 587
- setpayloaddN** function, 622
- setpayloaddNx** function, 622
- setpayloadf** function, 523, 587
- setpayloadfN** function, 622
- setpayloadfNx** function, 622
- setpayloadl** function, 523, 587
- setpayloadsig** function, 523, 560, 587
- setpayloadsigd128** function, 524, 587
- setpayloadsigd32** function, 524, 587
- setpayloadsigd64** function, 524, 587
- setpayloadsigdN** function, 622
- setpayloadsigdNx** function, 622
- setpayloadsigf** function, 523, 587
- setpayloadsigfN** function, 622
- setpayloadsigfNx** function, 622
- setpayloadsigl** function, 523, 587
- setvbuf** function, 339, 342, 343, 345, 347, 348, 537, 641
- shall, 9
- shift expression, 91
- shift sequence, 204
- shift state, 20
  - initial, 20
- short identifier
  - character, 22
- short int** type, 40, 110
- short int** type conversion, 47–49
- short** keyword, 55
- SHRT\_MAX** macro, 24, 512, 555
- SHRT\_MIN** macro, 24, 512, 555
- SHRT\_WIDTH** macro, 23, 512, 555
- side effect, 14, 43, 52, 76, 82, 96, 146, 162, 222, 225, 367, 368, 463, 566, 569, 570
- SIG** identifier prefix, 298, 489
- SIG\_** identifier prefix, 298, 489
- SIG\_ATOMIC\_MAX** macro, 338, 536
- SIG\_ATOMIC\_MIN** macro, 338, 536
- sig\_atomic\_t** type, 14, 298, 299, 338, 532, 632, 640
- SIG\_ATOMIC\_WIDTH** macro, 338, 536
- SIG\_DFL** macro, 298, 299, 532, 649
- SIG\_ERR** macro, 298, 299, 532, 640
- SIG\_IGN** macro, 298, 299, 532, 645
- SIGABRT** macro, 298, 299, 384, 532
- SIGFPE** macro, 250, 298, 299, 532, 640, 645, 654
- SIGILL** macro, 298, 299, 532, 640, 645
- SIGINT** macro, 298, 532
- sign bit, 44
- signal, 14, 21, 298
- signal** function, 15, 16, 141, 298, 299, 385, 386, 532, 640, 645, 649
- signal handler, 14, 21, 299, 300
- signal handling function, 298
- signal handling header, 298, 489
- signaling NaN, 25, 559
- signals, 298
- signbit** macro, 253, 512, 562, 563, 583
- signed char** type, 40
- signed character, 48
- signed integer type, 40, 48, 62
- signed integer types, 40
- signed** keyword, 55
- signed** type, 40, 110
- signed** type conversion, 47–49
- significand part, 64
- SIGSEGV** macro, 298, 299, 532, 640, 645
- SIGTERM** macro, 298, 532
- simple assignment, 97
- simple assignment operator (=), 97
- simple declaration, 105
- sin** function, 82, 256, 417, 419, 513, 565, 573
- sin** type-generic macro, 417
- sind128** function, 256, 518
- sind32** function, 256, 518
- sind64** function, 256, 518
- sindN** function, 613
- sindNx** function, 613
- sinf** function, 256, 513
- sinfN** function, 613
- sinfNx** function, 613
- single-byte character, 20
- single-byte/wide character conversion function, 477
- single-quote escape sequence (\ '), 69, 71

- singularity, 250
- sinh** function, 261, 417, 513, 565, 574
- sinh** type-generic macro, 417
- sinhd128** function, 261, 519
- sinhd32** function, 261, 519
- sinhd64** function, 261, 519
- sinhdN** function, 614
- sinhdNx** function, 614
- sinhf** function, 261, 513
- sinhfN** function, 614
- sinhfNx** function, 614
- sinhl** function, 261, 513
- sinl** function, 256, 513
- sinpi** function, 258, 259, 513, 565, 574
- sinpi** type-generic macro, 417
- sinpid128** function, 259, 518
- sinpid32** function, 259, 518
- sinpid64** function, 259, 518
- sinpidN** function, 613
- sinpidNx** function, 613
- sinpif** function, 258, 513
- sinpifN** function, 613
- sinpifNx** function, 613
- sinpil** function, 259, 513
- SIZE\_MAX** macro, 42, 404, 536, 673
- SIZE\_WIDTH** macro, 338, 536
- sizeof** keyword, 55
- sizeof** operator, 51, 86, 87
- snprintf** function, 363, 365, 374, 375, 436, 537, 679, 725
- snprintf\_s** function, 538, 679, 680
- snwprintf\_s** function, 550, 704, 705
- sorting utility function, 387, 688
- source character set, 11, 19
- source file, 11
  - name, 197, 199
- source file inclusion, 182
- source line, 11
- source text, 11
- space character ( ' '), 11, 20, 54, 218, 219, 483
- space* format flag, 349, 448
- spilling, 15
- sprintf** function, 363, 365, 537, 680
- sprintf\_s** function, 538, 679, 680
- sqrt** function, 158, 272, 417, 515, 560, 580, 585
- sqrt** type-generic macro, 417
- sqrtd128** function, 272, 520
- sqrtd32** function, 272, 419, 520
- sqrtd64** function, 272, 520
- sqrtdN** function, 617
- sqrtdNx** function, 617
- sqrtf** function, 272, 515
- sqrtfN** function, 617
- sqrtfNx** function, 617
- sqrtl** function, 272, 515
- square root rounded to narrower type, 289
- srand** function, 381, 539
- sscanf** function, 361, 363, 366, 537
- sscanf\_s** function, 538, 680, 681, 684
- standard attribute, 149
- standard embed parameter, 184
- standard error stream, 340, 342, 372
- standard floating type, 40
- standard header, 9, 204
- standard headers
  - <assert.h>, 186–188, 205, 208, 232, 507
  - <complex.h>, 25, 30, 108, 146, 201, 204, 205, 209–216, 416, 417, 488, 507, 508, 590, 591, 592, 607, 609, 644, 646, 654, 724, 725
  - <ctype.h>, 205, 217, 218–220, 488, 509
  - <errno.h>, 157, 205, 221, 488, 509, 672
  - <fcntl.h>, 9, 14, 16, 25, 30, 35, 98, 157, 205, 222, 224, 225, 227–234, 250, 488, 509, 559, 563, 566–569, 577, 581, 583, 725
  - <float.h>, 9, 10, 22, 24, 25, 29, 30, 108, 205, 235, 248, 353, 376, 452, 466, 488, 489, 510, 511, 555, 559, 565, 601, 602, 651, 724, 725
  - <inttypes.h>, 205, 236, 237, 238, 351, 450, 488, 511, 725
  - <iso646.h>, 9, 10, 205, 239, 511, 724, 726
  - <limits.h>, 9, 10, 22, 23, 24, 40, 41, 205, 240, 327, 328, 512, 555, 651
  - <locale.h>, 157, 205, 241, 242, 488, 512
  - <math.h>, ii, 9, 25, 30, 35, 77, 157, 205, 226–228, 247, 248, 250–253, 254–287, 288, 289–294, 295, 354, 416–418, 453, 488, 489, 512, 523, 524, 558, 559, 565, 571, 572, 577, 581, 583–587, 590, 591, 599, 609, 622–624, 634, 644, 646, 651, 654, 725
  - <setjmp.h>, 205, 296, 297, 532
  - <signal.h>, 205, 298, 300, 489, 532
  - <stdalign.h>, 9, 10, 205, 301, 532, 724
  - <stdarg.h>, 9, 10, 137, 205, 302, 303–305, 364–366, 458–460, 532, 681–684, 705–708
  - <stdatomic.h>, 201, 204, 205, 299, 307, 308, 311, 312, 314–316, 489, 532, 640, 724
  - <stdbit.h>, 9, 205, 318, 319–328, 489, 533, 721, 722
  - <stdbool.h>, 9, 10, 205, 330, 489, 536, 725
  - <stdckdint.h>, 205, 331, 489, 536
  - <stddef.h>, 9, 10, 52, 70, 72, 87, 88, 91, 149, 185, 186, 205, 238, 332, 333, 334, 361, 362, 536, 673
  - <stdint.h>, 9, 10, 22, 24, 179, 205, 236, 237, 327, 335, 337, 338, 350, 357, 358, 449, 455, 489, 536, 651, 673, 725
  - <stdio.h>, 16, 25, 30, 35, 58, 157, 182, 186,

- 205, 226, 228, **339**, 343–345, 347, 348,  
354, 356, 360–372, 433, 447, 453, 457–  
459, 461–464, **489**, **537**, 538, 646, **673**,  
674–685, 703, 705, 706, 724, 725
- <stdlib.h>, 9, 25, 30, 35, 205, 207, 226,  
228, **373**, 374, 375, 377, 379–391, **489**,  
**538**, 539, 540, 599, **624**, 625, **626**, **627**,  
646, 672, **686**, 687–692, 724
- <stdmchar.h>, **393**, 395, 396, 400, 403,  
404, **540**, 721
- <stdnoreturn.h>, 9, 10, 155, 205, **406**,  
**544**
- <string.h>, 9, 185, 188, 189, 205, **407**,  
408–415, **490**, **544**, 545, **693**, 694–699
- <tgmath.h>, 35, 205, **416**, 419, **545**, 558,  
**571**, **628**, 630, 725
- <threads.h>, 157, 201, 204, 205, **421**, 422–  
429, **490**, **546**, 724
- <time.h>, 205, 421, **431**, 432–437, 476, **490**,  
**547**, **699**, 700–702, 723
- <uchar.h>, 70, 72, 205, 394, **441**, 442–445,  
**547**, 722, 724
- <wchar.h>, 25, 30, 35, 157, 205, 226, 228,  
237, 340, 394, **446**, 447, 453, 457–465,  
467–480, **490**, **548**, 549, 550, **627**, **628**,  
646, **702**, 703–717, 724–726
- <wctype.h>, 205, **482**, 483–487, **490**, 550,  
724, 726
- standard input stream, 340, **342**
- standard integer type, **40**
- standard output stream, 340, **342**
- standard signed integer type, **40**
- standard unsigned integer type, **40**
- state-dependent encoding, **20**, 389, 690
- stateless function, **156**
- statement, **160**
  - break**, 55, **168**
  - compound, **162**
  - continue**, 55, 166, **167**, 168, 404
  - do**, 55, **166**
  - else**, 55, **163**
  - expression, **162**
  - for**, 55, **166**
  - goto**, 55, **166**
  - if**, 55, **163**
  - iteration, **165**
  - jump, **166**
  - labeled, **161**
  - null, **162**
  - return**, 55, **169**, **566**
  - selection, **162**
  - sequencing, **160**
  - switch**, 55, **164**
  - while**, 55, **166**
- static**, **106**
- static assertion, **149**
- static storage duration, **38**
- static** storage-class specifier, 37, 38, 55
- static**, in array declarator, 135, **137**
- static\_assert**, **149**
- static\_assert** storage-class specifier, 55
- STDC** pragma, 98, 158, **198**, **203**, **209**, 224–227,  
230, 232, 234, 251, 507, 509, 510, 512,  
567–569, 577, 581, 583, 590, 591, 638,  
648
- stdc\_** identifier prefix, 489
- stdc\_bit\_ceil** macro, 325, 535
- stdc\_bit\_ceil\_uc** function, 325, 535
- stdc\_bit\_ceil\_ui** function, 325, 535
- stdc\_bit\_ceil\_ul** function, 325, 535
- stdc\_bit\_ceil\_ull** function, 325, 535
- stdc\_bit\_ceil\_us** function, 325, 535
- stdc\_bit\_floor** macro, 325, 535
- stdc\_bit\_floor\_uc** function, 325, 535
- stdc\_bit\_floor\_ui** function, 325, 535
- stdc\_bit\_floor\_ul** function, 325, 535
- stdc\_bit\_floor\_ull** function, 325, 535
- stdc\_bit\_floor\_us** function, 325, 535
- stdc\_bit\_width** macro, 324, 535
- stdc\_bit\_width\_uc** function, 324, 535
- stdc\_bit\_width\_ui** function, 324, 535
- stdc\_bit\_width\_ul** function, 324, 535
- stdc\_bit\_width\_ull** function, 324, 535
- stdc\_bit\_width\_us** function, 324, 535
- STDC\_C16\_MAX** macro, **394**, **395**, 404, 540
- stdc\_c16nrtoc16n** function, 397, 542
- stdc\_c16nrtoc32n** function, 397, 542
- stdc\_c16nrtoc8n** function, 397, 542
- stdc\_c16nrtomcn** function, 397, 542
- stdc\_c16nrtomwcn** function, 397, 542
- stdc\_c16snrtoc16sn** function, 401, 544
- stdc\_c16snrtoc32sn** function, 401, 544
- stdc\_c16snrtoc8sn** function, 401, 544
- stdc\_c16snrtomcsn** function, 401, 543
- stdc\_c16snrtomwcsn** function, 401, 543
- STDC\_C32\_MAX** macro, **394**, **395**, 540
- stdc\_c32nrtoc16n** function, 397, 542
- stdc\_c32nrtoc32n** function, 397, 542
- stdc\_c32nrtoc8n** function, 397, 542
- stdc\_c32nrtomcn** function, 397, 542
- stdc\_c32nrtomwcn** function, 397, 542
- stdc\_c32snrtoc16sn** function, 401, 544
- stdc\_c32snrtoc32sn** function, 401, 544
- stdc\_c32snrtoc8sn** function, 401, 544
- stdc\_c32snrtomcsn** function, 401, 544
- stdc\_c32snrtomwcsn** function, 401, 544
- STDC\_C8\_MAX** macro, **394**, **395**, 540
- stdc\_c8nrtoc16n** function, 397, 542
- stdc\_c8nrtoc32n** function, 397, 542
- stdc\_c8nrtoc8n** function, 397, 541
- stdc\_c8nrtomcn** function, 396, 541
- stdc\_c8nrtomwcn** function, 396, 541

**stdc\_c8snrtoc16sn** function, 401, 404, 543  
**stdc\_c8snrtoc32sn** function, 401, 543  
**stdc\_c8snrtoc8sn** function, 401, 543  
**stdc\_c8snrtomcsn** function, 401, 543  
**stdc\_c8snrtomwcsn** function, 400, 543  
**stdc\_count\_ones** macro, 323, 534  
**stdc\_count\_ones\_uc** function, 323, 534  
**stdc\_count\_ones\_ui** function, 323, 534  
**stdc\_count\_ones\_ul** function, 323, 534  
**stdc\_count\_ones\_ull** function, 323, 534  
**stdc\_count\_ones\_us** function, 323, 534  
**stdc\_count\_zeros** macro, 323, 534  
**stdc\_count\_zeros\_uc** function, 323, 534  
**stdc\_count\_zeros\_ui** function, 323, 534  
**stdc\_count\_zeros\_ul** function, 323, 534  
**stdc\_count\_zeros\_ull** function, 323, 534  
**stdc\_count\_zeros\_us** function, 323, 534  
**stdc\_first\_leading\_one** macro, 321, 534  
**stdc\_first\_leading\_one\_uc** function, 321, 534  
**stdc\_first\_leading\_one\_ui** function, 321, 534  
**stdc\_first\_leading\_one\_ul** function, 321, 534  
**stdc\_first\_leading\_one\_ull** function, 321, 534  
**stdc\_first\_leading\_one\_us** function, 321, 534  
**stdc\_first\_leading\_zero** macro, 321, 534  
**stdc\_first\_leading\_zero\_uc** function, 321, 534  
**stdc\_first\_leading\_zero\_ui** function, 321, 534  
**stdc\_first\_leading\_zero\_ul** function, 321, 534  
**stdc\_first\_leading\_zero\_ull** function, 321, 534  
**stdc\_first\_leading\_zero\_us** function, 321, 534  
**stdc\_first\_trailing\_one** macro, 322, 534  
**stdc\_first\_trailing\_one\_uc** function, 322, 534  
**stdc\_first\_trailing\_one\_ui** function, 322, 534  
**stdc\_first\_trailing\_one\_ul** function, 322, 534  
**stdc\_first\_trailing\_one\_ull** function, 322, 534  
**stdc\_first\_trailing\_one\_us** function, 322, 534  
**stdc\_first\_trailing\_zero** macro, 322, 534  
**stdc\_first\_trailing\_zero\_uc** function, 322, 534  
**stdc\_first\_trailing\_zero\_ui** function, 322, 534  
**stdc\_first\_trailing\_zero\_ul** function, 322, 534  
**stdc\_first\_trailing\_zero\_ull** function, 322, 534  
**stdc\_has\_single\_bit** macro, 324, 535  
**stdc\_has\_single\_bit\_uc** function, 324, 534  
**stdc\_has\_single\_bit\_ui** function, 324, 534  
**stdc\_has\_single\_bit\_ul** function, 324, 534  
**stdc\_has\_single\_bit\_ull** function, 324, 535  
**stdc\_has\_single\_bit\_us** function, 324, 534  
**stdc\_leading\_ones** macro, 319, 533  
**stdc\_leading\_ones\_uc** function, 319, 533  
**stdc\_leading\_ones\_ui** function, 319, 533  
**stdc\_leading\_ones\_ul** function, 319, 533  
**stdc\_leading\_ones\_ull** function, 319, 533  
**stdc\_leading\_ones\_us** function, 319, 533  
**stdc\_leading\_zeros** macro, 319, 533  
**stdc\_leading\_zeros\_uc** function, 319, 533  
**stdc\_leading\_zeros\_ui** function, 319, 533  
**stdc\_leading\_zeros\_ul** function, 319, 533  
**stdc\_leading\_zeros\_ull** function, 319, 533  
**stdc\_leading\_zeros\_us** function, 319, 533  
**stdc\_load8\_aligned\_bes** function, 328, 535  
**stdc\_load8\_aligned\_beu** function, 328, 535  
**stdc\_load8\_aligned\_les** function, 328, 535  
**stdc\_load8\_aligned\_leu** function, 328, 535  
**stdc\_load8\_bes** function, 328, 535  
**stdc\_load8\_beu** function, 328, 535  
**stdc\_load8\_les** function, 328, 535  
**stdc\_load8\_leu** function, 328, 535  
**STDC\_MC\_MAX** macro, 394, 395, 541  
**stdc\_mcerr** type, 394, 395–397, 400–404, 540–544  
**stdc\_mcerr\_incomplete\_input** structure member, 393, 395, 396, 399, 403, 541  
**stdc\_mcerr\_insufficient\_output** structure member, 395, 399, 403, 541  
**stdc\_mcerr\_invalid** structure member, 393, 395, 396, 398, 399, 403, 541  
**stdc\_mcerr\_ok** structure member, 393, 395, 398, 399, 403, 404, 541  
**stdc\_mcnrtoc16n** function, 396, 541  
**stdc\_mcnrtoc32n** function, 396, 541  
**stdc\_mcnrtoc8n** function, 396, 403, 541  
**stdc\_mcnrtomcn** function, 396, 541  
**stdc\_mcnrtomwcn** function, 396, 541  
**stdc\_mcsnrtoc16sn** function, 400, 543  
**stdc\_mcsnrtoc32sn** function, 400, 543  
**stdc\_mcsnrtoc8sn** function, 400, 542  
**stdc\_mcsnrtomcsn** function, 400, 542  
**stdc\_mcsnrtomwcsn** function, 400, 542  
**stdc\_memreverse8u** function, 327, 535  
**STDC\_MWC\_MAX** macro, 394, 395, 541  
**stdc\_mwcnrtoc16n** function, 396, 541



- stdc\_mwcnrtoc32n** function, 396, 541
- stdc\_mwcnrtoc8n** function, 396, 541
- stdc\_mwcnrtomcn** function, 396, 541
- stdc\_mwcnrtomwcn** function, 396, 541
- stdc\_mwcsnrtoc16sn** function, 400, 543
- stdc\_mwcsnrtoc32sn** function, 400, 543
- stdc\_mwcsnrtoc8sn** function, 400, 543
- stdc\_mwcsnrtomcsn** function, 400, 543
- stdc\_mwcsnrtomwcn** function, 400, 543
- stdc\_rotate\_left** macro, 326, 535
- stdc\_rotate\_left\_uc** function, 326, 535
- stdc\_rotate\_left\_ui** function, 326, 535
- stdc\_rotate\_left\_ul** function, 326, 535
- stdc\_rotate\_left\_ull** function, 326, 535
- stdc\_rotate\_left\_us** function, 326, 535
- stdc\_rotate\_right** macro, 326, 535
- stdc\_rotate\_right\_uc** function, 326, 535
- stdc\_rotate\_right\_ui** function, 326, 535
- stdc\_rotate\_right\_ul** function, 326, 535
- stdc\_rotate\_right\_ull** function, 326, 535
- stdc\_rotate\_right\_us** function, 326, 327, 535
- stdc\_store8\_aligned\_bes** function, 329, 535
- stdc\_store8\_aligned\_beu** function, 329, 535
- stdc\_store8\_aligned\_les** function, 329, 535
- stdc\_store8\_aligned\_leu** function, 328, 535
- stdc\_store8\_bes** function, 329, 535
- stdc\_store8\_beu** function, 328, 535
- stdc\_store8\_les** function, 329, 535
- stdc\_store8\_leu** function, 328, 535
- stdc\_trailing\_ones** macro, 320, 534
- stdc\_trailing\_ones\_uc** function, 320, 534
- stdc\_trailing\_ones\_ui** function, 320, 534
- stdc\_trailing\_ones\_ul** function, 320, 534
- stdc\_trailing\_ones\_ull** function, 320, 534
- stdc\_trailing\_ones\_us** function, 320, 534
- stdc\_trailing\_zeros** macro, 320, 534
- stdc\_trailing\_zeros\_uc** function, 320, 533
- stdc\_trailing\_zeros\_ui** function, 320, 533
- stdc\_trailing\_zeros\_ul** function, 320, 533
- stdc\_trailing\_zeros\_ull** function, 320, 533
- stdc\_trailing\_zeros\_us** function, 320, 533
- stderr** stream, 197, 340, 341, 342, 347, 364, 459, 537, 653
- stdin** stream, 340, 341, 342, 347, 360–362, 368, 457, 461, 463, 537, 678, 679, 685, 710
- stdout** stream, 340, 341, 342, 347, 354, 362, 368, 453, 461, 464, 537
- storage duration, 38
- storage order of array, 79
- storage unit (bit-field), 43, 113
- storage-class specifier, 106, 203
  - \_Thread\_local**, 55
  - auto**, 55, 106, 107, 136, 137, 142, 143, 171, 203, 723
  - constexpr**, 37–39, 55, 84, 100, 101, 102, 105–110, 144, 203, 299, 404, 567, 568, 722
  - extern**, 37, 38, 55, 87, 102, 103, 106, 107, 121, 127–132, 136, 137, 141, 158, 165, 172–174, 188, 189, 207, 652, 653
  - register**, 55, 106
  - static**, 55, 106
  - thread\_local**, 38, 55, 106
  - typedef**, 55, 106, 140
- store and load, 15
- str** identifier prefix, 488–490
- strcat** function, 409, 544
- strcat\_s** function, 545, 695, 696
- strchr** macro, 411, 412, 490, 544, 723
- strcmp** function, 409, 410, 411, 544
- strcoll** function, 9, 242, 410, 411, 544
- strcpy** function, 188–190, 408, 544
- strcpy\_s** function, 545, 694
- strcspn** function, 412, 544
- strdup** function, 9, 408, 544, 721
- stream, 341, 385
  - binary, 341
  - fully buffered, 342
  - line buffered, 342
  - orientation, 341
  - standard error, 340, 342
  - standard input, 340, 342
  - standard output, 340, 342
  - text, 341
  - unbuffered, 342
- strerror** function, 9, 372, 414, 415, 545, 643, 644, 652
- strerror\_s** function, 415, 545, 698, 699
- strerrorlen\_s** function, 545, 699
- strfrom** family, 226, 228
- strfromd** family, 375
- strfromd** function, 374, 375, 464, 465, 538, 561
- strfromd128** function, 375, 539
- strfromd32** function, 375, 539
- strfromd64** function, 375, 539
- strfromdN** function, 625
- strfromdNx** function, 625
- strfromencbindN** function, 626
- strfromencdecN** function, 626
- strfromencf128** function, 625, 626
- strfromencfN** function, 626
- strfromf** function, 374, 375, 464, 538
- strfromfN** function, 625
- strfromfNx** function, 625
- strfroml** function, 374, 375, 538
- strftime** function, ii, 242, 435, 437, 440, 476,

- 547, 634, 641–643, 651, 700, 701, 721, 726
- stricter, 46
- strictly conforming program, 9
- string, 204
  - comparison function, 409
  - concatenation function, 409, 695
  - conversion function, 242
  - copying function, 407, 693
  - library function convention, 407
  - literal, 12, 19, 51, 71, 77, 145
  - miscellaneous function, 414, 698
  - numeric conversion function, 237, 373
  - search function, 411, 697
- string duplicate function, 408, 409
- string handling header, 407, 490, 693
- string literal
  - wide, 71
- stringizing, 193, 202
- stringizing argument, 193
- strlen** function, 409, 415, 545
- strncat** function, 409, 544
- strncat\_s** function, 545, 696, 697
- strncmp** function, 196, 409, 410, 411, 544
- strncpy** function, 408, 544
- strncpy\_s** function, 545, 694, 695
- strndup** function, 9, 409, 544, 721
- strnlen** function, ii, 415, 545, 721
- strnlen\_s** function, 545, 694–696, 699
- stronger, 46
- strpbrk** macro, 411, 412, 490, 545, 723
- strrchr** macro, 411, 412, 413, 490, 545, 723
- strspn** function, 413, 545
- strstr** macro, 411, 413, 490, 545, 723
- strto** family, 34, 226, 228, 377
- strtod** family, 377, 378
- strtod** function, 66, 280, 357, 358, 362, 374, 375, 464, 538, 561, 566, 567, 633, 650
- strtod128** function, 377, 539, 626, 650
- strtod32** function, 377, 539, 650
- strtod64** function, 377, 378, 539, 650
- strtodN** function, 625
- strtodNx** function, 625
- strtoencbindN** function, 627
- strtoencdecN** function, 627
- strtoencfN** function, 627
- strtof** function, 280, 362, 374, 375, 538, 633, 650
- strtofN** function, 625
- strtofNx** function, 625
- strtoimax** function, 237, 238, 511
- strtok** function, 9, 413, 414, 545, 643
- strtok\_s** function, 414, 545, 697, 698
- strtol** function, 238, 357, 358, 362, 374, 379, 380, 538
- strtold** function, 280, 362, 374, 375, 538, 633, 650
- strtoll** function, 238, 362, 374, 379, 380, 538
- strtoul** function, 238, 358, 362, 374, 379, 380, 538
- strtoull** function, 238, 362, 374, 379, 380, 538
- strtoumax** function, 237, 238, 511
- struct** keyword, 55
- structure
  - arrow operator (->), 81
  - content, 123
  - dot operator (.), 81
  - initialization, 145
  - member alignment, 114
  - member name space, 38
  - member operator (.), 51, 81
  - pointer operator (->), 81
  - specifier, 112
  - tag, 38, 123
  - type, 41, 112
- structure content, 123
- structure or union constant, 101
- strxfrm** function, 9, 242, 411, 544, 643
- subnormal floating-point number, 24
- subscripting, 79
- subtract and round to narrower type, 287
- subtraction assignment operator (-=), 98
- subtraction operator (-), 90, 591
- successful termination, 385
- suffix
  - floating literal, 65
  - integer literal, 62
- suffix** embed parameter, 184, 188, 189, 190
- switch body, 164
- switch **case** label, 161, 164
- switch **default** label, 161, 164
- switch** keyword, 55
- switch** statement, 164
- swprintf** function, 458, 459, 548, 704, 705
- swprintf\_s** function, 550, 704, 705
- swscanf** function, 458, 460, 548
- swscanf\_s** function, 550, 705, 708
- symbol, 3
- synchronization operation, 17
- synchronize, 156
- synchronize with, 17
- syntactic category, 36
- syntax notation, 36
- syntax rule precedence, 11
- syntax summary
  - language, 491
- system** function, 386, 387, 539, 643, 645, 650
- t format modifier, 350, 357, 449, 455
- tab character, 20, 54
- tag, 121
- tag compatibility, 44

- tag name space, 38
- tags, 38
- tan** function, 256, 257, 417, 513, 565, 573
- tan** type-generic macro, 417
- tand128** function, 256, 518
- tand32** function, 256, 518
- tand64** function, 256, 518
- tandN** function, 613
- tandNx** function, 613
- tanf** function, 256, 513
- tanfN** function, 613
- tanfNx** function, 613
- tanh** function, 261, 417, 513, 565, 575
- tanh** type-generic macro, 417
- tanhd128** function, 261, 519
- tanhd32** function, 261, 519
- tanhd64** function, 261, 519
- tanhdN** function, 614
- tanhdNx** function, 614
- tanhf** function, 261, 513
- tanhfN** function, 614
- tanhfNx** function, 614
- tanhl** function, 261, 514
- tanl** function, 256, 513
- tanpi** function, 259, 513, 565, 574
- tanpi** type-generic macro, 417
- tanpid128** function, 259, 518
- tanpid32** function, 259, 518
- tanpid64** function, 259, 518
- tanpidN** function, 613
- tanpidNx** function, 613
- tanpif** function, 259, 513
- tanpifN** function, 613
- tanpifNx** function, 613
- tanpil** function, 259, 513
- temporary lifetime, 39
- tentative definition, 173
- term, 3
- text stream, 341, 369–371
- Text transcoding utilities header, 393
- tgamma** function, 273, 274, 515, 580
- tgamma** type-generic macro, 417
- tgammaad128** function, 273, 520
- tgammaad32** function, 273, 520
- tgammaad64** function, 273, 520
- tgammaadN** function, 617
- tgammaadNx** function, 617
- tgammaaf** function, 273, 515
- tgammaafN** function, 617
- tgammaafNx** function, 617
- tgammaal** function, 273, 515
- thousands\_sep** structure member, 241, 243
- thrd\_** identifier prefix, 490
- thrd\_busy** constant, 422, 426, 546
- thrd\_create** function, 421, 426, 546
- thrd\_current** function, 426, 546
- thrd\_detach** function, 427, 546, 644
- thrd\_equal** function, 427, 547
- thrd\_error** constant, 422, 423–430, 546
- thrd\_exit** function, 426, 427, 547, 634
- thrd\_join** function, 427, 428, 547, 644
- thrd\_nomem** constant, 422, 423, 426, 546
- thrd\_sleep** function, 428, 547
- thrd\_start\_t** type, 421, 426, 546
- thrd\_success** constant, 422, 423–430, 546
- thrd\_t** type, 421, 426, 427, 546, 547
- thrd\_timeout** constant, 422, 424, 425, 546
- thrd\_yield** function, 428, 547
- thread, 16
- thread of execution, 16, 207, 222, 386, 688
- thread storage duration, 38, 222
- thread\_local**, 38, 106
- thread\_local** storage-class specifier, 38, 55
- threads header, 421, 490
- time
  - broken down, 432, 433–437, 700, 702
  - calendar, 431, 432–434, 436, 437, 701, 702
  - component, 431, 700
  - conversion function, 435, 700
    - wide character, 476
  - local, 431
  - manipulation function, 432
  - normalized broken down, 700
- time base, 431, 434
- time** function, 433, 434, 547, 634
- TIME\_** identifier prefix, 431, 490
- TIME\_ACTIVE** macro, 431, 435, 490, 547, 651
- TIME\_MONOTONIC** macro, 431, 434, 490, 547, 651
- time\_t** type, 431, 432, 433, 434, 436, 437, 547, 651, 701, 702
- TIME\_THREAD\_ACTIVE** macro, 431, 435, 490, 547, 651
- TIME\_UTC** macro, 424, 425, 428, 431, 434, 547, 651
- timegm** function, 433, 434, 547, 723
- timespec** structure type, 423, 425, 428, 431, 432, 434, 435, 546, 547
- timespec\_get** function, 431, 434, 435, 547
- timespec\_getres** function, 431, 435, 547, 723
- tm** structure type, 431, 432, 433–437, 446, 476, 547–549, 700, 701, 702
- tm\_hour** structure member, 432, 433, 436, 438, 700
- tm\_isdst** structure member, 432, 433
- tm\_mday** structure member, 432, 433, 434, 436, 438, 700
- tm\_min** structure member, 432, 433, 436, 438, 701
- tm\_mon** structure member, 432, 433, 434, 436–438, 700
- tm\_sec** structure member, 432, 433, 436, 438,

- 701
- tm\_wday** structure member, 432, 433, 434, 436–438, 700
- tm\_yday** structure member, 432, 433, 434, 438
- tm\_year** structure member, 432, 433, 434, 436–438, 701
- TMP\_MAX** macro, 340, 344, 345, 537
- TMP\_MAX\_S** macro, 538, 673, 674, 675
- tmpfile** function, 344, 385, 537
- tmpfile\_s** function, 538, 674, 675
- tmpnam** function, 340, 344, 345, 537, 675
- tmpnam\_s** function, 538, 673, 674, 675
- to** identifier prefix, 488, 490
- token, 12, 54
- token concatenation, 193
- token pasting, 193
- tolower** function, 219, 509
- totalorder** function, 523, 562, 585, 586
- totalorderd128** function, 523, 585
- totalorderd32** function, 523, 585
- totalorderd64** function, 523, 585
- totalorderdN** function, 621
- totalorderdNx** function, 621
- totalorderf** function, 523, 585
- totalorderfN** function, 621
- totalorderfNx** function, 621
- totalorderl** function, 523, 585
- totalordermag** function, 523, 562, 586
- totalordermagd128** function, 524, 586
- totalordermagd32** function, 523, 586
- totalordermagd64** function, 524, 586
- totalordermagdN** function, 621
- totalordermagdNx** function, 621
- totalordermagf** function, 523, 586
- totalordermagfN** function, 621
- totalordermagfNx** function, 621
- totalordermagl** function, 523, 586
- toupper** function, 219, 220, 509
- towctrans** function, 486, 487, 551, 644, 652
- tolower** function, 486, 487, 551
- towupper** function, 486, 487, 551
- translation environment, 11
- translation limit, 22
- translation phase, 11
- translation unit, 11, 171
- trigonometric function
  - complex, 210, 592
  - real, 254, 572
- true** keyword, 55
- true** predefined constant, 70
- trunc** function, 248, 277, 516, 559, 582
- trunc** type-generic macro, 417
- truncation, 49, 277, 342
- truncation toward zero, 89
- truncd128** function, 277, 521
- truncd32** function, 277, 521
- truncd64** function, 277, 521
- truncdN** function, 618
- truncdNx** function, 618
- truncf** function, 277, 516
- truncfN** function, 618
- truncfNx** function, 618
- truncl** function, 277, 516
- tss\_** identifier prefix, 490
- tss\_create** function, 428, 429, 547, 644
- tss\_delete** function, 429, 547, 634, 644
- TSS\_DTOR\_ITERATIONS** macro, 421, 427, 546
- tss\_dtor\_t** type, 421, 428, 546, 547
- tss\_get** function, 429, 547, 644
- tss\_set** function, 429, 430, 547, 644
- tss\_t** type, 421, 428, 429, 546, 547
- tv\_nsec** structure member, 432, 434
- tv\_sec** structure member, 432, 434
- two's complement, 44
- type, 39
  - aggregate, 42
  - arithmetic, 41
  - atomic, 14, 42, 43, 51, 98, 124, 201, 313
  - basic, 41
  - compatible, 44, 111, 128, 134
  - complex, 41
  - composite, 45
  - const qualified, 128
  - conversion, 47
  - derived, 41
  - derived declarator, 42
  - extended floating, 600
  - integer, 41
  - real, 41
  - restrict qualified, 128
  - scalar, 42
  - typeof, 124
  - volatile qualified, 128
- type category, 42
- type conversion, 47
- type definition, 140
- type domain, 41
- type inference, 142, 653
- type name, 139
- type punning, 81
- type qualifier, 127
- type specifier, 110
  - \_BitInt**, 23, 40, 48, 50, 51, 55, 62, 63, 110, 111, 326, 327, 722
  - \_Bool**, 55
  - \_Complex**, 26, 41, 50, 55, 108–111, 209, 590, 601, 607
  - \_Decimal128**, 55
  - \_Decimal32**, 55
  - \_Decimal64**, 55
  - bool**, 55
  - char**, 55



- double**, 55
- enum**, 55
- float**, 55
- int**, 55
- long**, 55
- short**, 55
- signed**, 55
- struct**, 55
- union**, 55
- unsigned**, 55
- void**, 55
- type-generic macro, 416
- type-generic math header, 416
- typedef**, 106, 140
- typedef** declaration, 140
- typedef** storage-class specifier, 55, 140
- typeof** keyword, 55
- typeof operator, 51, 124
- typeof specifier, 124
- typeof\_unqual** keyword, 55
- types
  - atomic, 81, 82
  - character, 145
  - complex, 589
- U** encoding prefix, 67–69, 71, 72
- u** encoding prefix, 67–69, 71, 72
- u8** encoding prefix, 67–69, 71, 72
- uabs** function, 388, 539
- UCHAR\_MAX** macro, 24, 512, 555
- UCHAR\_WIDTH** macro, 23, 512, 555
- ufromfp** function, 248, 277, 278, 516, 561, 565, 582
- ufromfp** function, 277
- ufromfp** type-generic macro, 417
- ufromfpd128** function, 277, 521
- ufromfpd32** function, 277, 521
- ufromfpd64** function, 277, 521
- ufromfpdN** function, 618
- ufromfpdNx** function, 618
- ufromfpf** function, 277, 516
- ufromfpfN** function, 618
- ufromfpfNx** function, 618
- ufromfpl** function, 277, 516
- ufromfpx** function, 248, 278, 516, 561, 565, 582
- ufromfpx** function, 278
- ufromfpx** type-generic macro, 417
- ufromfpxd128** function, 278, 521
- ufromfpxd32** function, 278, 521
- ufromfpxd64** function, 278, 521
- ufromfpxdN** function, 618
- ufromfpxdNx** function, 618
- ufromfpxf** function, 278, 516
- ufromfpxfN** function, 618
- ufromfpxfNx** function, 618
- ufromfpxl** function, 278, 516
- uimaxabs** function, 237, 511, 721
- UINT** identifier prefix, 337, 338, 489, 536
- uint** identifier prefix, 327, 335, 336, 489, 535, 536, 724
- UINTN\_C** macro, 338
- UINTN\_MAX** macro, 337
- uintN\_t** type, 335
- UINTN\_WIDTH** macro, 337
- UINT64\_C** macro, 338
- uint64\_t** type, 278
- UINT\_FAST** identifier prefix, 337, 536
- uint\_fast** identifier prefix, 336, 536, 724
- UINT\_LEAST** identifier prefix, 337, 536
- uint\_least** identifier prefix, 328, 329, 335, 336, 338, 535, 536
- UINT\_FASTN\_MAX** macro, 337
- UINT\_FASTN\_WIDTH** macro, 337
- uint\_fast16\_t** type, 313, 336
- uint\_fast32\_t** type, 313, 336
- uint\_fast64\_t** type, 313, 336
- uint\_fast8\_t** type, 313, 336
- uint\_fastN\_t** type, 336
- UINT\_LEASTN\_MAX** macro, 337
- uint\_leastN\_t** type, 335
- UINT\_LEASTN\_WIDTH** macro, 337
- uint\_least16\_t** type, 313, 335, 336, 441
- uint\_least32\_t** type, 313, 336, 441
- uint\_least64\_t** type, 313, 336, 338
- uint\_least8\_t** type, 313, 336
- UINT\_MAX** macro, 24, 110, 119, 179, 512, 555
- UINT\_WIDTH** macro, 23, 278, 512, 555
- UINTMAX\_C** macro, 338, 536
- UINTMAX\_MAX** macro, 237, 238, 337, 536, 565
- uintmax\_t** type, 24, 179, 237, 238, 313, 337, 338, 350, 357, 449, 455, 511, 536, 723, 725
- UINTMAX\_WIDTH** macro, 337, 536
- UINTPTR\_MAX** macro, 337, 536
- uintptr\_t** type, 313, 336, 536
- UINTPTR\_WIDTH** macro, 337, 536
- uabs** function, 388, 539
- u labs** function, 388, 539
- ULLONG\_MAX** macro, 24, 109, 119, 380, 469, 512, 555
- ULLONG\_WIDTH** macro, 23, 512, 555
- ULONG\_MAX** macro, 24, 380, 469, 512, 555
- ULONG\_WIDTH** macro, 23, 512, 555
- unary arithmetic operator, 87
- unary expression, 85
- unary minus operator (-), 87, 563
- unary operator, 85
- unary plus operator (+), 87
- unbuffered, 342
- unbuffered stream, 342
- undef**, 57, 175, 178, 195, 199, 206, 207, 638
- undef** preprocessing directive, 195, 206
- undefined behavior, 4, 9, 635

- underlying type, **116**
- underscore
  - leading
    - in identifier, **206**
- underspecified, **105**
- ungetc** function, **340**, **368**, 369–371, 489, 537, 633, 642, 653, 726
- ungetwc** function, **340**, **464**, 548, 633, 653
- Unicode, **441**
- Unicode encodings, **393**
- Unicode required set, **200**
- Unicode Standard
  - Annex, UAX #31, 727
  - Annex, UAX #44, 2
  - Derived Core Properties, 2
- Unicode utilities header, **441**
- union
  - arrow operator (->), **81**
  - content, **123**
  - dot operator (.), **81**
  - initialization, **145**
  - member alignment, **114**
  - member name space, **38**
  - member operator (.), 51, **81**
  - pointer operator (->), **81**
  - specifier, **112**
  - tag, 38, **123**
  - type, **41**, **112**
- union content, **123**
- union** keyword, 55
- universal character name, **58**, 553
- unnormalized floating-point number, **24**
- unprefixed octal literal, **60**, **62**
- unqualified type, **42**
- unqualified version of type, **42**
- unreachable, 333
- unreachable** macro, ix, **332**, **333**, 536, 641, 724
- unrecorded encoding error, **47**
- unsequenced, **14**, 76, 96, **156**
- unsequenced** attribute, 150, 155, **157**, 158, 319–325, 533–535, 638, 720, 722
- unsigned bit-precise integer suffix, **uwb** or **UWB**, **62**
- unsigned integer suffix, **u** or **U**, **62**
- unsigned integer type, **40**, 48, **62**
- unsigned** keyword, 55
- unsigned** type, **40**, 110, 349, 357, 449, 454
- unsigned** type conversion, **48**, **49**
- unspecified behavior, **4**, 9, 632
- unspecified value, **7**
- unsuccessful termination, **384**, **385**
- uppercase letter, **19**
- use of library function, **206**
- USHRT\_MAX** macro, **24**, 118, 119, 512, 555
- USHRT\_WIDTH** macro, **23**, 512, 555
- usual arithmetic conversion, **49**, **50**, 89, 90, 92–95
- UTF-8 string literal, *see* string literal
- UTF-16 character literal, **68**
- UTF-16 string literal, **71**
- UTF-32 character literal, **68**
- UTF-32 string literal, **71**
- UTF-8 character literal, **68**
- UTF-8 string literal, **71**
- utility
  - bit and byte, **318**
  - general, **373**, 489, **686**
  - wide string, **464**, **710**
  - text transcoding, **393**
  - Unicode, **441**
- va\_arg** function, 236, **302**, 303–305, 353, 364–366, 452, 458–460, 532, 640, 641, 682–684, 706, 708, 709
- va\_copy** function, 206, 302, **303**, 305, 532, 633, 640, 641, 726
- va\_end** function, 206, 302, **303**, 304, 305, 364–366, 458–460, 532, 633, 640–642, 682–684, 706, 708, 709
- va\_list** type, **302**, 303–305, 364–366, 458–460, 532, 537, 538, 548, 550, 640–642, 649, 681–684, 705–709
- va\_start** function, 302, **303**, 304, 305, 364–366, 458–460, 532, 640, 641, 682–684, 706, 708, 709, 723
- value, **7**
- value bit, **43**
- value of a string, **204**
- value of a wide string, **204**
- variable length array, 134, **135**, 201
- variably modified, **134**
- variably modified type, **134**, 135, 201
- variadic functions, **137**
- varying argument, **191**
- varying arguments, **137**
- varying arguments header, **302**
- vertical tab, **21**
- vertical-tab character, **20**, 54
- vertical-tab escape sequence (\v), **21**, 69, 219
- vfprintf** function, **340**, **363**, 364, 537, 642, 681
- vfprintf\_s** function, 538, **681**, 682–684
- vfscanf** function, **340**, **364**, 537, 642
- vfscanf\_s** function, 538, **682**, 683, 684
- vfwprintf** function, **340**, **458**, 459, 548, 642, 706
- vfwprintf\_s** function, 550, **705**, 706
- fwscanf** function, **340**, 458, **459**, 464, 548, 642
- fwscanf\_s** function, 550, **706**, 708, 709
- visibility of identifier, **36**
- visible, **36**
- visible side effect, **18**
- void expression, **52**

- void** function parameter, 137
- void** keyword, 55
- void** type, 52, 110
- void** type conversion, 52
- volatile**, 42
- volatile access, 14
- volatile** keyword, 55
- volatile** type qualifier, 127
- volatile-qualified type, 42, 128
- vprintf** function, 340, 364, 365, 537, 642, 682
- vprintf\_s** function, 538, 682, 683, 684
- vscanf** function, 340, 364, 365, 537, 642, 725
- vscanf\_s** function, 538, 682, 683–685
- vsprintf** function, 364, 365, 537, 642, 683
- vsprintf\_s** function, 538, 682, 683, 684
- vsnwprintf\_s** function, 550, 706, 707
- vsprintf** function, 364, 365, 366, 537, 642, 684
- vsprintf\_s** function, 538, 682, 683, 684
- vsscanf** function, 364, 366, 537, 642
- vsscanf\_s** function, 538, 682, 683, 684, 685
- vswprintf** function, 458, 459, 548, 642, 707
- vswprintf\_s** function, 550, 707
- vswscanf** function, 458, 460, 548, 642
- vswscanf\_s** function, 550, 706, 707, 708, 709
- wprintf** function, 340, 458, 460, 548, 642, 708
- wprintf\_s** function, 550, 708
- wscanf** function, 340, 458, 460, 464, 548, 642
- wscanf\_s** function, 550, 706, 708, 709
  
- warning, 12, 631
- WCHAR\_MAX** macro, 394, 446, 536, 540, 548
- WCHAR\_MIN** macro, 338, 394, 446, 536, 540, 548
- wchar\_t** type, 5, 46, 47, 68–71, 144, 200, 201, 238, 313, 332, 338, 350, 352, 353, 357, 359, 362, 373, 389–391, 393–397, 400, 401, 441, 446, 447, 449, 452–454, 456–465, 467–476, 478–480, 482, 511, 536, 538–544, 548–550, 625, 627, 628, 634, 644, 646, 690–692, 702–717
- WCHAR\_UTF16** macro, 394, 446, 540, 548
- WCHAR\_UTF32** macro, 394, 446, 540, 548
- WCHAR\_UTF8** macro, 394, 446, 540, 548
- WCHAR\_WIDTH** macro, 338, 394, 446, 536, 540, 548
- wrtomb** function, 343, 352, 353, 356, 362, 447, 456, 458, 479, 481, 549, 634, 692, 715, 718
- wrtomb\_s** function, 550, 715, 716
- wcs** identifier prefix, 488–490
- wscat** function, 470, 471, 549
- wscat\_s** function, 550, 712, 713
- wscchr** macro, 472, 473, 490, 549, 723
- wscmp** function, 471, 472, 549
- wscoll** function, 471, 472, 549
- wscpy** function, 469, 549
- wscpy\_s** function, 550, 710
- wscspn** function, 473, 549
- wcsftime** function, 242, 476, 549, 634, 641–643, 651
- wcslen** function, 471, 475, 549, 715
- wcsncat** function, 471, 549
- wcsncat\_s** function, 550, 713, 714
- wcsncmp** function, 471, 472, 549
- wcsncpy** function, 469, 470, 549
- wcsncpy\_s** function, 550, 710, 711
- wcsnlen** function, ii, 475, 476, 549, 721
- wcsnlen\_s** function, 550, 710–713, 715
- wcspbrk** macro, 472, 473, 490, 549, 723
- wcsrchr** macro, 472, 473, 474, 490, 549, 723
- wcsrtombs** function, 480, 481, 549, 716
- wcsrtombs\_s** function, 550, 716, 717, 718
- wcsspn** function, 474, 549
- wcssr** macro, 472, 474, 490, 549, 723
- wcsto** family, 34, 226, 228, 467
- wcstod** family, 467, 468
- wcstod** function, 454, 456, 458, 465, 548, 561, 633, 650
- wcstod128** function, 465, 467, 549, 650
- wcstod32** function, 465, 467, 549, 650
- wcstod64** function, 465, 467, 549, 650
- wcstodN** function, 625
- wcstodNx** function, 625
- wcstoencbindN** function, 628
- wcstoencdecN** function, 628
- wcstoencfN** function, 627
- wcstof** function, 458, 465, 548, 633, 650
- wcstofN** function, 625
- wcstofNx** function, 625
- wcstoimax** function, 238, 511
- wcstok** function, 474, 475, 549, 643, 644
- wcstok\_s** function, 550, 714, 715
- wcstol** function, 238, 454, 455, 458, 468, 469, 548
- wcstold** function, 458, 465, 548, 633, 650
- wcstoll** function, 238, 458, 468, 469, 548
- wcstombs** function, 391, 479, 539
- wcstombs\_s** function, 540, 692, 693
- wcstoul** function, 238, 455, 456, 458, 468, 469, 548
- wcstoull** function, 238, 458, 468, 469, 549
- wcstoumax** function, 238, 511
- wcsxfrm** function, 472, 549, 643
- wctob** function, 477, 482, 549
- wctomb** function, 389, 390, 391, 478, 539
- wctomb\_s** function, 540, 690, 691
- wctrans** function, 486, 487, 551, 644
- wctrans\_t** type, 482, 486, 487, 550, 551
- wctype** function, 485, 486, 551, 644
- wctype\_t** type, 482, 485, 486, 550, 551
- weaker, 46
- WEOF** macro, 446, 447, 461–464, 477, 482, 548, 550, 644
- wfN format modifier, 350, 357, 449, 455

- while** keyword, 55
- while** statement, 166
- white space, 11, 54, 176, 177, 219, 484
- white-space character, 54, 204
- White-space wide character, 204
- wide character, 5
  - case mapping function, 486
  - extensible, 486
  - classification function, 482
  - extensible classification function, 485
  - formatted input/output function, 447, 702
  - input function, 340
  - input/output function, 340, 461
  - literal, 68
  - output function, 340
  - single-byte conversion function, 477
- wide character classification and mapping utilities header, 482, 490
- wide character input functions, 340
- wide character input/output function, 340
- wide character literal, 68
- wide character output function, 340
- wide execution encoding, 47
- wide literal encoding, 46
- wide string, 204
- wide string comparison function, 471
- wide string concatenation function, 470, 712
- wide string copying function, 469, 710
- wide string literal, *see* string literal, 71
- wide string miscellaneous function, 475, 715
- wide string numeric conversion function, 238, 464
- wide string search function, 472, 714
- wide-oriented stream, 341
- width, 43, 44
- WINT\_MAX** macro, 536
- WINT\_MIN** macro, 338, 536
- wint\_t** type, 338, 350, 352, 446, 447, 449, 452, 461–464, 477, 482, 483–486, 548–551, 644
- WINT\_WIDTH** macro, 338, 536
- wmemchr** macro, 472, 475, 490, 549, 723
- wmemcmp** function, 472, 549
- wmemcpy** function, 470, 549
- wmemcpy\_s** function, 550, 711, 712
- wmemmove** function, 470, 549
- wmemmove\_s** function, 550, 712
- wmemset** function, 476, 549
- wN format modifier, 350, 357, 449, 455
- wprintf** function, 226, 228, 237, 340, 460, 461, 548, 561, 709
- wprintf\_s** function, 550, 709
- wraparound, 8
- write-read coherence, 19
- write-write coherence, 18
- wscanf** function, 226, 228, 340, 460, 461, 464, 548, 561
- wscanf\_s** function, 550, 709, 710
- XID\_Continue, 56
- XID\_Start, 56
- xor** macro, 239, 512
- xor\_eq** macro, 239, 512
- z format modifier, 350, 357, 449, 455
- zero, 589